

# Snowflake SnowPro Certification Exam Cheat Sheet by Jeno Yamma

## // Snowflake - General //

### Introduction

- Analytic data warehouse
- SaaS offering
  - No hardware (or installation or updates/patch)
  - No ongoing maintenance or tuning
  - Can't run privately (on-prem or hosted)
- Runs completely on the cloud
  - Either AWS, Azure or GCP (NO ON-PREM or HYBRID!)
  - Has its own VPC!
- Decoupled compute and storage (scaled compute does not need scaled storage)

### Pricing

- Unit costs for Credits and data storage determined by region (not cloud platform)
  - Cost of Warehouse used
    - Minimum 60s on resuming, per 1s after
  - Cost of storage (Temp, Transient, Perm Tables & time-travel data & fail-safe)
    - Calcs based on daily average
- Pricing model = on demand or discounted upfront
- Just remember:
  - >= Enterprise = 90 days time travel (default: 1 day for all) + materialized view + multi-cluster warehouse
  - >= Business Critical = lots more security (HIPAA, SOC 1&2, PCI DSS)
  - VPS edition = own cloud service layer (not shared with accounts)

### Supported Regions

- Multi-region account isn't supported, each SF in single region
- 17 Regions in total
  - AWS - 9 regions (Asia Pacific-3, EU-2, North America-4)
  - GCP - 1 (Asia Pacific-0, EU-0, North America-1) >> in preview
  - Azure - 7 (Asia Pacific-2, EU-1, North America-4)

## // Connecting to SF //

	Create	Clone	Drop	Modify	Load Data	Transfer Ownership	
Tables	✓	✓	✓		✓	✓	- Web-based UI (see above for usage, capabilities, restriction)
Views	✓		✓		✓		- Command line client (SnowSQL)
Schemas	✓	✓	✓		✓		- ODBC and JDBC (you have to download the driver!)
Stages	✓	✓	✓	✓	✓		- Native Connectors (Python, Spark & Kafka)
File Formats	✓	✓	✓	✓	✓		- Third party connectors (e.g. Matillion)
Sequences	✓	✓	✓	✓		✓	- Others (Node.js, .Net)

### Snowflake Address

[https://account\\_name\\_region\\_\(provider\).snowflakecomputing.com](https://account_name_region_(provider).snowflakecomputing.com)

Account\_name either:

- AWS = account\_name.region
- GCP/Azure = account\_name.region.gcp/azure

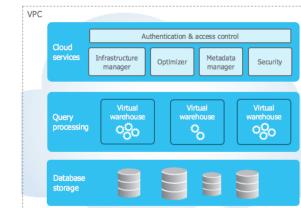
e.g. <https://pp12345.ap-southeast-2.snowflakecomputing.com>

## // Architecture //

### Overview

- Hybrid of shared-disk db and shared-nothing db
  - Uses central data repo for persisted storage - All compute nodes have data access
  - Using MPP clusters to process queries - Each node stores portion of data locally

### Architectural layers



### Storage

- All data stored as an internal, optimised, compressed columnar format (micro-partitions - represents logical structure of table)
- Can't access the data directly, only through Snowflake (SQL etc)

### Query Processing

- Where the processing of query happens
- Uses virtual warehouses - an MPP compute cluster
  - Each cluster = Independent - doesn't share resources with other vwh and doesn't impact others

### Cloud Services

- Command centre of SF - coordinates and ties all activities together within SF
- Gets provisioned by Snowflake and within AWS, Azure or GCP
  - You don't have access to the build of this or to do any modifications
  - Instance is shard to other SF accounts unless you have VPS SF Edition
- Services that this layer takes care of:
  - Authentication
  - Infrastructure management
  - Metadata management
  - Query parsing and optimisation
  - Access control

### Caches

- Snowflake Caches different data to improve query performance and assist in reducing cost

>Metadata Cache - Cloud Services layer

- Improves compile time for queries against commonly used tables

>Result Cache - Cloud Services layer

- Holds the query results
- If Customers run the exact same query within 24 hours, result cache is used and no warehouse is required to be active

>Local Disk Cache or Warehouse Cache - Storage Layer

- Caches the data used by the SQL query in its local SSD and memory
- This improves query performance if the same data was used (less time to fetch remotely)
- Cache is deleted when the Warehouse is suspended

## // Data Loading //

### File Location

- On Local
- On Cloud
  - AWS (Can load directly from S3 into SF)
  - Azure (Can load directly from Blob Storage into SF)
  - GCP (Can load directly from GCS into SF)

### File Type

- Structured
  - Delimited files (CSV, TSV etc)
- Sem-structured
  - JSON (SF can auto detect if Snappy compressed)
  - ORC (SF can auto detect if Snappy compressed or zlib)
  - Parquet (SF can auto detect if Snappy compressed)
  - XML (in preview)
- \* Note: If files = uncompressed, on load to SF it is gzip (can disable)

### >Compression

- SF auto compresses data from local fs to gzip - can change (AUTO\_COMPRESS)
- Specify compression type on loading compressed data

### >Encryption for load

- Loading unencrypted files
  - SF auto encrypts files using 128-bit keys! (or 256-keys - requires configuring)
- Loading encrypted files
  - Provide your key to SF on load

### Best Practice

#### >File Sizing

- ~10-100 MB file, compressed - Optimises parallel operations for data load
- Aggregate smaller files - Minimises processing overhead
- Split up larger files to small files - Distributes workload among server in active vwh
- # data files processed in parallel depends on capacity of servers in vwh
- Parquet: >3GB compressed could time out - split into 1GB chunks

#### >Semi Structured Sizing

- VARIANT data type has 16 MB compressed size limit per row
- For JSON or Avro, outer array structure can be removed using *STRIP\_OUTER\_ARRAY*.

#### >Continues Data Loads File Sizing

- Load new data within a minute after file notification sent, longer if file is large
  - Large compute is required (decompress, decrepit, transform etc).
- If > 1min to accumulate MBs of data in source app, create data file once per min
  - Lowers cost, improve performance (load latency)
- Creating smaller files and staging them in cloud storage >1 time per minute = disadvantage
  - Reduction in latency between staging and loading data != guaranteed
  - Overhead to manage file in internal load queue - increase utilisation cost
  - Overhead charges apply if per than 1000 files in queue

### Planning Data Load

#### Dedicating Separate Warehouse

- Load of large data = affect query performance
  - Use separate Warehouse
  - # data files processed in parallel determined by servers in the warehouse.
  - Split large data to scale linearly, use of small vwh should be sufficient

#### Staging Data

##### > Organising Data by Path - Snowflake Staging

- Both internal and external ref can include path (prefix) in the cloud storage
- SF recommends (for file organisation in Cloud Storage):
  1. Partition into logical path - includes identifying detail with date
  2. Organising file by path = allow copy fraction of partitioned data in SF in one command
    - Result: Easier to execute concurrent COPY statements - takes advantage of parallel operations
- Named stage operation:
  - Staging = A place where the location/path of the data is stored to assist in processing the upload of files
  - Remember: Files uploaded to snowflake Staging area using PUT are automatically encrypted with 128-bit or 256-bit (CLIENT\_ENCRYPTION\_KEY\_SIZE to specify) keys!

#### Loading Data

##### > COPY command - parallel execution

- Supports loading by Internal Stage or S3 Bucket path
- Specifying specific list of files to upload (1000 files max at a time)
- Identify files through pattern matching (regex)
  - Validating Data:
    - Use VALIDATION\_MODE - validate errors on load - does not load into table
    - ON\_ERROR to run actions to follow
- When COPY command runs, SF sets load status in the table's metadata
  - Prevents parallel COPY from loading the same file
  - On complete, SF sets the load status and to those that failed as well
  - Metadata (expires after 64 days):
    - Name of file
    - File size
    - ETag of the file
    - # rows parsed in file
    - Timestamp of last load
    - Information on errors during load
- SF recommends removing the Data from the Stage once the load is completed to avoid reloading again - use REMOVE command (and specify PURGE in COPY argument)
  - Improves performance - doesn't have to scan processed files

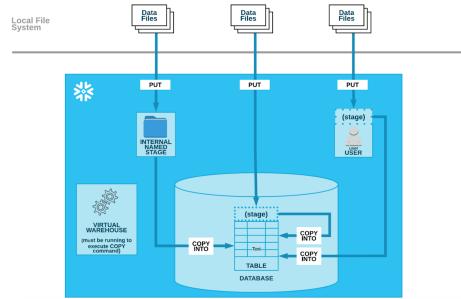
##### > Note for Semi-structured

- SF loads ss data to VARIANT type column
- Can load ss data into multiple columns but ss data must be stored as field in structured data
- Use FLATTEN to explode compounded values into multiple rows

#### >Data Transformation during load

- Supported
  - Sequence, substring, to\_binary, to\_decimal
- Commands not supported
  - WHERE, FLATTEN, JOIN, GROUP BY, DISTINCT (not fully supported)
  - VALIDATION\_MODE (if any aggregation applied)
  - CURRENT\_TIME (will result in the same time)

## Snowflake Stage



## Types of Stages

- Default: each table and user are allocated an internal named stage
  - Specify Internal Stage in PUT command when uploading file to SF
  - Specify the Same Stage in COPY INTO when loading data into a table

### > User Stage (Ref: @~)

- Accessed by single user but need copying to multiple tables
- Can't be altered or dropped
- Can't set file format, need to specify in COPY command to table

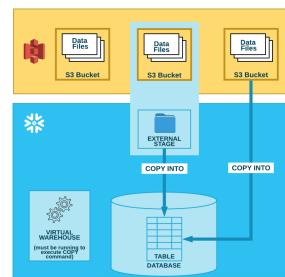
### > Table Stage (Ref: @%)

- Accessed by multiple users and copied to single tale
- Can't be altered or dropped
- Can't set file format, need to specify in COPY command to table
- No transformation while loading

### Internal Named Stages (Ref: @)

- A Database object
- Can load data into any tables (Needs user with privilege)
- Ownership of stage can be transferred

## AWS - Bulk Load Loading from S3



- SF uses S3 Gateway Endpoint. If region bucket = SF region - no route through public internet.

### > Securely access S3

1. Configure storage integration object - delegate authentication responsibility for external cloud storage to SF identity and IAM
2. Configure AWS IAM role -> allow S3 bucket access
3. Configure IAM user, provide Key and Secret Key

## Snowpipe - Incremental Load

### > Introduction

	COPY Command	Snowpipe
Authentication	Security options supported by the client for authenticating and initiating a user session.	<b>When calling the REST endpoints:</b> Key pair authentication with JSON Web Token (JWT). JWTs are signed using a public/private key pair with RSA encryption.
Transactions	Adds data to a table in transactions alongside any other SQL statements submitted manually by users.	Adds data to a table in transactions controlled by Snowflake with no opportunity to involve other statements in the transaction.
Load History	Stored in the metadata for the target table for 64 days. Available upon completion of the COPY statement as statement results.	Stored in the metadata for the pipe for 14 days. Must be requested from Snowflake (via a REST endpoint, SQL table function, or ACCOUNT_USAGE view).
Compute Resources	Requires a user-specified warehouse to execute COPY statements.	Uses Snowflake-supplied compute resources.
Cost	Billed for the amount of time each virtual warehouse is active.	Billed according to the compute resources used in the Snowpipe warehouse while loading the files.

- Enable you to loads data as soon as they are in stage (uses COPY command)
- Uses Pipe (first-class SF object) which contains the COPY command
  - Can DROP, CREATE, ALTER, DESCRIBE and SHOW PIPES
- Think of this as SF way of streaming data into the table as long as you have created a Named Stage.
- Generally loads older files first but doesn't guarantee - Snowpipe appends to queue.
- Has loading metadata to avoid duplicated loads

### > Billing and Usage

- User doesn't need to worry about managing vwh
- Charges based on resource usage - consume credit only when active
- View charges and usage
  - Web UI - Account > Billing & Usage
  - SQL - Information Schema > PIPE\_USAGE\_HISTORY

### > Automating Snowpipe

- Check out their AWS, GCP, Azure, and REST API article

## Querying from Staged File - Query data outside SF

- SF allows you to query data directly in external locations as long as you have created a stage for it
  - SF recommends using simple queries only for
- Performance = impacted since it's not compressed and columnised to SF standard

## // Monitoring //

### Resource Monitor

- Helps control the cost and unexpected spike in credit usage
- Set Actions (or triggers) to Notify and/or Suspend if credit usage is above certain threshold
- Set intervals of monitoring
- Created by users with admin roles (ACCOUNTADMIN)
- If monitor suspends, any warehouse assigned to the monitor cannot be resumed until
  1. Next interval starts
  2. Monitor is dropped
  3. Credit quota is increased
  4. Warehouse is no longer assigned to the monitor
  5. Credit threshold is suspended

## // Unloading Data //

### Introduction

- Allows you to export your data out of SF
- Uses COPY INTO <location> (either external or SF stage)
  - If SF then use GET command to download file
  - Use SELECT and other full SQL syntax to build data for export
  - Can specify MAX\_FILE\_SIZE to split into multiple files by chunks

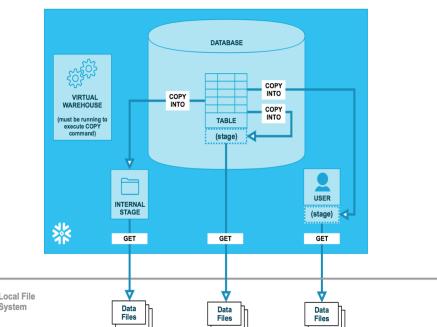
### Format and Restrictions

- Allowed formats for export (only UTF-8 as encoding allowed)
  - Delimited Files (CSV etc)
  - JSON
  - Parquet
- Allowed compression for export
  - gzip (default)
  - bzip2
  - Brotli
  - Zstandard
- Encryption for export
  - Internal Stage - 128-bit or 256-bit - gets decrypted when downloaded
  - External - Customer supply encryption key

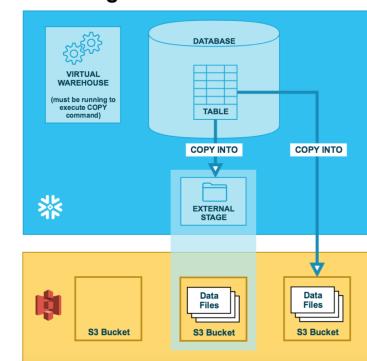
### Considerations

- Empty String and Null
  - FIELD\_OPTIONALLY\_ENCLOSED\_BY
  - EMPTY\_FIELD\_AS\_NULL
  - NULL\_IF
- Unloading a single file
  - MAX\_FILE\_SIZE default = 16mb
  - Max max is 5gb for AWS, GCP, and 256mb for Azure
- Unloading Relational table to JSON
  - User OBJECT\_OBSTRUCT obstrict in COPY command to convert rows of relational table to VARIANT column then unload the data as per usual

### Unloading in SF Stage



### Unloading into S3



## // Virtual Warehouse //

### Cluster of resource

- Provides CPU, Mem and Temp Storage
- Is active on usage of SELECT and DML (DELETE, INSERT, COPY INTO etc)
- Can be stopped at any time and resized at any time (even while running)
  - Running queries are do not get affected, only new queries
- Warehouse sizes = T-shirt sizes (generally query performance scales linearly with vwh size)
  - X-Small - 4X-Large
- When creating a vwh you can specify:
  - Auto-suspend (default: 10mins) - suspends warehouse if active after certain time
  - Auto-resume - auto-resumes warehouse whenever a statement that requires active vwh is required
- Query Caching
  - vwh maintains cache of table data - Improves query performance
  - The Larger the vwh, larger the cache - Cache dropped when vwh suspended

### Choosing VWH

- SF recommends experimenting by running the same queries on different vwh
- Monitoring Warehouse Load - WebUI: In Warehouse (view queued and running queries)

### Query Processing and Concurrency

- When queries submitted, vwh calls and reserves resource
  - Query is queued if vwh doesn't have enough resource
    - STATEMENT\_QUEUED\_TIMEOUT\_IN\_SECONDS
    - STATEMENT\_TIMEOUT\_IN\_SECONDS
  - Both can be used to control query processing and concurrency
- Size and complexity of query determines the concurrency (also # queries being executed)

### Multi-cluster warehouse

- Up to 10 server clusters
- Auto-suspend and auto-resume is of whole cluster not 1 server
- Can resize anytime
- Multi-cluster mode
  - Maximised - associate same min and max for # cluster
  - SF starts all clusters at start would be good for expected workload
- Auto-scale - different min and max for # cluster
  - To help control this, SF provides scaling policy

Policy	Description	Cluster Starts...	Cluster Shuts Down...
Standard (default)	Prevents/minimizes queuing by favoring starting additional clusters over conserving credits.	Immediately when either a query is queued or the system detects that there's one more query than the currently-running clusters can execute.	After 2 to 3 consecutive successful checks (performed at 1 minute intervals), which determine whether the load on the least-loaded cluster could be redistributed to the other clusters without spinning up the cluster again.
Economy	Conserves credits by favoring keeping running clusters fully-loaded rather than starting additional clusters, which may result in queries being queued and taking longer to complete.	Only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes.	After 5 to 6 consecutive successful checks (performed at 1 minute intervals), which determine whether the load on the least-loaded cluster could be redistributed to the other clusters without spinning up the cluster again.

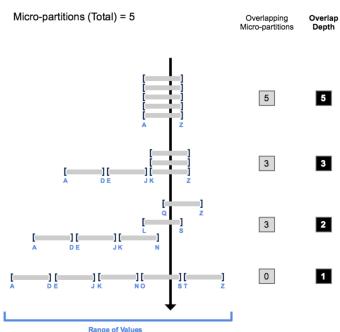
### Some rules:

- Maximized:
  - ↑ max & min: Specified number of clusters start immediately.
  - ↓ max & min: Specified number of clusters shut down when they finish executing statements and the auto-suspend period elapses.
- Auto-scale:
  - ↑ max: If `new_max_clusters > running_clusters`, no changes until additional clusters are needed.
  - ↓ max: If `new_max_clusters < running_clusters`, excess clusters shut down when they finish executing statements and the scaling policy conditions are met.
  - ↑ min: If `new_min_clusters > running_clusters`, additional clusters immediately started to meet the minimum.
  - ↓ min: If `new_min_clusters < running_clusters`, excess clusters shut down when they finish executing statements and the scaling policy conditions are met.

## // Tables in Snowflake //

### Micro-partition

- All of SF Tables are divided into micro-partition
- Each micro-partitions are compressed columnar data
  - Max size of 16mb compressed (50-500MB of uncompressed data)
  - Stored on logical hard-drive
  - Data only accessible via query through SF not directly
  - They are immutable, can't be changed
- Order of data ingestion are used to partition these data
- SF uses natural clustering to colocate column with the same value or similar range
  - Results to non-overlapping micro-partition and least depth
    - Improves query performance to avoid scanning unnecessarily micr-partitions



- Metadata of the micropartitions are stored in the Cloud Services Layer
  - Customers can query this without needing an Active vwh
    - Range of each columns
    - Number of Distinct Values
    - Count NULL

### Clustering Key

- SF recommends the use of Clustering key once your table grows really large (multi-terabyte range).
  - Especially when data does not cluster optimally.
- Clustering keys are subset of columns designed to colocate data in the table in the same micro-partitions
  - Improves query performance on large tables by skipping data that is not part of filtering predicate

### Zero-copy Cloning

- Allows Customers to Clone their table
- SF references the original data (micro-partition) - hence zero-copy (no additional storage cost)
  - When a change is made to the cloned table (new data added, deleted or updated) then a new micro-partition is created (incurs storage cost).
- Cloned object do not inherit the source's granted privileges

### Types of Tables (Internal)

#### >Temporary Tables

- Used to store data temporarily
- Non-permanent and exists only within the session - data is purged after session ends
- Not visible to other users and not recoverable
- Contributes to overall storage
- Belongs to DB and Schema - Can have the same name as another non-temp table within the same DB and Schema
- Default 1 day Time Travel

#### >Transient Tables (or DB and Schema)

- Persists until dropped
- Have all functionality as permanent table but no Fail-safe mode (no FS storage cost)
- Default 1 day Time Travel

#### >Permanent Table

- Have 7 days fail safe
- Default 1 day Time Travel (up to 90 days for >=Enterprise edition)

#### >External Table

- Allows access to data stored in External Stage

#### >Creating tables

- SF doesn't enforce any constraints (primary key, unique, foreign key) beside the NOT NULL constraint

### Types of Views

#### >Non-materialized views (views)

- Named definition of query
- Result are not stored

#### >Materialized views

- Result are stored
- Faster performance than normal views
- Contributes towards storage cost

	Performance Benefits	Security Benefits	Simplifies Query Logic	Supports Clustering	Uses Storage	Uses Credits for Maintenance	Notes
Regular table					✓	✓	
Regular view		✓	✓				
Cached query result		✓					Used only if data has not changed and if query only uses deterministic functions (e.g. not CURRENT_DATE).
Materialized view	✓	✓	✓	✓	✓	✓	Storage and maintenance requirements typically result in increased costs.

#### >Secure view

- Both non-materialized and materialized view can be defined as Secure view
- Improves data privacy - hides view definition and details from unauthorised viewers
- Performance is impacted as internal optimization is bypassed

## // Time Travel and Fail Safe //

### Introduction

## Continuous Data Protection Lifecycle

Standard operations allowed:  
Queries, DDL, DML, etc.

**Time Travel allowed:**  
SELECT ... AT|BEFORE ...  
CLONE ... AT|BEFORE ...  
UNDROP ...

No user operations allowed  
(data recoverable only by Snowflake)

Current Data Storage

Time Travel Retention  
(1-90 Days)

Fail-Safe  
(transient: 0 days,  
Permanent:7 days)

#### > Time Travel (0-90 days)

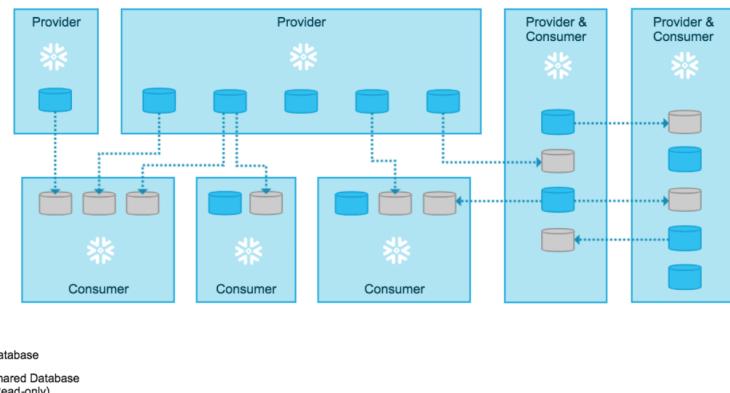
- Allows db to query, clone or restore historical data to tables, schema or db for up to 90 days
  - Can use this to access snapshot of data at a point in time
- Useful if data was deleted, dropped or updated.

#### > Fail safe (7 days)

- Allows disaster recovery of historical data - Only accessible by Snowflake!

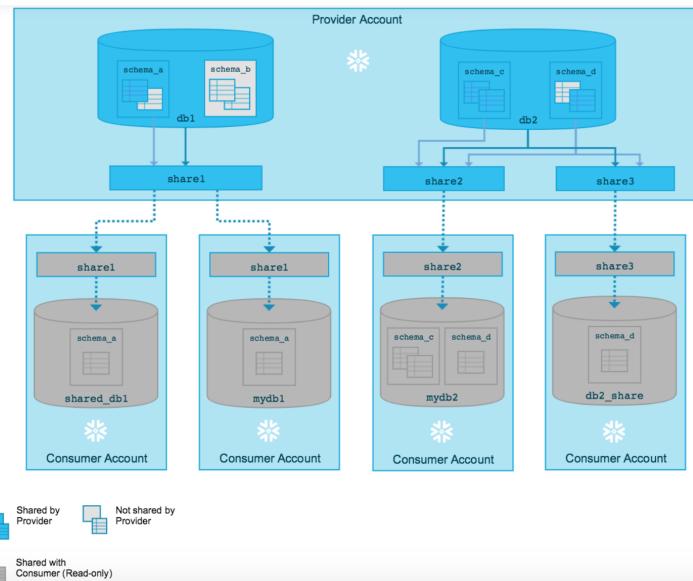
## // Data Sharing //

### Introduction



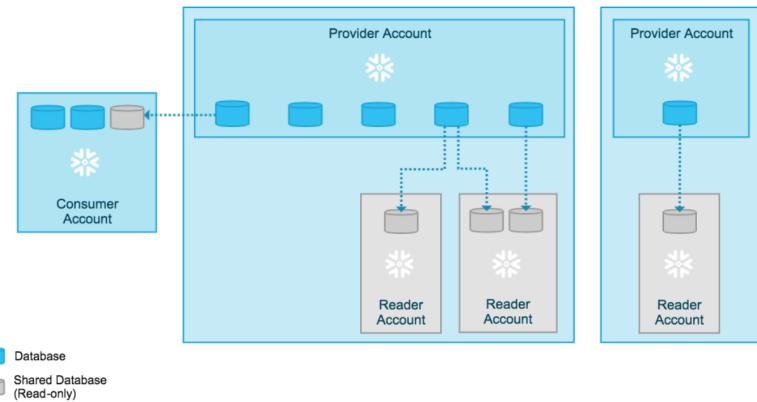
- Only ACCOUNTADMIN can provision share
- The sharer is the Provider while the user of the shared data is the Consumer
- Only ACCOUNTADMIN role can create this
- Secure data sharing enables account-to-account sharing of Database, Tables and Views
  - No actual data is copied or transferred
- Sharing accomplished using SF service layer and metadata store
- Charge is compute resources used to query shared data
- User creates share of DB then grants object level access
  - read-only db is created on consumer side
- Access can be revoked at any time
- Type of share includes Share and Reader
- No limit on number of shares or accounts but 1 DB per share

### Share



- Between SF accounts
- Each share consists of
  - Privilege that grant db access
  - Privilege that grant access to specific object (tables, views)
  - Consumer accounts with which db and object are shared
- Can perform DML operations

### Reader Account



- Share data with user who isn't on SF
- Objects can only be read not modified (no DML) - consumer can only consume the data

## // Access Control within Snowflake //

### Network Policy

- Allows access based on IP whitelist or restrictions to IP blacklist
  - Apply through SQL or WebUI
  - Only ACCOUNTADMIN or SECURITYADMIN can modify, drop or create these

### PrivateLink

- If time permits briefly read through these (AWS and Azure)

### MFA

- Powered by Duo System and enrolled on all accounts, Customers need to enable it
- Recommend enabling MFA on ACCOUNTADMIN
- Use with WebUI, SnowSQL, ODBC, JDBC, Python Connector

### Federated Auth and SSO

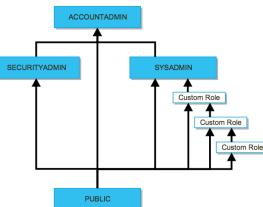
- User authenticate through external SAML 2.0-compliant identity provider (IdP)
  - Users don't have to log into Snowflake directly
- As per basic introduction to SSO, a token is passed to the application that the user is trying to login to authenticate which in turns will open access when verified

### Access Control Models

SF approach to access control uses the following models:

- Discretionary Access Control (DAC) - All object have an owner, owner can grant access to their objects
- Role-based Access Control (RBAC) - Privileges are assigned to roles, then roles to users
  - Roles are entities that contains granted privileges (level of access to object) which are

### System Defined Roles



### >ACCOUNTADMIN

- encapsulates SECURITYADMIN and SYSADMIN

### >SECURITYADMIN

- Creates, modify and drops user, roles, networks, monitor or grants

### >SYSADMIN

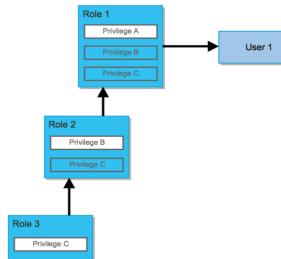
- Has privileges to create vwh, and db and its objects
- Recommend assigning all custom roles to SYSADMIN so it has control over all objects created (create role hierarchy based on this)

### >PUBLIC

- Automatically grant to all users
- Can own secured objects
- Used where explicit access control is not required

### >Custom Roles

- Created by SECURITYADMIN
- Assigned to SYSADMIN
  - If not assigned, only roles with MANAGE GRANTS can modify grants on it.
- Create custom roles with least privilege and role them up

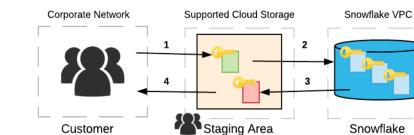


- Visit "Access Control Privileges" in SF Doc and have a quick skim over the privileges

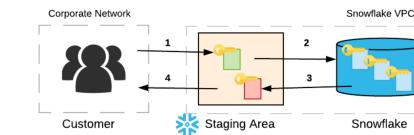
## // Data Security within Snowflake //

### > End-to-end Encryption

- Snowflake encrypts all data by default, no additional cost
- No one other than the customer or runtime components can read the data
  - Data is always protected as it is in an encrypted state
- Customer provides encrypted data to the external staging area (S3 etc)
  - Provide SF with the encryption master key when creating the Named Stage



- Client Side Encryption
  - Customer provides unencrypted data to SF internal staging area, SF will automatically encrypt the data.



### >Key Rotation

- Snowflake encrypts all data by default and keys are rotated regularly
- Uses Hierarchy Key Model for its Encryption Key Management - comprises of
  - Root Key
  - Account Master Key (auto-rotate if >30 days old)
  - Table Master Key (auto-rotate if >30 days old)
  - File Keys

### >Tri-Secret Secure and Customer-managed keys (Business Critical Edition)

- Combines the customer key with snowflake's maintained key
  - Creates composite master key then use it to encrypts all data in the account
  - If customer key or snowflake key is revoked, data cannot be decrypted

Question 1: Skipped

According to snowflake architecture, data storage is independent of compute

- TRUE  
**(Correct)**
- FALSE

### Explanation

Snowflake's architecture has been defined with the belief that a tight coupling of compute and storage will not be an effective approach of leveraging the scale of cloud. The idea is to be able scale each of this component independent of each other. Snowflake's **multi-cluster, shared data** architecture separates compute resource scaling from storage resources, thus enabling seamless, non-disruptive scaling.

For data storage, Snowflake leverages the power of the respective cloud providers object store(AWS S3, Azure BLOB, GCS) and for compute it leverages the cloud VM instances which is called '**VIRTUAL WAREHOUSE**' in snowflake. Virtual warehouses are nothing but a set of compute clusters abstracted using the underlying cloud VM machines.

Question 2: Skipped

All virtual warehouses in snowflake has access to all data

- TRUE  
**(Correct)**
- FALSE

### Explanation

Snowflake was built with a unique new architecture that makes it possible to solve this challenge. Snowflake's unique **multi-cluster, shared data architecture** makes it possible to allocate multiple independent, isolated clusters for processing while sharing the same data. Each cluster (we call them "virtual warehouses") can both read and write data, with full transactional consistency ensured by Snowflake's cloud services layer. The size and resources of each virtual warehouse can be chosen independently by the user based on the characteristics and performance requirements for the workload(s) that will run on each virtual warehouse.

<https://www.snowflake.com/blog/snowflake-challenge-concurrent-load-and-query/>

Question 3: Skipped

Which layer in the snowflake architecture contains virtual warehouse

- Cloud services
- Query processing  
**(Correct)**

- Database Storage
- None

### Explanation

Virtual warehouses are used to provide compute resources for processing queries. Hence it falls under query processing

Question 4: Skipped

Which layer in snowflake contains data in compressed, columnar format

- Cloud services
- Query processing
- Database storage  
**(Correct)**
- None

### Explanation

Snowflake stores data in compressed and encrypted format by default in the database storage layer. Data is stored in micropartitions in a columnar format.

Question 5: Skipped

Which of the below approach would result in a better performance through linear scaling of data ingestion workload

- Resize virtual warehouse
- Consider the practice of organizing data by granular path
- Consider the practice of splitting input file batch within the recommended range of 10 MB and 100 MB
- All of the above  
**(Correct)**

### Explanation

Resizing a warehouse to a larger size is useful when the operations being performed by the warehouse will benefit from more compute resources, including:

1. Improving the performance of large, complex queries against large data sets.
2. Improving performance while loading and unloading significant amounts of data.

Note that, as a general practice, while performing data loads for Snowflake ETL, it is important to partition the data in your Snowflake or external locations like S3 buckets or Azure containers using logical, granular paths. By creating a partition using details like location or application along with the date when this data was written, you are optimizing for a later data loading activity. When you load the data, you can simply copy any fraction of the partitioned data into Snowflake with a single command. You can copy data into Snowflake by the hour, day, month, or even year when you initially populate tables.

The best practice is to load smaller files of 10-100 MB to maximize parallel loading i.e, number of files will be distributed between available nodes for faster processing and performance.

Due to all these, the answer to this question is "All of the above"

Question 6: Skipped

Snowflake provides specific administration features and capabilities to support the following activities except?

- Manage databases and warehouses within snowflake account
- Manage roles and users within a snowflake account
- Monitor usage and manage resources to control costs in a snowflake account
- Manage 3rd party applications providing data to a snowflake account  
**(Correct)**

Question 7: Skipped

Select all characteristics of snowflake's multi-cluster environment

- Multiple virtual warehouses in a deployment  
**(Correct)**
- Individual warehouses scale out/in based on query activity  
**(Correct)**
- Multi-cluster warehouses support the same properties and actions as single-cluster warehouse  
**(Correct)**
- User must specify which cluster each query will use

### Explanation

Multiple virtual warehouse in a deployment means specifying more than 1 warehouse while creating a multi-cluster warehouse. In a multi-cluster warehouse, you can specify the

maximum number of server cluster(or virtual warehouse) greater than 1 and upto 10 and also a minimum number of clusters(equal to or less than the maximum(up to 10)).

Additionally, multi-cluster warehouses **support all the same properties and actions as single-cluster warehouses**, including:

1. Specifying a warehouse size.
2. Resizing a warehouse at any time.
3. Auto-suspending a running warehouse due to inactivity; note that this does not apply to individual clusters, but rather the entire warehouse.
4. Auto-resuming a suspended warehouse when new queries are submitted.

In a multi-cluster warehouse, individual warehouses can scale up or down based on query load which is not possible in a single cluster warehouse. However as mentioned above, **it supports the same properties and actions like the single cluster warehouse**

**Question 8:** Skipped

What are the two modes in which a multi-cluster warehouse can run?

- Static
- Dynamic
- Maximized  
**(Correct)**
- Auto-scale  
**(Correct)**
- None of the above

### Explanation

You can choose to run a multi-cluster warehouse in either of the following modes:

#### Maximized

This mode is enabled by specifying the **same** value for both maximum and minimum clusters (note that the specified value must be larger than 1). In this mode, when the warehouse is

started, Snowflake starts all the clusters so that maximum resources are available while the warehouse is running.

This mode is effective for statically controlling the available resources (i.e. servers), particularly if you have large numbers of concurrent user sessions and/or queries and the numbers do not fluctuate significantly.

### Auto-scale

This mode is enabled by specifying *different* values for maximum and minimum clusters. In this mode, Snowflake starts and stops clusters as needed to dynamically manage the load on the warehouse:

As the number of concurrent user sessions and/or queries for the warehouse increases, and queries start to queue due to insufficient resources, Snowflake automatically starts additional clusters, up to the maximum number defined for the warehouse.

Similarly, as the load on the warehouse decreases, Snowflake automatically shuts down clusters to reduce the number of running servers and, correspondingly, the number of credits used by the warehouse.

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, SCALING\_POLICY, that determines the scaling policy to use when automatically starting or shutting down additional clusters. For more information, see [Setting the Scaling Policy for a Multi-cluster Warehouse](#) (in this topic).

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

Question 9: Skipped

What does snowflake offer to help control the credits consumed by a multi-cluster warehouse running in Auto-scale mode?

- Auto scale
- Scaling policy  
**(Correct)**
- Maximum\_number\_of\_server\_clusters
- Minimum\_number\_of\_server\_clusters
- None of the above

### Explanation

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, **SCALING\_POLICY**, that determines the scaling policy to use when automatically starting or shutting down additional clusters

Question 10: Skipped

What scaling policies are supported by snowflake?

- Standard  
**(Correct)**
- Economy  
**(Correct)**
- Premium

### Explanation

#### Standard (default)

Prevents/minimizes queuing by favoring starting additional clusters over conserving credits.

#### Economy

Conserves credits by favoring keeping running clusters fully-loaded rather than starting additional clusters, which may result in queries being queued and taking longer to complete.

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 11: Skipped

You are sizing a snowflake warehouse, what factors would you consider?

- Number of users
- Number of concurrent queries  
**(Correct)**
- Number of tables being queried  
**(Correct)**
- Data size and composition  
**(Correct)**

### Explanation

Warehouse size does not depend on number of users. Number of concurrent queries is important, because if the warehouse does not have enough resource, queries will start queueing. Same is the case with number of tables and data size.

Question 12: Skipped

Compute usage billing in snowflake is based on

- Per-second basis, with a minimum of 60 seconds  
**(Correct)**
- Per-second basis only
- By number of users

**Explanation**

Snowflake compute utilizes per second billing after the first minute. what it means is that when a warehouse starts or resumes, it bills for the first minute whether you use it or not, after that the billing is based on per second. For example if a warehouse is resumed and suspended after 30 seconds, you will still have to pay for 60 seconds because that is the first minute after the warehouse resumed.

Question 13: Skipped

SQL functionality in snowflake can be extended by

- USER DEFINED FUNCTIONS USING SQL  
**(Correct)**
- USER DEFINED FUNCTIONS USING JAVASCRIPT  
**(Correct)**
- SESSION VARIABLES  
**(Correct)**
- NONE OF THE ABOVE

**Explanation**

Snowflake supports SQL/SESSION variables declared by the user. They have many uses, such as storing application-specific environment settings.

<https://docs.snowflake.com/en/sql-reference/session-variables.html#sql-variables>

User-defined functions (UDFs) let you extend the system to perform operations that are not available through the built-in, system-defined functions provided by Snowflake. Snowflake currently supports two types of UDFs, *SQL* and *JavaScript*

Example of SQL UDF

```
create function area_of_circle(radius float)
    returns float
```

```
as  
$$  
    pi() * radius * radius  
$$  
;
```

### Example of JAVASCRIPT UDF

```
CREATE OR REPLACE FUNCTION array_sort(a array)  
RETURNS array  
LANGUAGE JAVASCRIPT  
AS  
$$  
    return A.sort();  
$$  
;  
  
-- Call the UDF with a small array.  
SELECT ARRAY_SORT(PARSE_JSON('[2,4,5,3,1]'));
```

#### Question 14: Skipped

To provide failover protection, how many availability zones does snowflake replicate to

- one
- two **(Correct)**
- three
- as configured by user

#### Explanation

Providing an even higher degree of data protection and service resilience, within the same deployment region, Snowflake provides standard failover protection **across three availability zones (including the primary active zone)**. Your data and business are protected. As you ingest your data, it is synchronously and transparently replicated across availability zones. This protection is automatically extended from Snowflake to customers, at no added charge.

<https://www.snowflake.com/blog/how-to-make-data-protection-and-high-availability-for-analytics-fast-and-easy/>

### Additional note

The answer to this question has been corrected(from 3 to 2) based on the feedback from a student. The explanation above is correct, but the replication happens to two additional availability zones only, and the primary zone is the original location of the data.

Question 15: Skipped

Which of the following is not a characteristic of micro-partitions in snowflake?

- New partitions are created in logical properties  
**(Correct)**
- Avoids skews between partitions
- Partitioning is automatically completed on the natural ingestion order of the data
- The natural ingestion order maintains correlations between columns which could be useful for pruning

### Explanation

Snowflake uses physical properties and partitions can overlap ranges

Question 16: Skipped

The storage architecture of snowflake has two key features, they are

- Time travel  
**(Correct)**
- Replication
- Zero-copy cloning  
**(Correct)**
- Query tag

### Explanation

Snowflake Time Travel enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It serves as a powerful tool for performing the following tasks:

1. Restoring data-related objects (tables, schemas, and databases) that may have been accidentally or intentionally deleted.

2. Duplicating and backing up data from key points in the past.
3. Analyzing data usage/manipulation over specified periods of time

Cloning also referred to as “zero-copy cloning” creates a copy of a database, schema or table. A **snapshot of data** present in the source object is taken when the clone is created and is made available to the cloned object. The cloned object is writable and is independent of the clone source. That is, changes made to either the source object or the clone object are not part of the other. Cloning a database will clone all the schemas and tables within that database. Cloning a schema will clone all the tables in that schema

Question 17: Skipped

Which is not a characteristic of time travel?

- Protection from accidental data operations
- Previous versions of data automatically retained
- Retention period selected by customers (up to 90 days for Enterprise edition)
- Recover data with the cost of running backups  
**(Correct)**

Question 18: Skipped

The commands to load data into snowflake are

- COPY AND PUT  
**(Correct)**
- COPY AND INSERT  
**(Correct)**
- INSERT AND PUT  
**(Correct)**
- NONE OF THE ABOVE

### Explanation

PUT in snowflake is used to load data from your local system to snowflake stage

```
put file:///tmp/data/mydata.csv @my_int_stage;
```

COPY and INSERT are the operations through which you can load data into snowflake tables

Question 19: Skipped

In order to load data into snowflake, which of the below are required?

- Virtual Warehouse  
**(Correct)**
- Predefined target table  
**(Correct)**
- Staging location with data staged  
**(Correct)**
- File format  
**(Correct)**

**Explanation**

Virtual warehouse is required to provide compute resources

Target table must be pre-defined load does not have feature to auto create tables

Data must be staged through the PUT command

File format must be defined to indicate whether CSV, JSON, XML etc is being loaded

Question 20: Skipped

Which of the following is not a snowpipe feature?

- Snowpipe can load data from any internal or external stage
- It is a server-less compute model
- The service provides REST end points and uses snowflake provided compute resources to load the data and retrieve history reports
- Snowpipe loads data after it is in stage and use executes the LOADDATA command  
**(Correct)**

**Explanation**

There is no LOADDATA command and data is automatically loaded by snowpipe after the files are added to the stage

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro.html#introduction-to-snowpipe>

Question 21: Skipped

Which feature does not come with query profile?

- 
- 
- 
- Hints for improving the query performance  
**(Correct)**
- 
- 

### Explanation

Go to query profile in snowflake web UI and see if you can find HINT:)

Question 22: Skipped

Snowflake mechanism to limit the number of micro-partitions scanned by a query is called what?

- Pruning  
**(Correct)**
- 
- 
- 
- 

### Explanation

Snowflake makes extensive use of pruning to reduce the amount of data that has to be read from storage. In summary, this means that a query like

```
SELECT SUM(x) FROM T1 WHERE y=42
```

will not read Columns x and y completely from storage but will make use of the predicate y=42 and limit the table scan to the subset of all partitions that can potentially match this condition. Snowflake can do this because it stores the min and max value of each of the columns in the metadata store

Question 23: Skipped

Which type of data integration tools leverage snowflake's scalable compute for data transformation?

- ELT  
**(Correct)**
-

- ETL
- - Database replication
  - 
  - Streaming

### Explanation

Commonly referred to as ETL, data integration encompasses the following primary operations:

#### Extract

Exporting data from specified data sources.

#### Transform

Modifying the source data (as needed), using rules, merges, lookup tables or other conversion methods, to match the target.

#### Load

Importing the resulting transformed data into a target database.

More recent usage **references the term** **ELT**, **emphasizing that the transformation part of the process** does not necessarily need to be performed before loading, particularly in systems such as Snowflake that support transformation during or after loading.

In addition, the definition of data integration has expanded to include a wider range of operations, including:

Data preparation.

Data migration/movement and management.

Data warehouse automation.

<https://docs.snowflake.com/en/user-guide/ecosystem-etl.html>

Question 24: **Skipped**

Snowflake offers tools to extract data from source systems

- TRUE
- FALSE  
**(Correct)**

Question 25: **Skipped**

Which of the following are options when creating a virtual warehouse

- Auto Suspend  
**(Correct)**
- Auto Resume  
**(Correct)**
- Storage Size
- Server count

## Explanation

### Syntax

```
CREATE [ OR REPLACE ] WAREHOUSE [ IF NOT EXISTS ] <name>
    [ [ WITH ] objectProperties ]
    [ objectParams ]
```

Where:

```
objectProperties ::=

WAREHOUSE_SIZE = XSMALL | SMALL | MEDIUM | LARGE | XLARGE | XXLARGE | XXXLARGE
GE | X4LARGE

MAX_CLUSTER_COUNT = <num>

MIN_CLUSTER_COUNT = <num>

SCALING_POLICY = STANDARD | ECONOMY

AUTO_SUSPEND = <num> | NULL

AUTO_RESUME = TRUE | FALSE

INITIALLY_SUSPENDED = TRUE | FALSE

RESOURCE_MONITOR = <monitor_name>

COMMENT = '<string_literal>'

objectParams ::=

MAX_CONCURRENCY_LEVEL = <num>

STATEMENT_QUEUED_TIMEOUT_IN_SECONDS = <num>

STATEMENT_TIMEOUT_IN_SECONDS = <num>
```

Question 26: Skipped

The warehouse cache may be reset if a running warehouse is suspended and then resumes

- TRUE  
**(Correct)**
- FALSE

### Explanation

You do not need to remember this, just understand the concept. Warehouses are nothing but VM instances underneath. And the cache is nothing but SSD cache attached to the instances, If the warehouse is suspended and resumed again it may not get the same instance.

Question 27: Skipped

A table in snowflake can only be queried using the virtual warehouse used to load the data

- TRUE
- FALSE  
**(Correct)**

### Explanation

Again apply concept, storage and compute are isolated in snowflake. That is why we can run different workloads in different warehouses but they all can refer to the same data

Question 28: Skipped

Snowflake caches are automatically invalidated if the underlying data changes

- TRUE  
**(Correct)**
- FALSE

Question 29: Skipped

The query profiler view is only available for completed queries

- TRUE
- FALSE  
**(Correct)**

Question 30: Skipped

Which best describes zero-copy cloning?

- Metadata only operation
- No replication of data
- Unmodified data stored once and modified data stored as new micro partitions
- All of the above  
**(Correct)**

### Explanation

When a clone is created of a table, the clone utilizes no data storage because it shares all the existing micro-partitions of the original table at the time it was cloned; however, rows can

then be added, deleted, or updated in the clone independently from the original table. Each change to the clone results in new micro-partitions that are owned exclusively by the clone and are protected through CDP(continuous data protection).

Question 31: Skipped

Your business team runs a set of identical queries every day after the batch ETL run is complete. From the following actions, what is the best action that you will recommend.

- After the ETL run, execute the identical queries so that they remain in the result cache  
**(Correct)**
- After the ETL run, resize the warehouse to a larger warehouse
- After the ETL run, copy the tables to another schema for the business users to query

### Explanation

Please note the key word here which is **IDENTICAL** queries. When a query is run the first time, and the same query is run the second time it picks up the results from the query result cache and it does not cost you any compute. The query result cache is valid for 24 hours and it gets extended for another 24 hours, every time you access it(even if you access it at 23 hours 59 seconds).

### Try this out yourself

Using worksheet create a table and insert the rows as below

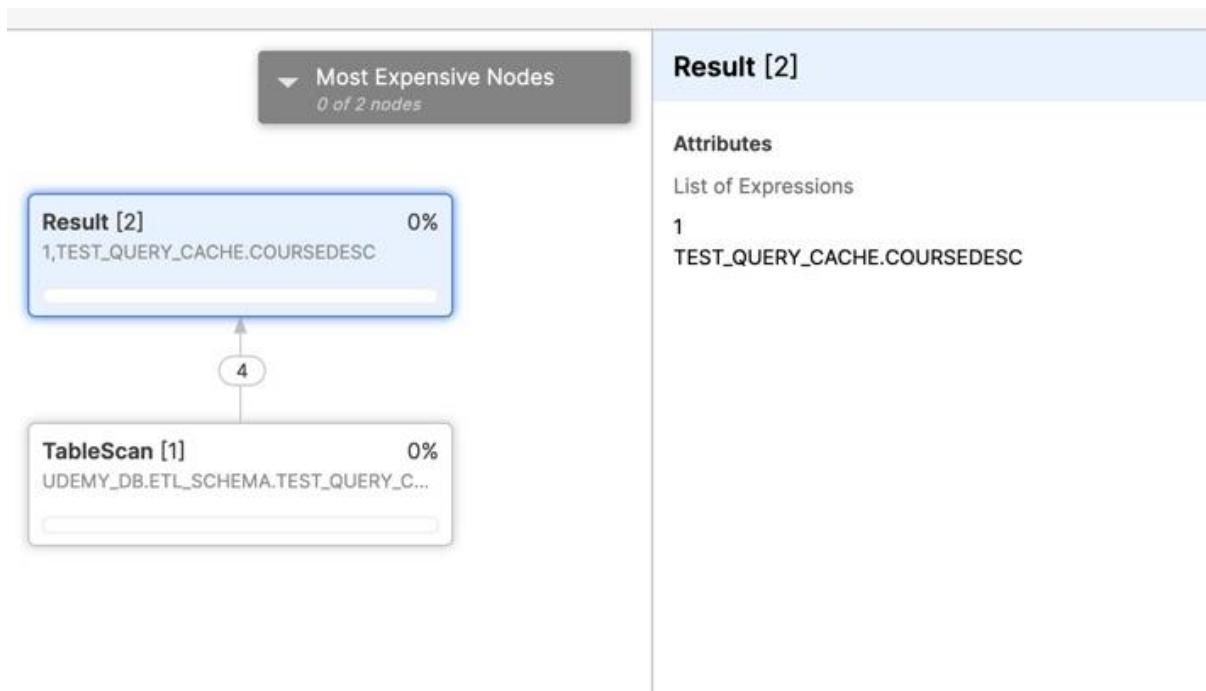
```
CREATE TABLE TEST_QUERY_CACHE(COURSEID NUMBER, COURSEDESC VARCHAR);

INSERT INTO TEST_QUERY_CACHE VALUES(1, 'SNOWFLAKE');
INSERT INTO TEST_QUERY_CACHE VALUES(1, 'C++');
INSERT INTO TEST_QUERY_CACHE VALUES(1, 'PYTHON');
```

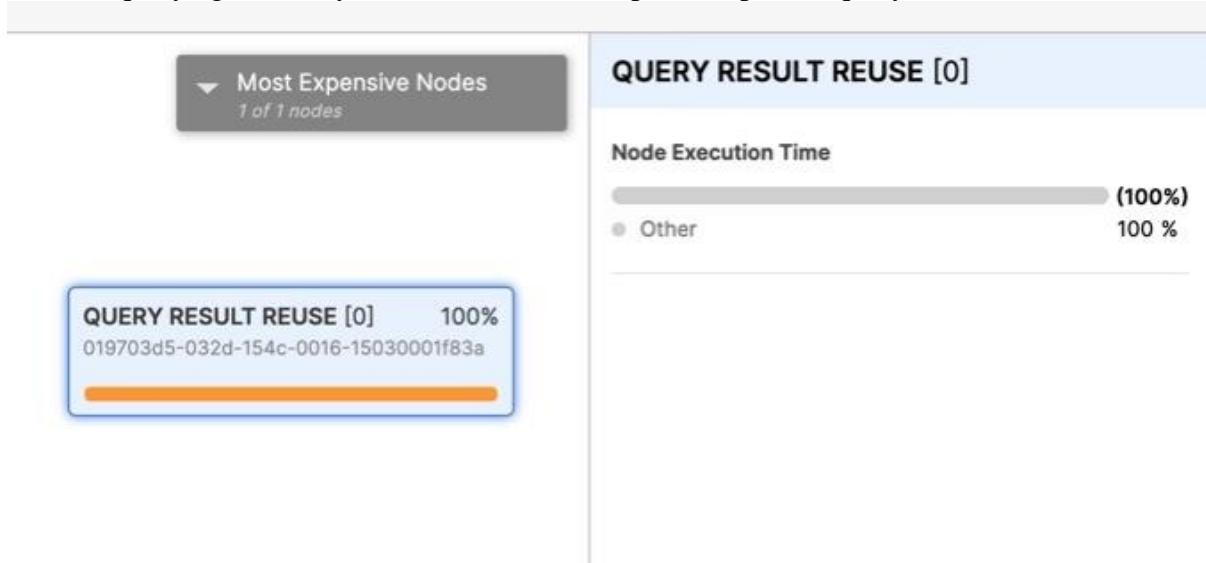
Run the below query and go to query profile

```
SELECT * FROM TEST_QUERY_CACHE;
```

You see that it is doing a table scan



Run the query again, and you will see that it has picked up from query cache



Next question:) How do you disable query cache. to know the answer look at the next question

Question 32: **Skipped**

If you would like to disable the query cache, what will you do?

- ALTER SESSION SET USE\_CACHED\_RESULT = FALSE;  
**(Correct)**
- ALTER SESSION SET USE\_CACHED\_RESULT = OFF;
-

```
ALTER SESSION SET USE_QUER_CACHE = TRUE;
```

### Explanation

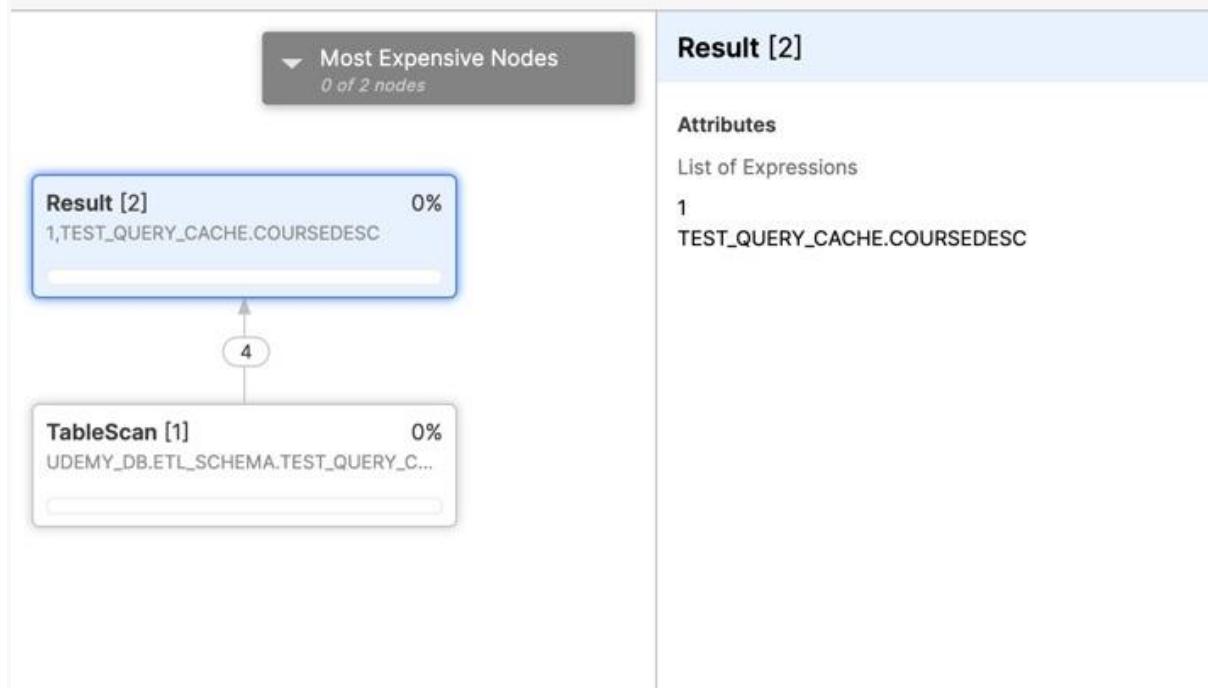
Try it again, run the below command

```
ALTER SESSION SET USE_CACHED_RESULT = FALSE;
```

Then run the query one more time, what do you see

```
SELECT * FROM TEST_QUERY_CACHE;
```

You will see the table scan coming back:)



Question 33: Skipped

A task can have a maximum of 100 *child* tasks

- TRUE  
**(Correct)**
- FALSE

### Explanation

This is something that you will need to remember

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state. An individual task in the tree is limited to a single predecessor task; however, a task can have a

maximum of 100 **child tasks** (i.e. other tasks that identify the task as a predecessor).

Question 34: Skipped

Snowflake guarantees that only one instance of a task with a defined **predecessor task** is running at a given time

- TRUE
- FALSE  
**(Correct)**

### Explanation

Snowflake cannot guarantee that only one instance of a task with a defined **predecessor task** is running at a given time.

Question 35: Skipped

You created an warehouse(ETL\_WH) which is sized at XSMALL, you want to resize the Warehouse to SMALL. How will you do that?

- 1. alter warehouse etl\_wh set warehouse\_size=SMALL;  
**(Correct)**
- You will need to recreate the warehouse with the new size
- Once created an warehouse can never be re-sized

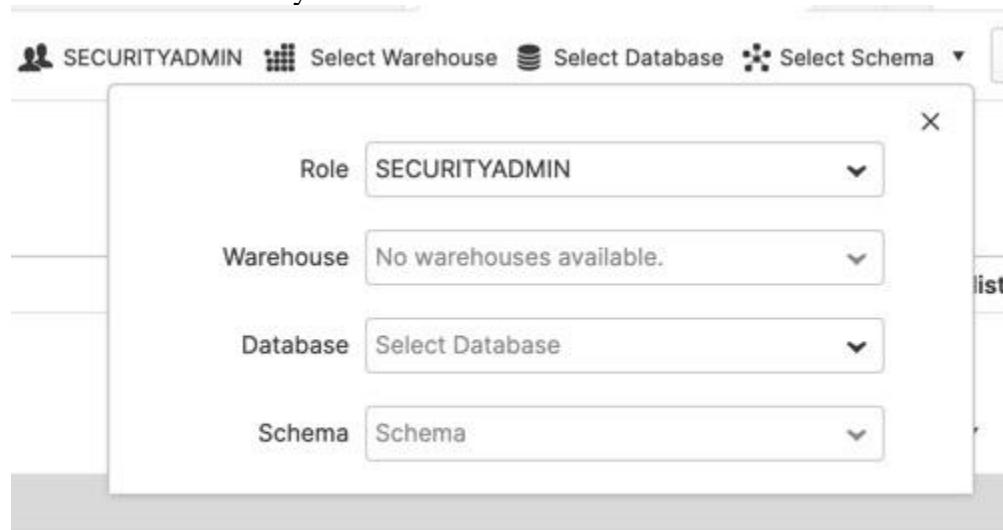
### Explanation

Ok, to remember it for the rest of your life, please log on to SNOWFLAKE and execute

```
ALTER WAREHOUSE ETL_WH SET WAREHOUSE_SIZE=SMALL;
```

The screenshot shows a Snowflake query editor interface. At the top, there is a toolbar with a 'Run' button, an 'All Queries' checkbox, and a 'Saved 2 seconds ago' message. Below the toolbar, a query history section displays a single query: '1 ALTER WAREHOUSE ETL\_WH SET WAREHOUSE\_SIZE=SMALL;'. The results section is titled 'Results' and shows a table with one row. The table has two columns: 'Row' and 'status'. The first row contains the value '1' under 'Row' and the text 'Statement executed successfully.' under 'status'. There are also buttons for 'Filter result...', 'Copy', and a download icon.

Ok just for fun, change your role to SECURITYADMIN, and try to resize the WAREHOUSE. Will you be able to do it?



No, by default SECURITYADMIN or any other role cannot change the warehouse size. You need to be SYSADMIN or ACCOUNTADMIN. Recommended approach is to use SYSADMIN .

```
1 ALTER WAREHOUSE ETL_WH SET WAREHOUSE_SIZE=SMALL;
```

Results Data Preview

✖ Query ID SQL 20ms

SQL compilation error: Warehouse 'ETL\_WH' does not exist or not authorized.

Question 36: Skipped

You have just loaded a file named student\_data.csv from your snowflake stage to snowflake table. You try to reload again using the COPY INTO command. You will be able to load the file.

- TRUE
- FALSE  
**(Correct)**

### Explanation

Let's understand why we will not be able to load the file again. Please also look at the **64 Days** mentioned below

Snowflake maintains detailed metadata for each table into which data is loaded, including:

1. Name of each file from which data was loaded
2. File size
3. ETag for the file
4. Number of rows parsed in the file
5. Timestamp of the last load for the file
6. Information about any errors encountered in the file during loading

This load metadata expires after 64 days. If the LAST\_MODIFIED date for a staged data file is less than or equal to 64 days, the COPY command can determine its load status for a given table and prevent reloading (and data duplication). The LAST\_MODIFIED date is the timestamp when the file was initially staged or when it was last modified, whichever is later.

### So, is there no way to reload it?

There is a way, you can use FORCE=TRUE

#### Reloading Files

Add **FORCE = TRUE** to a COPY command to reload (duplicate) data from a set of staged data files that have not changed (i.e. have the same checksum as when they were first loaded).

In the following example, the first command loads the specified files and the second command forces the same files to be loaded again (producing duplicate rows), even though the contents of the files have not changed:

```
copy into load1 from @%load1/data1/
  files=('test1.csv', 'test2.csv')

copy into load1 from @%load1/data1/
  files=('test1.csv', 'test2.csv')
  force=true;
```

Question 37: Skipped

For a particular snowflake session, you can specify more than one virtual warehouse

- TRUE
- FALSE  
**(Correct)**

### Explanation

A session can only use one warehouse but you can change the warehouse using USE WAREHOUSE <WAREHOUSE\_NAME>

## Warehouse Usage in Sessions

When a session is initiated in Snowflake, the session does not, by default, have a warehouse associated with it. Until a session has a warehouse associated with it, queries cannot be submitted within the session.

### Default Warehouse for Users

To facilitate querying immediately after a session is initiated, Snowflake supports specifying a default warehouse for each individual user. The default warehouse for a user is used as the warehouse for all sessions initiated by the user.

A default warehouse can be specified when creating or modifying the user, either through the web interface or using [CREATE USER/ALTER USER](#).

### Default Warehouse for Client Utilities/Drivers/Connectors

In addition to default warehouses for users, any of the Snowflake clients (SnowSQL, JDBC driver, ODBC driver, Python connector, etc.) can have a default warehouse:

SnowSQL supports both a configuration file and command line option for specifying a default warehouse.

The drivers and connectors support specifying a default warehouse as a connection parameter when initiating a session.

For more information, see [Connecting to Snowflake](#).

### Precedence for Warehouse Defaults

When a user connects to Snowflake and start a session, Snowflake determines the default warehouse for the session in the following order:

Default warehouse for the user,

» **overridden by...**

Default warehouse in the configuration file for the client utility (SnowSQL, JDBC driver, etc.) used to connect to Snowflake (if the client supports configuration files),

» **overridden by...**

Default warehouse specified on the client command line or through the driver/connector parameters passed to Snowflake.

Note

In addition, the default warehouse for a session can be changed at any time by executing the USE WAREHOUSE command within the session.

Question 38: Skipped

Time travel and fail safe requires additional storage which will cost you storage cost

- TRUE  
**(Correct)**
- FALSE

**Explanation**

Storage fees are incurred for maintaining historical data during both the Time Travel and Fail-safe periods.

**Storage Usage and Fees**

The fees are calculated for each 24-hour period (i.e. 1 day) from the time the data changed. The number of days historical data is maintained is based on the table type and the Time Travel retention period for the table.

Also, Snowflake minimizes the amount of storage required for historical data by maintaining only the information required to restore the individual table rows that were updated or deleted. As a result, storage usage is calculated as a percentage of the table that changed. Full copies of tables are only maintained when tables are dropped or truncated.

Question 39: Skipped

From which of the following stage locations can snowflake load data

- AWS S3  
**(Correct)**
- AZURE BLOB STORAGE  
**(Correct)**
- GOOGLE CLOUD STORAGE  
**(Correct)**

- INTERNAL STAGE  
**(Correct)**
- FILE SERVER
- IBM CLOUD STORAGE

### Explanation

#### External Stages

Loading data from any of the following cloud storage services is supported regardless of the [cloud platform](#) that hosts your Snowflake account:

Amazon S3  
Google Cloud Storage  
Microsoft Azure

Upload (i.e. *stage*) files to your cloud storage account using the tools provided by the cloud storage service.

A named external stage is a database object created in a schema. This object stores the URL to files in cloud storage, the settings used to access the cloud storage account, and convenience settings such as the options that describe the format of staged files. Create stages using the [CREATE STAGE](#) command.

#### Internal Stages

Snowflake maintains the following stage types in your account:

##### User

A user stage is allocated to each user for storing files. This stage type is designed to store files that are staged and managed by a single user but can be loaded into multiple tables. User stages cannot be altered or dropped.

##### Table

A table stage is available for each table created in Snowflake. This stage type is designed to store files that are staged and managed by one or more users but only loaded into a single table. Table stages cannot be altered or dropped.

Note that a table stage is not a separate database object; rather, it is an implicit stage tied to the table itself. A table stage has no grantable privileges of its own. To stage files to a table

stage, list the files, query them on the stage, or drop them, you must be the table owner (have the role with the OWNERSHIP privilege on the table).

## Named

A named internal stage is a database object created in a schema. This stage type can store files that are staged and managed by one or more users and loaded into one or more tables. Because named stages are database objects, the ability to create, modify, use, or drop them can be controlled using security access control privileges. Create stages using the [CREATE STAGE](#) command.

Upload files to any of the internal stage types from your local file system using the [PUT](#) command.

Question 40: Skipped

ACCOUNTADMIN ROLE encapsulates the SYSADMIN and SECURITYADMIN system-defined roles and ACCOUNTADMIN sits on top of the role hierarchy

- TRUE  
(Correct)
- FALSE

## Explanation

**ACCOUNTADMIN**

(aka Account Administrator)

Role that encapsulates the SYSADMIN and SECURITYADMIN system-defined roles. It is the top-level role in the system and should be granted only to a limited/controlled number of users in your account.

Question 41: Skipped

Select the two true statements about multi cluster warehouse

- These type of warehouses cannot be resumed or suspended
- Multi cluster warehouses can be created on snowflake hosted on AWS only
- As query demand decreases, clusters are removed from this warehouse  
(Correct)
- When query workload increases, clusters are automatically added  
(Correct)

## Explanation

In Auto-scale mode, a multi-cluster warehouse eliminates the need for resizing the warehouse or starting and stopping additional warehouses to handle fluctuating workloads. **Snowflake**  
**automatically starts and stops additional clusters as needed.**

In Maximized mode, you can control the capacity of the warehouse by increasing or decreasing the number of clusters as needed.

Question 42: **Skipped**

If you want to clone a table, your role must have which access on the Table

- SELECT  
**(Correct)**
- USAGE
- READ
- CLONE

### Explanation

To create a clone, your current role must have the following privilege(s) on the source object:

**Tables - SELECT**

**Pipes, Streams, Tasks - OWNERSHIP**

**Other objects - USAGE**

In addition, to clone a schema or an object within a schema, your current role must have required privileges on the container object(s) for both the source and the clone.

Question 43: **Skipped**

You have three worksheets open in Snowflake Web UI. You can have three different roles in these worksheets.

- TRUE  
**(Correct)**
- FALSE

### Explanation

Please do not believe me:) Try it out

Log on to snowflake, open three worksheets and try to use a different role. Run the below in worksheet and see if it works

Worksheet 1

```
USE ROLE SYSADMIN
```

## Worksheet 2

```
USE ROLE ACCOUNTADMIN
```

## Worksheet 3

```
USE ROLE PUBLIC
```

Question 44: Skipped

When will you consider creating a clustering key on a table. Please select all that applies

- The table houses multi-terabyte data  
**(Correct)**
- The table has variant columns
- The query performance on the table has degraded over time  
**(Correct)**

### Explanation

#### No need to mug this up

1. Clustering will be effective only when you have enough partitions in the table and the table is more than 1 TB

2. As you insert/update rows in the table, the partitions are moved around and hence it no longer remains clustered. With Auto clustering, Snowflake automatically reclusters the table based on a proprietary algorithm

Question 45: Skipped

Which of the below objects will contribute to your account level storage usage?

- Secure Views
- Internal Stage  
**(Correct)**
- External stage  
**(Correct)**
- Database and all tables inside it  
**(Correct)**
-

## Functions

### Explanation

SECURE VIEWS and FUNCTION definitions are stored in metadata store. They do not store any data physically.

Update to the explanation

---

One of my students asked a valid question that external stage is a pointer to the cloud provider's storage and hence from a snowflake perspective there is no cost. This is absolutely true. But when this question is presented to you and it is not explicitly asking 'snowflake account', you should tick external stage also. If the question specifically asks about 'snowflake account', then you should not tick this option. This is kind of a trick question so please pay attention to the question text.

### More on external stage

#### External stage

References data files stored in **a location outside of Snowflake**. Currently, the following cloud storage services are supported:

1. Amazon S3 buckets
2. Google Cloud Storage buckets
3. Microsoft Azure containers

The storage location can be either private/protected or public.

Question 46: Skipped

A consumer of shared database can add as many tables and views they want in the shared database

- FALSE  
**(Correct)**
- TRUE

## Explanation

Remember, we discussed earlier that SHARES are read-only and cannot be changed by consumer

Question 47: Skipped

To have multi cluster data virtual warehouse, what should be the minimum snowflake subscription

- ENTERPRISE  
**(Correct)**
- BUSINESS CRITICAL
- STANDARD

Question 48: Skipped

Once a new release has been deployed, Snowflake does not move all accounts to the release at the same time. It follows a staged release approach.

- TRUE  
**(Correct)**
- FALSE

## Explanation

### Staged Release

Once a new release has been deployed, Snowflake does not move all accounts to the release at the same time. Accounts are moved to the release using a three stage approach over two days. Accounts are moved to the new release in the following order, based on their [Snowflake Edition](#):

#### Day 1

Stage 1 (*early access*) for designated Enterprise accounts.

#### Day 1 or 2

Stage 2 (*regular access*) for all Standard Edition accounts.

#### Day 2

Stage 3 (*final*) for all Enterprise Edition and VPS accounts.

The minimum amount of elapsed time between the early access and final stages is 24 hours. This staged approach enables Snowflake to monitor activity as accounts are moved and respond to any issues that may occur.

Question 49: Skipped

When you are loading data through the COPY command, the table and the file being loaded must have the same order of columns

- FALSE  
**(Correct)**
- TRUE

### Explanation

Note that the actual field/column order in the data files can be different from the column order in the target table. It is only important that the SELECT list maps fields/columns in the data files to the **corresponding** columns in the table.

Question 50: Skipped

A reader account user will not be able to do the following

- LOAD DATA  
**(Correct)**
- INSERT  
**(Correct)**
- SELECT
- UPDATE  
**(Correct)**

### Explanation

A reader account is intended primarily for querying data shared by the provider of the account. Adding new data to the account and/or updating shared data in the account is not supported. As such, the following DML and DDL commands are not allowed:

INSERT

UPDATE

DELETE

MERGE

COPY INTO <table>

CREATE PIPE

CREATE SHARE

All other operations are allowed.

Question 51: Skipped

Which two modes are available in Multi cluster warehouse

- Auto-scale  
**(Correct)**
- Peak Scale
- Flexible
- Maximized  
**(Correct)**

### Explanation

Maximized vs. Auto-scale

You can choose to run a multi-cluster warehouse in either of the following modes:

#### Maximized

This mode is enabled by specifying the *same* value for both maximum and minimum clusters (note that the specified value must be larger than 1). In this mode, when the warehouse is started, Snowflake starts all the clusters so that maximum resources are available while the warehouse is running.

This mode is effective for statically controlling the available resources (i.e. servers), particularly if you have large numbers of concurrent user sessions and/or queries and the numbers do not fluctuate significantly.

#### Auto-scale

This mode is enabled by specifying *different* values for maximum and minimum clusters. In this mode, Snowflake starts and stops clusters as needed to dynamically manage the load on the warehouse:

As the number of concurrent user sessions and/or queries for the warehouse increases, and queries start to queue due to insufficient resources, Snowflake automatically starts additional clusters, up to the maximum number defined for the warehouse.

Similarly, as the load on the warehouse decreases, Snowflake automatically shuts down clusters to reduce the number of running servers and, correspondingly, the number of credits used by the warehouse.

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, SCALING\_POLICY, that determines the scaling policy to use when automatically starting or shutting down additional clusters. For more information, see [Setting the Scaling Policy for a Multi-cluster Warehouse](#) (in this topic).

Question 52: Skipped

Which of the below operations are allowed on an inbound share data?

- MERGE
- CREATE TABLE
- ALTER TABLE
- SELECT WITH GROUP BY  
**(Correct)**
- SELECT WITH JOIN  
**(Correct)**
- INSERT INTO

### Explanation

This is a trick question:) remember a share is read only, so you can only select data from a share

### Important

All database objects shared between accounts are ***read-only*** (i.e. the objects cannot be modified or deleted, including adding or modifying table data).

Question 53: Skipped

Select the most accurate statement regarding Snowflake's transaction support

- ACID Compliant  
**(Correct)**
- Requires special user configuration
- BASE compliant
- Works like a NoSQL database

Question 54: Skipped

Select the table type that is automatically deleted after the session is closed and hence it has no fail-safe or time travel option post closure of the session

- TRANSIENT

- TEMPORARY  
**(Correct)**

- PERMANENT
- EXTERNAL

### Explanation

Please create a temporary table and try it out

```
CREATE TEMPORARY TABLE T1(COL1 VARCHAR);
INSERT INTO T1 VALUES('TESTING TIME TRAVEL FOR T1');
```

Close the worksheet and open another worksheet and run the below query, are you able to time travel?

```
SELECT * FROM T1 AT(OFFSET => -60*1);
```

Question 55: Skipped

What does snowflake recommend for making a variant column accessible in a BI Tool?

- Create a View  
**(Correct)**
- Create a separate table
- Use Parse\_Json
- Convert into external table

Question 56: Skipped

When you load data into snowflake, snowflake reorganizes the data into its internal optimized compressed columnar format

- TRUE  
**(Correct)**
- FALSE

## Explanation

When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed, columnar format. Snowflake stores this optimized data in cloud storage.

Snowflake manages all aspects of how this data is stored — the organization, file size, structure, compression, metadata, statistics, and other aspects of data storage are handled by Snowflake. The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake.

Question 57: Skipped

When you clone a table, it does not include the load history of the source table, and hence the data files that were already loaded in source table can be reloaded in the clone table

- TRUE  
(Correct)
- FALSE

Question 58: Skipped

A snowflake account hosted on AWS cannot load data from a file that is staged in GCP or Azure

- FALSE  
(Correct)
- TRUE

Question 59: Skipped

If you clone a database, the internal snowflake stages in that database are also cloned

- FALSE  
(Correct)
- TRUE

## Explanation

### Cloning and Stages

The following rules apply to cloning stages or objects that contain stages (i.e. databases and schemas):

Individual external named stages can be cloned; internal named stages cannot be cloned.

When cloning a database or schema:

External named stages that were present in the source when the cloning operation started are cloned.

Tables are cloned, which means their internal stages are also cloned.

Internal named stages are **not** cloned.

Regardless of how a stage was cloned, the clone does **not** include any of the files from the source. i.e. all cloned stages are empty.

Question 60: Skipped

Which command is used to download data from Snowflake stage

- PUT
- GET  
**(Correct)**
- COPY
- DUMP

### Explanation

**GET**

Downloads data files from one of the following Snowflake stages to a local directory/folder on a client machine:

Named internal stage.

Internal stage for a specified table.

Internal stage for the current user.

Typically, this command is executed after using the `COPY INTO <location>` command to unload data from a table into a Snowflake stage.

Question 61: Skipped

AWS private link provides a secure connection from the customer's on premise data center to Snowflake

- True
- False  
**(Correct)**

### Explanation

<https://docs.snowflake.com/en/user-guide/admin-security-privatelink.html#what-is-aws-privatelink>

[AWS PrivateLink](#) is an AWS service for creating private VPC endpoints that allow direct, secure connectivity between AWS VPCs without traversing the public Internet.

In addition, if you have an on-premises environment (e.g. a non-hosted data center), you can choose to use [AWS Direct Connect](#), in conjunction with AWS PrivateLink, to connect all your virtual and physical environments in a single, private network.

Question 62: [Skipped](#)

Which of the below is an example of "**Increasing max number of clusters**"?

- Multi-Cluster Warehouse  
**(Correct)**
- Virtual Warehouse
- Dedicated Warehouse

### Explanation

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#what-is-a-multi-cluster-warehouse>

With multi-cluster warehouses, Snowflake supports allocating, either statically or dynamically, additional warehouses to make a larger pool of compute resources available. A multi-cluster warehouse is defined by specifying the following properties:

Maximum number of warehouses, greater than 1 (up to 10).

Minimum number of warehouses, equal to or less than the maximum (up to 10).

Question 63: [Skipped](#)

Reader accounts can extract data from shared data objects to use outside Snowflake?

- True
- False  
**(Correct)**

### Explanation

Reader accounts (formerly known as “read-only accounts”) provide a quick, easy, and cost-effective way to share data without requiring the consumer to become a Snowflake customer. Users in a reader account can query data that has been shared with them, but cannot perform any of the DML tasks that are allowed in a full account (data loading, insert, update, etc.)

Question 64: [Skipped](#)

Zero copy cloning requires additional licensing since it is considered as an additional feature in snowflake

-

True

- 

False

(Correct)

Question 65: Skipped

Some queries are answered by metadata cache and do not require active warehouse

- 

True

(Correct)

- 

False

### Explanation

Metadata caching is **maintained in Global Service Layer**. Metadata about Tables and micro-partitions are collected and managed by snowflake automatically. Snowflake does not use compute(and hence no warehouse) to provide Range values like MIN,MAX, Number of distinct values ,NULL count and ROW count and clustering information

Question 66: Skipped

For how many days Query history in Snowflake web Ui is kept?

- 

14 days

(Correct)

- 

30 days

- 

365 days

### Explanation

The **History** page allows you to view and drill into the details of all queries executed in the last 14 days

<https://docs.snowflake.com/en/user-guide/ui-history.html#using-the-history-page-to-monitor-queries>

Question 67: Skipped

Storage cost is incurred for both internal and external stage.

- 

True

(Correct)

- 

False

Question 68: Skipped

Fail-safe period can be disabled at account level

-

True

- 

False

(Correct)

### Explanation

Fail-safe provides a (non-configurable) 7-day period during which historical data may be recoverable by Snowflake. This period starts immediately after the Time Travel retention period ends. Note, however, that a long-running Time Travel query will delay moving any data and objects (tables, schemas, and databases) in the account into Fail-safe, until the query completes.

Question 69: Skipped

Reader account does not cost storage cost

- 

TRUE

(Correct)

- 

FALSE

### Explanation

Reader account does not COPY the data, in fact reader account does not have to be snowflake users also. A reader account can only consume data from the provider account that created it. A reader account is intended primarily for querying data shared by the provider of the account. Adding new data to the account and/or updating shared data in the account is not supported

Question 70: Skipped

Data Unloading compression applies to JSON and CSV only

- 

True

- 

False

(Correct)

### Explanation

It applies to parquet output also

Question 71: Skipped

A stage be created without specifying file format

- 

TRUE

(Correct)

- 

FALSE

### Explanation

If you do not specify the file format, it will create with CSV as default. Try this out to check it

```
CREATE STAGE TEST_STAGE;  
DESCRIBE STAGE TEST_STAGE;
```

Question 72: Skipped

Account storage consumption can be checked via

- Snowflake UI  
**(Correct)**
- ACCOUNT USAGE HISTORY
- ACCOUNT USAGE METRICS

### Explanation

There is nothing like ACCOUNT USAGE METRICS and ACCOUNT USAGE HISTORY

Question 73: Skipped

Micro-partition metadata is managed by which service layer

- Global Service Layer  
**(Correct)**
- Compute Layer
- Storage Layer
- Metadata layer

### Explanation

Metadata caching is maintained in Global Service Layer

Question 74: Skipped

4XL warehouse may have more lag time for provisioning compare to small warehouse

- True  
**(Correct)**
- False

Question 75: Skipped

A third party tool that supports JDBC or ODBC but has no snowflake specific driver will be unable to connect with snowflake

- True  
**(Correct)**

- False

Question 76: **Skipped**

Query IDs are unique across all snowflake deployments and can be used in communication with snowflake support to help troubleshoot issues

- True  
**(Correct)**

- False

Question 77: **Skipped**

Snowflake is designed for

- Concurrent workloads  
**(Correct)**
- OLAP  
**(Correct)**
- OLTP
- ON PREMISE

Question 78: **Skipped**

Which of the below interfaces can be used to create and/or manage virtual warehouses?

- Snowflake UI  
**(Correct)**
- SQL commands  
**(Correct)**
- Data ingestion tools
- All of the above

Question 79: **Skipped**

Credit consumption by the compute layer is based on

- Number of user

- Amount of data processed  
**(Correct)**

- Number of clusters

- Size of warehouse  
**(Correct)**

Question 80: Skipped

Which of the below are benefits of Snowflake's separation of compute and storage

- Multiple warehouses can access stored data without contention  
**(Correct)**

- Storage expands without requirement to add more compute  
**(Correct)**

- Compute can be scaled up and down without adding more storage  
**(Correct)**

- Growth of storage and compute are tightly coupled together

Question 81: Skipped

What happens when pipe is recreated using CREATE OR REPLACE pipe command

- Pipe load history is reset to empty  
**(Correct)**
- Refresh parameter set to true
- Previous loaded files will be ignored

### Explanation

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-manage.html#load-history>

[Continue](#)[Retake test](#)

Question 1: Skipped

**In a multi cluster warehouse, for each cluster there is a way we can specify different warehouse sizes**

- FALSE  
**(Correct)**
- TRUE

### Explanation

In a multi cluster warehouse, all clusters will have one size only. Refer Link -

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 2: Skipped

In auto scale mode, if `new_min_clusters < running clusters`, then excess clusters shut down when they finish executing statements and the scaling policy conditions are met

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 3: Skipped

**A warehouse was started and it was up and running for 40 seconds. For how many seconds will the warehouse be billed for?**

- 40 seconds
- No charge would be applied as it was up for less than 1 minute
- 60 seconds  
**(Correct)**

### Explanation

Snowflake utilizes per-second billing (with a 60-second minimum each time the warehouse starts) Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 4: Skipped

Which scaling policy, starts additional cluster only if the system estimates there is enough query load to keep the cluster busy for at least 6 minutes.

- Legacy
- Standard
- None

- Economy
- (Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 5: Skipped

**Scaling in is a manual process (e.g. Max Clusters to min clusters)**

- TRUE
- FALSE
- (Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 6: Skipped

**How many credits per hour is consumed by a 3X large warehouse**

- 64
- 32
- 16
- 128

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 7: Skipped

**Which of the following are options when creating a Virtual Warehouse?**

- Auto Resume
- Cache Size
- Storage Size
- Auto Suspend

### Explanation

Refer Link- <https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Question 8: Skipped

Which scaling policy, starts additional cluster immediately when either a query is queued or the system detects that there is one more query than the currently-running clusters can execute.

- Legacy
- Standard  
**(Correct)**
- Economy
- None

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 9: Skipped

**When the warehouse is resized (scaled up or out), then the additional servers that get added will be available for use by any queries that are queued or newly submitted**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 10: Skipped

**Scaling out of a warehouse can be done only in Enterprise edition and higher**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 11: Skipped

**If a server runs for 61 seconds, shuts down, and then restarts and runs for 30 seconds, it is billed for how many seconds?**

- 90
- 120
-

121

(Correct)

- 91

### Explanation

The minimum billing charge for provisioning a server is 1 minute (i.e. 60 seconds). Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 12: Skipped

### Snowflake Credit Charges for a virtual warehouse is calculated based on:

- The length of time each server in each cluster runs.  
**(Correct)**
- Number of queries getting executed
- Number of Users running the queries concurrently
- The number of servers per cluster (determined by warehouse size).  
**(Correct)**
- The number of clusters (if using multi-cluster warehouses).  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 13: Skipped

### In maximized mode, what happens if the number of clusters are decreased? (e.g. previously min & max was 4, but now min & max cluster size is changed to 2)

- Specified number of clusters shut down when they finish executing statements and the auto-suspend period elapses  
**(Correct)**
- The number of clusters in maximized mode cannot be changed
- Specified number of clusters shut down immediately, as soon as they finish executing statements

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 14: Skipped

### The columns in the query history include the QueryID, the SQL Text, the Warehouse name, the Warehouse Size, the Session ID and others. Which column is a good indicator of whether a Warehouse was used (and Compute costs incurred) by a query?

-

- Session ID
- 
- Warehouse Name
- 
- SQL Text
- 
- QueryID
- 
- Warehouse Size  
**(Correct)**
- 

### Explanation

The queries against which we have the warehouse size indicated are the one's that uses the compute power of a warehouse

Question 15: Skipped

**Scaling a warehouse DOWN will decrease the number of servers. (e.g. Medium to Small)**

- 
- TRUE  
**(Correct)**
- 
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 16: Skipped

**Which scaling policy, prevents/minimizes queuing by favoring starting additional clusters over conserving credits?**

- 
- None
- 
- Standard  
**(Correct)**
- 
- Economy
- 
- Legacy

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 17: Skipped

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Auto-scale mode for 3 hours, in which Cluster 1 runs**

**continuously, Cluster 2 runs continuously for the entire 2nd hour and 30 minutes in the 3rd hour, Cluster 3 runs for 30 minutes in the 3rd hour**

- 20  
**(Correct)**
- 24
- 16
- 18

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 18: Skipped

**In snowflake, while processing a query, the number of rows in the tables being queried has more impact than the overall size of the tables**

- FALSE  
**(Correct)**
- TRUE

### Explanation

The overall size of the tables being queried has more impact than the number of rows. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 19: Skipped

**What all resources are provided by Virtual Warehouse?**

- Metadata Storage
- CPU  
**(Correct)**
- Temporary Storage  
**(Correct)**
- Permanent Storage
- Memory  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses.html>

Question 20: Skipped

**Whenever the warehouse size is reduced or the number of clusters are decreased, the servers are always removed in reverse order of when they were added (LIFO - Last In First Out)**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Here is one resource blog from snowflake which gives this information -  
<https://resources.snowflake.com/snowflake-product/virtual-warehouse-billing-made-smarter>

Question 21: **Skipped**

**Multi-cluster warehouses are beneficial for improving the performance of slow-running queries or data loading**

- FALSE  
**(Correct)**
- TRUE

### Explanation

Multi-cluster warehouses are best utilized for scaling resources to improve concurrency for users/queries. They are not as beneficial for improving the performance of slow-running queries or data loading. For these types of operations, resizing the warehouse provides more benefits.

Question 22: **Skipped**

**Snowflake does not begin executing SQL statements submitted to a warehouse until all of the servers for the warehouse are successfully provisioned, unless any of the servers fail to provision**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Question 23: **Skipped**

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Maximized mode for 2 hours**

- 24  
**(Correct)**
-

- 12  
•  36  
•  18

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 24: Skipped

**When a warehouse is suspended, the idle servers in that warehouse are immediately shut down, but any servers that are executing queries are allowed to run until the queries are completed. Such servers that are waiting to shut down are considered to be in which mode?**

- Suspending
- Quiet
- Quiesce  
**(Correct)**
- Shutting Down

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Question 25: Skipped

**Multi-cluster warehouses is only available in Enterprise Edition and above**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 26: Skipped

**Warehouses can also be resized at any time, even while running, to accommodate the need for more or less compute resources**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 27: Skipped

**Scaling a warehouse OUT will increase the number of clusters. (e.g. Min to Max)**

- FALSE

- TRUE  
**(Correct)**

#### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 28: Skipped

Out of below four options, which options require a running warehouse when retrieving data from them?

- Metadata Cache
- Warehouse Cache  
**(Correct)**
- Results Cache
- Local Disk Cache  
**(Correct)**

#### Explanation

Metadata Cache and Result cache fetch results from services layer itself, and hence they don't need a running warehouse

Question 29: Skipped

**How many servers are there in 2X large warehouse?**

- 32  
**(Correct)**
- 64
- 128
- 16

#### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 30: Skipped

**A warehouse was started and it was up and running for 62 seconds. For how many seconds will the warehouse be billed for?**

- 120 seconds

- 60 seconds

- 62 seconds  
**(Correct)**

- 70 seconds

### Explanation

Snowflake utilizes per-second billing (with a 60-second minimum each time the warehouse starts) Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 31: Skipped

**There is no benefit to stopping a warehouse before the first 60-second period is over because the credits have already been billed for that period**

- TRUE  
**(Correct)**
- FALSE

### Explanation

The minimum billing charge for provisioning a server is 1 minute (i.e. 60 seconds). Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 32: Skipped

**Snowflake compute costs depend on which of the following?**

- The total number of warehouses in the account
- The number of rows returned in queries.
- The sizes of running warehouses.  
**(Correct)**
- The amount of time warehouses have run.  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 33: Skipped

A Snowflake session can have multiple current warehouse at a time

- TRUE

- FALSE  
**(Correct)**

### Explanation

A Snowflake session can have only one current warehouse at a time. The current warehouse for a session can be specified or changed at any time through the USE WAREHOUSE command.

Question 34: Skipped

### Default Size of the warehouse created using CREATE WAREHOUSE command is

- X-SMALL  
**(Correct)**
- MEDIUM
- X-LARGE
- SMALL

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 35: Skipped

### Scaling a warehouse UP will decrease the number of clusters. (e.g. Small to Medium)

- TRUE
- FALSE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 36: Skipped

### Which scaling policy, shut down the cluster after 2 to 3 consecutive successful checks (performed at 1 minute intervals), which determine whether the load on the least-loaded cluster could be redistributed to the other clusters without spinning up the cluster again

- Legacy
- None
- Economy
- Standard

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 37: Skipped

**In maximized multi-cluster warehouse same value is specified for both maximum & minimum clusters (except for the value 1)**

- FALSE
- TRUE

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 38: Skipped

**Scaling up of a warehouse can be done only in Enterprise edition and higher**

- FALSE
- TRUE

### Explanation

Scaling up (increasing size of warehouse) can be done in all editions of snowflake, however Scaling out (increasing number of clusters) can be done only in Enterprise edition and higher

Question 39: Skipped

In auto scale mode, if new\_max\_clusters < running clusters, then excess clusters shut down immediately

- FALSE
- TRUE

### Explanation

If new\_max\_clusters < running\_clusters, excess clusters shut down when they finish executing statements and the scaling policy conditions are met. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 40: Skipped

**Which of the following are true of Multi-Cluster warehouses? Select all that apply.**

- Multi-cluster warehouses will never auto-suspend
- Scales In when query activity slows
- 

(Correct)

Adds clusters automatically based on query activity  
**(Correct)**

- Sizes each cluster optimally based on the queries

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 41: Skipped

**Which scaling policy, conserves credits by favoring keeping running clusters fully-loaded rather than starting additional clusters, which may result in queries being queued and taking longer to complete?**

- None
- Legacy
- Standard
- Economy  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 42: Skipped

**If a 3X-Large multi-cluster warehouse runs 1 cluster for one full hour and then runs 2 clusters for the next full hour, the total number of credits billed would be**

- 128
- 64
- 256
- 192  
**(Correct)**

### Explanation

For a multi-cluster warehouse, the number of credits billed is calculated based on the number of servers per cluster and the number of clusters that run within the time period. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 43: Skipped

**The scaling policy for a multi-cluster warehouse only applies (make sense) if it is running in Auto-scale mode and not in Maximized Mode**

- FALSE

- TRUE
- (Correct)

### Explanation

The scaling policy for a multi-cluster warehouse only applies if it is running in Auto-scale mode. In Maximized mode, all clusters run concurrently so there is no need to start or shut down individual clusters. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 44: Skipped

**If a multi-cluster warehouse is resized, the new size applies to all the clusters for the warehouse, including clusters that are currently running and any clusters that are started after the warehouse is resized**

- FALSE
  - TRUE
- (Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 45: Skipped

**Scaling a warehouse IN will decrease the number of clusters. (e.g. Max to Min)**

- TRUE
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 46: Skipped

**In a multi cluster warehouse, auto-suspending a running warehouse due to inactivity, does not apply to individual clusters, but rather the entire warehouse.**

- TRUE
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 47: Skipped

**Scaling out is an automated process (e.g. Min clusters to max clusters)**

- TRUE
- (Correct)

- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 48: Skipped

**Maximum number of clusters that can be added in multi cluster warehouse is?**

- 20
- 15
- 5
- 10

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 49: Skipped

**A warehouse can be started at any time, excluding the initial creation time**

- FALSE
- TRUE

### Explanation

A warehouse can be started at any time, including on initial creation. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Question 50: Skipped

**Which are the 2 modes in which a multi cluster warehouse can run?**

- Maximized
- Auto-Scale
- Balanced
- Dynamic
-

Minimized

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 51: Skipped

**Multi-cluster warehouses are best utilized for scaling resources to improve concurrency for users/queries**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 52: Skipped

**Scaling a warehouse OUT will increase the number of servers. (e.g. Min to Max)**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 53: Skipped

**In a multi cluster warehouse, auto-suspend and auto-resume apply only to the entire warehouse and not to the individual clusters in the warehouse**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 54: Skipped

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Auto Scale mode for 2 hours, where Cluster 1 runs continuously for 2 hours, Cluster 2 runs continuously for 2nd hour only, Cluster 3 runs for 30 minutes only during 2nd hour**

- 16
- 14  
**(Correct)**

- 24
- 12

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 55: Skipped

### What happens when a suspended warehouse is resized (e.g. resized from small to medium)?

- Additional Servers are added to the warehouse immediately even in suspended state
- A suspended warehouse cannot be resized
- Additional servers are provisioned only when the warehouse is next resumed  
**(Correct)**
- When size of a suspended warehouse is increased it starts the warehouse, and then suspends it back after provisioning additional servers

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Question 56: Skipped

In auto scale mode, if new\_max\_clusters > running clusters, then no changes happen until additional clusters are needed

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 57: Skipped

### How is the query load calculated, which is displayed on the "Warehouse Load over Time" chart in Snowflake Web UI

- The formula is not known to the users, and its internally computed by snowflake only
- (Number of queries running + Number of queries in queue) in the given interval divided by the number of seconds in the interval
- Number of queries running in the given interval divided by the number of seconds in the interval

- 

Execution time + Queued Time (in seconds) of all queries in an interval by the total time (in seconds) for the interval

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-load-monitoring.html>

Question 58: Skipped

**In auto scale multi-cluster warehouse same value is specified for both maximum & minimum clusters (except for the value 1)**

- 

TRUE

- 

FALSE

(Correct)

### Explanation

In auto scale mode different value for minimum & maximum clusters are specified. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 59: Skipped

**Warehouses can be started and stopped at any time**

- 

TRUE

(Correct)

- 

FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 60: Skipped

**Scaling down is a automated process. (e.g. X-Large to Large)**

- 

FALSE

(Correct)

- 

TRUE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 61: Skipped

**If any of the servers for the warehouse fail to provision during start-up, Snowflake attempts to repair the failed server(s). During the repair process, the warehouse starts processing SQL statements only when a minimum of how much % of the requested servers are successfully provisioned**

- 

30

-

50

(Correct)

- 25
- 75

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Question 62: Skipped

### Which of the following terms are associated with the Compute/Warehouse Layer?

- Query Design
- Query Planning
- Query Optimization
- Query Compilation
- Query Processing

(Correct)

### Explanation

Only query processing happens in the compute/warehouse layer, rest all options (except query design) that are given are performed in the services layer. Query design is performed by the user executing the query

Question 63: Skipped

### Snowflake utilizes per-second billing (with a 60-second minimum each time the warehouse starts)

- FALSE
- TRUE

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 64: Skipped

### Default Size of the warehouse created using Snowflake Web UI is

- MEDIUM
- X-SMALL

- X-LARGE  
**(Correct)**

- SMALL

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 65: Skipped

**Which scaling policy, shut down the cluster after 5 to 6 consecutive successful checks (performed at 1 minute intervals), which determine whether the load on the least-loaded cluster could be redistributed to the other clusters without spinning up the cluster again**

- None
- Economy  
**(Correct)**

- Legacy
- Standard

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 66: Skipped

**When the warehouse is resized (scaled up or out), then the additional servers that get added will also improve the performance of already running queries**

- FALSE  
**(Correct)**

- TRUE

### Explanation

The additional servers do not impact any queries that are already running, but they are available for use by any queries that are queued or newly submitted. Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 67: Skipped

**In maximized mode, what happens if the number of clusters are increased? (e.g. previously min & max was 2, but now min & max cluster size is changed to 4)**

- Specified number of clusters would start after current queries that are executed are completed
- There will not be any immediate impact, and the additional clusters will start only when the query workload demands

- The number of clusters in maximized mode cannot be changed
- Specified number of clusters (in this case 4) would start immediately  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 68: Skipped

### Each server in a warehouse cluster has:

- A position in the warehouse that is maintained, even when the warehouse is suspended or resized  
**(Correct)**
- Track of each query that was executed on that server
- An internal timer that tracks when each server was started  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Question 69: Skipped

### How many credits will be consumed if a Small-size warehouse (2 servers per cluster) with 3 clusters runs in Maximized mode for 2 hours

- 36
- 18
- 24
- 12  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 70: Skipped

**A medium size warehouse runs in Auto-scale mode for 3 hours with a resize from Medium (4 servers per cluster) to Large (8 servers per cluster). Warehouse is resized from Medium to Large at 1:30 hours, Cluster 1 runs continuously, Cluster 2 runs continuously for the 2nd and 3rd hours, Cluster 3 runs for 15 minutes in the 3rd hour. How many total credits will be consumed**

- 40

- 30

- 34

(Correct)

- 32

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 71: Skipped

### Scaling up is a manual process. (e.g. Small to Medium)

- FALSE

- TRUE

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 72: Skipped

### What can be the maximum credits consumed by a medium size warehouse with 3 clusters in one hour?

- 12

(Correct)

- 6

- 8

- 24

### Explanation

The total number of servers for the warehouse is calculated by multiplying the warehouse size by the maximum number of clusters. This also indicates the maximum number of credits consumed by the warehouse per full hour of usage (i.e. if all clusters run during the hour).

Question 73: Skipped

In auto scale mode, if new\_min\_clusters > running clusters, then additional clusters immediately started to meet the minimum.

- TRUE

(Correct)

-

FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Question 74: Skipped

**True or false: A table in Snowflake can only be queried using the Virtual Warehouse used to load the data.**

- TRUE
- FALSE  
**(Correct)**

### Explanation

Any warehouse can be used to query data in any table as long as the user/role executing the query has access to the warehouse

Question 75: Skipped

**When configuring a Warehouse using a Snowflake edition that has Elastic Data**

**Warehousing enabled, what facets or components will you need to configure that are not needed in accounts where Elastic Data Warehousing is not enabled?**

- Auto-Resume
- Auto-Suspend
- Scaling Policy  
**(Correct)**
- Minimum and Maximum Servers
- Minimum and Maximum Clusters  
**(Correct)**

### Explanation

Minimum & Maximum servers cannot be specified in either single cluster or multi cluster warehouse. Auto Resume & Auto suspend is applicable for both single cluster and multi cluster warehouse. Only Scaling Policy and Min/Max cluster options are available in multi cluster warehouse which is not there in the single cluster warehouse

[Continue](#)

[Retake test](#)

Question 1: Skipped

**Data retention period of database is set to 2 days. Data retention period of schema inside that database is set to 4 days. Data retention of a table inside the schema is set to 7 days. Now if the database is dropped, the table that was present in that database can be restored back within how many days?**

- 
- Since the Database itself is dropped, the underlying objects cannot be restored
- 
- 7 days
- 
- 4 days
- 2 days

**(Correct)**

### Explanation

Currently, when a database is dropped, the data retention period for child schemas or tables, if explicitly set to be different from the retention of the database, is not honored. The child schemas or tables are retained for the same period of time as the database. Similarly, when a schema is dropped, the data retention period for child tables, if explicitly set to be different from the retention of the schema, is not honored. The child tables are retained for the same period of time as the schema. To honor the data retention period for these child objects (schemas or tables), drop them explicitly before you drop the database or schema.

Question 2: Skipped

**Restoring tables and schemas is only supported in the current schema or current database, even if a fully-qualified object name is specified.**

- 
- TRUE
- (Correct)**
- 
- FALSE

Question 3: Skipped

**After the retention period for an object has passed and the object has been purged, it is no longer displayed in the SHOW <object\_type> HISTORY output.**

- 
- TRUE
- (Correct)**
- 
- FALSE

Question 4: Skipped

**An individual task in the tree can have multiple predecessor task**

- 
- TRUE
-

FALSE  
**(Correct)**

### Explanation

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state. An individual task in the tree is limited to a single predecessor task; however, a task can have a maximum of 100 child tasks (i.e. other tasks that identify the task as a predecessor).

Question 5: Skipped

**Which of the following parameters are provided by Snowflake to be used with AT & BEFORE clause to pinpoint the exact historical data you wish to access**

- OFFSET (time difference in seconds from the present time)
- All options are correct  
**(Correct)**
- TIMESTAMP
- STATEMENT (identifier for statement, e.g. query ID)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 6: Skipped

**Check all true statements about Fail-safe:**

- Only a Snowflake employee can recover data from fail-safe storage  
**(Correct)**
- The data stored as part of fail-safe is part of storage costs charged to customers  
**(Correct)**
- Fail-safe is not available for tables that have Time Travel.
- Fail-safe is a reliable way to create Dev/Test/QA and other environments

Question 7: Skipped

**The data retention period for a database, schema, or table cannot be changed after they are created.**

- TRUE
- FALSE  
**(Correct)**

### Explanation

The data retention period for a database, schema, or table can be changed at any time.

Question 8: Skipped

**If you change the data retention period for a table, the new retention period impacts all data that is active, as well as any data currently in Time Travel**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 9: Skipped

**What is the standard time travel retention period which is enabled for all snowflake accounts?**

- 7 days
- 24 hours  
**(Correct)**
- 48 hours
- 14 days

### Explanation

The standard retention period is 1 day (24 hours) and is automatically enabled for all Snowflake accounts: For Snowflake Standard Edition, the retention period can be set to 0 (or unset back to the default of 1 day) at the account and object level (i.e. databases, schemas, and tables). For Snowflake Enterprise Edition (and higher): For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also true for temporary tables. For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days.

Question 10: Skipped

**When the owner role of a given task (i.e. the role with the OWNERSHIP privilege on the task) is deleted, what happens?**

- The roles granted to any task cannot be dropped unless all associated tasks are dropped
- The task is “re-possessed” by the role that dropped the owner role  
**(Correct)**
- The tasks is also dropped
-

The tasks is suspended immediately, and the task will not have any owner till explicitly assigned

### Explanation

When the owner role of a given task (i.e. the role with the OWNERSHIP privilege on the task) is deleted, the task is “re-possessed” by the role that dropped the owner role. This ensures that ownership moves to a role that is closer to the root of the role hierarchy. When a task is re-possessed, it is automatically paused, i.e., all executions currently in flight complete processing, but new executions will not be scheduled until the task is resumed explicitly by the new owner. The rationale for this is to prevent a user with access to a particular role from leaving behind tasks that suddenly execute with higher permissions when the role is removed. If the role that a running task is executing under is dropped while the task is running, the task completes processing under the dropped role.

Question 11: Skipped

**Changing the time travel retention period for your account or individual objects changes the value for all lower-level objects that do not have a retention period explicitly set**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Changing the retention period for your account or individual objects changes the value for all lower-level objects that do not have a retention period explicitly set. For example: If you change the retention period at the account level, all databases, schemas, and tables that do not have an explicit retention period automatically inherit the new retention period. If you change the retention period at the schema level, all tables in the schema that do not have an explicit retention period inherit the new retention period. Keep this in mind when changing the retention period for your account or any objects in your account because the change might have Time Travel consequences that you did not anticipate or intend. In particular, we do not recommend changing the retention period to 0 at the account level.

Question 12: Skipped

**True or false: You can point Snowflake at any S3 bucket to directly query the files in that bucket as long as the files are in Parquet or ORC format.**

- TRUE
- FALSE  
**(Correct)**

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 13: Skipped

**DATA\_RETENTION\_TIME\_IN\_DAYS paramete is defined as 1 day at account level, and the same parameter is set as 7 days at Database level. Now a table is created inside the given database without specifying any explicit**

**DATA\_RETENTION\_TIME\_IN\_DAYS** parameter. What would be the data retention period for table in this case?

- 1 day
- 0 day
- 7 days  
**(Correct)**

### Explanation

If a retention period is specified for a database or schema, the period is inherited by default for all objects created in the database/schema.

Question 14: Skipped

**Which parameter is used to set the time travel retention period?**

- DATA\_RETENTION\_TIME
- DATA\_RETENTION\_TIME\_IN\_DAYS  
**(Correct)**
- DATA\_RETENTION\_TIME\_IN\_HOURS
- DATA\_RETENTION\_TIME\_IN\_WEEKS

### Explanation

To specify the data retention period for Time Travel: The **DATA\_RETENTION\_TIME\_IN\_DAYS** object parameter can be used by users with the ACCOUNTADMIN role to set the default retention period for your account. The same parameter can be used to explicitly override the default when creating a database, schema, and individual table. The data retention period for a database, schema, or table can be changed at any time.

Question 15: Skipped

**What is the default table type when a CREATE TABLE is executed?**

- PERMANENT  
**(Correct)**
- TEMPORARY
- TRANSIENT
- MATERIALIZED

### Explanation

The default type for tables is permanent. To define a table as temporary or transient, you must explicitly specify the type during table creation: CREATE [ OR REPLACE ] [ TEMPORARY | TRANSIENT ] TABLE <name> ...

Question 16: Skipped

**If you have a table with a 10-day retention period and increase the period to 20 days, then the data that was moved to fail safe on 11th day, will be moved back to Time Travel, and will have more 9 days left for time travel retention end**

- TRUE
- FALSE

(Correct)

### Explanation

The data once moved to fail safe cannot move back to time travel retention period.

Question 17: Skipped

**You can view the amount of storage used by failsafe under the Account->Billing & Usage section**

- FALSE
- TRUE

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

Question 18: Skipped

**When a table data changes (for e.g. few records are updated and few are deleted) and the table is within the time travel retention period, which of the following statement is true?**

- Snowflake takes a full copy of the table every time any change is made to the table
- Snowflake maintains only the information pertaining to these changed records (i.e. updated & deleted ones)

(Correct)

### Explanation

Snowflake minimizes the amount of storage required for historical data by maintaining only the information required to restore the individual table rows that were updated or deleted. As a result, storage usage is calculated as a percentage of the table that changed. Full copies of tables are only maintained when tables are dropped or truncated.

Question 19: Skipped

**If you have a table with a 10-day retention period and you decrease the period to 1-day, then what happens to the data that is already in time travel for 4 days**

- The data would move to fail safe after 1 day

- The data would be removed and will not even be available in fail safe
- The data would continue to be in time travel, and new time travel retention of 1 day will be applied to only active data which is not already in time travel window
- The data would be moved to fail safe  
**(Correct)**

### Explanation

Reduces the amount of time data is retained in Time Travel: For active data modified after the retention period is reduced, the new shorter period applies. For data that is currently in Time Travel: If the data is still within the new shorter period, it remains in Time Travel. If the data is outside the new period, it moves into Fail-safe. For example, if you have a table with a 10-day retention period and you decrease the period to 1-day, data from days 2 to 10 will be moved into Fail-safe, leaving only the data from day 1 accessible through Time Travel.

However, the process of moving the data from Time Travel into Fail-safe is performed by a background process, so the change is not immediately visible. Snowflake guarantees that the data will be moved, but does not specify when the process will complete; until the background process completes, the data is still accessible through Time Travel

Question 20: Skipped

**The FLATTEN command will parse nested objects into separate rows. One version of the FLATTEN command uses a join and the other uses an object keyword. Select the two words that represent the options used with the FLATTEN command**

- LATERAL  
**(Correct)**
- SCHEMATA
- TABLE  
**(Correct)**
- SIDEWAYS

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 21: Skipped

**A stream object records data manipulation language (DML) changes made to tables, including inserts and updates but cannot record deletes made to tables**

- TRUE
- FALSE  
**(Correct)**

## Explanation

A stream object records data manipulation language (DML) changes made to tables, including inserts, updates, and deletes, as well as metadata about each change, so that actions can be taken using the changed data. This process is referred to as change data capture (CDC). An individual table stream tracks the changes made to rows in a source table. A table stream (also referred to as simply a “stream”) makes a “change table” available of what changed, at the row level, between two transactional points of time in a table. This allows querying and consuming a sequence of change records in a transactional fashion.

Question 22: Skipped

**Storage fees are incurred for maintaining historical data during both the Time Travel and Fail-safe periods**

- TRUE  
**(Correct)**
- FALSE

Question 23: Skipped

**Which two SQL Extension clause are provided by Snowflake to support time travel?**

- BEFORE  
**(Correct)**
- PRIOR
- DURING
- AT  
**(Correct)**

## Explanation

To support Time Travel, the following SQL extensions have been implemented: AT | BEFORE clause which can be specified in SELECT statements and CREATE ... CLONE commands (immediately after the object name). The clause uses one of the following parameters to pinpoint the exact historical data you wish to access: TIMESTAMP OFFSET (time difference in seconds from the present time) STATEMENT (identifier for statement, e.g. query ID) UNDROP command for tables, schemas, and databases

Question 24: Skipped

**A stream allows querying and consuming a set of changes to a table, at the row level, between two transactional points of time**

- TRUE  
**(Correct)**
- FALSE

## Explanation

A stream object records the delta of change data capture (CDC) information for a table (such as a staging table), including inserts and other data manipulation language (DML) changes. A stream allows querying and consuming a set of changes to a table, at the row level, between two transactional points of time.

Question 25: Skipped

**The time travel retention period can only be set at individual object level, and cannot be setup at account level**

- TRUE
- FALSE  
**(Correct)**

### Explanation

The retention period can be applied at Account level, and object level (i.e. databases, schemas, and tables).

Question 26: Skipped

**Which type of table stream is supported only on external tables?**

- Native
- Insert-Only  
**(Correct)**
- Append-Only
- Update-Only
- Standard

### Explanation

Insert-only Supported on external tables only. An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops). For example, in-between any two offsets, if File1 is removed from the cloud storage location referenced by the external table, and File2 is added, the stream returns records for the rows in File2 only. Unlike when tracking CDC data for standard tables, Snowflake cannot access the historical records for files in cloud storage. Overwritten files are essentially handled as new files: The old version of the file is removed from cloud storage, but the insert-only stream does not record the delete operation. The new version of the file is added to cloud storage, and the insert-only stream records the rows as inserts. The stream does not record the diff of the old and new file versions

Question 27: Skipped

**Which of the following Snowflake Object is used for a schedule execution of SQL statement?**

- Streams

- Tasks  
**(Correct)**

- Stored Procedure
- Snowpipe

### Explanation

User-defined tasks allow scheduled execution of SQL statements. Tasks run according to a specified execution configuration, using any combination of a set interval and/or a flexible schedule using a subset of familiar cron utility syntax. Tasks currently run in user-specified warehouses but will eventually run in a serverless environment using Snowflake-supplied compute resources.

Question 28: Skipped

### Which two table types don't have fail safe?

- TEMPORARY  
**(Correct)**
- TRANSIENT  
**(Correct)**
- PERMANENT

### Explanation

To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides two table types, temporary and transient, which do not incur the same fees as standard (i.e. permanent) tables: Transient tables can have a Time Travel retention period of either 0 or 1 day. Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon as the table is dropped or the session in which the table was created ends. Transient and temporary tables have no Fail-safe period. As a result, the maximum additional fees incurred for Time Travel and Fail-safe by these types of tables is limited to 1 day

Question 29: Skipped

### For which table type the retention period ends as soon as the the table is dropped or the session in which the table was created ends.

- TEMPORARY  
**(Correct)**
- TRANSIENT
- MATERIALIZED
-

## PERMANENT

### Explanation

To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides two table types, temporary and transient, which do not incur the same fees as standard (i.e. permanent) tables: Transient tables can have a Time Travel retention period of either 0 or 1 day. Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon as the table is dropped or the session in which the table was created ends. Transient and temporary tables have no Fail-safe period. As a result, the maximum additional fees incurred for Time Travel and Fail-safe by these types of tables is limited to 1 day

Question 30: Skipped

**A table named "MY\_TABLE" has a retention period of 7 days. The table is dropped, and a new table with the same name "MY\_TABLE" has been created. The new table that is created is also dropped. Now retrieving the original version (i.e. the 1st version) of the table is not possible. This is True or False?**

- TRUE
- FALSE

(Correct)

### Explanation

If an object has been dropped more than once, each version of the object is included as a separate row in the output, and hence any previous version of the object can be restored till the time it is in the retention period window

Question 31: Skipped

**Under what circumstances would Fail-Safe NOT be an effective method for data recovery?**

- To recover data that may have been lost or damaged due to extreme operational failures
- To protect data in the event of a system failure or other catastrophic event e.g. a hardware failure or security breach
- As a means for accessing historical data after the Time Travel retention period has ended

(Correct)

### Explanation

Fail Safe is something that is managed by Snowflake, and the end user cannot directly query data from Fail Safe.

In case for any table whose time travel period is over, and we need to still restore the data within 7 days of time travel getting over, then Snowflake Support team can help us get that data.

However, as per Snowflake "Fail-safe is **not** provided as a means for accessing historical data after the Time Travel retention period has ended. It is for use **only** by Snowflake to recover data that may have been lost or damaged due to extreme operational failures."

You can refer this link for the details - <https://docs.snowflake.com/en/user-guide/data-failsafe.html>

Question 32: Skipped

**What is the maximum number of child tasks, a given task can have?**

- Unlimited
- 1000
- 10000
- 10
- 100

(Correct)

#### Explanation

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state. An individual task in the tree is limited to a single predecessor task; however, a task can have a maximum of 100 child tasks (i.e. other tasks that identify the task as a predecessor).

Question 33: Skipped

**If the TIMESTAMP, OFFSET, or STATEMENT specified in the AT | BEFORE clause falls outside the data retention period for the table, the query fails and returns an error.**

- TRUE
- FALSE

(Correct)

#### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 34: Skipped

**Time Travel in Snowflake is useful for which of the following tasks**

- Restoring data in event of a system failure or other catastrophic event, e.g. a hardware failure or security breach.
- Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted
- Analyzing data usage/manipulation over specified periods of time

(Correct)

(Correct)

- 

Duplicating and backing up data from key points in the past  
**(Correct)**

### Explanation

Snowflake Time Travel enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It serves as a powerful tool for performing the following tasks: Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted. Duplicating and backing up data from key points in the past. Analyzing data usage/manipulation over specified periods of time. Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 35: Skipped

**What is the limit of number of tasks that can be included in a tree of tasks?**

- 

10000

- 

100

- 

Unlimited

- 

1000

**(Correct)**

### Explanation

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state. An individual task in the tree is limited to a single predecessor task; however, a task can have a maximum of 100 child tasks (i.e. other tasks that identify the task as a predecessor).

Question 36: Skipped

**What data structure types can be ingested into a VARIANT column in a Snowflake table? (Check all that apply)**

- 

HADOOP

- 

AVRO

**(Correct)**

- 

XML

**(Correct)**

- 

ORC

**(Correct)**

- 

JSON

**(Correct)**

- PARQUET  
**(Correct)**

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 37: Skipped

**An append-only table stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded**

- TRUE  
**(Correct)**
- FALSE

### Explanation

An append-only table stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded. For example, if 10 rows are inserted into a table and then 5 of those rows are deleted before the offset for an append-only stream is advanced, the stream records 10 rows. An append-only stream returns the appended rows only and therefore can be much more performant than a standard stream for extract, load, transform (ELT) and similar scenarios that depend exclusively on row inserts. For example, the source table can be truncated immediately after the rows in an append-only stream are consumed, and the record deletions do not contribute to the overhead the next time the stream is queried or consumed.

Question 38: Skipped

**After dropping an object, creating an object with the same name does not restore the object. Instead, it creates a new version of the object. The original, dropped version is still available and can be restored.**

- TRUE  
**(Correct)**
- FALSE

### Explanation

After dropping an object, creating an object with the same name does not restore the object. Instead, it creates a new version of the object. The original, dropped version is still available and can be restored. Restoring a dropped object restores the object in place (i.e. it does not create a new object).

Question 39: Skipped

**Which parameter is used to check if the stream contains change data for a table or not?**

- SYSTEM\$STREAM\_HAS\_DATA  
**(Correct)**
- SYSTEM\$STREAM\_HAS\_MODIFIED\_DATA
-

SYSTEM\$STREAM\_HAS\_NEW\_DATA

- SYSTEM\$STREAM\_HAS\_CHANGED\_DATA

### Explanation

Tasks may optionally use table streams to provide a convenient way to continuously process new or changed data. A task can transform new or changed rows that a stream surfaces. Each time a task is scheduled to run, it can verify whether a stream contains change data for a table (using SYSTEM\$STREAM\_HAS\_DATA) and either consume the change data or skip the current run if no change data exists.

Question 40: Skipped

### The VARIANT data type imposes, how much size limit on individual rows

- 64 MB Uncompressed
- 16 MB Compressed  
**(Correct)**
- 64 MB Compressed
- 16 MB Uncompressed

Question 41: Skipped

`select * from my_table before(statement => '8e5d0ca9-005e-44e6-b858-a8f5b37c5726');`  
This query will select historical data from "my\_table" up to, but not including any changes made by the specified statement

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 42: Skipped

### What is the full form of JSON?

- Java Semi Object Notation
- Java Object Notation
- Javascript Notation
- JavaScript Object Notation

(Correct)

Question 43: Skipped

**UNDROP command is applicable to which all objects**

- ROLES
- USERS
- TABLES  
**(Correct)**
- SCHEMA  
**(Correct)**
- DATABASES  
**(Correct)**

**Explanation**

Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 44: Skipped

**To restore the table by undrop command what all conditions should be met?**

- User must of CREATE privileges on the schema in which table is getting restored  
**(Correct)**
- User executing undrop command should have ownership privilege on the table  
**(Correct)**
- The table can be restored in any schema irrespective from whichever schema the table was dropped.
- The table can be restored in same schema only from which it has been dropped  
**(Correct)**

**Explanation**

Similar to dropping an object, a user must have OWNERSHIP privileges for an object to restore it. In addition, the user must have CREATE privileges on the object type for the database or schema where the dropped object will be restored. Restoring tables and schemas is only supported in the current schema or current database, even if a fully-qualified object name is specified Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html#data-retention-period>

Question 45: Skipped

**select \* from my\_table at(offset => -60\*5); This query will select historical data from "my\_table" as of -**

-

5 minutes ago  
**(Correct)**

- 5 days ago
- 5 seconds ago
- 5 hours ago

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 46: Skipped

**Snowflake minimizes the amount of storage required for historical data by maintaining only the information required to restore the individual table rows that were updated or deleted**

- FALSE
- TRUE

**(Correct)**

### Explanation

The fees are calculated for each 24-hour period (i.e. 1 day) from the time the data changed. The number of days historical data is maintained is based on the table type and the Time Travel retention period for the table. Also, Snowflake minimizes the amount of storage required for historical data by maintaining only the information required to restore the individual table rows that were updated or deleted. As a result, storage usage is calculated as a percentage of the table that changed. Full copies of tables are only maintained when tables are dropped or truncated.

Question 47: Skipped

**In Snowflake Enterprise edition, the maximum retention period for a transient or temporary table is?**

- 90 days
- Transient & Temporary table cannot be configured for time travel
- 7 days
- 1 day

**(Correct)**

### Explanation

For Snowflake Enterprise Edition (and higher): For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also

true for temporary tables. For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days

Question 48: Skipped

**Within a transaction multiple SQL statements will see the same set of records in a stream**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Streams support repeatable read isolation. In repeatable read mode, multiple SQL statements within a transaction see the same set of records in a stream. This differs from the read committed mode supported for tables, in which statements see any changes made by previous statements executed within the same transaction, even though those changes are not yet committed.

Question 49: Skipped

**To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides which two table types**

- TRANSIENT  
**(Correct)**
- TEMPORARY  
**(Correct)**
- MATERIALIZED
- PERMANENT

### Explanation

To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides two table types, temporary and transient, which do not incur the same fees as standard (i.e. permanent) tables: Transient tables can have a Time Travel retention period of either 0 or 1 day. Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon as the table is dropped or the session in which the table was created ends. Transient and temporary tables have no Fail-safe period. As a result, the maximum additional fees incurred for Time Travel and Fail-safe by these types of tables is limited to 1 day

Question 50: Skipped

**Which of the following currently uses user-specified warehouse, but eventually run in a serverless environment using Snowflake-supplied compute resources**

- Stored Procedure
-

Tasks  
**(Correct)**

- Snowpipe
- Streams

**Explanation**

Refer Link - <https://docs.snowflake.com/en/user-guide/tasks.html>

Question 51: Skipped

**In Snowflake Enterprise edition, the maximum retention period for a permanent table is?**

- 365 days
- 1 day
- 7 days
- 90 days

**(Correct)**

**Explanation**

For Snowflake Enterprise Edition (and higher): For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also true for temporary tables. For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days

Question 52: Skipped

**Fail Safe is not supported in Standard Snowflake edition**

- FALSE
- TRUE

**Explanation**

Fail Safe is a default feature available with all snowflake edition, without any extra license cost

Question 53: Skipped

**Which two table types can have a maximum retention period of 1 day?**

- MATERIALIZED
- PERMANENT

- TEMPORARY  
**(Correct)**

- TRANSIENT  
**(Correct)**

### Explanation

To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides two table types, temporary and transient, which do not incur the same fees as standard (i.e. permanent) tables: Transient tables can have a Time Travel retention period of either 0 or 1 day. Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon as the table is dropped or the session in which the table was created ends. Transient and temporary tables have no Fail-safe period. As a result, the maximum additional fees incurred for Time Travel and Fail-safe by these types of tables is limited to 1 day

Question 54: Skipped

**When a task is created, it is in resumed state by default?**

- TRUE
- FALSE  
**(Correct)**

### Explanation

Create a task using CREATE TASK. The task is suspended by default.

Question 55: Skipped

**If the role that a running task is executing under is dropped while the task is running, the task completes processing under the dropped role.**

- TRUE  
**(Correct)**
- FALSE

Question 56: Skipped

**Currently, streams cannot track changes in materialized views.**

- FALSE
- TRUE  
**(Correct)**

Question 57: Skipped

**Which file format option can be enabled in JSON or AVRO file to remove the outer array structure and load the records into separate table rows**

- STRIP\_OUTER\_ARRAY

(Correct)

- ELIMINATE\_OUTER\_ARRAY
- REMOVE\_OUTER\_ARRAY
- READ\_INTERNAL\_ARRAYS

### Explanation

In general, JSON and Avro data sets are a simple concatenation of multiple documents. The JSON or Avro output from some software is composed of a single huge array containing multiple records. There is no need to separate the documents with line breaks or commas, though both are supported. Instead, we recommend enabling the STRIP\_OUTER\_ARRAY file format option for the COPY INTO <table> command to remove the outer array structure and load the records into separate table rows: copy into <table> from @~/<file>.json file\_format = (type = 'JSON' strip\_outer\_array = true);

Question 58: Skipped

### Extended time travel of 90 days is supported starting which snowflake edition

- Virtual Private Snowflake
- Business Critical
- Enterprise
- Premier
- Standard

(Correct)

### Explanation

Standard Time Travel is 1 day. Extended Time Travel (up to 90 days) requires Snowflake Enterprise Edition.

Question 59: Skipped

### Table stream also contains actual table data

- TRUE
- FALSE

(Correct)

### Explanation

A stream itself does not contain any table data. A stream only stores the offset for the source table and returns CDC records by leveraging the versioning history for the source table. When the first stream for a table is created, a pair of hidden columns are added to the source

table and begin storing change tracking metadata. These columns consume a small amount of storage. The CDC records returned when querying a stream rely on a combination of the offset stored in the stream and the change tracking metadata stored in the table.

Question 60: Skipped

**JSON data is a hierarchical collection of name/value pairs grouped into objects and arrays**

- TRUE  
**(Correct)**
- FALSE

### Explanation

JSON data is a hierarchical collection of name/value pairs grouped into objects and arrays:

Colons : separate names and values in name/value pairs. Curly braces {} denote objects.

Square brackets [] denote arrays. Commas , separate entities in objects and arrays. Refer Link - <https://docs.snowflake.com/en/user-guide/semistructured-intro.html>

Question 61: Skipped

**You are a developer with in one company running a Snowflake data warehouse. You required to load JSON based data into a table. What is the best way to do it in snowflake?**

- Ask the source team to send data into CSV instead of JSON
- Write a program in Java or Python to convert the JSON into tabular format and then load in Snowflake table
- Create a table with variant column and load the json data into that table  
**(Correct)**

### Explanation

Snowflake provides capability to store the semi structured data into Variant column.

Question 62: Skipped

**What is the column datatype used for ingesting semi-structured data?**

- VARIANT  
**(Correct)**
- STRING
- VARCHAR
- DATETIME
- BINARY

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 63: Skipped

### Fail Safe and time travel are one and the same thing

- TRUE
- FALSE  
**(Correct)**

### Explanation

Separate and distinct from Time Travel, Fail-safe ensures historical data is protected in the event of a system failure or other catastrophic event, e.g. a hardware failure or security breach.

Question 64: Skipped

### Fail Safe is configurable by the user

- TRUE
- FALSE  
**(Correct)**

### Explanation

Fail-safe provides a (non-configurable) 7-day period during which historical data is recoverable by Snowflake. This period starts immediately after the Time Travel retention period ends.

Question 65: Skipped

### If a task is still running when the next scheduled execution time occurs, what happens?

- The next execution of the tasks starts as soon as the current execution gets over
- The tasks that is running fails and retriggers
- That schedule time is skipped  
**(Correct)**
- The tasks that is running fails and stops

### Explanation

Snowflake ensures only one instance of a task with a schedule (i.e. a standalone task or the root task in a tree of tasks) is executed at a given time. If a task is still running when the next scheduled execution time occurs, then that scheduled time is skipped.

Question 66: Skipped

A table named "MY\_TABLE" has a retention period of 7 days. The table is dropped, and a new table with the same name "MY\_TABLE" has been created. Now if the UNDROP MY\_TABLE command is executed what will happen?

- The UNDROP command will fail  
**(Correct)**
- The UNDROP command will restore the dropped version with a new name
- The UNDROP command will overwrite the already existing table with the previously dropped table

### Explanation

If an object with the same name already exists, UNDROP fails. You must rename the existing object, which then enables you to restore the previous version of the object.

Question 67: Skipped

### The cron expression in a task definition supports specifying a time zone

- TRUE  
**(Correct)**
- FALSE

### Explanation

The cron expression in a task definition supports specifying a time zone. A scheduled task runs according to the specified cron expression in the local time for a given time zone. Special care should be taken with regard to scheduling tasks for time zones that recognize daylight saving time. Tasks scheduled during specific times on days when the transition from standard time to daylight saving time (or the reverse) occurs can have unexpected behaviors.

Question 68: Skipped

### What features are provided by Snowflake to enable continuous data pipelines?

- Recurring Tasks
- Change Data Tracking
- All options are correct  
**(Correct)**
- Continous Data Loading

### Explanation

Refer document - <https://docs.snowflake.com/en/user-guide/data-pipelines-intro.html>

Question 69: Skipped

### Which type of table stream tracks all DML changes to the source table including inserts, updates, deletes and also truncate?

- Standard  
**(Correct)**

- Multifunction
- Insert-Only
- Native
- Append-Only

### Explanation

A standard (i.e. delta) table stream tracks all DML changes to the source table, including inserts, updates, and deletes (including table truncates). This stream type performs a join on inserted and deleted rows in the change set to provide the row level delta. As a net effect, for example, a row that is inserted and then deleted between two transactional points of time in a table is removed in the delta (i.e. is not returned when the stream is queried).

Question 70: Skipped

**True or false: A table in Snowflake can contain both a variant columns and a timestamp.**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 71: Skipped

**Which of the following roles can use SQL to view the task history within a specified date range**

- Any role can view
- Task owner (i.e. role that has the OWNERSHIP privilege on a task)  
**(Correct)**
- Any role that has the global MONITOR EXECUTION privilege  
**(Correct)**
- Public
- Accountadmin  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/tasks-intro.html>

Question 72: Skipped

### A single tasks in snowflake can execute multiple SQL Statements

- FALSE  
**(Correct)**
- TRUE

### Explanation

Currently, a task can execute a single SQL statement, including a call to a stored procedure. Tasks can be combined with table streams for continuous ELT workflows to process recently changed table rows. Streams ensure exactly once semantics for new or changed data in a table. Tasks can also be used independently to generate periodic reports by inserting or merging rows into a report table or perform other periodic work.

Question 73: Skipped

### Dropped tables, schemas, and databases can be listed using the SHOW commands with which keyword specified?

- HISTORY  
**(Correct)**
- OLD VERSIONS
- PREVIOUS VERSIONS
- DROPPED\_RECORDS

### Explanation

Dropped tables, schemas, and databases can be listed using the following commands with the HISTORY keyword specified: SHOW TABLES SHOW SCHEMAS SHOW DATABASES. The output includes all dropped objects and an additional DROPPED\_ON column, which displays the date and time when the object was dropped. If an object has been dropped more than once, each version of the object is included as a separate row in the output.

Question 74: Skipped

### Using Time Travel, you can perform which of the following actions within a defined period of time (i.e. within the time travel period of the object)

- Create clones of entire tables, schemas, and databases at or before specific points in the past
- All options are correct  
**(Correct)**
- Query data in the past that has since been updated or deleted

- 

Restore tables, schemas, and databases that have been dropped

### Explanation

Using Time Travel, you can perform the following actions within a defined period of time: Query data in the past that has since been updated or deleted. Create clones of entire tables, schemas, and databases at or before specific points in the past. Restore tables, schemas, and databases that have been dropped. Once the defined period of time has elapsed, the data is moved into Snowflake Fail-safe and these actions can no longer be performed.

Question 75: Skipped

### What is the default time limit on a single run of a task

- 

30 min

- 

60 min

(Correct)

- 

There is no limit as such

- 

10 min

### Explanation

There is a 60 minute default limit on a single run of a task. This limitation was implemented as a safeguard against non-terminating tasks. Query the TASK\_HISTORY table function. If the task was canceled or exceeded the window scheduled for the task, the cause is often an undersized warehouse. Review the warehouse size and consider increasing it to fit within the schedule window or the one-hour limit. Alternatively, consider increasing the timeout limit for the task by executing ALTER TASK ... SET USER\_TASK\_TIMEOUT\_MS = <num>.

Question 76: Skipped

### MAX\_DATA\_EXTENSION\_TIME\_IN\_DAYS - This parameter is used for what purpose?

- 

Maximum number of days for which Snowflake can extend the data retention period for fail safe of an object

- 

Maximum number of days for which Snowflake can extend the data retention period for time travel of an object

- 

Maximum number of days for which Snowflake can extend the data retention period for tables to prevent streams on the tables from becoming stale

(Correct)

### Explanation

Maximum number of days for which Snowflake can extend the data retention period for tables to prevent streams on the tables from becoming stale. By default, if the DATA\_RETENTION\_TIME\_IN\_DAYS setting for a source table is less than 14 days, and a stream has not been consumed, Snowflake temporarily extends this period to the stream's

offset, up to a maximum of 14 days, regardless of the Snowflake Edition for your account. The MAX\_DATA\_EXTENSION\_TIME\_IN\_DAYS parameter enables you to limit this automatic extension period to control storage costs for data retention or for compliance reasons.

Question 77: Skipped

### Multiple streams can be created on a single table

- TRUE  
**(Correct)**
- FALSE

Question 78: Skipped

**Snowflake Tasks executions are associated with the user who created the task, i.e. in the history database we can see that the task is getting executed under the username of the User who created that task**

- TRUE
- FALSE  
**(Correct)**

### Explanation

Snowflake runs tasks with the privileges of the task owner (i.e. the role that has OWNERSHIP privilege on the task), but task runs are not associated with a user. Instead, each run is executed by a system service. Tasks are decoupled from specific users to avoid complications that can arise when users are dropped, locked due to authentication issues, or have roles removed. Because task runs are decoupled from a user, the query history for task runs are associated with the system service. SYSTEM is not a user in the account; it is a behind-the-scenes service. As such, there are no user credentials for this service, and no individual (from Snowflake or in your account) can assume its identity. Activity for the system service is limited to your account. The same encryption protections and other security protocols are built into this service as are enforced for other operations.

Question 79: Skipped

### Which semi structure data types are available in Snowflake

- OBJECT  
**(Correct)**
- STRING
- VARIANT  
**(Correct)**
- ARRAY  
**(Correct)**

- SEMI STRUCTURE

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 80: Skipped

**If you have a table with a 10-day retention period and increase the period to 20 days, then what happens to the data that is currently in time travel for past 2 days.**

- The data will move to fail safe immediately
- The data would move to fail safe as per previous retention period only i.e. will move to fail safe after 10 days of time travel
- 
- The table will lose all time travel data, and the new time travel retention period will be applied going forward
- The data would move to fail safe as per new retention period i.e. will move to fail safe after 20 days of time travel

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Question 81: Skipped

**Which of the following is a common case for Cloning in Snowflake? (Select all that apply)**

- Data Encryption Protection
- "Point in Time" Snapshot  
(Correct)
- Agile Release and Development  
(Correct)
- Data Life Cycle Management  
(Correct)

Question 82: Skipped

**Many logical data models refer to entities, attributes and values. In JSON, and some other semi-structured types, the word "attribute" is often replaced with what word?**

- Columns
-

- Rows
- Relationships
- Values
- Keys  
**(Correct)**

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 83: Skipped

### The CDC records returned when querying a stream rely on a combination of -

- Offset Stored in Stream  
**(Correct)**
- The change tracking metadata stored in the table  
**(Correct)**
- Offset stored in Table
- Change Tracking Metadata Stored in Stream

### Explanation

Note that a stream itself does not contain any table data. A stream only stores the offset for the source table and returns CDC records by leveraging the versioning history for the source table. When the first stream for a table is created, a pair of hidden columns are added to the source table and begin storing change tracking metadata. These columns consume a small amount of storage. The CDC records returned when querying a stream rely on a combination of the offset stored in the stream and the change tracking metadata stored in the table.

Question 84: Skipped

### JSON arrays are written inside which type of brackets?

- Curly Braces
- Round Brackets
- Square Brackets  
**(Correct)**

### Explanation

JSON arrays are written inside square brackets. An array can contain multiple objects, separated by commas. For example: {"employees": [ {"firstName": "John",

```
"lastName":"Doe"}, {"firstName":"Anna", "lastName":"Smith"}, {"firstName":"Peter",  
"lastName":"Jones"} ] }
```

Question 85: Skipped

**Snowflake Tasks can be event based as well along with time based?**

- FALSE  
**(Correct)**
- TRUE

### Explanation

There is no event source that can trigger a task; instead, a task runs on a schedule, which can be defined when creating a task (using CREATE TASK) or later (using ALTER TASK). Snowflake ensures only one instance of a task with a schedule (i.e. a standalone task or the root task in a tree of tasks) is executed at a given time. If a task is still running when the next scheduled execution time occurs, then that scheduled time is skipped.

Question 86: Skipped

**Which of the following are true about the variant data type in Snowflake? Select all that apply.**

- Requires a custom mapping for each record type
- Stored in a separate file format from structured data
- Can be queried using JSON path notation  
**(Correct)**
- Optimized storage based on repeated elements  
**(Correct)**

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 87: Skipped

**Snowflake supports which of the following file formats for data loading? Select all that apply.**

- ORC  
**(Correct)**
- Parquet  
**(Correct)**
- PDF
-

CSV  
**(Correct)**

### Explanation

Refer Snowflake Document section of Semi Structure - <https://docs.snowflake.com/en/user-guide/semistructured-concepts.html>

Question 88: Skipped

**When working with unfamiliar semi-structured data, you might not know the key names in an object. You can use the FLATTEN function with which argument to return the list of distinct key names in all nested elements in an object?**

- RECURSIVE  
**(Correct)**
- NESTING
- GROUPING
- DISTINCT

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/semistructured-considerations.html>

Question 89: Skipped

**What are the additional metadata columns available in Stream to track changes?**

- METADATA\$ISDELETE
- METADATA\$ROW\_ID  
**(Correct)**
- METADATA\$ACTION  
**(Correct)**
- METADATA\$ISINSERT
- METADATA\$ISUPDATE  
**(Correct)**

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/streams.html>

Question 90: Skipped

**What is the data type provided by Snowflake to store semi structure data?**

- Varchar
-

Text

- Semi-structure
- Variant  
**(Correct)**

### Explanation

Snowflake loads semi-structured data into a single VARIANT column

Question 91: Skipped

**Which file format is an open-source data serialization and RPC framework originally developed for use with Apache Hadoop**

- XML
- JSON
- AVRO  
**(Correct)**
- ORC
- PARQUET

### Explanation

Avro is an open-source data serialization and RPC framework originally developed for use with Apache Hadoop. It utilizes schemas defined in JSON to produce serialized data in a compact binary format. The serialized data can be sent to any destination (i.e. application or program) and can be easily deserialized at the destination because the schema is included in the data

Question 92: Skipped

**What are the different types of streams available?**

- Native
- Insert-Only  
**(Correct)**
- Append-Only  
**(Correct)**
- Update-Only
-

Standard  
**(Correct)**

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guidestreams.html>

Question 93: Skipped

**Non-native values such as dates and timestamps are stored as strings when loaded into a VARIANT column**

- FALSE
- TRUE

**(Correct)**

### Explanation

Non-native values such as dates and timestamps are stored as strings when loaded into a VARIANT column, so operations on these values could be slower and also consume more space than when stored in a relational column with the corresponding data type.

Question 94: Skipped

**A task object defines a recurring schedule for executing a SQL statement, including statements that call stored procedures**

- TRUE
- FALSE

**(Correct)**

### Explanation

A task object defines a recurring schedule for executing a SQL statement, including statements that call stored procedures. Tasks can be chained together for successive execution to support more complex periodic processing.

Question 95: Skipped

**Storage fees are incurred for maintaining historical data during the Time Travel but not during Fail-safe periods as Fail safe is managed by Snowflake**

- TRUE
- FALSE

**(Correct)**

### Explanation

Storage cost is for both Time Travel & Fail Safe. Refer Link - <https://docs.snowflake.com/en/user-guide/data-cdp-storage-costs.html>

Question 96: Skipped

**Tasks cannot be chained together for successive execution to support more complex periodic processing**

- TRUE

- FALSE
- (Correct)

### Explanation

A task object defines a recurring schedule for executing a SQL statement, including statements that call stored procedures. Tasks can be chained together for successive execution to support more complex periodic processing.

Question 97: Skipped

**If multiple DML statements change a row, the stream records all the action taken on that row.**

- FALSE
- (Correct)
- TRUE

### Explanation

The stream maintains only the delta of the changes; if multiple DML statements change a row, the stream records only the latest action taken on that row

Question 98: Skipped

**Which command can be executed to get the version history of table "Test1" in schema "mytestdb.myschema" within the time travel period in snowflake?**

- show tables history like 'Test1' in mytestdb.myschema;
- show version of table like 'Test1' in mytestdb.myschema;
- show tables version like 'Test1' in mytestdb.myschema;
- show history of table like 'Test1' in mytestdb.myschema;

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-time-travel.html#data-retention-period>

Question 99: Skipped

**When a database or schema that contains a source table and stream is cloned, any unconsumed records in the stream (in the clone) are inaccessible**

- FALSE
- (Correct)
- TRUE

### Explanation

Currently, when a database or schema that contains a source table and stream is cloned, any unconsumed records in the stream (in the clone) are inaccessible. This behavior is consistent with Time Travel for tables. If a table is cloned, historical data for the table clone begins at the time/point when the clone was created

Question 100: Skipped

### What is the full form of ORC?

- Object Read Class
- Optimized Read Columnar
- Optimized Row Columnar  
**(Correct)**
- Object Row Columnar

### Explanation

Used to store Hive data, the ORC (Optimized Row Columnar) file format was designed for efficient compression and improved performance for reading, writing, and processing data over earlier Hive file formats. For more information about ORC, see <https://orc.apache.org/>.

Question 101: Skipped

Which of the following tables cannot be cloned as permanent table?

- Temporary Tables  
**(Correct)**
- Transient Tables  
**(Correct)**
- Permanent Tables

### Explanation

Refer Link for details - <https://docs.snowflake.com/en/user-guide/tables-temp-transient.html>

Refer the section Comparison of Table Types in the above link.

Question 102: Skipped

If a row is added and then updated in the current offset,then the METADATA\$ISUPDATE row records a TRUE value.

- FALSE  
**(Correct)**
- TRUE

### Explanation

Note that streams record the differences between two offsets. If a row is added and then updated in the current offset, the delta change is a new row. The METADATA\$ISUPDATE row records a FALSE value

[Continue](#)[Retake test](#)

Question 1: Skipped

**Which of the following is a separate database object?**

- User Stage
- Named Stage  
**(Correct)**
- Table Stage

Question 2: Skipped

**Copy options set in multiple locations are cumulative. Individual options set in one place override the same option set lower in the order of precedence.**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Notes at link - <https://docs.snowflake.com/en/user-guide/data-load-prepare.html>

Question 3: Skipped

**For Data unloading which command is used**

- COPY INTO <TABLE>
- UNLOAD DATA TO <LOCATION>
- UNLOAD DATA FROM <TABLE>
- COPY INTO <LOCATION>  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-unload-overview.html>

Question 4: Skipped

**Failover/Failback requires minimum which snowflake edition**

- Virtual Private Snowflake
- Standard
- Enterprise

- 

Business Critical  
**(Correct)**

### Explanation

Refer Link -<https://docs.snowflake.com/en/user-guide/database-replication-failover.html>

Question 5: Skipped

**Which of the following file format option is not supported in "COPY INTO TABLE (Copy Load)" command, for CSV type**

- 

SKIP\_BLANK\_LINES

- 

FIELD\_DELIMITER

- 

VALIDATE\_UTF8

- 

COMPRESSION

- 

FILE\_EXTENSION

**(Correct)**

- 

SKIP\_HEADER

### Explanation

FILE\_EXTENSION option is supported in case of Copy Unload (i.e. moving data out of the table), but not in case of Copy Load

Question 6: Skipped

**PIPE\_EXECUTION\_PAUSED parameter that is used to pause the execution of pipe can be set at which all levels?**

- 

Database

- 

Account

**(Correct)**

- 

Schema

**(Correct)**

- 

Role

- 

Pipe

**(Correct)**

### Explanation

At the pipe level, the object owner (or a parent role in a role hierarchy) can set the parameter to pause or resume an individual pipe. An account administrator (user with the ACCOUNTADMIN role) can set this parameter at the account level to pause or resume all pipes in the account. Likewise, a user with the MODIFY privilege on the schema can pause or resume pipes at the schema level. Note that this larger domain control only affects pipes for which the parameter was not already set at a lower level; e.g., by the owner at the object level.

Question 7: Skipped

**Which of the following conditions will result in the database replication operation to get failed?**

- The primary database is in an Enterprise (or higher) account and contains a masking policy but one or more of the accounts approved for replication are on lower editions  
**(Correct)**
- There are User Defined Functions (UDF) in the primary (source) database
- There are Materialized Views in the source database
- A masking policy contained in the primary database is applied to columns in a table or view in another database, or vice-versa  
**(Correct)**
- There are external tables in the primary database (source database)  
**(Correct)**

### Explanation

Refer link -<https://docs.snowflake.com/en/user-guide/database-replication-considerations.html>

Question 8: Skipped

**Account parameters are also replicated, when a database replication is setup**

- TRUE
- FALSE  
**(Correct)**

### Explanation

Account parameters are not replicated.

Question 9: Skipped

**When the JSON or Avro output from some software is composed of a single huge array containing multiple records, then which file format option should be enabled to load the records in multiple rows**

- ELIMINATE\_OUTER\_ARRAY
-

REMOVE\_OUTER\_ARRAY

- 
- STRIP\_OUT\_ARRAY
- STRIP\_OUTER\_ARRAY  
**(Correct)**

#### Explanation

In general, JSON and Avro data sets are a simple concatenation of multiple documents. The JSON or Avro output from some software is composed of a single huge array containing multiple records. There is no need to separate the documents with line breaks or commas, though both are supported. Instead, we recommend enabling the STRIP\_OUTER\_ARRAY file format option for the COPY INTO <table> command to remove the outer array structure and load the records into separate table rows

Question 10: Skipped

**All data files unloaded to Snowflake internal locations are automatically encrypted using 128-bit keys. 256-bit keys can be enabled (for stronger encryption); however, additional configuration is required.**

- 
- TRUE  
**(Correct)**
- FALSE

#### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/intro-summary-unloading.html>

Question 11: Skipped

**If file format options are specified in multiple locations, the load operation applies the options in the following order of precedence:**

- 
- Table Definition,COPY INTO TABLE statement, Stage Definition
- Stage Definition,COPY INTO TABLE statement, Table Definition
- Table Definition, Stage Definition,COPY INTO TABLE statement
- COPY INTO TABLE statement, Stage Definition, Table Definition  
**(Correct)**

#### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-prepare.html>

Question 12: Skipped

**What is the name of the warehouse in which the credit utilized by Snowpipe is captured?**

- SNOWPIPE

**(Correct)**

- SNOWPIPE\_CREDITS
- Snowpipe Credits are not captured and send as separate monthly bill to customer
- Any existing warehouse can be configured to log the SNOWPIPE charges

**Explanation**

Refer link - <https://docs.snowflake.com/en/user-guide/data-load-snowpipe-billing.html>

Question 13: Skipped

**When a copy command cannot definitively determine whether a file has been loaded already, then what is the default behaviour?**

- The file load is skipped
- The file is loaded
- Error message is generated, that load metadata for the file for given table doesn't exist

**Explanation**

When the COPY command cannot definitively determine whether a file has been loaded already if the LAST\_MODIFIED date is older than 64 days and the initial set of data was loaded into the table more than 64 days earlier (and if the file was loaded into the table, that also occurred more than 64 days earlier). In this case, to prevent accidental reload, the command skips the file by default. Refer link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html>

Question 14: Skipped

**User stages are referenced using?**

- "@%"
- "@"
- "@~"
- "%~"

**(Correct)**

**Explanation**

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-local-file-system-create-stage.html>

Question 15: Skipped

**Which is the default encoding considered in COPY command while loading delimited files like CSV, TSV?**

- UTF-16
- ANSI
- UCS-2
- UTF-8

**(Correct)**

### **Explanation**

For delimited files (CSV, TSV, etc.), the default character set is UTF-8. To use any other character sets, you must explicitly specify the encoding to use for loading. For the list of supported character sets, see link - <https://docs.snowflake.com/en/user-guide/intro-summary-loading.html>

Question 16: Skipped

**When the session parameter TWO\_DIGIT\_CENTURY\_START is set to 1980, then value 79 will be parsed as which year?**

- 1879
- 2079
- 2179
- 1979

**(Correct)**

### **Explanation**

Two-digit year, controlled by the TWO\_DIGIT\_CENTURY\_START session parameter, e.g. when set to 1980, values of 79 and 80 parsed as 2079 and 1980 respectively. Refer link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-prepare.html>

Question 17: Skipped

**Which types of database objects are currently not supported while doing database replication?**

- Tasks
- Temporary Tables
- 

**(Correct)**

**(Correct)**

## Permanent Tables

- Sequences
- Stages  
**(Correct)**
- Streams  
**(Correct)**

## Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 18: Skipped

**We have few clone tables in a database. Now for the database, replication is setup. In the secondary (target) database, the clone object will not be physically copied, and hence no storage cost will be incurred for them. True or False?**

- FALSE  
**(Correct)**
- TRUE

## Explanation

Cloned objects are replicated physically rather than logically to secondary databases. That is, cloned tables in a standard database do not contribute to the overall data storage unless or until DML operations on the clone add to or modify existing data. However, when a cloned table is replicated to a secondary database, the physical data is also replicated, increasing the data storage usage for your account.

Question 19: Skipped

The COPY INTO <location> command provides a copy option (SINGLE) for unloading data into a single file or multiple files. The default is?

- SINGLE = TRUE (i.e. data extracted in single file)
- SINGLE = FALSE (i.e. data extracted in multiple files)  
**(Correct)**

## Explanation

The COPY INTO <location> command provides a copy option (SINGLE) for unloading data into a single file or multiple files. The default is SINGLE = FALSE (i.e. unload into multiple files). Snowflake assigns each file a unique name. The location path specified for the command can contain a filename prefix that is assigned to all the data files generated. If a prefix is not specified, Snowflake prefixes the generated filenames with data\_.Snowflake appends a suffix that ensures each file name is unique across parallel execution threads; e.g. data\_stats\_0\_1\_0. When unloading data into multiple files, use the MAX\_FILE\_SIZE copy option to specify the maximum size of each file created.

Question 20: Skipped

**Table named mytable has a stage referenced as:**

- "@~mytable"
- "@mytable"
- "@%mytable"  
**(Correct)**
- "%~mytable"

**Explanation**

Table stages have the same name as the table; e.g. a table named mytable has a stage referenced as @%mytable

Question 21: Skipped

**What is the file size limit that can be loaded into a table using the Snowflake Web Interface**

- 10 MB
- 20 MB
- 100 MB
- 50 MB  
**(Correct)**

**Explanation**

The wizard is only intended for loading small numbers of files of limited size (up to 50 MB). This file size limit is intended to ensure better performance because browser performance varies from computer to computer and between different browser versions. Also, the memory consumption required to encrypt larger files might cause a browser to run out of memory and crash. For loading larger files or large numbers of files, we recommend using the Snowflake client, SnowSQL

Question 22: Skipped

**Which of the following options for loading data files from a stage using COPY command the slowest?**

- Specifying a list of specific files to load.
- By path (internal stages) / prefix (Amazon S3 bucket)
- Using pattern matching to identify specific files by pattern  
**(Correct)**

- 

There is no difference in the speed of loading for any option

### Explanation

In general, pattern matching using a regular expression is generally the slowest of the three options for identifying/specifying data files to load from a stage; however, this option works well if you exported your files in named order from your external application and want to batch load the files in the same order.

Question 23: Skipped

### Load Metadata for each table expires in how many days?

- 

365

- 

14

- 

64

(Correct)

- 

90

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html>

Question 24: Skipped

### For Downloading a file that is staged in Snowflake Internal Stage which command is used

- 

RETRIEVE

- 

GET

(Correct)

- 

DOWNLOAD

- 

PUT

### Explanation

Download the file from the stage: From a Snowflake stage, use the GET command to download the data file(s). From S3, use the interfaces/tools provided by Amazon S3 to get the data file(s). From Azure, use the interfaces/tools provided by Microsoft Azure to get the data file(s). Refer Link - <https://docs.snowflake.com/en/user-guide/data-unload-overview.html>

Question 25: Skipped

### A pipe is a named, first-class Snowflake object that contains an INSERT statement used by Snowpipe.

- 

TRUE

- FALSE
- (Correct)

### Explanation

A pipe is a named, first-class Snowflake object that contains a COPY statement used by Snowpipe. The COPY statement identifies the source location of the data files (i.e., a stage) and a target table. All data types are supported, including semi-structured data types such as JSON and Avro.

Question 26: Skipped

**What is the default value for MAX\_FILE\_SIZE parameter while unloading a file?**

- 25 MB
- 16 MB
- 32 MB
- 10 MB

### Explanation

By default, COPY INTO location statements separate table data into a set of output files to take advantage of parallel operations. The maximum size for each file is set using the MAX\_FILE\_SIZE copy option. The default value is 16000000 (16 MB) but can be increased to accommodate larger files. The maximum file size supported is 5 GB for Amazon S3, Google Cloud Storage, or Microsoft Azure stages.

Question 27: Skipped

**Which parameter is available in COPY command to load files using a regular expression?**

- REGEXP
- REGEX
- LIKE
- PATTERN

### Explanation

Sample copy command - copy into people\_data from @%people\_data/data1/  
pattern='.\*person\_data[^0-9{1,3}\$\$].csv';

Question 28: Skipped

**Snowpipes maintain the load history metadata for how many days**

-

365

- 14  
**(Correct)**
- 64
- 30

### Explanation

Snowflake uses file loading metadata to prevent reloading the same files (and duplicating data) in a table. Snowpipe prevents loading files with the same name even if they were later modified (i.e. have a different eTag). The file loading metadata is associated with the pipe object rather than the table. As a result: Staged files with the same name as files that were already loaded are ignored, even if they have been modified, e.g. if new rows were added or errors in the file were corrected. Truncating the table using the TRUNCATE TABLE command does not delete the Snowpipe file loading metadata. However, note that pipes only maintain the load history metadata for 14 days

Question 29: Skipped

**File format options set in multiple locations are not cumulative. Any options set in one place override all options (whether the same or different options) set lower in the order of precedence.**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Notes at link - <https://docs.snowflake.com/en/user-guide/data-load-prepare.html>

Question 30: Skipped

**In case of database replication, the charges for compute and data transfer is charged to which account?**

- There is no charges for database replication
- Target Account (Account with Secondary database that is getting refreshed)  
**(Correct)**
- Source Account (Account with Primary Database)

### Explanation

Charges based on database replication are divided into two categories: data transfer and compute resources. Both categories are billed on the target account (i.e. the account that stores the secondary database that is refreshed).

Question 31: Skipped

**To see files that have been uploaded to a Snowflake stage, which command should be used**

- VIEW
- STAGE\_VIEW
- LIST  
**(Correct)**
- STAGE\_LIST

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-local-file-system-stage.html>  
Question 32: Skipped

**Databases created from shares cannot be replicated.**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 33: Skipped

If the null values in JSON documents indicate missing values and have no other special meaning, Snowflake recommend setting which file format option to TRUE for the COPY INTO <table> command when loading the JSON files.

- ELIMINATE\_NULL\_VALUES
- STRIP\_NULL\_VALUES  
**(Correct)**
- RETAIN\_NULL\_VALUES
- REMOVE\_NULL\_VALUES

### Explanation

In a VARIANT column, NULL values are stored as a string containing the word “null” not the SQL NULL value. If the “null” values in your JSON documents indicate missing values and have no other special meaning, we recommend setting the file format option STRIP\_NULL\_VALUES to TRUE for the COPY INTO <table> command when loading the JSON files. Retaining the “null” values often wastes storage and slows query processing.

Question 34: Skipped

**Privileges granted on database objects are not replicated to a secondary database**

- FALSE
- TRUE  
**(Correct)**

#### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 35: Skipped

**External tables** enable querying existing data stored in external cloud storage for analysis without first loading it into Snowflake.

- FALSE
- TRUE  
**(Correct)**

#### Explanation

External tables enable querying existing data stored in external cloud storage for analysis without first loading it into Snowflake. The source of truth for the data remains in the external cloud storage. Data sets materialized in Snowflake via materialized views are read-only. This solution is especially beneficial to accounts that have a large amount of data stored in external cloud storage and only want to query a portion of the data; for example, the most recent data. Users can create materialized views on subsets of this data for improved query performance.

Question 36: Skipped

**Replication operations use Snowflake-provided compute resources to copy data between accounts across regions.**

- TRUE  
**(Correct)**
- FALSE

#### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-billing.html>

Question 37: Skipped

**Resource monitors provide control over virtual warehouse credit usage; however, you cannot use them to control credit usage for the Snowflake-provided warehouses, including the SNOWPIPE warehouse**

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/data-load-snowpipe-billing.html>

Question 38: Skipped

### Bulk loading relies on the user-provided virtual warehouses

- TRUE  
**(Correct)**
- FALSE

### Explanation

Bulk loading relies on user-provided virtual warehouses, which are specified in the COPY statement. Users are required to size the warehouse appropriately to accommodate expected loads.

Question 39: Skipped

### Object parameters are not replicated, except the parameter "DATA\_RETENTION\_TIME\_IN\_DAYS" for schemas & tables

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 40: Skipped

### How much size limit the VARIANT data type imposes on individual rows.

- 10 MB (compressed)
- 16 MB (compressed)  
**(Correct)**
- 16 MB (uncompressed)
- 10 MB (uncompressed)

### Explanation

Refer link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-prepare.html>

Question 41: Skipped

### When staging uncompressed files in a Snowflake stage, the files are automatically compressed using gzip, unless compression is explicitly disabled.

- TRUE  
**(Correct)**
-

FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/intro-summary-loading.html>

Question 42: Skipped

**Which all permissions snowflake requires on a S3 bucket and folder to be able to access files in the folder (and subfolders)**

- s3>ListBucket  
**(Correct)**

- s3GetObjectVersion  
**(Correct)**

- s3ReadObject

- s3GetObject  
**(Correct)**

### Explanation

Step 1: Configure Access Permissions for the S3 Bucket

AWS Access Control Requirements

Snowflake requires the following permissions on an S3 bucket and folder to be able to access files in the folder (and sub-folders):

- s3GetObject
- s3GetObjectVersion
- s3ListBucket

Reference Link for above snapshot - <https://docs.snowflake.com/en/user-guide/data-load-s3-config-storage-integration.html>

Question 43: Skipped

**Which Snowflake Objects are currently not supported for replication?**

- Users  
**(Correct)**

- Shares  
**(Correct)**

- Databases

- Resource Monitors  
**(Correct)**

- Roles  
**(Correct)**



Warehouses  
**(Correct)**

### Explanation

Currently, replication is supported for databases only. Other types of objects in an account cannot be replicated. This list includes: Users Roles Warehouses Resource monitors Shares Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 44: Skipped

### What is VALIDATE\_MODE parameter used for in COPY INTO command?



To Validate the Snowflake connectivity



To Validate Connectivity with the External Cloud Provider (AWS, GCP, Azure)



To validate syntax of copy command



To Validate that the data in the uploaded files in Stage will load correctly.

**(Correct)**

### Explanation

Before loading your data, you can validate that the data in the uploaded files will load correctly. To validate data in an uploaded file, execute COPY INTO <table> in validation mode using the VALIDATION\_MODE parameter. The VALIDATION\_MODE parameter returns any errors that it encounters in a file. You can then modify the data in the file to ensure it loads without error. In addition, the ON\_ERROR copy option for the COPY INTO <table> command indicates what action to perform if errors are encountered in a file during loading

Question 45: Skipped

### Historical usage data (i.e. data in views like LOGIN\_HISTORY, COPY\_HISTORY, QUERY\_HISTORY) for activity in the primary database is also replicated to secondary databases.



TRUE



FALSE

**(Correct)**

### Explanation

Historical usage data for activity in the primary database is not replicated to secondary databases. Each account has its own query history, login history, etc. Historical usage data includes the query data returned by the following Information Schema table functions or Account Usage views: COPY\_HISTORY LOGIN\_HISTORY QUERY\_HISTORY etc.

Question 46: Skipped

### What is the recommended size for data files to optimize the number of parallel operations for a load



10 MB to 100 MB uncompressed

- 10 MB to 100 MB compressed
- 100MB to 250 MB compressed  
**(Correct)**
- 100MB to 250 MB uncompressed

### Explanation

The number of load operations that run in parallel cannot exceed the number of data files to be loaded. To optimize the number of parallel operations for a load, we recommend aiming to produce data files roughly 100-250 MB (or larger) in size **compressed**.

Refer Link -<https://docs.snowflake.com/en/user-guide/data-load-considerations-prepare.html>

Question 47: Skipped

**What is the maximum number of file names that can be specified in the file parameter of COPY command to load into a table**

- Unlimited
- 10000
- 100
- 1000  
**(Correct)**

### Explanation

Of the three options for identifying/specifying data files to load from a stage, providing a discrete list of files is generally the fastest; however, the FILES parameter supports a maximum of 1,000 files, meaning a COPY command executed with the FILES parameter can only load up to 1,000 files. Refer link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html>

Question 48: Skipped

**Which function to be used combined with COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file**

- OBJECT\_TABLE
- OBJECT\_EXTRACT
- OBJECT\_CONSTRUCT

(Correct)

- OBJECT\_JSON

### Explanation

You can use the OBJECT\_CONSTRUCT function combined with the COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file

Question 49: Skipped

**When unloading data into multiple files, which copy option is used to specify the maximum size of each file created.**

- MAX\_FILE\_SIZE  
**(Correct)**
- MAX\_SIZE
- MAX\_NBR\_ROWS
- MAX\_FILE\_SIZE\_MB

### Explanation

When unloading data into multiple files, use the MAX\_FILE\_SIZE copy option to specify the maximum size of each file created.

Question 50: Skipped

**Which types of database objects are currently supported while doing database replication?**

- Materialized Views  
**(Correct)**
- Stored Procedures  
**(Correct)**
- Streams
- File Formats  
**(Correct)**
- Stages
- Tasks

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 51: Skipped

**Which of the following compression methods cannot be detected automatically by snowflake**

- brotli  
**(Correct)**
- gzip
- bzip2
- Zstandard  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/intro-summary-loading.html>

Question 52: Skipped

**Which type of objects will cause database replication to fail?**

- Non-materialized views that reference any object in another database (e.g. table columns, other views, UDFs, or stages), and when both the databases are replicated.
- Materialized views that reference any object in another database  
**(Correct)**
- A object in the primary database references a dropped object  
**(Correct)**
- A foreign key in the primary database references a primary key in another database, or vice-versa  
**(Correct)**

### Explanation

Refer Dangling References section at link - <https://docs.snowflake.com/en/user-guide/database-replication-considerations.html>

Question 53: Skipped

**Which parameter should be set to true if the requirement is to load the files whose load metadata has expired**

- LOAD\_UNCERTAIN\_FILES  
**(Correct)**
- LOAD\_EXPIRED\_FILES

- LOAD\_ALL\_FILES

### Explanation

To load files whose metadata has expired, set the LOAD\_UNCERTAIN\_FILES copy option to true. The copy option references load metadata, if available, to avoid data duplication, but also attempts to load files with expired load metadata.

Question 54: Skipped

### For File formats like JSON, Avro the only supported character set is UTF-8

- TRUE  
**(Correct)**
- FALSE

### Explanation

Apart from delimited files like CSV, TSV, all other supported file formats (JSON, Avro, etc.), the only supported character set is UTF-8.

Question 55: Skipped

### Snowpipe generally loads older files first, there is no guarantee that files are loaded in the same order they are staged.

- TRUE  
**(Correct)**
- FALSE

### Explanation

For each pipe object, Snowflake establishes a single queue to sequence data files awaiting loading. As new data files are discovered in a stage, Snowpipe appends them to the queue. However, multiple processes pull files from the queue; and so, while Snowpipe generally loads older files first, there is no guarantee that files are loaded in the same order they are staged.

Question 56: Skipped

### Only named stages (internal or external) and user stages are supported for COPY transformations

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-transform.html>

Question 57: Skipped

### How much does the Snowpipe charges for queued files?

- 0.06 credits per 10000 files queued^

- 0.6 credit per 1000 files queued
- 0.06 credits per 1000 files queued  
**(Correct)**
- 0.06 credits per 100 files queued

### Explanation

In addition to resource consumption, an overhead to manage files in the internal load queue is included in the utilization costs charged for Snowpipe. This overhead increases in relation to the number of files queued for loading. Snowpipe charges 0.06 credits per 1000 files queued.

Question 58: Skipped

### What is the file encoding of the output files generated by Snowflake when unloading data

- UTF-8  
**(Correct)**
- ANSI
- UCS2
- UTF-16

### Explanation

Output files are always encoded using UTF-8, regardless of the file format; no other character sets are supported. Refer link - <https://docs.snowflake.com/en/user-guide/intro-summary-unloading.html>

Question 59: Skipped

### The table stage is designed to store files that are staged and managed by one or more users but only loaded into a single table.

- TRUE  
**(Correct)**
- FALSE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-overview.html>

Question 60: Skipped

### copy into mytable file\_format = (type = csv field\_delimiter = '|' skip\_header = 1); This command will load file from which stage

- User Stage

- Named Stage
- Table Stage  
**(Correct)**
- Error out as from clause is not specified

### Explanation

Note that when copying data from files in a table stage, the FROM clause can be omitted because Snowflake automatically checks for files in the table stage. Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-local-file-system-copy.html>

Question 61: Skipped

### What is the default compression of the output files generated by Snowflake when unloading data

- brotli
- Zstandard
- bzip2
- gzip  
**(Correct)**

### Explanation

By default, all unloaded data files are compressed using gzip, unless compression is explicitly disabled or one of the other supported compression methods is explicitly specified.

Question 62: Skipped

### Which of the following stages cannot be altered or dropped?

- Named Internal Stage
- Named External Stage
- Table  
**(Correct)**
- User  
**(Correct)**

### Explanation

A user stage is allocated to each user for storing files. This stage type is designed to store files that are staged and managed by a single user but can be loaded into multiple tables. User

stages cannot be altered or dropped. A table stage is available for each table created in Snowflake. This stage type is designed to store files that are staged and managed by one or more users but only loaded into a single table. Table stages cannot be altered or dropped. Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-overview.html>. Refer Section "Internal Stages"

Question 63: Skipped

**Which Snowflake object records DML changes made to tables, including inserts, updates, deletes as well as metadata about each change, so that actions can be taken using the changed data**

- Task
- Kafka
- Snowpipe
- Stream

(Correct)

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/streams.html>

Question 64: Skipped

**Which option should be set to load all files in copy command, ignoring the load metadata if it exists**

- LOAD\_ALL\_FILES
- OVERWRITE
- IGNORE\_METADATA
- FORCE

(Correct)

### Explanation

To load files whose metadata has expired, set the LOAD\_UNCERTAIN\_FILES copy option to true. The copy option references load metadata, if available, to avoid data duplication, but also attempts to load files with expired load metadata. Alternatively, set the FORCE option to load all files, ignoring load metadata if it exists. Note that this option reloads files, potentially duplicating data in a table.

Question 65: Skipped

**A file staged in a user stage can be loaded in only one table**

- FALSE

(Correct)

- TRUE

### Explanation

A user stage is allocated to each user for storing files. This stage type is designed to store files that are staged and managed by a single user but can be loaded into multiple tables

Question 66: Skipped

**Snowflake maintains detailed metadata for each table into which data is loaded, including:**

- Name of each file from which data was loaded  
**(Correct)**
- Number of rows parsed in the file  
**(Correct)**
- Column Names
- File Size  
**(Correct)**
- Number of Columns
- Etag for the file  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html>

Question 67: Skipped

**Table stages do support transforming data while loading it (i.e. using a query as the source for the COPY command).**

- FALSE  
**(Correct)**
- TRUE

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-local-file-system-create-stage.html>

Question 68: Skipped

**Which of the following compression methods can snowflake automatically detect**

- brotli
-

deflate  
**(Correct)**

- gzip
- (Correct)**

- bzip2
- (Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/intro-summary-loading.html>

Question 69: Skipped

**Source database have a materialized view, which gets replicated to the target database as part of database replication activity. The automatic maintenance & refresh of the materialized view in the target database will be enabled by default. True or False?**

- TRUE
- FALSE  
**(Correct)**

### Explanation

In the primary database, Snowflake performs automatic background maintenance of materialized views. When a base table changes, all materialized views defined on the table are updated by a background service that uses compute resources provided by Snowflake. In addition, if Automatic Clustering is enabled for a materialized view, then the view is monitored and reclustered as necessary in the primary database. A refresh operation replicates the materialized view definitions to a secondary database; however, the materialized view data is not replicated, meaning some or all of the data in the materialized views could become out of date. To perform automatic background maintenance of materialized views in a secondary database, explicitly set

AUTO\_REFRESH\_MATERIALIZED\_VIEWS\_ON\_SECONDARY = TRUE on the secondary database either when you create the secondary database (using CREATE DATABASE ... AS REPLICA OF) or later (using ALTER DATABASE). If Automatic Clustering is enabled for a materialized view in the primary database, then setting AUTO\_REFRESH\_MATERIALIZED\_VIEWS\_ON\_SECONDARY = TRUE on a secondary database also enables automatic monitoring and reclustering of the materialized view in the secondary database

Question 70: Skipped

**Database replication is supported across regions and across cloud platforms.**

- FALSE
- TRUE  
**(Correct)**

### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/database-replication-intro.html>

Question 71: Skipped

**Which type of stage can store files that are staged and managed by one or more users and loaded into one or more tables**

- User
- Table
- Named  
**(Correct)**

#### Explanation

A named internal stage is a database object created in a schema. This stage type can store files that are staged and managed by one or more users and loaded into one or more tables. Because named stages are database objects, the ability to create, modify, use, or drop them can be controlled using security access control privileges. Create stages using the CREATE STAGE command.

Question 72: Skipped

**Snowflake supports loading data from AWS, Azure, GCP regardless of the cloud platform on which the snowflake account is hosted**

- TRUE  
**(Correct)**
- FALSE

#### Explanation

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-overview.html>

Question 73: Skipped

**A stage links to a storage integration using a hidden ID rather than the name of the storage integration**

- TRUE  
**(Correct)**
- FALSE

#### Explanation

Error: Integration {0} associated with the stage {1} cannot be found Loading data from an external stage may produce an error similar to the following: 003139=SQL compilation error:\nIntegration "{0}" associated with the stage "{1}" cannot be found. This error can occur when the association between the external stage and the storage integration linked to the stage has been broken. This happens when the storage integration object has been recreated (using CREATE OR REPLACE STORAGE INTEGRATION). A stage links to a storage integration using a hidden ID rather than the name of the storage integration. Behind the scenes, the CREATE OR REPLACE syntax drops the object and recreates it with a different hidden ID. If you must recreate a storage integration after it has been linked to one or more stages, you must reestablish the association between each stage and the storage

integration by executing ALTER STAGE stage\_name SET STORAGE\_INTEGRATION = storage\_integration\_name,

Question 74: Skipped

**Which file formats are supported in the COPY command?**

- Avro
- TSV  
**(Correct)**
- Unstructured
- CSV  
**(Correct)**
- JSON  
**(Correct)**

**Explanation**

Refer link - <https://docs.snowflake.com/en/user-guide/intro-summary-loading.html>

Question 75: Skipped

**Using the Snowflake Web interface only the files from local system can be loaded, and the files available on cloud storage like AWS, Azure cannot be loaded**

- TRUE
- FALSE  
**(Correct)**

**Explanation**

You can choose to load data from files on your local machine or files already staged in an existing cloud storage location on Snowflake, Amazon S3, Google Cloud Storage, or Microsoft Azure.

Question 76: Skipped

**Which option should be set in COPY command, to delete/remove the staged files from Snowflake Stage after the file is successfully loaded**

- PURGE  
**(Correct)**
- REMOVE
- DELETE

**Explanation**

Files that were loaded successfully can be deleted from the stage during a load by specifying the PURGE copy option in the COPY INTO <table> command.

Question 77: Skipped

**A file staged in a user stage can be managed by multiple users**

- FALSE  
**(Correct)**
- TRUE

**Explanation**

A user stage is allocated to each user for storing files. This stage type is designed to store files that are staged and managed by a single user but can be loaded into multiple tables

Question 78: Skipped

**Snowpipe uses compute resources provided by Snowflake (i.e. a serverless compute model)**

- FALSE
- TRUE  
**(Correct)**

**Explanation**

Snowpipe uses compute resources provided by Snowflake (i.e. a serverless compute model). These Snowflake-provided resources are automatically resized and scaled up or down as required, and are charged and itemized using per-second billing. Data ingestion is charged based upon the actual workloads.

Question 79: Skipped

**Which stage type do not support setting file format option, instead, must specify file format and copy options as part of the COPY INTO <table> command.**

- Named External Stage
- Named Internal Stage
- Table Stage  
**(Correct)**
- User Stage  
**(Correct)**

**Explanation**

Refer Link - <https://docs.snowflake.com/en/user-guide/data-load-local-file-system-create-stage.html>

Question 80: Skipped

**Recreating the pipe will not remove the history of the files that are loaded using that pipe**

- TRUE
- FALSE  
**(Correct)**

### Explanation

The file loading metadata is associated with the pipe object rather than the table. Recreating the pipe removes the history of files loaded. Ensure that files already loaded by Snowpipe are not accidentally resubmitted to the pipe and loaded into the target table again. To view the query history for a table, query the COPY\_HISTORY function.

Question 81: Skipped

**When database replication is setup, then querying tables and views in a secondary database using Time Travel will always produce exact same results as that of when executing the same query in the primary database.**

- FALSE  
**(Correct)**
- TRUE

### Explanation

Querying tables and views in a secondary database using Time Travel can produce different results than when executing the same query in the primary database. Historical Data  
Historical data available to query in a primary database using Time Travel is not replicated to secondary databases. For example, suppose data is loaded continuously into a table every 10 minutes using Snowpipe, and a secondary database is refreshed every hour. The refresh operation only replicates the latest version of the table. While every hourly version of the table within the retention window is available for query using Time Travel, none of the iterative versions within each hour (the individual Snowpipe loads) are available. Data Retention Period The data retention period for tables in a secondary database begins when the secondary database is refreshed with the DML operations (i.e. changing or deleting data) written to tables in the primary database.

Question 82: Skipped

**In a VARIANT column, NULL values are stored as a string containing the word "null" not the SQL NULL value.**

- TRUE  
**(Correct)**
- FALSE

Question 83: Skipped

**What is the maximum single file size that is supported for Amazon S3, while unloading the data?**

- 10GB

- 2GB

- 5GB

(Correct)

- 1GB

### Explanation

By default, COPY INTO location statements separate table data into a set of output files to take advantage of parallel operations. The maximum size for each file is set using the MAX\_FILE\_SIZE copy option. The default value is 16000000 (16 MB) but can be increased to accommodate larger files. The maximum file size supported is 5 GB for Amazon S3, Google Cloud Storage, or Microsoft Azure stages.

Question 84: Skipped

### Which of the transformations are supported while loading data into table using the COPY command

- Truncating text strings that exceed the target column length

(Correct)

- Casts

(Correct)

- Column Reordering

(Correct)

- Column Omission

(Correct)

- Aggregations

Question 85: Skipped

### Which of the following options for loading data files from a stage using COPY command the fastest?

- By path (internal stages) / prefix (Amazon S3 bucket)

- Specifying a list of specific files to load.

(Correct)

- Using pattern matching to identify specific files by pattern

- There is no difference in the speed of loading for any option

### **Explanation**

Of the three options for identifying/specifying data files to load from a stage, providing a discrete list of files is generally the fastest; however, the FILES parameter supports a maximum of 1,000 files, meaning a COPY command executed with the FILES parameter can only load up to 1,000 files. Refer link - <https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html>

[Continue](#)

[Retake test](#)

**Which interfaces can be used to create and/or manage Virtual Warehouses?**

- A. The Snowflake Web Interface (UI)**
- B. SQL commands**
- C. Data integration tools**
- D. All of the Above**

**Answer:** D

**Which of the following objects can be cloned?**

**Choose 4 answers**

- A. Tables**
- B. Named File Formats**
- C. Schemas**
- D. Shares**
- E. Databases**
- F. Users**

**Answer:**

**ABCE**

**Which of the following DML commands isn't supported by Snowflake? Select one.**

- A. UPSERT**
- B. MERGE**
- C. UPDATE**
- D. TRUNCATE**

**TABLE Answer:**

**A**

**Increasing the maximum number of clusters in a Multi-Cluster Warehouse is an example of:**

**Select one.**

- A. Scaling rhythmically**
- B. Scaling max**
- C. Scaling out**
- D. Scaling up**

**up Answer:**

**C**

**Account-level storage usage can be monitored via: Select one.**

- A. The Snowflake Web Interface (UI) in the Databases section.**
- B. The Snowflake Web Interface (UI) in the Account -> Billing & Usage section**
- C. The Information Schema -> ACCOUNT\_USAGE\_HISTORY View**
- D. The Account Usage Schema -> ACCOUNT\_USAGE\_METRICS**

**View Answer: B**

**What is the minimum Snowflake edition that provides multi-cluster warehouses and up to 90 days of Time Travel? Select one.**

- A. Standard**
- B. Premier**
- C. Enterprise**
- D. Business Critical Edition**

**Answer: C**

**True or False: Pipes can be suspended and resumed**

- A. True**
- B. False**

**Answer: A**

**Which of the following best describes where Snowflake's metadata is stored? Select one.**

- A. Within the data files**
- B. Inside the Virtual Warehouses**
- C. In the Cloud Services Layer**
- D. Within the**

**drivers Answer: C**

**Which of the following items does the Cloud Services Layer manage?**

**Choose 4 answers**

- A. User authentication**
- B. Metadata**
- C. Query compilation and optimization**
- D. External blob storage**
- E. Data security**

**Answer: ABCE**

**Which statement best describes Snowflake tables? Select one.**

- A. Snowflake tables are logical representations of underlying physical data
- B. Snowflake tables are the physical instantiation of data loaded into Snowflake
- C. Snowflake tables require that clustering keys be defined to perform optimally
- D. Snowflake tables are owned by a

**user Answer: A**

**Which formats are supported for unloading data from Snowflake?**

**Choose 2 answers**

- A. Delimited (CSV, TSV, etc.)
- B. Avro
- C. JSON
- D. ORC

**Answer:**

**AC**

**Which of the following are examples of operations that require a Virtual Warehouse to complete, assuming no queries have been executed previously?**

**Choose 3 answers**

- A. MIN(< column value >>)
- B. COPY
- C. SUM(<< column value >>)
- D. UPDATE

**Answer: BCD**

**Why would a customer size a Virtual Warehouse from an X-Small to a Medium? Select one.**

- A. To accommodate more queries.
- B. To accommodate more users.
- C. To accommodate fluctuations in workload.

**D. To accommodate a more complex workload.**

**Answer: D**

**True or False: Fail-safe can be disabled within a Snowflake account.**

- A. True**
- B. False**

**Answer: B**

**True or False: Data Consumers must execute the SHARE REFRESH; command in order to view new or modified data in shared objects.**

- A. True**
- B. False**

**Answer: B**

**Which of the following statements about Data Sharing is true? Select one.**

- A. Accounts can share with other accounts regardless of Snowflake edition, without requiring help from Snowflake support (e.g. Enterprise, Business Critical Edition)**
- B. Data sharing can cross regions, but not cloud providers**
- C. The Data Consumer can only see objects in the Data Provider's source database that have been explicitly added to the Share**
- D. A Data Provider can only share with other Snowflake customers**

**Answer: A**

**True or False: Although a user can be granted multiple roles, each session has a single current role which determines current privileges.**

- A. True**
- B. False**

**Answer: A**

**True or False: All Snowflake table types include fail-safe storage.**

- A. True
- B. False

**Answer: B**

**True or False: Loading data via the COPY command only allows error handling on a per-file basis—all files are loaded in their entirety or marked as error, there are no partial loads.**

- A. True
- B. False

**Answer: B**

**What technique does Snowflake recommend for determining which Warehouse size to select?**

**Select one.**

- A. Always start with an X-Small and size up if the query doesn't complete in 2 minutes
- B. Experiment by running the same queries against Warehouses of multiple sizes
- C. Use the default Snowflake chooses for you
- D. Use the Warehouse Size Selection Wizard in the Snowflake Web Interface (UI)

**Answer: B**

**Query results are stored in the Result Cache for how long after they are last accessed, assuming no data changes have occurred?**

**Select one.**

- A. 1 Hour
- B. 3 Hours
- C. 12 hours
- D. 24 hours

**Answer: D**

**What is the minimum Snowflake edition that provides data sharing? Select one.**

- A. Standard
- B. Premier
- C. Enterprise

#### D. Business Critical Edition

Answer: A

The Query History in the Snowflake Web Interface (UI) is kept for approximately:

Select one.

- A. 60 minutes
- B. 24 hours
- C. 14 days
- D. 30 days
- E. 1 year

Answer: C

\* To run a multi-cluster Warehouse in maximized mode, a user would: Select one.

- A. Configure the Maximum Clusters setting to "Maximum"
- B. Manually turn on the additional clusters after starting the Warehouse
- C. Set the Minimum Clusters and Maximum Clusters settings to the same value
- D. Set the Minimum Clusters and Maximum Clusters settings to the different values

Answer: A

True or False: The longer the data retention period, the higher the resulting storage costs.

- A. True
- B. False

Answer: A

Which statement best describes "clustering"? Select one.

- A. Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions
- B. The database administrator must define the clustering methodology for each Snowflake table.
- C. The clustering key must be included on the COPY command when loading data into Snowflake.
- D. Clustering can be disabled within a Snowflake account.

Answer: A

Which of the following are options when creating a Virtual Warehouse?  
Choose 2 answers

- A. Auto-drop
- B. Auto-resize
- C. Auto-resume
- D. Auto-suspend

**Answer:** CD

**Each incremental increase in Virtual Warehouse size (e.g. Medium to Large) generally results in what?**

Select one.

- A. More micro-partitions
- B. Better query scheduling
- C. Double the numbers of servers in the compute cluster
- D. Higher storage costs

**Answer:** C

**What command is used to load files into an Internal Stage within Snowflake? Select one.**

- A. PUT
- B. COPY INTO
- C. TRANSFER.
- D. INSERT

**Answer:** A

**30. True or False: When Snowflake is configured to use Single Sign-On (SSO), Snowflake receives the usernames and credentials from the SSO service and loads them into the customer's Snowflake account.**

- A. True
- B. False

**Answer**

:

**B**

**What happens when a Data Provider revokes privileges to a Share on an object in their source database?**

Select one.

- A. The object immediately becomes unavailable for all Data Consumers
- B. Any additional data arriving after this point in time will not be visible to Data Consumers
- C. The Data Consumers stop seeing data updates and become responsible for storage charges for the object
- D. A static copy of the object at the time the privilege was revoked is created in the Data Consumers' accounts

**Answer: A**

**Which of the following statements are true of Virtual Warehouses? Select all that apply below.**

- A. Customers can change the size of the Warehouse after creation
- B. A Warehouse can be resized while running
- C. A Warehouse can be configured to suspend after a period of inactivity
- D. A Warehouse can be configured to auto-resume when new queries are submitted

**Answer: ABCD**

**What is the most granular object that the Time Travel retention period can be defined on? Select one.**

- A. Account
- B. Database
- C. Schema
- D. Table

**Answer: D**

**When a Pipe is recreated using the CREATE OR REPLACE PIPE command: Select one.**

- A. The Pipe load history is reset to empty
- B. The REFRESH parameter is set to TRUE
- C. Previously loaded files will be ignored
- D. All of the above

**Answer: A**

**True or False: Each worksheet in the Snowflake Web Interface (UI) can be associated with different roles, databases, schemas, and Virtual Warehouses.**

- A. True
- B. False

**Answer:**

**A**

**A Virtual Warehouse's auto-suspend and auto-resume settings apply to: Select one.**

- A. The primary cluster in the Virtual Warehouse
- B. The entire Virtual Warehouse
- C. The database the Virtual Warehouse resides in
- D. The queries currently being run by the Virtual Warehouse

**Answer: B**

**To run a Multi-Cluster Warehouse in auto-scale mode, a user would:**

Select one.

- A. Configure the Maximum Clusters setting to "Auto-Scale"
- B. Set the Warehouse type to "Auto"
- C. Set the Minimum Clusters and Maximum Clusters settings to the same value
- D. Set the Minimum Clusters and Maximum Clusters settings to the different values

Answer: D

A deterministic query is run at 8am, takes 5 minutes, and the results are cached. Which of the following statements are true?

Choose 2 answers

- A. The exact query will **ALWAYS** return the precomputed result set for the RESULT\_CACHE\_ACTIVE = time period
- B. The same exact query will return the precomputed results if the underlying data hasn't changed and the results were last accessed within the previous 24 hour period
- C. The same exact query will return the precomputed results even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period
- D. The "24 hour" timer on the precomputed results gets renewed every time the exact query is executed

Answer: BD

How would you best determine the size of the virtual warehouse used for a task?

- A. Root task may be executed concurrently (i.e. multiple instances), it is recommended to leave some margins in the execution window to avoid missing instances of execution
- B. Querying (SELECT) the size of the stream content would help determine the warehouse size. For example, if querying large stream content, use a larger warehouse size.
- C. If using the stored procedure to execute multiple SQL statements, it's best to test run the stored procedure separately to size the compute resource first.
- D. Since task infrastructure is based on running the task body on schedule, it's recommended to configure the virtual warehouse for automatic concurrency handling using Multi-duster warehouse (MCW) to match the task schedule.

Answer: D

**40. . True or False: The following command will fail: GRANT SELECT, INSERT, UPDATE, DELETE ON MYTABLE TO USER JOHN;**

- A. True
- B. False

Answer

: A

**41.. True or False: Reader Accounts incur no additional Compute costs to the Data Provider since they are simply reading the shared data without making changes.**

- A. True
- B. False

**Answer**

:

**B**

**The FLATTEN function is used to query which type of data in Snowflake? Select one.**

- A. Structured data
- B. Semi-structured data
- C. Both of the above
- D. None of the above

**Answer: B**

**What is the default File Format used in the COPY command if one is not specified?**

**Select one.**

- A. CSV
- B. JSON
- C. Parquet
- D. XML

**Answer: A**

**Which of the following are true of multi-cluster Warehouses? Select all that apply below.**

- A. A multi-cluster Warehouse can add clusters automatically based on query activity
- B. A multi-cluster Warehouse can automatically turn itself off after a period of inactivity
- C. A multi-cluster Warehouse can scale down when query activity slows
- D. A multi-cluster Warehouse can automatically turn itself on when a query is executed against it

**Answer: ABD**

**The fail-safe retention period is how many days? Select one.**

- A. 1 day
- B. 7 days
- C. 45 days
- D. 90 days

**Answer: B**

**Which of the following accurately represents how a table fits into Snowflake's logical container hierarchy? Select one.**

- A. Account -> Schema -> Database -> Table
- B. Account -> Database -> Schema -> Table
- C. Database -> Schema -> Table -> Account
- D. Database -> Table -> Schema -> Account

**Answer: B**

**True or False: Snowflake charges a premium for storing semi-structured data.**

- A. True
- B. Fals

**e Answer:**

**B**

**Which of the following connectors are available in the Downloads section of the Snowflake Web Interface (UI)?**

**Choose 2 answers**

- A. SnowSQL
- B. ODBC
- C. R
- D. HIVE

**Answer:**

**AB**

**True or False: When active, a Pipe requires a dedicated Virtual Warehouse to execute.**

- A. True
- B. Fals

**e Answer:**

**B**

**True or False: The PUT command can be run through the Snowflake Web Interface (UI).**

- A. True
- B. Fals

**e Answer:**

**B**

**True or False: Users are able to see the result sets of queries executed by other users that share their same role.**

- A. True
- B. False

**Answer: B**

**True or False: You can query the files in an External Stage directly without having to load the data into a table.**

- A. True
- B. Fals

**e Answer:**

**A**

**What is the recommended Snowflake data type to store semi-structured data like JSON?**

**Select one.**

- A. VARCHAR
- B. RAW
- C. LOB
- D. VARIANT

**T Answer: D**

**Which of the following statements are true about Schemas in Snowflake?**

**Choose 2 answers**

- A. A Schema may contain one or many Databases
- B. A Database may contain one or many Schemas
- C. A Schema is a logical grouping of Database Objects
- D. Each schema is contained within a

**Warehouse Answer: BC**

**True or False: Snowflake charges additional fees to Data Providers for each Share they create.**

- A. True
- B. False

**Answer: B**

**When can a Virtual Warehouse start running queries? Select one.**

- A. 12am-5am
- B. Only during administrator defined time slots
- C. When its provisioning is complete
- D. After

**replication Answer:**

**C**

**As a best practice, clustering keys should only be defined on tables of which minimum size?**

**Select one.**

- A. Multi-Kilobyte (KB) Range
- B. Multi-Megabyte (MB) Range
- C. Multi-Gigabyte (G8) Range
- D. Multi-Terabyte (TB)

**Range Answer: D**

**In which layer of its architecture does Snowflake store its metadata statistics? Select one.**

- A. Storage Layer
- B. Compute Layer
- C. Database Layer
- D. Cloud Services Layer

**Answer: D**

**If a Small Warehouse is made up of 2 servers/cluster, how many servers/cluster make up a Medium Warehouse?**

**Select one.**

- A. 4
- B. 16
- C. 32
- D. 128

**Answer: A**

**Select the three types of tables that exist within Snowflake.  
Choose 3 answers**

- A. Temporary
- B. Transient
- C. Provisional
- D. Permanent

**Answer: ABD**

**Which item in the Data Warehouse migration process does not apply in Snowflake?**

Select one.

- A. Migrate Users
- B. Migrate Schemas
- C. Migrate Indexes
- D. Build the Data Pipeline

**Answer: C**

**65. What would you do to execute a series of tasks that run in sequence and are dependent on the previous task execution?**

- A. Create a series of tasks based on a series of schedules. Example Create one task that runs every 1 minutes, and a following task that runs every 2 minutes and so on: CREATE TASK task1 .. SCHEDULE = '1 minute' CREATE TASK task2 ... SCHEDULE = '2 minute'
- B. When creating a task, indicate a predecessor task that it should run after. In this way, you can string together a number of tasks, all dependent on one task or each other.
- C. Create a stored procedure for each task: CREATE PROCEDURE my\_task1\_spl()... CREATE PROCEDURE my\_task2\_sp2(); Then create a simple tree by specifying the stored procedure my\_task\_spl as the predecessor that triggers the new task my\_task2 when run successfully CREATE TASK my\_task2 WAREHOUSE = mywh AFTER my\_task1\_spl AS ...
- D. This is currently not possible using

**Snowflake. Answer: B**

**The PUT command is used to place local files into: Select one.**

- A. An External Stage
- B. A Stage Zone
- C. A Transient Stage
- D. An Internal Stage

**Answer: D**

**What are the three layers that make up Snowflake's architecture?**

**Choose 3 answers**

- A. Compute
- B. Tri-Secret Secure
- C. Storage
- D. Cloud Services

**Answer: ACD**

**True or False: A table in Snowflake can only be queried using the Virtual Warehouse that was used to load the data.**

- A. True
- B. Fals

e

**Answer:**

B

**Snowflake provides two mechanisms to reduce data storage costs for short-lived tables.**

**These mechanisms are:**

**Choose 2 answers**

- A. Temporary Tables
- B. Transient Tables
- C. Provisional Tables
- D. Permanent Tables

**Answer: AB**

**Which statements are true of micro-partitions?**

**Choose 2 answers**

- A. They are approximately 16MB in size
- B. They are stored compressed only if COMPRESS=TRUE on Table
- C. They are immutable
- D. They are only encrypted in the Enterprise edition and above

**Answer: AC**

**True or False: It is best practice to define a clustering key on every table.**

- A. True
- B. Fals

e **Answer:**

B

**Which of the following objects is not covered by Time Travel? Select one.**

- A. Tables
- B. Schemas
- C. Databases
- D. Stages

**Answer: D**

**The Information Schema and Account Usage Share provide storage information for which of the following objects?**

**Choose 3 answers**

- A. Users
- B. Tables
- C. Databases
- D. Internal Stages

**Answer: BCD**

**Data storage for individual tables can be monitored using which commands and/or object(s)?**

**Choose 2 answers**

- A. SHOW TABLES;
- B. SHOW STORAGE BY TABLE;
- C. Information Schema -> TABLE\_STORAGE\_METRICS
- D. Information Schema -> TABLE\_HISTORY

**Answer: AC**

**True or False: A Virtual Warehouse consumes Snowflake credits even when inactive.**

- A. True
- B. Fals

**e Answer:**

**B**

**For a multi-cluster Warehouse, the number of credits billed is calculated on: Select one.**

- A. The number of queries that ran using the Warehouse.
- B. The size of the Warehouse and the number of clusters that ran within a given time period.
- C. The size of the Warehouse and the maximum number of clusters configured for the Warehouse.

**D. The number of users who accessed the Warehouse.**

**Answer: B**

**What is the maximum compressed row size in Snowflake? Select one.**

- A. 8KB
- B. 16MB
- C. 50MB
- D. 4000GB

**Answer: B**

**What are the three things customers want most from their enterprise data warehouse solution?**

**Choose 3 answers**

- A. On-premise availability
- B. Simplicity
- C. Open source based
- D. Concurrency
- E. Performance

**Answer: BDE**

**Which of the following statements are true of Snowflake releases:**

**Choose 2 answers**

- A. They happen approximately weekly
- B. They roll up and release approximately monthly, but customers can request early release application
- C. During a release, new customer requests/queries/connections transparently move over to the newer version
- D. A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and the customer is upgraded [ \ ]

**Answer: AC**

**True or False: A customer using SnowSQL I native connectors will be unable to also use the Snowflake Web Interface (UI) unless access to the UI is explicitly granted by support.**

- A. True
- B. False

**e Answer:**

**B**

**Which of the following connectors allow Multi-Factor Authentication (MFA) authorization when connecting?**

Select all the below that apply.

- A. JDBC
- B. SnowSQL
- C. Snowflake Web Interface (UI)
- D. ODBC
- E. Python

**Answer:**

**ABCDE**

**Which of the following commands are not blocking operations?**

Choose 2 answers

- A. UPDATE
- B. INSERT
- C. MERGE
- D. COPY

**Answer:**

**BD**

**True or False: Micro-partition metadata enables some operations to be completed without requiring Compute.**

- A. True
- B. Fals

**e Answer:**

**A**

**Select the different types of Internal Stages:**

Choose 3 answers

- A. Named Stage
- B. User Stage
- C. Table Stage
- D. Schema Stage

**Answer: ABC**

**True or False: Snowpipe via REST API can only reference External Stages as source.**

- A. True
- B. Fals

**e Answer:**

**B**

**Snowflake Virtual Warehouses are part of which layer of the Snowflake architecture?**  
**Select one.**

- A. Compute Layer
- B. Storage Layer
- C. Database Layer
- D. Cloud Services

**Layer Answer: A**

**Which of the following terms best describes Snowflake's database architecture? Select one.**

- A. Columnar shared nothing
- B. Shared disk
- C. Multi-cluster, shared data
- D. Cloud-native shared

**memory Answer: C**

**Which of the following roles is recommended to be used to create and manage users and roles?**

**Select one.**

- A. SYSADMIN
- B. SECURITYADMIN
- C. PUBLIC
- D. ACCOUNTADMI

**N Answer: B**

**True or False: Every object is owned by one and only one role; therefore a user must be using that role to access the object.**

- A. True
- B. False

**Answer: B**

**What are the two ways Snowflake offers to create Virtual Warehouses?**

**Choose 2 answers**

- A. The Snowflake Web Interface (UI) using the Create button in the Warehouses tab
- B. The CREATE WAREHOUSE command
- C. The CREATE VIRTUAL\_WAREHOUSE command
- D. Asking Snowflake support to create the Warehouse for

**you Answer: AB**

**93. Which of the following is an example of an operation that can be completed without requiring Compute, assuming no queries have been executed previously? Select one.**

- A. SELECT SUM(ORDER\_AMT) FROM SALES;
- B. SELECT AVG(ORDER\_QTY) FROM SALES;
- C. SELECT MIN(ORDER\_AMT) FROM SALES;
- D. SELECT ORDER\_AMT \* ORDER\_QTY FROM

**SALES; Answer: C**

**Which of the following commands sets the Virtual Warehouse for a session? Select one.**

- A. COPY WAREHOUSE FROM <<config file>>;
- B. SET WAREHOUSE = <<warehouse name>>;
- C. USE WAREHOUSE <<warehouse name>>;
- D. USE VIRTUAL\_WAREHOUSE <<warehouse

**name>>; Answer: C**

**95. True or False: A third-party tool that supports standard JDBC or ODBC but has no Snowflake-specific driver will be unable to connect to Snowflake.**

- A. True
- B. Fals

**e Answer:**

**B**

**96. . Snowflake recommends, as a minimum, that all users with the following role(s) should be enrolled in Multi-Factor Authentication (MFA):**

Select one.

- A. SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN
- B. SECURITYADMIN, ACCOUNTADMIN, SYSADMIN
- C. SECURITYADMIN, ACCOUNTADMIN
- D. ACCOUNTADMI

**N Answer: D**

**Which of the following statements is true of Snowflake micro-partitioning? Select one.**

- A. Micro-partitioning has been known to introduce data skew
- B. Micro-partitioning: requires a partitioning schema to be defined up front
- C. Micro-partitioning is transparently completed using the ordering that occurs when the data is inserted/loaded
- D. Micro-partitioning can be disabled within a Snowflake

**account Answer: C**

**True or False: During data unloading, only JSON and CSV files can be compressed.**

- A. True
- B. Fals

**e Answer:**

**B**

**The number of queries that a Virtual Warehouse can concurrently process is determined by:**

**Choose 2 answers**

- A. The complexity of each query
- B. The CONCURRENT\_QUERY\_LIMIT parameter set on the Snowflake account
- C. The size of the data required for each query
- D. The tool that is executing the

**query Answer: AC**

**True or False: Snowflake requires that a single user run its ETL, reporting, and data science workloads on the same Virtual Warehouse.**

- A. True
- B. False

**Answer: B**

**True or False: You can define multiple columns within a clustering key on a table.**

- A. True

**B. False**

**Answer:**

**A**

**True or False: A 4X-Large Warehouse may, at times, take longer to provision than a X-Small Warehouse.**

- A. True
- B. False

**Answer: A**

**True or False: You can resize a Virtual Warehouse while queries are running.**

- A. True
- B. False

**Answer:**

**A**

**True or False: Fail-safe can be disabled within a Snowflake account.**

- A. True
- B. False

**Answer:**

**B**

**True or False: Some queries can be answered through the metadata cache and do not require an active Virtual Warehouse.**

- A. True
- B. False

**Answer:**

**A**

**When loading data into Snowflake, the COPY command supports**

**Choose 2 answers**

- A. Joins
- B. Filters
- C. Data type conversions
- D. Column reordering
- E. Aggregates

**Answer: CD**

**How would you execute a series of SQL statements using Task?**

- A. **Include the SQL statements in the body of the task** CREATE TASK mytask .. AS  
INSERT INTO target1 SELECT .. FROM stream\_s1 WHERE ..  
SELECT .. FROM stream\_s1 WHERE ..

- B. A stored procedure can have only one DML statement per stored procedure invocation and therefore you should sequence stored procedures calls in the task definition CREATE TASK mytask..... AS call stored\_proc1(); call stored\_proc2();
- C. A stored procedure executing multiple SQL statements and invoke the stored procedure from the task. CREATE TASK mytask .....AS call stored\_proc\_multiple\_statements\_inside();
- D. Create a task for each SQL statement (e.g. resulting in task1, task2, etc) and string the series of SQL statements by having a control task calling task 1, task 2 sequentially.

**Answer: C**

**Which of the following statements would be used to export/unload data from Snowflake?**  
**Select one.**

- A. COPY INTO @stage
- B. EXPORT TO @stage
- C. INSERT INTO @stage
- D. EXPORT\_TO\_STAGE(stage => @stage, select => 'select \* from t1');

**Answer: A**

**When should you consider disabling auto-suspend for a Virtual Warehouse?**

**Choose 2 answers**

- A. When users will be using compute at different times throughout a 24/7 period
- B. When managing a steady workload
- C. When the compute must be available with no delay or lag time
- D. When you don't want to have to manually turn on the Warehouse each time a user needs it

**Answer: AC**

**Which of the following statements is true of Virtual Warehouse resizing?**

**Select one.**

- A. A resize requires the Warehouse to be in suspended status.
- B. A resize can be completed at any time.
- C. A resize will affect running, queued, and new queries.
- D. A resize can only be completed once per day.

**Answer: B**

**What is the best approach for loading data into Snowflake?**

**Select one.**

- A. Merge files into one big file
- B. Many big files and a big Virtual Warehouse size
- C. Many small files with a high level of parallelism
- D. Few files and a big Virtual Warehouse size

**Answer: C**

**When a Pipe is recreated using the CREATE OR REPLACE PIPE command:**

Select one.

- A. The Pipe load history is reset to empty
- B. The REFRESH parameter is set to TRUE
- C. Previously loaded files will be ignored
- D. All of the above

**Answer: A**

**Which of the following languages can be used to implement Snowflake User Defined Functions (UDFs)?**

Choose 2 answers

- A. Java
- B. Javascript
- C. SQL
- D. Python

**Answer: BC**

**Which of the following statements are true?**

Choose 3 answers

- A. The COPY command must reference a Stage or cloud storage location
- B. A named File Format is optional when using the COPY command
- C. A File format option used in the COPY command will override the one in Stage
- D. A File Format object must be defined when creating a Stage

**Answer: ABC**

**True or False: A single schema can exist in more than one database.**

- A. True
- B. False

**ANSWER:B**

**Which types of data aren't natively supported in Snowflake?**

Choose 2 answers

- A. CSV
- B. XML
- C. JPEG
- D. MP4

**ANSWER:**

**CD**

**What is the optimal file size for non-continuous bulk data loading into Snowflake?**

**Select one.**

**A. 100 Kilobytes-1 Megabyte**

- B. 1 Megabyte - 10 Megabytes**
- C. 10 Megabytes - 100 Megabytes**
- D. 1 Gigabyte or above**

**ANSWER:C**

**Which object allows you to limit the number of credits consumed within a Snowflake account?**

**SELECT ONE.**

- A. Account Usage Tracking**
- B. Resource Monitor**
- C. Warehouse Limit Parameter**
- D. Credit Consumption Tracker**

**ANSWER:B**

**True or False: The user has to specify which cluster a query will run on in a multi-cluster Warehouse.**

- A. TRUE**
- B. FALSE**

**ANSWER:B**

**If auto-suspend is enabled for a Virtual Warehouse, the Warehouse is automatically suspended when:**

**SELECT ONE.**

- A. All Snowflake sessions using the Warehouse are terminated.**
- B. The last query using the Warehouse completes.**
- C. There are no users logged into Snowflake.**
- D. The Warehouse is inactive for a specified period of time.**

**ANSWER:D**

**Which is true of Snowflake network policies? A Snowflake network policy:**

**CHOOSE 2 ANSWERS**

- A. Is available to all Snowflake Editions**
- B. Is only available to customers with Business Critical Edition**
- C. Restricts or enables access to specific IP addresses**
- D. Is activated using an "ALTER DATABASE" command.**

**ANSWER:AC**

**Which of the following is true of Snowpipe via REST API?**

**CHOOSE 2 ANSWERS**

- A. You can only use it on Internal Stages**
- B. All COPY INTO options are available during pipe creation**
- C. Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands**
- D. Snowpipe keeps track of which files it has loaded**

**ANSWER:CD**

- A. TRUE**
- B. FALSE**

**ANSWER:B**

**When sizing a Virtual Warehouse up to the next available t-shirt size, the number of servers increase by a factor of what?**

**SELECT ONE.**

- A. 1**
- B. 2**
- C. 4**
- D. 8**

**ANSWE**

**R:B**

**Which of the following operations can be completed without requiring Compute?**

**SELECT ONE.**

- A. INSERT**
- B. UPDATE**
- C. DROP**
- D. MERGE**

**ANSWER:C**

**What is the default value in the Snowflake Web Interface (UI) for auto suspending a Virtual Warehouse?**

**SELECT ONE.**

- A. 1 minute**
- B. 5 minutes**
- C. 10 minutes**
- D. 15 minutes**

**ANSWER:C**

**When creating a user it is advisable to:**

**SELECT 2 ANSWERS**

- A. Set the user to be initially disabled**
- B. Force an immediate password change**
- C. Set a default role for the user**
- D. Set the number of minutes to unlock to 15 minutes**
- E. Set the user's access to expire within a specified timeframe**

**ANSWER:BC**

**True or False: Bulk unloading of data from Snowflake supports the use of a SELECT statement.**

- A. TRUE**
- B. FALSE**

**ANSWER:A**

**True or False: When a new Snowflake object is created, it is automatically owned by the user who created it.**

- A. TRUE**
- B. FALSE**

**ANSWER:B**

**What is the minimum Snowflake edition that customers planning on storing protected information in Snowflake should consider for regulatory compliance? Select one.**

- A. Standard**
- B. Premier**
- C. Enterprise**
- D. Business Critical Edition**

**ANSWER:C**

**How many shares can be consumed by a single Data Consumer?**

**SELECT ONE.**

- A. 1**
- B. 10**
- C. 100z but can be increased by contacting**
- D. Unlimited**

**ANSWER:D**

**True or False: Snowflake's data warehouse was built from the ground up for the cloud in lieu of using an existing database or a platform, like Hadoop, as a base.**

- A. TRUE
- B. FALSE

**ANSWER:B**

**True or False: Snowflake bills for a minimum of five minutes each time a Virtual Warehouse is started.**

- A. TRUE
- B. FALSE

**ANSWER:B**

**What parameter controls if the Virtual Warehouse starts immediately after the CREATE WAREHOUSE statement?**

**SELECT ONE.**

- A. INITIALLY\_SUSPENDED = TRUE/FALSE
- B. START\_AFTER\_CREATE = TRUE/FALSE
- C. START\_TIME = 60 11 (seconds from now)
- D. STARTTIME = CURRENT\_DATE()

**ANSWER:A**

**Which of the following are main sections of the top navigation of the Snowflake Web Interface (UI)?**

**CHOOSE 3 ANSWERS**

- A. Databases
- B. Tables
- C. Warehouses
- D. Worksheets

**ANSWER:ACD**

**When scaling up Virtual Warehouses by increasing Virtual Warehouse t-shirt size, you are primarily scaling for improved:**

**SELECT ONE.**

- A. Concurrency

**B. Performance**

**ANSWER:B**

**The default Time Travel retention period is:**

**SELECT ONE.**

- A. 1 day
- B. 7 day
- C. 45 days
- D. 90 days

**ANSWER:A**

**Which of the following statements is true of zero-copy cloning?**

**SELECT ONE.**

- A. Zero-copy clones increase storage costs as cloning the table requires storing its data twice
- B. All zero-copy clone objects inherit the privileges of their original objects
- C. Zero-copy cloning is licensed as an additional Snowflake feature
- D. At the instance/instant a clone is created, all micro-partitions in the original table and the clone are fully shared

**ANSWER:D**

**Which of the following are common use cases for zero-copy cloning?**

**CHOOSE 3 ANSWERS**

- A. Quick provisioning of Dev and Test/QA environments
- B. Data backups
- C. Point in time snapshots
- D. Performance optimization

**ANSWER:ABC**

**True or False: A single database can exist in more than one Snowflake account.**

- A. TRUE
- B. FALSE

**ANSWER:B**

**True or False: ACCOUNTADMINs are able to see the result sets of queries executed by all users within a Snowflake account.**

- A. TRUE
- B. FALSE

**ANSWER:B**

**True or False: Reader Accounts are able to extract data from shared data objects for use outside of Snowflake.**

- A. TRUE
- B. FALSE

**ANSWER:B**

**True or False: Once created, a micro-partition will never be changed.**

- A. TRUE
- B. FALSE

**ANSWER:A**

**Which of the following statements describes a benefit of Snowflake's separation of compute and storage? Select all the below that apply.**

- A. Growth of storage and compute are tightly coupled together
- B. Storage expands without the requirement to add more compute
- C. Compute can be scaled up or down without the requirement to add more storage
- D. Multiple compute clusters can access stored data without contention

**ANSWER:BCD**

**True or False: Multi-Factor Authentication (MFA) is available only to customers with Enterprise edition and above.**

- A. TRUE
- B. FALSE

**ANSWER:B**

**Which statement best describes Snowflake tables?**

Select one.

- A. Snowflake tables are logical representations of underlying physical data
- B. Snowflake tables are the physical instantiation of data loaded into
- C. Snowflake tables require that clustering keys be defined to perform

optimally

**D. Snowflake tables are owned by a user**

**Answer: A**

**What is the minimum Snowflake edition that provides data sharing?**

**Select one.**

- A. Standard**
- B. Premier**
- C. Enterprise**
- D. Business Critical Edition**

**ANSWER:A**

**How would you best determine the size of the virtual warehouse used for a task?**

- A. Root task may be executed concurrently (i.e. multiple instances), it is recommended to leave some margins in the execution window to avoid missing instances of execution**
- B. Querying (SELECT) the size of the stream content would help determine the warehouse size. For example, if querying large stream content, use a larger warehouse size.**
- C. If using the stored procedure to execute multiple SQL statements, it's best to test run the stored procedure separately to size the compute resource first.**
- D. Since task infrastructure is based on running the task body on schedule, it's recommended to configure the virtual warehouse for automatic onnoiirrpnrv handling ncinn Mi ilti-rli Ktpr warphnnv fMCW'l to matrh thp ta<:k <rhpdnlp**

**ANSWER:D**

**True or False: When a user creates a role, they are initially assigned ownership of the role and they maintain ownership until it is transferred to another user.**

- A. True**
- B. False**

**ANSWER:**

**B**

**True or False: An active warehouse is required to run a COPY INTO statement.**

- A. True**
- B. False**

**ANSWER:A**

**Fail-safe is unavailable on which table types?**

**Choose 2 answers**

- A. Temporary**
- B. Transient**
- C. Provisional**

**D. Permanent**

**Answer: AB**

**How are Snowpipe charges calculated?**

**Select one.**

- A. Per-second/per Warehouse size**
- B. Per-second/per-core granularity**
- C. Number of Pipes in account**
- D. Total storage bucket size**

**Answer: B**

**True or False: Multi-Factor Authentication (MFA) in Snowflake is only supported in conjunction with Single Sign-On (SSO).**

- A. True**
- B. False**

**Answer:**

**B**

**Which are true of Snowflake roles?**

**Choose 2 answers**

- A. All grants to objects are given to roles, and never to users**
- B. In order to do DML/DDL, a user needs to have selected a single role that has that specific access to the object and operation**
- C. The PUBLIC role controls all other roles**
- D. Roles are a subset of users and users own objects in Snowflake**

**Answer: AB**

**True or False: Snowflake allows its customers to directly access the micro-partition files that make up its tables.**

- A. True**
- B. False**

**Answer:**

**B**

**Which of the following statements about data sharing are true?**

**Choose 2 answers**

- A. New objects created by a Data Provider are automatically shared with existing Data Consumers & Reader Accounts**
- B. All database objects can be included in a shared database**
- C. Reader Accounts are created and funded by Data Providers**

**D. Shared databases are read-only**

**Answer: CD**

**Credit Consumption by the Compute Layer (Virtual Warehouses) is based on:**

**Choose 2 answers**

- A. Number of users**
- B. Warehouse size**
- C. Amount of data processed**
- D. # of clusters for the Warehouse**

**Answer: BD**

**Select one.**

**How are Snowpipe charges calculated?**

- A. Per-second/per Warehouse size**
- B. Per-second/per-core granularity**
- C. Number of Pipes in account**
- D. Total storage bucket size**

**Answer: B**

**What is the recommended method for loading data into Snowflake?**

**Select one.**

- A. Load row by row**
- B. Load data in batch**
- C. Load data by writing it in the Snowflake Web Interface (UI)**
- D. Load data via frequent, single row DML's**

**Answer: B**

**Which of the following are best practices for users with the SYSADMIN/ACCOUNTADMIN roles?**

**Choose 3 answers**

- A. Their default role should be set to SYSADMIN (the lower of the two)**
- B. They should not set up Multi-Factor Authentication (MFA)-as administrators they may need to change the MFA settings and those enrolled in MFA are unable to do so**
- C. They should only access and 'step into' the ACCOUNTADMIN role temporarily, as needed to complete a specific operation**
- D. They should ensure all database objects in the account are owned by the ACCOUNTADMIN role**
- E. They should use the SYSADMIN role to perform administrative work on database objects**

**Answer: ACE**

**True or False:** It is possible to use the VALIDATE option or operation to view all of the records that failed to load on a previous run.

- A. True
- B. False

**Answer:**

**A**

**True or False:** AWS Private Link provides a secure connection from the Customer's on-premise data center to the Snowflake Virtual Private Cloud (VPC).

- A. True
- B. False

**Answer:**

**B**

**True or False:** The COPY command supports regular expressions to filter the population of files to be loaded.

- A. True
- B. False

**Answer:**

**A**

**True or False:** Data Providers can share data with only one Data Consumer.

- A. True
- B. False

**Answer:**

**B**

**ACCOUNTADMIN** can execute which of the following commands on existing users?

Select one.

- A. Can SHOW users, DESCRIBE a given user or ALTER or DROP a user
- B. Can DEFINE users, DESCRIBE a given user or ALTER or DELETE a user
- C. Can SHOW users, INDEX a given user or ALTER or DELETE a user
- D. Can SHOW users, DEFINE a given user or ALTER, DROP or MODIFY a user

**Answer:** A

**True or False:** Data Providers can share data with only one Data Consumer.

- A. True
- B. False

**Answer:**

**B**

**What is the minimum duration charged when starting a virtual warehouse?**

Select one.

- A. 1 second
- B. 1 minute
- C. 1 hour
- D. 1 day

**Answer:** B

**Which of the following statements apply to Snowflake In terms of security?**

Choose 2 answers

- A. Snowflake leverages a role-based access control model
- B. Snowflake requires you to configure an IAM user in order to connect to the database
- C. All data in Snowflake Is encrypted
- D. Snowflake can run within a customer's own Virtual Private Cloud (VPC)

**Answer:** AC

**Which of the following statements is true of data loading?**

**Select one.**

- A. Resizing the Virtual Warehouse from X-Small to Small will process a single file twice as fast
- B. The ideal file size for loading is 16MB to match micro-partition size
- C. Many files In the 10-100MB range tend to land in the "sweet spot" for load parallelism
- D. Once loaded, there is no option to force a reload of an already loaded file

**Answer: C**

**SnowSQL can unload query results to a local file system in the following file formats:**

**Choose 3 answers**

- A. JSON
- B. Parquet
- C. ORC
- D. CSV and any other delimited formats

**Answer: ABD**

**True or False: Snowflake supports federated authentication in all editions.**

- A. True
- B. False

**Answer:**

**A**

**Which of the following are use cases that fit well in Snowflake?**

**Choose 3 answers**

- A. Data Lake
- B. Enterprise Data Warehouse
- C. Online Transaction Processing
- D. Datamart Consolidation

**Answer: ABD**

**True or False: It Is possible to unload structured data to semi-structured formats such as JSON and Parquet.**

- A. True
- B. False

**Answer:**

**A**

**True or False: When you create a custom role, it is a best practice to immediately grant that role to ACCOUNTADMIN.**

- A. True
- B. False

**Answer:**

**B**

**Snowflake provides a mechanism for Its customers to override its natural clustering algorithms. This method is:**

**Select one.**

- A. Micro-partitions
- B. Clustering keys
- C. Key partitions

**D. Clustered partitions**

**Answer: B**

**True or False: A Snowflake account Is charged for data stored in both Internal and External Stages.**

- A. True
- B. False

**Answer:**

**B**

**What are the two ways to create and manage Data Shares in Snowflake?**

**Choose 2 answers**

- A. Via the Snowflake Web Interface (UI)
- B. Via the DATA\_SHARE=TRUE parameter
- C. Via SQL commands
- D. Via Virtual Warehouses

**Answer: AC**

**What is the lowest Snowflake edition that offers Time Travel up to 90 days?**

**Select one.**

- A. Standard Edition
- B. Premier Edition
- C. Enterprise Edition
- D. Business Critical Edition

**Answer: C**

**True or False: Snowflake enforces unique, primary key, and foreign key constraints during DML operations.**

- A. True
- B. False

**Answer:**

**B**

**True or False: Query ID's are unique across all Snowflake deployments and can be used in communication with Snowflake Support to help troubleshoot issues.**

- A. True
- B. False

**Answer:**

**A**

**True or False: Snowflake's Global Services Layer gathers and maintains statistics on all columns In all micro-partitions.**

- A. True

**B. False**

**Answer:**

**A**

**Which type of table corresponds to a single Snowflake session?**

Select one.

- A. Temporary
- B. Transient
- C. Provisional
- D. Permanent

**Answer:** A

**True or False: Data in fail-safe can be deleted by a user or the Snowflake team before it expires.**

- A. True
- B. False

**Answer:**

**B**

**True or False: Every query should be expected to execute faster on a larger Virtual Warehouse than on a smaller Virtual Warehouse.**

- A. True
- B. False

**Answer:**

**B**

**Which of the following are valid approaches to loading data into a Snowflake table? Select all the below that apply.**

- A. Bulk copy from an External Stage
- B. Continuous load using Snowpipe REST API
- C. The Snowflake Web Interface (UI) data loading wizard
- D. Bulk copy from an Internal Stage

**Answer:** ABC

**When scaling out by adding clusters to a multi-duster warehouse, you are primarily scaling for improved:**

Select one.

- A. Concurrency
- B. Performance

**Answer:** A

**Increasing the size of a Virtual Warehouse from an X-Small to an X-Large is an example of:**

Select one.

- A. Scaling rhythmically
- B. Scaling max
- C. Scaling out
- D. Scaling up

**Answer:** D

**Which statement best describes Snowflake tables?**

**Select one.**

- A. Snowflake tables are logical representations of underlying physical data
- B. Snowflake tables are the physical instantiation of data loaded into Snowflake
- C. Snowflake tables require that clustering keys be defined to perform optimally
- D. Snowflake tables are owned by a user

**Answer: A**

**True or False: When a data share Is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data Consumers.**

- A. True
- B. False

**Answer:**

**B**

**Which of the following statements is true of Snowflake?**

**Select one.**

- A. It was built specifically for the cloud
- B. It was built as an on-premises solution and then ported to the cloud
- C. It was designed as a hybrid database to allow customers to store data either on premises or in the cloud
- D. It was built for Hadoop architecture
- E. It's based on an Oracle Architecture

**Answer: A**

**True or False: When a data share is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data Consumers.**

- A. True
- B. False

**Answer:**

**A**

**Snowflake supports External Stages against which Cloud Platforms?**

**Choose 3 answers**

- A. Amazon Web Services (AWS)
- B. Docker
- C. IBM Cloud
- D. Microsoft Azure
- E. Google Cloud Platform
- F. Oracle Cloud

**Answer: ADE**

**What services does Snowflake automatically provide for customers that they may have been responsible for with their on-premise system? Select all the below that apply.**

- A. Installing and configuring hardware
- B. Patching software
- C. Physical security
- D. Maintaining metadata and statistics

**Answer: ABCD**

**True or False: It is possible for a user to run a query against the query result cache without requiring an active Warehouse.**

- A. True
- B. False

**Answer:**

**A**

**Which Snowflake object enables loading data from files as soon as they are available in a cloud storage location?**

Select one.

- A. Pipe
- B. External Stage
- C. File Format
- D. VARIANT

**Answer: A**

**True or False: A Virtual Warehouse can only be resized while suspended.**

- A. True
- B. False

**Answer:**

**B**

**What privileges are required to create a task?**

- A. The GLOBAL privilege CREATE TASK is required to create a new task.
- B. Tasks are created at the Application level and can only be created by the Account Admin role.
- C. Many Snowflake DDLs are metadata operations only, and CREATE TASK DDL can be executed without virtual warehouse requirement or t; specific grants.
- D. The role must have access to the target schema and the CREATE TASK privilege on the schema itself.

**Answer: C**

**Which of the following statements is true of data loading?**

Select one.

- A. Resizing the Virtual Warehouse from X-Small to Small will process a single file twice as fast
- B. The ideal file size for loading is 16MB to match micro-partition size
- C. Many files in the 10-100MB range tend to land in the "sweet spot" for load parallelism
- D. Once loaded, there is no option to force a reload of an already loaded file

**Answer: C**

**Question**

**True or False: The COPY command must specify a File Format in order to execute.**

- A. True
- B. False

**Answer:**

**B**

**Question**

**True or False:** It is possible to query data from an Internal or named External stage without loading the data into Snowflake.

- A. True
- B. False

**Answer:**

**A**

**Question**

**On which of the following cloud platforms can a Snowflake account be hosted?**

**Choose 2 answers**

- A. Amazon web Services
- B. Private Virtual Cloud
- C. Grade Cloud
- D. Microsoft Azure Cloud

**Answer: AD**

**Question**

**A client uses a PUT command to upload a new file to an Internal Stage. At which point is the data encrypted?**

**Select one.**

- A. When it reaches the virtual warehouse
- B. When it gets micro-partitioned and stored in Cloud Storage
- C. On the client's machine before being sent
- D. Clients must encrypt files before sending to Snowflake

**Answer: D**

**Question**

**Snowflake is designed for which type of workloads?**

**Choose 2 answers**

- A. OLAP (Analytics) workloads
- B. OLTP (Transactional) workloads
- C. Concurrent workloads
- D. On-premise workloads

**Answer: AC**

**Question**

**Which of the following are valid Snowflake Virtual Warehouse Scaling Policies?**

**Choose 2 answers**

- A. Custom
- B. Economy
- C. Optimized
- D. Standard

**Answer: BD**

**Question**

**On which of the following cloud platforms can a Snowflake account be hosted?**

**Choose 2 answers**

- A. Amazon Web Services
- B. Private Virtual Cloud
- C. Oracle Cloud
- D. Microsoft Azure Cloud

**Answer:AD**

**Question**

**The PUT command:**

**Choose 2 answers**

- A. Automatically creates a File Format object
- B. Automatically uses the last Stage created

**C. Automatically compresses files using Gzip**

**D. Automatically encrypts files**

**Answer: CD**

**Question**

**True or False:** Loading data into Snowflake requires that source data files be no larger than 16MB.

- A. True
- B. False

**Answer:**

B

**Question**

**Which of the following are valid approaches to loading data into a Snowflake table? Select all the below that apply.**

- A. Bulk copy from an External Stage
- B. Continuous load using Snowpipe REST API
- C. The Snowflake Web Interface (UI) data loading wizard
- D. Bulk copy from an Internal Stage

**Answer:** ABCD

**Question**

**. Which of the following statements are true of Snowflake releases:**

**Choose 2 answers**

- A. They happen approximately weekly
- B. They roll up and release approximately monthly, but customers can request early release application
- C. During a release, new customer requests/queries/connections transparently move over to the newer version
- D. A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and the customer is upgraded

**Answer:** AC

**Question**

**Which of the following statements are true of VALIDATION\_MODE in Snowflake?**

**Choose 2 answers**

- A. The VALIDATION\_MODE option is used when creating an Internal Stage
- B. VALIDATION\_MODE=RETURN\_ALL\_ERRORS is a parameter of the COPY command
- C. The VALIDATION\_MODE option will validate data to be loaded by the COPY statement while completing the load and will return the rows that could not be loaded without error
- D. The VALIDATION\_MODE option will validate data to be loaded by the COPY statement without completing the load and will return possible errors

**Answer:** CD

**Question**

**Which of the following statements are true of Snowflake data loading?**

**Choose 3 answers**

- A. VARIANT "null" values are not the same as SQL NULL values
- B. It is recommended to do frequent, single row DMLs
- C. It is recommended to validate the data before loading into the Snowflake target table

**D. It is recommended to use staging tables to manage MERGE statements**

**Answer: ACD**

**Question**

**True or False:** An active warehouse is required to run a COPY INTO statement.

- A. True
- B. False

**Answer:**A

**Question**

**What privileges are required to execute a task?**

- A. To execute a task, you need the EXECUTE TASK.
- B. A task is an account-level object and can only be executed by the Account Admin role.
- C. Tasks run automatically and do run under any role.
- D. Tasks are stored in a special database called the utility database and require EXECUTE TASK –privilege in the utility database.

**Answer:**A

**Question**

**A role is created and owns 2 tables. This role is then dropped. Who will now own the two tables?**

**Select one.**

- A. The tables are now orphaned
- B. The user that deleted the role
- C. SYSADMIN
- D. The assumed role that dropped the role

**ANSWER:** C

**Question**

**True or False: When a data share is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data Consumers.**

- A. True
- B. False

**Answer:**

**B**

**Question**

**Snowflake collects which of the following metadata on each of its micro partitions? Select all the below that apply.**

- A. Range of values
- B. Number of distinct values
- C. MIN/MAX values
- D. NULL count

**Answer:** ABC

Background:

- Have been focused on reading <https://docs.snowflake.com/en/index.html>, have only minimal hands-on experience. Before this, have no exposure to database.
- Have very basic sql knowledge from self learning.
- Took first 3 mock tests on Udemy, all passed, and these tests are easy for me.
- Half way through Sara's SF training, contents that he hasn't covered yet is blurry to me just from reading documentation.

Result:

- Fail. Got 640 points, and 750 points are needed to pass.

My thinking:

- Compared with first 3 mock tests on Udemy, the real test questions are much less straightforward. (only < 50% of questions are straightforwardly defined on documentation, while the rest needs deeper understanding of how things actually work).
- To achieve required level of understanding, hands on lab is a must for my background.

Questions I remember:

1. Two languages that are used to execute external functions/UDFs/(some sort of function )? python/java/sql/javascript?
2. A single database can exist in more than one Snowflake account? (I confuse this with data sharing across accounts, so I am not sure)
3. A single schema can exist in more than one databases? (I confuse this with data sharing across accounts, so I am not sure. So what about tables?)
4. How to run a task? (detailed explanation) (selections contains procedures definition)
5. What is external function used for?
6. What is the maximum size for each row for any table in SF?
7. Up to at least which level of table size, is only recommended to use clustering key? kb(mb/tb)?
8. Which statement best describes column masking? (each description is very very long and none seemed correct to me, and I do understand column masking...)
9. You cannot run "PUT" in SF UI? yes
10. About Micro-partition - check all that is correct. Size of each micropartition, etc
11. How often is SF's new version, etc. weekly
12. Where to check data storage usage
13. Which sql command if executed twice, may not need computing resources? Sql command that calculates average/multiplication of two rows/minimum of a row/etc? - minimum because it is stored in metadata in cloud services layer
14. Which stage can be modified and dropped? external/internal/user/table stage?
15. Which role is recommended to manage roles and privileges? Account admin/security admin/public/ etc. selections didn't include useradmin
16. Which sql command doesn't have blocking? Choose two. select/insert/merge/update?
17. Snowflake doesn't support which command? upsert/select/XXX/truncate table?

18. Snowpipe via rest api only support staging files from external stages?
19. Grant XXXX privilege to user XXX will return an error?
20. When export files from stage, only json and csv could be in compressed zip format?
21. Which option yields best staging efficiency?
22. Security -
  23. When create or replace pipe is executed, what happens? What happens to files uploaded with previous pipe?
  24. Is file format object need to be defined for loading?
  25. When exciting a task, how to estimate/make it efficient? test/use multi cluster wh/...?
  26. Best file sizes for loading (is included in a larger multiple selection question)
  27. Commands used to load to table/ put files into stage/ etc
  28. During a query, how is warehouse(?) cache stored? Columns that is used in the query/rows that is used in the query/XXX?/result of query ?
  29. SF as a SaaS eliminates user from what? Choose all that apply - There is no hardware (virtual or physical) to select, install, configure, or manage / There is virtually no software to install, configure, or manage / Ongoing maintenance, management, upgrades, and tuning are handled by Snowflake / user no need to maintain metadata(?) ?
  30. Sharing is available starting from which SF edition?
  31. IP blocking list ect, white list/black list is available starting from which SF edition?
  32. Share can only happen within accounts of same region and same cloud provider?
  33. Select correct statements about SF releases?
  34. What is recommended practice for user with both account admin and csys admin roles?  
Select all that apply. They should only use account admin in rare cases when must / should try to make account admin be the owner of all objects / XXX / (something about MFA should or not be recommended) ?
  35. Single sign on allow provider to send user account info to SF to XXX ?
  36. What is supported? Bulk loading/snow pipe XXX ?
  37. WHat is true about materialized views? Select all that apply
  38. In what way is materialized views different from table? materialized view can be created from only two tables/cannot join materialized views with table/ XXX ... ?
  39. What is SF good for? Select two. OLTP/OLAP/concurrent processing/XXX ?
  40. Stored procedures (didn't understand the question, cannot remember answers)
  41. Time travel and fail safe time for different SF editions
  42. Semi-structured data loaded to variant
  43. Clone is good for? Select all that apply - QA analysis/XXX/XXX/XXX ?
  44. AWS private link is for AWS user to connect to SF private cloud? No
  45. What provides user a way to change how micropartitions are stored / arranged?  
Clustering key / some XXX optimization service / XXX / XXX ?
  46. What cannot be replicated? role/user/database/indexing/ XXX ?
  47. Are micro partitions compressed in storage?
  48. Encryption happens at client side before loading
  49. Data share incurs no storage cost for consumers because it is XXX ?
  50. You can query files from external stage even if they are not loaded?
  51. SF support standard XXX(some filtering rules) in PUT statement for data loading?

52. Single = true command when unloading
53. Scale in/out up/down of single warehouse/multi clustered warehouse?
54. System admin role can see all other user's query and query results in the account?
55. Users in same role can see each other's query results?
56. How to set a multi cluster warehouse in Maximum/auto scale mode?
57. All tables in SF have fail safe?
58. Keys are automatically rotated every ? days? 30
59. Share data consumers cannot download the shared data and use them in other applications/platforms? And can only query them. ?
60. Time travel can be set lowest upto which level? account/database/scheme/table?
61. Scale up a warehouse by one size will make a query around twice as fast?
62. Size limit of variant ?
63. Granular of snowpipe charge? Per sec per processor or something else?
64. ... All the rest of the questions are more straightforward, so I cannot remember

The Snowflake Pro Core examination took place on 23/10/2021.  
Medium hard when compared to Mock test.  
I scored 921 out of 1000.  
I completed the exam in 45 Min.  
These are my questions.

1	Which interfaces can be used to create and/or manage Virtual Warehouses?	
	The Snowflake Web Interface (UI)	
	SQL commands	
	Data integration tools	
	All of the Above	Answer: D
2	Which of the following DML commands isn't supported by Snowflake? Select one.	
	UPSERT	
	MERGE	
	UPDATE	
	TRUNCATE TABLE	A
3	Account-level storage usage can be monitored via: Select one.	
	The Snowflake Web Interface (UI) in the Databases section.	

	The Snowflake Web Interface (UI) in the Account -> Billing & Usage section	
	The Information Schema -> ACCOUNT_USAGE_HISTORY View	
	The Account Usage Schema -> ACCOUNT_USAGE_METRICS	View Answer: B
4	Which statement best describes Snowflake tables? Select one.	
	Snowflake tables are logical representations of underlying physical data	
	Snowflake tables are the physical instantiation of data loaded into Snowflake	
	Snowflake tables require that clustering keys be defined to perform optimally	
	Snowflake tables are owned by a user	Answer: A
5	True or False: Fail-safe can be disabled within a Snowflake account.	
	TRUE	
	FALSE	Answer: B

6	True or False: Data Consumers must execute the SHARE REFRESH; command in order to view new or modified data in shared objects.	
	TRUE	
	FALSE	Answer: B
7	Which of the following statements about Data Sharing is true? Select one.	
	Accounts can share with other accounts regardless of Snowflake edition, without requiring help from Snowflake support (e.g. Enterprise, Business Critical Edition)	
	Data sharing can cross regions, but not cloud providers	
	The Data Consumer can only see objects in the Data Provider's source database that have been explicitly added to the Share	
	A Data Provider can only share with other Snowflake customers	Answer: A
8	True or False: All Snowflake table types include fail-safe storage.	
	TRUE	
	False	Answer: B

9	True or False: Loading data via the COPY command only allows error handling on a per- file basis—all files are loaded in their entirety or marked as error, there are no partial loads.	
	TRUE	
	False	Answer: B
10	What is the minimum Snowflake edition that provides data sharing? Select one.	
	Standard	
	Premier	
	Enterprise	
	Business Critical Edition	Answer: A
11	To run a multi-cluster Warehouse in maximized mode, a user would: Select one.	
	Configure the Maximum Clusters setting to "Maximum"	
	Manually turn on the additional clusters after starting the Warehouse	
	Set the Minimum Clusters and Maximum Clusters settings to the same value	
	Set the Minimum Clusters and Maximum Clusters settings to the different values	Answer: C

12	Each incremental increase in Virtual Warehouse size (e.g. Medium to Large) generally results in what?	
	Select one.	
	More micro-partitions	
	Better query scheduling	
	Double the numbers of servers in the compute cluster	
	Higher storage costs	Answer: C
13	True or False: When Snowflake is configured to use Single Sign-On (SSO), Snowflake receives the usernames and credentials from the SSO service and loads them into the customer's Snowflake account.	
	TRUE	
	False	B
14	What happens when a Data Provider revokes privileges to a Share on an object in their source database?	
	Select one.	
	The object immediately becomes unavailable for all Data Consumers	
	Any additional data arriving after this point in time will not be visible to Data Consumers	

	The Data Consumers stop seeing data updates and become responsible for storage charges for the object	
	A static copy of the object at the time the privilege was revoked is created in the Data Consumers' accounts	Answer: A
15	What is the most granular object that the Time Travel retention period can be defined on? Select one.	
	Account	
	Database	
	Schema	
	Table	Answer: D
16	When a Pipe is recreated using the CREATE OR REPLACE PIPE command: Select one.	
	The Pipe load history is reset to empty	
	The REFRESH parameter is set to TRUE	
	Previously loaded files will be ignored	
	All of the above	Answer: A
17	A Virtual Warehouse's auto-suspend and auto-resume settings apply to: Select one.	

	The primary cluster in the Virtual Warehouse	
	The entire Virtual Warehouse	
	The database the Virtual Warehouse resides in	
	The queries currently being run by the Virtual Warehouse	Answer: B
18	To run a Multi-Cluster Warehouse in auto-scale mode, a user would:	
	Select one.	
	Configure the Maximum Clusters setting to "Auto-Scale"	
	Set the Warehouse type to "Auto"	
	Set the Minimum Clusters and Maximum Clusters settings to the same value	
	Set the Minimum Clusters and Maximum Clusters settings to the different values	Answer: D
19	A deterministic query is run at 8am, takes 5 minutes, and the results are cached. Which of the following statements are true?	
	Choose 2 answers	
	The exact query will ALWAYS return the precomputed result set for the RESULT_CACHE_ACTIVE = time period	

	The same exact query will return the precomputed results if the underlying data hasn't changed and the results were last accessed within the previous 24 hour period	
	The same exact query will return the precomputed results even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period	
	The "24 hour" timer on the precomputed results gets renewed every time the exact query is executed	Answer: BD
20	How would you best determine the size of the virtual warehouse used for a task?	
	Root task may be executed concurrently (i.e. multiple instances), it is recommended to leave some margins in the execution window to avoid missing instances of execution	
	Querying (SELECT) the size of the stream content would help determine the warehouse size. For example, if querying large stream content, use a larger warehouse size.	
	If using the stored procedure to execute multiple SQL statements, it's best to test run the stored procedure separately to size the compute resource first.	
	Since task infrastructure is based on running the task body on schedule, it's recommended to configure the virtual warehouse for automatic concurrency handling using Multi-duster warehouse (MCW) to match the task schedule.	Answer: D
21	True or False: The following command will fail: GRANT SELECT, INSERT, UPDATE, DELETE ON MYTABLE TO USER JOHN;	

	TRUE	
	False	Answer: A
22	True or False: Reader Accounts incur no additional Storage costs to the Data Provider since they are simply reading the shared data without making changes.	
	TRUE	
	False	Answer:A
23	Which of the following are true of multi-cluster Warehouses? Select all that apply below.	
	A multi-cluster Warehouse can add clusters automatically based on query activity	
	A multi-cluster Warehouse can automatically turn itself off after a period of inactivity	
	A multi-cluster Warehouse can scale down when query activity slows	
	A multi-cluster Warehouse can automatically turn itself on when a query is executed against it	Answer: ABCD
24	True or False: When active, a Pipe requires a dedicated Virtual Warehouse to execute.	
	TRUE	

	False	Answer: B
25	True or False: The PUT command can be run through the Snowflake Web Interface (UI).	
	TRUE	
	False	Answer: B
26	True or False: Users are able to see the result sets of queries executed by other users that share their same role.	
	TRUE	
	FALSE	Answer: B
27	True or False: You can query the files in an External Stage directly without having to load the data into a table.	
	TRUE	
	False	Answer: A
28	When can a Virtual Warehouse start running queries? Select one.	
	12am-5am	

	Only during administrator defined time slots	
	When its provisioning is complete	
	After replication	Answer: C
29	As a best practice, clustering keys should only be defined on tables of which minimum size?	
	Select one.	
	Multi-Kilobyte (KB) Range	
	Multi-Megabyte (MB) Range	
	Multi-Gigabyte (G8) Range	
	Multi-Terabyte (TB) Range	Answer: D
30	Which item in the Data Warehouse migration process does not apply in Snowflake? Select one.	
	Migrate Users	
	Migrate Schemas	
	Migrate Indexes	
	Build the Data Pipeline	Answer: C

31	What would you do to execute a series of tasks that run in sequence and are dependent on the previous task execution?	
	Create a series of tasks based on a series of schedules. Example Create one task that runs every 1 minutes, and a following task that runs every 2 minutes and so on: CREATE TASK task1 .. SCHEDULE = '1 minute' CREATE TASK task2 ... SCHEDULE = '2 minute'	
	When creating a task, indicate a predecessor task that it should run after. In this way, you can string together a number of tasks, all dependent on one task or each other.	
	Create a stored procedure for each task: CREATE PROCEDURE my_task1_sp1()... CREATE PROCEDURE my_task2_sp2(); Then create a simple tree by specifying the stored procedure my_task_sp1 as the predecessor that triggers the new task my_task2 when run successfully CREATE TASK my_task2 WAREHOUSE = mywh AFTER my_task1_sp1 AS ...	
	This is currently not possible using Snowflake.	Answer: B
32	Which statements are true of micro-partitions? Choose 2 answers	
	They are approximately 16MB in size	
	They are stored compressed only if COMPRESS=TRUE on Table	
	They are immutable	
	They are only encrypted in the Enterprise edition and above	Answer: AC

33	What is the maximum compressed row size in Snowflake? Select one.	
	8KB	
	16MB	
	50MB	
	4000GB	Answer: B
34	Which of the following statements are true of Snowflake releases: Choose 2 answers	
	They happen approximately weekly	
	They roll up and release approximately monthly, but customers can request early release application	
	During a release, new customer requests/queries/connections transparently move over to the newer version	
	A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and the customer is upgraded	Answer: AC
35	Which of the following commands are not blocking operations? Choose 2 answers	
	UPDATE	
	INSERT	

	MERGE	
	COPY	Answer: BD
36	True or False: Snowpipe via REST API can only reference External Stages as source.	
	TRUE	
	False	Answer: B
37	Which of the following terms best describes Snowflake's database architecture? Select one.	
	Columnar shared nothing	
	Shared disk	
	Multi-cluster, shared data	
	Cloud-native shared memory	Answer: C
38	Which of the following roles is recommended to be used to create and manage users and roles?	
	Select one.	
	SYSADMIN	
	SECURITYADMIN	

	PUBLIC	
	ACCOUNTADMIN	Answer: B
39	True or False: Every object is owned by one and only one role; therefore a user must be using that role to access the object.	
	TRUE	
	FALSE	Answer: B
40	Which of the following is an example of an operation that can be completed without requiring Compute, assuming no queries have been executed previously? Select one.	
	SELECT SUM(ORDER_AMT) FROM SALES;	
	SELECT AVG(ORDER_QTY) FROM SALES;	
	SELECT MIN(ORDER_AMT) FROM SALES;	
	SELECT ORDER_AMT * ORDER_QTY FROM SALES;	Answer: C
41	Which of the following commands sets the Virtual Warehouse for a session? Select one.	
	COPY WAREHOUSE FROM <>config file<>;	

	SET WAREHOUSE = <<warehouse name>>;	
	USE WAREHOUSE <<warehouse name>>;	
	USE VIRTUAL_WAREHOUSE <<warehouse name>>;	Answer: C
42	Which of the following statements is true of Snowflake micro-partitioning? Select one.	
	Micro-partitioning has been known to introduce data skew	
	Micro-partitioning: requires a partitioning schema to be defined up front	
	Micro-partitioning is transparently completed using the ordering that occurs when the data is inserted/loaded	
	Micro-partitioning can be disabled within a Snowflake account	Answer: C
43	True or False: During data unloading, only JSON and CSV files can be compressed.	
	TRUE	
	False	Answer: B
44	The number of queries that a Virtual Warehouse can concurrently process is determined by:	
	Choose 2 answers	

	The complexity of each query	
	The CONCURRENT_QUERY_LIMIT parameter set on the Snowflake account	
	The size of the data required for each query	
	The tool that is executing the query	Answer: AC
45	True or False: You can define multiple columns within a clustering key on a table.	
	TRUE	
	False	Answer: A
46	True or False: A 4X-Large Warehouse may, at times, take longer to provision than a X-Small Warehouse.	
	TRUE	
	False	Answer: A
47	How would you execute a series of SQL statements using Task?	
	Include the SQL statements in the body of the task CREATE TASK mytask .. AS INSERT INTO target1 SELECT .. FROM stream_si WHERE ..INSERT INTO target? SELECT .. FROM stream_si WHERE ..	

	A stored procedure can have only one DML statement per stored procedure invocation and therefore you should sequence stored procedures calls in the task definition CREATE TASK mytask	AS call stored_proc1(); call stored_proc2();
	A stored procedure executing multiple SQL statements and invoke the stored procedure from the task. CREATE TASK mytask	AS call
	stored_proc_multiple_statements_inside();	
	Create a task for each SQL statement (e.g. resulting in task1, task2, etc) and string the series of SQL statements by having a control task calling task 1, task 2 sequentially.	Answer: C
48	Which of the following statements would be used to export/unload data from Snowflake? Select one.	
	COPY INTO @stage	
	EXPORT TO @stage	
	INSERT INTO @stage	
	EXPORT_TO_STAGE(stage => @stage, select => 'select * from t1');	Answer: A
49	What is the best approach for loading data into Snowflake? Select one.	
	Merge files into one big file	
	Many big files and a big Virtual Warehouse size	

	Many small files with a high level of parallelism	
	Few files and a big Virtual Warehouse size	Answer: C
50	When a Pipe is recreated using the CREATE OR REPLACE PIPE command: Select one.	
	The Pipe load history is reset to empty	
	The REFRESH parameter is set to TRUE	
	Previously loaded files will be ignored	
	All of the above	Answer: A
51	Which of the following languages can be used to implement Snowflake User Defined Functions (UDFs)?	
	Choose 2 answers	
	Java	
	Javascript	
	SQL	
	Python	Answer: BC
52	Which of the following statements are true? Choose 3 answers	

	The COPY command must reference a Stage or cloud storage location	
	A named File Format is optional when using the COPY command	
	A File format option used in the COPY command will override the one in Stage	
	A File Format object must be defined when creating a Stage	Answer: ABC
53	True or False: A single schema can exist in more than one database.	
	TRUE	
	FALSE	ANSWER:B
54	Which is true of Snowflake network policies? A Snowflake network policy: CHOOSE 2 ANSWERS	
	Is available to all Snowflake Editions	
	Is only available to customers with Business Critical Edition	
	Restricts or enables access to specific IP addresses	
	Is activated using an "ALTER DATABASE" command.	ANSWER:AC
55	True or False: Bulk unloading of data from Snowflake supports the use of a SELECT statement.	

	TRUE	
	FALSE	ANSWER:A
56	True or False: When a new Snowflake object is created, it is automatically owned by the user who created it.	
	TRUE	
	FALSE	ANSWER:B
57	When scaling up Virtual Warehouses by increasing Virtual Warehouse t-shirt size, you are primarily scaling for improved:	
	SELECT ONE.	
	Concurrency	
	Performance	ANSWER:B
58	Which of the following statements is true of zero-copy cloning? SELECT ONE.	
	Zero-copy clones increase storage costs as cloning the table requires storing its data twice	
	All zero-copy clone objects inherit the privileges of their original objects	
	Zero-copy cloning is licensed as an additional Snowflake feature	

	At the instance/instant a clone is created, all micro-partitions in the original table and the clone are fully shared	ANSWER:D
59	Which of the following are common use cases for zero-copy cloning? CHOOSE 3 ANSWERS	
	Quick provisioning of Dev and Test/QA environments	
	Data backups	
	Point in time snapshots	
	Performance optimization	ANSWER:ABC
60	True or False: A single database can exist in more than one Snowflake account.	
	TRUE	
	FALSE	ANSWER:B
61	True or False: Reader Accounts are able to extract data from shared data objects for use outside of Snowflake.	
	TRUE	
	FALSE	ANSWER:B

62	True or False: ACCOUNTADMINs are able to see the result sets of queries executed by all users within a Snowflake account.	
	TRUE	
	FALSE	ANSWER:B
63	True or False: Multi-Factor Authentication (MFA) is available only to customers with Enterprise edition and above.	
	TRUE	
	FALSE	ANSWER:B
64	What is the minimum Snowflake edition that provides data sharing? Select one.	
	Standard	
	Premier	
	Enterprise	
	Business Critical Edition	ANSWER:A
65	True or False: When a user creates a role, they are initially assigned ownership of the role and they maintain ownership until it is transferred to another user.	
	TRUE	

	False	ANSWER:B
66	True or False: An active warehouse is required to run a COPY INTO statement.	
	TRUE	
	FALSE	ANSWER:A
67	How are Snowpipe charges calculated? Select one.	
	Per-second/per Warehouse size	
	Per-second/per-core granularity	
	Number of Pipes in account	
	Total storage bucket size	Answer: B
68	Which are true of Snowflake roles? Choose 2 answers	
	All grants to objects are given to roles, and never to users	
	In order to do DML/DDL, a user needs to have selected a single role that has that specific access to the object and operation	
	The PUBLIC role controls all other roles	
	Roles are a subset of users and users own objects in Snowflake	Answer: AB

69	True or False: Snowflake allows its customers to directly access the micro-partition files that make up its tables.	
	TRUE	
	False	Answer: B
70	Credit Consumption by the Compute Layer (Virtual Warehouses) is based on: Choose 2 answers	
	Number of users	
	Warehouse size	
	Amount of data processed	
	# of clusters for the Warehouse	Answer: BD
71	What is the recommended method for loading data into Snowflake? Select one.	
	Load row by row	
	Load data in batch	
	Load data by writing it in the Snowflake Web Interface (UI)	
	Load data via frequent, single row DML's	Answer: B

72	Which of the following are best practices for users with the SYSADMIN/ACCOUNTADMIN roles?	
	Choose 3 answers	
	Their default role should be set to SYSADMIN (the lower of the two)	
	They should not set up Multi-Factor Authentication (MFA)-as administrators they may need to change the MFA settings and those enrolled in MFA are unable to do so	
	They should only access and 'step into' the ACCOUNTADMIN role temporarily, as needed to complete a specific operation	
	They should ensure all database objects in the account are owned by the ACCOUNTADMIN role	
	They should use the SYSADMIN role to perform administrative work on database objects	Answer: ACE
73	True or False: AWS Private Link provides a secure connection from the Customer's on-premise data center to the Snowflake Virtual Private Cloud (VPC).	
	TRUE	
	False	Answer: B
74	True or False: The COPY command supports regular expressions to filter the population of files to be loaded.	
	TRUE	

	False	Answer: A
75	What is the minimum duration charged when starting a virtual warehouse? Select one.	
	1 second	
	1 minute	
	1 hour	
	1 day	Answer: B
76	Which of the following statements is true of data loading? Select one.	
	Resizing the Virtual Warehouse from X-Small to Small will process a single file twice as fast	
	The ideal file size for loading is 16MB to match micro-partition size	
	Many files In the 10-100MB range tend to land in the "sweet spot" for load parallelism	
	Once loaded, there is no option to force a reload of an already loaded file	Answer: C
77	Which of the following are use cases that fit well in Snowflake? Choose 3 answers	
	Data Lake	

	Enterprise Data Warehouse	
	Online Transaction Processing	
	Datamart Consolidation	Answer: ABD
78	Snowflake provides a mechanism for its customers to override its natural clustering algorithms. This method is:	
	Select one.	
	Micro-partitions	
	Clustering keys	
	Key partitions	
	Clustered partitions	Answer: B
79	True or False: A Snowflake account is charged for data stored in both Internal and External Stages.	
	TRUE	
	False	Answer: B
80	What are the two ways to create and manage Data Shares in Snowflake? Choose 2 answers	
	Via the Snowflake Web Interface (UI)	

	Via the DATA_SHARE=TRUE parameter	
	Via SQL commands	
	Via Virtual Warehouses	Answer: AC
81	True or False: Snowflake enforces unique, primary key, and foreign key constraints during DML operations.	
	TRUE	
	False	Answer: B
82	True or False: Snowflake's Global Services Layer gathers and maintains statistics on all columns in all micro-partitions.	
	TRUE	
	False	Answer: A
83	Which of the following are valid approaches to loading data into a Snowflake table? Select all that apply.	
	Bulk copy from an External Stage	
	Continuous load using Snowpipe REST API	
	The Snowflake Web Interface (UI) data loading wizard	

	Bulk copy from an Internal Stage	Answer: ABCD
84	When scaling out by adding clusters to a multi-cluster warehouse, you are primarily scaling for improved:	
	Select one.	
	Concurrency	
	Performance	Answer: A
85	True or False: When a data share is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data Consumers.	
	TRUE	
	False	Answer: B
86	What services does Snowflake automatically provide for customers that they may have been responsible for with their on-premise system? Select all that apply.	
	Installing and configuring hardware	
	Patching software	
	Physical security	
	Maintaining metadata and statistics	Answer: ABD

87	Which Snowflake object enables loading data from files as soon as they are available in a cloud storage location?	
	Select one.	
	Pipe	
	External Stage	
	File Format	
	VARIANT	Answer: A
88	True or False: The COPY command must specify a File Format in order to execute.	
	TRUE	
	False	Answer: B
89	True or False: It is possible to query data from an Internal or named External stage without loading the data into Snowflake.	
	TRUE	
	False	Answer: A

90	A client uses a PUT command to upload a new file to an Internal Stage. At which point is the data encrypted?	
	Select one.	
	When it reaches the virtual warehouse	
	When it gets micro-partitioned and stored in Cloud Storage	
	On the client's machine before being sent	
	Clients must encrypt files before sending to Snowflake	Answer: D
	Question	
91	Snowflake is designed for which type of workloads? Choose 2 answers	
	OLAP (Analytics) workloads	
	OLTP (Transactional) workloads	
	Concurrent workloads	
	On-premise workloads	Answer: AC
92	Tableau and looker are which type	Business Intelligence
93	What are the below that supports data science tools??	Data Robot

94	Data Replication reg question	
95	Data Market place question reg how can we req access to data from other publisher	
96	what are the type accounts involved in a share??	Reader/Consumer/Provider
	What is the minimum duration charged when starting a virtual warehouse? A . 1 second B . 1 minute C . 1 hour D . 1 day	B
	True or False: When you create a custom role, it is a practice to immediately grant that role to ACCOUNTADMIN. A . True B . False	B
	True or False: It is possible for a user to run a query against the query result cache without requiring an active Warehouse. A . True B . False	B
	What is the most granular object that the Time Travel retention period can be defined on? A . Account B . Database C . Schema D . Table	D

	<p>Snowflake recommends, as a minimize, that all users with the following roles(s) should be enrolled in Multi-Factor Authentication (MFA):</p> <ul style="list-style-type: none"> <li>A . SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN</li> <li>B . SECURITYADMIN ACCOUNTADMIN, SYSADMIN</li> <li>C . SECURITYADMIN, ACCOUNTADMIN</li> <li>D . ACCOUNTADMIN</li> </ul>	D
	<p>Which statement best describes Snowflake tables?</p> <p><b>A . Snowflake tables are logical representations of underlying physical data</b></p> <ul style="list-style-type: none"> <li>B . Snowflake tables ate the physical instantiation of data loaded Into Snowflake</li> <li>C . Snowflake tables require that clustering keys be defined to perform optimally</li> <li>D . Snowflake tables are owned by a user</li> </ul>	A
	<p>True or False: Users are able to see the result sets of queries executed by other users that share their same role.</p> <ul style="list-style-type: none"> <li>A . True</li> <li>B . False</li> </ul>	B
	<p>Which of the following commands sets the Virtual Warehouse for a session?</p> <ul style="list-style-type: none"> <li>A . COPT WAREHOUSE FROM &lt;&lt;Config file&gt; ;</li> <li>B . SET warehouse = &lt;&lt;warehouse name&gt;&gt;;</li> <li>C . USE WAREHOUSE &lt;&lt;warehouse name&gt;&gt;;</li> <li>D . USE VIRTUAL_WAREHOUSE &lt;&lt;warehouse name&gt;&gt;;</li> </ul>	C
	<p>Which of the following statements is true of Snowflake micro-partitioning?</p> <ul style="list-style-type: none"> <li>A . Micro-partitioning has been known to introduce data skew</li> <li>B . Micro-partitioning: requires a partitioning schema to be defined up front</li> <li><b>C . Micro-partitioning is transparently completed using the ordering that occurs when the data is inserted-loaded</b></li> <li>D . Micro-partitioning can be disabled within a Snowflake account</li> </ul>	C
	<p>What happens when a Data Provider revokes privileges to a Share on an object in their source database?</p> <ul style="list-style-type: none"> <li>A . The object immediately becomes unavailable for all Data Consumers</li> <li>B . Any additional data arriving after this point in time will not be visible to Data Consumers</li> <li>C . The Data Consumers stop seeing data updates and</li> </ul>	A

	become responsible for storage charges for the object D . A static copy of the object at the time the privilege was revoked is created In the Data Consumers' accounts	
	True or False: AWS Private Link provides a secure connection from the Customer's on-premise data center to the Snowflake. <b>A . True</b> <b>B . False</b>	
	True or False: Snowpipe via RFST API can only reference External Stages as source. <b>A . True</b> B . False	
	The Information Schema and Account Usage Share provide storage information for which of the following objects? (Choose three.) A . Users <b>B . Tables</b> <b>C . Databases</b> <b>D . Internal Stages</b>	
	True or False: During data unloading, only JSON and CSV files can be compressed. A . True <b>B . False</b>	
	Which of the following statements are true of Snowflake releases: (Choose two.) <b>A . They happen approximately weekly</b> B . They roll up and release approximately monthly, but customers can request early release application <b>C . During a release, new customer requests/queries/connections transparently move over to the newer version</b> D . A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and customer is upgraded	
	If auto-suspend is enable for a Virtual Warehouse, he Warehouse is automatically suspended when: A . All Snowflake sessions using the warehouse are terminated. B . The last query using the warehouse completes. C . There are no users loaned into Snowflake. <b>D . The Warehouse is inactive for a specified period of time.</b>	
	True or False: A 4X-Large Warehouse may, at times, take longer to provision than a X-Small Warehouse.	B

	A . True <b>B . False</b>	
	Which of the following commands are not blocking operations?  Choose 2 answers A . UPDATE <b>B . INSERT</b> C . MERGE <b>D . COPY</b>	
	True or False: It is possible to unload structured data to semi-structured formats such as JSON and parquet. <b>A . True</b> B . False	
	True or False: When a user creates a role, they are initially assigned ownership of the role and they maintain ownership until it is transferred to another user. <b>A . True</b> B . False	
	When can a Virtual Warehouse start running queries? A . 12am-5am B . Only during administrator defined time slots <b>C . When its provisioning is complete</b> D . After replication	
	True or False: When active, a pipe requires a dedicated Virtual Warehouse to execute. A . True <b>B . False</b>	
	Data storage for individual tables can be monitored using which commands and/or object(s)? Choose 2 answers <b>A . SHOW TABLES;</b> B . SHOW STORAGE BY TABLE; <b>C . Information Schema -&gt; TABLE_STORAGE_METRICS</b> D . Information Schema -> TABLE_HISTORY	
	Which of the following are examples of operations that require a Virtual Warehouse to complete, assuming no queries have been executed previously? Choose 3 answers A . MIN(< < column value>>) <b>B . COPY</b> <b>C . SUM(&lt;&lt; column value &gt;&gt;)</b> <b>D . UPDATE</b>	
	Which of the following statements is true of data loading? A . Resizing the virtual warehouse from x-Small to Small will process a single file twice as fast B . The "ideal file size for loading is 16MB to match micro-partition size <b>C . Many files in the 10-100MB range tend to land In the</b>	

	<p>'sweet spot" for load parallelism Once loaded, there is no option to force a reload of an already loaded file</p>	
	<p>How a Snowpipe charges calculated?</p> <p>A . Per-second/per Warehouse size</p> <p><b>B . Per-second/per-core granularity</b></p> <p>C . Number of Pipes in account</p> <p>D . Total storage bucket size</p>	
	<p>Which is true of Snowflake network policies? A Snowflake network policy: (Choose two.)</p> <p><b>A . Is available to all Snowflake Editions</b></p> <p>B . Is only available to customers with Business Critical Edition</p> <p><b>C . Restricts or enables access to specific IP addresses</b></p> <p>D . Is activated using an "ALTER DATABASE" command</p>	
	<p>True or False: Each worksheet in the Snowflake Web Interface (UI) can be associated with different roles, databases, schemas, and Virtual Warehouses.</p> <p><b>A . True</b></p> <p>B . False</p>	
	<p>True or False: Micro-partition metadata enables some operations to be completed without requiring Compute.</p> <p><b>A . True</b></p> <p>B . False</p>	
	<p>Which of the following statements about data sharing are true? choose 2 answers</p> <p>A . New objects created by a Data Provider are automatically shared with existing Data Consumers &amp; Reader Accounts</p> <p>B . All database objects can be included In a shared database</p> <p><b>C . Reader Accounts are created and funded by Data Prowlers</b></p> <p><b>D . Shared databases are read-only</b></p>	
	<p>The fail-safe retention period is how many days?</p> <p>A . 1 day</p> <p><b>B . 7 days</b></p> <p>C . 45 days</p> <p>D . 90 days</p>	
	<p>Account-level storage usage can be monitored via:</p> <p>A . The snowflake wet Interface (UI) in the Databases section.</p> <p>B . The Snowflake web interface (UI) in the Account - &gt; Billing a usage section</p>	

	<p>C . The Information Schema -&gt; ACCOUNT_USAGE_HISTORY View</p> <p><b>D . The Account usage Schema – &gt; ACCOUNT_USAGE_METRICS View</b></p>	
	<p>Query results are stored in the Result Cache for how long after they are last accessed, assuming no data changes have occurred?</p> <p>A . 1 Hour B . 3 Hours C . 12 hours <b>D . 24 hours</b></p>	
	<p>Which type of table corresponds to a single Snowflake session?</p> <p><b>A . Temporary</b> B . Transient C . Provisional D . Permanent</p>	
	<p>True or False: Snowflake allows its customers to directly access the micro-partition files that make up its tables.</p> <p>A . True <b>B . False</b></p>	
	<p>Which statement best describes " clustering"?</p> <p><b>A . Clustering represents the way data is grouped together and stored within snowflake's micro-partitions</b> B . The database administrator must define the clustering methodology for each Snowflake table. C . The clustering key must be included on the COPV command when loading data into Snowflake. D . Clustering can be disabled within a Snowflake account.</p>	
	<p>What is the recommended method for loading data into Snowflake?</p> <p>A . Load row by row <b>B . Load data in batch</b> C . Load data by writing it In the Snowflake Web Interface (UI) D . Load data via frequent, angle row DML's</p>	
	<p>Which of the following statements are true of Virtual Warehouses? (Choose all that apply.)</p> <p><b>A . Customers can change the size of the Warehouse after creation</b> <b>B . A Warehouse can be resized while running</b> <b>C . A Warehouse can be configured to suspend after a period of inactivity</b></p>	

	<p><b>D . A Warehouse can be configured to auto-resume when new queries are submitted</b></p> <p>HIDE ANSWER</p>	
	<p>True or False: You can resize a Virtual Warehouse while queries are running.</p> <p><b>A . True</b> B . False</p>	
	<p>What is the lowest Snowflake edition that offers Time Travel up to 90 days.</p> <p>A . standard Edition B . Premier Edition <b>C . Enterprise Edition</b> D . Business Critical Edition</p>	
	<p>True or False: The COPY command must specify a File Format in order to execute.</p> <p>A . True <b>B . False</b></p>	
	<p>How many shares can be consumed by single Data Consumer?</p> <p>A . 1 B . 10 C . 100, but can be increased by contacting support <b>D . Unlimited</b></p> <p>HIDE ANSWER</p>	
	<p>Which of the following statements are true about Schemas in Snowflake? (Choose two.)</p> <p>A . A Schema may contain one or more Databases <b>B . A Database may contain one or more Schemas</b> <b>C . A Schema is a logical grouping of Database Objects</b> D . Each Schema is contained within a Warehouse</p>	
	<p>What parameter controls if the Virtual warehouse starts immediately after the CREATE WAREHOUSE statement?</p> <p><b>A . INITIALLY_SUSPENDED = TRUE/FALSE</b> B . START_AFTCR_CREATE = TRUE/FALSE C . START_TTIME = 60 // (seconds from now) D . STAR E . TIME = CURREN F . DATE()</p>	

	<p>True or False: A single database can exist in more than one Snowflake account.</p> <p>A . True  <b>B . False</b></p>	
	<p>When loading data into Snowflake, the COPY command supports: Choose 2 answers</p> <p>A . Joins  B . Filters  <b>C . Data type conversions</b>  <b>D . Column reordering</b>  E . Aggregates</p>	
	<p>Snowflake is designed for which type of workloads? (Choose two.)</p> <p><b>A . OLAP (Analytics) workloads</b>  B . OLTP (Transactional) workloads  <b>C . Concurrent workloads</b>  D . On-premise workloads</p> <p>HIDE ANSWER</p>	
	<p>True or False: Data in Fail-safe can be deleted by a user or the Snowflake team before it expires.</p> <p>A . True  <b>B . False</b></p>	
	<p>To run a Multi-Cluster Warehouse in auto-scale mode, a user would:</p> <p>A . Configure the Maximum Clusters setting to “Auto-Scale”  B . Set the Warehouse type to “Auto”  C . Set the Minimum Clusters and Maximum Clusters settings to the same value  <b>D . Set the Minimum Clusters and Maximum Clusters settings to the different values</b></p>	
	<p>Which formats are supported for unloading data from Snowflake? Choose 2 answers</p> <p><b>A . Delimited (CSV, TSV, etc.)</b>  B . Avro  <b>C . JSON</b>  D . ORC</p>	
	<p>Select the different types of Internal Stages: (Choose three.)</p> <p><b>A . Named Stage</b>  <b>B . User Stage</b>  <b>C . Table Stage</b>  D . Schema Stage</p>	

	<p>What is the maximum compressed row size in Snowflake?</p> <p>A . 8KB  <b>B . 16MB</b>  C . 50MB  D . 4000GB</p>	
	<p>When a pipe is recreated using the CREATE or REPLACE PIPE command: Select one.</p> <p><b>A . The Pipe load history is reset to empty</b>  B . The REFRESH parameter is set to TRUE  C . Previously loaded files will be ignored  D . All of the above</p>	
	<p>Select the three types of tables that exist within Snowflake.</p> <p>Choose 3 answers</p> <p><b>A . Temporary</b>  <b>B . Transient</b>  C . Provisioned  <b>D . Permanent</b></p>	
	<p>Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)</p> <p>A . Custom  <b>B . Economy</b>  C . Optimized  <b>D . Standard</b></p>	

Which of the following are options when creating a virtual Warehouse? Choose 2 answers

- A . Auto-drop
- B . Auto resize
- C . Auto-resume**
- D . Auto-suspend**

What is the maximum compressed row size in Snowflake?

- A . 8KB
- B . 16MB**
- C . 50MB
- D . 4000GB

Which of the following statements is true of Snowflake?

- A . It was built specifically for the cloud**
- B . It was built as an on-premises solution and then ported to the cloud
- C . It was designed as a hybrid database to allow customers to store data either on premises or in the cloud

- D . It was built for Hadoop architecture
- E . It's based on an Oracle Architecture

True or False: Reader Accounts are able to extract data from shared data objects for use outside of Snowflake.

**A . True**

B . False

For a multi-cluster Warehouse, the number of credits billed is calculated on: Select one.

- A . The number of queries that ran using the Warehouse.
- B . The size of the Warehouse and the number of clusters that ran within a given time period.**
- C . The size of the Warehouse and the maximum number of clusters configured for the Warehouse.
- D . The number of users who accessed the Warehouse.

Which of the following accurately represents how a table fits into Snowflake's logical container hierarchy?

- A . Account -> Schema -> Database -> Table
- B . Account -> Database -> Schema -> Table**
- C . Database -> Schema -> Table -> Account
- D . Database -> Table -> Schema -> Account

Which of the following are examples of operations that require a Virtual Warehouse to complete, assuming no queries have been executed previously? Choose 3 answers

- A . MIN(< < column value>>)
- B . COPY**
- C . SUM(<< column value >>)**
- D . UPDATE**

Select the different types of Internal Stages: (Choose three.)

- A . Named Stage**
- B . User Stage**
- C . Table Stage**
- D . Schema Stage

The fail-safe retention period is how many days?

- A . 1 day
- B . 7 days**
- C . 45 days
- D . 90 days

True or False: You can resize a Virtual Warehouse while queries are running.

A . True

B . False

Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)

A . Custom

**B . Economy**

C . Optimized

**D . Standard**

HIDE ANSWER

The Information Schema and Account Usage Share provide storage information for which of the following objects? (Choose three.)

**A . Users**

**B . Tables**

**C . Databases**

D . Internal Stages

HIDE ANSWER

**Answer: A,B,C**

True or False: Snowflake allows its customers to directly access the micro-partition files that make up its tables.

A . True

**B . False**

HIDE ANSWER

**Answer: B**

Which of the following statements are true of Snowflake releases: (Choose two.)

A . They happen approximately weekly

B . They roll up and release approximately monthly, but customers can request early release application

**C . During a release, new customer requests/queries/connections transparently move**

**over to the newer version**

**D . A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and customer is upgraded**

HIDE ANSWER

**Answer: C,D**

What happens when a Data Provider revokes privileges to a Share on an object in their source database?

**A . The object immediately becomes unavailable for all Data Consumers**

B . Any additional data arriving after this point in time will not be visible to Data Consumers

C . The Data Consumers stop seeing data updates and become responsible for storage charges for the object

D . A static copy of the object at the time the privilege was revoked is created In the Data Consumers' accounts

HIDE ANSWER

**Answer: A**

Which of the following statements is true of data loading?

A . Resizing the virtual warehouse from x-Small to Small will process a single file twice as fast

B . The "deal file size for loading is 16MB to match micro-partition size

**C . Many files in the 10-100MB range tend to land In the 'sweet spot" for load parallelism Once loaded, there is no option to force a reload of an already loaded file**

HIDE ANSWER

**Answer: C**

September 25, 2021 [exams](#) [Leave a comment](#)

## Post navigation

Which of the following accurately represents how a table fits into Snowflake's logical container hierarchy?

A . Account -> Schema -> Database -> Table

**B . Account -> Database -> Schema -> Table**

- C . Database -> Schema -> Table -> Account
- D . Database -> Table -> Schema -> Account

HIDE ANSWER

**Answer: B**

Which of the following are options when creating a virtual Warehouse? Choose 2 answers

- A . Auto-drop
- B . Auto resize
- C . Auto-resume**
- D . Auto-suspend**

HIDE ANSWER

**Answer: C,D**

What is the lowest Snowflake edition that offers Time Travel up to 90 days.

- A . standard Edition
- B . Premier Edition
- C . Enterprise Edition**
- D . Business Critical Edition

HIDE ANSWER

**Answer: C**

To run a Multi-Cluster Warehouse in auto-scale mode, a user would:

- A . Configure the Maximum Clusters setting to “Auto-Scale”
- B . Set the Warehouse type to “Auto”
- C . Set the Minimum Clusters and Maximum Clusters settings to the same value
- D . Set the Minimum Clusters and Maximum Clusters settings to the different values**

HIDE ANSWER

**Answer: D**

True or False: The COPY command must specify a File Format in order to execute.

A . True

B . False

HIDE ANSWER

Answer: A

True or False: Reader Accounts are able to extract data from shared data objects for use outside of Snowflake.

A . True

B . False

HIDE ANSWER

Answer: A

Which of the following statements about data sharing are true? choose 2 answers

A . New objects created by a Data Provider are automatically shared with existing Data

Consumers & Reader Accounts

B . All database objects can be included In a shared database

C . Reader Accounts are created and funded by Data Prowlers

D . Shared databases are read-only

HIDE ANSWER

Answer: C,D

If auto-suspend is enable for a Virtual Warehouse, he Warehouse is automatically suspended when:

A . All Snowflake sessions using the warehouse are terminated.

B . The last query using the warehouse completes.

C . There are no users loaned into Snowflake.

D . The Warehouse is inactive for a specified period of time.

HIDE ANSWER

Answer: D

Which type of table corresponds to a single Snowflake session?

- A . **Temporary**
- B . Transient
- C . Provisional
- D . Permanent

HIDE ANSWER

**Answer: A**

What is the most granular object that the Time Travel retention period can be defined on?

- A . **Account**
- B . Database
- C . Schema
- D . Table

HIDE ANSWER

**Answer: A**

When loading data into Snowflake, the COPY command supports: Choose 2 answers

- A . Joins
- B . Filters
- C . **Data type conversions**
- D . **Column reordering**
- E . Aggregates

HIDE ANSWER

**Answer: C,D**

Which of the following commands sets the Virtual Warehouse for a session?

- A . COPT WAREHOUSE FROM <<Config file>> ;
- B . SET warehouse = <<warehouse name>>;
- C . **USE WAREHOUSE <<warehouse name>>;**
- D . USE VIRTUAL\_WAREHOUSE <<warehouse name>>;

HIDE ANSWER

**Answer: C**

True or False: Each worksheet in the Snowflake Web Interface (UI) can be associated with different roles, databases, schemas, and Virtual Warehouses.

A . True

**B . False**

HIDE ANSWER

**Answer: B**

True or False: AWS Private Link provides a secure connection from the Customer's on-premise data center to the Snowflake.

A . True

**B . False**

HIDE ANSWER

**Answer: B**

True or False: AWS Private Link provides a secure connection from the Customer's on-premise data center to the Snowflake.

A . True

**B . False**

HIDE ANSWER

**Answer: B**

Snowflake is designed for which type of workloads? (Choose two.)

**A . OLAP (Analytics) workloads**

**B . OLTP (Transactional) workloads**

C . Concurrent workloads

D . On-premise workloads

HIDE ANSWER

**Answer: A,B**

True or False: It is possible to unload structured data to semi-structured formats such as JSON and parquet.

A . True

B . False

HIDE ANSWER

Answer: A

Which is true of Snowflake network policies? A Snowflake network policy: (Choose two.)

A . Is available to all Snowflake Editions

B . Is only available to customers with Business Critical Edition

C . Restricts or enables access to specific IP addresses

D . Is activated using an “ALTER DATABASE” command

HIDE ANSWER

Answer: A,C

True or False: When you create a custom role, it is a practice to immediately grant that role to ACCOUNTADMIN.

A . True

B . False

HIDE ANSWER

Answer: B

Which of the following commands are not blocking operations?

Choose 2 answers

A . UPDATE

B . INSERT

C . MERGE

D . COPY

HIDE ANSWER

Answer: B,D

When can a Virtual Warehouse start running queries?

- A . 12am-5am
- B . Only during administrator defined time slots**
- C . When its provisioning is complete
- D . After replication

HIDE ANSWER

**Answer: B**

Which formats are supported for unloading data from Snowflake? Choose 2 answers

- A . Delimited (CSV, TSV, etc.)**
- B . Avro
- C . JSON**
- D . ORC

HIDE ANSWER

**Answer: A,C**

True or False: Data in Fail-safe can be deleted by a user or the Snowflake team before it expires.

- A . True
- B . False**

HIDE ANSWER

**Answer: B**

Which of the following statements is true of Snowflake?

- A . It was built specifically for the cloud**
- B . It was built as an on-premises solution and then ported to the cloud
- C . It was designed as a hybrid database to allow customers to store data either on premises or in the cloud
- D . It was built for Hadoop architecture
- E . It's based on an Oracle Architecture

HIDE ANSWER

**Answer: A**

True or False: It is possible for a user to run a query against the query result cache without requiring an active Warehouse.

A . True

**B . False**

HIDE ANSWER

**Answer: B**

Snowflake recommends, as a minimize, that all users with the following role(s) should be enrolled in Multi-Factor Authentication (MFA):

A . SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN

B . SECURITYADMIN ACCOUNTADMIN, SYSADMIN

C . SECURITYADMIN, ACCOUNTADMIN

**D . ACCOUNTADMIN**

HIDE ANSWER

**Answer: D**

Which statement best describes Snowflake tables?

**A . Snowflake tables are logical representations of underlying physical data**

B . Snowflake tables ate the physical instantiation of data loaded Into Snowflake

C . Snowflake tables require that clustering keys be defined to perform optimally

D . Snowflake tables are owned by a user

HIDE ANSWER

**Answer: A**

What parameter controls if the Virtual warehouse starts immediately after the CREATE WAREHOUSE statement?

**A . INITIALLY\_SUSPENDED = TRUE/FALSE**

B . START\_AFTCR\_CREATE = TRUE/FALSE

C . START\_TTIME = 60 // (seconds from now)

D . STAR

E . TIME = CURREN

F . DATE()

HIDE ANSWER

Answer: A



Show course content navigation

[Course](#) > [Contain...](#) > [✓ Che...](#) > ▾

✓ 1/1 point (graded)

Which of these are Snowflake table types?

Select all that apply.

Secure

Permanent

External

Materialized

Temporary

Transient



Submit

✓ Correct (1/1 point)

✓ 1/1 point (graded)

Which of the following are Snowflake view types?

Select all that apply.



Secure



Permanent



External



Standard



Materialized



Transient



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

What are the three Snowflake Stage types?

Select all that apply.



Secure



Permanent



User

Materialized

Table

Named



Submit



Correct (1/1 point)



1/1 point (graded)

Named stages come in two varieties, what are they?

Secure

Permanent

Internal

Materialized

External



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Which type of view is most like a table?

Standard

Secure

Materialized

External



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Which type of view has an extra layer of protection to hide the SQL code from unauthorized viewing?

Standard

Secure

Materialized

Permanent



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

In a Snowflake account named MX43210, you need to set a user's default namespace to a database called MYDB and the PUBLIC schema. Which of the following commands would you use?

set default\_namespace = MX43210.mydb.public

set default\_namespace = public-my-db-MX43210

set default\_namespace = mydb.public

set default\_namespace = mydb.MX43210, default schema = public



Submit

---

❶ Answers are displayed within the problem

---

✓ 1/1 point (graded)

Which statements are true about the Snowflake container hierarchy? Select all that apply.

Accounts contain databases which contain schemas.

Schemas contain databases which contain tables.

Schemas contain tables as well as views.

Databases contain users which are made up of roles.

Databases contain roles which have tables with sequences.



Submit



Correct (1/1 point)



1/1 point (graded)

In the Snowflake container hierarchy, what container is represented as a URL (for example: <https://HJ54364.snowflakecomputing.com>)?

Database

Schema

Role

Account



Submit

---

✓ Correct (1/1 point)

✓ 1/1 point (graded)

Fail-Safe is a seven-day history of data and is automatically available on which table types?

Permanent

Temporary

Transient

External



**Submit**

---

✓ Correct (1/1 point)

✓ 1/1 point (graded)

Each Snowflake account comes with two shared databases. One is a set of sample data and the other contains Account Usage information. Check all true statements about these shared databases.

SNOWFLAKE\_SAMPLE\_DATA contains a schema called ACCOUNT\_USAGE

SNOWFLAKE contains a table called ACCOUNT\_USAGE

SNOWFLAKE contains a schema called ACCOUNT\_USAGE

SNOWFLAKE\_SAMPLE\_DATA contains several schemas from TPC (tpc.org)

ACCOUNT\_USAGE is a schema filled with external tables

ACCOUNT USAGE is a schema filled with secure views



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Which of the following statements are true about Fail-safe?

Select all that apply.

Only a Snowflake employee can recover data from Fail-safe storage.

Fail-safe is a reliable way to create Dev/Test/QA and other environments.

The data stored as part of Fail-safe is part of storage costs charged to customers.

Fail-safe is not available for tables that have Time Travel.



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Time travel is available for which table types?

Select all that apply.



Permanent



Temporary



Transient



External



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Which types of stages are automatically available in Snowflake and do not need to be created or configured?

Select all that apply.



User



Table



Named Internal



Named External



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Which table type disappears after the close of the session and therefore has no Fail-safe, and no Time Travel options after the close of the session?

Permanent

Temporary

Transient

External



Submit

---

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

You set up a Snowflake account, choosing AWS as your cloud platform provider. What stages can you use to load data files?

Select all that apply.

USER



TABLE



NAMED INTERNAL



NAMED EXTERNAL - using S3 Buckets



NAMED EXTERNAL - using Azure BLOB storage



NAMED EXTERNAL - using GCS/GCP Buckets

Submit

✓ Correct (1/1 point)

---

✓ 1/1 point (graded)

Which of the following object types are stored within schemas?

Select all that apply.



Stages



File Formats



Sequences



Stored Procedures



User Defined Functions



Roles



What are advantages clones have over tables created with CREATE TABLE AS SELECT statement? (Choose two.)

- A. The clone always stays in sync with the original table.
- B. The clone has better query performance.
- C. The clone is created almost instantly.
- D. The clone will have time travel history from the original table.
- E. The clone saves space by not duplicating storage.

Which statement MOST accurately describes clustering in Snowflake?

- A. The database ACCOUNTADMIN must define the clustering methodology for each Snowflake table.
- B. Clustering is the way data is grouped together and stored within Snowflake micro-partitions.
- C. The clustering key must be included in the COPY command when loading data into Snowflake.
- D. Clustering can be disabled within a Snowflake account.

Which of the following can be used when unloading data from Snowflake? (Choose two.)

- A. When unloading semi-structured data, it is recommended that the STRIP\_OUTER\_ARRAY option be used.
- B. Use the ENCODING file format option to change the encoding from the default UTF-8.
- C. The OBJECT\_CONSTRUCT function can be used to convert relational data to semi-structured data.
- D. By using the SINGLE = TRUE parameter, a single file up to 5 GB in size can be exported to the storage layer.
- E. Use the PARSE\_JSON function to ensure structured data will be unloaded into the VARIANT data type.

What is the purpose of enabling Federated Authentication on a Snowflake account?

- A. Disables the ability to use key pair and basic authentication (e.g., username/password) when connecting
- B. Allows dual Multi-Factor Authentication (MFA) when connecting to Snowflake
- C. Forces users to connect through a secure network proxy

- D. Allows users to connect using secure single sign-on (SSO) through an external identity provider

Masking policies can be applied to which of the following Snowflake objects?  
(Choose two.)

- A. A materialized view
- B. A stored procedure
- C. A table
- D. A stream
- E. A pipe -  
E. A User-Defined Function (UDF)

Sr. No.	Question	Options	Answers	Comments/ Reference
1	Which privileges can be granted but cannot be revoked?	Ownership Manage grants SYSADMIN all of the above	<b>Ownership</b>	
2	Bulk grant on objects of type PIPE to ROLE is restricted.	Boolean	<b>TRUE</b>	Need to check again.
3	Accountadmin can view results of all the queries of that account?	Boolean	<b>FALSE</b>	
4	Snowflake stores all query results for 24 hours.	Boolean	<b>TRUE</b>	
5	A RESULT_SCAN is not guaranteed to return rows in the same order as the original query returned the rows.	Boolean	<b>TRUE</b>	
6	The VALIDATION_MODE parameter does not support COPY statements that transform data during a load.	Boolean	<b>TRUE</b>	
7	Table stage is supported for Copy into statement.	Boolean	<b>TRUE</b>	
8	COPY transformations do not support the FLATTEN function, or JOIN or GROUP BY (aggregate) syntax	Boolean	<b>TRUE</b>	Note specifically that COPY transformations do not support the FLATTEN function, or JOIN or GROUP BY (aggregate) syntax:
9	Users with a role granted the MONITOR USAGE global privilege can use the Snowflake web interface or SQL to view the credits billed to your Snowflake account.	Boolean	<b>TRUE</b>	
10	Copy statement used in snowpipe supports force attribute?	Boolean	<b>FALSE</b>	
11	Which metadata table captures the snowpipe Copy command history?	copy_history load_history pipe_history pipecopy_history	<b>Copy_history</b>	

				The number of data files that are processed in parallel is determined by the number and capacity of servers in a warehouse.  The number of data files that are processed in parallel is determined by the number and capacity of servers in a warehouse.
12	Boolean	TRUE		The number of data files that are processed in parallel is determined by the number and capacity of servers in a warehouse.  We recommend splitting large files by line to avoid records that span chunks.
13	Boolean	TRUE		STRIP_OUTER_ARRAY file format option will remove the outer array and load the records into separate table rows.  Instead, we recommend enabling the STRIP_OUTER_ARRAY file format option for the COPY INTO <table> command to remove the outer array structure and load the records into separate table rows:
14	hex base64 utf-8 binary	hex base64 utf-8		Snowflake supports following binary format:
15	Boolean	TRUE		You can increase or decrease the number of clusters for a warehouse at any time, even while the warehouse is running and executing statements.
16	Boolean	TRUE		Suspending a warehouse does not abort any queries being processed by

	the warehouse at the time it is suspended			
17	The default suspend time for a warehouse is	5 min 10 min 1 min 6 min	<b>10 min</b>	
18	Snowflake supports following unload formats	CSV JSON Parquet XML	<b>CSV JSON Parquet</b>	
19	What is true about cloning?	External named stages can be cloned. Internal named stages are not cloned. Clone does not include any of the files from the source. Table stages are cloned.	<b>External named stages can be cloned. Internal named stages are not cloned. Table stages are cloned.</b>	Individual external named stages can be cloned; internal named stages cannot be cloned.  When cloning a database or schema:  External named stages that were present in the source when the cloning operation started are cloned.  Tables are cloned, which means their internal stages are also cloned.
20	When a table with a column with a default sequence is cloned, the cloned table still references the original sequence object.	Boolean	<b>TRUE</b>	

21	Which of the global privileges are provided for performing all tasks related to shares?	Import Share Create Share Ownership all of the above	<b>Import Share Create Share</b>	
22	Which objects can be shared?	schemas tables secure views secure UDFs	<b>schemas tables secure views secure UDFs</b>	All these can be shared.
23	The data retention period for a database, schema, or table can be changed at any time.	Boolean	<b>TRUE</b>	
24	Data Retention is max 1 day for temporary table on Snowflake Enterprise edition (and higher)	Boolean	<b>TRUE</b>	For Snowflake Enterprise Edition (and higher):  For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the <b>default of 1 day</b> ). The same is also true for temporary tables
25	The default retention period for Standard edition is 1 day.	Boolean	<b>TRUE</b>	For Snowflake Standard Edition, the retention period can be set to 0 (or unset back to the default of 1 day) at the account and object level (i.e. databases, schemas, and tables).

				The Snowflake web interface provides a convenient wizard for loading limited amounts of data into a table from a small set of flat files
26	Snowflake web interface provides a UI to upload a file	Boolean	<b>TRUE</b>	
27	Snowflake web interface Worksheet data is stored in user stage	Boolean	<b>TRUE</b>	
28	Multi-cluster support is not available for standard edition	Boolean	<b>TRUE</b>	Its available from Enterprise, Business Critical and VPS
29	Following metadata columns are allowed to query data from stage files	METADATA\$FILE_NAME METADATA\$FILE_ROW_NUMBER METADATA\$COUNT METADATA\$FILE_SIZE	<b>METADATA\$FILENAME</b> <b>METADATA\$FILE_ROW_NUMBER</b>	
30	Which of the following Copy options Snowpipe does not support?	Force Files Purge all of the above	<b>all of the above</b>	All COPY options are supported except for the following:  FILES = ('file_name1' [ 'file_name2' , ... ])  PATTERN = 'regex_pattern'  ON_ERROR = ABORT_STA

				EMENT MATCH_BY_COLUMN_N AME = CASE_SENSITIVE   CASE_INSENSITIVE   NONE  SIZE_LIMIT = num  PURGE = TRUE   FALSE (i.e. automatic purging while loading)  FORCE = TRUE   FALSE
31	Snowflake datawarehouse is built for which workload? (Choose 2). 4th Option missing.	a)OLTP b)OLAP c) parallel processing	<b>b)OLAP</b> <b>c)parallel processing</b>	
32	All tables are clustered tables in snowflake-	True/False	<b>FALSE</b>	
33	can we create cluster key on more than one column?	True/False	<b>TRUE</b>	
34	Which type of table does not have fail safe?	a) Permanent b) Temporary c) Transient d) provisioned	<b>b) Temporary</b> <b>c) Transient</b>	
35	Recommended File Size to load in Snowflake	Boolean	10 MB to 100 MB	We recommend aiming to produce data files roughly 10 MB to 100 MB in size compressed.
36	Loading the Existing file consideration if we recreate Snowpipe			It does not load the previously loaded file.

37	Snowflake supports Landing data into Internal Stage	Boolean	<b>TRUE</b>	
38	Can Shared data can be cloned	Boolean	<b>FALSE</b>	No I guess
39	Billing related questions on the Shared Database			
40	Query Encryption is supported on Which Edition		<b>Business Critical/Enterprise for Sensitive</b>	
41	Time Travel for Temporary and Transient tables			1 Day for Temp and Transient.
42	Different regions required different accounts	Boolean	<b>TRUE</b>	
43	Data Sharing will be allowed on the same region	Boolean	<b>TRUE</b>	
44	Connectors available in Snowflake? Multiple choice	a. SnowsqlCLI b. Python c. R d. Hive	a. SnowsqlCLI b. Python	
45	Snowflake is built for? Multiple Choice any 2	a. OLAP workload b. OLTP workload c. On-Premise d. Concurrent workload	a. OLAP workload d. Concurrent workload	
46	Where does the metadata gets stored?	a. Global Services layer(Cloud services) b. Storage layer c. Compute layer	a. Global Services layer(Cloud services)	
47	User need to raise a request to the support team for accessing the snowflake web interface even when there exists SnowSQL CLI.	True or False?	FALSE	
48	Will credit be consumed when the virtual warehouse is in inactive state?	True or False?	FALSE	
49	Order of objects?	a. Database->Account->Schema->Tables b. Account->Schema->Database->Tables c. Account->Database->Schema>Tables d. Schema->Database->Tables->Account	c. Account->Database->Schema>Tables	

50	Can multiple Snowflake databases have same schema	True or False?	TRUE	
51	What are the stages available in Snowflake? Multiple Choices	a. Named Stage b. User Stage c. Table Stage d. Schema Stage	a. Named Stage b. User Stage c. Table Stage	
52	What are the two table types that Snowflake provides to help manage the storage costs associated with Time Travel and Fail-safe? Multiple Choices	a. Temporary table b. Transient Table c. Permanent Table d. Global Table	a. Temporary table b. Transient Table	
53	Does Temporary table have fail-safe?	True or False	FALSE	<ul style="list-style-type: none"> <li>• Transient tables can have a Time Travel retention period of either 0 or 1 day.</li> <li>• Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon as the table is dropped or the session in which the table was created ends.</li> <li>• Transient and temporary tables have no Fail-safe period.</li> </ul>
54	Question on Snowflake releases- Multiple choice		Weekly	No Downtime Transparent to the User

		a. USE WAREHOUSE; b. USE VIRTUAL_WAREHOUSE; c. SET WAREHOUSE; d. SET VIRTUAL_WAREHOUSE;		
55	Which command is used to set the warehouse	a. USE WAREHOUSE;		
56	Minimum edition required for setting up time travel for 90 days?	a. Premium b. Enterprise c. Business Critical d. All the above	b. Enterprise	
57	Can Network Policies be applied to all the editions?	True or False	<b>TRUE</b>	
58	Minimum Edition from which PHI, PCI data is supported?	a. Premium b. Enterprise c. Business Critical d. All the above	c. Business Critical	
59	Snowflake can be connected through?	a. Snowsql CLI b. Web UI c. ETL Integrators (ODBC, JDBC) d. All the above	d. All the above	
60	Will ETL Integration would work with Snowflake with its existing ODBC and JDBC drivers without downloading Snowflake ODBC/JDBC Drivers?	True or False	FALSE	
61	Can you disable Fail-Safe in Snowflake?	True or False	FALSE	
62	FLATTEN can be used to convert data from Semi-structured to Structured data?	True or False	<b>TRUE</b>	
63	Which DML operation is not supported in warehouse?	a. INSERT b. UPDATE c. UPSERT d. MERGE	c. UPSERT	
64	By which factor servers are added to the warehouse while resizing	a. 4 b. 2 c. 1 d. 8	b. 2	
65	If Small warehouse has 2 servers, how many clusters will medium size have?	a. 4 b. 2 c. 16 d. 8	a. 4	

66	Minimum time the warehouse is billed when the warehouse starts?	a. 6 seconds b. 60 seconds c. 5 minutes d. 2 minutes	b. 60 seconds	
67	Will the credits be calculated even when the warehouse is in inactive mode?	True or False	FALSE	
68	Choose the best practice on warehouse considerations		Experiment by running the same queries against warehouses of multiple sizes (e.g. X-Large, Large, Medium).	
69	What do we call the method of adding clusters to the warehouse?	a. Scale up b. Scale out c. Resize d. None of the above	b. Scale out	
70	Can the warehouse be resized when the query is running?	True or False	TRUE	
71	Syntax to create warehouse?	a. CREATE VIRTUAL_WAREHOUSE b. CREATE WAREHOUSE c. CREATE VT_WAREHOUSE d. CREATE VIRTUAL_WH	b. CREATE WAREHOUSE	
72	Which method is used to help control the credits consumed by a multi-cluster warehouse?	a. SCALING POLICY b. LEGACY POLICY c. NETWORK POLICY	a. SCALING POLICY	
73	What are the two scaling policies supported by Snowflake? Multiple choice	a. Standard b. Premium c. Economy d. Enterprise	a. Standard c. Economy	
74	Which mode is enabled by specifying different values for maximum and minimum clusters?	a. Maximized b. Auto-scale c. Both d. Standard	b. Auto-scale	
75	Which option is used to restart/resume the warehouse when the create command is executed?	a. START_TIME=60 seconds b. INITIALLY_SUSPENDED=TRUE/FALSE c. START_WH=60	b. INITIALLY_SUSPENDED=TRUE/FALSE	

		seconds d. Both a& c		
76	Can Pipe be suspended and resumed?	True or False	TRUE	
77	Which of the options is false about Snowpipe RESTAPIs?		Can be used to connect to External stages only	
78	Snowpipe required the warehouse to be active for the data ingestion?	True or False	FALSE	
79	Best practice to load the file to snowflake?		Split the bigger files into multiple small files of size <b>10-100MB</b> while loading	
80	Can cluster key have multiple columns?	True or False	TRUE	Ask snowflake support for the answer. Multiple columns can have 16 MB size in variant column.
81	What is the maximum size of a row in snowflake table?	a. 1 MB b. 8 MB c. 16 MB d. 16 GB	c. 16 MB	
82	What is the maximum size of uncompressed data that Micro Partitions can hold?	a. 10-100 MB b. 5-50 MB c. 50-500 MB d. 1-100 MB	c. 50-500 MB	
83	What metadata does Micro Partition hold?	a. Range of values for each column b. Number of distinct values c. MIN/MAX Values d. NULL	a. Range of values for each column b. Number of distinct values	
84	What is true about Micro Partitions?		Micro-partitioning is automatically performed on all Snowflake tables.  <b>Tables are transparently partitioned using the ordering of the data</b>	

			<b>as it is inserted/loaded.</b>
85	Is it good practice to cluster all the tables?	True or False	FALSE
86	Does all the tables in snowflake are clustered?	True or False	FALSE
87	Minimum Size of the table where cluster is required?	a. In MB Range b. In GB Range c. In TB Range d. In KB Range	c. In TB Range
88	Choose the functions to monitor the clustering metadata for a table? Multiple Choice Any 2	a. SYSTEM\$CLUSTERING_DEPTH b. SYSTEM\$CLUSTERING_USAGE c. SYSTEM\$CLUSTERING_PARTITIONDEPTH d. SYSTEM\$CLUSTERING_INFORMATION	a. SYSTEM\$CLUSTERING_DEPTH d. SYSTEM\$CLUSTERING_INFORMATION
89	Credits are calculated in Snowflake for – Multiple choice Choose all the correct answers	a. TimeTravel b. Fail-Safe c. Purging d. AutoScaling	a. TimeTravel b. Fail-Safe d. AutoScaling
90	Storage charge is incurred for Data provider when the data is consumed by Reader account?	True or False	TRUE
91	Computation charge is incurred for Data provider when the data is consumed by Reader account?	True or False	TRUE
92	How many shares can be created by data providers?	a. 1 b. 10 c. No limit d. 1000	c. No limit
93	Which objects can be cloned? Multiple choice Choose all the correct answers	a. Tables b. Database c. Schema d. Named File Format e. Internal Stage	a. Tables b. Database c. Schema d. Named File Format

Check with Vinod

However,  
the  
following  
object types  
are not  
cloned:

				External tables
				Internal (Snowflake) stages
94	Which one is true about cloning?		A cloned object inherits the definition of the source object at the time the statement is executed	
95	What is the maximum size of Variant data type?	a. 1 MB b. 8 MB c. 16 MB d. 16 GB	c. 16 MB	
96	Copy command compresses only CSV and JSON files?	True or False	FALSE	
97	COPY INTO command requires File Format must be specified ?	True or False	FALSE	IT IS OPTIONAL
98	Where clause can be applied to which command while loading data into snowflake?	a. COPY INTO b. INSERT c. PUT d. GET	b. INSERT	Try out in snowflake once.
99	What is true about PUT Command? Multiple choice Choose all the correct answers	a. Loads the data into default internal stage b. Automatically compresses the data c. Loads the data with the default file format d. Automatically encrypts the data	b. Automatically compresses the data d. Automatically encrypts the data	Check with Vinod
100	Can Filter be applied in COPY command?	True or False	FALSE	Check once in Snowflake environment by executing the command. Add where condition
101	Maximum number of hours the query results can be held in the result cache	a. 12 b. 10 c. 24 d. 48	c. 24	

102	Can result cache be read when the virtual warehouse is not in active state?	True or False	TRUE	
103	Can metadata cache be read when the virtual warehouse is not in active state?	True or False	TRUE	
104	Can result set of the query executed by one user be accessed by other users?	True or False	FALSE	
105	Does snowflake support the conversion of Structured data to Semi Structured data?	True or False	TRUE	
106	Is Federated Authentication mandatory for all the users?	True or False	FALSE	
107	Is worksheet associated with the users of same roles and objects?	True or False	FALSE	
108	Snowflake Stored Procedures can be written in Multiple Choice	a. Java b. SQL c. JavaScript d. Python	b. SQL c. JavaScript	
109	Information Schema and Account Usage stores information about? Multiple Choice choose all the correct answers	a. Tables b. Database c. Stage d. User	ALL	
110	What are the roles can be created within Snowflake? Multiple Choice choose all the correct answers	a. SYSADMIN b. SECURITYADMIN c. PUBLIC d. ACCOUNTADMIN	ALL	
111	(Manan Thakar) Which is the minimum version which supports federated authentication ?	a) Standard b) Premier c) Enterprise d) Business Critical	a) Standard	
112	Which Statement from below describes Snowflake in best way ?	a) It is built specifically for cloud. B) Initially it was built for on Premise. Then got migrated on Cloud. C) It was originally built for Oracle.	a) It is built specifically for cloud.	

		a) Database b) Schema c) Table d) Named File Format e) Share f) Warehouse	a) Database b) Schema c) Table d) Named File Format	
113	Which of these objects can be cloned in Snowflake (Select Four)	a) Flat files (CSV, TSV) b) JSON c) Parquet d) ORC	a) Flat files (CSV, TSV) b) JSON c) Parquet	
114	Data unloading is supported to which of these file formats ? (Select Two)	a) CSV files (unzipped) b) CSV files (GZipped) c) Parquet d) ORC	a) CSV files (GZipped)	
115	Data loading is most performance-efficient for which file format ? (select one)	a) VARCHAR b) OBJECT c) VARIANT d) ARRAY	c) VARIANT	
116	Which is the most suitable datatype to store semi-structured data ?	a) JSON b) XML c) GZIP	a) JSON b) XML	
117	VARIANT datatype supports which of the data formats ? (Select Two)	a) 16 MB b) 800 KB c) 50 MB d) 4000 GB	a) 16 MB	
118	What is the maximum row size in Snowflake ?	a) Complexity of query b) "Data size to be processed by that query" c) "Tool on which the query is getting executed" d) "CONCURRENCY parameter set at account level"	a, d	
119	How many concurrent queries are allowed - it depends on.. (Select Two)	a) Table b) Schema c) Database d) Account	a) Table	
120	Time Travel period can be set at which most granular level ?	a) ACCOUNTADMIN b) ACCOUNTADMIN, SECURITYADMIN, SYSADMIN	c) ACCOUNTADMIN, SECURITYADMIN, SYSADMIN	
121	MFA (multi-factor Authentication) is recommended at least for which roles ?			

		c) ACCOUNTADMIN, SECURITYADMIN, SYSADMIN d) ACCOUNTADMIN, SECURITYADMIN, SYSADMIN, PUBLIC		
122	Changing the warehouse size in t-shirt sizing - is example of which type of scaling ?	a) scale-out b) scale up c) scale across	b) scale up	
123	ACCOUNTADMIN can disable MFA for any user and allow him/her to re-enroll for MFA.	True/False	TRUE	
124	Which are the options available while warehouse creation ?	a) Auto-Suspend b) Auto-Resume c) SSD size d) Memory size	a) Auto-Suspend b) Auto-Resume	
125	Which of these option best describes Snowflake architecture ? (Select Three)	a) Compute engine can be increased/decreased without need to increase Storage. b) Storage can be increased without increasing the compute. c) Compute and Storage are tightly coupled with each other. d) The storage can be accessed from any compute engine.	a, b, d	
126	Below are the characteristics of micro-partitions (Select Two)	a) Immutable b) Fully managed by Snowflake c) Logical/virtual object of storage d) Customer can select in which format data should be stored in micro-partitions	a) Immutable b) Fully managed by Snowflake	

127	Micro-partitions are insert-only. Those cannot be updated.	True/False	TRUE	
128	Types of accounts involved in Data sharing (Select Three)	a) Data Provider b) Data Consumer c) Reader Accounts d) Writer	a) Data Provider b) Data Consumer c) Reader Accounts	
129	Table data in micro-partitions are stored in the natural ordering of row insertion.	TRUE/FALSE	TRUE	
130	Which of these commands are non-blocking commands ? (Select Two)	a) INSERT b) UPDATE c) COPY d) MERGE	a) INSERT c) COPY	
131	Which is the minimum account-type to enable data-encryption in internal stage?	a) Standard b) Premium c) Enterprise d) Business Critical	a) Standard	
132	It is possible to query data from external stage files without loading them in table.	True/False	TRUE	
133	How many days of Fail-safe is applicable to tables?	a) 1 b) 7 c) 30 d) 90	b) 7	
134	Below is true about query result cache. (Select Two)	a) The query result cache can be used within 24 hours if the no change occurred in the respective database objects b) If the query result cache is used, then 24 hours time period to expire result cache - is reset. 24 hours will be counted from the point when the result cache was referenced last time. c) Within 24 hours, it can be used even if the	a	

		underlying database objects are changed.		
135	It is possible to execute few queries on Snowflake without active Warehouse.	True/False	TRUE	
136	Which of the below functions can be executed without using compute engine ? (Select One)	a) Min (Column) b) Avg(Column) c) Concat(Columns ) d) SUM(Columns)	a) Min (Column)	
137	Which of the connection methods supported by Snowflake (Select Three)	a) ODBC b) SnowSQL c) Perl-Connector d) Go Connector	a) ODBC b) SnowSQL d) Go Connector	
138	From where can you monitor Storage utilization by Account? (Select One)	a) On Web UI - in Account tab - Storage utilization b) "STORAGE_USAG E in information schema"	a) On Web UI - in Account tab - Storage utilization	
139	Which all the transformations are supported for loading data into table using COPY command ? (Select Two)	a) Filtering b) Column Reordering c) DatatypeCasting d) Aggregation	b, c	
140	Using Snowflake COPY command to load files into table, either the file will be fully loaded. Or, file load will completely failed. There is no option for partial load.	True/False	File can be partially loaded	

		a) RETURN_ALL_ERR ORS is a value for VALIDATION_MODE parameter. B) If it is set, the command will return the error records and load the valid records. C) If it is set, the command will return the error records but it will not load the valid records.		
141	Describe VALIDATION_MODE for COPY command (Select Two)		a, c	
142	On Web UI, one can select different Database, schema, role, warehouse for different worksheet.	True/False	TRUE	
143	Below format types are supported for Data unloading using COPY command. (Select Two)	a) CSV b) JSON c) Avro d) ORC	a) CSV b) JSON	
144	What happens when data Provider revokes an object access for a share ?	a) Immediately that object becomes inaccessible for all the data consumers. b) Consumers will not be able to see any updates/inserts on that object after that instance. But the data at that point in time will be still accessible. c) The current data will be copied to Consumer storage automatically and now the storage cost will be beared by consumer.	a) Immediately that object becomes inaccessible for all the data consumers.	

145	When a database is shared, its all the child objects with all object types will be shared automatically.	True/False	Seems false, need to provid the usage on schema	
146	Data sharing is supported for which account types?	a) Enterprise onwards b) All except VPS c) All accounts d) Premier onwards.	b) All except VPS	
147	The users who access Snowflake by SnowSQL, cannot login on Web UI without raising a specific request to support team.	True/False	FALSE	
148	In Snowflake, MFA can be implemented only in conjunction with Single Sign-on (SSO) functionality.	True/False	FALSE	
149	Data unloading is supported to which of these file formats ? (Select Two)	a) JSON b) CSV , TSV c) ORC d) XML	a) JSON b) CSV , TSV	
150	What is the maximum row size in micropartition in Snowflake ?	a) 16 MB b) 8 MB c) 4000 GB d) 50 MB	a) 16 MB	
151	Options available while creating the warehouse ?	a) Auto Resume b) Auto suspend c) Stop d) memory Size	a) Auto Resume b) Auto suspend	
152	Micro-partitions are insert-only or can be updated	True/False	TRUE	
153	Table data in micro-partitions are stored in the natural ordering of row insertion.	(True/False)	TRUE	
154	It is possible to query data from external stage files without loading them in table.	(True/False)	TRUE	
155	Which of the query can be executed without using compute engine ? (Select One)	a) Avg b) sum c) min d) concat	c) min	
156	Can snowflake web UI , provide the option to select different Database, schema,	(True/False)	TRUE	

	role, warehouse for different worksheet.			
157	possible file formats supported for Data loading using COPY command.	CSV/Tab, Json Parquet XML Avro ORC		
158	there was question on data sharing like data can be shared in which snowflake accounts		Except VPS	
159	Snowflake security and authentication including data encryption using snowflake customer keys	True/false	True	
160	snowflake supports which two type of mode in warehouse creation	Maximized Auto scale		
161	SQL functionality can be extended by which udf	a) Javascript b) sql udf c) java d) none	a) Javascript b) sql udf	
162	Can we partition the external table in snowflake	True/False	TRUE	
163	Question on when to use clustering key	a) Size 1 TB b) Queries are slower		
164	Which is the most powerful role on SF?		Accountadmin	
165	Virtual warehouse supports which of the following parameters:	a) auto_scale b) stop c) auto_resume d) auto_suspend	c) auto_resume d) auto_suspend	
166	When a virtual warehouse is resized all currently executing queries will start performing better?	True/False	False	
167	It is possible to share data with a Snowflake customer whose Snowflake instance exists in a different Region than the provider?		Yes, we need to do the replication first.	
168	Can we suspend and resume task?	True/False	TRUE	
169	Billing for snowpipe usage		per-second/per-core granularity actual data load.	
170	Which tables cannot be recovered in failsafe?		Temporary and Transient	

171	Can external table be partitioned?	True/False	TRUE	
172	Encryption is applicable to which min edition of SF.		Standard	
173	Can we query stage without loading it into SF	True/False	TRUE	
174	Data and storage are tightly coupled in SF.	True/False	FALSE	
175	Snowflake supports landing data into internal stage from which cloud storage platform.	Azure Blob GCP AWS S3 Internal Stage		
176	Which is the minimum version which supports 90 days time travel for permanent table?	Enterprise		
177	Which data type is supported for semi structured data?	Variant		
178	Changing WH size from small to medium is example of	Scale up		
179	Which logical structuring is correct out of 4	Account->Database->Schema->Table		
180	Reader account belong to	Provider		
181	Can consumer clone the shared schema?	True/False	FALSE	
182	default failsafe for permanent tables	1		
183	que on query result cache validity	24 hours.. max 31 days		
184	If a tool is connected to SF via snowsql client, does it need to ask SF support to enable access to SF UI.	True/False	FALSE	
185	Data sharing is not possible for which SF edition?	VPS		
186	Multiple choice questions on Roles and Users			
187	Multiple choice Question on MFA			
188	Reader Account Billing Related Questions			1. If Consumer does not have the reader account then storage and compute will be beared by

				the Provider. 2. If the data is shared with person having snowflake account then consumer will bear the compute cost.
189	Few questions related to Web UI, Semi Structured data loading, data sharing, cloning, DDL, DML, Security policies, Clustering, Micro Partitioning, Multi cluster warehousing, Copy Command options, Resource monitor etc.			
190	Question around Schema and Database <b>Reference from Snowflake documentation:</b> <ul style="list-style-type: none"><li>• A database is a <b>logical</b> grouping of schemas. Each database belongs to a single Snowflake account.</li><li>• A schema is a <b>logical</b> grouping of database objects (tables, views, etc.). Each schema belongs to a single database.</li></ul>			
191	When sizing a warehouse which of the factors are considered			
192	question on snowpipe vs copy into and difference in pricing			
193	Question on json parsing, gave sample json and asked how to access the particular element via snowflake query /notation			

**Snowflake does not start executing any queries on a new virtual warehouse until all of the servers are provisioned.**

- 
- **False**
- 
- **True**
- **(Correct)**

### **Explanation**

Correct. The SQL execution starts only once all the servers are provisioned. If provisioning fails for any reason, Snowflake will attempt to fix the failed servers and SQL will start executing once 50% or more virtual servers are provisioned

Question 2: **Correct**

**You can view the storage used by failsafe under the Account->Billing & Usage section**

- 
- **True**
- **(Correct)**
- 
- **False**

### **Explanation**

<https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

Question 3: **Correct**

**Which of the following are options when creating a new virtual warehouse. Select all that apply**

- 
- **auto\_resume**
- **(Correct)**

- resume\_on\_demand**
  - suspend\_and\_resume**
  - auto\_suspend**
- (Correct)**

#### **Explanation**

Please see the below link for detailed syntax on virtual warehouse creation

<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

---

Question 4: **Incorrect**

**Which one of the following can not be cloned?**

- table stage**
  - internal named stage**
  - external named stage**
- (Correct)**

#### **Explanation**

Internal named stages can not be cloned

---

Question 5: **Correct**

**The Snowflake credit usage for multi cluster virtual warehouses is based on**

- **the number of clusters running during each hour for the warehouse is running**

**(Correct)**

- **the value that is configured as the maximum in the multi cluster virtual warehouse configuration**

### **Explanation**

The billing for multi cluster virtual warehouses is based on a) the size of the multicluster virtual warehouse b) the number of clusters that were active

---

Question 6: **Incorrect**

**Snowflake uses a unique architecture in which data and compute has been decoupled and both can be scaled independent of each other.**

- 

**Yes**

**(Correct)**

- 

**No**

**(Incorrect)**

---

Question 7: **Incorrect**

**Which of the following are options when created a new multicluster virtual warehouse. Select all that apply**

- 

**SCALING\_POLICY**

**(Correct)**

- 

**MIN\_CLUSTER\_COUNT**

**(Correct)**

- 

**MAX\_CLUSTER\_COUNT**

**(Correct)**

- 

**SUSPEND\_AND\_RESUME**

### **Explanation**

Please see the below link for detailed syntax on virtual warehouse creation

<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

---

Question 8: **Incorrect**

**What happens when a table is shared by a snowflake account to another snowflake account? Select all that applies**

- 

**Sharing is managed through Snowflake metadata services layer**

**(Correct)**

- 

**No actual data is copied or transferred between accounts**

**(Correct)**

- 

**the target snowflake account is charged for the shared data storage**

**(Incorrect)**

- 

**data is copied to the target snowflake account**

### **Explanation**

In Snowflake sharing, no data is copied. It is just the metadata that enables the sharing of data. Since no data is copied, the target snowflake account (also called the

consumer) is not charged for any storage <https://docs.snowflake.net/manuals/user-guide/data-sharing-intro.html#how-does-secure-data-sharing-work>

---

Question 9: **Incorrect**

**The "load metadata" (which maintains which files have already been loaded) for a table expires after how many days?**

- 64 **(Correct)**
- 365
- 30 **(Incorrect)**
- 1

### Explanation

The load metadata expires after 64 days. Please see below for more details  
<https://docs.snowflake.net/manuals/user-guide/data-load-considerations-load.html#load-metadata>

---

Question 10: **Incorrect**

**Which of the following is true regarding Roles in Snowflake?**

- **Users in Snowflake can only have a maximum of 10 Roles assigned.**
- **A Snowflake User can have one or many Roles.** **(Correct)**
-

**In addition to Roles, privileges may be directly assigned to a named User.**

**(Incorrect)**

- 

**For an active user session, only one Role can be the active Role at a given time.**

**(Correct)**

### **Explanation**

Users may be assigned one or more Roles, however, for a session only one role can be active at a time. Users can switch roles by using the USE ROLE <rolename> command.

Please see also

<https://docs.snowflake.com/en/user-guide/security-access-control-overview.html>

Question 11: **Correct**

**You have a virtual warehouse named reporting\_warehouse which is sized as Medium. You need the size of the virtual warehouse to be Large for future queries. What is the best way to make this change.**

- 

**Alter the existing virtual warehouse and change the size to Large**

**ALTER WAREHOUSE reporting\_warehouse WAREHOUSE\_SIZE=large**

**(Correct)**

- 

**Drop the existing virtual warehouse and create a new one with the required parameters**

- 

**It is not permitted to re-size a virtual warehouse once it has been created**

### **Explanation**

Please see the below link for detailed syntax <https://docs.snowflake.net/manuals/sql-reference/sql/alter-warehouse.html>

Question 12: **Incorrect**

You are a developer with a retail organization running a Snowflake data warehouse. You required to load JSON based data into a table.

What approach should you take?

- 

**Use Snowflake provided function to process JSON data while loading it into the table.**

**(Correct)**

- 

**Write a program using a programming language of your choice (Python, Java etc..) to process the JSON file into a CSV structure.**

**Load the CSV file into the table using the COPY command.**

**(Incorrect)**

- 

**Ask the source system to send you a CSV file rather than a JSON file as Snowflake doesn't support JSON.**

**Explanation**

Snowflake provides functionality to process JSON data, therefore the best course of action is to use the Snowflake built in functionality

Question 13: **Correct**

**For a snowflake session, more than one virtual warehouse can be specified at a time for executing a query.**

- 

**True**

- 

**False**

**(Correct)**

## Explanation

Incorrect, a snowflake session can only have one virtual warehouse specified, though you can change the current virtual warehouse for a session

---

Question 14: **Incorrect**

**Snowflake deploy new release at what frequency?**

- 

**monthly**

- 

**every week**

**(Correct)**

- 

**every 2nd week**

**(Incorrect)**

## Explanation

Snowflake releases & patch release are deployed every week

---

Question 15: **Incorrect**

In the Snowflake Web UI which of the following is true?

- 

**It doesn't matter how many sheets you open in the Snowflake Web UI, they share a single session**

- 

**Each open worksheet in the Snowflake Web UI only a single Role can be selected which applies to all worksheets**

- 

**Each worksheet in the Snowflake Web UI is a different user session**

**(Correct)**

-

**Each open worksheet in the Snowflake Web UI can have a different Role selected**

**(Correct)**

### **Explanation**

Each worksheet in Snowflake Web UI is an independent session. Therefore, each worksheet can have its combination of Database, Schema, Virtual Warehouse & Role selected

---

Question 16: **Incorrect**

**A virtual warehouse can resized at any time whether they are running or stopped**

- 

**True**

**(Correct)**

- 

**False**

**(Incorrect)**

### **Explanation**

Correct. The resize operation can be done any time, even when it is running queries.

---

Question 17: **Correct**

**When a virtual warehouse is provisioned there is no benefit in stopping the virtual warehouse with in the first 60 seconds because that time period has already been billed.**

- 

**False**

- 

**True**

**(Correct)**

### **Explanation**

Correct. When a virtual warehouse is provisioned the minimum billing is 1 minute, after which the virtual warehouses is billed on per second basis

---

Question 18: **Correct**

**On the WebUI which button on the top bar should you select to get to tabs that provide information about virtual warehouse usage?**

- 

**History**

- 

**Account**

**(Correct)**

- 

**Databases**

- 

**Warehouses**

- 

**Shares**

### **Explanation**

The Account button will take you to screen which has multiple sub tabs, one of them providing virtual warehouse compute usage information

---

Question 19: **Incorrect**

**What method does Snowflake use to limit the number of micro-partitions accessed during a query?**

- 

**Indexing**

- Pruning  
**(Correct)**
- 
- table scan**
- 
- In memory operations**  
**(Incorrect)**

### Explanation

Partition elimination or pruning is used to limit the number of micro partitions.

---

Question 20: **Incorrect**

To clone a table your current role must have which privilege on the source table?

- SELECT**  
**(Correct)**
- 
- WRITE**
- 
- USAGE**  
**(Incorrect)**

### Explanation

To clone a table you need a SELECT privileges on the source table and for all other objects that can be cloned you need the USAGE privilege

---

Question 21: **Incorrect**

The privileges provided by the SYSADMIN & SECURITYADMIN role are automatically contained in the ACCOUNTADMIN role since the ACCOUNTADMIN role sits on the top of the role hierarchy.

- 

**False**

**(Incorrect)**

- 

**True**

**(Correct)**

### **Explanation**

Correct. The ACCOUNTADMIN role is the all powerful role with all privileges.

---

Question 22: **Correct**

**Which role is most powerful role in a Snowflake system?**

- 

**PUBLIC**

- 

**SYSADMIN**

- 

**SECURITYADMIN**

- 

**ACCOUNTADMIN**

**(Correct)**

### **Explanation**

Account admin is the top level role with most power, it should ideally be shared with least amount of people

---

Question 23: **Correct**

**A consumer of a shared database can add new tables or views in the shared database?**

- 

**False**

**(Correct)**

- 

**True**

### **Explanation**

Shared objects are read-only and can not be modified

---

Question 24: **Incorrect**

**Snowflake can load from data staged in which of the following?**

- 

**Google Cloud Storage**

**(Correct)**

- 

**VMWare Storage**

- 

**Oracle Cloud Storage**

- 

**Internal Stage**

**(Correct)**

- 

**Azure Blob Storage**

**(Correct)**

-

**AWS S3**

**(Correct)**

### **Explanation**

Snowflake supports loading from Internal Stage, AWS S3, Azure Blob and Google Cloud Storage. Please see below link for more details

<https://docs.snowflake.net/manuals/user-guide/data-load-overview.html#loading-from-data-files-staged-on-other-cloud-platforms>

Question 25: **Correct**

**What is the lowest snowflake edition that allows customer dedicated virtual warehouses?**

- 

**Business Critical**

- 

**Standard**

**(Correct)**

- 

**Enterprise**

### **Explanation**

All editions of Snowflake provide customer dedicated virtual warehouses, therefore the lowest edition providing such functionality is Standard

Question 26: **Correct**

**On the WebUI which button on the top bar should you select to see the query history?**

- 

**Shares**

- 

**Databases**

- 

**Virtual Warehouses**

- -
- Account**
- -
- History**
- (Correct)**
- -
- Worksheets**
- 
- Question 27: **Correct**
- As a consumer you can create only one database per share?**
- -
- False**
- -
- True**
- (Correct)**
- Explanation**
- Correct. While as a consumer you may have multiple shares, however for each share you can create only one database.
- 
- Question 28: **Correct**
- When data is shared between Snowflake accounts, a database is created on the consumer side for sharing purposes. The type of this database is ?**
- -
- Open Access**
- -

### Writable

- 
- 

### Temporary

- 
- 

### Read-only

(Correct)

- 
- 

### Permanent

#### Explanation

The correct answer is read-only. The consumer creates a database from the share and the database is created as read-only

---

Question 29: **Incorrect**

**Multi cluster virtual warehouses are designed to handle ... Select all that apply**

- 

**Large number of concurrent users**

(Correct)

- 

**Large Complex Queries**

(Incorrect)

- 

**Large number of concurrent queries**

(Correct)

- 

**Queuing issues**

(Correct)

## Explanation

The primary purpose of Multi cluster virtual warehouses is to improve queueing and concurrency issues <https://docs.snowflake.net/manuals/user-guide/warehouses-considerations.html#multi-cluster-warehouses-improve-concurrency>

---

Question 30: **Incorrect**

**What is the correct way to find out virtual warehouse credit usage information in Snowflake?**

- 

**Query WAREHOUSE\_METERING\_HISTORY in the Information Schema**

**(Correct)**

- 

**Query WAREHOUSE\_CREDITS\_USED in the Account Usage Schema**

**(Incorrect)**

- 

**Go to Account -> Billing & Usage in the WebUI**

**(Correct)**

- 

**Query WAREHOUSE\_METERING\_HISTORY in the Account Usage Schema**

**(Correct)**

- 

**Query WAREHOUSE\_CREDITS\_USED in the Information Schema**

## Explanation

All 3 are valid ways of finding virtual warehouse credit usage

<https://docs.snowflake.net/manuals/user-guide/credits.html#viewing-warehouse-credit-usage-for-your-account>

---

Question 31: **Incorrect**

**When loading data through COPY command it is a requirement that your table and the file from which the data is being loaded should have same order of columns.**

- 

**True**

**(Incorrect)**

- 

**False**

**(Correct)**

### **Explanation**

The order & the number of columns in the file and the table can be different

---

Question 32: **Incorrect**

**Both time travel & failsafe require extra storage which has cost associated with it.**

- 

**False**

**(Incorrect)**

- 

**True**

**(Correct)**

### **Explanation**

Correct. Both options use additional storage which will result in additional costs

<https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

---

Question 33: **Correct**

As a data engineer you are developing jobs to load data into a snowflake table. You have an S3 stage defined which contains a single file containing 100 rows.

You managed to load those 100 rows but while performing further development you notice that your COPY command is executing successfully but is loading zero rows into the target table.

What could be the reason?

- 
- You have already loaded the single file that was in the Snowflake stage. Snowflake tracks if a file has already been loaded and will not load it again.**
- (Correct)**
- 
- The file format that you have defined for your data is incorrect, therefore no data is getting loaded.**
- 
- The file in the Snowflake stage is corrupt and therefore can't be loaded.**

#### Explanation

Since you have already loaded this file, Snowflake will track through metadata which files have been loaded already and will not process them again. If you want to retest your copy command you must place a new file or clear the metadata

---

Question 34: **Correct**

**A virtual warehouse can be suspended or resumed as required.**

- 
- No**
- 
- Yes**
- (Correct)**

#### Explanation

A virtual warehouse can be suspended and resumed as per required. This capability allows snowflake to have zero, one or more virtual warehouses executing against the same data.

---

Question 35: **Correct**

**Which of the following are valid parameters when create a new virtual warehouse?**

- **CPU\_COUNT**
- **RAM\_SIZE**
- **AUTO\_RESUME**  
**(Correct)**
- **WAREHOUSE\_SIZE**  
**(Correct)**
- **AUTO\_SUSPEND**  
**(Correct)**

### **Explanation**

Please see the below link for detailed syntax on virtual warehouse creation

<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

---

Question 36: **Incorrect**

**Snowflake database is based on the massively parallel shared nothing architecture used by data bases like Teradata, Greenplum.**

- 

**Yes**

**(Incorrect)**

- 

**No**

**(Correct)**

### Explanation

Snowflake architecture is an hybrid architecture, in which unlike shared nothing architectures it uses a single shared storage and multiple compute engines operating on the same storage.

---

Question 37: **Incorrect**

**Which of the following will contribute to your account level storage usage?**

- 

**Snowflake External Stages**

**(Incorrect)**

- 

**Views**

**(Incorrect)**

- 

**User Defined Functions**

**(Incorrect)**

- 

**Snowflake Internal Stage**

**(Correct)**

- 

**Database (and the tables in it)**

**(Correct)**

### Explanation

Storage used by a database (and its tables) and storage used by Snowflake internal stages contribute towards the overall storage usage

---

Question 38: **Incorrect**

**Multifactor Authentication can be enabled for which of the following? Select all that Apply?**

- 

**SnowPipe**

**(Incorrect)**

- 

**ODBC**

**(Correct)**

- 

**Python**

**(Correct)**

- 

**Snowflake WebUI**

**(Correct)**

### **Explanation**

Please see the below link for details <https://docs.snowflake.net/manuals/user-guide/security-mfa.html>

---

Question 39: **Incorrect**

**What is correct about multi cluster virtual warehouses?**

- 

**They will add additional clusters automatically when under high query workload which is beyond the current capacity.**

**(Correct)**

- 

**They will remove the additional clusters as query demand decreases.**

**(Correct)**

- 

**Multi cluster virtual warehouses can not be suspended or resumed.**

- 

**Only one multi cluster virtual warehouse can be created per Snowflake customer.**

#### **Explanation**

A multi-cluster virtual warehouse will scale up and scale down depending on the query & user demand.

---

Question 40: **Correct**

**On the WebUI which button on the top bar should you select to get to tabs that provide billing and credit information?**

- 

**Account**

**(Correct)**

- 

**Restricted**

- 

**Account Administrator**

#### **Explanation**

The Account button will take you to screen which has multiple sub tabs, one of them providing billing and credit information

---

Question 41: **Correct**

You are the data modeler at a large retail organization who store their transactional data in a Snowflake table called "Transactions". The daily revenue reports execute using the "Transactions" table and calculate the revenue for the current day.

The "Transactions" table had 500GB of data originally but it has now grown to be 5TB. You have noticed that over time the performance of the daily revenue reports has degraded.

What is the most efficient & cost-effective way of optimizing performance?

- 

**Increase the size of the virtual warehouse executing the daily reports. That will ensure that the reports execute much more faster.**

- 

**Cluster the Transactions table on the transaction date column. The daily report will only access one day at a time and therefore will only scan required data.**

**(Correct)**

- 

**Create tables for each year of data e.g. Transactions\_2017, Transactions\_2018 and Transactions\_2019. Insert the relevant data from the Transactions table into these \_year tables. Change your report to point to Transactions\_2019 table.**

**Now that the daily revenue report is accessing only a year worth of data, performance will be better.**

### **Explanation**

Clustering keys is the most efficient option. The query will only scan a single day worth of data using partition pruning. (there is a cost associated with reorganization of data when you add a cluster key)

Question 42: **Correct**

**It is not possible to suspend a specific cluster in a multi-cluster virtual warehouse but rather the whole virtual warehouse is suspended.**

- 

**True**

**(Correct)**

- 

**False**

### Explanation

True. Individual clusters in a virtual warehouse can not be controlled independently, but rather they appear logically as a virtual warehouse together.

---

Question 43: **Correct**

**Assuming you have multiple worksheets open in Snowflake Web UI with in the same browsing session, you can have different role specified in each of the worksheet?**

- 

**True**

**(Correct)**

- 

**False**

### Explanation

Correct, you can specify different roles in different worksheets

---

Question 44: **Incorrect**

**Users created in a reader account can not do which of the following?**

- 

**INSERT**

**(Correct)**

- 

**SELECT**

- 

**UPDATE**

**(Correct)**

- 

**LOAD DATA**

**(Correct)**

## Explanation

Users in a reader account can query shared data but can not perform any DML.

---

Question 45: **Correct**

**Which minimum Snowflake license allows Multicloud data virtual warehouse capability?**

- 

**Standard**

- 

**Business Critical**

- 

**Enterprise**

**(Correct)**

## Explanation

The multi cluster virtual warehouse capability is available starting from the Enterprise Edition. The capability is not available in Standard or Premier

---

Question 46: **Correct**

You are the solution architect for a large retail company running a Snowflake data warehouse. Your Snowflake implementation has just gone live with a single virtual warehouse and has been opened up to all the users across the organization.

One of your heaviest users are the Finance department, who have a large number of users and execute a large number of queries. The finance department has complained that the queries take a long time to execute for them.

What is the best immediate action should you take to improve their experience?

- 

**Introduce a dedicated virtual warehouse instance for the finance department and size is according to their needs.**

**Set the new virtual warehouse to auto suspend and auto resume.**

**(Correct)**

- 

**Increase the size of your single virtual warehouse to the maximum size available so that queries for all users (not just finance) execute faster.**

- 

**Limit the number of queries that the finance department can execute.**

### **Explanation**

A dedicated virtual warehouse is the best short term option to optimize performance

---

Question 47: **Correct**

**What are the virtual warehouse sizing approach recommended by Snowflake?**

- 

**Execute your queries on the largest available virtual warehouse and keep scaling down the size of the virtual warehouse, until you start seeing degradation in query response**

- 

**Experiment with different types of query and virtual warehouses to determine the virtual warehouse size that best meets your requirements**

**(Correct)**

- 

**Execute your queries on the smallest available virtual warehouse and keep scaling up the size of the virtual warehouse, until the query performance no longer shows improvements**

### **Explanation**

Snowflake recommended approach is to experiment to arrive at the best size

<https://docs.snowflake.net/manuals/user-guide/warehouses-considerations.html>

---

Question 48: **Correct**

**Which minimum Snowflake license provides up to 90 days of time travel?**

- 

**Business Critical**

- 
- Standard**
- 
- Enterprise**
- (Correct)**

### Explanation

Enterprise is the lowest edition that provides 90 days of time travel

---

Question 49: **Correct**

**For a multi cluster virtual warehouse the administrator set the minimum cluster count and the maximum cluster count to different values. This can be best described as ?**

- 
- Scaling up**
- 
- Auto Scaling**
- (Correct)**

### Explanation

Setting the minimum and maximum to different values will ensure that the multi cluster virtual warehouse scales up in periods of high demand and scale back down in periods of less demand

---

Question 50: **Incorrect**

**Under the Download section in Snowflake Web UI, what are the connectors & programs that can be downloaded? Select All that apply**

- 
- SnowSQL**
- (Correct)**
-

**ODBC driver**

**(Correct)**

- 

**JDBC driver**

**(Correct)**

- 

**Snowpipe Connector**

- 

**Load Connector**

### **Explanation**

Please see download page under Help

---

Question 51: **Incorrect**

**Snowflake uses a staged release process for new releases in which different account types are upgraded gradually.**

- 

**True**

**(Correct)**

- 

**False**

**(Incorrect)**

### **Explanation**

Correct. The staged release process through which Snowflake releases updates is first applied to early access account, then standard accounts and finally Enterprise & upwards

---

Question 52: **Correct**

Which of the following Snowflake Editions allow only a maximum of 1 day of time travel?

- **Standard**  
**(Correct)**

- **Business Critical**
- 

### **Enterprise** **Explanation**

Standard edition allows only 1 day of time travel.

<https://www.snowflake.com/pricing/pricing-guide/>

Question 53: **Incorrect**

**What are some general indicators that a clustering key is required on a table?  
Select all that apply.**

- **The size of the table is in multi terabytes**  
**(Correct)**
- **The query performance has slowed down over time**  
**(Correct)**
- **The table has large number of columns**  
**(Incorrect)**

---

Question 54: **Correct**

**You can increase or decrease the size of a virtual warehouse as your user base & query demand increase or decrease?**

-

**False**

- 

**True**

**(Correct)**

---

Question 55: **Incorrect**

**What are the resources that a virtual warehouse provides? Select all that apply.**

- 

**Temporary Storage**

**(Correct)**

- 

**CPU**

**(Correct)**

- 

**Memory**

**(Correct)**

- 

**User Management**

**Explanation**

Virtual warehouse provide CPU, Memory & temporary storage to process queries

---

Question 56: **Correct**

**When a file has been loaded into a table, Snowflake marks that file as loaded in metadata so that the file doesn't get processed again.**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Snowflake maintains file load metadata, through which it tracks which files have already been processed

---

Question 57: **Incorrect**

**Virtual Warehouse credits are charged on what basis?**

- 

**per minute of use**

- 

**per milli second of use**

- 

**per second of use**

**(Correct)**

- 

**per hour of use**

**(Incorrect)**

### **Explanation**

Snowflake credits are billed on per second basis of virtual warehouse use. Its important to note thought that when a virtual warehouse is started it will consume minimum 1 minute of snowflake credits and will change to per second billing after that

---

Question 58: **Correct**

**Files that have already been processed into the source table can be loaded again into a cloned table?**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Correct. The cloning copy structure, data and certain other attributes but doesn't copy the load metadata, therefore files that have already been processed for the source table can be re-processed into a cloned table.

Question 59: **Correct**

**If retrieval of data from fail safe is required, you must**

- 

**use SQL to retrieve the data yourself**

- 

**contact Snowflake support team to facilitate the retrieval**

**(Correct)**

### **Explanation**

Failsafe data can not be retrieved by Snowflake customer and support must be engaged to retrieve this data.

Question 60: **Correct**

**A database is created from a share by the consuming account, the access to this database is configurable using the role-based access control that Snowflake provides?**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Correct. The same role based access is applicable to a shared database

Question 1:

Skipped

When a virtual warehouse is resumed from suspended state what is the minimum amount of Snowflake credits that it will immediately consume?

- 

**1 second**

- 

**30 minutes**

- 

**1 minute**

**(Correct)**

- 

**30 seconds**

### Explanation

When a virtual warehouse is started it will consume minimum 1 minute of snowflake credits and will change to per second billing after that

---

Question 2:

Skipped

Which of the following are characteristics of Snowflake releases?

- 

**no down time**

**(Correct)**

- 

**transparent to user**

**(Correct)**

-

**allow user to choose whether to upgrade or not**

- 

**require restart**

### **Explanation**

Snowflake releases are transparent to the user and don't cause any downtime.

---

Question 3:

**Skipped**

Which of the following is required in order for a role to be able to successfully execute a Task?

- 

**The role should have the global EXECUTE TASK privilege on the Account**

**(Correct)**

- 

**The role should have the permissions required to run the SQL statement executed in the Task**

**(Correct)**

- 

**The role should have SYSADMIN privileges**

- 

**The role should have ACCOUNTADMIN privileges**

### **Explanation**

To execute a task, the role being used must have the global EXECUTE TASK privilege on the Account. Additionally, all the privileges that the SQL statement in the task requires must be available to the role.

---

<https://docs.snowflake.com/en/user-guide/tasks-intro.html>

---

Question 4:

**Skipped**

Which type of short lived Snowflake tables will continue to exist even if the session is closed?

- Permanent
- 
- Temporary
- 
- Clustered
- 
- Transient
- (Correct)

#### Explanation

Transient tables can be used as short lived tables for ETL work tables and are not dropped when the session is closed

---

Question 5:

Skipped

You are planning to utilize the multi-cluster virtual warehouse to provide auto-scaling and performance for your users. You want to make sure that the auto scaling doesn't start clusters immediately and is cost effective. What is the correct value for Scaling Policy to achieve this?

- 
- SCALING\_POLICY = PERFORMANCE
- 
- SCALING\_POLICY = STANDARD
- SCALING\_POLICY = ECONOMY
- (Correct)

#### Explanation

ECONOMY scaling policy will not auto scale immediately but will be conservative in scaling See link for further details <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#label-mcw-scaling-policies>

---

Question 6:

**Skipped**

**Snowflake storage capacity can be pre-purchased for a lower price?**

- 

**True**

**(Correct)**

- 

**False**

#### **Explanation**

True, you can pre-purchase storage capacity which will generally be of lower price compared to on demand pricing which is as per your usage

---

Question 7:

**Skipped**

**The COPY command in a Snowpipe definition supports the same transformation as provided by the usual COPY command.**

- 

**True**

**(Correct)**

- 

**False**

#### **Explanation**

Correct. The same transformations are available in a Snowpipe definition.

---

Question 8:

**When processing semi structured data into structured (i.e. a table) what is the correct way to cast a column into a data type? Assume the target column name is CustomerName and the data type is String**

- 

**SELECT col1:CustomerName::String**

**(Correct)**

- 

**SELECT CAST(col1:CustomerName AS String)**

- 

**SELECT col1 AS CustomerName WITH DATATYPE AS String**

### Explanation

Please see below link for more details <https://docs.snowflake.net/manuals/user-guide/semistructured-considerations.html>

Question 9:

**Under what circumstances results for a query will be fulfilled by the query result cache? Select all that apply**

- 

**The query results cache was generated was generated or last used less than 24 hours ago**

**(Correct)**

- 

**The underlying data for the tables in the query has not changed**

**(Correct)**

- 

**The query is being executed from the same virtual warehouse as the previously executed query**

-

**The query being run is syntactically identical to a previously executed query**

**(Correct)**

- 

**The query being run returns result set of less then 100 thousand rows**

### **Explanation**

The underlying data must not change and the query should be syntactically identical in order for the query result cache to be used. The cache must have been generated (or last used) less than 24 hours ago. <https://docs.snowflake.net/manuals/user-guide/querying-persisted-results.html#retrieval-optimization>

---

Question 10:

**Which of the following are valid Casting function available in Snowflake?**

- 

**TO\_TIMESTAMP**

**(Correct)**

- 

**TO\_NUMERIC**

**(Correct)**

- 

**TO\_VARIANT**

**(Correct)**

- 

**TO\_NUMBER**

**(Correct)**

- 

**TO\_BLOB**

Question 11:

Which of the following properties can be set for a multi cluster virtual warehouse? Select all that apply?

- **AUTO\_SCALE**
- **AUTO\_RESUME**  
**(Correct)**
- **MIN\_CLUSTER\_COUNT**  
**(Correct)**
- **MAX\_CLUSTER\_COUNT**  
**(Correct)**
- **AUTO\_SUSPEND**  
**(Correct)**

### Explanation

Auto scale is not a valid property. Other than auto scale all other given properties are valid

---

Question 12:

Snowflake allows which ways to scale a virtual warehouse? Select all that apply

- **Exponential Scaling**
- **Linear Scaling**

- 

**Scale Up (increasing the size of a virtual warehouse)**

**(Correct)**

- 

**Scale Out (adding clusters to a multi cluster virtual warehouse)**

**(Correct)**

### **Explanation**

The two scaling methods are scale up and scale out

---

Question 13:

**Snowflake SQL UDFs can return either Scalar or Tabular results?**

- 

**False**

- 

**True**

**(Correct)**

### **Explanation**

Either type of UDFs (JavaScript or SQL) can return Scalar or Tabular results

---

<https://docs.snowflake.net/manuals/sql-reference/user-defined-functions.html>

Question 14:

Multifactor Authentication can be enabled for which of the following?

Select all that Apply?

- 

**SnowPipe**

- 

**Snowflake WebUI**

**(Correct)**

- 

**JDBC**

**(Correct)**

- 

**SnowSQL**

**(Correct)**

### **Explanation**

Please see the below link for details <https://docs.snowflake.net/manuals/user-guide/security-mfa.html>

---

Question 15:

**You increased the size of a virtual warehouse to improve performance of a query? What is this an example of?**

- 

**Scale Out**

- 

**Scale Up**

**(Correct)**

### **Explanation**

Increasing the size of a virtual warehouse will enable faster processing for queries but will not necessarily accommodate more queries and more users, therefore this is an example of scaling up

---

Question 16:

**Snowflake patch releases are applied to all accounts at the same time.**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Correct. The patch releases are applied to all accounts at the same time. The staged released process is only used for new software releases.

---

Question 17:

**What is one of the way to improve performance in Snowflake?**

- 

**Query Hints**

- 

**Join Indices**

- 

**Secondary Indices**

- 

**Clustering Keys**

**(Correct)**

### **Explanation**

Clustering keys is the performance optimization option provided by Snowflake.

---

Question 18:

**You are the Snowflake Administrator for a large telecom company. You had a Medium sized virtual warehouse on which a complex query was executing, you change the virtual warehouse size to be Small. Will the query be impacted and perform slower?**

- 

**No. The currently executing query will keep on executing on the original cluster. Only new queries will execute on the Small sized cluster.**

**(Correct)**

-

**Yes. The query will now start using the Small virtual warehouse, therefore will be slower**

### **Explanation**

The query which is already executing will continue using the resources it has acquired, only new queries will be impacted by the change in size

---

Question 19:

**Under the Download section in Snowflake Web UI, what are the connectors & programs that can be downloaded? Select All that apply**

- **SnowSQL**  
**(Correct)**
- **JDBC driver**  
**(Correct)**
- **ODBC driver**  
**(Correct)**
- **Snowpipe Connector**
- **Load Connector**

### **Explanation**

Please see download page under Help

---

Question 20:

**What are some of the ways to improve performance in Snowflake? Select all that apply**

- 

**Query Result Caching**

**(Correct)**

- 

**Dedicated Virtual Warehouse**

**(Correct)**

- 

**Clustering Keys**

**(Correct)**

- 

**Multi Clustering Virtual Warehouse**

**(Correct)**

- 

**Secondary Indices**

---

Question 21:

**When a virtual warehouse is resized to a SMALLER size, the required servers are only removed when ?**

- 

**there are no SQL statements running on the servers**

**(Correct)**

- 

**immediately**

- 

**after 1 minute**

### **Explanation**

The servers are removed only when they are not being used by a statement  
<https://docs.snowflake.net/manuals/user-guide/warehouses-tasks.html>

---

Question 22:

**On the Snowflake Web UI, which of the following are buttons in the top bar?**  
**Select all that apply**

- 

**Roles**

- 

**Databases**

**(Correct)**

- 

**Security**

- 

**Shares**

**(Correct)**

---

Question 23:

**A virtual warehouse can be set to auto-suspend or auto-resume but not both.**

- 

**False**

**(Correct)**

- 

**True**

## Explanation

You can set a virtual warehouse to auto suspend and auto resume as per your needs.

Question 24:

**Which type of Snowflake table do not have a failsafe?**

- 

**Permanent**

- 

**Temporary**

**(Correct)**

- 

**Transient**

**(Correct)**

- 

**Clustered**

## Explanation

Transient and temporary tables don't have any failsafe, this is done to reduce storage costs for temporary and transient data. <https://docs.snowflake.net/manuals/user-guide/data-cdp-storage-costs.html>

Question 25:

**In a multi cluster virtual warehouse, if you set different values for minimum cluster count and the maximum cluster count**

- 

**Snowflake starts minimum required clusters in the multi-cluster virtual warehouse**

**(Correct)**

- 

**Snowflake starts all clusters in the multi-cluster virtual warehouse**

-

**The mode is called auto-scale**

**(Correct)**

- 

**The mode is called Maximized**

### **Explanation**

If minimum and maximum value are set to different values the multi cluster virtual warehouse will be in auto scale mode <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

---

Question 26:

**A Snowflake share ....**

- 

**is a container that contains objects to be shared**

**(Correct)**

- 

**is only available to enterprise customers**

- 

**has information about the consumer(s)**

**(Correct)**

### **Explanation**

A share contains USAGE privileges on the database & the schema to be shared, privileges on the tables, secure views which will be shared and the consumer account(s) to which the share will be available. It is available to all Snowflake editions

---

Question 27:

**On the WebUI which button on the top bar should you select to get to tabs that provide information about storage usage?**

- 

**Shares**

- 

### Virtual Warehouses

- 

### Worksheets

- 

**Account**

**(Correct)**

- 

### Databases

- 

### History

#### Explanation

The Account button will take you to screen which has multiple sub tabs, one of them providing account level storage usage information

---

Question 28:

**Micro partitions can overlap in their range of values?**

- 

**False**

- 

**True**

**(Correct)**

#### Explanation

Micro partitions can overlap with each other in their range of values. This can help reduce skew in partition sizes <https://docs.snowflake.net/manuals/user-guide/tables-clustering-micropartitions.html>

---

Question 29:

What are the types of transformations which are available when loading data into a table using the COPY command. Select all that apply.

- 

Pivot Rows into Columns

- 

Re-order Columns

(Correct)

- 

Cast Columns

(Correct)

- 

Concatenate Columns

(Correct)

- 

Omit Columns

(Correct)

Question 30:

What happens when different value is specified for the minimum & the maximum cluster? Select all that apply

- 

Snowflake starts all clusters in the multi-cluster virtual warehouse

- 

Snowflake starts minimum required clusters in the multi-cluster virtual warehouse that can accommodate the number of users & number of queries

(Correct)

- 

If the number of concurrent users and queries increase on the virtual warehouse, additional clusters are started, up to the defined maximum cluster value

(Correct)

- 

If the demand is decreased the number of clusters are gradually shutdown, up to the defined minimum cluster value

(Correct)

- 

The mode is called auto-scale

(Correct)

### Explanation

For an auto scale multi cluster virtual warehouse 1) only minimum required clusters are started 2&3) The running clusters may be scaled up and down within the maximum and minimum cluster defined <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

### Question 31:

Using Snowpipe customers can load data continuously from which of the following?

- 

External Tables

- 

Internal Stages

(Correct)

- 

External Stages

(Correct)

## Explanation

Snowpipes can be used to load data from external as well as internal stages.

---

Question 32:

**When a database or a schema is cloned, which of the following statements are true for the stages in that database?**

- 

**External named stages are NOT cloned**

- 

**External named stages are cloned**

**(Correct)**

- 

**Table internal stages are cloned**

**(Correct)**

- 

**named internal stages are NOT cloned**

**(Correct)**

- 

**named internal stages are cloned**

## Explanation

See rules for stage cloning at the below link

<https://docs.snowflake.net/manuals/user-guide/object-clone.html#rules-for-cloning-stages>

---

Question 33:

REST API calls can be used to trigger a Snowpipe regardless of whether it is defined over an Internal Stage or an External Stage

- - **False**
  - 
  - **True**
- (Correct)**

#### Explanation

SnowPipe can be used to load data from external as well as internal stages. The REST API which is used to trigger a SnowPipe applies to SnowPipe created over external stages as well as internal stages

---

Question 34:

If you used a specific virtual warehouse to load data into a table, you must use the same virtual warehouse to query that data.

- - **False**
  - **(Correct)**
- - **True**

#### Explanation

Since snowflake uses shared data architecture, you can use any virtual warehouse to load data and the same or any other virtual warehouse to query that data

---

Question 35:

Which of the following objects can be cloned? Select all that apply.

- 
- **Schema**
- **(Correct)**

- 
- **Virtual Warehouse**
- 
- **User**
- 
- **Share**
- 
- **Database**
- **(Correct)**
- 
- **Table**
- **(Correct)**

### Explanation

Of these virtual warehouses and shares can't be cloned

<https://docs.snowflake.net/manuals/sql-reference/sql/create-clone.html>

---

Question 36:

**To clone a database your current role must have which privilege on the source database?**

- 
- 
- **SELECT**
- 
- **WRITE**
- 
- **USAGE**
- **(Correct)**

### **Explanation**

To clone a table you need a SELECT privileges on the source table and for all other objects that can be cloned you need the USAGE privilege

---

Question 37:

**Which of the following objects can be shared? Select all that apply.**

- 

**Views**

- 

**Secure UDFs**

**(Correct)**

- 

**Schema**

**(Correct)**

- 

**Tables**

**(Correct)**

- 

**Secure Views**

**(Correct)**

- 

**Database**

**(Correct)**

### **Explanation**

Database, Schemas, Tables, Secure views & Secure UDFs can be shared.

---

Question 38:

Which of the following semi-structured file formats are supported by Snowflake ?  
Select all that apply

- 

**AVRO**

**(Correct)**

- 

**PARQUET**

**(Correct)**

- 

**JSON**

**(Correct)**

- 

**XML**

**(Correct)**

- 

**YAML**

- 

**ORC**

**(Correct)**

- 

**HTML**

### **Explanation**

Please see below link for details on the supported formats

<https://docs.snowflake.net/manuals/user-guide/semistructured-intro.html#supported-file-formats>

Question 39:

**Which of the following statement is true when sharing data with an organization which doesn't have a Snowflake account?**

- 

**A consumer account must be created for a non-snowflake customer to allow sharing**

- 

**A reader account must be created for a non-snowflake customer to allow sharing**

**(Correct)**

- 

**No extra configuration is required to share data with a non-snowflake user/organization**

- 

**Snowflake allows sharing of data to/from Snowflake Accounts only (so the consumer must have a Snowflake account or a reader account)**

**(Correct)**

### Explanation

Sharing with a non-snowflake user requires creation of a reader account as Snowflake allows sharing only between Snowflake accounts

---

Question 40:

**If you have multiple virtual warehouses in your Snowflake system, they will each be accessing the same shared data.**

- 

**True**

**(Correct)**

- 

**False**

### Explanation

True. Although Virtual warehouses are independent of each other and don't share the CPU or memory or cache but they do operate on the shared data.

---

Question 41:

You are planning to utilize the multi-cluster virtual warehouse to provide auto-scaling and performance for your users. You want to make sure that the user queries don't wait in queue and if auto-scaling is required the cluster should auto scale immediately. What is the correct value for Scaling Policy to achieve this?

- 

**SCALING\_POLICY = STANDARD**

**(Correct)**

- 

**SCALING\_POLICY = PERFORMANCE**

- 

**SCALING\_POLICY = ECONOMY**

### Explanation

STANDARD scaling policy will auto scale as soon as a query is queued or if it is detected that the number of clusters can't handle the current queries  
<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#label-mcw-scaling-policies>

---

Question 42:

A new custom role will automatically be assigned to all existing users.

- 

**True**

- 

**False**

**(Correct)**

### Explanation

A new role is not assigned to any user. You have to explicitly assign the role.

---

Question 43:

**Which minimum Snowflake license allows data sharing capability?**

- 
- **Standard**  
**(Correct)**
- 
- **Enterprise**
- 
- **Business Critical**

**Explanation**

All editions of Snowflake provide customer dedicated virtual warehouses, therefore the lowest edition providing such functionality is Standard

---

Question 44:

**Please select the 3 key services which are part of the Snowflake Architecture**

- **Query Processing**  
**(Correct)**
- **Cloud Services**  
**(Correct)**
- **Database Storage**  
**(Correct)**
- **Virtual Machines**

Question 45:

**Cloning a database will clone which of the following**

- 

**The database itself**

**(Correct)**

- 

**All tables ONLY in the public schema in the database**

- 

**ONLY the public schema in the database**

- 

**All tables with in every schema in that database**

**(Correct)**

- 

**All the schemas in the database**

**(Correct)**

### **Explanation**

Cloning a database clones all the schemas in the database, all the tables in the schemas and all other cloneable objects in that schema

---

Question 46:

**Which one of the following represents correctly the storage hierarchy in Snowflake?**

- 

**Account->Database->Schema->Table**

**(Correct)**

- 

**Database->Account->Schema->Table**

- 

**Account->Table->Database->Schema**

- 

**Schema->Account->Database->Table**

#### **Explanation**

In Snowflake the highest level is an Account. Within an account, you can have databases. Each database can contain one or more schemas. Objects such as tables exist within schemas

---

Question 47:

**If a multi cluster virtual warehouse is resized the new size applies to ?**

- 

**All the running clusters & any clusters that are started after the multi cluster virtual warehouse is resized**

**(Correct)**

- 

**You can not resize a multi cluster virtual warehouse**

#### **Explanation**

The resize will apply to all existing clusters which are in running state and any future clusters which may be added due to autoscaling

<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#multi-cluster-size-and-credit-usage>

---

Question 48:

**What is the maximum time travel period allowed for transient & temporary tables?**

- 

**7**

- 

**90**

-

0

- 

1

(Correct)

### Explanation

Transient and temporary tables can have a maximum of 1 days of time travel regardless of which snowflake edition you are using

<https://docs.snowflake.net/manuals/user-guide/data-cdp-storage-costs.html#temporary-and-transient-tables>

---

Question 49:

When a data provider stops sharing an object, that object immediately becomes unavailable to all its consumers.

- 

True

(Correct)

- 

False

### Explanation

Data Sharing in Snowflake is a metadata operation. If a producer stops sharing any object, it will immediately become unavailable for any consumer.

---

Question 50:

When a table is cloned, a snapshot of the data is taken at the time of the cloning and represents the state of the source data at the time of cloning.

- 

False

- 

True

(Correct)

### Explanation

Correct. A snapshot is taken which represents the state of the source data. The clone is based on that snapshot. <https://docs.snowflake.net/manuals/sql-reference/sql/create-clone.html#general-usage-notes>

---

Question 51:

You are required to run a series of SQL statements on a scheduled interval. The SQL statements must execute in a defined order. What is the best way to achieve this in Snowflake?

- **1.Create multiple Tasks each containing one SQL statement.**
- **2.Connect the Tasks in a parent-child relationship to define the execution order.**
- **3.Schedule the root Task.**

**(Correct)**

- **1.Create multiple scripts each containing one SQL statement.**
- **2.Create a master script that calls all the scripts in the execution order.**
- **3.Schedule the master script.**
- **1.Create a script containing all the statements.**
- **2.Schedule the script to run on a set interval.**

### Explanation

Tasks are the Snowflake way of executing a single SQL statement on a schedule. Multiple Tasks can be connected to create a tree of tasks that starts executing from the root task and executes all child Tasks.

See more at

<https://docs.snowflake.com/en/user-guide/tasks-intro.html>

---

Question 52:

**A user can enable multi factor authentication (MFA) using the Snowflake Web interface?**

- **True**

**(Correct)**

- 

**False**

---

Question 53:

Failsafe is provided as an alternate means for customers to access historical data once the time travel retention period has ended.

- 

**False**

**(Correct)**

- 

**True**

### **Explanation**

Incorrect. Failsafe can be used only in failure scenarios for data protection and can be accessed only by the Snowflake support <https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

---

Question 54:

**You are designing a snowflake based solution for a large retail chain. What should you consider for ensuring performance once the Snowflake based data warehouse goes live? Select all that apply**

- 

**Consider using multi clustering to auto spawn & auto shutdown virtual warehouse instances to cater for increase & decrease in user demand.**

**(Correct)**

- 

**Consider using dedicated virtual warehouses for different departments.**

**(Correct)**

-

**Minimize the amount of data stored in Snowflake tables. Only store last 3 months worth of data.**

- 

**Set the size of the virtual warehouses in your solution to maximum available size (4XL) to ensure best performance.**

---

Question 55:

**What is the Snowflake recommended approach to identifying the minimum and maximum values for a multi-cluster virtual warehouse? Select all that apply**

- 

**Start with the multi-cluster virtual warehouse in auto scale mode**

**(Correct)**

- 

**Start with a minimum of 1 & a maximum cluster value of 2 or 3**

**(Correct)**

- 

**Track the virtual warehouse usage over time and adjust the minimum and maximum values that provide the best configuration for your workload**

**(Correct)**

- 

**Start with the minimum and maximum value set to the same number.**

- 

**Start with a minimum of 5 & a maximum cluster value of 10**

### **Explanation**

Snowflake recommended approach is to start small, set to auto scale and monitor and adjust according to your workload <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

---

Question 56:

Which of the following will allow a user in consumer account to create a database from the share. Select two.

- 

**SECURITY ADMIN role**

- 

**ACCOUNTADMIN role**

**(Correct)**

- 

**SYSADMIN role**

- 

**IMPORT SHARE privileges**

**(Correct)**

### Explanation

A user in consumer account can create a database from the share if they have the ACCOUNTADMIN role OR the IMPORT SHARE privileges

---

Question 57:

To create a SHARE what is the minimum required role?

- 

**ACCOUNTADMIN**

**(Correct)**

- 

**SYSADMIN**

- 

**SECURITYADMIN**

### Explanation

ACCOUNTADMIN is the role required to create & manage a share.

---

Question 58:

In the Snowflake staged release process for new releases, which account types are applied updates the **LAST**?

- 

**Designated accounts opting for early access**

- 

**Standard Accounts**

- 

**Trial Accounts**

- 

**Enterprise & Higher Accounts**

**(Correct)**

### **Explanation**

The staged release process through which Snowflake releases updates is first applied to early access account who have opted for early access, then standard accounts and finally Enterprise & upwards

---

Question 59:

**All queries will perform better if the virtual warehouse size is increased?**

- 

**False**

**(Correct)**

- 

**True**

### **Explanation**

Incorrect. Simple and short queries will generally not benefit from the additional resources provided by larger virtual warehouses

Question 60:

**How are the columns stored in a Snowflake micro partition? (Also referred to as Columnar storage)**

- 

**Combined - columns for a given row are stored together**

- 

**Independently - each column stored on its own**

**(Correct)**

#### **Explanation**

Columns are stored independently so that they can be scanned more efficiently

<https://docs.snowflake.net/manuals/user-guide/tables-clustering-micropartitions.html>

---

Question 61:

**Which of the following statement is True regarding how Snowflake stores its data.**

- 

**Snowflake uses its own proprietary columnar format to store table data.**

**(Correct)**

- 

**Snowflake uses parquet file format to store the table data.**

- 

**Snowflake stores table data as simple comma separated files in cloud based storage.**

---

Question 62:

**Snowpipe is designed to load?.**

- 

**large volumes of data on a batch schedule**

- 

**small volumes of frequent data**

**(Correct)**

### **Explanation**

Snowpipe is used for continuously loading small volumes of data frequently

---

Question 63:

**Each snowflake account is limited to a maximum of 10 shares.**

- 

**True**

- 

**False**

**(Correct)**

### **Explanation**

Incorrect. There is no hard limit on the number of shares that can be created per account.

---

Question 64:

**Which minimum Snowflake license provides the failsafe feature?**

- 

**Standard**

**(Correct)**

- 

**Enterprise**

- 

**Business Critical**

### **Explanation**

All editions of snowflake provide the failsafe feature

---

Question 65:

**Which function is used to find the query id of the last executed query?**

- - SELECT max(query\_id)**
  - 
  - SELECT min(query\_id)**
  - 
  - SELECT last\_query\_id()**
- (Correct)**

#### Explanation

[https://docs.snowflake.net/manuals/sql-reference/functions/last\\_query\\_id.html](https://docs.snowflake.net/manuals/sql-reference/functions/last_query_id.html)

---

Question 66:

**Assume a virtual warehouse of size X-Large(128 servers) running for an hour.  
What is the amount of Snowflake credits that will be billed?**

- - 64**
  - 
  - 60**
  - 
  - 128**
- (Correct)**

#### Explanation

One snowflake credit is charged per hour per server, so a cluster consisting of 128 servers will be billed as 128 Snowflake credits.

---

Question 67:

To create a new Task a user should be using a role which..

- has the ACCOUNTADMIN privileges
- has the CREATE TASK privilege on the schema where the Task is being created  
**(Correct)**
- has the SYSADMIN privileges

#### Explanation

To create a task, the role being used by the creating user must have the CREATE TASK privilege on the schema in which the Task is being created.

<https://docs.snowflake.com/en/user-guide/tasks-intro.html>

---

Question 68:

Which of the following are valid scaling type for a multi cluster virtual warehouse? Select all that apply?

- SPEED
- STANDARD  
**(Correct)**
- PERFORMANCE
- ECONOMY  
**(Correct)**

### **Explanation**

The two scaling types are STANDARD & ECONOMY

<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

---

Question 69:

**Which of the following actions can not be performed by the consumer of a shared database?**

- **use shared data in complex queries**
- **insert data in a shared table**  
**(Correct)**
- **view the list of tables in a shared database**
- **create a new table in the shared database**  
**(Correct)**

### **Explanation**

For a consumer a shared database is read-only so they can not create new objects or modify/append data

---

Question 70:

**Snowflake automatically determines the most efficient algorithm to compress columns in a micro partitions**

- 

**True**

**(Correct)**

-

**False**

### **Explanation**

Correct. The columns are compressed in a micro-partition and the algorithm is determined automatically by Snowflake based on column characteristics  
<https://docs.snowflake.net/manuals/user-guide/tables-clustering-micropartitions.html>

Question 71:

**Using the COPY command you can load data from which of the following sources? Select ALL that apply.**

- **AWS S3**  
**(Correct)**
- **IBM Cloud**
- **Oracle Cloud**
- **Internal Stages**  
**(Correct)**
- **Azure Blob Storage**  
**(Correct)**
- **Google Cloud Storage**  
**(Correct)**

### **Explanation**

Snowflake supports loading from Internal Stage, AWS S3, Azure Blob and Google Cloud Storage. Please see below link for more details

<https://docs.snowflake.net/manuals/user-guide/data-load-overview.html#loading-from-data-files-staged-on-other-cloud-platforms>

---

Question 72:

**The default table type for new tables in Snowflake is..**

- Transient**
- Temporary**
- 

**Permanent**  
**(Correct)**

#### **Explanation**

The default type for a new table is Permanent. To create other types of tables you will need to explicitly specify <https://docs.snowflake.net/manuals/user-guide/data-cdp-storage-costs.html#considerations-for-using-temporary-and-transient-tables-to-manage-storage-costs>

---

Question 1:

**As an alternative to data loading, Snowflake allows creating external tables through which you can query external data without first loading it.**

- True**  
**(Correct)**
- False**

#### **Explanation**

An external table can be created & queried directly over a data stored in cloud storage without requiring a load first. This is useful generally if you would like to query a subset of the data stored in the cloud  
<https://docs.snowflake.net/manuals/user-guide/data-load-overview.html#external-tables>

---

Question 2:

If you suspend a virtual warehouse on which a query is running, the virtual warehouse will not enter the suspended state until the active query is finished.

- 

**True**

**(Correct)**

- 

**False**

**Explanation**

Correct. The virtual warehouse will wait for the query to finish before entering the suspended state.

---

Question 3:

Snowflake JavaScript UDFs can return either Scalar or Tabular results?

- 

**True**

**(Correct)**

- 

**False**

**Explanation**

Either type of UDFs (JavaScript or SQL) can return Scalar or Tabular results  
<https://docs.snowflake.net/manuals/sql-reference/user-defined-functions.html>

---

Question 4:

Snowflake doesn't support processing semi structured data.

- 

**True**

- 

False

(Correct)

#### Explanation

Snowflake has built in function to process and query semi structured data.

---

#### Question 5:

**Which of the below statements about Snowflake is True?**

- 

You can increase or decrease the compute resources in Snowflake without making any change to the storage.

(Correct)

- 

In Snowflake whenever the storage is increased the compute must be increased.

- 

Snowflake storage & compute are independent of each other.

(Correct)

- 

Snowflake uses a monolithic architecture in which compute and storage are closely coupled.

#### Question 6:

**Historical data in Temporary tables can not be recovered by Snowflake?**

- 

False

- 

True

(Correct)

### **Explanation**

Correct. Temporary tables don't have any failsafe therefore there is no way to recover historical data

---

### **Question 7:**

**For a permanent table Snowflake provides how many days of failsafe retention?**

- 
- 0**
- 
- 14**
- 
- 90**
- 
- 7**

**(Correct)**

### **Explanation**

Failsafe retention for permanent tables is 7 days

<https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

---

### **Question 8:**

**Snowflake customers can control what format Snowflake stores the data for a table.**

- 
- Yes**
- 
- No**

**(Correct)**

---

### **Question 9:**

**Snowflake allows two different styles for loading data. Please select the two correct options.**

- **Continuous Data Loading**  
**(Correct)**
- **Bulk Data Loading**  
**(Correct)**
- **Intermittent Data Loading**

### Explanation

There are two approaches to loading data in Snowflake. Bulk data load which is performed through the COPY command and continuous data load performed through Snowpipe

---

Question 10:

**On the Web UI for each worksheet you can select which of the following? Select all that Apply**

- **Schema**  
**(Correct)**
- **Database**  
**(Correct)**
- **Role**  
**(Correct)**
- **Table**
-

## User

- 

**Virtual Warehouse**

**(Correct)**

## Explanation

You must select Database, Schema, Role & the virtual warehouse for each Worksheet before you can run a query

---

### Question 11:

A cloned database doesn't inherit the source permissions, but schemas, tables & views contained inside the cloned database will inherit the source permissions.

- 

**False**

- 

**True**

**(Correct)**

## Explanation

<https://docs.snowflake.net/manuals/user-guide/object-clone.html#access-control-privileges-for-cloned-objects>

---

### Question 12:

It is possible to share data with a Snowflake customer whose Snowflake instance exists in a **different Region** than the provider?

- 

**Yes. Nothing special needs to be done to enable cross region data sharing.**

- 

**Yes, but to enable cross region data sharing you must enable replication first.**

**(Correct)**

- 

**No, it is not possible to share with customers in other regions.**

### **Explanation**

It is possible to share data with other regions, but the provider must enable replication and replicate your existing database to the other region.

---

**Question 13:**

You are a data engineer with a telecom company who use Snowflake as a data warehouse. The company requires all network signaling data to be loaded into a table in a near real-time manner.

The network signaling data is already landing into an S3 bucket every 1 minute. What is the course of action that you should take?

Select all that apply

- **Create a Snowpipe that loads data from a defined stage.**  
**(Correct)**
- **Configure notification event on the S3 bucket which trigger the Snowpipe.**  
**(Correct)**
- **Configure the Snowpipe to continuously check for new files in the S3 bucket.**
- **Create a stage on top of the S3 bucket where the real-time data is landing.**  
**(Correct)**
- **Run the Snowpipe.**

### **Explanation**

Explanation for incorrect answers A Snowpipe can not be run through a command, it must be triggered by REST API or through cloud storage notifications A Snowpipe itself can not watch for new files to arrive

---

**Question 14:**

It is possible to share data with a Snowflake customer whose Snowflake instance exists in a **different Cloud** than the provider?

- 

**Yes. Nothing special needs to be done to enable cross cloud data sharing.**

- 

**No, it is not possible to share with customers in other clouds.**

- 

**Yes, but to enable data sharing to a different cloud you must enable replication first.**

**(Correct)**

**Explanation**

It is possible to share data with other clouds but the provider must enable replication and replicate your existing database to the other cloud

---

**Question 15:**

**Cloning a schema will clone which of the following**

- 

**All other cloneable objects in the schema**

**(Correct)**

- 

**The schema itself**

**(Correct)**

- 

**All tables in the schema**

**(Correct)**

**Explanation**

Cloning a schema clones all the tables in the schemas and all other cloneable objects in that schema

---

Question 16:

If the cloned object is a database or a schema, the child tables & views in the database or schema will inherit all the original privileges.

- 

False

- 

True

(Correct)

#### Explanation

Correct. If the cloned object is a database or a schema, the cloned child tables & views will inherit the permissions that were applied on the source objects  
<https://docs.snowflake.net/manuals/user-guide/object-clone.html#access-control-privileges-for-cloned-objects>

Question 17:

A reader account can consume data from sources other than the producer that created the reader account.

- 

False

(Correct)

- 

True

#### Explanation

Incorrect. A reader account can only consume data from the producer account that created it

Question 18:

On the WebUI which button on the top bar should you select to create new compute clusters?

- 

Virtual Warehouses

(Correct)

- 
- Shares**
- 
- History**
- 
- Account**
- 
- Databases**
- 
- Worksheets**

#### **Explanation**

Compute Cluster or virtual warehouses can be created by selecting the "Virtual Warehouse" button from the top bar

---

Question 19:

An ACCOUNTADMIN or SECURITYADMIN can disable the MFA for a user and can allow the user to re-enroll MFA

- 
- True**
- (Correct)**
- 
- False**

#### **Explanation**

Correct. A user with security admin role can disable MFA for user(s) and also allow them to re-enroll

---

Question 20:

Which of the following is not a layer in Snowflake?

- 
- Database Storage**

- Virtual Warehouse
  - Cloud Services
  - Virtual Machines
- (Correct)
- 

Question 21:

It is possible to disable failsafe entirely for a Snowflake account.

- True
  - False
- (Correct)

### Explanation

Incorrect. Failsafe can not be disabled for the Account or a database or a table. It is a Snowflake feature that is always present. The only tables that don't have failsafe (by default, and this can't be configured either) are transient and temporary tables.

---

Question 22:

In a multi cluster virtual warehouse, if you set the same value for minimum cluster count and the maximum cluster count

- Snowflake starts all clusters in the multi-cluster virtual warehouse
- Snowflake starts minimum required clusters in the multi-cluster virtual warehouse

- 

The mode is called auto-scale

- 

The mode is called Maximized

(Correct)

### Explanation

If minimum and maximum value is set to same for a multi cluster virtual warehouse, Snowflake will start all clusters. The mode is called Maximized and is used to a) provide maximum performance b) statically control the resources  
<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

---

Question 23:

The Snowflake processing engine uses which of the following?

- 

Hive

- 

Map Reduce

- 

Spark

- 

SQL

(Correct)

### Explanation

The Snowflake processing is based on SQL

---

Question 24:

Which one of the following loading method will use Virtual Warehouse resources?

- 

Snowpipe

- 

**COPY command**

**(Correct)**

#### **Explanation**

COPY command uses virtual warehouse resources, Snowpipe is billed separately and doesn't use virtual warehouse resources

---

Question 25:

**An ACCOUNTADMIN can see the results of any query executed by any user in a Snowflake system?**

- 

**False**

**(Correct)**

- 

**True**

#### **Explanation**

"You can view results only for queries you have executed. If you have privileges to view queries executed by another user, the Query Detail page displays the details for the query, but, for data privacy reasons, the page does not display the actual query result." See note at <https://docs.snowflake.net/manuals/user-guide/ui-history.html>

---

Question 26:

**You are required to install Snowflake database engine on a server before you can start using it**

- 

**False**

**(Correct)**

- 

**True**

#### **Explanation**

Incorrect. Snowflake is a software as a service offering and doesn't require installation

---

Question 27:

The variant data type can hold up to 16MB of data per row?

- 

True

(Correct)

- 

False

#### Explanation

The VARIATNT can store up to 16 MB. Please see below link for more details

<https://docs.snowflake.net/manuals/sql-reference/data-types-semistructured.html#variant>

---

Question 28:

If a virtual warehouse is suspended it is likely that the results stored in its cache will be cleared.

- 

True

(Correct)

- 

False

#### Explanation

Correct. The cache is likely to be cleared on suspend

---

Question 29:

When setting up replication for cross cloud or cross region data sharing, the data provider must replicate data once for each data consumer.

- 

False

(Correct)

-

**True**

**Explanation**

Incorrect. Only one instance of data per cloud or region has to be replicated. Once the instance is replicated more than one consumer can make use of this data

---

**Question 30:**

**When a virtual warehouse is in a running state it consumes Snowflake credits regardless whether a query is being run on it or not.**

- 

**True**

**(Correct)**

- 

**False**

**Explanation**

As soon as the virtual warehouse is in a running state it starts consuming snowflake credits

---

**Question 31:**

**Which of the following mechanism allows a Snowflake customer to query data without loading it first?**

- 

**Snowpipe**

- 

**Virtual Table**

- 

**COPY**

- 

**External Table**

**(Correct)**

**Explanation**

External tables allow querying data in external storage without first loading it into a Snowflake table

Question 32:

Which of the below statements about Snowpipes are correct? Select all that apply.

- 

**Snowpipes are billed independent of the virtual warehouse usage.**

**(Correct)**

- 

**Snowpipes use serverless architecture and as such don't use virtual warehouse resources.**

**(Correct)**

- 

**Costs associated with Snowpipe usage are billed as virtual warehouse compute.**

- 

**For processing, Snowpipe share virtual warehouse compute resources.**

#### Explanation

Snowpipe uses a server less architecture and therefore doesn't use a virtual warehouse compute resources. Please see below link for further details  
<https://docs.snowflake.net/manuals/user-guide/data-load-snowpipe-intro.html>

Question 33:

**You can set your virtual warehouses to auto suspend and auto resume so that when the virtual warehouse is not being used for a set time period, it goes into suspended mode and resumes when a query is executed.**

- 

**False**

- 

**True**

**(Correct)**

#### Explanation

A virtual warehouse can be suspended and resumed as required.

Question 34:

Compared to a single cluster virtual warehouse, a multi-cluster virtual warehouse supports which of the following? Select all that apply

- 

Auto Suspend

(Correct)

- 

Stop

(Correct)

- 

Auto Resume

(Correct)

- 

Resize

(Correct)

Explanation

All are correct. A multi cluster virtual warehouse allows all operations that a single cluster virtual warehouse allows

Question 35:

The COPY command allows using simple transformations while loading data

- 

True

(Correct)

- 

False

Explanation

Correct. You can perform basic transformations such as column re-ordering, column omission, cast data types & truncating text data while load it through the COPY command

**Question 36:**

**Which of the following are options when creating a new multicloud virtual warehouse. Select all that apply**

- 

**MIN\_CLUSTER\_COUNT**

**(Correct)**

- 

**MAX\_CLUSTER\_COUNT**

**(Correct)**

- 

**SCALING\_POLICY**

**(Correct)**

- 

**SUSPEND\_AND\_RESUME**

**Explanation**

Please see the below link for detailed syntax on virtual warehouse creation

<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

**Question 37:**

**Assuming an administrator has instructed a virtual warehouse to stop, the virtual warehouse will stop immediately even if it is executing an SQL statement.**

- 

**True**

- 

**False**

**(Correct)**

**Explanation**

Incorrect, an executing SQL statement either must complete or be aborted before a virtual warehouse can be stopped

Question 38:

It is not possible to specify a specific cluster in a multi-cluster virtual warehouse to execute your query

- 

**True**

**(Correct)**

- 

**False**

#### Explanation

True. Individual clusters in a virtual warehouse can not be specified independently, but rather they appear logically as a virtual warehouse together. The assignment of the query to a specific cluster will be automatically managed by Snowflake and can not be controlled

Question 39:

Once a table type i.e. permanent, transient or temporary has been defined for a table it can not be changed.

- 

**False**

- 

**True**

**(Correct)**

#### Explanation

Correct. In order to change the type you will need to re-create the table with the desired table type

Question 40:

Which of the following objects can be cloned? Select all that apply.

- 

**User**

- 

**Stage**

**(Correct)**

- 

**File Format**

**(Correct)**

- 

**Sequence**

**(Correct)**

- 

**Virtual Warehouse**

- 

**Share**

#### **Explanation**

Of these virtual warehouses and shares can't be cloned

<https://docs.snowflake.net/manuals/sql-reference/sql/create-clone.html>

Question 41:

**Under what circumstances results for a query will be fulfilled by the query result cache? Select all that apply**

- 

**The query being run returns result set of less than 100 thousand rows**

- 

**The query is being executed from the same virtual warehouse as the previously executed query**

- 

**The query results cache was generated or last used less than 24 hours ago**

**(Correct)**

- 

**The role executing the query has privileges to use all the tables in the query**

**(Correct)**

- 

**The micro-partitions for the tables in the query have not changed**

**(Correct)**

### **Explanation**

The underlying data must not change and the query should be syntactically identical in order for the query result cache to be used. The cache must have been generated (or last used) less than 24 hours ago. <https://docs.snowflake.net/manuals/user-guide/querying-persisted-results.html#retrieval-optimization>

---

**Question 42:**

**Which of the following are true for Snowpipe?**

- 

**The resource for Snowpipe are automatically scaled up and down by Snowflake**

**(Correct)**

- 

**You must scale a virtual warehouse yourself in order to manage the compute available to Snowpipe**

- 

**Snowpipe uses serverless compute resources managed by Snowflake**

**(Correct)**

- 

**Snowpipe makes use of the active virtual warehouse for compute resources**

### **Explanation**

Snowpipe uses a server less compute which is managed by Snowflake and is scaled up and down automatically by Snowflake

---

**Question 43:**

**For data loading and transformation, the approach recommended by Snowflake is?**

-

**ELT (Extract, Load, Transform)**

**(Correct)**

- 

**ETL (Extract, Transform, Load)**

### **Explanation**

The ELT approach utilizes the processing power of Snowflake to Transform the data after it has been loaded.

---

**Question 44:**

**In the Snowflake WebUI you can open only ONE worksheet at a time?**

- 

**True**

- 

**False**

**(Correct)**

### **Explanation**

More than one worksheets can be opened on the Snowflake Web UI

---

**Question 45:**

**What are the types of caching in Snowflake?**

- 

**Metadata Caching**

**(Correct)**

- 

**Index Caching**

- 

**Query Result Caching**

**(Correct)**

-

## Warehouse Caching

(Correct)

- 

## Memory Caching

### Explanation

Metadata caching is used for queries which can be fulfilled directly from metadata, e.g. the row count of a table. Query Result Caching is for queries that have been executed already and the results can be accessed by a queryID. Warehouse caching is within the virtual warehouse instance and is usually based on queries that have already executed.

---

Question 46:

**When a virtual warehouse is resized to a smaller size, the cache associated with the servers is dropped?**

- 

**False**

- 

**True**

(Correct)

### Explanation

Correct. The cache will be invalidated

---

Question 47:

**You need to create a new Reader Account for sharing a table with a non Snowflake customer. You are logged in with a SYSADMIN role. When you run the query to create the new Reader Account you are met with a privilege error. What do you need to do be able to create the Reader Account.**

- 

**Login with or switch role to SECURTIYADMIN. Only SECURITYADMIN role is allowed to create new users and new reader accounts**

- 

**Login with or switch role to ACCOUNTADMIN. Only ACCOUNTADMIN is allowed to create new reader accounts.**

**(Correct)**

- 

**Use the WebUI to create the Reader Account.**

#### **Explanation**

Only ACCOUNTADMIN role is allowed to create new Reader Account, since it is deemed to be an account level activity.

---

**Question 48:**

**Clustering key help query performance by?**

- |

**Improving column compression**

**(Correct)**

- 

**Pruning un-necessary partitions**

**(Correct)**

- 

**distributing the data over multiple clusters**

- 

**Pre-calculating query results**

#### **Explanation**

<https://docs.snowflake.net/manuals/user-guide/tables-clustering-keys.html>

---

**Question 49:**

**Which of the following are valid virtual warehouse sizes?**

- 

**Medium**

**(Correct)**

- 

**X-Large**

**(Correct)**

- 

**Standard**

- 

**Small**

**(Correct)**

- 

**2X-Large**

**(Correct)**

- 

**Premium**

---

**Question 50:**

**The COPY command allows the following options for selecting files for loading data from a stage. Select all that apply**

- 

**Load specific files by providing exact file names**

**(Correct)**

- 

**Load all files in a specific path**

**(Correct)**

- 

**Load files that match a pattern**

**(Correct)**

### **Explanation**

All of these options are correct. COPY command allows you to load specific files, load all files or load files that match a pattern. Please below for further details.

<https://docs.snowflake.net/manuals/user-guide/data-load-considerations-load.html#options-for-selecting-staged-data-files>

---

Question 51:

Through Snowflake sharing a producer can share data with ?

- 

A non snowflake customer

(Correct)

- 

A Snowflake customer

(Correct)

- 

google driver users

- 

other databases

- 

multiple snowflake customers

(Correct)

### Explanation

You can share data with multiple consumers and these consumers can be snowflake customers, non snowflake customers or a mix of both

---

Question 52:

Snowflake offers the following four editions.

- 

Standard, Enterprise, Business Critical, Virtual Private Snowflake

(Correct)

- 

Personal, Standard, Enterprise, On Premise,

- 
- 

### Basic, Enterprise, Snowball, On Premise

#### Explanation

Please see details at <https://www.snowflake.com/pricing/pricing-guide/>

---

Question 53:

The Account button on the top bar in the Snowflake Web UI is available to which role(s)? Select all that apply

- 
- 

**SECURITYADMIN**

(Correct)

- 

**ACCOUNTADMIN**

(Correct)

- 

**PUBLIC**

- 

**SNOWFLAKE**

- 

**SYSADMIN**

#### Explanation

Both SECURITYADMIN & the ACCOUNTADMIN roles can see & access this button on the Snowflake Web UI. The account level functions in the sub tabs will be ONLY available to ACCOUNTADMIN though.

---

Question 54:

An administrator must grant explicit privileges on a cloned object as the cloning doesn't copy the original privileges of an object however there is no need to grant privileges to child objects because they inherit the original permissions.

- 
- 

**False**

- 

True

(Correct)

#### Explanation

<https://docs.snowflake.net/manuals/user-guide/object-clone.html#access-control-privileges-for-cloned-objects>

Question 55:

You added an additional cluster to a multi cluster virtual warehouse. What is this an example of?

- 

Scale Out

(Correct)

- 

Scale Up

#### Explanation

Adding additional clusters will allow more users and more queries to execute which is an example of scaling out

Question 56:

Which of the following roles can create & manage users? Select all that apply

- 

SYSADMIN

- 

ACCOUNTADMIN

(Correct)

- 

SECURITYADMIN

(Correct)

-

## SNOWFLAKE

### Explanation

Both SecurityAdmin and AccountAdmin can manage users, but the recommended practice is to use the SecurityAdmin role to manage users

---

Question 57:

**It is possible to disable failsafe for certain databases, schemas or tables.**

- True**
  - False**
- (Correct)**

### Explanation

Incorrect. Failsafe can not be disabled for the Account or a database or a table. It is a Snowflake feature that is always present. The only tables that don't have failsafe (by default, and this can't be configured either) are transient and temporary tables.

---

Question 58:

**If you create more than one virtual warehouse they will share the memory and CPU resources.**

- True**
  - False**
- (Correct)**

### Explanation

Incorrect. Virtual warehouses are independent of each other and don't share the CPU or memory or cache. They do operate on the shared data though.

---

Question 59:

**Any snowflake account regardless of the edition, can produce and consume shared data.**

-

**True, except for VPS (virtual private snowflake)**

**(Correct)**

- 

**False, the sharing capability is disabled by default and has to be turned on by Snowflake support**

- 

**False, the sharing capability is not available to the Standard and Premier accounts.**

- 

**True. Any edition can produce and consume data**

**Explanation**

All snowflake editions can produce and consume shared data, except the Virtual Private Snowflake, which by nature is designed to be secure and can not share data

---

**Question 60:**

**When a virtual warehouse is resized all currently executing queries will start performing better?**

- 

**False**

**(Correct)**

- 

**True**

**Explanation**

Incorrect. The resize operation does not impact existing running queries

---

**Question 1:**

**Unloading of data from Snowflake requires a running virtual warehouse.**

- 

**False**

-

**True**

**(Correct)**

### **Explanation**

A running virtual warehouse is required for unloading data from Snowflake. A running virtual warehouse will be needed even in cases when you are unloading using a query, where the query results are cached. This is because a running virtual warehouse is needed to be able to write the rows to the output file(s).

<https://docs.snowflake.net/manuals/user-guide/data-unload-snowflake.html>

---

**Question 2:**

**Which of the following are valid Snowflake stage types?**

- 

**External Stage**

**(Correct)**

- 

**Table Stage**

**(Correct)**

- 

**Internal Stage**

**(Correct)**

- 

**Database Stage**

### **Explanation**

The three types of stages supported by Snowflake are Internal Stages - in which data is stored internally by Snowflake, and External Stages - where data is stored in Cloud Storage & Table stage - which is automatically created for each table.

---

**Question 3:**

**When Unload data from a Snowflake table the results can be written into ...**

-

**Multiple files**

**(Correct)**

- 

**Another table**

- 

**A single file**

**(Correct)**

### **Explanation**

Snowflake allows unloading data into a single or multiple files. This can be controlled by changing the COPY command option SINGLE from TRUE to FALSE.  
<https://docs.snowflake.net/manuals/user-guide/data-unload-overview.html#bulk-unloading-into-single-or-multiple-files>

**Question 4:**

**When converting a character value to Boolean, which of the following values will result in "TRUE".**

- 

**'yes'**

**(Correct)**

- 

**'1'**

**(Correct)**

- 

**'t'**

**(Correct)**

- 

**'y'**

**(Correct)**

-

'true'

(Correct)

- 

'on'

(Correct)

### Explanation

All of these will convert to a value of TRUE <https://docs.snowflake.net/manuals/sql-reference/data-types-logical.html#boolean-conversion>

### Question 5:

**Which of the following are supported file formats when UNLOADING data from Snowflake?**

- 

**Delimited Text**

(Correct)

- 

**Parquet**

(Correct)

- 

**AVRO**

- 

**Tar.gz**

- 

**Zip**

- 

**JSON**

(Correct)

### Explanation

For data unloading or extraction the supported formats are delimited text (i.e. CSV etc..), JSON and Parquet. <https://docs.snowflake.net/manuals/user-guide/data-unload-prepare.html#supported-file-formats>

---

Question 6:

To unload a relational table into a JSON file which function should you use to convert the data ?

- 

The OBJECT\_CONSTRUCT will help you convert the data into a JSON structure

(Correct)

- 

Perform string concatenation and complex SQL to produce JSON string of your choice

- 

The unloaded data can not be in the JSON format

#### Explanation

The Object\_Construct function will help you in converting the data into the JSON format <https://docs.snowflake.net/manuals/user-guide/data-unload-considerations.html#unloading-a-relational-table-to-json>  
[https://docs.snowflake.net/manuals/sql-reference/functions/object\\_construct.html](https://docs.snowflake.net/manuals/sql-reference/functions/object_construct.html)

---

Question 7:

To subtract or add a period of time from a date, time or time stamp you should ....

- 

use INTERVAL constant

(Correct)

- 

convert date/times to integer data type to perform addition / subtraction and then convert back to date / time

#### Explanation

Snowflake provides the INTERVAL constant to enable you to subtract periods of time from a date or time data type e.g. `SELECT TO_DATE('2020-02-15') + INTERVAL '1`

YEAR'; <https://docs.snowflake.net/manuals/sql-reference/data-types-datetime.html#interval-constants>

---

Question 8:

**Can an external table can be partitioned?**

- 

**False**

- 

**True**

**(Correct)**

**Explanation**

External tables can be partitioned. This is dependent on if your data is residing in a logical folder structure which allows effective partitioning.  
<https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html#partitioned-external-tables>

---

Question 9:

**When unloading (or extracting) data using a Query, you are allowed only a limited set of SQL options?**

- 

**False**

**(Correct)**

- 

**True**

**Explanation**

When using a query to unload data from Snowflake, you are allowed the full SQL capabilities, including joining with other tables, grouping etc.  
<https://docs.snowflake.net/manuals/user-guide/data-unload-overview.html#bulk-unloading-using-queries>

---

Question 10:

**When converting a numeric value to Boolean, which of the following values will result in "TRUE".**

-

2

(Correct)

- 

-1

(Correct)

- 

1

(Correct)

- 

0

### Explanation

Any non Zero value will be converted to TRUE. Zero will always convert to FALSE.  
<https://docs.snowflake.net/manuals/sql-reference/data-types-logical.html#boolean-conversion>

---

Question 11:

Snowflake Internal Stages can be defined as Permanent or Temporary?

- 
- **False**
- 
- **True**

(Correct)

### Explanation

Snowflake internal stages can be temporary or permanent. A temporary internal stage is automatically removed at the end of the session

---

Question 12:

When unloading (or extracting) data using a Query, you can extract data from ...

- **A View**  
**(Correct)**
- **A database**
- **A Single table**  
**(Correct)**
- **Multiple tables joined together**  
**(Correct)**

### Explanation

When using a query to unload data from Snowflake, you are allowed the full SQL capabilities, including joining with other tables, grouping etc. therefore you can extract from a single table, or multiple tables joined together or a view, or even multiple views joined together <https://docs.snowflake.net/manuals/user-guide/data-unload-overview.html#bulk-unloading-using-queries>

---

Question 13:

**In Snowflake all data at rest is always encrypted.**

- **True**  
**(Correct)**
- **False**

### Explanation

All data at rest i.e. data in tables, data in internal stages etc.. Is encrypted automatically <https://docs.snowflake.net/manuals/user-guide/security-encryption.html>

---

Question 14:

Non-encrypted data is loaded from a client machine to an internal stage. At what point in time the data will be encrypted?

- 

The data is encrypted on the client machine before being transferred to the internal stage

(Correct)

- 

The data is encrypted once they are in the internal stage

#### Explanation

When copying data to an internal stage the files are encrypted on the client machine before they are transferred to the internal stage.

<https://docs.snowflake.net/manuals/user-guide/security-encryption.html#end-to-end-encryption>

Question 15:

When unloading data from Snowflake, what are the places where an output format of the file can be specified?

- 

In the definition of the named stage

(Correct)

- 

In the definition of a table itself

(Correct)

- 

Specified as option directly in the COPY command

(Correct)

#### Explanation

All 3 are correct options. You can specify either in the table definition it self <https://docs.snowflake.net/manuals/sql-reference/sql/create-table.html#format-type-options-formattypeoptions> You can define a named file format and use that in the copy command, or you can directly specify your file format in the copy command

<https://docs.snowflake.net/manuals/user-guide/data-unload-prepare.html#specifying-file-format-options>

Question 16:

Data in a Snowflake table is automatically encrypted by Snowflake using ...?

- 

MD5 Hash

- 

**AES-256 strong encryption**

**(Correct)**

- 

SSL

- 

AES-128 bit encryption

#### Explanation

Data in Snowflake tables is encrypted automatically using AES 256 strong encryption

Question 17: Which of the following statements are true when data is UNLOADED into a Snowflake stage (internal or external).

- 

If the data is unloaded into an INTERNAL stage the data IS automatically encrypted

**(Correct)**

- 

If the data is unloaded into an EXTERNAL stage the data is NOT automatically encrypted however the customer can choose to encrypt as an option

**(Correct)**

-

**When UNLOADING the data is never automatically encrypted, regardless if it is an internal or external stage**

- 

**The customer downloads data (files) from the INTERNAL stage and decrypts the data on their own machines.**

**(Correct)**

### **Explanation**

When UNLOADING, snowflake will auto encrypt if the stage is INTERNAL, if the stage is EXTERNAL then the customer can optionally choose to encrypt. The files downloaded from an internal stage must be decrypted on the client side.  
<https://docs.snowflake.net/manuals/user-guide/security-encryption.html>

---

**Question 18: Which Snowflake feature allows you to query data in an external stage without loading it first into Snowflake?**

- 

**Snowpipe**

- 

**Views**

- 

**COPY**

- 

**External Tables**

**(Correct)**

### **Explanation**

External table is the easiest way to query data in an external stage without loading it first into Snowflake. <https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html>

---

Question 19:

**Snowflake supports which of the following for security?**

- 

**MFA (Multi factor authentication)**

**(Correct)**

- 

**Role based access control (RBAC)**

**(Correct)**

- 

**Discretionary Access control (DAC)**

**(Correct)**

### **Explanation**

Snowflake supports all 3 of these. MFA can be enabled for a user to provide multi factor authentication support. DAC is the access control provided by the owner of the object and RBAC is role based control over various objects in the Snowflake instance.

---

Question 20:

**Which of the following statements are true about Encryption in Snowflake?**

- 

**All files stored in external stages automatically encrypted**

- 

**All files stored in internal stages automatically encrypted**

**(Correct)**

- 

**All data inside Snowflake is automatically encrypted**

**(Correct)**

-

**Supports encryption using customer provided keys**

**(Correct)**

### **Explanation**

All data at rest i.e. data in tables, data in internal stages etc.. Is encrypted automatically. Snowflake internal stages are also automatically encrypted. Since data in external stages is usually managed by a customer, automatic encryption is not applied on files on external stages, how ever they can be encrypted optionally <https://docs.snowflake.net/manuals/user-guide/security-encryption.html>

---

**Question 21:**

**A string constant in Snowflake can be enclosed by**

- double quotes i.e. "Sydney"
- \$\$ symbol i.e. \$\$Sydney\$\$

**(Correct)**

- single quotes i.e. 'Sydney'

**(Correct)**

### **Explanation**

String constants can be enclosed in single quotes or \$ symbols..

<https://docs.snowflake.net/manuals/sql-reference/data-types-text.html#single-quoted-string-constants> <https://docs.snowflake.net/manuals/sql-reference/data-types-text.html#dollar-quoted-string-constants>

---

**Question 22:**

**When a temporary "External" Stage is dropped then ..**

- All files in the stage are removed
- Only files that have not been loaded are retained, rest are dropped
-

**None of the files in the stage are removed**

**(Correct)**

### **Explanation**

For temporary external stages, only the stage is dropped, the files being external are not removed.

---

Question 23:

**Non-encrypted data is loaded from a cloud storage to a Snowflake table. At what point in time the data will be encrypted?**

- 

**The data is encrypted as it loaded into the Snowflake table**

**(Correct)**

- 

**The data is encrypted on the cloud storage before being transferred to the Snowflake table**

### **Explanation**

When copying data to a Snowflake table from cloud storage, the data is encrypted as soon as it is loaded into the Snowflake table.

<https://docs.snowflake.net/manuals/user-guide/security-encryption.html#end-to-end-encryption>

---

Question 24:

**Which of the following data types are treated as VARCHAR "under the covers" by Snowflake?**

- 

**STRING**

**(Correct)**

- 

**TEXT**

**(Correct)**

- 

**CHAR**

**(Correct)**

- 

**CHARACTER**

**(Correct)**

#### **Explanation**

All of these are treated as VARCHAR data type by Snowflake, with the following caveats. CHAR & CHARACTER default to size "1" if no size is specified  
<https://docs.snowflake.net/manuals/sql-reference/data-types-text.html#char>  
<https://docs.snowflake.net/manuals/sql-reference/data-types-text.html#string-text>

---

**Question 25:**

**Snowflake supports querying data in an internal or external stage using Standard SQL?**

- 

**Yes. You can query data in both internal and external stages using Standard SQL.**

**(Correct)**

- 

**No. Standard SQL can only be used to query data in an internal stage**

- 

**No. Standard SQL can only be used to query data in an external stage**

#### **Explanation**

Snowflake allows querying data in both internal and external stages using standard SQL syntax <https://docs.snowflake.net/manuals/user-guide/querying-stage.html>

---

**Question 26:**

**Select the correct statement about Snowflake external tables.**

- 

**External table support all the file formats that the COPY command supports.**

**(Correct)**

- 

**External table can only be created on top of CSV data**

### Explanation

External tables support all file formats that are supported by the COPY command  
<https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html>

---

Question 27:

Select the correct statement about Snowflake external tables.

- 

Querying data in an external table is faster than querying a native Snowflake table

- 

Querying data in an external table is slower than querying a native Snowflake table

(Correct)

### Explanation

External table being outside of Snowflake, don't have the same optimizations as a native Snowflake table, therefore querying an external table is slower.

<https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html>

---

Question 28:

Snowflake doesn't have a Logical data type to store True & False information. You need to use an Integer type and store 1/0 to represent the True / False.

- 

Correct. There is no way to store Boolean information in Snowflake

- 

Incorrect. There is a BOOLEAN data type

(Correct)

### Explanation

There is a Boolean data type which can store TRUE / FALSE information

---

Question 29:

Snowpipe prevents loading a single file multiple times by...

- 

Snowpipe can not prevent loading a single file multiple times, you must delete the files after it has been loaded by Snowpipe

- deleting the files in the (internal or external) stage after it has successfully loaded the file.
  - maintaining metadata for each file loaded and ignoring files that were previously loaded
- (Correct)

### Explanation

Snowflake prevents data duplication by maintaining metadata and ignoring files that have already been loaded. <https://docs.snowflake.net/manuals/user-guide/data-load-snowpipe-intro.html#data-duplication>

---

Question 30:

While Snowpipe usually loads data files in the order that they arrive, there is no guarantee that files will be loaded in the order they arrived in the stage.

- - False
  -
- True
- (Correct)

### Explanation

Correct. Since multiple processes are reading from the queue, there is no guarantee that the files will be loaded in the exact order that they arrived.

<https://docs.snowflake.net/manuals/user-guide/data-load-snowpipe-intro.html#load-order-of-data-files>

---

Question 31:

Snowflake supports which of the following types of stages?

- 
- Unload Stage
- 
- Virtual Stage
-

### External Stage

(Correct)

- 

### Internal Stage

(Correct)

- 

### Load Stage

#### Explanation

The two types of stages supported by Snowflake are Internal Stages - in which data is stored internally by Snowflake, and External Stages - where data is stored in Cloud Storage Additionally a third type of stage is table stage (not mentioned in this question) which is automatically created for each table.

<https://docs.snowflake.net/manuals/sql-reference/ddl-stage.html#stage-management>

---

Question 32:

**The VARCHAR data type can store maximum 16MB of uncompressed character data.**

- 

True

(Correct)

- 

False

#### Explanation

The maximum size of data that a VARCHAR column can store is 16 MB uncompressed. The maximum number of characters that can be stored in a VARCHAR depends on if the character is a single byte character or 2 bytes or multibyte. In case of single byte the maximum length of VARCHAR is 1677216

<https://docs.snowflake.net/manuals/sql-reference/data-types-text.html#varchar>

---

Question 33:

**When uncompressed files are staged in a Snowflake internal stage, the data is automatically compressed using gzip.**

-

**False**

- 

**True**

**(Correct)**

### **Explanation**

Snowflake will automatically compress data stored in an external stage (unless compression is disabled) <https://docs.snowflake.net/manuals/user-guide/intro-summary-loading.html#compression-of-staged-files>

---

**Question 34:**

**In Snowflake the Boolean data type can store Unknown in addition to True / False**

- 

**Correct**

**(Correct)**

- 

**Incorrect**

### **Explanation**

Correct. The Boolean data type allows storing of True / False & Unknown values

---

**Question 35:**

**Which of the following are supported data types in Snowflake, for storing date or time information?**

- 

**DATE**

**(Correct)**

- 

**TIME**

**(Correct)**

- 

**TIMESTAMP**

**(Correct)**

### Explanation

All are examples of date / time data type. You can find more date/time data types at <https://docs.snowflake.net/manuals/sql-reference/data-types-datetime.html>

---

Question 36:

**When a temporary "Internal" Stage is dropped then ..**

- 

**Only files that have not been loaded are retained, rest are dropped**

- 

**All files in the stage are removed**

**(Correct)**

- 

**None of the files in the stage are removed**

### Explanation

When a temporary internal stage is dropped all files in that stage are dropped regardless of their load status. This is done primarily to conserve storage costs as files in temporary internal stages also count towards your Snowflake storage

---

Question 37:

**For efficient & parallelized loading of data, Snowflake recommends which of the following**

- 

**Split large files into smaller files**

**(Correct)**

- 

**One large file will load better than multiple smaller files**

- 

**Files should be 10 - 100 MB in size (compressed)**

**(Correct)**

-

## Files should be 1GB in size (compressed)

### Explanation

Split large files into smaller files of 10-100 MB after compression. The number of files will generally be proportional to the number of servers in the virtual warehouse being used to load the data . <https://docs.snowflake.net/manuals/user-guide/data-load-considerations-prepare.html#file-sizing-best-practices-and-limitations>

---

Question 38:

**When converting a character value to Boolean, which of the following values will result in "FALSE".**

- 

'no'

(Correct)

- 

'0'

(Correct)

- 

'f'

(Correct)

- 

'off'

(Correct)

- 

'false'

(Correct)

- 

'n'

(Correct)

### Explanation

All of these will convert to a value of FALSE

<https://docs.snowflake.net/manuals/sql-reference/data-types-logical.html#boolean-conversion>

---

Question 39:

The **BINARY** data type can store what amount of data?

- Unlimited
- 16MB
- 8MB
- (Correct)
- 32MB

#### Explanation

The BINARY data type can store 8MB of data

<https://docs.snowflake.net/manuals/sql-reference/data-types-text.html#binary>

---

Question 40:

The load history for Snowpipe is stored in the metadata for how many days?

- 
- 14 days
- (Correct)
- 
- 64 days
- 
- 28 days

#### Explanation

The SnowPipe load history is retained for 14 days and must be queried through REST API or the ACCOUNT\_USAGE view. <https://docs.snowflake.net/manuals/user-guide/snowpipe.html#load-history>

guide/data-load-snowpipe-intro.html#how-is-snowpipe-different-from-the-copy-command

---

Question 41:

**Snowflake supports query data in an internal or external snowflake stage. Can you use functions in your queries when querying data in a snowflake stage?**

- 

**Yes, you can use functions that are supported**

**(Correct)**

- 

**Yes, you can use any function that you can use in SQL.**

- 

**No. You can not use functions when querying data in internal or external stages**

#### Explanation

You can use only the "supported" functions when querying data in a Snowflake stage. <https://docs.snowflake.net/manuals/user-guide/querying-stage.html#supported-functions>

---

Question 42:

**What are some of the ways to improve performance of an external table?**

- 

**Apply compression on the table columns**

- 

**Create materialized views on top of the external table**

**(Correct)**

- 

**Create the external table in columnar format**

- 

**Partition the external table (if the logical folder structure permits)**

**(Correct)**

### **Explanation**

External tables can be partitioned. This is dependent on if your data is residing in a logical folder structure which allows effective partitioning. If you have Enterprise license (&above) you can create materialized views on top of the external table  
Incorrect: Columns in external table can not be compressed or stored as columnar format <https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html#partitioned-external-tables> <https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html#materialized-views-over-external-tables>

---

**Question 43:**

**When unloading files to a local folder on your computer, before the files can be downloaded locally they are first unloaded to ..**

- 
- 

**Snowflake internal Stage**

**(Correct)**

- 

**Another table**

- 

**Cloud storage**

### **Explanation**

The files are first unloaded to an internal stage from where they can be downloaded locally using the GET command <https://docs.snowflake.net/manuals/user-guide/intro-summary-unloading.html#output-data-file-details>

---

**Question 44:**

**If data is loaded into an Internal Snowflake Stage it will be automatically encrypted.**

- 
- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

When copying data to an internal stage the files are encrypted on the client machine before they are transferred to the internal stage.

<https://docs.snowflake.net/manuals/user-guide/security-encryption.html#end-to-end-encryption>

---

Question 45:

You can unload (or extract) data from Snowflake using which of the following mechanisms?

- 

Unload/Extract using a Query

(Correct)

- 

Unload/Extract a complete database

- 

Unload/Extract from a Table

(Correct)

### Explanation

In Snowflake, you can unload or extract data directly from a table or alternatively you can specify a Query using which the data is extracted.

<https://docs.snowflake.net/manuals/user-guide/data-unload-overview.html>

---

Question 46:

Select what are the features of External tables in Snowflake?

- 

Views can be created on top of external tables

(Correct)

- 

They can be joined with other tables

(Correct)

- 

They are temporary and are dropped at the end of the session

- 

**They are Read Only**

**(Correct)**

### **Explanation**

External tables are read-only and don't allow modification of data. You can use them in JOIN conditions. Views can be created on top of external tables. They are not dropped at the end of the session. <https://docs.snowflake.net/manuals/user-guide/tables-external-intro.html>

---

**Question 47:**

**What are the valid properties of the VARIANT data type?**

- 

**Any data type can be cast to the VARIANT data type**

**(Correct)**

- 

**Only STRING data type can be cast to the VARIANT data type**

- 

**VARIANT data type can store unlimited amount of data**

- 

**VARIANT data type can store up to 16MB of compressed data**

**(Correct)**

### **Explanation**

VARIANT can store up to 16 MB of compressed data and any data type can be cast to VARIANT <https://docs.snowflake.net/manuals/sql-reference/data-types-semistructured.html#variant>

---

**Question 48:**

**When un-encrypted files are staged in a Snowflake internal stage, the files are encrypted by Snowflake automatically.**

- 

**True**

**(Correct)**

- 

**False**

#### **Explanation**

Snowflake will automatically encrypt un-encrypted files when they are added to an internal stage. <https://docs.snowflake.net/manuals/user-guide/intro-summary-loading.html#encryption-of-staged-files>

---

**Question 49:**

**Which of the following actions can be taken by a Snowflake administrator regarding MFA (multi factor authentication) for a user?**

- 

**Temporarily allow bypass of MFA**

**(Correct)**

- 

**Disable MFA for that user completely**

**(Correct)**

- 

**Enroll the user in MFA**

#### **Explanation**

An administrator can allow a temporary bypass of MFA for a user (in cases if they misplace their MFA device i.e. the phone). The administrator can also disable MFA, allowing the user to re-enroll. The enrollment must be done by the user themselves, administrators can not enroll other users.

<https://docs.snowflake.net/manuals/user-guide/security-mfa.html#managing-mfa-for-your-account-and-users>

---

**Question 50:**

**What is true about the Snowflake Data Loading Wizard? Select all that applies**

- 

**The wizard uses the COPY command behind the scenes to load the data**

**(Correct)**

- The data loading wizard is intended to be used for extremely large files
- The data loading wizard doesn't need an active virtual warehouse.
- The data loading wizard is intended to be used for small files (up to 50 MB)  
**(Correct)**

#### Explanation

The data loading wizard being browser based, is intended for loading for small files. The data loading wizard uses the COPY & PUT command behind the scenes, therefore it requires an active virtual warehouse connection.  
<https://docs.snowflake.net/manuals/user-guide/data-load-web-ui.html>

# Account and Security

*Q 1. You are the owner of a table T1 which is in schema S1. The schema is in database D1. In order to grant read-only permissions of this table to a newly created role R1, you will need to... (select all that apply)*

- A. Grant ‘USAGE’ on database D1
- B. Grant ‘SELECT’ on database D1
- C. Grant ‘USAGE’ on schema S1
- D. Grant ‘SELECT’ on schema S1
- E. Grant ‘USAGE’ on table T1
- F. Grant ‘SELECT’ on table T1

## Correct Answers: A, C, and F

To allow a role to use your database objects, you, as the owner must grant appropriate privileges on the database, schema, and objects. First up, you will need to grant USAGE privilege on the parent database and schema. USAGE privilege enables the grantee to refer to the database and schema in the query. Then, you will need to grant SELECT privilege on the table. The SQL statements are given below:

Grant usage on the database:

```
grant usage on database D1 to role R1;
```

Grant usage on the schema:

```
grant usage on schema D1.S1 to role R1;
```

Grant the ability to query an existing table:

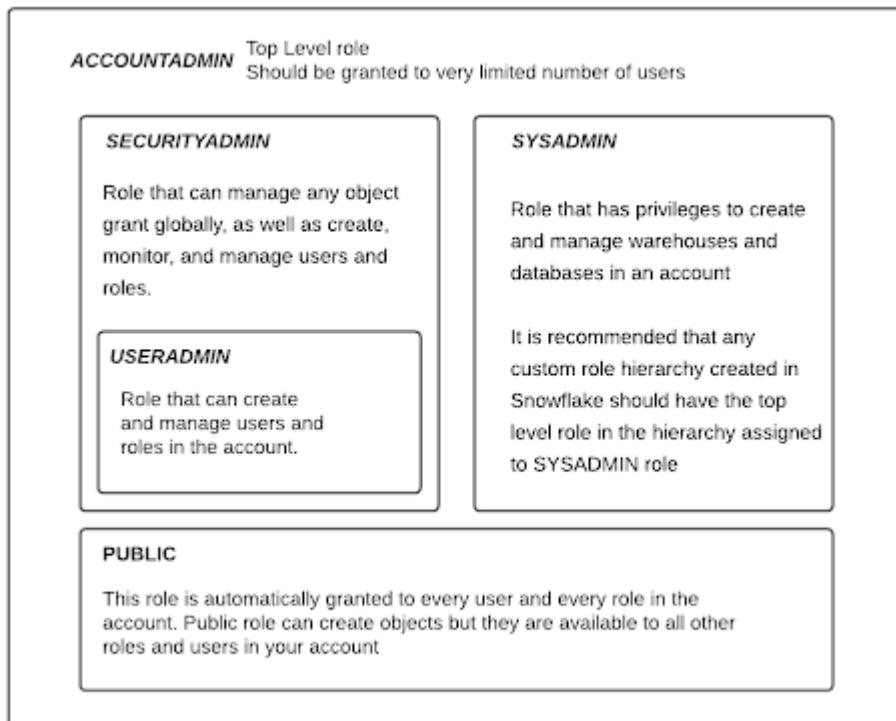
```
grant select on table D1.S1.T1 to role R1;
```

*Q 2. Identify system-defined roles in Snowflake from the roles given below. (select four)*

- A. AccountAdmin
- B. SysAdmin
- C. Auditor
- D. RoleAdmin
- E. SecurityAdmin
- F. Public

## Correct Answers: A, B, E, and F

A role in snowflake is essentially a container of privileges on objects. Roles are assigned to users to allow them to perform actions on the objects. A role can be directly assigned to the user, or a role can be assigned to a different role leading to the creation of role hierarchies. The role at a higher level in the hierarchy inherits the privileges of the role at a lower level in the hierarchy. Snowflake defined 5 system-defined roles. These roles, the key features, and the hierarchy of system-defined roles are shown in the diagram below.



account admin snowflake

## Snowflake Virtual Warehouses

*Q 3. What are the minimum and the maximum number of clusters in a multi-cluster warehouse?*

- A. Minimum: 1, Maximum: 99
- B. Minimum: 1, Maximum: 100
- C. Minimum: 1, Maximum: 10
- D. Minimum: 1, Maximum: unlimited

**Correct Answer: C**

- **Answer choice-C is correct.** In a multi-cluster warehouse, the minimum number of clusters can be 1 and the maximum number of clusters can be 10.

*Q 4. When a virtual warehouse is started or resized or resumed, the minimum billing charge is 1 minute.*

- A. True
- B. False

**Correct Answer: A**

The above statement is True. When a new warehouse is started afresh, or when you resize an existing running warehouse or when a warehouse is resumed (automatically or manually), the warehouse is billed for a minimum of 1 minute's worth of usage. After the 1st minute has elapsed, all subsequent billing is charged per second.

► **Practical Info** – There is no benefit in stopping a warehouse before the first 1st minute is over because you are anyway going to pay for the first 60-seconds, as the discussion given above shows.

► **Exam Tip** – It is NOT necessary to stop a virtual warehouse to resize it. You can resize a virtual warehouse even when it is executing a query.

*Q 5 . When a multi-cluster warehouse is suspended, which of the following Snowflake cache will be purged?*

- A. Metadata cache
- B. Remote disk cache (Resultset cache)
- C. Local disk cache
- D. All of the above

#### **Correct Answer: C**

- When a multi-cluster warehouse (or a single cluster warehouse) is suspended, only the local disk cache is dropped. This cache is also known as warehouse cache. This is because this cache is essentially the fast SSD of the warehouse. So when the warehouse is gone, this cache is also gone with it. **Answer choice-C is correct.**
- Metadata cache contains Snowflake account object information and statistics. It is always on and it is never dropped. **Answer choice-A is incorrect.**
- Remote disk cache (resultset cache) caches exact query results including aggregations etc. and it persists the exact resultset for 24-hours after the query is executed irrespective of the state of the warehouse. **Answer choice-C is incorrect.**

## **Snowflake Overview and Architecture**

*Q 6. From the given options, select the cloud services provider supported by Snowflake architecture. (three choices)*

- A. Amazon Web Services (AWS)
- B. Google (GCP)
- C. HP Virtual Private Cloud
- D. Cloud Foundry
- E. Microsoft Azure

#### **Correct Answers: A, B, and E**

Snowflake is a SaaS product. It supports three public cloud platforms given below.

1. Amazon Web Services (AWS)
2. Microsoft Azure (Azure)
3. Google Cloud Platform (GCP)

Snowflake does NOT support any other cloud platform, any private cloud deployment, or on-prem deployment. Therefore, option choices – HP Virtual Private Cloud and Cloud Foundry – are not correct.

► **Practical Info:** It is not necessary that you have to host Snowflake on the same cloud platform where your organization's other IT applications are hosted. For example, it is perfectly okay to have Snowflake hosted on Azure and your other IT applications in your organization hosted on AWS/GCP.

► **Exam Tip:** It must be noted that each Snowflake account is hosted in one cloud platform and one region, and you must choose cloud platform and region while opening an account. You cannot change it thereafter. If you want to change region and/or cloud platform at a later point, you will need to open a new Snowflake account in that region and/or cloud platform.

Your organization may open multiple Snowflake accounts (such as Dev, Test, Prod), and they can all reside on the same cloud platform and the same region, OR you may choose to open them on different cloud platforms and/or regions.

*Q 7 . What are the key considerations for choosing a specific region for your Snowflake account? (Select all that apply)*

- A. Cost
- B. Proximity to the end-users
- C. Data sovereignty
- D. Regulatory requirements

**Correct Answers: A, B, C, and D**

- **Option 1 is correct.** Snowflake offers different pricing for computing and storage depending upon the region you choose. For example, a Snowflake instance running in the Mumbai region costs \$3.30/credit, whereas the same instance running in the Frankfurt region costs \$4.00/credit. So if you are looking to save costs, you would choose a low-cost region.
- **Option 2 is correct.** The closer you are to your end-users, data transfer latency would be lower so you can transfer and access data faster. Hence, it is always a good practice to set up your Snowflake account in the region closer to your end-users.
- **Option 3 and 4 are correct.** As is often the case in most countries, there are data sovereignty and regulatory restrictions on where the data of users is stored. For example, in the EU, certain sensitive user data is now allowed to cross the EU boundary. India also has similar restrictions enforced by the law. Hence, if you build software that will store such sensitive user data, you will need to check the regulatory requirements and choose the Snowflake region accordingly.

*Q 8. Informatica, Matillion, Azure data factory are examples of ..... partners of the Snowflake partner ecosystem.*

- A. Data Management
- B. Business Intelligence
- C. Data Integration
- D. Data Science

## Correct Answer: C

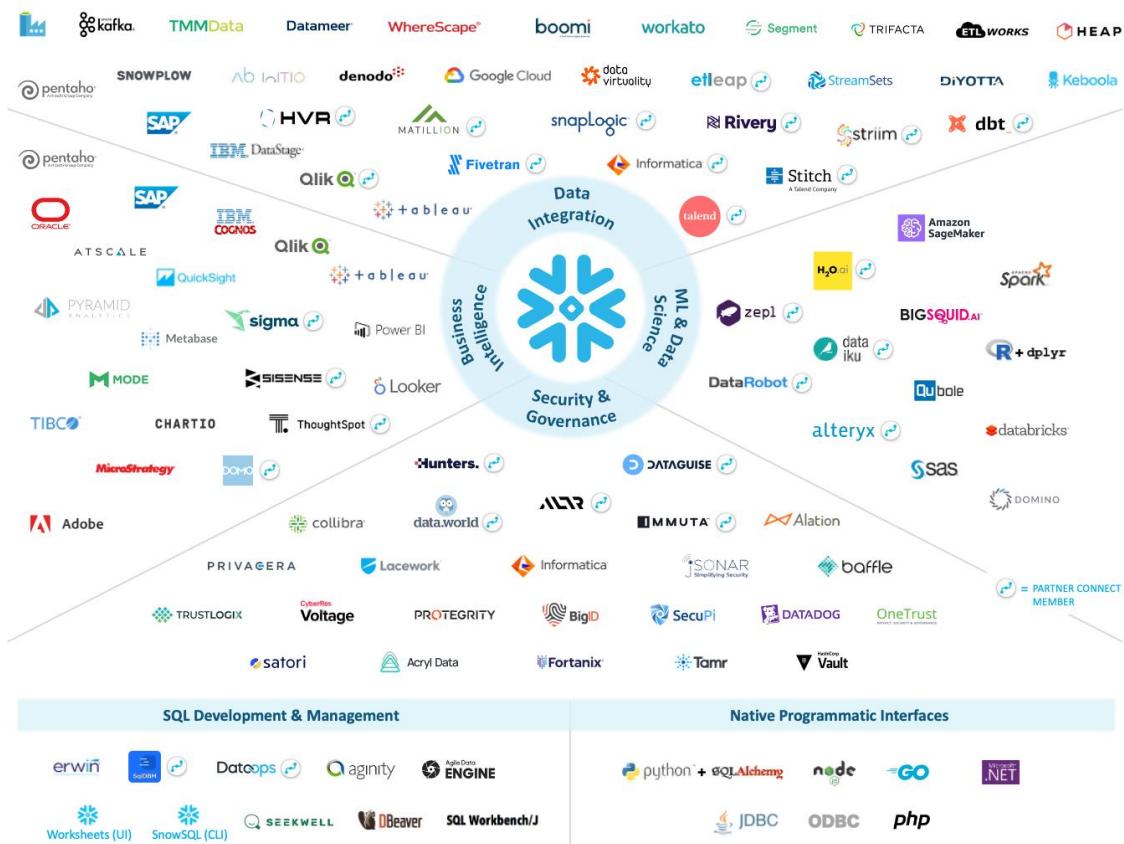
Informatica, Matillion, and Talend are **data integration** partners that provide the following functionalities:

Extract – Exporting data from the data source.

Transform – Modifying data according to business rules.

Load – Loading data into Snowflake.

The below diagram shows Snowflake partner ecosystem players divided based on the functional categories.



overview

Snowflake partner eco-system (image source: Snowflake official documentation)

*Q 9 . Which of the following programming languages are supported in Snowflake to write user-defined functions (UDFs)?*

- A. SQL
- B. Java
- C. JavaScript
- D. Python

ecosystem-

## **Correct Answers: A and C**

User-defined functions (UDFs) can be used in Snowflake to write custom business logic which is hard to encapsulate within individual SQL statements. UDFs (together with Stored Procedures) enable database-level programming in Snowflake.

At present, Snowflake supports two programming languages to write UDFs: Javascript UDFs and SQL UDFs.

► **Practical Info**— Stored procedures in Snowflake are written in Javascript. The javascript code can issue SQL statements to process data in the Snowflake database. An example of a Snowflake stored procedure that clears a log table using SQL is given below: (this is only for your understanding, we are not expecting coding related questions in the exam)

```
create or replace procedure clear_log()
returns int not null
language javascript
as
$$
var delete_log_sql_command = "delete from user_log";
var return_value = 0;
var statement1 = snowflake.createStatement( {sqlText: delete_log_sql_command} );
var result_set1 = statement1.execute();
return_value = 1;
return return_value;
$$
```

*Q 10. Which Snowflake edition supports transmitting data in encrypted form over the network between VPCs (virtual private cloud)?*

- A. All editions
- B. Enterprise edition and above
- C. Business Critical edition and Above
- D. All except for the Virtual Private Snowflake (VPS) edition

## **Correct Answer: C**

A Snowflake account on AWS (or Azure) is implemented as a VPC. There are two ways to establish communication between your Snowflake VPC and other VPCs (e.g. your organization's VPC). One is to transmit the traffic over the public internet. Other (and safer) option is to establish an exclusive, highly secure network between your Snowflake account and your other AWS VPCs (in the same AWS region), fully protected from unauthorized access. To implement this secure channel of communication between VPCs, AWS supports a feature called AWS PrivateLink (Azure also supports a similar feature called Azure PrivateLink). Snowflake offers support for AWS PrivateLink (and Azure PrivateLink) based communication in Business Critical Edition and above.

*Q 11. More than one clustering key can co-exist in a Snowflake table*

- A. True
- B. False

### **Correct Answer: B**

This statement is false. You can define at most one clustering key in a Snowflake table to organize micro-partitions. When you define a clustering key, Snowflake will reorganize the naturally clustered micro-partitions and will relocate related rows to the same micro-partition and group them according to the clustering key. This process is called **Reclustering**.

► **Practical Info** – Reclustering happens automatically once a clustering key is defined for a table. The process consumes credits. So be cognizant of the cost when you go for reclustering.

*Q 12. Which of the following statements will you use to change the warehouse for workload processing to a warehouse named ‘COMPUTE\_WH\_XL’?*

- A. SET CURRENT\_WAREHOUSE = COMPUTE\_WH\_XL
- B. USE WAREHOUSE COMPUTE\_WH\_XL;
- C. USE CURRENT\_WAREHOUSE(‘COMPUTE\_WH\_XL’);
- D. SET CURRENT\_WAREHOUSE = COMPUTE\_WH, SIZE = XL;

### **Correct Answer: B**

A session context in Snowflake consists of 4 objects:

1. Role
2. Warehouse
3. Database
4. Schema

You can set appropriate session context using a set of SQL statements given below. These statements specify the role, warehouse, database, or schema to use for the current session:

You can set appropriate session context using a set of SQL statements given below. These statements specify the role, warehouse, database, or schema to use for the current session:

Use warehouse <>warehouse-name<>

Use role <>role-name<>

Use database <>database-name<>

Use schema <>schema-name<>

So the correct answer of this question is:

Use warehouse COMPUTE\_WH\_XL;

*Q 13. In the case of a Snowflake account created on AWS, ..... is responsible for the management of Availability Zones?*

- A. Customer
- B. Snowflake
- C. Cloud Provider
- D. It is a shared responsibility

### **Correct Answer: C**

An Availability Zone (AZ) is essentially a group of one or more physically separated data centers with redundant power, networking, and connectivity and located within a single cloud region. Each cloud region has multiple AZs (most have 3 or more). The AZs are connected with high-bandwidth, low-latency network infrastructure and support synchronous replication. Due to the automatic, synchronous replication and physical separation plus isolation, AZs enable your applications and databases to be highly available. When one AZ within a region fails, another one should remain active and the switch between the two zones is silent and transparent to customers. The cloud providers are responsible for the maintenance of AZ infrastructure and replication of your databases and switching to alternate AZ in case of any failure.

## **Snowflake Storage and Protection**

*Q 14. Once the time-travel period has expired, it is possible to request Snowflake support to retrieve historical data for a period of*

- A. Day
- B. Days
- C. Days
- D. It depends on the Snowflake edition
- E. It is user-configurable

### **Correct Answer: B**

After the time travel data retention period is over, you can use Snowflake's fail-safe feature to recover your data. The duration of the fail-safe period is 7 days (Answer choice 2 is correct). This is a fixed duration and cannot be changed. Only Snowflake support personnel can help recover data during the fail-safe period. The fail-safe feature is available to all customers irrespective of the Snowflake edition.

The below diagram succinctly summarizes key differences between Snowflake's two important data protection features – time-travel and fail-safe.



*Q 15. Which of the following statements are TRUE concerning a stream object in Snowflake? (select all that apply)*

- A. A stream object provides a record of DML changes (inserts, updates, deletes) made to a table at row level.
- B. A stream object can keep track of DML changes for the entire life of a table.
- C. Streams on materialized views are not supported.
- D. Streams on external tables are not supported.

### **Correct Answers: A and C**

**Answer choice-A is correct.** A stream object provides change tracking over a source table. It records DML changes made to tables, (inserts, updates, and deletes) as well as metadata about each change. This is referred to as Change Data Capture (CDC), and this feature is extensively used in data warehousing scenarios to create data pipelines. Please note that the stream object itself does not store this data. It relies on the version history of source data maintained in the metadata layer.

**Answer choice-B is incorrect.** Stream object keeps track of DML changes of a source table up until the data retention period of the source table. After that, the DML changes are no longer accessible.

**Answer choice-C is correct.** Currently, Snowflake does not support creating stream objects on materialized views.

**Answer choice-D is incorrect.** Snowflake supports creating insert-only stream objects on external tables.

*Q 16. Only one stream object can be created on a source table*

- A. True
- B. False

### **Correct Answer: B**

- **The above statement is false.** You can create any number of streams on a source table. These streams can have the same or different offset positions. One example of creating multiple streams is when you want to report month-on-month changes, week-on-week changes, and day-on-day changes happening in a product inventory table of your POS database. In this case, you may create three streams on the table to record monthly, weekly and daily changes. All three streams exist independently of each other with their respective offset positions.

*Q 17. Snowflake replicates the following layer(s) across availability zones (select two)*

- A. Cloud Services Layer
- B. Warehouse Layer
- C. Storage Layer
- D. Cloud Agnostic Layer
- E. Data Sharing Layer

### **Correct Answers: A and C**

Snowflake replicates the cloud services layer (**Answer choice-A**) and the storage layer (**Answer choice-C**) across the availability zones of your cloud and region.

## **Data loading/unloading**

*Q 18. Which Snowflake edition supports Search Optimization Service to improve performance of point lookup queries?*

- A. All editions
- B. Enterprise edition and above
- C. Business Critical edition and Above
- D. All except Virtual Private Snowflake (VPS) edition

### **Correct Answer: B**

Search Optimization Service in Snowflake is a background service that is designed to improve the performance of point lookup queries. A point lookup query is essentially a SELECT statement that returns only a small number of distinct rows from a large dataset. The service runs transparently from an end-user's standpoint. Only the Enterprise edition and above provide this feature.

*Q 19. Consider the XML given below. The XML file is loaded in a User Stage:*

```
<bookshelf>
<book>
<name>Famous Five</name>
```

```

<author> Roald dahl</author>
</book>
<book>
<name>Secret Seven</name>
<author> Enid Blayton</author>
</book>
<book>
<name>Mahashweta</name>
<author> Sudha Murthy</author>
</book>
</bookshelf>

```

*To remove the top-level element <bookshelf/> and load <book/> elements into separate rows in a Snowflake table, which of the following file format options will be used in the COPY INTO command?*

- A. STRIP\_TOP\_ELEMENT
- B. STRIP\_OUTER\_ELEMENT
- C. STRIP\_OUTER\_ARRAY
- D. STRIP\_FIRST\_ELEMENT

### **Correct Answer: B**

Enabling the STRIP\_OUTER\_ELEMENT file format option for the COPY INTO <table> command to remove the outer element and load the <book/> records into separate table rows:

```

copy into <table> from @~/<file>.xml
file_format = (type = 'XML' strip_outer_element = true);

```

► **Exam Tip:** If the semi-structured file format is JSON (instead of XML), you will need to enable STRIP\_OUTER\_ARRAY file format option for the COPY INTO command.

## **Performance and Tuning**

*Q 20. History' tab in the Snowflake Web Interface supports viewing of user queries submitted in the past..... day(s)*

- A. 24 hours or 1 day
- B. 7 days
- C. 14 days
- D. 90 days

## **Correct Answer: C**

‘Query History’ view in Snowflake Web UI stores queries executed and query profile information for **all queries, executed by all users, using any interface (e.g., Web UI, SnowSQL, Business Application... or any other interface) for a period of 14 days**

► **Exam Tip:** In addition to storing queries and profiler, it also allows users to download the resultset of queries within 24 hours. However, the most important point here is: You cannot view other people’s query results if your role and the role of the user who executed the query are different.

*Q 21. What is the optimal file sizing recommendation for loading data into Snowflake?*

- A. 10 MB to 100 MB, uncompressed
- B. 100 MB to 250 MB, compressed
- C. 100 MB to 250 MB, uncompressed
- D. 10 GB to 100 GB, compressed

## **Correct Answer: B**

Snowflake recommends that to optimize the number of parallel operations for a load, data files should **roughly be 100-250 MB in size, compressed**. In case input data files are smaller than this, it is better to aggregate multiple input files. On the other hand, if the input data files are larger than this, it is better to split them into smaller files to match this guidance.

Snowflake also recommends against loading very large files (e.g., 100 GB+ sizes) without splitting them as splitting of files will take advantage of parallelism.

► **Practical Info** – The recommendation applies to both modes of data loading – bulk load using SnowSQL (COPY INTO) and continuous/micro-batch using Snowpipe.

*Q 22. Selecting the right clustering key can dramatically improve query performance. Following are the good yardsticks to choose an appropriate clustering key:*

- A. Use columns that are most frequently found in WHERE clause
- B. A date column is usually not a good candidate to be in the clustering key
- C. A join column can be considered as clustering key
- D. Using expressions in the clustering key may help reduce cardinality.

## **Correct Answers: A, C, and D**

**Answer choice-A is correct.** Choosing the clustering keys from dimensions (i.e. columns in the WHERE clause) used in often repeating queries will significantly improve the performance of these queries.

**Answer choice-B is incorrect.** If you have queries that filter the data based on date criteria, you can use the date column in the clustering key. However, you will need to check the cardinality of the date column before using it in the clustering key.

**Answer choice-C is correct.** Columns used in the join clauses with other tables are worthy candidates to be in the clustering key.

**Answer choice-D is correct.** In case you would like to use a particular column in a clustering key but the column has very high cardinality, using expressions may help reduce cardinality. For example, you can use SUBSTR expression to get a relevant portion of a string column for use in the clustering key.

*Q 23. Selecting the right clustering key can dramatically improve query performance.  
Following are the good yardsticks to choose an appropriate clustering key:*

- A. Use columns that are most frequently found in WHERE clause
- B. A date column is usually not a good candidate to be in the clustering key
- C. A join column can be considered as clustering key
- D. Using expressions in the clustering key may help reduce cardinality.

### **Correct Answers: A, C, and D**

**Answer choice-A is correct.** Choosing the clustering keys from dimensions (i.e. columns in the WHERE clause) used in often repeating queries will significantly improve the performance of these queries.

**Answer choice-B is incorrect.** If you have queries that filter the data based on date criteria, you can use the date column in the clustering key. However, you will need to check the cardinality of the date column before using it in the clustering key.

**Answer choice-C is correct.** Columns used in the join clauses with other tables are worthy candidates to be in the clustering key.

**Answer choice-D is correct.** In case you would like to use a particular column in a clustering key but the column has very high cardinality, using expressions may help reduce cardinality. For example, you can use SUBSTR expression to get a relevant portion of a string column for use in the clustering key.

*Q 24. Which cache does the query optimizer refer to first? Second? And Third?*

- A. First: warehouse cache. Second: results cache. Third: metadata cache
- B. First: warehouse Cache. Second: metadata Cache. Third: results cache
- C. First: metadata Cache. Second: results Cache. Third: warehouse cache
- D. First: metadata Cache. Second: warehouse Cache. Third: results cache

### **Correct Answer: C**

The query optimizer in the Snowflake cloud services layer examines the metadata cache first. Thereafter it refers to the results cache next and finally the warehouse cache.

## **Semi-Structured data**

*Q 25. Which of the following data types are supported in Snowflake for semi-structured data? (select all that apply)*

- A. VARIANT
- B. OBJECT
- C. TEXT
- D. CLOB
- E. ARRAY

### **Correct Answers: A, B, and E**

Snowflake supports three data types for processing semi-structured data.

**VARIANT** – The basic (and most used) data type for storing semi-structured data.

**OBJECT** – Used to store key-value pairs, where the key is a non-empty string, and value is semi-structured data in VARIANT data type.

**ARRAY** – Used to store arrays of semi-structured data where array index is non-negative integer and value is semi-structured data in VARIANT data type.

When are you planning to take SnowFlake SnowPro Core Exam? Do share your feedback in comments and help other learners with real exam experience. Whizlabs [practice tests](#) are based on real exam pattern; we have received good feedback from exam takers.

### **Reference links:**

- <https://www.youtube.com/channel/UCs10x-muRrTQMJ4Ya-fmIlw>
- <https://docs.snowflake.com/en/>

## Domain : Snowflake Overview and Architecture

*Q26 : Following are the key reasons that make Snowflake a ‘‘True SaaS’’ product.*

- A.** Snowflake completely runs on cloud infrastructure and is accessible over the internet
- B.** There’s no hardware or software to install, configure, maintain or upgrade at the user’s end
- C.** Snowflake supports key properties of as-a-service offering such as on-demand and instance elasticity, unlimited scalability of compute and storage and pay-per-use
- D.** All of the above

### **Correct Answer: D**

### **Explanation**

Snowflake is fully managed, cloud-native, on-demand software offered as a service (SaaS – software as a service) over the internet. It is built from the ground up for the cloud. It is fully hosted on public cloud infrastructure and uses cloud resources for compute and storage. At the user’s end, there’s no hardware or software installation required. There are no patches or upgrades to be done by the user. New releases are rolled

out transparently and don't require user intervention. The software is available on-demand over the internet. You can log on to [snowflake.com](http://snowflake.com), open an account, and you are ready to use Snowflake in your organization.

Snowflake offers unlimited and instant scalability of compute and storage resources. You can scale up and down on-demand based on your need. You need to only pay for the Snowflake resources that you use. Owing to these unique design elements, Snowflake is currently the most popular cloud data warehouse solution for enterprises worldwide.

**Further Reading:** [Key Concepts & Architecture — Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q27 : Following programming languages/drivers, support is available out-of-the-box (natively) in Snowflake (select all that apply)*

- A.** JDBC/ODBC driver
- B.** Google Flutter
- C.** Go, Node.js driver
- D.** Scala connector
- E.** Python connector
- F.** Spark connector

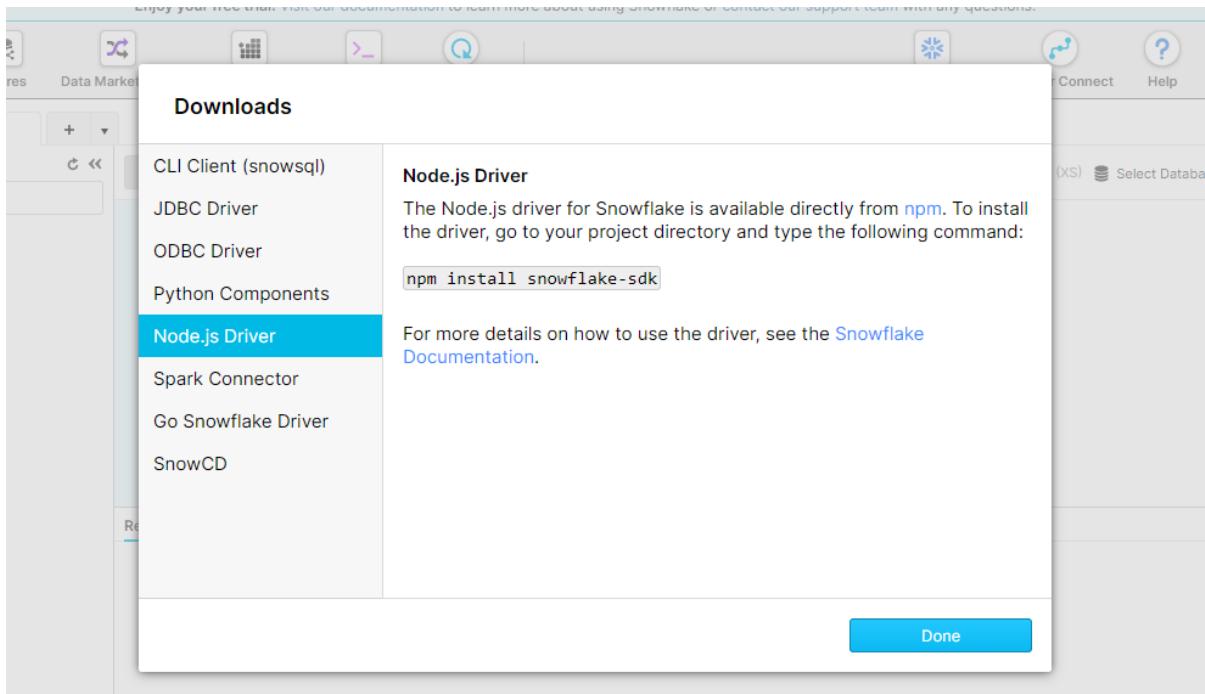
**Correct Answers: A, C, E and F**

## Explanation

Snowflake provides two types of native clients to develop business applications.

1. Connectors – For Python and Spark
2. Drivers – For JDBC, ODBC, Node.js, Go

These components are available for download in Snowflake Web UI. Go to Help → Downloads in the Web UI.



**Further Reading:** Native Programmatic Interfaces — Snowflake Documentation

## Domain : Snowflake Overview and Architecture

*Q28 : To execute a query in Snowflake, you need to provide context. The context consists of the following. (select four)*

- A.** Virtual Warehouse (VW)
- B.** Database
- C.** Schema
- D.** Role
- E.** User
- F.** Session

**Correct Answers: A, B, C and D**

## Explanation

A query context in Snowflake consists of 4 objects:

1. Role – Every query you execute should have a role that it uses to access the objects referred to in the query.
2. Warehouse – You need to assign a warehouse to use to a query that will do the necessary computation required for the query and return a resultset
3. Database – Every query must be assigned a database.

4. Schema – Every query must be assigned a specific schema within the database. Objects referred to in the query are looked up in the schema.

In Snowflake WebUI, you can set the context at the worksheet level by selecting appropriate values in the context drop-down as shown below.

The screenshot shows the Snowflake WebUI interface. On the left, there's a sidebar with a 'Run' button and a message indicating 'All Queries | Saved 3 seconds ago'. The main area has a code editor with the following SQL:

```
1 //your query goes here
2 select * from databases;
```

A blue box highlights the top navigation bar, specifically the 'Context' dropdown menu. This menu is open, showing several options: 'SYSADMIN', 'COMPUTE\_WH (XS)', 'DEMO\_DB', 'INFORMATION\_SC...', and three other items represented by small icons. A blue bracket is drawn around this entire context menu area.

Alternatively, you can set the context using a set of SQL statements given below:

Use warehouse <><warehouse-name>>

Use role <><role-name>>

Use database <><database-name>>

Use schema <><schema-name>>

## Domain : Snowflake Virtual Warehouses

*Q29 : As a Snowflake data engineer, you have a task to ingest data generated by sensors located on electric poles. As more sensors are getting installed in the city, the ingestion workload is increasing continuously. You have been tasked to support the growth in data loading volume while optimizing the cost of ingestion. Which best practices would you employ to increase the performance of your data ingestion workload? (select all options that apply)*

- A.** Increase the size of the virtual warehouse (S >> M >> L)
- B.** Split input files into small batches
- C.** Using SnowPipe as SnowPipe method is better in ingesting continuous micro-batch workloads
- D.** Use SnowSQL as the SnowSQL method is better in ingesting continuous micro-batch workloads

**Correct Answers: A, B and C**

**Explanation**

**Option choice 1 is correct** because Scaling up a virtual warehouse (i.e., resizing the virtual warehouse to increase its computing power) helps in faster processing of ingestion workload. This is because as you increase the size of the warehouse, the performance increases **by a factor of two**. That means your data ingestion will run twice as fast.

**Option choice 2 is correct** because when you split the file into batches and submit them for processing, Snowflake will take advantage of multiple cores of the CPUs and process the files in parallel. Snowflake recommends that instead of submitting 1 large file for processing, you should submit multiple small files in 10-100 MB. Refer to the video lesson in the further reading section for an experiment that shows this phenomenon.

**Option choice 3 is correct, and Option choice 4 is not correct** because the SnowPipe feature in Snowflake is designed to support the ingestion of a continuous stream of micro-batches. Because it has per-second billing, it is more cost-effective than SnowSQL.

**Further Reading:** [Concurrency in Snowflake – YouTube](#), [Getting Started – Introduction to Virtual Warehouses – YouTube](#)

## Domain : Snowflake Overview and Architecture

*Q30 : Namespace in Snowflake is comprised of:*

- A.** Database and schema
- B.** Account, database, and schema
- C.** Database, schema, and table
- D.** Warehouse, database, and schema

**Correct Answer: A**

### Explanation

When you want to perform any operation (e.g., DDL or DML) on Snowflake's database objects, you need to specify the namespace. Namespace in Snowflake consists of the name of the database and the schema's name within that database. **Therefore answer choice-A is correct.** There are two ways you can specify the namespace. 1) By setting the database and schema in the session context. In this case, Snowflake will infer the namespace from the context. This is referred to as an 'inferred namespace' 2) By explicitly specifying the namespace in the query. This is called an 'explicit namespace.'

Inferred namespace:

Explicit namespace :

The screenshot shows the Snowflake UI interface. At the top, there are tabs for 'heet' and 'New Worksheet'. Below the tabs, there are buttons for 'Run' (highlighted in blue), 'All Queries' (checkbox), and 'Saved 0 seconds ago'. To the right, there are user icons for 'SYSADMIN' and 'COMPUTE\_WH (XS)'. On the far right, there are buttons for 'Select Database' and 'Select Schema'.

```
select * from snowflake_sample_data.information_schema.databases;
```

This query explicitly specifies the database and schema (this is an example of explicit namespace)

► **Practical Info** – In Snowflake, The concept of namespace makes it super easy to write a query that processes data across databases within the same account. You just need to specify the right namespace of the database objects you are interested in querying. Many traditional databases may require you to write expensive and effort-intensive ETL routines to do such operations.

See below example where an explicit namespace is used to join tables across two separate databases – Student\_db and course\_db.

```
select * from student_db.public.students_master student, course_db.public.course_master course,  
course_db.public.student_course_mapping_table where
```

```
student.rollno = mapping_table.student_id and mapping_table.course_id = course.course_id;
```

### Further Reading: [Database, Schema, & Share DDL — Snowflake Documentation](#)

The YouTube video given below explains how Snowflake simplifies querying across databases:

- [Query Multiple Databases Within Snowflake – YouTube](#)

## Domain : Snowflake Overview and Architecture

*Q31 : A BOOLEAN data type in Snowflake can represent (select all that apply)*

- A. TRUE**
- B. FALSE**
- C. NULL**
- D. BOOLEAN datatype is not supported**

### Correct Answers: A, B and C

### Explanation

Snowflake supports BOOLEAN data type. You can define a column in a table that can store boolean values. See the example given below.

```
Create table german_language_students_table (rollno int, name string, knows_basic_german boolean);
```

BOOLEAN data type here can have one of the three values: TRUE, FALSE, or NULL. NULL represents ‘unknown.’

— the student knows basic german

```
insert into german_language_students_table values (1, 'amit khanna', TRUE);
```

— the student does not know basic german

```
insert into german_language_students_table values (2, 'shardul mandaloi', FALSE);
```

— it is not known whether the student knows german

```
insert into german_language_students_table values (2, 'siddharth shah', NULL);
```

#### » Exam Tip:

STRING and NUMBER datatype can be implicitly cast into BOOLEAN datatype as follows:

A string ‘true’ is implicitly cast into boolean TRUE, ‘false’ is implicitly cast into boolean FALSE. The numeral 0 is implicitly cast into boolean FALSE. Any non-zero numeral is implicitly cast into boolean TRUE.

**Further Reading:** [Logical Data Types — Snowflake Documentation](#)

## Domain : Snowflake Virtual Warehouses

*Q32 : It is a good practice to create multiple virtual warehouses of various sizes in Snowflake based on the workload.*

**A.** True

**B.** False

**Correct Answer: A**

#### Explanation

A snowflake virtual warehouse is one or more compute clusters that enable users to load data, execute queries and perform other operations on the data.

Snowflake has a unique pay-per-use model to charge the cost of compute. The charge is based on the size of virtual warehouses and the time they are run. When you create a new virtual warehouse, Snowflake only stores

the meta-data of the virtual warehouse. Only when this warehouse is in a running state, you are charged. Hence, creating multiple virtual warehouses doesn't cost money.

Secondly, the workload processing of each warehouse is completely isolated from the workload processing of other warehouses. Therefore, if an L size warehouse is loading data and an M size warehouse is processing user queries, each workload's performance is completely unaffected by the presence of the other.

Owing to these reasons, it is a good practice to tailor each warehouse size to the expected workload. Therefore, the answer to this question is A (True).

**Further reading:** A nice 6-min video lesson by Snowflake on key features of virtual warehouses. [Getting Started – Introduction to Virtual Warehouses – YouTube](#)

## Domain : Snowflake Virtual Warehouses

*Q33 : Following are the two modes in which you can run a multi-cluster virtual warehouse in Snowflake.*

- A.** Dynamic
- B.** Maximized
- C.** Auto-Scale
- D.** Multi-threaded

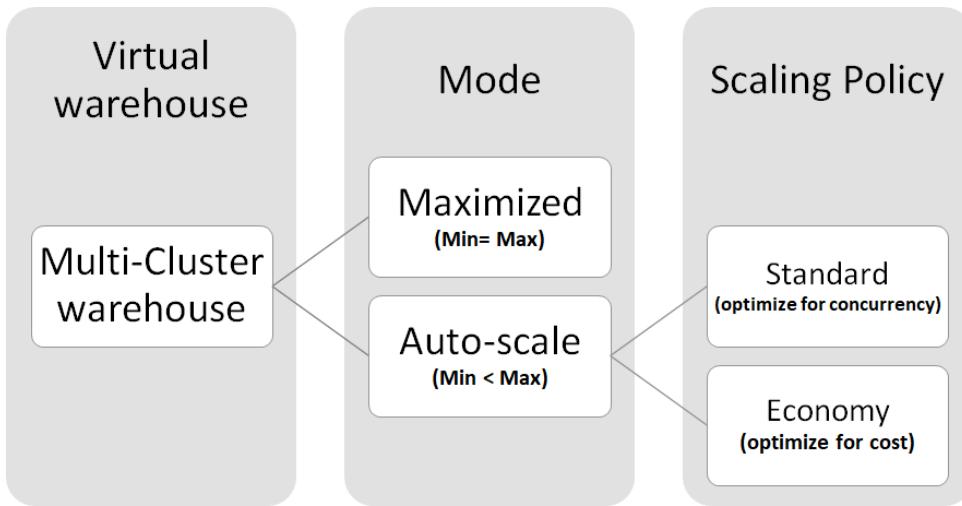
### Correct Answers: B and C

### Explanation

A multi-cluster virtual warehouse consists of a group of compute clusters provisioned either based on the user/query load. There are two modes in which you can create a multi-cluster warehouse. These modes govern the way compute Snowflake provisions clusters.

- Maximized – In this mode, Snowflake starts all compute clusters within the warehouse when the warehouse is started. This mode is enabled by setting the **same value for Min. and Max. clusters**
- Auto-Scale – In this mode, Snowflake starts the warehouse with Min. number of clusters and spins up new clusters incrementally and automatically as and when the user/query load increases. The maximum number of clusters Snowflake can add is governed by the value given in Max. clusters. This mode is enabled by setting different values for Min. and Max. clusters.

» **Exam Tip:** To provide further control over the way the 'Auto-scale' mode operates, Snowflake provides a property called 'SCALING\_POLICY'. This property can have two values: 1) Standard 2) Economy.



**Further Reading:** [Multi-cluster Warehouses — Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q34 : The following object in Snowflake is also known as the ‘data dictionary’ of a database.*

- A.** PUBLIC schema
- B.** Data Catalog
- C.** INFORMATION\_SCHEMA schema
- D.** None of the above

**Correct Answer: C**

### Explanation

When you create a database in Snowflake, two schemas get automatically created within the database. 1. INFORMATION\_SCHEMA schema and 2. PUBLIC schema.

INFORMATION\_SCHEMA is a read-only schema. It contains only views. The views have all meta-data information about the database objects e.g., all tables of that database, all columns, all stages, all sequences etc. These views can be queried to retrieve meta-data information about the database. Hence, the INFORMATION\_SCHEMA schema is also known as the data dictionary of a database.

MYDATABASE

INFORMATION\_SCHEMA

No Tables in this Schema

Views

- APPLICABLE\_ROLES
- COLUMNS
- DATABASES
- ENABLED\_ROLES
- EXTERNAL\_TABLES
- FILE\_FORMATS
- FUNCTIONS
- INFORMATION\_SCHEMA\_CATALOG...
- LOAD\_HISTORY
- OBJECT\_PRIVILEGES
- PIPES
- PROCEDURES
- REFERENTIAL\_CONSTRAINTS
- REPLICATION\_DATABASES

» **Exam Tip** – Another potential exam question: Name the two schemas that every database has in Snowflake.

**Further Reading:** [Information Schema — Snowflake Documentation](#)

## Domain : Snowflake Virtual Warehouses

*Q35 : After a warehouse is suspended by a credit monitor action, the warehouse will not resume until the following event(s) occur. (select all that apply)*

- A. Snowflake account administrator manually resumes the warehouse
- B. Credit quota is increased in the resource monitor configuration
- C. Credit threshold is increased in the resource monitor configuration
- D. Next monitoring cycle (monitoring interval) starts

**Correct Answers: B, C and D**

### Explanation

If a monitor has a ‘Suspend’ or ‘Suspend Immediately’ action defined and its used credits reach the threshold for the action, all warehouses assigned to the monitor are suspended and cannot be resumed until (any) one of the following event(s) take place.

You alter the resource monitor and increase the credit quota value (**Answer choice-B**) or You can increase the credit threshold for the ‘suspend’ action to a higher value. This will allow warehouses to function again (**Answer choice-C**).

Alternatively, the warehouses will become available again for query processing when the new monitoring cycle (monitoring interval) starts. The new cycle will reset the credit quota. Therefore the warehouses can now be run again. (**Answer choice-D**)

**Answer choice-A is invalid.** A warehouse suspended due to resource monitor action will not resume until the resource monitor configuration is changed.

» **Exam Tip –** In addition to the above, you can also take any of the below actions.

- You alter the resource monitor configuration and remove the warehouse monitoring. This will stop the monitoring of the credit consumption, or
- You drop the resource monitor altogether. This will release the credit quota threshold. Therefore, the warehouse can run again.

**Further Reading:** [Working with Resource Monitors — Snowflake Documentation](#)

## Domain : Data loading/unloading

*Q36 : Snowflake provides the following types of internal stages to load data. (select all that apply)*

- A. Table Stage**
- B. Data Stage**
- C. User Stage**
- D. Named Stage**

**Correct Answers: A, C and D**

## Explanation

A stage is essentially a location in the cloud where you will place your data to load in Snowflake. A stage can be internal (i.e., within Snowflake) or external (outside Snowflake in supported cloud environments i.e., AWS, Azure or GCP). Snowflake provides 3 types of internal stages to load data: table stage, user stage, named stage. A comparison of these stages and corresponding features are given below. You may get a variety of questions in the exam about the features of each stage.

	<b>Table stage</b>	<b>User Stage</b>	<b>Named stage</b>
Purpose	This is a table-specific stage. Automatically available for every Snowflake table. Can copy data into that specific	This is a user-specific stage. Automatically available for every user. Can copy data into any table to which the user has	Most flexible of the three. Can copy data into any table and can be used by any user

	table only. Can be used by any user.	access. One user cannot access another user's stage.	having the privileges to access that stage.
Name	Has the same name as the table name.	Has the same name as the user name	first-class database objects and can have any user-defined name
Referenced as	@%<stage-name>	@~<stage-name>	@<stage-name>
Data transformations	Not supported	Not supported	Supported
Web UI	Cannot be seen in the Stages tab of Web UI	Cannot be seen in the Stages tab of Web UI	Can be seen in the Stages tab of Web UI

**Further Reading:** [Choosing a Stage for Local Files — Snowflake Documentation](#), [Copying Data from an Internal Stage — Snowflake Documentation](#)

## Domain : Data loading/unloading

*Q37 : While loading data into Snowflake, which of the following transformations are supported? (select all that apply)*

- A.** Column reordering
- B.** Column omission
- C.** Derived columns
- D.** Cast operation
- E.** Truncation
- F.** Transpose

**Correct Answers: A, B, D and E**

### Explanation

Snowflake supports 4 types of simple transformations while doing data loading.

The supported transformations are – reordering of columns, omission of columns, data type conversion (also called casting or typecasting) and truncating text strings that exceed the length in the target column. Derived columns and transpose or any other complex data transformation operations are not supported while data loading.

## Domain : Data loading/unloading

*Q38 : It is NOT possible to query data stored in the cloud in Snowflake without first loading data into a Snowflake table.*

- A.** True
- B.** False

### Correct Answer: B

### Explanation

**The above statement is false.** Snowflake supports direct querying of data stored in a cloud location.

Imagine a large organization that has invested time and energy to build a data lake on the Azure cloud (or Amazon cloud or GCP). The data lake contains petabytes of data accumulated over many years across multiple channels. Now, the organization wants to query a portion of this data to derive useful insights without reloading data in Snowflake. For these kind of scenarios, Snowflake provides a feature called ‘External Tables’.

An external table in Snowflake references to your data files in external cloud storage (e.g., data lake). External tables store file-level metadata about the data files in the data lake, such as the filename, a version identifier and related properties. Now when you query an external table, Snowflake executes that query against the data stored in external cloud storage, as if it were inside a database.

► **Practical Info** – When you query an external table, it may appear as if data is being fetched from the database. But in reality, the data is fetched directly from external cloud storage every time you run a query on an external table. Therefore the source of truth for the data is always the external cloud store.

**Further Reading:** [Overview of Data Loading — Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q39 : From the statements given below, select the statements which are TRUE for reader accounts. (select all that apply)*

- A.** A reader account facilitates data sharing with consumers that do not have a Snowflake account
- B.** One reader account can be used to access multiple shares provided by different provider accounts
- C.** A reader account consumes credit of the provider account
- D.** A reader account has the same access to Snowflake support as the provider account

### Correct Answers: A and C

### Explanation

**Answer choice-A is correct.** If consumers of your share do not have a Snowflake account (e.g., a consumer using other database products such as Oracle or PostgreSQL), Snowflake provides a functionality

called ‘reader account’ which can be used to access the share. A reader is a special Snowflake account created by the provider of the share for the sole purpose of providing access to the share to a consumer. The provider fully manages the reader accounts.

**Answer choice-B is incorrect.** Each reader account belongs to the provider account that created it. This means that if you need to access multiple shares from different providers, you will need to have multiple reader accounts – one reader account per provider.

**Answer choice-C is correct.** A reader account will require a warehouse to execute queries on the data shared with it. As the reader account is not a licensed user of Snowflake, the warehouse needs to be provisioned by the provider account. Consequently, the provider will be charged for the credits consumed by the queries executed by the reader accounts it has created.

**Answer choice-D is incorrect.** The reader account does not have access to Snowflake support. The provider that creates and manages the reader account will be able to route questions to Snowflake support on behalf of the reader account.

► **Practical Info:** Technically, a reader account can consume an unlimited number of credits, which will be charged to the provider account. Therefore, to monitor and control the credit consumption by the reader accounts, the provider account should create resource monitors.

#### Further Reading: [Managing reader accounts](#)

## Domain : Snowflake Storage and Protection

*Q40 : Match the appropriate data protection feature given on the right with the scenario given on the left. (Assume Business-critical edition of Snowflake)*

	Scenario	Data Protection Feature
	<i>The latest release rolled out today morning has corrupted data in the production environment.</i>	
S1	<i>A data issue in the production database was discovered 3 days after time travel has expired.</i>	1. Fail-Safe
S2		2. Triple Redundancy
S3	<i>Snowflake instance has crashed, but the cloud region is available.</i>	3. Time-Travel
S4	<i>One availability zone in the cloud region where Snowflake was hosted has failed.</i>	4. Database replication
S5	<i>The entire cloud region where Snowflake was hosted has failed (for example, due to some natural calamity or network outage)</i>	5. Automatic AZ failover

- A.** S1 – 2, S2 – 3, S3 – 5, S4 – 1, S5 – 4  
**B.** S1 – 3, S2 – 2, S3 – 1, S4 – 5, S5 – 4

**C.** S1 – 4, S2 – 1, S3 – 5, S4 – 4, S5 – 3

**D.** S1 – 3, S2 – 1, S3 – 2, S4 – 5, S5 – 4

## Correct Answer: D

### Explanation

**S1 –** Latest release rolled out today morning has corrupted data in the production environment. **Answer:**

#### Time-Travel

As data got corrupted in the morning, you have time to recover this data yourself by using the Snowflake time-travel feature. Time-travel is available for every Snowflake edition. For the Standard Edition, you have 1 day (24 hours) of default time travel. For Enterprise edition and above, you get up to 90 days of time travel.

**S2 –** A data issue in the production database was discovered 3 days after time travel has expired

As it is given in this scenario that the time-travel window has expired, you can rely on the fail-safe feature of Snowflake and retrieve the older version of your data. The fail-safe period is 7 days after the expiry of time travel. **Answer: Fail-safe**

**S3 –** Snowflake instance has crashed, but the cloud region is available.

This scenario indicates that the Snowflake instance has crashed, but the underlying cloud provider/region is up. Therefore this is a Snowflake-specific problem. Snowflake provides a triple redundancy feature that will get activated in such scenarios to mitigate the impact on client workloads. This is completely transparent to the end-users. **Answer: Triple Redundancy**

**S4 –** One availability zone in the cloud region where Snowflake was hosted has failed.

Snowflake provides failover protection across three availability zones of your cloud provider region. Your data is synchronously replicated across the three availability zones. Therefore when your primary availability zone fails, the workload is automatically and transparently moved to one of the other availability zones. **Answer:**

#### Automatic AZ Fail-over

**S5 –** The entire cloud region where Snowflake was hosted has failed.

To ensure resiliency in this scenario, You will use the ‘Database replication and failover/failback’ feature available in the Business Critical edition. This feature allows you to create a secondary database in another cloud provider/region. The secondary database is automatically synced with the primary database. In case of a cloud region failure, you can promote your secondary database to serve as the primary database and continue to process data. **Answer: Database replication**

**Further Reading:** [How Snowflake protects your data and services](#)

## Domain : Performance and Tuning

*Q41 : Select the statements that apply to clustering depth in a micro-partition. (select all that apply)*

- A.** Clustering depth is the average depth of overlapping micro-partitions for a specified set of columns
- B.** The lower the clustering depth, the better clustered the table is
- C.** The higher the clustering depth, the better clustered the table is
- D.** Is an indication of whether the table needs an explicit clustering key

### **Correct Answers: A, B and D**

### **Explanation**

**Answer choice-A is correct.** Clustering depth is the average depth of overlapping micro-partitions. Clustering depth indicates how well the data is dispersed across micro-partitions.

**Answer choice-B is correct and Answer choice-C is incorrect.** The ideal clustering depth is 1 and it indicates that no micro-partitions overlap for a given set of columns. Therefore, when a search query is executed, Snowflake can directly find the right micro-partition. Any number greater than 1 indicates that when a search query is executed, Snowflake will find several micro-partitions and it will need to process these micro-partitions to prepare the resultset. For an empty table, the clustering depth is 0.

**Answer choice-D is correct.** In general, Snowflake organizes your data in micro-partitions in such a way that clustering depth is low. However, over time, as insert/updates are performed on the data, the data may no longer stay clustered optimally and clustering depth becomes large. This is specifically true for very large tables. When this happens, the search query performance becomes slower over time.

To improve the clustering of the data in micro-partitions, Snowflake allows you to designate one or more table columns/expressions as an explicit clustering key for the table. A clustering key co-locates rows in the same micro-partitions thereby improving search efficiency. Hence, high clustering depth indicates that the queries on that table may benefit by defining an explicit clustering key.

**Further Reading:** [Clustering depth illustration – Snowflake documentation](#)

## Domain : Snowflake Virtual Warehouses

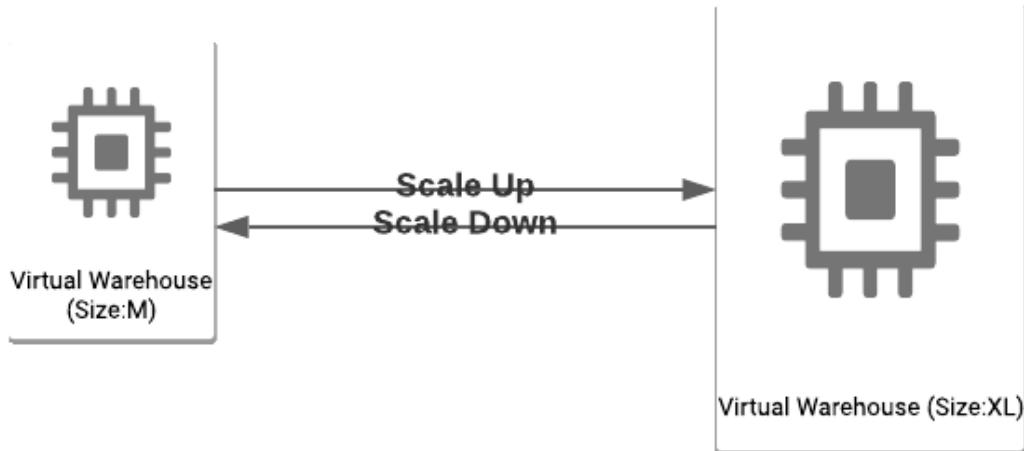
*Q42 : Select statement(s) which are TRUE from the given set of statements concerning Snowflake virtual warehouses. (select all that apply)*

- A.** Scaling up a virtual warehouse is an automatic process
- B.** Scaling up a virtual warehouse is a manual process
- C.** Scaling down a virtual warehouse is an automatic process
- D.** Scaling down a virtual warehouse is a manual process
- E.** Scaling out a virtual warehouse is an automatic process
- F.** Scaling out a virtual warehouse is a manual process
- G.** Scaling back a virtual warehouse is an automatic process
- H.** Scaling back a virtual warehouse is a manual process

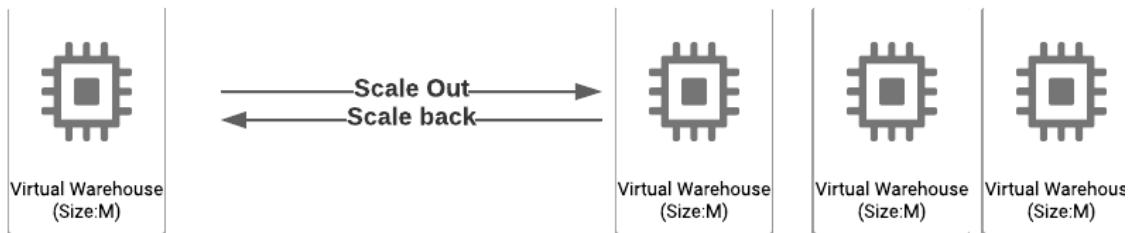
## Correct Answers: B, D, E and G

### Explanation

Scaling up or scaling down a virtual warehouse is also known as warehouse re-sizing. When you re-size a warehouse, you change the T-shirt size of the warehouse (e.g. from S to L or from XL to M). This is a manual activity. You may want to increase the T-shirt size of your warehouse (i.e. scale up) if your query performance or data loading performance is slow. You will want to decrease the T-shirt size of your warehouse (i.e. scale down) if your query workload is low.



Scaling out or scaling back a virtual warehouse is adding or removing new warehouse clusters (without regard to the size of the warehouse). This is an automatic process managed transparently by Snowflake. Snowflake will automatically increase the number of warehouse clusters as the number of queries increases to prevent queries from queuing. When Snowflake determines that additional clusters are no longer needed, it will shut them down. In summary, Snowflake ensures that a multi-cluster warehouse dynamically adapts to increase or decrease in the number of queries without any user intervention.



Summary:

Scaling model	How?	For?
Scale-Up/Down	Manual	Query/ Data Loading Performance
Scale-Out/In	Automatic	Query concurrency/

**Further Reading:** [Working with virtual warehouses in Snowflake – YouTube video](#)

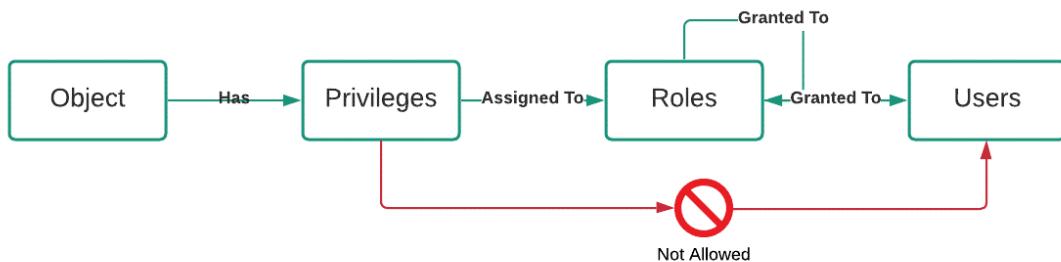
## Domain : Account and Security

*Q43 : In Snowflake, For each securable object, there is a set of privileges. The privileges provide fine-grained access control on the object. Each securable object has an owner that can assign these privileges directly to a user or a group of users.*

- A.** True
- B.** False

### Correct Answer: B

The above statement is false. This is because the owner of a securable object CANNOT grant privileges over the object directly to a user. The privileges must be assigned to roles. The roles can be granted to other roles creating a role hierarchy OR the role can be granted to a user or a group of users.



**Further Reading:** [Access Control Framework – Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q44 : Snowflake being a SaaS software, which of the following activities of an on-prem/hosted data warehouse are not required in Snowflake from the user's perspective? (select all that apply)*

- A.** Hardware sizing, purchase, or configuration
- B.** Hardware scaling
- C.** User and access management
- D.** Software installation, maintenance, or upgrades
- E.** Database Tuning
- F.** Site-level disaster recovery (due to loss of data center)

## **Correct Answers: A, B, D, E and F**

### **Explanation**

Snowflake is a true SaaS, cloud-native, data warehouse product. As it completely runs on public cloud infrastructure, it does not require end-users to purchase, configure or manage any on-prem hardware (**answer choice A**). Practically unlimited scalability available on tap means end users will not need to carry out costly and effort-intensive hardware scaling by themselves (**answer choice B**). All software maintenance, bug fixes, feature updates, etc. are rolled out completely by Snowflake without any end-user intervention (**answer choice D**). Database tuning Is no longer a task for the DBA in Snowflake as Snowflake does not support the creation of indexes, database partitions, (**Answer choice-E**). Your data is automatically replicated across three availability zones in the cloud region without any user intervention (**Answer choice-F**). Out of the given list of answer choices, only 'User and access management is a user responsibility.

**Further Reading:** [Snowflake data warehouse pros and cons – Snowflake community blog](#)

## Domain : Snowflake Storage and Protection

*Q45 : Zero-copy cloning operation is supported for the following objects in Snowflake.  
(select all that apply)*

- A.** A temporary Table
- B.** A transient Table
- C.** An external table
- D.** A permanent Table
- E.** A database created from a Share
- F.** Internal (named) Stage

## **Correct Answers: B and D**

### **Explanation**

- Zero-copy cloning is available only for the permanent and transient table types. It is not available for temporary and external table types.
- You cannot clone a database that is created from a Share.
- Also, zero-copy cloning is not available for stage objects.

## Domain : Snowflake Virtual Warehouses

*Q46 : A resource monitor can be created by...*

- A.** Owner of the virtual warehouse
- B.** ACCOUNTADMIN role
- C.** SYSADMIN role
- D.** USERADMIN role
- E.** Any of the above

**Correct Answer: B**

### **Explanation**

Resource monitors can only be created by account administrators (i.e. users with the ACCOUNTADMIN role). However, account administrators can choose to grant MONITOR and MODIFY privileges on resource monitors to the users with other roles as needed.

**Further Reading:** [Access Control Privileges for Resource Monitors – Snowflake Documentation](#)

**Domain : Data loading/unloading**

*Q47 : You have a business-critical edition of Snowflake on AWS cloud. The data lake implementation of your organization is in Azure and utilizes Azure Data Lake Gen2 service. As your Snowflake account and your data lake are on different clouds, it is not possible to do bulk loading from Azure data lake into Snowflake using the COPY command.*

- A.** True
- B.** False

**Correct Answer: B**

### **Explanation**

This statement is not correct. Bulk loading using COPY INTO supports data loading into Snowflake from the data files on your local file system or in cloud storage external to Snowflake (Amazon S3, Microsoft Azure, or Google Cloud Storage) irrespective of where your Snowflake account is hosted.

**Further Reading:** [Introduction to Data Loading – Youtube Video](#)

## Domain : Semi Structured data

*Q48 : Consider a table vehicle\_inventory that stores vehicle information of all vehicles in your dealership. The table has only one VARIANT column called vehicle\_data which stores information in JSON format. The data is given below:*

```
{  
  "date_of_arrival": "2021-04-28",  
  "supplier_name": "Hillside Honda",  
  "contact_person": {  
    "name": "Derek Larssen",  
    "phone": "8423459854"  
  },  
  "vehicle": [  
    {  
      "make": "Honda",  
      "model": "Civic",  
      "variant": "GLX",  
      "year": "2020"  
    }  
  ]  
}
```

*Which of the following are valid SQL queries that retrieve supplier\_name? (select all that apply)*

- A.** select vehicle\_data.supplier\_name::string from vehicle\_inventory
- B.** select vehicle\_data.supplier\_name::string from vehicle\_inventory
- C.** select vehicle\_data(supplier\_name(string)) from vehicle\_inventory
- D.** select vehicle\_data(supplier\_name')::string from vehicle\_inventory

**Correct Answer: B**

### Explanation

In this example, supplier\_name is the first level element in the JSON. To access this in SQL you will need to use the below syntax in the SQL if you are using . (dot) notation.

```
<column_name>.<key_name>::<cast_datatype>
```

Therefore, the correct SQL statements would be –

```
select vehicle_data:supplier_name::string from vehicle_inventory
```

**Further Reading:** [Querying Semi-Structured Data – Snowflake Documentation](#)

## Domain : Snowflake Virtual Warehouses

*Q49 : Which of the following statements are true about multi-cluster Warehouses in Snowflake? (Select all that apply).*

- A.** Multi-cluster warehouses support all properties of a single-cluster warehouse
- B.** Increasing the number of compute resources in a multi-cluster warehouse is an example of scaling up
- C.** A single cluster warehouse can be reconfigured to turn into a multi-cluster warehouse but vice-versa is not true
- D.** None of the above

**Correct Answer: A**

### Explanation

**Answer choice-A is correct.** Multi-cluster warehouses support all the same properties and actions as single warehouses, including:

- Specifying warehouse size ( e.g. XS, S, M, L...)
- Resizing a warehouse at any time.
- Auto-suspending a running warehouse due to inactivity;
- Auto-resuming a suspended warehouse when new queries are submitted.

**Answer choice-B is incorrect.** Increasing the number of compute resources in a multi-cluster warehouse is an example of scaling out. Increasing the size of compute clusters is an example of scaling up.

**Answer choice-C is incorrect.** A single cluster warehouse can be reconfigured to turn into a multi-cluster warehouse. To do this, you will have to set the max\_cluster\_count to a value greater than 1 as shown in the example below:

```
ALTER WAREHOUSE “COMPUTE_WH” SET MAX_CLUSTER_COUNT = 3;
```

Similarly, a multi-cluster warehouse can be reconfigured to turn into a single cluster warehouse. To do this, you will have to set the max\_cluster\_count to a value equal to 1 as shown in the example below:

```
ALTER WAREHOUSE “COMPUTE_WH” SET MAX_CLUSTER_COUNT = 1;
```

## Domain : Snowflake Overview and Architecture

*Q50 : Using materialized views can be beneficial in the following scenarios. (select all that apply)*

- A.** when you have a complex query that is frequently used
- B.** When the underlying data changes frequently
- C.** when the query results are smaller than the base table
- D.** When the query includes an external table

**Correct Answers: A, C and D**

### Explanation

**Answer choice-A is correct.** Encapsulating a frequently used, complex query that normally takes a long time to execute into a materialized view can improve performance. This is because the query is executed automatically in the background and the computed result set is stored in the materialized view. Querying the view, then, becomes significantly faster compared to re-running the query.

**Answer choice-B is incorrect.** A background service in Snowflake refreshes the materialized view after changes are made to the base table. This service consumes credits. If the underlying data changes frequently, the background service has to be executed repeatedly which leads to higher credit consumption. Hence, it is not recommended to create materialized views on a base table that changes frequently.

**Answer choice-C is correct.** When the query results are smaller than the base table, the background service that refreshes the materialized view needs to compute results for a small set of records. This results in lower consumption of credits to keep the view up to date.

**Answer choice-D is correct.** Querying an external table is likely to be slower than querying native database tables as the data in an external table is stored in an external stage in the cloud (outside of Snowflake). Creating a materialized view brings the resultset locally within your Snowflake account thereby improving query performance.

### Question 1

To run a Multi-Cluster Warehouse in auto-scale mode, a user would:

- **A.** Configure the Maximum Clusters setting to “Auto-Scale”
- **B.** Set the Warehouse type to “Auto”
- **C.** Set the Minimum Clusters and Maximum Clusters settings to the same value
- **D.** Set the Minimum Clusters and Maximum Clusters settings to the different values

Answer : **D**

Reference:

[https://help.pentaho.com/Documentation/9.1/Products/Modify\\_Snowflake\\_warehouse](https://help.pentaho.com/Documentation/9.1/Products/Modify_Snowflake_warehouse)

Next Question

### Question 2

Which of the following terms best describes Snowflake's database architecture?

- **A.** Columnar shared nothing
- **B.** Shared disk
- **C.** Multi-cluster, shared data
- **D.** Cloud-native shared memory

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

Next Question

### Question 3

Which of the following are options when creating a Virtual Warehouse? (Choose two.)

- **A.** Auto-drop
- **B.** Auto-resize
- **C.** Auto-resume
- **D.** Auto-suspend

Answer : **CD**

Reference:

[https://help.pentaho.com/Documentation/9.1/Products/Create\\_Snowflake\\_warehouse](https://help.pentaho.com/Documentation/9.1/Products/Create_Snowflake_warehouse)

Next Question

### Question 4

A Virtual Warehouse's auto-suspend and auto-resume settings apply to:

- **A.** The primary cluster in the Virtual Warehouse
- **B.** The entire Virtual Warehouse
- **C.** The database the Virtual Warehouse resides in
- **D.** The queries currently being run by the Virtual Warehouse

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Next Question

### Question 5

Fail-safe is unavailable on which table types? (Choose two.)

- **A.** Temporary
- **B.** Transient
- **C.** Provisional
- **D.** Permanent

Answer : **AB**

Reference:

<https://docs.snowflake.com/en/user-guide/tables-temp-transient.html>

Next Question

### Question 6

Which of the following objects is not covered by Time Travel?

- **A.** Tables
- **B.** Schemas
- **C.** Databases
- **D.** Stages

Answer : **D**

Reference:

<https://docs.snowflake.com/en/user-guide/data-time-travel.html>

Next Question

### Question 7

True or False: Micro-partition metadata enables some operations to be completed without requiring Compute.

- **A.** True
- **B.** False

Answer : **A**

Reference:

<https://blog.ippon.tech/innovative-snowflake-features-caching/>

Next Question

### Question 8

Which of the following commands are not blocking operations? (Choose two.)

- A. UPDATE
- B. INSERT
- C. MERGE
- D. COPY

Answer : **BD**

Next Question

### Question 9

Which of the following is true of Snowpipe via REST API? (Choose two.)

- A. You can only use it on Internal Stages
- B. All COPY INTO options are available during pipe creation
- C. Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands
- D. Snowpipe keeps track of which files it has loaded

Answer : **BD**

Next Question

### Question 10

Snowflake recommends, as a minimum, that all users with the following role(s) should be enrolled in Multi-Factor Authentication (MFA):

- A. SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN
- B. SECURITYADMIN, ACCOUNTADMIN, SYSADMIN
- C. SECURITYADMIN, ACCOUNTADMIN
- D. ACCOUNTADMIN

Answer : **D**

Reference:

<https://docs.snowflake.com/en/user-guide/security-mfa.html>

Next Question

### Question 11

When can a Virtual Warehouse start running queries?

- A. 12am-5am

- **B.** Only during administrator defined time slots
- **C.** When its provisioning is complete
- **D.** After replication

Answer : **B**

Next Question

### **Question 12**

True or False: Users are able to see the result sets of queries executed by other users that share their same role.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/security-access-control-considerations.html>

Next Question

### **Question 13**

True or False: The user has to specify which cluster a query will run on in a multi-cluster Warehouse.

- **A.** True
- **B.** False

Answer : **B**

Next Question

### **Question 14**

True or False: Pipes can be suspended and resumed.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro.html>

Next Question

### **Question 15**

Which of the following languages can be used to implement Snowflake User Defined Functions (UDFs)? (Choose two.)

- **A.** Java
- **B.** Javascript
- **C.** SQL
- **D.** Python

Answer : **BC**

Reference:

<https://docs.snowflake.com/en/sql-reference/user-defined-functions.html>

### **Question 1 ( Topic 1 )**

Snowflake provides a mechanism for its customers to override its natural clustering algorithms. This method is:

- **A.** Micro-partitions
- **B.** Clustering keys
- **C.** Key partitions
- **D.** Clustered partitions

Answer : **B**

Reference:

<https://community.snowflake.com/s/article/Snowflake-What-the-Cluster>

Next Question

### **Question 2 ( Topic 1 )**

Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)

- **A.** Custom
- **B.** Economy
- **C.** Optimized
- **D.** Standard

Answer : **BD**

Reference:

<https://community.snowflake.com/s/article/Snowflake-Visualizing-Warehouse-Performance>

Next Question

### **Question 3 ( Topic 1 )**

True or False: A single database can exist in more than one Snowflake account.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://docs.snowflake.com/en/sql-reference/ddl-database.html>

Next Question

### **Question 4 ( Topic 1 )**

Which of the following roles is recommended to be used to create and manage users and roles?

- **A.** SYSADMIN
- **B.** SECURITYADMIN
- **C.** PUBLIC
- **D.** ACCOUNTADMIN

Answer : **D**

Reference:

<https://docs.snowflake.com/en/user-guide/security-access-control-considerations.html>

Next Question

### **Question 5 ( Topic 1 )**

True or False: Bulk unloading of data from Snowflake supports the use of a SELECT statement.

- **A.** True
- **B.** False

Answer : **A**

### **Question 6 ( Topic 1 )**

Select the different types of Internal Stages: (Choose three.)

- **A.** Named Stage
- **B.** User Stage
- **C.** Table Stage
- **D.** Schema Stage

Answer : **ABC**

Reference:

<https://dwgeek.com/type-of-snowflake-stages-how-to-create-and-use.html/#Snowflake-Internal-Named-Stage>

Next Question

### **Question 7 ( Topic 1 )**

True or False: A customer using SnowSQL / native connectors will be unable to also use the Snowflake Web Interface (UI) unless access to the UI is explicitly granted by support.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/connecting.html>

Next Question

### **Question 8 ( Topic 1 )**

Account-level storage usage can be monitored via:

- **A.** The Snowflake Web Interface (UI) in the Databases section
- **B.** The Snowflake Web Interface (UI) in the Account -> Billing & Usage section
- **C.** The Information Schema -> ACCOUNT\_USAGE\_HISTORY View
- **D.** The Account Usage Schema -> ACCOUNT\_USAGE\_METRICS View

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/admin-usage-billing.html>

Next Question

### **Question 9 ( Topic 1 )**

Credit Consumption by the Compute Layer (Virtual Warehouses) is based on: (Choose two.)

- **A.** Number of users
- **B.** Warehouse size
- **C.** Amount of data processed
- **D.** # of clusters for the Warehouse

Answer : **BC**

Next Question

### Question 10 ( Topic 1 )

Which statement best describes clustering?

- **A.** Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions
- **B.** The database administrator must define the clustering methodology for each Snowflake table
- **C.** The clustering key must be included on the COPY command when loading data into Snowflake
- **D.** Clustering can be disabled within a Snowflake account

Answer : **A**

Reference:

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html>

### Question 11 ( Topic 1 )

True or False: The COPY command must specify a File Format in order to execute.

- **A.** True
- **B.** False

Answer : **A**

Reference:

<https://interworks.com/blog/hcalder/2018/07/26/the-basics-of-loading-data-into-snowflake/>

Next Question

### Question 12 ( Topic 1 )

Which of the following commands sets the Virtual Warehouse for a session?

- **A.** COPY WAREHOUSE FROM <<config file>>;
- **B.** SET WAREHOUSE = <<warehouse name>>;
- **C.** USE WAREHOUSE <<warehouse name>>;
- **D.** USE VIRTUAL\_WAREHOUSE <<warehouse name>>;

Answer : **C**

Reference:

<https://docs.snowflake.com/en/user-guide/warehouses-tasks.html>

Next Question

### Question 13 ( Topic 1 )

Which of the following objects can be cloned? (Choose four.)

- **A.** Tables

- **B.** Named File Formats
- **C.** Schemas
- **D.** Shares
- **E.** Databases
- **F.** Users

Answer : **ACEF**

Next Question

#### **Question 14 ( Topic 1 )**

Which object allows you to limit the number of credits consumed within a Snowflake account?

- **A.** Account Usage Tracking
- **B.** Resource Monitor
- **C.** Warehouse Limit Parameter
- **D.** Credit Consumption Tracker

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/resource-monitors.html>

Next Question

#### **Question 15 ( Topic 1 )**

Snowflake is designed for which type of workloads? (Choose two.)

- **A.** OLAP (Analytics) workloads
- **B.** OLTP (Transactional) workloads
- **C.** Concurrent workloads
- **D.** On-premise workloads

Answer : **AB**

Reference:

<https://blog.couchbase.com/its-the-workload-stupid/>

#### **Question 16 ( Topic 1 )**

What are the three layers that make up Snowflake's architecture? (Choose three.)

- **A.** Compute
- **B.** Tri-Secret Secure
- **C.** Storage
- **D.** Cloud Services

Answer : **ACD**

Reference:

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

Next Question

### **Question 17 ( Topic 1 )**

Why would a customer size a Virtual Warehouse from an X-Small to a Medium?

- **A.** To accommodate more queries
- **B.** To accommodate more users
- **C.** To accommodate fluctuations in workload
- **D.** To accommodate a more complex workload

Answer : **A**

Reference:

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html>

Next Question

### **Question 18 ( Topic 1 )**

True or False: Reader Accounts incur no additional Compute costs to the Data Provider since they are simply reading the shared data without making changes.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://interworks.com/blog/bdu/2020/02/05/zero-to-snowflake-secure-data-sharing/>

Next Question

### **Question 19 ( Topic 1 )**

Which of the following connectors allow Multi-Factor Authentication (MFA) authorization when connecting? (Choose all that apply.)

- **A.** JDBC
- **B.** SnowSQL
- **C.** Snowflake Web Interface (UI)
- **D.** ODBC
- **E.** Python

Answer : **ABCDE**

Reference:

<https://docs.snowflake.com/en/user-guide/security-mfa.html>

Next Question

### **Question 20 ( Topic 1 )**

True or False: Snowflake charges a premium for storing semi-structured data.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://snowflakecommunity.force.com/s/question/0D50Z00008ckwNuSAI/does-snowflake-charges-premium-for-storing-semi-structured-data>

### **Question 21 ( Topic 1 )**

Which of the following statements describes a benefit of Snowflake™'s separation of compute and storage? (Choose all that apply.)

- **A.** Growth of storage and compute are tightly coupled together
- **B.** Storage expands without the requirement to add more compute
- **C.** Compute can be scaled up or down without the requirement to add more storage
- **D.** Multiple compute clusters can access stored data without contention

Answer : **BCD**

Next Question

### **Question 22 ( Topic 1 )**

True or False: It is possible to unload structured data to semi-structured formats such as JSON and Parquet.

- **A.** True
- **B.** False

Answer : **A**

Reference:

<https://docs.snowflake.com/en/user-guide/data-unload-prepare.html>

Next Question

### **Question 23 ( Topic 1 )**

In which layer of its architecture does Snowflake store its metadata statistics?

- **A.** Storage Layer
- **B.** Compute Layer
- **C.** Database Layer
- **D.** Cloud Services Layer

Answer : **D**

Reference:

<https://hevodata.com/blog/snowflake-architecture-cloud-data-warehouse/>

Next Question

### **Question 24 ( Topic 1 )**

True or False: Data in fail-safe can be deleted by a user or the Snowflake team before it expires.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://blog.knoldus.com/ksnow-time-travel-and-fail-safe-in-snowflake/>

Next Question

### **Question 25 ( Topic 1 )**

True or False: Snowflake™'s data warehouse was built from the ground up for the cloud in lieu of using an existing database or a platform, like Hadoop, as a base.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

### **Question 26 ( Topic 1 )**

Which of the following statements are true of Virtual Warehouses? (Choose all that apply.)

- **A.** Customers can change the size of the Warehouse after creation
- **B.** A Warehouse can be resized while running
- **C.** A Warehouse can be configured to suspend after a period of inactivity

- **D.** A Warehouse can be configured to auto-resume when new queries are submitted

Answer : **BC**

Reference:

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html>

Next Question

### **Question 27 ( Topic 1 )**

The PUT command: (Choose two.)

- **A.** Automatically creates a File Format object
- **B.** Automatically uses the last Stage created
- **C.** Automatically compresses files using Gzip
- **D.** Automatically encrypts files

Answer : **CD**

Reference:

<https://docs.snowflake.com/en/sql-reference/sql/put.html>

Next Question

### **Question 28 ( Topic 1 )**

Which type of table corresponds to a single Snowflake session?

- **A.** Temporary
- **B.** Transient
- **C.** Provisional
- **D.** Permanent

Answer : **A**

Next Question

### **Question 29 ( Topic 1 )**

Which interfaces can be used to create and/or manage Virtual Warehouses?

- **A.** The Snowflake Web Interface (UI)
- **B.** SQL commands
- **C.** Data integration tools
- **D.** All of the above

Answer : **B**

Reference:

<https://docs.snowflake.com/en/user-guide/warehouses.html>

Next Question

### Question 30 ( Topic 1 )

When a Pipe is recreated using the CREATE OR REPLACE PIPE command:

- **A.** The Pipe load history is reset to empty
- **B.** The REFRESH parameter is set to TRUE
- **C.** Previously loaded files will be ignored
- **D.** All of the above

Answer : **B**

### Question 31 ( Topic 1 )

What is the minimum Snowflake edition that customers planning on storing protected information in Snowflake should consider for regulatory compliance?

- **A.** Standard
- **B.** Premier
- **C.** Enterprise
- **D.** Business Critical Edition

Answer : **D**

Reference:

<https://docs.snowflake.com/en/user-guide/intro-editions.html>

Next Question

### Question 32 ( Topic 1 )

Select the three types of tables that exist within Snowflake. (Choose three.)

- **A.** Temporary
- **B.** Transient
- **C.** Provisional
- **D.** Permanent

Answer : **ABD**

Reference:

<https://community.snowflake.com/s/article/Making-Transient-table-by-Default>

Next Question

### **Question 33 ( Topic 1 )**

True or False: Snowpipe via REST API can only reference External Stages as source.

- **A.** True
- **B.** False

Answer : **B**

Reference:

<https://community.snowflake.com/s/article/Making-Transient-table-by-Default>

Next Question

### **Question 34 ( Topic 1 )**

True or False: A third-party tool that supports standard JDBC or ODBC but has no Snowflake-specific driver will be unable to connect to Snowflake.

- **A.** True
- **B.** False

Answer : **A**

Reference:

<https://docs.snowflake.com/en/user-guide-connecting.html>

Next Question

### **Question 35 ( Topic 1 )**

True or False: It is possible to load data into Snowflake without creating a named File Format object.

- **A.** True
- **B.** False

Answer : **A**

Reference:

<https://docs.snowflake.com/en/user-guide/data-load-external-tutorial-create-file-format.html>

From Udemy Course:

General Usage Notes A clone is writable and is independent of its source (i.e. changes made to the source or clone are not reflected in the other object). To create a clone, your current role must have the following privilege(s) on the source object: Tables SELECT Pipes, Streams, Tasks OWNERSHIP Other objects USAGE In addition, to clone a schema or an object within a schema, your current role must have required privileges on the container object(s) for both the source and the clone. For tables, Snowflake only supports cloning permanent and transient tables; temporary tables cannot be

#### Question 9:

When cloning a table, if the COPY GRANTS keywords are not included in the CREATE <object> statement, then the new object does not inherit any explicit access privileges granted on the original table but does inherit any future grants defined for the object type in the schema.

Ans : True

#### Question 1:

If you clone a permanent table(employees) into a transient table(employees\_new) using the below command

```
create transient table employees_new clone employees copy grants;
```

What will happen to the partitions?

Old partitions will not affected, however new partitions added to the clone will follow the transient lifecycle.

Question 1: Skipped

According to snowflake architecture, data storage is independent of compute

- TRUE  
**(Correct)**
- FALSE

### Explanation

Snowflake's architecture has been defined with the belief that a tight coupling of compute and storage will not be an effective approach of leveraging the scale of cloud. The idea is to be able scale each of this component independent of each other. Snowflake's **multi-cluster, shared data** architecture separates compute resource scaling from storage resources, thus enabling seamless, non-disruptive scaling.

For data storage, Snowflake leverages the power of the respective cloud providers object store(AWS S3, Azure BLOB, GCS) and for compute it leverages the cloud VM instances which is called '**VIRTUAL WAREHOUSE**' in snowflake. Virtual warehouses are nothing but a set of compute clusters abstracted using the underlying cloud VM machines.

Question 2: Skipped

All virtual warehouses in snowflake has access to all data

- TRUE  
**(Correct)**
- FALSE

### Explanation

Snowflake was built with a unique new architecture that makes it possible to solve this challenge. Snowflake's unique **multi-cluster, shared data architecture** makes it possible to allocate multiple independent, isolated clusters for processing while sharing the same data. Each cluster (we call them "virtual warehouses") can both read and write data, with full transactional consistency ensured by Snowflake's cloud services layer. The size and resources of each virtual warehouse can be chosen independently by the user based on the characteristics and performance requirements for the workload(s) that will run on each virtual warehouse.

<https://www.snowflake.com/blog/snowflake-challenge-concurrent-load-and-query/>

Question 3: Skipped

Which layer in the snowflake architecture contains virtual warehouse

- Cloud services
- Query processing  
**(Correct)**

- Database Storage
- None

### Explanation

Virtual warehouses are used to provide compute resources for processing queries. Hence it falls under query processing

Question 4: Skipped

Which layer in snowflake contains data in compressed, columnar format

- Cloud services
- Query processing
- Database storage  
**(Correct)**
- None

### Explanation

Snowflake stores data in compressed and encrypted format by default in the database storage layer. Data is stored in micropartitions in a columnar format.

Question 5: Skipped

Which of the below approach would result in a better performance through linear scaling of data ingestion workload

- Resize virtual warehouse
- Consider the practice of organizing data by granular path
- Consider the practice of splitting input file batch within the recommended range of 10 MB and 100 MB
- All of the above  
**(Correct)**

### Explanation

Resizing a warehouse to a larger size is useful when the operations being performed by the warehouse will benefit from more compute resources, including:

1. Improving the performance of large, complex queries against large data sets.
2. Improving performance while loading and unloading significant amounts of data.

Note that, as a general practice, while performing data loads for Snowflake ETL, it is important to partition the data in your Snowflake or external locations like S3 buckets or Azure containers using logical, granular paths. By creating a partition using details like location or application along with the date when this data was written, you are optimizing for a later data loading activity. When you load the data, you can simply copy any fraction of the partitioned data into Snowflake with a single command. You can copy data into Snowflake by the hour, day, month, or even year when you initially populate tables.

The best practice is to load smaller files of 10-100 MB to maximize parallel loading i.e, number of files will be distributed between available nodes for faster processing and performance.

Due to all these, the answer to this question is "All of the above"

Question 6: Skipped

Snowflake provides specific administration features and capabilities to support the following activities except?

- Manage databases and warehouses within snowflake account
- Manage roles and users within a snowflake account
- Monitor usage and manage resources to control costs in a snowflake account
- Manage 3rd party applications providing data to a snowflake account  
**(Correct)**

Question 7: Skipped

Select all characteristics of snowflake's multi-cluster environment

- Multiple virtual warehouses in a deployment  
**(Correct)**
- Individual warehouses scale out/in based on query activity  
**(Correct)**
- Multi-cluster warehouses support the same properties and actions as single-cluster warehouse  
**(Correct)**
- User must specify which cluster each query will use

### Explanation

Multiple virtual warehouse in a deployment means specifying more than 1 warehouse while creating a multi-cluster warehouse. In a multi-cluster warehouse, you can specify the

maximum number of server cluster(or virtual warehouse) greater than 1 and upto 10 and also a minimum number of clusters(equal to or less than the maximum(up to 10)).

Additionally, multi-cluster warehouses **support all the same properties and actions as single-cluster warehouses**, including:

1. Specifying a warehouse size.
2. Resizing a warehouse at any time.
3. Auto-suspending a running warehouse due to inactivity; note that this does not apply to individual clusters, but rather the entire warehouse.
4. Auto-resuming a suspended warehouse when new queries are submitted.

In a multi-cluster warehouse, individual warehouses can scale up or down based on query load which is not possible in a single cluster warehouse. However as mentioned above, **it supports the same properties and actions like the single cluster warehouse**

**Question 8:** Skipped

What are the two modes in which a multi-cluster warehouse can run?

- Static
- Dynamic
- Maximized  
**(Correct)**
- Auto-scale  
**(Correct)**
- None of the above

### Explanation

You can choose to run a multi-cluster warehouse in either of the following modes:

#### Maximized

This mode is enabled by specifying the **same** value for both maximum and minimum clusters (note that the specified value must be larger than 1). In this mode, when the warehouse is

started, Snowflake starts all the clusters so that maximum resources are available while the warehouse is running.

This mode is effective for statically controlling the available resources (i.e. servers), particularly if you have large numbers of concurrent user sessions and/or queries and the numbers do not fluctuate significantly.

### Auto-scale

This mode is enabled by specifying *different* values for maximum and minimum clusters. In this mode, Snowflake starts and stops clusters as needed to dynamically manage the load on the warehouse:

As the number of concurrent user sessions and/or queries for the warehouse increases, and queries start to queue due to insufficient resources, Snowflake automatically starts additional clusters, up to the maximum number defined for the warehouse.

Similarly, as the load on the warehouse decreases, Snowflake automatically shuts down clusters to reduce the number of running servers and, correspondingly, the number of credits used by the warehouse.

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, SCALING\_POLICY, that determines the scaling policy to use when automatically starting or shutting down additional clusters. For more information, see [Setting the Scaling Policy for a Multi-cluster Warehouse](#) (in this topic).

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

Question 9: Skipped

What does snowflake offer to help control the credits consumed by a multi-cluster warehouse running in Auto-scale mode?

- Auto scale
- Scaling policy  
**(Correct)**
- Maximum\_number\_of\_server\_clusters
- Minimum\_number\_of\_server\_clusters
- None of the above

### Explanation

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, **SCALING\_POLICY**, that determines the scaling policy to use when automatically starting or shutting down additional clusters

Question 10: Skipped

What scaling policies are supported by snowflake?

- Standard  
**(Correct)**
- Economy  
**(Correct)**
- Premium

### Explanation

#### Standard (default)

Prevents/minimizes queuing by favoring starting additional clusters over conserving credits.

#### Economy

Conserves credits by favoring keeping running clusters fully-loaded rather than starting additional clusters, which may result in queries being queued and taking longer to complete.

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 11: Skipped

You are sizing a snowflake warehouse, what factors would you consider?

- Number of users
- Number of concurrent queries  
**(Correct)**
- Number of tables being queried  
**(Correct)**
- Data size and composition  
**(Correct)**

### Explanation

Warehouse size does not depend on number of users. Number of concurrent queries is important, because if the warehouse does not have enough resource, queries will start queueing. Same is the case with number of tables and data size.

Question 12: Skipped

Compute usage billing in snowflake is based on

- Per-second basis, with a minimum of 60 seconds  
**(Correct)**
- Per-second basis only
- By number of users

**Explanation**

Snowflake compute utilizes per second billing after the first minute. what it means is that when a warehouse starts or resumes, it bills for the first minute whether you use it or not, after that the billing is based on per second. For example if a warehouse is resumed and suspended after 30 seconds, you will still have to pay for 60 seconds because that is the first minute after the warehouse resumed.

Question 13: Skipped

SQL functionality in snowflake can be extended by

- USER DEFINED FUNCTIONS USING SQL  
**(Correct)**
- USER DEFINED FUNCTIONS USING JAVASCRIPT  
**(Correct)**
- SESSION VARIABLES  
**(Correct)**
- NONE OF THE ABOVE

**Explanation**

Snowflake supports SQL/SESSION variables declared by the user. They have many uses, such as storing application-specific environment settings.

<https://docs.snowflake.com/en/sql-reference/session-variables.html#sql-variables>

User-defined functions (UDFs) let you extend the system to perform operations that are not available through the built-in, system-defined functions provided by Snowflake. Snowflake currently supports two types of UDFs, *SQL* and *JavaScript*

Example of SQL UDF

```
create function area_of_circle(radius float)
    returns float
```

```
as  
$$  
    pi() * radius * radius  
$$  
;
```

### Example of JAVASCRIPT UDF

```
CREATE OR REPLACE FUNCTION array_sort(a array)  
RETURNS array  
LANGUAGE JAVASCRIPT  
AS  
$$  
    return A.sort();  
$$  
;  
  
-- Call the UDF with a small array.  
SELECT ARRAY_SORT(PARSE_JSON('[2,4,5,3,1]'));
```

#### Question 14: Skipped

To provide failover protection, how many availability zones does snowflake replicate to

- one
- two **(Correct)**
- three
- as configured by user

#### Explanation

Providing an even higher degree of data protection and service resilience, within the same deployment region, Snowflake provides standard failover protection **across three availability zones (including the primary active zone)**. Your data and business are protected. As you ingest your data, it is synchronously and transparently replicated across availability zones. This protection is automatically extended from Snowflake to customers, at no added charge.

<https://www.snowflake.com/blog/how-to-make-data-protection-and-high-availability-for-analytics-fast-and-easy/>

### Additional note

The answer to this question has been corrected(from 3 to 2) based on the feedback from a student. The explanation above is correct, but the replication happens to two additional availability zones only, and the primary zone is the original location of the data.

Question 15: Skipped

Which of the following is not a characteristic of micro-partitions in snowflake?

- New partitions are created in logical properties  
**(Correct)**
- Avoids skews between partitions
- Partitioning is automatically completed on the natural ingestion order of the data
- The natural ingestion order maintains correlations between columns which could be useful for pruning

### Explanation

Snowflake uses physical properties and partitions can overlap ranges

Question 16: Skipped

The storage architecture of snowflake has two key features, they are

- Time travel  
**(Correct)**
- Replication
- Zero-copy cloning  
**(Correct)**
- Query tag

### Explanation

Snowflake Time Travel enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It serves as a powerful tool for performing the following tasks:

1. Restoring data-related objects (tables, schemas, and databases) that may have been accidentally or intentionally deleted.

2. Duplicating and backing up data from key points in the past.
3. Analyzing data usage/manipulation over specified periods of time

Cloning also referred to as “zero-copy cloning” creates a copy of a database, schema or table. A **snapshot of data** present in the source object is taken when the clone is created and is made available to the cloned object. The cloned object is writable and is independent of the clone source. That is, changes made to either the source object or the clone object are not part of the other. Cloning a database will clone all the schemas and tables within that database. Cloning a schema will clone all the tables in that schema

Question 17: Skipped

Which is not a characteristic of time travel?

- Protection from accidental data operations
- Previous versions of data automatically retained
- Retention period selected by customers (up to 90 days for Enterprise edition)
- Recover data with the cost of running backups  
**(Correct)**

Question 18: Skipped

The commands to load data into snowflake are

- COPY AND PUT  
**(Correct)**
- COPY AND INSERT  
**(Correct)**
- INSERT AND PUT  
**(Correct)**
- NONE OF THE ABOVE

### Explanation

PUT in snowflake is used to load data from your local system to snowflake stage

```
put file:///tmp/data/mydata.csv @my_int_stage;
```

COPY and INSERT are the operations through which you can load data into snowflake tables

Question 19: Skipped

In order to load data into snowflake, which of the below are required?

- Virtual Warehouse  
**(Correct)**
- Predefined target table  
**(Correct)**
- Staging location with data staged  
**(Correct)**
- File format  
**(Correct)**

**Explanation**

Virtual warehouse is required to provide compute resources

Target table must be pre-defined load does not have feature to auto create tables

Data must be staged through the PUT command

File format must be defined to indicate whether CSV, JSON, XML etc is being loaded

Question 20: Skipped

Which of the following is not a snowpipe feature?

- Snowpipe can load data from any internal or external stage
- It is a server-less compute model
- The service provides REST end points and uses snowflake provided compute resources to load the data and retrieve history reports
- Snowpipe loads data after it is in stage and use executes the LOADDATA command  
**(Correct)**

**Explanation**

There is no LOADDATA command and data is automatically loaded by snowpipe after the files are added to the stage

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro.html#introduction-to-snowpipe>

Question 21: Skipped

Which feature does not come with query profile?

- 
- 
- 
- Hints for improving the query performance  
**(Correct)**
- 
- 

### Explanation

Go to query profile in snowflake web UI and see if you can find HINT:)

Question 22: Skipped

Snowflake mechanism to limit the number of micro-partitions scanned by a query is called what?

- Pruning  
**(Correct)**
- 
- 
- 
- 

### Explanation

Snowflake makes extensive use of pruning to reduce the amount of data that has to be read from storage. In summary, this means that a query like

```
SELECT SUM(x) FROM T1 WHERE y=42
```

will not read Columns x and y completely from storage but will make use of the predicate y=42 and limit the table scan to the subset of all partitions that can potentially match this condition. Snowflake can do this because it stores the min and max value of each of the columns in the metadata store

Question 23: Skipped

Which type of data integration tools leverage snowflake's scalable compute for data transformation?

- ELT  
**(Correct)**
-

- ETL
- - Database replication
  - 
  - Streaming

### Explanation

Commonly referred to as ETL, data integration encompasses the following primary operations:

#### Extract

Exporting data from specified data sources.

#### Transform

Modifying the source data (as needed), using rules, merges, lookup tables or other conversion methods, to match the target.

#### Load

Importing the resulting transformed data into a target database.

More recent usage **references the term** **ELT**, **emphasizing that the transformation part of the process** does not necessarily need to be performed before loading, particularly in systems such as Snowflake that support transformation during or after loading.

In addition, the definition of data integration has expanded to include a wider range of operations, including:

Data preparation.

Data migration/movement and management.

Data warehouse automation.

<https://docs.snowflake.com/en/user-guide/ecosystem-etl.html>

Question 24: **Skipped**

Snowflake offers tools to extract data from source systems

- TRUE
- FALSE  
**(Correct)**

Question 25: **Skipped**

Which of the following are options when creating a virtual warehouse

- Auto Suspend  
**(Correct)**
- Auto Resume  
**(Correct)**
- Storage Size
- Server count

## Explanation

### Syntax

```
CREATE [ OR REPLACE ] WAREHOUSE [ IF NOT EXISTS ] <name>
    [ [ WITH ] objectProperties ]
    [ objectParams ]
```

Where:

```
objectProperties ::=

WAREHOUSE_SIZE = XSMALL | SMALL | MEDIUM | LARGE | XLARGE | XXLARGE | XXXLARGE
GE | X4LARGE

MAX_CLUSTER_COUNT = <num>

MIN_CLUSTER_COUNT = <num>

SCALING_POLICY = STANDARD | ECONOMY

AUTO_SUSPEND = <num> | NULL

AUTO_RESUME = TRUE | FALSE

INITIALLY_SUSPENDED = TRUE | FALSE

RESOURCE_MONITOR = <monitor_name>

COMMENT = '<string_literal>'

objectParams ::=

MAX_CONCURRENCY_LEVEL = <num>

STATEMENT_QUEUED_TIMEOUT_IN_SECONDS = <num>

STATEMENT_TIMEOUT_IN_SECONDS = <num>
```

Question 26: Skipped

The warehouse cache may be reset if a running warehouse is suspended and then resumes

- TRUE  
**(Correct)**
- FALSE

### Explanation

You do not need to remember this, just understand the concept. Warehouses are nothing but VM instances underneath. And the cache is nothing but SSD cache attached to the instances, If the warehouse is suspended and resumed again it may not get the same instance.

Question 27: Skipped

A table in snowflake can only be queried using the virtual warehouse used to load the data

- TRUE
- FALSE  
**(Correct)**

### Explanation

Again apply concept, storage and compute are isolated in snowflake. That is why we can run different workloads in different warehouses but they all can refer to the same data

Question 28: Skipped

Snowflake caches are automatically invalidated if the underlying data changes

- TRUE  
**(Correct)**
- FALSE

Question 29: Skipped

The query profiler view is only available for completed queries

- TRUE
- FALSE  
**(Correct)**

Question 30: Skipped

Which best describes zero-copy cloning?

- Metadata only operation
- No replication of data
- Unmodified data stored once and modified data stored as new micro partitions
- All of the above  
**(Correct)**

### Explanation

When a clone is created of a table, the clone utilizes no data storage because it shares all the existing micro-partitions of the original table at the time it was cloned; however, rows can

then be added, deleted, or updated in the clone independently from the original table. Each change to the clone results in new micro-partitions that are owned exclusively by the clone and are protected through CDP(continuous data protection).

Question 31: Skipped

Your business team runs a set of identical queries every day after the batch ETL run is complete. From the following actions, what is the best action that you will recommend.

- After the ETL run, execute the identical queries so that they remain in the result cache  
**(Correct)**
- After the ETL run, resize the warehouse to a larger warehouse
- After the ETL run, copy the tables to another schema for the business users to query

### Explanation

Please note the key word here which is **IDENTICAL** queries. When a query is run the first time, and the same query is run the second time it picks up the results from the query result cache and it does not cost you any compute. The query result cache is valid for 24 hours and it gets extended for another 24 hours, every time you access it(even if you access it at 23 hours 59 seconds).

### Try this out yourself

Using worksheet create a table and insert the rows as below

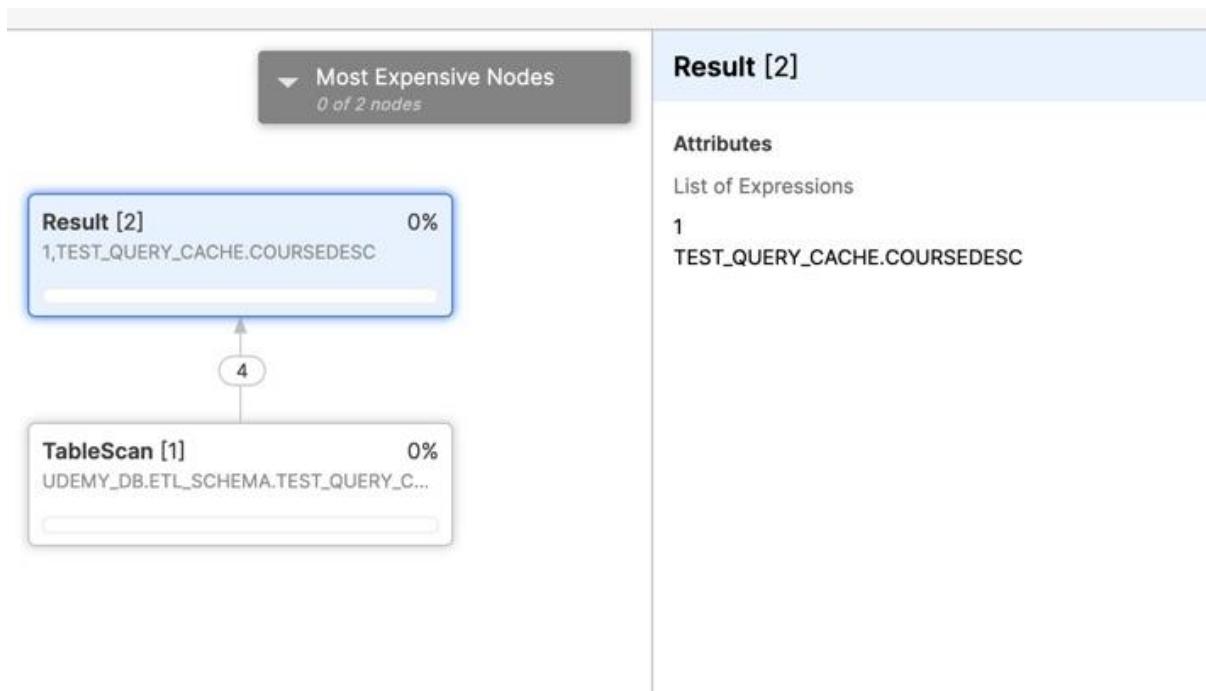
```
CREATE TABLE TEST_QUERY_CACHE(COURSEID NUMBER, COURSEDESC VARCHAR);

INSERT INTO TEST_QUERY_CACHE VALUES(1, 'SNOWFLAKE');
INSERT INTO TEST_QUERY_CACHE VALUES(1, 'C++');
INSERT INTO TEST_QUERY_CACHE VALUES(1, 'PYTHON');
```

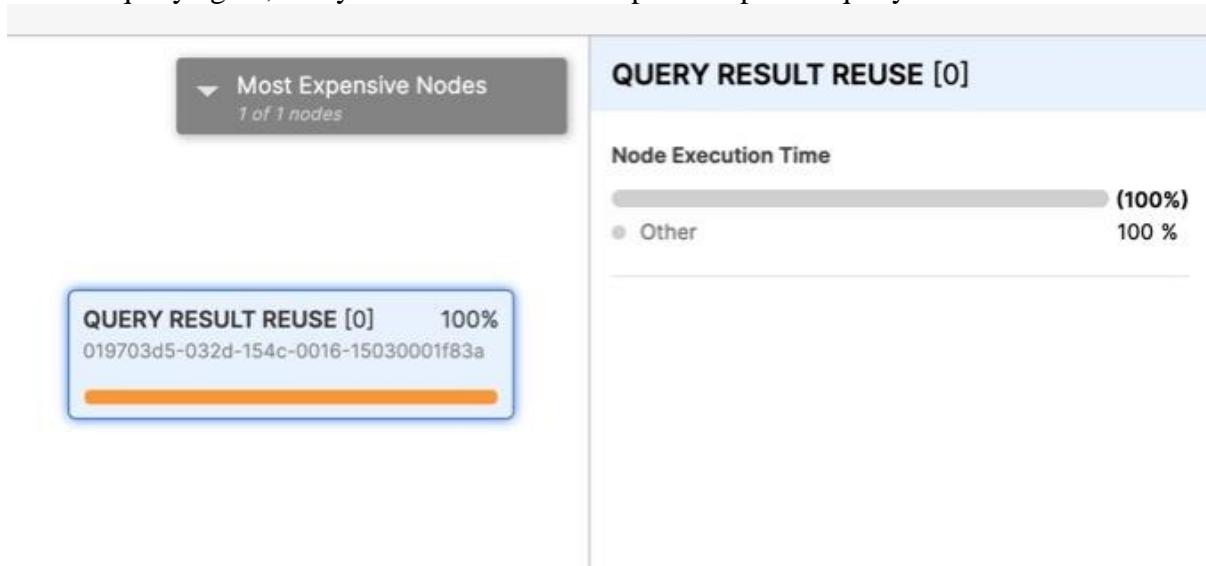
Run the below query and go to query profile

```
SELECT * FROM TEST_QUERY_CACHE;
```

You see that it is doing a table scan



Run the query again, and you will see that it has picked up from query cache



Next question:) How do you disable query cache. to know the answer look at the next question

Question 32: **Skipped**

If you would like to disable the query cache, what will you do?

- ALTER SESSION SET USE\_CACHED\_RESULT = FALSE;  
**(Correct)**
- ALTER SESSION SET USE\_CACHED\_RESULT = OFF;
-

```
ALTER SESSION SET USE_QUER_CACHE = TRUE;
```

### Explanation

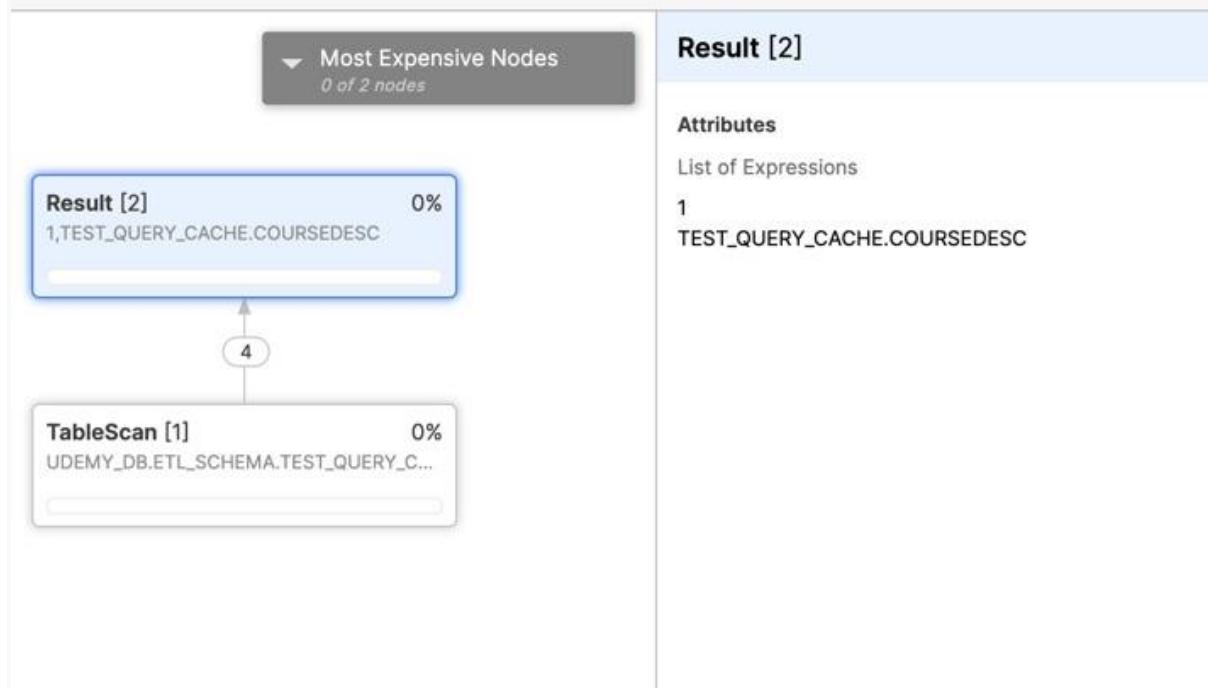
Try it again, run the below command

```
ALTER SESSION SET USE_CACHED_RESULT = FALSE;
```

Then run the query one more time, what do you see

```
SELECT * FROM TEST_QUERY_CACHE;
```

You will see the table scan coming back:)



Question 33: Skipped

A task can have a maximum of 100 *child* tasks

- TRUE  
**(Correct)**
- FALSE

### Explanation

This is something that you will need to remember

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state. An individual task in the tree is limited to a single predecessor task; however, a task can have a

maximum of 100 **child tasks** (i.e. other tasks that identify the task as a predecessor).

Question 34: Skipped

Snowflake guarantees that only one instance of a task with a defined **predecessor task** is running at a given time

- TRUE
- FALSE  
**(Correct)**

### Explanation

Snowflake cannot guarantee that only one instance of a task with a defined **predecessor task** is running at a given time.

Question 35: Skipped

You created an warehouse(ETL\_WH) which is sized at XSMALL, you want to resize the Warehouse to SMALL. How will you do that?

- 1. alter warehouse etl\_wh set warehouse\_size=SMALL;  
**(Correct)**
- You will need to recreate the warehouse with the new size
- Once created an warehouse can never be re-sized

### Explanation

Ok, to remember it for the rest of your life, please log on to SNOWFLAKE and execute

```
ALTER WAREHOUSE ETL_WH SET WAREHOUSE_SIZE=SMALL;
```

Run    All Queries | Saved 2 seconds ago

1 ALTER WAREHOUSE ETL\_WH SET WAREHOUSE\_SIZE=SMALL;

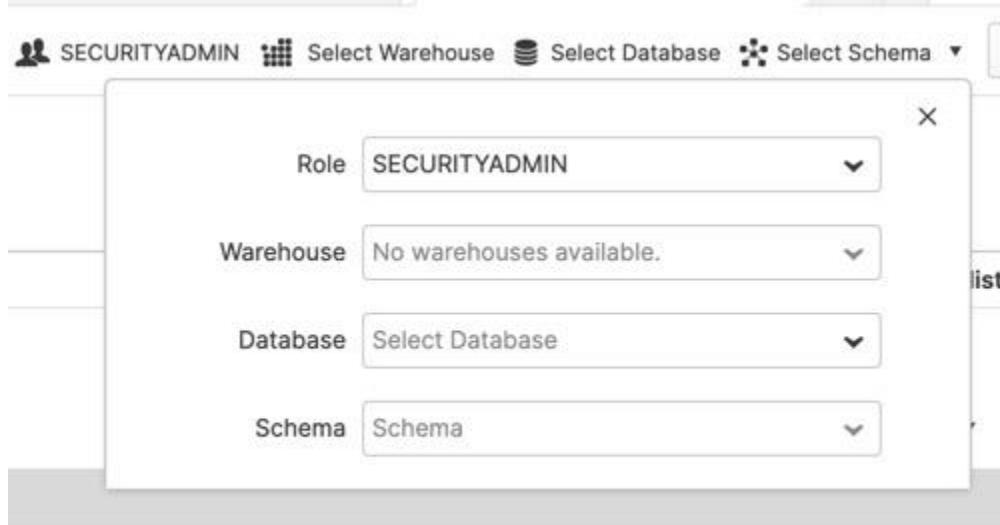
Results Data Preview

✓ Query\_ID SQL 84ms 1 rows

Filter result...

Row	status
1	Statement executed successfully.

Ok just for fun, change your role to SECURITYADMIN, and try to resize the WAREHOUSE. Will you be able to do it?



No, by default SECURITYADMIN or any other role cannot change the warehouse size. You need to be SYSADMIN or ACCOUNTADMIN. Recommended approach is to use SYSADMIN .

```
1 ALTER WAREHOUSE ETL_WH SET WAREHOUSE_SIZE=SMALL;
```

Results Data Preview

✖ Query ID SQL 20ms

SQL compilation error: Warehouse 'ETL\_WH' does not exist or not authorized.

Question 36: Skipped

You have just loaded a file named student\_data.csv from your snowflake stage to snowflake table. You try to reload again using the COPY INTO command. You will be able to load the file.

- TRUE
- FALSE  
**(Correct)**

Explanation

Let's understand why we will not be able to load the file again. Please also look at the **64 Days** mentioned below

Snowflake maintains detailed metadata for each table into which data is loaded, including:

1. Name of each file from which data was loaded
2. File size
3. ETag for the file
4. Number of rows parsed in the file
5. Timestamp of the last load for the file
6. Information about any errors encountered in the file during loading

This load metadata expires after 64 days. If the LAST\_MODIFIED date for a staged data file is less than or equal to 64 days, the COPY command can determine its load status for a given table and prevent reloading (and data duplication). The LAST\_MODIFIED date is the timestamp when the file was initially staged or when it was last modified, whichever is later.

### So, is there no way to reload it?

There is a way, you can use FORCE=TRUE

#### Reloading Files

Add **FORCE = TRUE** to a COPY command to reload (duplicate) data from a set of staged data files that have not changed (i.e. have the same checksum as when they were first loaded).

In the following example, the first command loads the specified files and the second command forces the same files to be loaded again (producing duplicate rows), even though the contents of the files have not changed:

```
copy into load1 from @%load1/data1/
  files=('test1.csv', 'test2.csv')

copy into load1 from @%load1/data1/
  files=('test1.csv', 'test2.csv')
  force=true;
```

Question 37: Skipped

For a particular snowflake session, you can specify more than one virtual warehouse

- TRUE
- FALSE  
**(Correct)**

### Explanation

A session can only use one warehouse but you can change the warehouse using USE WAREHOUSE <WAREHOUSE\_NAME>

## Warehouse Usage in Sessions

When a session is initiated in Snowflake, the session does not, by default, have a warehouse associated with it. Until a session has a warehouse associated with it, queries cannot be submitted within the session.

### Default Warehouse for Users

To facilitate querying immediately after a session is initiated, Snowflake supports specifying a default warehouse for each individual user. The default warehouse for a user is used as the warehouse for all sessions initiated by the user.

A default warehouse can be specified when creating or modifying the user, either through the web interface or using [CREATE USER/ALTER USER](#).

### Default Warehouse for Client Utilities/Drivers/Connectors

In addition to default warehouses for users, any of the Snowflake clients (SnowSQL, JDBC driver, ODBC driver, Python connector, etc.) can have a default warehouse:

SnowSQL supports both a configuration file and command line option for specifying a default warehouse.

The drivers and connectors support specifying a default warehouse as a connection parameter when initiating a session.

For more information, see [Connecting to Snowflake](#).

### Precedence for Warehouse Defaults

When a user connects to Snowflake and start a session, Snowflake determines the default warehouse for the session in the following order:

Default warehouse for the user,

» **overridden by...**

Default warehouse in the configuration file for the client utility (SnowSQL, JDBC driver, etc.) used to connect to Snowflake (if the client supports configuration files),

» **overridden by...**

Default warehouse specified on the client command line or through the driver/connector parameters passed to Snowflake.

Note

In addition, the default warehouse for a session can be changed at any time by executing the USE WAREHOUSE command within the session.

Question 38: Skipped

Time travel and fail safe requires additional storage which will cost you storage cost

- TRUE  
**(Correct)**
- FALSE

**Explanation**

Storage fees are incurred for maintaining historical data during both the Time Travel and Fail-safe periods.

**Storage Usage and Fees**

The fees are calculated for each 24-hour period (i.e. 1 day) from the time the data changed. The number of days historical data is maintained is based on the table type and the Time Travel retention period for the table.

Also, Snowflake minimizes the amount of storage required for historical data by maintaining only the information required to restore the individual table rows that were updated or deleted. As a result, storage usage is calculated as a percentage of the table that changed. Full copies of tables are only maintained when tables are dropped or truncated.

Question 39: Skipped

From which of the following stage locations can snowflake load data

- AWS S3  
**(Correct)**
- AZURE BLOB STORAGE  
**(Correct)**
- GOOGLE CLOUD STORAGE  
**(Correct)**

- INTERNAL STAGE  
**(Correct)**
- FILE SERVER
- IBM CLOUD STORAGE

### Explanation

#### External Stages

Loading data from any of the following cloud storage services is supported regardless of the [cloud platform](#) that hosts your Snowflake account:

Amazon S3  
Google Cloud Storage  
Microsoft Azure

Upload (i.e. *stage*) files to your cloud storage account using the tools provided by the cloud storage service.

A named external stage is a database object created in a schema. This object stores the URL to files in cloud storage, the settings used to access the cloud storage account, and convenience settings such as the options that describe the format of staged files. Create stages using the [CREATE STAGE](#) command.

#### Internal Stages

Snowflake maintains the following stage types in your account:

##### User

A user stage is allocated to each user for storing files. This stage type is designed to store files that are staged and managed by a single user but can be loaded into multiple tables. User stages cannot be altered or dropped.

##### Table

A table stage is available for each table created in Snowflake. This stage type is designed to store files that are staged and managed by one or more users but only loaded into a single table. Table stages cannot be altered or dropped.

Note that a table stage is not a separate database object; rather, it is an implicit stage tied to the table itself. A table stage has no grantable privileges of its own. To stage files to a table

stage, list the files, query them on the stage, or drop them, you must be the table owner (have the role with the OWNERSHIP privilege on the table).

## Named

A named internal stage is a database object created in a schema. This stage type can store files that are staged and managed by one or more users and loaded into one or more tables. Because named stages are database objects, the ability to create, modify, use, or drop them can be controlled using security access control privileges. Create stages using the [CREATE STAGE](#) command.

Upload files to any of the internal stage types from your local file system using the [PUT](#) command.

Question 40: Skipped

ACCOUNTADMIN ROLE encapsulates the SYSADMIN and SECURITYADMIN system-defined roles and ACCOUNTADMIN sits on top of the role hierarchy

- TRUE  
(Correct)
- FALSE

## Explanation

**ACCOUNTADMIN**

(aka Account Administrator)

Role that encapsulates the SYSADMIN and SECURITYADMIN system-defined roles. It is the top-level role in the system and should be granted only to a limited/controlled number of users in your account.

Question 41: Skipped

Select the two true statements about multi cluster warehouse

- These type of warehouses cannot be resumed or suspended
- Multi cluster warehouses can be created on snowflake hosted on AWS only
- As query demand decreases, clusters are removed from this warehouse  
(Correct)
- When query workload increases, clusters are automatically added  
(Correct)

## Explanation

In Auto-scale mode, a multi-cluster warehouse eliminates the need for resizing the warehouse or starting and stopping additional warehouses to handle fluctuating workloads. **Snowflake**  
**automatically starts and stops additional clusters as needed.**

In Maximized mode, you can control the capacity of the warehouse by increasing or decreasing the number of clusters as needed.

Question 42: **Skipped**

If you want to clone a table, your role must have which access on the Table

- SELECT  
**(Correct)**
- USAGE
- READ
- CLONE

### Explanation

To create a clone, your current role must have the following privilege(s) on the source object:

**Tables - SELECT**

**Pipes, Streams, Tasks - OWNERSHIP**

**Other objects - USAGE**

In addition, to clone a schema or an object within a schema, your current role must have required privileges on the container object(s) for both the source and the clone.

Question 43: **Skipped**

You have three worksheets open in Snowflake Web UI. You can have three different roles in these worksheets.

- TRUE  
**(Correct)**
- FALSE

### Explanation

Please do not believe me:) Try it out

Log on to snowflake, open three worksheets and try to use a different role. Run the below in worksheet and see if it works

Worksheet 1

```
USE ROLE SYSADMIN
```

## Worksheet 2

```
USE ROLE ACCOUNTADMIN
```

## Worksheet 3

```
USE ROLE PUBLIC
```

Question 44: Skipped

When will you consider creating a clustering key on a table. Please select all that applies

- The table houses multi-terabyte data  
**(Correct)**
- The table has variant columns
- The query performance on the table has degraded over time  
**(Correct)**

### Explanation

#### No need to mug this up

1. Clustering will be effective only when you have enough partitions in the table and the table is more than 1 TB

2. As you insert/update rows in the table, the partitions are moved around and hence it no longer remains clustered. With Auto clustering, Snowflake automatically reclusters the table based on a proprietary algorithm

Question 45: Skipped

Which of the below objects will contribute to your account level storage usage?

- Secure Views
- Internal Stage  
**(Correct)**
- External stage  
**(Correct)**
- Database and all tables inside it  
**(Correct)**
-

## Functions

### Explanation

SECURE VIEWS and FUNCTION definitions are stored in metadata store. They do not store any data physically.

Update to the explanation

---

One of my students asked a valid question that external stage is a pointer to the cloud provider's storage and hence from a snowflake perspective there is no cost. This is absolutely true. But when this question is presented to you and it is not explicitly asking 'snowflake account', you should tick external stage also. If the question specifically asks about 'snowflake account', then you should not tick this option. This is kind of a trick question so please pay attention to the question text.

### More on external stage

#### External stage

References data files stored in **a location outside of Snowflake**. Currently, the following cloud storage services are supported:

1. Amazon S3 buckets
2. Google Cloud Storage buckets
3. Microsoft Azure containers

The storage location can be either private/protected or public.

Question 46: Skipped

A consumer of shared database can add as many tables and views they want in the shared database

- FALSE  
**(Correct)**
- TRUE

## Explanation

Remember, we discussed earlier that SHARES are read-only and cannot be changed by consumer

Question 47: Skipped

To have multi cluster data virtual warehouse, what should be the minimum snowflake subscription

- ENTERPRISE  
**(Correct)**
- BUSINESS CRITICAL
- STANDARD

Question 48: Skipped

Once a new release has been deployed, Snowflake does not move all accounts to the release at the same time. It follows a staged release approach.

- TRUE  
**(Correct)**
- FALSE

## Explanation

### Staged Release

Once a new release has been deployed, Snowflake does not move all accounts to the release at the same time. Accounts are moved to the release using a three stage approach over two days. Accounts are moved to the new release in the following order, based on their [Snowflake Edition](#):

#### Day 1

Stage 1 (*early access*) for designated Enterprise accounts.

#### Day 1 or 2

Stage 2 (*regular access*) for all Standard Edition accounts.

#### Day 2

Stage 3 (*final*) for all Enterprise Edition and VPS accounts.

The minimum amount of elapsed time between the early access and final stages is 24 hours. This staged approach enables Snowflake to monitor activity as accounts are moved and respond to any issues that may occur.

Question 49: Skipped

When you are loading data through the COPY command, the table and the file being loaded must have the same order of columns

- FALSE  
**(Correct)**
- TRUE

### Explanation

Note that the actual field/column order in the data files can be different from the column order in the target table. It is only important that the SELECT list maps fields/columns in the data files to the **corresponding** columns in the table.

Question 50: Skipped

A reader account user will not be able to do the following

- LOAD DATA  
**(Correct)**
- INSERT  
**(Correct)**
- SELECT
- UPDATE  
**(Correct)**

### Explanation

A reader account is intended primarily for querying data shared by the provider of the account. Adding new data to the account and/or updating shared data in the account is not supported. As such, the following DML and DDL commands are not allowed:

INSERT

UPDATE

DELETE

MERGE

COPY INTO <table>

CREATE PIPE

CREATE SHARE

All other operations are allowed.

Question 51: Skipped

Which two modes are available in Multi cluster warehouse

- Auto-scale  
**(Correct)**
- Peak Scale
- Flexible
- Maximized  
**(Correct)**

### Explanation

Maximized vs. Auto-scale

You can choose to run a multi-cluster warehouse in either of the following modes:

#### Maximized

This mode is enabled by specifying the *same* value for both maximum and minimum clusters (note that the specified value must be larger than 1). In this mode, when the warehouse is started, Snowflake starts all the clusters so that maximum resources are available while the warehouse is running.

This mode is effective for statically controlling the available resources (i.e. servers), particularly if you have large numbers of concurrent user sessions and/or queries and the numbers do not fluctuate significantly.

#### Auto-scale

This mode is enabled by specifying *different* values for maximum and minimum clusters. In this mode, Snowflake starts and stops clusters as needed to dynamically manage the load on the warehouse:

As the number of concurrent user sessions and/or queries for the warehouse increases, and queries start to queue due to insufficient resources, Snowflake automatically starts additional clusters, up to the maximum number defined for the warehouse.

Similarly, as the load on the warehouse decreases, Snowflake automatically shuts down clusters to reduce the number of running servers and, correspondingly, the number of credits used by the warehouse.

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, SCALING\_POLICY, that determines the scaling policy to use when automatically starting or shutting down additional clusters. For more information, see [Setting the Scaling Policy for a Multi-cluster Warehouse](#) (in this topic).

Question 52: Skipped

Which of the below operations are allowed on an inbound share data?

- MERGE
- CREATE TABLE
- ALTER TABLE
- SELECT WITH GROUP BY  
**(Correct)**
- SELECT WITH JOIN  
**(Correct)**
- INSERT INTO

### Explanation

This is a trick question:) remember a share is read only, so you can only select data from a share

### Important

All database objects shared between accounts are ***read-only*** (i.e. the objects cannot be modified or deleted, including adding or modifying table data).

Question 53: Skipped

Select the most accurate statement regarding Snowflake's transaction support

- ACID Compliant  
**(Correct)**
- Requires special user configuration
- BASE compliant
- Works like a NoSQL database

Question 54: Skipped

Select the table type that is automatically deleted after the session is closed and hence it has no fail-safe or time travel option post closure of the session

- TRANSIENT

- TEMPORARY  
**(Correct)**

- PERMANENT
- EXTERNAL

### Explanation

Please create a temporary table and try it out

```
CREATE TEMPORARY TABLE T1(COL1 VARCHAR);
INSERT INTO T1 VALUES('TESTING TIME TRAVEL FOR T1');
```

Close the worksheet and open another worksheet and run the below query, are you able to time travel?

```
SELECT * FROM T1 AT(OFFSET => -60*1);
```

Question 55: Skipped

What does snowflake recommend for making a variant column accessible in a BI Tool?

- Create a View  
**(Correct)**
- Create a separate table
- Use Parse\_Json
- Convert into external table

Question 56: Skipped

When you load data into snowflake, snowflake reorganizes the data into its internal optimized compressed columnar format

- TRUE  
**(Correct)**
- FALSE

## Explanation

When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed, columnar format. Snowflake stores this optimized data in cloud storage.

Snowflake manages all aspects of how this data is stored — the organization, file size, structure, compression, metadata, statistics, and other aspects of data storage are handled by Snowflake. The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake.

Question 57: Skipped

When you clone a table, it does not include the load history of the source table, and hence the data files that were already loaded in source table can be reloaded in the clone table

- TRUE  
(Correct)
- FALSE

Question 58: Skipped

A snowflake account hosted on AWS cannot load data from a file that is staged in GCP or Azure

- FALSE  
(Correct)
- TRUE

Question 59: Skipped

If you clone a database, the internal snowflake stages in that database are also cloned

- FALSE  
(Correct)
- TRUE

## Explanation

### Cloning and Stages

The following rules apply to cloning stages or objects that contain stages (i.e. databases and schemas):

Individual external named stages can be cloned; internal named stages cannot be cloned.

When cloning a database or schema:

External named stages that were present in the source when the cloning operation started are cloned.

Tables are cloned, which means their internal stages are also cloned.

Internal named stages are **not** cloned.

Regardless of how a stage was cloned, the clone does **not** include any of the files from the source. i.e. all cloned stages are empty.

Question 60: Skipped

Which command is used to download data from Snowflake stage

- PUT
- GET  
**(Correct)**
- COPY
- DUMP

### Explanation

**GET**

Downloads data files from one of the following Snowflake stages to a local directory/folder on a client machine:

Named internal stage.

Internal stage for a specified table.

Internal stage for the current user.

Typically, this command is executed after using the `COPY INTO <location>` command to unload data from a table into a Snowflake stage.

Question 61: Skipped

AWS private link provides a secure connection from the customer's on premise data center to Snowflake

- True
- False  
**(Correct)**

### Explanation

<https://docs.snowflake.com/en/user-guide/admin-security-privatelink.html#what-is-aws-privatelink>

[AWS PrivateLink](#) is an AWS service for creating private VPC endpoints that allow direct, secure connectivity between AWS VPCs without traversing the public Internet.

In addition, if you have an on-premises environment (e.g. a non-hosted data center), you can choose to use [AWS Direct Connect](#), in conjunction with AWS PrivateLink, to connect all your virtual and physical environments in a single, private network.

Question 62: [Skipped](#)

Which of the below is an example of "**Increasing max number of clusters**"?

- Multi-Cluster Warehouse  
**(Correct)**
- Virtual Warehouse
- Dedicated Warehouse

### Explanation

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#what-is-a-multi-cluster-warehouse>

With multi-cluster warehouses, Snowflake supports allocating, either statically or dynamically, additional warehouses to make a larger pool of compute resources available. A multi-cluster warehouse is defined by specifying the following properties:

Maximum number of warehouses, greater than 1 (up to 10).

Minimum number of warehouses, equal to or less than the maximum (up to 10).

Question 63: [Skipped](#)

Reader accounts can extract data from shared data objects to use outside Snowflake?

- True
- False  
**(Correct)**

### Explanation

Reader accounts (formerly known as “read-only accounts”) provide a quick, easy, and cost-effective way to share data without requiring the consumer to become a Snowflake customer. Users in a reader account can query data that has been shared with them, but cannot perform any of the DML tasks that are allowed in a full account (data loading, insert, update, etc.)

Question 64: [Skipped](#)

Zero copy cloning requires additional licensing since it is considered as an additional feature in snowflake

-

True

- 

False

(Correct)

Question 65: Skipped

Some queries are answered by metadata cache and do not require active warehouse

- 

True

(Correct)

- 

False

### Explanation

Metadata caching is **maintained in Global Service Layer**. Metadata about Tables and micro-partitions are collected and managed by snowflake automatically. Snowflake does not use compute(and hence no warehouse) to provide Range values like MIN,MAX, Number of distinct values ,NULL count and ROW count and clustering information

Question 66: Skipped

For how many days Query history in Snowflake web Ui is kept?

- 

14 days

(Correct)

- 

30 days

- 

365 days

### Explanation

The **History** page allows you to view and drill into the details of all queries executed in the last 14 days

<https://docs.snowflake.com/en/user-guide/ui-history.html#using-the-history-page-to-monitor-queries>

Question 67: Skipped

Storage cost is incurred for both internal and external stage.

- 

True

(Correct)

- 

False

Question 68: Skipped

Fail-safe period can be disabled at account level

-

True

- 

False

(Correct)

### Explanation

Fail-safe provides a (non-configurable) 7-day period during which historical data may be recoverable by Snowflake. This period starts immediately after the Time Travel retention period ends. Note, however, that a long-running Time Travel query will delay moving any data and objects (tables, schemas, and databases) in the account into Fail-safe, until the query completes.

Question 69: Skipped

Reader account does not cost storage cost

- 

TRUE

(Correct)

- 

FALSE

### Explanation

Reader account does not COPY the data, in fact reader account does not have to be snowflake users also. A reader account can only consume data from the provider account that created it. A reader account is intended primarily for querying data shared by the provider of the account. Adding new data to the account and/or updating shared data in the account is not supported

Question 70: Skipped

Data Unloading compression applies to JSON and CSV only

- 

True

- 

False

(Correct)

### Explanation

It applies to parquet output also

Question 71: Skipped

A stage be created without specifying file format

- 

TRUE

(Correct)

- 

FALSE

### Explanation

If you do not specify the file format, it will create with CSV as default. Try this out to check it

```
CREATE STAGE TEST_STAGE;  
DESCRIBE STAGE TEST_STAGE;
```

Question 72: Skipped

Account storage consumption can be checked via

- Snowflake UI  
**(Correct)**
- ACCOUNT USAGE HISTORY
- ACCOUNT USAGE METRICS

### Explanation

There is nothing like ACCOUNT USAGE METRICS and ACCOUNT USAGE HISTORY

Question 73: Skipped

Micro-partition metadata is managed by which service layer

- Global Service Layer  
**(Correct)**
- Compute Layer
- Storage Layer
- Metadata layer

### Explanation

Metadata caching is maintained in Global Service Layer

Question 74: Skipped

4XL warehouse may have more lag time for provisioning compare to small warehouse

- True  
**(Correct)**
- False

Question 75: Skipped

A third party tool that supports JDBC or ODBC but has no snowflake specific driver will be unable to connect with snowflake

- True  
**(Correct)**

- False

Question 76: **Skipped**

Query IDs are unique across all snowflake deployments and can be used in communication with snowflake support to help troubleshoot issues

- True  
**(Correct)**

- False

Question 77: **Skipped**

Snowflake is designed for

- Concurrent workloads  
**(Correct)**
- OLAP  
**(Correct)**
- OLTP
- ON PREMISE

Question 78: **Skipped**

Which of the below interfaces can be used to create and/or manage virtual warehouses?

- Snowflake UI  
**(Correct)**
- SQL commands  
**(Correct)**
- Data ingestion tools
- All of the above

Question 79: **Skipped**

Credit consumption by the compute layer is based on

- Number of user

- Amount of data processed  
**(Correct)**

- Number of clusters

- Size of warehouse  
**(Correct)**

Question 80: Skipped

Which of the below are benefits of Snowflake's separation of compute and storage

- Multiple warehouses can access stored data without contention  
**(Correct)**

- Storage expands without requirement to add more compute  
**(Correct)**

- Compute can be scaled up and down without adding more storage  
**(Correct)**

- Growth of storage and compute are tightly coupled together

Question 81: Skipped

What happens when pipe is recreated using CREATE OR REPLACE pipe command

- Pipe load history is reset to empty  
**(Correct)**
- Refresh parameter set to true
- Previous loaded files will be ignored

### Explanation

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-manage.html#load-history>

[Continue](#) [Retake test](#)

Question 1: Skipped

With respect to snowflake security, Tri-secret requires that customers manage their keys

- TRUE  
**(Correct)**
- FALSE

### Explanation

Tri-Secret Secure lets you control access to your data using a master encryption key that you maintain in the key management service for the cloud provider that hosts your Snowflake account:

**AWS:** [AWS Key Management Service \(KMS\)](#)

**Google Cloud:** [Cloud Key Management Service \(Cloud KMS\)](#)

**Microsoft Azure:** [Azure Key Vault](#)

Please note that to use tri-secret, you will need business critical edition or more

<https://docs.snowflake.com/en/user-guide/security-encryption.html#tri-secret-secure-and-customer-managed-keys>

Question 2: Skipped

The security information of snowflake are store in which layer

- Storage
- Compute
- Service  
**(Correct)**

### Explanation

Remember, what the cloud services layer is for? Authentication and access control.

The cloud services layer is a collection of services that coordinate activities across Snowflake. These services tie together all of the different components of Snowflake in order to process user requests, from login to query dispatch. The cloud services layer also runs on compute instances provisioned by Snowflake from the cloud provider.

Among the services in this layer:

Authentication

Infrastructure management

Metadata management

Query parsing and optimization

Access control

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html#cloud-services>

Question 3: Skipped

Snowflake supports multi-factor authentication (MFA) to provide increased login security for users connecting to Snowflake. Which statements are true about MFA security?

- MFA is not enabled by default. it is available to all, but user need to activate it  
**(Correct)**
- MFA is an integrated feature powered by DUO Security Service  
**(Correct)**
- MFA can be used to connect to SNOWFLAKE using Snowflake JDBC Driver  
**(Correct)**
- MFA Login is designed to connect to snowflake only through the web interface

### Explanation

Snowflake supports multi-factor authentication (MFA) to provide increased login security for users connecting to Snowflake. MFA support is provided as an integrated Snowflake feature, powered by the [Duo Security](#) service, which is managed completely by Snowflake. Users do not need to separately sign up with Duo or perform any tasks, other than installing the Duo Mobile application, which is supported on multiple smart phone platforms (iOS, Android, Windows, etc.). See the [Duo User Guide](#) for more information about supported platforms/devices and how Duo multi-factor authentication works.

MFA is enabled on a per-user basis; however, at this time, users are not automatically enrolled in MFA. To use MFA, users must enroll themselves.

**Previously, users could only be enrolled in MFA by submitting a request to Snowflake Support. This is no longer required.**

<https://docs.snowflake.com/en/user-guide/security-mfa.html#multi-factor-authentication-mfa>

Question 4: Skipped

Which of the below activities can be done leveraging time-travel feature of snowflake?

- Restoring data-related objects (tables, schemas, and databases) that may have been accidentally or intentionally deleted  
**(Correct)**
- Duplicating and backing up data from key points in the past  
**(Correct)**
- Analyzing data usage/manipulation over specified periods of time  
**(Correct)**
- To backup data for compliance and regulatory requirements

### Explanation

Time travel is a beautiful concept in snowflake and you must understand this feature in details. A key component of Snowflake Time Travel is the data retention period. The standard retention period is 1 day (24 hours) and is automatically enabled for all Snowflake accounts

For **Snowflake Standard Edition**, the **retention period can be set to 0** (or unset back to the default of 1 day) at the account and object level (i.e. databases, schemas, and tables).

### For Snowflake Enterprise Edition (and higher):

1. For **transient databases, schemas, and tables, the retention period can be set to 0** (or unset back to the default of 1 day). The same is also true for temporary tables.
2. For permanent databases, schemas, and tables, the retention period **can be set to any value from 0 up to 90 days**.

You can use this feature to restore objects if you have accidentally deleted one, you can duplicate or back up data from a certain point in the past. You can also analyze data usage over a period of time

<https://docs.snowflake.com/en/user-guide/data-time-travel.html#understanding-using-time-travel>

Question 5: Skipped

One of your users have accidentally dropped a table T1 in production. What can you do to restore the table?

- Execute **UNDROP TABLE T1;**  
**(Correct)**
- Execute **UNDELETE TABLE T1;**
- Call Snowflake team and request them to recover the data from cloud storage
- Nothing can be done, whatever is lost is lost

**Explanation**

UNDROP is a very useful feature in Snowflake. It makes it easier to retrieve your data and that is because snowflake stores data and objects in encrypted form for 24 hours by default. You have an option to purchase even longer retention period, if required.

Any lost data can be instantaneously recovered with the UNDROP features. It is a single short command

***UNDROP TABLE <tablename>***

In addition to this, you can also do **UNDROP SCHEMA** and **UNDROP DATABASE** if someone messes it up even bigger.

<https://docs.snowflake.com/en/sql-reference/sql/undrop-table.html#undrop-table>

Question 6: Skipped

In your organization, you use a 2XL warehouse to load data into Snowflake as part of your ETL workload. Your BI team reached out to you and wanted access to the data stored in Snowflake. You gave them access to the tables required for them to query the data. What will be your recommendation with respect to the warehouse that they will need to use?

- Since you already have a 2XL warehouse for loading data, you will let them use that warehouse
- Warehouse is not required for querying data in snowflake
-

It is recommended to use a multi-warehouse approach for workload isolation in snowflake. Since you would not like the load workload to interfere with the read workload of the BI team, you will create a separate warehouse for them

(Correct)

### Explanation

Snowflake's unique [multi-cluster, shared data architecture](#) makes it possible to allocate multiple independent, isolated clusters for processing while sharing the same data. In snowflake, each cluster can both read and write data complying to transactional consistency which is ensured by the cloud services layer. The compute clusters in Snowflake are called virtual warehouses. The size and resources of each virtual warehouse can be chosen independently by the user based on the characteristics and performance requirements for the workload(s) that will run on each virtual warehouse. It is, therefore, recommended to isolate workloads and have each workload use its own warehouse. It helps to track usage better and also through this means you will ensure that one workload does not interfere with the other workload. Hence you will create a separate workload for the BI TEAM

Question 7: Skipped

When you clone a table in snowflake, the storage requirements are doubled as it needs to now store two copies of data.

- TRUE
- FALSE

(Correct)

### Explanation

A massive benefit of zero-copy cloning of snowflake is that the underlying data is never duplicated. Only the metadata/pointers of the micropartitions are replicated for the new cloned table. Hence, when you clone a table actual data is never copied over to the new cloned table.

Question 8: Skipped

To ingest semi structured data in snowflake, which data type will you use?

- NUMBER
- VARCHAR
- VARIANT

(Correct)

### Explanation

VARIANT data type allows Snowflake to ingest semi-structured data without having to pre-define the schema.

<https://docs.snowflake.com/en/sql-reference/data-types-semistructured.html#variant>

Question 9: Skipped

Select the ones that are true about data sharing through data shares

- There is only one copy of data, which lives in the data provider's account  
**(Correct)**
- Shared data is always live, real-time and immediately available to consumers  
**(Correct)**
- Providers can establish revocable, fine-grained access grants to shares  
**(Correct)**
- Data Sharing is supported between accounts in the same or different Snowflake Provider and Region

### Explanation

Data sharing is currently supported between accounts in the same Snowflake Provider and Region. If the accounts are in different region, please consider data replication

Question 10: Skipped

Which factors influence the unit cost of snowflake credits and data storage?

- Snowflake Edition  
**(Correct)**
- The region where your Snowflake account is located  
**(Correct)**
- Whether it is an *On Demand* or *Capacity* account  
**(Correct)**
- Number of users of snowflake

### Explanation

The Snowflake Edition that your organization chooses determines the unit costs for the credits and the data storage you use. Other factors that impact unit costs are the *region* where your Snowflake account is located and whether it is an *On Demand* or *Capacity* account:

1. On Demand: Usage-based pricing with no long-term licensing requirements.
2. Capacity: Discounted pricing based on an up-front Capacity commitment.

Question 11: Skipped

What was Business Critical Edition formerly known as

- Standard Edition
- Enterprise Edition

- Business Edition
- Enterprise for Sensitive Data (ESD)  
**(Correct)**

### Explanation

Business Critical Edition, formerly known as Enterprise for Sensitive Data (ESD), offers even higher levels of data protection to support the needs of organizations with extremely sensitive data, particularly PHI data that must comply with HIPAA regulations.

It includes all the features and services of [Enterprise Edition](#), with the addition of enhanced security and data protection. In addition, database failover/failback adds support for business continuity and disaster recovery.

Question 12: [Skipped](#)

Before any PHI data can be stored in Snowflake, a signed business associate agreement (BAA) must be in place between your agency/organization and Snowflake Inc

- TRUE  
**(Correct)**
- FALSE

### Explanation

As required by HIPAA regulations, before any PHI data can be stored in Snowflake, a signed business associate agreement (BAA) must be in place between your agency/organization and Snowflake Inc.

Question 13: [Skipped](#)

Your organization requires a snowflake instance that is a completely separate Snowflake environment, isolated from all other Snowflake accounts. Which edition will you suggest?

- Standard
- Premier
- Virtual Private Snowflake (VPS)  
**(Correct)**

### Explanation

Virtual Private Snowflake offers our highest level of security for organizations that have the strictest requirements, such as financial institutions and any other large enterprises that collect, analyze, and share highly sensitive data.

It includes all the features and services of [Business Critical Edition](#), but in a completely separate Snowflake environment, isolated from all other Snowflake accounts (i.e. VPS accounts do not share any resources with accounts outside the VPS).

<https://docs.snowflake.com/en/user-guide/intro-editions.html#virtual-private-snowflake-vps>

Question 14: Skipped

Extended Time Travel (up to 90 days) is available in which editions

- Standard
- Enterprise  
**(Correct)**
- Business Critical  
**(Correct)**
- VPS  
**(Correct)**

### Explanation

Please go through this link. there may be questions on feature availability

<https://docs.snowflake.com/en/user-guide/intro-editions.html#security-data-protections>

Question 15: Skipped

Your organization has a need for Column-level security by applying masking policies to columns in tables or views. Which editions will be suitable for this requirement?

- Standard
- Enterprise  
**(Correct)**
- Business Critical  
**(Correct)**
- VPS  
**(Correct)**

### Explanation

Dynamic masking is currently a preview feature and is available to all accounts(as of Aug 2020), it requires enterprise edition or above.

Question 16: Skipped

Which of the below features are only available to Business Critical edition and above?

- Customer-managed encryption keys through Tri-Secret Secure.  
**(Correct)**

- Support for secure, direct proxy to your other virtual networks or on-premises data centers using AWS PrivateLink or Azure Private Link.  
**(Correct)**
- Support for PHI data (in accordance with HIPAA regulations).  
**(Correct)**
- Support for PCI DSS.  
**(Correct)**
- Extended Time Travel (up to 90 days).

### Explanation

Extended Time Travel (up to 90 days) is available in Enterprise and above editions

Question 17: Skipped

Which feature is only available in VPS edition?

- Dedicated metadata store and pool of virtual servers (used in virtual warehouses).  
**(Correct)**
- SOC 2 Type II certification.
- Automatic encryption of all data.

Question 18: Skipped

In load metadata, Snowflake maintains detailed metadata for each table into which data is loaded. What else is stored in this metadata

- Name of each file from which data was loaded  
**(Correct)**
- File size  
**(Correct)**
- Number of columns in the file
- ETag for the file  
**(Correct)**
- Number of rows parsed in the file  
**(Correct)**
-

Timestamp of the last load for the file  
**(Correct)**

- Information about any errors encountered in the file during loading  
**(Correct)**

### Explanation

Snowflake maintains detailed metadata for each table into which data is loaded, including:

1. Name of each file from which data was loaded
2. File size
3. ETag for the file
4. Number of rows parsed in the file
5. Timestamp of the last load for the file
6. Information about any errors encountered in the file during loading

This load metadata expires after 64 days. If the LAST\_MODIFIED date for a staged data file is less than or equal to 64 days, the COPY command can determine its load status for a given table and prevent reloading (and data duplication). The LAST\_MODIFIED date is the timestamp when the file was initially staged or when it was last modified, whichever is later.

<https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html#load-metadata>

Question 19: **Skipped**

After how many days does the load metadata expire

- 30
- 60
- 64  
**(Correct)**

### Explanation

This load metadata expires after 64 days. If the LAST\_MODIFIED date for a staged data file is less than or equal to 64 days, the COPY command can determine its load status for a given table and prevent reloading (and data duplication). The LAST\_MODIFIED date is the timestamp when the file was initially staged or when it was last modified, whichever is later.

Question 20: **Skipped**

Which OPTION will you set on the copy if you would like to ignore the load metadata?

- FORCE  
**(Correct)**

- LOAD\_UNCERTAIN\_FILES
- IGNORE\_METADATA

### Explanation

To load files whose metadata has expired, set the LOAD\_UNCERTAIN\_FILES copy option to true. The copy option references load metadata, if available, to avoid data duplication, but also attempts to load files with expired load metadata.

Alternatively, set the FORCE option to load all files, ignoring load metadata if it exists. Note that this option reloads files, potentially duplicating data in a table.

<https://docs.snowflake.com/en/user-guide/data-load-considerations-load.html#workarounds>

Question 21: Skipped

Standard Time Travel (up to 1 day) for accessing/restoring modified and deleted data is available in all editions

- FALSE
- TRUE  
(Correct)

Question 22: Skipped

Each object has an owner, who can in turn grant access to that object. What is this model called?

- Discretionary Access Control (DAC)  
(Correct)
- Role-based Access Control (RBAC)
- Object ownership model

### Explanation

Snowflake's approach to access control combines aspects from both of the following models:

- 1. Discretionary Access Control (DAC):** Each object has an owner, who can in turn grant access to that object.
- 2. Role-based Access Control (RBAC):** Access privileges are assigned to roles, which are in turn assigned to users.

<https://docs.snowflake.com/en/user-guide/security-access-control-overview.html#access-control-framework>

Question 23: Skipped

An entity to which access can be granted. Unless allowed by a grant, access will be denied. in snowflake terms, what is this called?

- Securable object  
**(Correct)**
- Role
- Privilege
- User

**Explanation**

The key concepts to understanding access control in Snowflake are:

**Securable object:** An entity to which access can be granted. Unless allowed by a grant, access will be denied.

**Role:** An entity to which privileges can be granted. Roles are in turn assigned to users. Note that roles can also be assigned to other roles, creating a role hierarchy.

**Privilege:** A defined level of access to an object. Multiple distinct privileges may be used to control the granularity of access granted.

**User:** A user identity recognized by Snowflake, whether associated with a person or program.

<https://docs.snowflake.com/en/user-guide/security-access-control-overview.html#access-control-framework>

Question 24: Skipped

Which are the system defined roles in snowflake?

- ACCOUNTADMIN  
**(Correct)**
- SECURITYADMIN  
**(Correct)**
- USERADMIN  
**(Correct)**
- SYSADMIN  
**(Correct)**



## POWERUSER

### **Explanation**

Below are all the system defined roles

### **ACCOUNTADMIN**

(aka Account Administrator)

Role that encapsulates the SYSADMIN and SECURITYADMIN system-defined roles. It is the top-level role in the system and should be granted only to a limited/controlled number of users in your account.

### **SECURITYADMIN**

(aka Security Administrator)

Role that can manage any object grant globally, as well as create, monitor, and manage users and roles. More specifically, this role:

Is granted the MANAGE GRANTS security privilege to be able to modify any grant, including revoking it.

Inherits the privileges of the USERADMIN role via the system role hierarchy (e.g. USERADMIN role is granted to SECURITYADMIN).

### **USERADMIN**

(aka User and Role Administrator)

Role that is dedicated to user and role management only. More specifically, this role:

Is granted the CREATE USER and CREATE ROLE security privileges.

Can create and manage users and roles in the account (assuming that ownership of those roles or users has not been transferred to another role).

### **SYSADMIN**

(aka System Administrator)

Role that has privileges to create warehouses and databases (and other objects) in an account.

If, as [recommended](#), you create a role hierarchy that ultimately assigns all custom roles to the SYSADMIN role, this role also has the ability to grant privileges on warehouses, databases, and other objects to other roles.

### **PUBLIC**

Pseudo-role that is automatically granted to every user and every role in your account. The PUBLIC role can own securable objects, just like any other role; however, the objects owned by the role are, by definition, available to every other user and role in your account.

This role is typically used in cases where explicit access control is not needed and all users are viewed as equal with regard to their access rights.

Question 25: Skipped

In snowflake, you want to compare whether two expressions are equal. The expressions can have null values. What is the best way to compare in this case

- Use EQUAL\_NULL  
(Correct)
- Run a query with WHERE EXPR1 = EXPR2
- Run a query with WHERE EXPR2=EXPR1

### Explanation

This is an important concept to understand in Snowflake. This has caused me pains in production. Even we do not follow this, you will have incorrect data in production.

EQUAL comparison operator ( $=$ ), treats NULLs as unknown values. Like most SQL languages, comparing  $\text{NULL} = \text{NULL}$  does not return  $\text{TRUE}$ . In SnowFlake, it returns  $\text{NULL}$

EQUAL\_NULL function is NULL-safe, meaning it treats NULLs as known values for comparing equality.

[https://docs.snowflake.com/en/sql-reference/functions/equal\\_null.html#equal-null](https://docs.snowflake.com/en/sql-reference/functions/equal_null.html#equal-null)

Question 26: Skipped

PUT command in snowflake can upload files from a local folder/directory of a client machines to the below locations

- Named internal stage  
(Correct)
- Internal stage for a specified table  
(Correct)
- Internal stage for the current user  
(Correct)
- External Stage

## Explanation

PUT does **not** support uploading files to external stages. To upload files to external stages, use the utilities provided by the cloud service.

The following Snowflake clients do not support PUT:

1. Go Snowflake Driver
2. .NET Driver
3. Node.js Driver

The [ODBC driver](#) supports PUT with Snowflake accounts hosted on the following platforms:

1. Amazon Web Services (using ODBC Driver Version 2.17.5 and higher).
2. Google Cloud Platform (using ODBC Driver Version 2.21.5 and higher).
3. Microsoft Azure (using ODBC Driver Version 2.20.2 and higher).

Question 27: [Skipped](#)

If you attempt to PUT a file but cannot because a file with the same name already exists in the stage, what can you do to load the file?

- Wait until the existing file's data is loaded, then retry PUT  
**(Correct)**
- Rename the file that you want to PUT  
**(Correct)**
- Set **OVERWRITE = TRUE**  
**(Correct)**
- Set **FORCE COPY = TRUE**

## Explanation

**OVERWRITE = TRUE | FALSE**

Specifies whether Snowflake overwrites an existing file with the same name during upload:

**TRUE**: An existing file with the same name is overwritten.

**FALSE**: An existing file with the same name is not overwritten.

If you attempt to PUT a file but cannot because a file with the same name already exists in the stage, you can do any of the following:

Wait until the existing file's data is loaded, then retry PUT.

Rename the file that you want to PUT.

Set `OVERWRITE = TRUE`. Do this only if it is actually safe to overwrite a file with data that might not yet have been loaded into Snowflake.

Question 28: Skipped

OVERWRITE option for PUT is supported in which cloud providers

- Amazon AWS  
**(Correct)**
- Microsoft Azure  
**(Correct)**
- Google cloud platform

### Explanation

The following platforms support the OVERWRITE option:

Amazon AWS

Microsoft Azure.

GCP, the Google cloud platform, does not support `OVERWRITE=FALSE`. On GCP, the PUT command always uploads the file, even if a file with the same name exists and OVERWRITE is set to false.

Question 29: Skipped

PUT command can be run from Snowflake worksheets

- TRUE
- FALSE  
**(Correct)**

### Explanation

The PUT command cannot be executed from the **Worksheets** page in the Snowflake web interface; instead, use the SnowSQL client to upload data files, or check the documentation for the specific Snowflake client to verify support for this command.

<https://docs.snowflake.com/en/sql-reference/sql/put.html#usage-notes>

Question 30: Skipped

A COUNT(\*) and COUNT(<COLUMN\_NAME>) on a table will always give you the same result

- TRUE
- FALSE  
**(Correct)**

### Explanation

A COUNT(\*) will return a total count of rows in the table, while COUNT(<column\_name>) will return a count of rows with a non-NULL value in that particular column. COUNT(A,B) only counts the rows that have no NULL values in either the A or B column while COUNT(<ALIAS>.\* ) can be used to count all the rows containing no NULL columns.

Question 31: **Skipped**

How did you validate the data after unloading it using COPY INTO command?

- Load the data in a relational table and validate the rows
- Load the data in a CSV and validate the rows
- Use validation\_mode='RETURN\_ROWS'  
**(Correct)**

### Explanation

Validating Data to be Unloaded (from a Query)

Execute COPY in validation mode to return the result of a query and view the data that will be unloaded from the **orderstiny** table if COPY is executed in normal mode:

```
copy into @my_stage
from (select * from orderstiny limit 5)
validation_mode='RETURN_ROWS';

-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
C1 | C2 | C3 | C4 | C5 | C6 | C7 | C8 |
C9 |           |           |           |           |           |           |
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
1 | 36901 | 0 | 173665.47 | 1996-01-02 | 5-LOW | Clerk#000000951 | 0 |
nstructions sleep furiously among |
2 | 78002 | 0 | 46929.18 | 1996-12-01 | 1-URGENT | Clerk#000000880 | 0 |
foxes. pending accounts at the pending\, silent asymptot |
3 | 123314 | F | 193846.25 | 1993-10-14 | 5-LOW | Clerk#000000955 | 0 |
sly final accounts boost. carefully regular ideas cajole carefully. depos |
```

4		136777		0		32151.78		1995-10-11		5-LOW		Clerk#000000124		0		
sits.	slyly	regular	warthogs	cajole.	regular\,	regular	theodolites	acro								
5		44485		F		144659.20		1994-07-30		5-LOW		Clerk#000000925		0		
quickly.																

<https://docs.snowflake.com/en/sql-reference/sql/copy-into-location.html#validating-data-to-be-unloaded-from-a-query>

Question 32: Skipped

How do you remove files from a stage?

- REMOVE  
**(Correct)**
- DELETE
- PURGE
- TRUNCATE

### Explanation

#### REMOVE

Removes files that have been staged (i.e. uploaded from a local file system or unloaded from a table) in one of the following Snowflake internal stages:

Named internal stage.

Stage for a specified table.

Stage for the current user.

Note that using the command to remove files from an external stage might work but is not officially supported.

REMOVE can be abbreviated to RM.

<https://docs.snowflake.com/en/sql-reference/sql/remove.html#remove>

Question 33: Skipped

To improve performance materialized view can created on top of external tables without any additional cost

- TRUE
-

False  
**(Correct)**

### Explanation

Materialized views are designed to improve query performance for workloads composed of common, repeated query patterns. However, materializing intermediate results **incurs additional costs**. As such, before creating any materialized views, you should consider whether the costs are offset by the savings from re-using these results frequently enough.

Question 34: Skipped

Which COPY INTO <table> copy options are not supported by SNOWPIPE?

- FILES = ( 'file\_name1' [ , 'file\_name2', ... ] )**  
**(Correct)**
- ON\_ERROR = ABORT\_STATEMENT**  
**(Correct)**
- PURGE = TRUE | FALSE**  
**(Correct)**
- FORCE = TRUE | FALSE**  
**(Correct)**
- FILE\_FORMAT =( { FORMAT\_NAME = '[<namespace>.]<file\_format\_name>' | TYPE = { CSV | JSON | AVRO | ORC | PARQUET | XML } [ formatTypeOptions ] } )**

### Explanation

#### Usage Notes

All **COPY INTO <table>** copy options are supported *except for* the following:

**FILES = ( 'file\_name1' [ , 'file\_name2', ... ] )**

**ON\_ERROR = ABORT\_STATEMENT**

**SIZE\_LIMIT = num**

**PURGE = TRUE | FALSE** (i.e. automatic purging while loading)

**MATCH\_BY\_COLUMN\_NAME = CASE\_SENSITIVE | CASE\_INSENSITIVE | NONE**

**FORCE = TRUE | FALSE**

Note that you can manually remove files from an internal (i.e. Snowflake) stage (after they've been loaded) using the **REMOVE** command.

`RETURN_FAILED_ONLY = TRUE | FALSE`

`VALIDATION_MODE = RETURN_n_ROWS | RETURN_ERRORS | RETURN_ALL_ERRORS`

Please go through this link thoroughly

<https://docs.snowflake.com/en/sql-reference/sql/create-pipe.html#usage-notes>

Please note that `PATTERN = 'regex_pattern'` copy option is provided as a **preview feature**. If this is listed as one of the option, it will fall under not supported option.

Question 35: Skipped

A stream itself does **not** contain any table data

- TRUE  
(Correct)
- FALSE

### Explanation

Note that a stream itself does **not** contain any table data. A stream only stores the offset for the source table and returns CDC records by leveraging the versioning history for the source table. When the first stream for a table is created, a pair of hidden columns are added to the source table and begin storing change tracking metadata. These columns consume a small amount of storage. The CDC records returned when querying a stream rely on a combination of the *offset* stored in the stream and the *change tracking metadata* stored in the table.

Question 36: Skipped

A DML statement that selects from a stream consumes all of the change data in the stream as long as the transaction commits successfully

- TRUE  
(Correct)
- FALSE

### Explanation

A **DML statement** that selects from a stream consumes all of the change data in the stream as long as the transaction commits successfully. To ensure multiple statements access the same change records in the stream, surround them with an explicit transaction statement (**BEGIN .. COMMIT**). This locks the stream. DML updates to the source table in parallel transactions are tracked by the change tracking system but do not update the stream until the explicit transaction statement is committed and the existing change data is consumed.

This is very important to know, if you are given a TASK which consumes from the stream and a set of query statements and then if you are asked which query statement will consume all the change data then please select the DML query statement(like INSERT query)

Question 37: Skipped

By executing the 'SHOW TABLES' command, we can list all the tables in all the schemas even if we do not have access to all the tables

- TRUE
- FALSE  
(Correct)

### Explanation

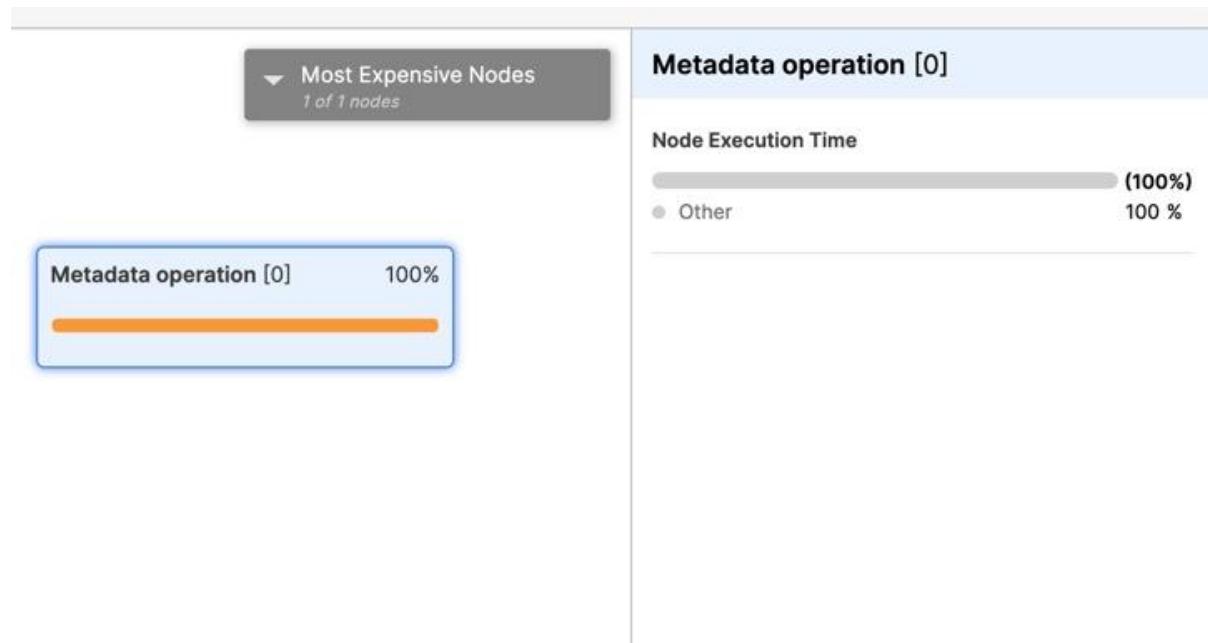
#### SHOW TABLES

Lists the tables for **which you have access privileges**, including dropped tables that are still within the Time Travel retention period and, therefore, can be undropped. The command can be used to list tables for the current/specify database or schema, or across your entire account.

The output returns table metadata and properties, ordered lexicographically by database, schema, and table name (see [Output](#) in this topic for descriptions of the output columns). This is important to note if you wish to filter the results using the provided filters.

Please also note the comment on **table metadata**

The SHOW COMMAND does not cost you any compute resource as it does not use warehouse. It gets the result from metadata store. Please try it yourself. Run this command and go to query profile and see yourself



Question 38: Skipped

Let's say that you have two JSONs as below

1. {"stuId":2000, "stuName":"Amy"}
2. {"stuId":2000, "stuCourse":"Snowflake"}

How will you write a query that will check if stuId in JSON in #1 is also there in JSON in #2

```
1. with stu_demography
2. as
3. (select parse_json(column1) as src, src:stuId as ID
4. from values('{"stuId":2000, "stuName":"Amy"}')),
5. stu_course
6. as
7. (select parse_json(column1) as src, src:stuId as ID
8. from values('{"stuId":2000, "stuCourse":"Snowflake"}'))
9. select
10. case
11. when stdemo.ID in(select ID from stu_course)
12. then 'True'
13. else 'False'
14. end as result
15. from stu_demography stdemo;
```

**(Correct)**

```
1. with stu_demography
2. as
3. (select parse_json(column1) as src, src['stuId'] as ID
4. from values('{"stuId":2000, "stuName":"Amy"}')),
5. stu_course
6. as
7. (select parse_json(column1) as src, src['stuId'] as ID
8. from values('{"stuId":2000, "stuCourse":"Snowflake"}'))
9. select
10. case
11. when stdemo.ID in(select ID from stu_course)
12. then 'True'
13. else 'False'
14. end as result
15. from stu_demography stdemo;
```

**(Correct)**

```
SELECT CONTAINS('{"stuId":2000,
"stuName":"Amy"}', '{"stuId":2000, "stuCourse":"Snowflake"}');
```

```
1. with stu_demography
2. as
3. (select parse_json(column1) as src, src['STUID'] as ID
4. from values('{"stuId":2000, "stuName":"Amy"}')),
5. stu_course
6. as
7. (select parse_json(column1) as src, src['stuId'] as ID
8. from values('{"stuId":2000, "stuCourse":"Snowflake"}'))
```

```

9. select
10. case
11. when stdemo.ID in(select ID from stu_course)
12. then 'True'
13. else 'False'
14. end as result
15. from stu_demography stdemo;

```

## Explanation

I would like you to try this out in your snowflake instance and find that out

Please note that this may not be the way the question will appear in the certification exam, but why we are still learning this?

1. When you take this course, I prepare you not just for the certification but also to be prepared to work on a project
2. With this hands-on, what will you learn
  - a. That there are two ways to query JSON, dot notation(src:stuId) and bracket notation src['stuId']
  - b. The JSON Value is case sensitive, if you run the 4th query with src['STUID'], it will return false because src['STUID'] will return you a NULL since there is a mismatch in the case.

In the certification you may be asked question on a), b) - but do you agree if you learn this way you will never forget the concept:)

## Question 39: Skipped

Create a task and a stream following the below steps. So, when the system\$stream\_has\_data('rawstream1') condition returns false, what will happen to the task ?

1. -- Create a landing table to store raw JSON data.
2. -- Snowpipe could load data **into this** table.
3. create or replace table raw (var variant);
- 4.
5. -- Create a stream to capture inserts to the landing table.
6. -- A task will consume a **set** of columns **from this** stream.
7. create or replace stream rawstream1 on table raw;
- 8.
9. -- Create a second stream to capture inserts to the landing table.
10. -- A second task will consume another **set** of columns **from this** stream.
11. create or replace stream rawstream2 on table raw;
- 12.
13. -- Create a table that stores the names of office visitors identified **in** the raw data.
14. create or replace table names (id int, first\_name string, last\_name string);
- 15.
16. -- Create a table that stores the visitation dates of office visitors identified **in** the raw data.
17. create or replace table visits (id int, dt date);
- 18.
19. -- Create a task that inserts **new** name records **from** the rawstream1 stream **into** the names table

```

20. -- every minute when the stream contains records.
21. -- Replace the 'etl_wh' warehouse with a warehouse that your role has USAGE privilege on.
22. create or replace task raw_to_names
23. warehouse = etl_wh
24. schedule = '1 minute'
25. when
26. system$stream_has_data('rawstream1')
27. as
28. merge into names n
29.   using (select var:id id, var:fname fname, var:lname lname from rawstream1) r1 on
      n.id = to_number(r1.id)
30.   when matched then update set n.first_name = r1.fname, n.last_name = r1.lname
31.   when not matched then insert (id, first_name, last_name) values (r1.id, r1.fname,
      r1.lname)
32. ;
33.
34. -- Create another task that merges visitation records from the rawstream1 stream in
   to the visits table
35. -- every minute when the stream contains records.
36. -- Records with new IDs are inserted into the visits table;
37. -- Records with IDs that exist in the visits table update the DT column in the table.
38. -- Replace the 'etl_wh' warehouse with a warehouse that your role has USAGE privilege on.
39. create or replace task raw_to_visits
40. warehouse = etl_wh
41. schedule = '1 minute'
42. when
43. system$stream_has_data('rawstream2')
44. as
45. merge into visits v
46.   using (select var:id id, var:visit_dt visit_dt from rawstream2) r2 on v.id = to_number(r2.id)
47.   when matched then update set v.dt = r2.visit_dt
48.   when not matched then insert (id, dt) values (r2.id, r2.visit_dt)
49. ;
50.
51. -- Resume both tasks.
52. alter task raw_to_names resume;
53. alter task raw_to_visits resume;
54.
55. -- Insert a set of records into the landing table.
56. insert into raw
57.   select parse_json(column1)
58.   from values
59.   ('{"id": "123", "fname": "Jane", "lname": "Smith", "visit_dt": "2019-09-17"}'),
60.   ('{"id": "456", "fname": "Peter", "lname": "Williams", "visit_dt": "2019-09-17"}');
61.
62. -- Query the change data capture record in the table streams
63. select * from rawstream1;
64. select * from rawstream2;

```

- Task will be executed but no rows will be merged

- Task will return an warning message

- Task will be skipped  
**(Correct)**

## Explanation

I would like you to complete the above steps and then run the below query to see it yourself. This way you will never forget this:)

```
select *
  from table(information_schema.task_history())
 order by scheduled_time;
```

Do you see that the task has been skipped?

MA_NAME	QUERY_TEXT	CONDITION_TEXT	STATE	ERROR_CODE	ERROR_MESSAGE	SCHEDULED_TIME	QUERY_START
IV_SC...	merge into v...	system\$stream_has_data('rawstream2')	SUCCEEDED	NULL	NULL	2020-09-16 ...	2020-09-16
IV_SC...	merge into n...	system\$stream_has_data('rawstream1')	SKIPPED	0040003	Conditional expression for task evaluated to false.	2020-09-16 ...	NUI
IV_SC...	merge into v...	system\$stream_has_data('rawstream2')	SKIPPED	0040003	Conditional expression for task evaluated to false.	2020-09-16 ...	NUI
IV_SC...	merge into n...	system\$stream_has_data('rawstream1')	SKIPPED	0040003	Conditional expression for task evaluated to false.	2020-09-16 ...	NUI
IV_SC...	merge into v...	system\$stream_has_data('rawstream2')	SKIPPED	0040003	Conditional expression for task evaluated to false.	2020-09-16 ...	NUI
IV_SC...	merge into n...	system\$stream_has_data('rawstream1')	SKIPPED	0040003	Conditional expression for task evaluated to false.	2020-09-16 ...	NUI
IV_SC...	merge into v...	system\$stream_has_data('rawstream2')	SKIPPED	0040003	Conditional expression for task evaluated to false.	2020-09-16 ...	NUI

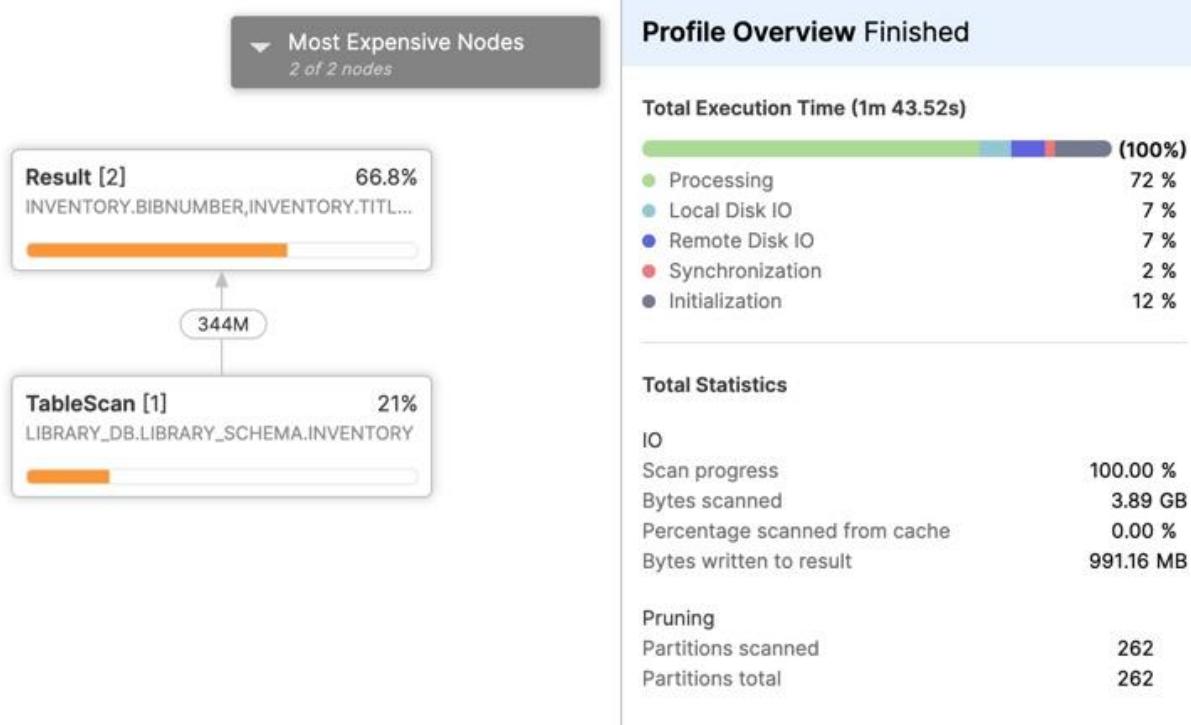
Question 40: Skipped

You ran the below query. I have a warehouse with auto suspend set at 5 seconds. Also result cache is disabled.

```
SELECT * FROM INVENTORY;
```

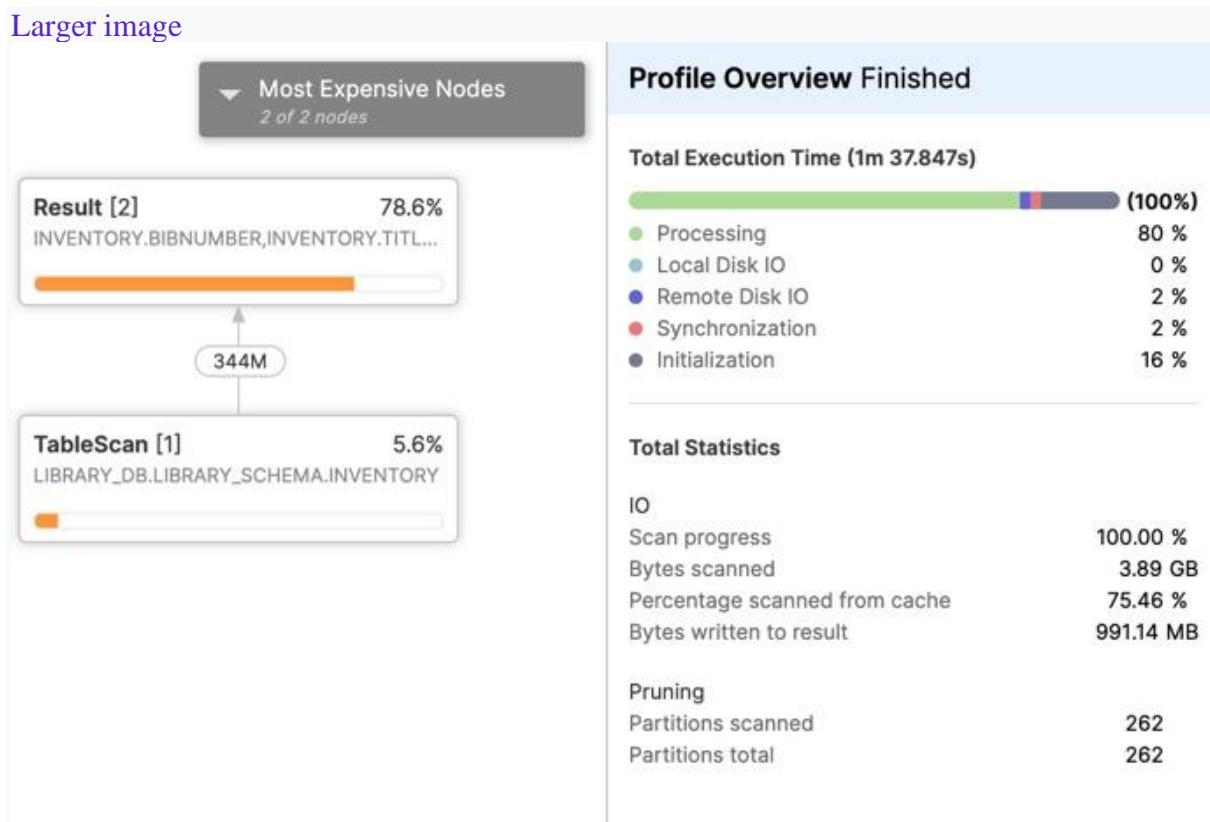
The query profile looks like as below. Please see below 'Percentage scanned from cache' is 0%

Larger image



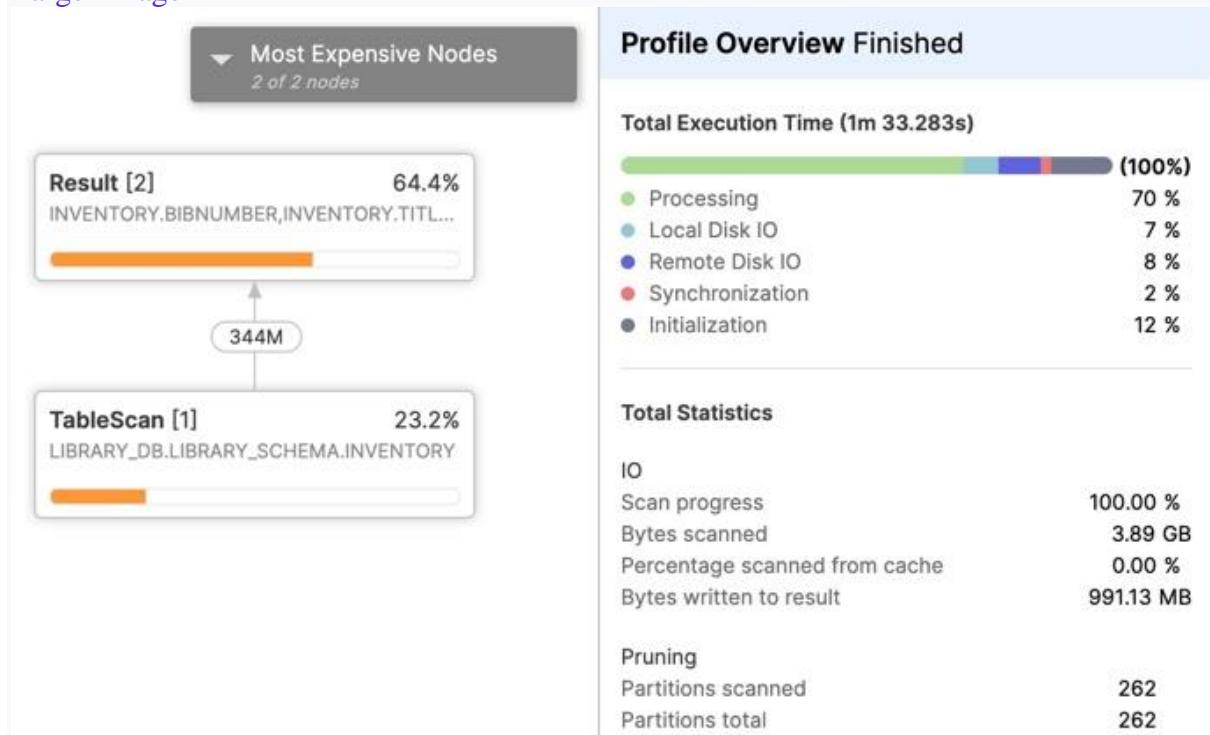
You ran the query again before 5 seconds has elapsed and the query profile looks as below. Look at the 'Percentage scanned for cache', it is 75%

[Larger image](#)



You ran the query again after 5 seconds. The query profile looks as below. Look at the 'Percentage scanned from cache', it is zero again.

[Larger image](#)



Why is this happening?

- The second run of the query used data cache to retrieve part of the result since it ran before the warehouse was suspended  
**(Correct)**
- The second run of the query used query result cache
- The third run of the query used query result cache

### Explanation

This is a very important concept to understand. There may be many different questions on this concept.

### Lets, understand what is going on in here

Virtual warehouses are an abstraction on the compute instances of the cloud provider(in case of AWS, it is EC2 instances). Each Virtual warehouse is a cluster of these compute instances(or EC2 in case of AWS). The compute instances has local SSD attached to them. When you ran the query for the first time, the results of the query were retrieved from the remote storage(which is the object store of the cloud provider, S3 in case of AWS), part of the results also got cached in the local SSD storage of the compute instance. So, when we ran the query second time, part of the results got retrieved from the SSD cache also known as **Data Cache**.

Ok, if that is the case why did not it retrieve from data cache the third time. The third run of the query happened after 5 seconds. The virtual warehouse had a auto suspend setting of 5 seconds. So, since there were no activity for 5 seconds, the warehouse suspended itself. When the warehouse is suspended, it loses the data cache. Why? because when it resumes it may be a completely new set of compute instances which were not used in the earlier runs.

Hope this clarifies and you probably will never forget this concept now. Please remember SSD cache is also called **Data cache**

In real work scenario, how can you leverage the understanding of this concept. Well if you a nightly ETL that runs against several tables, group the ETLs which use similar tables to run at the same time. This will help you to utilize the data cache effectively. This strategy, I named it as ETL-Colocation:)

Question 41: Skipped

Snowflake adds the system defined roles to each new account. Check the system defined roles below.

- SYSADMIN  
**(Correct)**

- ACCOUNTADMIN  
**(Correct)**
- SECURITYADMIN  
**(Correct)**
- PUBLIC  
**(Correct)**
- USERADMIN  
**(Correct)**
- SUPERUSER

### Explanation

There are a small number of system-defined roles in a Snowflake account. Users with appropriate access can alter the system-defined roles and can also create custom roles.

#### System-Defined Roles

##### **ACCOUNTADMIN**

(aka Account Administrator)

Role that encapsulates the SYSADMIN and SECURITYADMIN system-defined roles. It is the top-level role in the system and should be granted only to a limited/controlled number of users in your account.

##### **SECURITYADMIN**

(aka Security Administrator)

Role that can manage any object grant globally, as well as create, monitor, and manage users and roles. More specifically, this role:

Is granted the MANAGE GRANTS security privilege to be able to modify any grant, including revoking it.

Inherits the privileges of the USERADMIN role via the system role hierarchy (e.g. USERADMIN role is granted to SECURITYADMIN).

##### **USERADMIN**

(aka User and Role Administrator)

Role that is dedicated to user and role management only. More specifically, this role:

Is granted the CREATE USER and CREATE ROLE security privileges.

Can create and manage users and roles in the account (assuming that ownership of those roles or users has not been transferred to another role).

## SYSADMIN

(aka System Administrator)

Role that has privileges to create warehouses and databases (and other objects) in an account.

If, as [recommended](#), you create a role hierarchy that ultimately assigns all custom roles to the SYSADMIN role, this role also has the ability to grant privileges on warehouses, databases, and other objects to other roles.

## PUBLIC

Pseudo-role that is automatically granted to every user and every role in your account. The PUBLIC role can own securable objects, just like any other role; however, the objects owned by the role are, by definition, available to every other user and role in your account.

This role is typically used in cases where explicit access control is not needed and all users are viewed as equal with regard to their access rights.

Question 42: [Skipped](#)

You are running your query in the worksheet using a large warehouse. You want to change the warehouse. How do you do that? Choose 2 answers

- USE WAREHOUSE ETL\_WH;  
**(Correct)**
- SET WAREHOUSE CONTEXT=ETL\_WH
- Update the Warehouse field in the Context Menu located above the worksheet  
**(Correct)**
- Go to the Worksheet properties page and set the Warehouse field to ETL\_WH

## Explanation

Let us not mug this up. Logon to your snowflake account and open a worksheet

## Option#1

```

1 USE WAREHOUSE ETL_WH;

```

**Results Data Preview**

✓ Query ID SQL 37ms 1 rows

Filter result...

Row	status
1	Statement executed successfully.

## Option#2

New Worksheet ✓ WAREHOUSE + >

SYSADMIN ETL\_WH (XS) Select Database Select Schema

Role SYSADMIN Change

Warehouse ETL\_WH (XS) Suspended

Database Select Database

Schema Schema

Question 43: Skipped

The COPY command is more performant than the INSERT statement?

- TRUE  
**(Correct)**
- FALSE

## Explanation

Let us understand why the answer is TRUE here.

Do you remember that Snowflake stores data physically into immutable micro-partitions? So, how does insert/update work in snowflake?

Since the micropartitions are immutable, any insert or update will have to copy the entire micro-partition into a new copy and then insert the new record in that partition. For update

the old record will be marked for deletion and the new record will be added. so, imagine if you do insert one by one, how many micropartitions will get created.

Snowflake's architecture is **optimized for bulk load**. Otherwise, small DMLs (and at high frequency) cause issue of triggering the engine to create a lot of small partition files (this is because each DML creates a new (micro)partition file; and a lot of small partition files are bad for data processing and the database engine has to regularly compact the small partition files into larger ones for more optimal query performance).

Hope this clarifies the concept.

Question 44: Skipped

To help control the usage of credits in Auto-scale mode, Snowflake provides a property, \_\_\_\_\_, that determines the scaling policy to use when automatically starting or shutting down additional clusters. Select the property

- Auto Scale
- Scaling Policy  
**(Correct)**
- Max no of cluster
- Min no of cluster

### Explanation

#### Setting the Scaling Policy for a Multi-cluster Warehouse

To help control the credits consumed by a multi-cluster warehouse running in Auto-scale mode, Snowflake provides scaling policies, which are used to determine when to start or shut down a cluster.

The scaling policy for a multi-cluster warehouse only applies if it is running in Auto-scale mode. In Maximized mode, all clusters run concurrently so there is no need to start or shut down individual clusters.

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

Question 45: Skipped

Which three objects did we explicitly refer to using the COPY INTO command to load data on using external stages?

- Table  
**(Correct)**
- File Format

(Correct)

- View

- Stage

(Correct)

### Explanation

Try this out and you will never forget

1. create a table as below

```
CREATE OR REPLACE TABLE EMPLOYEE(FIRST_NAME VARCHAR, LAST_NAME  
VARCHAR,EMPLOYEE_ID NUMBER, SALARY NUMBER,JOINING_DATE TIMESTAMP);
```

2. Create a file format and then a stage

```
CREATE OR REPLACE file format load_csv_format  
TYPE = 'CSV'  
FIELD_DELIMITER = ','  
FIELD_OPTIONALLY_ENCLOSED_BY = '\"'  
;  
CREATE OR REPLACE stage employee_stage  
file_format = load_csv_format  
;
```

3. Use SNOWSQL to load a file to a stage using PUT

```
put file:///Users/Udemy/Documents/EMPLOYEE.csv @employee_stage; --replace with  
location of your csv
```

Populate the CSV with at least 1 or 2 rows(First Name, Last Name, EMPLOYEE ID, SALARY, JOINING DATE)

4. Run the below command. I had to do SKIP\_HEADER because my CSV had the column headers also.

```
COPY INTO EMPLOYEE  
FROM @employee_stage  
FILE_FORMAT=(TYPE=CSV SKIP_HEADER=1);
```

The screenshot shows the Data Preview interface for a copy operation. The SQL query is:

```
COPY INTO EMPLOYEE  
FROM @employee_stage  
FILE_FORMAT=(TYPE=CSV SKIP_HEADER=1);
```

The results table has the following columns:

Row	file	status	rows_parsed	rows_loaded	error_limit	errors_seen	first_error	first_error_line	first_error_charac	first_error_column
1	employee_stag...	LOADED	1	1	1	0	NULL	NULL	NULL	NULL

Question 46: Skipped

Mark all the true statements

- Snowflake Stages can be defined as either External or Internal.  
**(Correct)**
- Snowflake Stages can be defined as either FTP or SFTP.
- External Stages require a cloud storage provider.  
**(Correct)**
- FTP Stages require three cloud storage providers.
- Cloud storage providers must be linked to internal stage objects.

Question 47: Skipped

What is the column datatype used for ingesting semi-structured data?

- VARCHAR
- BINARY
- VARIANT  
**(Correct)**

### Explanation

#### Try this out

```
create or replace table mytable
(
    src variant -- this is the datatype
);

insert into mytable
select parse_json(column1)
from values
('{
    "a": "1",
    "b": "2",
    "c": null
}')
, ('{
```

```
"a": "1",
"b": "2",
"c": "3"
}');
```

A bonus one:) Try this out as well

```
select strip_null_value(src:c) from mytable;
```

strip\_null\_value converted json null to SQL NULL. Remember this.

Question 48: Skipped

JSON and PARQUET files can be loaded to columns in the same table

- TRUE
- FALSE  
**(Correct)**

### Explanation

**Let us try to answer some logical questions to arrive at the answer for this question**

1. How do you specify the file format when you try to load a file to a snowflake table

```
[ FILE_FORMAT = ( { FORMAT_NAME = '[<namespace>.]<file_format_name>' |
                    TYPE = { CSV | JSON | AVRO | ORC | PARQUET | XML } [ form
                    atTypeOptions ] } ) ]
```

2. Can you specify more than one file format while loading a table?

No

And there is your answer:)

Question 49: Skipped

```
create or replace table result_scan_table_1 like result_scan_table;
```

Will the above query cost compute credits?

- No, since it is a metadata operation only  
**(Correct)**
- Yes, it will need compute as the table structure is also getting created
- It is hard to say without looking at query profile

### Explanation

Why do you think it will not cost any compute credit? The reason behind this is that the **LIKE clause** only creates the definition of the table(or the structure of the table) in the metadata repository. It does not require any compute credit for any type of metadata operations.

How can you prove that it did not cost you any compute. Look at the history tab after you run the query.

Status	Query ID	SQL Text	U...	Warehouse	Clust...	Size	Session ID	Start Time	End Time	Total Duration	Bytes Scanned
✓	0196fa14...	create or replace table...	J...	ETL_WH			94841851937	4:16:45 PM	4:16:45 PM	146ms	

The size column is empty. If it would have cost you compute, the size column would be populated with the warehouse size.

Now please try this out. Run the below query

```
SHOW TABLES LIKE 'result_scan%';
```

Will this cost you compute? (HINT - It is a metadata operation)

If you are not able to answer this, please ping me on UDEMY.

#### Question 50: Skipped

To avoid unexpected task executions due to daylight saving time, you will do one of the below activities.

- Do not schedule tasks to run at a specific time between 1 AM and 3 AM (daily, or on days of the week that include Sundays)  
**(Correct)**
- Manually adjust the cron expression for tasks scheduled during those hours twice each year to compensate for the time change due to daylight saving time  
**(Correct)**
- Add the USE DAYLIGHT SAVINGS parameter for task

#### Explanation

The cron expression in a task definition supports specifying a time zone. A scheduled task runs according to the specified cron expression in the local time for a given time zone. Special care should be taken with regard to scheduling tasks for time zones that recognize daylight saving time. Tasks scheduled during specific times on days when the transition from standard time to daylight saving time (or the reverse) occurs can have unexpected behaviors.

For example:

During the autumn change from daylight saving time to standard time, a task scheduled to start at 1 AM in the America/Los\_Angeles time zone (i.e. **0 1 \* \* \***)

`America/Los_Angeles`) would run *twice*: once at 1 AM and then again when 1:59:59 AM shifts to 1:00:00 AM local time. That is, there are two points in time when the local time is 1 AM.

During the spring change from standard time to daylight saving time, a task scheduled to start at 2 AM in the America/Los\_Angeles time zone (i.e. `0 2 * * * America/Los_Angeles`) would **not run at all** because the local time shifts from 1:59:59 AM to 3:00:00 AM. That is, there is no point during that day when the local time is 2 AM.

To avoid unexpected task executions due to daylight saving time, *either*:

1. Do not schedule tasks to run at a specific time between 1 AM and 3 AM (daily, or on days of the week that include Sundays), or
2. Manually adjust the cron expression for tasks scheduled during those hours twice each year to compensate for the time change due to daylight saving time.

Question 51: Skipped

The query profiler view is only available for completed queries

- TRUE
- FALSE  
**(Correct)**

### Explanation

The query profiler is available for query in progress also.

Question 52: Skipped

You have created a secure view `view1`. Now you want to convert it back to a regular view.

Which of the below can be used for the conversion?

- 1. `alter view view1 set type = unsecure;`
- 1. `alter view view1 unset secure;`  
**(Correct)**
- 1. `alter view view1 set secure;`
- None of the above

### Explanation

Revert a secure view to a regular view:

```
alter view view1 unset secure;
```

Please look at the other examples here

<https://docs.snowflake.com/en/sql-reference/sql/alter-view.html#examples>

Question 53: Skipped

You want to load a JSON file but would like to remove the outer array structure and load the records into separate table rows. Which file format option would you use?

- STRIP\_OUTER\_ARRAY  
**(Correct)**
- TRUNCATE\_OUTER\_ARRAY
- UNSET\_OUTER\_ARRAY
- DROP\_OUTER\_ARRAY

**Explanation**

In general, JSON and Avro data sets are a simple concatenation of multiple documents. The JSON or Avro output from some software is composed of a single huge array containing multiple records. There is no need to separate the documents with line breaks or commas, though both are supported.

Instead, snowflake recommends enabling the STRIP\_OUTER\_ARRAY file format option for the **COPY INTO <table>** command to remove the outer array structure and load the records into separate table rows:

```
copy into <table>
from @~/<file>.json
file_format = (type = 'JSON' strip_outer_array = true);
```

Question 54: Skipped

Which two options mentioned below will you choose to concat a column named 'FRUIT\_DESCRIPTION' between two % signs?

- '%'||FRUIT\_DESCRIPTION||'%'  
**(Correct)**
- CONCAT('%',CONCAT(FRUIT\_DESCRIPTION,'%'))  
**(Correct)**
- %+FRUIT\_DESCRIPTION+%
- SUBSTR(%,FRUIT\_DESCRIPTION,%)

**Explanation**

**Try it yourself**

```
CREATE OR REPLACE TABLE FRUITS(FRUIT_NUMBER NUMBER, FRUIT_DESCRIPTION VARCHAR,
AVAILABILITY VARCHAR);
```

```

INSERT INTO FRUITS VALUES(1,'APPLE',NULL);
INSERT INTO FRUITS VALUES(2,'MANGO',NULL);
INSERT INTO FRUITS VALUES(3,'ORANGE',NULL);
INSERT INTO FRUITS VALUES(4,'KIWI',NULL);
INSERT INTO FRUITS VALUES(5,'MELON',NULL);
INSERT INTO FRUITS VALUES(6,'PEARS',NULL);
INSERT INTO FRUITS VALUES(7,'BANANA',NULL);
INSERT INTO FRUITS VALUES(8,'PLUM',NULL);
INSERT INTO FRUITS VALUES(9,'APRICOT',NULL);
INSERT INTO FRUITS VALUES(10,'WATERMELON',NULL);

SELECT '%' || FRUIT_DESCRIPTION || '%' FROM FRUITS;
SELECT CONCAT('%',CONCAT(FRUIT_DESCRIPTION,'%')) FROM FRUITS;

```

Question 55: Skipped

Snowflake has three types of stages USER, TABLE, NAMED. Named stage is again divided into External and Internal. Which of the below two stages are automatically created and does not need explicit configuration by the user

- TABLE  
**(Correct)**
- EXTERNAL
- INTERNAL
- USER  
**(Correct)**

Question 56: Skipped

When you create internal stage, you must link it to the cloud storage providers

- FALSE  
**(Correct)**
- TRUE

### Explanation

That is the beauty of internal stage. Snowflake manages it for you. You do not have to do anything with the underlying cloud provider. Snowflake also automatically encrypts the data for you in internal stage.

**Question 57:** Skipped

Data providers must sign an agreement before their data is made public on the data marketplace of snowflake

- TRUE  
**(Correct)**
- FALSE

**Explanation**

Before you can list a data set on the Snowflake Data Marketplace, you must become an approved data provider.

The approval is a one-time process that ensures that:

1. Provider is a Snowflake customer in the supported regions.
2. Provider has signed the Snowflake Provider agreement.
3. Provider's data meets Snowflake's data requirements. To meet the requirements, shared data must be:
  - a. Fresh — the data is near real-time or updated on a regular basis.
  - b. Real — not sample data.
  - c. Legally shareable — the provider must own the data or have the right to share it.

**Question 58:** Skipped

Snowflake supports the below security features

- TRI-SECRET SECURE ENCRYPTION  
**(Correct)**
- ROLE-BASED ACCESS  
**(Correct)**
- MULTI-FACTOR AUTHENTICATION  
**(Correct)**
- NONE OF THE ABOVE

**Question 59:** Skipped

Snowflake supports the following performance optimization techniques

- SSD CACHING/DATA CACHING  
**(Correct)**

- QUERY RESULT CACHING  
**(Correct)**
- B-TREE INDEXES

### Explanation

Please read the below article

<https://community.snowflake.com/s/article/Caching-in-Snowflake-Data-Warehouse>

Question 60: **Skipped**

Select two options which are true about variant data type in snowflake?

- Supports querying using JSON path notation  
**(Correct)**
- Optimizes storage based on repeated elements  
**(Correct)**
- A separate file format is used to store variant data
- Requires custom mapping for each type of record

[Continue](#)

[Retake test](#)

Question 1: Skipped

Which role in Snowflake allows a user to administer users and manage all database objects?

- ACCOUNTADMIN  
**(Correct)**
- SYSADMIN
- SECURITYADMIN
- ROOT

### Explanation

The account administrator (ACCOUNTADMIN) role is the most powerful role in the system. This role alone is responsible for configuring parameters at the account level. Users with the ACCOUNTADMIN role can view and operate on all objects in the account, can view and manage Snowflake billing and credit data, and can stop any running SQL statements.

In the default access control hierarchy, both of the other administrator roles are owned by this role:

The security administrator (SECURITYADMIN) role includes the privileges to create and manage users and roles.

The system administrator (SYSADMIN) role includes the privileges to create warehouses, databases, and all database objects (schemas, tables, etc.).

<https://docs.snowflake.com/en/user-guide/security-access-control-considerations.html#using-the-accountadmin-role>

Question 2: Skipped

Which transformations are available when using the COPY INTO command to load data files into Snowflake from a stage? (select all that apply)

- Column data type conversion  
**(Correct)**
- Column concatenation  
**(Correct)**
- Filters

- Aggregates

### Explanation

Filtering the results of a `FROM` clause using a `WHERE` clause is not supported.

The `VALIDATION_MODE` parameter does not support `COPY` statements that transform data during a load.

<https://docs.snowflake.com/en/user-guide/data-load-transform.html#supported-functions>

Question 3: Skipped

Snowflake offers tools to extract data from source systems

- TRUE
- FALSE  
**(Correct)**

Question 4: Skipped

Select the layers which are part of snowflake(select 3)

- STORAGE  
**(Correct)**
- DATA CATALOG
- VIRTUAL WAREHOUSE  
**(Correct)**
- CLOUD SERVICES  
**(Correct)**

### Explanation

Snowflake's novel design physically separates but logically integrates storage, compute and services like security and metadata; we call it multi-cluster, shared data and it consists of 3 components:

1. Storage: the persistent storage layer for data stored in Snowflake
2. Compute: a collection of independent compute resources that execute data processing tasks required for queries

3. Services: a collection of system services that handle infrastructure, security, metadata, and optimization across the entire Snowflake system

<https://www.snowflake.com/wp-content/uploads/2014/10/A-Detailed-View-Inside-Snowflake.pdf>

Question 5: Skipped

When data is staged to a Snowflake internal staging area using the PUT command, the data is encrypted on the client's machine

- TRUE  
**(Correct)**
- FALSE

#### Explanation

Uploaded files are automatically encrypted with 128-bit or 256-bit keys. The CLIENT\_ENCRYPTION\_KEY\_SIZE parameter specifies the size key used to encrypt the files

<https://docs.snowflake.com/en/sql-reference/sql/put.html#usage-notes>

Question 6: Skipped

Which snowflake features are available for enabling continuous data pipelines

- Continuous data loading  
**(Correct)**
- Change data tracking  
**(Correct)**
- Recurring tasks  
**(Correct)**
- Table Pipes

#### Explanation

Snowflake provides the following features to enable continuous data pipelines:

#### Continuous data loading

Options for continuous data loading include the following:

[Snowpipe](#)

## Snowflake Connector for Kafka

Third-party data integration tools

### Change data tracking

A *stream* object records the delta of change data capture (CDC) information for a table (such as a staging table), including inserts and other data manipulation language (DML) changes. A stream allows querying and consuming a set of changes to a table, at the row level, between two transactional points of time.

In a continuous data pipeline, table streams record when staging tables and any downstream tables are populated with data from business applications using continuous data loading and are ready for further processing using SQL statements.

For more information, see [Change Tracking Using Table Streams](#).

### Recurring tasks

A *task* object defines a recurring schedule for executing a SQL statement, including statements that call stored procedures. Tasks can be chained together for successive execution to support more complex periodic processing.

Tasks may optionally use table streams to provide a convenient way to continuously process new or changed data. A task can transform new or changed rows that a stream surfaces. Each time a task is scheduled to run, it can verify whether a stream contains change data for a table (using `SYSTEM$STREAM_HAS_DATA`) and either consume the change data or skip the current run if no change data exists.

Users can define a simple tree-like structure of tasks that executes consecutive SQL statements to process data and move it to various destination tables.

<https://docs.snowflake.com/en/user-guide/data-pipelines-intro.html#introduction-to-data-pipelines>

Question 7: Skipped

Select the statements which are true for an external table

- External tables are read-only  
**(Correct)**
- External tables can be used for query and join operations  
**(Correct)**
- Views can be created against external tables

(Correct)

- Data can be updated in external tables

### Explanation

In a typical table, the data is stored in the database; however, in an external table, the data is stored in files in an external stage. External tables store file-level metadata about the data files, such as the filename, a version identifier and related properties. This enables querying data stored in files in an external stage as if it were inside a database. External tables can access data stored in any format supported by `COPY INTO <table>` statements.

External tables are read-only, therefore no DML operations can be performed on them; however, external tables can be used for query and join operations. Views can be created against external tables.

Querying data stored external to the database is likely to be slower than querying native database tables; however, materialized views based on external tables can improve query performance.

<https://docs.snowflake.com/en/user-guide/tables-external-intro.html>

Question 8: Skipped

What are the two data loading approaches in snowflake

- BULK LOADING  
**(Correct)**
- CONTINUOUS LOADING  
**(Correct)**
- INGEST LOADING

### Explanation

BULK LOADING and CONTINUOUS LOADING are the two approaches. For bulk loading you can use the COPY command and select a relevant warehouse to perform the COPY. In this case compute is user managed

For continuous loading you can use snowpipe which is a serverless way of loading data in micro batches. Compute is managed by snowflake.

Question 9: Skipped

Which of the below are considered as best practices while loading data into snowflake?

-

Isolate data loading workload into its own virtual warehouse  
**(Correct)**

- Split large files into smaller files  
**(Correct)**
- Compress the source files  
**(Correct)**
- If format is in CSV, convert them into ORC

### Explanation

Having a dedicated warehouse for load workload ensures that it will not be interrupted by any other workloads. Splitting the large files is a good practice because it enables snowflake to parallelize the load operation. Since the files are sent over the wire, it is always better to compress them first.

Question 10: Skipped

Load performance in snowflake is fastest for which file format

- CSV  
**(Correct)**
- ORC
- AVRO
- PARQUET

Question 11: Skipped

COPY and INSERT operations in snowflake are non-blocking

- TRUE  
**(Correct)**
- FALSE

### Explanation

COPY and INSERT do not block any other operations on the table

Question 12: Skipped

Organizing input data by granular path can improve load performance

- TRUE  
**(Correct)**
-

FALSE

Question 13: Skipped

Which are the key concepts that will need to be considered while loading data into snowflake

- STAGE OBJECT  
**(Correct)**
- FILE FORMAT  
**(Correct)**
- TRANSFORMATION AND ERROR VALIDATION  
**(Correct)**
- FILE SIZE

### Explanation

Copying the file to a stage object is recommended while loading data into snowflake.

File format is used to identify the data format (CSV, JSON etc) of the source file.

Minor transformation and validations can be done as part of loading data

Question 14: Skipped

Which approach would result in improved performance through linear scaling of data ingestion workload?

- Resize virtual warehouse
- Consider practice of organizing data by granular path
- Consider practice of splitting input file batch within the recommended size of 10 MB to 100 MB
- All of the above  
**(Correct)**

Question 15: Skipped

Which are the two variant columns available in a snowflake table loaded by kafka connector

- RECORD\_CONTENT  
**(Correct)**
- RECORD\_METADATA  
**(Correct)**

- RECORD\_KEY

### Explanation

Every Snowflake table loaded by the Kafka connector has a schema consisting of two VARIANT columns:

RECORD\_CONTENT. This contains the Kafka message.

RECORD\_METADATA. This contains metadata about the message, for example, the topic from which the message was read.

If Snowflake creates the table, then the table contains only these two columns. If the user creates the table for the Kafka Connector to add rows to, then the table can contain more than these two columns (any additional columns must allow NULL values because data from the connector does not include values for those columns).

<https://docs.snowflake.com/en/user-guide/kafka-connector-overview.html#schema-of-topics-for-kafka-topics>

Question 16: Skipped

The RECORD\_METADATA contains which information

- Topic  
**(Correct)**
- Partition  
**(Correct)**
- Key  
**(Correct)**
- CreateTime / LogAppendTime  
**(Correct)**
- Value

### Explanation

The RECORD\_METADATA column contains the following information by default:

Topic

Partition

Offset

CreateTime / LogAppendTime

key

schema\_id

headers

The value of the message is in RECORD\_CONTENT

<https://docs.snowflake.com/en/user-guide/kafka-connector-overview.html#schema-of-topics-for-kafka-topics>

Question 17: Skipped

If multiple instances of the kafka connector is started on the same topic or partitions, duplicate records may flow into snowflake table

- TRUE  
**(Correct)**
- FALSE

### Explanation

Instances of the Kafka connector do not communicate with each other. If you start multiple instances of the connector on the same topics or partitions, then multiple copies of the same row might be inserted into the table. This is not recommended; each topic should be processed by only one instance of the connector.

Question 18: Skipped

Select the ones that are true for snowflake kafka connector

- Kafka connector guarantees exactly-once delivery  
**(Correct)**
- Kafka connector guarantees that rows are inserted in the order
- Kafka connector guarantees reprocessing of messages

### Explanation

Although the Kafka connector guarantees exactly-once delivery, it does **not** guarantee that rows are inserted in the order that they were originally published.

Question 19: Skipped

Which are true with respect to SMT(simple message transformation) when neither `key.converter` or `value.converter` is set

- All SMTs are supported
- Most SMTs are supported  
**(Correct)**
- regex.router is not supported  
**(Correct)**

### Explanation

Single Message Transformations (SMTs) are applied to messages as they flow through Kafka Connect. When you configure the [Kafka Configuration Properties](#), if you set **either** `key.converter` or `value.converter` to one of the following values, then SMTs are not supported on the corresponding key or value:

```
com.snowflake.kafka.connector.records.SnowflakeJsonConverter
```

```
com.snowflake.kafka.connector.records.SnowflakeAvroConverter
```

```
com.snowflake.kafka.connector.records.SnowflakeAvroConverterWithoutSchema
Registry
```

When neither `key.converter` or `value.converter` is set, then most SMTs are supported, with the current exception of `regex.router`.

<https://docs.snowflake.com/en/user-guide/kafka-connector-overview.html#kafka-connector-limitations>

Question 20: [Skipped](#)

Which are the supported data types for a JSON name/value pair?

- A number (integer or floating point)  
**(Correct)**
- A string (in double quotes)  
**(Correct)**
- A Boolean (true or false)  
**(Correct)**
- An array (in square brackets)

**(Correct)**

- An object (in curly braces)  
**(Correct)**
- Null  
**(Correct)**
- complex datatype

### Explanation

A value in a name/value pair can be:

1. A number (integer or floating point)
2. A string (in double quotes)
3. A Boolean (true or false)
4. An array (in square brackets)
5. An object (in curly braces)
6. Null

<https://docs.snowflake.com/en/user-guide/semistructured-intro.html#supported-data-types>

Question 21: Skipped

Which of the below are binary formats?

- AVRO  
**(Correct)**
- JSON
- ORC  
**(Correct)**
- PARQUET  
**(Correct)**

### Explanation

What is ORC?

Used to store Hive data, the ORC (Optimized Row Columnar) file format was designed for efficient compression and improved performance for reading, writing, and processing data over earlier Hive file formats. For more information about ORC, see <https://orc.apache.org/>.

Snowflake reads ORC data into a single VARIANT column. You can query the data in a VARIANT column just as you would JSON data, using similar commands and functions.

Alternatively, you can extract select columns from a staged ORC file into separate table columns using a CREATE TABLE AS SELECT statement.

### **ORC is a binary format.**

#### **What is parquet?**

Parquet is a compressed, efficient columnar data representation designed for projects in the Hadoop ecosystem. The file format supports complex nested data structures and uses Dremel record shredding and assembly algorithms. For more information, see [parquet.apache.org/documentation/latest/](http://parquet.apache.org/documentation/latest/).

Snowflake reads Parquet data into a single VARIANT column. You can query the data in a VARIANT column just as you would JSON data, using similar commands and functions.

Alternatively, you can extract select columns from a staged Parquet file into separate table columns using a CREATE TABLE AS SELECT statement.

### **Parquet is a binary format.**

#### **What is AVRO?**

Avro is an open-source data serialization and RPC framework originally developed for use with Apache Hadoop. It utilizes schemas defined in JSON to produce serialized data in a compact binary format. The serialized data can be sent to any destination (i.e. application or program) and can be easily deserialized at the destination because the schema is included in the data.

An Avro schema consists of a JSON string, object, or array that defines the type of schema and the data attributes (field names, data types, etc.) for the schema type. The attributes differ depending on the schema type. Complex data types such as arrays and maps are supported.

Snowflake reads Avro data into a single VARIANT column. You can query the data in a VARIANT column just as you would JSON data, using similar commands and functions.

<https://docs.snowflake.com/en/user-guide/semistructured-intro.html#what-is-parquet>

<https://docs.snowflake.com/en/user-guide/semistructured-intro.html#what-is-orc>

<https://docs.snowflake.com/en/user-guide/semistructured-intro.html#what-is-avro>

Question 22: Skipped

Which constraints are enforced in snowflake

- Referential integrity constraints
- NOT NULL constraint  
(Correct)
- UNIQUE Constraint

### Explanation

Referential integrity constraints in Snowflake are informational and, with the exception of NOT NULL, not enforced. Constraints other than NOT NULL are created as disabled.

However, constraints provide valuable metadata. The primary keys and foreign keys enable members of your project team to orient themselves to the schema design and familiarize themselves with how the tables relate with one another.

<https://docs.snowflake.com/en/user-guide/table-considerations.html#referential-integrity-constraints>

Question 23: Skipped

You want to get the DDL statement of a snowflake table. What is the command that you will use?

- 1. `select get_ddl('table', 'mydb.public.salesorders');`  
(Correct)
- 1. `show table 'mydb.public.salesorders';`
- 1. `show table like 'mydb.public.salesorders';`

### Explanation

Query the `GET_DDL` function to retrieve a DDL statement that could be executed to recreate the specified table. The statement includes the constraints currently set on a table.

Question 24: Skipped

You have a small table in snowflake which has only 10,000 rows. Specifying a clustering key will further improved the queries that run on this table

- FALSE  
**(Correct)**
- TRUE

### Explanation

Specifying a **clustering key** is not necessary for most tables. Snowflake performs automatic tuning via the optimization engine and micro-partitioning. In many cases, data is loaded and organized into micro-partitions by date or timestamp, and is queried along the same dimension.

When should you specify a clustering key for a table? First, note that clustering a small table typically doesn't improve query performance significantly.

For larger data sets, you might consider specifying a clustering key for a table when:

The order in which the data is loaded does not match the dimension by which it is most commonly queried (e.g. the data is loaded by date, but reports filter the data by ID). If your existing scripts or reports query the data by both date **and** ID (and potentially a third or fourth column), you may see some performance improvement by creating a multi-column clustering key.

**Query Profile** indicates that a significant percentage of the total duration time for typical queries against the table is spent scanning. This applies to queries that filter on one or more specific columns.

Note that reclustering rewrites existing data with a different order. The previous ordering is stored for 7 days to provide Fail-safe protection. Reclustering a table incurs compute costs that correlate to the size of the data that is reordered.

<https://docs.snowflake.com/en/user-guide/table-considerations.html#when-to-set-a-clustering-key>

Question 25: Skipped

There is no query performance difference between a column with a maximum length declaration (e.g. `VARCHAR(16777216)`), and a smaller precision. Still it is recommended to define an appropriate column length because of the below reasons

- Data loading operations are more likely to detect issues such as columns loaded out of order, e.g. a 50-character string loaded erroneously into a `VARCHAR(10)` column. Such issues produce errors  
**(Correct)**
-

When the column length is unspecified, some third-party tools may anticipate consuming the maximum size value, which can translate into increased client-side memory usage or unusual behavior

(Correct)

- Data unloading will be performant if appropriate column lengths are defined

### Explanation

Snowflake compresses column data effectively; therefore, creating columns larger than necessary has minimal impact on the size of data tables. Likewise, there is no query performance difference between a column with a maximum length declaration (e.g. `VARCHAR(16777216)`), and a smaller precision.

However, when the size of your column data is predictable, we do recommend defining an appropriate column length, for the following reasons:

Data loading operations are more likely to detect issues such as columns loaded out of order, e.g. a 50-character string loaded erroneously into a `VARCHAR(10)` column. Such issues produce errors.

When the column length is unspecified, some third-party tools may anticipate consuming the maximum size value, which can translate into increased client-side memory usage or unusual behavior.

<https://docs.snowflake.com/en/user-guide/table-considerations.html#when-to-specify-column-lengths>

Question 26: Skipped

You want to convert an existing permanent table to a transient table (or vice versa) while preserving data and other characteristics such as column defaults and granted privileges. What is the best way to do it?

- Run an `ALTER TABLE` command to convert the tables
- Unload the data from the existing table into a CSV file. Create the new table and then load the data back in
- Create a new table and use the `COPY GRANTS` clause

(Correct)

### Explanation

Currently, it isn't possible to change a permanent table to a `transient` table using the `ALTER TABLE` command. The `TRANSIENT` property is set at table creation and cannot be modified.

Similarly, it isn't possible to directly change a transient table to a permanent table.

To convert an existing permanent table to a transient table (or vice versa) while preserving data and other characteristics such as column defaults and granted privileges, you can create a new table and use the **COPY GRANTS** clause, then copy the data:

```
create transient table my_new_table like my_old_table copy grants;  
insert into my_new_table select * from my_old_table;
```

<https://docs.snowflake.com/en/user-guide/table-considerations.html#converting-a-permanent-table-to-a-transient-table-or-vice-versa>

Question 27: Skipped

A transient table can be cloned to a permanent table

- TRUE
- FALSE  
**(Correct)**

### Explanation

You can't clone a transient table to a permanent table.

Question 28: Skipped

If you clone a permanent table(bar) into a transient table(foo) using the below command

1. create **transient** table foo clone bar copy grants;

What will happen to the partitions?

- Old partitions will **not** be affected, but new partitions added to the clone will follow the transient lifecycle  
**(Correct)**
- All the partitions will be affected
- Only the old partitions will be affected

### Explanation

Another way to make a copy of a table (but change the lifecycle from permanent to transient) is to **CLONE** the table, for example:

```
create transient table foo clone bar copy grants;
```

Old partitions will **not** be affected (i.e. won't become transient), but new partitions added to the clone will follow the transient lifecycle.

Question 29: Skipped

You want to identify the potential performance bottlenecks and improvement opportunities of a query. What will you do?

- Use Query Profile  
**(Correct)**
- Use Explain plan
- Call snowflake support

### Explanation

Query Profile is a powerful tool for understanding the mechanics of queries. It can be used whenever you want or need to know more about the performance or behavior of a particular query. It is designed to help you spot typical mistakes in SQL query expressions to identify potential performance bottlenecks and improvement opportunities.

<https://docs.snowflake.com/en/user-guide/ui-query-profile.html#when-to-use-query-profile>

Question 30: Skipped

What does snowflake use for monitoring network traffic and use activity?

- Lacework  
**(Correct)**
- Sumo logic
- Threat Stack

### Explanation

Snowflake uses lacework for behavioral monitoring of production infrastructure which includes network traffic and user activity. It uses Sumo Logic and Threat Stack to monitor failed logins, file integrity monitoring and unauthorized system modifications.

Question 31: Skipped

In which scenarios would you consider to use materialized views

- The query results contain a small number of rows and/or columns relative to the base table  
**(Correct)**
- Query results contain results that require significant processing  
**(Correct)**

- Query is on an external table  
**(Correct)**
- View's base table does not change frequently  
**(Correct)**
- None of the above

### Explanation

Materialized views are particularly useful when:

1. Query results contain a small number of rows and/or columns relative to the base table (the table on which the view is defined).
2. Query results contain results that require significant processing, including:
3. Analysis of semi-structured data.
4. Aggregates that take a long time to calculate.
5. The query is on an external table (i.e. data sets stored in files in an external stage), which might have slower performance compared to querying native database tables.
6. The view's base table does not change frequently.

<https://docs.snowflake.com/en/user-guide/views-materialized.html#when-to-use-materialized-views>

Question 32: Skipped

What will the below query return

```
SELECT TOP 10 GRADES FROM STUDENT;
```

- The top 10 highest grades
- The 10 lowest grades
- Non-deterministic list of 10 grades  
**(Correct)**

### Explanation

An **ORDER BY** clause is not required; however, without an **ORDER BY** clause, the **results are non-deterministic** because results within a result set are not necessarily in any particular order. To control the results returned, use an **ORDER BY** clause.

**n** must be a non-negative integer constant.

[https://docs.snowflake.com/en/sql-reference/constructs/top\\_n.html#usage-notes](https://docs.snowflake.com/en/sql-reference/constructs/top_n.html#usage-notes)

Question 33: Skipped

Loading data using snowpipe REST API is supported for external stage only

- TRUE
- FALSE  
**(Correct)**

### Explanation

Snowpipe supports loading from the following stage types:

1. Named internal (Snowflake) or external (Amazon S3, Google Cloud Storage, or Microsoft Azure) stages
2. Table stages

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-rest-gs.html#step-1-create-a-stage-if-needed>

Question 34: Skipped

With default settings, how long will a query run on snowflake

- Snowflake will cancel the query if it runs more than 48 hours  
**(Correct)**
- Snowflake will cancel the query if it runs more than 24 hours
- Snowflake will cancel the query if the warehouse runs out of memory
- Snowflake will cancel the query if the warehouse runs out of memory and hard disk storage

### Explanation

**STATEMENT\_TIMEOUT\_IN\_SECONDS**

This parameter tells Snowflake how long can a SQL statement run before the system cancels it. The default value is 172800 seconds (48 hours)

This is both a session and object type parameter. As a session type, it can be applied to the account, a user or a session. As an object type, it can be applied to warehouses. If set at both levels, the lowest value is used.

Question 35: Skipped

You have created a TASK in snowflake. How will you resume it?

- No need to resume, the creation operation automatically enables the task
- ALTER TASK mytask1 RESUME;  
**(Correct)**
- ALTER TASK mytask1 START;

### Explanation

It is important to remember that a Task that has just been created will be suspended by default. It is necessary to manually enable this task by “altering” the task as follows:

ALTER TASK mytask1 RESUME;

Question 36: Skipped

What technique does snowflake use to limit the number of micro-partitions scanned by each query

- Indexing
- Pruning  
**(Correct)**
- Map Reduce
- B-tree

### Explanation

The micro-partition metadata maintained by Snowflake enables precise pruning of columns in micro-partitions at query run-time, including columns containing semi-structured data. In other words, a query that specifies a filter predicate on a range of values that accesses 10% of the values in the range should ideally only scan 10% of the micro-partitions.

For example, assume a large table contains one year of historical data with date and hour columns. Assuming uniform distribution of the data, a query targeting a particular hour would ideally scan 1/8760th of the micro-partitions in the table and then only scan the portion of the micro-partitions that contain the data for the hour column; Snowflake uses columnar scanning of partitions so that an entire partition is not scanned if a query only filters by one column.

In other words, the closer the ratio of scanned micro-partitions and columnar data is to the ratio of actual data selected, the more efficient is the pruning performed on the table.

For time-series data, this level of pruning enables potentially sub-second response times for queries within ranges (i.e. “slices”) as fine-grained as one hour or even less.

Not all predicate expressions can be used to prune. For example, Snowflake does not prune micro-partitions based on a predicate with a subquery, even if the subquery results in a constant.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html#query-pruning>

Question 37: Skipped

Which of the two statements are true about the variant data type in SnowFlake?

- Optimized storage based on repeated elements  
**(Correct)**
- Stored in a seperate file format from structured data
- Can be queried using json path notation  
**(Correct)**
- Requires a custom mapping for each record type

### Explanation

When Snowflake loads semi-structured data, it optimizes how it stores that data internally by automatically discovering the attributes and structure that exist in the data, and using that knowledge to optimize how the data is stored. Snowflake also looks for repeated attributes across records, organizing and storing those repeated attributes separately. This enables better compression and faster access, similar to the way that a columnar database optimizes storage of columns of data.

Question 38: Skipped

What is the recommended approach for making a variant column accessible in a BI tool

- A pre-defined mapping
- A view  
**(Correct)**
- Leveraging a json parser
- BI tool cannot access json

Question 39: Skipped

You can map snowflake to any s3 bucket and can query the data directly as long as the data is in Parquet or ORC format

- TRUE
- FALSE  
**(Correct)**

Question 40: Skipped

The following factors affect data load rates

- Physical location of the stage  
**(Correct)**
- Virtual warehouse RAM
- Gzip compression efficiency  
**(Correct)**
- Thread size

Question 41: Skipped

What are the two mechanisms to detect if new stage file is there in a snowpipe?

- Automating Snowpipe using cloud messaging  
**(Correct)**
- Calling Snowpipe REST endpoints  
**(Correct)**
-

Calling the custom APIs exposed through AWS EKS

## Explanation

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro.html#how-does-snowpipe-work>

Snowpipe loads data from files **as soon as they are available in a stage**. The data is loaded according to the COPY statement defined in a referenced pipe.

A pipe is a named, first-class Snowflake object that contains a COPY statement used by Snowpipe. The COPY statement identifies the source location of the data files (i.e., a stage) and a target table. All data types are supported, including semi-structured data types such as JSON and Avro.

Different mechanisms for detecting the staged files are available:

### Automating Snowpipe using cloud messaging

Automated data loads leverage event notifications for cloud storage to inform Snowpipe of the arrival of new data files to load. Snowpipe copies the files into a queue, from which they are loaded into the target table in a continuous, serverless fashion based on parameters defined in a specified pipe object.

Snowflake currently supports the following storage account types:

Amazon Web Services (AWS)

Amazon S3

Microsoft Azure

Blob storage

Data Lake Storage Gen2 – Supported as a [preview](#) feature.

General-purpose v2

For more information, see [Automating Continuous Data Loading Using Cloud Messaging](#).

### Calling Snowpipe REST endpoints

Your client application calls a public REST endpoint with the name of a pipe object and a list of data filenames. If new data files matching the list are discovered in the stage referenced by the pipe object, they are queued for loading. Snowflake-provided compute resources load data from the queue into a Snowflake table based on parameters defined in the pipe.

Amazon Web Services (AWS)

Amazon S3

Google Cloud Platform

Cloud Storage

Microsoft Azure

Blob storage

Data Lake Storage Gen2 – Supported as a [preview](#) feature.

General-purpose v2

For more information, see [Calling Snowpipe REST Endpoints to Load Data](#).

Question 42: [Skipped](#)

When you load data using Snowpipe, loads are always performed in a single transaction

- FALSE  
**(Correct)**
- TRUE

### Explanation

It is important to know the difference between the two techniques, so not just on this transaction related questions, please also go through the other differences here

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro.html#how-is-snowpipe-different-from-bulk-data-loading>

Transactions

### Bulk data load

Loads are always performed in a single transaction. Data is inserted into table alongside any other SQL statements submitted manually by users.

### Snowpipe

Loads are combined or split into a single or multiple transactions based on the number and size of the rows in each data file. Rows of partially loaded files (based on the ON\_ERROR copy option setting) can also be combined or split into one or more transactions.

Question 43: Skipped

Snowpipe does not guarantee loading of files in the order that they are staged

- FALSE

- TRUE

(Correct)

### Explanation

Very important to remember this

## Load Order of Data Files

For each pipe object, Snowflake establishes a single queue to sequence data files awaiting loading. As new data files are discovered in a stage, Snowpipe appends them to the queue. However, multiple processes pull files from the queue; and so, while Snowpipe generally loads older files first, there is no guarantee that files are loaded in the same order they are staged.

Question 44: Skipped

What command will you run to pause a pipe?

- 1. alter pipe <pipe name> set pipe\_execution\_paused = true;

(Correct)

- 1. alter pipe <pipe name> set pipe\_execution\_paused = stop;

- 1. alter pipe <pipe name> set pipe\_execution\_paused = halt;

### Explanation

Pause the `mypipe` pipe:

```
alter pipe mypipe set pipe_execution_paused = true;
```

<https://docs.snowflake.com/en/sql-reference/sql/alter-pipe.html#examples>

Question 45: Skipped

All of the below are valid executionState of a snowpipe except:

- RUNNING
- STOPPED\_FEATURE\_DISABLED
- STALLED\_EXECUTION\_ERROR

- PAUSED
- STOPPED  
**(Correct)**

### Explanation

`executionState`

Current execution state of the pipe; could be any one of the following:

`RUNNING` (i.e. everything is normal; Snowflake may or may not be actively processing files for this pipe)

`STOPPED_FEATURE_DISABLED`

`STOPPED_STAGE_DROPPED`

`STOPPED_FILE_FORMAT_DROPPED`

`STOPPED_MISSING_PIPE`

`STOPPED_MISSING_TABLE`

`STALLED_COMPILATION_ERROR`

`STALLED_INITIALIZATION_ERROR`

`STALLED_EXECUTION_ERROR`

`STALLED_INTERNAL_ERROR`

`PAUSED`

`PAUSED_BY_SNOWFLAKE_ADMIN`

`PAUSED_BY_ACCOUNT_ADMIN`

Question 46: `Skipped`

How do you set a return value in a task?

- 1. `create task set_return_value`
  2. `warehouse=return_task_wh`
  3. `schedule='1 minute' as`

```
4. call system$set_return_value('The quick brown fox jumps over the lazy dog'
    );
(Correct)
```

- TRUE
  - 1. create task set\_return\_value  
2. warehouse=return\_task\_wh  
3. schedule='1 minute' as  
4. call system\$set\_return\_code('The quick brown fox jumps over the lazy dog')
 ;
  - FALSE
- 1. create task set\_return\_value  
2. warehouse=return\_task\_wh  
3. schedule='1 minute' as  
4. call set\_return\_value('The quick brown fox jumps over the lazy dog');

### Explanation

#### SYSTEM\$SET\_RETURN\_VALUE

Explicitly sets the return value for a task.

In a tree of tasks, a task can call this function to set a return value. Another task that identifies this task as the predecessor task (using the **AFTER** keyword in the task definition) can retrieve the return value set by the predecessor task.

[https://docs.snowflake.com/en/sql-reference/functions/system\\_set\\_return\\_value.html#examples](https://docs.snowflake.com/en/sql-reference/functions/system_set_return_value.html#examples)

Question 47: Skipped

Query load is calculated by dividing the execution time (in seconds) of all queries in an interval by the total time (in seconds) for the interval.

- TRUE
- (Correct)
- FALSE

### Explanation

<https://docs.snowflake.com/en/user-guide/warehouses-load-monitoring.html#how-query-load-is-calculated>

Question 48: Skipped

Resource monitors can be used to control credit usage for the Snowflake-provided warehouses, including the snowpipe warehouse

- TRUE
- (Correct)
- FALSE

### Explanation

Resource monitors provide control over virtual warehouse credit usage; however, you cannot use them to control credit usage for the Snowflake-provided warehouses, including the SNOWPIPE warehouse.

Question 49: Skipped

This snowpipe rest API Fetches a report about ingested files whose contents have been added to table

- `loadHistoryScan and insertReport`  
**(Correct)**
- `insertPipeReport`
- `insertFiles`

### Explanation

Endpoint: `loadHistoryScan`

Fetches a report about ingested files whose contents have been added to table. Note that for large files, this may only be part of the file. This endpoint differs from `insertReport` in that it views the history between two points in time. There is a maximum of 10,000 items returned, but multiple calls can be issued to cover the desired time range.

### Additional explanation

Please note there was a mistake in this question earlier. The earlier question's answer selected only `loadHistoryScan` but for this question both `loadHistoryScan` and `insertReport` are correct. Please see the explanation above to see the difference between `loadHistoryScan` and `insertReport`

Thanks to Rupa who caught this mistake

Question 50: Skipped

To help avoid exceeding the rate limit (error code 429), snowflake recommends relying more heavily on `insertReport` than `loadHistoryScan`

- `TRUE`  
**(Correct)**
- `FALSE`

## Explanation

loadHistoryScan endpoint is rate limited to avoid excessive calls. To help avoid exceeding the rate limit (error code 429), we recommend relying more heavily on `insertReport` than `loadHistoryScan`. When calling `loadHistoryScan`, specify the most narrow time range that includes a set of data loads. For example, reading the last 10 minutes of history every 8 minutes would work well. Trying to read the last 24 hours of history every minute will result in 429 errors indicating a rate limit has been reached. The rate limits are designed to allow each history record to be read a handful of times.

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-rest-apis.html#endpoint-loadhistoryscan>

Question 51: Skipped

Select two true statement related to streams

- Stream itself does *not* contain any table data  
**(Correct)**
- A stream only stores the offset for the source table  
**(Correct)**
- The hidden columns used by a stream does not consume any storage

## Explanation

Note that a stream itself does *not* contain any table data. A stream only stores the offset for the source table and returns CDC records by leveraging the versioning history for the source table. When the first stream for a table is created, a pair of hidden columns are added to the source table and begin storing change tracking metadata. `These columns consume a small amount of storage`. The CDC records returned when querying a stream rely on a combination of the `offset` stored in the stream and the `change tracking metadata` stored in the table.

Question 52: Skipped

You have a Snowflake table which is defined as below

```
CREATE OR REPLACE TABLE FRUITS(FRUIT_NUMBER NUMBER, FRUIT_DESCRIPTION  
VARCHAR, AVAILABILITY VARCHAR);
```

If you would like to convert the fruit\_number column to be a decimal with a certain precision and scale, which command will you run?

- SELECT FRUIT\_NUMBER::DECIMAL(10,5) FROM FRUITS;  
**(Correct)**

- SELECT FRUIT\_NUMBER(DECIMAL(10,5)) FROM FRUITS;
- SELECT FRUIT\_NUMBER AS DECIMAL(10,5) FROM FRUITS;
- SELECT FRUIT\_NUMBER.DECIMAL(10,5) FROM FRUITS;

### Explanation

This question on CAST may come in many forms, but just remember that you can CAST to a specific datatype using '::' or by using the CAST command

**CAST**, 

Converts a value of one data type into another data type. The semantics of CAST are the same as the semantics of the corresponding TO\_  conversion functions. If the cast is not possible, an error is raised. For more details, see the individual TO\_  conversion functions.

The  operator provides alternative syntax for CAST.

I could have also written the query as below

```
SELECT CAST(FRUIT_NUMBER AS DECIMAL(10,5)) FROM FRUITS;
```

Question 53: 

You ran a query in snowflake and went to query history tab. The query history shows you the below columns

1. QueryID
2. SQL TEXT
3. WAREHOUSE NAME
4. WAREHOUSE SIZE
5. SESSION ID
6. START TIME
7. END TIME

Which of the above column will indicate if an compute cost was incurred to run the query?

- WAREHOUSE NAME
- CREDIT
- WAREHOUSE SIZE  
**(Correct)**
- SESSION ID

### Explanation

Anytime a query incurs compute cost, you will see the warehouse size mentioned as shown below

Status	Query ID	SQL Text	User	Warehouse	Clust...	Size	Session ID	Start Time	End Time	Total Duration	Bytes Scanned	Client Info	Row
✓	019755bd...	SELECT CAST(IFRUIT...	JOY...	ETL_WH	1	Small	6215552117...	11:21:10 PM	11:21:11 PM	510ms	1KB	Snowflake UI 202010...	10

Just fo an experiment as below

1. Run the query **SHOW TABLES**

2. Go to query history

Do you see the ware house size? No, because SHOW TABLES is a metadata query and does not incur any compute cost

Question 54: **Skipped**

The FLATTEN command in snowflake has two version. One version uses a join and the other version uses an object keyword. Please select two words that represent the options used with the command?

- OBJECT\_CONSTRUCT
- TABLE  
**(Correct)**
- TRY\_CAST
- LATERAL

(Correct)

### Explanation

An example of the command is as below

```
select * from table(flatten(input => parse_json('[1, ,77]'))) f;
```

Question 55: Skipped

You have two types of named stages, one is external stage and the other one is internal stage. External stage will always require a cloud storage provider

- TRUE  
(Correct)
- FALSE

### Explanation

Ok, this is an easy question. But when you are working for a customer, what will you suggest him? When should we use external stage? Please read below

It is preferred to use internal stage because SNOWFLAKE automatically encrypts the data in internal stage. SNOWFLAKE is responsible for that encryption. If you use external stage, it will be your responsibility to encrypt the data. But does it mean that you will never use external stage. Well no, there may be use cases where the data is not much and is coming from an external stakeholder and you do not want to store that data into your snowflake tables, in such cases go ahead and use external stage. But always do an analysis of the use case in hand

Question 56: Skipped

When cloning a table, if the COPY GRANTS keywords are not included in the CREATE *<object>* statement, then the new object does **not** inherit any explicit access privileges granted on the original table but does inherit any future grants defined for the object type in the schema

- TRUE  
(Correct)
- FALSE

### Explanation

#### General Usage Notes

A clone is writable and is independent of its source (i.e. changes made to the source or clone are not reflected in the other object).

To create a clone, your current role must have the following privilege(s) on the source object:

## Tables

SELECT

## Pipes, Streams, Tasks

OWNERSHIP

## Other objects

USAGE

In addition, to clone a schema or an object within a schema, your current role must have required privileges on the container object(s) for both the source and the clone.

For tables, Snowflake only supports cloning permanent and transient tables; temporary tables cannot be cloned.

For databases and schemas, cloning is recursive:

Cloning a database clones all the schemas and other objects in the database.

Cloning a schema clones all the contained objects in the schema.

However, the following object types are *not* cloned:

External tables

Internal (Snowflake) stages

For databases, schemas, and tables, a clone does *not* contribute to the overall data storage for the object until operations are performed on the clone that modify existing data or add new data, such as:

Adding, deleting, or modifying rows in a cloned table.

Creating a new, populated table in a cloned schema.

Cloning a table replicates the structure, data, and certain other properties (e.g. `STAGE FILE FORMAT`) of the source table. A cloned table does *not* include the load history of the source table. Data files that were loaded into a source table can be loaded again into its clones.

When cloning tables, the `CREATE <object>` command syntax includes the `COPY GRANTS` keywords:

If the COPY GRANTS keywords are not included in the CREATE <object> statement, then the new object does not inherit any explicit access privileges granted on the original table but does inherit any future grants defined for the object type in the schema (using the GRANT <privileges> ... TO ROLE ... ON FUTURE syntax).

If the COPY GRANTS option is specified in the CREATE <object> statement, then the new object inherits any explicit access privileges granted on the original table but does **not** inherit any future grants defined for the object type in the schema.

Question 57: Skipped

Scaling down a virtual warehouse(e.g from a large warehouse to a small one) is an automated process.

- TRUE
- FALSE  
**(Correct)**

### Explanation

To use a different warehouse, you will need to use the below command

```
USE WAREHOUSE <WAREHOUSE_NAME>
```

Question 58: Skipped

When working with a cloned table, you can use the below SQL statements

- SELECT
- DROP
- SHOW
- ALL OF THE ABOVE  
**(Correct)**

### Explanation

Clone table is just another table, only thing is that it shares micro partitions with the table it has been cloned from.

Question 59: Skipped

Select the options that differentiates a Partner Connect partner from a regular partner

- Connect with snowflake through a wizard  
**(Correct)**

- Includes a partner trial account signup  
**(Correct)**
- Can be connected from the WEB UI  
**(Correct)**
- Includes automated role, user and staging database setup  
**(Correct)**
- None of the above

### Explanation

Partner Connect lets you easily create trial accounts with selected Snowflake business partners and integrate these accounts with Snowflake. This feature provides a convenient option for trying additional tools and services, and then adopting the ones that best meet your business needs.

<https://docs.snowflake.com/en/user-guide/ecosystem-partner-connect.html#snowflake-partner-connect>

Question 60: Skipped

To improve the query performance, which of the below techniques can be used in snowflake

- Indexes
- Distribution keys
- Query hints
- Cluster keys/Reclustering  
**(Correct)**

Question 61: Skipped

Which of the following best describes Snowflake's processing engine

- EMR(Elastic Map Reduce)
- Spark Engine
- Presto

- 
- 

Native SQL Database engine  
**(Correct)**

### Explanation

Snowflake's data warehouse is not built on an existing database or "big data" software platform such as Hadoop. The Snowflake data warehouse uses a new SQL database engine with a unique architecture designed for the cloud. To the user, Snowflake has many similarities to other enterprise data warehouses, but also has additional functionality and unique capabilities

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html#key-concepts-architecture>

Question 62: **Skipped**

Which of the below are automatically provided by snowflake compared to other databases?

- 
- 

Installation and Hardware Configurations  
**(Correct)**

- 
- 

Patch releases  
**(Correct)**

- 
- 

Physical Security  
**(Correct)**

- 
- 

Metadata and Collection statistics documentation

Question 63: **Skipped**

Client has ODBC or JDBC available in their system but they do not have snowflake drivers. Client will still be able to connect to snowflake?

- 
- 

TRUE

- 
- 

FALSE  
**(Correct)**

### Explanation

You will need the Snowflake ODBC or JDBC driver to connect to Snowflake, just having JDBC and ODBC will not solve the problem

## JDBC

<https://docs.snowflake.com/en/user-guide/jdbc.html#jdbc-driver>

Snowflake provides a JDBC type 4 driver that supports core JDBC functionality. The JDBC driver must be installed in a 64-bit environment and requires Java 1.8 (or higher).

## ODBC

<https://docs.snowflake.com/en/user-guide/odbc.html#odbc-driver>

Snowflake provides a driver for connecting to Snowflake using ODBC-based client applications.

The ODBC driver has different prerequisites depending on the platform where it is installed. For details, see the individual installation and configuration instructions for each platform.

In addition, different versions of the ODBC driver support the **GET** and **PUT** commands, depending on the cloud service that hosts your Snowflake account:

Amazon Web Services: Version 2.17.5 (and higher)

Google Cloud Platform: Version 2.21.5 (and higher)

Microsoft Azure: Version 2.20.2 (and higher)

Question 64: Skipped

Auto clustering can be switched off at an account level

- FALSE  
**(Correct)**
- TRUE

### Explanation

Auto clustering cannot be switched off at database or account level, it will need to be done at the table level.

Question 65: Skipped

A user can be defaulted to a role which user does not have access to

- TRUE
- FALSE  
**(Correct)**

### Explanation

You will be able to create a user with a default role to which the user does not have access. However, the user will not be able to logon to snowflake if he/she does not

have access to the default role. Hence a user cannot be defaulted to a role which he/she does not have access to

Question 66: Skipped

With respect to Snowflake UI, which of the following is true?

- A single session can be shared between multiple worksheets
- Every worksheet can have a different role, warehouse and a database  
**(Correct)**
- Worksheets cannot have different role, warehouse and database
- Every worksheet has its own session  
**(Correct)**

Question 67: Skipped

A network policy includes values in both allowed and blocked IP address lists, snowflake applies the blocked IP address list first.

- TRUE  
**(Correct)**
- FALSE

### Explanation

When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the **blocked** IP address list first.

Do **not** add `0.0.0.0/0` to the blocked IP address list. `0.0.0.0/0` is interpreted to be “all IPv4 addresses on the local machine”. Because Snowflake resolves this list first, this would block your own access. Also, note that it is not necessary to include this IP address in the allowed IP address list.

<https://docs.snowflake.com/en/user-guide/network-policies.html#managing-account-level-network-policies>

Question 68: Skipped

How are virtual warehouse credits charged?

- per minute
- per second

- per-second, with a 60-second (i.e. 1-minute) minimum:  
**(Correct)**
- per hour

Question 69: Skipped

Tri-Secret Secure option is available in which snowflake edition

- Business critical or higher  
**(Correct)**
- Enterprise Edition
- All editions

Question 70: Skipped

A warehouse can be assigned to a single resource monitor only

- TRUE  
**(Correct)**
- FALSE

### Explanation

Assignment of Resource Monitors

A single monitor can be set at the account level to control credit usage for all warehouses in your account.

In addition, a monitor can be assigned to one or more warehouses, thereby controlling the credit usage for each assigned warehouse. Note, however, that a warehouse can be assigned to only a single resource monitor.

<https://docs.snowflake.com/en/user-guide/resource-monitors.html#assignment-of-resource-monitors>

Question 1: Skipped

Cloud services layer is one of the layers in snowflake. Which services are provided by this layer

- Authentication  
**(Correct)**
- Infrastructure management  
**(Correct)**
- Metadata management  
**(Correct)**
- Query parsing and optimization  
**(Correct)**
- Access control  
**(Correct)**
- Query processing

### Explanation

The cloud services layer is a collection of services that coordinate activities across Snowflake. These services tie together all of the different components of Snowflake in order to process user requests, from login to query dispatch. The cloud services layer also runs on compute instances provisioned by Snowflake from the cloud provider.

Among the services in this layer:

Authentication

Infrastructure management

Metadata management

Query parsing and optimization

Access control

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html#cloud-services>

Question 2: Skipped

Two virtual warehouses can access the same data at the same time without any contention issues

- TRUE  
**(Correct)**

- FALSE

Question 3: Skipped

You are speaking to a CTO of an organization and would like to define the architecture of snowflake to him in two words. What will you use?

- Shared Disk
- Shared Nothing
- Shared Data  
**(Correct)**
- Shared Memory

**Explanation**

Snowflake's architecture is a hybrid of traditional shared-disk database architectures and shared-nothing database architectures. Similar to shared-disk architectures, Snowflake uses a central data repository for persisted data that is accessible from all compute nodes in the data warehouse. But similar to shared-nothing architectures, Snowflake processes queries using MPP (massively parallel processing) compute clusters where each node in the cluster stores a portion of the entire data set locally. This approach offers the data management simplicity of a shared-disk architecture, but with the performance and scale-out benefits of a shared-nothing architecture.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

Question 4: Skipped

Which statements are true about a micro partition in snowflake

- Contiguous unit of storage  
**(Correct)**
- Each micro partition contains between 50 MB and 500 MB of uncompressed data  
**(Correct)**
- Organized in a columnar fashion  
**(Correct)**
-

Micro partitions can be updated using snowflake query

### Explanation

Micro partitions are immutable which means once created they cannot be changed. If a row is updated in snowflake, the micro partition holding the row is copied into a new micro partition and the updated row is inserted in that micro partition. The old micro partition is marked for deletion.

All data in Snowflake tables is automatically divided into micro-partitions, which are contiguous units of storage. Each micro-partition contains between 50 MB and 500 MB of uncompressed data (note that the actual size in Snowflake is smaller because data is always stored compressed). Groups of rows in tables are mapped into individual micro-partitions, organized in a columnar fashion. This size and structure allows for extremely granular pruning of very large tables, which can be comprised of millions, or even hundreds of millions, of micro-partitions.

Snowflake stores metadata about all rows stored in a micro-partition, including:

The range of values for each of the columns in the micro-partition.

The number of distinct values.

Additional properties used for both optimization and efficient query processing.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html>

Question 5: Skipped

What are the available partner categories in snowflake?

- Data integration  
**(Correct)**
- Business Intelligence  
**(Correct)**
- Security and governance  
**(Correct)**
- ML and data science  
**(Correct)**
- Data Warehousing

### Explanation

<https://docs.snowflake.com/en/user-guide/ecosystem.html>

Question 6: Skipped

While sharing data, consumer account and provider account must be in the same region

- TRUE  
**(Correct)**
- FALSE

### Explanation

This is an important point to note. Data sharing works within the same region. However, you can replicate your account to another region and then share data from that replicated account within that account's region. This is also true across cloud platforms.

### Additional explanation

I am adding this based on a question from a student. In the exam if this question comes, please answer 'TRUE'.

*Data Sharing only works within a single region. But, it is possible to replicate your account to another region and then share data from that replicated account within that account's region. This is also true across cloud platforms.*

So, technically it is true that you can share data across accounts in different regions. But that is data replication. Data sharing between accounts is possible only if they are in the same region

Question 7: Skipped

Who can setup a data share

- SECURITYADMIN
- SYSADMIN
- ACCOUNTADMIN or a role granted the CREATE SHARES global privilege  
**(Correct)**

### Explanation

You must use the ACCOUNTADMIN role (or a role granted the CREATE SHARES global privilege) to perform these tasks. For more details about the CREATE SHARES privilege

<https://docs.snowflake.com/en/user-guide/data-sharing-provider.html>

Question 8: Skipped

Select the statements that are true for a data shar

- Shares are named Snowflake objects  
**(Correct)**
-

Share consists of privileges that grant access to the database(s) and the schema containing the objects to share

(Correct)

- Share consists of privileges that grant access to the specific objects in the database
- Share consists of consumer accounts with which the database and its objects are shared

(Correct)

### Explanation

Shares are named Snowflake objects that encapsulate all of the information required to share a database. Each share consists of:

The privileges that grant access to the database(s) and the schema containing the objects to share.

The privileges that grant access to the specific objects in the database.

The consumer accounts with which the database and its objects are shared.

<https://docs.snowflake.com/en/user-guide/data-sharing-intro.html#what-is-a-share>

Question 9: Skipped

Select the statements that are true for a shared database

- Shared databases are read-only
- A clone can be created from a shared database
- Time travel can be done on a shared database
- Can edit the comments for a shared database

### Explanation

Shared databases have the following limitations for consumers:

Shared databases are read-only. Users in a consumer account can view/query data, but cannot insert or update data, or create any objects in the database.

The following actions are not supported:

Creating a clone of a shared database or any schemas/tables in the database.

Time Travel for a shared database or any schemas/tables in the database.

Editing the comments for a shared database.

<https://docs.snowflake.com/en/user-guide/data-share-consumers.html#general-limitations-for-shared-databases>

Question 10: Skipped

What are the product offerings for secure data sharing?

- DIRECT SHARE  
(Correct)
- DATA MARKETPLACE  
(Correct)
- DATA EXCHANGE  
(Correct)
- DATA STORE

### Explanation

Snowflake provides three product offerings for data sharing that utilize [Snowflake Secure Data Sharing](#) to connect providers of data with consumers.

[Direct Share](#)

[Snowflake Data Marketplace](#)

[Data Exchange](#)

<https://docs.snowflake.com/en/user-guide/data-sharing-product-offerings.html#direct-share>

Question 11: Skipped

Virtual warehouses cannot be resized while they are running

- TRUE
- FALSE  
(Correct)

### Explanation

Warehouses can be started and stopped at any time. They can also be resized at any time, even while running, to accommodate the need for more or less compute resources, based on the type of operations being performed by the warehouse.

<https://docs.snowflake.com/en/user-guide/warehouses-overview.html>

Question 12: Skipped

If a transaction is running on a session and the session disconnects abruptly, the transaction will be in a detached state. It will not allow the transaction to commit or roll back. The locks on the resources will also not be released. If you do not do anything and let snowflake abort the transaction, how long will you need to wait?

- A 4 hours  
**(Correct)**
- B 15 minutes
- C 1 hour

**Explanation**

If a transaction is running in a session and the session disconnects abruptly, preventing the transaction from committing or rolling back, the transaction is left in a detached state, including any locks that the transaction is holding on resources. If this happens, you might need to abort the transaction.

To abort a running transaction, the user who started the transaction or an account administrator can call the system function, **SYSTEM\$ABORT\_TRANSACTION**.

If the transaction is left open, Snowflake typically rolls back the transaction after it has been idle for four hours.

<https://docs.snowflake.com/en/sql-reference/transactions.html#aborting-transactions>

Question 13: Skipped

Please look at the below two queries

- A. SELECT \* FROM TABLE1 WHERE YEAR = 2015 AND UPPER(COL1) = 'COLVAL';  
B. SELECT \* FROM TABLE1 WHERE YEAR = 2015 AND COL1 = 'COLVAL';

Which one of these will scan more partitions?

Note: COL1 stores values both in upper and lower case

- A  
**(Correct)**
- B

- Both will scan the same number of partitions
- 

### Explanation

This probably will not come in the certification exam, but it is important to know this when you are working on a project.

Snowflake stores the metadata for each column in a metadata services layer which includes the range of values of the columns in the micro partition. In this case, the range of values stored in the metadata layer for COL1 is stored in its raw form which is a mix of upper and lower case. Because of this when you apply the UPPER function, the metadata services layer cannot use the information stored in the layer to fetch the micro partitions for the range of the rows

Question 14: Skipped

You are using a 2XL warehouse which has 32 nodes and 8 cores in each node. When you submit a query to this warehouse, snowflake will try to split the query into how many partitions? (Provided that the query is eligible for splitting)

- 256  
(Correct)
- 32
- 8

### Explanation

There are total  $8 \times 32 = 256$  cores in 2xl ware house. When a query is submitted to this warehouse snowflake will try to split the query into 256 parts and each of the cores will try to scan 1/256 part of the table.

Please note that this is an ideal situation, the query must be splittable and the warehouse cores must be available to do these parallelization

Question 15: Skipped

What factors influence the credits charged for a warehouse?

- The number of servers per cluster (determined by warehouse size)  
(Correct)
- The number of clusters (if using multi-cluster warehouses)  
(Correct)
-

The size of the query that runs on the warehouse

- The length of time each server in each cluster runs  
**(Correct)**

### Explanation

It is very important to understand this concept

Credit charges are calculated based on:

1. The number of servers per cluster (determined by warehouse size).
2. The number of clusters (if using [multi-cluster warehouses](#)).
3. The length of time each server in each cluster runs.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html#how-are-credits-charged-for-warehouses>

Question 16: Skipped

When creating a warehouse, what are the two most critical factors to consider from cost and performance perspective

- Warehouse size (i.e. number of servers per cluster)  
**(Correct)**
- Manual vs automated management (for starting/resuming and suspending warehouses)  
**(Correct)**
- Cloud provider selection

### Explanation

When creating a warehouse, the two most critical factors to consider, from a cost and performance perspective, are:

1. Warehouse size (i.e. number of servers per cluster).
2. Manual vs automated management (for starting/resuming and suspending warehouses).

Ideally, bigger the warehouse faster your query will run. But this may not be true, if your query itself is not big enough to be split up into enough number of partitions to take advantage of the more number of cores in the warehouse. Please be aware of this.

Setting the auto suspend time is very important. You must calculate this based on the query patterns. If the warehouse is used for adhoc queries and most of the queries finish in less than a minute, it will be a waste of money if your auto suspend is set to 15 minutes. In this case,

even if the query ends in a minute, the warehouse will keep running for 15 minutes before it is suspended

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html#how-are-credits-charged-for-warehouses>

Question 17: Skipped

You want to scale out query performance. Which strategy will you follow with respect to your warehouses?

- Resize the warehouse
- Add clusters to the warehouse  
**(Correct)**
- Create another warehouse

### Explanation

Adding cluster to a multi-cluster warehouse helps scale out if the query load increases. Snowflake does this scale out automatically

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html#scaling-up-vs-scaling-out>

Question 18: Skipped

You need to handle a huge number of concurrent users for your snowflake data warehouse. All of them are from the same team and needs same type of access. Which is the most appropriate warehouse strategy will you recommend?

- Create one warehouse for each user
- Create one warehouse for each group of user
- Use multi-cluster warehouse  
**(Correct)**

### Explanation

Multi-cluster warehouse is the most appropriate answer for this.

Multi-cluster warehouses are designed specifically for handling queuing and performance issues related to large numbers of concurrent users and/or queries. In addition, multi-cluster warehouses can help automate this process if your number of users/queries tend to fluctuate.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html#multi-cluster-warehouses-improve-concurrency>

Question 19: Skipped

Select all the options which are true about the data that is stored in a micro partition metadata

- The range of values for each of the columns in the micro-partition  
**(Correct)**
- Mean and standard deviation of each column
- The number of distinct values  
**(Correct)**
- Additional properties used for both optimization and efficient query processing  
**(Correct)**

### Explanation

Snowflake stores metadata about all rows stored in a micro-partition, including:

1. The range of values for each of the columns in the micro-partition.
2. The number of distinct values.
3. Additional properties used for both optimization and efficient query processing.

<https://docs.snowflake.com/en/user-guide/tables-clustering-micropartitions.html>

Question 20: Skipped

You want to get the clustering information for a table T1 that is clustered on column C1.

Which of the below queries will you run to get the details?

- 1. `select system$clustering_information('T1')`  
**(Correct)**
- 1. `select clustering('T1') from information_schema.clusteringinfo;`
- 1. `Show clustering info for T1;`

### Explanation

Returns clustering information, including average clustering depth, for a table based on one or more columns in the table

[https://docs.snowflake.com/en/sql-reference/functions/system\\_clustering\\_information.html#system-clustering-information](https://docs.snowflake.com/en/sql-reference/functions/system_clustering_information.html#system-clustering-information)

Question 21: Skipped

Each time a warehouse is started or resized to a larger size, the warehouse is billed for 1 minute's worth of usage. After 1 minute, all subsequent billing is per-second

- TRUE  
**(Correct)**

- FALSE

### Explanation

Warehouse credits roll over to the first minute and then the credits are charged per second. What this means is, if your warehouse is resumed at time T0 and it was up for 30 seconds( $T0+30\ SECS$ ), you will still be paying for a minute. If it was up for 62 seconds, then you will be paying for 62 seconds(60 seconds + 2 seconds).

<https://docs.snowflake.com/en/user-guide/credits.html#virtual-warehouse-credit-usage>

Question 22: Skipped

Cloud services also use compute. Does snowflake charge you for cloud services.

- Yes, if it exceeds 10% of the daily usage of the compute resources  
**(Correct)**

- Yes, irrespective of how much compute resources are used

- Snowflake never charges for cloud services

### Explanation

Similar to virtual warehouse usage, [Snowflake credits](#) are used to pay for the usage of the cloud services that exceeds 10% of the daily usage of the compute resources.

Usage for cloud-services is charged only if the daily consumption of cloud services exceeds 10% of the daily usage of the compute resources. The charge is calculated daily (in the UTC time zone). This ensures that the 10% adjustment is accurately applied each day, at the credit price for that day

<https://docs.snowflake.com/en/user-guide/credits.html#cloud-services-credit-usage>

Question 23: Skipped

If you are on Snowflake Enterprise Edition (and higher), the time travel retention period can be set to any value from 0 up to 90 days for databases, schemas, and tables

- TRUE  
**(Correct)**

- FALSE

### Explanation

Data Retention Period

A key component of Snowflake Time Travel is the data retention period.

When data in a table is modified, including deletion of data or dropping an object containing data, Snowflake preserves the state of the data before the update. The data retention period specifies the number of days for which this historical data is preserved and, therefore, Time Travel operations (SELECT, CREATE ... CLONE, UNDROP) can be performed on the data.

The standard retention period is 1 day (24 hours) and is automatically enabled for all Snowflake accounts:

For Snowflake Standard Edition, the retention period can be set to 0 (or unset back to the default of 1 day) at the account and object level (i.e. databases, schemas, and tables).

For Snowflake Enterprise Edition (and higher):

For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also true for temporary tables.

For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days.

<https://docs.snowflake.com/en/user-guide/data-time-travel.html#data-retention-period>

Question 24: Skipped

Time travel can be disabled at the account level

- TRUE
- FALSE  
**(Correct)**

#### Explanation

Time Travel cannot be disabled for an account; however, it can be disabled for individual databases, schemas, and tables by specifying `DATA_RETENTION_TIME_IN_DAYS` with a value of 0 for the object.

<https://docs.snowflake.com/en/user-guide/data-time-travel.html#enabling-and-disabling-time-travel>

-----  
Additional clarification

-----  
In snowflake, in one place it has been written that time-travel cannot be disabled for an account, in another place it has been written that we can specify

DATA\_RETENTION\_TIME\_IN\_DAYS to 0. This part of the documentation in snowflake is little confusing. however what the document is trying to say that Time Travel is enabled by default for all accounts and cannot be taken out. It however can be set to zero at the account level which is actually not the same as disabling it for the account

Question 25: Skipped

What is the encryption algorithm used by snowflake?

- AES 256-bit encryption  
**(Correct)**
- DES
- BLOWFISH

### Explanation

All Snowflake customer data is encrypted by default using the latest security standards and best practices. Snowflake uses strong AES 256-bit encryption with a hierarchical key model rooted in a hardware security module. Keys are automatically rotated on a regular basis by the Snowflake service, and data can be automatically re-encrypted (“rekeyed”) on a regular basis. Data encryption and key management is entirely transparent and requires no configuration or management.

<https://docs.snowflake.com/en/user-guide/security-encryption.html#encryption-key-management>

Question 26: Skipped

Choose the answer which true about key rotation in snowflake

- Account and table master keys are automatically rotated by Snowflake when they are more than 30 days old  
**(Correct)**
- Only account master key is rotated by snowflake when it is required
- Only table master key is rotated when it is required
- Key rotation is an expensive operation hence it is done manually when required

### Explanation

Account and table master keys are automatically rotated by Snowflake when they are more than 30 days old. Active keys are retired, and new keys are created. When Snowflake

determines the retired key is no longer needed, the key is automatically destroyed. When active, a key is used to encrypt data and is available for usage by the originator. When retired, the key is used solely to decrypt data and is only available for usage by the recipient. When wrapping child keys in the key hierarchy, or when inserting data into a table, only the current, active key is used to encrypt data. When a key is destroyed, it is not used for either encryption or decryption. Regular key rotation limits the lifecycle for the keys to a limited period of time.

<https://docs.snowflake.com/en/user-guide/security-encryption.html#encryption-key-rotation>

Question 27: Skipped

Which roles create, alter, or drop network policies?

- ACCOUNTADMIN  
**(Correct)**
- SECURITYADMIN  
**(Correct)**
- SYSADMIN

### Explanation

Only account administrators and security administrators (i.e. users with the ACCOUNTADMIN or SECURITYADMIN role) can create, alter, or drop network policies.

Question 28: Skipped

When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the **blocked** IP address list first

- TRUE  
**(Correct)**
- FALSE

### Explanation

When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the **blocked** IP address list first.

Do **not** add `0.0.0.0/0` to the blocked IP address list. `0.0.0.0/0` is interpreted to be “all IPv4 addresses on the local machine”. Because Snowflake resolves this list first, this would block your own access. Also, note that it is not necessary to include this IP address in the allowed IP address list.

Question 29: Skipped

What is the recommended file size to optimize the number of parallel operations for a load?

- 100-250 MB(or larger) in size *compressed*  
**(Correct)**

- 
- 10 MB to 100 MB in size un *compressed*.
- 
- less than 1 MB compressed

### Explanation

The number of load operations that run in parallel cannot exceed the number of data files to be loaded. To optimize the number of parallel operations for a load, we recommend aiming to produce data files roughly 10 MB to 100 MB in size **compressed**. Aggregate smaller files to minimize the processing overhead for each file. Split larger files into a greater number of smaller files to distribute the load among the servers in an active warehouse. The number of data files that are processed in parallel is determined by the number and capacity of servers in a warehouse. We recommend splitting large files by line to avoid records that span chunks.

If your source database does not allow you to export data files in smaller chunks, you can use a third-party utility to split large CSV files.

<https://docs.snowflake.com/en/user-guide/data-load-considerations-prepare.html#file-sizing-best-practices-and-limitations>

Question 30: Skipped

What is the size limit on the variant data type in snowflake

- 16 MB  
**(Correct)**
- 64 MB
- 32 MB

### Explanation

The VARIANT data type imposes a 16 MB (compressed) size limit on individual rows.

<https://docs.snowflake.com/en/user-guide/data-load-considerations-prepare.html#semi-structured-data-size-limitations>

Question 31: Skipped

The interactions with data are initialized through the services layer

- TRUE  
**(Correct)**
- FALSE

### Explanation

Cloud services tie together all of the different components of Snowflake in order to **process user requests, from login to query dispatch**

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html#cloud-services>

Question 32: Skipped

You have many queries that need to run at the same time on a warehouse and you need high concurrency when they execute. You will go for a multi-cluster warehouse in this case.

- TRUE  
**(Correct)**
- FALSE

### Explanation

This question can be framed in many ways in the exam. But please remember only two things

- 1. Multi Cluster warehouse can scale compute resources**
- 2. The scaling is based on concurrent users and queries**

Multi-cluster warehouses enable you to scale compute resources to manage your user and query concurrency needs as they change, such as during peak and off hours.

<https://docs.snowflake.com/en/user-guide/warehouses-multicloud.html#multi-cluster-warehouses>

Question 33: Skipped

MFA or multi factor authentication is part of snowflake's security and authentication

- TRUE  
**(Correct)**
- FALSE

### Explanation

Here, you need to remember one more thing. MFA is integrated in snowflake and powered by DUO. Also users are not automatically enrolled to MFA, they must enroll themselves

Snowflake supports multi-factor authentication (MFA) to provide increased login security for users connecting to Snowflake. MFA support is provided as an integrated Snowflake feature, powered by the Duo Security service, which is managed completely by Snowflake. Users do not need to separately sign up with Duo or perform any tasks, other than installing the Duo Mobile application, which is supported on multiple smart phone platforms (iOS, Android, Windows, etc.). See the [Duo User Guide](#) for more information about supported platforms/devices and how Duo multi-factor authentication works.

MFA is enabled on a per-user basis; however, at this time, users are not automatically enrolled in MFA. To use MFA, users must enroll themselves.

Question 34: Skipped

Snowflake has administration settings for resource consumption to help with below activities

- Help control costs associated with unexpected warehouse credit usage  
**(Correct)**
- Manage access to snowflake for specific users
- Manage availability of the product

### Explanation

I remembered that I have not included resource monitors in any of the previous practice tests. resource monitor is important for the exam

### What is a resource monitor?

To help control costs and avoid unexpected credit usage caused by running warehouses, Snowflake provides *resource monitors*.

A virtual warehouse consumes Snowflake credits while it runs. Resource monitors can be used to impose limits on the number of credits that are consumed by:

1. User-managed virtual warehouses
2. Virtual warehouses used by cloud services

<https://docs.snowflake.com/en/user-guide/resource-monitors.html#working-with-resource-monitors>

Question 35: Skipped

When a resource limit has hit the limit or is about to hit the limit, what can the resource monitor do

- Suspend the warehouse  
**(Correct)**
- Send alert notifications  
**(Correct)**
- Kill the query that is running  
**(Correct)**

### Explanation

Limits can be set for a specified interval or date range. When these limits are reached and/or are approaching, the resource monitor can trigger various actions, **such as sending alert notifications and/or suspending the warehouses.**

## **DO SUSPEND | SUSPEND\_IMMEDIATE | NOTIFY**

Specifies the action performed by the trigger when the threshold is reached:

**SUSPEND**: Suspend all assigned warehouses while allowing currently running queries to complete. No new queries can be executed by the warehouses until the credit quota for the resource monitor is increased. In addition, this action sends a notification to all users who have enabled notifications for themselves.

**SUSPEND\_IMMEDIATE**: Suspend all assigned warehouses immediately and cancel any currently running queries or statements using the warehouses. In addition, this action sends a notification to all users who have enabled notifications for themselves.

**NOTIFY**: Send an alert (to all users who have enabled notifications for themselves), but do not take any other action.

<https://docs.snowflake.com/en/sql-reference/sql/create-resource-monitor.html#optional-parameters>

Question 36: **Skipped**

Snowflake provides specific administration features and capabilities to support the following activities

- Managing database and warehouses within a snowflake account  
**(Correct)**
- Managing roles and users within a snowflake account  
**(Correct)**
- Monitor usage and manage resources to control costs in a snowflake account  
**(Correct)**
- Manage 3rd party applications providing data to snowflake

### **Explanation**

3rd party applications are not managed by snowflake

Question 37: **Skipped**

What is a point in time snapshot of data which can be updated by users is called

- Snowflakes' cloning  
**(Correct)**
- Time travel
- Fail safe

## Explanation

This is very important to know

### What is cloning in snowflake?

Cloning also referred to as “zero-copy cloning” creates a copy of a database, schema or table. A **snapshot of data** present in the source object is taken when the clone is created and is made available to the cloned object. The cloned object is **writable** and is independent of the clone source. That is, changes made to either the source object or the clone object are not part of the other. Cloning a database will clone all the schemas and tables within that database. Cloning a schema will clone all the tables in that schema. This topic provides examples of cloning tables, databases, and schemas in a script, with explanatory comments.

Question 38: Skipped

Federated authentication in snowflake is compliant with SAML2.0

- TRUE  
**(Correct)**
- FALSE

## Explanation

Federated authentication enables your users to connect to Snowflake using secure SSO (single sign-on). With SSO enabled, your users authenticate through an external, **SAML 2.0-compliant identity provider (IdP)**

<https://docs.snowflake.com/en/user-guide/admin-security-fed-auth.html#federated-authentication-sso>

Question 39: Skipped

All security related information is stored in which layer of snowflake architecture

- Service  
**(Correct)**
- Storage
- Compute
- All of the above

## Explanation

Service(or Cloud service) layer is the one which stores all security related information

The services included in this layer are

1. Authentication

2. Infrastructure management
3. Metadata management
4. Query parsing and optimization
5. Access control

Question 40: Skipped

What is the benefit of client-side encryption in snowflake?

- It provides a secure system for managing data in cloud storage  
**(Correct)**
- It helps to decrypt data faster at a later period
- It reduces the cost of encryption by the cloud provider

#### Explanation

Client-side encryption in general, is the most secure form of managing data on a cloud provider. With client-side encryption, the data is encrypted on the client before it is uploaded. That means, the cloud provider only stores the encrypted version of the data and never sees data in the clear.

Question 41: Skipped

Multi-factor Authentication can be used for connecting to Snowflake via the Snowflake JDBC driver

- TRUE  
**(Correct)**
- FALSE

Question 42: Skipped

Having separate accounts in snowflake enables users to have

- Different editions of snowflake and different regions
- Billing can be done at account level
- Database object deployment between environments become simpler
- All of the above  
**(Correct)**

Question 43: Skipped

Semi structured data strings in snowflake are stored in a column with the below datatype

- VARIANT  
**(Correct)**
- OBJECT
- VARCHAR
- NONE OF THE ABOVE

### Explanation

VARIANT is a tagged universal type, which can store values of any other type, including OBJECT and ARRAY, up to a maximum size of 16 MB compressed.

Question 44: Skipped

When you size a warehouse, which of the below factors are not required to be considered

- Number of users  
**(Correct)**
- Number of concurrent queries
- Number of tables queried
- Data size and composition
- All of the above

### Explanation

Number of concurrent queries will tell you whether to size the warehouse bigger(or smaller) or to use multi-cluster warehouse

Warehouse size will also depend on number of tables queried, data size and composition.

Question 45: Skipped

This property of resource monitor specifies whether the resource monitor is used to monitor the credit usage for your entire Account (i.e. all warehouses in the account) or a specific set of individual warehouses

- CREDIT QUOTA
- MONITOR LEVEL  
**(Correct)**

- SCHEDULE

### Explanation

Please go through all the properties of resource monitor

### Monitor Level

This property specifies whether the resource monitor is used to monitor the credit usage for your entire Account (i.e. all warehouses in the account) or a specific set of individual warehouses.

If this property is not set, the resource monitor doesn't monitor any credit usage. It simply remains dormant.

<https://docs.snowflake.com/en/user-guide/resource-monitors.html#monitor-level>

Question 46: Skipped

Resource monitors support the below actions

- Notify & Suspend  
**(Correct)**
- Notify & Suspend Immediately  
**(Correct)**
- Notify  
**(Correct)**
- Notify and automatically call the account administrator

### Explanation

Resource monitors support the following actions:

#### Notify & Suspend

Send a notification (to all account administrators with notifications enabled) and suspend all assigned warehouses after all statements being executed by the warehouse(s) have completed.

#### Notify & Suspend Immediately

Send a notification (to all account administrators with notifications enabled) and suspend all assigned warehouses immediately, which cancels any statements being executed by the warehouses at the time.

#### Notify

Perform no action, but send an alert notification (to all account administrators with notifications enabled).

<https://docs.snowflake.com/en/user-guide/resource-monitors.html#actions>

Question 47: Skipped

A warehouse can be assigned to only a single resource monitor.

- TRUE  
**(Correct)**
- FALSE

### Explanation

A single monitor can be set at the account level to control credit usage for all warehouses in your account.

In addition, a monitor can be assigned to one or more warehouses, thereby controlling the credit usage for each assigned warehouse. **Note, however, that a warehouse can be assigned to only a single resource monitor.**

Question 48: Skipped

Resource monitor notifications can be received by account administrators through the following means

- Web interface  
**(Correct)**
- Email  
**(Correct)**
- Pager
- Mobile

### Explanation

Notifications can be received by account administrators through the web interface and/or email; however, by default, notifications are *not* enabled:

To receive notifications, each account administrator **must** explicitly enable notifications through their preferences in the web interface.

In addition, if an account administrator chooses to receive email notifications, they must provide a valid email address (and verify the address) before they will receive any emails.

**Question 49:** Skipped

Which access control privileges must be granted on specified resource monitors for roles other than account administrators to view and modify resource monitors as needed

- MONITOR  
**(Correct)**
- CREATE
- MODIFY  
**(Correct)**

**Explanation**

By default, resource monitors can only be created by account administrators and, therefore, can only be viewed and maintained by them.

However, Snowflake supports enabling other users to view and modify resource monitors as needed. This functionality is implemented through the following access control privileges, which can be granted on specified resource monitors to specified roles:

1. MONITOR

2. MODIFY

**Question 50:** Skipped

Which properties of a resource monitor can be modified?

- Increase or decrease the credit quota for the monitor  
**(Correct)**
- If the monitor is monitoring your account, convert it to monitor individual warehouses  
**(Correct)**
- If the monitor is monitoring individual warehouses:
  1. Add or remove warehouses from the list.
  2. Convert it to monitor your account  
**(Correct)**
- Customize the schedule (frequency, start timestamp, and end timestamp) for the monitor  
**(Correct)**
- Add or remove actions, or modify the threshold percentages for existing actions  
**(Correct)**

**Explanation**

This is important to know

You can modify the following properties for an existing resource monitor:

1. Increase or decrease the credit quota for the monitor.
2. If the monitor is monitoring your account, convert it to monitor individual warehouses.
3. If the monitor is monitoring individual warehouses:
  - Add or remove warehouses from the list.
  - Convert it to monitor your account.
4. Customize the schedule (frequency, start timestamp, and end timestamp) for the monitor.
5. Add or remove actions, or modify the threshold percentages for existing actions.

Question 51: Skipped

Select the TABLE function which helps to convert semi-structured data to a relational representation

- FLATTEN  
(Correct)
- CHECK\_JSON
- PARSE\_JSON

#### Explanation

Flattens (explodes) compound values into multiple rows.

FLATTEN is a table function that takes a VARIANT, OBJECT, or ARRAY column and produces a lateral view (i.e. an inline view that contains correlation referring to other tables that precede it in the **FROM** clause).

FLATTEN can be used to convert semi-structured data to a relational representation.

<https://docs.snowflake.com/en/sql-reference/functions/flatten.html#flatten>

Question 52: Skipped

Output of a FLATTEN query has the below columns. Select three.

- SEQ  
(Correct)

- PATH  
**(Correct)**
- VALUE  
**(Correct)**
- RECORD\_METADATA

### Explanation

The returned rows from a FLATTEN table function consist of a fixed set of columns:

SEQ	KEY	PATH	INDEX	VALUE	THIS

#### SEQ

A unique sequence number associated with the input record; the sequence is not guaranteed to be gap-free or ordered in any particular way.

#### KEY

For maps or objects, this column contains the key to the exploded value.

#### PATH

The path to the element within a data structure which needs to be flattened.

#### INDEX

The index of the element, if it is an array; otherwise NULL.

#### VALUE

The value of the element of the flattened array/object.

#### THIS

The element being flattened (useful in recursive flattening).

Question 53: **Skipped**

Micropartitions can be modified after creation of the table

- TRUE
- FALSE  
**(Correct)**

### Explanation

In snowflake data is physically stored into micro-partitions in proprietary columnar format. As used you do not have direct access to micropartitions. Snowflake manages the micropartitions. Also micro partitions are immutable which means that once a micro partition is created, it cannot be changed. When a row is updated, the micro partition containing that row is copied to a new micro partition and then the updated row is inserted into that partition. The old micro-partition is marked for deletion.

As used, you however have control to create clustering on your table which helps to re-arrange the micro-partitions based on the cluster keys for faster query retrieval.

Question 54: Skipped

You want to update the value of a specific column in a table. But the updation is not based on any specific condition. You just need to ensure that 80% of the rows in the table are updated with the new value. What snowflake feature will you use

- SAMPLE  
**(Correct)**
- CONCAT
- AVERAGE
- MAX

### Explanation

#### SAMPLE / TABLESAMPLE

Returns a subset of rows sampled randomly from the specified table. The following sampling methods are supported:

Sample a fraction of a table, with a specified probability for including a given row. The number of rows returned depends on the size of the table and the requested probability. A seed can be specified to make the sampling deterministic.

Sample a fixed, specified number of rows. The exact number of specified rows is returned unless the table contains fewer rows.

SAMPLE and TABLESAMPLE are synonymous and can be used interchangeably.

#### Also try a hands on as below

```
CREATE OR REPLACE TABLE FRUITS(FRUIT_NUMBER NUMBER, FRUIT_DESCRIPTION VARCHAR,  
AVAILABILITY VARCHAR);
```

```

INSERT INTO FRUITS VALUES(1,'APPLE',NULL);
INSERT INTO FRUITS VALUES(2,'MANGO',NULL);
INSERT INTO FRUITS VALUES(3,'ORANGE',NULL);
INSERT INTO FRUITS VALUES(4,'KIWI',NULL);
INSERT INTO FRUITS VALUES(5,'MELON',NULL);
INSERT INTO FRUITS VALUES(6,'PEARS',NULL);
INSERT INTO FRUITS VALUES(7,'BANANA',NULL);
INSERT INTO FRUITS VALUES(8,'PLUM',NULL);
INSERT INTO FRUITS VALUES(9,'APRICOT',NULL);
INSERT INTO FRUITS VALUES(10,'WATERMELON',NULL);

SET COUNT=(SELECT COUNT(*)*0.8 FROM FRUITS);

UPDATE FRUITS SET AVAILABILITY = 'YES'
FROM (SELECT FRUIT_NUMBER FROM FRUITS SAMPLE($COUNT ROWS)) T1
WHERE FRUITS.FRUIT_NUMBER = T1.FRUIT_NUMBER;

```

Question 55: Skipped

Select two true statements for FLATTEN ?

- Takes a VARIANT, OBJECT, or ARRAY column and produces a lateral view  
**(Correct)**
- Can be used to convert semi-structured data to a relational representation  
**(Correct)**
- Cannot be used on permanent tables

### Explanation

#### FLATTEN

Flattens (explodes) compound values into multiple rows.

FLATTEN is a table function that **takes a VARIANT, OBJECT, or ARRAY column and produces a lateral view** (i.e. an inline view that contains correlation referring to other tables that precede it in the **FROM** clause).

**FLATTEN can be used to convert semi-structured data to a relational representation.**

Question 56: Skipped

When migrating from Teradata to snowflake, you will also need to migrate the indexes.

- TRUE
- FALSE  
**(Correct)**

### Explanation

Snowflake has an optimizer built from the ground up and architected for MPP and the cloud. Snowflake understands parallel execution plans and automatically optimizes them. This means **there are no indexes in Snowflake** so you don't have to migrate Indexes from Teradata

Question 57: Skipped

Choose the continuous data loading options from below

- SNOWPIPE  
**(Correct)**
- SNOWFLAKE CONNECTOR FOR KAFKA  
**(Correct)**
- Third-party data integration tools  
**(Correct)**
- BULK COPY

### Explanation

#### Continuous data loading

Options for continuous data loading include the following:

1. Snowpipe
2. Snowflake Connector for Kafka
3. Third-party data integration tools

Question 58: Skipped

A *stream* object records the delta of change data capture (CDC) information for a table (such as a staging table), including inserts and other data manipulation language (DML) changes.

- TRUE  
**(Correct)**
- FALSE

## Explanation

### Change data tracking

A *stream* object records the delta of change data capture (CDC) information for a table (such as a staging table), including inserts and other data manipulation language (DML) changes. A stream allows querying and consuming a set of changes to a table, at the row level, between two transactional points of time.

In a continuous data pipeline, table streams record when staging tables and any downstream tables are populated with data from business applications using continuous data loading and are ready for further processing using SQL statements.

For more information, see [Change Tracking Using Table Streams](#).

Question 59: Skipped

Snowpipe continuously loads micro-batches of data from an external stage location (Amazon S3, Google Cloud Storage, or Microsoft Azure) into a staging table.

- TRUE  
**(Correct)**
- FALSE

## Explanation

### Snowpipe

Snowpipe continuously loads micro-batches of data from an external stage location (Amazon S3, Google Cloud Storage, or Microsoft Azure) into a staging table.

Question 60: Skipped

The Kafka connector continuously loads records from one or more Apache Kafka topics into an internal (Snowflake) stage and then into a staging table using Snowpipe.

- TRUE  
**(Correct)**
- FALSE

## Explanation

### Snowflake Connector for Kafka

The Kafka connector continuously loads records from one or more Apache Kafka topics into an internal (Snowflake) stage and then into a staging table using Snowpipe.

Question 61: Skipped

The amount of storage used by failsafe can be viewed under Account->Billing & Usage section

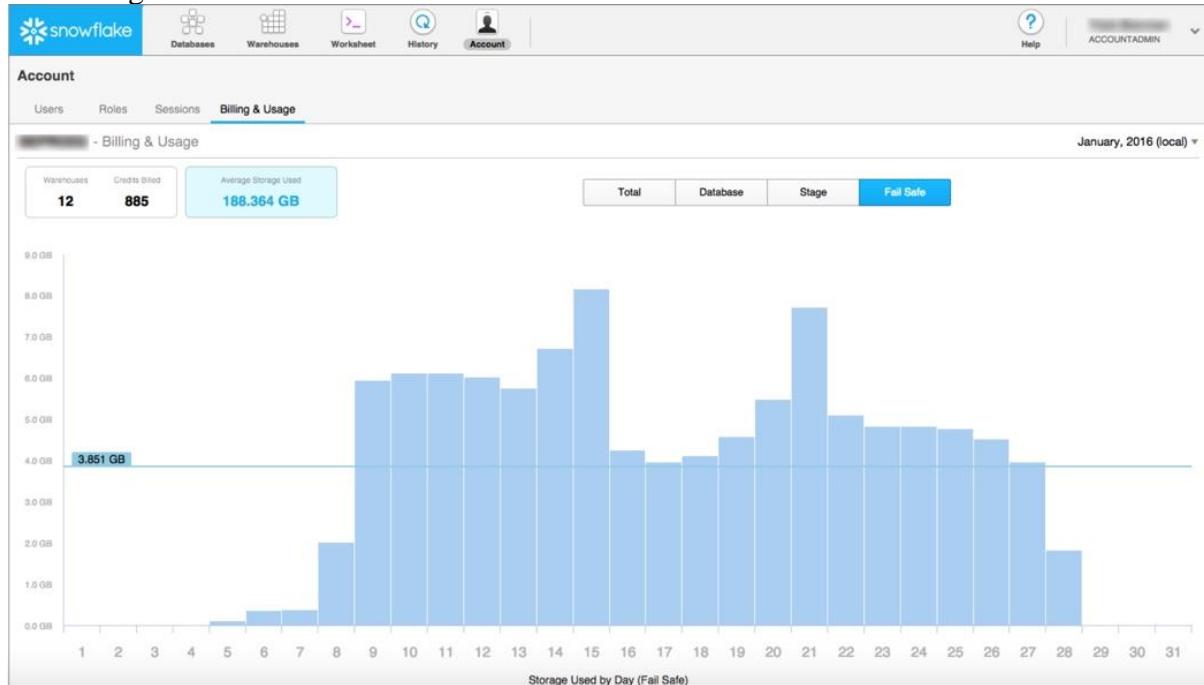
- TRUE

(Correct)

- FALSE

### Explanation

In the web interface, account administrators can view the total data storage for their account, including historical data in Fail-safe. The page also provides a filter for viewing only Fail-safe storage:



Question 62: Skipped

A stream stores data in the same shape as the source table (i.e. the same column names and ordering). What additional columns does it store?

- METADATA\$ACTION  
**(Correct)**
- METADATA\$ISUPDATE  
**(Correct)**
- METADATA\$ROW\_ID  
**(Correct)**
- METADATA\$STREAM\_ID

### Explanation

Stream Columns

A stream stores data in the same shape as the source table (i.e. the same column names and ordering) with the following additional columns:

## METADATA\$ACTION

Indicates the DML operation (INSERT, DELETE) recorded.

## METADATA\$ISUPDATE

Indicates whether the operation was part of an UPDATE statement. Updates to rows in the source table are represented as a pair of DELETE and INSERT records in the stream with a metadata column METADATA\$ISUPDATE values set to TRUE.

Note that streams record the differences between two offsets. If a row is added and then updated in the current offset, the delta change is a new row. The METADATA\$ISUPDATE row records a FALSE value.

## METADATA\$ROW\_ID

Specifies the unique and immutable ID for the row, which can be used to track changes to specific rows over time.

Question 63: Skipped

COPY transformations are supported only for named stages(both internal and external) and user stages

- TRUE  
**(Correct)**
- FALSE

### Explanation

Supported Stages

Only named stages (internal or external) and user stages are supported for COPY transformations.

<https://docs.snowflake.com/en/user-guide/data-load-transform.html#usage-notes>

Question 64: Skipped

The PIPE\_EXECUTION\_PAUSED parameter can be set at which levels

- DATABASE
- ACCOUNT  
**(Correct)**
- SCHEMA  
**(Correct)**

- PIPE
- (Correct)

### Explanation

#### Transferring Pipe Ownership

Complete the following steps to transfer ownership of a pipe:

Set the `PIPE_EXECUTION_PAUSED` parameter to TRUE.

This parameter enables pausing or resuming a pipe. The parameter is supported at the following levels:

Account

Schema

Pipe

At the pipe level, the object owner (or a parent role in a role hierarchy) can set the parameter to pause or resume an individual pipe.

An account administrator (user with the ACCOUNTADMIN role) can set this parameter at the account level to pause or resume all pipes in the account. Likewise, a user with the MODIFY privilege on the schema can pause or resume pipes at the schema level. Note that this larger domain control only affects pipes for which the parameter was not already set at a lower level; e.g., by the owner at the object level.

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-manage.html#transferring-pipe-ownership>

Question 65: Skipped

You want to delete the snowpipe file loading metadata. You decided to truncate the table to be able to do this. This will remove the metadata from the pipe

- TRUE
  - FALSE
- (Correct)

### Explanation

#### Unable to Reload Modified Data, Modified Data Loaded Unintentionally

Snowflake uses file loading metadata to prevent reloading the same files (and duplicating data) in a table. Snowpipe prevents loading files with the same name even if they were later modified (i.e. have a different eTag).

The file loading metadata is associated with the **pipe object** rather than the table. As a result:

Staged files with the same name as files that were already loaded are ignored, even if they have been modified, e.g. if new rows were added or errors in the file were corrected.

Truncating the table using the TRUNCATE TABLE command does **not** delete the Snowpipe file loading metadata.

However, note that pipes only maintain the load history metadata for **14 days**.

Question 66: Skipped

Which cache type gets purged in a predictable way so that it ends up empty of all cached information?

- Warehouse Cache  
**(Correct)**
- Results Cache
- Metadata Cache

### Explanation

Lets investigate why this answer is Warehouse cache

First what is metadata cache, this is the cache which is in the global service layer and stores Object Information and statistics of the objects(like count, max and min). As a user you do not have any control on this cache

Results Cache - This stores the exact results from the exact queries and last for 24 hours. But any time the query is run, this time gets extended by another 24 hours so there is no way you can predict when it will be empty

Warehouse cache - This contains the raw data from the table, this is cached in the SSD disk attached to the warehouse. When a warehouse is suspended, this data is lost and you completely lose the cache.

Now you see why the warehouse cache is a predictable way to empty the cache

**Question 67:** Skipped

Which cache runs on a 24 hour clock?

- Result Cache  
**(Correct)**
- Metadata cache
- Warehouse cache

**Question 68:** Skipped

You are a snowflake architect and you are investigating a long running query to optimize the query. You want to know if the query retrieved data from long-term centralized storage. If you are looking at the Query History table in the history area of the WebUI, where will you look for this information?

- Look in the Bytes Scanned Column for a green bar  
**(Correct)**
- Look in the Bytes scanned column for the acronym LTCS
- Look in the Rows column for a red bar
- Look in the Rows column for the acronym LTCS

**Explanation**

Status	Query ID	SQL Text	U...	Warehouse	Clust...	Size	Session ID	Start Time	End Time	Total Duration	Bytes Scanned	Client Info	Rows
✓	0197e598...	SELECT * FROM inve...	J...	ETL_WH	1	Small	6215552118...	1:08:46 PM	1:09:44 PM	58.0s	3.9GB	Snowflake UI 202010...	

**Question 69:** Skipped

Which is the preferred way to store a generated identity and access management (IAM) entity for external cloud storage

- Access Token
- Stage
- Storage integration  
**(Correct)**
- File Format

**Explanation**

A storage integration is a Snowflake object that stores a generated identity and access management (IAM) entity for your external cloud storage, along with an optional set of

allowed or blocked storage locations (Amazon S3, Google Cloud Storage, or Microsoft Azure). Cloud provider administrators in your organization grant permissions on the storage locations to the generated entity. This option allows users to avoid supplying credentials when creating stages or when loading or unloading data.

A single storage integration can support multiple external stages. The URL in the stage definition must align with the storage location specified for the STORAGE\_ALLOWED\_LOCATIONS parameter.

<https://docs.snowflake.com/en/sql-reference/sql/create-storage-integration.html#create-storage-integration>

Question 70: Skipped

What is the default option for ON\_ERROR clause of COPY command for bulk loading?

- CONTINUE**
- SKIP\_FILE**
- SKIP\_FILE\_10**
- ABORT\_STATEMENT**  
**(Correct)**

### Explanation

```
ON_ERROR = CONTINUE | SKIP_FILE | SKIP_FILE_num | SKIP_FILE_num% |  
ABORT_STATEMENT
```

### Definition

String (constant) that specifies the action to perform when an error is encountered while loading data from a file:

### Values

Supported Values

Notes

**CONTINUE**

Continue loading the file. The COPY statement returns an error message for a maximum of one error encountered per data file. Note that the difference between the ROWS\_PARSED and ROWS\_LOADED column values represents the number of rows that include detected errors. However, each of these rows could include multiple errors. To view all errors in the data files, use the VALIDATION\_MODE parameter or query the **VALIDATE** function.

### **SKIP\_FILE**

Skip file if any errors encountered in the file.

**SKIP\_FILE\_num** (e.g. **SKIP\_FILE\_10**)

Skip file when the number of errors in the file is equal to or exceeds the specified number.

**SKIP\_FILE\_num%** (e.g. **SKIP\_FILE\_10%**)

Skip file when the percentage of errors in the file exceeds the specified percentage.

### **ABORT\_STATEMENT**

Abort the load operation if any error is encountered in a data file. Note that the load operation is **not** aborted if the data file cannot be found (e.g. because it does not exist or cannot be accessed).

## **Default**

**Bulk loading using COPY :** **ABORT\_STATEMENT**

**Snowpipe:** **SKIP\_FILE**

<https://docs.snowflake.com/en/sql-reference/sql/copy-into-table.html#copy-options-copyoptions>

Question 71: **Skipped**

Select the limitations with materialized views(Select 3)

- Snowflake’s “Time Travel” feature is not supported on materialized views  
**(Correct)**
- using CURRENT\_TIME or CURRENT\_TIMESTAMP is not permitted  
**(Correct)**
- A materialized view cannot include Window functions  
**(Correct)**
- Clustering keys cannot be defined in Materialized views

## **Explanation**

<https://docs.snowflake.com/en/user-guide/views-materialized.html#limitations-on-creating-materialized-views>

The following limitations apply to creating materialized views:

A materialized view can query only a single table.

Joins, including self-joins, are not supported.

A materialized view cannot query:

1. A materialized view.
2. A non-materialized view.
3. A UDTF (user-defined table function).

A materialized view cannot include:

1. UDFs (this limitation applies to all types of user-defined functions, including external functions).
2. Window functions.
3. HAVING clauses.
4. ORDER BY clause.
5. LIMIT clause.
6. GROUP BY keys that are not within the SELECT list. All GROUP BY keys in a materialized view must be part of the SELECT list.
7. GROUP BY GROUPING SETS.
8. GROUP BY ROLLUP.
9. GROUP BY CUBE.
10. Nesting of subqueries within a materialized view.

**Many aggregate functions are not allowed in a materialized view definition.**

The aggregate functions that are *supported* in materialized views are:

[APPROX\\_COUNT\\_DISTINCT \(HLL\)](#).

[AVG](#) (except when used in [PIVOT](#)).

[BITAND\\_AGG](#).

[BITOR\\_AGG](#).

`BITXOR_AGG.`

`COUNT.`

`MIN.`

`MAX.`

`STDDEV.`

`STDDEV_POP.`

`STDDEV_SAMP.`

`SUM.`

`VARIANCE (VARIANCE_SAMP, VAR_SAMP).`

`VARIANCE_POP (VAR_POP).`

The other aggregate functions are ***not supported*** in materialized views.

### Note

Aggregate functions that are allowed in materialized views still have some restrictions:

Aggregate functions cannot be nested.

Aggregate functions used in complex expressions (e.g. `(sum(salary)/10)`) can only be used in the outer-most level of a query, not in a subquery or an in-line view.

For example, the following is allowed:

```
create materialized view mv1 as
    select
        sum(x) + 100
    from t;
```

The following is ***not*** allowed:

```
create materialized view mv2 as
    select
        y + 10
    from (
        select
            sum(x) as y
```

```
    from t  
);
```

DISTINCT cannot be combined with aggregate functions.

In a materialized view, the aggregate functions **AVG**, **COUNT**, **MIN**, **MAX**, and **SUM** can be used as aggregate functions but not as window functions. In a materialized view, these functions cannot be used with the **OVER** clause:

```
over ( [ partition by <expr1> ] [ order by <expr2> ] )
```

If an aggregate function is in a subquery, the materialized view cannot create an expression on top of the aggregated column(s) from that subquery. For example, consider the following materialized view definition:

```
create or replace materialized view mv1 as  
select c1 + 10 as c1new, c2  
from (select sum(c1) as c1, c2 from t group by c2);
```

The expression “ $c1 + 10$ ” is an expression on top of an aggregate function in a subquery, and therefore causes an error message.

Note that even an equality operator counts as an expression, which means that **CASE** expressions using columns that represent aggregate functions in a subquery are also prohibited.

To work around this limitation, create a materialized view without the expression, and then create a non-materialized view that includes the expression, for example:

```
create or replace materialized view mv1 as  
select c1, c2  
from (select sum(c1) as c1, c2 from t group by c2);  
  
create or replace view expr_v1 as  
select c1 + 10 as c1new, c2  
from (select c1, c2 from mv1);
```

Functions used in a materialized view must be deterministic. For example, using **CURRENT\_TIME** or **CURRENT\_TIMESTAMP** is not permitted.

A materialized view should not be defined using a function that produces different results for different settings of parameters, such as the session-level parameter **TIMESTAMP\_TYPE\_MAPPING**.

**Question 72:** Skipped

Which of the below columns are usually a good choice for clustering key.

- UUID column in a Customer fact table
- Gender with Male/Female values in a multi-terabyte table
- Timestamp column in a 10TB order fact table
- Store\_Id in a 2 TB Item Store Sales fact table  
**(Correct)**

**Explanation**

What is the reason behind this?

Remember that a cluster key should not have a very high cardinality nor it should have a very low cardinality. If it has high cardinality, there will be lot of partitions that will be created, with low cardinality very few partitions will be created and hence query performance will not be able take advantage of partition pruning. We need to take a middle ground.

**Question 73:** Skipped

An empty table with no micro-partition will have a clustering depth of

- 1
- 0  
**(Correct)**
- 1
- 99

**Explanation**

The clustering depth for a populated table measures the average depth (1 or greater) of the overlapping micro-partitions for specified columns in a table. The smaller the average depth, the better clustered the table is with regards to the specified columns.

Clustering depth can be used for a variety of purposes, including:

1. Monitoring the clustering “health” of a large table, particularly over time as DML is performed on the table.
2. Determining whether a large table would benefit from explicitly defining a [clustering key](#).

**A table with no micro-partitions (i.e. an unpopulated/empty table) has a clustering depth of  0.**

Question 74: Skipped

Why would you not recommend a high cardinality column for including in the clustering key? (Select 2)

- Extremely high cardinality column results in skewed distribution across micro-partitions and hence will impact query pruning
- High cardinality will lead to low clustering depth, so as you add more data query performance will degrade
- Cluster maintenance with high cardinality keys is expensive  
**(Correct)**
- Point lookup queries will only benefit, whereas range queries will not be effectively utilize the clustering key  
**(Correct)**

### Explanation

In general, if a column (or expression) has higher cardinality, then maintaining clustering on that column is more expensive.

The cost of clustering on a unique key might be more than the benefit of clustering on that key, especially if point lookups are not the primary use case for that table.

Question 75: Skipped

Temporary tables cannot be created with a clustering key defined

- TRUE
- FALSE  
**(Correct)**

### Explanation

Run the below command in snowflake and see if you can create it

```
CREATE TEMPORARY TABLE TEMP_TABLE(ITEM_ID INT) CLUSTER BY(ITEM_ID);
```

Question 76: Skipped

In which of the below scenarios, you will use an external table.

- You have data on the cloud providers's object store(AWS/GCS/AZURE Blob) but the data cannot be copied or moved to any other location due to compliance regulations  
**(Correct)**

- You have high volume of data on the cloud providers's object store(AWS/GCS/AZURE Blob) but only a part of the data is needed in snowflake  
**(Correct)**
- You have data on the cloud providers's object store(AWS/GCS/AZURE Blob) which needs to be updated by snowflake after applying required transformations
- You have XML data on cloud provider's object store

### Explanation

Please use elimination method to answer this questions

1. External tables are read-only, hence #3 is not correct
2. External tables do not support XML and hence #4 is not correct

Question 77: **Skipped**

What is the recommended maximum data size limitations for Parquet files?

- 100 MB
- 10 MB
- 1 GB  
**(Correct)**
- 3 GB

### Explanation

Parquet Data Size Limitations

Currently, data loads of large Parquet files (e.g. greater than 3 GB) could time out. Split large files into files 1 GB in size (or smaller) for loading.

Question 78: **Skipped**

Integration object is a first-class database object

- True
- False  
**(Correct)**

### Explanation

Integration object is an account level object and not a database object

Question 79: **Skipped**

Which of the below function returns the name of the warehouse of the current session?

- ACTIVE\_WAREHOUSE()
- SESSION\_WAREHOUSE()
- CURRENT\_WAREHOUSE()  
**(Correct)**
- WAREHOUSE()

#### Explanation

Run the below query to see it yourself

```
SELECT CURRENT_WAREHOUSE();
```

Question 80: **Skipped**

Warehouse cache size changes with warehouse size

- True  
**(Correct)**
- False

#### Explanation

The size of the cache is determined by the number of servers in the warehouse (i.e. the larger the warehouse and, therefore, the number of servers in the warehouse), the larger the cache.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html#how-does-warehouse-caching-impact-queries>

Question 81: **Skipped**

A task can execute multiple SQL statement.

- TRUE
- FALSE  
**(Correct)**

#### Explanation

Currently, a task can execute a single SQL statement, including a call to a stored procedure.

Tasks can be combined with **table streams** for continuous ELT workflows to process recently changed table rows. Streams ensure exactly once semantics for new or changed data in a table.

Tasks can also be used independently to generate periodic reports by inserting or merging rows into a report table or perform other periodic work.

#### Question 82: Skipped

What are the recommendations from Snowflake to effectively and efficiently use a warehouse.

- Experiment with different types of queries and different warehouse sizes to determine the combinations that best meet your specific query needs and workload  
**(Correct)**
- Don't focus on warehouse size. Snowflake utilizes per-second billing, so you can run larger warehouses (Large, X-Large, 2X-Large, etc.) and simply suspend them when not in use  
**(Correct)**
- Be very careful on what size of warehouse you are picking up. The credits charged are based on number of queries running on the warehouse and the size of the warehouse

#### Explanation

The keys to using warehouses effectively and efficiently are:

1. Experiment with different types of queries and different warehouse sizes to determine the combinations that best meet your specific query needs and workload.
2. Don't focus on warehouse size. Snowflake utilizes per-second billing, so you can run larger warehouses (Large, X-Large, 2X-Large, etc.) and simply suspend them when not in use.

<https://docs.snowflake.com/en/user-guide/warehouses-considerations.html#warehouse-considerations>

Remember one more thing. The credits are not dependent on number of queries running in your warehouse, it depends on how long your warehouse is active.

#### Question 83: Skipped

In a stream, if a row is added and then updated in the current offset, how will the delta change be represented

- The delta change is a new row  
**(Correct)**
- The delta change will be a pair of delete and insert record
- The delta change will be a pair of insert and update record

#### Explanation

**METADATA\$ISUPDATE**

Indicates whether the operation was part of an UPDATE statement. Updates to rows in the source table are represented as a pair of DELETE and INSERT records in the stream with a metadata column METADATA\$ISUPDATE values set to TRUE.

Note that streams record the differences between two offsets. If a row is added and then updated in the current offset, the delta change is a new row. The METADATA\$ISUPDATE row records a FALSE value.

In the above highlighted, please also note how METADATA\$ISUPDATE will be handled, it will record a FALSE value. This may be asked in the exam.

Question 84: Skipped

What are the two types of snowflake data listings which define how data is shared and consumed via data marketplace?

- Standard Data Listings  
**(Correct)**
- Personalized Data Listings  
**(Correct)**
- Business Critical Data Listings
- Enterprise Data Listings

## Explanation

### Types of Data Listings

The Snowflake Data Marketplace offers two types of data listings which define how data is shared and consumed.

Once published, both types of listings are displayed to consumers in the selected regions. The key difference between standard and personalized listings is the ability to access the data share. Personalized listings allow you to control which consumers can access the data.

### Standard Data Listings

A standard listing provides instant access to a published data set. This type of listing is best for providing generic, aggregated, or non-customer-specific data. Each listing includes details about the shared data, sample queries, and information about the data provider.

To access data from a standard data listing, consumers click **Get**, opening a prompt where they name the database for Snowflake users, permission access among Snowflake users, and agree to the provider's terms of use and Snowflake's consumer terms. Clicking **Create**

**Database** immediately creates the database inside their Snowflake account in the Databases tab.

Only users that have a role with the IMPORT SHARE privilege can create databases from the Snowflake Data Marketplace.

### Personalized Data Listings

A personalized listing allows customers to request specific data sets.

This can be premium data that a provider charges for or data that is specific to each consumer. Each listing includes details about the shared data, sample usage examples, and information about the data provider.

To access data from a personalized listing, consumers must submit a request by clicking the **Request** button, and provide their contact information. Once a request is submitted, the data provider is notified. The provider then contacts the consumer.

Each data provider can have different commercial terms. Once those are agreed to, then the personalized data set is created and shared with the consumer.

Question 85: Skipped

Snowflake is a SaaS offering

- TRUE  
(Correct)
- FALSE

### Explanation

SaaS is software as a service. So, snowflake is not IaaS like Amazon Redshift since you do not need to manage infrastructure, it is not a PaaS as well as it is not a platform offered by Snowflake. It is a data warehousing software offered as a service on all three leading cloud providers(Google, AWS, Azure)

Question 86: Skipped

How will you check if a share is inbound or outbound?

- Execute SHOW SHARES and look at kind column  
(Correct)
- Execute describe share and look at location column
- Execute DISPLAY SHARE and look at kind column

### Explanation

Use **SHOW SHARES** to confirm the share. The output of the command lists the details of the share. The **kind** column indicates that the share is OUTBOUND, meaning this share is

sharing a database with other Snowflake accounts. The `to` column lists all accounts to which the share has been made available:

Question 87: Skipped

A share cannot be shared

- TRUE  
**(Correct)**
- FALSE

Question 88: Skipped

What are the two types of data consumer accounts available in snowflake

- FULL ACCOUNT  
**(Correct)**
- READER ACCOUNT  
**(Correct)**
- SHARED ACCOUNT
- QUERY ONLY ACCOUNT

### Explanation

Reader Accounts

Data sharing is only supported between Snowflake accounts. As a data provider, you might wish to share data with a consumer who does not already have a Snowflake account and/or is not ready to become a licensed Snowflake customer.

To facilitate sharing data with these consumers, Snowflake supports providers creating reader accounts. Reader accounts (formerly known as “read-only accounts”) provide a quick, easy, and cost-effective way to share data without requiring the consumer to become a Snowflake customer.

Users in a `reader account` can query data that has been shared with it, but cannot perform any of the DML tasks that are allowed in a `full account (data loading, insert, update, etc.).`

Question 89: Skipped

What is the maximum number of child tasks can a task have?

-

50

- 100  
**(Correct)**
- 1000
- 10

### Explanation

A task is limited to a single predecessor task; however, **a task can have a maximum of 100 child tasks** (i.e. other tasks that identify the task as a predecessor); in addition, a simple tree of tasks is limited to a maximum of 1000 tasks **total** (including the root task) in either resumed (i.e. in a ‘Started’ state) or suspended.

For example, task T2 is limited to a single predecessor task, e.g. T1; however, T1 can serve as the predecessor task for tasks T2, T3, T4, etc.

Read the task chapter thoroughly

<https://docs.snowflake.com/en/sql-reference/sql/create-task.html>

Question 90: Skipped

A simple tree of tasks is limited to a maximum of 1000 tasks **total**

- TRUE  
**(Correct)**
- FALSE

### Explanation

A simple tree of tasks is limited to a maximum of **1000 tasks total** (including the root task) in either resumed (i.e. in a ‘Started’ state) or suspended.

Question 91: Skipped

What are the three features provided by Snowflake for continuous data pipeline

- Continuous data loading  
**(Correct)**
- Change data tracking  
**(Correct)**
- Recurring tasks

(Correct)

- Cloud functions

### Explanation

Snowflake provides the following features to enable continuous data pipelines:

#### Continuous data loading

Options for continuous data loading include the following:

[Snowpipe](#)

[Snowflake Connector for Kafka](#)

Third-party data integration tools

#### Change data tracking

A *stream* object records the delta of change data capture (CDC) information for a table (such as a staging table), including inserts and other data manipulation language (DML) changes. A stream allows querying and consuming a set of changes to a table, at the row level, between two transactional points of time.

In a continuous data pipeline, table streams record when staging tables and any downstream tables are populated with data from business applications using continuous data loading and are ready for further processing using SQL statements.

For more information, see [Change Tracking Using Table Streams](#).

#### Recurring tasks

A *task* object defines a recurring schedule for executing a SQL statement, including statements that call stored procedures. Tasks can be chained together for successive execution to support more complex periodic processing.

Tasks may optionally use table streams to provide a convenient way to continuously process new or changed data. A task can transform new or changed rows that a stream surfaces. Each time a task is scheduled to run, it can verify whether a stream contains change data for a table (using [SYSTEM\\$STREAM\\_HAS\\_DATA](#)) and either consume the change data or skip the current run if no change data exists.

Users can define a simple tree-like structure of tasks that executes consecutive SQL statements to process data and move it to various destination tables.

Question 92: [Skipped](#)

Multiple streams can be created for the same table and consumed by different tasks

- TRUE

(Correct)

- FALSE

### Explanation

This concept is very important for understand not just for the exam but for real world implementations as well.

Multiple tasks that consume change data **from a single table stream retrieve different deltas**. When a task consumes the change data in a stream using a DML statement, the stream advances the offset. The change data is no longer available for the next task to consume. Currently, we recommend that only a single task consumes the change data from a stream. **Multiple streams can be created for the same table and consumed by different tasks.**

<https://docs.snowflake.com/en/user-guidestreams.html#stream-consumption-using-tasks>

Question 93: **Skipped**

The cron expression in a task definition supports specifying a time zone

- TRUE  
**(Correct)**
- FALSE

### Explanation

#### Task Scheduling and Daylight Saving Time

**The cron expression in a task definition supports specifying a time zone.** A scheduled task runs according to the specified cron expression in the local time for a given time zone. Special care should be taken with regard to scheduling tasks for time zones that recognize daylight saving time. Tasks scheduled during specific times on days when the transition from standard time to daylight saving time (or the reverse) occurs can have unexpected behaviors.

For example:

During the autumn change from daylight saving time to standard time, a task scheduled to start at 1 AM in the America/Los\_Angeles time zone (i.e. **0 1 \* \* \*** **America/Los\_Angeles**) would run *twice*: once at 1 AM and then again when 1:59:59 AM shifts to 1:00:00 AM local time. That is, there are two points in time when the local time is 1 AM.

During the spring change from standard time to daylight saving time, a task scheduled to start at 2 AM in the America/Los\_Angeles time zone (i.e. **0 2 \* \* \* America/Los\_Angeles**)

would ***not run at all*** because the local time shifts from 1:59:59 AM to 3:00:00 AM. That is, there is no point during that day when the local time is 2 AM.

To avoid unexpected task executions due to daylight saving time, ***either***:

Do not schedule tasks to run at a specific time between 1 AM and 3 AM (daily, or on days of the week that include Sundays), ***or***

Manually adjust the cron expression for tasks scheduled during those hours twice each year to compensate for the time change due to daylight saving time.

<https://docs.snowflake.com/en/user-guide/tasks-intro.html#task-scheduling-and-daylight-saving-time>

Question 94: Skipped

In snowflake data storage cost depends on which factors

- Cloud service provider  
**(Correct)**
- Region  
**(Correct)**
- Type of account we chose (on-demand/capacity)  
**(Correct)**
- Edition

### Explanation

The storage cost is determined by cloud provider, region and the type of account. Please go the below site and try out the options

<https://www.snowflake.com/pricing/>

Question 95: Skipped

Your business team wants to load JSON data into a table in snowflake. What will you recommend them?

- Create a table with a variant column and load the JSON data in that column  
**(Correct)**
- Tell the team to convert JSON to CSV
- Construct a python program to convert the JSON into CSV and then load the data.

### Explanation

This is very important to note and this is also a performance optimization technique. I have seen people creating columns as VARCHAR and then load JSON data in it. Never do that. Always create a variant columns and then load JSON data. Why? If you use VARCHAR, you will have to explicitly use parse\_json for every row in the table which is additional processing. If you use VARIANT, you can use DOT or BRACKET notation to retrieve the data

<https://docs.snowflake.com/en/user-guide/semistructured-intro.html#loading-semi-structured-data>

<https://docs.snowflake.com/en/sql-reference/data-types-semistructured.html#variant>

Question 96: Skipped

JSON data is a hierarchical collection of name/value pairs grouped into objects and arrays

- TRUE  
**(Correct)**
- FALSE

### Explanation

Basic JSON Syntax

JSON data is a hierarchical collection of name/value pairs grouped into objects and arrays:

Colons `:` separate names and values in name/value pairs.

Curly braces `{}` denote objects.

Square brackets `[]` denote arrays.

Commas `,` separate entities in objects and arrays.

### Name/Value Pairs

JSON name/value pairs consist of a field name (in double quotes), followed by a colon, then a value.

For example:

```
{"firstName": "John", "empid": 45611}
```

Question 97: Skipped

Which parameter controls the maximum number of days for which Snowflake can extend the data retention period for tables to prevent streams on the tables from becoming stale.

- MAX\_DATA\_EXTENSION\_TIME\_IN\_DAYS**  
(Correct)
- MAX\_DATA\_RETENTION\_TIME\_IN\_DAYS
- MAX\_DATA\_EXTENSION\_TIME\_IN\_MONTHS

### Explanation

#### MAX\_DATA\_EXTENSION\_TIME\_IN\_DAYS

##### Type

Object (for databases, schemas, and tables) — Can be set for Account » Database » Schema » Table

##### Data Type

Integer

##### Description

Maximum number of days for which Snowflake can extend the data retention period for tables to prevent streams on the tables from becoming stale. By default, if the [DATA\\_RETENTION\\_TIME\\_IN\\_DAYS](#) setting for a source table is less than 14 days, and a stream has not been consumed, Snowflake temporarily extends this period to the stream's offset, up to a maximum of 14 days, regardless of the [Snowflake Edition](#) for your account. The MAX\_DATA\_EXTENSION\_TIME\_IN\_DAYS parameter enables you to limit this automatic extension period to control storage costs for data retention or for compliance reasons.

This parameter can be set at the account, database, schema, and table levels. Note that setting the parameter at the account or schema level only affects tables for which the parameter has not already been explicitly set at a lower level (e.g. at the table level by the table owner). A value of [0](#) effectively disables the automatic extension for the specified database, schema, or table. For more information about streams and staleness, see [Change Tracking Using Table Streams](#).

##### Values

[0](#) to [14](#) (i.e. 14 days) — a value of [0](#) disables the automatic extension of the data retention period. To increase the maximum value for tables in your account, contact [Snowflake Support](#).

## Default

14

<https://docs.snowflake.com/en/sql-reference/parameters.html#max-data-extension-time-in-days>

Question 98: Skipped

What is the size limit imposed by VARIANT data type.

- 100 MB compressed
- 16 MB compressed  
**(Correct)**
- 16 MB uncompressed
- 100 MB compressed

## Explanation

VARIANT

A tagged universal type, which can store values of any other type, including OBJECT and ARRAY, up to a maximum size of 16 MB compressed.

Question 99: Skipped

Non-native values such as dates and timestamps are stored as strings when loaded into a VARIANT column

- TRUE  
**(Correct)**
- FALSE

## Explanation

For data that is mostly regular and uses only native JSON types (strings and numbers, not timestamps), both storage and query performance for operations on relational data and data in a VARIANT column is very similar. Non-native values such as dates and timestamps are stored as strings when loaded into a VARIANT column, so operations on these values could be slower and also consume more space than when stored in a relational column with the corresponding data type.

Question 100: Skipped

Streams cannot track changes in materialized views

-

TRUE  
**(Correct)**

- FALSE

### Explanation

#### Overview of Table Streams

When created, a table stream logically takes an initial snapshot of every row in the source table by initializing a point in time (called an *offset*) as the current transactional version of the table. The change tracking system utilized by the stream then records information about the DML changes after this snapshot was taken. Change records provide the state of a row before and after the change. Change information mirrors the column structure of the tracked source table and includes additional metadata columns that describe each change event.

Note that a stream itself does **not** contain any table data. A stream only stores the offset for the source table and returns CDC records by leveraging the versioning history for the source table. When the first stream for a table is created, a pair of hidden columns are added to the source table and begin storing change tracking metadata. These columns consume a small amount of storage. The CDC records returned when querying a stream rely on a combination of the *offset* stored in the stream and the *change tracking metadata* stored in the table.

It may be useful to think of a stream as a bookmark, which indicates a point in time in the pages of a book (i.e. the source table). A bookmark may be thrown away and other bookmarks inserted in different places in a book. So too, a stream may be dropped and other streams created at the same or different points of time (either by creating the streams consecutively at different times or by using [Time Travel](#)) to consume the change records for a table at the same or different offsets.

One example of a consumer of CDC records is a [data pipeline](#), in which only the data in staging tables that has changed since the last extraction is transformed and copied into other tables.

Currently, streams cannot track changes in materialized views.

[Continue](#)

[Retake test](#)

Question 1: Skipped

Which of the following best describes snowflakes processing engine?

- Leverages Apache Spark
- Leverages Map Reduce
- A derivative of presto
- Native SQL  
**(Correct)**

Question 2: Skipped

Which factors affect your snowflake data loading rate?

- Physical location of the stage  
**(Correct)**
- RAM on the virtual warehouse
- GZIP compression efficiency  
**(Correct)**
- Thread size

**Explanation**

<https://www.snowflake.com/blog/how-to-load-terabytes-into-snowflake-speeds-feeds-and-techniques/>

Load rates for your own data files may differ based on a number of factors:

*Location of your S3 buckets* – For our test, both our Snowflake deployment and S3 buckets were located in us-west-2

*Number and types of columns* – A larger number of columns may require more time relative to number of bytes in the files.

*Gzip Compression efficiency* – More data read from S3 per uncompressed byte may lead to longer load times.

Question 3: Skipped

You ran a query in snowflake and it took 2 minute 3 seconds to run, you ran the query again and it returned the results in less than a second. What might have happened?

- When the query ran for the second time, it used the persisted query results from the query result cache

**(Correct)**

- The query used the meta data cache
- All the partitions were scanned from the SSD cache

### **Explanation**

When a query is executed, the result is persisted (i.e. cached) for a period of time (currently 24 hours). At the end of the time period, the result is purged from the system.

Snowflake uses persisted query results to avoid re-generating results when nothing has changed (i.e. “retrieval optimization”).

In addition, you can use persisted query results to post-process the results (e.g. layering a new query on top of the results already calculated).

<https://docs.snowflake.com/en/user-guide/querying-persisted-results.html#using-persisted-query-results>

**Question 4:** Skipped

You are loading data to a snowflake internal stage area using the PUT Command, the data will be encrypted on the client's machine and you will not be able to turn off the encryption.

- TRUE  
**(Correct)**
- FALSE

### **Explanation**

One of the biggest worries people have about moving to the cloud is security. One key piece of providing enterprise class security is the ability to encrypt the data in your data warehouse environment. With Snowflake, your data is automatically encrypted by default.

<https://www.snowflake.com/blog/automatic-encryption-data/>

**Question 5:** Skipped

Please select the ones that are true for a snowflake task

- A task can execute a single SQL Statement  
**(Correct)**
- A task can execute a call to a store procedure  
**(Correct)**
- A task can execute multiple SQL statements

- A task can execute a function

### Explanation

Currently, a task can execute a single SQL statement, including a call to a stored procedure.

<https://docs.snowflake.com/en/user-guide/tasks-intro.html#introduction-to-tasks>

Question 6: Skipped

Which object in snowflake records data manipulation language changes made to tables, including inserts, updates, and deletes, as well as metadata change about each change?

- Tasks
- Table stream  
**(Correct)**
- Snowpipe
- Change Data Capture

### Explanation

A stream object records data manipulation language (DML) changes made to tables, including inserts, updates, and deletes, as well as metadata about each change, so that actions can be taken using the changed data. This process is referred to as change data capture (CDC). An individual table stream tracks the changes made to rows in a *source table*. A table stream (also referred to as simply a “stream”) makes a “change table” available of what changed, at the row level, between two transactional points of time in a table. This allows querying and consuming a sequence of change records in a transactional fashion.

Question 7: Skipped

A table stream(also known as stream) itself does **not** contain any table data

- TRUE  
**(Correct)**
- FALSE

### Explanation

This is very important to know.

Note that a stream itself does **not** contain any table data. A stream only stores the **offset** for the source table and returns CDC records by leveraging the versioning history for the source table. When the first stream for a table is created, a pair of hidden columns are added to the source table and begin storing change tracking metadata. These columns consume a small amount of storage. The CDC records returned when querying a stream rely on a combination of the *offset* stored in the stream and the *change tracking metadata* stored in the table.

<https://docs.snowflake.com/en/user-guidestreams.html#overview-of-table-streams>

Question 8: Skipped

The current offset for a stream can be determined by querying which function?

- SYSTEM\$STREAM\_GET\_TABLE\_TIMESTAMP  
**(Correct)**
- SYSTEM\$DISABLE\_BEHAVIOR\_CHANGE\_BUNDLE
- SYSTEM\$TASK\_DEPENDENTS\_ENABLE
- SYSTEM\$USER\_TASK\_CANCEL\_ONGOING\_EXECUTIONS

**Explanation**

**SYSTEM\$STREAM\_GET\_TABLE\_TIMESTAMP** returns a timestamp indicating the transactional point when the stream contents were last consumed using a DML statement.

[https://docs.snowflake.com/en/sql-reference/functions/system\\_stream\\_get\\_table\\_timestamp.html#system-stream-get-table-timestamp](https://docs.snowflake.com/en/sql-reference/functions/system_stream_get_table_timestamp.html#system-stream-get-table-timestamp)

Question 9: Skipped

Streams support repeatable read isolation

- TRUE  
**(Correct)**
- FALSE

**Explanation**

Streams support repeatable read isolation. In repeatable read mode, multiple SQL statements within a transaction see the same set of records in a stream. This differs from the read committed mode supported for tables, in which statements see any changes made by previous statements executed within the same transaction, even though those changes are not yet committed.

<https://docs.snowflake.com/en/user-guidestreams.html#repeatable-read-isolation>

Question 10: Skipped

A stream stores data in the same shape as the source table with some additional columns. Which are those additional columns?

- METADATA\$ACTION  
**(Correct)**
-

### METADATA\$ISUPDATE (Correct)

- METADATA\$ROW\_ID  
(Correct)
- METADATA\$COLUMN\_ID

### Explanation

A stream stores data in the same shape as the source table (i.e. the same column names and ordering) with the following additional columns:

### METADATA\$ACTION

Indicates the DML operation (INSERT, DELETE) recorded.

### METADATA\$ISUPDATE

Indicates whether the operation was part of an UPDATE statement. Updates to rows in the source table are represented as a pair of DELETE and INSERT records in the stream with a metadata column METADATA\$ISUPDATE values set to TRUE.

Note that streams record the differences between two offsets. If a row is added and then updated in the current offset, the delta change is a new row. The METADATA\$ISUPDATE row records a FALSE value.

### METADATA\$ROW\_ID

Specifies the unique and immutable ID for the row, which can be used to track changes to specific rows over time.

Question 11: Skipped

What are the various types of table streams?

- Standard  
(Correct)
- Append-only  
(Correct)
- Insert-only  
(Correct)
- Upsert

### Explanation

<https://docs.snowflake.com/en/user-guidestreams.html#types-of-streams>

## Please note insert-only stream is in preview mode

The following stream types are available based on the metadata recorded by each:

### Standard

A standard (i.e. delta) table stream tracks all DML changes to the source table, including inserts, updates, and deletes (including table truncates). This stream type performs a join on inserted and deleted rows in the change set to provide the row level delta. As a net effect, for example, a row that is inserted and then deleted between two transactional points of time in a table is removed in the delta (i.e. is not returned when the stream is queried).

### Append-only

An append-only table stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded. For example, if 10 rows are inserted into a table and then 5 of those rows are deleted before the offset for an append-only stream is advanced, the stream records 10 rows.

An append-only stream returns the appended rows only and therefore can be much more performant than a standard stream for extract, load, transform (ELT) and similar scenarios that depend exclusively on row inserts. For example, the source table can be truncated immediately after the rows in an append-only stream are consumed, and the record deletions do not contribute to the overhead the next time the stream is queried or consumed.

### Insert-only

**Supported on external tables only.** An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops). For example, in-between any two offsets, if File1 is removed from the cloud storage location referenced by the external table, and File2 is added, the stream returns records for the rows in File2 only. Unlike when tracking CDC data for standard tables, Snowflake cannot access the historical records for files in cloud storage.

Question 12: Skipped

How will you know if a stream has become stale?

- Execute `DESCRIBE STREAM` or `SHOW STREAMS`  
**(Correct)**
- Execute `SHOW TABLE STREAMS`
- Execute `SHOW PIPES`

### Explanation

To determine whether a stream has become stale, execute the **DESCRIBE STREAM** or **SHOW STREAMS** command. In the command output, when the STALE column value for the stream is TRUE, the stream is stale.

Question 13: Skipped

What are the two ways available in snowflake before change tracking metadata is recorded for a table?

- Change tracking is enabled on the table (using `ALTER TABLE ... CHANGE_TRACKING = TRUE`)  
**(Correct)**
- A stream is created for the table (using `CREATE STREAM`)  
**(Correct)**
- Enable CDC on the table

### Explanation

Currently, *either* of the following must be true before change tracking metadata is recorded for a table:

1. Change tracking is enabled on the table (using `ALTER TABLE ... CHANGE_TRACKING = TRUE`).
2. A stream is created for the table (using `CREATE STREAM`).

Either option adds a pair of hidden columns to the table and begins storing change tracking metadata. The columns consume a small amount of storage.

No change tracking metadata for the table is available for the period before one of these conditions is satisfied.

Question 14: Skipped

Which role permissions are required to create and manage streams ? Please specify the OBJECT and the PERMISSION.

- DATABASE - USAGE  
**(Correct)**
- SCHEMA - USAGE, CREATE STREAM  
**(Correct)**
- TABLE - SELECT  
**(Correct)**
- TABLE - DELETE

### Explanation

<https://docs.snowflake.com/en/user-guidestreams.html#required-access-privileges>

Question 15: Skipped

Streams ensure **exactly once semantics** for new or changed data in a table

- TRUE  
**(Correct)**
- FALSE

### Explanation

Tasks can be combined with **table streams** for continuous ELT workflows to process recently changed table rows. **Streams ensure exactly once semantics for new or changed data in a table.**

<https://docs.snowflake.com/en/user-guide/tasks-intro.html#introduction-to-tasks>

Question 16: Skipped

Lets say that you have scheduled a task. At time T0, the task has started executing. If it is still executing when the next scheduling execution time occurs, what happens to that task

- It will be skipped  
**(Correct)**
- It will wait for the previous task to complete
- It will be aborted

### Explanation

Snowflake ensures only one instance of a task with a schedule (i.e. a standalone task or the root task in a tree of tasks) is executed at a given time. If a task is still running when the next scheduled execution time occurs, then that scheduled time is skipped.

Question 17: Skipped

Which role or privileges are required to view **TASK\_HISTORY**?

- Account administrator (i.e. users with the ACCOUNTADMIN role)  
**(Correct)**
- Task owner (i.e. role that has the OWNERSHIP privilege on a task)  
**(Correct)**
- Any role that has the global MONITOR EXECUTION privilege  
**(Correct)**
- Any one with SYSADMIN role

### Explanation

The following roles (or roles with the specified privileges) can use SQL to view the task history within a specified date range:

Account administrator (i.e. users with the ACCOUNTADMIN role).

Task owner (i.e. role that has the OWNERSHIP privilege on a task).

Any role that has the global MONITOR EXECUTION privilege.

Question 18: Skipped

You want to query the history of task usage within a specified date range. What will you do?

- You will query the TASK\_HISTORY table function  
**(Correct)**
- You will run '**SHOW TASK HISTORY**'
- You will run '**'DESCRIBE TASKS'**'

### Explanation

#### TASK\_HISTORY

This table function can be used to query the history of task usage within a specified date range. The function returns the history of task usage for your entire Snowflake account or a specified task.

```
TASK_HISTORY(  
    [ SCHEDULED_TIME_RANGE_START => <constant_expr> ]  
    [, SCHEDULED_TIME_RANGE_END => <constant_expr> ]  
    [, RESULT_LIMIT => <integer> ]  
    [, TASK_NAME => '<string>' ] )
```

Retrieve the execution history for tasks in the account within a specified 30 minute block of time within the past 7 days:

```
select *  
from table(information_schema.task_history(  
    scheduled_time_range_start=>to_timestamp_ltz('2018-11-9 12:00:00.000 -0700  
'),
```

```
scheduled_time_range_end=>to_timestamp_ltz('2018-11-9 12:30:00.000 -0700')  
));
```

Question 19: Skipped

Which SNOWFLAKE object stores a generated identity and access management (IAM) entity for your external cloud storage, along with an optional set of allowed or blocked storage locations (Amazon S3, Google Cloud Storage, or Microsoft Azure)?

- STORAGE INTEGRATION**  
**(Correct)**
- STORAGE GROUP**
- STORAGE CLUSTERS**

### Explanation

A storage integration is a Snowflake object that stores a generated identity and access management (IAM) entity for your external cloud storage, along with an optional set of allowed or blocked storage locations (Amazon S3, Google Cloud Storage, or Microsoft Azure). Cloud provider administrators in your organization grant permissions on the storage locations to the generated entity. This option allows users to avoid supplying credentials when creating stages or when loading or unloading data

Question 20: Skipped

A single storage integration can support multiple external stages

- TRUE**  
**(Correct)**
- FALSE**

### Explanation

A single storage integration can support multiple external stages. The URL in the stage definition must align with the storage location specified for the **STORAGE\_ALLOWED\_LOCATIONS** parameter.

Please see example below

```
create storage integration s3_int  
  type = external_stage  
  storage_provider = s3  
  storage_aws_role_arn = 'arn:aws:iam::001234567890:role/myrole'  
  enabled = true  
  storage_allowed_locations = ('s3://mybucket1/path1', 's3://mybucket2/path2/  
');  
create storage integration gcs_int
```

```

type = external_stage
storage_provider = gcs
enabled = true
storage_allowed_locations = ('gcs://mybucket1/path1/', 'gcs://mybucket2/path2/');
create storage integration azure_int
type = external_stage
storage_provider = azure
enabled = true
azure_tenant_id = '<tenant_id>'
storage_allowed_locations = ('azure://myaccount.blob.core.windows.net/mycontainer/path1/', 'azure://myaccount.blob.core.windows.net/mycontainer/path2/');

```

Question 21: Skipped

What will you use to query Snowflake data loading history for the last 365 days (1 year)?

- COPY\_HISTORY View  
**(Correct)**
- LOAD\_HISTORY View from INFORMATION SCHEMA
- INSERT\_HISTORY View

### Explanation

#### COPY\_HISTORY View

This Account Usage view can be used to query Snowflake data loading history for the last 365 days (1 year). The view displays load activity for both `COPY INTO <table>` statements and continuous data loading using `Snowpipe`. The view avoids the 10,000 row limitation of the `LOAD_HISTORY` View.

[https://docs.snowflake.com/en/sql-reference/account-usage/copy\\_history.html#copy-history-view](https://docs.snowflake.com/en/sql-reference/account-usage/copy_history.html#copy-history-view)

#### LOAD\_HISTORY View from Information schema

This Information Schema view enables you to retrieve the history of data loaded into tables using the `COPY INTO <table>` command within the last 14 days. The view displays one row for each file loaded.

[https://docs.snowflake.com/en/sql-reference/info-schema/load\\_history.html#load-history-view](https://docs.snowflake.com/en/sql-reference/info-schema/load_history.html#load-history-view)

**Please also read about the LOAD\_HISTORY view from ACCOUNT USAGE schema**

[https://docs.snowflake.com/en/sql-reference/account-usage/load\\_history.html#load-history-view](https://docs.snowflake.com/en/sql-reference/account-usage/load_history.html#load-history-view)

Question 22: Skipped

Which feature in snowflake allows you to load data in micro batches?

- SNOWPIPE  
**(Correct)**
- COPY INTO COMMAND
- LOAD COMMAND

#### Explanation

Snowpipe enables loading data from files as soon as they're available in a stage. This means you can load data from files in **micro-batches**, making it available to users within minutes, rather than manually executing COPY statements on a schedule to load larger batches.

Question 23: Skipped

Load history for bulk data load is stored in the metadata of the target table for 64 days

- TRUE  
**(Correct)**
- FALSE

#### Explanation

#### Bulk data load

Stored in the metadata of the target table for 64 days. Available upon completion of the COPY statement as the statement output.

#### Snowpipe

Stored in the metadata of the pipe for 14 days. Must be requested from Snowflake via a REST endpoint, SQL table function, or ACCOUNT\_USAGE view.

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-intro.html#load-history>

Question 24: Skipped

Which one of the below requires an user-specified warehouse to execute COPY statements

- Bulk data load  
**(Correct)**

- Snowpipe
- Both

## Explanation

### Bulk data load

Requires a user-specified warehouse to execute COPY statements.

### Snowpipe

Uses Snowflake-supplied compute resources.

Question 25: Skipped

When you remove the predecessor of a child task, what may happen to the child task. Select two possibilities?

- The child task is removed from the system
- The child task may become the root task  
**(Correct)**
- The child task may become a stand-alone task  
**(Correct)**

## Explanation

Tasks is a very important topic for the exam. I have added a lot of questions now on tasks across the question sets. Please ensure that you read the task chapter thoroughly.

### Severing the Link Between Predecessor and Child Tasks

If the predecessor for a child task is removed (using [ALTER TASK ... REMOVE AFTER](#)), then the former child task becomes either a standalone task or a root task, depending on whether other tasks identify this former child task as their predecessor.

If a predecessor task is dropped (using [DROP TASK](#)), or if ownership of a predecessor task is transferred to another role (using [GRANT OWNERSHIP](#)), then all former child tasks that identified this task as the predecessor become either standalone tasks or root tasks, depending on whether other tasks identify these former child tasks as their predecessor.

Question 26: Skipped

While loading data through the COPY command, you can transform the data. Which of the below transformations are allowed?

- TRUNCATECOLUMNS  
**(Correct)**

- REORDER COLUMNS  
**(Correct)**
- OMIT COLUMNS  
**(Correct)**
- CAST  
**(Correct)**
- AGGREGATE

### Explanation

Snowflake supports transforming data while loading it into a table using the `COPY INTO <table>` command, dramatically simplifying your ETL pipeline for basic transformations. This feature helps you avoid the use of temporary tables to store pre-transformed data when reordering columns during a data load.

The COPY command supports:

Column reordering, column omission, and casts using a SELECT statement. There is no requirement for your data files to have the same number and ordering of columns as your target table.

The `ENFORCE_LENGTH | TRUNCATECOLUMNS` option, which can truncate text strings that exceed the target column length.

Currently, these copy options support CSV data only.

Question 27: Skipped

Which of the below is true about snowflake architecture

- One node shared data
- Multi Cluster Shared data  
**(Correct)**
- Multi Node distributed data

### Explanation

If you miss this question, it will hurt me a lot:(

You do not need to remember this. Ask yourself the question. How does compute work in Snowflake? Virtual Warehouse, right? What are Virtual Warehouses? They are a cluster of nodes(like EC2 instances if on AWS). Each virtual warehouse is an MPP compute cluster

composed of **multiple compute nodes allocated by Snowflake from a cloud provider**. A Large WAREHOUSE has 8 nodes in a cluster. Each node has 8 cores so total 64 cores. All these nodes refer to a single source of data and that data is shared by all the nodes. Snowflake is architected in such a way that, all the nodes have access to the same data and then access it without stepping on each others toes.

Below is a list of all the warehouse sizes and their details.

Warehouse Size	Servers / Cluster	Credits / Hour	Credits / Second	Notes
X-Small	1	1	0.0003	Default size for warehouses created using CREATE WAREHOUSE.
Small	2	2	0.0006	
Medium	4	4	0.0011	
Large	8	8	0.0022	
X-Large	16	16	0.0044	Default for warehouses created in the web interface.
2X-Large	32	32	0.0089	
3X-Large	64	64	0.0178	
4X-Large	128	128	0.0356	

Question 28: **Skipped**

Snowflake supports multi-factor authentication (MFA) to provide increased login security for users connecting to Snowflake. At this time, Snowflake users are not automatically enrolled in MFA.

- TRUE  
**(Correct)**
- FALSE

### Explanation

MFA is enabled on a per-user basis; however, at this time, users are not automatically enrolled in MFA. To use MFA, users must enroll themselves.

<https://docs.snowflake.com/en/user-guide/security-mfa.html#multi-factor-authentication-mfa>

Question 29: **Skipped**

Select two true statements about fail safe from the following

- Fail safe is same as time travel
- Fail-safe provides a (non-configurable) 7-day period during which historical data is recoverable by Snowflake

**(Correct)**

- Fail-safe provides a (non-configurable) 90-day period during which historical data is recoverable by Snowflake
- Fail-safe ensures historical data is protected in the event of a system failure or other catastrophic event, e.g. a hardware failure or security breach

**(Correct)**

### Explanation

Separate and distinct from Time Travel, Fail-safe ensures historical data is protected in the event of a system failure or other catastrophic event, e.g. a hardware failure or security breach.

Fail-safe provides a (non-configurable) 7-day period during which historical data is recoverable by Snowflake. This period starts immediately after the Time Travel retention period ends.

<https://docs.snowflake.com/en/user-guide/data-failsafe.html#understanding-viewing-fail-safe>

### Very important to Note:

**Fail-safe is *not* provided as a means for accessing historical data after the Time Travel retention period has ended. It is for use *only* by Snowflake to recover data that may have been lost or damaged due to extreme operational failures.**

Question 30: Skipped

When you load a table from a file through COPY command, the file should have the same number and order of columns as the table

- TRUE
- FALSE

**(Correct)**

### Explanation

Note that the actual field/column order in the data files can be different from the column order in the target table. It is only important that the SELECT list maps fields/columns in the data files to the *corresponding* columns in the table.

<https://docs.snowflake.com/en/sql-reference/sql/copy-into-table.html#transformation-parameters>

Try it out yourself. You will never forget it.

1. Load a snowflake stage with a file using a PUT command from SNOWSQL

2. In my case I have already loaded an employee file into stage and it looks like as below

Row	name	size	md5	last_modified
1	employee_stage/EMPLOYEE.csv.gz	128	13df091dbb5ac8015bcb4b8f8b20a792	Tue, 8 Sep 2020 21:10:14 GMT

3. Do a select on the stage. Please note how I wrote the select. This is how you do select from stage

Row	\$1	\$2	\$3	\$4	\$5
1	FIRST_NAME	LAST_NAME	EMPLOYEE_ID	SALARY	JOINING_DATE
2	NISHA	JOSE	102	12000	10/11/01

4. I will now create a table with only first name and last name and load this data into the table

```
CREATE OR REPLACE TABLE EMPLOYEE_FIRST_LAST(FNAME VARCHAR, LNAME VARCHAR);
```

```
COPY INTO EMPLOYEE_FIRST_LAST
FROM (SELECT $1,$2 FROM @EMPLOYEE_STAGE );
```

Row	file	status	rows_parsed	rows_loaded	error_limit	errors_seen	first_error	first_error_line	first_error_character	first_error_column
1	employee_stag...	LOADED	2	2	1	0	NULL	NULL	NULL	NULL

Question 31: Skipped

Failsafe period for temporary and transient table is zero

- TRUE  
**(Correct)**
- FALSE

### Explanation

This is a question that you need to remember. temporary table lives till the session is alive, so no point in having fail safe, transient tables also do not have fail safe and hence the name:)

To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides two table types, temporary and transient, which do not incur the same fees as standard (i.e. permanent) tables:

Transient tables can have a Time Travel retention period of either 0 or 1 day.

Temporary tables can also have a Time Travel retention period of 0 or 1 day; however, this retention period ends as soon as the table is dropped or the session in which the table was created ends.

Transient and temporary tables have no Fail-safe period.

Question 32: Skipped

When the snowflake staged release happens, which accounts get the updates first

- Enterprise or above
- Trial accounts
- Premier accounts
- Accounts who has opted for early access  
**(Correct)**

### Explanation

If you have multiple Enterprise Edition (or higher) accounts, you can designate one or more of these accounts to take advantage of the 24-hour period between the early access and final stages for new releases. This can be particularly useful if you maintain separate accounts for development/testing and production.

To designate an account for early access, please contact [Snowflake Support](#).

Once you have designated one or more accounts for early access, you can implement a testing framework similar to the following:

Use `CURRENT_VERSION` (or a UDF that returns similar results) to verify when your early access account(s) are on the new release.

Use your early access account(s) to test your production workloads against the new release.

If any issues are encountered, notify Snowflake Support, who can work with you to prevent the issues from disrupting your other accounts.

*Early access is not required or recommended for all organizations with Enterprise Edition accounts; Snowflake's rigorous release testing and monitoring during deployments is usually sufficient to prevent most issues. Early access is intended primarily for organizations that desire added certainty that their production accounts will not be affected by changes in new releases.*

Question 33: Skipped

After you have granted a share to a consumer, you added more objects to the share. You have provided the required access by using the GRANT <privilege> ... TO SHARE command.

Please select the true statement

- The added objects will be accessible to the consumers immediately  
**(Correct)**
- You will need to recreate the share

### Explanation

<https://docs.snowflake.com/en/user-guide/data-sharing-provider.html#adding-objects-to-a-share>

### Adding Objects to a Share

You can add objects to an existing share at any time using the **GRANT <privilege> ... TO SHARE** command. Any objects that you add to a share are instantly available to the consumers accounts who have created databases from the share. For example, if you add a table to a share, users in consumer accounts can query the data in the table as soon as the table is added to the share

Question 34: Skipped

For federated authentication and MFA, what is the minimum snowflake edition required?

- Business critical
- Enterprise
- Premier
- Standard  
**(Correct)**

Question 35: Skipped

A warehouse can be assigned to a single resource monitor only

- TRUE  
**(Correct)**
- FALSE

### Explanation

#### Assignment of Resource Monitors

A single monitor can be set at the account level to control credit usage for all warehouses in your account.

In addition, a monitor can be assigned to one or more warehouses, thereby controlling the credit usage for each assigned warehouse. Note, however, **that a warehouse can be assigned to only a single resource monitor.**

Question 36: Skipped

Select the objects which do not have a fail safe period

- Transient table  
**(Correct)**
- Temporary table  
**(Correct)**
- Permanent table

Question 37: Skipped

When staging uncompressed files in a Snowflake stage, the files are automatically compressed using gzip, unless compression is explicitly disabled.

- TRUE  
**(Correct)**
- FALSE

## Explanation

Feature	Supported	Notes
Uncompressed files	gzip	When staging uncompressed files in a Snowflake stage, the files are automatically compressed using gzip, unless compression is explicitly disabled.
Already-compressed files	gzip bzip2  deflate  raw_deflate	Snowflake can automatically detect any of these compression methods or you can explicitly specify the method that was used to compress the files.
	Brotli  Zstandard	Auto-detection is <b>not</b> yet supported for these methods; when staging or loading files compressed with either of these methods, you must explicitly specify the compression method that was used.

Question 38: Skipped

COPY does not purge loaded files from the location. What option will you specify to purge files after successful load of a table from table stage

- PURGE=TRUE  
**(Correct)**
-

DELETE=TRUE

- 
- REMOVE=TRUE
- 
- TRUNCATE=TRUE

### Explanation

#### Purging Files After Loading

Load files from a table's stage into the table and purge files after loading. By default, COPY does not purge loaded files from the location. To purge the files after loading:

Set **PURGE=TRUE** for the table to specify that all files successfully loaded into the table are purged after loading:

```
alter table mytable set stage_copy_options = (purge = true);  
  
copy into mytable;
```

You can also override any of the copy options directly in the COPY command:

```
copy into mytable purge = true;
```

Question 39: Skipped

You have specified PURGE=TRUE in your COPY command. Due to a network failure the data files could not be purged after successful load. You will get an error returned by snowflake due to this failure.

- 
- TRUE
- 
- FALSE  
**(Correct)**

### Explanation

**PURGE = TRUE | FALSE**

Boolean that specifies whether to remove the data files from the stage automatically after the data is loaded successfully.

If this option is set to **TRUE**, note that a best effort is made to remove successfully loaded data files. **If the purge operation fails for any reason, no error is returned currently**. We recommend that you list staged files periodically (using **LIST**) and manually remove successfully loaded files, if any exist.

Default: **FALSE**

<https://docs.snowflake.com/en/sql-reference/sql/copy-into-table.html#copy-options-copyoptions>

Question 40: Skipped

You can specify the below action to perform when an error is encountered while loading data from a file

- CONTINUE**  
**(Correct)**
- SKIP\_FILE**  
**(Correct)**
- SKIP\_FILE\_<num>**  
**(Correct)**
- SKIP\_FILE\_<num>%**  
**(Correct)**
- ABORT\_STATEMENT**  
**(Correct)**
- ERROR\_STATEMENT**

### Explanation

**Copy Options (copyOptions)**

You can specify one or more of the following copy options (separated by blank spaces, commas, or new lines):

```
ON_ERROR = CONTINUE | SKIP_FILE | SKIP_FILE_num | SKIP_FILE_num% |  
ABORT_STATEMENT
```

String (constant) that specifies the action to perform when an error is encountered while loading data from a file:

Supported Values	Notes
<code>CONTINUE</code>	Continue loading the file. The COPY statement returns an error message for a maximum of one error encountered per data file. Note that the difference between the ROWS_PARSED and ROWS_LOADED column values represents the number of rows that include detected errors. However, each of these rows could include multiple errors. To view all errors in the data files, use the VALIDATION_MODE parameter or query the VALIDATE function.
<code>SKIP_FILE</code>	Skip file if any errors encountered in the file.
<code>SKIP_FILE_&lt;num&gt;</code> (e.g. <code>SKIP_FILE_10</code> )	Skip file when the number of errors in the file is equal to or exceeds the specified number.
<code>SKIP_FILE_&lt;num&gt;%</code> (e.g. <code>SKIP_FILE_10%</code> )	Skip file when the percentage of errors in the file exceeds the specified percentage.
<code>ABORT_STATEMENT</code>	Abort the load operation if any error is encountered in a data file. Note that the load operation is <i>not</i> aborted if the data file cannot be found (e.g. because it does not exist or cannot be accessed).

Question 41: Skipped

What is the default action to perform when an error is encountered while loading data from a file using SNOWPIPE?

- 
- ABORT\_STATEMENT
- 
- SKIP\_FILE  
**(Correct)**
- 
- CONTINUE

### Explanation

please remember this. This is different for SNOWPIPE and BULK LOAD using COPY

Default:

### Bulk loading using COPY

`ABORT_STATEMENT`

### Snowpipe

`SKIP_FILE`

Question 42: Skipped

Snowpipe is a serverless function

- 
- TRUE  
**(Correct)**
- 
- FALSE

### Explanation

Automated data loads leverage event notifications for cloud storage to inform Snowpipe of the arrival of new data files to load. Snowpipe copies the files into a queue, from which they are loaded into the **target table in a continuous, serverless fashion** based on parameters defined in a specified pipe object.

Question 43: Skipped

Which character below identifies a table stage

- 1. @%
- 1. @
- 1. %

### Explanation

Unloading Data to a Table Stage

Use the **COPY INTO <location>** command to unload all the rows from a table into one or more files in the stage for the table. The following example unloads data files to the stage using the named **my\_csv\_unload\_format** file format created in [Preparing to Unload Data](#).

The statement prefixes the unloaded file(s) with **unload/** to organize the files in the stage:

For example:

```
copy into @%mytable/unload/ from mytable file_format = (format_name = 'my_csv_unload_format' compression = none);
```

Note that the @% character combination identifies a table stage.

Question 44: Skipped

Which permissions are required by Snowflake in AWS S3 to unload data to S3?

- s3:DeleteObject**  
**(Correct)**
- s3:PutObject**  
**(Correct)**
- s3>ListObject**

### Explanation

**Configuring an S3 Bucket for Unloading Data**

Snowflake requires the following permissions on an S3 bucket and folder to create new files in the folder (and any sub-folders):

**s3:DeleteObject**

### s3:PutObject

As a best practice, Snowflake recommends configuring a storage integration object to delegate authentication responsibility for external cloud storage to a Snowflake identity and access management (IAM) entity.

Question 45: Skipped

If you recreate a storage integration after it has been linked to one or more stages, you must reestablish the association between each stage by running the ALTER STAGE command.

- TRUE  
**(Correct)**
- FALSE

### Explanation

If you must recreate a storage integration after it has been linked to one or more stages, you must reestablish the association between each stage and the storage integration by executing `ALTER STAGE stage_name SET STORAGE_INTEGRATION = storage_integration_name`, where:

*stage\_name* is the name of the stage.

*storage\_integration\_name* is the name of the storage integration.

Question 46: Skipped

A single storage integration can support multiple external stages

- TRUE  
**(Correct)**
- FALSE

### Explanation

Creates a new storage integration in the account or replaces an existing integration.

A storage integration is a Snowflake object that stores a generated identity and access management (IAM) entity for your external cloud storage, along with an optional set of allowed or blocked storage locations (Amazon S3, Google Cloud Storage, or Microsoft Azure). Cloud provider administrators in your organization grant permissions on the storage locations to the generated entity. This option allows users to avoid supplying credentials when creating stages or when loading or unloading data.

**A single storage integration can support multiple external stages.** The URL in the stage definition must align with the storage location specified for the STORAGE\_ALLOWED\_LOCATIONS parameter.

Question 47: Skipped

What is the storage hierarchy in snowflake?

- Account->Database->Schemas->Objects  
**(Correct)**
- Organization->folder->project->resource
- Database->Schemas->Objects

### Explanation

Snowflake storage hierarchy starts with an account. A customer can have as many accounts as possible, but each account will have a single URL and each account is deployed in a single cloud provider platform. Also an account cannot be multi-region, it is deployed on a single region and a single snowflake edition is attached to the account

Databases are within accounts and each database belongs to a single account. While they can be replicated to other accounts they cannot span accounts

Schemas are contained within a database. A schema can be attached to a single database only. Again while schema can be replicated to other database and accounts it cannot span accounts or databases

All objects including tables, views, file formats, sequences, UDFs and stored procedures are contained within the schema

Question 48: Skipped

If you logon to snowflake using <https://mycompany.snowflakecomputing.com/>, then what is the cloud provider and region you are logging into?

- AWS and us-west-2  
**(Correct)**
- Azure and us-east-1
- AWS and us-east-2

### Explanation

If an account has just the account name followed by snowflakecomputing.com, it means that the cloud provider is AWS and the region is us-west-2. This is the first cloud provider and region where snowflake was deployed

Question 49: Skipped

What table types are available in snowflake

- PERMANENT

(Correct)

- TEMPORARY  
**(Correct)**
- TRANSIENT  
**(Correct)**
- EXTERNAL  
**(Correct)**
- INTERNAL

### Explanation

Snowflake supports the below type of tables. It is important to know when to use which table type

1. PERMANENT - This is the default standard table type. It has time travel turned on. It also has a fail safe period of 7 days which is non-configurable
2. TEMPORARY - Only available for the session in which it is created. No time travel or fail safe period
3. TRANSIENT - It is same as PERMANENT except it does not have fail safe period. Time travel upto 1 day is possible
4. EXTERNAL - Stored in external stages and is read only. No cloning, time travel or fail safe is possible.

<https://docs.snowflake.com/en/user-guide/tables-temp-transient.html>

<https://docs.snowflake.com/en/user-guide/tables-external-intro.html>

Question 50: Skipped

What types of view are available in snowflake?

- REGULAR  
**(Correct)**

- SECURE  
**(Correct)**
- MATERIALIZED  
**(Correct)**
- HASH VIEW
- TABLE VIEW

### Explanation

REGULAR Views are the default view type. The underlying DDL is accessible to any role who has access to the view

SECURE VIEWS are those where the query definition and the details are accessible to authorized users only. Please note that since the underlying DDL is not visible to the optimizer, this view type will be slower than a normal view

Materialized View - These are more like tables, results are stored like it is in a table. Snowflake can auto refresh a materialized view. Please take extra care while using materialized views, it may increase your cost.

### Question 51: Skipped

You are developing an ETL query. One of your ETL logic requires you to store the intermediate data temporarily. This data will be used only by your ETL. The data will not be required outside the ETL. If you want to optimize storage cost, what type of table will you create to store this data

- PERMANENT
- TRANSIENT
- TEMPORARY  
**(Correct)**
- EXTERNAL

### Explanation

Please note that temporary tables are accessible within the session. It gets automatically destroyed when the session is destroyed. You create temporary tables if you have a use case to store data temporarily for a session only

### Question 52: Skipped

Tick the statements which are true for fail safe

- Fail safe period is configurable

- Data can be recovered from a fail-safe storage only by snowflake employee  
**(Correct)**
- Using fail-safe you can create dev/test environments
- Tables with fail-safe turned on incurs additional storage costs compared to tables where fail-safe is not turned on  
**(Correct)**

### Explanation

When fail-safe is turned on a table, the table requires additional storage to store historical data. That additional data storage cost is charged to the customer. Also to restore data from fail-safe storage, customer will need to contact snowflake

Question 53: Skipped

What are the various stages available in Snowflake

- USER  
**(Correct)**
- TABLE  
**(Correct)**
- NAMED INTERNAL  
**(Correct)**
- NAMED EXTERNAL  
**(Correct)**
- ACCOUNT STAGE
- SCHEMA STAGE

### Explanation

SNOWFLAKE supports three type of stages

1. USER
2. TABLE
3. NAMED

Named stages are again of two types internal and external

USER and Table stages are automatically available in snowflake and are not needed to be created or configured

<https://docs.snowflake.com/en/user-guide/data-load-local-file-system-create-stage.html>

Question 54: Skipped

Which of the following object types are stored within schema

- Tables  
**(Correct)**
- Views  
**(Correct)**
- User defined functions  
**(Correct)**
- File formats  
**(Correct)**
- Sequences  
**(Correct)**
- Roles

### Explanation

Roles are at the account level and not at schema level

Question 55: Skipped

One of the government agencies has decided to use Snowflake for their datawarehouse. They require compliance with US federal privacy and security standards, such as FIPS 140-2and FedRAMP (Moderate Impact) . Which cloud providers can you choose to setup the snowflake instance?

- AWS  
**(Correct)**
- AZURE  
**(Correct)**
- GOOGLE
- IBM

### Explanation

For government agencies that require compliance with US federal privacy and security standards, such as [FIPS 140-2](#) and [FedRAMP \(Moderate Impact\)](#), Snowflake supports the following regions in the US:

AWS - US East (Commercial Gov - N. Virginia)

AZURE - US Gov Virginia

<https://docs.snowflake.com/en/user-guide/intro-regions.html#government-regions>

Question 56: [Skipped](#)

Which AWS regions are supported by snowflake in US and Canada

- us-west-2  
**(Correct)**

- us-east-2  
**(Correct)**

- us-east-1  
**(Correct)**

- ca-central-1  
**(Correct)**

- us-central-2

### Explanation

<https://docs.snowflake.com/en/user-guide/intro-regions.html#general-regions>

Question 57: [Skipped](#)

Which GCP regions are supported by Snowflake in US and Canada

- us-central1  
**(Correct)**

- us-central2

- us-east1

### Explanation

Only US Central1 (Iowa) is supported in GCP

<https://docs.snowflake.com/en/user-guide/intro-regions.html#general-regions>

Question 58: Skipped

Which Azure regions are supported by Snowflake in US and Canada

- west-us-2  
**(Correct)**
- east-us-2  
**(Correct)**
- canada-central  
**(Correct)**
- west-us-1

### Explanation

<https://docs.snowflake.com/en/user-guide/intro-regions.html#general-regions>

Question 59: Skipped

Your customer deals with extremely sensitive data like PHI data. They need to be compliant with HIPAA. What is the minimum edition that you will recommend

- STANDARD
- PREMIUM
- Business Critical Edition  
**(Correct)**
- Virtual Private Snowflake

### Explanation

Business Critical Edition, formerly known as Enterprise for Sensitive Data (ESD), offers even higher levels of data protection to support the needs of organizations with extremely sensitive data, particularly PHI data that must comply with HIPAA regulations.

It includes all the features and services of [Enterprise Edition](#), with the addition of enhanced security and data protection. In addition, database failover/failback adds support for business continuity and disaster recovery.

<https://docs.snowflake.com/en/user-guide/intro-editions.html#business-critical-edition>

Question 60: Skipped

Your customer required the highest level of data protection and needs a completely isolated environment for snowflake. Which edition will you recommend

- STANDARD

- PREMIER
- Virtual Private Snowflake  
**(Correct)**
- Business Critical Edition

### Explanation

Virtual Private Snowflake offers our highest level of security for organizations that have the strictest requirements, such as financial institutions and any other large enterprises that collect, analyze, and share highly sensitive data.

It includes all the features and services of [Business Critical Edition](#), but in a completely separate Snowflake environment, isolated from all other Snowflake accounts (i.e. VPS accounts do not share any resources with accounts outside the VPS).

<https://docs.snowflake.com/en/user-guide/intro-editions.html#virtual-private-snowflake-vps>

Question 61: Skipped

Which of the below semi structured data types are supported by snowflake?

- VARIANT  
**(Correct)**
- OBJECT  
**(Correct)**
- ARRAY  
**(Correct)**
- GEOGRAPHY

### Explanation

Snowflake supports three types of semi structured data types

VARIANT - A tagged universal type, which can store values of any other type, including OBJECT and ARRAY, up to a maximum size of 16 MB compressed.

OBJECT - Used to represent collections of key-value pairs, where the key is a non-empty string, and the value is a value of VARIANT type. Snowflake does not currently support explicitly-typed objects.

**ARRAY** - Used to represent dense or sparse arrays of arbitrary size, where index is a non-negative integer (up to  $2^{31}-1$ ), and values have VARIANT type. Snowflake does not currently support fixed-size arrays or arrays of elements of a specific non-VARIANT type.

<https://www.udemy.com/instructor/course/3459188/manage/practice-tests>

Question 62: Skipped

You have a multi-threaded client program that interacts with snowflake. If you are following snowflake recommendations what will you do

- Use a separate connection for each thread  
**(Correct)**
- Execute the threads synchronously  
**(Correct)**
- Execute the threads asynchronously
- Use one connection for all the threads

### Explanation

Although multiple sessions cannot share the same transaction, multiple **threads** using a single connection share the same session, and thus share the same transaction. This can lead to unexpected results, such as one thread rolling back work done in another thread.

This situation can occur when a client application using a Snowflake driver (such as the Snowflake JDBC Driver) or connector (such as the Snowflake Connector for Python) is multi-threaded. If two or more threads share the same connection, then those threads also share the current transaction in that connection. A BEGIN, COMMIT, or ROLLBACK by one thread affects all threads using that shared connection. If the threads are running asynchronously, the results can be unpredictable.

Similarly, changing the AUTOCOMMIT setting in one thread affects the AUTOCOMMIT setting in all other threads that use the same connection.

Note that changing the state of the AUTOCOMMIT setting commits and existing transaction.

Snowflake recommends that multi-threaded client programs do at least one of the following:

Use a separate connection for each thread.

Note that even with separate connections, your code can still hit race conditions that generate unpredictable output; for example, one thread might delete data before another thread tries to update it.

Execute the threads synchronously rather than asynchronously, to control the order in which steps are performed.

<https://docs.snowflake.com/en/sql-reference/transactions.html#transactions-and-multi-threading>

Question 63: Skipped

Which database objects can be shared using snowflake's secure data sharing

- Tables  
**(Correct)**
- External tables  
**(Correct)**
- Secure views  
**(Correct)**
- Secure materialized views  
**(Correct)**
- Secure UDFs  
**(Correct)**
- Roles

### Explanation

Secure Data Sharing enables sharing selected objects in a database in your account with other Snowflake accounts. The following Snowflake database objects can be shared:

1. Tables
2. External tables
3. Secure views
4. Secure materialized views
5. Secure UDFs

Snowflake enables the sharing of databases through *shares*, which are created by data providers and “imported” by data consumers

<https://docs.snowflake.com/en/user-guide/data-sharing-intro.html#introduction-to-secure-data-sharing>

## **Additional Explanation based on a student's question**

Only secure views or secure materialized views can be shared. This is because of security reasons and that is why it is also called secure data sharing. All objects must be secure except for the tables. With replication, any objects contained within a database can be replicated to another account. The disadvantage with this over data sharing is that you are actually replicating data which will incur storage cost, compute cost to replicate data and also data egress if the replication is happening to a different region. The other thing is that data sharing will allow you to see immediate changes to the data, data replication will have latency based on the frequency of replication. Data sharing is the recommended approach if you are in the same region or deployment.

You can also share a schema. You need to include a schema in the share as the objects reside in schema.

Question 64: Skipped

Your customer needs to setup Failover/Failback mechanism in their snowflake instance. Which edition of Snowflake will you suggest?

- STANDARD
- PREMIER
- BUSINESS CRITICAL(OR HIGHER)  
**(Correct)**

### **Explanation**

Failover/Failback requires Business Critical (or higher)

<https://docs.snowflake.com/en/user-guide/database-replication-failover.html>

Question 65: Skipped

Currently in snowflake, replication is supported for which type of objects

- Databases  
**(Correct)**
- Users
- Roles
- Warehouses

## Explanation

Currently, replication is supported for *databases only*. Other types of objects in an account **cannot** be replicated. This list includes:

1. Users
2. Roles
3. Warehouses
4. Resource monitors
5. Shares

<https://docs.snowflake.com/en/user-guide/database-replication-intro.html#other-objects-in-an-account>

### Question 66: Skipped

You have stored the below JSON in a table named car\_sales as a variant column

```
1. {
2.   "customer": [
3.     {
4.       "address": "San Francisco, CA",
5.       "name": "Joyce Ridgely",
6.       "phone": "1650437889"
7.     }
8.   ],
9.   "date": "2017-04-28",
10.  "dealership": "Valley View Auto Sales",
11.  "salesperson": {
12.    "id": "55",
13.    "name": "Frank Beasley"
14.  },
15.  "vehicle": [
16.    {
17.      "extras": [
18.        "ext warranty",
19.        "paint protection"
20.      ],
21.      "make": "Honda",
22.      "model": "Civic",
23.      "price": "20275",
24.      "year": "2017"
25.    }
26.  ]
27. }
```

How will you query the table to get the dealership data?

- select src:dealership from car\_sales;  
**(Correct)**
-

```
select src:$dealership from car_sales;
```

- select dealership from car\_sales;

### Explanation

Insert a colon `:` between the VARIANT column name and any first-level element: `<column>:<level1_element>`

<https://docs.snowflake.com/en/user-guide/querying-semistructured.html#traversing-semi-structured-data>

Question 67: Skipped

One of your team members accidentally truncated a table in your snowflake database? What feature of snowflake will you use to restore data into the table

- TIME TRAVEL  
**(Correct)**
- FAIL SAFE
- UNDROP TABLE

### Explanation

Snowflake Time Travel enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It serves as a powerful tool for performing the following tasks:

1. Restoring data-related objects (tables, schemas, and databases) that may have been accidentally or intentionally deleted.
2. Duplicating and backing up data from key points in the past.
3. Analyzing data usage/manipulation over specified periods of time.

<https://docs.snowflake.com/en/user-guide/data-time-travel.html#understanding-using-time-travel>

### Additional explanation:

One of my students pointed out why is undrop not a correct option for this question. So, in the exam you may get some tricky questions. Pay careful attention to what is being asked in the question. In this question, the table is truncated and not dropped. So , your table is still there, you just lost the data. The most appropriate answer in this case is to use time travel and get the data back.

Question 68: Skipped

What happens when you execute the below query

**CREATE OR REPLACE TABLE MyTable\_V2 CLONE MyTable**

- MyTable\_V2 is created and all the data from MyTable is copied into MyTable\_V2
- MyTable\_V2 is created and a batch job is submitted to copy data from MyTable
- Snowflake creates a new entry in the metadata store to keep track of the new clone. The existing micro partitions of table MyTable are also mapped to the new table. This is called zero copy clone.  
**(Correct)**

### Explanation

Clone creates a copy of an existing object in the system. This command is primarily used for creating **zero-copy clones** of databases, schemas, and non-temporary tables; however, it can also be used to quickly/easily create clones of other schema objects (i.e. external stages, file formats, and sequences).

<https://docs.snowflake.com/en/sql-reference/sql/create-clone.html#create-object-clone>

Question 69: Skipped

Which of the below statements are true about snowflake data sharing?

- Consumers can query shared tables in the same query as their own tables  
**(Correct)**
- Producers can share data to other databases such as Postgres and MySQL
- You don't need a Snowflake Virtual Warehouse to query shared data
- Data Sharing is integrated with role-based access controls  
**(Correct)**

### Explanation

<https://docs.snowflake.com/en/user-guide/data-sharing-intro.html#how-does-secure-data-sharing-work>

Question 70: Skipped

What are the default roles that snowflake automatically adds for any new account?

- AccountAdmin  
**(Correct)**
- SecurityAdmin  
**(Correct)**

- SysAdmin  
**(Correct)**
- Public  
**(Correct)**
- Security Admin

### Explanation

By default Snowflake offers the following [default roles](#):

1. AccountAdmin
2. SecurityAdmin
3. SysAdmin
4. Public

These are a good starting point, but won't be sufficient for most implementations, so we end up defining [custom roles](#). Snowflake's recommendation is to create a hierarchy of custom roles with the top-most custom role assigned to the system role SYSADMIN. That way the system administrators will be able to manage all warehouses and databases while maintaining management of user and roles restricted to users granted the SECURITYADMIN or ACCOUNTADMIN roles.

<https://support.snowflake.net/s/article/Quickly-Visualize-Snowflake-s-Roles-Grants-and-Privileges>

Question 71: [Skipped](#)

What are the different caching mechanisms that are available in snowflake?

- Warehouse caching  
**(Correct)**
- Index Caching
- Metadata caching  
**(Correct)**
- Query result caching  
**(Correct)**

### Explanation

Snowflake supports three types of caching

1. Metadata cache -
2. Query result cache
3. Warehouse cache

**further reading:**

<https://blog.ippon.tech/innovative-snowflake-features-caching/#:~:text=Metadata%20Cache,-Fully%20Managed%20in&text=Snowflake%20automatically%20collects%20and%20manages,Table%20Size%20in%20Bytes>

Question 72: Skipped

One of your user complained that a query is taking a long time to return any result. What techniques would you consider to improve performance?

- Create an index on the search columns
- Create cluster keys and turn on auto clustering on the table  
**(Correct)**
- Enable query hints on the table
- Define the partition keys for the table

**Explanation**

Automatic Clustering is the Snowflake service that seamlessly and continually manages all reclustering, as needed, of clustered tables.

Automatic Clustering eliminates the need for performing any of the following tasks:

1. Monitoring the state of clustered tables.
2. Instead, as DML is performed on these tables, Snowflake monitors and evaluates the tables to determine whether they would benefit from reclustering, and automatically reclusters them, as needed.
3. Designating warehouses in your account to use for reclustering.
4. Snowflake performs automatic reclustering in the background, and you do not need to specify a warehouse to use.

All you need to do is define a clustering key for each table (if appropriate) and Snowflake manages all future maintenance.

<https://docs.snowflake.com/en/user-guide/tables-auto-reclustering.html>

Question 73: Skipped

Which parameter will you use while creating a STORAGE integration to limit external stages that use the integration to reference one or more storage locations.

- STORAGE\_ALLOWED\_LOCATIONS**  
**(Correct)**
- STORAGE\_ALLOWED\_SITES**
- STORAGE\_OBJECTS\_ALLOWED\_LOCATIONS**

**Explanation**

**`STORAGE_ALLOWED_LOCATIONS = ('<cloud_specific_url>')`**

Explicitly limits external stages that use the integration to reference one or more storage locations (i.e. S3 bucket, GCS bucket, or Azure container). Supports a comma-separated list of URLs for existing buckets and, optionally, paths used to store data files for loading/unloading. Alternatively supports the `*` wildcard, meaning “allow access to all buckets and/or paths”.

<https://docs.snowflake.com/en/sql-reference/sql/create-storage-integration.html#required-parameters>

Question 74: Skipped

Which are the required parameters while creating a storage integration?

- Name  
**(Correct)**
- Type  
**(Correct)**
- Enabled  
**(Correct)**
- Storage\_Allowed\_Locations  
**(Correct)**
- Storage\_blocked\_locations

**Explanation**

<https://docs.snowflake.com/en/sql-reference/sql/create-storage-integration.html#required-parameters>

## Required Parameters

### `name`

String that specifies the identifier (i.e. name) for the integration; must be unique in your account.

In addition, the identifier must start with an alphabetic character and cannot contain spaces or special characters unless the entire identifier string is enclosed in double quotes (e.g. `"My object"`). Identifiers enclosed in double quotes are also case-sensitive.

For more details, see [Identifier Requirements](#).

### `TYPE = EXTERNAL_STAGE`

Specify the type of integration:

`EXTERNAL_STAGE`: Creates an interface between Snowflake and an external cloud storage location.

### `ENABLED = TRUE / FALSE`

Specifies whether this storage integration is available for usage in stages.

`TRUE` allows users to create new stages that reference this integration. Existing stages that reference this integration function normally.

`FALSE` prevents users from creating new stages that reference this integration. Existing stages that reference this integration cannot access the storage location in the stage definition.

### `STORAGE_ALLOWED_LOCATIONS = ('cloud_specific_url')`

Explicitly limits external stages that use the integration to reference one or more storage locations (i.e. S3 bucket, GCS bucket, or Azure container). Supports a comma-separated list of URLs for existing buckets and, optionally, paths used to store data files for loading/unloading. Alternatively supports the `*` wildcard, meaning “allow access to all buckets and/or paths”.

## Amazon S3

### `STORAGE_ALLOWED_LOCATIONS = ('s3://bucket/path/','s3://bucket/path/')`

`bucket` is the name of an S3 bucket that stores your data files (e.g. `mybucket`).

`path` is an optional case-sensitive path for files in the cloud storage location (i.e. files have names that begin with a common string) that limits access to a set of files. Paths are alternatively called *prefixes* or *folders* by different cloud storage services.

## Google Cloud Storage

```
STORAGE_ALLOWED_LOCATIONS = ('gcs://bucket/path/', 'gcs://bucket/path/')
```

`bucket` is the name of a GCS bucket that stores your data files (e.g. `mybucket`).

`path` is an optional case-sensitive path for files in the cloud storage location (i.e. files have names that begin with a common string) that limits access to a set of files. Paths are alternatively called *prefixes* or *folders* by different cloud storage services.

## Microsoft Azure

```
STORAGE_ALLOWED_LOCATIONS =
('azure://account.blob.core.windows.net/container/path/',
'azure://account.blob.core.windows.net/container/path/')
```

`account` is the name of the Azure account (e.g. `myaccount`). Use the `blob.core.windows.net` endpoint for all supported types of Azure blob storage accounts, including Data Lake Storage Gen2.

`container` is the name of a Azure blob storage container that stores your data files (e.g. `mycontainer`).

`path` is an optional case-sensitive path for files in the cloud storage location (i.e. files have names that begin with a common string) that limits access to a set of files. Paths are alternatively called *prefixes* or *folders* by different cloud storage services.

Question 75: Skipped

Which command will you run to list all users and roles to which a role has been granted.

- 1. SHOW GRANTS OF ROLE <role\_name>  
**(Correct)**
- 1. SHOW GRANTS TO ROLE <role\_name>
- 1. SHOW GRANTS ON ACCOUNT

## Explanation

<https://docs.snowflake.com/en/sql-reference/sql/show-grants.html#syntax>

SHOW GRANTS OF...

ROLE *role\_name*

Lists all users and roles to which the role has been granted.

Question 76: Skipped

If you plan to regularly unload similarly-formatted data, it is recommended to use named file formats

- TRUE  
**(Correct)**
- FALSE

## Explanation

Named file formats are optional, but are *recommended* when you plan to regularly unload similarly-formatted data.

Question 77: Skipped

Which command will you run to view a list of files that have been unloaded to a stage named mystage.

- 1. list @mystage;  
**(Correct)**
- 1. show @mystage;
- 1. display @mystage;

## Explanation

Use the **LIST** command to view a list of files that have been unloaded to the stage:

```
list @mystage;
```

name	last_modified	size	md5
mystage/unload/data_0_0_0.csv.gz	Mon, 11 Sep 2017 18:13:07 GMT	112	6f77daba007a643bdff4eae10de5bed3

+	-	-	-	-	-	-	-	-
-	-	-	-	-	-	-	-	-

Question 78: Skipped

Which character combination identifies a user stage

- @~
- @~

(Correct)

- ~
- ~

- @
- @

### Explanation

Use the `COPY INTO <location>` command to unload all the rows from a table into one or more files in your stage. The following example unloads data files to your user stage using the named `my_csv_unload_format` file format created in [Preparing to Unload Data](#). The statement prefixes the unloaded file(s) with `unload/` to organize the files in the stage:

For example:

```
copy into @~/unload/ from mytable file_format = (format_name = 'my_csv_unload_format' compression = none);
```

Note that the `@~` character combination identifies a user stage.

Question 79: Skipped

The Snowflake Web UI data loading wizard is only intended for loading small numbers of files of limited size (up to 50 MB).

- TRUE
- TRUE

- FALSE
- FALSE

### Explanation

The wizard is only intended for loading small numbers of files of limited size (up to 50 MB). This file size limit is intended to ensure better performance because browser performance varies from computer to computer and between different browser versions. Also, the memory consumption required to encrypt larger files might cause a browser to run out of memory and crash.

For loading larger files or large numbers of files, snowflake recommend using the Snowflake client, SnowSQL

Question 80: Skipped

If a query is running slow, to check if this is for an overloaded warehouse where in the Snowflake UI you will go to check the same

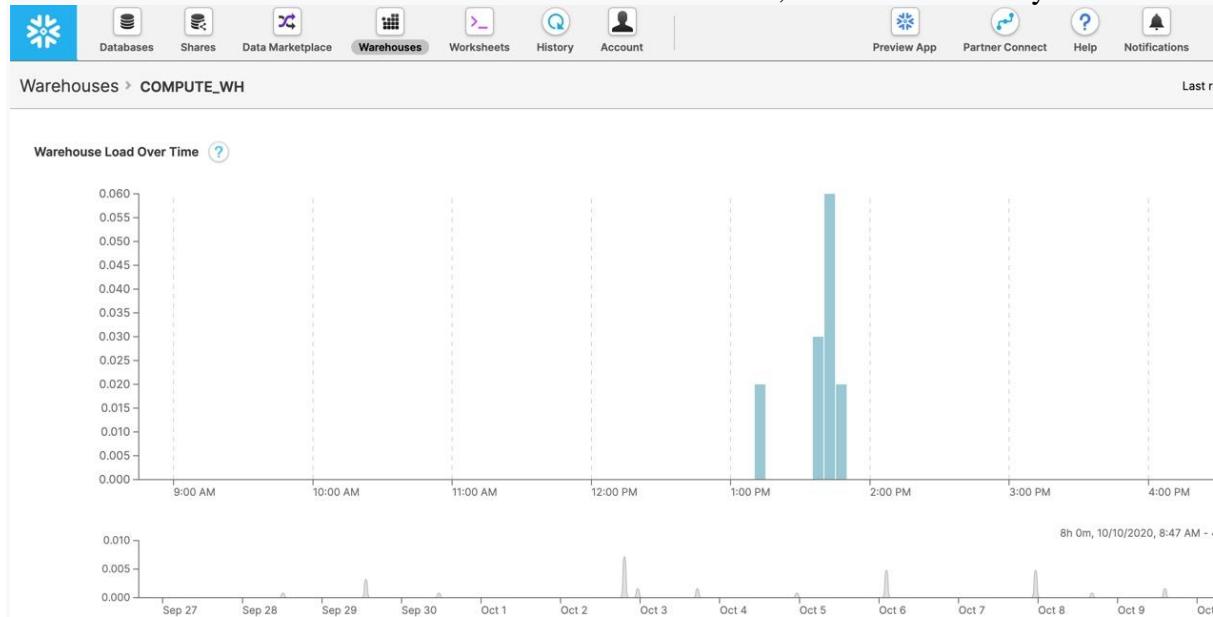
-

Warehouse Tab  
**(Correct)**

- Database Tab
- Query Tab

**Explanation**

It will be shown under the warehouse tab. if it is overloaded, the bars will be in yellow color



Question 81: Skipped

Which of the below mentioned clustering metadata for the micro-partitions is maintained by snowflake in a table?

- The total number of micro-partitions that comprise the table  
**(Correct)**
- The number of micro-partitions containing values that overlap with each other  
**(Correct)**
- The depth of the overlapping micro-partitions  
**(Correct)**
- None of the above

**Explanation**

**Clustering Information Maintained for Micro-partitions**

Snowflake maintains clustering metadata for the micro-partitions in a table, including:

1. The total number of micro-partitions that comprise the table.

2. The number of micro-partitions containing values that overlap with each other (in a specified subset of table columns).
3. The depth of the overlapping micro-partitions.

Question 82: Skipped

What would you check to see if a large table will benefit from explicitly defining a clustering key

- Clustering depth  
**(Correct)**
- Clustering percentage
- Clustering ratio

### Explanation

#### Clustering Depth

The clustering depth for a populated table measures the average depth (1 or greater) of the overlapping micro-partitions for specified columns in a table. The smaller the average depth, the better clustered the table is with regards to the specified columns.

Clustering depth can be used for a variety of purposes, including:

1. Monitoring the clustering “health” of a large table, particularly over time as DML is performed on the table.
2. Determining whether a large table would benefit from explicitly defining a [clustering key](#).

A table with no micro-partitions (i.e. an unpopulated/empty table) has a clustering depth of 0.

Question 83: Skipped

If query performance degrades over time, the table is likely no longer well-clustered and may benefit from clustering

- TRUE  
**(Correct)**
- FALSE

### Explanation

The clustering depth for a table is *not* an absolute or precise measure of whether the table is well-clustered. Ultimately, query performance is the best indicator of how well-clustered a table is:

1. If queries on a table are performing as needed or expected, the table is likely well-clustered.
2. If query performance degrades over time, the table is likely no longer well-clustered and may benefit from clustering.

Question 84: Skipped

Which system functions are available in snowflake to view/monitor the clustering metadata for a table

- SYSTEM\$CLUSTERING\_DEPTH  
**(Correct)**
- SYSTEM\$CLUSTERING\_INFORMATION (including clustering depth)  
**(Correct)**
- CLUSTERING\_DEPTH\_VIEW

### Explanation

#### Monitoring Clustering Information for Tables

To view/monitor the clustering metadata for a table, Snowflake provides the following system functions:

SYSTEM\$CLUSTERING\_DEPTH

SYSTEM\$CLUSTERING\_INFORMATION (including clustering depth)

Question 85: Skipped

Clustering is generally most cost-effective for tables that are queried frequently and do not change frequently

- TRUE  
**(Correct)**
- FALSE

### Explanation

The more frequently a table is queried, the more benefit clustering provides. However, the more frequently a table changes, the more expensive it will be to keep it clustered. Therefore, clustering is generally most cost-effective for tables that are queried frequently and do not change frequently.

Question 86: Skipped

Which of the below will you consider while choosing a cluster key

- Columns that are most actively used in selective filters  
**(Correct)**

- Columns frequently used in join predicates  
**(Correct)**
- Columns with extremely high cardinality
- Columns with extremely low cardinality

### Explanation

Snowflake recommends prioritizing keys in the order below:

1. Cluster columns that are most actively used in selective filters. For many fact tables involved in date-based queries (for example “WHERE invoice\_date > x AND invoice\_date <= y”), choosing the date column is a good idea. For event tables, event type might be a good choice, if there are a large number of different event types. (If your table has only a small number of different event types, then see the comments on cardinality below before choosing an event column as a clustering key.)
2. If there is room for additional cluster keys, then consider columns frequently used in join predicates, for example “FROM table1 JOIN table2 ON table2.column\_A = table1.column\_B”.

If you typically filter queries by two dimensions

(e.g. `application_id` and `user_status` columns), then clustering on both columns can improve performance.

The number of distinct values (i.e. cardinality) in a column/expression is a critical aspect of selecting it as a clustering key. It is important to choose a clustering key that has:

A large enough number of distinct values to enable effective pruning on the table.

A small enough number of distinct values to allow Snowflake to effectively group rows in the same micro-partitions.

A column with very low cardinality (e.g. a column that indicates only whether a person is male or female) might yield only minimal pruning. At the other extreme, a column with very high cardinality (e.g. a column containing UUID or nanosecond timestamp values) is also typically **not** a good candidate to use as a clustering key directly.

Question 87: Skipped

Reclustering in Snowflake requires manual configuration

- False  
**(Correct)**
- True

### Explanation

Reclustering in Snowflake is automatic; no maintenance is needed.

<https://docs.snowflake.com/en/user-guide/tables-auto-reclustering.html>

Question 88: Skipped

Automatic reclustering in snowflake is triggered only if/when the table would benefit from the operation

- TRUE  
(Correct)
- FALSE

### Explanation

Optimal Efficiency

With Automatic Clustering, Snowflake internally manages the state of clustered tables, as well as the resources (servers, memory, etc.) used for all automated clustering operations. This allows Snowflake to dynamically allocate resources as needed, resulting in the most efficient and effective reclustering.

Also, Automatic Clustering does not perform any unnecessary reclustering.  
Reclustering is triggered only if/when the table would benefit from the operation.

Question 89: Skipped

Which of the below privileges are required to add clustering on a table

- USAGE or Ownership on the schema and database  
(Correct)
- MONITOR USAGE on the table
- ALTER privilege on the table

### Explanation

To add clustering to a table, you must also have USAGE or OWNERSHIP privileges on the schema and database that contain the table.

Question 90: Skipped

The data stored in temporary table does not contribute to storage cost

- FALSE  
(Correct)
-

TRUE

### Explanation

#### Data Storage Usage for Temporary Tables

For the duration of the existence of a temporary table, the data stored in the table contributes to the overall storage charges that Snowflake bills your account. To prevent any unexpected storage changes, particularly if you create large temporary tables in sessions that you maintain for periods longer than 24 hours, Snowflake recommends explicitly dropping these tables once they are no longer needed. You can also explicitly exit the session in which the table was created to ensure no additional charges are accrued.

Question 91: Skipped

You have executed a task. Due to some coding issue, the task went into a loop. How long the task will run before it is terminated by Snowflake

- 24 hours
- 60 minutes  
**(Correct)**
- 30 minutes

### Explanation

**There is a 60 minute default limit on a single run of a task.** This limitation was implemented as a safeguard against non-terminating tasks. Query the [TASK\\_HISTORY](#) table function. If the task was canceled or exceeded the window scheduled for the task, the cause is often an undersized warehouse. Review the warehouse size and consider increasing it to fit within the schedule window or the one-hour limit.

Alternatively, consider increasing the timeout limit for the task by executing [ALTER TASK ... SET USER\\_TASK\\_TIMEOUT\\_MS = <num>](#).

A larger warehouse may not help if there are query parallelization issues. Consider looking at alternate ways to rewrite the SQL statement run by the task.

Question 92: Skipped

Snowflake supports below type of streams. Please select three

- Standard  
**(Correct)**
- Update-only
- Append-only

(Correct)



- Insert-only

(Correct)

## Explanation

The following stream types are available based on the metadata recorded by each:

### Standard

A standard (i.e. delta) table stream tracks all DML changes to the source table, including inserts, updates, and deletes (including table truncates). This stream type performs a join on inserted and deleted rows in the change set to provide the row level delta. As a net effect, for example, a row that is inserted and then deleted between two transactional points of time in a table is removed in the delta (i.e. is not returned when the stream is queried).

### Append-only

An append-only table stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded. For example, if 10 rows are inserted into a table and then 5 of those rows are deleted before the offset for an append-only stream is advanced, the stream records 10 rows.

An append-only stream returns the appended rows only and therefore can be much more performant than a standard stream for extract, load, transform (ELT) and similar scenarios that depend exclusively on row inserts. For example, the source table can be truncated immediately after the rows in an append-only stream are consumed, and the record deletions do not contribute to the overhead the next time the stream is queried or consumed.

**Insert-only** - This is currently a preview feature

**Supported on external tables only.** An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops). For example, in-between any two offsets, if File1 is removed from the cloud storage location referenced by the external table, and File2 is added, the stream returns records for the rows in File2 only. Unlike when tracking CDC data for standard tables, Snowflake cannot access the historical records for files in cloud storage.

Overwritten files are essentially handled as new files: The old version of the file is removed from cloud storage, but the insert-only stream does not record the delete operation. The new version of the file is added to cloud storage, and the insert-only stream records the rows as inserts. The stream does not record the diff of the old and new file versions.

Question 93: Skipped

You have a schema with a table and a stream. You cloned the schema. What will happen to the stream



- Any unconsumed records in the stream (in the clone) are inaccessible

(Correct)

- Any unconsumed records in the stream (in the clone) will also be accessible in the clone
- Any unconsumed records in the stream (in the clone) will automatically flow into the table

### Explanation

Currently, when a database or schema that contains a source table and stream is cloned, any unconsumed records in the stream (in the clone) are inaccessible. This behavior is consistent with [Time Travel](#) for tables. If a table is cloned, historical data for the table clone begins at the time/point when the clone was created.

Question 94: [Skipped](#)

You have a materialized view and you want to track changes in the view. You can use streams to track that change

- TRUE
- FALSE  
**(Correct)**

### Explanation

Currently, streams cannot track changes in materialized views.

Question 95: [Skipped](#)

A Stream has been created on a table. A row in the table got inserted and subsequently got updated. The stream will capture both the events.

- TRUE
- FALSE  
**(Correct)**

### Explanation

A stream maintains a point of time into the transactional versioned timeline of the source table, called an *offset*, which starts at the transactional point when the stream contents were last consumed using a DML statement. The stream provides the set of changes from the current offset to the current *transactional time* of the source table (i.e. the current version of the table). **The stream maintains only the delta of the changes; if multiple DML statements change a row, the stream records only the latest action taken on that row.**

Question 96: [Skipped](#)

You would like to execute a merge statement in snowflake every one hour. If you want to schedule this query using a snowflake object what will you use?

- TASK  
**(Correct)**
-

## STREAM

- SNOWPIPE

### Explanation Task Scheduling

There is no event source that can trigger a task; instead, a task runs on a schedule, which can be defined when creating a task (using [CREATE TASK](#)) or later (using [ALTER TASK](#)).

Snowflake ensures only one instance of a task with a schedule (i.e. a standalone task or the root task in a tree of tasks) is executed at a given time. If a task is still running when the next scheduled execution time occurs, then that scheduled time is skipped.

Question 97: [Skipped](#)

Which roles can use SQL to view the task history within a specified date range?

- Account administrator (i.e. users with the ACCOUNTADMIN role).  
**(Correct)**
- Task owner (i.e. role that has the OWNERSHIP privilege on a task).  
**(Correct)**
- Any role that has the global MONITOR EXECUTION privilege.  
**(Correct)**
- SECURITYADMIN

### Explanation

#### Viewing the Task History for Your Account

The following roles (or roles with the specified privileges) can use SQL to view the task history within a specified date range:

1. Account administrator (i.e. users with the ACCOUNTADMIN role).
2. Task owner (i.e. role that has the OWNERSHIP privilege on a task).
3. Any role that has the global MONITOR EXECUTION privilege.

To view the task history:

## SQL

Query the [TASK\\_HISTORY](#) table function (in the [Information Schema](#)).

Question 98: [Skipped](#)

Insert-only streams are supported on external tables only and it is currently a preview feature

- TRUE  
**(Correct)**
- FALSE

### Explanation

#### Insert-only

**Supported on external tables only.** An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops). For example, in-between any two offsets, if File1 is removed from the cloud storage location referenced by the external table, and File2 is added, the stream returns records for the rows in File2 only. Unlike when tracking CDC data for standard tables, Snowflake cannot access the historical records for files in cloud storage.

Overwritten files are essentially handled as new files: The old version of the file is removed from cloud storage, but the insert-only stream does not record the delete operation. The new version of the file is added to cloud storage, and the insert-only stream records the rows as inserts. The stream does not record the diff of the old and new file versions.

<https://docs.snowflake.com/en/user-guidestreams.html#types-of-streams>

Question 99: **Skipped**

When the owner role of a given task (i.e. the role with the OWNERSHIP privilege on the task) is deleted, what happens to the task?

- The task is “re-possessed” by the role that dropped the owner role  
**(Correct)**
- The task is deleted immediately
- The task is suspended immediately

### Explanation

#### Dropping a Task Owner Role

When the owner role of a given task (i.e. the role with the OWNERSHIP privilege on the task) is deleted, the task is “re-possessed” by the role that dropped the owner role. This ensures that ownership moves to a role that is closer to the root of the role hierarchy. When a task is re-possessed, it is automatically paused, i.e., all executions currently in flight complete processing, but new executions will not be scheduled until the task is resumed explicitly by the new owner. The rationale for this is to prevent a user with access to a particular role from leaving behind tasks that suddenly execute with higher permissions when the role is removed.

If the role that a running task is executing under is dropped while the task is running, the task completes processing under the dropped role.

Question 100: Skipped

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state

- TRUE  
**(Correct)**
- FALSE

### Explanation

A simple tree of tasks is limited to a maximum of 1000 tasks total (including the root task) in a resumed state. An individual task in the tree is limited to a single predecessor task; however, a task can have a maximum of 100 *child* tasks (i.e. other tasks that identify the task as a predecessor).

Currently, we cannot guarantee that only one instance of a task with a defined *predecessor task* is running at a given time.

[Continue](#)

[Retake test](#)

Question 1: Skipped

Snowflake is available on premise

- True
- False

(Correct)

### Explanation

Snowflake's data warehouse is a true SaaS offering. More specifically:

There is no hardware (virtual or physical) for you to select, install, configure, or manage.

There is no software for you to install, configure, or manage.

Ongoing maintenance, management, and tuning is handled by Snowflake.

Snowflake runs completely on cloud infrastructure. All components of Snowflake's service (other than an optional command line client), run in a public cloud infrastructure.

Snowflake uses virtual compute instances for its compute needs and a storage service for persistent storage of data. Snowflake cannot be run on private cloud infrastructures (on-premises or hosted).

Snowflake is not a packaged software offering that can be installed by a user. Snowflake manages all aspects of software installation and updates.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

Question 2: Skipped

Pick the true statement for snowflake architecture

- Shared nothing architecture
- Shared disk architecture
- Multi-Cluster Shared Data architecture

(Correct)

### Explanation

Snowflake's architecture is a hybrid of traditional shared-disk database architectures and shared-nothing database architectures. Similar to shared-disk architectures, Snowflake uses a central data repository for persisted data that is accessible from all compute nodes in the data warehouse. But similar to shared-nothing architectures, Snowflake processes queries using MPP (massively parallel processing) compute clusters where each node in the cluster stores a portion of the entire data set locally. This approach offers the data management simplicity of

a shared-disk architecture, but with the performance and scale-out benefits of a shared-nothing architecture.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

**Question 3:** Skipped

The three key layers of snowflake are

- Extraction, Ingestion, Load
- Database storage, Query Processing, Cloud Services  
**(Correct)**
- Database, Virtual Warehouse, Data Experience

#### **Explanation**

Snowflake's unique architecture consists of three key layers:

Database Storage

Query Processing

Cloud Services

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

**Question 4:** Skipped

How is data loaded into snowflake?

- Snowflake loads the data in parquet format on the underlying cloud storage
- Snowflake loads the data in JSON format on the underlying cloud storage
- Snowflake reorganizes the data into its internal optimized, compressed, columnar format at stores on the underlying cloud storage  
**(Correct)**

#### **Explanation**

When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed, columnar format. Snowflake stores this optimized data in cloud storage.

Snowflake manages all aspects of how this data is stored — the organization, file size, structure, compression, metadata, statistics, and other aspects of data storage are handled by Snowflake. The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

Question 5: Skipped

Query processing in snowflake is done by...

- 
- 
- 
- Snowflake process queries using 'Virtual Warehouses'  
**(Correct)**

### Explanation

Query execution is performed in the processing layer. Snowflake processes queries using “virtual warehouses”. Each virtual warehouse is an MPP compute cluster composed of multiple compute nodes allocated by Snowflake from a cloud provider.

Each virtual warehouse is an independent compute cluster that does not share compute resources with other virtual warehouses. As a result, each virtual warehouse has no impact on the performance of other virtual warehouses.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html>

Question 6: Skipped

Which of the below services are provided by cloud services

- Metadata management  
**(Correct)**
- Authentication  
**(Correct)**
- Infrastructure management  
**(Correct)**
- Query execution

### Explanation

The cloud services layer is a collection of services that coordinate activities across Snowflake. These services tie together all of the different components of Snowflake in order to process user requests, from login to query dispatch. The cloud services layer also runs on compute instances provisioned by Snowflake from the cloud provider.

Among the services in this layer:

Authentication

Infrastructure management

Metadata management

Query parsing and optimization

Access control

Question 7: Skipped

How will you store JSON data in snowflake

- Using a column with datatype as JSON
- Using a column with datatype as VARCHAR
- Using a column with datatype as VARIANT  
**(Correct)**

### Explanation

Snowflakes variant datatype allows you to store semi-structured non relational data like JSON, AVRO, XML

Question 8: Skipped

You have two virtual warehouses in your snowflake instance. You have updated the data in the storage layer using one of the warehouses. When will the other warehouse be able to see the data

- Immediately  
**(Correct)**
- After 30 minutes once snowflake completes data synchronization
- You will need to trigger the data synchronization process for the other warehouse to see the data

### Explanation

Virtual warehouses uses the same data storage layer. So If one of the warehouses updates the data, it is immediately available to all the warehouses

Question 9: Skipped

Zero-copy cloning in snowflake is powered by which service?

- Metadata store of the service layer  
**(Correct)**
- SSD cache of Virtual warehouse
- Query result cache

### **Explanation**

Metadata store is a key component of the services layer. It supports zero copy cloning, data sharing and time travel

Question 10: Skipped

What influences snowflake pricing?

- Amount of data queried from snowflake
- Amount of data scanned during querying snowflake
- Snowflake pricing is based on usage and it charges only for storage and compute  
**(Correct)**

### **Explanation**

There are only two const components for which Snowflake charges its customers. They are storage and compute. This is very important to remember when you optimize snowflake for cost.

Question 11: Skipped

Compute cost in snowflake depends on

- The actual query execution time
- The query execution time and the time the query waits for the resource
- The warehouse size and how long the warehouse runs  
**(Correct)**

### **Explanation**

This is the beauty of snowflake. The cost does not depend on how many queries you run on the warehouse. It depends on which warehouse size you have chose and how long the warehouse was alive.

Question 12: Skipped

You are an account administrator and you want to use a third party product in snowflake.

Where will you go to enable the third party product?

- The third party product's webpage and contact them through contact us
- Call snowflake support to enable the product
- Enable the product through Partner Connect in Snowflake web console  
**(Correct)**

### **Explanation**

You can enable 3rd party products by navigating to partner connect on Snowflake web console

Question 13: Skipped

Compute in snowflake is provided by

- 
- The cloud providers VM instances
- 
- Only EC2 instances on AWS
- 
- Virtual warehouses  
**(Correct)**

### Explanation

Snowflake is a SaaS and runs on all the three major cloud providers AWS, Azure and GCP. It abstracts the computing resources through its virtual warehouse concept. Virtual warehouses are one or more cluster of servers that provide compute resources

Question 14: **Skipped**

Once you have selected a warehouse size, you cannot resize the size of the warehouse

- 
- True
- 
- False  
**(Correct)**

### Explanation

You can always resize a warehouse to make it bigger or smaller. However the queries that are already in process will not be able to use the resized warehouse. The resized warehouse will be used by all future queries

Question 15: **Skipped**

What is the easiest way to monitor the queries that are run on snowflake?

- 
- Create a tableau dashboard and connect to snowflake
- 
- All queries go to cloudwatch and use cloudwatch to monitor
- 
- Click on the History tab and monitor all queries that are executed in the last 14 days  
**(Correct)**

### Explanation

The history page displays all queries that are executed in last 14 days.

Question 16: **Skipped**

You are a snowflake architect hired by an organization. They want you to design their warehouse strategy. Which one of the strategy will you pick up.

- 
- Both loading and analysis of data will be done by a single warehouse to reduce cost
- 
- You do not need to use warehouse for loading data, only query analysis will require a warehouse

- 

You will recommend to use a multi-warehouse strategy. The load workload will be done by one warehouse and the query analysis workload will be done by another warehouse  
**(Correct)**

### Explanation

Multi-warehouse strategy allows you to isolate workloads by the workload type. This also helps manage access, cost accounting in a better way. Hence you should recommend multi-warehouse strategy. Please note that cost does not depend on how many warehouses you have. Compute cost depends on the size of the warehouse and the amount of time they are running.

Question 17: **Skipped**

In your snowflake environment, you have a medium warehouse. The medium warehouse is used by business to run adhoc queries. The warehouse has auto suspend set at 15 minutes. You have noticed that all of the queries that run on this warehouse finishes within a minute. What will you do to optimize cost of compute in this case?

- 

Since all the queries are completing in a minute, you should delete this warehouse and tell the users to use another existing warehouse

- 

You will tell the users that they do not need snowflake to run their queries, they should load their data in another on premise database and run query from there

- 

You will recommend to reduce the auto suspend time to 1 minute

**(Correct)**

### Explanation

All virtual warehouses can be setup with 'Auto Suspend' setting. This setting suspends the warehouse if it is idle for the time that has been setup for the configuration. In this case, all the queries are completing in 1 minute but since the auto suspend is turned on, the warehouse will keep on running for 15 minutes. Compute cost depends on how long the warehouse is running. So, you should reduce the auto suspend time setting

Question 18: **Skipped**

When a warehouse does not have enough resources available to process the queries, what happens to the incoming queries?

- 

Queries are aborted immediately

- 

Snowflake automatically resizes the warehouse

- 

The queries are queued and then executed when the resources are available again

**(Correct)**

### Explanation

If the warehouse is already overloaded with queries, the new queries are queued for execution. As and when the resources get available, the warehouse executes the query. You can monitor if your warehouse is overloaded by tracking the query overload parameter or going to the warehouse view in the snowflake web UI.

**Question 19:** Skipped

In your organization, you have Snowflake Enterprise addition. You notice that consistently, your queries are getting queued on the warehouses and delaying your ETL process. What are the possible solution options that you can think of?

- Resize the warehouse  
**(Correct)**
- Use multi-cluster warehouse  
**(Correct)**
- Set the auto-resize parameter of the warehouse
- Contact Snowflake support to increase the number of servers for the warehouse

**Explanation**

When you see that you are experiencing query overload on a consistent basis, you know that the warehouse is getting overworked. In this case, you can either resize your warehouse to make it bigger or use a multi-cluster warehouse. In a multi-cluster warehouse, you can mention the minimum and maximum servers to instantiate. Multi-cluster warehouses are an enterprise edition feature.

**Question 20:** Skipped

You can load data using PUT command through worksheets in Snowflake web UI

- True
- False  
**(Correct)**

**Explanation**

You cannot use worksheets to load data into Snowflake using PUT. PUT only works from SNOWSQL.

**Question 21:** Skipped

What are the valid data loading options in Snowflake?

- Using snowsql and sql  
**(Correct)**
- Using snowpipe  
**(Correct)**
- Using 3rd party ETL tool  
**(Correct)**
-

Using the cloud providers data upload tools

### Explanation

You can bulk load data into snowflake using SNOWSQL, Snowpipe or 3rd party ETL tools like Informatica, Talend, Pentaho etc.

Question 22: Skipped

Snowflake allows only loading of structured data

- True
- False

(Correct)

### Explanation

In snowflake, you can load and store both structured and semi-structured data like XML, JSON, AVRO, Parquet and ORC. Variant datatype is usually used store semi-structured data

Question 23: Skipped

What are the usual data loading steps in snowflake?

- Source -> Snowflake Stage -> Snowflake table
- Source -> Cloud storage -> Snowflake stage -> Snowflake table
- Source -> Snowflake temp table -> Snowflake transient table-> Snowflake permanent table

### Explanation

Data from source is usually batch loaded to Snowflake stage which is a snowflake object that loads the data as raw files in the underlying cloud storage object. From the snowflake stage, the data is copied into snowflake tables for further processing.

Question 24: Skipped

You have several CSV files loaded into your named snowflake internal stage. You want to load files from the stage into a table using pattern matching to only load uncompressed CSV files whose names include the string **sales**. Which is the command that you will use to do the same?

- ```
1. copy into mytable
2.   from @my_int_stage
3.   file_format = (format_name = myformat)
4.   pattern='.*sales.*[.]csv';
```

(Correct)

- ```
1. copy into mytable
2.   from @my_int_stage
3.   regex='.*sales.*[.]csv';
```

- ```
1. copy into mytable
2.   from @my_int_stage
```

```
3.    match_pattern='.*sales.*[.]csv';
```

### Explanation

To load data from staged files to an existing table, the files must already be staged in one of the following locations:

1. Named internal stage (or table/user stage). Files can be staged using the **PUT** command.
2. Named external stage that references an external location (Amazon S3, Google Cloud Storage, or Microsoft Azure).
3. External location (Amazon S3, Google Cloud Storage, or Microsoft Azure).

In this case we have the file staged in a named internal stage. To load files matching a pattern, we will need to use the **PATTERN** keyword.

```
[ PATTERN = '<regex_pattern>' ]
```

<https://docs.snowflake.com/en/sql-reference/sql/copy-into-table.html>

Question 25: Skipped

Snowflake is compliant with which certifications?

- HIPAA, PCI DSS, HIPAA, PCI DSS and FedRAMP  
**(Correct)**
- HIPAA, PCI DSS
- Only HIPAA
- Only FedRAMP

Question 26: Skipped

When you sign up for a snowflake trial account, what are three snowflake editions that are offered

- Free-Tier
- Standard  
**(Correct)**
- Enterprise Data Security
- Enterprise  
**(Correct)**

- Business critical  
**(Correct)**

Question 27: **Skipped**

What are the snowflake editions available as of today?

- Standard Edition,Enterprise Edition,Business Critical Edition,Virtual Private Snowflake  
**(Correct)**
- Standard Edition,Enterprise Edition,ESD,Business Critical Edition,Virtual Private Snowflake
- Standard Edition,Enterprise Edition

#### Explanation

Snowflake offers multiple editions to choose from.

The Snowflake Edition that your organization chooses determines the unit costs for the credits and the data storage you use. Other factors that impact unit costs are the regions where your Snowflake account is located and whether it is an *On Demand* or *Capacity* account:

On Demand: Usage-based pricing with no long-term licensing requirements.

Capacity: Discounted pricing based on an up-front Capacity commitment.

<https://docs.snowflake.com/en/user-guide/intro-editions.html#>

Question 28: **Skipped**

Which snowflake edition supports private communication between Snowflake and your other VPCs through AWS PrivateLink

- Standard
- Premier
- Enterprise
- Business critical  
**(Correct)**

#### Explanation

Only business critical and above supports private communication between Snowflake and your other VPCs through AWS PrivateLink or your other VNets through Azure PrivateLink

Question 29: **Skipped**

Which cloud providers are supported by snowflake

-

Google Cloud Platform  
**(Correct)**

- Azure  
**(Correct)**
- AWS  
**(Correct)**

### Explanation

A Snowflake account can be hosted on any of the following cloud platforms:

[Amazon Web Services \(AWS\)](#)

[Google Cloud Platform \(GCP\)](#)

[Microsoft Azure \(Azure\)](#)

On each platform, Snowflake provides one or more [regions](#) where the account is provisioned.

**Note : GCP support has started from 2020** - <https://www.snowflake.com/news/snowflake-announces-general-availability-on-google-cloud/>

<https://docs.snowflake.com/en/user-guide/intro-cloud-platforms.html>

Question 30: [Skipped](#)

Snowflake automatically stores data in encrypted form in all editions

- True  
**(Correct)**
- False

### Explanation

Protecting customer data is one of Snowflake's highest priorities. Snowflake encrypts all customer data by default, using the latest security standards, at no additional cost. Snowflake provides best-in-class key management, which is entirely transparent to customers. This makes Snowflake one of the easiest to use and most secure data warehouses

<https://docs.snowflake.com/en/user-guide/security-encryption.html>

Question 31: [Skipped](#)

You have created a lot of shares with role as ACCOUNTADMIN. You want to create the shares with SYSADMIN as role, so you have granted 'CREATE SHARE' privilege to SYSADMIN. How do you change the ownership of the existing shares?

-

Execute 'GRANT OWNERSHIP ON <SHARE NAME> TO ROLE SYSADMIN'

- Execute 'GRANT OWNERSHIP TO ROLE SYSADMIN ON <SHARE NAME>'
- The only option is to drop and recreate the share. Ownership of a share cannot be granted to another role  
**(Correct)**

### Explanation

The ACCOUNTADMIN role has the privileges to create a share.

You can also grant CREATE SHARE and IMPORT SHARE to other roles, enabling the tasks to be delegated to other users in the account. For more information, see Enabling non-ACCOUNTADMIN Roles to Perform Data Sharing Tasks.

Ownership of a share cannot be transferred to another role. It needs to be deleted and recreated. This is by design and is done for security reasons. Users with this role can expose any object they own.

Ownership of a share, as well as the objects in the share, may be either through a direct grant to the role or inherited from a lower-level role in the role hierarchy. For more details, see Role Hierarchy and Privilege Inheritance.

Also, it's possible for the same role to own a share and the objects in the share.

<https://docs.snowflake.com/en/user-guide/security-access-privileges-shares.html#enabling-non-accountadmin-roles-to-perform-data-sharing-tasks>

Question 32: Skipped

Select all options that are true for ORDER BY in snowflake

- All data is sorted according to the numeric byte value of each character in the ASCII table. UTF-8 encoding is supported  
**(Correct)**
- For numeric values, leading zeros before the decimal point and trailing zeros (0) after the decimal point have no effect on sort order.  
**(Correct)**
- Unless specified otherwise, NULL values are considered to be higher than any non-NULL values. As a result, the ordering for NULLS depends on the sort order:
  1. If the sort order is ASC, NULLS are returned last; to force NULLS to be first, use NULLS FIRST.
  2. If the sort order is DESC, NULLS are returned first; to force NULLS to be last, use NULLS LAST.  
**(Correct)**

### Explanation

You should at least remember the below two points. Even if this does not appear in certification, you need to know this when you will work with snowflake

1. If the sort order is **ASC**, NULLS are returned last; to force NULLS to be first, use **NULLS FIRST**.
2. If the sort order is **DESC**, NULLS are returned first; to force NULLS to be last, use **NULLS LAST**.

<https://docs.snowflake.com/en/sql-reference/constructs/order-by.html#usage-notes>

Question 33: Skipped

Column level security in Snowflake allows the application of a masking policy to a column within a table or view. Which two features are related to column level security

- Dynamic data masking  
**(Correct)**
- External tokenization  
**(Correct)**
- Data loss prevention(DLP)

### Explanation

Column level security is still in preview feature and is available in enterprise edition or above. This is a very useful feature and you will need this while designing your solution on snowflake

<https://docs.snowflake.com/en/user-guide/security-column.html#column-level-security>

Question 34: Skipped

In snowflake, what are the two major cost categories

- Storage  
**(Correct)**
- Compute  
**(Correct)**
- Cloud services
- Network Traffic

### Explanation

This is a very important thing to remember. In snowflake you pay only for storage and compute. So, everything that you do in snowflake, you must keep this in perspective. For example while designing queries, check if your query can be satisfied by the metadata store. Metadata store in snowflake is in the services layer and does not charge you for compute or storage. Running MIN() and MAX() functions on your field doesn't require any warehouse compute as it is stored in the metadata store, hence it does not cost you anything. Also queries like 'SHOW COLUMNS' does not require a running warehouse to execute as the data comes from the metadata store

Question 35: Skipped

What are the installment options available for snowflake?

- Snowflake hosted accounts on AWS  
**(Correct)**
- Hybrid on-premise and cloud
- On-premise
- Snowflake hosted accounts on Azure  
**(Correct)**

### Explanation

This is very important to know. Snowflake does not have an on-premise option. Since last year it was available on AWS and Azure. From this year onwards, it is also available on GCP

Snowflake's data warehouse is a true SaaS offering. More specifically:

1. There is no hardware (virtual or physical) for you to select, install, configure, or manage.
2. There is no software for you to install, configure, or manage.
3. Ongoing maintenance, management, and tuning is handled by Snowflake.

Snowflake **runs completely on cloud infrastructure**. All components of Snowflake's service (other than an optional command line client), run in a public cloud infrastructure.

<https://docs.snowflake.com/en/user-guide/intro-key-concepts.html#data-warehouse-as-a-cloud-service>

Question 36: Skipped

How is the data storage cost computed for snowflake?

- Based on amount of compressed data  
**(Correct)**

- Based on amount of un-compressed data
- Based on daily average of data stored  
**(Correct)**
- Amount stored on first day of month
- Amount stored on last day of month

### Explanation

Usage for **data storage** is calculated on the daily average amount of data (in bytes) stored in the system for:

1. Files staged for bulk data loading/unloading (can be stored compressed or uncompressed).
2. Database tables, including historical data for Time Travel (always compressed by Snowflake).
3. Fail-safe for database tables (always compressed by Snowflake).
4. Clones of database tables that reference data deleted in the table that owns the clones.

Please look at point #3 above. Snowflake has a storage cost for fail-safe. Now by this time you may already know that Transient tables do not have fail-safe. So for use cases where it makes sense to use transient tables, you can choose transient table to save some costs. Note that storage costs are not that much so weigh in on what you are gaining by going with transient table.

Question 37: Skipped

Which type of data incur snowflake storage cost?

- Data stored in permanent tables  
**(Correct)**
- Data retained to enable data recovery(fail-safe and time travel)  
**(Correct)**
- Cached results
- Semi structured data stored in the cloud storage(like AWS S3, GCS)

## Explanation

Storage cost is charged for below type of data

1. Files stored in Snowflake locations (i.e. user and table stages or internal named stages) for bulk data loading/unloading. The user who stages a file can choose whether or not to compress the file to reduce storage.
2. Data stored in database tables, including historical data maintained for Time Travel. Snowflake automatically compresses all data stored in tables and uses the compressed file size to calculate the total storage used for an account.
3. Historical data maintained for Fail-safe.

Question 38: Skipped

How often does snowflake release new feature

- Daily
- Weekly  
**(Correct)**
- Monthly
- Annually

## Explanation

Snowflake is committed to providing a seamless, always up-to-date experience for our users while also delivering ever-increasing value through rapid development and continual innovation.

To meet this commitment, we deploy new releases and patch releases each week. This allows us to regularly deliver service improvements in the form of new features, enhancements, and fixes. The deployments happen transparently in the background; users experience no downtime or disruption of service, and are always assured of running on the most-recent release with access to the latest features.

<https://docs.snowflake.com/en/user-guide/intro-releases.html#snowflake-releases>

Question 39: Skipped

Snowflake can carry out transformations after loading files staged by partner software (ELT)

- TRUE  
**(Correct)**
- FALSE

## Explanation

Snowflake, the cloud data platform, offers secure data sharing that eliminates the need for data extraction or transformation between departments, geographies, or partners. For primary data source loading, Snowflake works with a range of data integration partners and allows users to choose either ETL or transform data after loading (ELT). Snowflake removes the worry from data integration and allows you to focus on results.

Question 40: Skipped

From a snowflake perspective what is common about Fivetran, Informatica, Matillion, Segment, Stitch and Talend

- They are all snowflake data integration partner  
**(Correct)**
- They are all snowflake competitors
- They are all programming interface patterns

### Explanation

These products are data integration partner of snowflake and can be used to load data into snowflake

Question 41: Skipped

Looker is a snowflake business intelligence partner

- TRUE  
**(Correct)**
- FALSE

Question 42: Skipped

What is the name of the Snowflake's Command Line Interface tool?

- SnowCLI
- SnowSQL  
**(Correct)**
- SnowCMD
- SnowSpark

### Explanation

SnowSQL is the next-generation command line client for connecting to Snowflake to execute SQL queries and perform all DDL and DML operations, including loading data into and unloading data out of database tables.

SnowSQL (`snowsql` executable) can be run as an interactive shell or in batch mode through `stdin` or using the `-f` option.

<https://docs.snowflake.com/en/user-guide/snowsql.html#snowsql-cli-client>

Question 43: Skipped

What is the best way to get the latest ODBC connector for use with Snowflake?

- Download it from snowflake web UI  
**(Correct)**
- Search google and download any version of ODBC
- Compile in .NET

Question 44: Skipped

Snowflake has scalar and tabular user-defined functions (UDFs), with support for both SQL and JavaScript.

- TRUE  
**(Correct)**
- FALSE

### Explanation

Snowflake currently supports two types of UDFs, *SQL* and *JavaScript*:

A SQL UDF evaluates an arbitrary SQL expression and returns either scalar or tabular results.

A JavaScript UDF lets you use the JavaScript programming language to manipulate data and return either scalar or tabular results.

<https://docs.snowflake.com/en/sql-reference/user-defined-functions.html#udfs-user-defined-functions>

Question 45: Skipped

This object in snowflake returns a set of rows instead of a single, scalar value, and can be accessed in the FROM clause of a query

- UDTF  
**(Correct)**
- UDF
- Stored procedure

## Explanation

A UDTF returns a set of rows instead of a single, scalar value, and can be accessed in the FROM clause of a query. Snowflake supports both SQL and JavaScript UDTFs. This topic covers SQL UDTFs.

<https://docs.snowflake.com/en/sql-reference/udf-table-functions.html#sql-udtfs-user-defined-table-functions>

Question 46: Skipped

By default, result reuse is enabled in snowflake. if you want to disable it what will you do?

- Execute `ALTER SESSION SET USE_CACHED_RESULT=FALSE`  
**(Correct)**
- Execute `ALTER SESSION SET QUERY_CACHE=FALSE`
- Execute `ALTER SESSION SET QUERY_RESULT_CACHE =FALSE`

## Explanation

You will need to change it with the alter session query. This may not come in certification. But you will probably using it a lot when you actually work on a snowflake project. You will use it to test the performance once you have optimized your query. For example if a table is not partitioned well and you want to check the query performance after properly partitioning the table(either through clustering or reloading the table using natural order)

Question 47: Skipped

Each time a persisted result for a query is reused, Snowflake resets the 24-hour retention period for the result, up to a maximum of 31 days from the date and time that the query was first executed. After 31 days, the result is purged and the next time the query is submitted, a new result is generated and persisted.

- TRUE  
**(Correct)**
- FALSE

## Explanation

This is a very important point to remember. The query result of an already run query is persisted for 24 hours and every time the query is run the 24 hour window is reset for another 24 hours till 31 days

Question 48: Skipped

Snowpipe guarantees that files are loaded in the same order they are staged

- TRUE
- FALSE  
**(Correct)**

## Explanation

For each pipe object, Snowflake establishes a single queue to sequence data files awaiting loading. As new data files are discovered in a stage, Snowpipe appends them to the queue. However, multiple processes pull files from the queue; and so, while Snowpipe generally loads older files first, there is no guarantee that files are loaded in the same order they are staged.

Question 49: Skipped

Snowpipe is recommended to load data in which of the below scenario?

- You have small volume of frequent data  
**(Correct)**
- You have a huge volume of data generated as part of a batch schedule

## Explanation

Please note the key word 'SMALL VOLUME' and 'FREQUENT'. This question can be framed in multiple ways.

Question 50: Skipped

What should be your minimum snowflake edition for data sharing capability

- Standard  
**(Correct)**
- Enterprise
- Business critical

## Explanation

Data Sharing

| Feature/Service                                                                                                               | Standard | Enterprise | Business Critical | VPS |
|-------------------------------------------------------------------------------------------------------------------------------|----------|------------|-------------------|-----|
| As a data provider, securely share data with other accounts.                                                                  | ✓        | ✓          | ✓                 | ✓   |
| As a data consumer, query data shared with your account by data providers.                                                    | ✓        | ✓          | ✓                 | ✓   |
| Secure data sharing across regions and cloud platforms (through data replication)                                             | ✓        | ✓          | ✓                 | ✓   |
| Snowflake Data Marketplace, where providers and consumers meet to securely sharing data.                                      | ✓        | ✓          | ✓                 | ✓   |
| Data Exchange, a private hub of administrators, providers, and consumers that you invite to securely collaborate around data. | ✓        | ✓          | ✓                 | ✓   |

Question 51: Skipped

Select the snowflake edition that allow only a maximum of 1 day of time travel

- Standard  
**(Correct)**
- Enterprise

- VPS
- Business Critical

### Explanation

Standard [Time Travel](#) (up to 1 day) for accessing/restoring modified and deleted data is available for all editions. [Extended Time Travel](#) (up to 90 days) is available for all editions except Standard

Question 52: [Skipped](#)

Files that are already copied from the stage to the source table can be loaded again into a table cloned from the source table

- TRUE  
**(Correct)**
- FALSE

### Explanation

You do not have to mug up the answer.

Just remember what prevents the loaded files to be reloaded again. It is the load metadata. Snowflake maintains detailed metadata for each table into which data is loaded. This data expires after 64 days. Through this data, Snowflake knows that the file has been already loaded to the table. When you clone the table, this metadata is not cloned and hence you will be able to load the data again to the cloned table.

### Refresh your memory on load metadata

#### Load Metadata

Snowflake maintains detailed metadata for each table into which data is loaded, including:

Name of each file from which data was loaded

1. File size
2. ETag for the file
3. Number of rows parsed in the file
4. Timestamp of the last load for the file
5. Information about any errors encountered in the file during loading

This load metadata expires after 64 days. If the LAST\_MODIFIED date for a staged data file is less than or equal to 64 days, the COPY command can determine its load status for a given table and prevent reloading (and data duplication). The LAST\_MODIFIED date is the timestamp when the file was initially staged or when it was last modified, whichever is later.

If the LAST\_MODIFIED date is older than 64 days, the load status is still known if *either* of the following events occurred less than or equal to 64 days prior to the current date:

The file was loaded successfully.

The initial set of data for the table (i.e. the first batch after the table was created) was loaded.

However, the COPY command cannot definitively determine whether a file has been loaded already if the LAST\_MODIFIED date is older than 64 days *and* the initial set of data was loaded into the table more than 64 days earlier (*and* if the file was loaded into the table, that also occurred more than 64 days earlier). In this case, to prevent accidental reload, the command skips the file by default.

Question 53: Skipped

What is the technique called which snowflake uses to limit the number of micro-partitions retrieved as part of a query?

- Pruning  
**(Correct)**
- Selective Filter
- Indexing
- Clustering

### Explanation

#### Query Pruning

The micro-partition metadata maintained by Snowflake enables precise pruning of columns in micro-partitions at query run-time, including columns containing semi-structured data. In other words, a query that specifies a filter predicate on a range of values that accesses 10% of the values in the range should ideally only scan 10% of the micro-partitions.

For example, assume a large table contains one year of historical data with date and hour columns. Assuming uniform distribution of the data, a query targeting a particular hour would ideally scan 1/8760th of the micro-partitions in the table and then only scan the portion of the micro-partitions that contain the data for the hour column; Snowflake uses columnar scanning of partitions so that an entire partition is not scanned if a query only filters by one column.

In other words, the closer the ratio of scanned micro-partitions and columnar data is to the ratio of actual data selected, the more efficient is the pruning performed on the table.

For time-series data, this level of pruning enables potentially sub-second response times for queries within ranges (i.e. “slices”) as fine-grained as one hour or even less.

Not all predicate expressions can be used to prune. For example, Snowflake does not prune micro-partitions based on a predicate with a subquery, even if the subquery results in a constant.

Question 54: Skipped

Select two choices that are true about snowflake roles

- Snowflake users has a limit on the number of roles that they can assume
- Snowflake user can have one or more roles  
**(Correct)**
- Privileges can be directly assigned to a user
- For a particular session, only one role can be active at a given time  
**(Correct)**

Question 55: Skipped

Snowflake waits till all the servers are provisioned for a new virtual warehouse, before it executes query on that warehouse

- TRUE  
**(Correct)**
- FALSE

**Explanation**

Snowflake does not begin executing SQL statements submitted to a warehouse until all of the servers for the warehouse are successfully provisioned, unless any of the servers fail to provision:

1. If any of the servers for the warehouse fail to provision during start-up, Snowflake attempts to repair the failed server(s).
2. During the repair process, the warehouse starts processing SQL statements once 50% or more of the requested servers are successfully provisioned.

Question 56: Skipped

If you want customer dedicated virtual warehouse, which is the lowest snowflake edition that you should opt for

- Business Critical
- Enterprise
- Standard

(Correct)

### Explanation

Customer dedicated virtual warehouse is available in all editions

Question 57: Skipped

You need to contact Snowflake Support team if retrieval of data is required from fail safe

- TRUE
- FALSE

(Correct)

### Explanation

There is nothing to explain here, you need to remember this.

Question 58: Skipped

After how many days do the load metadata of a table expire?

- 64
- 14
- 7
- 365

(Correct)

### Explanation

This load metadata **expires after 64 days**. If the LAST\_MODIFIED date for a staged data file is less than or equal to 64 days, the COPY command can determine its load status for a given table and prevent reloading (and data duplication). The LAST\_MODIFIED date is the timestamp when the file was initially staged or when it was last modified, whichever is later.

Question 59: Skipped

Let us imagine you have a permanent table named EMPLOYEE, you have dropped the table. Then you created another table with the same name. If you execute the UNDROP command now, what will happen?

- UNDROP command will fail
- It will work similar to insert overwrite
- The new table will be renamed with a suffix of v1.

(Correct)

### Explanation

Let us do a hand's on and learn it ourselves.

```
CREATE OR REPLACE TABLE EMPLOYEE(EMPLOYEE_NAME VARCHAR);

INSERT INTO EMPLOYEE VALUES ('AKSHAY'),('RUPAK'),('JEETU'); -- Please note the
syntax, this is bulk insert.

DROP TABLE EMPLOYEE;

CREATE OR REPLACE TABLE EMPLOYEE(EMPLOYEE_NAME VARCHAR);

UNDROP TABLE EMPLOYEE;
```

You will get the below error

**SQL compilation error: Object 'EMPLOYEE' already exists.**

**Question 60:** Skipped

Which of the below are true with respect to snowflake web ui?

- Each worksheet can have different role, warehouse, schema and database  
**(Correct)**
- Each worksheet is a separate session  
**(Correct)**
- You will be able to run a 'PUT' command from worksheet
- You must use the same role and warehouse for all worksheets

### **Explanation**

These are the easy set of questions in the exam. You should not miss it. Please log on to the WEB UI and validate each statement. Try to run the below. Can you run it

```
CREATE STAGE CSV_STAGE;
PUT 'D:/MYFILE.CVS' @CSV_STAG;
```

This will not run and you will get below error

**SQL compilation error: The command is not supported from the UI: PUT**

Question 61: Skipped

When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the **blocked** IP address list first.

- TRUE  
**(Correct)**
- FALSE

### Explanation

When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the **blocked** IP address list first.

Do **not** add `0.0.0.0/0` to the blocked IP address list. `0.0.0.0/0` is interpreted to be “all IPv4 addresses on the local machine”. Because Snowflake resolves this list first, this would block your own access. Also, note that it is not necessary to include this IP address in the allowed IP address list.

Question 62: Skipped

Lets say you executed a transaction in a snowflake session. Due to some reason the session disconnects and the transaction now is in a detached state. The transaction cannot be committed or rolled back. The object on which the transaction was applied is also now locked. if you do not do anything and let snowflake eventually abort the transaction, how long will you need to wait?

- 60 minutes
- 15 minutes
- 4 hours  
**(Correct)**

### Explanation

If a transaction is running in a session and the session disconnects abruptly, preventing the transaction from committing or rolling back, the transaction is left in a detached state, including any locks that the transaction is holding on resources. If this happens, you might need to abort the transaction.

To abort a running transaction, the user who started the transaction or an account administrator can call the system function, `SYSTEM$ABORT_TRANSACTION`.

If the transaction is left open, Snowflake typically rolls back the transaction after it has been idle for four hours.

Question 63: Skipped

Snowflake does not support nested transactions

- TRUE  
**(Correct)**
- FALSE

### Explanation

Overlapping Transactions

This section describes overlapping transactions.

A stored procedure that contains a transaction can be called from within another transaction. The outer transaction can be in an outer stored procedure or can be outside any stored procedure.

The inner transaction is **not** treated as nested; instead, the inner transaction is **a separate transaction**. Snowflake calls these “autonomous scoped transactions” (or simply “scoped transactions”), because each transaction executes in a conceptually **independent scope**.

Note

Terminology note:

The terms “inner” and “outer” are commonly used when describing nested operations, such as nested stored procedure calls. Although Snowflake supports nested procedure calls, **Snowflake does not support nested transactions**;

<https://docs.snowflake.com/en/sql-reference/transactions.html#overlapping-transactions>

Question 64: **Skipped**

You have cloned a table. Which of the below queries will work on the cloned table?

- DROP TABLE <TABLE\_NAME>
- SELECT \* FROM <TABLE\_NAME>
- SHOW TABLES LIKE '<TABLE\_NAME>'  
**(Correct)**
- ALL OF THE ABOVE

### Explanation

Cloning creates a copy of an existing object in the system. This command is primarily used for creating **zero-copy clones** of databases, schemas, and non-temporary tables; however, it can also be used to quickly/easily create clones of other schema objects (i.e. external stages, file formats, and sequences).

The command is a variation of the object-specific `CREATE <object>` commands with the addition of the `CLONE` keyword.

When you clone a table, it is actually a regular table which shares the micro-partitions with the table from which it has been cloned

Question 65: Skipped

Select the term that is associated with compute layer?

- Query optimization
- Query planning
- Query processing  
**(Correct)**

### Explanation

Compute layer has virtual warehouses which are responsible for query processing

Question 66: Skipped

You can use the query profiler view only for completed queries

- TRUE
- FALSE  
**(Correct)**

Question 67: Skipped

You need snowflake instances in different regions. You will need to have separate account for each region

- TRUE  
**(Correct)**
- FALSE

Question 68: Skipped

In a multi-cluster warehouse, you will need to manually scale the warehouses(minimum cluster to maximum cluster)

- TRUE
- FALSE  
**(Correct)**

### Explanation

Multi-cluster warehouses enable you to scale compute resources to manage your user and query concurrency needs as they change, such as during peak and off hours. And it is 100% automated.

Question 69: Skipped

Which of the below statement will you use to recreate a specified object?

- GET\_DDL  
**(Correct)**
- GET\_SCHEMA
- SHOW SCHEMA

### Explanation

#### GET\_DDL

Returns a DDL statement that can be used to recreate the specified object. For databases and schemas, GET\_DDL is recursive, i.e. it returns the DDL statements for recreating all supported objects within the specified database/schema.

GET\_DDL currently supports the following object types:

Databases (see [CREATE DATABASE](#))

Schemas (see [CREATE SCHEMA](#))

Tables (see [CREATE TABLE](#)) (not external tables)

Views (see [CREATE VIEW](#))

Streams (see [CREATE STREAM](#))

Tasks (see [CREATE TASK](#))

Sequences (see [CREATE SEQUENCE](#))

File formats (see [CREATE FILE FORMAT](#))

Pipes (see [CREATE PIPE](#))

UDFs (see [CREATE FUNCTION](#))

Stored Procedures (see [CREATE PROCEDURE](#))

### Try it out

```
CREATE OR REPLACE TABLE EMPLOYEE(EMPLOYEE_NAME VARCHAR);
select get_ddl('TABLE', 'EMPLOYEE', true);
```

What do you get?

Question 70: Skipped

This command can be used to list streams for the current/specify database or schema, or across your entire account.

- SHOW STREAMS  
**(Correct)**
- DISPLAY STREAMS
- LIST STREAMS

### Explanation

#### SHOW STREAMS

Lists the streams for which you have access privileges. The command can be used to list streams for the current/specify database or schema, or across your entire account.

The output returns stream metadata and properties, ordered lexicographically by database, schema, and stream name (see [Output](#) in this topic for descriptions of the output columns). This is important to note if you wish to filter the results using the provided filters.

<https://docs.snowflake.com/en/sql-reference/sql/show-streams.html#show-streams>

Question 71: Skipped

When you use this parameter with SHOW STREAMS, The output also includes an additional [dropped\\_on](#) column, which displays

1. Date and timestamp (for dropped streams).
2. [NULL](#) (for active streams).

- HISTORY  
**(Correct)**
- TERSE
- DROPPED=TRUE

### Explanation

#### Parameters

[TERSE](#)

Returns only a subset of the output columns:

`created_on`

`name`

`kind` (rename of `type` column in full set of columns)

`database_name`

`schema_name`

`tableOn` (rename of `table_name` column in full set of columns)

## HISTORY

Optionally includes dropped streams that have not yet been purged (i.e. they are still within their respective Time Travel retention periods). If multiple versions of a dropped stream exist, the output displays a row for each version. The output also includes an additional `dropped_on` column, which displays:

Date and timestamp (for dropped streams).

`NULL` (for active streams).

Default: No value (dropped streams are *not* included in the output)

<https://docs.snowflake.com/en/sql-reference/sql/show-streams.html#parameters>

Question 72: Skipped

The table functions in INFORMATION\_SCHEMA can be used to return account-level usage and historical information for storage, warehouses, user logins, and queries

- TRUE  
**(Correct)**
- FALSE

## Explanation

List of Table Functions

The table functions in INFORMATION\_SCHEMA can be used to return account-level usage and historical information for storage, warehouses, user logins, and queries:

<https://docs.snowflake.com/en/sql-reference/info-schema.html#list-of-table-functions>

Question 73: Skipped

What can you expect if the filters specified in an INFORMATION\_SCHEMA query are not sufficiently selective?

- AN ERROR  
(Correct)
- WILL SHOW RESULTS WITH AN WARNING
- WILL SWITCH TO A LARGE WAREHOUSE

### Explanation

To prevent performance issues, the following error is returned if the filters specified in an INFORMATION\_SCHEMA query are not sufficiently selective:

Information schema query returned too much data. Please repeat query with more selective predicates.

<https://docs.snowflake.com/en/sql-reference/info-schema.html#general-usage-notes>

Question 74: Skipped

Information Schema Views will require a warehouse to execute whereas SHOW command does not require one.

- TRUE  
(Correct)
- FALSE

### Explanation

The INFORMATION\_SCHEMA views provides a SQL interface to the same information provided by the `SHOW <objects>` commands. You can use the views to replace these commands; however, there are some key differences to consider before switching:

| Considerations             | SHOW Commands                                                      | Information Schema Views                                                                                                                                                                                                                                                                                                                                     |
|----------------------------|--------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Warehouses                 | Not required to execute.                                           | Warehouse must be running and currently in use to query the views.                                                                                                                                                                                                                                                                                           |
| Pattern matching/filtering | Case-insensitive (when filtering using LIKE).                      | Standard (case-sensitive) SQL semantics. Snowflake automatically converts unquoted, case-insensitive identifiers to uppercase internally, so unquoted object names must be queried in uppercase in the Information Schema views.                                                                                                                             |
| Query results              | Most SHOW commands limit results to the current schema by default. | Views display all objects in the current/specify database. To query against a particular schema, you must use a filter predicate (e.g. <code>... WHERE table_schema = CURRENT_SCHEMA() ...</code> ). Note that Information Schema queries lacking sufficiently selective filters return an error and do not execute (see General Usage Notes in this topic). |

Question 75: Skipped

When unloading data into multiple files, you will use this copy option to specify the maximum size of each file created

- MAX\_FILE\_SIZE  
**(Correct)**
- MAX\_SIZE
- MAX\_FILE\_BYTES

### Explanation

#### Bulk Unloading into Single or Multiple Files

The `COPY INTO <location>` command provides a copy option (SINGLE) for unloading data into a single file or multiple files. The default is SINGLE = FALSE (i.e. unload into multiple files).

Snowflake assigns each file a unique name. The location path specified for the command can contain a filename prefix that is assigned to all the data files generated. If a prefix is not specified, Snowflake prefixes the generated filenames with `data_`.

Snowflake appends a suffix that ensures each file name is unique across parallel execution threads; e.g. `data_stats_0_1_0`.

**When unloading data into multiple files, use the `MAX_FILE_SIZE` copy option to specify the maximum size of each file created.**

Question 76: Skipped

How do you truncate a date (from a timestamp) down to the year, month, and day

- Use concatenation of the date part
- Split the date into year, month and day
- Use `date_trunc` and execute a query as below

```
1. select to_date('2015-05-08T23:39:20.123-07:00') as "DATE1",
2.        date_trunc('YEAR', "DATE1") as "TRUNCATED TO YEAR",
3.        date_trunc('MONTH', "DATE1") as "TRUNCATED TO MONTH",
4.        date_trunc('DAY', "DATE1") as "TRUNCATED TO DAY";
```

**(Correct)**

### Explanation

You do not need to remember the query(just remember the function `date_trunc`) and it will most probably not come in your certification, but this will be an useful function when you actually start working on snowflake

[https://docs.snowflake.com/en/sql-reference/functions/date\\_trunc.html#examples](https://docs.snowflake.com/en/sql-reference/functions/date_trunc.html#examples)

Question 77: Skipped

You are trying to set a variable by suing the set variable function `SET MY_VARIABLE = 'XXXX'`. You got an error as below

Assignment to 'MY\_VARIABLE' not done because value exceeds size limit of variables. Its size is 312; the limit is 256

What is the reason for this error?

- The size of string or binary variables is limited to 256 bytes  
**(Correct)**
- This is a temporary error due to insufficient memory in VM instance
- Your warehouse is not big enough to accomodate this

### Explanation

Variables can be set by executing the SQL statement `SET` or by setting the variables in the connection string when you connect to Snowflake.

The size of string or binary variables is limited to 256 bytes.

<https://docs.snowflake.com/en/sql-reference/session-variables.html#initializing-variables>

Question 78: Skipped

Snowflake supports specifying a SELECT statement instead of a table in the COPY INTO <location> command

- TRUE  
**(Correct)**
- FALSE

### Explanation

#### Bulk Unloading Using Queries

Snowflake supports specifying a SELECT statement instead of a table in the COPY INTO <location> command. The results of the query are written to one or more files as specified in the command and the file(s) are stored in the specified location (internal or external).

SELECT queries in COPY statements support the full syntax and semantics of Snowflake SQL queries, including JOIN clauses, which enables downloading data from multiple tables.

<https://docs.snowflake.com/en/user-guide/data-unload-overview.html#bulk-unloading-using-queries>

Question 79: Skipped

Using **COPY INTO <location>** command, you can unload data from a table (or query) into the below locations.

- Named internal stage (or table/user stage)  
**(Correct)**
- Named external stage  
**(Correct)**
- External location (Amazon S3, Google Cloud Storage, or Microsoft Azure)  
**(Correct)**
- Local drive

### Explanation

#### **COPY INTO <location>**

Unloads data from a table (or query) into one or more files in one of the following locations:

1. Named internal stage (or table/user stage). The files can then be downloaded from the stage/location using the **GET command**.
2. Named external stage that references an external location (Amazon S3, Google Cloud Storage, or Microsoft Azure).
3. External location (Amazon S3, Google Cloud Storage, or Microsoft Azure).

Question 80: Skipped

To download files from the stage/location loaded through **COPY INTO <LOCATION>** command, you will use

- GET  
**(Correct)**
- UNLOAD
- COPY INTO

### Explanation

#### **GET**

Downloads data files from one of the following Snowflake stages to a local directory/folder on a client machine:

Named internal stage.

Internal stage for a specified table.

Internal stage for the current user.

Typically, this command is executed after using the **COPY INTO <location>** command to unload data from a table into a Snowflake stage.

<https://docs.snowflake.com/en/sql-reference/sql/get.html#get>

Question 81: Skipped

GET does **not** support downloading files from external stages

- FALSE
- TRUE  
**(Correct)**

### Explanation

**GET does not support downloading files from external stages.** To download files from external stages, use the utilities provided by the cloud service.

The following Snowflake clients do not support GET:

Go Snowflake Driver

.NET Driver

Node.js Driver

**The ODBC driver supports GET with Snowflake accounts hosted on the following platforms:**

Amazon Web Services (using ODBC Driver Version 2.17.5 and higher).

Google Cloud Platform (using ODBC Driver Version 2.21.5 and higher).

Microsoft Azure (using ODBC Driver Version 2.20.2 and higher).

Question 82: Skipped

You will use this parameter to specify the number of threads to use for downloading the files using GET command

- PARALLEL = <INTEGER>  
**(Correct)**

- DOWNLOAD\_PARALLEL=<INTEGER>
- DUMP=<INTEGER>

## Explanation Optional Parameters

### **PARALLEL = *integer***

Specifies the number of threads to use for downloading the files. The granularity unit for downloading is one file.

Increasing the number of threads can improve performance when downloading large files.

Supported values: Any integer value from **1** (no parallelism) to **99** (use 99 threads for downloading files).

Default: **10**

### **PATTERN = '*regex\_pattern*'**

Specifies a regular expression pattern for filtering files to download.

Default: No value (all files in the specified stage are downloaded)

<https://docs.snowflake.com/en/sql-reference/sql/get.html#optional-parameters>

Question 83: Skipped

Following commands cannot be executed from worksheets

- PUT  
**(Correct)**
- GET  
**(Correct)**
- SHOW
- LIST <STAGE>

Question 84: Skipped

If file format options are specified in multiple locations, the load operation applies the options in the following order of precedence.

- 1. COPY INTO TABLE statement.

- 2. Stage definition.
  - 3. Table definition.
- (Correct)**

- 1. Stage definition.  
2. COPY INTO TABLE statement.  
3. Table definition.
- 1. COPY INTO TABLE statement.  
2. Table definition.  
3. Stage definition.

### Explanation

If file format options are specified in multiple locations, the load operation applies the options in the following order of precedence:

1. COPY INTO TABLE statement.
2. Stage definition.
3. Table definition.

Question 85: **Skipped**

Which are the two metadata columns for staged files

- METADATA\$FILENAME**  
**(Correct)**
- METADATA\$FILEFORMAT**
- METADATA\$FILE\_ROW\_NUMBER**  
**(Correct)**

### Explanation

#### Metadata Columns

Currently, the following metadata columns can be queried or copied into tables:

#### **METADATA\$FILENAME**

Name of the staged data file the current row belongs to. Includes the path to the data file in the stage.

#### **METADATA\$FILE\_ROW\_NUMBER**

Row number for each record in the container staged data file.

<https://docs.snowflake.com/en/user-guide/querying-metadata.html#metadata-columns>

Question 86: Skipped

Metadata columns for staged files can only be queried by name

- TRUE  
**(Correct)**
- FALSE

### Explanation

#### Query Limitations

Metadata cannot be inserted into existing table rows.

Metadata columns can only be queried by name; as such, they are not included in the output of any of the following statements:

`SELECT *`

`SHOW <objects>`

`DESCRIBE <object>`

Queries on `INFORMATION_SCHEMA` views

<https://docs.snowflake.com/en/user-guide/querying-metadata.html#query-limitations>

Question 87: Skipped

Only named stages (internal or external) and user stages are supported for `COPY` transformations

- TRUE  
**(Correct)**
- FALSE

### Explanation

#### Usage Notes

This section provides usage information for transforming staged data files during a load.

#### Supported Stages

Only named stages (internal or external) and user stages are supported for `COPY` transformations.

#### Supported File Formats

The following file format types are supported for COPY transformations:

CSV

JSON

Avro

ORC

Parquet

XML

Question 88: Skipped

The VALIDATION\_MODE parameter does not support COPY statements that transform data during a load.

- FALSE
- TRUE  
**(Correct)**

### Explanation

VALIDATION\_MODE Parameter

The VALIDATION\_MODE parameter does not support COPY statements that transform data during a load.

<https://docs.snowflake.com/en/user-guide/data-load-transform.html#validation-mode-parameter>

Question 89: Skipped

Following transformations are not supported in a COPY command. Select three.

- FLATTEN  
**(Correct)**
- JOIN  
**(Correct)**
- GROUP BY  
**(Correct)**
- NVL

### Explanation

Note that COPY transformations do ***not*** support the FLATTEN function, or JOIN or GROUP BY (aggregate) syntax.

This is a very important topic, please read it

<https://docs.snowflake.com/en/user-guide/data-load-transform.html#transforming-data-during-a-load>

Question 90: Skipped

Snowflake recommends to compress your data files when you are loading large data sets.

- TRUE  
**(Correct)**
- FALSE

### Explanation

We recommend that you compress your data files when you are loading large data sets. See CREATE FILE FORMAT for the compression algorithms supported for each data type.

When loading compressed data, specify the compression method for your data files. The COMPRESSION file format option describes how your data files are ***already*** compressed in the stage. Set the COMPRESSION option in one of the following ways:

1. As a file format option specified directly in the COPY INTO <table> statement.
2. As a file format option specified for a named file format or stage object. The named file format/stage object can then be referenced in the COPY INTO <table> statement.

Question 91: Skipped

Snowpipe is a serverless function

- TRUE  
**(Correct)**
- FALSE

### Explanation

Automated data loads leverage event notifications for cloud storage to inform Snowpipe of the arrival of new data files to load. Snowpipe copies the files into a queue, from which they are loaded into the target table in a continuous, serverless fashion based on parameters defined in a specified pipe object.

Question 92: Skipped

When calling the rest endpoints in snowpipe, below authentication methods are supported

- key pair authentication with JSON Web Token  
**(Correct)**

- user id and password
- Both

### Explanation

Authentication

### Bulk data load

Relies on the security options supported by the client for authenticating and initiating a user session.

### Snowpipe

**When calling the REST endpoints:** Requires key pair authentication with JSON Web Token (JWT). JWTs are signed using a public/private key pair with RSA encryption.

Question 93: Skipped

Load history of Snowpipe expires after how many days.

- 30
- 14 (Correct)
- 64

### Explanation

Load History

### Bulk data load

Stored in the metadata of the target table for 64 days. Available upon completion of the COPY statement as the statement output.

### Snowpipe

Stored in the metadata of the pipe for 14 days. Must be requested from Snowflake via a REST endpoint, SQL table function, or ACCOUNT\_USAGE view.

Question 94: Skipped

SnowPipe can load a file with same name if it has been modified later.

- TRUE
- FALSE

(Correct)

## Explanation

### Data Duplication

Snowpipe uses file loading metadata associated with each *pipe object* to prevent reloading the same files (and duplicating data) in a table. This metadata stores the path (i.e. prefix) and name of each loaded file, and prevents loading files with the same name even if they were later modified (i.e. have a different eTag).

Question 95: Skipped

SNOWPIPE REST APIs support both internal and external stage

- TRUE  
(Correct)
- FALSE

## Explanation

Snowpipe supports both internal (Snowflake) stages and external stages, i.e. S3 buckets. Please note that this is true for REST APIs. SNOWPIPE Auto ingest supports only external stage

Question 96: Skipped

Which of the below are SNOWPIPE REST APIs. Select three.

- `insertFiles`  
(Correct)
- `insertReport`  
(Correct)
- `loadHistoryScan`  
(Correct)
- `loadHistoryProgress`

## Explanation

SNOWPIPE has three APIs

### Data File Ingestion

1. Endpoint: `insertFiles`

### Load History Reports

1. Endpoint: `insertReport`

## 2. Endpoint: `loadHistoryScan`

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-rest-apis.html#snowpipe-rest-api>

Question 97: Skipped

A successful response from the SNOWPIPE insertFiles API means that the files are ingested.

- TRUE
- FALSE  
**(Correct)**

### Explanation

Endpoint: `insertFiles`

Informs Snowflake about the files to be ingested into a table. A successful response from this endpoint means that Snowflake has recorded the list of files to add to the table. **It does not necessarily mean the files have been ingested**

Question 98: Skipped

The insertReport SNOWPIPE API can retrieve file ingestion events and report it. The events are retained for infinite time until deleted manually.

- TRUE
- FALSE  
**(Correct)**

### Explanation

The SNOWPIPE REST API topic is relevant for both SNOWPRO CORE and SNOWPRO ARCHITECT certifications. Please read this chapter thoroughly

<https://docs.snowflake.com/en/user-guide/data-load-snowpipe-rest-apis.html#snowpipe-rest-api>

Endpoint: `insertReport`

Retrieves a report of files submitted via `insertFiles` whose contents were recently ingested into a table. Note that for large files, this may only be part of the file.

Note the following limitations for this endpoint:

1. The 10,000 most recent events are retained.
2. Events are retained for a maximum of 10 minutes.

An event occurs when data from a file submitted via `insertFiles` has been committed to the table and is available to queries. The `insertReport` endpoint can be thought of like the UNIX command tail. By calling this command repeatedly, it is possible to see the full history of events on a pipe over time. Note that the command must be called often enough to not miss events. How often depends on the rate files are sent to `insertFiles`.

Question 99: Skipped

To help avoid exceeding the rate limit (error code 429), snowflake recommends relying more heavily on `insertReport` than `loadHistoryScan`

- TRUE  
**(Correct)**
- FALSE

### Explanation

Endpoint: `loadHistoryScan`

Fetches a report about ingested files whose contents have been added to table. Note that for large files, this may only be part of the file. This endpoint differs from `insertReport` in that it views the history between two points in time. There is a maximum of 10,000 items returned, but multiple calls can be issued to cover the desired time range.

### Important

This endpoint is rate limited to avoid excessive calls. To help avoid exceeding the rate limit (error code 429), we recommend relying more heavily on `insertReport` than `loadHistoryScan`. When calling `loadHistoryScan`, specify the most narrow time range that includes a set of data loads. For example, reading the last 10 minutes of history every 8 minutes would work well. Trying to read the last 24 hours of history every minute will result in 429 errors indicating a rate limit has been reached. The rate limits are designed to allow each history record to be read a handful of times.

Question 100: Skipped

You have a warehouse. It ran for 62 seconds and then was suspended. After that it resumed and ran for 20 seconds. For how many seconds will you be billed

- 122  
**(Correct)**
- 20
- 92

### Explanation

When a virtual warehouse starts for the first time, the first one minute is anyway charged after that the billing is per second

Warehouses are only billed for credit usage when they are running. When a warehouse is suspended, it does not accrue any credit usage.

The credit numbers shown here are for a full hour of usage; however, credits are billed per-second, with a 60-second (i.e. 1-minute) minimum:

Each time a warehouse is started or resized to a larger size, the warehouse is billed for 1 minute's worth of usage based on the hourly rate shown above.

After 1 minute, all subsequent billing is per-second.

Stopping and restarting a warehouse within the first minute does not change the amount billed; the minimum billing charge is 1 minute.

When a warehouse is increased in size, credits are billed only for the **additional** servers that are provisioned. For example, changing from Small (2) to Medium (4) results in billing charges for 1 minute's worth of 2 credits.

<https://docs.snowflake.com/en/user-guide/credits.html#virtual-warehouse-credit-usage>

[Continue](#)

[Retake test](#)

**Question 2: Incorrect**

Snowflake network policies currently support both Internet Protocol versions 4 and 6 (i.e., IPv4 as well as IPv6). (True/False)

FALSE

(Correct)

TRUE

(Incorrect)

**Explanation**

Network policies currently support only Internet Protocol version 4 (i.e. IPv4) addresses.

**Question 4: Incorrect**

Which roles can use SQL to view the task history within a specified date range?  
(Select all that apply)

Role that has the global MONITOR EXECUTION privilege

(Correct)

Account Administrator (ACCOUNTADMIN)

(Correct)

Task Owner having OWNERSHIP privilege on a task

(Correct)

**Explanation**

All of these roles can use SQL to view the task history within a specified date range. To view the run history for a single task: Query the TASK\_HISTORY table function (in the Snowflake Information Schema). To view details on a DAG run that is currently scheduled or is executing: Query the CURRENT\_TASK\_GRAPHS table function (in the Snowflake Information Schema). To view the history for DAG runs that executed successfully, failed, or were canceled in the past 60 minutes: Query the COMPLETE\_TASK\_GRAPHS table function (in the Snowflake Information Schema). Once the COMPLETE\_TASK\_GRAPHS

Question 6: **Incorrect**

What actions can a consumer perform on a share? (Select 2)

- |                                                                                                                      |             |
|----------------------------------------------------------------------------------------------------------------------|-------------|
| <input type="checkbox"/> Re-share the share                                                                          |             |
| <input type="checkbox"/> Execute Time Travel on a share                                                              |             |
| <input checked="" type="checkbox"/> Copy shared data into another table in their own account with<br>CREATE TABLE AS | (Correct)   |
| <input checked="" type="checkbox"/> Clone a share                                                                    | (Incorrect) |
| <input checked="" type="checkbox"/> Query the shared data and join it with an existing table in their<br>own account | (Correct)   |
| <input checked="" type="checkbox"/> Import the same share to more than one database                                  | (Incorrect) |

**Explanation**

Shared databases are **read-only**. A consumer cannot UPDATE a share. However, the consumer can do a CREATE TABLE AS to make a point-in-time copy of the data that's been shared. The consumer cannot clone and re-share a share or forward it. And also, time travel data on a share is not available to the consumer. A share can be imported into one database.

**Note:** In the exam, you may be asked for Reader Account as well.

Question 7: **Incorrect**

You have a dashboard that connects to Snowflake via JDBC. The dashboard is refreshed hundreds of times per day. The data is very stable, only changing once or twice per day. The query run by the dashboard user never changes. How will Snowflake manage changing and non-changing data? Mark all true statements.

- |                                                                                                                                                     |             |
|-----------------------------------------------------------------------------------------------------------------------------------------------------|-------------|
| <input checked="" type="checkbox"/> Snowflake will re-use data from the Results Cache as long as it is<br>still the most up-to-date data available. | (Correct)   |
| <input checked="" type="checkbox"/> Snowflake will show the most up-to-date data each time the<br>dashboard is refreshed.                           | (Correct)   |
| <input checked="" type="checkbox"/> Snowflake will compile result cache data from all user results, so<br>no warehouse is needed.                   | (Incorrect) |
| <input type="checkbox"/> Snowflake will spin up a warehouse only if the underlying data has<br>changed.                                             | (Correct)   |

**Explanation**

Until data has not changed and the query is the same - Snowflake reuses the data from the cache. Please note, Each time the persisted result for a query is reused, Snowflake resets the 24-hour retention period for the result up to a maximum of 31 days from the date and time that the query was first executed. After 31 days, the result is purged, and the next time the query is submitted, a new result is generated and persisted.

Question 9: **Incorrect**

John wants to load data files from an external stage to Snowflake. He has split the large file into smaller 100 - 250 MB data files, and there is a total of 16 smaller data files. What warehouse size would you recommend him to use for loading these data files quickly and cost-effectively?

L

S

(Correct)

XS

(Incorrect)

XL

M

**Explanation**

XS sized warehouse can load eight files parallelly. S sized warehouse can load sixteen files parallelly. M sized warehouse can load thirty-two files parallelly. L sized warehouse can load sixty-four files parallelly. XL sized warehouse can load one hundred twenty-eight files parallelly and so on.

Question 11: **Incorrect**

Which object parameter can users with the ACCOUNTADMIN role use to set the minimum retention period for their account?

MIN\_DATA\_RETENTION\_TIME\_IN\_HOURS

DATA\_RETENTION\_TIME\_IN\_DAYS

(Incorrect)

MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS

(Correct)

DATA\_RETENTION\_TIME\_IN\_MIN\_DAYS

**Explanation**

The MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS account parameter can be set by users with the ACCOUNTADMIN role to set a minimum retention period for the account. This parameter does not alter or replace the DATA\_RETENTION\_TIME\_IN\_DAYS parameter value. However, it may change the effective data retention time. When this parameter is set at the account level, the effective minimum data retention period for an object is determined by MAX(DATA\_RETENTION\_TIME\_IN\_DAYS, MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS).

Question 14: **Incorrect**

UDF runs with either the caller's or the owner's rights. (TRUE / FALSE)

TRUE

(Incorrect)

FALSE

(Correct)

**Explanation**

UDF only runs as the function owner. A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time.

A caller's rights stored procedure runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query.

An owner's rights stored procedure runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table.

At the time that the stored procedure is created, the creator specifies whether the procedure runs with the owner's rights or the caller's rights. The default is owner's rights.

Question 15: **Incorrect**

Suppose you have an auto-scaling mode setup with a Standard policy. In what situation does Snowflake spin up an additional cluster?

Only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes.

(Incorrect)

The first cluster starts immediately when either a query is queued or the system detects that there's one more query than the currently-running clusters can execute.

(Correct)

**Explanation**

In Standard Scaling policy, the first cluster starts immediately when either a query is queued, or the system detects that there's one more query than the currently-running clusters can execute.

Each successive cluster waits to start 20 seconds after the prior one has started. For example, if your warehouse is configured with ten max clusters, it can take 200+ seconds to start all 10 clusters.

Question 20: **Incorrect**

Which of these Sampling methods does Snowflake support?

Sample a fixed, specified number of rows (Correct)

Sample exact rows of the table with the specified sequence keys (Incorrect)

Sample a fraction of the table with a specified probability of including a given row (Correct)

**Explanation**

SAMPLE / TABLESAMPLE returns a subset of rows sampled randomly from the specified table. The following sampling methods are supported: **Sample a fraction of a table**, with a specified probability for including a given row. The number of rows returned depends on the size of the table and the requested probability. A seed can be specified to make the sampling deterministic. **Sample a fixed, specified number of rows**. The exact number of specified rows is returned unless the table contains fewer rows. SAMPLE and TABLESAMPLE are synonymous and can be used interchangeably.

Question 22: **Incorrect**

Which Snowsight interface does help in setting up Multi-factor authentication (MFA)?

Left Nav interface (Incorrect)

Account Selector Interface

Admin Interface

User Menu Interface (Correct)

You can not setup Multi-factor authentication (MFA) using Snowsight interface

**Explanation**

There are three interfaces in Snowsight. Left Nav, User Menu, and Account Selector.

**Left Navigation** consists of Worksheets, Dashboards, Data, Marketplace, Activity, Admin, Help & Support.

**User Menu** lets you Switch Roles, Profile including multi-factor authentication (MFA), Partner Connect, Documentation, Support and Sign Out.

The **account selector**, located at the bottom of the left nav, lets you sign in to other Snowflake accounts.

Question 23: **Incorrect**

A role inherits all the privileges of those higher in the hierarchy. (True / False)

FALSE

(Correct)

TRUE

(Incorrect)

**Explanation**

A role inherits all the privileges of its underlying roles (those "lower" in the hierarchy).

- ACCOUNTADMIN inherits privileges from SECURITYADMIN
- USERADMIN, SYSADMIN, and PUBLIC.
- SECURITYADMIN inherits privileges from USERADMIN and PUBLIC.
- USERADMIN and SYSADMIN inherit privileges from PUBLIC
- PUBLIC inherits nothing.

Question 27: **Incorrect**

Snowflake stores metadata about all rows stored in a micro-partition, including (Select 3)

The number of distinct values

(Correct)

Additional properties used for both optimization and efficient query processing

(Correct)

The range of values for each of the columns in the micro-partition

(Correct)

The number of similar values

The range of values for the first column in the micro-partition

(Incorrect)

**Explanation**

Micro-partitioning is automatically performed on all Snowflake tables. Tables are transparently partitioned using the Ordering of the data as inserted/loaded. Snowflake stores metadata about all rows stored in a micro-partition, including:

- The range of values for each of the columns in the micro-partition.
- The number of distinct values.
- Additional properties used for both optimization and efficient query processing.

Question 31: **Incorrect**

Is it possible to create a user without a password?

No

(Incorrect)

Yes

(Correct)

**Explanation**

Yes, it is possible to create a user in Snowflake without a password. We cannot use the Snowflake web interface to create users with no passwords or remove passwords from existing users, and we must use CREATE USER or ALTER USER. Without a password in Snowflake, a user cannot log in using Snowflake authentication and must use federated authentication instead.

Question 33: **Incorrect**

What key insights can we get from the Explain plan in Snowflake? (Select 3)

Estimated Query Time

(Incorrect)

Join Types

(Correct)

Join Ordering

(Correct)

Exact Query Time

Partition Pruning

(Correct)

**Explanation**

The key insights that the explain plan gives us in its results output are information on **partition pruning, join ordering, and join types**.

The explain plan is a useful tool for determining the efficiency of your query. It's a command that compiles your query to figure out all the steps Snowflake would have to work through if it were actually to run the query.

Question 38: **Incorrect**

What size of the virtual warehouse needs to be created by the sysadmin while loading using Snowpipe? (Select the best answer)

4XL Size

None of these

(Correct)

XS Size

(Incorrect)

L Size

M Size

**Explanation**

Snowpipe uses compute resources provided and managed by Snowflake (i.e. a serverless compute model). These Snowflake-provided resources are automatically resized and scaled up or down as required, and are charged and itemized using per-second billing. Data ingestion is charged based upon the actual workloads. User doesn't need to create any warehouse as it is taken care by Snowflake.

Question 39: **Incorrect**

Federated Authentication is supported by:

Standard Edition

(Correct)

Enterprise Edition

(Correct)

VPS

(Correct)

Business Critical

(Correct)

**Explanation**

Federated authentication is supported by all of the Snowflake editions.

Question 42: **Incorrect**

What are the security layers that Snowflake takes care of?

|                                                     |           |
|-----------------------------------------------------|-----------|
| <input checked="" type="checkbox"/> Authentication  | (Correct) |
| <input type="checkbox"/> Access                     | (Correct) |
| <input checked="" type="checkbox"/> Authorization   | (Correct) |
| <input checked="" type="checkbox"/> Data Protection | (Correct) |
| <input type="checkbox"/> Infrastructure             |           |

Explanation

Infrastructure Security is managed by the cloud provider.

Question 43: **Incorrect**

The automatic refresh of metadata of the directory table in the cloud storage does not incur any charges to Snowflake Customers. (True/False)

|                                       |             |
|---------------------------------------|-------------|
| <input checked="" type="radio"/> TRUE | (Incorrect) |
| <input type="radio"/> FALSE           | (Correct)   |

Explanation

Snowflake customers' charges include an overhead to manage event notifications for automatically refreshing directory table metadata. This overhead increases in relation to the number of files added in cloud storage for customers' stages that include directory tables. Snowflake charges 0.06 credits per 1000 event notifications received.

Question 45: **Incorrect**

In a Snowflake federated environment, Snowflake serves as the Identity provider (IdP). (True/False)

|                                       |             |
|---------------------------------------|-------------|
| <input type="radio"/> FALSE           | (Correct)   |
| <input checked="" type="radio"/> TRUE | (Incorrect) |

Explanation

In a Snowflake federated environment, Snowflake serves as the Service Provider (SP). The external, independent entity like Okta serves as the Identify Provider (IdP)

Question 52: **Incorrect**

Snowflake supports SQL UDFs that return a set of rows. Which keyword in CREATE FUNCTION statement does need to be specified to enable UDF (i.e., UDTF) to return a set of rows?

MULTIPLE

ROWS

SCALAR

(Incorrect)

TABLE

(Correct)

**Explanation**

TABLE keyword after RETURNS needs to be specified to create a UDTF (user-defined table function). Example :

```
create function t()
returns table(msg varchar)
as
$$
select 'Hello'
union
select 'World'
$$;
```

Remember - UDF returns a singular scalar value or if defined as a TABLE function, a set of rows. If you see UDTF in the exam, that simply means UDF that returns a set of rows.

Question 54: **Incorrect**

**What is the best way to analyze the optimum warehouse size?**

- Execute queries of widely-varying size and/or complexity on the same warehouse (Incorrect)
- Execute relatively homogeneous queries (size, complexity, data sets, etc.) on the same warehouse (Correct)

#### Explanation

To achieve the best results, **try to execute relatively homogeneous queries (size, complexity, data sets, etc.) on the same warehouse**; executing queries of widely-varying size and/or complexity on the same warehouse makes it more difficult to analyze warehouse load, which can make it more difficult to select the best size to match the size, composition, and number of queries in your workload.

Question 58: **Incorrect**

**What are all operations performed using Snowflake SQL API?**

- Calling stored procedures that returns a table (Incorrect)
- Cancel the execution of a statement (Correct)
- Check the status of the execution of a statement (Correct)
- Fetch query results concurrently (Correct)
- Submit SQL statements for execution (Correct)

#### Explanation

The Snowflake SQL API provides operations that we can use to:

- Submit SQL statements for execution.
- Check the status of the execution of a statement.
- Cancel the execution of a statement.
- Fetch query results concurrently.

Currently, Snowflake SQL API has limitations for the call command with stored procedures that return a table (stored procedures with the RETURNS TABLE clause).

Question 60: **Incorrect**

Which of these SQL functions helps extract the path of a staged file relative to its location in the stage using the stage name and absolute file path in cloud storage as inputs?

GET\_PRESIGNED\_URL

BUILD\_STAGE\_FILE\_URI

GET\_ABSOLUTE\_PATH

(Incorrect)

GET\_RELATIVE\_PATH

(Correct)

GET\_STAGE\_LOCATION

BUILD\_SCOPED\_FILE\_URL

**Explanation**

GET\_RELATIVE\_PATH extracts the path of a staged file relative to its location in the stage using the stage name and absolute file path in cloud storage as inputs.

Question 64: **Incorrect**

Scoped URL is ideal for

Ideal for custom applications that require access to unstructured data files

None of these

Ideal for business intelligence applications or reporting tools  
that need to display the unstructured file contents

(Incorrect)

Ideal for use in custom applications, providing unstructured data  
to other accounts via a share

(Correct)

**Explanation**

**Scoped URL:** Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. **Ideal for use in custom applications, providing unstructured data to other accounts via a share, or for downloading and ad hoc analysis of unstructured data via Snowsight.** **File URL:** URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. **Ideal for custom applications that require access to unstructured data files.**

**Pre-signed URL:** Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable. **Ideal for business intelligence applications or reporting tools that need to display unstructured file contents.**

Question 65: **Incorrect**

Which data types are not supported by the Search Optimization Service? (Select 4)

- |                                                                      |             |
|----------------------------------------------------------------------|-------------|
| <input type="checkbox"/> VARCHAR                                     |             |
| <input checked="" type="checkbox"/> Floating-point data types        | (Correct)   |
| <input checked="" type="checkbox"/> Semi-structured data types       | (Correct)   |
| <input type="checkbox"/> Fixed-point numbers (e.g. INTEGER, NUMERIC) |             |
| <input checked="" type="checkbox"/> DATE, TIME, and TIMESTAMP        | (Incorrect) |
| <input checked="" type="checkbox"/> BINARY                           | (Incorrect) |

Explanation

The search optimization service currently supports equality predicate and IN list predicate searches for the following data types: Fixed-point numbers (e.g. INTEGER, NUMERIC). DATE, TIME, and TIMESTAMP. VARCHAR. BINARY. Currently, the search optimization service does not support floating point data types, semi-structured data types, or other data types not listed above.

*[Important term for exam: EQUALITY for Search Optimization Service]*

Question 67: **Incorrect**

John has a SECURITYADMIN role. He created a custom DBA\_ROLE and granted the SYSADMIN role to DBA\_ROLE. Then, John created a user, 'Monica.' John then granted DBA\_ROLE to Monica. Monica creates a Database Monica\_DB. Monica then created a Table T1 in Monica\_DB under the PUBLIC schema. What should John do to access Table T1, created by Monica?

- |                                                                                                                 |             |
|-----------------------------------------------------------------------------------------------------------------|-------------|
| <input checked="" type="radio"/> USE ROLE SECURITYADMIN; USE DATABASE monica_db; Select * from t1;              | (Incorrect) |
| <input type="radio"/> GRANT ROLE DBA_ROLE TO John; USE DATABASE monica_db; Select * from t1;                    |             |
| <input type="radio"/> USE ROLE dba_role; USE DATABASE monica_db; Select * from t1;                              |             |
| <input type="radio"/> GRANT ROLE DBA_ROLE TO John; USE ROLE DBA_ROLE; USE DATABASE monica_db; Select * from t1; | (Correct)   |

Explanation

It does not matter if John has created the DBA\_ROLE. If John wants to access the object created by DBA\_ROLE, he needs to grant DBA\_ROLE to himself.

Question 68: **Incorrect**

Snowflake supports \_\_\_\_\_

SOAP for unstructured data

REST API for unstructured data

(Correct)

Both of these

(Incorrect)

None of these

Explanation

Currently, Snowflake only supports REST API for unstructured data.

Question 70: **Incorrect**

Which of these are types of the stream? (Select 3)

External

Update-only

(Incorrect)

Append-only

(Correct)

Insert-only

(Correct)

Standard

(Correct)

Explanation

The following stream types are available based on the metadata recorded by each:

**Standard** - Supported for streams on tables, directory tables, or views. A standard (i.e., delta) stream tracks all DML changes to the source object, including inserts, updates, and deletes (including table truncates).

**Append-only** - Supported for streams on standard tables, directory tables, or views. An append-only stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded.

**Insert-only** - Supported for streams on external tables only. An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e., no-ops).

Question 71: **Incorrect**

Fail-safe helps access historical data after the Time Travel retention period has ended.  
(True/False)

FALSE

(Correct)

TRUE

(Incorrect)

**Explanation**

**Fail-safe is not provided as a means for accessing historical data after the Time Travel retention period has ended.** It is for use only by Snowflake to recover data that may have been lost or damaged due to extreme operational failures. Data recovery through Fail-safe may take from several hours to several days to complete.

Question 73: **Incorrect**

If you want to create a warehouse that remains in a suspended state initially, which property do you need to set for that warehouse?

AUTO\_SUSPEND = 0

AUTO\_RESUME = TRUE

AUTO\_SUSPEND = TRUE

(Incorrect)

AUTO\_RESUME = FALSE

INITIALLY\_SUSPENDED = TRUE

(Correct)

**Explanation**

**INITIALLY\_SUSPENDED = TRUE | FALSE** Specifies whether the warehouse is created initially in the 'Suspended' state. The valid values are TRUE and FALSE.  
TRUE: The warehouse is created, but suspended.      FALSE: The warehouse starts running after it is created.    Default is FALSE

Question 74: **Incorrect**

What does OVERWRITE parameter do with the INSERT command?

It drops the table, recreates, and inserts the values into the table.

It specifies that the target table should be truncated before inserting the values into the table.

(Correct)

It de-duplicates while inserting and skips the insert if there is an exact similar record in the table.

(Incorrect)

It helps ignore any errors while inserting the values into the table.

**Explanation**

**OVERWRITE** specifies that the target table should be truncated before inserting the values into the table. Note that specifying this option does not affect the access control privileges on the table.

Question 76: **Incorrect**

Which is the fastest option for selecting staged data files to load from a stage?

- By path (internal stages) / prefix (Amazon S3 bucket)
- Using pattern matching to identify specific files by pattern (Incorrect)
- Specifying a list of specific files to load (Correct)

**Explanation**

Of the three options for identifying/specifying data files to load from a stage, **providing a discrete list of files is generally the fastest**; however, the FILES parameter supports a maximum of 1,000 files, meaning a COPY command executed with the FILES parameter can only load up to 1,000 files. Example: copy into load1 from @%load1/data1/files=('test1.csv', 'test2.csv', 'test3.csv', 'test4.csv')

Question 82: **Incorrect**

VARIANT is used to FLATTEN hierarchical data. (True / False)

- TRUE (Incorrect)
- FALSE (Correct)

**Explanation**

VARIANT is a data type that can hold a value of any other data type (including ARRAY and OBJECT). VARIANT is used to build and store hierarchical data. VARIANT is not a function to FLATTEN. FLATTEN is a table function that is used to produce a lateral view of a VARIANT, OBJECT, or ARRAY column.

Question 83: **Incorrect**

You can create a user-level network policy using \_\_\_\_\_

Snowsight

Classic Web Interface

SQL

(Correct)

Only Snowflake Support can create the Account level Network Policy

(Incorrect)

**Explanation**

You can set a policy at the user level, but this can only be done through the SQL commands.

To activate a network policy for an individual user, set the NETWORK\_POLICY parameter for the user using ALTER USER.

Question 88: **Incorrect**

Which are the required parameters for creating a Network Policy? (Select 2)

Blocked IP Addresses

(Incorrect)

Allowed IP Addresses

(Correct)

Policy Name

(Correct)

Comment

**Explanation**

**Policy Name and Allowed IP Addresses are mandatory for creating a Network Policy.**

Blocked IP Addresses and Comment Parameters are optional.

**Please Note - If you provide both Allowed IP Addresses and Blocked IP Addresses, Snowflake applies the Blocked List first.**

Question 89: **Incorrect**

Choose the false statements. (Select 2)

Results are stored in the Cloud Storage layer in the case of Local Disk Cache. (Correct)

Group and Execute similar queries on the same virtual warehouse to maximize local disk cache reuse, for performance and cost optimization.

If Virtual Warehouse is suspended, then results in Local Disk Cache will be lost

Group and Execute similar queries on the different virtual warehouses to maximize local disk cache reuse, for performance and cost optimization. (Correct)

Results are stored in SSD in Virtual Warehouse in case of Local Disk Cache. (Incorrect)

#### Explanation

As a best practice, Group and Execute similar queries on the same virtual warehouse to maximize local disk cache reuse for performance and cost optimization. The results get stored in the SSD of the Virtual Warehouse. So, if the Virtual Warehouse gets suspended, then results get lost.

Question 90: **Incorrect**

Snowflake prunes micro-partitions based on a predicate with a subquery, even if the subquery result is constant. (TRUE/FALSE)

TRUE (Incorrect)

FALSE (Correct)

#### Explanation

Please note, not all predicate expressions can be used to prune. Snowflake does not prune micro-partitions based on a predicate with a subquery, even if the subquery results in a constant.

Question 91: **Incorrect**

During data loading using COPY INTO <table> command, if the string exceeds the target column length, what options do you have to truncate the string? (Select 2)

ENFORCE\_LENGTH = TRUE

(Incorrect)

ENFORCE\_LENGTH = FALSE

(Correct)

TRUNCATECOLUMNS = FALSE

TRUNCATECOLUMNS = TRUE

(Correct)

**Explanation**

**ENFORCE\_LENGTH:**

- If TRUE, the COPY statement produces an error if a loaded string exceeds the target column length.
- If FALSE, strings are automatically truncated to the target column length.

**TRUNCATECOLUMNS:**

- If TRUE, strings are automatically truncated to the target column length.
- If FALSE, the COPY statement produces an error if a loaded string exceeds the target column length.

Question 93: **Incorrect**

ACCOUNTADMIN role should not be used to create objects in Snowflake?  
(True/False)

FALSE

(Incorrect)

TRUE

(Correct)

**Explanation**

The ACCOUNTADMIN role is intended for performing initial setup tasks in the system and managing account-level objects and tasks on a day-to-day basis. It should not be used to create objects in your account unless you absolutely need these objects to have the highest level of secure access.

Question 94: **Incorrect**

What are the key benefits of The Data Cloud? (Select 3)

|                                                |             |
|------------------------------------------------|-------------|
| <input checked="" type="checkbox"/> Governance | (Correct)   |
| <input checked="" type="checkbox"/> Backup     | (Incorrect) |
| <input type="checkbox"/> Action                | (Correct)   |
| <input checked="" type="checkbox"/> Access     | (Correct)   |
| <input type="checkbox"/> Maintenance           |             |

**Explanation**

The benefits of The Data Cloud are **Access, Governance, and Action**.

**Access** means that organizations can easily discover data and share it internally or with third parties without regard to geographical location.

**Governance** is about setting policies and rules and protecting the data in a way that can unlock new value and collaboration while maintaining the highest levels of security and compliance.

**Action** means you can empower every part of your business with data to build better products, make faster decisions, create new revenue streams and realize the value of your greatest untapped asset, your data.

Question 98: **Incorrect**

What value will be return by the following query? `SELECT * FROM TABLE(FLATTEN(input => parse_json('[]'), outer => true)) f;`

|                                     |             |
|-------------------------------------|-------------|
| <input type="radio"/> 0             |             |
| <input type="radio"/> NULL          | (Correct)   |
| <input checked="" type="radio"/> [] | (Incorrect) |

**Explanation**

The OUTER => TRUE argument with FLATTEN generates exactly one row for zero-row expansions (with NULL in the KEY, INDEX, and VALUE columns).

Question 99: **Incorrect**

Monica is confused about which sampling method she should use with one of the very large tables, considering better performance. Which sampling method would you recommend from BERNOULLI | ROW and SYSTEM | BLOCK?

SYSTEM | BLOCK

(Correct)

BERNOULLI | ROW

(Incorrect)

**Explanation**

SYSTEM | BLOCK sampling is often faster than BERNOULLI | ROW sampling. Also, BERNOULLI | ROW method is good for Smaller Tables, and SYSTEM | BLOCK method is for Larger Tables.

**Snowflake does not start executing any queries on a new virtual warehouse until all of the servers are provisioned.**

- 

**False**

- 

**True**

**(Correct)**

### **Explanation**

Correct. The SQL execution starts only once all the servers are provisioned. If provisioning fails for any reason, Snowflake will attempt to fix the failed servers and SQL will start executing once 50% or more virtual servers are provisioned

**Question 2: Correct**

**You can view the storage used by failsafe under the Account->Billing & Usage section**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

<https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

**Question 3: Correct**

**Which of the following are options when creating a new virtual warehouse.**

**Select all that apply**

- 

**auto\_resume**

**(Correct)**

- **resume\_on\_demand**
  - **suspend\_and\_resume**
  - **auto\_suspend**
- (Correct)**

#### Explanation

Please see the below link for detailed syntax on virtual warehouse creation  
<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

---

Question 4: **Incorrect**

**Which one of the following can not be cloned?**

- **table stage**
  - **internal named stage**
- (Correct)**
- **external named stage**
- (Incorrect)**

#### Explanation

Internal named stages can not be cloned

---

Question 5: **Correct**

**The Snowflake credit usage for multi cluster virtual warehouses is based on**

- 

**the number of clusters running during each hour for the warehouse is running**

**(Correct)**

- 

**the value that is configured as the maximum in the multi cluster virtual warehouse configuration**

### **Explanation**

The billing for multi cluster virtual warehouses is based on a) the size of the multicluster virtual warehouse b) the number of clusters that were active

---

**Question 6: Incorrect**

**Snowflake uses a unique architecture in which data and compute has been decoupled and both can be scaled independent of each other.**

- 

**Yes**

**(Correct)**

- 

**No**

**(Incorrect)**

---

**Question 7: Incorrect**

**Which of the following are options when created a new multicluster virtual warehouse. Select all that apply**

- 

**SCALING\_POLICY**

**(Correct)**

- 

**MIN\_CLUSTER\_COUNT**

**(Correct)**

- 

**MAX\_CLUSTER\_COUNT**

**(Correct)**

- 

**SUSPEND\_AND\_RESUME**

### **Explanation**

Please see the below link for detailed syntax on virtual warehouse creation

<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

---

Question 8: **Incorrect**

**What happens when a table is shared by a snowflake account to another snowflake account? Select all that applies**

- 

**Sharing is managed through Snowflake metadata services layer**

**(Correct)**

- 

**No actual data is copied or transferred between accounts**

**(Correct)**

- 

**the target snowflake account is charged for the shared data storage**

**(Incorrect)**

- 

**data is copied to the target snowflake account**

### **Explanation**

In Snowflake sharing, no data is copied. It is just the metadata that enables the sharing of data. Since no data is copied, the target snowflake account (also called the

consumer) is not charged for any storage <https://docs.snowflake.net/manuals/user-guide/data-sharing-intro.html#how-does-secure-data-sharing-work>

---

Question 9: **Incorrect**

**The "load metadata" (which maintains which files have already been loaded) for a table expires after how many days?**

•

**64**

**(Correct)**

•

**365**

•

**30**

**(Incorrect)**

•

**1**

### Explanation

The load metadata expires after 64 days. Please see below for more details  
<https://docs.snowflake.net/manuals/user-guide/data-load-considerations-load.html#load-metadata>

---

Question 10: **Incorrect**

**Which of the following is true regarding Roles in Snowflake?**

•

**Users in Snowflake can only have a maximum of 10 Roles assigned.**

•

**A Snowflake User can have one or many Roles.**

**(Correct)**

•

**In addition to Roles, privileges may be directly assigned to a named User.**

**(Incorrect)**

- 

**For an active user session, only one Role can be the active Role at a given time.**

**(Correct)**

### **Explanation**

Users may be assigned one or more Roles, however, for a session only one role can be active at a time. Users can switch roles by using the USE ROLE <rolename> command.

Please see also

<https://docs.snowflake.com/en/user-guide/security-access-control-overview.html>

Question 11: **Correct**

**You have a virtual warehouse named reporting\_warehouse which is sized as Medium. You need the size of the virtual warehouse to be Large for future queries. What is the best way to make this change.**

- 

**Alter the existing virtual warehouse and change the size to Large**

**ALTER WAREHOUSE reporting\_warehouse WAREHOUSE\_SIZE=large**

**(Correct)**

- 

**Drop the existing virtual warehouse and create a new one with the required parameters**

- 

**It is not permitted to re-size a virtual warehouse once it has been created**

### **Explanation**

Please see the below link for detailed syntax <https://docs.snowflake.net/manuals/sql-reference/sql/alter-warehouse.html>

Question 12: **Incorrect**

You are a developer with a retail organization running a Snowflake data warehouse. You required to load JSON based data into a table.

What approach should you take?

- 

**Use Snowflake provided function to process JSON data while loading it into the table.**

**(Correct)**

- 

**Write a program using a programming language of your choice (Python, Java etc..) to process the JSON file into a CSV structure.**

**Load the CSV file into the table using the COPY command.**

**(Incorrect)**

- 

**Ask the source system to send you a CSV file rather then a JSON file as Snowflake doesn't support JSON.**

**Explanation**

Snowflake provides functionality to process JSON data, therefore the best course of action is to use the Snowflake built in functionality

Question 13: **Correct**

**For a snowflake session, more then one virtual warehouse can be specified at a time for executing a query.**

- 

**True**

- 

**False**

**(Correct)**

## Explanation

Incorrect, a snowflake session can only have one virtual warehouse specified, though you can change the current virtual warehouse for a session

---

Question 14: **Incorrect**

**Snowflake deploy new release at what frequency?**



**monthly**



**every week**

**(Correct)**



**every 2nd week**

**(Incorrect)**

## Explanation

Snowflake releases & patch release are deployed every week

---

Question 15: **Incorrect**

In the Snowflake Web UI which of the following is true?



**It doesn't matter how many sheets you open in the Snowflake Web UI, they share a single session**



**Each open worksheet in the Snowflake Web UI only a single Role can be selected which applies to all worksheets**



**Each worksheet in the Snowflake Web UI is a different user session**

**(Correct)**

- 

**Each open worksheet in the Snowflake Web UI can have a different Role selected**

**(Correct)**

### **Explanation**

Each worksheet in Snowflake Web UI is an independent session. Therefore, each worksheet can have its combination of Database, Schema, Virtual Warehouse & Role selected

---

Question 16: **Incorrect**

**A virtual warehouse can resized at any time whether they are running or stopped**

- 

**True**

**(Correct)**

- 

**False**

**(Incorrect)**

### **Explanation**

Correct. The resize operation can be done any time, even when it is running queries.

---

Question 17: **Correct**

**When a virtual warehouse is provisioned there is no benefit in stopping the virtual warehouse with in the first 60 seconds because that time period has already been billed.**

- 

**False**

- 

**True**

**(Correct)**

### Explanation

Correct. When a virtual warehouse is provisioned the minimum billing is 1 minute, after which the virtual warehouses is billed on per second basis

---

Question 18: **Correct**

**On the WebUI which button on the top bar should you select to get to tabs that provide information about virtual warehouse usage?**



**History**



**Account**

**(Correct)**



**Databases**



**Warehouses**



**Shares**

### Explanation

The Account button will take you to screen which has multiple sub tabs, one of them providing virtual warehouse compute usage information

---

Question 19: **Incorrect**

**What method does Snowflake use to limit the number of micro-partitions accessed during a query?**

- 

**Indexing**

- 

**Pruning**

**(Correct)**

- 

**table scan**

- 

**In memory operations**

**(Incorrect)**

### **Explanation**

Partition elimination or pruning is used to limit the number of micro partitions.

---

Question 20: **Incorrect**

**To clone a table your current role must have which privilege on the source table?**

- 

**SELECT**

**(Correct)**

- 

**WRITE**

- 

**USAGE**

**(Incorrect)**

### **Explanation**

To clone a table you need a SELECT privileges on the source table and for all other objects that can be cloned you need the USAGE privilege

---

Question 21: **Incorrect**

**The privileges provided by the SYSADMIN & SECURITYADMIN role are automatically contained in the ACCOUNTADMIN role since the ACCOUNTADMIN role sits on the top of the role hierarchy.**

- 

**False**

**(Incorrect)**

- 

**True**

**(Correct)**

### **Explanation**

Correct. The ACCOUNTADMIN role is the all powerful role with all privileges.

---

Question 22: **Correct**

**Which role is most powerful role in a Snowflake system?**

- 

**PUBLIC**

- 

**SYSADMIN**

- 

**SECURITYADMIN**

- 

**ACCOUNTADMIN**

**(Correct)**

### **Explanation**

Account admin is the top level role with most power, it should ideally be shared with least amount of people

---

Question 23: **Correct**

**A consumer of a shared database can add new tables or views in the shared database?**

- 

**False**

**(Correct)**

- 

**True**

### **Explanation**

Shared objects are read-only and can not be modified

---

Question 24: **Incorrect**

**Snowflake can load from data staged in which of the following?**

- 

**Google Cloud Storage**

**(Correct)**

- 

**VMWare Storage**

- 

**Oracle Cloud Storage**

- 

**Internal Stage**

**(Correct)**

- 

**Azure Blob Storage**

**(Correct)**

- 

**AWS S3**

**(Correct)**

### **Explanation**

Snowflake supports loading from Internal Stage, AWS S3, Azure Blob and Google Cloud Storage. Please see below link for more details

<https://docs.snowflake.net/manuals/user-guide/data-load-overview.html#loading-from-data-files-staged-on-other-cloud-platforms>

---

Question 25: **Correct**

**What is the lowest snowflake edition that allows customer dedicated virtual warehouses?**

- 

**Business Critical**

- 

**Standard**

**(Correct)**

- 

**Enterprise**

### **Explanation**

All editions of Snowflake provide customer dedicated virtual warehouses, therefore the lowest edition providing such functionality is Standard

---

Question 26: **Correct**

**On the WebUI which button on the top bar should you select to see the query history?**

- 

**Shares**

- 

**Databases**

- 

**Virtual Warehouses**

- 

**Account**

- 

**History**

**(Correct)**

- 

**Worksheets**

---

Question 27: **Correct**

**As a consumer you can create only one database per share?**

- 

**False**

- 

**True**

**(Correct)**

### **Explanation**

Correct. While as a consumer you may have multiple shares, however for each share you can create only one database.

---

Question 28: **Correct**

**When data is shared between Snowflake accounts, a database is created on the consumer side for sharing purposes. The type of this database is ?**

- Open Access
- 
- Writable**
- 
- Temporary**
- Read-only**  
**(Correct)**
- 

### **Permanent**

#### **Explanation**

The correct answer is read-only. The consumer creates a database from the share and the database is created as read-only

---

Question 29: **Incorrect**

**Multi cluster virtual warehouses are designed to handle ... Select all that apply**

- Large number of concurrent users**  
**(Correct)**
- Large Complex Queries**  
**(Incorrect)**
- Large number of concurrent queries**  
**(Correct)**

- 

**Queuing issues**

**(Correct)**

### **Explanation**

The primary purpose of Multi cluster virtual warehouses is to improve queueing and concurrency issues <https://docs.snowflake.net/manuals/user-guide/warehouses-considerations.html#multi-cluster-warehouses-improve-concurrency>

---

Question 30: **Incorrect**

**What is the correct way to find out virtual warehouse credit usage information in Snowflake?**

- 

**Query WAREHOUSE\_METERING\_HISTORY in the Information Schema**

**(Correct)**

- 

**Query WAREHOUSE\_CREDITS\_USED in the Account Usage Schema**

**(Incorrect)**

- 

**Go to Account -> Billing & Usage in the WebUI**

**(Correct)**

- 

**Query WAREHOUSE\_METERING\_HISTORY in the Account Usage Schema**

**(Correct)**

- 

**Query WAREHOUSE\_CREDITS\_USED in the Information Schema**

### **Explanation**

All 3 are valid ways of finding virtual warehouse credit usage  
<https://docs.snowflake.net/manuals/user-guide/credits.html#viewing-warehouse-credit-usage-for-your-account>

---

Question 31: **Incorrect**

**When loading data through COPY command it is a requirement that your table and the file from which the data is being loaded should have same order of columns.**

- 

**True**

**(Incorrect)**

- 

**False**

**(Correct)**

### **Explanation**

The order & the number of columns in the file and the table can be different

---

Question 32: **Incorrect**

**Both time travel & failsafe require extra storage which has cost associated with it.**

- 

**False**

**(Incorrect)**

- 

**True**

**(Correct)**

### **Explanation**

Correct. Both options use additional storage which will result in additional costs  
<https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

---

Question 33: **Correct**

As a data engineer you are developing jobs to load data into a snowflake table. You have an S3 stage defined which contains a single file containing 100 rows.

You managed to load those 100 rows but while performing further development you notice that your COPY command is executing successfully but is loading zero rows into the target table.

What could be the reason?



**You have already loaded the single file that was in the Snowflake stage. Snowflake tracks if a file has already been loaded and will not load it again.**

**(Correct)**



**The file format that you have defined for your data is incorrect, therefore no data is getting loaded.**



**The file in the Snowflake stage is corrupt and therefore can't be loaded.**

**Explanation**

Since you have already loaded this file, Snowflake will track through metadata which files have been loaded already and will not process them again. If you want to retest your copy command you must place a new file or clear the metadata

Question 34: **Correct**

**A virtual warehouse can be suspended or resumed as required.**



**No**



**Yes**

**(Correct)**

### **Explanation**

A virtual warehouse can be suspended and resumed as per required. This capability allows snowflake to have zero, one or more virtual warehouses executing against the same data.

---

Question 35: **Correct**

**Which of the following are valid parameters when create a new virtual warehouse?**

- **CPU\_COUNT**
- **RAM\_SIZE**
- **AUTO\_RESUME**  
**(Correct)**
- **WAREHOUSE\_SIZE**  
**(Correct)**
- **AUTO\_SUSPEND**  
**(Correct)**

### **Explanation**

Please see the below link for detailed syntax on virtual warehouse creation

<https://docs.snowflake.net/manuals/sql-reference/sql/create-warehouse.html>

---

Question 36: **Incorrect**

**Snowflake database is based on the massively parallel shared nothing architecture used by data bases like Teradata, Greenplum.**

- Yes  
**(Incorrect)**
- No  
**(Correct)**

### Explanation

Snowflake architecture is an hybrid architecture, in which unlike shared nothing architectures it uses a single shared storage and multiple compute engines operating on the same storage.

Question 37: **Incorrect**

**Which of the following will contribute to your account level storage usage?**

- Snowflake External Stages

**(Incorrect)**

- Views

**(Incorrect)**

- User Defined Functions

**(Incorrect)**

- Snowflake Internal Stage

**(Correct)**

-

### Database (and the tables in it)

(Correct)

#### Explanation

Storage used by a database (and its tables) and storage used by Snowflake internal stages contribute towards the over all storage usage

---

Question 38: **Incorrect**

**Multifactor Authentication can be enabled for which of the following? Select all that Apply?**

- 

**SnowPipe**

(Incorrect)

- 

**ODBC**

(Correct)

- 

**Python**

(Correct)

- 

**Snowflake WebUI**

(Correct)

#### Explanation

Please see the below link for details <https://docs.snowflake.net/manuals/user-guide/security-mfa.html>

---

Question 39: **Incorrect**

**What is correct about multi cluster virtual warehouses?**

-

**They will add additional clusters automatically when under high query workload which is beyond the current capacity.**

**(Correct)**

- 

**They will remove the additional clusters as query demand decreases.**

**(Correct)**

- 

**Multi cluster virtual warehouses can not be suspended or resumed.**

- 

**Only one multi cluster virtual warehouse can be created per Snowflake customer.**

### **Explanation**

A multi-cluster virtual warehouse will scale up and scale down depending on the query & user demand.

---

Question 40: **Correct**

**On the WebUI which button on the top bar should you select to get to tabs that provide billing and credit information?**

- 

**Account**

**(Correct)**

- 

**Restricted**

- 

**Account Administrator**

### **Explanation**

The Account button will take you to screen which has multiple sub tabs, one of them providing billing and credit information

---

Question 41: **Correct**

You are the data modeler at a large retail organization who store their transactional data in a Snowflake table called "Transactions". The daily revenue reports execute using the "Transactions" table and calculate the revenue for the current day.

The "Transactions" table had 500GB of data originally but it has now grown to be 5TB. You have noticed that over time the performance of the daily revenue reports has degraded.

What is the most efficient & cost-effective way of optimizing performance?

- 

**Increase the size of the virtual warehouse executing the daily reports. That will ensure that the reports execute much more faster.**

- 

**Cluster the Transactions table on the transaction date column. The daily report will only access one day at a time and therefore will only scan required data.**

**(Correct)**

- 

**Create tables for each year of data e.g. Transactions\_2017, Transactions\_2018 and Transactions\_2019. Insert the relevant data from the Transactions table into these \_year tables. Change your report to point to Transactions\_2019 table.**

**Now that the daily revenue report is accessing only a year worth of data, performance will be better.**

### **Explanation**

Clustering keys is the most efficient option. The query will only scan a single day worth of data using partition pruning. (there is a cost associated with reorganization of data when you add a cluster key)

Question 42: **Correct**

**It is not possible to suspend a specific cluster in a multi-cluster virtual warehouse but rather the whole virtual warehouse is suspended.**

-

**True**

**(Correct)**

- 

**False**

### **Explanation**

True. Individual clusters in a virtual warehouse can not be controlled independently, but rather they appear logically as a virtual warehouse together.

---

Question 43: **Correct**

**Assuming you have multiple worksheets open in Snowflake Web UI with in the same browsing session, you can have different role specified in each of the worksheet?**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Correct, you can specify different roles in different worksheets

---

Question 44: **Incorrect**

**Users created in a reader account can not do which of the following?**

- 

**INSERT**

**(Correct)**

- 

**SELECT**

- 

**UPDATE**

**(Correct)**

- 

**LOAD DATA**

**(Correct)**

### **Explanation**

Users in a reader account can query shared data but can not perform any DML.

---

Question 45: **Correct**

**Which minimum Snowflake license allows Multicloud data virtual warehouse capability?**

- 

**Standard**

- 

**Business Critical**

- 

**Enterprise**

**(Correct)**

### **Explanation**

The multi cluster virtual warehouse capability is available starting from the Enterprise Edition. The capability is not available in Standard or Premier

---

Question 46: **Correct**

You are the solution architect for a large retail company running a Snowflake data warehouse. Your Snowflake implementation has just gone live with a single virtual warehouse and has been opened up to all the users across the organization.

One of your heaviest users are the Finance department, who have a large number of users and execute a large number of queries. The finance department has complained that the queries take a long time to execute for them.

What is the best immediate action should you take to improve their experience?

-

**Introduce a dedicated virtual warehouse instance for the finance department and size is according to their needs.**

**Set the new virtual warehouse to auto suspend and auto resume.**

**(Correct)**

- 

**Increase the size of your single virtual warehouse to the maximum size available so that queries for all users (not just finance) execute faster.**

- 

**Limit the number of queries that the finance department can execute.**

### **Explanation**

A dedicated virtual warehouse is the best short term option to optimize performance

Question 47: **Correct**

**What are the virtual warehouse sizing approach recommended by Snowflake?**

- 

**Execute your queries on the largest available virtual warehouse and keep scaling down the size of the virtual warehouse, until you start seeing degradation in query response**

- 

**Experiment with different types of query and virtual warehouses to determine the virtual warehouse size that best meets your requirements**

**(Correct)**

- 

**Execute your queries on the smallest available virtual warehouse and keep scaling up the size of the virtual warehouse, until the query performance no longer shows improvements**

### **Explanation**

Snowflake recommended approach is to experiment to arrive at the best size  
<https://docs.snowflake.net/manuals/user-guide/warehouses-considerations.html>

Question 48: **Correct**

Which minimum Snowflake license provides up to 90 days of time travel?

- 

**Business Critical**

- 

**Standard**

- 

**Enterprise**

**(Correct)**

#### Explanation

Enterprise is the lowest edition that provides 90 days of time travel

---

Question 49: **Correct**

For a multi cluster virtual warehouse the administrator set the minimum cluster count and the maximum cluster count to different values. This can be best described as ?

- 

**Scaling up**

- 

**Auto Scaling**

**(Correct)**

#### Explanation

Setting the minimum and maximum to different values will ensure that the multi cluster virtual warehouse scales up in periods of high demand and scale back down in periods of less demand

---

Question 50: **Incorrect**

Under the Download section in Snowflake Web UI, what are the connectors & programs that can be downloaded? Select All that apply

-

**SnowSQL**

**(Correct)**

- 

**ODBC driver**

**(Correct)**

- 

**JDBC driver**

**(Correct)**

- 

**Snowpipe Connector**

- 

**Load Connector**

### **Explanation**

Please see download page under Help

---

Question 51: **Incorrect**

**Snowflake uses a staged release process for new releases in which different account types are upgraded gradually.**

- 

**True**

**(Correct)**

- 

**False**

**(Incorrect)**

### **Explanation**

Correct. The staged release process through which Snowflake releases updates is first applied to early access account, then standard accounts and finally Enterprise & upwards

---

Question 52: **Correct**

Which of the following Snowflake Editions allow only a maximum of 1 day of time travel?

- 

**Standard**

**(Correct)**

- 

**Business Critical**

- 

**Enterprise**

**Explanation**

Standard edition allows only 1 day of time travel.

<https://www.snowflake.com/pricing/pricing-guide/>

Question 53: **Incorrect**

**What are some general indicators that a clustering key is required on a table?  
Select all that apply.**

- 

**The size of the table is in multi terabytes**

**(Correct)**

- 

**The query performance has slowed down over time**

**(Correct)**

- 

**The table has large number of columns**

**(Incorrect)**

---

Question 54: **Correct**

You can increase or decrease the size of a virtual warehouse as your user base & query demand increase or decrease?

- 

**False**

- 

**True**

**(Correct)**

---

Question 55: **Incorrect**

What are the resources that a virtual warehouse provides? Select all that apply.

- 

**Temporary Storage**

**(Correct)**

- 

**CPU**

**(Correct)**

- 

**Memory**

**(Correct)**

- 

**User Management**

### **Explanation**

Virtual warehouse provide CPU, Memory & temporary storage to process queries

---

Question 56: **Correct**

**When a file has been loaded into a table, Snowflake marks that file as loaded in metadata so that the file doesn't get processed again.**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Snowflake maintains file load metadata, through which it tracks which files have already been processed

---

Question 57: **Incorrect**

**Virtual Warehouse credits are charged on what basis?**

- 

**per minute of use**

- 

**per milli second of use**

- 

**per second of use**

**(Correct)**

- 

**per hour of use**

**(Incorrect)**

### **Explanation**

Snowflake credits are billed on per second basis of virtual warehouse use. Its important to note thought that when a virtual warehouse is started it will consume

minimum 1 minute of snowflake credits and will change to per second billing after that

---

Question 58: **Correct**

**Files that have already been processed into the source table can be loaded again into a cloned table?**

- 

**True**

**(Correct)**

- 

**False**

#### **Explanation**

Correct. The cloning copy structure, data and certain other attributes but doesn't copy the load metadata, therefore files that have already been processed for the source table can be re-processed into a cloned table.

---

Question 59: **Correct**

**If retrieval of data from fail safe is required, you must**

- 

**use SQL to retrieve the data yourself**

- 

**contact Snowflake support team to facilitate the retrieval**

**(Correct)**

#### **Explanation**

Failsafe data can not be retrieved by Snowflake customer and support must be engaged to retrieve this data.

Question 60: **Correct**

**A database is created from a share by the consuming account, the access to this database is configurable using the role-based access control that Snowflake provides?**

-

**True**

**(Correct)**

- 

**False**

#### **Explanation**

Correct. The same role based access is applicable to a shared database

---

**Question 1:**

**Skipped**

**When a virtual warehouse is resumed from suspended state what is the minimum amount of Snowflake credits that it will immediately consume?**

- 

**1 second**

- 

**30 minutes**

- 

**1 minute**

**(Correct)**

- 

**30 seconds**

#### **Explanation**

When a virtual warehouse is started it will consume minimum 1 minute of snowflake credits and will change to per second billing after that

---

**Question 2:**

**Skipped**

**Which of the following are characteristics of Snowflake releases?**

- no down time**  
**(Correct)**
- transparent to user**  
**(Correct)**
- allow user to choose whether to upgrade or not**
- require restart**

### Explanation

Snowflake releases are transparent to the user and don't cause any downtime.

---

Question 3:

**Skipped**

Which of the following is required in order for a role to be able to successfully execute a Task?

- The role should have the global EXECUTE TASK privilege on the Account**  
**(Correct)**
- The role should have the permissions required to run the SQL statement executed in the Task**  
**(Correct)**
- The role should have SYSADMIN privileges**

- 

**The role should have ACCOUNTADMIN privileges**

#### **Explanation**

To execute a task, the role being used must have the global EXECUTE TASK privilege on the Account. Additionally, all the privileges that the SQL statement in the task requires must be available to the role.

<https://docs.snowflake.com/en/user-guide/tasks-intro.html>

---

Question 4:

**Skipped**

**Which type of short lived Snowflake tables will continue to exist even if the session is closed?**

- 

**Permanent**

- 

**Temporary**

- 

**Clustered**

- 

**Transient**

**(Correct)**

#### **Explanation**

Transient tables can be used as short lived tables for ETL work tables and are not dropped when the session is closed

---

Question 5:

**Skipped**

**You are planning to utilize the multi-cluster virtual warehouse to provide auto-scaling and performance for your users. You want to make sure that the auto**

**scaling doesn't start clusters immediately and is cost effective. What is the correct value for Scaling Policy to achieve this?**

- 
- **SCALING\_POLICY = PERFORMANCE**
- 
- **SCALING\_POLICY = STANDARD**
- **SCALING\_POLICY = ECONOMY**  
**(Correct)**

#### **Explanation**

ECONOMY scaling policy will not auto scale immediately but will be conservative in scaling See link for further details <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#label-mcw-scaling-policies>

---

Question 6:

**Skipped**

**Snowflake storage capacity can be pre-purchased for a lower price?**

- 
- **True**  
**(Correct)**
- 
- **False**

#### **Explanation**

True, you can pre-purchase storage capacity which will generally be of lower price compared to on demand pricing which is as per your usage

---

Question 7:

**Skipped**

**The COPY command in a Snowpipe definition supports the same transformation as provided by the usual COPY command.**

- True
- (Correct)
- False

### Explanation

Correct. The same transformations are available in a Snowpipe definition.

---

Question 8:

**When processing semi structured data into structured (i.e. a table) what is the correct way to cast a column into a data type? Assume the target column name is CustomerName and the data type is String**

- `SELECT col1:CustomerName::String`
- (Correct)
- `SELECT CAST(col1:CustomerName AS String)`
- `SELECT col1 AS CustomerName WITH DATATYPE AS String`

### Explanation

Please see below link for more details <https://docs.snowflake.net/manuals/user-guide/semistructured-considerations.html>

---

Question 9:

**Under what circumstances results for a query will be fulfilled by the query result cache? Select all that apply**

-

**The query results cache was generated or last used less than 24 hours ago**

**(Correct)**

- 

**The underlying data for the tables in the query has not changed**

**(Correct)**

- 

**The query is being executed from the same virtual warehouse as the previously executed query**

- 

**The query being run is syntactically identical to a previously executed query**

**(Correct)**

- 

**The query being run returns result set of less than 100 thousand rows**

### **Explanation**

The underlying data must not change and the query should be syntactically identical in order for the query result cache to be used. The cache must have been generated (or last used) less than 24 hours ago. <https://docs.snowflake.net/manuals/user-guide/querying-persisted-results.html#retrieval-optimization>

---

**Question 10:**

**Which of the following are valid Casting function available in Snowflake?**

- 

**TO\_TIMESTAMP**

**(Correct)**

- 

**TO\_NUMERIC**

**(Correct)**

- 

**TO\_VARIANT**

**(Correct)**

- 

**TO\_NUMBER**

**(Correct)**

- 

**TO\_BLOB**

Question 11:

**Which of the following properties can be set for a multi cluster virtual warehouse? Select all that apply?**

- 

**AUTO\_SCALE**

- 

**AUTO\_RESUME**

**(Correct)**

- 

**MIN\_CLUSTER\_COUNT**

**(Correct)**

- 

**MAX\_CLUSTER\_COUNT**

**(Correct)**

-

**AUTO\_SUSPEND**

**(Correct)**

### **Explanation**

Auto scale is not a valid property. Other than auto scale all other given properties are valid

---

Question 12:

**Snowflake allows which ways to scale a virtual warehouse? Select all that apply**

- 

**Exponential Scaling**

- 

**Linear Scaling**

- 

**Scale Up (increasing the size of a virtual warehouse)**

**(Correct)**

- 

**Scale Out (adding clusters to a multi cluster virtual warehouse)**

**(Correct)**

### **Explanation**

The two scaling methods are scale up and scale out

---

Question 13:

**Snowflake SQL UDFs can return either Scalar or Tabular results?**

- 

**False**

- 

**True**

**(Correct)**

### Explanation

Either type of UDFs (JavaScript or SQL) can return Scalar or Tabular results

<https://docs.snowflake.net/manuals/sql-reference/user-defined-functions.html>

---

Question 14:

Multifactor Authentication can be enabled for which of the following?

Select all that Apply?

- 

**SnowPipe**

- 

**Snowflake WebUI**

**(Correct)**

- 

**JDBC**

**(Correct)**

- 

**SnowSQL**

**(Correct)**

### Explanation

Please see the below link for details <https://docs.snowflake.net/manuals/user-guide/security-mfa.html>

---

Question 15:

**You increased the size of a virtual warehouse to improve performance of a query? What is this an example of?**

-

## Scale Out

- 
- Scale Up**
- (Correct)**

### Explanation

Increasing the size of a virtual warehouse will enable faster processing for queries but will not necessarily accommodate more queries and more users, therefore this is an example of scaling up

---

Question 16:

**Snowflake patch releases are applied to all accounts at the same time.**

- - True**
  - (Correct)**
- - False**

### Explanation

Correct. The patch releases are applied to all accounts at the same time. The staged released process is only used for new software releases.

---

Question 17:

**What is one of the way to improve performance in Snowflake?**

- 
- Query Hints**
- 
- Join Indices**
- 
- Secondary Indices**

- 

**Clustering Keys**

**(Correct)**

### **Explanation**

Clustering keys is the performance optimization option provided by Snowflake.

---

Question 18:

**You are the Snowflake Administrator for a large telecom company. You had a Medium sized virtual warehouse on which a complex query was executing, you change the virtual warehouse size to be Small. Will the query be impacted and perform slower?**

- 

**No. The currently executing query will keep on executing on the original cluster. Only new queries will execute on the Small sized cluster.**

**(Correct)**

- 

**Yes. The query will now start using the Small virtual warehouse, therefore will be slower**

### **Explanation**

The query which is already executing will continue using the resources it has acquired, only new queries will be impacted by the change in size

---

Question 19:

**Under the Download section in Snowflake Web UI, what are the connectors & programs that can be downloaded? Select All that apply**

- 

**SnowSQL**

**(Correct)**

- 

**JDBC driver**

**(Correct)**

- **ODBC driver**
- **Snowpipe Connector**
- **Load Connector**

#### **Explanation**

Please see download page under Help

---

Question 20:

**What are some of the ways to improve performance in Snowflake? Select all that apply**

- **Query Result Caching**
- **Dedicated Virtual Warehouse**
- **Clustering Keys**

**(Correct)**

- 

**Multi Clustering Virtual Warehouse**

**(Correct)**

- 

**Secondary Indices**

---

Question 21:

**When a virtual warehouse is resized to a SMALLER size, the required servers are only removed when ?**

- 

**there are no SQL statements running on the servers**

**(Correct)**

- 

**immediately**

- 

**after 1 minute**

### **Explanation**

The servers are removed only when they are not being used by a statement

<https://docs.snowflake.net/manuals/user-guide/warehouses-tasks.html>

---

Question 22:

**On the Snowflake Web UI, which of the following are buttons in the top bar?**

**Select all that apply**

- 

**Roles**

-

## Databases

(Correct)

- 

## Security

- 

## Shares

(Correct)

---

Question 23:

A virtual warehouse can be set to auto-suspend or auto-resume but not both.

- 

False

(Correct)

- 

True

## Explanation

You can set a virtual warehouse to auto suspend and auto resume as per your needs.

---

Question 24:

Which type of Snowflake table do not have a failsafe?

- 

Permanent

- 

Temporary

(Correct)

-

**Transient**

**(Correct)**

- 

**Clustered**

### **Explanation**

Transient and temporary tables don't have any failsafe, this is done to reduce storage costs for temporary and transient data. <https://docs.snowflake.net/manuals/user-guide/data-cdp-storage-costs.html>

---

Question 25:

**In a multi cluster virtual warehouse, if you set different values for minimum cluster count and the maximum cluster count**

- 

**Snowflake starts minimum required clusters in the multi-cluster virtual warehouse**

**(Correct)**

- 

**Snowflake starts all clusters in the multi-cluster virtual warehouse**

- 

**The mode is called auto-scale**

**(Correct)**

- 

**The mode is called Maximized**

### **Explanation**

If minimum and maximum value are set to different values the multi cluster virtual warehouse will be in auto scale mode <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

---

Question 26:

**A Snowflake share ....**

- is a container that contains objects to be shared**  
**(Correct)**
- is only available to enterprise customers**
- has information about the consumer(s)**  
**(Correct)**

### Explanation

A share contains USAGE privileges on the database & the schema to be shared, privileges on the tables, secure views which will be shared and the consumer account(s) to which the share will be available. It is available to all Snowflake editions

Question 27:

**On the WebUI which button on the top bar should you select to get to tabs that provide information about storage usage?**

- Shares**
- Virtual Warehouses**
- 
- Worksheets**
- Account**  
**(Correct)**
-

## Databases

- 

## History

### Explanation

The Account button will take you to screen which has multiple sub tabs, one of them providing account level storage usage information

---

Question 28:

**Micro partitions can overlap in their range of values?**

- 

**False**

- 

**True**

**(Correct)**

### Explanation

Micro partitions can overlap with each other in their range of values. This can help reduce skew in partition sizes <https://docs.snowflake.net/manuals/user-guide/tables-clustering-micropartitions.html>

---

Question 29:

**What are the types of transformations which are available when loading data into a table using the COPY command. Select all that apply.**

- 

**Pivot Rows into Columns**

- 

**Re-order Columns**

**(Correct)**

-

**Cast Columns**

**(Correct)**

- 

**Concatenate Columns**

**(Correct)**

- 

**Omit Columns**

**(Correct)**

Question 30:

**What happens when different value is specified for the minimum & the maximum cluster? Select all that apply**

- 

**Snowflake starts all clusters in the multi-cluster virtual warehouse**

- 

**Snowflake starts minimum required clusters in the multi-cluster virtual warehouse that can accommodate the number of users & number of queries**

**(Correct)**

- 

**If the number of concurrent users and queries increase on the virtual warehouse, additional clusters are started, up to the defined maximum cluster value**

**(Correct)**

- 

**If the demand is decreased the number of clusters are gradually shutdown, up to the defined minimum cluster value**

**(Correct)**

- 

**The mode is called auto-scale**

**(Correct)**

### **Explanation**

For an auto scale multi cluster virtual warehouse 1) only minimum required clusters are started 2&3) The running clusters may be scaled up and down within the maximum and minimum cluster defined <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

---

Question 31:

Using Snowpipe customers can load data continuously from which of the following?

- 

**External Tables**

- 

**Internal Stages**

**(Correct)**

- 

**External Stages**

**(Correct)**

### **Explanation**

Snowpipes can be used to load data from external as well as internal stages.

---

Question 32:

**When a database or a schema is cloned, which of the following statements are true for the stages in that database?**

- 

**External named stages are NOT cloned**

- 

**External named stages are cloned**

**(Correct)**

- 

**Table internal stages are cloned**

**(Correct)**

- 

**named internal stages are NOT cloned**

**(Correct)**

- 

**named internal stages are cloned**

### **Explanation**

See rules for stage cloning at the below link

<https://docs.snowflake.net/manuals/user-guide/object-clone.html#rules-for-cloning-stages>

**Question 33:**

REST API calls can be used to trigger a Snowpipe regardless of whether it is defined over an Internal Stage or an External Stage

- 

**False**

- 

**True**

**(Correct)**

### **Explanation**

SnowPipe can be used to load data from external as well as internal stages. The REST API which is used to trigger a SnowPipe applies to SnowPipe created over external stages as well as internal stages

---

Question 34:

If you used a specific virtual warehouse to load data into a table, you must use the same virtual warehouse to query that data.

- 

**False**

**(Correct)**

- 

**True**

### Explanation

Since snowflake uses shared data architecture, you can use any virtual warehouse to load data and the same or any other virtual warehouse to query that data

---

Question 35:

Which of the following objects can be cloned? Select all that apply.

- 

**Schema**

**(Correct)**

- 

**Virtual Warehouse**

- 

**User**

- 

**Share**

- 

**Database**

**(Correct)**

- 

**Table**

**(Correct)**

### **Explanation**

Of these virtual warehouses and shares can't be cloned

<https://docs.snowflake.net/manuals/sql-reference/sql/create-clone.html>

---

Question 36:

**To clone a database your current role must have which privilege on the source database?**

- 

**SELECT**

- 

**WRITE**

- 

**USAGE**

**(Correct)**

### **Explanation**

To clone a table you need a SELECT privileges on the source table and for all other objects that can be cloned you need the USAGE privilege

---

Question 37:

**Which of the following objects can be shared? Select all that apply.**

- 

**Views**

- 

**Secure UDFs**

**(Correct)**

- **Schema**  
**(Correct)**

- **Tables**  
**(Correct)**

- **Secure Views**  
**(Correct)**
- **Database**  
**(Correct)**

### Explanation

Database, Schemas, Tables, Secure views & Secure UDFs can be shared.

---

Question 38:

Which of the following semi-structured file formats are supported by Snowflake ?  
Select all that apply

- **AVRO**  
**(Correct)**

- **PARQUET**  
**(Correct)**

-

**JSON**

**(Correct)**

- 

**XML**

**(Correct)**

- 

**YAML**

- 

**ORC**

**(Correct)**

- 

**HTML**

#### **Explanation**

Please see below link for details on the supported formats

<https://docs.snowflake.net/manuals/user-guide/semistructured-intro.html#supported-file-formats>

---

Question 39:

**Which of the following statement is true when sharing data with an organization which doesn't have a Snowflake account?**

- 

**A consumer account must be created for a non-snowflake customer to allow sharing**

- 

**A reader account must be created for a non-snowflake customer to allow sharing**

**(Correct)**

-

**No extra configuration is required to share data with a non-snowflake user/organization**

- 

**Snowflake allows sharing of data to/from Snowflake Accounts only (so the consumer must have a Snowflake account or a reader account)**

**(Correct)**

### **Explanation**

Sharing with a non-snowflake user requires creation of a reader account as Snowflake allows sharing only between Snowflake accounts

---

Question 40:

**If you have multiple virtual warehouses in your Snowflake system, they will each be accessing the same shared data.**

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

True. Although Virtual warehouses are independent of each other and don't share the CPU or memory or cache but they do operate on the shared data.

---

Question 41:

**You are planning to utilize the multi-cluster virtual warehouse to provide auto-scaling and performance for your users. You want to make sure that the user queries don't wait in queue and if auto-scaling is required the cluster should auto scale immediately. What is the correct value for Scaling Policy to achieve this?**

- 

**SCALING\_POLICY = STANDARD**

**(Correct)**

- 

**SCALING\_POLICY = PERFORMANCE**

- 

**SCALING\_POLICY = ECONOMY**

#### **Explanation**

STANDARD scaling policy will auto scale as soon as a query is queued or if it is detected that the number of clusters can't handle the current queries  
<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#label-mcw-scaling-policies>

---

Question 42:

**A new custom role will automatically be assigned to all existing users.**

- 

**True**

- 

**False**

**(Correct)**

#### **Explanation**

A new role is not assigned to any user. You have to explicitly assign the role.

---

Question 43:

**Which minimum Snowflake license allows data sharing capability?**

- 

**Standard**

**(Correct)**

- 

**Enterprise**

- 

**Business Critical**

### **Explanation**

All editions of Snowflake provide customer dedicated virtual warehouses, therefore the lowest edition providing such functionality is Standard

---

Question 44:

**Please select the 3 key services which are part of the Snowflake Architecture**

- 

**Query Processing**

**(Correct)**

- 

**Cloud Services**

**(Correct)**

- 

**Database Storage**

**(Correct)**

- 

**Virtual Machines**

Question 45:

**Cloning a database will clone which of the following**

- 

**The database itself**

**(Correct)**

- 

**All tables ONLY in the public schema in the database**

- 

**ONLY the public schema in the database**

- 

All tables with in every schema in that database

(Correct)

- 

All the schemas in the database

(Correct)

### Explanation

Cloning a database clones all the schemas in the database, all the tables in the schemas and all other cloneable objects in that schema

---

Question 46:

Which one of the following represents correctly the storage hierarchy in Snowflake?

- 

Account->Database->Schema->Table

(Correct)

- 

Database->Account->Schema->Table

- 

Account->Table->Database->Schema

- 

Schema->Account->Database->Table

### Explanation

In Snowflake the highest level is an Account. Within an account, you can have databases. Each database can contain one or more schemas. Objects such as tables exist within schemas

---

Question 47:

If a multi cluster virtual warehouse is resized the new size applies to ?

-

All the running clusters & any clusters that are started after the multi cluster virtual warehouse is resized

(Correct)

- 

You can not resize a multi cluster virtual warehouse

### Explanation

The resize will apply to all existing clusters which are in running state and any future clusters which may be added due to autoscaling

<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#multi-cluster-size-and-credit-usage>

Question 48:

What is the maximum time travel period allowed for transient & temporary tables?

- 

7

- 

90

- 

0

- 

1

(Correct)

### Explanation

Transient and temporary tables can have a maximum of 1 days of time travel regardless of which snowflake edition you are using

<https://docs.snowflake.net/manuals/user-guide/data-cdp-storage-costs.html#temporary-and-transient-tables>

Question 49:

When a data provider stops sharing an object, that object immediately becomes unavailable to all its consumers.

- 

**True**

**(Correct)**

- 

**False**

### **Explanation**

Data Sharing in Snowflake is a metadata operation. If a producer stops sharing any object, it will immediately become unavailable for any consumer.

---

Question 50:

**When a table is cloned, a snapshot of the data is taken at the time of the cloning and represents the state of the source data at the time of cloning.**

- 

**False**

- 

**True**

**(Correct)**

### **Explanation**

Correct. A snapshot is taken which represents the state of the source data. The clone is based on that snapshot. <https://docs.snowflake.net/manuals/sql-reference/sql/create-clone.html#general-usage-notes>

---

Question 51:

You are required to run a series of SQL statements on a scheduled interval. The SQL statements must execute in a defined order. What is the best way to achieve this in Snowflake?

- 

**1.Create multiple Tasks each containing one SQL statement.**

**2.Connect the Tasks in a parent-child relationship to define the execution order.**

**3.Schedule the root Task.**

**(Correct)**

- 

**1.Create multiple scripts each containing one SQL statement.  
2.Create a master script that calls all the scripts in the execution order.  
3.Schedule the master script.**

- 

**1.Create a script containing all the statements.  
2.Schedule the script to run on a set interval.**

### **Explanation**

Tasks are the Snowflake way of executing a single SQL statement on a schedule. Multiple Tasks can be connected to create a tree of tasks that starts executing from the root task and executes all child Tasks.

See more at

<https://docs.snowflake.com/en/user-guide/tasks-intro.html>

---

Question 52:

**A user can enable multi factor authentication (MFA) using the Snowflake Web interface?**

- 

**True**

**(Correct)**

- 

**False**

---

Question 53:

**Failsafe is provided as an alternate means for customers to access historical data once the time travel retention period has ended.**

- 

**False**

**(Correct)**

- 

**True**

### **Explanation**

Incorrect. Failsafe can be used only in failure scenarios for data protection and can be accessed only by the Snowflake support <https://docs.snowflake.net/manuals/user-guide/data-failsafe.html>

---

Question 54:

**You are designing a snowflake based solution for a large retail chain. What should you consider for ensuring performance once the Snowflake based data warehouse goes live? Select all that apply**

- 

**Consider using multi clustering to auto spawn & auto shutdown virtual warehouse instances to cater for increase & decrease in user demand.**

**(Correct)**

- 

**Consider using dedicated virtual warehouses for different departments.**

**(Correct)**

- 

**Minimize the amount of data stored in Snowflake tables. Only store last 3 months worth of data.**

- 

**Set the size of the virtual warehouses in your solution to maximum available size (4XL) to ensure best performance.**

Question 55:

**What is the Snowflake recommended approach to identifying the minimum and maximum values for a multi-cluster virtual warehouse? Select all that apply**

-

**Start with the multi-cluster virtual warehouse in auto scale mode**

**(Correct)**

- 

**Start with a minimum of 1 & a maximum cluster value of 2 or 3**

**(Correct)**

- 

**Track the virtual warehouse usage over time and adjust the minimum and maximum values that provide the best configuration for your workload**

**(Correct)**

- 

**Start with the minimum and maximum value set to the same number.**

- 

**Start with a minimum of 5 & a maximum cluster value of 10**

### **Explanation**

Snowflake recommended approach is to start small, set to auto scale and monitor and adjust according to your workload <https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#maximized-vs-auto-scale>

---

Question 56:

**Which of the following will allow a user in consumer account to create a database from the share. Select two.**

- 

**SECURITY ADMIN role**

- 

**ACCOUNTADMIN role**

**(Correct)**

- 
- SYSADMIN role**
- 
- IMPORT SHARE privileges**
- (Correct)**

#### **Explanation**

A user in consumer account can create a database from the share if they have the ACCOUNTADMIN role OR the IMPORT SHARE privileges

---

Question 57:

**To create a SHARE what is the minimum required role?**

- 
- ACCOUNTADMIN**
- (Correct)**
- 
- SYSADMIN**
- 
- SECURITYADMIN**

#### **Explanation**

ACCOUNTADMIN is the role required to create & manage a share.

---

Question 58:

In the Snowflake staged release process for new releases, which account types are applied updates the **LAST**?

- 
- Designated accounts opting for early access**
- 
- Standard Accounts**

- 

**Trial Accounts**

- 

**Enterprise & Higher Accounts**

**(Correct)**

### **Explanation**

The staged release process through which Snowflake releases updates is first applied to early access account who have opted for early access, then standard accounts and finally Enterprise & upwards

---

Question 59:

**All queries will perform better if the virtual warehouse size is increased?**

- 

**False**

**(Correct)**

- 

**True**

### **Explanation**

Incorrect. Simple and short queries will generally not benefit from the additional resources provided by larger virtual warehouses

Question 60:

**How are the columns stored in a Snowflake micro partition? (Also referred to as Columnar storage)**

- 

**Combined - columns for a given row are stored together**

- 

**Independently - each column stored on its own**

**(Correct)**

### **Explanation**

Columns are stored independently so that they can be scanned more efficiently  
<https://docs.snowflake.net/manuals/user-guide/tables-clustering-micropartitions.html>

---

Question 61:

**Which of the following statement is True regarding how Snowflake stores its data.**

- 

**Snowflake uses its own proprietary columnar format to store table data.**

**(Correct)**

- 

**Snowflake uses parquet file format to store the table data.**

- 

**Snowflake stores table data as simple comma separated files in cloud based storage.**

---

Question 62:

**Snowpipe is designed to load?.**

- 

**large volumes of data on a batch schedule**

- 

**small volumes of frequent data**

**(Correct)**

### **Explanation**

Snowpipe is used for continuously loading small volumes of data frequently

---

Question 63:

**Each snowflake account is limited to a maximum of 10 shares.**

- 

**True**

- 

**False**

**(Correct)**

### **Explanation**

Incorrect. There is no hard limit on the number of shares that can be created per account.

---

Question 64:

**Which minimum Snowflake license provides the failsafe feature?**

- 

**Standard**

**(Correct)**

- 

**Enterprise**

- 

**Business Critical**

### **Explanation**

All editions of snowflake provide the failsafe feature

---

Question 65:

**Which function is used to find the query id of the last executed query?**

- 

**SELECT max(query\_id)**

- 

**SELECT min(query\_id)**

- 

**SELECT last\_query\_id()**

**(Correct)**

### **Explanation**

[https://docs.snowflake.net/manuals/sql-reference/functions/last\\_query\\_id.html](https://docs.snowflake.net/manuals/sql-reference/functions/last_query_id.html)

---

Question 66:

**Assume a virtual warehouse of size X-Large(128 servers) running for an hour.  
What is the amount of Snowflake credits that will be billed?**

- 

**64**

- 

**60**

- 

**128**

**(Correct)**

### **Explanation**

One snowflake credit is charged per hour per server, so a cluster consisting of 128 servers will be billed as 128 Snowflake credits.

---

Question 67:

To create a new Task a user should be using a role which..

- 

**has the ACCOUNTADMIN privileges**

- 

**has the CREATE TASK privilege on the schema where the Task is being created**

**(Correct)**

- 

**has the SYSADMIN privileges**

**Explanation**

To create a task, the role being used by the creating user must have the CREATE TASK privilege on the schema in which the Task is being created.

<https://docs.snowflake.com/en/user-guide/tasks-intro.html>

---

Question 68:

**Which of the following are valid scaling type for a multi cluster virtual warehouse? Select all that apply?**

- 

**SPEED**

- 

**STANDARD**

**(Correct)**

- 

**PERFORMANCE**

- 

**ECONOMY**

**(Correct)**

**Explanation**

The two scaling types are STANDARD & ECONOMY

<https://docs.snowflake.net/manuals/user-guide/warehouses-multicloud.html#setting-the-scaling-policy-for-a-multi-cluster-warehouse>

---

Question 69:

**Which of the following actions can not be performed by the consumer of a shared database?**

-

## use shared data in complex queries

- **insert data in a shared table**  
**(Correct)**
- **view the list of tables in a shared database**
- **create a new table in the shared database**  
**(Correct)**

### Explanation

For a consumer a shared database is read-only so they can not create new objects or modify/append data

---

Question 70:

**Snowflake automatically determines the most efficient algorithm to compress columns in a micro partitions**

- **True**  
**(Correct)**
- **False**

### Explanation

Correct. The columns are compressed in a micro-partition and the algorithm is determined automatically by Snowflake based on column characteristics  
<https://docs.snowflake.net/manuals/user-guide/tables-clustering-micropartitions.html>

---

Question 71:

**Using the COPY command you can load data from which of the following sources? Select ALL that apply.**

- 
- **AWS S3**  
**(Correct)**
- 
- **IBM Cloud**
- 
- **Oracle Cloud**
- 
- **Internal Stages**  
**(Correct)**
- 
- **Azure Blob Storage**  
**(Correct)**
- 
- **Google Cloud Storage**  
**(Correct)**

### Explanation

Snowflake supports loading from Internal Stage, AWS S3, Azure Blob and Google Cloud Storage. Please see below link for more details

<https://docs.snowflake.net/manuals/user-guide/data-load-overview.html#loading-from-data-files-staged-on-other-cloud-platforms>

---

Question 72:

**The default table type for new tables in Snowflake is..**

- 

**Transient**

- **Temporary**
- **Permanent**  
**(Correct)**

### **Explanation**

The default type for a new table is Permanent. To create other types of tables you will need to explicitly specify <https://docs.snowflake.net/manuals/user-guide/data-cdp-storage-costs.html#considerations-for-using-temporary-and-transient-tables-to-manage-storage-costs>

---

Question 1: **Incorrect**

**Which of these roles is dedicated to user and role management only?**

- 
- **ORGADMIN**
- 
- **ACCOUNTADMIN**
- 
- **SYSADMIN**
- **(Incorrect)**
- 
- **USERADMIN**
- **(Correct)**
- 
- **SECURITYADMIN**

### **Explanation**

USERADMIN role is dedicated to user and role management only. More specifically, this role:

- Is granted the CREATE USER and CREATE ROLE security privileges.
- Can create users and roles in the account.
- This role can also manage users and roles that it owns.

Only the role with the OWNERSHIP privilege on an object (i.e. user or role), or a higher role, can modify the object properties.

Question 2: **Incorrect**

**Snowflake blocks certain IPs by default to ensure that customer is getting the highest level of Network security. (TRUE / FALSE)**

- 
- **TRUE**

**(Incorrect)**

- 

**FALSE**

**(Correct)**

#### **Explanation**

**By default, Snowflake allows users to connect to the service from any computer or device IP address.** A security administrator (or higher) can create a network policy to allow or deny access to a single IP address or a list of addresses.

Question 3: **Incorrect**

**Micro-partitioning is the on-demand feature of Snowflake. It is required to be enabled explicitly by ACCOUNTADMIN. (True / False)**

- 

**TRUE**

**(Incorrect)**

- 

**FALSE**

**(Correct)**

#### **Explanation**

Micro-partitioning is automatically performed on all Snowflake tables. Tables are transparently partitioned using the Ordering of the data as inserted or loaded.

Question 4: **Correct**

**The Snowflake Information Schema includes table functions you can query to retrieve information about your directory tables. Which table function can be used to query the history of data files registered in the metadata of specified objects and the credits billed for these operations?**

- 

**STAGE\_STORAGE\_USAGE\_HISTORY**

- 

**AUTO\_REFRESH\_REGISTRATION\_HISTORY**

**(Correct)**

- 

**STAGE\_DIRECTORY\_FILE\_REGISTRATION\_HISTORY**

- 

**DATABASE\_REFRESH\_HISTORY**

### **Explanation**

AUTO\_REFRESH\_REGISTRATION\_HISTORY table function can be used to query the history of data files registered in the metadata of specified objects and the credits billed for these operations. The table function returns the billing history within a specified date range for your entire Snowflake account. This function returns billing activity within the last 14 days. Please note,

STAGE\_DIRECTORY\_FILE\_REGISTRATION\_HISTORY table function can be used to query information about the metadata history for a directory table, including:

- Files added or removed automatically as part of a metadata refresh.
- Any errors found when refreshing the metadata.

Question 5: **Incorrect**

**Select the correct statements for Table Clustering. (Select 3)**

- 

**Automatic clustering can not be suspended or resumed**

- 

**Automatic Clustering doesn't consume credit**

**(Incorrect)**

- 

**Tables in multi-terabytes range are good candidate for clustering keys**

**(Correct)**

- 

**Clustering keys are not for every table**

**(Correct)**

- 

**Snowflake recommends a maximum of three or four columns (or expressions) per key**

**(Correct)**

- 

**Snowflake doesn't charge for Reclustering**

**(Incorrect)**

### **Explanation**

Clustering keys are not for every table. Tables in the multi-terabyte range are good candidates for clustering keys. Both automatic clustering and reclustering consume credit. A single clustering key can contain one or more columns or expressions. Snowflake recommends a maximum of three or four columns (or expressions) per key for most tables. Adding more than 3-4 columns tends to increase costs more than benefits.

Question 6: **Incorrect**

**What is the purpose of VALIDATION\_MODE in the COPY INTO <table> command?**

- 

**VALIDATION\_MODE is used to validate the load file and load it into the specified table if there is no error.**

**(Incorrect)**

- 

**VALIDATION\_MODE is used to validate the load file, skip the errored data and then load it into the specified table.**

- 

**VALIDATION\_MODE is used to validate the load file for errors instead of loading it into the specified table.**

**(Correct)**

### **Explanation**

VALIDATION\_MODE instructs the COPY command to validate the data files instead of loading them into the specified table; i.e., the COPY command tests the files for errors but does not load them. The command validates the data to be loaded and returns results based on the validation option specified: Syntax : VALIDATION\_MODE = RETURN\_n\_ROWS | RETURN\_ERRORS | RETURN\_ALL\_ERRORS RETURN\_n\_ROWS (e.g. RETURN\_10\_ROWS) - Validates the specified number of rows, if no errors are encountered; otherwise, fails at the first error encountered in the rows.

RETURN\_ERRORS - Returns all errors (parsing, conversion, etc.) across all files specified in the COPY statement. RETURN\_ALL\_ERRORS - Returns all errors across all files specified in the COPY statement, including files with errors that were partially loaded during an earlier load because the ON\_ERROR copy option was set to CONTINUE during the load.

Question 7: **Correct**

**Suppose you have an auto-scaling mode setup with a Economy policy. In what situation does Snowflake spin up an additional cluster?**

- 

**Only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes.**

**(Correct)**

- 

**The first cluster starts immediately when either a query is queued or the system detects that there's one more query than the currently-running clusters can execute.**

### **Explanation**

In Economy Scaling policy, Snowflake spins up an additional cluster only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes

Question 8: **Correct**

**A stored procedure can simultaneously run the caller's and the owner's rights.**

**(True / False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. A **caller's rights stored procedure** runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query. An **owner's rights stored procedure** runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table. At the time that the stored procedure is created, the creator specifies whether the procedure runs with owner's rights or caller's rights. The default is owner's rights.

Question 9: **Incorrect**

**Which object parameter can users with the ACCOUNTADMIN role use to set the default retention period for their account?**

- -
- DATA\_RETENTION\_TIME\_MAX**
- -
- DATA\_RETENTION\_IN\_TIME\_TRAVEL**
- (Incorrect)**
- -
- DATA\_RETENTION\_TIME\_IN\_HOURS**
- -
- DATA\_RETENTION\_TIME\_IN\_DAYS**
- (Correct)**

### **Explanation**

Users can use the DATA\_RETENTION\_TIME\_IN\_DAYS object parameter with the ACCOUNTADMIN role to set the default retention period for their account.

Question 10: **Incorrect**

**Which copyOptions can help load a file with expired metadata (if the LAST\_MODIFIED date is older than 64 days and the initial set of data was loaded into the table more than 64 days earlier (and if the file was loaded into the table, that also occurred more than 64 days earlier))? (Select 2)**

- **LOAD\_UNCERTAIN\_FILES = TRUE**  
**(Correct)**
- **LOAD\_CERTAIN\_FILES = TRUE**
- **LOAD\_FILES = TRUE**
- **ON\_ERROR = CONTINUE**  
**(Incorrect)**
- **FORCE = TRUE**
- **FORCE = FALSE**  
**(Correct)**

### Explanation

To load files whose metadata has expired, set the LOAD\_UNCERTAIN\_FILES copy option to true. The copy option references load metadata, if available, to avoid data duplication, but also attempts to load files with expired load metadata. Alternatively, set the FORCE option to load all files, ignoring load metadata if it exists. Note that this option reloads files, potentially duplicating data in a table.

Question 11: **Incorrect**

**Dynamic Data Masking is supported by:**

- 
- **Standard Edition**
- **VPS**

**(Correct)**

- 

**Enterprise Edition**

**(Correct)**

- 

**Business Critical**

**(Correct)**

#### **Explanation**

**Dynamic Data Masking** features require Enterprise Edition (or higher).

Question 12: **Incorrect**

**What will happen if a policy is assigned to a user who is already signed in?**

- 

**The user can continue running the SQL queries in the currently opened session.**

**(Incorrect)**

- 

**The user can't do anything else until signed in and signed back in again.**

**(Correct)**

- 

**There will be no interruption until the user logoffs and signs in again.**

#### **Explanation**

If a policy is assigned to a user who already signed in, they can't do anything else until they sign and signed back in again to make use of the new policy

Question 13: **Incorrect**

**Which schema can be used to find out about storage, compute, and objects in a Snowflake account?**

-

## **RESOURCE\_SCHEMA**

- 

## **SNOWFLAKE\_SCHEMA**

- 

## **INFORMATION\_SCHEMA**

**(Correct)**

- 

## **USAGE\_SCHEMA**

**(Incorrect)**

### **Explanation**

INFORMATION\_SCHEMA can be used to find out about storage, compute, and objects in a Snowflake account. Every database that you create on Snowflake has a schema called INFORMATION\_SCHEMA that's automatically created, and inside that schema, you can find views and table functions that provide metadata information about objects in your account.

Question 14:

**Skipped**

**Which privileges are provided with a share by the provider? (Select 2)**

- 

**Grant access(USAGE) to the database and the schema containing the tables to share**

**(Correct)**

- 

**Grant access(SELECT) to the specific tables in the database**

**(Correct)**

- 

**Grant access(USAGE) to the specific tables in the database**

- 

**Grant access(OPERATE) to the database and the schema containing the tables to share**

- 

**Grant access(MODIFY) to the specific tables in the database**

### Explanation

Shares are named Snowflake objects that encapsulate all of the information required to share a database. Each share consists of:

- The privileges that grant access to the database(s) and the schema containing the objects to share.
- The privileges that grant access to the specific objects in the database.
- The consumer accounts with which the database and its objects are shared.

example: CREATE SHARE "SHARED\_DATA" COMMENT=""; GRANT USAGE ON DATABASE "DEMO\_DB" TO SHARE "SHARED\_DATA"; GRANT USAGE ON SCHEMA "DEMO\_DB"."TWITTER\_DATA" TO SHARE "SHARED\_DATA"; GRANT SELECT ON VIEW "DEMO\_DB"."TWITTER\_DATA"."FOLLOWERS" TO SHARE "SHARED\_DATA";

Question 15:

### Skipped

**What value will be return by the following query? SELECT \* FROM TABLE(FLATTEN(input => parse\_json('[]')) f;**

- 

[]

- 

**NULL**

- 

**0**

- 

**nothing will return / output of the input row will be omitted**

**(Correct)**

### Explanation

If you don't specify OUTER argument with FLATTEN, it would be defaulted to FALSE. The OUTER => FALSE argument with FLATTEN omits the output of the input rows

that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries.

Question 16:

**Skipped**

If you create a Network Policy by providing both 'Allowed IP Addresses' and 'Blocked IP Addresses', which is applied first by Snowflake while validating the access?

- 

**Allowed IP Addresses**

- 

**Blocked IP Addresses**

**(Correct)**

**Explanation**

If you provide both Allowed IP Addresses and Blocked IP Addresses, **Snowflake applies the Blocked List first**.

Question 17:

**Skipped**

Select the type of function that can operate on a subset of rows within the set of input rows.

- 

**Window Function**

**(Correct)**

- 

**System Function**

- 

**Scalar Function**

- 

**Table Function**

-

## Aggregate Function

- 

## User-Defined Function

### Explanation

A window function is any function that operates over a window of rows.

Question 18:

**Skipped**

**An account-level resource monitor overrides the resource monitor assignment for individual warehouses. (True/False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### Explanation

An account-level resource monitor does not override resource monitor assignments for individual warehouses. If either the account resource monitor or the warehouse resource monitor reaches its defined threshold and a suspend action has been defined, the warehouse is suspended.

Question 19:

**Skipped**

**Monica ran a SELECT query on a large table t1. The query took longer than expected. She looked into the query profile and found that 'Bytes spilled to local storage' and 'Bytes spilled to remote storage' are very high. What advice will you give to her to improve the query performance? (Select 3)**

- 

**Trying to split the processing into several steps**

- 

**Processing data in smaller batches**

**(Correct)**

- 

**Using a larger warehouse (effectively increasing the available memory/local disk space for the operation)**

**(Correct)**

- 

**Increasing the number of parallel queries running in the warehouse**

**(Correct)**

- 

**Processing data in larger batches**

### **Explanation**

When Snowflake warehouse cannot fit an operation in memory, it starts spilling (storing) data first to the local disk of a warehouse node, and then to remote storage. In such a case, Snowflake first tries to temporarily store the data on the warehouse's local disk. As this means extra IO operations, any query that requires spilling will take longer than a similar query running on similar data that is capable to fit the operations in memory. Also, if the local disk is not sufficient to fit the spilled data, Snowflake further tries to write to the remote cloud storage, which will be shown in the query profile as "Bytes spilled to remote storage".

The spilling can't always be avoided, especially for large batches of data, but it can be decreased by: Reducing the amount of data processed.

For example, by trying to improve partition pruning, or projecting only the columns that are needed in the output. Decreasing the number of parallel queries running in the warehouse. Trying to split the processing into several steps (for example by replacing the CTEs with temporary tables). Using a larger warehouse - effectively means more memory and more local disk space.

Question 20:

**Skipped**

**Which algorithm does Snowflake use to estimate the approximate number of distinct values in a data set?**

- 

**HyperMedianLog**

- 
- HyperEstimateLog**
- 
- HyerAccumulateLog**
- 
- HyperMeanLog**
- 
- HyperLogLog**
- (Correct)**

### Explanation

Snowflake uses HyperLogLog to estimate the approximate number of distinct values in a data set. HyperLogLog is a state-of-the-art cardinality estimation algorithm, capable of estimating distinct cardinalities of trillions of rows with an average relative error of a few percent.

Question 21:

### Skipped

**While transforming Semi-structure data, If you want expansion for all the sub-elements recursively using FLATTEN function, what argument would you need to set with FLATTEN function?**

- 
- OUTER => TRUE**
- 
- RECURSIVE => FALSE**
- 
- OUTER => FALSE**
- 
- RECURSIVE => TRUE**

**(Correct)**

### **Explanation**

The expansion is performed for all sub-elements recursively by argument RECURSIVE => TRUE. Only the element referenced by PATH is expanded BY RECURSIVE => FALSE. The OUTER argument is used to handle the input rows that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries.

Question 22:

**Skipped**

**Which of these system-defined roles can manage operations at the organization level?**

- 

**SYSADMIN**

- 

**USERADMIN**

- 

**ORGADMIN**

**(Correct)**

- 

**ACCOUNTADMIN**

- 

**SECURITYADMIN**

### **Explanation**

ORGADMIN role manages operations at the organizational level. More specifically, this role:

- Can create accounts in the organization.
- Can view all accounts in the organization (using SHOW ORGANIZATION ACCOUNTS) and all regions enabled for the organization (using SHOW REGIONS).
- Can view usage information across the organization.

Question 23:

**Skipped**

**Which data does not fit into a predefined data model or schema?**

- 
- **Semi-Structured Data**
- 
- **Structured-data**
- 
- **All of these**
- 
- **Unstructured Data**

**(Correct)**

### **Explanation**

Unstructured data is information that does not fit into a predefined data model or schema. Typically text-heavy, such as form responses and social media conversations, unstructured data also encompasses images, video, and audio. Industry-specific file types such as VCF (genomics), KDF (semiconductors), or HDF5 (aeronautics) are included in this category.

Question 24:

**Skipped**

**If you recreate a pipe using CREATE OR REPLACE PIPE command. What does happen to load history if the Snowpipe gets recreated?**

- 
- **The recreated Pipe still has tracks of the files loaded by the old Pipe**
- 
- **The pipe can not be recreated**
- 
- **Snowflake still keeps load history**
-

**The load history gets reset to empty**

**(Correct)**

### **Explanation**

When you recreate a pipe, if you do CREATE OR REPLACE PIPE, that load history is reset to empty, so Snowflake doesn't know which files we've already loaded.

Question 25:

**Skipped**

**What would happen if we suspend the warehouse while it is executing the SQL statement?**

- 

**When trying to suspend the warehouse, we will get an error while the same warehouse is executing SQL statements.**

- 

**All compute resources of the warehouse will be up until the statement is complete.**

- 

**Only idle compute resources of the warehouse will be shut down, allowing any compute resources executing statements to continue until the statement is complete.**

**(Correct)**

- 

**All the compute resources of the warehouse will be shut down immediately, and the running statement will be canceled.**

### **Explanation**

When we suspend a warehouse, Snowflake immediately shuts down all idle compute resources for the warehouse. However, it allows any compute resources executing statements to continue until the statements are complete. At this time, the resources are shut down, and the warehouse status changes to "Suspended". Compute resources waiting to shut down are considered to be in "quiesce" mode.

Question 26:

**Skipped**

**Monica has successfully created a task with the 5 minutes schedule. It has been 30 minutes, but the task did not run. What could be the reason?**

- Task schedule should not be less than 60 minutes
- Monica should run the ALTER TASK command to RESUME the task  
**(Correct)**
- Monica doesn't have the authority to run the task
- Monica should run the ALTER TASK command to SUSPEND the task, and then again run the ALTER TASK command to RESUME the task

### Explanation

The first time we create the TASK, we need to run the ALTER TASK command to RESUME the task.

Question 27:

**Skipped**

File URL is ideal for

- 
- Ideal for use in custom applications, providing unstructured data to other accounts via a share
- Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents
- 
- None of these
- Ideal for custom applications that require access to unstructured data files

**(Correct)**

### **Explanation**

Scoped URL: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. Ideal for use in custom applications, providing unstructured data to other accounts via a share, or for downloading and ad hoc analysis of unstructured data via Snowsight. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. Ideal for custom applications that require access to unstructured data files. Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable. Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents.

Question 28:

**Skipped**

**A task can execute any one of the following types of SQL code: (Select 3)**

- 

**Call to a stored procedure**

**(Correct)**

- 

**Single SQL Statement**

**(Correct)**

- 

**Procedural logic using Snowflake Scripting**

**(Correct)**

- 

**Multiple SQL statements**

### **Explanation**

A task can execute any one of the following types of SQL code:

- Single SQL statement
- Call to a stored procedure

- Procedural logic using Snowflake Scripting.

Question 29:

**Skipped**

**Which of the Snowflake editions provides a federated authorization feature?**

- **Virtual Private Snowflake(VPS)**
- 
- **Business Critical**
- 
- **All of the Snowflake Editions**
- **(Correct)**
- 
- **Enterprise**
- 
- **Standard**

**Explanation**

All Snowflake Editions (Standard, Enterprise, Business Critical, Virtual Private Snowflake) provide Federated Authentication.

Question 30:

**Skipped**

**Which of these SQL functions helps returns the absolute path of a staged file using the stage name and path of the file relative to its location in the stage as inputs.?**

- 
- **GET\_ABSOLUTE\_PATH**
- **(Correct)**
- 
- **GET\_RELATIVE\_PATH**

- 
- **BUILD\_STAGE\_FILE\_URI**
- 
- **GET\_STAGE\_LOCATION**
- 
- **GET\_PRESIGNED\_URL**
- 
- **BUILD\_SCOPED\_FILE\_URL**

### Explanation

GET\_ABSOLUTE\_PATH returns the absolute path of a staged file using the stage name and path of the file relative to its location in the stage as inputs.

Question 31:

**Skipped**

**Which services are managed by Snowflake's cloud services layer? (Select all that apply)**

- 
- **Access Control**  
**(Correct)**
- 
- **Only Infrastructure Management**
- 
- **Metadata Management**  
**(Correct)**
- 
- **Infrastructure Management**  
**(Correct)**

- 

### Authentication

(Correct)

- 

### Query Parsing and Optimization

(Correct)

### Explanation

The cloud services layer is a collection of services that coordinate activities across Snowflake. These services tie together all of the different components of Snowflake in order to process user requests, from login to query dispatch.

The cloud service layer manages Authentication, Infrastructure Management, Metadata Management, Query parsing and optimization, and Access control services.

Question 32:

Skipped

**Which database objects can be shared using the Snowflake Secure Data Sharing feature? (Select all that apply)**

- 

### Tables

(Correct)

- 

### External Tables

(Correct)

- 

### Roles

- 

### Secure Materialized View

(Correct)

- 

### Secure UDFs

(Correct)

- 

### Secure Views

(Correct)

## Explanation

**Secure Data Sharing** enables sharing selected objects in a database in your account with other Snowflake accounts. The following Snowflake database objects can be shared:

- Tables
- External tables
- Secure views
- Secure materialized views
- Secure UDFs

Snowflake enables the sharing of databases through shares created by data providers and “imported” by data consumers.

Question 33:

Skipped

**How many maximum columns (or expressions) are recommended for a cluster key?**

- 

**7 to 8**

- 

**3 to 4**

(Correct)

- 

**Higher the number of columns (or expressions) in the key, better will be the performance**

- 

## 12 to 16

### Explanation

A single clustering key can contain one or more columns or expressions. **Snowflake recommends a maximum of 3 or 4 columns (or expressions) per key for most tables.** Adding more than 3-4 columns tends to increase costs more than benefits.

Question 34:

### Skipped

**A user can be assigned multiple roles. (True / False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### Explanation

Roles are the entities to which privileges on securable objects can be granted and revoked. Roles are assigned to users to allow them to perform actions required for business functions in their organization. **A user can be assigned multiple roles.** It allows users to switch roles (i.e., choose which role is active in the current Snowflake session) to perform different actions using separate sets of privileges.

Question 35:

### Skipped

**Snowflake automatically and transparently maintains materialized views.**

**(True/False)**

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### Explanation

**Snowflake automatically and transparently maintains materialized views.** A background service updates the materialized view after changes to the base table. This is more efficient and less error-prone than manually maintaining the equivalent of a materialized view at the application level.

Question 36:

**Skipped**

**You have a table with a 30-day retention period. If you decrease the data period to 20 days, how would it affect the data that would have been removed after 30 days?**

- 
- 

**The data will still retain for 30-day before moving to Fail-safe**

- 
- 

**The data will now retain for a shorter period of 20 days**

**(Correct)**

### **Explanation**

Decreasing Retention reduces the amount of time data is retained in Time Travel: - For active data modified after the retention period is reduced, the new shorter period applies. - For data that is currently in Time Travel: - If the data is still within the new shorter period, it remains in Time Travel. - If the data is outside the new period, it moves into Fail-safe. For example, if you have a table with a 30-day retention period and you decrease the period to 20-day, data from days 21 to 30 will be moved into Fail-safe, leaving only the data from day 1 to 20 accessible through Time Travel. However, the process of moving the data from Time Travel into Fail-safe is performed by a background process, so the change is not immediately visible. Snowflake guarantees that the data will be moved, but does not specify when the process will complete; until the background process completes, the data is still accessible through Time Travel.

Question 37:

**Skipped**

**Which SQL command determines whether a network policy is set on the account or for a specific user?**

- 
- 

**SHOW PARAMETER**

- 
-

## SHOW PARAMETERS

(Correct)

- 

**SHOW NETWORK\_POLICIES**

- 

**SHOW POLICIES**

- 

**SHOW POLICY**

### Explanation

The **SHOW PARAMETERS** command determines whether a network policy is set on the account or for a specific user.

**For Account level:** SHOW PARAMETERS LIKE 'network\_policy' IN ACCOUNT;

**For User level :** SHOW PARAMETERS LIKE 'network\_policy' IN USER <username>;

**Example** - SHOW PARAMETERS LIKE 'network\_policy' IN USER john;

Question 38:

### Skipped

Monica wants to delete all the data from table t1. She wants to keep the table structure, so she does not need to create the table again. Which command will be appropriate for her need?

- 

**REMOVE**

- 

**UNDROP**

- 

**DELETE**

-

## TRUNCATE

(Correct)

- 

## DROP

### Explanation

TRUNCATE will delete all of the data from a single table. So, once Monica truncates table t1, table t1's structure remains, but the data will be deleted. DELETE is usually used for deleting single rows of data.

Question 39:

Skipped

**Which is generally the slowest option for selecting staged data files to load from a stage?**

- 

**By path (internal stages) / prefix (Amazon S3 bucket)**

- 

**Using pattern matching to identify specific files by pattern**

(Correct)

- 

**Specifying a list of specific files to load**

### Explanation

Pattern matching using a regular expression is generally the slowest of the three options for identifying/specifying data files to load from a stage; however, this option works well if you exported your files in named order from your external application and want to batch load the files in the same order

Question 40:

Skipped

A user cannot view the result set from a query that another user executed except for the ACCOUNTADMIN role. (True / False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

**A user cannot view the result set from a query that another user executed.** This behavior is intentional. For security reasons, only the user who executed a query can access the query results. This behavior is not connected to the Snowflake access control model for objects. **Even a user with the ACCOUNTADMIN role cannot view the results for a query run by another user.**

Question 41:

**Skipped**

**There are two modes to set up a multi-cluster warehouse. Select those from the given choices.**

- 

**Minimized mode**

- 

**Maximum mode**

- 

**Minimum mode**

- 

**Auto-scaling mode**

**(Correct)**

- 

**Maximized mode**

**(Correct)**

### **Explanation**

There are two ways to set up a multi-cluster warehouse: in maximized mode, or auto-scaling mode. With maximized mode, you simply set your minimum equal to your maximum, and those values are something greater than one.

Question 42:

**Skipped**

**If you drop or disable a user in Snowflake in an Okta IdP federated environment, the user can still access Snowflake login through Okta.**  
**(True/False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

**Explanation**

Users who are dropped or disabled in Snowflake are still able to log into their Okta accounts, but they will receive an error message when they attempt to connect to Snowflake. You must recreate or enable the user before they can log in.

Question 43:

**Skipped**

**Both external (external cloud storage) and internal (i.e., Snowflake) stages support unstructured data. (True / False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

**Explanation**

True, both external (external cloud storage, such as, Amazon S3, Google Cloud Storage, Azure Blob Storage etc.) and internal (i.e. Snowflake) stages support unstructured data.

Question 44:

**Skipped**

**In which of the cloud platforms a Snowflake account can be hosted? (Select 3)**

-

## AZURE

(Correct)

- 

## IBM Cloud

- 

## AWS

(Correct)

- 

## Oracle Cloud

- 

## GCP

(Correct)

### Explanation

A Snowflake account can be hosted on any of the following cloud platforms: Amazon Web Services (AWS), Google Cloud Platform (GCP), Microsoft Azure (Azure). On each platform, Snowflake provides one or more regions where the account is provisioned.

Question 45:

Skipped

If a user is logged in to Snowflake in a federated environment and IdP times out, what does happen to the user's snowflake session?

- 

The Snowflake web interface is disabled, and the prompt for IdP authentication is displayed.

- 

It does not affect the user's Snowflake sessions. However, to initiate any new Snowflake sessions, the user must log into the IdP again.

(Correct)

## Explanation

After a specified period of time (defined by the IdP), a user's session in the IdP automatically times out, but this does not affect their Snowflake sessions. Any Snowflake sessions that are active at the time remain open and do not require re-authentication. However, to initiate any new Snowflake sessions, the user must log into the IdP again.

Question 46:

### Skipped

**Which objects are not available for replication in the Standard Edition of Snowflake? (Select 3)**

- 

**Users**

- 

**Integrations**

- 

**Shares**

**(Correct)**

- 

**Database**

**(Correct)**

- 

**Roles**

## Explanation

Database and share replication are available in all editions, including the Standard edition. Replication of all other objects is only available for Business Critical Edition (or higher).

Question 47:

### Skipped

**Which of these functions helps generate the FILE URL to access the unstructured data file?**

-

### **GET\_PRESIGNED\_URL**

- 

### **BUILD\_SCOPED\_FILE\_URL**

- 

### **BUILD\_STAGE\_FILE\_URI**

**(Correct)**

- 

### **GET\_RELATIVE\_PATH**

- 

### **GET\_ABSOLUTE\_PATH**

- 

### **GET\_STAGE\_LOCATION**

#### **Explanation**

BUILD\_STAGE\_FILE\_URL generates a Snowflake-hosted file URL to a staged file using the stage name and relative file path as inputs. A file URL permits prolonged access to a specified file. That is, the file URL does not expire. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files.

Question 48:

**Skipped**

**How can you unload the data from Snowflake using COPY INTO location statements in a Single file?**

- 

**By specifying copy option MULTIPLE=FALSE**

- 

**By specifying copy option SINGLE=TRUE**

**(Correct)**

- 

**By specifying copy option ONE\_FILE=TRUE**

- 

**By specifying copy option MULTIPLE\_FILES=FALSE**

#### **Explanation**

To unload data to a single output file (at the potential cost of decreased performance), specify the SINGLE = true copy option in your statement. You can optionally specify a name for the file in the path.

Question 49:

**Skipped**

**At what frequency does Snowflake rotate the object keys?**

- 

**1 Year**

- 

**60 Days**

- 

**16 Days**

- 

**30 Days**

**(Correct)**

#### **Explanation**

All Snowflake-managed keys are automatically **rotated by Snowflake when they are more than 30 days** old. Active keys are retired, and new keys are created. When Snowflake determines the retired key is no longer needed, the key is automatically destroyed. When active, a key is used to encrypt data and is available for usage by the customer. When retired, the key is used solely to decrypt data and is only available for accessing the data.

Question 50:

**Skipped**

**Snowflake Query history page allows you to view the details of all the queries executed in the last 31 days. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

Snowflake Query history page allows you to view the details of all the queries executed in the last 14 days. You can query the Query\_History view in Snowflake's Account Usage schema for older queries.

Question 51:

**Skipped**

**How can we turn off the query result cache?**

- 

**Setting the parameter USE\_CACHED\_RESULT to FALSE**

**(Correct)**

- 

**Setting the parameter USE\_QUERY\_CACHED to FALSE**

- 

**Query result cache can be turned off.**

- 

**Setting the parameter USE\_CACHED\_INFO to FALSE**

### **Explanation**

We can turn off the query result cache by setting the parameter **USE\_CACHED\_RESULT to FALSE**. Though the only reason we would really want to do this is if we are doing performance testing.

Question 52:

**Skipped**

**Which is not the DML (Data Manipulation Language) command?**

-

## UPDATE

- 

## INSERT

- 

## UNDROP

(Correct)

- 

## TRUNCATE

- 

## DELETE

- 

## MERGE

### Explanation

UNDROP is Snowflake's DDL (Data Definition Language) command.

Question 53:

### Skipped

**When deciding whether to suspend a warehouse or leave it running, what should you consider?**

- 

**Consider suspending the warehouse if the warehouse is large and there are no active queries.**

- 

**Consider the trade-off between saving credits by suspending the warehouse versus maintaining the cache of data from the previous queries to help with performance.**

(Correct)

-

**Consider the trade-off between saving credits by suspending the warehouse versus the operational cost of resuming the warehouse when needed.**

### Explanation

Consider the trade-off between saving credits by suspending a warehouse versus maintaining the cache of data from previous queries to help with performance.

Question 54:

### Skipped

**In what situations should you consider User-Managed Tasks over Serverless Tasks? (Select 2)**

- 

**Consider when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources.**

**(Correct)**

- 

**Consider when adherence to the schedule interval is less important.**

**(Correct)**

- 

**Consider when adherence to the schedule interval is highly important.**

- 

**Consider when you cannot fully utilize a warehouse because too few tasks run concurrently or they run to completion quickly (in less than 1 minute).**

### Explanation

**User-managed Tasks** is recommended when you **can fully utilize a single warehouse** by scheduling multiple concurrent tasks to take advantage of available compute resources. Also, recommended when adherence to the schedule interval is less critical. **Serverless Tasks** is recommended when you **cannot fully utilize a warehouse** because too few tasks run concurrently or they run to completion quickly (in less than 1 minute). Also, recommended when adherence to the schedule interval is critical.

Question 55:

## Skipped

Which of these IdP vendors provides native Snowflake support For federated authentication and SSO?

- 

**Microsoft Azure Active Directory**

- 

**Onelogin**

- 

**Microsoft ADFS**

**(Correct)**

- 

**Google G Suite**

- 

**Okta**

**(Correct)**

## Explanation

Okta and Microsoft ADFS provide native Snowflake support for federated authentication and SSO.

Question 56:

## Skipped

A user's default role is

- 

**the name used to log in to the Snowflake WebUI.**

- 

**the role a user gets set to each time the user logs in to Snowflake.**

**(Correct)**

-

**changed each time the user logs in to Snowflake.**

- 

**always the default PUBLIC role.**

### **Explanation**

A user's default role is the role a user gets set to each time the user logs in to Snowflake. Snowflake uses roles to control the objects (virtual warehouses, databases, tables, etc.) that users can access:

- Snowflake provides a set of predefined roles, as well as a framework for defining a hierarchy of custom roles.
- All Snowflake users are automatically assigned the predefined PUBLIC role, which enables login to Snowflake and basic object access.
- In addition to the PUBLIC role, each user can be assigned additional roles, with one of these roles designated as their default role.
- A user's default role determines the role used in the Snowflake sessions initiated by the user; however, this is only a default. Users can change roles within a session at any time.
- Roles can be assigned at user creation or afterward.

Question 57:

**Skipped**

**Which stream type is supported for streams on the external table only?**

- 

**Update-only**

- 

**Standard**

- 

**Insert-only**

**(Correct)**

- 

**Append-only**

- 

**External**

## Explanation

**Insert-only** is supported for streams on external tables only. An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops).

Question 58:

### Skipped

**Which of the following languages does Snowflake support for writing UDFs (User-Defined Functions)?**

- 

**Python**

**(Correct)**

- 

**C#**

- 

**JAVA**

**(Correct)**

- 

**SQL**

**(Correct)**

- 

**JavaScript**

**(Correct)**

- 

**GO**

## Explanation

User-defined functions (UDFs) let you extend the system to perform operations that are not available through the built-in, system-defined functions provided by Snowflake. Snowflake currently supports the following languages for writing UDFs:

Java: A Java UDF lets you use the Java programming language to manipulate data and return either scalar or tabular results. JavaScript: A JavaScript UDF lets you use the JavaScript programming language to manipulate data and return either scalar or tabular results. Python: A Python UDF lets you use the Python programming language to manipulate data and return either scalar or tabular results. SQL: A SQL UDF evaluates an arbitrary SQL expression and returns either scalar or tabular results.

Question 59:

**Skipped**

**Permissions on database objects such as databases or tables are granted to:**

- 
- Users**
- 
- Virtual Warehouses**
- 
- Roles**
- (Correct)**
- 
- Schemas**

**Explanation**

**Snowflake supports Role-Based Access control.** Permissions on database objects such as databases or tables are granted to Roles.

Question 60:

**Skipped**

**If an account has federated authentication enabled. Can Snowflake admins still maintain user IDs and passwords in Snowflake?**

- 
- Yes**
- (Correct)**
- 
- No**

### **Explanation**

With federated authentication enabled on an account, Snowflake still allows maintaining and using Snowflake user credentials (login name and password). In other words:

- Account and security administrators can still create users with passwords maintained in Snowflake.
- Users can still log into Snowflake using their Snowflake credentials.

However, if federated authentication is enabled for an account, Snowflake does not recommend maintaining user passwords in Snowflake. Instead, user passwords should be maintained solely in your IdP.

Question 61:

**Skipped**

**Which of these are kind of Cache in Snowflake?**

- 

**Query Result Cache**

- 

**All of these**

**(Correct)**

- 

**Metadata Cache**

- 

**Data/Local Disk Cache**

### **Explanation**

Snowflake has three types of cache.

- The **metadata cache** that lives in the cloud services layer.
- The **data cache/local disk cache** that lives on the SSD drives in the virtual warehouses, and
- The **query result cache**. If a result is small, it will be stored in the cloud services layer, but larger results are going to be stored in the storage layer.

Question 62:

**Skipped**

**Multi-cluster warehouses are beneficial in improving the performance of slow-running queries or data loading. (True/False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### **Explanation**

Multi-cluster warehouses are best utilized for scaling resources to improve concurrency for users/queries. They are not as beneficial for improving the performance of slow-running queries or data loading. For these types of operations, resizing the warehouse provides more benefits.

Question 63:

**Skipped**

**Snowflake is available in four editions. Which are those? (Select 4)**

- 

**Standard**

**(Correct)**

- 

**Virtual Private Snowflake (VPS)**

**(Correct)**

- 

**Enterprise**

**(Correct)**

- 

**Professional Plus**

-

## **Business Critical**

**(Correct)**

- 

## **Professional**

### **Explanation**

Snowflake is available in four editions: Standard, Enterprise, Business Critical, and Virtual Private Snowflake (VPS). Standard comes with most of the available features. Enterprise adds on to Standard with things like: extra days of time travel, materialized view support, and data masking. Business Critical brings to the table: HIPAA support, Tri-secret Secure, and more. And Virtual Private Snowflake is everything that Business Critical has, but with the ability to have customer-dedicated metadata stores and customer-dedicated virtual service.

Question 64:

**Skipped**

**Search optimization is a Database-level property applied to all the tables within the database with supported data types. (True/False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### **Explanation**

**Search optimization** is a table-level property and applies to all columns with supported data types. The search optimization service aims to significantly improve the performance of selective point lookup queries on tables. A point lookup query returns only one or a small number of distinct rows. A user can register one or more tables to the search optimization service.

Question 65:

**Skipped**

**Which view in the Account Usage Schema can be used to query the replication history for a specified database?**

- 

**DATABASE\_REFRESH\_HISTORY**

- 

### **REPLICATION\_USAGE\_HISTORY**

**(Correct)**

- 

### **REPLICATION\_GROUP\_REFRESH\_HISTORY**

- 

### **DATA\_TRANSFER\_HISTORY**

#### **Explanation**

This REPLICATION\_USAGE\_HISTORY view in the Account Usage Schema can be used to query the replication history for a specified database. The returned results include the database name, credits consumed, and bytes transferred for replication. Usage data is retained for 365 days (1 year).

Question 66:

**Skipped**

**What all options are available for data transformation while loading data into a table using the COPY command? (Select that all apply)**

- 

#### **Column reordering**

**(Correct)**

- 

#### **Join**

- 

#### **Casts**

**(Correct)**

- 

#### **Truncation of Text Strings**

**(Correct)**

- 

### Column omission

(Correct)

#### Explanation

Snowflake supports transforming data while loading it into a table using the COPY command. Options include: Column reordering Column omission Casts Truncating text strings that exceed the target column length

Question 67:

Skipped

The major benefits of defining Clustering: (Select 2)

- 

To help in faster data sharing

- 

To help improve query performance

(Correct)

- 

To help optimize table maintenance

(Correct)

- 

To help in organizing small tables (<1 GB)

#### Explanation

Defining clustering keys for very large tables (in the multi-terabyte range) **helps optimize table maintenance and query performance**. Small tables are not a good candidate for clustering.

Question 68:

Skipped

Which of these are Snowgrid's capabilities? (Select all that apply)

- 

Live, ready to query data

**(Correct)**

- 

**Secure, governed data sharing**

**(Correct)**

- 

**ETL dependent**

- 

**Zero-copy cloning**

**(Correct)**

- 

**Share internally with private data exchange or externally with public data exchange**

**(Correct)**

### **Explanation**

**Snowgrid** allows you to use Secure Data Sharing features to provide access to live data, without any ETL or movement of files across environments.

Question 69:

**Skipped**

**Which primary tool loads data to Snowflake from a local file system?**

- 

**Snowflake UI**

- 

**SnowSQL**

**(Correct)**

- 

**External Stage**

- ○

## ETL tools

### Explanation

SnowSQL is the primary tool used to load data to Snowflake from a local file system. You can run it in either interactive shell or batch mode.

Question 70:

**Skipped**

**What all locations do Snowflake support for staging the data?**

- 

**Snowflake Internal Stages**

**(Correct)**

- 

**Google Cloud Storage**

**(Correct)**

- 

**Amazon S3**

**(Correct)**

- 

**Oracle Cloud Storage**

- 

**Microsoft Azure Blob Storage**

**(Correct)**

### Explanation

Snowflake supports loading data from files staged in any of the following locations, regardless of the cloud platform for your Snowflake account: Internal (i.e. Snowflake) stages, Amazon S3, Google Cloud Storage, Microsoft Azure blob storage

Question 71:

**Skipped**

**What would you create (UDF or Stored procedure) if you need a function that can be called as part of a SQL statement and must return a value that will be used in the statement?**

- 

**UDF**

**(Correct)**

- 

**Stored Procedure**

### **Explanation**

An UDF evaluates to a value, and can be used in contexts in which a general expression can be used (e.g. SELECT my\_function() ...). A stored procedure does not evaluate to a value, and cannot be used in all contexts in which a general expression can be used. For example, you cannot execute SELECT my\_stored\_procedure()....

Question 72:

**Skipped**

**Which of these SQL functions does Snowflake support?**

- 

**Scalar**

**(Correct)**

- 

**Table**

**(Correct)**

- 

**User-Defined**

**(Correct)**

- 

**System**

**(Correct)**

- 

**Aggregate**

**(Correct)**

- 

**Window**

**(Correct)**

### **Explanation**

Snowflake Supports all these SQL functions.

Question 73:

**Skipped**

**The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake. (True/False)**

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### **Explanation**

Snowflake manages all aspects of how this data is stored — the organization, file size, structure, compression, metadata, statistics, and other aspects of data storage are handled by Snowflake. The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake.

Question 74:

**Skipped**

For which object the Kafka connector does create a topic?

- 

**One pipe to ingest the data files for each topic partition**

- All of these  
**(Correct)**
- One internal stage to temporarily store data files for each topic
- One table for each topic. If the table specified for each topic does not exist

### Explanation

The connector creates the following objects for each topic:

- One internal stage to temporarily store data files for each topic.
- One pipe to ingest the data files for each topic partition.
- One table for each topic. If the table specified for each topic does not exist, the connector creates it; otherwise, the connector creates the RECORD\_CONTENT and RECORD\_METADATA columns in the existing table and verifies that the other columns are nullable (and produces an error if they are not).

Question 75:

### Skipped

Which command will list the pipes for which you have access privileges?

- LIST PIPES;
- DESCRIBE PIPES;
- LIST PIPES();
- SHOW PIPES();
-

**SHOW PIPES;**

**(Correct)**

### **Explanation**

**SHOW PIPES** Command lists the pipes for which you have access privileges. This command can list the pipes for a specified database or schema (or the current database/schema for the session), or your entire account.

Question 76:

**Skipped**

**Which systems function can help find the overlap depth of a table's micro-partitions?**

- 

**SYSTEM\$CLUSTERING\_WEIGHT**

- 

**SYSTEM\$CLUSTERING\_INFORMATION**

**(Correct)**

- 

**SYSTEM\$CLUSTERING\_INFO**

- 

**SYSTEM\$CLUSTERING\_ALL**

- 

**SYSTEM\$CLUSTERING\_DEPTH**

**(Correct)**

### **Explanation**

For example, if you have an EMPLOYEE table - you can run any of these queries to find the depth - `SELECT SYSTEM$CLUSTERING_INFORMATION('EMPLOYEE');` `SELECT SYSTEM$CLUSTERING_DEPTH('EMPLOYEE');`

Question 77:

**Skipped**

**The VALIDATION\_MODE parameter does not support COPY statements that transform data during a load. (True / False)**

- FALSE
- TRUE

**(Correct)**

### Explanation

True. "VALIDATION\_MODE instructs the COPY command to validate the data files instead of loading them into the specified table; i.e., the COPY command tests the files for errors but does not load them. The command validates the data to be loaded and returns results based on the validation option specified: Syntax :

VALIDATION\_MODE = RETURN\_n\_ROWS | RETURN\_ERRORS | RETURN\_ALL\_ERRORS  
RETURN\_n\_ROWS (e.g. RETURN\_10\_ROWS) - Validates the specified number of rows, if no errors are encountered; otherwise, fails at the first error encountered in the rows. RETURN\_ERRORS - Returns all errors (parsing, conversion, etc.) across all files specified in the COPY statement. RETURN\_ALL\_ERRORS - Returns all errors across all files specified in the COPY statement, including files with errors that were partially loaded during an earlier load because the ON\_ERROR copy option was set to CONTINUE during the load."

Question 78:

### Skipped

**select \* from t1 sample row(100); What would the above query return?**

- Return an empty sample
- samplingMethod is not applied in the query. The query will result in an error.
- Return a sample of a table in which each row has a 10% probability of being included in the sample
- Return an entire table, including all rows in the table

**(Correct)**

### **Explanation**

Return an entire table, including all rows in the table. Sampling method is optional. If no method is applied after sample keyword, the default it takes is BERNOULLI.

Question 79:

**Skipped**

**What happens to the data when the retention period ends for an object?**

- 
- Data can be restored by increasing the retention period**
- 
- Data is permanently lost**
- 
- Data is moved to Snowflake Fail-safe**

**(Correct)**

- 
- SYSADMIN can restore the data from Fail-safe**

### **Explanation**

When the retention period ends for an object, the historical data is moved into Snowflake Fail-safe. Snowflake support needs to be contacted to get the data restored from Fail-safe.

Question 80:

**Skipped**

**What authentication methods does Snowflake support for REST API authentication?**

- 
- Authentication is not required in case Snowflake SQL API**
- 
- Snowflake Account User ID and Password**
-

## OAuth

(Correct)

- 

## Key Pair Authentication

(Correct)

### Explanation

Snowflake SQL API supports Oauth, and Key Pair authentication.

Question 81:

Skipped

**How can we add a Directory table explicitly to a stage to store a catalog of staged files?**

- 

## Using CREATE STAGE command

(Correct)

- 

## Using CREATE DIRECTORY TABLES command and then add to the stage by ALTER STAGE command

- 

## Using CREATE DIRECTORY TABLE command and then add to the stage by ALTER STAGE command

### Explanation

A Directory table is not a separate database object; it stores a catalog of staged files in cloud storage. Roles with sufficient privileges can query a directory table to retrieve file URLs to access the staged files and other metadata. A directory table can be added explicitly to a stage when the stage is created (using CREATE STAGE) or later (using ALTER STAGE) with supplying directoryTableParams.

directoryTableParams (for internal stages) ::= [ DIRECTORY = ( ENABLE = { TRUE | FALSE } [ REFRESH\_ON\_CREATE = { TRUE | FALSE } ] ) ] ENABLE = TRUE | FALSE

Specifies whether to add a directory table to the stage. When the value is TRUE, a directory table is created with the stage.

Question 82:

Skipped

**John is trying to load JSON data sets with a huge array containing multiple records. Considering the VARIANT data type imposed size of 16 MB, what do you recommend to John for optimally loading the data?**

- Enable the STRIP\_OUTER\_ARRAY file format option for the COPY INTO <table> command**  
**(Correct)**
- Enable VARIANT\_OUTER\_ARRAY file format for the COPY INTO <table> command**
- Separate the documents with line break of commas**
- No need to remove the outer array structure as Snowflake Intelligent Engine will take care of that**

#### **Explanation**

If the data exceeds 16 MB, enable the STRIP\_OUTER\_ARRAY file format option for the COPY INTO <table> command to remove the outer array structure and load the records into separate table rows: copy into <table> from @~/<file>.json file\_format = (type = 'JSON' strip\_outer\_array = true);

Question 83:

**Skipped**

**Which of these compression methods are supported by Snowflake? (Select all that apply)**

- deflate**  
**(Correct)**
- gzip**

**(Correct)**

- 

**raw\_deflate**

**(Correct)**

- 

**Zstandard**

**(Correct)**

- 

**bzip2**

**(Correct)**

- 

**Brotli**

**(Correct)**

### **Explanation**

All of these are supported by Snowflake. Snowflake can automatically detect any of these compression methods except Brotli and Zstandard.

Question 84:

**Skipped**

**You can create an account level network policy using \_\_\_\_ (Select all that apply)**

- 

**SQL**

**(Correct)**

- 

**Classic Web Interface**

**(Correct)**

- 

**Snowsight**

**(Correct)**

- 

**Only Snowflake Support can create the Account level Network Policy**

### **Explanation**

Only security administrators (i.e., users with the SECURITYADMIN role) or higher or a role with the global CREATE NETWORK POLICY privilege can create network policies using Snowsight, Classic Web Interface, and SQL.

Question 85:

**Skipped**

**Which of these Snowflake features does enable accessing historical data (i.e., data that has been changed or deleted) at any point within a defined period?**

- 

**Search Optimization Service**

- 

**Time Travel**

**(Correct)**

- 

**Zero Copy Cloning**

- 

**Data Sharing**

### **Explanation**

Snowflake Time Travel enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It serves as a powerful tool for performing the following tasks: -Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted. - Duplicating and backing up data from key points in the past. - Analyzing data usage/manipulation over specified periods of time.

Question 86:

## Skipped

How can an ACCOUNTADMIN view the billing for Automatic Clustering? (Select all that apply)

- 

**Classic Web Interface: Click on Account > Billing & Usage under storage named 'AUTOMATIC\_CLUSTERING'**

- 

**Query - AUTOMATIC\_CLUSTERING\_HISTORY View (in Account Usage)**

- 

**Query - AUTOMATIC\_CLUSTERING\_HISTORY table function (in the Snowflake Information Schema)**

**(Correct)**

- 

**Classic Web Interface: Click on Account > Billing & Usage under warehouse named 'AUTOMATIC\_CLUSTERING'**

**(Correct)**

- 

**There is no way to check the Automatic Clustering billing without contacting Snowflake Support Team**

**(Correct)**

- 

**Snowsight: Select Admin > Usage**

**(Correct)**

## Explanation

Users with the ACCOUNTADMIN role can view the billing for Automatic Clustering using Snowsight, the classic web interface, or SQL: Snowsight: Select Admin » Usage. Classic Web Interface: Click on Account tab » Billing & Usage The billing for Automatic Clustering shows up as a separate Snowflake-provided warehouse named

AUTOMATIC\_CLUSTERING. SQL:Query either of the following:

AUTOMATIC\_CLUSTERING\_HISTORY table function (in the Snowflake Information Schema). AUTOMATIC\_CLUSTERING\_HISTORY View (in Account Usage).

Question 87:

**Skipped**

**Time Travel can be disabled for an account by ACCOUNTADMIN. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

Time Travel cannot be disabled for an account. A user with the ACCOUNTADMIN role can set DATA\_RETENTION\_TIME\_IN\_DAYS to 0 at the account level, which means that all databases (and subsequently all schemas and tables) created in the account have no retention period by default; however, this default can be overridden at any time for any database, schema, or table.

Question 88:

**Skipped**

**If we make any changes to the original table, then**

- 

**The changes get immediately reflected in the cloned table**

- 

**The cloned table data get refreshed with the entire new data of the source table**

- 

**The changes do not reflect in the cloned table**

**(Correct)**

### **Explanation**

Zero-copy cloning allows us to make a snapshot of any table, schema, or database without actually copying data. A clone is writable and is independent of its source

(i.e., changes made to the source or clone are not reflected in the other object). A new clone of a table points to the original table's micro partitions, using no data storage. If we make any changes in the cloned table, then only its changed micro partitions are written to storage.

Question 89:

**Skipped**

**Which of these are not supported by the Search Optimization Service? (Select all that apply)**

- 

**Column Concatenation**

**(Correct)**

- 

**Analytical Expressions**

**(Correct)**

- 

**Columns defined with COLLATE clause**

**(Correct)**

- 

**External Tables**

**(Correct)**

- 

**Casts on table columns**

**(Correct)**

- 

**Materialized Views**

**(Correct)**

**Explanation**

**None of these** are currently supported by the **Search Optimization Service**.

Additionally, Tables and views protected by row access policies cannot be used with the Search Optimization Search.

Question 90:

**Skipped**

**Only the user who generated the scoped URL can use the URL to access the referenced file. (True/False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### **Explanation**

True, only the user who generated the scoped URL can use the URL to access the referenced file. In case of File URL, any role that has sufficient privileges on the stage can access the file.

Question 91:

**Skipped**

**Which of these types of VIEW does Snowflake support? (Select 3)**

- 

**SECURE VIEW**

**(Correct)**

- 

**EXTERNAL VIEW**

- 

**MATERIALIZED VIEW**

**(Correct)**

- 

**TEMPORARY VIEW**

- 

## STANDARD VIEW

(Correct)

- 

## PERMANENT VIEW

### Explanation

Snowflake supports three types of views.

Standard View, Secure View, and Materialized View.

**Standard View:** It is a default view type. Its underlying DDL is available to any role with access to the view. When you create a standard view, Snowflake saves a definition of the view. Snowflake does not run the query. When someone accesses the view, that is when the query is run. The standard view will always execute as the owning role.

**Secure View:** The secure view is exactly like a standard view, except users cannot see how that view was defined. Sometimes a secure view will run a little slower than a standard view to protect the information in a secure view. Snowflake may bypass some of the optimizations.

**Materialized View:** A materialized view is more like a table. Unlike a standard or secure view, Snowflake runs the query right away when you create a materialized view. It takes the results set and stores that result set as a table in Snowflake. Because Snowflake is storing that materialized view as a table, creating micro partitions. Snowflake is creating metadata about those micro partitions. So when you query a materialized view, if you put a filter on the view, you get the same benefit of micro partition pruning that you would get from a table. With Snowflake, the materialized view is automatically refreshed every time there is a transaction against the base table. So it is always going to be in sync. If you want, you can also create a secure materialized view, which again will hide the logic from the user. A note about materialized views, because Snowflake is auto-refreshing them in the background, they use some credits, so there is a little bit of a cost there. Moreover, there is some storage, and Snowflake stores the result set as a table in Snowflake. So materialized views use more storage and compute than standard or secure views.

Question 92:

**Skipped**

**How long do results remain in the Query results cache?**

- 24 hours**  
**(Correct)**
- 1 hours**
- 12 hours**
- 31 hours**
- 16 hours**

#### Explanation

Results are **retained for 24 hours** in Query Result Cache. Snowflake resets the 24-hour retention period for the result, **up to a maximum of 31 days** from the date and time that the query was first executed. After 31 days, the result is purged and the next time the query is submitted, a new result is generated and persisted.

Question 93:

#### Skipped

**The suspended warehouse cannot be resized until they resume. (True / False)**

- TRUE**
- FALSE**

**(Correct)**

#### Explanation

The suspended warehouse can be easily resized. Resizing a suspended warehouse does not provision any new compute resources for the warehouse. It simply instructs Snowflake to provision the additional compute resources when the warehouse is next

resumed, at which time all the usage and credit rules associated with starting a warehouse apply.

Question 94:

**Skipped**

**John has to create a PIPE that will be triggered for loading by calling the Snowpipe REST endpoints. What parameter does he need to specify in CREATE PIPE statement?**

- API\_INGEST = FALSE**
  - API\_INGEST = TRUE**
  - AUTO\_INGEST = TRUE**
  - AUTO\_INGEST = FALSE**
- (Correct)**

### **Explanation**

AUTO\_INGEST = TRUE enables automatic data loading. Snowpipe supports loading from external stages (Amazon S3, Google Cloud Storage, or Microsoft Azure).

AUTO\_INGEST = FALSE disables automatic data loading. You must make calls to the Snowpipe REST API endpoints to load data files.

Question 95:

**Skipped**

**Which of the following file format is not supported by Snowflake?**

- JSON**
- CSV**
-

## **PARQUET**

- 

## **AVRO**

- 

## **ORC**

- 

## **EDI**

**(Correct)**

### **Explanation**

Snowflake supports - CSV, TSV, JSON, AVRO, ORC, PARQUET. Snowflake also supports XML which is a Preview feature as of now. EDI format is not supported by Snowflake.

Question 96:

**Skipped**

### **What is the expiration period of a File URL?**

- 

**The URL expires when the persisted query result period ends**

- 

**Length of time specified in the expiration\_time argument**

- 

**It is Permanent**

**(Correct)**

### **Explanation**

The expiration period of Scoped URL: The URL expires when the persisted query result period ends. The expiration period of the File URL: It is permanent. The expiration period of Pre-Signed URL: Length of time specified in the expiration\_time argument.

Question 97:

**Skipped**

**David ran a query that took around 30 mins to complete. He referred to the Query profiler and found the 'Bytes spilled to local storage' has a big number. What could be the issue?**

- **David is using a comparatively larger warehouse.**
  - **David is using a comparatively smaller warehouse.**
- (Correct)**
- **David should contact Snowflake Personnel.**
  - **Warehouse size has no impact on Bytes spilling.**

#### **Explanation**

If a node has insufficient memory to complete its portion of a query, it will "spill" to local SSD storage. This can negatively impact performance but is sometimes acceptable. If a node has insufficient local SSD storage to complete its portion of a query, it will "spill" to remote cloud storage. This is almost always very bad for performance. The solution, in either case, is to simplify the SQL query or increase the warehouse size.

Question 98:

#### **Skipped**

**Readers accounts enable providers to share data with consumers who are not already Snowflake customers without requiring the consumers to become Snowflake Customers. Which role can create the Reader account?**

- 
- ACCOUNTADMIN**
- (Correct)**
- **SECURITYADMIN**

- 

**USERADMIN**

- 

**SYSADMIN**

**Explanation**

ACCOUNTADMIN role (or a role granted the CREATE ACCOUNT global privilege) only can create the Reader account.

Question 99:

**Skipped**

**The user access history can be found by querying the**

- 

**Account Usage ACCESS\_HISTORY view**

**(Correct)**

- 

**Account Usage ACCESS\_REPORT view**

- 

**Information Schema ACCESS\_REPORT view**

- 

**Information Schema ACCESS\_HISTORY view**

**Explanation**

Access History in Snowflake refers to when the user query reads column data and when the SQL statement performs a data write operation, such as INSERT, UPDATE, and DELETE, along with variations of the COPY command, from the source data object to the target data object. **The user access history can be found by querying the Account Usage ACCESS\_HISTORY view.**

Question 100:

**Skipped**

**Which of these Snowflake Editions automatically store data in encrypted state?**

-

**Standard**

- 

**Enterprise**

- 

**All of the Snowflake Editions**

**(Correct)**

- 

**Business Critical**

- 

**Virtual Private Snowflake(VPS)**

#### **Explanation**

All of the Snowflake Editions (Standard, Enterprise, Business Critical, Virtual Private Snowflake) automatically store data in encrypted state.

---

Question 1:

#### **Skipped**

**ACCOUNTADMIN role cannot view the results for a query run by another user.(True/ False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

#### **Explanation**

**A user cannot view the result set from a query that another user executed.** This behavior is intentional. For security reasons, only the user who executed a query can access the query results. **This behavior is not connected to the Snowflake access control model for objects.** Even a user with the ACCOUNTADMIN role cannot view the results for a query run by another user.

Question 2:

**Skipped**

**Snowflake supports multi-factor authentication (i.e., MFA) to provide increased login security for users connecting to Snowflake. Which role is strongly recommended for using MFA?**

- 

**USERADMIN**

- 

**ACCOUNTADMIN**

**(Correct)**

- 

**SYSADMIN**

- 

**USERADMIN**

### **Explanation**

Snowflake strongly recommends that all users with the **ACCOUNTADMIN role** be required to use MFA.

Question 3:

**Skipped**

**Which of these SQL functions helps retrieve the URL for an external or internal named stage using the stage name as the input?**

- 

**GET\_STAGE\_LOCATION**

**(Correct)**

- 

**GET\_ABSOLUTE\_PATH**

- 

**GET\_PRESIGNED\_URL**

- 
- **BUILD\_STAGE\_FILE\_URI**
- 
- **GET\_RELATIVE\_PATH**
- 
- **BUILD\_SCOPED\_FILE\_URL**

### Explanation

**GET\_STAGE\_LOCATION** retrieves the URL for an external or internal named stage using the stage name as the input.

Question 4:

#### Skipped

**Snowflake supports the secured access of unstructured data files in cloud storage. What types of URLs are available to access files in cloud storage? (Select 3)**

- **Scoped URL**  
**(Correct)**
- **Descoped URL**
- **Unstructured URL**
- **Pre-signed URL**  
**(Correct)**
- **File URL**  
**(Correct)**

- 

## Signed URL

### Explanation

The following types of URLs are available to access files in cloud storage: **Scoped URL**: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. **File URL**: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. **Pre-signed URL**: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable.

Question 5:

### Skipped

**Which features of Snowflake provide Column-level security? (Select 2)**

- 

## Dynamic Data Masking

**(Correct)**

- 

## Internal Tokenization

- 

## External Tokenization

**(Correct)**

- 

## Column Masking

### Explanation

**Dynamic Data Masking** is a Column-level Security feature that uses masking policies to selectively mask plain-text data in table and view columns at query time.

**External Tokenization** enables accounts to tokenize data before loading it into Snowflake and detokenize the data at query runtime. Tokenization is the process of

removing sensitive data by replacing it with an undecipherable token. External Tokenization makes use of masking policies with external functions.

Question 6:

**Skipped**

**What is the expiration period of a Pre-signed URL?**

- 

**The URL expires when the persisted query result period ends**

- 

**Length of time specified in the expiration\_time argument**

**(Correct)**

- 

**It is Permanent**

**Explanation**

**The expiration period of Pre-Signed URL: Length of time specified in the expiration\_time argument.**

The expiration period of Scoped URL: The URL expires when the persisted query result period ends. The expiration period of the File URL: It is permanent.

Question 7:

**Skipped**

**What is the recommended compressed size of data files for optimal bulk data loads?**

- 

**100-250 GB**

- 

**10-50 GB**

- 

**100-250 MB**

**(Correct)**

- 

**10-50 MB**

### **Explanation**

To optimize the number of parallel operations for a load, Snowflake recommends aiming to produce data files roughly **100-250 MB (or larger) in size compressed**. The number of load operations that run in parallel cannot exceed the number of data files to be loaded.

Question 8:

**Skipped**

**How many network policies can be activated for a user at a time?**

- 

**1**

**(Correct)**

- 

**100**

- 

**16**

- 

**99**

- 

**31**

### **Explanation**

**Only a single network policy can be activated for each user at a time;** however, different network policies can be activated for different users for granular control. Associating a network policy with a user automatically removes the currently-associated network policy (if any)

Question 9:

**Skipped**

**Both non-materialized and materialized views can be defined as secure. (True / False)**

- FALSE
- TRUE

**(Correct)**

**Explanation**

**Yes, both non-materialized and materialized views can be defined as secure.**

Question 10:

**Skipped**

**How long does Snowflake keep batch load history (from Stage) using COPY statement?**

- 1 day
- 31 days
- 30 days
- 64 days
- 14 days

**(Correct)**

**Explanation**

**Snowflake keeps the batch load history for 64 days.**

Question 11:

### Skipped

Cloning a table replicates the source table's structure, data, load history, and certain other properties (e.g., STAGE FILE FORMAT). (True/False)

- 
- **TRUE**
- **FALSE**

**(Correct)**

### Explanation

Cloning a table replicates the source table's structure, data, and certain other properties (e.g., STAGE FILE FORMAT). A cloned table **does not include the load history** of the source table. One consequence is that data files loaded into a source table can be loaded again into its clones.

Question 12:

### Skipped

How can you view the data storage across your entire Snowflake account?

(Select 2)

- 
- **Using Classic Web Interface: Click on Account > Resource Monitors > Average Storage Used**
- 
- **Using Snowsight: Select Data > Usage > Storage**
- 

**Using Classic Web Interface: Click on Account > Billing & Usage > Average Storage Used**

**(Correct)**

- 
- **Using Snowsight: Select Admin > Usage > Storage**

**(Correct)**

### Explanation

Suppose you have been assigned the ACCOUNTADMIN role (i.e., you serve as the top-level administrator for your Snowflake account). In that case, you can use Snowsight or the classic web interface to view data storage across your entire account:

- **Using Snowsight:** Select Admin > Usage > Storage
- **Using Classic Web Interface:** Click on Account > Billing & Usage > Average Storage Used

Question 13:

### Skipped

Snowflake stores data into its

- **internal optimized, uncompressed, row format**
  - **internal optimized, compressed, columnar format**
- (Correct)**
- **internal optimized, compressed, row format**
  - **internal optimized, uncompressed, columnar format**

### Explanation

When data is loaded into Snowflake, Snowflake reorganizes that data into its **internal optimized, compressed columnar format**. Snowflake stores this optimized data in cloud storage.

Question 14:

### Skipped

Snowflake supports multiple ways of connecting to the service. (Select 3)

- 
- **Only ODBC**
- **A web-based user interface**

**(Correct)**

- 

**ODBC and JDBC drivers**

**(Correct)**

- 

**Command line clients (e.g. SnowSQL)**

**(Correct)**

- 

**Only JDBC**

### **Explanation**

Snowflake supports the following multiple ways of connecting to its service:

**A web-based user interface** from which all aspects of managing and using Snowflake can be accessed.

**Command line clients (e.g. SnowSQL)** can also access all aspects of managing and using Snowflake.

**ODBC and JDBC drivers** that can be used by other applications (e.g. Tableau) to connect to Snowflake.

**Native connectors (e.g. Python, Spark)** that can be used to develop applications for connecting to Snowflake.

**Third-party connectors** that can be used to connect applications such as ETL tools (e.g. Informatica) and BI tools (e.g. ThoughtSpot) to Snowflake.

Question 15:

**Skipped**

**Network policies allow restricting access to your account based on\_\_\_\_ (Select all that apply)**

- 

**IP address**

**(Correct)**

- 

**User Operating System Type (example - MAC, Windows)**

- 

**CIDR Notation based IP ranges**

**(Correct)**

### **Explanation**

Network policies allow restricting access to your account based on user IP address. Effectively, a network policy enables you to create an IP allowed list, as well as an IP blocked list, if desired.

Question 16:

**Skipped**

**What is an "object" in semi-structured data?**

- 

**Collection of values in an array**

- 

**Collection of key-value pairs**

**(Correct)**

- 

**Collection of semi-structure data files**

### **Explanation**

OBJECT is also called a “dictionary”, “hash”, or “map” in many languages. This contains key-value pairs.

Question 17:

**Skipped**

**Snowflake data providers can share data from one database per share. Data from multiple databases can not be shared with a share. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

#### **Explanation**

**Snowflake data providers can share data that resides in different databases by using secure views.** A secure view can reference objects such as schemas, tables, and other views from one or more databases, as long as these databases belong to the same account.

Question 18:

**Skipped**

**How does Snowflake store semi-structured data, such as JSON, Avro, Parquet, ORC, and XML?**

- 

**Stores as VARCHAR data type**

- 

**Stores as VARIANT data type**

**(Correct)**

- 

**Stores as JSON data type**

- 

**Stores as FLATTEN data type**

#### **Explanation**

Snowflake stores semi-structured data, such as JSON, Avro, Parquet, ORC, and XML, as **VARIANT data type**. A VARIANT can store a value of any other type, including OBJECT and ARRAY.

Question 19:

**Skipped**

UDF does not support SQL DDL / DML? (True/Fales)

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### **Explanation**

UDF does not support SQL DDL / DML. That means you can select from a table, but you can't create or modify tables inside of a UDF.

Question 20:

**Skipped**

**Which of the Snowflake editions provides HIPPA Support feature? (Select 2)**

- 

**Virtual Private Snowflake (VPS)**

**(Correct)**

- 

**Standard**

- 

**Business Critical**

**(Correct)**

- 

**Enterprise**

- 

**All of the Snowflake Editions**

### **Explanation**

**Business Critical and Virtual Private Snowflake (VPS) editions provide HIPPA support.**

Question 21:

**Skipped**

**You have a table t1 with a column j that gets populated by a sequence s1. s1 is defined to start from 1 and with an increment of 1. create or replace sequence s1 start = 1 increment = 1 ; create or replace table t1 ( i int, j int default**

`s1.nextval ); You inserted 3 records in table t1: insert into t1 values (1,s1.nextval), (2,s1.nextval), (3,s1.nextval); After that insert statement, you altered the sequence s1 to set the increment to -4: alter sequence s1 set increment = -4; You again inserted 2 records in table t1: insert into t1 values (4,s1.nextval), (5,s1.nextval); What would be the result of the following query? select j from t1 where i = 4;`

- 0
- -1
- 5
- 3
- 4

**(Correct)**

#### **Explanation**

**ALTER SEQUENCE command takes effect after the second use of the sequence after executing the ALTER SEQUENCE command.**

So, if you fetch row where  $i = 5$ , you will find  $j = 0$  [row 4 value of  $j$  i.e.,  $4 + (-4) = 0$ ]

Question 22:

**Skipped**

**Tasks require compute resources to execute code. Either Snowflake-managed or User-managed compute models can be chosen for individual tasks. (True / False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

#### **Explanation**

**True, Tasks require compute resources to execute SQL code.** Either of the following compute models can be chosen for individual tasks: Snowflake-managed (i.e. serverless compute model), User-managed (i.e. virtual warehouse)

Question 23:

**Skipped**

**Will these queries be considered as same to get the benefit of the Query Result cache? Query 1 : SELECT \* FROM t1; Query 2 : select \* FROM t1;**

- 

**YES**

- 

**NO**

**(Correct)**

#### **Explanation**

These queries will not be considered as same to get the benefit of the Query Result cache. First Query has 'SELECT' is upper case, and the Second query has 'select' in lower case.

Question 24:

**Skipped**

How much-uncompressed data does a micro-partition contain in Snowflake?

- 

**Between 1 MB to 100 MB**

- 

**Between 1 GB to 10 GB**

- 

**Between 5 MB to 50 MB**

- 

**Between 50 MB to 500 MB**

**(Correct)**

### **Explanation**

Each micro-partition contains **between 50 MB and 500 MB of uncompressed data** (Note that the actual size in Snowflake is smaller because data is always stored compressed.). Groups of rows in tables are mapped into individual micro-partitions, organized in a columnar fashion. This size is between 50 MB and 500 MB, and **the structure allows for extremely granular pruning of very large tables**, which can be comprised of millions, or even hundreds of millions, of micro-partitions. It enables extremely efficient DML and fine-grained pruning for faster queries.

Question 25:

**Skipped**

**Which role is inherited to every other role in the account?**

- 

**SECURITYADMIN**

- 

**USERADMIN**

- 

**ACCOUNTADMIN**

- 

**SYSADMIN**

- 

**PUBLIC**

**(Correct)**

### **Explanation**

The **PUBLIC role** is Pseudo-role, which is automatically granted to every user and every role in your account. This role is typically used in cases where explicit access control is not needed, and all users are viewed as equal with regard to their access rights.

Question 26:

**Skipped**

**Which of these are unstructured data? (Select 2)**

- 

**XML**

- 

**Relational Data**

- 

**JSON**

- 

**Images**

**(Correct)**

- 

**Videos**

**(Correct)**

**Explanation**

**Unstructured data is information that does not fit into a predefined data model or schema.** Typically text-heavy, such as form responses and social media conversations, unstructured data also encompasses images, video, and audio. Industry-specific file types such as VCF (genomics), KDF (semiconductors), or HDF5 (aeronautics) are included in this category.

Question 27:

**Skipped**

**What is the minimum billing charge for provisioning compute resources?**

- 

**1 second**

- 

**120 seconds**

- 30 seconds
- 60 seconds

(Correct)

#### Explanation

The minimum billing charge for provisioning compute resources is 1 minute (i.e. 60 seconds). There is no benefit to stopping a warehouse before the first 60-second period is over because the credits have already been billed for that period.

Question 28:

Skipped

#### When should we use HyperLogLog?

- Whenever the input is potentially large, and an approximate result is acceptable.
- Whenever the input is potentially small, and an approximate result is not acceptable.
- Whenever the input is potentially small, and an approximate result is acceptable.
- Whenever the input is potentially large, and an approximate result is not acceptable.

#### Explanation

Snowflake recommends using HyperLogLog whenever the input is potentially large, and an approximate result is acceptable.

Question 29:

Skipped

**Which of the following Data Types are supported by Snowflake? (Select all that apply)**

- 

**BOOL**

- 

**FLOAT**

**(Correct)**

- 

**VARCHAR**

**(Correct)**

- 

**CHAR**

**(Correct)**

- 

**INTEGER**

**(Correct)**

- 

**NUMERIC**

**(Correct)**

### **Explanation**

All of these data types are supported by Snowflake except BOOL. BOOLEAN is the correct data type.

Question 30:

**Skipped**

**Suppose we have a table t1. We drop the table t1 and then create a new table t1 again. What will happen if we execute the UNDROP command to restore dropped t1 table now?**

- The dropped table t1 will be restored with a new arbitrary name set by Snowflake
- The dropped table t1 will be restored with name t1
- UNDROP command will fail  
**(Correct)**

#### Explanation

If an object with the same name already exists, UNDROP fails. We must rename the existing object, which then enables us to restore the previous version of the object.

Question 31:

Skipped

While choosing the clustering key, what should we consider? (Select 3)

- Columns which are more often used in where clause  
**(Correct)**
- Columns which are more less used in join conditions
- Columns which are more often used in join conditions  
**(Correct)**
- Columns which are less often used in where clause
- Ordering the columns from highest cardinality to lowest cardinality

- 

**Ordering the columns from lowest cardinality to highest cardinality**

**(Correct)**

### **Explanation**

Best Practices for choosing clustering key:

1. Columns which are more often used in where clause
2. Columns that are more often used in join conditions
3. Order you specify the clustering key is important.

As a general rule, Snowflake recommends ordering the lowest to highest cardinality columns.

Question 32:

**Skipped**

**During Warehouse provisioning, if any of the compute resources fail to provision, then Snowflake kills the entire warehouse provisioning and tries to provision a new warehouse of the same requested size. (True/False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### **Explanation**

Snowflake does not begin executing SQL statements submitted to a warehouse until all of the compute resources for the warehouse are successfully provisioned unless any of the resources fail to provision:

**If any of the compute resources for the warehouse fail to provision during start-up, Snowflake attempts to repair the failed resources.** During the repair process, the warehouse starts processing SQL statements once 50% or more of the requested compute resources are successfully provisioned.

Question 33:

**Skipped**

**Which command is used to create a security integration to enable an HTTP client that supports OAuth to redirect users to an authorization page and generate access tokens for access to the REST API endpoint?**

- 

**CREATE SECURITY INTEGRATION**

**(Correct)**

- 

**CREATE INTEGRATION**

- 

**CREATE SECURITY API**

- 

**CREATE OAUTH INTEGRATION**

#### **Explanation**

**CREATE SECURITY INTEGRATION** command is used to create a security integration that supports OAuth to redirect users to an authorization page and generate access tokens for access to the REST API endpoint.

Question 34:

**Skipped**

Select the type of function that returns multiple rows for each individual input.

- 

**System Function**

- 

**Aggregate Function**

- 

**User-Defined Function**

- 

**Window Function**

- 

### Scalar Function

- 

### Table Function

(Correct)

#### Explanation

**A table function returns a set of rows for each input row.** The returned set can contain zero, one, or more rows. Each row can contain one or more columns. Table functions are sometimes called “tabular functions”.

Question 35:

#### Skipped

**User-managed tasks are recommended when you cannot fully utilize a warehouse because only a few tasks run concurrently or they run to completion quickly (in less than 1 minute). (True / False)**

- 

**FALSE**

(Correct)

- 

**TRUE**

#### Explanation

**Serverless Tasks** is recommended when you cannot fully utilize a warehouse because too few tasks run concurrently or they run to completion quickly (in less than 1 minute).

**User-managed Tasks** is recommended when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources.

Question 36:

#### Skipped

**Which of these system-defined roles encapsulates the SYSADMIN and SECURITYADMIN roles?**

-

**SYSADMIN**

- 

**ACCOUNTADMIN**

**(Correct)**

- 

**SECURITYADMIN**

- 

**USERADMIN**

- 

**ORGADMIN**

#### **Explanation**

**ACCOUNTADMIN** role encapsulates the **SYSADMIN** and **SECURITYADMIN** system-defined roles.

It is the top-level role in the system and should be granted only to a limited/controlled number of users in your account.

Question 37:

**Skipped**

**Python UDFs and tabular Python UDFs can read and process unstructured data in staged files using SnowflakeFile class. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

#### **Explanation**

**Java UDFs and tabular Java UDFs** can read and process unstructured data in staged files using either the **SnowflakeFile** class or the **InputStream** class in the UDF code.

Question 38:

Skipped

Which function can be used in combination with COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file?

- 

**VARIANT\_CONSTRUCT**

- 

**COPY\_CONSTRUCT**

- 

**OBJECT\_CONSTRUCT**

**(Correct)**

- 

**UNLOAD\_CONSTRUCT**

### Explanation

The **OBJECT\_CONSTRUCT function** can be used in combination with the COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file.

Question 39:

Skipped

What is the best practice after creating a custom role in a Snowflake account?

- 

**Grant ownership of the role to itself so a member of the role can control access to the role.**

- 

**Grant the custom role to the SYSADMIN role so administrators can manage all objects in the account.**

**(Correct)**

-

**Grant the role to the USERADMIN role so users can be added to the role.**

- 

**Grant PUBLIC to the role so all database objects owned by PUBLIC will be available to the new role.**

### **Explanation**

The custom role gets mainly created for specific access to specific objects. **As per best practice, grant ownership to SYSADMIN or a child role under SYSADMIN so that SYSADMIN can control access to the objects.**

Question 40:

### **Skipped**

**LIST command returns a list of files that have been staged. Which of these stages supports the LIST command?**

- 

**Named external stage.**

- 

**Stage for a specified table.**

- 

**Named internal stage.**

- 

**Stage for the current user.**

- 

**All of these**

**(Correct)**

### **Explanation**

**LIST command returns a list of files staged from all of these specified snowflake stages.**

Question 41:

### **Skipped**

**Which of these Snowflake Connectors are available? (Select all that apply)**

- **Snowflake Connector for ODBC**
- **Snowflake Connector for Kafka**  
**(Correct)**
- **Snowflake Connector for JDBC**
- **Snowflake Connector of Python**  
**(Correct)**
- **Snowflake Connector for Spark**  
**(Correct)**

### Explanation

**ODBC and JDBC are drivers. Connectors available for Snowflake are Python, Kafka, and Spark.** Snowflake also provides several drivers like ODBC, JDBC, Node.js, Go,.Net, and PHP PDO. The Snowflake SQL API is a REST API that you can use to access and update data in a Snowflake database.

Question 42:

**Skipped**

**While transforming Semi-structure data, what argument would you need to set with FLATTEN function to omit the output of the input rows that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries?**

- **OUTER => FALSE**  
**(Correct)**

**OUTER => TRUE**

- 

**MODE => OBJECT**

- 

**RECURSIVE => TRUE**

- 

**RECURSIVE => FALSE**

#### **Explanation**

The **OUTER => FALSE argument with FLATTEN omits the output of the input rows that cannot be expanded**, either because they cannot be accessed in the path or because they have zero fields or entries. The **OUTER => TRUE argument with FLATTEN generates exactly one row for zero-row expansions** (with NULL in the KEY, INDEX, and VALUE columns).

RECURSIVE is used to instruct if only the element referenced by PATH is expanded or expansion is performed for all sub-elements recursively MODE Specifies whether only objects, arrays, or both should be flattened.

Question 43:

#### **Skipped**

**The data retention period for a database, schema, or table can not be changed once ACCOUNTADMIN sets it at the account level. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

#### **Explanation**

The data retention period for a database, schema, or table can be changed at any time. DATA\_RETENTION\_TIME\_IN\_DAYS parameter can be used to explicitly override

the default when creating a database, schema, and individual table. For example:

```
CREATE TABLE t1 (c1 int) DATA_RETENTION_IN_DAYS=90;
```

Question 44:

**Skipped**

**Which of these functions helps generate the Pre-signed URL to access the unstructured data file?**

- 

**GET\_STAGE\_LOCATION**

- 

**GET\_RELATIVE\_PATH**

- 

**BUILD\_STAGE\_FILE\_URI**

- 

**GET\_ABSOLUTE\_PATH**

- 

**BUILD\_SCOPED\_FILE\_URL**

- 

**GET\_PRESIGNED\_URL**

**(Correct)**

### **Explanation**

**GET\_PRESIGNED\_URL generates the pre-signed URL to a staged file using the stage name and relative file path as inputs.** Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable.

Question 45:

**Skipped**

**The snowflake data warehouse is not built on an existing database or “big data” software platform like Hadoop.(True/False)**

-

**TRUE**

**(Correct)**

- 

**FALSE**

**Explanation**

**Snowflake is a 100% cloud-native data platform.**

Question 46:

**Skipped**

**UDF runs with either the caller's or the owner's rights. (True / False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

**Explanation**

**UDF only runs as the function owner. A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. A caller's rights stored procedure runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query.**

**An owner's rights stored procedure** runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table.

At the time that the stored procedure is created, the creator specifies whether the procedure runs with the owner's rights or the caller's rights. The default is owner's rights.

Question 47:

**Skipped**

If **DATA\_RETENTION\_TIME\_IN\_DAYS** is set to a value of 0, and **MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS** is set higher at the account level and is greater than 0, which value (0 or higher) setting takes precedence?

- - 0 (set in DATA\_RETENTION\_TIME\_IN\_DAYS)**
  - 
  - Higher value (set in MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS)**
- (Correct)**

#### Explanation

If **DATA\_RETENTION\_TIME\_IN\_DAYS** is set to a value of 0, and **MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS** is set at the account level and is greater than 0, the higher value setting takes precedence. The data retention period for an object is determined by **MAX(DATA\_RETENTION\_TIME\_IN\_DAYS, MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS)**.

Question 48:

**Skipped**

**Snowflake architecture is**

- 
- Hybrid of Shared-disk and Shared-nothing database architectures**
- (Correct)**
- 
- Shared-disk architecture**
- 
- Shared-nothing architecture**
- 
- None of these**

#### Explanation

Snowflake's architecture **is a hybrid of traditional shared-disk and shared-nothing database architectures**. Like shared-disk architectures, Snowflake uses a central

data repository for persisted data accessible from all compute nodes in the platform. But similar to shared-nothing architectures, Snowflake processes queries using MPP (massively parallel processing) compute clusters where each node in the cluster stores a portion of the entire data set locally. This approach offers the data management simplicity of a shared-disk architecture but with the performance and scale-out benefits of a shared-nothing architecture. It is also termed as Multi-Cluster Shared Data Architecture.

Question 49:

**Skipped**

**Which of these sampling method keywords are used to specify which method to use? (Select 2)**

- 

**SYSTEM | ROW**

- 

**SYSTEM | BLOCK**

**(Correct)**

- 

**BERNOULLI | ROW**

**(Correct)**

- 

**BERNOULLI | BLOCK**

**Explanation**

**BERNOULLI | ROW and SYSTEM | BLOCK are used to specify the sampling method in SELECT query.** **BERNOULLI (or ROW):** Includes each row with a <probability> of p/100. Similar to flipping a weighted coin for each row. **SYSTEM (or BLOCK):** Includes each block of rows with a <probability> of p/100. Similar to flipping a weighted coin for each block of rows. This method does not support fixed-size sampling. Sampling method is optional. If no method is specified, the default is BERNoulli. **Example : select \* from t1 tablesample bernoulli (25); This query will return a sample of a table in which each row has a 25% probability of being included in the sample**

Question 50:

**Skipped**

**Which of the Snowflake shared view can be used to query the Snowflake Query History?**

- **QUERY\_HISTORY view in INFORMATION\_SCHEMA**

- **QUERY\_HISTORY view in ACCOUNT\_USAGE Schema**

**(Correct)**

- **QUERY\_HISTORY\_VIEW view in INFORMATION\_USAGE Schema**

- **QUERY\_HISTORY\_VIEW view in ACCOUNT\_USAGE Schema**

### **Explanation**

This Account Usage view can be used to query Snowflake query history by various dimensions (time range, session, user, warehouse, etc.) **within the last 365 days (1 year)**.

Question 51:

**Skipped**

**The best use of Snowpipe is to load large volumes of data and incrementally make them available for analysis. (True/False)**

- **TRUE**
- **FALSE**

**(Correct)**

### **Explanation**

Snowpipe is **designed to load small volumes of data** (i.e., micro-batches) and incrementally make them available for analysis.

Question 52:

**Skipped**

**Which of these columns gets appended on creating a stream on a table? (Select 3)**

- 

**METADATA\$ISUPDATE**

**(Correct)**

- 

**METADATA\$ISDELETE**

- 

**METADATA\$ACTION**

**(Correct)**

- 

**METADATA\$ROW\_ID**

**(Correct)**

- 

**METADATA\$ISINSERT**

#### **Explanation**

**Adding a stream to a table appends three metadata columns:  
METADATA\$ACTION, METADATA\$ISUPDATE, METADATA\$ROW\_ID.**

These columns track the CDC records and their type: appends, deletes, or both (updates = inserts + deletes).

- METADATA\$ACTION** - Indicates the DML operation (INSERT, DELETE) recorded.

- METADATA\$ISUPDATE** - Indicates whether the operation was part of an UPDATE statement.
-

- METADATA\$ROW\_ID - Specifies the unique and immutable ID for the row, which can be used to track changes to specific rows over time.

Question 53:

**Skipped**

**Which of these configurations will set up a warehouse in maximized mode?**

- **Minimum Clusters = 1 and Maximum Clusters = 10**
- **Minimum Clusters = 1 and Maximum Clusters = 1**
- **Minimum Clusters = 9 and Maximum Clusters = 10**
- **Minimum Clusters = 6 and Maximum Clusters = 6**

**(Correct)**

**Explanation**

**Maximized mode is enabled by specifying the same value for both the maximum and a minimum number of clusters (note that the specified value must be larger than 1).** In this mode, when the warehouse is started, Snowflake starts all the clusters so that maximum resources are available while the warehouse is running. This mode is effective for statically controlling the available compute resources, particularly if you have large numbers of concurrent user sessions and/or queries and the numbers do not fluctuate significantly.

Question 54:

**Skipped**

**A row access policy contains an expression that can specify Snowflake database objects (e.g., table or view) and use functions to determine which rows should be visible in a given context. Which of these functions are used in determining row access policy? (Select 2)**

- 
- Metadata Functions**
-

## Context Functions

(Correct)

- 

## Conversion Functions

- 

## Conditional Expression Functions

(Correct)

## Explanation

A row access policy uses **Conditional Expression Functions and Context Functions** to determine which rows should be visible in a given context.

Context Functions such as CURRENT\_USER(), CURRENT\_ROLE(), and CURRENT\_ACCOUNT(), act as dynamic filters and are commonly used with secure views to limit row access in a table.

Question 55:

**Skipped**

A Pre-Signed URL is ideal for

- 

**business intelligence applications or reporting tools that need to display the unstructured file contents**

(Correct)

- 

**None of these**

- 

**custom applications that require access to unstructured data files**

- 

**use in custom applications, providing unstructured data to other accounts via a share**

**Explanation**

**Pre-signed URL:** Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable. **Ideal for business intelligence applications or reporting tools that need to display unstructured file contents.**

**Scoped URL:** Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. **Ideal for use in custom applications, providing unstructured data to other accounts via a share, or for downloading and ad hoc analysis of unstructured data via Snowsight.**

**File URL:** URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. **Ideal for custom applications that require access to unstructured data files.**

Question 56:

Skipped

Which command will help you to view the current permissions on a Schema?

- 

**SHOW GRANTS OF SCHEMA <database\_name>.<schema\_name>;**

- 

**SHOW ALL PREIVILEGE ON SCHEMA <database\_name>.<schema\_name>;**

- 

**SHOW ROLES ON SCHEMA <database\_name>.<schema\_name>;**

- 

**SHOW GRANTS ON SCHEMA <database\_name>.<schema\_name>;**

**(Correct)**

**Explanation**

To view the current set of privileges granted on an object, you can execute the **SHOW GRANTS** command. To view the current permissions on a schema,

execute the following command: SHOW GRANTS ON SCHEMA <database\_name>.<schema\_name>;

Question 57:

**Skipped**

**Which of these Snowflake tasks can be performed by Time Travel? (Select 3)**

- 

**Create clones of entire tables, schemas, and databases at or before specific points in the past**

**(Correct)**

- 

**Restore tables, schemas, and databases that have been dropped.**

**(Correct)**

- 

**Share the restored data objects over a specified period of time**

- 

**Query data in the past that has since been updated or deleted**

**(Correct)**

### **Explanation**

Using Time Travel, you can perform the following actions within a defined period:

- Query data in the past that has since been updated or deleted.
- Create clones of entire tables, schemas, and databases at or before specific points in the past.
- Restore tables, schemas, and databases that have been dropped.

Question 58:

**Skipped**

**Loading into Snowflake from a local file system is a straightforward affair.**

**Which command is used to grab files from the local system, compress them and encrypt them, and then it copies them to Snowflake?**

- 

**PUT command**

**(Correct)**

- 

**COPY INTO <table>**

- 

**MOVE command**

- 

**GET command**

### **Explanation**

**Using the PUT command in SnowSQL.** It grabs the file or files, compresses them, encrypts them, and then copies them up into the stage you chose. Once in the stage, you can use a COPY INTO command to load the data from the stage into Snowflake tables.

Question 59:

**Skipped**

**The closer the ratio of scanned micro-partitions and columnar data is to the ratio of actual data selected, the more efficient is the pruning performed on the table. (TRUE/FALSE)**

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### **Explanation**

Snowflake uses columnar scanning of partitions so that an entire partition is not scanned if a query only filters by one column. The closer the ratio of scanned micro-partitions and columnar data is to the ratio of actual data selected, the more efficient is the pruning performed on the table

Question 60:

**Skipped**

**Which of these objects do not clone? (Select 2)**

-

## Internal (Snowflake) stages

(Correct)

- 

## Databases

- 

## External Table

(Correct)

- 

## Schemas

### Explanation

Databases and Schemas can be cloned. External Table and Internal (Snowflake) stages do not get cloned.

Question 61:

Skipped

Which privilege is required to change a warehouse's state (stop, start, suspend, resume)?

- 

## OPERATE

(Correct)

- 

## MONITOR

- 

## USAGE

- 

## MODIFY

### Explanation

Virtual Warehouse Privileges:

**OPERATE** - Enables changing the state of a warehouse (**stop, start, suspend, resume**). In addition, enables viewing current and past queries executed on a warehouse and aborting any executing queries.

**MODIFY** - Enables altering any properties of a warehouse, including changing its size. Required assigning a warehouse to a resource monitor. Note that only the ACCOUNTADMIN role can assign warehouses to resource monitors.

**MONITOR** - Enables viewing of current and past queries executed on a warehouse as well as usage statistics on that warehouse.

**USAGE** - Enables using a virtual warehouse and, as a result, executing queries on the warehouse. If the warehouse is configured to auto-resume when a SQL statement (e.g. query) is submitted to it, the warehouse resumes automatically and executes the statement.

**OWNERSHIP** - Grants full control over a warehouse. Only a single role can hold this privilege on a specific object at a time.

**ALL** [ PRIVILEGES ] - Grants all privileges, except OWNERSHIP, on the warehouse.

Question 62:

**Skipped**

**Which of the following is the correct hierarchy for the Snowflake objects?**

- 
- 

**ORGANIZATION > ACCOUNT > DATABASE > SCHEMA > TABLE > STAGE**

- 
- 

**ORGANIZATION > ACCOUNT > ROLE > USER > DATABASE > SCHEMA > STAGE > TABLE**

- 

**ORGANIZATION > ACCOUNT > DATABASE > SCHEMA > TABLE**

**(Correct)**

- 

**ACCOUNT > ORGANIZATION > ROLE > USER > DATABASE > SCHEMA > TABLE**

### **Explanation**

The top-most container is the customer organization. Securable objects such as tables, views, functions, and stages are contained in a schema object, which are in turn contained in a database. All databases for your Snowflake account are contained in the account object. USER, ROLE, DATABASE, WAREHOUSE are at same level and contained in a Snowflake Account Object.

Question 63:

**Skipped**

**Which options for selecting staged data files are supported by COPY command in Snowflake? (Select that all apply)**

- 

**By path (internal stages) / prefix (Amazon S3 bucket)**

**(Correct)**

- 

**Specifying a list of specific files to load**

**(Correct)**

- 

**Using pattern matching to identify specific files by pattern**

**(Correct)**

### **Explanation**

**All of these are supported.** This allows you to execute concurrent COPY statements that match a subset of files, taking advantage of parallel operations.

Question 64:

**Skipped**

**Suppose files downloaded from an internal stage are corrupted. What should be verified with the stage creator to determine why the downloaded file is corrupted?**

- 

**Verify if ENCRYPTION = (TYPE = 'SNOWFLAKE\_SSE') set for the stage**

**(Correct)**

- 

**Verify if ENCRYPTION = (TYPE = 'SNOWFLAKE\_FULL') set for the stage**

### **Explanation**

If files downloaded from an internal stage are corrupted, verify with the stage creator that **ENCRYPTION = (TYPE = 'SNOWFLAKE\_SSE')** is set for the stage.

Question 65:

**Skipped**

**What is the maximum data retention period for permanent databases, schemas, and tables for Snowflake Enterprise Edition (and higher)?**

- 

**0 days**

- 

**90 days**

**(Correct)**

- 

**30 days**

- 

**1 day**

### **Explanation**

**For Snowflake Enterprise Edition (and higher):**

- **For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days.**

- For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also true for temporary tables.

Question 66:

**Skipped**

**If you have privileges to view queries executed by another user, the Query Detail page displays the details for the query and the actual query result. (TRUE/FALSE)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

**Explanation**

If you have privileges to view queries executed by another user, the Query Detail page displays the details for the query, but, for data privacy reasons, **the page does not display the actual query result.**

Question 67:

**Skipped**

**The search optimization service speeds only equality searches. (True/False)**

- 

**TRUE**

**(Correct)**

- 

**FALSE**

**Explanation**

**True, The search optimization service speeds only equality searches.** However, this applies to all the columns of supported types in a table with search optimization enabled.

Question 68:

### Skipped

When loading data using **COPY INTO <table>** command, it is a must that the data files have the same number and order of columns as your target table.  
(True/False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### Explanation

There is no requirement for your data files to have the same number and ordering of columns as your target table.

Question 69:

### Skipped

What is a key benefit of scaling up a warehouse?

- 

**Scaling up improves concurrency.**

- 

**Scaling up improves performance.**

**(Correct)**

### Explanation

**Resizing a warehouse generally improves query performance**, particularly for larger, more complex queries. It can also help reduce the queuing that occurs if a warehouse does not have enough compute resources to process all the queries that are submitted concurrently. **Note that warehouse resizing is not intended for handling concurrency issues**; instead, use additional warehouses to handle the workload or use a multi-cluster warehouse.

Question 70:

### Skipped

How can a directory table metadata be refreshed automatically and efficiently to synchronize the metadata with the latest associated files in the external stage and path?

- 
- Using Cloud event notification service**  
**(Correct)**
- 
- Using both Tasks and Stream**
  - 
  - It is a manual process and can't be automatically refreshed**
  -
- 
- Using Stream**
- 
- Using Tasks**

### Explanation

The metadata for a directory table can be refreshed automatically **using the event notification service** for your cloud storage service. The refresh operation synchronizes the metadata with the latest set of associated files in the external stage and path, i.e.:

- New files in the path are added to the table metadata.
- Changes to files in the path are updated in the table metadata.
- Files no longer in the path are removed from the table metadata.

Question 71:

**Skipped**

**Which methods can be used to check the status of a COPY INTO command?**

- 
- Use the Query History tab in the Snowflake user interface.**  
**(Correct)**
- 
- Write a SQL query against the ACCOUNT\_USAGE LOAD\_HISTORY view.**
-

**Write a SQL query against the INFORMATION\_SCHEMA.LOAD\_HISTORY view.**

**(Correct)**

- 

**Use the resource monitor.**

### **Explanation**

The status of COPY INTO command can be checked from the Resource Monitors tab in the Snowflake user interface, as well as by querying the INFORMATION\_SCHEMA.LOAD\_HISTORY view.

Question 72:

**Skipped**

**What is the default compression method used by Snowflake while unloading data?**

- 

**bzip2**

- 

**Brotli**

- 

**gzip**

**(Correct)**

- 

**Zstandard**

### **Explanation**

By default, all unloaded data files are compressed using gzip, unless compression is explicitly disabled or one of the other supported compression methods is explicitly specified.

Question 73:

**Skipped**

**Which is not the DDL (Data Definition Language) command?**

-

**ALTER**

- 

**UNDROP**

- 

**TRUNCATE**

**(Correct)**

- 

**CREATE**

- 

**DROP**

- 

**SHOW SHARES**

#### **Explanation**

**TRUNCATE** is DML (Data Manipulation Language) command.

Question 74:

**Skipped**

**Which of these is not a valid authentication method supported by Snowflake?**

- 

**Username and one-time generated pin**

**(Correct)**

- 

**SCIM (System for Cross-domain Identity Management specification)**

- 

**Federated Authentication (SAM 2.0)**

-

## Muti-factor authentication

- 
- 
- 

## Username and password

- 
- 
- 

## Oauth, Key Pair

### Explanation

Snowflake supports the following authentication methods: Username and password, MFA, SAML, and other authentication methods such as OAuth, Key-pair, and SCIM.

Question 75:

### Skipped

**David wants to load a JSON file using the COPY INTO <table> command. He found that there are null values in the data for missing values and have no other special meaning. What file format option would you recommend him to use with COPY INTO <table> command to handle the JSON null values?**

- 
- 
- 

**David should use STRIP\_OUTER\_ELEMENT = TRUE**

- 
- 
- 

**David should use STRIP\_NULL\_VALUES = TRUE**

**(Correct)**

- 
- 
- 

**David should use STRIP\_NULL\_VALUES = FALSE**

- 
- 
- 

**David should use STRIP\_OUTER\_ELEMENT = FALSE**

- 
- 
- 

**David should use REPLACE\_INVALID\_CHARACTERS = TRUE**

### Explanation

In a VARIANT column, NULL values are stored as a string containing the word "null," not the SQL NULL value. If the "null" values in your JSON documents indicate missing

values and have no other special meaning, you should **recommend setting the file format option STRIP\_NULL\_VALUES to TRUE for the COPY INTO <table> command when loading the JSON files**. Retaining the “null” values often wastes storage and slows query processing.

Question 76:

**Skipped**

All files stored in internal stages for data loading and unloading are automatically encrypted using AES-256 strong encryption. (True/False)

- 

**TRUE**

**(Correct)**

- 

**FALSE**

**Explanation**

All files stored in internal stages for data loading and unloading automatically encrypted using AES-256 strong encryption.

Question 77:

**Skipped**

**There are multiple ways you can trigger the loading of files from the stage into Snowpipe. Select two ways generally used to trigger the loading with Snowpipe.**

- 

**By executing START PIPE <pipe\_name>**

- 

**Calling Snowpipe SOAP endpoints**

- 

**Automating Snowpipe using cloud messaging (notification)**

**(Correct)**

- 

**Calling Snowpipe REST endpoints**

**(Correct)**

### Explanation

Automating Snowpipe using cloud messaging (notification) and Calling Snowpipe REST endpoints are the mostly used methods for triggering loading with Snowpipe.

Question 78:

**Skipped**

**What of the following is the default character set for delimited files (CSV, TSV, etc.)?**

- 

**UTF-16**

- 

**UTF-16BE**

- 

**UTF-8**

**(Correct)**

- 

**UTF-32LE**

- 

**UTF-32**

### Explanation

**UTF-8 is the default character set for delimited files (CSV, TSV, etc.).** Snowflake also supports all others in the provided options, but you must explicitly specify the encoding to use for loading.

Question 79:

**Skipped**

**What types of accounts are involved in data sharing? (Select 3)**

- 

**Shared Accounts**

-

## Data Providers

(Correct)

- 

## Reader Accounts

(Correct)

- 

## Data Consumers

(Correct)

- 

## Data Publishers

### Explanation

There are three types of accounts involved in data sharing.

**Data Providers:** Share data with others

**Data Consumers:** Accesses shared data with their own Snowflake account.

**Reader Accounts:** Query data using compute from the data provider's account. Reader Accounts are what you can use to share data with somebody who does not already have a Snowflake account.

Question 80:

Skipped

Which command can be used to suspend Automatic Clustering for a table?

- 

**ALTER TABLE**

(Correct)

- 

**DROP CLUSTERING**

-

## SUSPEND RECLUSTER

- 
- 

## STOP TABLE

### Explanation

Example - **ALTER TABLE** EMPLOYEE SUSPEND RECLUSTER; please note, SUSPEND RECLUSTER is a clause here, not a command.

Question 81:

### Skipped

## Which of these are types of Scaling policies?

- 

### Business

- 

### Minimized

- 

### Standard

**(Correct)**

- 

### Economy

**(Correct)**

- 

### Maximized

### Explanation

There are two different scaling policies, one is the standard policy, and one is the economy policy.

Question 82:

### Skipped

A stored procedure can run both the caller's and the owner's rights simultaneously. (TRUE / FALSE)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. A **caller's rights** stored procedure runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query. An **owner's rights** stored procedure runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table. At the time that the stored procedure is created, the creator specifies whether the procedure runs with owner's rights or caller's rights. The default is owner's rights.

Question 83:

### **Skipped**

**Which privilege is required to enable altering any properties of the resource monitor, such as changing the monthly credit quote?**

- 

**MODIFY**

**(Correct)**

- 

**MONITOR**

- 

**USAGE**

### **Explanation**

- MODIFY - Enables altering any properties of a resource monitor, such as changing the monthly credit quota.**

- MONITOR - Enables viewing a resource monitor.

Question 84:

**Skipped**

**If a warehouse runs for 61 seconds, shuts down, and then restarts and runs for less than 60 seconds, for how much duration will the billing be charged?**



**180 seconds**



**120 seconds**



**60 seconds**



**61 seconds**



**121 seconds**

**(Correct)**

### **Explanation**

It will be billed for 121 seconds ( $60 + 1 + 60$ ). Please note - minimum billing if for 60 secs every time the warehouse starts, even if the warehouse is suspended within 60 seconds.

Question 85:

**Skipped**

**When the Virtual Warehouse data cache gets filled up, in which fashion does the data get flushed out from the data cache?**



**Last In Last Out (LIFO)**



## First In First Out (FIFO)

- 

## LEAST-RECENTLY USED (LRU)

(Correct)

- 

## MOST-RECENTLY USED (MRU)

### Explanation

If the cache does fill up, it is flushed out in a **least-recently used** fashion.

Question 86:

Skipped

Which capabilities are available in Snowsight (the new Snowflake web interface)? (Select all that apply)

- 

Creating and managing users and other account-level objects

(Correct)

- 

The smart autocompletes feature suggests SQL or object syntax to insert

(Correct)

- 

Snowflake Marketplace is not available with Snowsight currently

- 

Sharing data with other Snowflake accounts

(Correct)

- 

You can display visual statistics on columns (SUM, MIN, MAX, etc.) without re-running the query

**(Correct)**

### **Explanation**

Snowsight is the new Snowflake Web Interface. It can be used to perform the following operations:

- Building and running queries.
- Loading data into tables.
- Monitoring query performance and copy history.
- Creating and managing users and other account-level objects.
- Creating and using virtual warehouses.
- Creating and modifying databases and all database objects.
- Sharing data with other Snowflake accounts.
- Exploring and using the Snowflake Marketplace.

One of the cool features is the smart autocomplete, which suggests SQL or object syntax to insert.

Question 87:

**Skipped**

**Which of these security features are supported in Snowflake? (Select all that apply)**

- 

**Network Policy**

**(Correct)**

- 

**Multi-Factor Authentication**

**(Correct)**

- 

**Tri-Secret Secure Encryption**

**(Correct)**

-

## Role-Based Access Control

(Correct)

### Explanation

Snowflake is a highly secured platform and provides multi-level security like Multi-Factor Authentication (MFA), provision to set up Network policy to block access by unwanted IPs, Single Sign On (SSO), Role-Based Access Control, and Tri Secret Secure, and so on.

Tri-Secret Secure is the combination of a Snowflake-maintained key and a customer-managed key in the cloud provider platform that hosts your Snowflake account to create a composite master key to protect your Snowflake data.

Question 88:

Skipped

**The BI group is complaining about their queries taking too long to run. Checking the virtual warehouse information shows that the queued time is pretty high. What is the best way to fix this issue?**

- Determine which users have the high priority queries and set the other users**
- Increase the size of the virtual warehouse**
- Set the STATEMENT\_QUEUED\_TIMEOUT\_IN\_SECONDS parameter to a low value to cancel those queries if they get in the queue**
- Provide a virtual warehouse for every user in the group**
- Increase the virtual warehouse MAX\_CLUSTER\_COUNT property**

(Correct)

### Explanation

**Queuing can be solved by SCALE-OUT (provision new clusters)**, i.e., increasing MAX\_CLUSTER\_COUNT helps in additional cluster provisioning to handle the concurrent workloads.

Question 89:

**Skipped**

**Materialized views are particularly useful when:**

- 

**The view's base table does not change frequently.**

**(Correct)**

- 

**The view's base table changes frequently.**

- 

**Query results contain a small number of rows and/or columns relative to the base table (the table on which the view is defined).**

**(Correct)**

- 

**Query results contain results that require simple processing.**

- 

**Query results contain results that require significant processing.**

**(Correct)**

**Explanation**

Materialized views are advantageous when:

- **Query results contain a small number of rows and/or columns** relative to the base table (the table on which the view is defined).
- Query results contain **results that require significant processing**, including

- Analysis of semi-structured data.
  - Aggregates that take a long time to calculate.
- 
- The query is on an external table (i.e., data sets stored in files in an external stage), which might perform slower than querying native database tables.

- **The view's base table does not change frequently.**

Question 90:

**Skipped**

**Which role has privileges to create warehouses and databases (and other objects) in an account?**

- 

**ORGADMIN**

- 

**SECURITYADMIN**

- 

**USERADMIN**

- 

**ACCOUNTADMIN**

- 

**SYSADMIN**

**(Correct)**

**Explanation**

**SYSADMIN** role has privileges to create warehouses and databases (and other objects) in an account. This role also has the ability to grant privileges on warehouses, databases, and other objects to other roles.

Question 91:

**Skipped**

**Which of these stages can not be dropped or altered? (Select 2)**

- 

### Named Stage

- 

### Table Stage

**(Correct)**

- 

### User Stage

**(Correct)**

#### Explanation

**User Stage:** User stages cannot be altered or dropped. A user stage is allocated to each user for storing files. This stage type is designed to store staged and managed files by a single user but can be loaded into multiple tables.

**Table Stage:** Table stages cannot be altered or dropped. A table stage is available for each table created in Snowflake. This stage type is designed to store staged and managed files by one or more users but only loaded into a single table. Note that a table stage is not a separate database object but an implicit stage tied to the table itself. A table stage has no grantable privileges of its own.

**Named Stage:** A named internal stage is a database object created in a schema. This stage type can store files staged and managed by one or more users and loaded into one or more tables. **Because named stages are database objects, the ability to create, modify, use, or drop them can be controlled using security access control privileges.**

Question 92:

**Skipped**

**SQL clause that helps define the clustering key:**

- 

**CLUSTERING BY**

- 

**CLUSTER BY**

**(Correct)**

- 

**CLUSTER ON**

- 

**CLUSERTING ON**

### **Explanation**

Example - create or replace table t1 (c1 date, c2 string, c3 number) **cluster by (c1, c2);**

Question 93:

**Skipped**

**Which database objects are currently not supported for replication? (Select 3)**

- 

**Transient tables**

- 

**Views**

- 

**Stages**

**(Correct)**

- 

**Temporary tables**

**(Correct)**

- 

**Streams**

**(Correct)**

### **Explanation**

**Temporary tables, stages, tasks, streams, pipes, and external tables are not currently supported for replication.**

Question 94:

**Skipped**

Select the type of function that returns one value per invocation (one value per row).

- 

**User-Defined Function**

- 

**System Function**

- 

**Window Function**

- 

**Table Function**

- 

**Aggregate Function**

- 

**Scalar Function**

**(Correct)**

**Explanation**

A scalar function is a function that returns one value per invocation; in most cases, you can think of this as returning one value per row. This contrasts with Aggregate Functions, which return one value per group of rows. Scalar functions take every row in your table, perform some calculations on that row and give you another value back.

Question 95:

**Skipped**

What will happen if you add 0.0.0.0/0 to BLOCKED\_IP\_LIST and your IP address to ALLOWED\_IP\_LIST of a Network policy?

- 

You will be able to access the Snowflake account from your IP address

- 

**You will not be able to access the Snowflake account from your IP address.**

**(Correct)**

### **Explanation**

**If you provide both Allowed IP Addresses and Blocked IP Addresses, Snowflake applies the Blocked List first.** This would block your own access. Additionally, in order to block all IP addresses except a select list, you only need to add IP addresses to ALLOWED\_IP\_LIST. Snowflake automatically blocks all IP addresses not included in the allowed list.

Question 96:

**Skipped**

**Which objects are the securable objects in Snowflake?**

- 

**Database**

**(Correct)**

- 

**Table**

**(Correct)**

- 

**File Format**

**(Correct)**

- 

**Warehouse**

**(Correct)**

### **Explanation**

All of these are securable objects in Snowflake. Securable Object is an entity to which access can be granted. Unless allowed by a grant, access will be denied.

Question 97:

**Skipped**

In a federated Snowflake environment, can a user still log into Snowflake using their Snowflake credentials?

- 

**Yes**

**(Correct)**

- 

**No**

**Explanation**

**Users can still log into Snowflake using their Snowflake credentials.**

Question 98:

**Skipped**

Which of these SQL Queries can be answered completely by Metadata? (Select 3)

- 

**SHOW Commands**

**(Correct)**

- 

**None of these**

- 

**MIN, MAX (integers and dates)**

**(Correct)**

- 

**COUNT**

**(Correct)**

- 

**AVG**

### Explanation

**SHOW Commands, MIN, MAX (integers and dates), and COUNT SQL queries take advantage of the Metadata cache and do not require the virtual warehouse, but you still have some cloud service charges.**

Question 99:

### Skipped

If you create a user with **MUST\_CHANGE\_PASSWORD = TRUE** in a Snowflake federated environment, will that user be forced to change the password while logging through IdP the first time?

- 

**No**

**(Correct)**

- 

**Yes**

### Explanation

The **MUST\_CHANGE\_PASSWORD** user property does not apply for federated authentication and should not be used. In particular, if you choose not to maintain passwords in Snowflake for users, ensure this property is set to FALSE for these users.

Question 100:

### Skipped

Suppose we resize a warehouse to a larger size while it is executing SQL statements. In that case, the already running SQL statements will finish faster.  
**(True / False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### Explanation

Resizing a warehouse doesn't have any impact on statements that are currently being executed by the warehouse. When resizing to a larger size, the new compute resources, once fully provisioned, are used only to execute statements that are

already in the warehouse queue, as well as all future statements submitted to the warehouse.

Question 1:

Skipped

Which is generally the slowest option for selecting staged data files to load from a stage?

- Specifying a list of specific files to load
- By path (internal stages) / prefix (Amazon S3 bucket)
- Using pattern matching to identify specific files by pattern

(Correct)

Explanation

Pattern matching using a regular expression is generally the slowest of the three options for identifying/specifying data files to load from a stage; however, this option works well if you exported your files in named order from your external application and want to batch load the files in the same order

Question 2:

Skipped

Which copyOptions can help load a file with expired metadata (if the LAST\_MODIFIED date is older than 64 days and the initial set of data was loaded into the table more than 64 days earlier (and if the file was loaded into the table, that also occurred more than 64 days earlier))? (Select 2)

- ON\_ERROR = CONTINUE
- FORCE = FALSE
- LOAD\_UNCERTAIN\_FILES = TRUE

(Correct)

- LOAD\_FILES = TRUE
- LOAD\_CERTAIN\_FILES = TRUE

**FORCE = TRUE**

**(Correct)**

### **Explanation**

To load files whose metadata has expired, set the LOAD\_UNCERTAIN\_FILES copy option to true. The copy option references load metadata, if available, to avoid data duplication, but also attempts to load files with expired load metadata. Alternatively, set the FORCE option to load all files, ignoring load metadata if it exists. Note that this option reloads files, potentially duplicating data in a table.

Question 3:

**Skipped**

**select \* from t1 sample row(100); What would the above query return?**

- 

**Return a sample of a table in which each row has a 10% probability of being included in the sample**

- 

**Return an empty sample**

- 

**samplingMethod is not applied in the query. The query will result in an error.**

- 

**Return an entire table, including all rows in the table**

**(Correct)**

### **Explanation**

Return an entire table, including all rows in the table. Sampling method is optional. If no method is applied after sample keyword, the default it takes is BERNoulli.

Question 4:

**Skipped**

**Which of the following file format is not supported by Snowflake?**

- 

**CSV**

- 

**PARQUET**

- 

**JSON**

-

**EDI**

**(Correct)**

- 

**ORC**

- 

**AVRO**

### **Explanation**

Snowflake supports - CSV, TSV, JSON, AVRO, ORC, PARQUET. Snowflake also supports XML which is a Preview feature as of now. EDI format is not supported by Snowflake.

Question 5:

**Skipped**

**File URL is ideal for**

- 

**None of these**

- 

**Ideal for use in custom applications, providing unstructured data to other accounts via a share**

- 

**Ideal for custom applications that require access to unstructured data files**

**(Correct)**

- 

**Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents**

### **Explanation**

Scoped URL: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. Ideal for use in custom applications, providing unstructured data to other accounts via a share, or for downloading and ad hoc analysis of unstructured data via Snowsight. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. Ideal for custom applications that require access to unstructured data files. Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable. Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents.

Question 6:

**Skipped**

Which of the following languages does Snowflake support for writing UDFs (User-Defined Functions)? (Select 4)

- 

**Python**

**(Correct)**

- 

**JAVA**

**(Correct)**

- 

**JavaScript**

**(Correct)**

- 

**GO**

- 

**SQL**

**(Correct)**

- 

**C#**

### Explanation

User-defined functions (UDFs) let you extend the system to perform operations that are not available through the built-in, system-defined functions provided by Snowflake. Snowflake currently supports the following languages for writing UDFs: **Java**: A Java UDF lets you use the Java programming language to manipulate data and return either scalar or tabular results. **JavaScript**: A JavaScript UDF lets you use the JavaScript programming language to manipulate data and return either scalar or tabular results. **Python**: A Python UDF lets you use the Python programming language to manipulate data and return either scalar or tabular results. **SQL**: A SQL UDF evaluates an arbitrary SQL expression and returns either scalar or tabular results.

Question 7:

**Skipped**

**What is the expiration period of a File URL?**

- 

**It is Permanent**

**(Correct)**

- 
- Length of time specified in the expiration\_time argument**
- 
- The URL expires when the persisted query result period ends**

### Explanation

The expiration period of Scoped URL: The URL expires when the persisted query result period ends. The expiration period of the File URL: It is permanent. The expiration period of Pre-Signed URL: Length of time specified in the expiration\_time argument.

Question 8:

**Skipped**

**Loading into Snowflake from a local file system is a straightforward affair. Which command is used to grab files from the local system, compress them and encrypt them, and then it copies them to Snowflake?**

- 
- COPY INTO <table>**
- 
- GET command**
- 
- MOVE command**
- 
- PUT command**

**(Correct)**

### Explanation

Using the PUT command in SnowSQL. It grabs the file or files, compresses them, encrypts them, and then copies them up into the stage you chose. Once in the stage, you can use a COPY INTO command to load the data from the stage into Snowflake tables.

Question 9:

**Skipped**

**Which of these sampling method keywords are used to specify which method to use? (Select 2)**

- 
- SYSTEM | BLOCK**
- 
- (Correct)**
- 
- SYSTEM | ROW**

- 

**BERNOULLI | ROW**

**(Correct)**

- 

**BERNOULLI | BLOCK**

### **Explanation**

BERNOULLI | ROW and SYSTEM | BLOCK are used to specify the sampling method in SELECT query. BERNOULLI (or ROW): Includes each row with a <probability> of p/100. Similar to flipping a weighted coin for each row. SYSTEM (or BLOCK): Includes each block of rows with a <probability> of p/100. Similar to flipping a weighted coin for each block of rows. This method does not support fixed-size sampling. Sampling method is optional. If no method is specified, the default is BERNOULLI. Example : select \* from t1 tablesample bernoulli (25); This query will return a sample of a table in which each row has a 25% probability of being included in the sample

Question 10:

**Skipped**

**Scoped URL is ideal for**

- 

**Ideal for use in custom applications, providing unstructured data to other accounts via a share**

**(Correct)**

- 

**Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents**

- 

**Ideal for custom applications that require access to unstructured data files**

- 

**None of these**

### **Explanation**

Scoped URL: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. Ideal for use in custom applications, providing unstructured data to other accounts via a share, or for downloading and ad hoc analysis of unstructured data via Snowsight. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. Ideal for custom applications that require access to unstructured data files. Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is

configurable. Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents.

Question 11:

**Skipped**

Which of these SQL functions helps retrieve the URL for an external or internal named stage using the stage name as the input?

- BUILD\_STAGE\_FILE\_URI**
  - 
  - GET\_PRESIGNED\_URL**
  - 
  - GET\_RELATIVE\_PATH**
  - 
  - GET\_STAGE\_LOCATION**
- (Correct)**
- - GET\_ABSOLUTE\_PATH**
  - 
  - BUILD\_SCOPED\_FILE\_URL**

### **Explanation**

GET\_STAGE\_LOCATION retrieves the URL for an external or internal named stage using the stage name as the input.

Question 12:

**Skipped**

Select the type of function that returns one value per group of rows (for example - AVG, MAX, MIN)

- - Table Function**
  - 
  - System Function**
  - 
  - Aggregate Function**
- (Correct)**

- 

## Window Function

- 

## Scalar Function

- 

## User-Defined Function

### Explanation

Aggregate functions operate on values across rows to perform mathematical calculations such as sum, average, counting, minimum/maximum values, standard deviation, and estimation, as well as some non-mathematical operations. An aggregate function takes multiple rows (actually, zero, one, or more rows) as input and produces a single output. In contrast, scalar functions take one row as input and produce one row (one value) as output. An aggregate function always returns exactly one row, even when the input contains zero rows. Typically, if the input contained zero rows, the output is NULL. However, an aggregate function could return 0, an empty string, or some other value when passed zero rows.

Question 13:

### Skipped

**Both external (external cloud storage) and internal (i.e., Snowflake) stages support unstructured data. (True / False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### Explanation

True, both external (external cloud storage, such as, Amazon S3, Google Cloud Storage, Azure Blob Storage etc.) and internal (i.e. Snowflake) stages support unstructured data.

Question 14:

### Skipped

**While transforming Semi-structure data, If you want expansion for all the sub-elements recursively using FLATTEN function, what argument would you need to set with FLATTEN function?**

- 

**OUTER => FALSE**

- 

**OUTER => TRUE**

- 

**RECURSIVE => TRUE**

**(Correct)**

- 

**RECURSIVE => FALSE**

### **Explanation**

The expansion is performed for all sub-elements recursively by argument RECURSIVE => TRUE. Only the element referenced by PATH is expanded BY RECURSIVE => FALSE. The OUTER argument is used to handle the input rows that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries.

Question 15:

**Skipped**

**Which of these functions helps generate the Pre-signed URL to access the unstructured data file?**

- 

**GET\_STAGE\_LOCATION**

- 

**GET\_ABSOLUTE\_PATH**

- 

**BUILD\_SCOPED\_FILE\_URL**

- 

**GET\_RELATIVE\_PATH**

- 

**BUILD\_STAGE\_FILE\_URI**

- 

**GET\_PRESIGNED\_URL**

**(Correct)**

### **Explanation**

GET\_PRESIGNED\_URL generates the pre-signed URL to a staged file using the stage name and relative file path as inputs. Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable.

Question 16:

**Skipped**

**Which of these SQL functions helps extract the path of a staged file relative to its location in the stage using the stage name and absolute file path in cloud storage as inputs?**

- 
- GET\_PRESIGNED\_URL**
- 
- GET\_RELATIVE\_PATH**
- (Correct)**
- 
- GET\_ABSOLUTE\_PATH**
- 
- BUILD\_STAGE\_FILE\_URI**
- 
- BUILD\_SCOPED\_FILE\_URL**
- 
- GET\_STAGE\_LOCATION**

#### **Explanation**

GET\_RELATIVE\_PATH extracts the path of a staged file relative to its location in the stage using the stage name and absolute file path in cloud storage as inputs.

Question 17:

**Skipped**

**How can a directory table metadata be refreshed automatically and efficiently to synchronize the metadata with the latest associated files in the external stage and path?**

- - Using Tasks**
  - 
  - It is a manual process and cant be automatically refreshed**
  - 
  - Using both Tasks and Stream**
  - 
  - Using Cloud event notification service**
- (Correct)**

- 

## Using Stream

### Explanation

The metadata for a directory table can be refreshed automatically using the event notification service for your cloud storage service. The refresh operation synchronizes the metadata with the latest set of associated files in the external stage and path, i.e.: - New files in the path are added to the table metadata. - Changes to files in the path are updated in the table metadata. - Files no longer in the path are removed from the table metadata.

Question 18:

**Skipped**

**When should we use HyperLogLog?**

- 

**Whenever the input is potentially large, and an approximate result is not acceptable.**

- 

**Whenever the input is potentially small, and an approximate result is acceptable.**

- 

**Whenever the input is potentially small, and an approximate result is not acceptable.**

- 

**Whenever the input is potentially large, and an approximate result is acceptable.**

**(Correct)**

### Explanation

Snowflake recommends using HyperLogLog whenever the input is potentially large, and an approximate result is acceptable.

Question 19:

**Skipped**

**Select the type of function that can operate on a subset of rows within the set of input rows.**

- 

**Aggregate Function**

- 

**Scalar Function**

- 

**Table Function**

- 

**System Function**

- 

### Window Function

(Correct)

- 

### User-Defined Function

#### Explanation

A window function is any function that operates over a window of rows.

Question 20:

Skipped

**John wants to load data files from an external stage to Snowflake. He has split the large file into smaller 100 - 250 MB data files, and there is a total of 16 smaller data files. What warehouse size would you recommend him to use for loading these data files quickly and cost-effectively?**

- 

XS

- 

M

- 

L

- 

S

(Correct)

- 

XL

#### Explanation

XS sized warehouse can load eight files parallelly. S sized warehouse can load sixteen files parallelly. M sized warehouse can load thirty-two files parallelly. L sized warehouse can load sixty-four files parallelly. XL sized warehouse can load one hundred twenty-eight files parallelly and so on.

Question 21:

Skipped

Pre-Signed URL is ideal for

- 

None of these

-

**Ideal for custom applications that require access to unstructured data files**

- 

**Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents**

**(Correct)**

- 

**Ideal for use in custom applications, providing unstructured data to other accounts via a share**

### **Explanation**

Scoped URL: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. Ideal for use in custom applications, providing unstructured data to other accounts via a share, or for downloading and ad hoc analysis of unstructured data via Snowsight. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. Ideal for custom applications that require access to unstructured data files. Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable. Ideal for business intelligence applications or reporting tools that need to display the unstructured file contents.

Question 22:

**Skipped**

**What does OVERWRITE parameter do with INSERT command?**

- 

**It specifies that the target table should be truncated before inserting the values into the table.**

**(Correct)**

- 

**It drops the table, recreates, and inserts the values into the table.**

- 

**It de-duplicates while inserting and skips the insert if there is an exact similar record in the table.**

- 

**It helps ignore any errors while inserting the values into the table.**

### **Explanation**

OVERWRITE specifies that the target table should be truncated before inserting the values into the table. Note that specifying this option does not affect the access control privileges on the table.

Question 23:

**Skipped**

Suppose files downloaded from an internal stage are corrupted. What should be verified with the stage creator to determine why the downloaded file is corrupted?

- 

**Verify if ENCRYPTION = (TYPE = 'SNOWFLAKE\_SSE') set for the stage**

**(Correct)**

- 

**Verify if ENCRYPTION = (TYPE = 'SNOWFLAKE\_FULL') set for the stage**

**Explanation**

If files downloaded from an internal stage are corrupted, verify with the stage creator that ENCRYPTION = (TYPE = 'SNOWFLAKE\_SSE') is set for the stage.

Question 24:

**Skipped**

What all options are available for data transformation while loading data into a table using the COPY command? (Select all that apply)

- 

**Truncation of Text Strings**

**(Correct)**

- 

**Casts**

**(Correct)**

- 

**Column omission**

**(Correct)**

- 

**Column reordering**

**(Correct)**

- 

**Join**

**Explanation**

Snowflake supports transforming data while loading it into a table using the COPY command. Options include:

- Column reordering
- Column omission
- Casts
- Truncating text strings that exceed the target column length

Question 25:

Skipped

Select the type of function that returns multiple rows for each individual input.

- Scalar Function
- Window Function
- Table Function
- (Correct) User-Defined Function
- System Function
- Aggregate Function

### Explanation

A table function returns a set of rows for each input row. The returned set can contain zero, one, or more rows. Each row can contain one or more columns. Table functions are sometimes called “tabular functions”.

Question 26:

Skipped

A Directory table is a separate database object that stores a catalog of staged files in cloud storage. (True/False)

- TRUE
- FALSE

(Correct)

### Explanation

A Directory table is not a separate database object; it stores a catalog of staged files in cloud storage. Roles with sufficient privileges can query a directory table to retrieve file URLs to access the staged files and other metadata.

Question 27:

Skipped

Which function can be used in combination with COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file?

- 
- **VARIANT\_CONSTRUCT**
- 
- **OBJECT\_CONSTRUCT**
- **(Correct)**
- 
- **COPY\_CONSTRUCT**
- 
- **UNLOAD\_CONSTRUCT**

#### Explanation

The OBJECT\_CONSTRUCT function can be used in combination with the COPY command to convert the rows in a relational table to a single VARIANT column and unload the rows into a file.

Question 28:

Skipped

Which of these are unstructured data? (Select 2)

- **XML**
- **Relational Data**
- **JSON**
- **Videos**
- **(Correct)**
- **Images**

**(Correct)**

### Explanation

Unstructured data is information that does not fit into a predefined data model or schema.

Typically text-heavy, such as form responses and social media conversations, unstructured data also encompasses images, video, and audio. Industry-specific file types such as VCF (genomics), KDF (semiconductors), or HDF5 (aeronautics) are included in this category.

Question 29:

**Skipped**

**Only the user who generated the scoped URL can use the URL to access the referenced file.**

**(True/False)**

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### Explanation

True, only the user who generated the scoped URL can use the URL to access the referenced file. In case of File URL, any role that has sufficient privileges on the stage can access the file.

Question 30:

**Skipped**

**Which is the fastest option for selecting staged data files to load from a stage?**

- 

**Using pattern matching to identify specific files by pattern**

- 

**Specifying a list of specific files to load**

**(Correct)**

- 

**By path (internal stages) / prefix (Amazon S3 bucket)**

### Explanation

Of the three options for identifying/specifying data files to load from a stage, providing a discrete list of files is generally the fastest; however, the FILES parameter supports a maximum of 1,000 files, meaning a COPY command executed with the FILES parameter can only load up to 1,000 files. Example: copy into load1 from @%load1/data1/ files=('test1.csv', 'test2.csv', 'test3.csv', 'test4.csv')

Question 31:

**Skipped**

**Which copy option is used to delete the file from the Snowflake stage when data from staged files are loaded successfully?**

- 

**DELETE = TRUE**

- 

**PURGE = TRUE**

**(Correct)**

- 

**REMOVE = TRUE**

- 

**DEL = TRUE**

### **Explanation**

Staged files can be deleted from a Snowflake stage (user stage, table stage, or named stage) using the following methods: 1- Files that were loaded successfully can be deleted from the stage during a load by specifying the PURGE copy option in the COPY INTO <table> command. 2- After the load completes, use the REMOVE command to remove the files in the stage. Please note, DELETE or REMOVE are not COPY command options. REMOVE is a different DML command which is used to remove files in the stage.

Question 32:

**Skipped**

**Select the type of function that is used to execute an action in the system or return information about the system.**

- 

**Scalar Function**

- 

**System Function**

**(Correct)**

- 

**Aggregate Function**

- 

**Table Function**

- 

**Window Function**

## User-Defined Function

### Explanation

System function that is used to execute an action in the system or return information about the system. Snowflake provides the following types of system functions: Control functions that allow you to execute actions in the system (e.g. aborting a query). Information functions that return information about the system (e.g. calculating the clustering depth of a table). Information functions that return information about queries (e.g. information about EXPLAIN plans).

Question 33:

**Skipped**

**Snowflake supports the secured access of unstructured data files in cloud storage. What types of URLs are available to access files in cloud storage? (Select 3)**

- 

**Descoped URL**

- 

**Signed URL**

- 

**Unstructured URL**

- 

**Scoped URL**

**(Correct)**

- 

**File URL**

**(Correct)**

- 

**Pre-signed URL**

**(Correct)**

### Explanation

The following types of URLs are available to access files in cloud storage: Scoped URL: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files. Pre-signed URL: Simple HTTPS URL used to access a file via a web browser. A file is temporarily accessible to users via this URL using a pre-signed access token. The expiration time for the access token is configurable.

Question 34:

### Skipped

Monica is confused about which sampling method she should use with one the very large tables,, considering better performance. Which sampling method would you recommend from BERNOULLI | ROW and SYSTEM | BLOCK?

- 

SYSTEM | BLOCK

(Correct)

- 

BERNOULLI | ROW

### Explanation

SYSTEM | BLOCK sampling is often faster than BERNOULLI | ROW sampling. Also, BERNOULLI | ROW method is good for Smaller Tables and SYSTEM | BLOCK method for Larger Tables.

Question 35:

### Skipped

What size limit does VARIANT data type impose on individual rows?

- 

16 GB

- 

10 GB

- 

16 MB

(Correct)

- 

10 MB

- 

100 MB

### Explanation

The VARIANT data type imposes a 16 MB size limit on individual rows.

Question 36:

### Skipped

Which of these SQL functions helps generate the Scoped URL to access the unstructured data file?

- 

GET\_ABSOLUTE\_PATH

- 

**GET\_PRESIGNED\_URL**

- 

**BUILD\_STAGE\_FILE\_URI**

- 

**GET\_STAGE\_LOCATION**

- 

**BUILD\_SCOPED\_FILE\_URL**

**(Correct)**

- 

**GET\_RELATIVE\_PATH**

### **Explanation**

BUILD\_SCOPED\_FILE\_URL generates a scoped Snowflake-hosted URL to a staged file using the stage name and relative file path as inputs. A scoped URL is encoded and permits access to a specified file for a limited period of time. Scoped URL: Encoded URL that permits temporary access to a staged file without granting privileges to the stage. The URL expires when the persisted query result period ends (i.e., the results cache expires), which is currently 24 hours.

Question 37:

### **Skipped**

While transforming Semi-structure data, what argument would you need to set with FLATTEN function to omit the output of the input rows that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries?

- 

**OUTER => TRUE**

- 

**OUTER => FALSE**

**(Correct)**

- 

**RECURSIVE => FALSE**

- 

**RECURSIVE => TRUE**

-

## MODE => OBJECT

### Explanation

The OUTER => FALSE argument with FLATTEN omits the output of the input rows that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries. The OUTER => TRUE argument with FLATTEN generates exactly one row for zero-row expansions (with NULL in the KEY, INDEX, and VALUE columns). RECURSIVE is used to instruct if only the element referenced by PATH is expanded or expansion is performed for all sub-elements recursively. MODE Specifies whether only objects, arrays, or both should be flattened.

Question 38:

Skipped

How can we add a Directory table explicitly to a stage to store a catalog of staged files?

- 
- Using CREATE DIRECTORY TABLES command and then add to the stage by ALTER STAGE command**
- 
- Using CREATE DIRECTORY TABLE command and then add to the stage by ALTER STAGE command**
- 
- Using CREATE STAGE command**

(Correct)

### Explanation

A Directory table is not a separate database object; it stores a catalog of staged files in cloud storage. Roles with sufficient privileges can query a directory table to retrieve file URLs to access the staged files and other metadata. A directory table can be added explicitly to a stage when the stage is created (using CREATE STAGE) or later (using ALTER STAGE) with supplying directoryTableParams. directoryTableParams (for internal stages) ::= [ DIRECTORY = ( ENABLE = { TRUE | FALSE } [ REFRESH\_ON\_CREATE = { TRUE | FALSE } ] ) ] ENABLE = TRUE | FALSE Specifies whether to add a directory table to the stage. When the value is TRUE, a directory table is created with the stage.

Question 39:

Skipped

Which algorithm does Snowflake use to estimate the approximate number of distinct values in a data set?

- 
- HyperEstimateLog**
- 
- HyerAccumulateLog**
-

## HyperLogLog

(Correct)

- 

## HyperMeanLog

- 

## HyperMedianLog

### Explanation

Snowflake uses HyperLogLog to estimate the approximate number of distinct values in a data set. HyperLogLog is a state-of-the-art cardinality estimation algorithm, capable of estimating distinct cardinalities of trillions of rows with an average relative error of a few percent.

Question 40:

Skipped

Which of these compression methods are supported by Snowflake? (Select all that apply)

- 

Brotli

(Correct)

- 

gzip

(Correct)

- 

raw\_deflate

(Correct)

- 

deflate

(Correct)

- 

bzip2

(Correct)

- 

Zstandard

**(Correct)**

### Explanation

All of these are supported by Snowflake. Snowflake can automatically detect any of these compression methods except Brotli and Zstandard.

Question 41:

**Skipped**

**When loading data using COPY INTO <table> command, it is a must that the data files have the same number and order of columns as your target table. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### Explanation

There is no requirement for your data files to have the same number and ordering of columns as your target table.

Question 42:

**Skipped**

**The VALIDATION\_MODE parameter does not support COPY statements that transform data during a load. (True / False)**

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### Explanation

True. "VALIDATION\_MODE instructs the COPY command to validate the data files instead of loading them into the specified table; i.e., the COPY command tests the files for errors but does not load them. The command validates the data to be loaded and returns results based on the validation option specified: Syntax : VALIDATION\_MODE = RETURN\_n\_ROWS | RETURN\_ERRORS | RETURN\_ALL\_ERRORS RETURN\_n\_ROWS (e.g. RETURN\_10\_ROWS) - Validates the specified number of rows, if no errors are encountered; otherwise, fails at the first error encountered in the rows. RETURN\_ERRORS - Returns all errors (parsing, conversion, etc.) across all files specified in the COPY statement. RETURN\_ALL\_ERRORS - Returns all errors across all files specified in the COPY statement, including files with errors that were partially loaded during an earlier load because the ON\_ERROR copy option was set to CONTINUE during the load."

Question 43:

**Skipped**

Which methods can be used to check the status of a COPY INTO command? (Select 2)

- **Use the resource monitor.**
  - **Write a SQL query against the ACCOUNT\_USAGE.LOAD\_HISTORY view.**
  - **Write a SQL query against the INFORMATION\_SCHEMA.LOAD\_HISTORY view.**
- (Correct)**
- **Use the Query History tab in the Snowflake user interface.**
- (Correct)**

#### Explanation

The status of COPY INTO command can be checked from the Resource Monitors tab in the Snowflake user interface, as well as querying the INFORMATION\_SCHEMA.LOAD\_HISTORY view.

Question 44:

**Skipped**

What is the default compression method used by Snowflake while unloading data?

- **gzip**
- 
- **bzip2**
- 
- **Zstandard**
- 
- **Brotli**

#### Explanation

By default, all unloaded data files are compressed using gzip, unless compression is explicitly disabled or one of the other supported compression methods is explicitly specified.

Question 45:

**Skipped**

What is the preferred way to distinguish empty strings from NULLs while unloading in CSV files?

-

### **Set EMPTY\_FIELD\_AS\_NULL to TRUE**

- 

**Enclose strings in quotes by setting the FIELD\_OPTIONALLY\_ENCLOSED\_BY option.**

**(Correct)**

- 

**Leave string fields unenclosed by setting the FIELD\_OPTIONALLY\_ENCLOSED\_BY option to NONE.**

- 

### **Set EMPTY\_FIELD\_AS\_NULL to FALSE**

#### **Explanation**

An empty string is typically represented by a quoted empty string (e.g. "") to indicate that the string contains zero characters. The preferred way is to enclose strings in quotes by setting the FIELD\_OPTIONALLY\_ENCLOSED\_BY option, to distinguish empty strings from NULLs in output CSV files.

Question 46:

**Skipped**

**The best use of Snowpipe is to load large volumes of data and incrementally make them available for analysis. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

#### **Explanation**

Snowpipe is designed to load small volume of data (i.e., micro-batches) and incrementally make them available for analysis.

Question 47:

**Skipped**

**Snowflake supports \_\_\_\_\_**

- 

**Both of these**

- 

**SOAP for unstructured data**

-

**None of these**

- 

**REST API for unstructured data**

**(Correct)**

**Explanation**

Currently, Snowflake only supports REST API for unstructured data.

Question 48:

**Skipped**

**How does Snowflake store semi-structured data, such as JSON, Avro, Parquet, ORC, and XML?**

- 

**Stores as VARCHAR data type**

- 

**Stores as VARIANT data type**

**(Correct)**

- 

**Stores as FLATTEN data type**

- 

**Stores as JSON data type**

**Explanation**

Snowflake stores semi-structured data, such as JSON, Avro, Parquet, ORC, and XML, as VARIANT data type. A VARIANT can store a value of any other type, including OBJECT and ARRAY.

Question 49:

**Skipped**

**What is the purpose of VALIDATION\_MODE in the COPY INTO <table> command?**

- 

**VALIDATION\_MODE is used to validate the load file and load it into the specified table if there is no error.**

- 

**VALIDATION\_MODE is used to validate the load file for errors instead of loading it into the specified table.**

**(Correct)**

-

**VALIDATION\_MODE** is used to validate the load file, skip the errored data and then load it into the specified table.

### Explanation

VALIDATION\_MODE instructs the COPY command to validate the data files instead of loading them into the specified table; i.e., the COPY command tests the files for errors but does not load them. The command validates the data to be loaded and returns results based on the validation option specified: Syntax : VALIDATION\_MODE = RETURN\_n\_ROWS | RETURN\_ERRORS | RETURN\_ALL\_ERRORS RETURN\_n\_ROWS (e.g. RETURN\_10\_ROWS) - Validates the specified number of rows, if no errors are encountered; otherwise, fails at the first error encountered in the rows. RETURN\_ERRORS - Returns all errors (parsing, conversion, etc.) across all files specified in the COPY statement. RETURN\_ALL\_ERRORS - Returns all errors across all files specified in the COPY statement, including files with errors that were partially loaded during an earlier load because the ON\_ERROR copy option was set to CONTINUE during the load.

Question 50:

Skipped

**VARIANT** is used to FLATTEN hierarchical data. (True / False)

- 

**FALSE**

(Correct)

- 

**TRUE**

### Explanation

VARIANT is a data type that can hold a value of any other data type (including ARRAY and OBJECT). VARIANT is used to build and store hierarchical data. VARIANT is not a function to FLATTEN. FLATTEN is a table function that is used to produce a lateral view of a VARIANT, OBJECT, or ARRAY column.

Question 51:

Skipped

**Which options for selecting staged data files are supported by COPY command in Snowflake?**  
**(Select that all apply)**

- 

**Specifying a list of specific files to load**

(Correct)

- 

**Using pattern matching to identify specific files by pattern**

(Correct)

- 

**By path (internal stages) / prefix (Amazon S3 bucket)**

**(Correct)**

### Explanation

All of these are supported. These options enable you to copy a fraction of the staged data into Snowflake with a single command. This allows you to execute concurrent COPY statements that match a subset of files, taking advantage of parallel operations.

Question 52:

**Skipped**

**What is the default encoding used by Snowflake while unloading data?**

- UTF-16**
- UTF-8**
- UTF**
- UTF-32**

### Explanation

Output files are always encoded using UTF-8, regardless of the file format; no other character sets are supported.

Question 53:

**Skipped**

**What value will be return by the following query? SELECT \* FROM TABLE(FLATTEN(input => parse\_json('[]')) f;**

- 
- NULL**
- []**
- 0**
- nothing will return / output of the input row will be omitted**

**(Correct)**

## Explanation

If you don't specify OUTER argument with FLATTEN, it would be defaulted to FALSE. The OUTER => FALSE argument with FLATTEN omits the output of the input rows that cannot be expanded, either because they cannot be accessed in the path or because they have zero fields or entries.

Question 54:

### Skipped

A stored procedure can run both the caller's and the owner's rights simultaneously. (TRUE / FALSE)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

## Explanation

A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. A caller's rights stored procedure runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query. An owner's rights stored procedure runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table. At the time that the stored procedure is created, the creator specifies whether the procedure runs with owner's rights or caller's rights. The default is owner's rights.

Question 55:

### Skipped

If a file in a stage has the LAST\_MODIFIED date older than 64 days and the initial set of data was loaded into the table more than 64 days earlier. In this case, to prevent any data loss, the COPY command loads the file by default. (True / False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

## Explanation

The COPY command cannot definitively determine whether a file has been loaded already if the LAST\_MODIFIED date is older than 64 days and the initial set of data was loaded into the table more than 64 days earlier (and if the file was loaded into the table, that also occurred more than 64 days earlier). In this case, to prevent accidental reload, the command skips the file by default.

Question 56:

## Skipped

What all locations do Snowflake support for staging the data? (Select all that apply)

- 

**Oracle Cloud Storage**

- 

**Microsoft Azure Blob Storage**

**(Correct)**

- 

**Google Cloud Storage**

**(Correct)**

- 

**Amazon S3**

**(Correct)**

- 

**Snowflake Internal Stages**

**(Correct)**

## Explanation

Snowflake supports loading data from files staged in any of the following locations (**except Oracle Cloud Storage**), regardless of the cloud platform for your Snowflake account:

- Internal (i.e. Snowflake) stages
- Amazon S3
- Google Cloud Storage
- Microsoft Azure blob storage

Question 57:

## Skipped

Which of these Sampling methods does Snowflake support? (Select 2)

- 

**Sample a fixed, specified number of rows**

**(Correct)**

- 

**Sample exact rows of the table with the specified sequence keys**

-

## Sample a fraction of the table with a specified probability of including a given row

(Correct)

### Explanation

SAMPLE / TABLESAMPLE returns a subset of rows sampled randomly from the specified table. The following sampling methods are supported: **Sample a fraction of a table** with a specified probability for including a given row. The number of rows returned depends on the size of the table and the requested probability. A seed can be specified to make the sampling deterministic.

**Sample a fixed, specified number of rows.** The exact number of specified rows is returned unless the table contains fewer rows. SAMPLE and TABLESAMPLE are synonymous and can be used interchangeably.

Question 58:

Skipped

During data loading using COPY INTO <table> command, if the string exceeds the target column length, what options do you have to truncate the string? (Select 2)

- 

ENFORCE\_LENGTH = FALSE

(Correct)

- 

ENFORCE\_LENGTH = TRUE

- 

TRUNCATECOLUMNS = FALSE

- 

TRUNCATECOLUMNS = TRUE

(Correct)

### Explanation

ENFORCE\_LENGTH: - If TRUE, the COPY statement produces an error if a loaded string exceeds the target column length. - If FALSE, strings are automatically truncated to the target column length. TRUNCATECOLUMNS: - If TRUE, strings are automatically truncated to the target column length. - If FALSE, the COPY statement produces an error if a loaded string exceeds the target column length.

Question 59:

Skipped

Monica wants to share a UDF with other users. She wants to permit other users to use it, but she doesn't want them to be able to see how it was defined or the underlying logic behind it. What would you recommend to Monica?

-

**Monica should create an unsecure UDF and then block the logic in the sql statement**

- 
- Underlying logic can not be hidden with UDF

- 
- Monica should create a secure UDF

**(Correct)**

### **Explanation**

We can create a user-defined function to be unsecure or secure. A secure user-defined function means if we permit someone else to use this UDF, they will not be able to see how it was defined or the underlying logic behind it.

Question 60:

**Skipped**

Select the type of function that returns one value per invocation (one value per row).

- 
- Aggregate Function

- 
- Window Function

- 
- User-Defined Function

- 
- Table Function

- 
- System Function

- 
- Scalar Function

**(Correct)**

### **Explanation**

A scalar function is a function that returns one value per invocation; in most cases, you can think of this as returning one value per row. This contrasts with Aggregate Functions, which return one value per group of rows. Scalar functions take every row in your table, perform some calculation on that row and give you another value back.

Question 61:

**Skipped**

John has to create a PIPE that will be triggered for loading by calling the Snowpipe REST endpoints. What parameter does he need to specify in CREATE PIPE statement?

- AUTO\_INGEST = FALSE  
**(Correct)**
- API\_INGEST = FALSE
- AUTO\_INGEST = TRUE
- API\_INGEST = TRUE

#### Explanation

AUTO\_INGEST = TRUE enables automatic data loading. Snowpipe supports loading from external stages (Amazon S3, Google Cloud Storage, or Microsoft Azure). AUTO\_INGEST = FALSE disables automatic data loading. You must make calls to the Snowpipe REST API endpoints to load data files.

Question 62:

**Skipped**

What are all operations performed using Snowflake SQL API? (Select 4)

- Check the status of the execution of a statement  
**(Correct)**
- Fetch query results concurrently  
**(Correct)**
- Cancel the execution of a statement  
**(Correct)**
- Calling stored procedures that returns a table

## Submit SQL statements for execution

(Correct)

### Explanation

The Snowflake SQL API provides operations that we can use to: - Submit SQL statements for execution. - Check the status of the execution of a statement. - Cancel the execution of a statement. - Fetch query results concurrently. Currently, Snowflake SQL API has limitation for the call command with stored procedures that return a table (stored procedures with the RETURNS TABLE clause).

Question 63:

Skipped

What value will be return by the following query? `SELECT * FROM TABLE(FLATTEN(input => parse_json('[]'), outer => true)) f;`

- 
- []
- 
- **NULL**

(Correct)

- 
- 0

### Explanation

The OUTER => TRUE argument with FLATTEN generates exactly one row for zero-row expansions (with NULL in the KEY, INDEX, and VALUE columns).

Question 64:

Skipped

Which command is used to create a security integration to enable an HTTP client that supports OAuth to redirect users to an authorization page and generate access tokens for access to the REST API endpoint?

- 
- **CREATE SECURITY API**
- 
- **CREATE OAUTH INTEGRATION**
- 
- **CREATE SECURITY INTEGRATION**

(Correct)

-

## CREATE INTEGRATION

### Explanation

CREATE SECURITY INTEGRATION command is used to create a security integration that supports OAuth to redirect users to an authorization page and generate access tokens for access to the REST API endpoint.

Question 65:

Skipped

An HTTP client that sends a URL (either scoped URL or file URL) to the REST API must be configured to allow redirects. (True/False)

- 

TRUE

(Correct)

- 

FALSE

### Explanation

True, An HTTP client that sends a URL (either scoped URL or file URL) to the REST API must be configured to allow redirects.

Question 66:

Skipped

UDF runs with either the caller's or the owner's rights. (TRUE / FALSE)

- 

FALSE

(Correct)

- 

TRUE

### Explanation

UDF only runs as the function owner. A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. A caller's rights stored procedure runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query. An owner's rights stored procedure runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table. At the time that the stored procedure is created, the creator specifies whether the procedure runs with owner's rights or caller's rights. The default is owner's rights.

Question 67:

Skipped

How can you unload the data from Snowflake using COPY INTO location statements in a Single file?

- By specifying copy option SINGLE=TRUE  
**(Correct)**
- By specifying copy option MULTIPLE=FALSE
- By specifying copy option MULTIPLE\_FILES=FALSE
- By specifying copy option ONE\_FILE=TRUE

#### Explanation

To unload data to a single output file (at the potential cost of decreased performance), specify the SINGLE = true copy option in your statement. You can optionally specify a name for the file in the path.

Question 68:

**Skipped**

Which of these functions helps generate the FILE URL to access the unstructured data file?

- GET\_ABSOLUTE\_PATH
- BUILD\_STAGE\_FILE\_URI  
**(Correct)**
- GET\_PRESIGNED\_URL
- BUILD\_SCOPED\_FILE\_URL
- GET\_RELATIVE\_PATH
- GET\_STAGE\_LOCATION

## Explanation

BUILD\_STAGE\_FILE\_URL generates a Snowflake-hosted file URL to a staged file using the stage name and relative file path as inputs. A file URL permits prolonged access to a specified file. That is, the file URL does not expire. File URL: URL that identifies the database, schema, stage, and file path to a set of files. A role that has sufficient privileges on the stage can access the files.

Question 69:

Skipped

How can you produce a lateral view of a VARIANT, OBJECT or ARRAY Column?

- 
- Using SPLIT\_TO\_TABLE table function**
- 
- Using RESULT\_SCAN table function**
- 
- Using INFER\_SCHEMA table function**
- 
- Using FLATTEN table function**

(Correct)

## Explanation

FLATTEN is a table function that produces a lateral view of a VARIANT, OBJECT, or ARRAY column. INFER\_SCHEMA table function is used to detect the file metadata schema in a set of staged data files that contain semi-structured data and retrieves the column definitions. RESULT SCAN returns the result set of a previous command (within 24 hours of when you executed the query) as if the result was a table. SPLIT\_TO\_TABLE table function splits a string (based on a specified delimiter) and flattens the results into rows.

Question 70:

Skipped

What is the expiration period of a Pre-signed URL?

- 
- Length of time specified in the expiration\_time argument**
- 
- It is Permanent**
- 
- The URL expires when the persisted query result period ends**

## Explanation

The expiration period of Scoped URL: The URL expires when the persisted query result period ends. The expiration period of the File URL: It is permanent. The expiration period of Pre-Signed URL: Length of time specified in the expiration\_time argument.

Question 71:

**Skipped**

**Python UDFs and tabular Python UDFs can read and process unstructured data in staged files using SnowflakeFile class. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

**Explanation**

Java UDFs and tabular Java UDFs can read and process unstructured data in staged files using either the SnowflakeFile class or the InputStream class in the UDF code.

Question 72:

**Skipped**

**The automatic refresh of metadata of the directory table in the cloud storage does not incur any charges to Snowflake Customers. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

**Explanation**

Snowflake customers' charges include an overhead to manage event notifications for automatically refreshing directory table metadata. This overhead increases in relation to the number of files added in cloud storage for customers' stages that include directory tables. Snowflake charges 0.06 credits per 1000 event notifications received.

Question 73:

**Skipped**

**Snowflake supports various actions for Unstructured Data. Which one is not supported by Snowflake?**

- 

**Extract actual data from PDF and load it into Snowflake tables using Snowflake WebUI out of the box option.**

**(Correct)**

- Load file access URLs and other file metadata into Snowflake tables.
- Securely access data files located in cloud storage.
- Share file access URLs with collaborators and partners.

### Explanation

Snowflake supports the following actions for Unstructured data:

- Securely access data files located in cloud storage.
- Share file access URLs with collaborators and partners.
- Load file access URLs and other file metadata into Snowflake tables.

Question 74:

### Skipped

David wants to load a JSON file using the COPY INTO <table> command. He found that there are null values in the data for missing values and have no other special meaning. What file format option would you recommend him to use with COPY INTO <table> command to handle the JSON null values?

- David should use REPLACE\_INVALID\_CHARACTERS = TRUE
- David should use STRIP\_OUTER\_ELEMENT = TRUE
- David should use STRIP\_NULL\_VALUES = TRUE  
**(Correct)**
- David should use STRIP\_NULL\_VALUES = FALSE
- David should use STRIP\_OUTER\_ELEMENT = FALSE

### Explanation

In a VARIANT column, NULL values are stored as a string containing the word “null,” not the SQL NULL value. If the “null” values in your JSON documents indicate missing values and have no other special meaning, you should recommend setting the file format option STRIP\_NULL\_VALUES to TRUE for the COPY INTO <table> command when loading the JSON files. Retaining the “null” values often wastes storage and slows query processing.

Question 75:

### Skipped

The Snowflake Information Schema includes table functions you can query to retrieve information about your directory tables. Which table function can be used to query the history of data files registered in the metadata of specified objects and the credits billed for these operations?

- 
- AUTO\_REFRESH\_REGISTRATION\_HISTORY**
- (Correct)**
- 
- STAGE\_DIRECTORY\_FILE\_REGISTRATION\_HISTORY**
- 
- STAGE\_STORAGE\_USAGE\_HISTORY**
- 
- DATABASE\_REFRESH\_HISTORY**

#### Explanation

AUTO\_REFRESH\_REGISTRATION\_HISTORY table function can be used to query the history of data files registered in the metadata of specified objects and the credits billed for these operations. The table function returns the billing history within a specified date range for your entire Snowflake account. This function returns billing activity within the last 14 days. Please note, STAGE\_DIRECTORY\_FILE\_REGISTRATION\_HISTORY table function can be used to query information about the metadata history for a directory table, including:

- Files added or removed automatically as part of a metadata refresh.
- Any errors found when refreshing the metadata.

Question 76:

#### Skipped

There are multiple ways you can trigger the loading of files from the stage into Snowpipe. **Select two ways** generally used to trigger the loading with Snowpipe.

- Calling Snowpipe REST endpoints**
- (Correct)**
- Automating Snowpipe using cloud messaging (notification)**
- (Correct)**
- By executing START PIPE <pipe\_name>**
-

## Calling Snowpipe SOAP endpoints

### Explanation

Automating Snowpipe using cloud messaging (notification) and Calling Snowpipe REST endpoints are the mostly uses methods for triggering loading with Snowpipe.

Question 77:

### Skipped

John is trying to load JSON data sets with a huge array containing multiple records. Considering the VARIANT data type imposed size of 16 MB, what do you recommend to John for optimally loading the data?

- 

Separate the documents with line break or commas

- 

No need to remove the outer array structure as Snowflake Intelligent Engine will take care of that

- 

Enable the STRIP\_OUTER\_ARRAY file format option for the COPY INTO <table> command

(Correct)

- 

Enable VARIANT\_OUTER\_ARRAY file format for the COPY INTO <table> command

### Explanation

If the data exceeds 16 MB, enable the STRIP\_OUTER\_ARRAY file format option for the COPY INTO <table> command to remove the outer array structure and load the records into separate table rows: copy into <table> from @~/<file>.json file\_format = (type = 'JSON' strip\_outer\_array = true);

Question 78:

### Skipped

Which of these SQL functions does Snowflake support?

- 

Scalar

(Correct)

- 

Window

(Correct)

- 

Aggregate

**(Correct)**

- 

**User-Defined**

**(Correct)**

- 

**Table**

**(Correct)**

- 

**System**

**(Correct)**

### **Explanation**

Snowflake Supports all these SQL functions.

Question 79:

**Skipped**

**What would you create (UDF or Stored procedure) if you need a function that can be called as part of a SQL statement and must return a value that will be used in the statement?**

- 

**Stored Procedure**

- 

**UDF**

**(Correct)**

### **Explanation**

An UDF evaluates to a value, and can be used in contexts in which a general expression can be used (e.g. SELECT my\_function() ...). A stored procedure does not evaluate to a value, and cannot be used in all contexts in which a general expression can be used. For example, you cannot execute SELECT my\_stored\_procedure()....

Question 80:

**Skipped**

**Which data does not fit into a predefined data model or schema?**

- 

**Semi-Structured Data**

- 

**All of these**

- 

**Structured-data**

- 

**Unstructured Data**

**(Correct)**

### **Explanation**

Unstructured data is information that does not fit into a predefined data model or schema. Typically text-heavy, such as form responses and social media conversations, unstructured data also encompasses images, video, and audio. Industry-specific file types such as VCF (genomics), KDF (semiconductors), or HDF5 (aeronautics) are included in this category.

Question 81:

**Skipped**

**UDF does not support sql DDL / DML? (True/Fales)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### **Explanation**

UDF does not support sql DDL / DML. That means you can select from a table, but you can't create or modify tables inside of a UDF.

Question 82:

**Skipped**

**What is the expiration period of a Scoped URL?**

- 

**Length of time specified in the expiration\_time argument**

- 

**The URL expires when the persisted query result period ends**

**(Correct)**

- 

**The URL never expires. It is permanent**

### **Explanation**

The expiration period of Scoped URL: The URL expires when the persisted query result period ends. The expiration period of the File URL: It is permanent. The expiration period of Pre-Signed URL: Length of time specified in the expiration\_time argument.

Question 83:

Skipped

What is the recommended compressed size of data files for optimal bulk data loads?

- 10-50 MB
- 10-50 GB
- 100-250 GB
- 100-250 MB

(Correct)

### Explanation

The number of load operations that run in parallel cannot exceed the number of data files to be loaded. To optimize the number of parallel operations for a load, we recommend aiming to produce data files roughly 100-250 MB (or larger) in size compressed.

Question 84:

Skipped

When staging uncompressed files in a Snowflake stage, Snowflake automatically compresses the files unless compression is explicitly disabled. Which of the options is used by Snowflake for compressing the file?

- bzip2
- gzip
- (Correct)
- Zstandard
- deflate
-

## Brotli

### Explanation

When staging uncompressed files in a Snowflake stage, the files are automatically compressed using gzip, unless compression is explicitly disabled.

Question 85:

Skipped

What is an "object" in semi-structured data?

- 
- Collection of semi-structure data files**
- 
- Collection of key-value pairs**
- (Correct)**
- 
- Collection of values in an array**

### Explanation

OBJECT is also called a “dictionary”, “hash”, or “map” in many languages. This contains key-value pairs.

Question 86:

Skipped

What size of the virtual warehouse needs to be created by the sysadmin while loading using Snowpipe? (Select the best answer)

- 
- 4XL Size**
- 
- L Size**
- 
- M Size**
- 
- None of these**
- (Correct)**
- 
- XS Size**

### Explanation

Snowpipe uses compute resources provided by Snowflake (i.e. a serverless compute model). These Snowflake-provided resources are automatically resized and scaled up or down as required, and are charged and itemized using per-second billing. Data ingestion is charged based upon the actual workloads. User doesn't need to create any warehouse as it is taken care by Snowflake.

Question 87:

**Skipped**

**What of the following is the default character set for delimited files (CSV, TSV, etc.)?**

- 
- UTF-32LE**
- 
- UTF-8**
- (Correct)**
- 
- UTF-16BE**
- 
- UTF-32**
- 
- UTF-16**

**Explanation**

UTF-8 is the default character set for delimited files (CSV, TSV, etc.). Snowflake also supports all others in the provided options, but you must explicitly specify the encoding to use for loading.

Question 88:

**Skipped**

**What authentication methods does Snowflake support for REST API authentication? (Select 2)**

- Key Pair Authentication**
- (Correct)**
- Authentication is not required in case Snowflake SQL API**
- Snowflake Account User ID and Password**
- OAuth**

**(Correct)**

### Explanation

Snowflake SQL API supports **Oauth, and Key Pair authentication**.

Question 89:

**Skipped**

Which of these SQL functions helps returns the absolute path of a staged file using the stage name and path of the file relative to its location in the stage as inputs.?

- 
- **BUILD\_SCOPED\_FILE\_URL**
- 
- **GET\_PRESIGNED\_URL**
- 
- **GET\_RELATIVE\_PATH**
- 
- **BUILD\_STAGE\_FILE\_URI**
- 
- **GET\_STAGE\_LOCATION**
- 
- **GET\_ABSOLUTE\_PATH**

**(Correct)**

### Explanation

GET\_ABSOLUTE\_PATH returns the absolute path of a staged file using the stage name and path of the file relative to its location in the stage as inputs.

Question 90: **Incorrect**

If you recreate a pipe using CREATE OR REPLACE PIPE command. What does happen to load history if the Snowpipe gets recreated?

- 
- **The recreated Pipe still has tracks of the files loaded by the old Pipe**
- 
- **Snowflake still keeps load history**

**(Incorrect)**

**The load history gets reset to empty**

**(Correct)**

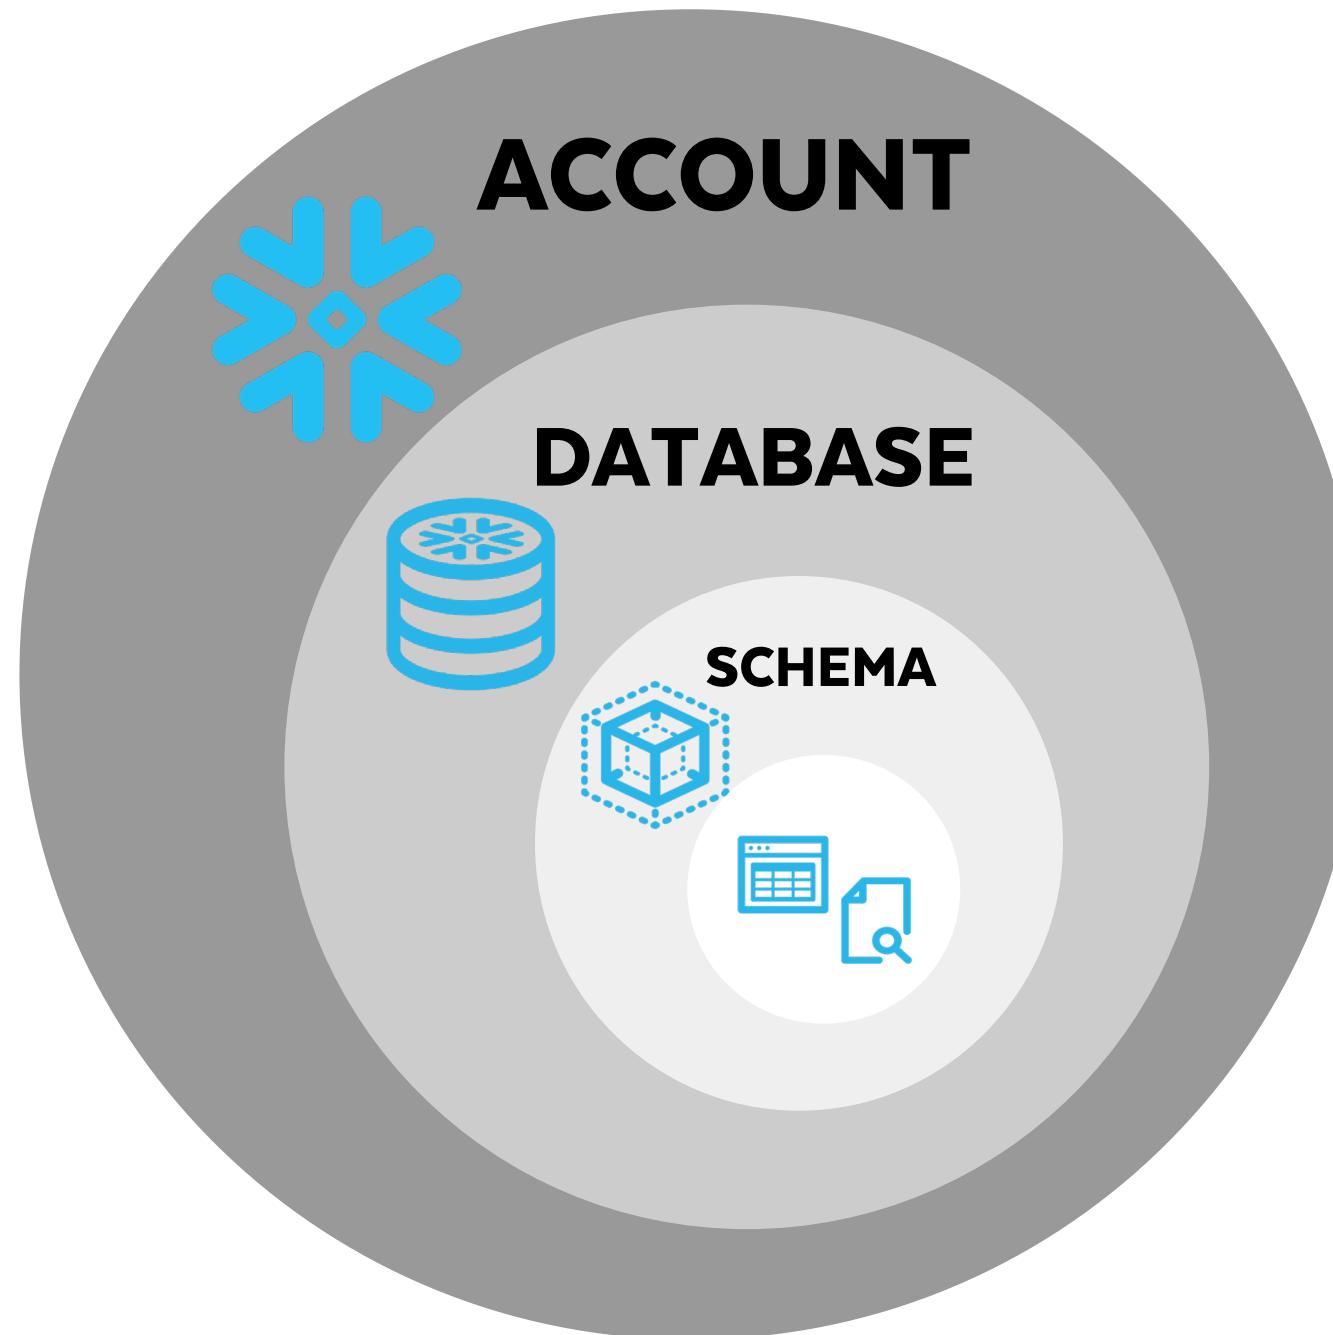
- 

**The pipe can not be recreated**

#### **Explanation**

When you recreate a pipe, if you do CREATE OR REPLACE PIPE, that load history is reset to empty, so Snowflake doesn't know which files we've already loaded.

# STORAGE HIERARCHY



## Snowflake Accounts

- A customer can have as many accounts as they want
- Each account has its own URL – which uses the name of the account.
- Each account is deployed on a single cloud provider platform (AWS, Azure, GCP)
- Each account exists in a single geographic region
- Each account exists with a single Snowflake Edition (e.g. Standard, Premier, etc.)

## Accounts Contain Databases

- Each database belongs to a single Snowflake account
- Databases can be replicated to other accounts, but they cannot SPAN multiple accounts

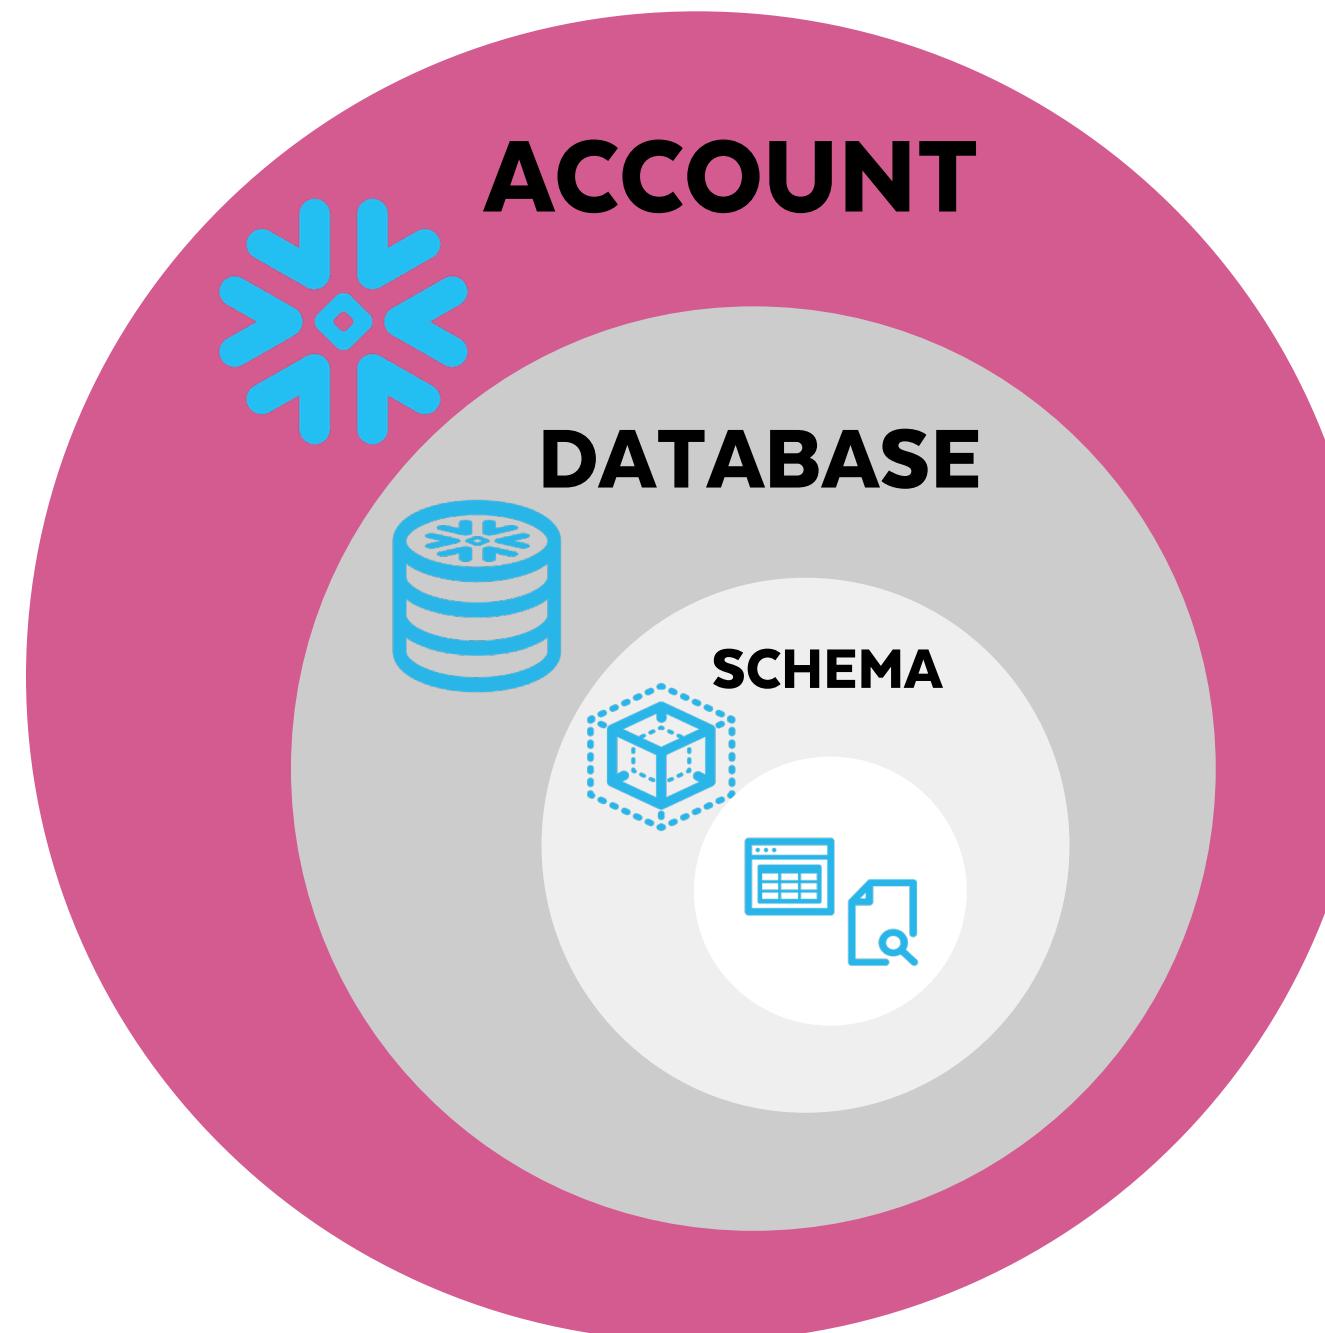
## Databases Contain Schemas

- Each schema belongs to a single Snowflake database, in a single account
- Schemas can be replicated to other accounts or databases, but they cannot SPAN accounts or databases

## Schemas Contain Other Objects

- Objects include tables, views, file formats, sequences, UDFs, & stored procs
- Objects belong to a single schema, in a single database, in a single account

# STORAGE CONTAINER HIERARCHY



**READ MORE ABOUT REGIONS HERE:**

<https://docs.snowflake.net/manuals/user-guide/intro-regions.html>

## Snowflake Accounts

The account shown here is the account NS61174. Most trial accounts start with two letters followed by 5 numbers. Trial accounts can be converted to paid accounts.

Customers can request a “vanity” address like:

<https://robertsfishandchips.snowflakecomputing.com>

When an account has just the account name, followed by “snowflakecomputing.com”, we know the cloud provider is **AWS** and the region is **us-west-2**. This is because this was the first cloud provider and region on which Snowflake deployed.

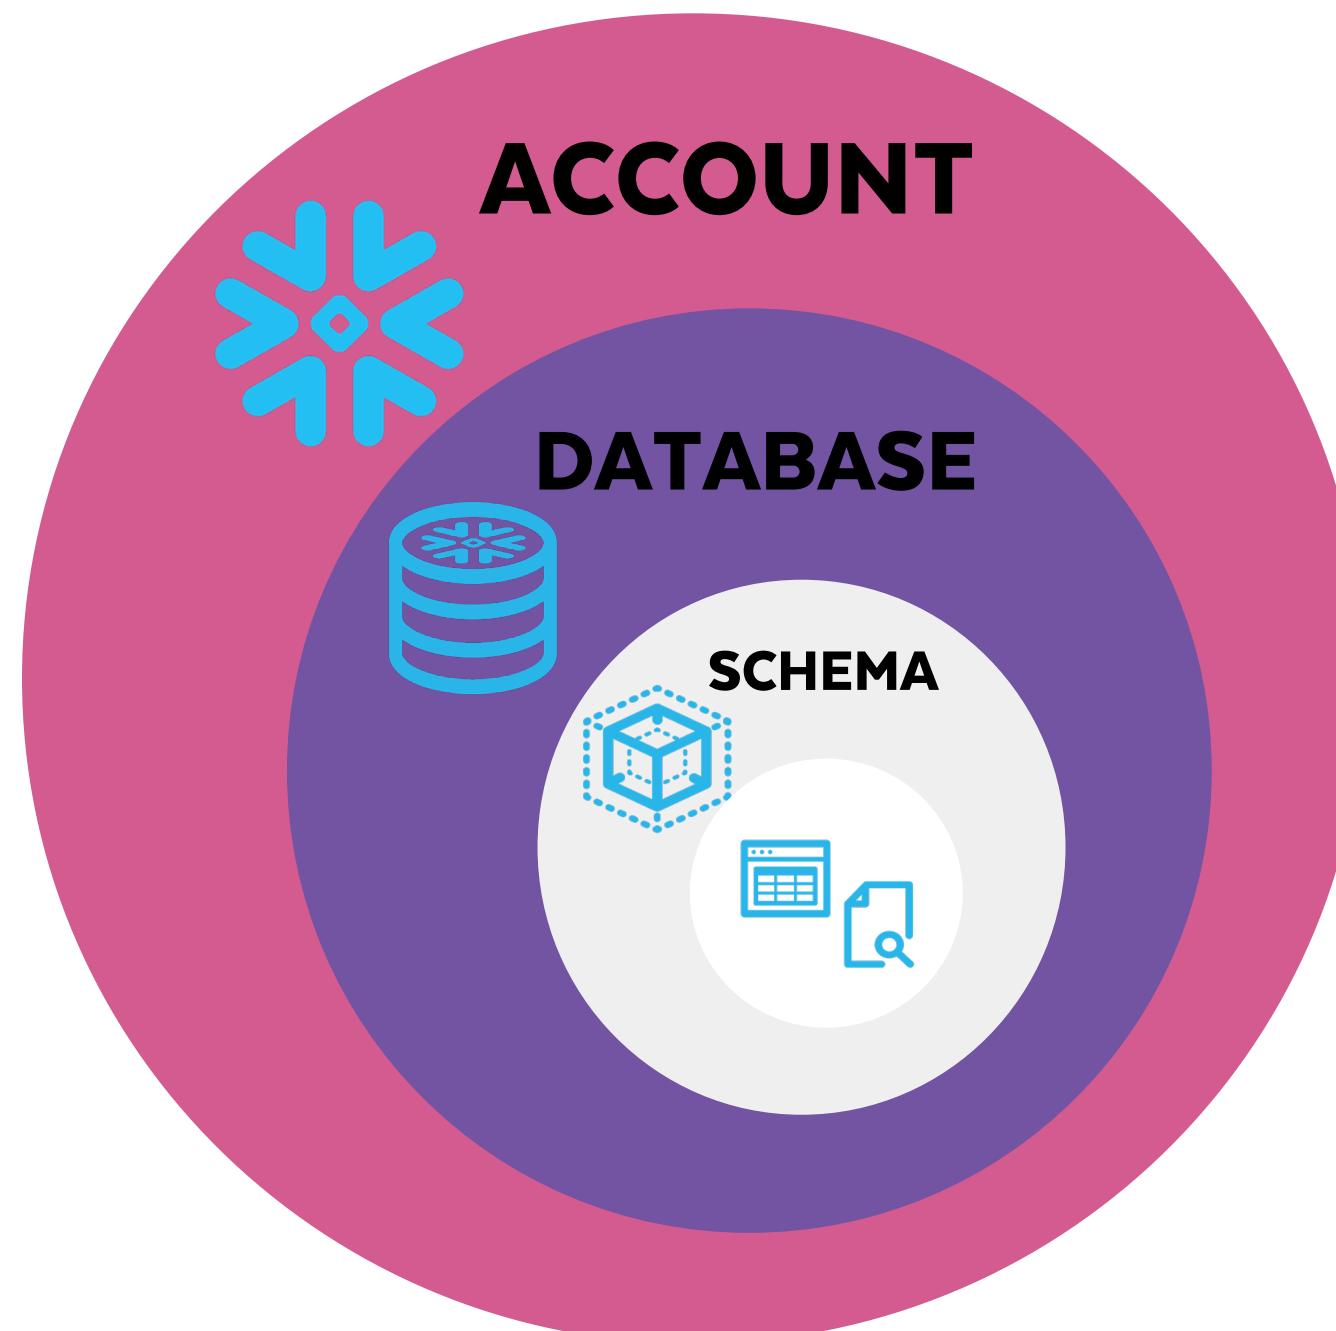
Later, when Snowflake expanded to other **regions** and **providers**, the region name was added to the account URL. For example, an address like:

<https://xy12345.east-us-2.azure.snowflakecomputing.com>

Databases

Create... Clone... Drop... Transfer Ownership

# STORAGE CONTAINER HIERARCHY



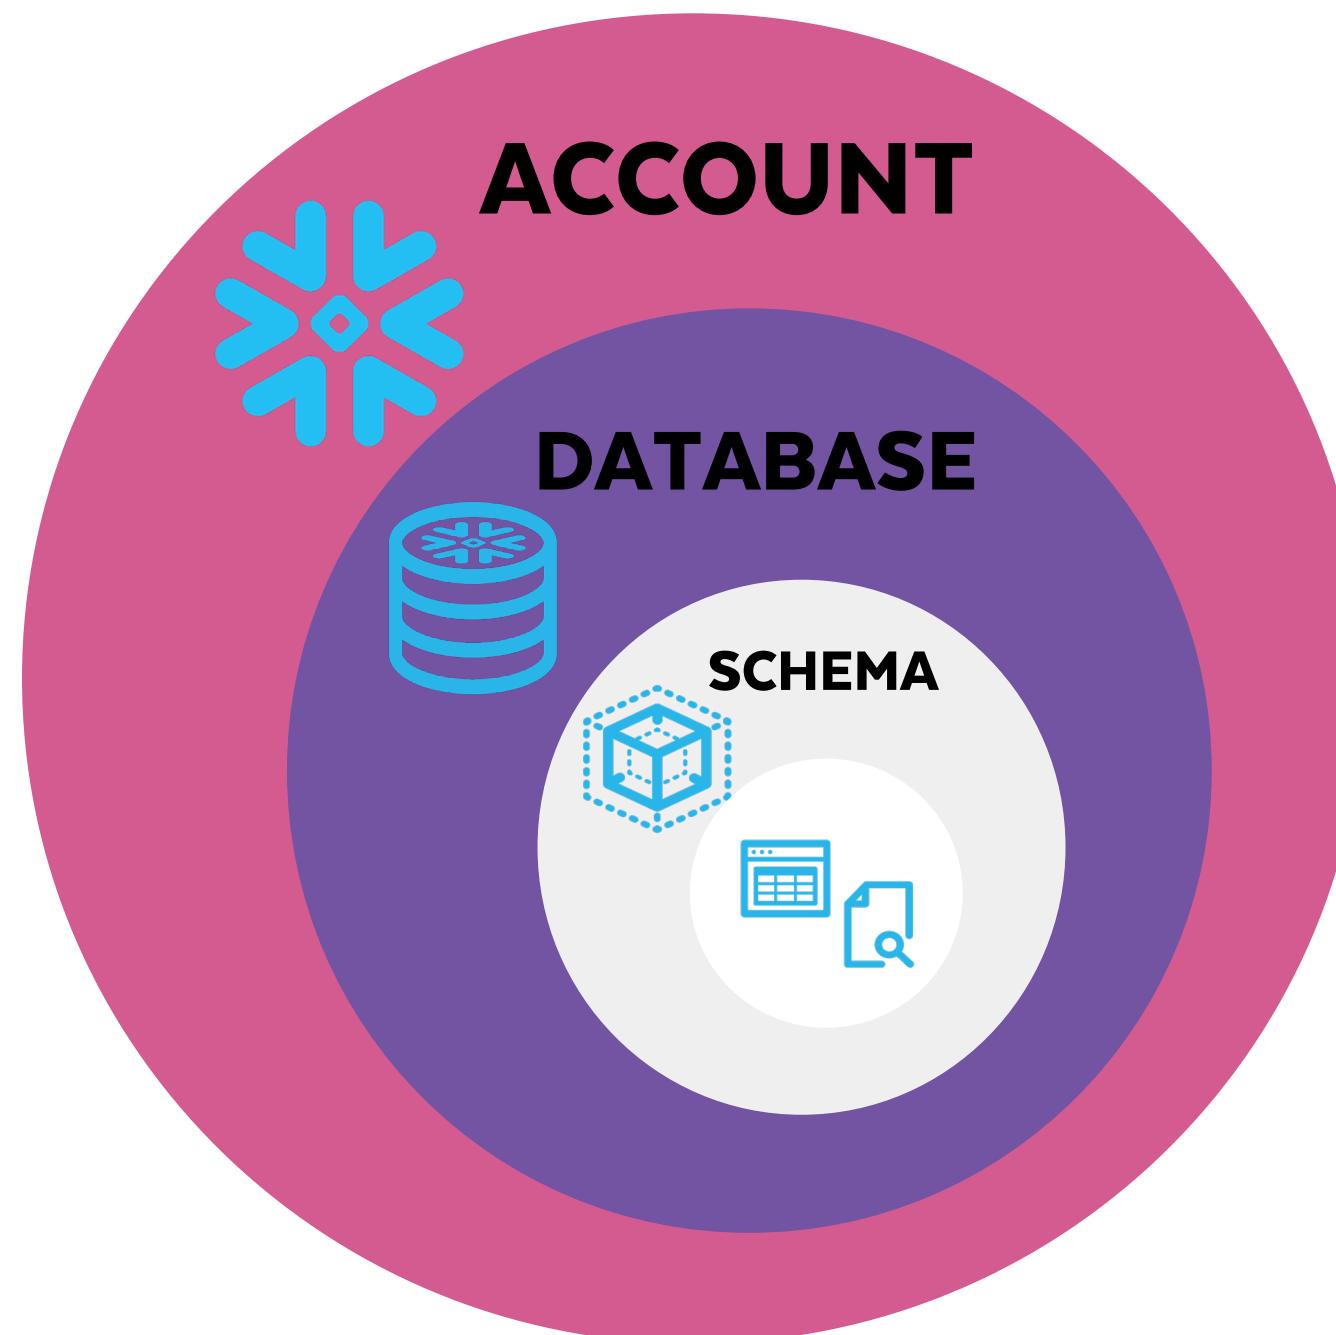
## Accounts Contain Databases

- In this example, account NS61174 has three databases.
- The **SNOWFLAKE** database contains information about account usage. It is automatically added by Snowflake to each new account. It is sometimes called the “Account Usage Share” because it is shared by Snowflake so customers can check usage.
- The **SNOWFLAKE\_SAMPLE\_DATA** is also shared by Snowflake. This database allows for performance testing and tutorials. Customers are not charged for storing this data., because they are NOT storing it. The database is a shared database.

Screenshot of the Snowflake web console showing the 'Databases' page for account ns61174. The URL in the browser bar is ns61174.snowflakecomputing.com/console#/data/databases. The page displays three databases: MY\_DATABASE, SNOWFLAKE\_SAMPLE\_DATA, and SNOWFLAKE. The SNOWFLAKE database is highlighted with a purple box.

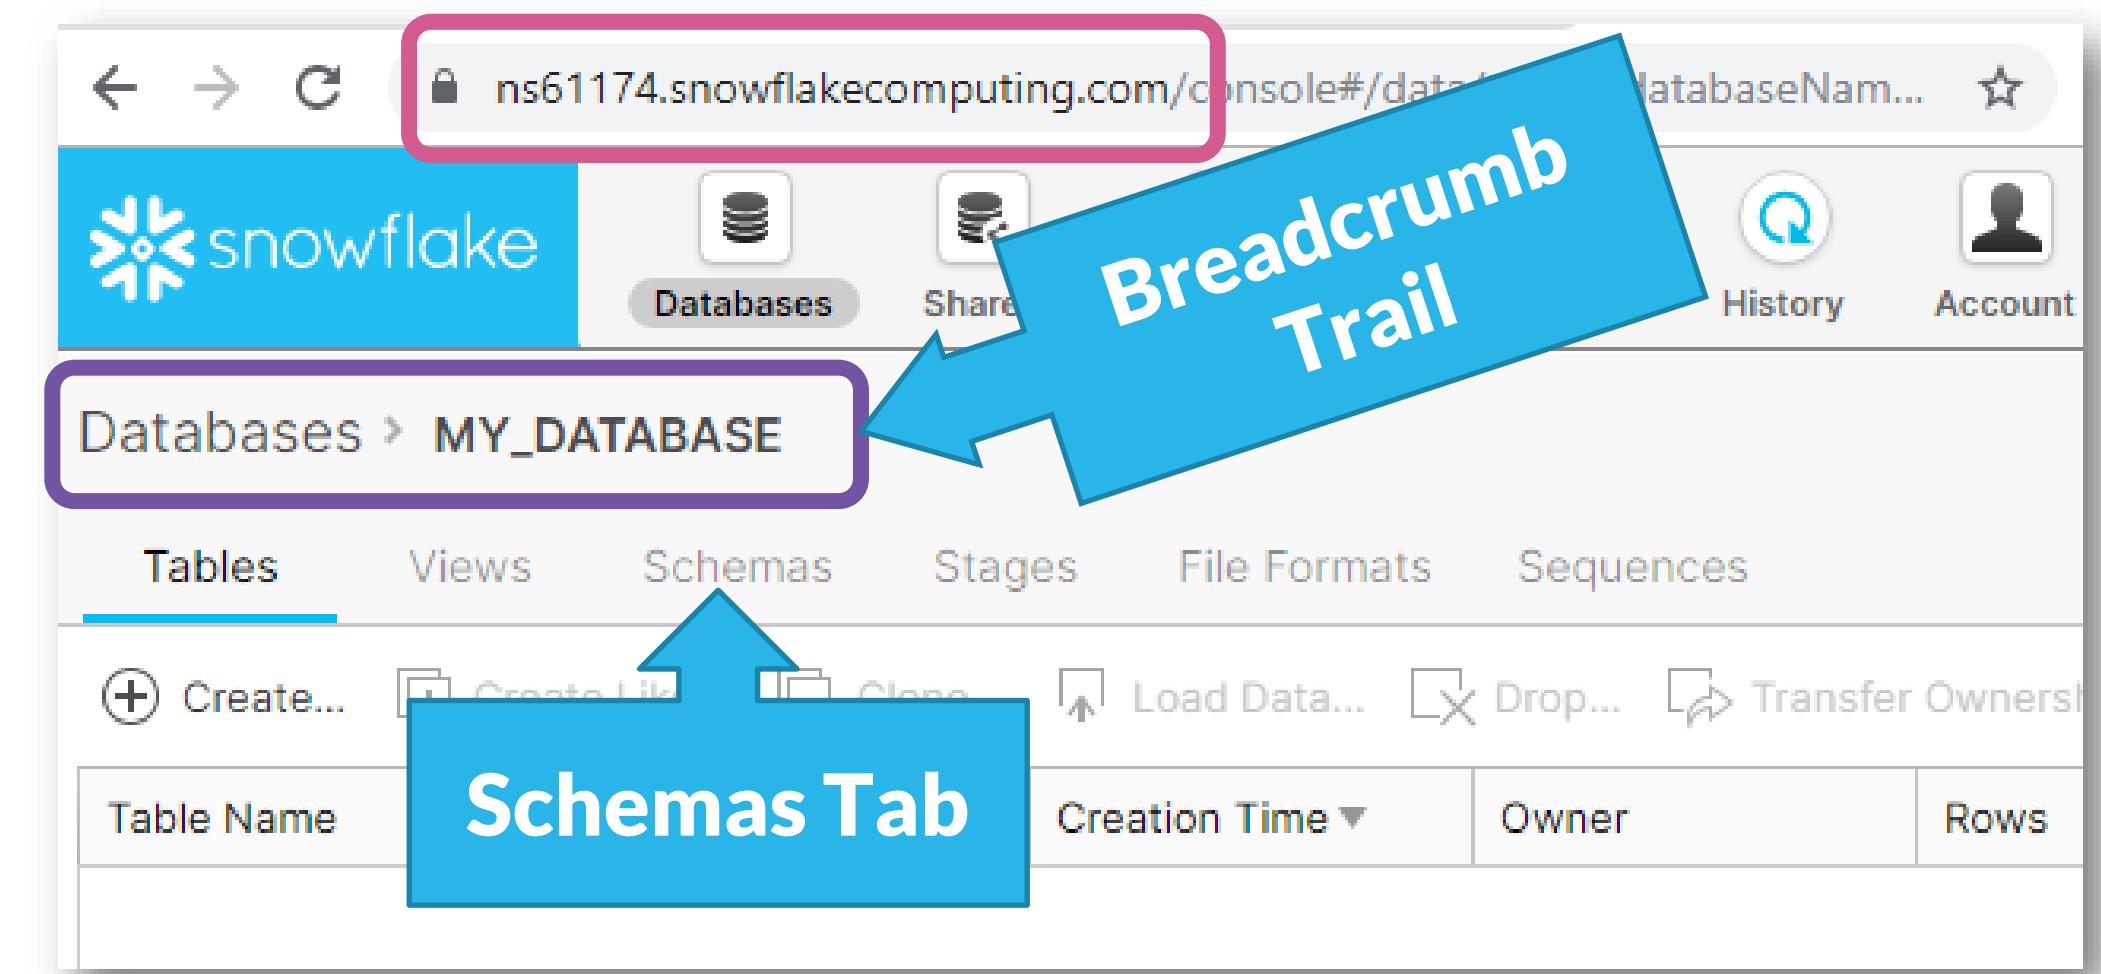
| Database              | Origin             | Creation Time      | Owner        |
|-----------------------|--------------------|--------------------|--------------|
| MY_DATABASE           |                    | 7:06:25 AM         | ACCOUNTADMIN |
| SNOWFLAKE_SAMPLE_DATA | SFC_SAMPLES.SA...  | 6/1/19 10:51:59 PM | ACCOUNTADMIN |
| SNOWFLAKE             | SNOWFLAKE.ACCTO... | 6/1/19 10:51:47 PM |              |

# STORAGE CONTAINER HIERARCHY

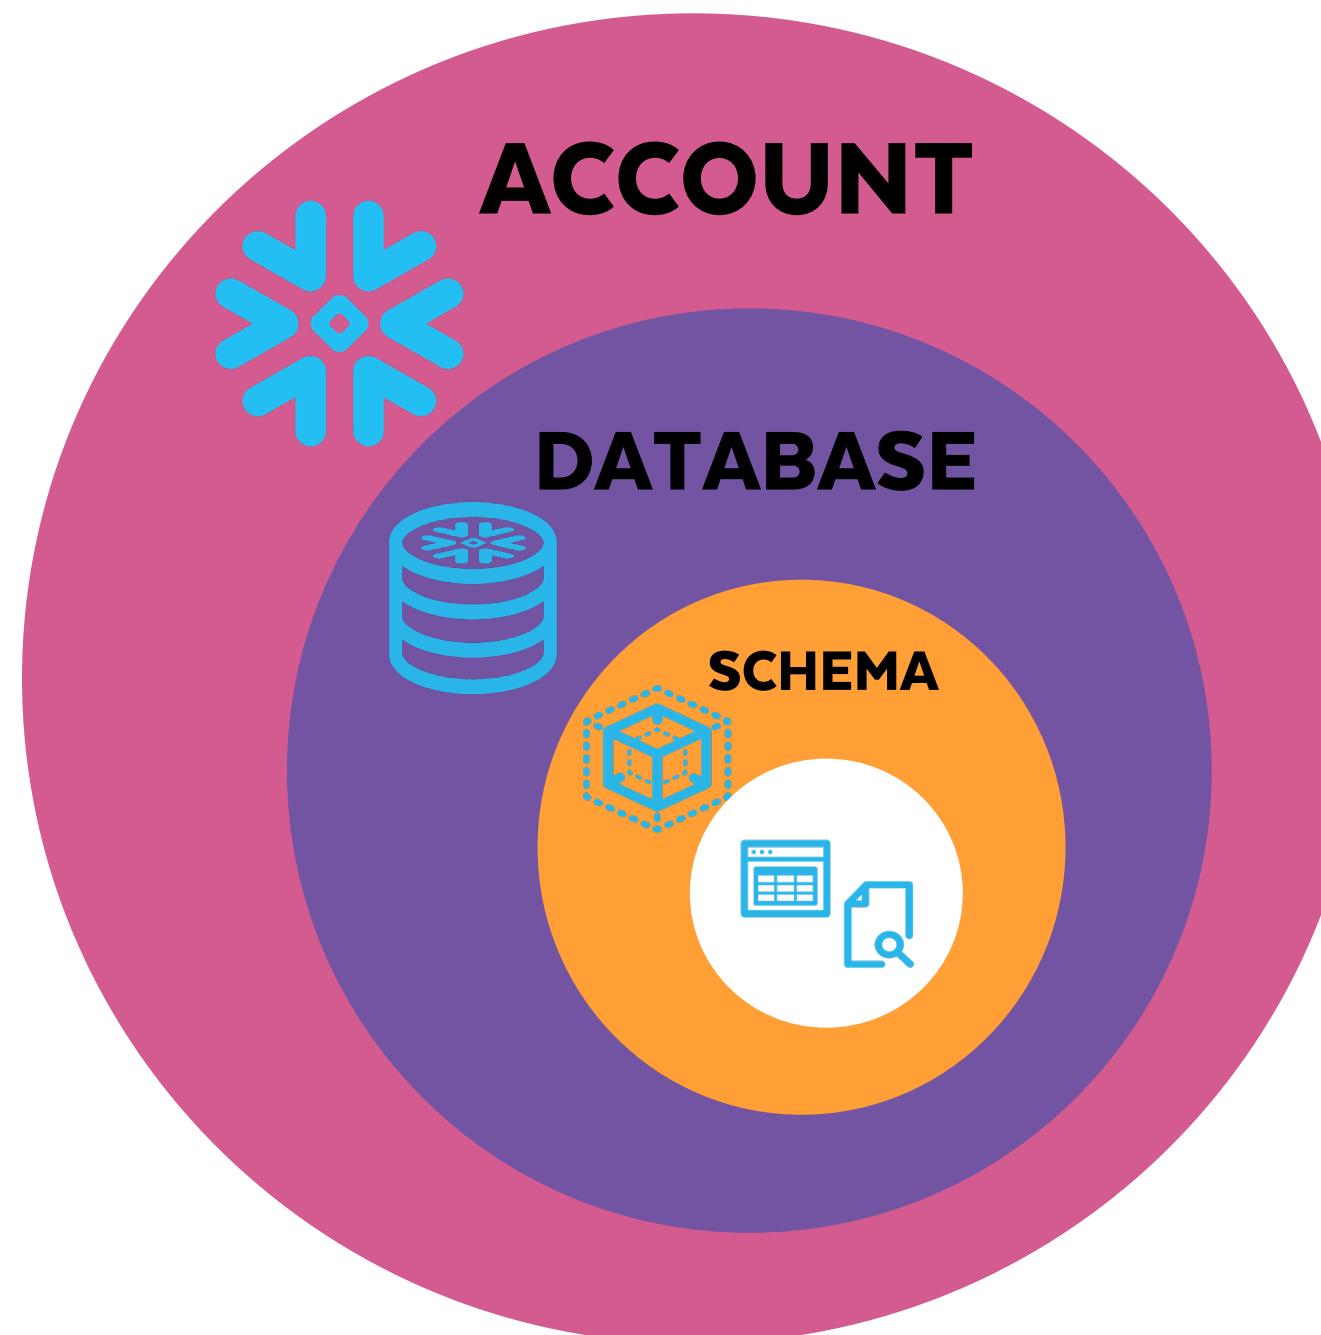


## Accounts Contain Databases

When we drill into any database, the name of the database gets added to the breadcrumb trail and a set of tabs appears showing Schemas and other object types.



# STORAGE CONTAINER HIERARCHY



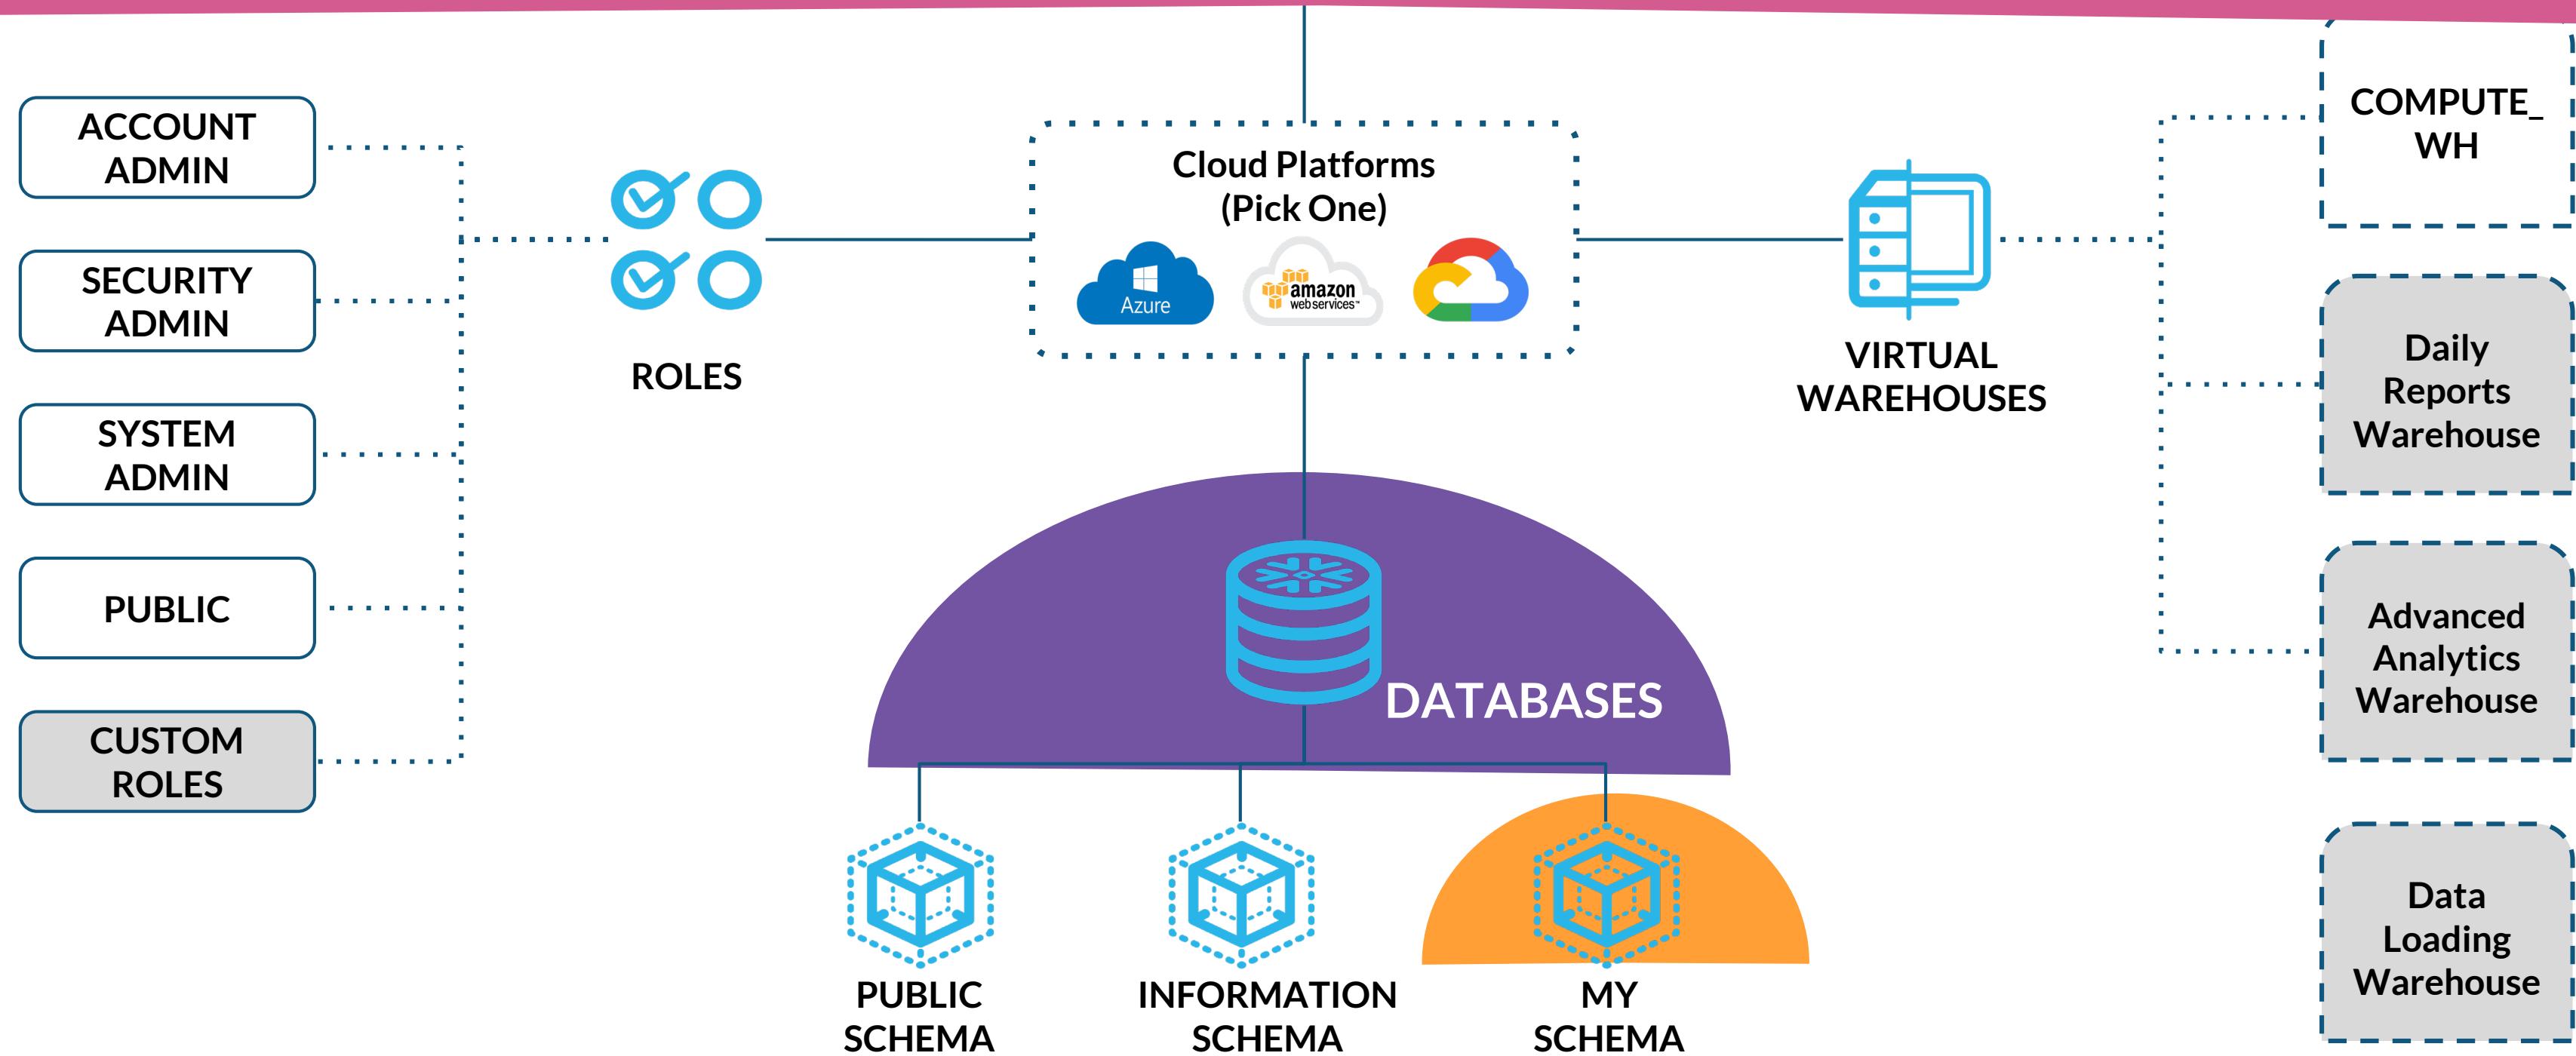
**Database.Schema = Namespace**

## Databases Contain Schemas

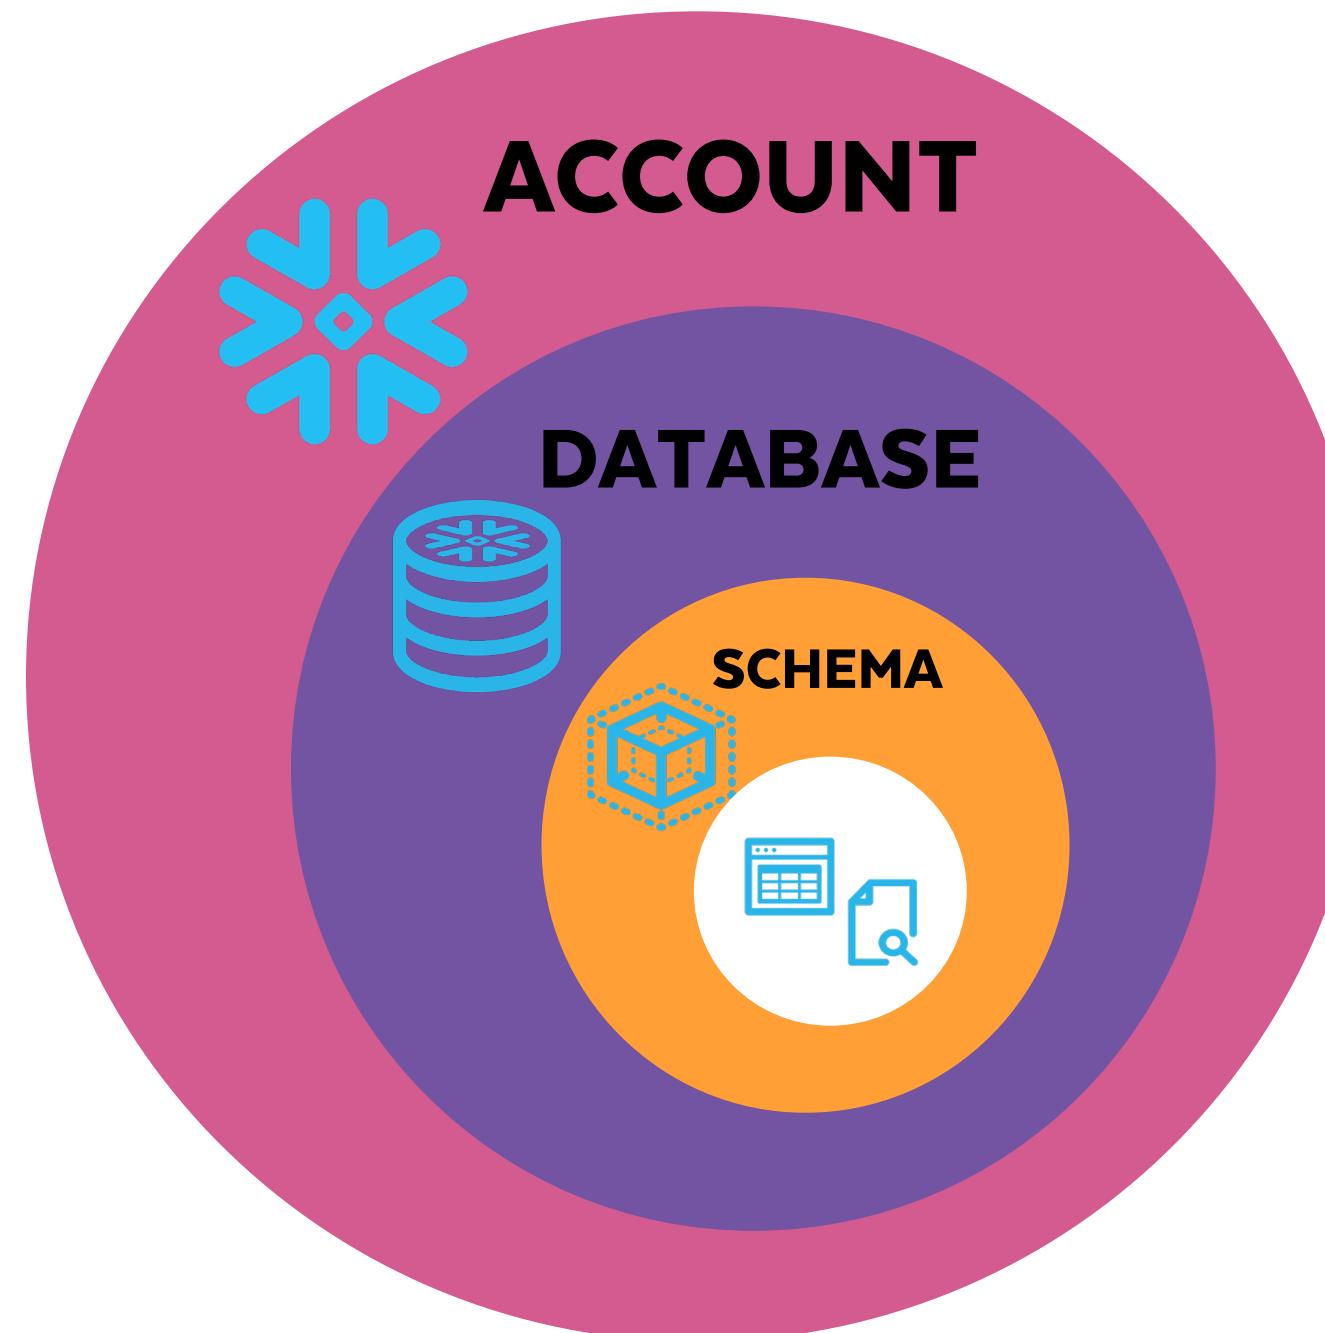
- The set of tabs below the breadcrumb trail makes it look as though Schemas are at the same level as Tables, Views, Stages, File Formats, and Sequence but this is not an accurate way to think of Schemas.
- Schemas are a logical grouping one level above the other objects listed.

| Schema             | Creation Time | Owner        | Managed Access | Comment |
|--------------------|---------------|--------------|----------------|---------|
| INFORMATION_SCHEMA | 7:21:55 AM    |              |                |         |
| MY_SCHEMA          | 7:06:39 AM    | ACCOUNTADMIN |                |         |
| PUBLIC             | 7:06:25 AM    | ACCOUNTADMIN |                |         |

# SNOWFLAKE ACCOUNT



# STORAGE CONTAINER HIERARCHY



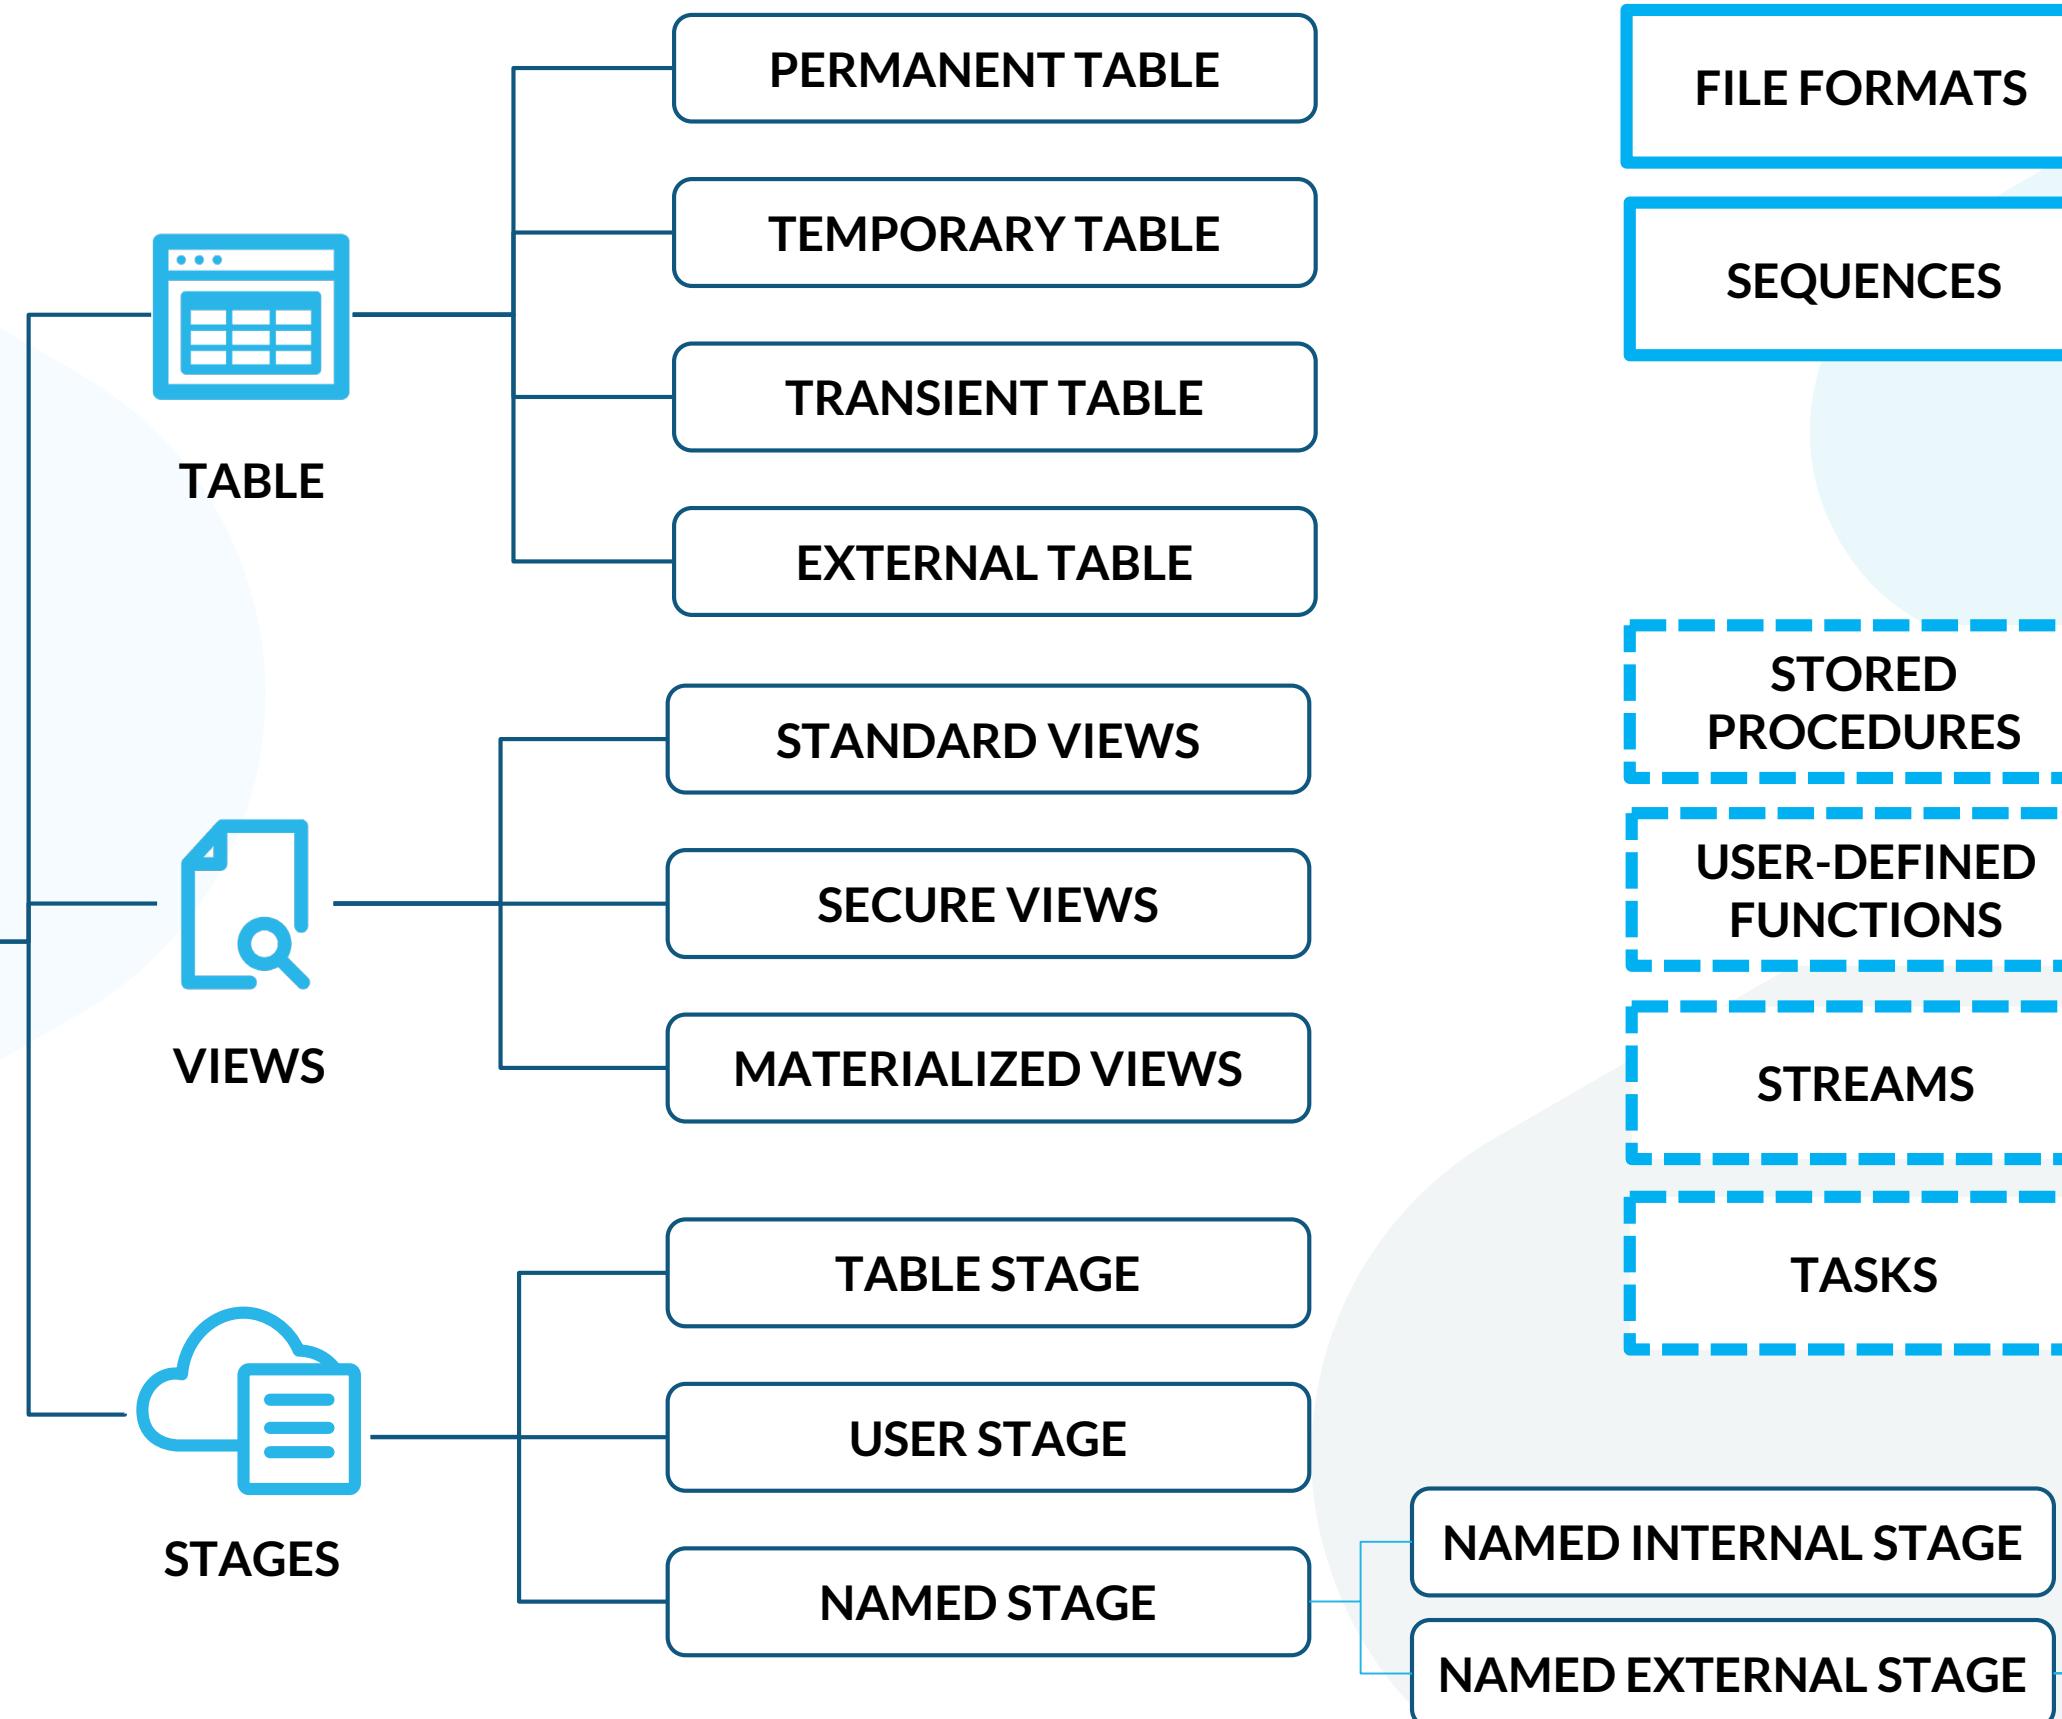
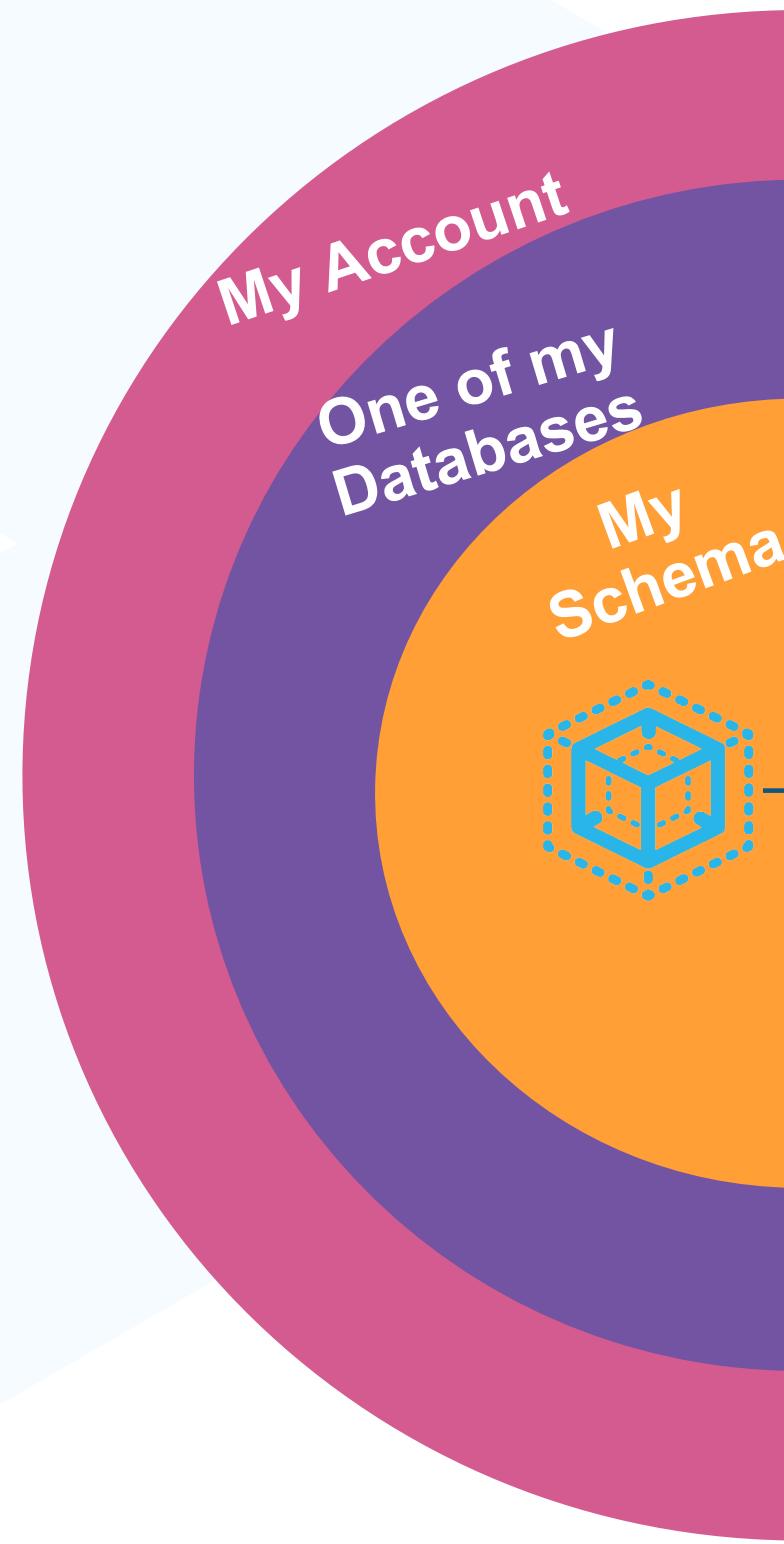
**Tables, Views, Stages, File Formats, and Sequences are situated within Schemas**

- The tables shown below are in different schemas.
- Snowflake allows you to bypass the schema during the drill-down process, making it seem as if tables and other objects are not within a schema, but they are.

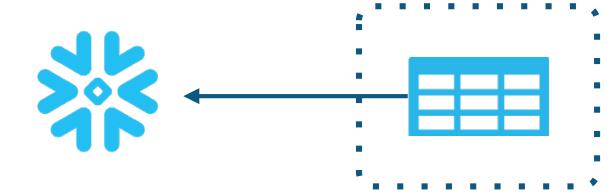
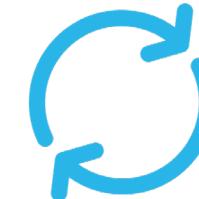
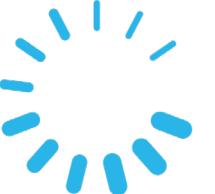
The screenshot shows the Snowflake web interface. At the top, there's a navigation bar with the Snowflake logo, followed by tabs for 'Databases', 'Shares', 'Warehouses', and 'Workshops'. Below this, a breadcrumb navigation shows 'Databases > MY\_DATABASE'. The main area has tabs for 'Tables', 'Views', 'Schemas', 'Stages', 'File Formats', and 'Sequences', with 'Tables' currently selected. Below the tabs is a row of action buttons: '+ Create...', '+ Create Like...', 'Clone...', 'Load Data...', and 'Drop'. A table lists three tables: 'MY\_SECOND\_TABLE' (Schema: MY\_SCHEMA, Creation Time: 9:00:11 AM), 'MY\_FIRST\_TABLE' (Schema: PUBLIC, Creation Time: 8:59:55 AM), and a third unnamed table. The 'Schema' column for the second table is highlighted with an orange border.

| Table Name      | Schema    | Creation Time |
|-----------------|-----------|---------------|
| MY_SECOND_TABLE | MY_SCHEMA | 9:00:11 AM    |
| MY_FIRST_TABLE  | PUBLIC    | 8:59:55 AM    |

# SCHEMA OBJECTS



# TABLE TYPES



## PERMANENT

- Standard, default table type.
- Most commonly used.
- Default Time-Travel
- 7 Day Fail-Safe – non-configurable

## TEMPORARY

- Lasts single session
- Can be used for moving data from one permanent to another with some transformation within the table
- No costs associated with failsafe storage

## TRANSIENT

- No costs associated with failsafe storage
- Usually less costly than permanent tables.
- Time travel up to one day is available

## EXTERNAL

- Stored in External Stages
- Read Only
- DML SQL works: Selects, Aggregations, Joins
- Can use in a View
- No cloning, time travel or fail-safe

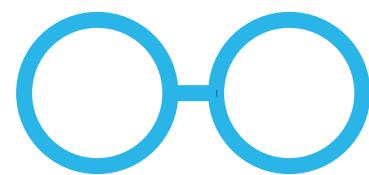
Time-Travel



Fail-Safe

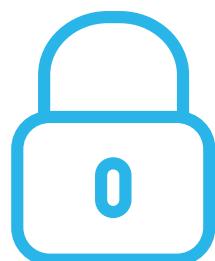


# VIEW TYPES



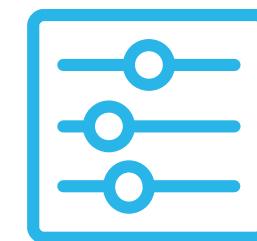
## Standard/Regular View

- Default View type
- Most commonly used
- Underlying DDL available to any role with access to the view



## Secure View

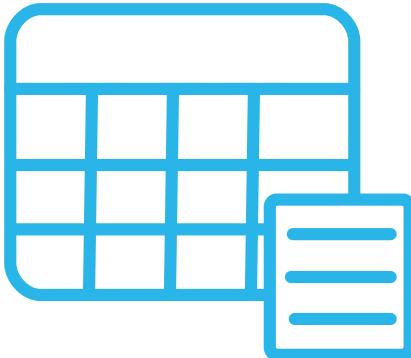
- The underlying query definition and details are only visible to authorized users
- Snowflake query optimizer behaves differently when running the SQL contained in a secure view so that it doesn't accidentally reveal things about the underlying data
- May not run as fast as a standard view, because the optimizer isn't used the same way



## Materialized View

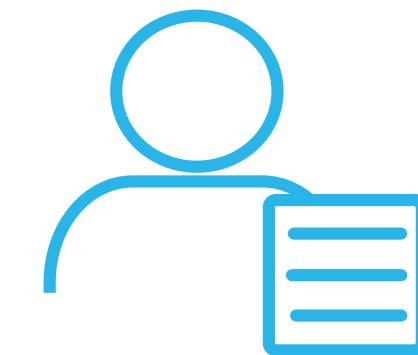
- It has “view” in the name, but it behaves more like a table
- Results are stored just as a table would be
- In Snowflake, these are auto-refreshed
- Can provide cost savings but are inflexible and somewhat limited in use

# STAGE TYPES



## Table Stage @%[TABLE\_NAME]

- Automatically defined – No set up needed
- Not viewable on Stages tab of WebUI
- You must specify certain file format-like parameters as part of the command during each COPY INTO event

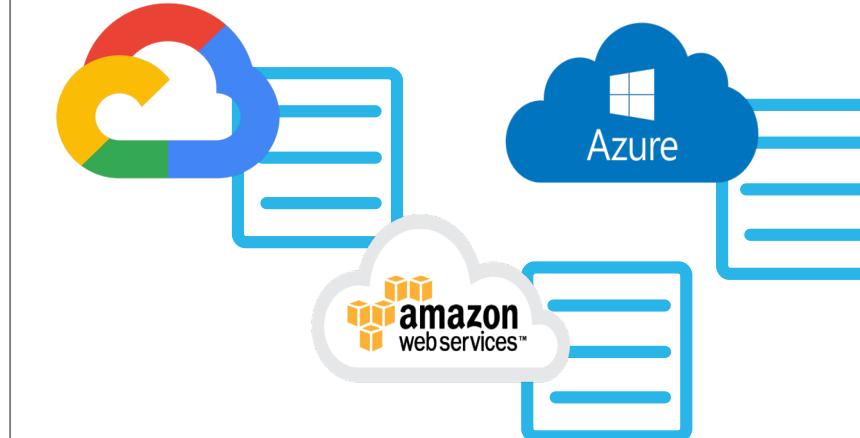


## User Stage @~[LOGIN]



## Named Internal Stage @[STAGE\_NAME]

- Definable via the WebUI
- More flexible than Table and User Stages
- Viewable on Stages tab of Database Area of WebUI
- Have embedded default File Format definitions or you can specify named File Formats for use



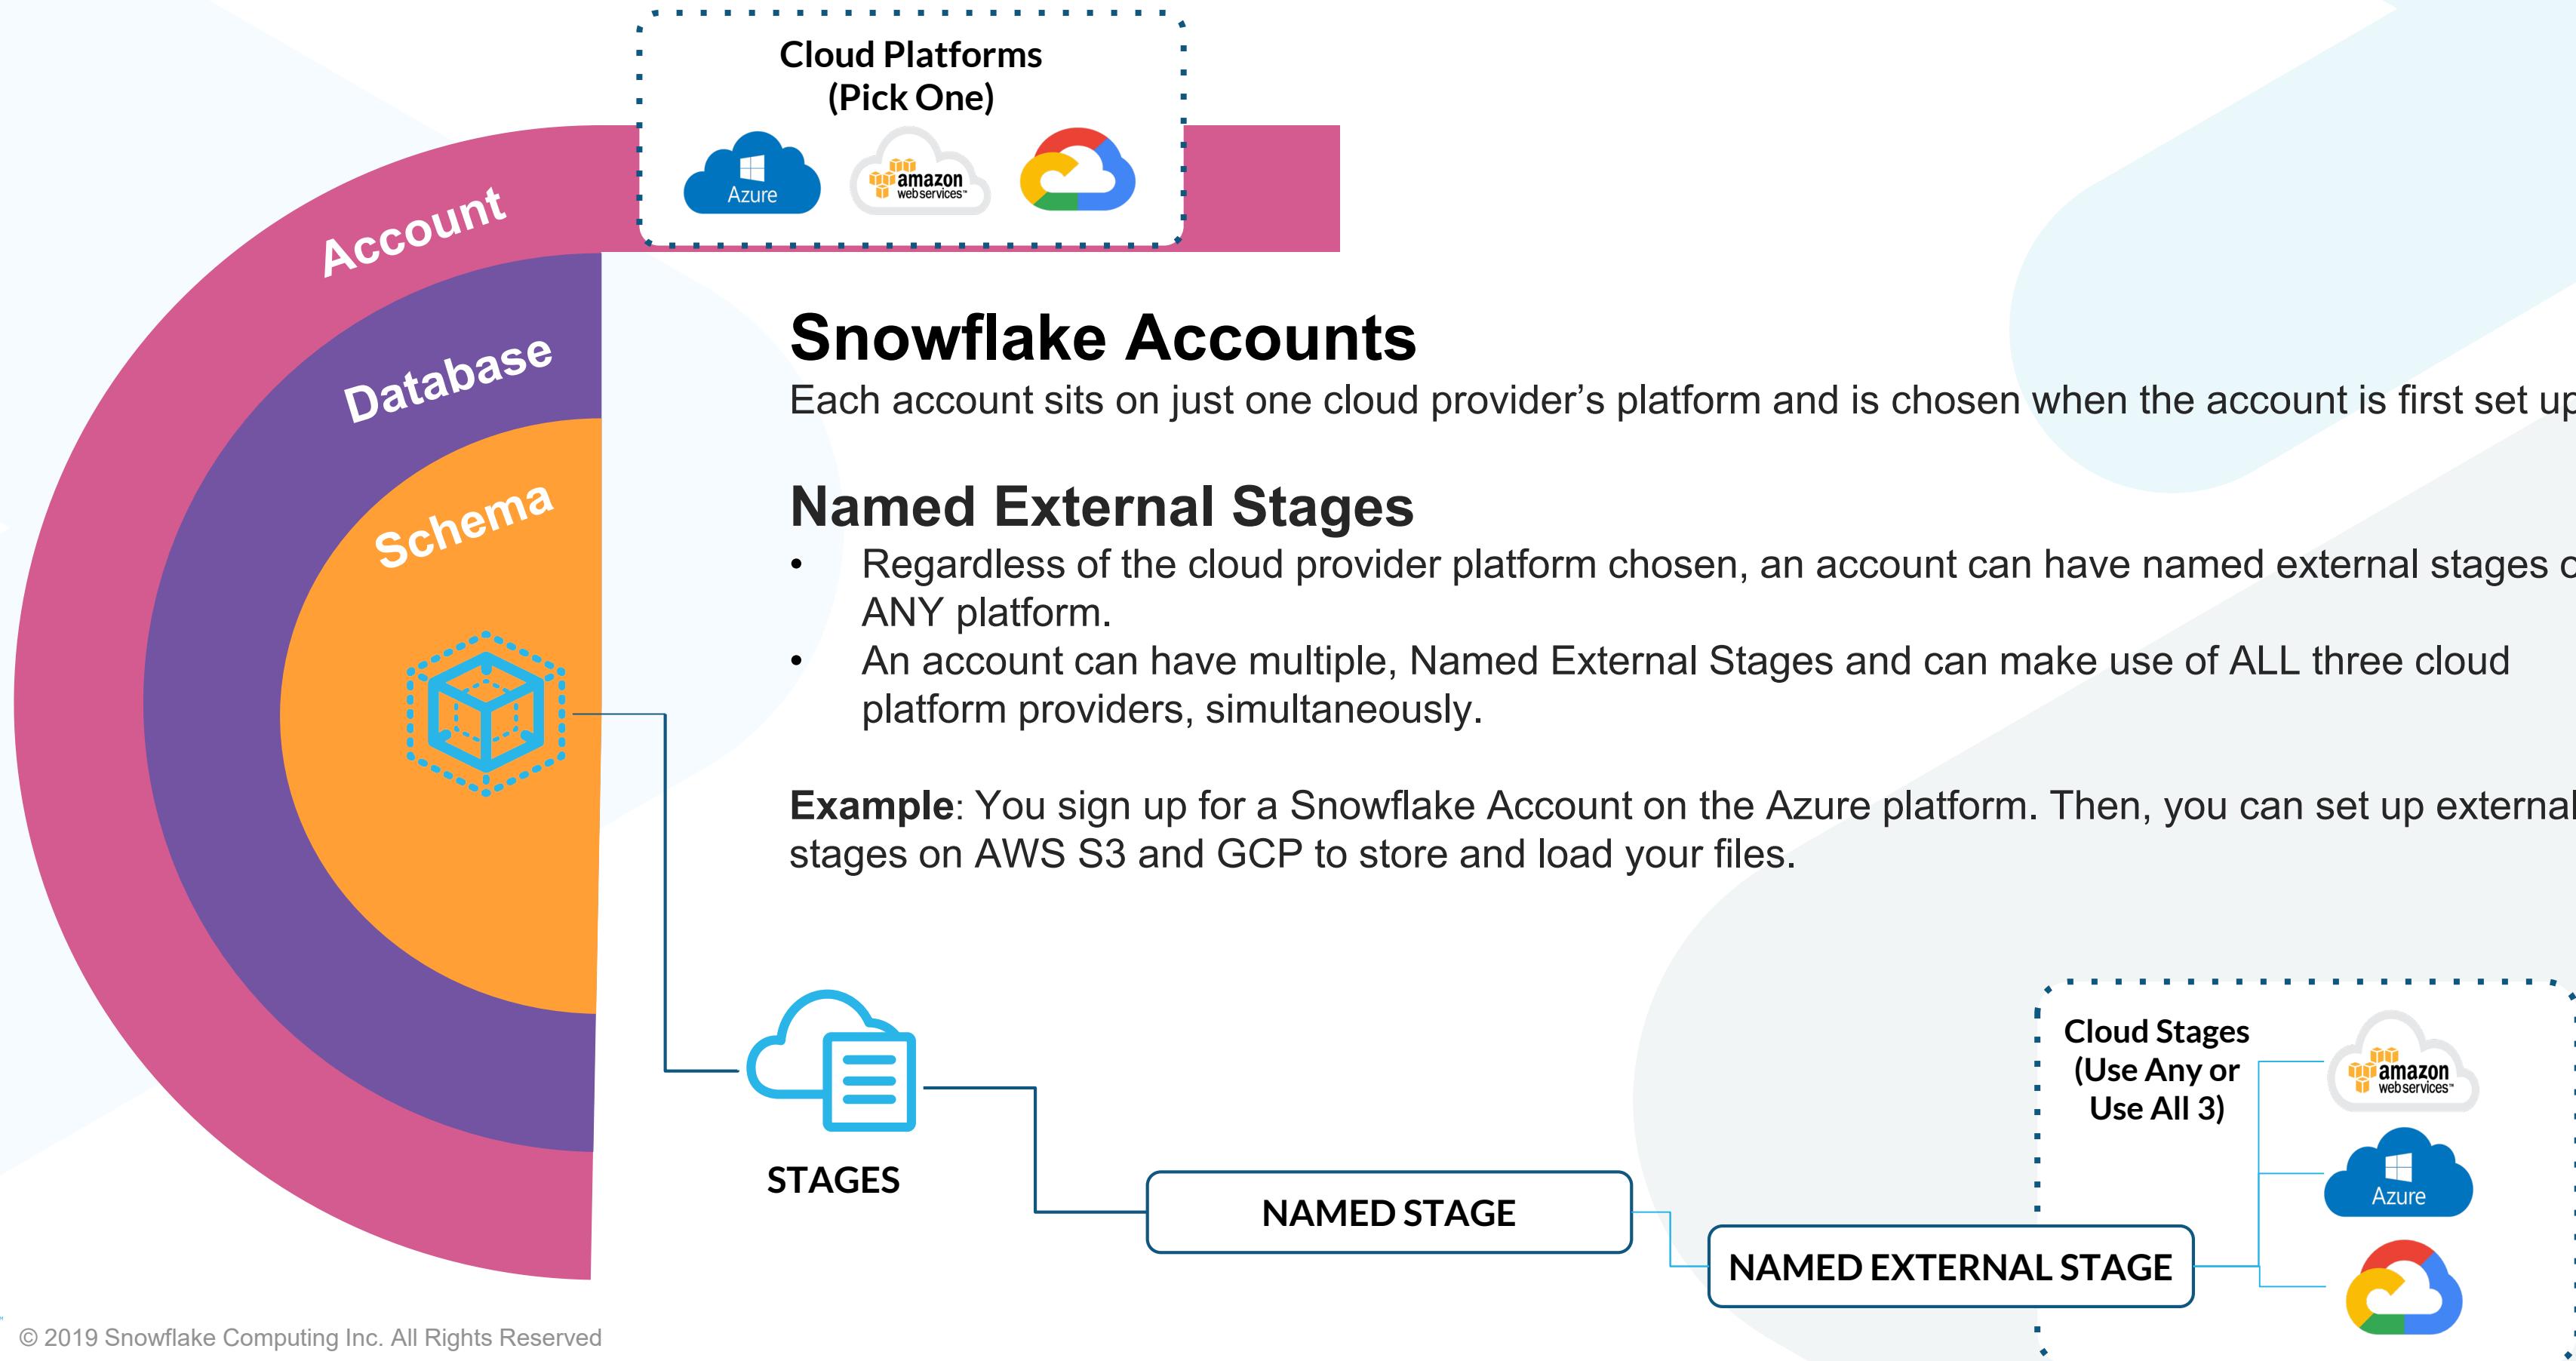
## Named External Stage @[STAGE\_NAME]

### GETTING FILES INTO AND OUT OF STAGES

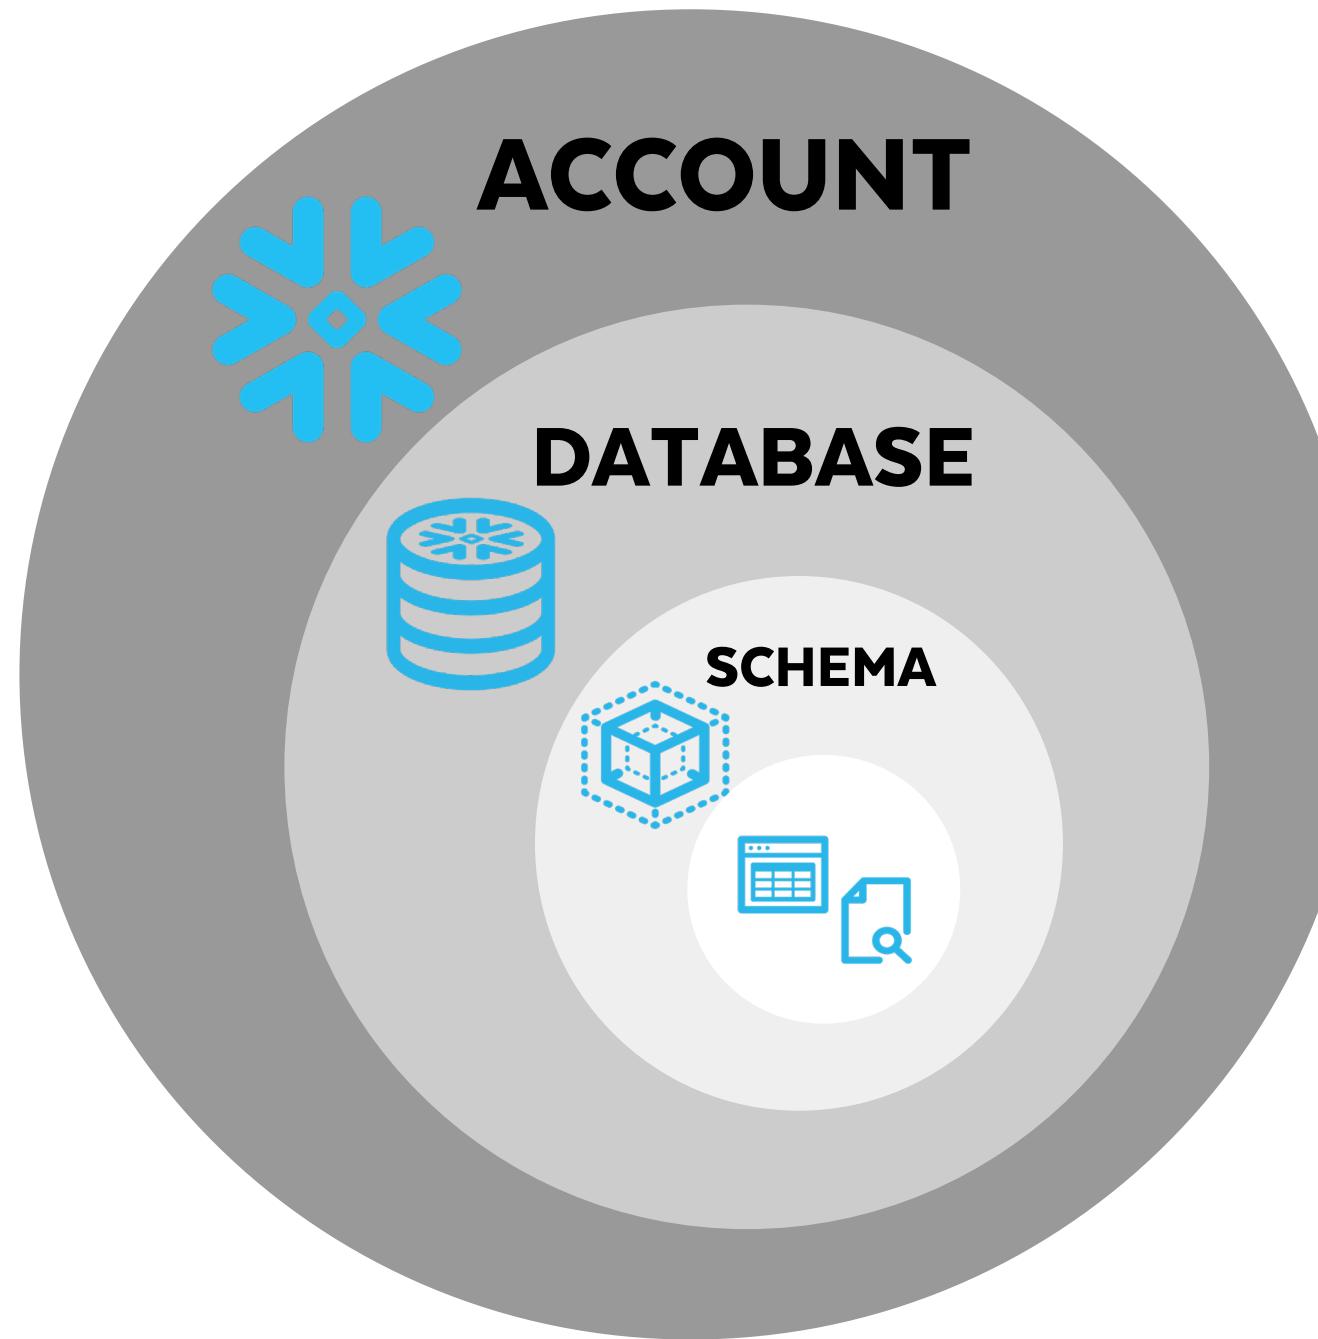
- Use the PUT command to load files into any type of stage\*
- Use the COPY INTO command to move files from any stage into a table
- Use the COPY INTO command to unload files from tables and put them into any kind of stage.

\* The PUT command cannot be run from WebUI Worksheet. The PUT command can be run from SnowSQL, the Snowflake CLI client. The PUT command can also be run via other programmatic interfaces. You can also load files into a stage by going to the AWS, GCP, or Azure console and dragging them into storage or uploading them via the tools offered by the cloud platform providers.

# PLATFORM VS. STAGES



# STORAGE HIERARCHY



## Further Reading:

### DATABASES AND SCHEMAS

<https://docs.snowflake.net/manuals/sql-reference/ddl-database.html#database-schema-share-ddl>

### TABLES

<https://docs.snowflake.net/manuals/user-guide/tables-temp-transient.html#comparison-of-table-types>

### VIEWS

<https://docs.snowflake.net/manuals/user-guide/views-introduction.html>

### STAGES

<https://docs.snowflake.net/manuals/user-guide/data-load-local-file-system-create-stage.html#types-of-stages>

**What is the recommended Snowflake data type to store semi-structured data like JSON?**

Options are :

- *VARCHAR*
- *RAW*
- *LOB*
- *VARIANT (Correct)*

**Answer :**

\**VARIANT (Correct)*

**If a Small Warehouse is made up of 2 servers/cluster, how many servers/cluster make up a Medium Warehouse?**

Options are :

- *4 (Correct)*
- *16*
- *32*
- *128*

**Answer :**

\**4 (Correct)*

**True or False: Snowflake charges additional credits base on usage to Data Providers for the SHARE they create for a Consumer who are not Snowflake Customer?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 8**

**Snowflake provides a mechanism for its customers to override its natural clustering algorithms. This method is:**

Options are :

- *Micro-partitions*
- *Clustering keys (Correct)*
- *Key partitions*
- *Clustered partitions*

**Answer :**

\*Clustering keys (Correct)

**Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)**

Options are :

- *Custom*
- *Economy (Correct)*
- *Optimized*
- *Standard (Correct)*

**Answer :**

\*Economy (Correct)

\*Standard (Correct)

**True or False: A single database can exist in more than one Snowflake account.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 12**

**Which of the following roles is recommended to be used to create and manage users and roles?**

Options are :

- *SYSADMIN*

- *SECURITYADMIN* (*Correct*)
- *PUBLIC*
- *ACCOUNTADMIN*

**Answer :**

\**SECURITYADMIN* (*Correct*)

**Which statement best describes ◊clustering◊?**

Options are :

- *Clustering represents the way data is grouped together and stored within Snowflake◊s micro-partitions* (*Correct*)
- *The database administrator must define the clustering methodology for each Snowflake table*
- *The clustering key must be included on the COPY command when loading data into Snowflake*
- *Clustering can be disabled within a Snowflake account*

**Answer :**

\**Clustering represents the way data is grouped together and stored within Snowflake◊s micro-partitions* (*Correct*)

**True or False: The COPY command must specify a File Format in order to execute.**

Options are :

- *TRUE*
- *FALSE* (*Correct*)

**Answer :**

\**FALSE* (*Correct*)

### Snowflake Certification Preparation Set 16

**Query results are stored in the Result Cache for how long after they are last accessed, assuming no data changes have occurred?**

Options are :

- *1 Hour*

- 3 Hours
- 12 hours
- 24 hours (Correct)

**Answer :**

\*24 hours (Correct)

**Which of the following connectors are available in the Downloads section of the Snowflake Web Interface (UI)? (Choose two.)**

Options are :

- SnowSQL (Correct)
- ODBC (Correct)
- R
- HIVE

**Answer :**

\*SnowSQL (Correct)

\*ODBC (Correct)

**Which of the following DML commands isn't supported by Snowflake?**

Options are :

- UPSERT (Correct)
- MERGE
- UPDATE
- TRUNCATE TABLE

**Answer :**

\*UPSERT (Correct)

**Snowflake SnowPro Core Certification Practice Test Set 3**

**Which of the following statements is true of zero-copy cloning?**

Options are :

- Zero-copy clones increase storage costs as cloning the table requires storing its data twice
- All zero-copy clone objects inherit the privileges of their original objects
- Zero-copy cloning is licensed as an additional Snowflake feature
- At the instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)

**Answer :**

\*At the instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)

**A Virtual Warehouse's auto-suspend and auto-resume settings apply to:**

Options are :

- The primary cluster in the Virtual Warehouse
- The entire Virtual Warehouse (Correct)
- The database the Virtual Warehouse resides in
- The queries currently being run by the Virtual Warehouse

**Answer :**

\*The entire Virtual Warehouse (Correct)

**Which of the following is true of Snowpipe REST API Option? (Choose two.)**

Options are :

- You can only use it on Internal Stages
- All COPY INTO options are available during pipe creation (Correct)
- Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands (Correct)
- Snowpipe keeps track of which files it has loaded (Correct)

**Answer :**

\*All COPY INTO options are available during pipe creation (Correct)

\*Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands (Correct)

\*Snowpipe keeps track of which files it has loaded (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 12**

**Snowflake recommends, as a minimum, that all users with the following role(s) should be enrolled in Multi-Factor Authentication (MFA):**

Options are :

- *SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN*
- *SECURITYADMIN, ACCOUNTADMIN, SYSADMIN*
- *SECURITYADMIN, ACCOUNTADMIN*
- *ACCOUNTADMIN (Correct)*

**Answer :**

\**ACCOUNTADMIN (Correct)*

**When can a Virtual Warehouse start running queries?**

Options are :

- *12am-5am*
- *Only during administrator defined time slots*
- *When its provisioning is complete (Correct)*
- *After replication*

**Answer :**

\**When its provisioning is complete (Correct)*

**True or False: Users are able to see the result sets of queries executed by other users that share their same role.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 19**

**True or False: Can a Snowflake user pick a cluster to run his query in a multi-cluster Warehouse?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**True or False: Snowflake Snow Pipes can be Paused and resumed?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**When should you consider disabling auto-suspend for a Virtual Warehouse? (Choose two.)**

Options are :

- *When users will be using compute at different times throughout a 24/7 period*
- *When managing a steady workload (Correct)*
- *When the compute must be available with no delay or lag time (Correct)*
- *When you do not want to have to manually turn on the Warehouse each time a user needs it*

**Answer :**

\**When managing a steady workload (Correct)*

\**When the compute must be available with no delay or lag time (Correct)*

### **Snowflake SnowPro Core Certification Practice Test Set 3**

**Which of the following simple transformations can be used while loading data through the COPY command?**

Options are :

- *Cast (Correct)*
- *Transpose*

- *Pivot*
- *Truncate (Correct)*
- *Re-order columns (Correct)*
- *Omit columns (Correct)*

**Answer :**

- \*Cast (Correct)
- \*Truncate (Correct)
- \*Re-order columns (Correct)
- \*Omit columns (Correct)

**Which of the following scaling type would result in Snowflake preserving credits over performance?**

Options are :

- *Economy (Correct)*
- *Performance*
- *Cost*
- *Standard*

**Answer :**

- \*Economy (Correct)

**True or False: Snowflake maintains the partitions using the order of the data in which it is inserted**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 32](#)

**True or False: If you have multiple virtual warehouses in your Snowflake system, they will each be accessing the same shared data.**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\**TRUE (Correct)*

**True or False: A new custom role will automatically be assigned to all existing users.**

Options are :

- *TRUE*

- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which of the following are valid Casting function available in Snowflake?**

Options are :

- *TO\_VARIANT (Correct)*

- *TO\_BLOB*

- *TO\_TIMESTAMP (Correct)*

- *TO\_NUMERIC (Correct)*

- *TO\_NUMBER (Correct)*

**Answer :**

\**TO\_VARIANT (Correct)*

\**TO\_TIMESTAMP (Correct)*

\**TO\_NUMERIC (Correct)*

\**TO\_NUMBER (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 23**

**What is the maximum time travel period allowed for transient & temporary tables?**

Options are :

- *7*

- *90*

- 0
- 1 (Correct)

**Answer :**

\*1 (Correct)

**Which role is most powerful role in a Snowflake system?**

Options are :

- PUBLIC
- SYSADMIN
- SECURITYADMIN
- ACCOUNTADMIN (Correct)
- USERADMIN

**Answer :**

\*ACCOUNTADMIN (Correct)

**True or False: REST API calls can be used to trigger a Snowpipe regardless of whether it is defined over an Internal Stage or an External Stage.**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

### Snowflake Snowpro Certification Practice Test Set 1

**When a database or a schema is cloned, which of the following statements are true for the stages in that database?**

Options are :

- Table Internal stages are cloned (Correct)
- External named stages are cloned (Correct)
- Named internal stages are NOT closed (Correct)

- *Named internal stages are cloned*
- *External named stages are Not cloned*

**Answer :**

- \*Table Internal stages are cloned (Correct)
- \*External named stages are cloned (Correct)
- \*Named internal stages are NOT closed (Correct)

**Which of the following objects can be shared? Select all that apply**

Options are :

- *Secure UDFs (Correct)*
- *Database (Correct)*
- *Tables (Correct)*
- *Schema (Correct)*
- *Secure Views (Correct)*
- *Views*

**Answer :**

- \*Secure UDFs (Correct)
- \*Database (Correct)
- \*Tables (Correct)
- \*Schema (Correct)
- \*Secure Views (Correct)

**An Engineer increased the size of a virtual warehouse to improve performance of a query? This is a case of?**

Options are :

- *Scale Out*
- *Scale UP (Correct)*
- *Scale High*
- *Scale Max*

**Answer :**

- \*Scale UP (Correct)

**Snowflake Snowpro Certification Practice Test Set 5**

**To create a new Task a user should be using a role which**

Options are :

- Has the CREATE TASK privilege on the schema where the Task is being created (Correct)
- Has the SYSADMIN privilege
- Has the ACCOUNTADMIN privilege
- Has the SECURITYADMIN privilege

**Answer :**

\*Has the CREATE TASK privilege on the schema where the Task is being created (Correct)

**Snowflake automatically determines the most efficient algorithm to compress columns in a micro partitions**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Which of the following is required in order for a role to be able to successfully execute a Task?**

Options are :

- The role should have the permissions required to run the SQL statement executed in the Task (Correct)
- The role should have ACCOUNTADMIN privillages
- The role should have SYSADMIN privillages
- The role should have the global EXECUTE TASK privillage on the Account.

**Answer :**

\*The role should have the permissions required to run the SQL statement executed in the Task (Correct)

**Snowflake Snowpro Certification Practice Test Set 5**

**Snowflake SQL UDFs can return either Scalar or Tabular results?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Assume a virtual warehouse of size X-Large (128 servers) running for an hour. What is the amount of Snowflake credits that will be billed?**

Options are :

- *64*
- *60*
- *256*
- *128 (Correct)*

**Answer :**

\**128 (Correct)*

**Compared to a single cluster virtual warehouse, a multi-cluster virtual warehouse supports which of the following? Select all that apply**

Options are :

- *Auto Suspend (Correct)*
- *Auto Resume (Correct)*
- *Suspend (Correct)*
- *Resize (Correct)*

**Answer :**

\**Auto Suspend (Correct)*  
\**Auto Resume (Correct)*  
\**Suspend (Correct)*  
\**Resize (Correct)*

**[SnowFlake SnowPro Advanced Architect Practice Exam Set 7](#)**

**A cloned database doesn't inherit the source permissions, but schemas, tables & views contained inside the cloned database will inherit the source permissions.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which of the following is not a layer in Snowflake Architecture?**

Options are :

- *Cloud services*
- *Virtual warehouse*
- *Database Storage*
- *Virtual Machines (Correct)*

**Answer :**

\**Virtual Machines (Correct)*

**Snowflake Clustering key helps query performance by?**

Options are :

- *Distributing the data over multiple cluster*
- *Pre-calculating query results*
- *Improving column compression*
- *Pruning unnecessary partitions (Correct)*

**Answer :**

\**Pruning unnecessary partitions (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 4**

**True or False: Except VPS, any snowflake account regardless of the edition, can produce and consume shared data.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**True or False:** If you create more then one virtual warehouse they will share the memory and CPU resources.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Snowflake javaScript UDFs can return either scalar or Tabular results?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**What is the account identifier inside this Snowflake URL ' <https://jz96798.us-east-1.azure.snowflakecomputing.com> ' ?**

Options are :

- *jz96798.us-east-1.azure*
- *us-east-1.azure*
- *jz96798 (Correct)*

- *snowflakecomputing.com*

- *Azure*

**Answer :**

\*jz96798 (Correct)

**Which the right options to convert 'AMOUNT' column to DECIMAL(10, 2) inside a select statement**

Options are :

- *SELECT AMOUNT::DECIMAL(10,5) (Correct)*
- *SELECT AMOUNT( DECIMAL(10,2) )*
- *SELECT AMOUNT as DECIMAL(10,2) (Correct)*
- *SELECT AMOUNT.DECIMAL(10,2)*

**Answer :**

\*SELECT AMOUNT::DECIMAL(10,5) (Correct)

\*SELECT AMOUNT as DECIMAL(10,2) (Correct)

**What factors impact credit consumption of the compute layer ? Select all that apply**

Options are :

- *Amount of data processed*
- *Size of warehouse (Correct)*
- *Number of clusters (Correct)*
- *Number of user*

**Answer :**

\*Size of warehouse (Correct)

\*Number of clusters (Correct)

### Snowflake Certification Preparation Set 7

**Which of the following Snowflake objects allows to query data without loading it first inside Snowflake ?**

Options are :

- *Snowpipe*

- *COPY*
- *External table (Correct)*
- *Virtual table*

**Answer :**

\*External table (Correct)

A client as an important number of users that access a BI report ever working day at 9AM. Every time a user refresh the report it requests Snowflake for the data. The report performances are not optimal.\n What suggestion could you make to improve the performance ? Select the best option

Options are :

- *Create a single warehouse for each user*
- *Regroup the user by age and associate a warehouse to each group*
- *Use a multi-cluster warehouse (Correct)*
- *None of the above*

**Answer :**

\*Use a multi-cluster warehouse (Correct)

Which type of view has an extra layer of protection to hide the SQL code from unauthorized viewing ?

Options are :

- *Standard*
- *Secure (Correct)*
- *Materialized*
- *Permanent*

**Answer :**

\*Secure (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 2

In which Snowflake editions is it possible to extended Time Travel (up to 90 days) ? (Select two or more)

Options are :

- *Standard Edition*
- *Enterprise Edition (Correct)*
- *Business Critical Edition (Correct)*
- *Virtual Private Snowflake (Correct)*

**Answer :**

- \*Enterprise Edition (Correct)
- \*Business Critical Edition (Correct)
- \*Virtual Private Snowflake (Correct)

**Which of the following are types of cache use by Snowflake ? Select all that apply**

Options are :

- *Metadata cache (Correct)*
- *Index cache*
- *Query result cache (Correct)*
- *Data cache (Correct)*
- *None of the above*

**Answer :**

- \*Metadata cache (Correct)
- \*Query result cache (Correct)
- \*Data cache (Correct)

**How many clusters are in standard Medium warehouse ?**

Options are :

- *One (Correct)*
- *Two*
- *Three*
- *Four*

**Answer :**

- \*One (Correct)

**Snowflake SnowPro Core Certification Exam Questions Set 5**

**External stages do not require a cloud provider. TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**In which mode Snowflake multi-cluster warehouses can run ? Select all that apply.**

Options are :

- *Maximized (Correct)*
- *Auto-resume*
- *Auto-suspending*
- *Auto-scale (Correct)*

**Answer :**

\**Maximized (Correct)*  
\**Auto-scale (Correct)*

**What kind of metadata does Snowflake stores in micro-partition for a table**

Options are :

- *MIN (Correct)*
- *MAX (Correct)*
- *Table's privileges*
- *Table's owner*
- *Number of distinct values (Correct)*
- *None of the above*

**Answer :**

\**MIN (Correct)*  
\**MAX (Correct)*  
\**Number of distinct values (Correct)*

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 36

**Which cache runs on a 24 hour clock?**

Options are :

- *Result cache (Correct)*
- *Metadata Cache*
- *Warehouse cache*

**Answer :**

\*Result cache (Correct)

**Select the two terms that refer to the same Snowflake layer ?**

Options are :

- *Metadata Layer*
- *Virtual Storage Layer*
- *Catalog Layer*
- *Data Layer (Correct)*
- *Storage Layer (Correct)*

**Answer :**

\*Data Layer (Correct)

\*Storage Layer (Correct)

**Which file formats is not supported by Snowflake to load data ?**

Options are :

- *Delimited files*
- *JSON*
- *Avro*
- *ORC*
- *Parquet*
- *PDF (Correct)*

**Answer :**

\*PDF (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 1

**What stages can you use when setting up your Snowflake account with AWS as your cloud platform provider ? Select all that apply**

Options are :

- *Named External - Using S3 Buckets (Correct)*
- *Named External - Hadoop Cloudera File System*
- *Named External - Using Azure Blob storage (Correct)*
- *Named External - Using GCS/GCP Buckets (Correct)*
- *Named External - AWS Redshift tables*
- *None of the above*

**Answer :**

\*Named External - Using S3 Buckets (Correct)  
\*Named External - Using Azure Blob storage (Correct)  
\*Named External - Using GCS/GCP Buckets (Correct)

**It is possible to install Snowflake on-premise. TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following techniques can be used in Snowflake to improve query performance?**

Options are :

- *Distribution keys*
- *Index definition*
- *Clustering keys (Correct)*

- *Reclustering (Correct)*
- *None of the above*

**Answer :**

- \*Clustering keys (Correct)
- \*Reclustering (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 37](#)

**Which installment option version of Snowflake are available ?**

Options are :

- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)*
- *Microsoft Cloud Native Accounts*
- *Hybrid On-Premise + Cloud Installation*
- *Enterprise In-House VPC Installation*
- *On-Premise Custom Installation*

**Answer :**

- \*Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)
- \*Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)

**With Economy scaling, how long might queries be queued before another cluster is added to the warehouse ?**

Options are :

- *3 minutes*
- *6 minutes (Correct)*
- *9 minutes*
- *12 minutes*
- *None of the above*

**Answer :**

- \*6 minutes (Correct)

**Which type of stages are automatically available in Snowflake and do not need to be created or configured ?**

Options are :

- *User (Correct)*
- *Table (Correct)*
- *Named Internal*
- *Named External*
- *None of the above*

**Answer :**

\*User (Correct)  
\*Table (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 17](#)

**Select the three terms refering to the same layer.**

Options are :

- *Compute Layer (Correct)*
- *Virtual Warehouse Layer (Correct)*
- *Catalog Layer*
- *Metadata Layer*
- *Query Processing Layer (Correct)*

**Answer :**

\*Compute Layer (Correct)  
\*Virtual Warehouse Layer (Correct)  
\*Query Processing Layer (Correct)

**Which scaling policy, starts additional cluster immediately when either a query is queued or the system detects that there is one more query than the currently-running clusters can execute.**

Options are :

- *Economy*
- *Standard (Correct)*

**Answer :**

\*Standard (Correct)

**Which compute resource can be used by Snowflake for data loading jobs ? Select all that apply**

Options are :

- *User-managed virtual warehouse (Correct)*
- *Hardware provisionned by user directly from cloud provider*
- *Snowflake-managed serverless compute (Correct)*
- *Cloud Provider Runtime Environment*

**Answer :**

\*User-managed virtual warehouse (Correct)  
\*Snowflake-managed serverless compute (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 9](#)

**Select all statements that are TRUE about FILE FORMAT.**

Options are :

- *Stores data from a loaded file in a specific, named format*
- *Is securable (Correct)*
- *Stores information needed to parse files during load/unload (Correct)*
- *Can only be of type CSV*

**Answer :**

\*Is securable (Correct)  
\*Stores information needed to parse files during load/unload (Correct)

**Which table disappears after the close of the session and has no fail-safe, no time travel options after the close of the session ?**

Options are :

- *Permanent*
- *Temporary (Correct)*
- *Transient*

- *External*
- *None of the above*

**Answer :**

\*Temporary (Correct)

**Which factors affect your snowflake data loading rate ?**

Options are :

- *Physical location of the stage (Correct)*
- *RAM on the virtual warehouse*
- *Thread size*
- *GZIP compression efficiency (Correct)*

**Answer :**

\*Physical location of the stage (Correct)

\*GZIP compression efficiency (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 17**

**What Snowflake objects can be shared ? (Select all that apply)**

Options are :

- *Tables (Correct)*
- *Views*
- *Roles*
- *Secure UDFs (Correct)*
- *Materialized views*
- *None of the above*

**Answer :**

\*Tables (Correct)

\*Secure UDFs (Correct)

**It is possible to set a storage quota at the account level. TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**What is the maximum number of child a single task can have in Snowflake ?**

Options are :

- *1*
- *100 (Correct)*
- *1000*
- *10000*

**Answer :**

\**100 (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 29**

**What are the two mechanisms that can be used with Snowpipe for detecting that staged files are available ?**

Options are :

- *Calling Snowpipe REST endpoints (Correct)*
- *Calling Snowpipe SOAP endpoints*
- *Automating Snowpipe using CloudWatch*
- *Automating Snowpipe using cloud messaging (Correct)*

**Answer :**

\**Calling Snowpipe REST endpoints (Correct)*

\**Automating Snowpipe using cloud messaging (Correct)*

**What is the maximum size for a VARIANT data type column ?**

Options are :

- *64MB uncompresssed*
- *16MB uncompressed (Correct)*

- 32MB uncompressed
- 128MB uncompressed

**Answer :**

\*16MB uncompressed (Correct)

**Multiple warehouses can access the same data simultaneously without resource contention issues. TRUE or FALSE ?**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

#### Snowflake Practice Questions - SnowPro Core Certified Set 4

**Materialized view can be created on top of external tables to improve query performance without any additional cost. TRUE or FALSE ?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Which one of these actions is an automated process**

Options are :

- Scaling OUT (Correct)
- Scaling UP

**Answer :**

\*Scaling OUT (Correct)

**The expected average run time is the average difference between the scheduled and completed times for a task, excluding any time period in which the task was queued. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 25

**Scaling a warehouse IN will decrease the number of clusters. (e.g. Max to Min). TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which the most accurate statement on Snowflake's transaction support. Select the best option**

Options are :

- *ACID compliant (Correct)*
- *Requires special user configuration*
- *Base compliant*
- *Works like a NoSQL database*

**Answer :**

\**ACID compliant (Correct)*

**There is no benefit to stopping a warehouse before the first 60-second period is over because the credits have already been billed for that period**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 7

**Snowflake data storage costs are calculated based on**

Options are :

- *Uncompressed size*
- *Compressed size (Correct)*
- *Amount stored on Last day of Month*
- *Amount stored on First day of Month*
- *Amout stored - Daily average (Correct)*

**Answer :**

\*Compressed size (Correct)

\*Amout stored - Daily average (Correct)

**When you share data with a consumer account, who pays for compute time when data is access by the consumer?**

Options are :

- *The data provider does*
- *The data consumer does. (Correct)*
- *Both provider and consumer does*
- *None of the above*

**Answer :**

\*The data consumer does. (Correct)

**Snowflake has scalar and tabular user-defined functions (UDFs), with support for both SQL, JavaScript, Java. TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

## Snowflake SnowPro Core Certification Practice Test Set 4

**How does Snowflake billed for compute ?**

Options are :

- *Per hour, with one-minute minimum*
- *Per second, with one-minute minimum (Correct)*
- *Per day*
- *Per month*

**Answer :**

\*Per second, with one-minute minimum (Correct)

**What attributes make Snowflake a true SaaS solution ?**

Options are :

- *No hardware to purchase or configure (Correct)*
- *No creation of user accounts or roles is required*
- *No maintenance upgrades or patches to install (Correct)*
- *No data storage costs*
- *No query processing costs*
- *Transparent releases don't required user intervention (Correct)*

**Answer :**

\*No hardware to purchase or configure (Correct)

\*No maintenance upgrades or patches to install (Correct)

\*Transparent releases don't required user intervention (Correct)

**When loading JSON file inside Snowflake table, which file format option remove the outer array structure and load the records into separate table rows ? (**

Options are :

- *STRIP\_NULL\_VALUES = TRUE*

- *STRIP\_OUTER\_ARRAY = TRUE* (*Correct*)
- *TRIM\_SPACE = TRUE*
- *None of the above*

**Answer :**

\**STRIP\_OUTER\_ARRAY = TRUE* (*Correct*)

**Virtual Warehouse credits are charged on what basis?**

Options are :

- *Per mili second of use*
- *Per second of use (Correct)*
- *Per hour of use*
- *Per minute of use*

**Answer :**

\**Per second of use (Correct)*

**What are the virtual warehouse sizing approach recommended by Snowflake?**

Options are :

- *Experiment with diffrent types of query and virtual warehouse to determine the virtual warehouse seze that best meets your requirments. (Correct)*
- *Execute your queries on the smallest available virtual warehouse and keep scaling up the size of the virtual warehouse, until the query performance no longer shows improvements.*
- *Execute your queries on the largest available virtual warehouse and keep scaling down the size of the virtual warehouse, until you start seeing degradation in query response.*

**Answer :**

\**Experiment with diffrent types of query and virtual warehouse to determine the virtual warehouse seze that best meets your requirments. (Correct)*

**Which of the following Snowflake Editions allow only a maximum of 1 day of time travel?**

Options are :

- *Standard (Correct)*
- *Business Critical*
- *Enterprise*
- *VPS*

**Answer :**

\*Standard (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 15

**You can increase or decrease the size of a virtual warehouse as your user base & query demand increase or decrease?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**On the WebUI which button on the top bar should you select to get to tabs that provide information about virtual warehouse usage?**

Options are :

- *Virtual Warehouse*
- *Account (Correct)*
- *History*
- *Shares*
- *Databases*

**Answer :**

\*Account (Correct)

**Files that have already been processed into the source table can be loaded again into a cloned table?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### Snowflake Certification Preparation Set 18

**When a virtual warehouse is provisioned there is no benefit in stopping the virtual warehouse with in the first 60 seconds because that time period has already been billed.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**What is one of the way to improve performance in Snowflake?**

Options are :

- *Secondary Indices*
- *Query Hints*
- *Clustering Keys (Correct)*
- *Join Indices*

**Answer :**

\**Clustering Keys (Correct)*

**When a table is cloned, a snapshot of the data is taken at the time of the cloning and represents the state of the source data at the time of cloning.**

Options are :

- *Independently - each column stored on its own (Correct)*

- Combined - columns for a given row are stored together.

**Answer :**

\*Independently - each column stored on its own (Correct)

### Snowflake Snowpro Certification Practice Test Set 1

**What happens when different value is specified for the minimum & the maximum cluster? Select all that apply.**

Options are :

- *Snowflake starts all clusters in the multi-cluster virtual warehouse.*
- *If the demand is decreased the number of cluster are gradually shutdown, up to the defined minimum cluster value. (Correct)*
- *If the number of concurrent users and queries increase on the virtual warehouse, additional clusters are started, up to the defined maximum cluster value. (Correct)*
- *The mode is called auto-scale (Correct)*
- *Snowflake starts minimum required clusters in the multi-cluster virtual warehouse that can accommodate the number of users & number of queries. (Correct)*

**Answer :**

\*If the demand is decreased the number of cluster are gradually shutdown, up to the defined minimum cluster value. (Correct)

\*If the number of concurrent users and queries increase on the virtual warehouse, additional clusters are started, up to the defined maximum cluster value. (Correct)

\*The mode is called auto-scale (Correct)

\*Snowflake starts minimum required clusters in the multi-cluster virtual warehouse that can accommodate the number of users & number of queries. (Correct)

**Which of the following are valid scaling type for a multi cluster virtual warehouse? Select all that apply.**

Options are :

- *PERFORMANCE*
- *STANDARD (Correct)*
- *SPEED*
- *ECONOMY (Correct)*

**Answer :**

- \*STANDARD (Correct)
- \*ECONOMY (Correct)

**All queries will perform better if the virtual warehouse size is increased?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

### Snowflake Certification Preparation Set 11

**Under what circumstances results for a query will be fulfilled by the query result cache? Select all that apply.**

Options are :

- *The query is being executed from the same virtual warehouse as the previously executed query.*
- *The query results cache was generated or last used less than 24 hours ago. (Correct)*
- *The underlying data for the tables in the query has not changed. (Correct)*
- *The query being run is syntactically identical to a previously executed query. (Correct)*
- *The query being run returns result set of less than 100 thousand rows.*

**Answer :**

- \*The query results cache was generated was generated or last used less than 24 hours ago. (Correct)
- \*The underlying data for the tables in the query has not changed. (Correct)
- \*The query being run is syntactically identical to a previously executed query. (Correct)

**There is a cost associated with maintaining the partitions associated with the clustering keys?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**When a virtual warehouse is resized to a SMALLER size, the required servers are only removed when?**

Options are :

- *Immediately*
- *After 1 minute*
- *After 60 minute*
- *There are no SQL statements running on the servers (Correct)*

**Answer :**

\*There are no SQL statements running on the servers (Correct)

### Snowflake Snowpro Certification Practice Test Set 1

**What is the Snowflake recommended approach to identifying the minimum and maximum values for a multi-cluster virtual warehouse? Select all that apply.**

Options are :

- *Start with a minimum of 1 & a maximum cluster value of 2 or 3 (Correct)*
- *Track the virtual warehouse usage over time and adjust the minimum and maximum values that provide the best configuration for your work load. (Correct)*
- *Start with the multi-cluster virtual warehouse in auto scale mode. (Correct)*
- *Start with a minimum of 5 & a maximum cluster value of 10*
- *Start with the minimum and maximum value set to the same number.*

**Answer :**

\*Start with a minimum of 1 & a maximum cluster value of 2 or 3 (Correct)

\*Track the virtual warehouse usage over time and adjust the minimum and maximum values that provide the best configuration for your work load. (Correct)

\*Start with the multi-cluster virtual warehouse in auto scale mode. (Correct)

**You are planning to utilise the multi-Cluster virtual warehouse to provide auto-scaling and performance for your users. You want to make sure that the user queries don't wait in queue and if auto-scaling is required the cluster should auto scale immediately. What is the correct value for scaling policy to achieve this?**

Options are :

- *SCALING\_POLICY = PERFORMANCE*
- *SCALING\_POLICY = STANDARD (Correct)*
- *SCALING\_POLICY = ECONOMY*
- *SCALING\_POLICY = AUTO*

**Answer :**

\**SCALING\_POLICY = STANDARD (Correct)*

**Which of the following statement is true when sharing data with an organisation which doesn't have a Snowflake account?**

Options are :

- *Snowflake allows sharing of data to/from Snowflake Accounts only (so the consumer must have a Snowflake account or a reader account) (Correct)*
- *A consumer account must be created for a non-snowflake customer to allow sharing*
- *A reader account must be created for a non-snowflake customer to allow sharing (Correct)*
- *No extra configuration is required to share data with a non-snowflake user/organization*

**Answer :**

\**Snowflake allows sharing of data to/from Snowflake Accounts only (so the consumer must have a Snowflake account or a reader account) (Correct)*

\**A reader account must be created for a non-snowflake customer to allow sharing (Correct)*

### **[SnowPro Core Certification 3 Full Practice Exams 2021 Set 24](#)**

**A new custom role will automatically be assigned to all existing users.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**If no parameters are provided to the last\_query\_id function it defaults to which value?**

Options are :

- -1 (*Correct*)
- 0
- 1
- None of the above

**Answer :**

\*-1 (Correct)

**An external table in Snowflake...**

Options are :

- contains meta-data about the data files located in cloud storage.
- is a read-only table.
- is usually slower in performance compared to a database table.
- All of the Above (*Correct*)

**Answer :**

\*All of the Above (Correct)

### **Snowflake Snowpro Certification Practice Test Set 3**

**What is the optimal file sizing recommendation for loading data into Snowflake?**

Options are :

- 10 MB to 100 MB, uncompressed
- 100 MB to 250 MB, compressed (*Correct*)
- 100 MB to 250 MB, uncompressed
- 10 GB to 100 GB, compressed

**Answer :**

\*100 MB to 250 MB, compressed (Correct)

**In relation to the data sharing features in Snowflake, which of the following statements are TRUE?**

Options are :

- Any user can set up a share for the objects owned by that user.
- All database objects shared are always read-only for the consumer of the share. (Correct)
- When you share data with other Snowflake accounts, Snowflake copies data from the provider's Snowflake account to the consumer's Snowflake account.
- A single Share can contain objects from different databases within the same provider account.
- Data sharing is supported between Snowflake accounts only. (Correct)

**Answer :**

\*All database objects shared are always read-only for the consumer of the share. (Correct)

\*Data sharing is supported between Snowflake accounts only. (Correct)

**Select appropriate database objects which can be part of a share. (select all that apply)**

Options are :

- A table (Correct)
- A view
- A secure view (Correct)
- A materialized view (Correct)
- A secure UDF (Correct)
- A stored procedure

**Answer :**

\*A table (Correct)

\*A secure view (Correct)

\*A materialized view (Correct)

\*A secure UDF (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 32**

**From the statements given below, select the statements which are TRUE for reader accounts. (select all that apply)**

Options are :

- A reader account facilitates data sharing with consumers that do not have a Snowflake account. (Correct)
- One reader account can be used to access multiple shares provided by different provider accounts.
- A reader account consumes credit of the provider account. (Correct)
- A reader account has the same access to Snowflake support as the provider account.

**Answer :**

\*A reader account facilitates data sharing with consumers that do not have a Snowflake account. (Correct)  
 \*A reader account consumes credit of the provider account. (Correct)

**Snowflake provides a share called ACCOUNT\_USAGE Share. Following are the unique features of ACCOUNT\_USAGE share. ... (select all that apply)**

Options are :

- It is available to every Snowflake account. (Correct)
- It cannot be dropped or deleted. (Correct)
- Default access to this share is available to the ACCOUNTADMIN role. (Correct)
- It can be shared with data consumers to share metadata, statistics and metrics.
- It can be used to find metadata about the objects which are dropped. (Correct)
- Data retention period of metadata stored in the ACCOUNT\_USAGE is technically unlimited.

**Answer :**

\*It is available to every Snowflake account. (Correct)  
 \*It cannot be dropped or deleted. (Correct)  
 \*Default access to this share is available to the ACCOUNTADMIN role. (Correct)  
 \*It can be used to find metadata about the objects which are dropped. (Correct)

**Data storage cost in Snowflake is calculated based on: (select all that apply)**

Options are :

- Daily average of persistent data stored in databases in compressed format (Correct)
- Daily average of persistent data stored in databases in uncompressed format
- Data maintained by Snowflake for time-travel and fail-safe (data recovery) (Correct)

- *The volume of data cache*
- *The volume of metadata*

**Answer :**

- \*Daily average of persistent data stored in databases in compressed format (Correct)
- \*Data maintained by Snowflake for time-travel and fail-safe (data recovery) (Correct)

### [Snowflake SnowPro Core Certification Practice Test Set 1](#)

**From the statements given below, select statements which are TRUE concerning database replication and failover/failback features in Snowflake.**

Options are :

- *Database replication and failover/failback are supported in Business Critical edition and above. (Correct)*
- *Database replication is cost-free in Snowflake.*
- *Snowflake supports database replication to non-Snowflake database products (e.g., Oracle/PostgreSQL, etc.)*
- *Database replication is supported for all objects in the primary account.*
- *Database created from Shares cannot be replicated. (Correct)*

**Answer :**

- \*Database replication and failover/failback are supported in Business Critical edition and above. (Correct)
- \*Database created from Shares cannot be replicated. (Correct)

**Which of the following data types are supported in Snowflake for semi- structured data? (select all that apply)**

Options are :

- *VARIANT (Correct)*
- *OBJECT (Correct)*
- *TEXT*
- *CLOB*
- *ARRAY (Correct)*

**Answer :**

- \**VARIANT (Correct)*

\*OBJECT (Correct)

\*ARRAY (Correct)

**Your application receives JSON data files from a sensor every 15 min. You decide to store the files in a database table in a VARIANT column. What is the maximum size supported by the VARIANT column?**

Options are :

- 16 MB, compressed (Correct)
- 8 MB, compressed
- 16 MB, uncompressed
- 8 MB, uncompressed

**Answer :**

\*16 MB, compressed (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 3

**Identify statements that are TRUE from the statements given below concerning the semi-structured data processing. (Select all that apply)**

Options are :

- *Snowflake supports JSON path notation in SQL query to access semi-structured data.* (Correct)
- *For a JSON file that contains mainly strings and integers, the storage requirements and query performance of semi- structured data are very similar to relational data processing.* (Correct)
- *FLATTEN function will remove the outer array structure of a JSON file and load the records in separate table rows.*
- *The LATERAL FLATTEN function can be used to parse arrays in a JSON file. It returns a row for each array object. LATERAL modifier joins the data with any information outside of the object.* (Correct)
- *VARIANT columns in a relational table are stored as separate physical columns.* (Correct)

**Answer :**

\*Snowflake supports JSON path notation in SQL query to access semi-structured data. (Correct)

\*For a JSON file that contains mainly strings and integers, the storage requirements and query performance of semi- structured data are very similar to relational data

processing. (Correct)

\*The LATERAL FLATTEN function can be used to parse arrays in a JSON file. It returns a row for each array object. LATERAL modifier joins the data with any information outside of the object. (Correct)

\*VARIANT columns in a relational table are stored as separate physical columns. (Correct)

**You are supporting a POS database that contains a table named t\_customer. For the SQL statement given below, which of the following options will hold? (select all that apply) Select count(\*) from t\_customer;**

Options are :

- *Snowflake will spin up a warehouse when the query is issued for the first time. After that, the results will be cached.*
- *As Count can keep changing, every time this query is issued, Snowflake will have to spin up a warehouse to retrieve the latest Count statistics.*
- *Snowflake does not spin up a warehouse to execute this query because Count statistics are stored in resultset Cache.*
- *Snowflake does not spin up a warehouse to execute this query because Count statistics are stored in metadata Cache. (Correct)*

**Answer :**

\*Snowflake does not spin up a warehouse to execute this query because Count statistics are stored in metadata Cache. (Correct)

**You are supporting a POS database that contains a table named t\_customer. Every day, about 20+ business users connect to this Snowflake database and execute a query on this table to pull the customer data. Once the data is downloaded, the users will then start analyzing the data using their favorite data analysis tools. An XS-sized warehouse is assigned to process this query workload. How will this scenario be managed in Snowflake? (select all that apply)**

Options are :

- *Every time the query is executed, Snowflake will spin up the warehouse (or use an already running warehouse) and retrieve data from the database to ensure the most up-to-date data is provided to the user.*
- *If the underlying data has not changed, Snowflake will not spin up the warehouse. It will re-use data from the Resultset Cache provided that the exact same query is executed by the user having the exact same role. (Correct)*
- *If the underlying data has changed, Snowflake will automatically spin up the warehouse (or use an already running warehouse) and retrieve data from the database. (Correct)*

- Snowflake will aggregate the results of queries issued by all the users so that when the same query comes back again, no warehouse will be needed.

**Answer :**

- \*If the underlying data has not changed, Snowflake will not spin up the warehouse. It will re-use data from the Resultset Cache provided that the exact same query is executed by the user having the exact same role. (Correct)
- \*If the underlying data has changed, Snowflake will automatically spin up the warehouse (or use an already running warehouse) and retrieve data from the database. (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 5

If you want to view the query history older than 14 days, you will...

Options are :

- Not be able to do this as Snowflake stores only 14 days worth of query history.
- Go to tab ? Select and specify a time-period.
- SNOWFLAKE (database) ? ACCOUNT\_USAGE (schema) ? QUERY\_HISTORY view (Correct)
- Both B & C

**Answer :**

\*SNOWFLAKE (database) ? ACCOUNT\_USAGE (schema) ? QUERY\_HISTORY view (Correct)

Which of the following are valid statements in terms of Security and Access control in Snowflake? (select all that apply)

Options are :

- Access control in Snowflake is based on a role-based access control model. (Correct)
- Access control in Snowflake is based on the discretionary access control model. (Correct)
- The ACCOUNTADMIN role can access, modify or drop any object in the account.
- As a security best practice, all custom roles should roll up into ACCOUNTADMIN role.
- Objects created using the PUBLIC role are accessible to all users. (Correct)

**Answer :**

\*Access control in Snowflake is based on a role-based access control model. (Correct)  
\*Access control in Snowflake is based on the discretionary access control model.

(Correct)

\*Objects created using the PUBLIC role are accessible to all users. (Correct)

**Select the correct statements concerning the views in Snowflake. (select all that apply)**

Options are :

- *Using secure views, you can implement row-level security in your database. (Correct)*
- *A materialized view cannot be created as a secure view.*
- *The materialized views always stay current. (Correct)*
- *Materialized view creation is advisable when you have an often used, complex query on underlying data that will not change very frequently. (Correct)*
- *A secure view cannot be part of a data Share.*

**Answer :**

\*Using secure views, you can implement row-level security in your database. (Correct)

\*The materialized views always stay current. (Correct)

\*Materialized view creation is advisable when you have an often used, complex query on underlying data that will not change very frequently. (Correct)

**Snowflake Certification Preparation Set 14**

**Which of the following statements are TRUE concerning the storage in Snowflake? (select all that apply)**

Options are :

- *The snowflake storage layer is multi-tenant. (Correct)*
- *The snowflake storage layer can be characterized as a compressed, hybrid columnar data structure. (Correct)*
- *ACCOUNTADMIN can directly access the underlying physical files (i.e., micro partitions) in the data storage layer.*
- *Data at rest in Snowflake is encrypted with AES 256-bit encryption. (Correct)*
- *The primary cost factor for enterprises adopting Snowflake is the cost of compute which is significantly higher than the cost of storage. (Correct)*

**Answer :**

\*The snowflake storage layer is multi-tenant. (Correct)

\*The snowflake storage layer can be characterized as a compressed, hybrid columnar data structure. (Correct)

\*Data at rest in Snowflake is encrypted with AES 256-bit encryption. (Correct)

\*The primary cost factor for enterprises adopting Snowflake is the cost of compute which is significantly higher than the cost of storage. (Correct)

**Which of the following statements are TRUE concerning a stream object in Snowflake? (select all that apply)**

Options are :

- *A stream object provides a record of DML changes (inserts, updates, deletes) made to a table at row level. (Correct)*
- *A stream object can keep track of DML changes for the entire life of a table.*
- *Streams on materialized views are not supported. (Correct)*
- *Streams on external tables are not supported.*

**Answer :**

\*A stream object provides a record of DML changes (inserts, updates, deletes) made to a table at row level. (Correct)

\*Streams on materialized views are not supported. (Correct)

**A stream is considered to be ♦stale♦ when...**

Options are :

- *The underlying table is dropped.*
- *The CDC information contained in the stream is consumed.*
- *The stream is not consumed, and the data retention period of the underlying table has expired. (Correct)*
- *All of the above*

**Answer :**

\*The stream is not consumed, and the data retention period of the underlying table has expired. (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 31

**Only one stream object can be created on a source table.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following statements are TRUE concerning a task object in Snowflake? (select all that apply)**

Options are :

- *A task object can execute a single SQL statement based on a schedule. (Correct)*
- *A task object can execute a single SQL statement based on an event trigger.*
- *If the first execution (E1) of a task is still running and if the second execution (E2) of that task gets triggered, E2 will wait until E1 is complete.*
- *A task object cannot call a stored procedure.*

**Answer :**

\*A task object can execute a single SQL statement based on a schedule. (Correct)

**Which of the following are true regarding a task tree?**

Options are :

- *Each task node in the task tree has its own schedule.*
- *Each task node in the task tree can have only one predecessor task. (Correct)*
- *A task tree can have an unlimited number of task objects.*
- *One parent task in a task tree can have an unlimited number of child tasks.*
- *None of the above*

**Answer :**

\*Each task node in the task tree can have only one predecessor task. (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 23**

**Following are the key features of Snowflake Partner Connect. (select all that apply)**

Options are :

- *It lets you easily create trial accounts with the Partner connect business partners and integrate these accounts with your Snowflake environment. (Correct)*
- *It automates the process of integrating your Snowflake account with the Partner Connect account. (Correct)*

- It Is accessible through Snowflake Web UI. (Correct)
- Snowflake has a pre-negotiated license agreement for each Partner Connect vendor for Snowflake Enterprise customers.
- Only the ACCOUNTADMIN role can use Partner Connect. (Correct)
- Partner connect features only the data integration category partners.

**Answer :**

- \*It lets you easily create trial accounts with the Partner connect business partners and integrate these accounts with your Snowflake environment. (Correct)
- \*It automates the process of integrating your Snowflake account with the Partner Connect account. (Correct)
- \*It Is accessible through Snowflake Web UI. (Correct)
- \*Only the ACCOUNTADMIN role can use Partner Connect. (Correct)

**Which of the following statements are TRUE concerning Data Marketplace in Snowflake (select all that apply)?**

Options are :

- Data Marketplace is supported on all 3 cloud providers (Azure, GCP, and AWS).
- For a data provider, the data marketplace enables to monetize data quickly and easily. (Correct)
- For a data consumer, it offers real-time access to datasets that can be combined with existing data to derive new insights without the complexity of data integration. (Correct)
- Data Marketplace is available in Snowflake Enterprise edition and above.

**Answer :**

- \*For a data provider, the data marketplace enables to monetize data quickly and easily. (Correct)
- \*For a data consumer, it offers real-time access to datasets that can be combined with existing data to derive new insights without the complexity of data integration. (Correct)

**Which of the following workload(s) is NOT supported by Snowflake? (select all that apply)**

Options are :

- Data ingestion and integration
- Enterprise data warehouse
- Enterprise data lake

- *Online Transaction Processing (OLTP) (Correct)*
- *Data engineering and science*
- *Data sharing with partners and consumers*

**Answer :**

\*Online Transaction Processing (OLTP) (Correct)

**Which of the following are NOT among the functions performed by the Cloud Services layer? (select all that apply)**

Options are :

- *Authentication and access control*
- *Query design (Correct)*
- *Query planning*
- *Query processing (Correct)*
- *Query parsing and optimization*
- *Metadata management*

**Answer :**

\*Query design (Correct)

\*Query processing (Correct)

**Which of the following statements are TRUE concerning Snowflake architecture (select all that apply)**

Options are :

- *The virtual warehouse layer in Snowflake is responsible for query compilation.*
- *In Snowflake, two virtual warehouses can concurrently access the same data in the storage layer (Correct)*
- *When you resize a virtual warehouse, the new compute capacity is instantly available for use by both current and subsequent queries.*
- *The cloud services layer in Snowflake is responsible for managing your data across availability zones in your cloud region.*

**Answer :**

\*In Snowflake, two virtual warehouses can concurrently access the same data in the storage layer (Correct)

**Snowflake features remain consistent across cloud providers and regions however the cost of compute and storage may vary depending on where the account is located and whether the account is a capacity account or an on- demand account.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### **Snowflake SnowPro Core Certification Exam Questions Set 3**

**Consider a scenario: You are a data warehouse designer. You come across a tool called ◆DesignEdge◆. The tool automates a large part of your schema design work. You want to connect this tool to your Snowflake account. You look up the partner products in the Snowflake ecosystem (<https://docs.snowflake.com/en/user-guide/ecosystem.html>). You dont find this tool listed there. What would be your next step?**

Options are :

- *Decide against using this tool as it is not part of the Snowflake ecosystem.*
- *Raise a ticket with Snowflake support to verify if this tool is safe to use and if so, Snowflake can add support for this tool.*
- *Look for a 3rd party library that may be able to connect to Snowflake*
- *Check if this tool supports JDBC. If yes, you can use the Snowflake JDBC driver (Correct)*
- *Check if this tool supports ODBC. If yes, you can use the Snowflake ODBC driver (Correct)*

**Answer :**

\*Check if this tool supports JDBC. If yes, you can use the Snowflake JDBC driver (Correct)

\*Check if this tool supports ODBC. If yes, you can use the Snowflake ODBC driver (Correct)

**Which Snowflake edition supports transmitting data in encrypted form over the network between VPCs (virtual private cloud)?**

Options are :

- *All editions*
- *Enterprise edition and above*
- *Business Critical edition and Above (Correct)*
- *All except Virtual Private Snowflake (VPS) edition*

**Answer :**

\*Business Critical edition and Above (Correct)

**Which Snowflake edition(s) supports secure data sharing across regions and cloud platforms?**

Options are :

- *All editions*
- *Enterprise edition and above*
- *Business Critical edition and Above*
- *All except Virtual Private Snowflake (VPS) edition (Correct)*

**Answer :**

\*All except Virtual Private Snowflake (VPS) edition (Correct)

### Snowflake Certification Preparation Set 13

**Which Snowflake edition supports Search Optimization Service to improve performance of point lookup queries?**

Options are :

- *All editions*
- *Enterprise edition and above (Correct)*
- *Business Critical edition and Above*
- *All except Virtual Private Snowflake (VPS) edition*

**Answer :**

\*Enterprise edition and above (Correct)

**When a query is submitted to Snowflake, the following is the two step process that Snowflake follows to return the resultset:**

Options are :

- 1: Prune the tables 2: Prune the micro-partitions
- 1: Prune the micro-partitions 2: Prune the columns (Correct)
- 1: Prune the micro-partitions 2: Prune the rows
- 1: Prune the columns 2: Prune the rows
- 1: Prune the cluster-keys 2: Prune the non-cluster key columns

**Answer :**

\*1: Prune the micro-partitions 2: Prune the columns (Correct)

**Select the statements that apply to clustering depth in a micro-partition. (select all that apply)**

Options are :

- Clustering depth is the average depth of overlapping micro-partitions for a specified set of columns (Correct)
- The lower the clustering depth, the better clustered the table is. (Correct)
- The higher the clustering depth, the better clustered the table is
- Is an indication of whether the table needs an explicit clustering key (Correct)

**Answer :**

\*Clustering depth is the average depth of overlapping micro-partitions for a specified set of columns (Correct)

\*The lower the clustering depth, the better clustered the table is. (Correct)

\*Is an indication of whether the table needs an explicit clustering key (Correct)

### Tackling Snowflake Certification - Practice Questions Set 1

**A clustering key in Snowflake table...(select all that apply):**

Options are :

- Should be used in scenarios where clustering depth is large (Correct)
- Is primarily created to improve query performance for large tables (Correct)
- Once a clustering key is defined for a table, it cannot be altered
- You can define a clustering key on materialized views (Correct)

**Answer :**

- \*Should be used in scenarios where clustering depth is large (Correct)
- \*Is primarily created to improve query performance for large tables (Correct)
- \*You can define a clustering key on materialized views (Correct)

**More than one clustering keys can co-exist in a Snowflake table**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**A column with very high cardinality is generally a good candidate for a clustering key.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

### Snowflake SnowPro Core Certification Practice Tests Set 13

**In a multi-column clustering key, it is recommended that the number of columns in the clustering key should not exceed ..... and the columns are organized from .....**

Options are :

- *Up to 2 columns, lower to higher cardinality*
- *Up to 2 columns, higher to higher cardinality*
- *3 to 4 columns, lower to higher cardinality (Correct)*
- *3 to 4 columns, higher to lower cardinality*
- *Up to 6 columns, lower to higher cardinality*
- *Up to 6 columns, higher to lower cardinality*

**Answer :**

- \**3 to 4 columns, lower to higher cardinality (Correct)*

**Select TRUE statements from the given list of statements below concerning reclustering process in Snowflake (select all that apply):**

Options are :

- *The reclustering process can be initiated by the user as and when it is observed that the same set of queries has become sluggish over time.*
- *It can be suspended and resumed by the user. (Correct)*
- *User will need to provision a warehouse to be used for reclustering*
- *Reclustering may slow down query performance for a temporary period while the process is being carried out by Snowflake in the background*

**Answer :**

\*It can be suspended and resumed by the user. (Correct)

**Selecting the right clustering key can dramatically improve query performance. Which of the following are the right statements to choose an appropriate clustering key:**

Options are :

- *Use columns that are most frequently found in WHERE clause (Correct)*
- *A date column is usually not a good candidate to be in the clustering key*
- *A join column can be considered as clustering key (Correct)*
- *Using expressions in the clustering key may help reduce cardinality. (Correct)*

**Answer :**

\*Use columns that are most frequently found in WHERE clause (Correct)

\*A join column can be considered as clustering key (Correct)

\*Using expressions in the clustering key may help reduce cardinality. (Correct)

### Snowflake Certification Preparation Set 8

**Change in size of a warehouse being used for a query (e.g. M to XL or L to S) will have ...? (select all that apply)**

Options are :

- *Linear impact on the credit usage per hour (Correct)*
- *Linear impact on data loading performance*
- *Linear impact in performance of queries (Correct)*

- *Linear impact on concurrency of queries*

**Answer :**

- \*Linear impact on the credit usage per hour (Correct)
- \*Linear impact in performance of queries (Correct)

**Select statement(s) which are TRUE from the given set of statements concerning Snowflake virtual warehouses. (select all that apply)**

Options are :

- *Scaling up a virtual warehouse is an automatic process*
- *Scaling up a virtual warehouse is a manual process (Correct)*
- *Scaling down a virtual warehouse is an automatic process*
- *Scaling down a virtual warehouse is a manual process (Correct)*
- *Scaling out a virtual warehouse is an automatic process (Correct)*
- *Scaling back a virtual warehouse is an automatic process (Correct)*

**Answer :**

- \*Scaling up a virtual warehouse is a manual process (Correct)
- \*Scaling down a virtual warehouse is a manual process (Correct)
- \*Scaling out a virtual warehouse is an automatic process (Correct)
- \*Scaling back a virtual warehouse is an automatic process (Correct)

**In your production environment, there is a batch process that runs every night from 0200 hrs to 0330 hrs. This batch process requires a cluster of 3 Small sized warehouses. The AUTO\_SUSPEND property for this cluster is set to 5 minutes. The cluster is exclusively used for batch processing purposes. Select valid statement(s) for this scenario from the statements given below:**

Options are :

- *If AUTO\_RESUME is not enabled, the batch process will fail as the cluster is in a suspended state. (Correct)*
- *If AUTO\_RESUME is not enabled, an alert will be sent to the SYSADMIN to remind about resuming the cluster.*
- *If AUTO\_RESUME is enabled, the warehouses will automatically resume servicing the batch processing workload. (Correct)*
- *As soon as you cancel auto\_suspend, the cluster will be ready to service the batch processing workload*

**Answer :**

- \*If AUTO\_RESUME is not enabled, the batch process will fail as the cluster is in a

suspended state. (Correct)

\*If AUTO\_RESUME is enabled, the warehouses will automatically resume servicing the batch processing workload. (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 24**

**When a multi-cluster warehouse is suspended, which of the following Snowflake cache will be purged?**

Options are :

- *Metadata cache*
- *Remote disk cache (Resultset cache)*
- *Local disk cache (Correct)*
- *All of the above*

**Answer :**

\*Local disk cache (Correct)

**Which of the following statements are TRUE concerning network policy features in Snowflake.**

Options are :

- *By configuring a network policy, you specify an allowed list of IPs that cannot access your Snowflake account (Correct)*
- *By configuring a network policy, you specify a blocked list of IPs that can access your Snowflake account (Correct)*
- *You will need the SYSADMIN role to be able to create a network policy*
- *It is possible to bypass the network policy. (Correct)*

**Answer :**

\*By configuring a network policy, you specify an allowed list of IPs that cannot access your Snowflake account (Correct)

\*By configuring a network policy, you specify a blocked list of IPs that can access your Snowflake account (Correct)

\*It is possible to bypass the network policy. (Correct)

**From the list of objects given below, select the objects which are stored within the Schema object. (select all that apply)**

Options are :

- *User*
- *User Stage (Correct)*
- *User Role*
- *Sequence (Correct)*
- *Share*
- *Warehouse*

**Answer :**

\*User Stage (Correct)  
 \*Sequence (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 20**

**Which of the following are TRUE about object relationships in Snowflake?**

Options are :

- *Snowflake tables are stored within schemas (Correct)*
- *Snowflake schemas are stored within databases (Correct)*
- *Users and Roles are stored within databases*
- *Stored procedures and UDFs are stored within warehouses*
- *Warehouses are stored within databases*
- *Databases are stored within accounts (Correct)*

**Answer :**

\*Snowflake tables are stored within schemas (Correct)  
 \*Snowflake schemas are stored within databases (Correct)  
 \*Databases are stored within accounts (Correct)

**In Snowflake, For each securable object, there is a set of privileges. The privileges provide fine-grained access control on the object. Each securable object has an owner that can assign these privileges directly to a user or a group of users.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Every active user session in Snowflake has a role which is referred to as the ♦current role♦ of the user. Which of the following statements are TRUE? (select all that apply)**

Options are :

- Every user logging into Snowflake can specify a role while making the connection request. If this role has been granted to the user, it becomes the ♦current role♦ of the user. (Correct)
- Every user is assigned a default role when the user was created in Snowflake. The next time when the user logs in, it becomes the ♦current role♦ of the user. (Correct)
- Every user logging into Snowflake automatically gets a default role called ♦PUBLIC♦ (Correct)
- Every user logging into Snowflake automatically gets the highest role granted to him in the role hierarchy
- Every user logging into Snowflake automatically gets the lowest role granted to him in the role hierarchy

**Answer :**

\*Every user logging into Snowflake can specify a role while making the connection request. If this role has been granted to the user, it becomes the ♦current role♦ of the user. (Correct)

\*Every user is assigned a default role when the user was created in Snowflake. The next time when the user logs in, it becomes the ♦current role♦ of the user. (Correct)

\*Every user logging into Snowflake automatically gets a default role called ♦PUBLIC♦ (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 23

**Which of the given statements will you use to find out the name of the cloud provider which hosts your Snowflake account?**

Options are :

- Get current\_cloud();
- Get current\_region();
- Return current\_cloud();
- Return current\_region();

- *Select current\_cloud();*
- *Select current\_region(); (Correct)*

**Answer :**

\**Select current\_region(); (Correct)*

**Which of the following statements will you use to change the warehouse for workload processing to a warehouse named `COMPUTE_WH_XL`?**

Options are :

- *SET CURRENT\_WAREHOUSE = COMPUTE\_WH\_XL*
- *USE WAREHOUSE COMPUTE\_WH\_XL; (Correct)*
- *USE CURRENT\_WAREHOUSE(`COMPUTE_WH_XL`);*
- *SET CURRENT\_WAREHOUSE = COMPUTE\_WH, SIZE = XL;*

**Answer :**

\**USE WAREHOUSE COMPUTE\_WH\_XL; (Correct)*

**If you create a sequence in Snowflake with a starting value of 2 and an increment of 3, what would be the output?**

Options are :

- *2, 5, 8, 11, 14... (Correct)*
- *3, 6, 9, 12, 15...*
- *2, 6, 10, 14, 17...*
- *None of above*

**Answer :**

\**2, 5, 8, 11, 14... (Correct)*

### Tackling Snowflake Certification - Practice Questions Set 5

**You are writing a series of SQL statements in Snowflake Web UI. You want to insert remarks before each statement to explain the business logic of the query. You will type the following in front of your remark text to qualify as a comment.**

Options are :

- */\* (a slash and an asterisk)*

- \*\* (a double asterisks) (Correct)
- // (a double slash) (Correct)
- -- (a double dash) (Correct)

**Answer :**

\*// (a double slash) (Correct)  
\*-- (a double dash) (Correct)

**Snowflake being a SaaS software, which of the following activities of an on-prem/hosted data warehouse are not required in Snowflake from the user's perspective? (select all that apply)**

Options are :

- Hardware sizing, purchase, or configuration (Correct)
- Hardware scaling (Correct)
- User and access management
- Software installation, maintenance, or upgrades (Correct)
- Database Tuning (Correct)
- Site-level disaster recovery (due to loss of data center) (Correct)

**Answer :**

\*Hardware sizing, purchase, or configuration (Correct)  
\*Hardware scaling (Correct)  
\*Software installation, maintenance, or upgrades (Correct)  
\*Database Tuning (Correct)  
\*Site-level disaster recovery (due to loss of data center) (Correct)

**In the case of a Snowflake account created on AWS, ..... is responsible for the management of Availability Zones?**

Options are :

- Customer
- Snowflake
- Cloud Provider (Correct)
- It is a shared responsibility

**Answer :**

\*Cloud Provider (Correct)

## Snowflake SnowPro Core Certification Practice Test Set 2

From the URL given below, which of the following information can be derived? (select all that apply) <https://wh83035.ca-central-1.aws.snowflakecomputing.com/>

Options are :

- *Cloud provider (Correct)*
- *Deployment region (Correct)*
- *Snowflake Edition*
- *Account Name (Correct)*

**Answer :**

\*Cloud provider (Correct)  
\*Deployment region (Correct)  
\*Account Name (Correct)

In a Snowflake account named wh83035, there is a user named BJW. You need to set the default namespace of this user to a database called WK\_DB and Schema WK\_SCHEMA. Which of the following SQL commands would you use?

Options are :

- *ALTER USER BJW SET DEFAULT\_NAMESPACE = WK\_DB.WK\_SCHEMA; (Correct)*
- *SET DEFAULT\_NAMESPACE = wh83085.WK\_DB.WK\_SCHEMA*
- *USE DEFAULT\_NAMESPACE WK\_DB.WK\_SCHEMA*
- *SET DEFAULT\_NAMESPACE (WK\_DB.WK\_SCHEMA?)*
- *None of the above*

**Answer :**

\**ALTER USER BJW SET DEFAULT\_NAMESPACE = WK\_DB.WK\_SCHEMA; (Correct)*

**A temporary table in Snowflake... (select all that apply)**

Options are :

- *Has session-level data retention (Correct)*
- *Does not support fail-safe (Correct)*
- *Does not support zero-copy cloning (Correct)*

- Offers slower query performance compared to the other table types.

**Answer :**

- \*Has session-level data retention (Correct)
- \*Does not support fail-safe (Correct)
- \*Does not support zero-copy cloning (Correct)

## Snowflake Certification Preparation Set 4

**The fail-Safe feature is available for which table types? (Select all that apply)**

Options are :

- Permanent (Correct)
- Temporary
- Transient
- External

**Answer :**

- \*Permanent (Correct)

**Which of the following statements are true concerning zero-copy cloning in Snowflake? (select all that apply)**

Options are :

- Zero-copy cloning is a metadata-only operation and does not entail creating a copy of the data (Correct)
- The original object and the cloned object are independent of each other. (Correct)
- The cloned object created by zero-copy cloning is read-only
- Zero-copy cloning is recursive (Correct)
- Zero-copy cloning feature allows you to backup your data across your Snowflake accounts

**Answer :**

- \*Zero-copy cloning is a metadata-only operation and does not entail creating a copy of the data (Correct)
- \*The original object and the cloned object are independent of each other. (Correct)
- \*Zero-copy cloning is recursive (Correct)

**Zero-copy cloning operation is supported for the following objects in Snowflake. (select all that apply)**

Options are :

- *A temporary Table*
- *A transient Table (Correct)*
- *An external table*
- *A permanent Table (Correct)*
- *A database created from a Share*
- *Internal (named) Stage*

**Answer :**

- \* *A transient Table (Correct)*
- \* *A permanent Table (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 23](#)

**What do resource monitors in Snowflake measure? (select all that apply)**

Options are :

- *Credits consumed by user-managed virtual warehouses (Correct)*
- *Credits consumed by snowflake-provisioned virtual warehouses (i.e. warehouses used by the cloud services layer) (Correct)*
- *Storage consumed by all user-created Snowflake objects (e.g. table, stage, etc.)*
- *Storage consumed by metadata layer in Snowflake*
- *Storage consumed by all Snowflake Caches (i.e. metadata cache, warehouse cache, resultset cache)*

**Answer :**

- \* *Credits consumed by user-managed virtual warehouses (Correct)*
- \* *Credits consumed by snowflake-provisioned virtual warehouses (i.e. warehouses used by the cloud services layer) (Correct)*

**What actions can a resource monitor trigger?**

Options are :

- *Send an email or a notification on the web to the account administrators when the credit quota threshold is reached (Correct)*

- Send an email or a notification on the web to the active warehouse user(s) when the credit quota threshold is reached
- Suspend a warehouse when the credit quota threshold is reached (Correct)
- Suspend a warehouse immediately when the credit quota threshold is reached (Correct)
- Suspend a warehouse at a specific date and time or time interval

**Answer :**

\*Send an email or a notification on the web to the account administrators when the credit quota threshold is reached (Correct)  
 \*Suspend a warehouse when the credit quota threshold is reached (Correct)  
 \*Suspend a warehouse immediately when the credit quota threshold is reached (Correct)

**At the following frequency interval, the resource monitor quota will be reset. (select all that apply)**

Options are :

- Daily (Correct)
- Weekly (Correct)
- Monthly (Correct)
- Quarterly
- Yearly (Correct)
- Never (Correct)

**Answer :**

\*Daily (Correct)  
 \*Weekly (Correct)  
 \*Monthly (Correct)  
 \*Yearly (Correct)  
 \*Never (Correct)

**Snowflake Certification Preparation Set 1**

**A resource monitor can be created by.....**

Options are :

- Owner of the virtual warehouse

- *ACCOUNTADMIN role (Correct)*
- *SYSADMIN role*
- *USERADMIN role*
- *Any of the above*

**Answer :**

\**ACCOUNTADMIN role (Correct)*

**When a resource monitor action is set to ◊Suspend Immediately◊ and its credit quota threshold is reached, Snowflake will immediately cancel any statement being executed by the warehouse and also any statements which are in the queue. This ensures that no additional credits would be consumed.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which of the following security features are available in all Snowflake editions? 1. Data encryption ◊at rest◊ using strong AES-256 bit encryption 2. Private communication between Snowflake and your other VPCs through AWS PrivateLink (or Azure PrivateLink). 3. MFA (multi-factor authentication) for account access by users. 4. Support for user SSO (single sign-on). 5. Support for encrypting data using customer-managed keys. 6. HIPAA and PCI-DSS security standards compliance**

Options are :

- *1, 3 and 4 (Correct)*
- *1, 3, 4 and 5*
- *All except 6*
- *All of the Above*

**Answer :**

\**1, 3 and 4 (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 28](#)

**You have written a complex query that will be used in a CXO level end-of-day report. The query uses several tables and performs complex computational operations. The execution time of this query is 10 minutes. You executed the query and checked the output; it was as expected. In the**

**evening, you receive an email from the user asking for the query output. What would be the most optimal approach to share the query results with the user?**

Options are :

- Since you didn't save the query, you rewrite the query and re-run it. Then, download the results from the results pane of the worksheet in Web UI.
- You go to Query\_History View in Web UI ? Locate the query ? Copy the query. Re-run the query in the worksheet and download the results.
- You Go to Query\_History View in Web UI ? Locate the query ? Click on Query ID ? Use  button (Correct)
- You believe that Snowflake cache might have the query results. So you contact SYSADMIN to access the Query\_History view and put a request to download the query results.

**Answer :**

\*You Go to Query\_History View in Web UI ? Locate the query ? Click on Query ID ? Use  button (Correct)

The  information in Snowflake... (select all that apply)

Options are :

- Shows execution details of a query in a graphical format along with execution time and statistics of each component of the query. (Correct)
- Can be accessed using Snowflake Web UI (Correct)
- Is available only for completed queries (Correct)
- Is available for 24 hours
- Is accessible only to the ACCOUNT\_ADMIN role

**Answer :**

\*Shows execution details of a query in a graphical format along with execution time and statistics of each component of the query. (Correct)

\*Can be accessed using Snowflake Web UI (Correct)

\*Is available only for completed queries (Correct)

**Which statements are true of micro-partitions? (Choose two.)**

Options are :

- They are approximately 16MB in size (Correct)

- They are stored compressed only if `COMPRESS=TRUE` on Table
- They are immutable (Correct)
- They are only encrypted in the Enterprise edition and above

**Answer :**

- \*They are approximately 16MB in size (Correct)
- \*They are immutable (Correct)

**When should you consider disabling auto-suspend for a Virtual Warehouse? (Choose two.)**

Options are :

- When users will be using compute at different times throughout a 24/7 period
- When managing a steady workload (Correct)
- When the compute must be available with no delay or lag time (Correct)
- When you do not want to have to manually turn on the Warehouse each time a user needs it

**Answer :**

- \*When managing a steady workload (Correct)
- \*When the compute must be available with no delay or lag time (Correct)

**In which layer of its architecture does Snowflake store its metadata statistics?**

Options are :

- Storage Layer
- Compute Layer
- Database Layer
- Cloud Services Layer (Correct)

**Answer :**

- \*Cloud Services Layer (Correct)

**Snowflake is designed for which type of workloads? (Choose two.)**

Options are :

- OLAP (Analytics) workloads (Correct)
- OLTP (Transactional) workloads
- Concurrent workloads (Correct)
- On-premise workloads

**Answer :**

\*OLAP (Analytics) workloads (Correct)  
\*Concurrent workloads (Correct)

## Snowflake Certification Preparation Set 16

**Which statement best describes Snowflake Clustering?**

Options are :

- Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions (Correct)
- The database administrator must define the clustering methodology for each Snowflake table
- The clustering key must be included on the COPY command when loading data into Snowflake
- Clustering can be disabled within a Snowflake account

**Answer :**

\*Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions (Correct)

**Snowflake provides a mechanism for its customers to override its natural clustering algorithms.**  
**This method is:**

Options are :

- Micro-partitions
- Clustering keys (Correct)
- Key partitions
- Clustered partitions

**Answer :**

\*Clustering keys (Correct)

**In the following scenario, for how much time will you be billed for compute ? \* Your ran a query for 72 seconds and then shut-down the virtual warehouse \* After few hours, you run another query for 20 seconds and shut-down the virtual warehouse**

Options are :

- 92 seconds
- 102 seconds
- 122 seconds
- 132 seconds (Correct)
- None of the above

**Answer :**

\*132 seconds (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 1

**Which Snowflake Edition limit Time Travel to maximum 24 hours ?**

Options are :

- Standard Edition (Correct)
- Business Critical Edition
- Virtual Private Snowflake
- Enterprise Edition
- None of the above

**Answer :**

\*Standard Edition (Correct)

**What kind of solution is Snowflake**

Options are :

- Infrastructure as a Service (IaaS)
- Platform as a Service (PaaS)

- *Software as a Service (SaaS) (Correct)*

**Answer :**

\*Software as a Service (SaaS) (Correct)

**In which mode Snowflake multi-cluster warehouses can run ? Select all that apply.**

Options are :

- *Maximized (Correct)*
- *Auto-resume*
- *Auto-suspending*
- *Auto-scale (Correct)*

**Answer :**

\*Maximized (Correct)

\*Auto-scale (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 12](#)

**How long does Bulk data load 'Load history' last ?**

Options are :

- 14
- 28
- 32
- 64 (*Correct*)

**Answer :**

\*64 (Correct)

**Which table do not have Fail-safe period ? Select all that apply**

Options are :

- *Permanent*
- *Temporary (Correct)*
- *Transient (Correct)*

- *External*
- *None of the above*

**Answer :**

- \*Temporary (Correct)
- \*Transient (Correct)

**Which of the following terms describes Snowflake's architecture ?**

Options are :

- *Multi-cluster shared data (Correct)*
- *Shared disk*
- *Columnar shared nothing*
- *Cloud-native shared memory*

**Answer :**

- \*Multi-cluster shared data (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 6

**What services are managed by Cloud Services layer ? (Select all that apply)**

Options are :

- *Authentication (Correct)*
- *Infrastructure Management (Correct)*
- *Metadata Management (Correct)*
- *Query parsing and optimization*
- *Data storage*
- *Query processing*

**Answer :**

- \*Authentication (Correct)
- \*Infrastructure Management (Correct)
- \*Metadata Management (Correct)

**Which of the following terms are associated with the Cloud Services Layer ?**

Options are :

- *Query Processing*
- *Query Planning (Correct)*
- *Query Optimization (Correct)*
- *Query Design (Correct)*
- *Query Compilation (Correct)*

**Answer :**

- \*Query Planning (Correct)
- \*Query Optimization (Correct)
- \*Query Design (Correct)
- \*Query Compilation (Correct)

**Which cache runs on a 24 hour clock?**

Options are :

- *Result cache (Correct)*
- *Metadata Cache*
- *Warehouse cache*

**Answer :**

- \*Result cache (Correct)

### [Snowflake Practice Questions - SnowPro Core Certified Set 7](#)

**What Snowflake objects can be shared ? (Select all that apply)**

Options are :

- *Tables (Correct)*
- *Views*
- *Roles*
- *Secure UDFs (Correct)*
- *Materialized views*
- *None of the above*

**Answer :**

- \*Tables (Correct)
- \*Secure UDFs (Correct)

**Virtual warehouses cannot be resized while they are running. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**How many credits consume a Large-size multi-cluster warehouse with 3 warehouses which runs in Maximized mode for 2 hours**

Options are :

- 24
- 48 (*Correct*)
- 16
- 56

**Answer :**

- \**48 (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 5**

**A Snowflake session can NOT have more than one virtual warehouse. TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \**TRUE (Correct)*

**Which type of view has an extra layer of protection to hide the SQL code from unauthorized viewing ?**

Options are :

- *Standard*
- *Secure (Correct)*
- *Materialized*
- *Permanent*

**Answer :**

\*Secure (Correct)

**Which of the following have drivers/connectors (or information about where to find them) available via Help->Downloads in the Snowflake WebUI?**

Options are :

- *Go (Correct)*
- *R*
- *Node.Js (Correct)*
- *JDBC (Correct)*
- *Hive*
- *Spark*

**Answer :**

\*Go (Correct)  
\*Node.Js (Correct)  
\*JDBC (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 2

**Scaling a warehouse IN will decrease the number of clusters. (e.g. Max to Min). TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### **Which Snowflake Edition can a user choose when creating a trial account**

Options are :

- *Standard Edition (Correct)*
- *Enterprise Edition (Correct)*
- *Free-trial Edition*
- *Virtual Private Snowflake*
- *Business Critical Edition (Correct)*

### **Answer :**

\*Standard Edition (Correct)  
\*Enterprise Edition (Correct)  
\*Business Critical Edition (Correct)

### **What is called the action of adding more servers to the same cluster of a virtual warehouse ?**

Options are :

- *Scale UP (Correct)*
- *Scale DOWN*
- *Scale IN*
- *Scale OUT*
- *None of the above*

### **Answer :**

\*Scale UP (Correct)

### **Tackling Snowflake Certification - Practice Questions Set 4**

#### **Who can use the metadata cache ?**

Options are :

- *Only the user who execute the query*
- *Only user with the same role*
- *Only account administrator*

- Everyone (Correct)
- None of the above

**Answer :**

\*Everyone (Correct)

**How does Snowflake bill for compute ?**

Options are :

- Per hour, with one-minute minimum
- Per second, with one-minute minimum (Correct)
- Per day
- Per month

**Answer :**

\*Per second, with one-minute minimum (Correct)

**What is the account identifier inside this Snowflake URL ' <https://jz96798.us-east-1.azure.snowflakecomputing.com> ' ?**

Options are :

- jz96798.us-east-1.azure
- us-east-1.azure
- jz96798 (Correct)
- snowflakecomputing.com
- Azure

**Answer :**

\*jz96798 (Correct)

**SnowFlake SnowPro Advanced Architect Practice Exam Set 1**

**Users created in a reader account CAN NOT do which of the following actions ?**

Options are :

- INSERT (Correct)

- *SELECT*
- *LOAD DATA (Correct)*
- *UPDATE (Correct)*

**Answer :**

- \**INSERT (Correct)*
- \**LOAD DATA (Correct)*
- \**UPDATE (Correct)*

**Which syntax allows you to concat the column named 'EMPLOYEE' between two % signs ? Select all that apply**

Options are :

- *%' || EMPLOYEE || '%' (Correct)*
- *% + EMPLOYEE + %*
- *CONCAT('%', EMPLOYEE, '%')*
- *concat('%', concat(EMPLOYEE, '%')) (Correct)*
- *All of the above*

**Answer :**

- \**%' || EMPLOYEE || '%' (Correct)*
- \**concat('%', concat(EMPLOYEE, '%')) (Correct)*

**In snowflake container hierarchy, what container is represented as a URL ( e.g. <https://jz96798.us-east-1.azure.snowflakecomputing.com> ) ?**

Options are :

- *Database*
- *Schema*
- *Role*
- *Account (Correct)*
- *Region*

**Answer :**

- \**Account (Correct)*

**[Tackling Snowflake Certification - Practice Questions Set 1](#)**

### **What attributes make Snowflake a true SaaS solution ?**

Options are :

- *No hardware to purchase or configure (Correct)*
- *No creation of user accounts or roles is required*
- *No maintenance upgrades or patches to install (Correct)*
- *No data storage costs*
- *No query processing costs*
- *Transparent releases don't required user intervention (Correct)*

**Answer :**

- \*No hardware to purchase or configure (Correct)
- \*No maintenance upgrades or patches to install (Correct)
- \*Transparent releases don't required user intervention (Correct)

### **How long does Snowpipe History Load last ?**

Options are :

- *7 days*
- *14 days (Correct)*
- *21 days*
- *28 days*

**Answer :**

- \*14 days (Correct)

### **When loading JSON file inside Snowflake table, which file format option remove the outer array structure and load the records into separate table rows ? (**

Options are :

- *STRIP\_NULL\_VALUES = TRUE*
- *STRIP\_OUTER\_ARRAY = TRUE (Correct)*
- *TRIM\_SPACE = TRUE*
- *None of the above*

**Answer :**

\*STRIP\_OUTER\_ARRAY = TRUE (Correct)

## Tackling Snowflake Certification - Practice Questions Set 2

**What is the maximum size for a VARIANT data type column ?**

Options are :

- 64MB uncompressed
- 16MB uncompressed (Correct)
- 32MB uncompressed
- 128MB uncompressed

**Answer :**

\*16MB uncompressed (Correct)

**Scaling policy for a multi-cluster warehouse ONLY apply if it is running in Auto-scale mode. TRUE or FALSE**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Which one of the below requires an user-specified warehouse to execute COPY statements**

Options are :

- Bulk data load (Correct)
- Snowpipe
- All of the above
- None of the above

**Answer :**

\*Bulk data load (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 2

If a length is not specified when defined a VARCHAR column the default length is ?

Options are :

- 256 characters
- 1 character
- 0 character
- Maximum length possible (Correct)
- None of the above

**Answer :**

\*Maximum length possible (Correct)

Select the three terms refering to the same layer.

Options are :

- Compute Layer (Correct)
- Virtual Warehouse Layer (Correct)
- Catalog Layer
- Metadata Layer
- Query Processing Layer (Correct)

**Answer :**

\*Compute Layer (Correct)

\*Virtual Warehouse Layer (Correct)

\*Query Processing Layer (Correct)

**It is possible to install Snowflake on-premise. TRUE or FALSE**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

## Snowflake SnowPro Advanced Architect Practice Exam Set 3

**Which role or privileges are required to view TASK\_HISTORY ? Select all that apply**

Options are :

- Account administrator (i.e users with ACCOUNTADMIN role) (Correct)
- Task Owner (i.e. role that has the ownership privilege on a task) (Correct)
- Anyone with SYSADMIN role
- Any role that has the global MONITOR EXECUTION privilege
- Any role in the account

**Answer :**

\*Account administrator (i.e users with ACCOUNTADMIN role) (Correct)

\*Task Owner (i.e. role that has the ownership privilege on a task) (Correct)

**Streams support repeatable read isolation. TRUE or FALSE ?**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Which the right options to convert 'AMOUNT' column to DECIMAL(10, 2) inside a select statement**

Options are :

- SELECT AMOUNT::DECIMAL(10,5) (Correct)
- SELECT AMOUNT( DECIMAL(10,2) )
- SELECT AMOUNT as DECIMAL(10,2) (Correct)
- SELECT AMOUNT.DECIMAL(10,2)

**Answer :**

\*SELECT AMOUNT::DECIMAL(10,5) (Correct)

\*SELECT AMOUNT as DECIMAL(10,2) (Correct)

## Snowflake Certification Preparation Set 4

**In Snowflake Web UI Console, in which tab can user check for Query performance ?**

Options are :

- Warehouses Tab
- Databases Tab
- History Tab (Correct)
- Account Tab
- None of the above

**Answer :**

\*History Tab (Correct)

**Which file formats is not supported by Snowflake to load data ?**

Options are :

- Delimited files
- JSON
- Avro
- ORC
- Parquet
- PDF (Correct)

**Answer :**

\*PDF (Correct)

**When data is staged to an internal stage using PUT command, the data is encrypted on client machine. TRUE or FALSE**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

### Snowflake Certification Preparation Set 23

**Which of the following are types of cache use by Snowflake ? Select all that apply**

Options are :

- *Metadata cache (Correct)*
- *Index cache*
- *Query result cache (Correct)*
- *Data cache (Correct)*
- *None of the above*

**Answer :**

\*Metadata cache (Correct)

\*Query result cache (Correct)

\*Data cache (Correct)

**Which Snowflake option enable to query data stored in external cloud storage without loading it into Snowflake**

Options are :

- *External views*
- *Materialized views*
- *External Tables (Correct)*

**Answer :**

\*External Tables (Correct)

**What attributes make Snowflake a true SaaS solution?**

Options are :

- *No hardware to install or configure (Correct)*

- *No maintenance upgrades or patches to install (Correct)*
- *Transparent releases don't require user intervention (Correct)*
- *No creation of user accounts or roles is required*
- *No query processing costs*
- *No data storage costs*

**Answer :**

- \*No hardware to install or configure (Correct)
- \*No maintenance upgrades or patches to install (Correct)
- \*Transparent releases don't require user intervention (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 12

**What kind of stream can be created on external table ?**

Options are :

- *Standard*
- *Append-only*
- *Insert-only (Correct)*
- *None of the above*

**Answer :**

- \*Insert-only (Correct)

**Multi-cluster warehouses is only available in Enterprise Edition and above**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**The expected average run time is the average difference between the scheduled and completed times for a task, excluding any time period in which the task was queued. TRUE or FALSE ?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 22

If you find a data-related tool that is not listed as part of the Snowflake ecosystem, what industry standard options could you check for as a way to easily connect to Snowflake?

Options are :

- Check see if the tool can connect to other solution via JDBC (Correct)
- Check see if the tool can connect to other solution via ODBC (Correct)
- Check see if there is a petition in the community to create a driver
- Check see if you can develop a driver and put it on github

**Answer :**

\*Check see if the tool can connect to other solution via JDBC (Correct)

\*Check see if the tool can connect to other solution via ODBC (Correct)

Which command allows to upload data files from a local directory/folder on a client machine to Snowflake stages ?

Options are :

- COPY
- PUT (Correct)
- UPLOAD
- GET

**Answer :**

\*PUT (Correct)

A task tree can have task spanning multiple schemas ?

Options are :

- TRUE

- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**When you share data with a consumer account, who pays for compute time when data is access by the consumer?**

Options are :

- *The data provider does*
- *The data consumer does. (Correct)*
- *Both provider and consumer does*
- *None of the above*

**Answer :**

\*The data consumer does. (Correct)

**Which installment option version of Snowflake are NOT available ? Select all that apply.**

Options are :

- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure)*
- *Microsoft Cloud Native Accounts (Correct)*
- *Hybrid On-Premise + Cloud Installation (Correct)*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure)*
- *Enterprise In-House VPC Installation (Correct)*
- *On-Premise Custom Installation (Correct)*

**Answer :**

\*Microsoft Cloud Native Accounts (Correct)  
\*Hybrid On-Premise + Cloud Installation (Correct)  
\*Enterprise In-House VPC Installation (Correct)  
\*On-Premise Custom Installation (Correct)

**Your organization wants to use Column-level security to apply masking policies to columns in tables or views. Which Snowflake's edition offers this features ? Select all that apply.**

Options are :

- *Enterprise Edition (Correct)*
- *Business Critical Edition (Correct)*
- *Virtual Private Snowflake (Correct)*
- *Standard*
- *None of the above*

**Answer :**

- \*Enterprise Edition (Correct)
- \*Business Critical Edition (Correct)
- \*Virtual Private Snowflake (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 36

**Multiple warehouses can access the same data simultaneously without resource contention issues.  
TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**What ROLE should create users and roles in Snowflake ? Select all that apply**

Options are :

- *ACCOUNTADMIN*
- *USERADMIN (Correct)*
- *SYSADMIN*
- *SECURITYADMIN (Correct)*
- *PUBLIC*

**Answer :**

- \**USERADMIN (Correct)*
- \**SECURITYADMIN (Correct)*

**What are the two major costs categories in Snowflake ? (Select the best options)**

Options are :

- *Storage (Correct)*
- *Compute (Correct)*
- *Data transfer*
- *Cloud services*
- *None of the above*

**Answer :**

\*Storage (Correct)  
\*Compute (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 5](#)

**Which of the following terms are NOT associated with the Cloud Services Layer ?**

Options are :

- *Query Planning*
- *Query Optimization*
- *Query Processing (Correct)*
- *Query Design*
- *Query Compilation*

**Answer :**

\*Query Processing (Correct)

**Via REST API, Snowpipe can only reference external stages as a source.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following techniques can be used in Snowflake to improve query performance?**

Options are :

- *Distribution keys*
- *Index definition*
- *Clustering keys (Correct)*
- *Reclustering (Correct)*
- *None of the above*

**Answer :**

\*Clustering keys (Correct)  
\*Reclustering (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 8](#)

**Snowflake can automatically scale your warehouse UP when needed (make your warehouse larger). TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**A company operating in the health sector wants to move to Snowflake. The company asks you about the edition of Snowflake to choose to comply with HIPAA regulation. Select the minimum edition of Snowflake to recommend ?**

Options are :

- *Enterprise Edition*
- *Business Critical Edition (Correct)*
- *Standard Edition*
- *Virtual Private Snowflake*

**Answer :**

\*Business Critical Edition (Correct)

**What factors impact credit consumption of the compute layer ? Select all that apply**

Options are :

- *Amount of data processed*
- *Size of warehouse (Correct)*
- *Number of clusters (Correct)*
- *Number of user*

**Answer :**

\*Size of warehouse (Correct)  
\*Number of clusters (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 3

**It is possible to set a storage quota at the account level. TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**A Snowflake session can have more than one virtual warehouse. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Snowflake compute costs depends on : (Select all that apply) ?**

Options are :

- *The amount of time warehouse have run (Correct)*
- *The number of rows returned in queries*

- *The sizes of running warehouses (Correct)*
- *The amount of data stored in the account*
- *None of the above*

**Answer :**

- \*The amount of time warehouse have run (Correct)
- \*The sizes of running warehouses (Correct)

### Snowflake Certification Preparation Set 15

**Scaling a warehouse UP will decrease the number of clusters. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**How many task can we have in a simple tree ?**

Options are :

- *10*
- *100*
- *1000 (Correct)*
- *10000*

**Answer :**

- \**1000 (Correct)*

**Which transformations can be applied on data while loading it into a table using COPY command ?  
Select all that apply**

Options are :

- *column reordering (Correct)*
- *Aggregation*
- *Casts (Correct)*

- *Filtering*
- *Column omission (Correct)*
- *Truncating text strings that exceed the target column length (Correct)*

**Answer :**

- \*column reordering (Correct)
- \*Casts (Correct)
- \*Column omission (Correct)
- \*Truncating text strings that exceed the target column length (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 22**

**Which of the following data types are treated as VARCHAR by Snowflake ?**

Options are :

- *TEXT (Correct)*
- *STRING (Correct)*
- *CHAR (Correct)*
- *VARCHAR (Correct)*
- *NVARCHAR*

**Answer :**

- \*TEXT (Correct)
- \*STRING (Correct)
- \*CHAR (Correct)
- \*VARCHAR (Correct)

**Which cloud platforms can a Snowflake account be hosted on ?**

Options are :

- *Google Cloud Platform (GCP) (Correct)*
- *OVHcloud*
- *Microsoft Azure (Azure) (Correct)*
- *Alibaba Cloud*
- *Amazon Web Services (AWS) (Correct)*
- *None of the above*

**Answer :**

- \*Google Cloud Platform (GCP) (Correct)
- \*Microsoft Azure (Azure) (Correct)
- \*Amazon Web Services (AWS) (Correct)

**What is the name of the Snowflake-produced Command Line Interface tool?**

Options are :

- *SnowCLI*
- *SnowCommand*
- *SnowSQL (Correct)*
- *SnowSpark*
- *None of the above*

**Answer :**

- \**SnowSQL (Correct)*

### [Snowflake Practice Questions - SnowPro Core Certified Set 8](#)

**A client has set Time Travel to keep a table data for 27 days. 35 days after dropping the table, their want to get it back. They ask you if it is possible to recover the data of that table.**

Options are :

- *Yes, it is possible*
- *No, it is not possible (Correct)*

**Answer :**

- \**No, it is not possible (Correct)*

**A company operating in the health sector wants to move to Snowflake. The company asks you about the edition of Snowflake to choose to comply with HIPAA regulation. Select the minimum edition of Snowflake to recommend ?**

Options are :

- *Enterprise Edition*
- *Business Critical Edition (Correct)*
- *Standard Edition*

- *Virtual Private Snowflake*

**Answer :**

\*Business Critical Edition (Correct)

**Snowflake account hosted on AWS cannot load data from a file that is staged on GCP or Azure.  
TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### [Snowflake Certification Preparation Set 23](#)

**What are the two mechanisms that can be used with Snowpipe for detecting that staged files are available ?**

Options are :

- *Calling Snowpipe REST endpoints (Correct)*
- *Calling Snowpipe SOAP endpoints*
- *Automating Snowpipe using CloudWatch*
- *Automating Snowpipe using cloud messaging (Correct)*

**Answer :**

\*Calling Snowpipe REST endpoints (Correct)

\*Automating Snowpipe using cloud messaging (Correct)

**What are the two scaling policies for a multi-cluster warehouse ?**

Options are :

- *Standard (Correct)*
- *Enterprise*
- *Economy (Correct)*
- *Business Critical*

**Answer :**

- \*Standard (Correct)
- \*Economy (Correct)

**Each Snowflake account is limited to a maximum of 10 shares**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

### [Snowflake SnowPro Core Certification Exam Questions Set 11](#)

**What kind of metadata does Snowflake stores in micro-partition for a table**

Options are :

- *MIN (Correct)*
- *MAX (Correct)*
- *Table's privileges*
- *Table's owner*
- *Number of distinct values (Correct)*
- *None of the above*

**Answer :**

- \**MIN (Correct)*
- \**MAX (Correct)*
- \**Number of distinct values (Correct)*

**Which of the following objects can be shared ? Select all that apply**

Options are :

- *Secure views (Correct)*
- *Views*
- *Schemas (Correct)*

- *Table (Correct)*
- *Databases (Correct)*
- *Secure UDFs (Correct)*

**Answer :**

- \*Secure views (Correct)
- \*Schemas (Correct)
- \*Table (Correct)
- \*Databases (Correct)
- \*Secure UDFs (Correct)

**What privileges does SYSADMIN have by default?**

Options are :

- *CREATE DATABASE (Correct)*
- *CREATE SHARE*
- *CREATE WAREHOUSE (Correct)*
- *CREATE USER*
- *CREATE ROLE*

**Answer :**

- \**CREATE DATABASE (Correct)*
- \**CREATE WAREHOUSE (Correct)*

## Snowflake Certification Preparation Set 12

**Which installment option version of Snowflake are available ?**

Options are :

- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)*
- *Microsoft Cloud Native Accounts*
- *Hybrid On-Premise + Cloud Installation*
- *Enterprise In-House VPC Installation*
- *On-Premise Custom Installation*

**Answer :**

- \*Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)
- \*Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)

**Which compute resource can be used by Snowflake for data loading jobs ? Select all that apply**

Options are :

- *User-managed virtual warehouse (Correct)*
- *Hardware provisionned by user directly from cloud provider*
- *Snowflake-managed serverless compute (Correct)*
- *Cloud Provider Runtime Environment*

**Answer :**

- \*User-managed virtual warehouse (Correct)
- \*Snowflake-managed serverless compute (Correct)

**What is the region identifier inside this Snowflake URL << https://jz96798.us-east-1.azure.snowflakecomputing.com >> ?**

Options are :

- *jz96798.us-east-1.azure*
- *us-east-1.azure (Correct)*
- *jz96798*
- *snowflakecomputing.com*

**Answer :**

- \**us-east-1.azure (Correct)*

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 4**

**In which mode Snowflake multi-cluster warehouses can run ? Select all that apply.**

Options are :

- *Maximized (Correct)*
- *Auto-resume*
- *Auto-suspended*

- *Auto-scale (Correct)*
- *None of the above*

**Answer :**

- \*Maximized (Correct)
- \*Auto-scale (Correct)

**Is it possible to disable Time Travel for a Snowflake Account ?**

Options are :

- *True*
- *False (Correct)*

**Answer :**

- \*False (Correct)

**Which cache type gets purged in a predictable way so that it ends up empty of all cached information ?**

Options are :

- *Metadata cache*
- *Result cache*
- *Warehouse cache (Correct)*

**Answer :**

- \*Warehouse cache (Correct)

**Query load is calculated by dividing the execution time (in seconds) of all queries in an interval by the total time (in seconds) for the interval. TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*

- FALSE

**Answer :**

\*TRUE (Correct)

**What are the elements of a worksheet context ?**

Options are :

- *Role, Database, Schema, Warehouse (Correct)*
- *User, Role, Database, Schema*
- *Account, User, Role, Warehouse*
- *Account, Database, Schema, Warehouse*

**Answer :**

\*Role, Database, Schema, Warehouse (Correct)

**In the Query History tab, how can you identify if a query pulled data from long-term centralized storage. Select the best option**

Options are :

- *Look in the Bytes scanned column for a green bar (Correct)*
- *Look in the Bytes scanned column for a green bar*
- *Look in the Total Duration column for a orange bar*
- *Look in the Total Duration column for a green bar*

**Answer :**

\*Look in the Bytes scanned column for a green bar (Correct)

### Snowflake Certification Preparation Set 1

**What does this cron expression (USING CRON 0 0-23 \* \* \*) express ?**

Options are :

- *Every minute of every hour of every day*
- *At 00:00 am on the 1st of every month*
- *The first minute of every hour of every day (Correct)*

- At 00:00 am on the 1st of every month

**Answer :**

\*The first minute of every hour of every day (Correct)

**Which of the following Snowflake objects allows to query data without loading it first inside Snowflake ?**

Options are :

- *Snowpipe*
- *COPY*
- *External table (Correct)*
- *Virtual table*

**Answer :**

\*External table (Correct)

**Which Tech Partner types are available from in-account menu items ?**

Options are :

- *Partner Connect (Correct)*
- *Certified ANSI*
- *Programmatic Interfaces (Correct)*
- *Founding Partnership Companies*
- *HIPAA Certified*

**Answer :**

\*Partner Connect (Correct)

\*Programmatic Interfaces (Correct)

### [SnowFlake SnowPro Advanced Architect Practice Exam Set 6](#)

**Which type of stages are automatically available in Snowflake and do not need to be created or configured ?**

Options are :

- *User (Correct)*

- *Table (Correct)*
- *Named Internal*
- *Named External*
- *None of the above*

**Answer :**

\*User (Correct)  
\*Table (Correct)

#### What influences Snowflake pricing ?

Options are :

- *Storage and Compute (Correct)*
- *Amount of data queried from Snowflake*
- *Amount of data scanned during querying Snowflake*
- *None of the above*

**Answer :**

\*Storage and Compute (Correct)

**Which scaling policy, starts additional cluster immediately when either a query is queued or the system detects that there is one more query than the currently-running clusters can execute.**

Options are :

- *Economy*
- *Standard (Correct)*

**Answer :**

\*Standard (Correct)

#### SnowFlake SnowPro Advanced Architect Practice Exam Set 4

**It is possible to the latest ODBC connector for use with Snowflake through Snowflake Web UI console. TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

**Snowflake can be run on private cloud infrastructures (on-premises or hosted). TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**What is the maximum number of child a single task can have in Snowflake ?**

Options are :

- *1*
- *100 (Correct)*
- *1000*
- *10000*

**Answer :**

\*100 (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 39**

**Scaling OUT is an automated process (e.g. Min clusters to max clusters). TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which factors affect your snowflake data loading rate ?**

Options are :

- *Physical location of the stage (Correct)*
- *RAM on the virtual warehouse*
- *Thread size*
- *GZIP compression efficiency (Correct)*

**Answer :**

\*Physical location of the stage (Correct)  
\*GZIP compression efficiency (Correct)

**Select the three Snowpipe REST API endpoints.**

Options are :

- *insertFiles (Correct)*
- *insertData*
- *insertReport (Correct)*
- *loadHistoryScan (Correct)*
- *loadHistory*

**Answer :**

\*insertFiles (Correct)  
\*insertReport (Correct)  
\*loadHistoryScan (Correct)

### **Snowflake Practice Questions - SnowPro Core Certified Set 3**

**Which statement is most accurate regarding Snowflake's transaction support?**

Options are :

- *It does not require updates*
- *Snowflake's transaction support is ACID-compliant (Correct)*
- *Snowflake transaction support allows only single statement*
- *Snowflake transaction support requires special user configuration*

**Answer :**

\*Snowflake's transaction support is ACID-compliant (Correct)

**How much does the Snowpipe charges for queued files?**

Options are :

- *0.06 credits per 1000 files queued (Correct)*
- *0.06 credits per 100 files queued*
- *0.6 credit per 1000 files queued*
- *0.06 credits per 10000 files queued`*

**Answer :**

\**0.06 credits per 1000 files queued (Correct)*

**Which statement accurately describes the Snowflake Cloud Services layer?**

Options are :

- *A collection of tightly coupled database management features*
- *Managed and scaled by the user*
- *A collection of independent, scalable, and stateless services providing crucial data management capabilities (Correct)*
- *A set of Cloud providers' general services made available to users via the internet*

**Answer :**

\**A collection of independent, scalable, and stateless services providing crucial data management capabilities (Correct)*

### Snowflake SnowPro Core Certification Exam Questions Set 1

**External tables enable querying existing data stored in external cloud storage for analysis without first loading it into Snowflake.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**When data is shared between Snowflake accounts, a database is created on the consumer side for sharing purposes. The type of this database is ?**

Options are :

- *Read & Update Only*
- *Read & Insert Only*
- *Read Only (Correct)*
- *Read & Delete Only*

**Answer :**

\*Read Only (Correct)

**Resource monitor notifications can be received by account administrators through which the following ways**

Options are :

- *Push Notification*
- *Web Interface (Correct)*
- *Email (Correct)*
- *SMS*

**Answer :**

\*Web Interface (Correct)

\*Email (Correct)

#### **SnowFlake SnowPro Advanced Architect Practice Exam Set 4**

**When running a SELECT COUNT(\*) on a table, which of the following statements is true?**

Options are :

- *No Warehouse will be needed because Count statistics are stored in the Metadata Cache (Correct)*
- *Only an XS Warehouse will be needed because all Roles share Count statistics*
- *Only an XS Warehouse will be needed because Count statistics are stored in the Results Cache.*
- *No Warehouse will be needed because all Users share Count statistics*

**Answer :**

\*No Warehouse will be needed because Count statistics are stored in the Metadata Cache (Correct)

**How does Snowflake's support for High Availability work?**

Options are :

- *High Availability capability is limited to the storage layer*
- *It does not support fault tolerance*
- *Snowflake's support for High Availability provides built-in, resilient capabilities in all three architectural layers (Correct)*
- *The feature requires users to manage provisioning & replication of the Cloud provider's virtual instances*

**Answer :**

\*Snowflake's support for High Availability provides built-in, resilient capabilities in all three architectural layers (Correct)

**What is the file size limit that can be loaded into a table using the Snowflake Web Interface**

Options are :

- *50 MB (Correct)*
- *10 MB*
- *20 MB*
- *100 MB*

**Answer :**

\*50 MB (Correct)

**Snowflake Practice Questions - SnowPro Core Certified Set 4****Select the statements that are true for a shared database**

Options are :

- *Consumer can clone the objects present in shared database*
- *Time Travel can be done by consumer on shared database*
- *Consumer can edit the comments of the shared database*

- *Users in a consumer account can view/query data, but cannot insert or update data, or create any objects in the database. (Correct)*
- *Shared databases are read-only (Correct)*

**Answer :**

\*Users in a consumer account can view/query data, but cannot insert or update data, or create any objects in the database. (Correct)  
\*Shared databases are read-only (Correct)

**In snowflake, while processing a query, the number of rows in the tables being queried has more impact than the overall size of the tables**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**Which access control privileges must be granted on specified resource monitors for roles other than account administrators to view and modify resource monitors as needed?**

Options are :

- *UPDATE*
- *CREATE*
- *MODIFY (Correct)*
- *MONITOR (Correct)*

**Answer :**

\*MODIFY (Correct)  
\*MONITOR (Correct)

### **Snowflake Certification Preparation Set 14**

**True or False: Periodic rekeying requires Enterprise Edition and higher**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

We have few clone tables in a database. Now for the database, replication is setup. In the secondary (target) database, the clone object will not be physically copied, and hence no storage cost will be incurred for them. True or False?

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

Load Metadata for each table expires in how many days?

Options are :

- 64 (Correct)
- 90
- 365
- 14

**Answer :**

\*64 (Correct)

## Snowflake SnowPro Advanced Architect Practice Exam Set 2

You can view the amount of storage used by failsafe under the Account->Billing & Usage section

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

Which Snowflake object records DML changes made to tables, including inserts, updates, deletes as well as metadata about each change, so that actions can be taken using the changed data

Options are :

- *Kafka*
- *Snowpipe*
- *Task*
- *Stream (Correct)*

**Answer :**

\**Stream (Correct)*

**True or false:** **Snowflake only works with cloud-based tools.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

### [Snowflake Certification Preparation Set 23](#)

**Databases based on most inbound shares can be dropped**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which type of Data Integration tools leverage Snowflake's scalable compute for data transformation? Select the best answer.**

Options are :

- *ETL*
- *ELT (Correct)*
- *Streaming*

- *Database Replication*

**Answer :**

\*ELT (Correct)

**What are some general indicators that a clustering key is required on a table?**

Options are :

- *The table is being used in many queries*
- *The query performance using that table has slowed down over a period of time (Correct)*
- *The table has many columns*
- *The size of the table is in multi terabytes (Correct)*

**Answer :**

\*The query performance using that table has slowed down over a period of time (Correct)

\*The size of the table is in multi terabytes (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 17

**Snowflake uses a staged release process for new releases in which different account types are upgraded gradually.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Snowflake database is just based on the massively parallel shared nothing architecture used by data bases like Netezza, Teradata, Greenplum.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**With Time Travel, a Snowflake administrator can, create clones of entire tables, schemas, and databases at or before specific points in the past**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 33

**Privileges granted on database objects are not replicated to a secondary database**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**What happens when a table is shared by a snowflake account to another snowflake account? Select all that applies**

Options are :

- *Target snowflake account is charged for data storage*
- *Data is copied to the target snowflake account*
- *Sharing of data is managed via Snowflake Metadata Services Layer (Correct)*
- *No actual data is copied or transferred between the accounts (Correct)*

**Answer :**

\**Sharing of data is managed via Snowflake Metadata Services Layer (Correct)*

\**No actual data is copied or transferred between the accounts (Correct)*

**Account parameters are also replicated, when a database replication is setup**

Options are :

- *FALSE (Correct)*

- *TRUE*

**Answer :**

\**FALSE* (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 1

**Restoring tables and schemas is only supported in the current schema or current database, even if a fully-qualified object name is specified.**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\**TRUE* (Correct)

**On the WebUI which button on the top bar should you select to get to tabs that provide billing and credit information?**

Options are :

- *Administrator*
- *Billing Information*
- *Credit Charge*
- *Account (Correct)*

**Answer :**

\**Account* (Correct)

**Source database have a materialized view, which gets replicated to the target database as part of database replication activity. The automatic maintenance & refresh of the materialized view in the target database will be enabled by default. True or False?**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE* (Correct)

## Snowflake Practice Questions - SnowPro Core Certified Set 5

**Which of the following is a common case for Cloning in Snowflake? (Select all that apply)**

Options are :

- *Data Encryption Protection*
- *Data Life Cycle Management (Correct)*
- *Point in Time" Snapshot (Correct)*
- *Agile Release and Development (Correct)*

**Answer :**

- \**Data Life Cycle Management (Correct)*
- \**Point in Time" Snapshot (Correct)*
- \**Agile Release and Development (Correct)*

**When a table which has time traveled enabled is dropped, and if a table with same name is created after that, what happens?**

Options are :

- *The table with same name cannot be created till the time time travel retention period of dropped object is over*
- *It creates a new version of the table (Correct)*
- *The table with same name cannot be created till the time time travel retention period and fail safe period of dropped object is over*
- *It restores the previous dropped version of the table*

**Answer :**

- \**It creates a new version of the table (Correct)*

**All security related information is stored in which layer of snowflake architecture**

Options are :

- *Storage*

- *Cloud Service (Correct)*
- *None*
- *Compute*

**Answer :**

\*Cloud Service (Correct)

### **Snowflake Snowpro Certification Practice Test Set 5**

**If file format options are specified in multiple locations, the load operation applies the options in the following order of precedence:**

Options are :

- *COPY INTO TABLE statement, Stage Definition, Table Definition (Correct)*
- *Stage Definition,COPY INTO TABLE statement, Table Definition*
- *Table Definition,COPY INTO TABLE statement, Stage Definition*
- *Table Definition, Stage Definition,COPY INTO TABLE statement*

**Answer :**

\*COPY INTO TABLE statement, Stage Definition, Table Definition (Correct)

**Which one of the following stage can not be cloned?**

Options are :

- *Table Stage*
- *External Named Stage*
- *Internal Named Stage (Correct)*

**Answer :**

\*Internal Named Stage (Correct)

**Which statements are true about storage relationships?**

Options are :

- *Snowflake Tables are stored within Schemas. (Correct)*
- *Snowflake Warehouses are stored within Data Marts*

- *Snowflake Schemas are stored within Databases (Correct)*
- *Snowflake Databases are stored within Warehouses*
- *Snowflake Schemas are stored within Warehouses*
- *Snowflake Warehouses are stored within Databases*

**Answer :**

\*Snowflake Tables are stored within Schemas. (Correct)  
\*Snowflake Schemas are stored within Databases (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 14](#)

**What is the maximum data retention period that an enterprise account can have?**

Options are :

- *90 Days (Correct)*
- *120 Days*
- *1 Day*
- *30 Days*

**Answer :**

\*90 Days (Correct)

**File format options set in multiple locations are not cumulative. Any options set in one place override all options (whether the same or different options) set lower in the order of precedence.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Replication operations use Snowflake-provided compute resources to copy data between accounts across regions.**

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

### Snowflake SnowPro Core Certification Practice Tests Set 17

**Snowpipe generally loads older files first, there is no guarantee that files are loaded in the same order they are staged.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**What happens If the TIMESTAMP, OFFSET, or STATEMENT specified in the AT | BEFORE clause of a SELECT statement falls outside the data retention period for the table?**

Options are :

- *The query returns no data*
- *The query fails and returns an error (Correct)*
- *The query returns the oldest timestamp data that is available*
- *The query fetches data from fail safe*

**Answer :**

\**The query fails and returns an error (Correct)*

**Which command can be executed to get the version history of table "Test1" in schema "mytestdb.myschema" within the time travel period in snowflake?**

Options are :

- *show tables version like 'Test1' in mytestdb.myschema;*
- *show version of table like 'Test1' in mytestdb.myschema;*
- *show tables history like 'Test1' in mytestdb.myschema; (Correct)*
- *show history of table like 'Test1' in mytestdb.myschema;*

**Answer :**

\*show tables history like 'Test1' in mytestdb.myschema; (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 18

**To clone a table your current role must have which privilege on the source table?**

Options are :

- *INSERT*
- *SELECT (Correct)*
- *UPDATE*
- *DELETE*

**Answer :**

\**SELECT* (Correct)

**Resource monitor notifications can be received only by account administrators**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE* (Correct)

**Failover/Failback requires minimum which snowflake edition**

Options are :

- *Virtual Private Snowflake*
- *Business Critical (Correct)*
- *Standard*
- *Enterprise*

**Answer :**

\**Business Critical* (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 9

**Which of the following will not contribute to your account level storage usage?**

Options are :

- *Databases (and tables in it)*
- *Snowflake Internal Stage*
- *UDF (Correct)*
- *Views (Normal Views & Not Materialized Views) (Correct)*
- *Snowflake External Stages (Correct)*

**Answer :**

- \*UDF (Correct)
- \*Views (Normal Views & Not Materialized Views) (Correct)
- \*Snowflake External Stages (Correct)

**Which of the following terms or phrases can also be used to describe Snowflake?**

Options are :

- *Hybrid Columnar (Correct)*
- *Native SQL (Correct)*
- *Hadoop-Compliant*
- *Multi-cluster (Correct)*
- *Built from the ground up for the cloud (Correct)*
- *Oracle derived*

**Answer :**

- \*Hybrid Columnar (Correct)
- \*Native SQL (Correct)
- \*Multi-cluster (Correct)
- \*Built from the ground up for the cloud (Correct)

**The table stage is designed to store files that are staged and managed by one or more users but only loaded into a single table.**

Options are :

- *FALSE*

- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

#### SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

**Which type of objects will cause database replication to fail?**

Options are :

- A object in the primary database references a dropped object (Correct)
- Non-materialized views that reference any object in another database (e.g. table columns, other views, UDFs, or stages)
- A foreign key in the primary database references a primary key in another database, or vice-versa (Correct)
- Materialized views that reference any object in another database (Correct)

**Answer :**

\*A object in the primary database references a dropped object (Correct)

\*A foreign key in the primary database references a primary key in another database, or vice-versa (Correct)

\*Materialized views that reference any object in another database (Correct)

**Databases created from shares cannot be replicated.**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Historical usage data (i.e. data in views like LOGIN\_HISTORY, COPY\_HISTORY, QUERY\_HISTORY) for activity in the primary database is also replicated to secondary databases.**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 2

**Which of the following best describes Snowflake's processing engine?**

Options are :

- *Native SQL (Correct)*
- *Leverages Apache Spark*
- *A derivative of Presto*
- *Based on Map Reduce*

**Answer :**

\*Native SQL (Correct)

**Which property of resource monitor specifies whether the resource monitor is used to monitor the credit usage for your entire Account (i.e. all warehouses in the account), or a specific set of individual warehouses?**

Options are :

- *SINGLE\_MULTIPLE*
- *None*
- *MONITOR LEVEL (Correct)*
- *CREDIT QUOTA*

**Answer :**

\*MONITOR LEVEL (Correct)

**Snowflake maintains detailed metadata for each table into which data is loaded, including:**

Options are :

- *Name of each file from which data was loaded (Correct)*
- *Column Names*
- *Number of Columns*
- *Etag for the file (Correct)*

- *File Size (Correct)*
- *Number of rows parsed in the file (Correct)*

**Answer :**

- \*Name of each file from which data was loaded (Correct)
- \*Etag for the file (Correct)
- \*File Size (Correct)
- \*Number of rows parsed in the file (Correct)

## Snowflake SnowPro Core Certification Practice Test Set 2

**Materialized views can be created on top of external tables in Snowflake**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \**TRUE (Correct)*

**True or False: As a consumer you can create only one database per share?**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \**TRUE (Correct)*

**Maximum historical data that can be maintained of Transient & Temporary table is?**

Options are :

- *1 Day (Correct)*
- *90 Days*
- *0 Day*
- *7 Days*

**Answer :**

\*1 Day (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 26

**Users created in a reader account can not do which of the following?**

Options are :

- *SELECT*
- *UPDATE (Correct)*
- *DELETE (Correct)*
- *INSERT (Correct)*

**Answer :**

\**UPDATE (Correct)*

\**DELETE (Correct)*

\**INSERT (Correct)*

**In the Snowflake Web UI which of the following is true?**

Options are :

- *For all open worksheets same role, warehouse and database apply*
- *Each open worksheet is a different session (Correct)*
- *Each open worksheet can have different role, warehouse and database selected (Correct)*
- *For all open worksheet the session is same*

**Answer :**

\**Each open worksheet is a different session (Correct)*

\**Each open worksheet can have different role, warehouse and database selected (Correct)*

**Which of the following object types are stored within schemas?**

Options are :

- *Stages (Correct)*
- *Roles*

- *User Defined Functions (Correct)*
- *File Formats (Correct)*
- *Stored Procedures (Correct)*
- *Sequences (Correct)*

**Answer :**

- \*Stages (Correct)
- \*User Defined Functions (Correct)
- \*File Formats (Correct)
- \*Stored Procedures (Correct)
- \*Sequences (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 18](#)

**Which of the following cannot be executed on SHARED data by a consumer?**

Options are :

- *Use a REPLACE function in a SELECT statement*
- *Create a table in the INBOUND SHARE's database (Correct)*
- *Use a GROUP BY clause in a SELECT statement*
- *JOIN the SHARED data to your own data*
- *Add the data to existing tables in INBOUND SHARE (Correct)*
- *Run an UPDATE statement on the data (Correct)*

**Answer :**

- \*Create a table in the INBOUND SHARE's database (Correct)
- \*Add the data to existing tables in INBOUND SHARE (Correct)
- \*Run an UPDATE statement on the data (Correct)

**What common tasks for traditional on-premise database and IT staff, are not required with Snowflake?**

Options are :

- *Maintaining database objects*
- *Maintaining the physical security of a server room (key cards, door locks, etc) (Correct)*
- *Maintaining statistics (Correct)*

- *Maintaining metadata (Correct)*

**Answer :**

- \*Maintaining the physical security of a server room (key cards, door locks, etc) (Correct)
- \*Maintaining statistics (Correct)
- \*Maintaining metadata (Correct)

**Which of the following terms are associated with the Compute/Warehouse Layer?**

Options are :

- *Query Processing (Correct)*
- *Query Design*
- *Query Planning*
- *Query Compilation*
- *Query Optimization*

**Answer :**

- \*Query Processing (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 31**

**Which installment option versions of Snowflake are available?**

Options are :

- *Microsoft Cloud Native Accounts*
- *Hybrid On-Premise + Cloud Installation*
- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)*
- *On-Premise Custom Installation*
- *Enterprise In-House VPC Installation*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)*

**Answer :**

- \*Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)
- \*Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)

**Which cache stores object definitions and statistics?**

Options are :

- *Results Cache*
- *Metadata Cache (Correct)*
- *Warehouse Cache*

**Answer :**

\*Metadata Cache (Correct)

**select \* from my\_table at(offset => -60\*5) . This statement will return which period of historical data for the table?**

Options are :

- *5 hours ago*
- *10 minutes ago*
- *60 minutes ago*
- *5 minutes ago (Correct)*

**Answer :**

\*5 minutes ago (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 34

**Three of these terms refer to the same layer. Select the correct 3.**

Options are :

- *Virtual Warehouse (Correct)*
- *Compute (Correct)*
- *Catalog*
- *Query Processing (Correct)*
- *Metadata*

**Answer :**

\*Virtual Warehouse (Correct)

\*Compute (Correct)

\*Query Processing (Correct)

**Table named mytable has a stage referenced as:**

Options are :

- @~mytable
- %~mytable
- @mytable
- @%mytable" (Correct)

**Answer :**

\*@%mytable" (Correct)

**Data Sharing (without replication) is only supported between provider and consumer account in same region**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 5

**What is the name of the warehouse in which the credit utilized by Snowpipe is captured?**

Options are :

- SNOWPIPE\_CREDITS
- Any existing warehouse can be configured to log the SNOWPIPE charges
- SNOWPIPE (Correct)
- Snowpipe Credits are not captured and send as separate montly bill to customer

**Answer :**

\*SNOWPIPE (Correct)

**Which of the following tables cannot be cloned as permanent table?**

Options are :

- Permanent Tables
- Temporary Tables (Correct)
- Transient Tables (Correct)

**Answer :**

- \*Temporary Tables (Correct)
- \*Transient Tables (Correct)

A pipe is a named, first-class Snowflake object that contains an INSERT statement used by Snowpipe.

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

- \*FALSE (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 10

Which type of view is most like a table?

Options are :

- External
- Secure
- Standard
- Materialized (Correct)

**Answer :**

- \*Materialized (Correct)

Which stage type do not support setting file format option, instead, must specify file format and copy options as part of the COPY INTO <table> command.

Options are :

- User Stage (Correct)
- Named External Stage

- *Named Internal Stage*
- *Table Stage (Correct)*

**Answer :**

\*User Stage (Correct)  
\*Table Stage (Correct)

**An inbound share with no database can be queried from the worksheet**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### Snowflake Certification Preparation Set 21

**Copy options set in multiple locations are cumulative. Individual options set in one place override the same option set lower in the order of precedence.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**The "load metadata" (which maintains which files have already been loaded) for a table expires after how many days?**

Options are :

- 7
- 14
- 64 (*Correct*)
- 365

**Answer :**

\*64 (Correct)

**When a table which has time traveled enabled is dropped, and if a new table with same name is created after that, then the dropped table cannot be restored back ever?**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

### Snowflake Practice Questions - SnowPro Core Certified Set 3

**Default Size of the warehouse created using Snowflake Web UI is**

Options are :

- *MEDIUM*
- *X-LARGE (Correct)*
- *SMALL*
- *X-SMALL*

**Answer :**

\**X-LARGE (Correct)*

**Tri-Secret Secure option is available in which Snowflake Edition**

Options are :

- *Enterprise & higher*
- *All editions*
- *Business Critical & higher (Correct)*

**Answer :**

\**Business Critical & higher (Correct)*

**Databases based on most inbound shares can be renamed**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 40

**True or False: A Provider can share data from multiple databases.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**When database replication is setup, then querying tables and views in a secondary database using Time Travel will always produce exact same results as that of when executing the same query in the primary database.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**When configuring a Warehouse using a Snowflake edition that has Elastic Data Warehousing enabled, what facets or components will you need to configure that are not needed in accounts where Elastic Data Warehousing is not enabled?**

Options are :

- *Auto-Resume*
- *Minimum and Maximum Clusters (Correct)*
- *Minimum and Maximum Servers*

- *Auto-Suspend*
- *Scaling Policy (Correct)*

**Answer :**

\*Minimum and Maximum Clusters (Correct)  
\*Scaling Policy (Correct)

**Snowflake Snowpro Certification Practice Test Set 4**

**Which type of view has an extra layer of protection to hide the SQL code from unauthorized viewing?**

Options are :

- *Permanent*
- *Materialized*
- *Standard*
- *Secure (Correct)*

**Answer :**

\*Secure (Correct)

**After how many days account and table master keys are automatically rotated by Snowflake?**

Options are :

- *30 (Correct)*
- *60*
- *45*
- *90*

**Answer :**

\*30 (Correct)

**Which types of database objects are currently not supported while doing database replication?**

Options are :

- *Streams (Correct)*
- *Sequences*

- *Temporary Tables (Correct)*

- *Tasks (Correct)*

- *Stages (Correct)*

- *Permanent Tables*

**Answer :**

\*Streams (Correct)

\*Temporary Tables (Correct)

\*Tasks (Correct)

\*Stages (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 3

**Resource monitors support the below actions**

Options are :

- *Notify & Suspend (Correct)*

- *Notify & Abort*

- *Notify & raise snowflake support ticket*

- *Notify & Suspend Immediately (Correct)*

- *Notify (Correct)*

- *Notify & Call accountadmins*

**Answer :**

\*Notify & Suspend (Correct)

\*Notify & Suspend Immediately (Correct)

\*Notify (Correct)

**When a table which has time traveled enabled is dropped, and if a new table with same name is created after that, then the undrop table command is executed, what will happen?**

Options are :

- *The newly created table will be overwritten with the restored version of the table*

- *The table will be restored with a different name*

- *The undrop command fails (Correct)*

**Answer :**

\*The undrop command fails (Correct)

**In case of database replication, the charges for compute and data transfer is charged to which account?**

Options are :

- *There is no charges for database replication*
- *Source Account (Account with Primary Database)*
- *Target Account (Account with Secondary database that is getting refreshed) (Correct)*

**Answer :**

\*Target Account (Account with Secondary database that is getting refreshed) (Correct)

**Snowflake SnowPro Core Certification Exam Questions Set 11**

**What attributes make Snowflake a true SaaS solution?**

Options are :

- *No maintenance upgrades or patches to install. (Correct)*
- *Transparent releases don't require user intervention. (Correct)*
- *No hardware to purchase or configure. (Correct)*
- *No data storage costs.*
- *No creation of user accounts or roles is required.*
- *No query processing costs.*

**Answer :**

\*No maintenance upgrades or patches to install. (Correct)

\*Transparent releases don't require user intervention. (Correct)

\*No hardware to purchase or configure. (Correct)

**Snowflake supports loading data from AWS, Azure, GCP regardless of the cloud platform on which the snowflake account is hosted**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which types of database objects are currently supported while doing database replication?**

Options are :

- *Tasks*
- *File Formats (Correct)*
- *Stored Procedures (Correct)*
- *Materialized Views (Correct)*
- *Stages*
- *Streams*

**Answer :**

\*File Formats (Correct)

\*Stored Procedures (Correct)

\*Materialized Views (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 11

**The Snowflake Architecture consists of:**

Options are :

- *None of the above*
- *A monolithic, integrated stack*
- *Tightly coupled storage and compute layers*
- *Three layers: Cloud Services, Virtual Warehouse, Hybrid Columnar Storage (Correct)*

**Answer :**

\*Three layers: Cloud Services, Virtual Warehouse, Hybrid Columnar Storage (Correct)

**Database replication is supported across regions and across cloud platforms.**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which of the following terms best describes Snowflake's database architecture?**

Options are :

- *Multi-cluster shared data (Correct)*
- *Shared disk*
- *Columnar shared nothing*
- *Cloud-native shared memory*

**Answer :**

\*Multi-cluster shared data (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 40

**Resource monitors provide control over virtual warehouse credit usage; however, you cannot use them to control credit usage for the Snowflake-provided warehouses, including the SNOWPIPE warehouse**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Which of the following is not a layer in Snowflake's Architecture?**

Options are :

- *Cloud Services*
- *Storage*
- *Data Catalog (Correct)*
- *Virtual Warehouses*

**Answer :**

\*Data Catalog (Correct)

**True or false: Snowflake offers tools to extract data from source systems like SAP, Salesforce?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### **Tackling Snowflake Certification - Practice Questions Set 2**

**How many number of consumer accounts can be added to a single share?**

Options are :

- *100*
- *Unlimited (Correct)*
- *10*
- *50*

**Answer :**

\*Unlimited (Correct)

**Multiple streams can be created on a single table**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Not all Snowflake Editions have Elastic Data Warehousing. Check all Snowflake editions that have Elastic Data Warehousing enabled**

Options are :

- *Enterprise Edition (Correct)*
- *Virtual Private Snowflake (Correct)*
- *Business Critical Edition (formerly Enterprise for Sensitive Data) (Correct)*
- *Standard Edition*
- *Premium Edition*

**Answer :**

- \*Enterprise Edition (Correct)
- \*Virtual Private Snowflake (Correct)
- \*Business Critical Edition (formerly Enterprise for Sensitive Data) (Correct)

### Snowflake Certification Preparation Set 11

**What makes a Partner Connect Partner different from other partners?**

Options are :

- *Includes a streamlined Partner Trial Account Signup (Correct)*
- *Can be connected from within the WebUI (Correct)*
- *Requires Enterprise Edition*
- *Can be connected to Snowflake using a streamlined wizard (Correct)*
- *Includes automated role, user and staging database set up (Correct)*

**Answer :**

- \*Includes a streamlined Partner Trial Account Signup (Correct)
- \*Can be connected from within the WebUI (Correct)
- \*Can be connected to Snowflake using a streamlined wizard (Correct)
- \*Includes automated role, user and staging database set up (Correct)

**Which of the following term is associated with the Compute/Warehouse Layer?**

Options are :

- *Query Optimization*
- *Query Compilation*
- *Query Processing (Correct)*
- *Query Design*

- *Query Planning*

**Answer :**

\*Query Processing (Correct)

**Which of the following conditions are required for sharing data in Snowflake?**

Options are :

- *Consumer accounts must be in same Snowflake region as the provider account (Correct)*
- *Data providers with ACCOUNTADMIN role can set up shares (Correct)*
- *Each data share must contain a single database (Correct)*
- *Any object in a share can be from different databases*
- *Secure views are not required when query performance is priority*

**Answer :**

\*Consumer accounts must be in same Snowflake region as the provider account (Correct)

\*Data providers with ACCOUNTADMIN role can set up shares (Correct)

\*Each data share must contain a single database (Correct)

## Tackling Snowflake Certification - Practice Questions Set 2

**Which two table types can have a maximum retention period of 1 day?**

Options are :

- *TEMPORARY (Correct)*
- *MATERIALIZED*
- *TRANSIENT (Correct)*
- *PERMANENT*

**Answer :**

\*TEMPORARY (Correct)

\*TRANSIENT (Correct)

**When a data is shared with an existing snowflake customer, the compute is charged to**

Options are :

- *Consumer (Correct)*
- *No Charges are applied*
- *Provider*

**Answer :**

\*Consumer (Correct)

**What features are provided by Snowflake to enable continuous data pipelines?**

Options are :

- *Continuous Data Loading*
- *All options are correct (Correct)*
- *Change Data Tracking*
- *Recurring Tasks*

**Answer :**

\*All options are correct (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 1

**To restore the table by undrop command what all conditions should be met?**

Options are :

- *The table can be restored in any schema irrespective from whichever schema the table was dropped.*
- *User executing undrop command should have ownership privilege on the table (Correct)*
- *The table can be restored in same schema only from which it has been dropped (Correct)*
- *User must of CREATE privileges on the schema in which table is getting restored (Correct)*

**Answer :**

\*User executing undrop command should have ownership privilege on the table (Correct)

\*The table can be restored in same schema only from which it has been dropped (Correct)

\*User must of CREATE privileges on the schema in which table is getting restored (Correct)

**External Tokenization is not supported on which Cloud Platforms (as of year 2020)**

Options are :

- AWS
- Above all cloud provide the required facility
- GCP (Correct)
- Azure (Correct)

**Answer :**

\*GCP (Correct)  
\*Azure (Correct)

**Data Consumers can be set up to consume data using which two types of consuming accounts?**

Options are :

- Query Accounts & Full Accounts
- Sharing Accounts & Query Accounts
- Full Accounts & Reader Accounts (Correct)
- Sharing Accounts & Providing Accounts

**Answer :**

\*Full Accounts & Reader Accounts (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 2

**True or false: A table in Snowflake can contain both a variant columns and a timestamp.**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Which of the following Snowflake Editions automatically store data in an encrypted state?**

Options are :

- *Enterprise (Correct)*
- *Standard (Correct)*
- *Business Critical (Correct)*

**Answer :**

- \*Enterprise (Correct)
- \*Standard (Correct)
- \*Business Critical (Correct)

**What is the maximum number of consumer accounts that can be added to a Share object?**

Options are :

- *100*
- *10*
- *One*
- *Unlimited (Correct)*

**Answer :**

- \*Unlimited (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 7

**After the retention period for an object has passed and the object has been purged, it is no longer displayed in the SHOW <object\_type> HISTORY output.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**What kind of offering is Snowflake?**

Options are :

- *Platform as a service (PaaS)*
- *Infrastructure as a service (IaaS)*

- *Software as a service (SaaS) (Correct)*
- *On Premise Application*

**Answer :**

\*Software as a service (SaaS) (Correct)

**Snowflake Instances in different regions require separate accounts**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 35**

**Which of the following conditions will result in the database replication operation to get failed?**

Options are :

- *A masking policy contained in the primary database is applied to columns in a table or view in another database, or vice-versa (Correct)*
- *There are User Defined Functions (UDF) in the primary (source) database*
- *There are external tables in the primary database (source database) (Correct)*
- *There are Materialized Views in the source database*
- *The primary database is in an Enterprise (or higher) account and contains a masking policy but one or more of the accounts approved for replication are on lower editions (Correct)*

**Answer :**

\*A masking policy contained in the primary database is applied to columns in a table or view in another database, or vice-versa (Correct)

\*There are external tables in the primary database (source database) (Correct)

\*The primary database is in an Enterprise (or higher) account and contains a masking policy but one or more of the accounts approved for replication are on lower editions (Correct)

**Each Snowflake account comes with two shared databases. One is a set of sample data and the other contains Account Usage information. Check all true statements about these shared databases.**

Options are :

- *SNOWFLAKE\_SAMPLE\_DATA* contains several schemas from TPC (tpc.org) (Correct)
- *ACCOUNT USAGE* is a schema filled with secure views (Correct)
- *ACCOUNT\_USAGE* is a schema filled with external tables
- *SNOWFLAKE* contains a schema called *ACCOUNT\_USAGE* (Correct)
- *SNOWFLAKE\_SAMPLE\_DATA* contains a schema called *ACCOUNT\_USAGE*
- *SNOWFLAKE* contains a table called *ACCOUNT\_USAGE*

**Answer :**

\**SNOWFLAKE\_SAMPLE\_DATA* contains several schemas from TPC (tpc.org) (Correct)

\**ACCOUNT USAGE* is a schema filled with secure views (Correct)

\**SNOWFLAKE* contains a schema called *ACCOUNT\_USAGE* (Correct)

**When configuring a Warehouse using a Snowflake edition that has Elastic Data Warehousing enabled, what facets or components will you need to configure that are not needed in accounts where Elastic Data Warehousing is not enabled**

Options are :

- *Scaling Policy* (Correct)
- *Auto-Suspend*
- *Auto-Resume*
- *Minimum and Maximum Clusters* (Correct)
- *Minimum and Maximum Servers*

**Answer :**

\**Scaling Policy* (Correct)

\**Minimum and Maximum Clusters* (Correct)

#### [Snowflake Snowpro Certification Practice Test Set 4](#)

**How often does Snowflake release new features?**

Options are :

- *Yearly*
- *Bi-Annually*
- *Monthly*
- *Weekly (Correct)*
- *Never*

**Answer :**

\*Weekly (Correct)

**True or False: You cannot share a share**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What is the default table type when a CREATE TABLE is executed?**

Options are :

- *MATERIALIZED*
- *TRANSIENT*
- *PERMANENT (Correct)*
- *TEMPORARY*

**Answer :**

\*PERMANENT (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 2**

**Snowflake cannot be run on private cloud infrastructures (on-premises or hosted)**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**The data retention period for a database, schema, or table cannot be changed after they are created.**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**Which of the following roles can use SQL to view the task history within a specified date range**

Options are :

- Any role that has the global MONITOR EXECUTION privilege (Correct)
- Task owner (i.e. role that has the OWNERSHIP privilege on a task) (Correct)
- Accountadmin (Correct)
- Public
- Any role can view

**Answer :**

\*Any role that has the global MONITOR EXECUTION privilege (Correct)

\*Task owner (i.e. role that has the OWNERSHIP privilege on a task) (Correct)

\*Accountadmin (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 12

**External Stages require a cloud storage provider**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**If a table is cloned with COPY GRANTS option specified, then the clone object inherits any explicit access privileges granted on the original table, but does not inherit any future grants defined for the table in the schema**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**How many number of shares can be created in a Snowflake Account?**

Options are :

- Unlimited (Correct)
- 10
- 100
- 50

**Answer :**

\*Unlimited (Correct)

#### Snowflake Certification Preparation Set 4

**What data structure types can be ingested into a VARIANT column in a Snowflake table? (Check all that apply)**

Options are :

- PARQUET (Correct)
- HADOOP
- XML (Correct)
- JSON (Correct)
- ORC (Correct)
- AVRO (Correct)

**Answer :**

- \*PARQUET (Correct)
- \*XML (Correct)
- \*JSON (Correct)
- \*ORC (Correct)
- \*AVRO (Correct)

**Providers can list data on the Marketplace without contacting Snowflake**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

- \**FALSE (Correct)*

**After dropping an object, creating an object with the same name does not restore the object. Instead, it creates a new version of the object. The original, dropped version is still available and can be restored.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \**TRUE (Correct)*

### SnowFlake SnowPro Advanced Architect Practice Exam Set 4

**An append-only table stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \**TRUE (Correct)*

**When a data is shared with an non snowflake customer, the compute is charged to**

Options are :

- *Provider (Correct)*
- *No Charges applied*
- *Consumer*

**Answer :**

\*Provider (Correct)

**Different editions of snowflake Instances require separate accounts**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 3**

**When adding Consumers to Outbound Shares, what types of Consumer Accounts can be selected?**

Options are :

- *GCP*
- *Full (Correct)*
- *AWS*
- *Reader (Correct)*
- *Azure*

**Answer :**

\*Full (Correct)

\*Reader (Correct)

**A table named "MY\_TABLE" has a retention period of 7 days. The table is dropped, and a new table with the same name "MY\_TABLE" has been created. Now if the UNDROP MY\_TABLE command is executed what will happen?**

Options are :

- The UNDROP command will restore the dropped version with a new name
- The UNDROP command will fail (Correct)
- The UNDROP command will overwrite the already existing table with the previously dropped table

**Answer :**

\*The UNDROP command will fail (Correct)

**While cloning the entire database, the Internal (Snowflake) stages inside that database are also cloned**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 11

**Snowflake is supported on which cloud providers**

Options are :

- Alibaba Cloud
- GCP (Correct)
- IBM Cloud
- Azure (Correct)
- AWS (Correct)

**Answer :**

\*GCP (Correct)

\*Azure (Correct)

\*AWS (Correct)

**How much Snowpipe charge for the files in Queue**

Options are :

- 0.06 credits per 100 files queued

- 0.06 credits per 10000 files queued
- 0.06 credits per 1000 files queued (Correct)
- No charges for files in Queue

**Answer :**

\*0.06 credits per 1000 files queued (Correct)

**A stream allows querying and consuming a set of changes to a table, at the row level, between two transactional points of time**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### [SnowFlake SnowPro Advanced Architect Practice Exam Set 4](#)

**Snowflake has which type of architecture**

Options are :

- None of the above
- Shared Nothing
- Shared Disk
- Shared Data & Multi Cluster (Correct)

**Answer :**

\*Shared Data & Multi Cluster (Correct)

**Which of the following terms are associated with the Cloud Services Layer?**

Options are :

- Query Design
- Query Planning (Correct)
- Query Optimization (Correct)

- *Query Processing*
- *Query Compilation (Correct)*

**Answer :**

- \*Query Planning (Correct)
- \*Query Optimization (Correct)
- \*Query Compilation (Correct)

**If you find a data-related tool that is not listed as part of the Snowflake ecosystem, what industry standard options could you check for as a way to easily connect to Snowflake?**

Options are :

- *Check to see if there is a petition in the community to create a driver*
- *Check to see if you can develop a driver and put it on Github*
- *Check to see if the tool can connect to other solutions via JDBC (Correct)*
- *Check to see if the tool can connect to other solutions via ODBC (Correct)*

**Answer :**

- \*Check to see if the tool can connect to other solutions via JDBC (Correct)
- \*Check to see if the tool can connect to other solutions via ODBC (Correct)

### Snowflake Certification Preparation Set 17

**A table named "MY\_TABLE" has a retention period of 7 days. The table is dropped, and a new table with the same name "MY\_TABLE" has been created. The new table that is created is also dropped. Now retrieving the original version (i.e. the 1st version) of the table is not possible. This is True or False?**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

- \*FALSE (Correct)

**Which of the following statements are true about Snowflake Data sharing? Select all that apply.**

Options are :

- *Data Sharing is integrated with role-based access controls (Correct)*

- You don't need a Snowflake Virtual Warehouse to query shared data
- Consumers can query shared tables in the same query as their own tables (Correct)
- Producers can share data to other databases such as Postgres and MySQL

**Answer :**

- \*Data Sharing is integrated with role-based access controls (Correct)
- \*Consumers can query shared tables in the same query as their own tables (Correct)

**Object parameters are not replicated, except the parameter "DATA\_RETENTION\_TIME\_IN\_DAYS" for schemas & tables**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

- \*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 35

**In Snowflake Enterprise edition, the maximum retention period for a transient or temporary table is?**

Options are :

- 7 days
- 90 days
- Transient & Temporary table cannot be configured for time travel
- 1 day (Correct)

**Answer :**

- \*1 day (Correct)

**With Snowflake, the customer needs to perform hardware maintenance, upgrades or regular patches**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Which type of table stream is supported only on external tables?**

Options are :

- *Update-Only*
- *Native*
- *Standard*
- *Append-Only*
- *Insert-Only (Correct)*

**Answer :**

\*Insert-Only (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 1

**A stream object records data manipulation language (DML) changes made to tables, including inserts and updates but cannot record deletes made to tables**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**If a table is cloned without COPY GRANTS option specified, then the clone object does not inherits any explicit access privileges granted on the original table, but does inherit any future grants defined for the table in the schema**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Which of the following are performed by the Cloud Services layer?**

Options are :

- *Metadata Storage (Correct)*
- *Data Security (Correct)*
- *Metadata Management (Correct)*
- *Availability Zone Management*
- *User Authentication (Correct)*

**Answer :**

\*Metadata Storage (Correct)  
\*Data Security (Correct)  
\*Metadata Management (Correct)  
\*User Authentication (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 11

**Which out of below do you not specify while setting up a snowflake account**

Options are :

- *Cloud Infrastructure Provider*
- *Snowflake Edition*
- *Geographic Deployment Region*
- *Availability Zone (Correct)*

**Answer :**

\*Availability Zone (Correct)

**Which parameter is used to check if the stream contains change data for a table or not?**

Options are :

- *SYSTEM\$STREAM\_HAS\_MODIFIED\_DATA*
- *SYSTEM\$STREAM\_HAS\_DATA (Correct)*
- *SYSTEM\$STREAM\_HAS\_NEW\_DATA*
- *SYSTEM\$STREAM\_HAS\_CHANGED\_DATA*

**Answer :**

\*SYSTEM\$STREAM\_HAS\_DATA (Correct)

**Which of the following Snowflake Object is used for a schedule execution of SQL statement?**

Options are :

- *Snowpipe*
- *Tasks (Correct)*
- *Stored Procedure*
- *Streams*

**Answer :**

\*Tasks (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 39](#)

**When working with unfamiliar semi-structured data, you might not know the key names in an object. You can use the FLATTEN function with which argument to return the list of distinct key names in all nested elements in an object?**

Options are :

- *GROUPING*
- *DISTINCT*
- *NESTING*
- *RECURSIVE (Correct)*

**Answer :**

\*RECURSIVE (Correct)

**Which of the following Snowflake Editions encrypt all data transmitted over the network within a Virtual Private Cloud (VPC)?**

Options are :

- *Enterprise*
- *Business Critical (Correct)*
- *Premier*

- *Standard*

**Answer :**

\*Business Critical (Correct)

**Which of the following database objects can be cloned in Snowflake?**

Options are :

- *Databases*
- *All of the above (Correct)*
- *Schemas*
- *Tables*

**Answer :**

\*All of the above (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 17

**When a database or schema that contains a source table and stream is cloned, any unconsumed records in the stream (in the clone) are inaccessible**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which type of table stream tracks all DML changes to the source table including inserts, updates, deletes and also truncate?**

Options are :

- *Insert-Only*
- *Standard (Correct)*
- *Native*
- *Append-Only*
- *Multifunction*

**Answer :**

\*Standard (Correct)

**Snowflake has which architecture type**

Options are :

- *None of the above*
- *Shared Disk*
- *Hybrid of Shared disk & shared nothing (Correct)*
- *Shared Nothing*

**Answer :**

\*Hybrid of Shared disk & shared nothing (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 26**

**PUT Command can be executed from Snowflake Web UI**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Shares take place between which two groups?**

Options are :

- *Providers & Consumers (Correct)*
- *Sharers & Consumers*
- *Data Writers & Data Readers*
- *Full Accounts & Reading Accounts*

**Answer :**

\*Providers & Consumers (Correct)

**Which SQL statements can work on a cloned table?**

Options are :

- *All of the above (Correct)*
- *SHOW command*
- *DROP table command*
- *SELECT command*

**Answer :**

\*All of the above (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 2

**You are a developer with in one company running a Snowflake data warehouse. You required to load JSON based data into a table. What is the best way to do it in snowflake?**

Options are :

- *Ask the source team to send data into CSV instead of JSON*
- *Create a table with variant column and load the json data into that table (Correct)*
- *Write a program in Java or Python to convert the JSON into tabular format and then load in Snowflake table*

**Answer :**

\*Create a table with variant column and load the json data into that table (Correct)

**Data retention period of database is set to 2 days. Data retention period of schema inside that database is set to 4 days. Data retention of a table inside the schema is set to 7 days. Now if the database is dropped, the table that was present in that database can be restored back within how many days?**

Options are :

- *2 days (Correct)*
- *Since the Database itself is dropped, the underlying objects cannot be restored*
- *4 days*

- 7 days

**Answer :**

\*2 days (Correct)

**JSON data is a hierarchical collection of name/value pairs grouped into objects and arrays**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 5

**The Snowflake Data Marketplace has two types of listings. What are they?**

Options are :

- Standard and Business Critical
- Business Critical and Premier
- De-Personalized and Premier
- De-Personalized and Personalized
- Standard and Personalized (Correct)

**Answer :**

\*Standard and Personalized (Correct)

**Snowflake Administrators should utilize resource monitors to help control costs and avoid unexpected credit usage. Which of the following actions can Snowflake's resource monitor triggers initiate automatically? (Check all that apply)**

Options are :

- Impose limits on the number of credits that warehouses consume each month (Correct)
- Roll over query executions to under utilized warehouses
- Trigger warehouse suspension for high usage (Correct)

- Trigger alert notifications for high usage (Correct)

**Answer :**

\*Impose limits on the number of credits that warehouses consume each month (Correct)

\*Trigger warehouse suspension for high usage (Correct)

\*Trigger alert notifications for high usage (Correct)

**When choosing a geographic deployment region, what factors might be considered?**

Options are :

- Number of availability zones within a region (Correct)
- End-user perceptions of glamorous or trendy geographic locations
- Proximity to the point of service (Correct)
- Additional fees charged for regions with geo-political unrest

**Answer :**

\*Number of availability zones within a region (Correct)

\*Proximity to the point of service (Correct)

**Snowflake SnowPro Core Certification Practice Tests Set 15**

**What are the names of the three Snowflake Editions offered when signing up for a trial account?**

Options are :

- Ultra
- Free-Tier Basic
- Business Critical (Correct)
- Standard (Correct)
- Enterprise (Correct)

**Answer :**

\*Business Critical (Correct)

\*Standard (Correct)

\*Enterprise (Correct)

**An agreement must be signed by Data Providers before a listing will be made public on Data Marketplace**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Cloning the table does not include the load history of the source table, and hence the data files that were loaded in source table can be loaded again into its clone tables**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

#### Snowflake Snowpro Certification Practice Test Set 4

**Providers can prep data for the Marketplace but will need Snowflake's guidance and approval at some point**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which of the below is not a role available when Snowflake account is created initially**

Options are :

- *Sysadmin*
- *Useradmin*
- *Developer (Correct)*
- *Accountadmin*

**Answer :**

\*Developer (Correct)

**Which of the following is a separate database object?**

Options are :

- *Table Stage*
- *Named Stage (Correct)*
- *User Stage*

**Answer :**

\*Named Stage (Correct)

### Snowflake Certification Preparation Set 8

**DATA\_RETENTION\_TIME\_IN\_DAYS parameter is defined as 1 day at account level, and the same parameter is set as 7 days at Database level. Now a table is created inside the given database without specifying any explicit DATA\_RETENTION\_TIME\_IN\_DAYS parameter. What would be the data retention period for table in this case?**

Options are :

- *7 days (Correct)*
- *0 day*
- *1 day*

**Answer :**

\*7 days (Correct)

**What command (partial) is used to convert a SECURE VIEW to a "regular" VIEW?**

Options are :

- *UNSET SECURE (Correct)*
- *SET INSECURE*
- *SET SECURE=False*
- *SET View\_Type=Regular*

**Answer :**

\*UNSET SECURE (Correct)

**For databases, schemas, and tables, a clone does not contribute to the overall data storage until operations are performed on the clone that modify existing data or add new data**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### Snowflake SnowPro Core Certification Practice Test Set 2

**For which table type the retention period ends as soon as the the table is dropped or the session in which the table was created ends.**

Options are :

- *PERMANENT*
- *MATERIALIZED*
- *TEMPORARY (Correct)*
- *TRANSIENT*

**Answer :**

\**TEMPORARY (Correct)*

**Which of the following are Snowflake view types?**

Options are :

- *Secure (Correct)*
- *Transient*
- *Standard (Correct)*
- *Materialized (Correct)*
- *External*
- *Permanent*

**Answer :**

\**Secure (Correct)*

\*Standard (Correct)

\*Materialized (Correct)

#### **Table stream also contains actual table data**

Options are :

- *FALSE (Correct)*

- *TRUE*

**Answer :**

\**FALSE (Correct)*

#### **Snowflake SnowPro Core Certification Practice Test Set 4**

If the **TIMESTAMP**, **OFFSET**, or **STATEMENT** specified in the **AT | BEFORE** clause falls outside the data retention period for the table, the query fails and returns an error.

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

#### **The Cloning feature in Snowflake requires less storage because:**

Options are :

- *The cloned data is constantly updated to synchronize with original table data*
- *Only metadata is copied; no physical data is copied (Correct)*
- *Data is replicated to ensure integrity of data in the original table data*

**Answer :**

\**Only metadata is copied; no physical data is copied (Correct)*

#### **Which of these are Snowflake table types?**

Options are :

- *Permanent (Correct)*

- *Temporary (Correct)*
- *External (Correct)*
- *Secure*
- *Transient (Correct)*
- *Materialized*

**Answer :**

- \*Permanent (Correct)
- \*Temporary (Correct)
- \*External (Correct)
- \*Transient (Correct)

### Snowflake Certification Preparation Set 18

**Tasks cannot be chained together for successive execution to support more complex periodic processing**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

- \**FALSE (Correct)*

**Query execution is performed by which layers of Snowflake architecture**

Options are :

- *Cloud Services Layer*
- *Query Processing Layer (Correct)*
- *Storage Layer*

**Answer :**

- \**Query Processing Layer (Correct)*

**What snowflake objects can be shared?**

Options are :

- *External (Correct)*

- *Regular Views*
- *Secure Views (Correct)*
- *Secure Materialized Views (Correct)*
- *Tables (Correct)*
- *Secure UDFs (Correct)*

**Answer :**

- \*External (Correct)
- \*Secure Views (Correct)
- \*Secure Materialized Views (Correct)
- \*Tables (Correct)
- \*Secure UDFs (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 9

**Data can be shared READ-ONLY across Snowflake accounts and can also be cloned**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**Snowflake on GCP is not available in which region currently (as of Year 2021)**

Options are :

- *Europe*
- *Asia Pacific (Correct)*
- *North America*

**Answer :**

- \**Asia Pacific (Correct)*

**To be listed on the Marketplace, data must meet few criteria. Check all the options that are applicable -**

Options are :

- Compliant according to PCI, HIPPA and other standards (Correct)
- If sample data, over 2GBs
- Legally Distributable (Correct)
- Fresh & Non-Static (Correct)

**Answer :**

- \*Compliant according to PCI, HIPPA and other standards (Correct)
- \*Legally Distributable (Correct)
- \*Fresh & Non-Static (Correct)

### Snowflake Certification Preparation Set 18

**Cloud Services layer of Snowflake provide which all functionalities**

Options are :

- Query Processing
- Database Storage
- Metadata Management (Correct)
- Infra Management (Correct)
- Authentication (Correct)

**Answer :**

- \*Metadata Management (Correct)
- \*Infra Management (Correct)
- \*Authentication (Correct)

**In Snowflake Enterprise edition, the maximum retention period for a permanent table is?**

Options are :

- 365 days
- 90 days (Correct)
- 7 days
- 1 day

**Answer :**

- \*90 days (Correct)

**A medium size warehouse runs in Auto-scale mode for 3 hours with a resize from Medium (4 servers per cluster) to Large (8 servers per cluster). Warehouse is resized from Medium to Large at 1:30 hours, Cluster 1 runs continuously, Cluster 2 runs continuously for the 2nd and 3rd hours, Cluster 3 runs for 15 minutes in the 3rd hour. How many total credits will be consumed**

Options are :

- 34 (*Correct*)
- 30
- 40
- 32

**Answer :**

\*34 (Correct)

### Snowflake Certification Preparation Set 21

**Under what circumstances would Fail-Safe NOT be an effective method for data recovery?**

Options are :

- As a means for accessing historical data after the Time Travel retention period has ended (*Correct*)
- To recover data that may have been lost or damaged due to extreme operational failures
- To protect data in the event of a system failure or other catastrophic event e.g. a hardware failure or security breach

**Answer :**

\*As a means for accessing historical data after the Time Travel retention period has ended (*Correct*)

**Which are the 2 modes in which a multi cluster warehouse can run?**

Options are :

- Maximized (*Correct*)
- Balanced
- Auto-Scale (*Correct*)
- Dynamic

- *Minimized*

**Answer :**

- \*Maximized (Correct)
- \*Auto-Scale (Correct)

**Which statement is the most accurate regarding the Snowflake Metadata Service?**

Options are :

- *Snowflake Metadata services can be managed by the user*
- *Snowflake Metadata services capabilities do not impact data sharing, time travel or cloning capabilities*
- *Snowflake Metadata services is a separate product offering from Snowflake*
- *It provides crucial, built-in management services for all metadata from capabilities such as query optimization and security (Correct)*

**Answer :**

- \*It provides crucial, built-in management services for all metadata from capabilities such as query optimization and security (Correct)

### [Snowflake SnowPro Core Certification Practice Test Set 6](#)

**Which parameter is used to set the time travel retention period?**

Options are :

- *DATA\_RETENTION\_TIME\_IN\_HOURS*
- *DATA\_RETENTION\_TIME\_IN\_WEEKS*
- *DATA\_RETENTION\_TIME*
- *DATA\_RETENTION\_TIME\_IN\_DAYS (Correct)*

**Answer :**

- \**DATA\_RETENTION\_TIME\_IN\_DAYS (Correct)*

**Using Time Travel, you can perform which of the following actions within a defined period of time (i.e. within the time travel period of the object)**

Options are :

- *All options are correct (Correct)*

- Query data in the past that has since been updated or deleted
- Restore tables, schemas, and databases that have been dropped
- Create clones of entire tables, schemas, and databases at or before specific points in the past

**Answer :**

\*All options are correct (Correct)

**The time travel retention period can only be set at individual object level, and cannot be setup at account level**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 13](#)

**If retrieval of data from fail safe is required, you must**

Options are :

- Download it directly from the underlying cloud storage of Snowflake
- Contact Snowflake Support (Correct)
- Retrieval from Fail Safe is not possible
- Query the Fail Safe DB and retrieve the data by yourself

**Answer :**

\*Contact Snowflake Support (Correct)

**Which of the following Snowflake Editions allow only a maximum of 1 day of time travel?**

Options are :

- Business Critical
- Enterprise
- VPS

- Standard (Correct)

**Answer :**

\*Standard (Correct)

**What is the benefit of client-side encryption in snowflake?**

Options are :

- It provides a secure system for managing data in cloud storage (Correct)
- None
- Saves the cloud encryption cost
- Helps in decrypting data faster when needed

**Answer :**

\*It provides a secure system for managing data in cloud storage (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 5

If a transaction is running on a session and the session disconnects abruptly, the transaction will be in a detached state. It will not allow the transaction to commit or roll back. The locks on the resources will also not be released. If you do not do anything and let snowflake abort the transaction, how long will you need to wait?

Options are :

- 15 Minutes
- 30 Minutes
- 4 Hours (Correct)
- 2 Hours
- 1 Hour

**Answer :**

\*4 Hours (Correct)

**Which two table types don't have fail safe?**

Options are :

- TEMPORARY (Correct)

- *PERMANENT*
- *TRANSIENT (Correct)*

**Answer :**

\*TEMPORARY (Correct)  
\*TRANSIENT (Correct)

**There is no additional storage charges with rekeying of data in Fail-safe.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 6

**If you have a table with a 10-day retention period and you decrease the period to 1-day, then what happens to the data that is already in time travel for 4 days**

Options are :

- *The data would be removed and will not even be available in fail safe*
- *The data would continue to be in time travel, and new time travel retention of 1 day will be applied to only active data which is not already in time travel window*
- *The data would be moved to fail safe (Correct)*
- *The data would move to fail safe after 1 day*

**Answer :**

\*The data would be moved to fail safe (Correct)

**Micro-partitions are immutable which means once created they cannot be changed.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Cloud storage providers must be linked to internal stage objects**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which two modes are available for multi cluster warehouse**

Options are :

- *Dynamic*
- *Auto Scale (Correct)*
- *Maximized (Correct)*
- *Minimized*

**Answer :**

\**Auto Scale (Correct)*

\**Maximized (Correct)*

**Both time travel & failsafe require extra storage which has cost associated with it**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Which two SQL Extension clause are provided by Snowflake to support time travel?**

Options are :

- *DURING*
- *BEFORE (Correct)*
- *AT (Correct)*

- *PRIOR*

**Answer :**

- \*BEFORE (Correct)
- \*AT (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 10

**A stage links to a storage integration using a hidden ID rather than the name of the storage integration**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**Warehouses can be started and stopped at any time**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**To help manage the storage costs associated with Time Travel and Fail-safe, Snowflake provides which two table types**

Options are :

- *PERMANENT*
- *MATERIALIZED*
- *TRANSIENT (Correct)*
- *TEMPORARY (Correct)*

**Answer :**

- \*TRANSIENT (Correct)
- \*TEMPORARY (Correct)

## Snowflake Certification Preparation Set 20

**select \* from my\_table at(offset => -60\*5); This query will select historical data from "my\_table" as of-**

Options are :

- *5 seconds ago*
- *5 minutes ago (Correct)*
- *5 hours ago*
- *5 days ago*

**Answer :**

\*5 minutes ago (Correct)

**Snowflake stores metadata about all rows stored in a micro-partition, including**

Options are :

- *Additional properties used for both optimization and efficient query processing (Correct)*
- *The number of distinct values (Correct)*
- *Mean & Standard deviation of each column*
- *The range of values for each of the columns in the micro-partition (Correct)*

**Answer :**

\*Additional properties used for both optimization and efficient query processing (Correct)

\*The number of distinct values (Correct)

\*The range of values for each of the columns in the micro-partition (Correct)

**True or false: Snowflake caches are automatically invalidated if the underlying data changes.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 33

**With Time Travel, a Snowflake administrator can query data in the past that has since been updated or deleted**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Fail-Safe is a seven-day history of data and is automatically available on which table types?**

Options are :

- *Permanent (Correct)*
- *External*
- *Transient*
- *Temporary*

**Answer :**

\**Permanent (Correct)*

**Which of the following compression methods can snowflake automatically detect**

Options are :

- *Zstandard*
- *brotli*
- *gzip (Correct)*
- *bzip2 (Correct)*

**Answer :**

\**gzip (Correct)*

\**bzip2 (Correct)*

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 2

**Extended time travel of 90 days is supported starting which snowflake edition**

Options are :

- *Enterprise (Correct)*
- *Standard*
- *Premier*
- *Business Critical*
- *Virtual Private Snowflake*

**Answer :**

\*Enterprise (Correct)

**Multi-cluster warehouses is only available in Enterprise Edition and above**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Time Travel and Fail safe are one and the same thing?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Snowflake SnowPro Core Certification Practice Tests Set 11**

**In maximized mode, what happens if the number of clusters are decreased? (e.g. previously min & max was 4, but now min & max cluster size is changed to 2)**

Options are :

- Specified number of clusters shut down when they finish executing statements and the auto-suspend period elapses (Correct)
- The number of clusters in maximized mode cannot be changed
- Specified number of clusters shut down immediately, as soon as they finish executing statements

**Answer :**

\*Specified number of clusters shut down when they finish executing statements and the auto-suspend period elapses (Correct)

**Which encryption algorithm used by snowflake?**

Options are :

- DES
- MD5
- RSA
- AES 256 bit encryption (Correct)

**Answer :**

\*AES 256 bit encryption (Correct)

**What technique does Snowflake use to limit the number of micro-partitions scanned by each query?**

Options are :

- Pruning (Correct)
- Indexing
- Map Reduce
- B-Tree

**Answer :**

\*Pruning (Correct)

**Snowflake Certification Preparation Set 10**

**If a task is still running when the next scheduled execution time occurs, what happens?**

Options are :

- *The tasks that is running fails and retriggers*
- *The tasks that is running fails and stops*
- *The next execution of the tasks starts as soon as the current execution gets over*
- *That schedule time is skipped (Correct)*

**Answer :**

\*That schedule time is skipped (Correct)

**Snowpipes maintain the load history metadata for how many days**

Options are :

- 365
- 14 (*Correct*)
- 64
- 30

**Answer :**

\*14 (Correct)

**What is the default compression of the output files generated by Snowflake when unloading data**

Options are :

- *gzip (Correct)*
- *bzip2*
- *brotli*
- *Zstandard*

**Answer :**

\**gzip (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 33**

**If you change the data retention period for a table, the new retention period impacts all data that is active, as well as any data currently in Time Travel**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Which parameter is required to be set to modify the time travel period of any object in snowflake?**

Options are :

- DATA\_RETENTION\_TIME\_IN\_MINUTES
- DATA\_RETENTION\_TIME\_IN\_HOURS
- DATA\_RETENTION\_TIME\_IN\_DAYS (Correct)
- DATA\_RETENTION\_TIME\_IN\_WEEKS

**Answer :**

\*DATA\_RETENTION\_TIME\_IN\_DAYS (Correct)

**MAX\_DATA\_EXTENSION\_TIME\_IN\_DAYS - This parameter is used for what purpose?**

Options are :

- Maximum number of days for which Snowflake can extend the data retention period for tables to prevent streams on the tables from becoming stale (Correct)
- Maximum number of days for which Snowflake can extend the data retention period for fail safe of an object
- Maximum number of days for which Snowflake can extend the data retention period for time travel of an object

**Answer :**

\*Maximum number of days for which Snowflake can extend the data retention period for tables to prevent streams on the tables from becoming stale (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 19**

**When a table data changes (for e.g. few records are updated and few are deleted) and the table is within the time travel retention period, which of the following statement is true?**

Options are :

- Snowflake maintains only the information pertaining to these changed records (i.e. updated & deleted ones) (Correct)

- Snowflake takes a full copy of the table every time any change is made to the table

**Answer :**

\*Snowflake maintains only the information pertaining to these changed records (i.e. updated & deleted ones) (Correct)

**The maximum number of server cluster that can be defined for a multi cluster warehouse is**

Options are :

- 15
- 10 (Correct)
- 5
- 20

**Answer :**

\*10 (Correct)

**With Time Travel, a Snowflake administrator can, restore tables, schemas, and databases that have been dropped after the retention period lapses**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

## Snowflake Certification Preparation Set 2

**A table has a data retention policy of 10 days. Few records in the table were modified 5 days earlier, and therefore that data is already into 5 days of the time travel. Now if the data retention of the table is reduced to 3 days, what happens to the data that was in 5 days of time travel?**

Options are :

- The data would be removed completely (i.e. not even moved to fail safe)
- That data will remain in time travel for 10 days even if the retention period is reduced
- None is true

- *The data will move to fail safe (Correct)*

**Answer :**

\*The data will move to fail safe (Correct)

**Which scaling policy, shut down the cluster after 2 to 3 consecutive successful checks (performed at 1 minute intervals), which determine whether the load on the least-loaded cluster could be redistributed to the other clusters without spinning up the cluster again**

Options are :

- *Economy*
- *Legacy*
- *Standard (Correct)*
- *None*

**Answer :**

\*Standard (Correct)

**Changing the time travel retention period for your account or individual objects changes the value for all lower-level objects that do not have a retention period explicitly set**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 1

**Changing the retention period for your account or individual objects changes the value for all lower-level objects that do not have a retention period explicitly set. For example: If you change the retention period at the account level, all databases, schemas, and tables that do not have an explicit retention period automatically inherit the new retention period.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**What is correct about multi cluster virtual warehouses?**

Options are :

- *Multi Cluster warehouse will spin up additional clusters when under high query workload which is beyond the current capacity (Correct)*
- *Maximum of 5 multi cluster warehouse can be created in one snowflake account*
- *Multi Cluster warehouses cannot be suspended*
- *Multi Cluster will remove the additional clusters as query demand decreases (Correct)*

**Answer :**

\*Multi Cluster warehouse will spin up additional clusters when under high query workload which is beyond the current capacity (Correct)

\*Multi Cluster will remove the additional clusters as query demand decreases (Correct)

**In a multi cluster warehouse, for each cluster there is a way we can specify different warehouse sizes**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 38**

**Which scaling policy, starts additional cluster only if the system estimates there is enough query load to keep the cluster busy for at least 6 minutes.**

Options are :

- *None*
- *Economy (Correct)*
- *Standard*
- *Legacy*

**Answer :**

\*Economy (Correct)

**In maximized mode, what happens if the number of clusters are increased? (e.g. previously min & max was 2, but now min & max cluster size is changed to 4)**

Options are :

- *There will not be any immediate impact, and the additional clusters will start only when the query workload demands*
- *Specified number of clusters (in this case 4) would start immediately (Correct)*
- *Specified number of clusters would start after current queries that are executed are completed*
- *The number of clusters in maximized mode cannot be changed*

**Answer :**

\*Specified number of clusters (in this case 4) would start immediately (Correct)

**Fail-safe is not provided as a means for accessing historical data after the Time Travel retention period has ended. It is for use only by Snowflake to recover data that may have been lost or damaged due to extreme operational failures**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

## Snowflake Certification Preparation Set 10

**Which of the following compression methods cannot be detected automatically by snowflake**

Options are :

- *bzip2*
- *Zstandard (Correct)*
- *brotli (Correct)*
- *gzip*

**Answer :**

- \*Zstandard (Correct)
- \*brotli (Correct)

**Snowflake data storage costs include which types of data?**

Options are :

- *Metadata*
- *Cached results*
- *Persistent data stored in permanent tables (Correct)*
- *Data retained to enable data recovery (time travel and fail-safe) (Correct)*
- *Semi-structured data - additional fees*

**Answer :**

- \*Persistent data stored in permanent tables (Correct)
- \*Data retained to enable data recovery (time travel and fail-safe) (Correct)

**Select all of the answers that describe Snowflake micro-partitions. Micro-partitions:**

Options are :

- *Are written to a customer configured cloud storage location*
- *Enable horizontal and vertical query pruning (Correct)*
- *Will be updated as their row & column values are updated*
- *Are the physical data files that comprise Snowflake logical tables (Correct)*

**Answer :**

- \*Enable horizontal and vertical query pruning (Correct)
- \*Are the physical data files that comprise Snowflake logical tables (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 16](#)

**Snowflake minimizes the amount of storage required for historical data by maintaining only the information required to restore the individual table rows that were updated or deleted**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Which scaling policy, shut down the cluster after 5 to 6 consecutive successful checks (performed at 1 minute intervals), which determine whether the load on the least-loaded cluster could be redistributed to the other clusters without spinning up the cluster again**

Options are :

- *None*
- *Standard*
- *Economy (Correct)*
- *Legacy*

**Answer :**

\*Economy (Correct)

**What is the default time limit on a single run of a task**

Options are :

- *10 min*
- *30 min*
- *60 min (Correct)*
- *There is no limit as such*

**Answer :**

\*60 min (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 4

**Multi-cluster warehouses are best utilized for scaling resources to improve concurrency for users/queries**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which of the following parameters are provided by Snowflake to be used with AT & BEFORE clause to pinpoint the exact historical data you wish to access**

Options are :

- *TIMESTAMP*
- *OFFSET (time difference in seconds from the present time)*
- *STATEMENT (identifier for statement, e.g. query ID)*
- *All options are correct (Correct)*

**Answer :**

\*All options are correct (Correct)

**When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed, columnar format**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which Snowflake Edition can a user choose when creating a trial account**

Options are :

- *Standard Edition (Correct)*
- *Enterprise Edition (Correct)*
- *Free-trial Edition*
- *Virtual Private Snowflake*
- *Business Critical Edition (Correct)*

**Answer :**

\*Standard Edition (Correct)

\*Enterprise Edition (Correct)

\*Business Critical Edition (Correct)

**Tri-Secret Secure option is available in which snowflake edition ?**

Options are :

- *Business critical (or higher) (Correct)*
- *Enterprise edition*
- *All editions*

**Answer :**

\*Business critical (or higher) (Correct)

**How many credits per hour is consumed by a 2X-Large warehouse**

Options are :

- *16*
- *32 (Correct)*
- *64*
- *128*

**Answer :**

\*32 (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 40**

**What kind of stream can be created on external table ?**

Options are :

- *Standard*
- *Append-only*
- *Insert-only (Correct)*
- *None of the above*

**Answer :**

\*Insert-only (Correct)

**Multi-cluster warehouses is only available in Enterprise Edition and above**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**What privileges does SYSADMIN have by default?**

Options are :

- *CREATE DATABASE (Correct)*
- *CREATE SHARE*
- *CREATE WAREHOUSE (Correct)*
- *CREATE USER*
- *CREATE ROLE*

**Answer :**

\**CREATE DATABASE (Correct)*

\**CREATE WAREHOUSE (Correct)*

### [Snowflake SnowPro Core Certification Practice Tests Set 14](#)

**What are the two major costs categories in Snowflake ? (Select the best options)**

Options are :

- *Storage (Correct)*
- *Compute (Correct)*
- *Data transfer*
- *Cloud services*
- *None of the above*

**Answer :**

\**Storage (Correct)*

\**Compute (Correct)*

**Snowflake Web Interface data loading wizard is designed to load small amounts of data where each data file is smaller than 50MB in size. TRUE or FALSE load**

Options are :

- *TRUE (Correct)*
- *FLASE*

**Answer :**

\**TRUE (Correct)*

**Which syntax allows you to concat the column named 'EMPLOYEE' between two % signs ? Select all that apply**

Options are :

- *%' || EMPLOYEE || '%' (Correct)*
- *% + EMPLOYEE + %*
- *CONCAT('%', EMPLOYEE, '%')*
- *concat('%', concat(EMPLOYEE, '%')) (Correct)*
- *All of the above*

**Answer :**

*\*%' || EMPLOYEE || '%' (Correct)*

*\*concat('%', concat(EMPLOYEE, '%')) (Correct)*

## [Snowflake Certification Preparation Set 6](#)

**Which of the following terms describes Snowflake's architecture ?**

Options are :

- *Multi-cluster shared data (Correct)*
- *Shared disk*
- *Columnar shared nothing*
- *Coud-native shared memory*

**Answer :**

*\*Multi-cluster shared data (Correct)*

**Scaling a warehouse UP will decrease the number of clusters. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**A client has set Time Travel to keep a table data for 27 days. 35 days after dropping the table, their want to get it back. They ask you if it is possible to recover the data of that table.**

Options are :

- *Yes, it is possible*
- *No, it is not possible (Correct)*

**Answer :**

\**No, it is not possible (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 30

**Who can use the metadata cache ?**

Options are :

- *Only the user who execute the query*
- *Only user with the same role*
- *Only account administrator*
- *Everyone (Correct)*
- *None of the above*

**Answer :**

\**Everyone (Correct)*

**How long does the query result cache last ?**

Options are :

- *72 hours*

- 48 hours
- 24 hours (Correct)
- 1 hour

**Answer :**

\*24 hours (Correct)

**A company operating in the health sector wants to move to Snowflake. The company asks you about the edition of Snowflake to choose to comply with HIPAA regulation. Select the minimum edition of Snowflake to recommend ?**

Options are :

- Enterprise Edition
- Business Critical Edition (Correct)
- Standard Edition
- Virtual Private Snowflake

**Answer :**

\*Business Critical Edition (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 23

**Which is Snowflake Architecture layer ? (Select all that apply)**

Options are :

- Database storage Layer (Correct)
- Account Layer
- Cloud services Layer (Correct)
- Query processing (Correct)
- Administration Layer

**Answer :**

\*Database storage Layer (Correct)  
\*Cloud services Layer (Correct)  
\*Query processing (Correct)

**Select the term that is associated with compute layer?**

Options are :

- *Query planning*
- *Query optimization*
- *Query processing (Correct)*
- *Query Design*
- *None of the above*

**Answer :**

\*Query processing (Correct)

**Which of the following terms are NOT associated with the Cloud Services Layer ?**

Options are :

- *Query Planning*
- *Query Optimization*
- *Query Processing (Correct)*
- *Query Design*
- *Query Compilation*

**Answer :**

\*Query Processing (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 30**

**Which of the following objects are stored within schemas ?**

Options are :

- *Stages (Correct)*
- *File formats (Correct)*
- *Sequences (Correct)*
- *Stored procedures (Correct)*
- *Virtual Warehouses*

- *Roles*

**Answer :**

- \*Stages (Correct)
- \*File formats (Correct)
- \*Sequences (Correct)
- \*Stored procedures (Correct)

**Which role in Snowflake allows a user to view and operate all objects in the account, manage Snowflake billing and credit data, and stop any running SQL statements ?**

Options are :

- *USERADMIN*
- *SECURITYADMIN*
- *ACCOUNTADMIN (Correct)*
- *SYSADMIN*
- *None of the above*

**Answer :**

- \**ACCOUNTADMIN (Correct)*

**A client ask your support to choose a geographic deployment region for a Snowflake account. What factors should you consider ? Select the two best options.**

Options are :

- *Proximity to the point of service (Correct)*
- *Additional fees charged for regions with geo-political unrest*
- *Number of availability zones within a region (Correct)*
- *End-user perceptions of glamorous or trendy geographic locations*

**Answer :**

- \**Proximity to the point of service (Correct)*
- \**Number of availability zones within a region (Correct)*

### [Snowflake SnowPro Core Certification Exam Questions Set 8](#)

**When running a `SELECT COUNT(*)` on a table, which of the following statements is TRUE ?**

Options are :

- Only XS warehouse is needed because Count statistics are stored in Results Cache
- No warehouse is needed because Count statistics are stored in Metadata Cache (Correct)
- No warehouse will be needed because all users share count statistics
- Only XS warehouse will be needed because all roles share count statistics

**Answer :**

\*No warehouse is needed because Count statistics are stored in Metadata Cache (Correct)

**Snowflake account hosted on AWS cannot load data from a file that is staged on GCP or Azure.**  
**TRUE or FALSE**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Each Snowflake account is limited to a maximum of 10 shares**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 20**

**What ROLE should create users and roles in Snowflake ? Select all that apply**

Options are :

- ACCOUNTADMIN
- USERADMIN (Correct)
- SYSADMIN

- *SECURITYADMIN* (*Correct*)
- *PUBLIC*

**Answer :**

- \**USERADMIN* (*Correct*)
- \**SECURITYADMIN* (*Correct*)

**Select all Snowflake stages.**

Options are :

- *Internal Named Stage* (*Correct*)
- *External Named Stage* (*Correct*)
- *User Stage* (*Correct*)
- *Table Stage* (*Correct*)
- *Account Stage* (*Correct*)

**Answer :**

- \**Internal Named Stage* (*Correct*)
- \**External Named Stage* (*Correct*)
- \**User Stage* (*Correct*)
- \**Table Stage* (*Correct*)
- \**Account Stage* (*Correct*)

**What are the two scaling policies for a multi-cluster warehouse ?**

Options are :

- *Standard* (*Correct*)
- *Enterprise*
- *Economy* (*Correct*)
- *Business Critical*

**Answer :**

- \**Standard* (*Correct*)
- \**Economy* (*Correct*)

**Snowflake Practice Questions - SnowPro Core Certified Set 2**

**What kind of solution is Snowflake**

Options are :

- *Infrastructure as a Service (IaaS)*
- *Platform as a Service (PaaS)*
- *Software as a Service (SaaS) (Correct)*

**Answer :**

\*Software as a Service (SaaS) (Correct)

**The COPY command supports column reordering, column omission, and CAST using a SELECT statement**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**When deciding whether to use bulk loading or Snowpipe, which factors should you consider ?**

Options are :

- *How often you will load data (Correct)*
- *Location of data (on-premise or cloud) (Correct)*
- *Data format (structured or semi-structured)*
- *Number of files to load at once (Correct)*

**Answer :**

\*How often you will load data (Correct)  
\*Location of data (on-premise or cloud) (Correct)  
\*Number of files to load at once (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 24](#)

**A task tree can have task spanning multiple schemas ?**

Options are :

- *TRUE*

- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**How can you disable < auto-suspend > option for a virtual warehouse**

Options are :

- Choose 'Never' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI (Correct)
- Choose 'Disable' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI
- ALTER WAREHOUSE set AUTO\_SUSPEND = 0 (Correct)
- ALTER WAREHOUSE set AUTO\_SUSPEND = NULL (Correct)
- ALTER WAREHOUSE set AUTO\_SUSPEND = -1
- None of the above

**Answer :**

\*Choose 'Never' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI (Correct)

\*ALTER WAREHOUSE set AUTO\_SUSPEND = 0 (Correct)

\*ALTER WAREHOUSE set AUTO\_SUSPEND = NULL (Correct)

**Scaling policy for a multi-cluster warehouse ONLY apply if it is running in Auto-scale mode. TRUE or FALSE**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

### Snowflake Certification Preparation Set 13

**In Snowflake Web UI Console, in which tab can user check for Query performance ?**

Options are :

- Warehouses Tab

- *Databases Tab*
- *History Tab (Correct)*
- *Account Tab*
- *None of the above*

**Answer :**

\*History Tab (Correct)

**Which of the following have drivers/connectors (or information about where to find them) available via Help->Downloads in the Snowflake WebUI?**

Options are :

- *Go (Correct)*
- *R*
- *Node.Js (Correct)*
- *JDBC (Correct)*
- *Hive*
- *Spark*

**Answer :**

\*Go (Correct)  
\*Node.Js (Correct)  
\*JDBC (Correct)

**To which of these Snowflake locations you CANNOT upload files from a local machine with PUT command ? Select all that apply.**

Options are :

- *Named internal stage*
- *Internal stage for specified table*
- *Named External stage (Correct)*
- *Internal stage for the current user*

**Answer :**

\*Named External stage (Correct)

## Snowflake Practice Questions - SnowPro Core Certified Set 2

**Is it possible to disable Time Travel for a Snowflake Account ?**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**Fail-safe is a 7-day history of data and is automatically available on which table types ?**

Options are :

- *Permanent (Correct)*
- *Temporary*
- *Transient*
- *External*

**Answer :**

\*Permanent (Correct)

**Which of the following terms describes Snowflake's Architecture ?**

Options are :

- *Shared Disk*
- *Shared Nothing*
- *Shared Data (Correct)*
- *Shared Memory*

**Answer :**

\*Shared Data (Correct)

## Snowflake Certification Preparation Set 8

**Your organization wants to use Column-level security to apply masking policies to columns in tables or views. Which Snowflake's edition offers this features ? Select all that apply.**

Options are :

- *Enterprise Edition (Correct)*
- *Business Critical Edition (Correct)*
- *Virtual Private Snowflake (Correct)*
- *Standard*
- *None of the above*

**Answer :**

\*Enterprise Edition (Correct)  
\*Business Critical Edition (Correct)  
\*Virtual Private Snowflake (Correct)

**What is the recommended file size to optimize the number of parallel operations for a data load ?  
(Select the best option)**

Options are :

- *100-250MB in size compressed (Correct)*
- *100-250GB in size compressed*
- *100-250TB in size compressed*

**Answer :**

\*100-250MB in size compressed (Correct)

**Query load is calculated by dividing the execution time (in seconds) of all queries in an interval by the total time (in seconds) for the interval. TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 13**

**What influences Snowflake pricing ?**

Options are :

- *Storage and Compute (Correct)*
- *Amount of data queried from Snowflake*
- *Amount of data scanned during querying Snowflake*
- *None of the above*

**Answer :**

\*Storage and Compute (Correct)

**What attributes make Snowflake a true SaaS solution?**

Options are :

- *No hardware to install or configure (Correct)*
- *No maintenance upgrades or patches to install (Correct)*
- *Transparent releases don't require user intervention (Correct)*
- *No creation of user accounts or roles is required*
- *No query processing costs*
- *No data storage costs*

**Answer :**

\*No hardware to install or configure (Correct)

\*No maintenance upgrades or patches to install (Correct)

\*Transparent releases don't require user intervention (Correct)

**Snowflake Web UI data loading wizard can be used to load a 1GB file. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FLASE (Correct)*

**Answer :**

\*FLASE (Correct)

**When staging uncompressed files in a Snowflake stage, the files are automatically compressed using gzip, unless compression is explicitly disabled.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Snowflake Tasks can be event based as well along with time based?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**What are the additional metadata columns available in Stream to track changes?**

Options are :

- *METADATA\$ROW\_ID (Correct)*
- *METADATA\$ISDELETE*
- *METADATA\$ISINSERT*
- *METADATA\$ISUPDATE (Correct)*
- *METADATA\$ACTION (Correct)*

**Answer :**

\**METADATA\$ROW\_ID (Correct)*  
\**METADATA\$ISUPDATE (Correct)*  
\**METADATA\$ACTION (Correct)*

**SnowFlake SnowPro Advanced Architect Practice Exam Set 3**

**What is the standard time travel retention period which is enabled for all snowflake accounts?**

Options are :

- 7 days
- 24 hours (Correct)
- 48 hours
- 14 days

**Answer :**

\*24 hours (Correct)

**Storage fees are incurred for maintaining historical data during the Time Travel but not during Fail-safe periods as Fail safe is managed by Snowflake**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Maximum number of clusters that can be added in multi cluster warehouse is?**

Options are :

- 5
- 15
- 10 (Correct)
- 20

**Answer :**

\*10 (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 6

**Using the Snowflake Web interface only the files from local system can be loaded, and the files available on cloud storage like AWS, Azure cannot be loaded**

Options are :

- FALSE (Correct)

- TRUE

**Answer :**

\*FALSE (Correct)

**Fail Safe and time travel are one and the same thing**

Options are :

- FALSE (Correct)

- TRUE

**Answer :**

\*FALSE (Correct)

**Time travel can be disabled at account level by any user?**

Options are :

- TRUE

- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### **Tackling Snowflake Certification - Practice Questions Set 3**

**The standard retention period for time travel is of how many days, which is automatically enabled for all Snowflake accounts**

Options are :

- 7

- 1 (Correct)

- 15

- 30

**Answer :**

\*1 (Correct)

**Time Travel in Snowflake is useful for which of the following tasks**

Options are :

- *Analyzing data usage/manipulation over specified periods of time (Correct)*
- *Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted (Correct)*
- *Duplicating and backing up data from key points in the past (Correct)*
- *Restoring data in event of a system failure or other catastrophic event, e.g. a hardware failure or security breach.*

**Answer :**

- \*Analyzing data usage/manipulation over specified periods of time (Correct)
- \*Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted (Correct)
- \*Duplicating and backing up data from key points in the past (Correct)

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Auto-scale mode for 3 hours, in which Cluster 1 runs continuously, Cluster 2 runs continuously for the entire 2nd hour and 30 minutes in the 3rd hour, Cluster 3 runs for 30 minutes in the 3rd hour**

Options are :

- 18
- 24
- 20 (*Correct*)
- 16

**Answer :**

- \*20 (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 23](#)

**If a 3X-Large multi-cluster warehouse runs 1 cluster for one full hour and then runs 2 clusters for the next full hour, the total number of credits billed would be**

Options are :

- 128
- 192 (*Correct*)
- 64

• 256

**Answer :**

\*192 (Correct)

**Which of the following tables have no fail safe period?**

Options are :

- *Transient (Correct)*
- *Temporary (Correct)*
- *Permanent*

**Answer :**

\*Transient (Correct)  
\*Temporary (Correct)

**What happens to the data of a table that is currently in time travel, when the retention period of a table is increased? E.g. The retention period is increased from 10 days to 20 days.**

Options are :

- *When data retention is changed, all the data in time travel moves to fail safe by default*
- *None is true*
- *The data in time travel will be retained for the extended period of retention period i.e. for 20 days (Correct)*
- *The data in time travel will be retained as per previous retention period only i.e. for 10 days only*

**Answer :**

\*The data in time travel will be retained for the extended period of retention period i.e. for 20 days (Correct)

### Snowflake Certification Preparation Set 23

**If you have a table with a 10-day retention period and increase the period to 20 days, then what happens to the data that is currently in time travel for past 2 days.**

Options are :

- *The data would move to fail safe as per new retention period i.e. will move to fail safe after 20 days of time travel (Correct)*

- The data will move to fail safe immediately
- The data would move to fail safe as per previous retention period only i.e. will move to fail safe after 10 days of time travel
- The table will lose all time travel data, and the new time travel retention period will be applied going forward

**Answer :**

\*The data would move to fail safe as per new retention period i.e. will move to fail safe after 20 days of time travel (Correct)

**Multi cluster virtual warehouses are designed to handle -**

Options are :

- Queuing issues (Correct)
- Large number of concurrent users (Correct)
- Large Complex Queries
- Large number of concurrent queries (Correct)

**Answer :**

\*Queuing issues (Correct)

\*Large number of concurrent users (Correct)

\*Large number of concurrent queries (Correct)

**The Snowflake credit usage for multi cluster virtual warehouses is based on**

Options are :

- For multi cluster credits are always charged based on the maximum cluster size that is defined
- Total Number of clusters running and for the time period they are running (Correct)
- For multi cluster credits are always charged based on the average cluster size that is defined

**Answer :**

\*Total Number of clusters running and for the time period they are running (Correct)

**SnowFlake SnowPro Advanced Architect Practice Exam Set 7**

**select \* from my\_table before(statement => '8e5d0ca9-005e-44e6-b858-a8f5b37c5726');** This query will select historical data from "my\_table" up to, but not including any changes made by the specified statement

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

If a row is added and then updated in the current offset, then the METADATA\$ISUPDATE row records a TRUE value.

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

Storage fees are incurred for maintaining historical data during both the Time Travel and Fail-safe periods

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 4](#)

A warehouse can be started at any time, excluding the initial creation time

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**What is the maximum compressed size limit for an individual row in the VARIANT data type column**

Options are :

- 10 MB
- 100 MB
- 16 MB (Correct)
- 18 MB

**Answer :**

\*16 MB (Correct)

**Time travel is available for which table types?**

Options are :

- Permanent (Correct)
- Transient (Correct)
- External
- Temporary (Correct)

**Answer :**

\*Permanent (Correct)  
\*Transient (Correct)  
\*Temporary (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 7

**Which of the following are true of Multi-Cluster warehouses? Select all that apply.**

Options are :

- Adds clusters automatically based on query activity (Correct)
- Multi-cluster warehouses will never auto-suspend
- Scales down when query activity slows (Correct)
- Sizes each cluster optimally based on the queries

**Answer :**

- \*Adds clusters automatically based on query activity (Correct)
- \*Scales down when query activity slows (Correct)

**Fail Safe is configurable by the user**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**The method that Snowflake uses to limit the number of micro-partitions accessed during a query is called :**

Options are :

- *Pruning (Correct)*
- *Cache Operations*
- *Index Scan*
- *Indexing*

**Answer :**

- \**Pruning (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 43](#)

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Auto Scale mode for 2 hours, where Cluster 1 runs continuously for 2 hours, Cluster 2 runs continuously for 2nd hour only, Cluster 3 runs for 30 minutes only during 2nd hour**

Options are :

- 24
- 12
- 14 (*Correct*)
- 16

**Answer :**

\*14 (Correct)

**Which table type disappears after the close of the session and therefore has no fail-safe, and no time travel options after the close of the session?**

Options are :

- *Temporary (Correct)*
- *Transient*
- *Permanent*
- *External*

**Answer :**

\*Temporary (Correct)

**Which of the following are types of caching use by Snowflake? Select all that apply**

Options are :

- *Index Caching*
- *Metadata caching (Correct)*
- *Warehouse caching (Correct)*
- *Query result caching (Correct)*

**Answer :**

\*Metadata caching (Correct)

\*Warehouse caching (Correct)

\*Query result caching (Correct)

### Snowflake Certification Preparation Set 21

**The cron expression in a task definition supports specifying a time zone**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**What is the maximum data retention period for Fail-Safe access? (Select one option only)**

Options are :

- 1 Day
- 90 Days
- 30 Days
- 7 Days (Correct)

**Answer :**

\*7 Days (Correct)

**Check all true statements about Fail-safe:**

Options are :

- Only a Snowflake employee can recover data from fail-safe storage (Correct)
- Fail-safe is a reliable way to create Dev/Test/QA and other environments
- The data stored as part of fail-safe is part of storage costs charged to customers (Correct)
- Fail-safe is not available for tables that have Time Travel.

**Answer :**

\*Only a Snowflake employee can recover data from fail-safe storage (Correct)

\*The data stored as part of fail-safe is part of storage costs charged to customers (Correct)

### Tackling Snowflake Certification - Practice Questions Set 4

**Fail Safe period starts right after time travel retention period ends**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**The Snowflake metadata repository contains which of the following? Select all that apply**

Options are :

- *Tracking of all versions of the table data within the data retention window (Correct)*
- *Aggregated data from the table*
- *References to all of the micropartition files for that table (Correct)*
- *Table definitions (Correct)*

**Answer :**

\*Tracking of all versions of the table data within the data retention window (Correct)

\*References to all of the micropartition files for that table (Correct)

\*Table definitions (Correct)

**If you have a table with a 10-day retention period and increase the period to 20 days, then the data that was moved to fail safe on 11th day, will be moved back to Time Travel, and will have more 9 days left for time travel retention end**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

### Snowflake Certification Preparation Set 5

**In maximized multi-cluster warehouse same value is specified for both maximum & minimum clusters (except for the value 1)**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Data from Fail-Safe can only be accessed / restored by Snowflake Employees**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which scaling policy, starts additional cluster immediately when either a query is queued or the system detects that there is one more query than the currently-running clusters can execute.**

Options are :

- *Legacy*
- *Standard (Correct)*
- *Economy*
- *None*

**Answer :**

\*Standard (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 10

**Fail Safe is not supported in Standard Snowflake edition**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**Snowflake data storage costs are calculated based on**

Options are :

- *Compressed Size (Correct)*
- *Amount Stored on First Day of Month*
- *Amount Stored - Daily Average (Correct)*
- *Uncompressed Size*
- *Amount Stored on Last Day of Month*

**Answer :**

- \*Compressed Size (Correct)
- \*Amount Stored - Daily Average (Correct)

**With Time Travel, a Snowflake administrator can, analyze data usage/manipulation over specified periods of time**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \**TRUE (Correct)*

### Snowflake SnowPro Core Certification Exam Questions Set 2

**Which of the following could be considered synonymous with the term "Warehouse Cache"? (Choose 3)**

Options are :

- *Metadata Cache*
- *SSD Cache (Correct)*
- *Raw Data Cache (Correct)*
- *Data Cache (Correct)*
- *Result Set Cache*

**Answer :**

- \**SSD Cache (Correct)*
- \**Raw Data Cache (Correct)*
- \**Data Cache (Correct)*

**Each server in a warehouse cluster has:**

Options are :

- *A position in the warehouse that is maintained, even when the warehouse is suspended or resized (Correct)*
- *Track of each query that was executed on that server*
- *An internal timer that tracks when each server was started (Correct)*

**Answer :**

- \*A position in the warehouse that is maintained, even when the warehouse is suspended or resized (Correct)
- \*An internal timer that tracks when each server was started (Correct)

**What all resources are provided by Virtual Warehouse?**

Options are :

- *Memory (Correct)*
- *Metadata Storage*
- *Permanent Storage*
- *CPU (Correct)*
- *Temporary Storage (Correct)*

**Answer :**

- \*Memory (Correct)
- \*CPU (Correct)
- \*Temporary Storage (Correct)

**Scaling out of a warehouse can be done only in Enterprise edition and higher**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**Snowpipe uses compute resources provided by Snowflake (i.e. a serverless compute model)**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

**There is no benefit to stopping a warehouse before the first 60-second period is over because the credits have already been billed for that period**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 3

**Which minimum Snowflake license allows Multicloud data virtual warehouse capability?**

Options are :

- *VPS*
- *Business Critical*
- *Standard*
- *Enterprise (Correct)*

**Answer :**

\*Enterprise (Correct)

**A Snowflake session can have multiple current warehouse at a time**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**What is the column datatype used for ingesting semi-structured data?**

Options are :

- *VARCHAR*
- *STRING*
- *VARIANT (Correct)*
- *BINARY*
- *DATETIME*

**Answer :**

\**VARIANT (Correct)*

#### Tackling Snowflake Certification - Practice Questions Set 4

**True or false: Virtual Warhouses cannot be resized while queries are running.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Virtual Warehouse credits are charged on what basis?**

Options are :

- *per milli second use*
- *per minute use*
- *per hour use*
- *per second use with no minimum charge*
- *per second use (with minimum charge of 1 minute) (Correct)*

**Answer :**

\**per second use (with minimum charge of 1 minute) (Correct)*

**Which cache type runs on a 24 hour clock?**

Options are :

- *Warehouse Cache*
- *Results Cache (Correct)*
- *Metadata Cache*

**Answer :**

\*Results Cache (Correct)

### **SnowFlake SnowPro Advanced Architect Practice Exam Set 2**

**Scaling a warehouse UP will decrease the number of clusters. (e.g. Small to Medium)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Scaling in is a manual process (e.g. Max Clusters to min clusters)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Scaling a warehouse OUT will increase the number of clusters. (e.g. Min to Max)**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

## Snowflake Certification Preparation Set 8

**Scaling a warehouse DOWN will decrease the number of servers. (e.g. Medium to Small)**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Which of the following are valid parameters when create a new virtual warehouse?**

Options are :

- Auto Suspend (Correct)
- Type of EC2
- Auto Resume (Correct)
- RAM Size
- CPU Counts
- Warehouse Size (Correct)

**Answer :**

\*Auto Suspend (Correct)  
\*Auto Resume (Correct)  
\*Warehouse Size (Correct)

**When a virtual warehouse is provisioned there is no benefit in stopping the virtual warehouse with in the first 60 seconds because that time period has already been billed.**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 10

**For File formats like JSON, Avro the only supported character set is UTF-8**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**User A and User B can access one another's result sets from the Results Cache, as long as which of the following are true? (Choose two)**

Options are :

- *They use the same Role. (Correct)*
- *They run with 59 minutes of one another*
- *They use the same Warehouse.*
- *They run the exact same SQL Text/Query. (Correct)*
- *They sign in to the same Session*

**Answer :**

\**They use the same Role. (Correct)*

\**They run the exact same SQL Text/Query. (Correct)*

**Warehouses can also be resized at any time, even while running, to accommodate the need for more or less compute resources**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 42**

**In auto scale mode, if new\_max\_clusters < running clusters, then excess clusters shut down immediately**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Out of below four options, which options require a running warehouse when retrieving data from them?**

Options are :

- *Metadata Cache*
- *Results Cache*
- *Warehouse Cache (Correct)*
- *Local Disk Cache (Correct)*

**Answer :**

\**Warehouse Cache (Correct)*

\**Local Disk Cache (Correct)*

**The columns in the query history include the QueryID, the SQL Text, the Warehouse name, the Warehouse Size, the Session ID and others. Which column is a good indicator of whether a Warehouse was used (and Compute costs incurred) by a query?**

Options are :

- *Warehouse Name*
- *SQL Text*
- *Warehouse Size (Correct)*
- *QueryID*
- *Session ID*

**Answer :**

\**Warehouse Size (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 28**

**What is the size limit on the variant data type in snowflake?**

Options are :

- 16 MB (uncompressed)
- 16 MB (compressed) (Correct)
- 10 MB (compressed)
- 10 MB (uncompressed)

**Answer :**

\*16 MB (compressed) (Correct)

**Whenever the warehouse size is reduced or the number of clusters are decreased, the servers are always removed in reverse order of when they were added (LIFO - Last In First Out)**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**In auto scale multi-cluster warehouse same value is specified for both maximum & minimum clusters (except for the value 1)**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

## Snowflake Certification Preparation Set 21

**When a warehouse is suspended, the idle servers in that warehouse are immediately shut down, but any servers that are executing queries are allowed to run until the queries are completed. Such servers that are waiting to shut down are considered to be in which mode?**

Options are :

- Shutting Down

- *Suspending*
- *Quiesce (Correct)*
- *Quiet*

**Answer :**

\*Quiesce (Correct)

**How much size limit the VARIANT data type imposes on individual rows.**

Options are :

- *10 MB (compressed)*
- *16 MB (uncompressed)*
- *10 MB (uncompressed)*
- *16 MB (compressed) (Correct)*

**Answer :**

\*16 MB (compressed) (Correct)

**JSON arrays are written inside which type of brackets?**

Options are :

- *Round Brackets*
- *Curly Braces*
- *Square Brackets (Correct)*

**Answer :**

\*Square Brackets (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 5**

**What are the resources that a virtual warehouse provides?**

Options are :

- *Temporary Storage (Correct)*
- *Long term data storage*
- *User & Role Management*

- *CPU (Processing Power) (Correct)*
- *Memory (Correct)*

**Answer :**

- \*Temporary Storage (Correct)
- \*CPU (Processing Power) (Correct)
- \*Memory (Correct)

**If a multi-cluster warehouse is resized, the new size applies to all the clusters for the warehouse, including clusters that are currently running and any clusters that are started after the warehouse is resized**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**When a warehouse is resized, which queries make use of the new size?**

Options are :

- *Both current and subsequent queries*
- *Only subsequent queries (Correct)*
- *Only currently running queries*

**Answer :**

- \*Only subsequent queries (Correct)

### Snowflake Certification Preparation Set 11

**Snowflake does not begin executing SQL statements submitted to a warehouse until all of the servers for the warehouse are successfully provisioned, unless any of the servers fail to provision**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**What action causes a Warehouse's cache to be purged?**

Options are :

- *Resuming the Warehouse*
- *A change to metadata Cache size*
- *Suspending the warehouse (Correct)*
- *The passing of 24 hours*

**Answer :**

\*Suspending the warehouse (Correct)

**When the warehouse is resized (scaled up or out), then the additional servers that get added will be available for use by any queries that are queued or newly submitted**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

## [Snowflake Certification Preparation Set 6](#)

**Which of the following are true about the variant data type in Snowflake? Select all that apply.**

Options are :

- *Optimized storage based on repeated elements (Correct)*
- *Requires a custom mapping for each record type*
- *Stored in a separate file format from structured data*
- *Can be queried using JSON path notation (Correct)*

**Answer :**

\*Optimized storage based on repeated elements (Correct)

\*Can be queried using JSON path notation (Correct)

**How many credits per hour is consumed by a 3X large warehouse**

Options are :

- 32
- 64 (*Correct*)
- 16
- 128

**Answer :**

\*64 (Correct)

**In auto scale mode, if new\_max\_clusters > running clusters, then no changes happen until additional clusters are needed**

Options are :

- FALSE
- TRUE (*Correct*)

**Answer :**

\*TRUE (Correct)

## Tackling Snowflake Certification - Practice Questions Set 2

**What is the full form of ORC?**

Options are :

- Optimized Row Columnar (*Correct*)
- Object Read Class
- Object Row Columnar
- Optimized Read Columnar

**Answer :**

\*Optimized Row Columnar (Correct)

**A warehouse was started and it was up and running for 62 seconds. For how many seconds will the warehouse be billed for?**

Options are :

- 70 seconds
- 120 seconds
- 62 seconds (Correct)
- 60 seconds

**Answer :**

\*62 seconds (Correct)

The FLATTEN command will parse nested objects into separate rows. One version of the FLATTEN command uses a join and the other uses an object keyword. Select the two words that represent the options used with the FLATTEN command

Options are :

- LATERAL (Correct)
- SCHEMATA
- TABLE (Correct)
- SIDEWAYS

**Answer :**

\*LATERAL (Correct)

\*TABLE (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 12

What can be the maximum credits consumed by a medium size warehouse with 3 clusters in one hour?

Options are :

- 8
- 24
- 12 (Correct)
- 6

**Answer :**

\*12 (Correct)

Scaling down is a automated process. (e.g. X-Large to Large)

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**Snowflake utilizes per-second billing (with a 60-second minimum each time the warehouse starts)**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**In a VARIANT column, NULL values are stored as a string containing the word "null" not the SQL NULL value.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Out of below four options, which options require a running warehouse when retrieving data from them?**

Options are :

- *Local Disk Cache (Correct)*
- *Warehouse Cache (Correct)*
- *Results Cache*
- *Metadata Cache*

**Answer :**

- \*Local Disk Cache (Correct)
- \*Warehouse Cache (Correct)

**In a multi cluster warehouse, auto-suspend and auto-resume apply only to the entire warehouse and not to the individual clusters in the warehouse**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \**TRUE (Correct)*

**When staging uncompressed files in a Snowflake stage, the files are automatically compressed using gzip, unless compression is explicitly disabled.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \**TRUE (Correct)*

**Snowflake Tasks can be event based as well along with time based?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**What are the additional metadata columns available in Stream to track changes?**

Options are :

- *METADATA\$ROW\_ID (Correct)*
- *METADATA\$ISDELETE*

- *METADATA\$ISINSERT*
- *METADATA\$ISUPDATE (Correct)*
- *METADATA\$ACTION (Correct)*

**Answer :**

- \**METADATA\$ROW\_ID (Correct)*
- \**METADATA\$ISUPDATE (Correct)*
- \**METADATA\$ACTION (Correct)*

### Snowflake SnowPro Advanced Architect Practice Exam Set 3

**What is the standard time travel retention period which is enabled for all snowflake accounts?**

Options are :

- *7 days*
- *24 hours (Correct)*
- *48 hours*
- *14 days*

**Answer :**

- \**24 hours (Correct)*

**Storage fees are incurred for maintaining historical data during the Time Travel but not during Fail-safe periods as Fail safe is managed by Snowflake**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**Maximum number of clusters that can be added in multi cluster warehouse is?**

Options are :

- *5*
- *15*

- 10 (Correct)

- 20

**Answer :**

\*10 (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 6

**Using the Snowflake Web interface only the files from local system can be loaded, and the files available on cloud storage like AWS, Azure cannot be loaded**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**Fail Safe and time travel are one and the same thing**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**Time travel can be disabled at account level by any user?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**The standard retention period for time travel is of how many days, which is automatically enabled for all Snowflake accounts**

Options are :

- 7
- 1 (Correct)
- 15
- 30

**Answer :**

\*1 (Correct)

**Time Travel in Snowflake is useful for which of the following tasks**

Options are :

- Analyzing data usage/manipulation over specified periods of time (Correct)
- Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted (Correct)
- Duplicating and backing up data from key points in the past (Correct)
- Restoring data in event of a system failure or other catastrophic event, e.g. a hardware failure or security breach.

**Answer :**

\*Analyzing data usage/manipulation over specified periods of time (Correct)

\*Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted (Correct)

\*Duplicating and backing up data from key points in the past (Correct)

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Auto-scale mode for 3 hours, in which Cluster 1 runs continuously, Cluster 2 runs continuously for the entire 2nd hour and 30 minutes in the 3rd hour, Cluster 3 runs for 30 minutes in the 3rd hour**

Options are :

- 18
- 24
- 20 (Correct)
- 16

**Answer :**

\*20 (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 23

If a 3X-Large multi-cluster warehouse runs 1 cluster for one full hour and then runs 2 clusters for the next full hour, the total number of credits billed would be

Options are :

- 128
- 192 (*Correct*)
- 64
- 256

**Answer :**

\*192 (Correct)

Which of the following tables have no fail safe period?

Options are :

- *Transient* (*Correct*)
- *Temporary* (*Correct*)
- *Permanent*

**Answer :**

\**Transient* (*Correct*)

\**Temporary* (*Correct*)

What happens to the data of a table that is currently in time travel, when the retention period of a table is increased? E.g. The retention period is increased from 10 days to 20 days.

Options are :

- When data retention is changed, all the data in time travel moves to fail safe by default
- None is true
- The data in time travel will be retained for the extended period of retention period i.e. for 20 days (*Correct*)

- *The data in time travel will be retained as per previous retention period only i.e. for 10 days only*

**Answer :**

\*The data in time travel will be retained for the extended period of retention period i.e. for 20 days (Correct)

### Snowflake Certification Preparation Set 23

**If you have a table with a 10-day retention period and increase the period to 20 days, then what happens to the data that is currently in time travel for past 2 days.**

Options are :

- *The data would move to fail safe as per new retention period i.e. will move to fail safe after 20 days of time travel (Correct)*
- *The data will move to fail safe immediately*
- *The data would move to fail safe as per previous retention period only i.e. will move to fail safe after 10 days of time travel*
- *The table will lose all time travel data, and the new time travel retention period will be applied going forward*

**Answer :**

\*The data would move to fail safe as per new retention period i.e. will move to fail safe after 20 days of time travel (Correct)

**Multi cluster virtual warehouses are designed to handle -**

Options are :

- *Queuing issues (Correct)*
- *Large number of concurrent users (Correct)*
- *Large Complex Queries*
- *Large number of concurrent queries (Correct)*

**Answer :**

\*Queuing issues (Correct)

\*Large number of concurrent users (Correct)

\*Large number of concurrent queries (Correct)

**The Snowflake credit usage for multi cluster virtual warehouses is based on**

Options are :

- For multi cluster credits are always charged based on the maximum cluster size that is defined
- Total Number of clusters running and for the time period they are running (Correct)
- For multi cluster credits are always charged based on the average cluster size that is defined

**Answer :**

\*Total Number of clusters running and for the time period they are running (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 7

**select \* from my\_table before(statement => '8e5d0ca9-005e-44e6-b858-a8f5b37c5726');** This query will select historical data from "my\_table" up to, but not including any changes made by the specified statement

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**If a row is added and then updated in the current offset, then the METADATA\$ISUPDATE row records a TRUE value.**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Storage fees are incurred for maintaining historical data during both the Time Travel and Fail-safe periods**

Options are :

- FALSE

- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

#### SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

**A warehouse can be started at any time, excluding the initial creation time**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**What is the maximum compressed size limit for an individual row in the VARIANT data type column**

Options are :

- *10 MB*
- *100 MB*
- *16 MB (Correct)*
- *18 MB*

**Answer :**

\*16 MB (Correct)

**Time travel is available for which table types?**

Options are :

- *Permanent (Correct)*
- *Transient (Correct)*
- *External*
- *Temporary (Correct)*

**Answer :**

\*Permanent (Correct)

\*Transient (Correct)  
\*Temporary (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 7

**Which of the following are true of Multi-Cluster warehouses? Select all that apply.**

Options are :

- Adds clusters automatically based on query activity (Correct)
- Multi-cluster warehouses will never auto-suspend
- Scales down when query activity slows (Correct)
- Sizes each cluster optimally based on the queries

**Answer :**

\*Adds clusters automatically based on query activity (Correct)  
\*Scales down when query activity slows (Correct)

**Fail Safe is configurable by the user**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**The method that Snowflake uses to limit the number of micro-partitions accessed during a query is called :**

Options are :

- Pruning (Correct)
- Cache Operations
- Index Scan
- Indexing

**Answer :**

\*Pruning (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 43

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Auto Scale mode for 2 hours, where Cluster 1 runs continuously for 2 hours, Cluster 2 runs continuously for 2nd hour only, Cluster 3 runs for 30 minutes only during 2nd hour**

Options are :

- 24
- 12
- 14 (*Correct*)
- 16

**Answer :**

\*14 (Correct)

**Which table type disappears after the close of the session and therefore has no fail-safe, and no time travel options after the close of the session?**

Options are :

- *Temporary* (*Correct*)
- *Transient*
- *Permanent*
- *External*

**Answer :**

\**Temporary* (Correct)

**Which of the following are types of caching use by Snowflake? Select all that apply**

Options are :

- *Index Caching*
- *Metadata caching* (*Correct*)
- *Warehouse caching* (*Correct*)
- *Query result caching* (*Correct*)

**Answer :**

- \*Metadata caching (Correct)
- \*Warehouse caching (Correct)
- \*Query result caching (Correct)

## Snowflake Certification Preparation Set 21

**The cron expression in a task definition supports specifying a time zone**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

- \*TRUE (Correct)

**What is the maximum data retention period for Fail-Safe access? (Select one option only)**

Options are :

- 1 Day
- 90 Days
- 30 Days
- 7 Days (Correct)

**Answer :**

- \*7 Days (Correct)

**Check all true statements about Fail-safe:**

Options are :

- Only a Snowflake employee can recover data from fail-safe storage (Correct)
- Fail-safe is a reliable way to create Dev/Test/QA and other environments
- The data stored as part of fail-safe is part of storage costs charged to customers (Correct)
- Fail-safe is not available for tables that have Time Travel.

**Answer :**

- \*Only a Snowflake employee can recover data from fail-safe storage (Correct)

\*The data stored as part of fail-safe is part of storage costs charged to customers  
(Correct)

### Tackling Snowflake Certification - Practice Questions Set 4

**Fail Safe period starts right after time travel retention period ends**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**The Snowflake metadata repository contains which of the following? Select all that apply**

Options are :

- Tracking of all versions of the table data within the data retention window (Correct)
- Aggregated data from the table
- References to all of the micropartition files for that table (Correct)
- Table definitions (Correct)

**Answer :**

\*Tracking of all versions of the table data within the data retention window (Correct)  
\*References to all of the micropartition files for that table (Correct)  
\*Table definitions (Correct)

**If you have a table with a 10-day retention period and increase the period to 20 days, then the data that was moved to fail safe on 11th day, will be moved back to Time Travel, and will have more 9 days left for time travel retention end**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

### Snowflake Certification Preparation Set 5

**In maximized multi-cluster warehouse same value is specified for both maximum & minimum clusters (except for the value 1)**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Data from Fail-Safe can only be accessed / restored by Snowflake Employees**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which scaling policy, starts additional cluster immediately when either a query is queued or the system detects that there is one more query than the currently-running clusters can execute.**

Options are :

- *Legacy*
- *Standard (Correct)*
- *Economy*
- *None*

**Answer :**

\**Standard (Correct)*

**Snowflake SnowPro Core Certification Exam Questions Set 10**

**Fail Safe is not supported in Standard Snowflake edition**

Options are :

- *FALSE (Correct)*

- *TRUE*

**Answer :**

\*FALSE (Correct)

**Snowflake data storage costs are calculated based on**

Options are :

- *Compressed Size (Correct)*
- *Amount Stored on First Day of Month*
- *Amount Stored - Daily Average (Correct)*
- *Uncompressed Size*
- *Amount Stored on Last Day of Month*

**Answer :**

\*Compressed Size (Correct)

\*Amount Stored - Daily Average (Correct)

**With Time Travel, a Snowflake administrator can, analyze data usage/manipulation over specified periods of time**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 2](#)

**Which of the following could be considered synonymous with the term "Warehouse Cache"? (Choose 3)**

Options are :

- *Metadata Cache*
- *SSD Cache (Correct)*
- *Raw Data Cache (Correct)*

- *Data Cache (Correct)*

- *Result Set Cache*

**Answer :**

- \*SSD Cache (Correct)
- \*Raw Data Cache (Correct)
- \*Data Cache (Correct)

**Each server in a warehouse cluster has:**

Options are :

- *A position in the warehouse that is maintained, even when the warehouse is suspended or resized (Correct)*
- *Track of each query that was executed on that server*
- *An internal timer that tracks when each server was started (Correct)*

**Answer :**

- \*A position in the warehouse that is maintained, even when the warehouse is suspended or resized (Correct)
- \*An internal timer that tracks when each server was started (Correct)

**What all resources are provided by Virtual Warehouse?**

Options are :

- *Memory (Correct)*
- *Metadata Storage*
- *Permanent Storage*
- *CPU (Correct)*
- *Temporary Storage (Correct)*

**Answer :**

- \*Memory (Correct)
- \*CPU (Correct)
- \*Temporary Storage (Correct)

**Scaling out of a warehouse can be done only in Enterprise edition and higher**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Snowpipe uses compute resources provided by Snowflake (i.e. a serverless compute model)**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**There is no benefit to stopping a warehouse before the first 60-second period is over because the credits have already been billed for that period**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**SnowFlake SnowPro Advanced Architect Practice Exam Set 3**

**Which minimum Snowflake license allows Multicloud data virtual warehouse capability?**

Options are :

- *VPS*
- *Business Critical*
- *Standard*
- *Enterprise (Correct)*

**Answer :**

\*Enterprise (Correct)

**A Snowflake session can have multiple current warehouse at a time**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**What is the column datatype used for ingesting semi-structured data?**

Options are :

- *VARCHAR*
- *STRING*
- *VARIANT (Correct)*
- *BINARY*
- *DATETIME*

**Answer :**

\**VARIANT (Correct)*

### Tackling Snowflake Certification - Practice Questions Set 4

**True or false: Virtual Warhouses cannot be resized while queries are running.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Virtual Warehouse credits are charged on what basis?**

Options are :

- *per milli second use*
- *per minute use*
- *per hour use*
- *per second use with no minimum charge*
- *per second use (with minimum charge of 1 minute) (Correct)*

**Answer :**

\*per second use (with minimum charge of 1 minute) (Correct)

**Which cache type runs on a 24 hour clock?**

Options are :

- *Warehouse Cache*
- *Results Cache (Correct)*
- *Metadata Cache*

**Answer :**

\*Results Cache (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 2

**Scaling a warehouse UP will decrease the number of clusters. (e.g. Small to Medium)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Scaling in is a manual process (e.g. Max Clusters to min clusters)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Scaling a warehouse OUT will increase the number of clusters. (e.g. Min to Max)**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### **Snowflake Certification Preparation Set 8**

**Scaling a warehouse DOWN will decrease the number of servers. (e.g. Medium to Small)**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Which of the following are valid parameters when create a new virtual warehouse?**

Options are :

- Auto Suspend (Correct)
- Type of EC2
- Auto Resume (Correct)
- RAM Size
- CPU Counts
- Warehouse Size (Correct)

**Answer :**

\*Auto Suspend (Correct)

\*Auto Resume (Correct)

\*Warehouse Size (Correct)

**When a virtual warehouse is provisioned there is no benefit in stopping the virtual warehouse with in the first 60 seconds because that time period has already been billed.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

### [Snowflake SnowPro Core Certification Exam Questions Set 10](#)

**For File formats like JSON, Avro the only supported character set is UTF-8**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**User A and User B can access one another's result sets from the Results Cache, as long as which of the following are true? (Choose two)**

Options are :

- *They use the same Role. (Correct)*
- *They run with 59 minutes of one another*
- *They use the same Warehouse.*
- *They run the exact same SQL Text/Query. (Correct)*
- *They sign in to the same Session*

**Answer :**

\**They use the same Role. (Correct)*

\**They run the exact same SQL Text/Query. (Correct)*

**Warehouses can also be resized at any time, even while running, to accommodate the need for more or less compute resources**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

**In auto scale mode, if new\_max\_clusters < running clusters, then excess clusters shut down immediately**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Out of below four options, which options require a running warehouse when retrieving data from them?**

Options are :

- *Metadata Cache*
- *Results Cache*
- *Warehouse Cache (Correct)*
- *Local Disk Cache (Correct)*

**Answer :**

\**Warehouse Cache (Correct)*

\**Local Disk Cache (Correct)*

**The columns in the query history include the QueryID, the SQL Text, the Warehouse name, the Warehouse Size, the Session ID and others. Which column is a good indicator of whether a Warehouse was used (and Compute costs incurred) by a query?**

Options are :

- *Warehouse Name*

- *SQL Text*
- *Warehouse Size (Correct)*
- *QueryID*
- *Session ID*

**Answer :**

\*Warehouse Size (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 28

**What is the size limit on the variant data type in snowflake?**

Options are :

- *16 MB (uncompressed)*
- *16 MB (compressed) (Correct)*
- *10 MB (compressed)*
- *10 MB (uncompressed)*

**Answer :**

\*16 MB (compressed) (Correct)

**Whenever the warehouse size is reduced or the number of clusters are decreased, the servers are always removed in reverse order of when they were added (LIFO - Last In First Out)**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**In auto scale multi-cluster warehouse same value is specified for both maximum & minimum clusters (except for the value 1)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

## Snowflake Certification Preparation Set 21

**When a warehouse is suspended, the idle servers in that warehouse are immediately shut down, but any servers that are executing queries are allowed to run until the queries are completed. Such servers that are waiting to shut down are considered to be in which mode?**

Options are :

- *Shutting Down*
- *Suspending*
- *Quiesce (Correct)*
- *Quiet*

**Answer :**

\*Quiesce (Correct)

**How much size limit the VARIANT data type imposes on individual rows.**

Options are :

- *10 MB (compressed)*
- *16 MB (uncompressed)*
- *10 MB (uncompressed)*
- *16 MB (compressed) (Correct)*

**Answer :**

\*16 MB (compressed) (Correct)

**JSON arrays are written inside which type of brackets?**

Options are :

- *Round Brackets*
- *Curly Braces*
- *Square Brackets (Correct)*

**Answer :**

\*Square Brackets (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 5

**What are the resources that a virtual warehouse provides?**

Options are :

- *Temporary Storage (Correct)*
- *Long term data storage*
- *User & Role Management*
- *CPU (Processing Power) (Correct)*
- *Memory (Correct)*

**Answer :**

\*Temporary Storage (Correct)

\*CPU (Processing Power) (Correct)

\*Memory (Correct)

**If a multi-cluster warehouse is resized, the new size applies to all the clusters for the warehouse, including clusters that are currently running and any clusters that are started after the warehouse is resized**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**When a warehouse is resized, which queries make use of the new size?**

Options are :

- *Both current and subsequent queries*
- *Only subsequent queries (Correct)*
- *Only currently running queries*

**Answer :**

\*Only subsequent queries (Correct)

## Snowflake Certification Preparation Set 11

**Snowflake does not begin executing SQL statements submitted to a warehouse until all of the servers for the warehouse are successfully provisioned, unless any of the servers fail to provision**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**What action causes a Warehouse's cache to be purged?**

Options are :

- *Resuming the Warehouse*
- *A change to metadata Cache size*
- *Suspending the warehouse (Correct)*
- *The passing of 24 hours*

**Answer :**

\**Suspending the warehouse (Correct)*

**When the warehouse is resized (scaled up or out), then the additional servers that get added will be available for use by any queries that are queued or newly submitted**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

## Snowflake Certification Preparation Set 6

**Which of the following are true about the variant data type in Snowflake? Select all that apply.**

Options are :

- Optimized storage based on repeated elements (Correct)
- Requires a custom mapping for each record type
- Stored in a separate file format from structured data
- Can be queried using JSON path notation (Correct)

**Answer :**

\*Optimized storage based on repeated elements (Correct)  
\*Can be queried using JSON path notation (Correct)

#### How many credits per hour is consumed by a 3X large warehouse

Options are :

- 32
- 64 (Correct)
- 16
- 128

**Answer :**

\*64 (Correct)

In auto scale mode, if new\_max\_clusters > running clusters, then no changes happen until additional clusters are needed

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

#### Tackling Snowflake Certification - Practice Questions Set 2

#### What is the full form of ORC?

Options are :

- Optimized Row Columnar (Correct)
- Object Read Class
- Object Row Columnar
- Optimized Read Columnar

**Answer :**

\*Optimized Row Columnar (Correct)

**A warehouse was started and it was up and running for 62 seconds. For how many seconds will the warehouse be billed for?**

Options are :

- 70 seconds
- 120 seconds
- 62 seconds (Correct)
- 60 seconds

**Answer :**

\*62 seconds (Correct)

**The FLATTEN command will parse nested objects into separate rows. One version of the FLATTEN command uses a join and the other uses an object keyword. Select the two words that represent the options used with the FLATTEN command**

Options are :

- LATERAL (Correct)
- SCHEMATA
- TABLE (Correct)
- SIDEWAYS

**Answer :**

\*LATERAL (Correct)  
\*TABLE (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 12

**What can be the maximum credits consumed by a medium size warehouse with 3 clusters in one hour?**

Options are :

- 8
- 24
- 12 (*Correct*)
- 6

**Answer :**

\*12 (Correct)

**Snowflake utilizes per-second billing (with a 60-second minimum each time the warehouse starts)**

Options are :

- FALSE (*Correct*)
- TRUE

**Answer :**

\*FALSE (Correct)

**Snowflake utilizes per-second billing (with a 60-second minimum each time the warehouse starts)**

Options are :

- FALSE
- TRUE (*Correct*)

**Answer :**

\*TRUE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 16**

**In a VARIANT column, NULL values are stored as a string containing the word "null" not the SQL NULL value.**

Options are :

- TRUE (*Correct*)
- FALSE

**Answer :**

\*TRUE (Correct)

**Out of below four options, which options require a running warehouse when retrieving data from them?**

Options are :

- *Local Disk Cache (Correct)*
- *Warehouse Cache (Correct)*
- *Results Cache*
- *Metadata Cache*

**Answer :**

\*Local Disk Cache (Correct)

\*Warehouse Cache (Correct)

**In a multi cluster warehouse, auto-suspend and auto-resume apply only to the entire warehouse and not to the individual clusters in the warehouse**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**When the warehouse is resized (scaled up or out), then the additional servers that get added will also improve the performance of already running queries**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**If a server runs for 61 seconds, shuts down, and then restarts and runs for 30 seconds, it is billed for how many seconds?**

Options are :

- 91
- 121 (*Correct*)
- 90
- 120

**Answer :**

\*121 (Correct)

**Which scaling policy, conserves credits by favoring keeping running clusters fully-loaded rather than starting additional clusters, which may result in queries being queued and taking longer to complete?**

Options are :

- *Economy* (*Correct*)
- *Legacy*
- *Standard*
- *None*

**Answer :**

\*Economy (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 1

**What is the full form of JSON?**

Options are :

- *JavaScript Object Notation* (*Correct*)
- *Java Semi Object Notation*
- *Java Object Notation*
- *Javascript Notation*

**Answer :**

\*JavaScript Object Notation (Correct)

**When the JSON or Avro output from some software is composed of a single huge array containing multiple records, then which file format option should be enabled to load the records in multiple rows**

Options are :

- *REMOVE\_OUTER\_ARRAY*
- *STRIP\_OUTER\_ARRAY*
- *STRIP\_OUTER\_ARRAY (Correct)*
- *ELIMINATE\_OUTER\_ARRAY*

**Answer :**

\**STRIP\_OUTER\_ARRAY (Correct)*

**You are the solution architect for a large manufacturing company running a Snowflake data warehouse. Your Snowflake implementation has just gone live with a single virtual warehouse and has been opened up to all the users across the organization. One of your heaviest users are the marketing department, who have a large number of users and execute a large number of queries. Marketing users have complained that the queries take a long time to execute. What is the best immediate action should you take to improve their experience?**

Options are :

- *Have a dedicated warehouse for marketing team, and set the capacity of this warehouse as per the needs of marketing team (Correct)*
- *Limit the number of queries that marketing team can execute*
- *Study each and every query of the marketing team and optimize them*
- *Increase the size of your single warehouse to the maximum size possible*

**Answer :**

\**Have a dedicated warehouse for marketing team, and set the capacity of this warehouse as per the needs of marketing team (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 2](#)

**In auto scale mode, if `next_max_clusters > running clusters`, then no changes happen until additional clusters are needed**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**A warehouse was started and it was up and running for 40 seconds. For how many seconds will the warehouse be billed for?**

Options are :

- *No charge would be applied as it was up for less than 1 minute*
- *60 seconds (Correct)*
- *40 seconds*

**Answer :**

\*60 seconds (Correct)

**If any of the servers for the warehouse fail to provision during start-up, Snowflake attempts to repair the failed server(s). During the repair process, the warehouse starts processing SQL statements only when a minimum of how much % of the requested servers are successfully provisioned**

Options are :

- *30*
- *50 (Correct)*
- *25*
- *75*

**Answer :**

\*50 (Correct)

### [Snowflake SnowPro Core Certification Practice Test Set 2](#)

**A warehouse can be assigned to only a single resource monitor.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**What is the correct way to find out virtual warehouse credit usage information in Snowflake?**

Options are :

- Go to Account --> Billing & Usage in Web UI (Correct)
- Query Warehouse\_Credits table in ACCOUNT\_USAGE Schema
- Query WAREHOUSE\_METERING\_HISTORY in INFORMATION\_schema (Correct)
- Query WAREHOUSE\_METERING\_HISTORY in account usage schema (Correct)
- Query Warehouse\_Credits table in INFORMATION\_SCHEMA

**Answer :**

\*Go to Account --> Billing & Usage in Web UI (Correct)  
\*Query WAREHOUSE\_METERING\_HISTORY in INFORMATION\_schema (Correct)  
\*Query WAREHOUSE\_METERING\_HISTORY in account usage schema (Correct)

**The VARIANT data type imposes, how much size limit on individual rows**

Options are :

- 64 MB Compressed
- 16 MB Uncompressed
- 64 MB Uncompressed
- 16 MB Compressed (Correct)

**Answer :**

\*16 MB Compressed (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 19**

**Scaling a warehouse OUT will increase the number of servers. (e.g. Min to Max)**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Which of the following are options when creating a Virtual Warehouse?**

Options are :

- *Storage Size*
- *Auto Resume (Correct)*
- *Auto Suspend (Correct)*
- *Cache Size*

**Answer :**

\*Auto Resume (Correct)  
\*Auto Suspend (Correct)

**How many credits will be consumed if a Medium-size warehouse (4 servers per cluster) with 3 clusters runs in Maximized mode for 2 hours**

Options are :

- 18
- 24 (*Correct*)
- 36
- 12

**Answer :**

\*24 (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 10**

**Scaling up of a warehouse can be done only in Enterprise edition and higher**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**How many servers are there in 2X large warehouse?**

Options are :

- 32 (*Correct*)
- 64
- 16
- 128

**Answer :**

\*32 (Correct)

**True or false: A table in Snowflake can only be queried using the Virtual Warehouse used to load the data.**

Options are :

- *FALSE* (*Correct*)
- *TRUE*

**Answer :**

\**FALSE* (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 10

**What is the recommended approached for making a variant column accessible in a BI tool?**

Options are :

- *A pre-defined mapping*
- *A View* (*Correct*)
- *There is no way to do this*
- *Leveraging a JSON parser*

**Answer :**

\**A View* (Correct)

**How many credits will be consumed if a Small-size warehouse (2 servers per cluster) with 3 clusters runs in Maximized mode for 2 hours**

Options are :

- 12 (*Correct*)
- 36

- 18

- 24

**Answer :**

\*12 (Correct)

**Scaling up is a manual process. (e.g. Small to Medium)**

Options are :

- FALSE

- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### Snowflake Certification Preparation Set 23

**What happens when a suspended warehouse is resized (e.g. resized from small to medium)?**

Options are :

- When size of a suspended warehouse is increased it starts the warehouse, and then suspends it back after provisioning additional servers

- A suspended warehouse cannot be resized

- Additional servers are provisioned only when the warehouse is next resumed (Correct)

- Additional Servers are added to the warehouse immediately even in suspended state

**Answer :**

\*Additional servers are provisioned only when the warehouse is next resumed (Correct)

**Snowflake allows for custom column separators (a.k.a. column delimiters). In order to use a caret symbol as a custom column separator, what option must you first choose?**

Options are :

- Other (Correct)

- Comma

- None

- *Tab*
- *Vertical Bar*

**Answer :**

\*Other (Correct)

Many logical data models refer to entities, attributes and values. In JSON, and some other semi-structured types, the word "attribute" is often replaced with what word?

Options are :

- *Values*
- *Columns*
- *Relationships*
- *Rows*
- *Keys (Correct)*

**Answer :**

\*Keys (Correct)

### Snowflake Snowpro Certification Practice Test Set 5

Scaling a warehouse IN will decrease the number of clusters. (e.g. Max to Min)

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

Which cache type gets purged regularly?

Options are :

- *Warehouse Cache*
- *Results Cache (Correct)*
- *Metadata Cache*

**Answer :**

\*Results Cache (Correct)

**For a single snowflake session, more than one virtual warehouse can be specified at a time for executing a query.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

### Snowflake Certification Preparation Set 15

**You have a virtual warehouse named my\_warehouse which is sized as Medium. You need to resize it to X-LARGE for future queries. What is the best way to make this change.**

Options are :

- *Drop and recreate the same warehouse with X-LARGE size*
- *Alter the existing warehouse either by using an ALTER command or through the Snowflake Web UI to change the size (Correct)*
- *It is not permitted to resize the warehouse once its created*

**Answer :**

\**Alter the existing warehouse either by using an ALTER command or through the Snowflake Web UI to change the size (Correct)*

**What is the data type provided by Snowflake to store semi structure data?**

Options are :

- *Varchar*
- *Text*
- *Variant (Correct)*
- *Semi-structure*

**Answer :**

\**Variant (Correct)*

**You can increase or decrease the size of a virtual warehouse as your user base & query demand increase or decrease?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### [Snowflake SnowPro Core Certification Exam Questions Set 11](#)

**Which of the following currently uses user-specified warehouse, but eventually run in a serverless environment using Snowflake-supplied compute resources**

Options are :

- *Streams*
- *Stored Procedure*
- *Tasks (Correct)*
- *Snowpipe*

**Answer :**

\**Tasks (Correct)*

**Which file format option can be enabled in JSON or AVRO file to remove the outer array structure and load the records into separate table rows**

Options are :

- *REMOVE\_OUTER\_ARRAY*
- *READ\_INTERNAL\_ARRAYS*
- *STRIP\_OUTER\_ARRAY (Correct)*
- *ELIMINATE\_OUTER\_ARRAY*

**Answer :**

\**STRIP\_OUTER\_ARRAY (Correct)*

**Cloud services also use compute. Does snowflake charge you for cloud services.**

Options are :

- *Cloud services are charged irrespective of the compute resource usage*
- *Cloud services are charged only if it exceeds 20% of the daily usage of the compute resources*
- *Snowflake never charges for cloud services*
- *Cloud services are charged only if it exceeds 10% of the daily usage of the compute resources (Correct)*

**Answer :**

\*Cloud services are charged only if it exceeds 10% of the daily usage of the compute resources (Correct)

### [Snowflake Certification Preparation Set 17](#)

**Snowflake Credit Charges for a virtual warehouse is calculated based on:**

Options are :

- *The number of clusters (if using multi-cluster warehouses). (Correct)*
- *Number of queries getting executed*
- *The length of time each server in each cluster runs. (Correct)*
- *The number of servers per cluster (determined by warehouse size). (Correct)*
- *Number of Users running the queries concurrently*

**Answer :**

\*The number of clusters (if using multi-cluster warehouses). (Correct)

\*The length of time each server in each cluster runs. (Correct)

\*The number of servers per cluster (determined by warehouse size). (Correct)

**Which semi structure data types are available in Snowflake**

Options are :

- *VARIANT (Correct)*
- *ARRAY (Correct)*
- *SEMI STRUCTURE*
- *OBJECT (Correct)*

- *STRING*

**Answer :**

- \*VARIANT (Correct)
- \*ARRAY (Correct)
- \*OBJECT (Correct)

**In auto scale mode, if new\_min\_clusters < running clusters, then excess clusters shut down when they finish executing statements and the scaling policy conditions are met**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 29

**True or false: You can point Snowflake at any S3 bucket to directly query the files in that bucket as long as the files are in Parquet or ORC format.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

- \*FALSE (Correct)

**Default Size of the warehouse created using CREATE WAREHOUSE command is**

Options are :

- *X-SMALL (Correct)*
- *X-LARGE*
- *MEDIUM*
- *SMALL*

**Answer :**

- \*X-SMALL (Correct)

**Storage can increase or decrease without any effect on virtual warehouse sizes.**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 10](#)

**Two Virtual Warehouses can access the same data at the same time without causing contention issues.**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**In auto scale mode, if new\_min\_clusters > running clusters, then additional clusters immediately started to meet the minimum.**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Which file format is an open-source data serialization and RPC framework originally developed for use with Apache Hadoop**

Options are :

- XML
- AVRO (Correct)
- JSON

- *PARQUET*

- *ORC*

**Answer :**

\*AVRO (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 11

**True or false:** The warehouse cache may be reset if a running warehouse is suspended and then resumed.

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Compute can be scaled up, down, out, or in and there is no effect on storage used.**

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Non-native values such as dates and timestamps are stored as strings when loaded into a VARIANT column**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

**The scaling policy for a multi-cluster warehouse only applies (make sense) if it is running in Auto-scale mode and not in Maximized Mode**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**It is not possible to suspend a specific cluster in a multi-cluster virtual warehouse but rather the whole virtual warehouse is suspended.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Snowflake compute costs depend on which of the following?**

Options are :

- *The sizes of running warehouses. (Correct)*
- *The number of rows returned in queries.*
- *The amount of time warehouses have run. (Correct)*
- *The total number of warehouses in the account*

**Answer :**

\**The sizes of running warehouses. (Correct)*

\**The amount of time warehouses have run. (Correct)*

## [Snowflake Certification Preparation Set 7](#)

**Snowflake data warehouse is built on top of Hadoop Big Data Software Platform**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which scaling policy, prevents/minimizes queuing by favoring starting additional clusters over conserving credits?**

Options are :

- *Legacy*
- *None*
- *Standard (Correct)*
- *Economy*

**Answer :**

\*Standard (Correct)

**In a multi cluster warehouse, auto-suspending a running warehouse due to inactivity, does not apply to individual clusters, but rather the entire warehouse.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 8

**Scaling out is an automated process (e.g. Min clusters to max clusters)**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**If the null values in JSON documents indicate missing values and have no other special meaning, Snowflake recommend setting which file format option to TRUE for the COPY INTO <table> command when loading the JSON files.**

Options are :

- *ELIMINATE\_NULL\_VALUES*
- *REMOVE\_NULL\_VALUES*
- *STRIP\_NULL\_VALUES (Correct)*
- *RETAIN\_NULL\_VALUES*

**Answer :**

\**STRIP\_NULL\_VALUES* (Correct)

**To see files that have been uploaded to a Snowflake stage, which command should be used**

Options are :

- *STAGE\_VIEW*
- *VIEW*
- *LIST (Correct)*
- *STAGE\_LIST*

**Answer :**

\**LIST* (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 38**

**Multi-factor Authentication can be used for connecting to Snowflake via the Snowflake JDBC driver**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE* (Correct)

**MFA can be used for connecting to Snowflake via the Snowflake JDBC driver**

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Which type of stage can store files that are staged and managed by one or more users and loaded into one or more tables**

Options are :

- *User*
- *Table*
- *Named (Correct)*

**Answer :**

\**Named (Correct)*

### Snowflake Certification Preparation Set 13

**What is the maximum number of child tasks, a given task can have?**

Options are :

- *10*
- *1000*
- *100 (Correct)*
- *10000*
- *Unlimited*

**Answer :**

\**100 (Correct)*

**Recreating the pipe will not remove the history of the files that are loaded using that pipe**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which of the following stages cannot be altered or dropped?**

Options are :

- *User (Correct)*
- *Table (Correct)*
- *Named Internal Stage*
- *Named External Stage*

**Answer :**

\*User (Correct)  
\*Table (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 43**

**What is the default value for MAX\_FILE\_SIZE parameter while unloading a file?**

Options are :

- *16 MB (Correct)*
- *32 MB*
- *25 MB*
- *10 MB*

**Answer :**

\*16 MB (Correct)

**Which of the following is the correct syntax for the REPLACE FUNCTION?**

Options are :

- *REPLACE( [,] )*
- *REPLACE( [,] ) (Correct)*
- *REPLACE( [,] )*

**Answer :**

\**REPLACE( [,] ) (Correct)*

**Snowflake encrypts all customer data by default, using the latest security standards, and charges customer for this encryption**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

#### SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

**Which parameter should be set to true if the requirement is to load the files whose load metadata has expired**

Options are :

- *LOAD\_ALL\_FILES*
- *LOAD\_EXPIRED\_FILES*
- *LOAD\_UNCERTAIN\_FILES (Correct)*

**Answer :**

\**LOAD\_UNCERTAIN\_FILES (Correct)*

**When the files is uploaded to snowflake internal stage, files are automatically encrypted by the client on the local machine prior to being transmitted to the internal stage.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**In Tri-Secret Secure option customer manages the master key**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

## Snowflake Certification Preparation Set 6

**Which of the transformations are supported while loading data into table using the COPY command**

Options are :

- *Column Reordering (Correct)*
- *Column Omission (Correct)*
- *Casts (Correct)*
- *Aggregations*
- *Truncating text strings that exceed the target column length (Correct)*

**Answer :**

- \*Column Reordering (Correct)
- \*Column Omission (Correct)
- \*Casts (Correct)
- \*Truncating text strings that exceed the target column length (Correct)

**Within a transaction multiple SQL statements will see the same set of records in a stream**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

**Only named stages (internal or external) and user stages are supported for COPY transformations**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

[SnowPro Core Certification 3 Full Practice Exams 2021 Set 23](#)

**Which of the following file format option is not supported in "COPY INTO TABLE (Copy Load)" command, for CSV type**

Options are :

- *FILE\_EXTENSION* (*Correct*)
- *COMPRESSION*
- *SKIP\_HEADER*
- *VALIDATE\_UTF8*
- *FIELD\_DELIMITER*
- *SKIP\_BLANK\_LINES*

**Answer :**

\**FILE\_EXTENSION* (*Correct*)

**Which file format option should be enabled while loading a JSON file to remove the outer most array structure and load the records in separate table rows**

Options are :

- *STRIP\_OUTER\_ARRAY* (*Correct*)
- *REMOVE\_OUTER\_ARRAY*
- *READ\_INNER\_ARRAY*
- *ELIMINATE\_OUTER\_ARRAY*

**Answer :**

\**STRIP\_OUTER\_ARRAY* (*Correct*)

**What is the file encoding of the output files generated by Snowflake when unloading data**

Options are :

- *UCS2*
- *ANSI*
- *UTF-8* (*Correct*)
- *UTF-16*

**Answer :**

\*UTF-8 (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 25

**User stages are referenced using?**

Options are :

- @~" (Correct)
- @%
- @
- %~

**Answer :**

\*@~" (Correct)

**When the owner role of a given task (i.e. the role with the OWNERSHIP privilege on the task) is deleted, what happens?**

Options are :

- *The task is re-possessed by the role that dropped the owner role (Correct)*
- *The roles granted to any task cannot be dropped unless all associated tasks are dropped*
- *The tasks is suspended immediately, and the task will not have any owner till explicitly assigned*
- *The tasks is also dropped*

**Answer :**

\*The task is re-possessed by the role that dropped the owner role (Correct)

**Which option should be set to load all files in copy command, ignoring the load metadata if it exists**

Options are :

- *FORCE (Correct)*
- *LOAD\_ALL\_FILES*
- *OVERWRITE*

- *IGNORE\_METADATA*

**Answer :**

\*FORCE (Correct)

### Snowflake Certification Preparation Set 15

**A file staged in a user stage can be loaded in only one table**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**A single tasks in snowflake can execute multiple SQL Statements**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**What can be typed in front of a line of text in the SQL pane to make it into a comment or note rather than have it treated as code?**

Options are :

- *A double dash: -- (Correct)*
- *A double slash: // (Correct)*
- *A dash, slash and dash pattern: -/-*
- *A slash, dash and slash pattern: /-/*

**Answer :**

\*A double dash: -- (Correct)

\*A double slash: // (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 19

If we chose to use DOUBLE QUOTES as the enclosing character in a file, and in our Snowflake file format definition, which of the following values would cause an error during loading?

Options are :

- *Snowflake's Query Results*
- *Snowflake "Virtual" Warehouse (Correct)*
- *Apple, Mango, & Cherry*

**Answer :**

\*Snowflake "Virtual" Warehouse (Correct)

**Multi-cluster warehouses are beneficial for improving the performance of slow-running queries or data loading**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**For Downloading a file that is staged in Snowflake Internal Stage which command is used**

Options are :

- *PUT*
- *GET (Correct)*
- *RETRIEVE*
- *DOWNLOAD*

**Answer :**

\**GET (Correct)*

### [Snowflake SnowPro Core Certification Practice Test Set 1](#)

**Named stages come in two varieties, what are they?**

Options are :

- *Internal (Correct)*

- *Materialized*

- *Secure*

- *External (Correct)*

- *Permanent*

**Answer :**

\*Internal (Correct)

\*External (Correct)

**The following factors affect data load rates: (select all that apply)**

Options are :

- *Gzip compression efficiency (Correct)*

- *Thread size*

- *RAM on the Virtual Warehouse*

- *Physical location of the Stage (Correct)*

**Answer :**

\*Gzip compression efficiency (Correct)

\*Physical location of the Stage (Correct)

**What ensures that all customer data, new and old, is encrypted with the latest security technology**

Options are :

- *This has to be done by contacting snowflake support*

- *Key Rotation*

- *Periodic Rekeying (Correct)*

- *None*

**Answer :**

\*Periodic Rekeying (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 26**

**You are using a XL warehouse which has 16 nodes and 4 cores in each node. When you submit a query to this warehouse, snowflake maximum split of the query can be into how many partitions? (Provided that the query is eligible for maximum splitting)**

Options are :

- 16
- 4
- 128
- 64 (*Correct*)

**Answer :**

\*64 (Correct)

**copy into mytable file\_format = (type = csv field\_delimiter = '|' skip\_header = 1); This command will load file from which stage**

Options are :

- *User Stage*
- *Named Stage*
- *Table Stage (Correct)*
- *Error out as from clause is not specified*

**Answer :**

\*Table Stage (Correct)

**When a task is created, it is in resumed state by default?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**SnowFlake SnowPro Advanced Architect Practice Exam Set 7**

**Which statements about Data Integration Tech Partners are true?**

Options are :

- *Snowflake can carry out transformations after loading files staged by partner software (ELT). (Correct)*
- *Data Integration Tech Partner software can be used to carry out transformations. (Correct)*
- *Data Integration Tech Partner software should be used to deliver data to stages, Snowflake is then used to load the data (Correct)*
- *Data Integration Tech Partner software can be used to extract data from other systems. (Correct)*
- *Snowflake can be used to extract data from other databases but Data Integration Tech Partner software must be used to do transformations.*
- *Snowflake must be used to extract data from other databases but Data Integration Tech Partner software can load data*

**Answer :**

\*Snowflake can carry out transformations after loading files staged by partner software (ELT). (Correct)

\*Data Integration Tech Partner software can be used to carry out transformations. (Correct)

\*Data Integration Tech Partner software should be used to deliver data to stages, Snowflake is then used to load the data (Correct)

\*Data Integration Tech Partner software can be used to extract data from other systems. (Correct)

**Snowflake Tasks executions are associated with the user who created the task, i.e. in the history database we can see that the task is getting executed under the username of the User who created that task**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**How is the query load calculated, which is displayed on the "Warehouse Load over Time" chart in Snowflake Web UI**

Options are :

- Number of queries running in the given interval divided by the number of seconds in the interval
- Execution time (in seconds) of all queries in an interval by the total time (in seconds) for the interval (Correct)
- The formula is not known to the users, and its internally computed by snowflake only
- $(\text{Number of queries running} + \text{Number of queries in queue})$  in the given interval divided by the number of seconds in the interval

**Answer :**

\*Execution time (in seconds) of all queries in an interval by the total time (in seconds) for the interval (Correct)

**Which role is most powerful role in a Snowflake system?**

Options are :

- ACCOUNTADMIN (Correct)
- PUBLIC
- SYSADMIN
- SECURITYADMIN

**Answer :**

\*ACCOUNTADMIN (Correct)

**Snowflake is HIPAA compliant?**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the blocked IP address list first**

Options are :

- FALSE

- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 1

**A user can upload one or more data files to an external stage. If the stage is a customer-managed container in a cloud storage service (like S3 in AWS), the user may or may not encrypt the data files using client-side encryption**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Federated authentication in snowflake is compliant with SAML 2.0**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Which is the default encoding considered in COPY command while loading delimited files like CSV, TSV?**

Options are :

- *ANSI*
- *UCS-2*
- *UTF-16*
- *UTF-8 (Correct)*

**Answer :**

\*UTF-8 (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 20

**A file staged in a user stage can be managed by multiple users**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**Which of the following techniques can be used in Snowflake to improve query performance?**

Options are :

- *Cluster Keys/Reclustering (Correct)*
- *Query Hints*
- *Distribution keys*
- *Index definition*

**Answer :**

\**Cluster Keys/Reclustering (Correct)*

**What are the three Snowflake Stage types?**

Options are :

- *Materialized*
- *Table (Correct)*
- *Permanent*
- *Named (Correct)*
- *User (Correct)*
- *Secure*

**Answer :**

\**Table (Correct)*

\**Named (Correct)*

\**User (Correct)*

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 30

**Multifactor Authentication can be enabled for which of the following? Select all that Apply?**

Options are :

- *Snowflake WebUI (Correct)*
- *Snowpipe*
- *ODBC (Correct)*
- *Python (Correct)*

**Answer :**

\*Snowflake WebUI (Correct)

\*ODBC (Correct)

\*Python (Correct)

**Bulk loading relies on the user-provided virtual warehouses**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Table stages do support transforming data while loading it (i.e. using a query as the source for the COPY command).**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 17

**Which all transformation options are supported while loading data into a table using COPY command**

Options are :

- *CAST (Correct)*
- *Aggregations*
- *Column Reordering (Correct)*
- *Column Omission (Correct)*
- *Truncating text strings (Correct)*
- *Filter*

**Answer :**

- \**CAST (Correct)*
- \**Column Reordering (Correct)*
- \**Column Omission (Correct)*
- \**Truncating text strings (Correct)*

**Snowpipe uses compute resources provided by user provided Virtual Warehouses**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**Which option should be set in COPY command, to delete/remove the staged files from Snowflake Stage after the file is successfully loaded**

Options are :

- *DELETE*
- *REMOVE*
- *PURGE (Correct)*

**Answer :**

- \**PURGE (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 23**

**When unloading data into multiple files, which copy option is used to specify the maximum size of each file created.**

Options are :

- *MAX\_SIZE*
- *MAX\_FILE\_SIZE (Correct)*
- *MAX\_NBR\_ROWS*
- *MAX\_FILE\_SIZE\_MB*

**Answer :**

\**MAX\_FILE\_SIZE (Correct)*

**To configure an External Stage you need to define the stage object in Snowflake. You also need two additional components. What are they?**

Options are :

- *Cloud access credentials (e.g. IAM User and Policy) (Correct)*
- *Cloud storage location (e.g. S3 Bucket) (Correct)*
- *Stage Admin Password*
- *Stage Admin Username*
- *Snowflake External Stage Password Authentication Certificate*

**Answer :**

\**Cloud access credentials (e.g. IAM User and Policy) (Correct)*

\**Cloud storage location (e.g. S3 Bucket) (Correct)*

**The change data capture (CDC) records returned when querying a stream rely on a combination of**

-

Options are :

- *Offset stored in Table*
- *Offset Stored in Stream (Correct)*
- *Change Tracking Metadata Stored in Stream*
- *The change tracking metadata stored in the table (Correct)*

**Answer :**

- \*Offset Stored in Stream (Correct)
- \*The change tracking metadata stored in the table (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**Which transformations are available when using the COPY INTO command to load data files into Snowflake from a stage? (select all that apply)**

Options are :

- *Column data type conversion (Correct)*
- *Filters*
- *Column concatenation (Correct)*
- *Aggregates*

**Answer :**

- \*Column data type conversion (Correct)
- \*Column concatenation (Correct)

**True or false: The Query Profile view is only available for completed queries.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**What is the recommended file size to optimize the number of parallel operations for a load?**

Options are :

- *100 MB to 250 MB (compressed) (Correct)*
- *1 MB to 10 MB (compressed)*
- *10 MB to 100 MB (uncompressed)*
- *10 MB to 100 MB (compressed)*

**Answer :**

- \**100 MB to 250 MB (compressed) (Correct)*

## Snowflake SnowPro Core Certification Practice Tests Set 14

You are viewing the **Query History** table in the **History** area of the WebUI. You want to see if a query pulled data from long-term centralized storage. Where will you look and what will you look for?

Options are :

- Look in the Bytes Scanned column for the acronym LTCS
- Look in the Rows column for the acronym LTCS
- Look in the Rows column for a red bar
- Look in the Bytes Scanned column for a green bar (Correct)

**Answer :**

\*Look in the Bytes Scanned column for a green bar (Correct)

**Tri-secret requires that customers manage their own keys**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Which file format can be used for both loading and unloading data from snowflake**

Options are :

- CSV (Correct)
- AVRO
- XML
- ORC
- PARQUET (Correct)
- JSON (Correct)

**Answer :**

\*CSV (Correct)

\*PARQUET (Correct)

\*JSON (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 5

**Which of the following options would result in a column named MY\_COL being sandwiched between two percent signs? (Select all that apply)**

Options are :

- `%'||MY_COL||'%'` (Correct)
- `MERGESTRING('%',MY_COL,'%')`
- `CONCATN('%',MY_COL,'%')`
- `CONCAT('%',CONCAT(MY_COL,'%'))` (Correct)

**Answer :**

\*`%'||MY_COL||'%'` (Correct)

\*`CONCAT('%',CONCAT(MY_COL,'%'))` (Correct)

**PIPE\_EXECUTION\_PAUSED parameter that is used to pause the execution of pipe can be set at which all levels?**

Options are :

- *Role*
- *Pipe* (Correct)
- *Database*
- *Account* (Correct)
- *Schema* (Correct)

**Answer :**

\*Pipe (Correct)

\*Account (Correct)

\*Schema (Correct)

**Snowpipe uses compute resources provided by Snowflake (i.e. serverless compute model)**

Options are :

- *TRUE* (Correct)

- FALSE

**Answer :**

\*TRUE (Correct)

### Snowflake Certification Preparation Set 23

**A task object defines a recurring schedule for executing a SQL statement, including statements that call stored procedures**

Options are :

- TRUE (Correct)

- FALSE

**Answer :**

\*TRUE (Correct)

**When a copy command cannot definitively determine whether a file has been loaded already, then what is the default behaviour?**

Options are :

- Error message is generated, that load metadata for the file for given table doesn't exist
- The file is loaded
- The file load is skipped (Correct)

**Answer :**

\*The file load is skipped (Correct)

**For Data unloading which command is used**

Options are :

- UNLOAD DATA FROM

- COPY INTO

- UNLOAD DATA TO

- *COPY INTO (Correct)*

**Answer :**

- \**COPY INTO (Correct)*

### **Snowflake Certification Preparation Set 19**

***What are the different types of streams available?***

*Options are :*

- *Append-Only (Correct)*
- *Standard (Correct)*
- *Insert-Only (Correct)*
- *Update-Only*
- *Native*

**Answer :**

- \**Append-Only (Correct)*

- \**Standard (Correct)*

- \**Insert-Only (Correct)*

***Which of the below two stages are automatically created and need not be configured explicitly by the user***

*Options are :*

- *Internal Stage*
- *Role Stage*
- *S3 Stage*
- *External Stage*
- *Table Stage (Correct)*
- *User Stage (Correct)*

**Answer :**

- \**Table Stage (Correct)*

- \**User Stage (Correct)*

***The COPY INTO <location> command provides a copy option (SINGLE) for unloading data into a single file or multiple files. The default is?***

*Options are :*

- *SINGLE = FALSE (i.e. data extracted in multiple files) (Correct)*
- *SINGLE = TRUE (i.e. data extracted in single file)*

**Answer :**

*\*SINGLE = FALSE (i.e. data extracted in multiple files) (Correct)*

### **Snowflake SnowPro Core Certification Exam Questions Set 7**

***What is the recommended size for data files to optimize the number of parallel operations for a load***

*Options are :*

- *100MB to 250 MB compressed (Correct)*
- *10 MB to 100 MB uncompressed*
- *100MB to 250 MB uncompressed*
- *10 MB to 100 MB compressed*

**Answer :**

*\*100MB to 250 MB compressed (Correct)*

***As an ETL developer you are developing jobs to load data into a snowflake table. You have an S3 stage defined which contains a single file containing 1000 rows. You managed to load those 1000 rows using the copy command of snowflake, but while performing further executions of those copy command, you notice that your COPY command is executing successfully but is loading zero rows into the target table. What could be the reason?***

*Options are :*

- *The file might have got corrupted*
- *Snowflake Tracks the file that have been loaded and will not load the same file again (Correct)*
- *The file that was loaded on first run got deleted, as with every copy command the associated files are removed from S3*

**Answer :**

*\*Snowflake Tracks the file that have been loaded and will not load the same file again (Correct)*

***Which of the following options for loading data files from a stage using COPY command the fastest?***

*Options are :*

- *By path (internal stages) / prefix (Amazon S3 bucket)*
- *Using pattern matching to identify specific files by pattern*
- *Specifying a list of specific files to load. (Correct)*
- *There is no difference in the speed of loading for any option*

**Answer :**

*\*Specifying a list of specific files to load. (Correct)*

### **Snowflake Certification Preparation Set 3**

***UNDROP command is applicable to which all objects***

*Options are :*

- *USERS*
- *DATABASES (Correct)*
- *TABLES (Correct)*
- *ROLES*
- *SCHEMA (Correct)*

**Answer :**

*\*DATABASES (Correct)*

*\*TABLES (Correct)*

*\*SCHEMA (Correct)*

***Which of the following is true regarding Roles in Snowflake?***

*Options are :*

- *Privileges can be assigned to role and also directly to individual users*
- *Maximum number of roles that can be assigned to a user is 5*
- *A role can be assigned to another role (Correct)*
- *A user can be assigned multiple roles (Correct)*

**Answer :**

*\*A role can be assigned to another role (Correct)*

*\*A user can be assigned multiple roles (Correct)*

**Snowflake can load from data staged in which of the following?**

*Options are :*

- *Azure Blob Storage (Correct)*
- *Google Cloud Storage (Correct)*
- *Snowflake Internal Stage (Correct)*
- *AWS S3 (Correct)*
- *Oracle Cloud Storage*
- *VMWare Storage*

**Answer :**

- \**Azure Blob Storage (Correct)*
- \**Google Cloud Storage (Correct)*
- \**Snowflake Internal Stage (Correct)*
- \**AWS S3 (Correct)*

#### **Snowflake Practice Questions - SnowPro Core Certified Set 2**

**A CAST command will force a value to be output as a certain datatype. Which of the following symbols will result in the "castcolumn" being output using the VARCHAR datatype?**

*Options are :*

- *SELECT castcolumn as VARCHAR*
- *SELECT castcolumn::VARCHAR (Correct)*
- *SELECT castcolumn//VARCHAR*
- *SELECT VARCHAR(castcolumn)*

**Answer :**

- \**SELECT castcolumn::VARCHAR (Correct)*

**If multiple DML statements change a row, the stream records all the action taken on that row.**

*Options are :*

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**An individual task in the tree can have multiple predecessor task**

*Options are :*

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

#### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 14**

**Which types of stages are automatically available in Snowflake and do not need to be created or configured?**

*Options are :*

- Table (Correct)
- Named External
- Named Internal
- User (Correct)

**Answer :**

\*Table (Correct)

\*User (Correct)

**True of False: Snowflake Stages can be defined as either External or Internal.**

*Options are :*

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Snowflake supports which of the following file formats for data loading? Select all that apply.**

*Options are :*

- *PDF*
- *ORC (Correct)*
- *Parquet (Correct)*
- *CSV (Correct)*

**Answer :**

- \**ORC (Correct)*
- \**Parquet (Correct)*
- \**CSV (Correct)*

**All data files unloaded to Snowflake internal locations are automatically encrypted using 128-bit keys. 256-bit keys can be enabled (for stronger encryption); however, additional configuration is required.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Which Snowflake Objects are currently not supported for replication?**

Options are :

- *Roles (Correct)*
- *Users (Correct)*
- *Shares (Correct)*
- *Databases*
- *Resource Monitors (Correct)*
- *Warehouses (Correct)*

**Answer :**

- \**Roles (Correct)*
- \**Users (Correct)*
- \**Shares (Correct)*
- \**Resource Monitors (Correct)*
- \**Warehouses (Correct)*

**Which command is used to download the data from Snowflake Stage**

Options are :

- *DOWNLOAD*
- *PUT*
- *GET (Correct)*
- *RETRIEVE*

**Answer :**

\**GET (Correct)*

### Tackling Snowflake Certification - Practice Questions Set 5

**Mark the types of Stages available in Snowflake**

Options are :

- *S3 Stage*
- *External Stage (Correct)*
- *Role Stage*
- *Internal Stage (Correct)*
- *User Stage (Correct)*
- *Table Stage (Correct)*

**Answer :**

\**External Stage (Correct)*  
\**Internal Stage (Correct)*  
\**User Stage (Correct)*  
\**Table Stage (Correct)*

**If the role that a running task is executing under is dropped while the task is running, the task completes processing under the dropped role.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Please look at the below two queries: Option 1: SELECT \* FROM TABLE\_NAME WHERE YEAR = 2015 AND UPPER(COL\_NAME) = 'VALUE1'; Option 2: SELECT \* FROM TABLE\_NAME WHERE YEAR = 2015 AND COL\_NAME = 'VALUE1'; Which one of these will scan less partitions? Note: COL\_NAME stores values both in upper and lower case**

Options are :

- *Option 2 (Correct)*
- *Scanning of partition is random, so no concrete logic can be applied on scanning logic*
- *Option 1*
- *Both will scan same number of partitions*

**Answer :**

\*Option 2 (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 5

**Currently, streams cannot track changes in materialized views.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**The number of load operations that run in parallel cannot exceed the number of data files to be loaded**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**What is the maximum number of file names that can be specified in the file parameter of COPY command to load into a table**

Options are :

- *1000 (Correct)*
- *10000*
- *100*
- *Unlimited*

**Answer :**

\**1000 (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 21

**What is the recommended file size, to optimize the number of parallel operation for a load**

Options are :

- *50 MB to 150 MB (Compressed)*
- *1 MB to 10 MB (Compressed)*
- *None of the above*
- *100-250 MB (Compressed) (Correct)*

**Answer :**

\**100-250 MB (Compressed) (Correct)*

**True or False: When data is staged to an internal stage using PUT command, the data is encrypted on client machine**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**Which of the below is role is automatically granted to every user and role in the account**

Options are :

- *Sysadmin*

- *Public (Correct)*
- *Developer*
- *Useradmin*

**Answer :**

\*Public (Correct)

#### **Snowflake SnowPro Core Certification Practice Test Set 4**

**Snowflake account hosted on AWS cannot load data from a file that is staged on GCP or Azure**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**Choose the answer which true about key rotation in snowflake**

Options are :

- *Account & Table master keys are automatically rotated by Snowflake when they are more than 30 days old (Correct)*
- *Only table master key is rotated by Snowflake*
- *Key rotation is an expensive operation, hence its done only on need basis*
- *Only account master key is rotated by snowflake*

**Answer :**

\*Account & Table master keys are automatically rotated by Snowflake when they are more than 30 days old (Correct)

**Which parameter is available in COPY command to load files using a regular expression?**

Options are :

- *LIKE*
- *PATTERN (Correct)*
- *REGEXP*

- REGEX

**Answer :**

\*PATTERN (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**On the WebUI which button on the top bar should you select to see the query history?**

Options are :

- Databases
- Warehouses
- Shares
- History (Correct)
- Account

**Answer :**

\*History (Correct)

**Which of the following options for loading data files from a stage using COPY command the slowest?**

Options are :

- There is no difference in the speed of loading for any option
- Specifying a list of specific files to load.
- By path (internal stages) / prefix (Amazon S3 bucket)
- Using pattern matching to identify specific files by pattern (Correct)

**Answer :**

\*Using pattern matching to identify specific files by pattern (Correct)

**Snowflake's hierarchical key model consists of which four levels of keys:**

Options are :

- Table Master Keys (Correct)
- Role Keys

- *User Keys*
- *Root Key (Correct)*
- *Account Master Keys (Correct)*
- *File Keys (Correct)*

**Answer :**

- \*Table Master Keys (Correct)
- \*Root Key (Correct)
- \*Account Master Keys (Correct)
- \*File Keys (Correct)

**Snowflake SnowPro Core Certification Practice Test Set 1**

**Which all permissions snowflake requires on a S3 bucket and folder to be able to access files in the folder (and subfolders)**

Options are :

- *s3:GetObjectVersion (Correct)*
- *s3>ListBucket (Correct)*
- *s3:ReadObject*
- *s3:GetObject (Correct)*

**Answer :**

- \**s3:GetObjectVersion (Correct)*
- \**s3>ListBucket (Correct)*
- \**s3:GetObject (Correct)*

**Which file formats are supported in the COPY command?**

Options are :

- *Avro (Correct)*
- *Unstructured*
- *JSON (Correct)*
- *CSV (Correct)*
- *TSV (Correct)*

**Answer :**

- \*Avro (Correct)
- \*JSON (Correct)
- \*CSV (Correct)
- \*TSV (Correct)

**What is the limit of number of tasks that can be included in a tree of tasks?**

Options are :

- 1000 (Correct)
- 100
- Unlimited
- 10000

**Answer :**

- \*1000 (Correct)

### Snowflake Certification Preparation Set 11

**Which security feature is supported in Snowflake?**

Options are :

- Multi-Factor Authentication
- All of the above (Correct)
- Tri-Secret Secure Encryption
- Role-Based Access Control

**Answer :**

- \*All of the above (Correct)

**The following performance optimizing query methods are supported by Snowflake:**

Options are :

- B-tree type indexes
- Caching techniques (Correct)
- Retrieving results of previous query from cache (Correct)

**Answer :**

- \*Caching techniques (Correct)
- \*Retrieving results of previous query from cache (Correct)

**What is the maximum single file size that is supported for Amazon S3, while unloading the data?**

Options are :

- 2GB
- 1GB
- 10GB
- 5GB (Correct)

**Answer :**

- \*5GB (Correct)

### Snowflake Certification Preparation Set 19

**Which of the following statements is true about data loading? Select one.**

Options are :

- Files in the 10-100MB range tend to land in the 'sweet spot' for load parallelism (Correct)
- Resizing the Virtual Warehouse from X-Small to Small will process a single file twice as fast
- Once loaded, there is no option to force a reload of an already loaded file
- The ideal file size for loading is 16MB to match micro-partition size

**Answer :**

- \*Files in the 10-100MB range tend to land in the 'sweet spot' for load parallelism (Correct)

**True or False: a user can define his own micro-partitions**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**Which of the following commands are not blocking operations? Choose 2 answers.**

Options are :

- *MERGE*
- *UPDATE*
- *COPY (Correct)*
- *INSERT (Correct)*

**Answer :**

\**COPY (Correct)*

\**INSERT (Correct)*

### Snowflake SnowPro Core Certification Exam Questions Set 12

**When loading data into Snowflake, the COPY command supports: Choose 2 answers**

Options are :

- *Aggregates*
- *Joins*
- *Data type conversion (Correct)*
- *Column reordering (Correct)*
- *Filters*

**Answer :**

\**Data type conversion (Correct)*

\**Column reordering (Correct)*

**Which of the following would be used to export / unload data from Snowflake?**

Options are :

- *INSERT INTO @stage*
- *EXPORT INTO @stage*
- *COPY INTO @ stage (Correct)*

- *EXPORT\_TO\_STAGE(stage => @Stage, select => 'select \* from t1');*

**Answer :**

\*COPY INTO @ stage (Correct)

**What is the default autosuspend for a warehouse created with the web UI**

Options are :

- *1 min*
- *5 min*
- *15 min*
- *10 min (Correct)*

**Answer :**

\*10 min (Correct)

### Snowflake Certification Preparation Set 1

**Which statement best describes "clustering"?**

Options are :

- *Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions (Correct)*
- *The clusterign key must be included in the COPY command when loading data into Snowflake*
- *The database administrator must define the clustering methodology for each Snowflake table*
- *Clustering can be disable within a Snowflake account*

**Answer :**

\*Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions (Correct)

**On which of the following cloud platforms can a Snowflake accoutn be hosted? Choose 2**

Options are :

- *Amazon Web Services (Correct)*
- *Oracle Cloud*

- *Private Virtual Cloud*
- *Microsoft Azure Cloud (Correct)*

**Answer :**

- \*Amazon Web Services (Correct)
- \*Microsoft Azure Cloud (Correct)

**A Virtual Warehouse auto-suspend and auto-resume setting apply to:**

Options are :

- *The queries currently being run by the Virtual Warehouse*
- *The primary cluster in the Virtual Warehouse*
- *The entire Virtual Warehouse (Correct)*
- *The database the Virtual Warehouse resides in*

**Answer :**

- \*The entire Virtual Warehouse (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 12

**Snowflake provides a method for its customers to override its clustering algorythm. This method is:**

Options are :

- *Cluster partitions*
- *Key partitions*
- *Micro-paritions*
- *Clusterign keys (Correct)*

**Answer :**

- \*Clusterign keys (Correct)

**When creating a user it is advisable to: Choose 2 answers**

Options are :

- *Set the user access to expire within a predefined timeframe*
- *Set a default role for the user (Correct)*

- Set the number of minutes to unlock to 15 minutes
- Set the user to be initially disabled
- Force an immediate password change (Correct)

**Answer :**

- \*Set a default role for the user (Correct)
- \*Force an immediate password change (Correct)

**Which statement best describes Snowflake tables?**

Options are :

- Snowflake tables are owned by a user
- Snowflake tables are the physical instantiation of data loaded into Snowflake
- Snowflake tables require a clustering key to be defined to perform optimally
- Snowflake tables are logical representation of underlying physical data (Correct)

**Answer :**

- \*Snowflake tables are logical representation of underlying physical data (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 7

**Which of the following commands set the Virtual Warehouse for a session?**

Options are :

- SET WAREHOUSE ;
- COPY WAREHOUSE FROM ;
- USE VIRTUAL\_WAREHOUSE ;
- USE WAREHOUSE ; (Correct)

**Answer :**

- \*USE WAREHOUSE ; (Correct)

**Increasing the maximum number of cluster in a Multi-Cluster Warehouse is an example of:**

Options are :

- Scaling max

- *Scaling up*
- *Scaling rythmically*
- *Scaling out (Correct)*

**Answer :**

\*Scaling out (Correct)

**What are three things customers want most from their enterprise data warehouse solution?  
Choose 3 answers**

Options are :

- *Concurrency (Correct)*
- *Simplicity (Correct)*
- *Performance (Correct)*
- *Open source based*
- *On-premise availability*

**Answer :**

\*Concurrency (Correct)  
\*Simplicity (Correct)  
\*Performance (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 17](#)

**True or False: When you create a custome role, it is a practice to immediately grant that role to ACCOUNT ADMIN**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**True or False: Multi Factor Authentication (MFA) is only available in conjunction with Single Sign On (SSO)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**True or False:** When data share is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data Consumers.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 5

**True or False:** It is possible to unload semi-structure formats such as JSON and parquet.

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**What is the fastest file format to upload files on snowflake**

Options are :

- *CSV (Correct)*
- *AVRO*
- *ORC*
- *parquet*

**Answer :**

\**CSV (Correct)*

**Which item in the Data Warehouse migration process does not apply in Snowflake?**

Options are :

- *Migrate Users*
- *Migrate Indexes (Correct)*
- *Migrate Schemas*
- *Build the Data pipeline*

**Answer :**

\*Migrate Indexes (Correct)

**True or False: There are storage charges fo both Internal and External Stages**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**Which of the following statement is true of zero copy cloning?**

Options are :

- *Zero copy cloning is licensed as an additional Snowflake feature*
- *At the time/instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)*
- *Zero copy cloning object does not inherit any privileges*
- *All zero copy cloning objects inherit privileges of their original objects*

**Answer :**

\*At the time/instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)

**Snowflake provides 2 mechanisms to reduce data storage cost for short-lived tables. These meachanisms are:**

Options are :

- *Transient Tables (Correct)*

- *Temporary Tables (Correct)*
- *Permanent Tables*
- *Provisional Tables*

**Answer :**

- \*Transient Tables (Correct)
- \*Temporary Tables (Correct)

### [Snowflake Practice Questions - SnowPro Core Certified Set 3](#)

**The number of queries that a Warehouse can concurrently process is determined by: Choose 2 answers.**

Options are :

- *The complexity of each query (Correct)*
- *The size of data required for each query (Correct)*
- *The size of data required for each query*
- *The CONCURRENT\_QUERY\_UNIT parameter set on the Snowflake account*

**Answer :**

- \*The complexity of each query (Correct)
- \*The size of data required for each query (Correct)

**After how long the history of queries accessed via web UI expires**

Options are :

- *14 days (Correct)*
- *1 hour*
- *30 days*
- *24 hours*

**Answer :**

- \*14 days (Correct)

**What privileges are required to execute a task?**

Options are :

- Tasks are stored in a special database called the utility database and require EXECUTE TASK privilege in the utility database
- A task is an account-level object and can only be executed by the ACCOUNT ADMIN role
- Tasks run automatically and do not run under any role
- To execute a task you need the EXECUTE TASK (Correct)

**Answer :**

\*To execute a task you need the EXECUTE TASK (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 8

**How is snowpipe credit cost calculated**

Options are :

- from the credit usage of the compute resources used (Correct)
- from the number of pipes

**Answer :**

\*from the credit usage of the compute resources used (Correct)

**What is the minimum duration charged when starting a virtual warehouse**

Options are :

- 1 day
- 1 second
- 1 hour
- 1 minute (Correct)

**Answer :**

\*1 minute (Correct)

**True or False: MFA can be used independently of SSO**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 4

**True or False: All storage table types include fail-safe storage**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**When scaling up you are primarily scaling for**

Options are :

- *concurrency*
- *performance (Correct)*

**Answer :**

\*performance (Correct)

**True or false to execute a task you must have OWNERSHIP privilege and also EXECUTE TASK**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 8

**How many shares can be consumed by a single Data Consumer?**

Options are :

- *1*

- Unlimited (Correct)
- 10
- 100, but can be increased by contacting support

**Answer :**

\*Unlimited (Correct)

**Which of the following are best practices for users with the SYSADMIN/ACCOUNTADMIN roles? Choose 3 answers**

Options are :

- They should not set up Multi Factor Authentication (MFA) as administrators. As administrator they might need to change the MFA settings and those enrolled in MFA are not able to do so
- They should only access and 'step into' the role ACCOUNTADMIN temporarily, as needed to complete a specified operation (Correct)
- They should ensure all the database objects in the account are owned by the ACCOUNTADMIN role
- The default role should be set to SYSTEMADMIN (the lower of the two) (Correct)
- They should use the SYSADMIN role to perform administrative work on database objects (Correct)

**Answer :**

\*They should only access and 'step into' the role ACCOUNTADMIN temporarily, as needed to complete a specified operation (Correct)  
 \*The default role should be set to SYSTEMADMIN (the lower of the two) (Correct)  
 \*They should use the SYSADMIN role to perform administrative work on database objects (Correct)

**A deterministic query is run at 8am, takes 5 minutes, and the results are cached. Which of the following statements are true? Choose 2 answers**

Options are :

- The same exact query will return the precomputed result if the underlying data hasn't changed and the results were last accessed within the previous 24 hour period (Correct)
- The exact query will ALWAYS return the precomputed result set for RESULT\_CACHE\_ACTIVE = time period
- The "24 hour" on the precomputed result gets renewed every time the exact query is executed (Correct)

- *The same exact query will return the precomputed result even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period.*

**Answer :**

\*The same exact query will return the precomputed result if the underlying data hasn't changed and the results were last accessed within the previous 24 hour period

(Correct)

\*The "24 hour" on the precomputed result gets renewed every time the exact query is executed (Correct)

### Snowflake Certification Preparation Set 16

#### **which file type can you use for unloading**

Options are :

- *parquet (Correct)*
- *AVRO*
- *JSON (Correct)*
- *CSV (Correct)*

**Answer :**

\*parquet (Correct)

\*JSON (Correct)

\*CSV (Correct)

#### **True or False: The Global Service Layer Stores statistical information on micro-partitions; used to build query execution plan**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

#### **True or False a single schema can exist in multiple databases**

Options are :

- *TRUE (Correct)*

- FALSE

**Answer :**

\*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 17

**Which of the following is are valid approaches to loading data into a Snowflake table? Select all that apply**

Options are :

- *The Snowflake Web Interface (UI) data loading wizard (Correct)*
- *Bulk copy from external storage (Correct)*
- *Continuous loading using Snowpiper REST API (Correct)*
- *Bulk copy from an internal stage (Correct)*

**Answer :**

\*The Snowflake Web Interface (UI) data loading wizard (Correct)

\*Bulk copy from external storage (Correct)

\*Continuous loading using Snowpiper REST API (Correct)

\*Bulk copy from an internal stage (Correct)

**True or False: Users are able to see the result sets of queries executed by other users who share the same role.**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**Which of the following are true of multi-cluster Warehouses? Select all that apply**

Options are :

- *A multi-cluster Wrehouse can add clusters automatically based on query activity (Correct)*
- *A multi-cluster Wrehouse can scale down when query activity slow (Correct)*

- A multi-cluster Wrehouse can automatically turn itself on when a query is executed against it (Correct)
- A multi-cluster Wrehouse can automatically turn itself off after a certain period of inactivity (Correct)

**Answer :**

- \*A multi-cluster Wrehouse can add clusters automatically based on query activity (Correct)
- \*A multi-cluster Wrehouse can scale down when query activity slow (Correct)
- \*A multi-cluster Wrehouse can automatically turn itself on when a query is executed against it (Correct)
- \*A multi-cluster Wrehouse can automatically turn itself off after a certain period of inactivity (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 36](#)

**Which type of table corresponds to a single Snowflake session**

Options are :

- *Transient*
- *Temporary (Correct)*
- *Provisioned*
- *Permanent*

**Answer :**

- \*Temporary (Correct)

**True or False: Snowpipe via REST API can only reference External Stages as source.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

- \**FALSE (Correct)*

**True or False: Snowflake charges additional fees to Data providers for each share they create.**

Options are :

- FALSE (Correct)

- TRUE

**Answer :**

\*FALSE (Correct)

## Snowflake Snowpro Certification Practice Test Set 2

**True or False: You can resize a Virtual Warehouse while queries are running.**

Options are :

- FALSE

- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Fail-safe is unavailable on which table types?**

Options are :

- Transient (Correct)

- Permanent

- Temporary (Correct)

- Provisional

**Answer :**

\*Transient (Correct)

\*Temporary (Correct)

**Which of the following statements about data sharing are true? Choose 2 answers.**

Options are :

- Reader Accounts are created and funded by Data Providers (Correct)

- New objects created by Data Provider are automatically shared with Data Consumers and Reader accounts

- Shared databases are read only (Correct)

- All database objects can be included in a shared database

**Answer :**

- \*Reader Accounts are created and funded by Data Providers (Correct)
- \*Shared databases are read only (Correct)

**Snowflake Snowpro Certification Practice Test Set 1**

**When should you consider disabling auto-suspend for a Virtual Warehouse?**

Options are :

- When users will be using compute at different times through a 24/7 period
- When you don't want to manually turn on the Warehouse each time a user needs it
- When compute must be available with no delay or lag time (Correct)
- When managing a steady workload (Correct)

**Answer :**

- \*When compute must be available with no delay or lag time (Correct)
- \*When managing a steady workload (Correct)

**Snowflake recommends, as a minimum, that all users with the following roles should be enrolled in Multi Factor Authentication (MFA): Select One**

Options are :

- SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN
- ACCOUNTADMIN (Correct)
- SECURITYADMIN, ACCOUNTADMIN
- SECURITYADMIN, ACCOUNTADMIN, SYSADMIN

**Answer :**

- \*ACCOUNTADMIN (Correct)

**UDF can be written in which language**

Options are :

- Java
- Python
- Javascript (Correct)

- *SQL (Correct)*

**Answer :**

- \*Javascript (Correct)
- \*SQL (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 21

**True or False: The longer the retention period, the higher the resulting storage cost.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**Which of the following DML commands isn't supported by Snowflake?**

Options are :

- *MERGE*
- *UPSERT (Correct)*
- *TRUNCATE TABLE*
- *UPDATE*

**Answer :**

- \*UPSERT (Correct)

**True or False: When active, a pipe requires a dedicated Virtual Warehouse to execute**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \*FALSE (Correct)

### Snowflake Certification Preparation Set 3

**Account level storage can be monitored via:**

Options are :

- *The Snowflake Web Interface (UI) in the Database section*
- *The snowflake Web Interface (UI) in the Account -> Billing and usage section (Correct)*
- *The Account Usage Schema -> ACCOUNT\_USAGE\_METRICS View*
- *The Information Schema -> account usage history View*

**Answer :**

\*The snowflake Web Interface (UI) in the Account -> Billing and usage section (Correct)

**Which of the following languages can be used to implement User Defined Functions (UDFs)?  
Choose 2 answers**

Options are :

- *Java*
- *Python*
- *Javascript (Correct)*
- *SQL (Correct)*

**Answer :**

\*Javascript (Correct)

\*SQL (Correct)

**True or False: A customer using SnowSQL / native connectors will be unable to also use the Snowflake Web Interface (UI) unless access to the UI is explicitly granted by support.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 7**

**What is the recommended file size for bulk upload?**

Options are :

- *files between 100 - 250 mb compressed and processed in sequence*
- *files between 100 - 250 mb compressed and processed in parallel (Correct)*
- *files between 100 - 250 mb non compressed and processed in parallel*
- *files between 100 - 250 mb non compressed and processed in sequence*

**Answer :**

\*files between 100 - 250 mb compressed and processed in parallel (Correct)

**True or False:** you can view which records returned errors for past loads using the VALIDATE command

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Micro-partition are sized within range of**

Options are :

- *terabytes*
- *megabytes (Correct)*
- *kilobites*
- *gigabytes*

**Answer :**

\*megabytes (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 11

**What is the lowest Snowflake edition that offers Time Travel for up to 90 days?**

Options are :

- *Premier edition*

- *Business Critical*
- *Enterprise edition (Correct)*
- *Standard edition*

**Answer :**

\*Enterprise edition (Correct)

**True or False:** A third party tool that supports JDBC or ODBC but has no Snowflake specific driver will be unable to connect to Snowflake.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of these are blocking statements**

Options are :

- *DELETE (Correct)*
- *COPY*
- *INSERT*
- *UPDATE (Correct)*
- *MERGE (Correct)*

**Answer :**

\*DELETE (Correct)

\*UPDATE (Correct)

\*MERGE (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 12

**True or False a single database can exist in multiple accounts**

Options are :

- *FALSE*

- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

**Which formats are supported for unloading data from Snowflake? Choose 2 answers**

Options are :

- ORC
- JSON (Correct)
- AVRO
- Delimited (CSV, TST, etc.) (Correct)

**Answer :**

\*JSON (Correct)

\*Delimited (CSV, TST, etc.) (Correct)

**Which of the following accurately represents how a table fits into Snowflake's logical container hierarchy?**

Options are :

- Database -> Schema -> Table -> Account
- Database -> Table -> Schema -> Account
- Account -> Schema -> Database -> Table
- Account -> Database -> Schema -> Table (Correct)

**Answer :**

\*Account -> Database -> Schema -> Table (Correct)

**What is the maximum compressed row size in Snowflake?**

Options are :

- 4000GB
- 50MB
- 8KB
- 16MB (Correct)

**Answer :**

\*16MB (Correct)

**The FLATTEN function is used to query which type of data in Snowflake?**

Options are :

- *None of the above*
- *Semi-structured data (Correct)*
- *Structured data*
- *Both above*

**Answer :**

\*Semi-structured data (Correct)

**Which of the following connectors are available in the Downloads section of the Snowflake Web Interface (UI)? Choose 2 answers.**

Options are :

- *Snow SQL (Correct)*
- *HIVE*
- *R*
- *ODBC (Correct)*

**Answer :**

\*Snow SQL (Correct)

\*ODBC (Correct)

#### [Snowflake SnowPro Core Certification Practice Test Set 4](#)

**What command is used to load files into an internal storage in Snowflake? Select one.**

Options are :

- *TRANSFER*
- *INSERT*
- *COPY INTO*
- *PUT (Correct)*

**Answer :**

\*PUT (Correct)

**What is the MAX\_FILE\_SIZE option that can be set when unloading data from Amazon S3?**

Options are :

- 1TB
- 500MB
- 50GB
- 5GB (Correct)

**Answer :**

\*5GB (Correct)

**Which object allows you to limit the number of credits consumed within a Snowflake account?**

Options are :

- Account usage tracking
- Credit consumption tracker
- Resource Monitor (Correct)
- Warehouse limit parameter

**Answer :**

\*Resource Monitor (Correct)

### Tackling Snowflake Certification - Practice Questions Set 5

**True or False: You can pause and resume a pipe**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**True or False: It is possible to query data from an internal or Named External stage without loading it first into Snowflake**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which statements are true of micropartitions? Choose 2 answers.**

Options are :

- *They are approximately 16MB in size (Correct)*
- *They are immutable (Correct)*
- *They are only encrypted in the Enterprise edition and above*
- *They are stored compressed only if COMPRESSED=TRUE on the table*

**Answer :**

\**They are approximately 16MB in size (Correct)*

\**They are immutable (Correct)*

### [Snowflake Snowpro Certification Practice Test Set 5](#)

**True or False: Snowflake enforces unique primary key, and foreign key constraints during DML operations**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**Which of the following stament is true of Snowflake: Select one**

Options are :

- *It was built as an on premise solution and then plotted to the cloud*
- *It was built specifically for the cloud (Correct)*

- It's based on Oracle architecture
- It was built for Hadoop Architecture
- It was designed as a hybrid database to allow customers to store data either on premises or in the cloud

**Answer :**

\*It was built specifically for the cloud (Correct)

**You can connect via MFA with**

Options are :

- Python (Correct)
- SnowSQL (Correct)
- Web Interface (Correct)
- ODBC (Correct)
- JDBC (Correct)

**Answer :**

\*Python (Correct)  
\*SnowSQL (Correct)  
\*Web Interface (Correct)  
\*ODBC (Correct)  
\*JDBC (Correct)

## SnowFlake SnowPro Advanced Architect Practice Exam Set 7

**Select three types of tables that exist within Snowflake**

Options are :

- Provisioned
- Transient (Correct)
- Temporary (Correct)
- Permanent (Correct)

**Answer :**

\*Transient (Correct)  
\*Temporary (Correct)  
\*Permanent (Correct)

**Snowflake natively supports which file formats**

Options are :

- *JPG*
- *MP4*
- *CSV (Correct)*
- *JSON (Correct)*

**Answer :**

\**CSV (Correct)*  
\**JSON (Correct)*

**How can you create a warehouse**

Options are :

- *with the command create WAREHOUSE (Correct)*
- *Through web interface "Account -> create warehouse"*
- *with the command create VIRTUAL\_WAREHOUSE*
- *through the web interface "Warehouses -> Create" (Correct)*

**Answer :**

\**with the command create WAREHOUSE (Correct)*  
\**through the web interface "Warehouses -> Create" (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 36**

**True or False: To load data you must specify a Named File Format**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**Which of the following statements is true about Snowflake micro-partitioning?**

Options are :

- Micro-partitioning can be disabled within a Snowflake account
- Micro-partitioning requires a partitioning schema to be defined upfront
- Micro-partitioning has been known to introduce data skew
- Micro-partitioning is transparently completed using the ordering that occurs when the data is inserted/loaded (Correct)

**Answer :**

\*Micro-partitioning is transparently completed using the ordering that occurs when the data is inserted/loaded (Correct)

**Which of the following statements are true of Snowflake data loading? Choose 3 answers.**

Options are :

- It is recommended to do frequent, single row DML
- It is recommended to validate the data before loading it into the Snowflake target table (Correct)
- It is recommended to use staging tables to manage MERGE statements (Correct)
- VARIANT Null values are not the same as SQL Null values (Correct)

**Answer :**

\*It is recommended to validate the data before loading it into the Snowflake target table (Correct)

\*It is recommended to use staging tables to manage MERGE statements (Correct)

\*VARIANT Null values are not the same as SQL Null values (Correct)

## Snowflake Certification Preparation Set 12

**Increasing the size of a warehouse from an X-Small to an X-Large is an example of:**

Options are :

- Scaling up (Correct)
- Scaling rhythmically
- Scaling max
- Scaling out

**Answer :**

\*Scaling up (Correct)

**True or False: Micro-partitions metadata enables some operations to be completed without requiring Compute**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Which of the following is true of Snowpipe via REST API?**

Options are :

- *You can only use it on internal stages*
- *All COPY INTO options are available during pipe creation*
- *Snowpipe keeps track of which files it has loaded (Correct)*
- *Snowflake Automatically manage the compute required to execute the Pipe's copy into command (Correct)*

**Answer :**

\*Snowpipe keeps track of which files it has loaded (Correct)

\*Snowflake Automatically manage the compute required to execute the Pipe's copy into command (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 17](#)

**Which are true of Snowflake roles**

Options are :

- *The public role controls all other roles*
- *Roles are subsets of users and users own objects in Snowflake*
- *All grants to objects are given to roles, never to users (Correct)*
- *In order to do DML/DQL, a user needs to have selected a single role that has specific access to the object and operation (Correct)*

**Answer :**

\*All grants to objects are given to roles, never to users (Correct)

\*In order to do DML/DOL, a user needs to have selected a single role that has specific access to the object and operation (Correct)

**True or False: It is best practice to define a clustering key on every table**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**As a best practice, clustering keys should only be defined on a table of which minimum size?**

Options are :

- Multi-terabyte (TB) Range (Correct)
- Multi-megabyte (MB) Range
- Multi-gigabyte (GB) Range
- Multi-kilobyte (KB) Range

**Answer :**

\*Multi-terabyte (TB) Range (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 4**

**If auto-suspend is enabled for a Virtual Warehouse, the warehouse is automatically suspended when:**

Options are :

- The last query using the Warehouse completes
- There are no users logged into Snowflake
- All Snowflake sessions using the Warehouse are terminated
- The warehouse is inactive for a specified period of time (Correct)

**Answer :**

\*The warehouse is inactive for a specified period of time (Correct)

**True or false: Each worksheet in the Snowflake Web Interface (UI) can be associated with different roles, databases, schema, and Virtual Warehouses**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**True or False: Snowflake charges a premium for semi-structured data.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

### [Snowflake SnowPro Core Certification Practice Tests Set 15](#)

**What parameter controls if the Virtual Warehouse starts immediately after the CREATE WAREHOUSE statement? Select one**

Options are :

- *INITIALLY\_SUSPENDED = TRUE/FALSE (Correct)*
- *START\_TIME = 60 // (seconds from now)*
- *START\_AFTER\_CREATE = TRUE/FALSE*
- *START.TIME = CURRENT.DATE()*

**Answer :**

\**INITIALLY\_SUSPENDED = TRUE/FALSE (Correct)*

**True or False: Snowflake bills a minimum of 5 minutes each time a Virtual Warehouse is started**

Options are :

- *FALSE (Correct)*

- *TRUE*

**Answer :**

\**FALSE* (Correct)

**In which layer of its architecture does Snowflake store its metadata statistics? Select one.**

Options are :

- *Compute Layer*
- *Cloud Service Layer (Correct)*
- *Database Layer*
- *Storage Layer*

**Answer :**

\**Cloud Service Layer (Correct)*

### **SnowFlake SnowPro Advanced Architect Practice Exam Set 6**

**True or False: a user has to specify which cluster will run on in multi-clustering Warehouse.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**Which security features does snowflake support**

Options are :

- *Data is always stored encrypted (Correct)*
- *Can run on a customers private cloud*
- *Role-based access control (Correct)*
- *Multi factor Authentication mandatory for all users*

**Answer :**

\**Data is always stored encrypted (Correct)*

\**Role-based access control (Correct)*

**When a pipe is recreated using CREATE or REPLACE command: Select one**

Options are :

- *All of the above*
- *Previously loaded files will be ignored*
- *The REFRESH parameter is set to TRUE*
- *The Pipe load history is reset to empty (Correct)*

**Answer :**

\*The Pipe load history is reset to empty (Correct)

**Snowflake SnowPro Core Certification Exam Questions Set 9**

**How would you exclude a serie of SQL statements using Task?**

Options are :

- *Include SQL statement in the body of the task, create task my task ↗ as insert into target1 select \* from stream\_s1 where ↗ INSERT into target2 SELECT from stram\_s1 where ↗*
- *A stored procedure executing multiple SQL statements and invoke the stored procedure from the task. CREATE TASK mytask AS CALL stored\_prc\_multiple\_statements\_inside(); (Correct)*
- *Create a task for each SQL statement (e.g. resulting in task1, task, etc.) and string the SQL statements by having a control task calling task1, task2, sequentially.*
- *Stored procedure can have only one DML statement per stored procedure invocation and therefore you should sequence stored procedure calls in the task definition CREATE TASK mytask ↗ AS CALL stored\_prc1(); call stored\_prc2();*

**Answer :**

\*A stored procedure executing multiple SQL statements and invoke the stored procedure from the task. CREATE TASK mytask AS CALL stored\_prc\_multiple\_statements\_inside(); (Correct)

**Which of the following are examples of operations which require a Virtual Warehouse to complete, assuming no queries have been executed previously? Choose 3 answers.**

Options are :

- *UPDATE (Correct)*

- *COUNT*
- *SUM (Correct)*
- *MIN (Correct)*

**Answer :**

- \**UPDATE (Correct)*
- \**SUM (Correct)*
- \**MIN (Correct)*

**Which of the following best describes where Snowflake stores metadata?**

Options are :

- *In the Cloud Services Layer (Correct)*
- *Inside the Virtual Warehouses*
- *Within the drivers*
- *Within data files*

**Answer :**

- \**In the Cloud Services Layer (Correct)*

### **Snowflake SnowPro Core Certification Exam Questions Set 6**

**True or False: when a user creates a role, they are initially assigned ownership of the role and they maintain ownership until it is transferred to another user**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \**TRUE (Correct)*

**For a Multi-cluster Warehouse the number of credits billed is calculated on: Select one.**

Options are :

- *The size of the Warehouse and the number of clusters that run within a given time period (Correct)*
- *The number of queries that run using the Warehouse*

- *The number of users who accessed the Warehouse*
- *The size of the Warehouse and the maximum number of clusters configured for the Warehouse*

**Answer :**

\*The size of the Warehouse and the number of clusters that run within a give time period (Correct)

**True or False: You can define multiple columns within a clustering key on a table**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### Snowflake Certification Preparation Set 16

**True or False: When Snowflake is configured using Single Sign in (SSO) Snowflake receives the username and credentials from the sso service and loads them into the customer's Snowflake account.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Each incremental increase in Virtual Warehouse size (eg, Medium to Large) generally results in what?**

Options are :

- *Better query scheduling*
- *More micro-partitions*
- *Higher storage cost*
- *Double the number of servers in the compute cluster (Correct)*

**Answer :**

\*Double the number of servers in the compute cluster (Correct)

**Which of the following items does the Cloud Service Layer manage?**

Options are :

- *Data security (Correct)*
- *External blob storage*
- *Query compilation and optimization (Correct)*
- *User authentication (Correct)*
- *Metadata (Correct)*

**Answer :**

\*Data security (Correct)

\*Query compilation and optimization (Correct)

\*User authentication (Correct)

\*Metadata (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 14

**True or False: Once created a micro-partition will never be changed.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Data storage for individual tables can be monitored using which commands and/or objects?  
Choose 2 answers**

Options are :

- *SHOW TABLES; (Correct)*
- *Information Schema -> TABLE HISTORY (Correct)*
- *SHOW STORAGE BY TABLE;*
- *Information Schema -> TABLE\_STORAGE\_METRICS*

**Answer :**

- \*SHOW TABLES; (Correct)
- \*Information Schema -> TABLE HISTORY (Correct)

**True or False: Data in Fail-Safe can be deleted by a user of the Snowflake Team before it expires**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

- \*FALSE (Correct)

**What is the recommended Snowflake data type to store semi-structured data like JSON?**

Options are :

- BLOB
- RAW
- VARCHAR
- VARIANT (Correct)

**Answer :**

- \*VARIANT (Correct)

**What is the most granular object that the Time Travel retention can be defined on?**

Options are :

- Database
- Table (Correct)
- Account
- Schema

**Answer :**

- \*Table (Correct)

**Which of the following cases are common uses for zero-copy cloning? Choose 3 answers**

Options are :

- *Data Backups (Correct)*
- *Point in time snapshots (Correct)*
- *Quick provisioning of Dev and Test/QA environments (Correct)*
- *Performance optimization*

**Answer :**

\*Data Backups (Correct)  
\*Point in time snapshots (Correct)  
\*Quick provisioning of Dev and Test/QA environments (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 6

**What is the minimum Snowflake edition that provides data sharing?**

Options are :

- *Business Critical*
- *Standard edition*
- *Enterprise (Correct)*
- *Premier*

**Answer :**

\*Enterprise (Correct)

**What is the recommended method for loading data into Snowflake?**

Options are :

- *Load data via frequent, single row DML's*
- *Load data in batch (Correct)*
- *Load by row*
- *Load data by writing it in the Snowflake Web Interface (UI)*

**Answer :**

\*Load data in batch (Correct)

**What are the three layers that make up the Snowflake architecture? Choose 3 answers**

Options are :

- *Compute (Correct)*
- *Tri-Secret Secure*
- *Storage (Correct)*
- *Cloud Services (Correct)*

**Answer :**

\*Compute (Correct)  
\*Storage (Correct)  
\*Cloud Services (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 24

**Query results are stored in the result cache for how long after they have last been accessed, assuming no data changes have occurred?**

Options are :

- *12 Hours*
- *24 Hours (Correct)*
- *3 Hours*
- *1 Hour*

**Answer :**

\*24 Hours (Correct)

**Which of the following are options when creating a Virtual Warehouse? Choose 2 answers**

Options are :

- *Auto-resize*
- *Auto-drop*
- *Auto-resume (Correct)*
- *Auto-suspend (Correct)*

**Answer :**

\*Auto-resume (Correct)  
\*Auto-suspend (Correct)

**True or False: Snowflake allows its customers to directly access the micro-partition files that make up its tables.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 35

**Which of the following are main sections of the top navigation of the Snowflake Web Interface (UI)?**

Options are :

- *Worksheet (Correct)*
- *Database (Correct)*
- *Tables*
- *Warehouses (Correct)*

**Answer :**

\**Worksheet (Correct)*  
\**Database (Correct)*  
\**Warehouses (Correct)*

**What happens when a data provider revokes privileges to a Share on an object in their source database?**

Options are :

- *The Object immediately becomes unavailable for all data consumers (Correct)*
- *A static copy of the object at the time the privilege was revoked is created in the Data Consumer account*
- *Any additional data arriving after this point in time will not be visible to Data Consumers*
- *The Data Consumers stop seeing data updates and become responsible for storage charges for the object*

**Answer :**

\*The Object immediately becomes unavailable for all data consumers (Correct)

**What service does Snowflake automatically provide for customers that they might have been responsible for with their on premise system? Select all below that apply.**

Options are :

- *Maintaining metadata and statistics (Correct)*
- *Patching software (Correct)*
- *Physical security*
- *Installing and configuring hardware (Correct)*

**Answer :**

\*Maintaining metadata and statistics (Correct)

\*Patching software (Correct)

\*Installing and configuring hardware (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 13](#)

**When scaling up warehouses by increasing Virtual Warehouse t-shirt size, you are primarily scaling for improved:**

Options are :

- *Performance (Correct)*
- *Concurrency*

**Answer :**

\*Performance (Correct)

**When scaling out you are primarily scaling for:**

Options are :

- *concurrency (Correct)*
- *performance*

**Answer :**

\*concurrency (Correct)

**Which of the following are options when creating a virtual warehouse**

Options are :

- *Local SSD size*
- *User count*
- *Auto-suspend (Correct)*
- *Auto-resume (Correct)*

**Answer :**

\*Auto-suspend (Correct)  
\*Auto-resume (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 32](#)

**How many credits per hour is consumed by a 2X-Large warehouse**

Options are :

- 16
- 32 (*Correct*)
- 64
- 128

**Answer :**

\*32 (Correct)

**Materialized view can be created on top of external tables to improve query performance without any additional cost. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Snowflake compute costs depends on : (Select all that apply) ?**

Options are :

- *The amount of time warehouse have run (Correct)*
- *The number of rows returned in queries*
- *The sizes of running warehouses (Correct)*
- *The amount of data stored in the account*
- *None of the above*

**Answer :**

- \*The amount of time warehouse have run (Correct)
- \*The sizes of running warehouses (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 35

**What stages can you use when setting up your Snowflake account with AWS as your cloud platform provider ? Select all that apply**

Options are :

- *Named External - Using S3 Buckets (Correct)*
- *Named External - Hadoop Cloudera File System*
- *Named External - Using Azure Blob storage (Correct)*
- *Named External - Using GCS/GCP Buckets (Correct)*
- *Named External - AWS Redshift tables*
- *None of the above*

**Answer :**

- \*Named External - Using S3 Buckets (Correct)
- \*Named External - Using Azure Blob storage (Correct)
- \*Named External - Using GCS/GCP Buckets (Correct)

**Warehouse cache size changes with warehouse size. TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**A client has set Time Travel to keep a table data for 27 days. 35 days after dropping the table, their want to get it back. They ask you if it is possible to recover the data of that table.**

Options are :

- Yes, it is possible
- No, it is not possible (Correct)

**Answer :**

\*No, it is not possible (Correct)

### Tackling Snowflake Certification - Practice Questions Set 1

**Which syntax allows you to concat the column named 'EMPLOYEE' between two % signs ? Select all that apply**

Options are :

- '%' || EMPLOYEE || '%' (Correct)
- % + EMPLOYEE + %
- CONCAT('%', EMPLOYEE, '%')
- concat('%', concat(EMPLOYEE, '%')) (Correct)
- All of the above

**Answer :**

\*%' || EMPLOYEE || '%' (Correct)  
\*concat('%', concat(EMPLOYEE, '%')) (Correct)

**Snowpipe is best suited to load what kind of data ? Select the best option**

Options are :

- Small volume of frequent data (Correct)
- Large volume of data on a batch schedule
- Any kind of volume small or batch or continuous
- None of the above

**Answer :**

\*Small volume of frequent data (Correct)

**Scaling a warehouse UP will decrease the number of clusters. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

### **Snowflake Certification Preparation Set 14**

**Which Snowflake Edition limit Time Travel to maximum 24 hours ?**

Options are :

- *Standard Edition (Correct)*
- *Business Critical Edition*
- *Virtual Private Snowflake*
- *Enterprise Edition*
- *None of the above*

**Answer :**

\**Standard Edition (Correct)*

**How many servers are in standard Large warehouse ?**

Options are :

- *2*
- *4*
- *8 (Correct)*
- *16*

**Answer :**

\**8 (Correct)*

**Which transformations can be applied on data while loading it into a table using COPY command ?  
Select all that apply**

Options are :

- *column reordering (Correct)*
- *Aggregation*
- *Casts (Correct)*
- *Filtering*
- *Column omission (Correct)*
- *Truncating text strings that exceed the target column length (Correct)*

**Answer :**

- \*column reordering (Correct)
- \*Casts (Correct)
- \*Column omission (Correct)
- \*Truncating text strings that exceed the target column length (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 6

**Which of the following terms describes Snowflake's Architecture ?**

Options are :

- *Shared Disk*
- *Shared Nothing*
- *Shared Data (Correct)*
- *Shared Memory*

**Answer :**

- \*Shared Data (Correct)

**In which mode Snowflake multi-cluster warehouses can run ? Select all that apply.**

Options are :

- *Maximized (Correct)*
- *Auto-resume*
- *Auto-suspended*
- *Auto-scale (Correct)*

- *None of the above*

**Answer :**

- \*Maximized (Correct)
- \*Auto-scale (Correct)

**Snowflake Web UI data loading wizard can be use to load a 1GB file. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FLASE (Correct)*

**Answer :**

- \*FLASE (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 36](#)

**Which of the following are types of cache use by Snowflake ? Select all that apply**

Options are :

- *Metadata cache (Correct)*
- *Index cache*
- *Query result cache (Correct)*
- *Data cache (Correct)*
- *None of the above*

**Answer :**

- \*Metadata cache (Correct)
- \*Query result cache (Correct)
- \*Data cache (Correct)

**Snowpipe◆s charges are calculated per-second and per core**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which type of view has an extra layer of protection to hide the SQL code from unauthorized viewing ?**

Options are :

- *Standard*
- *Secure (Correct)*
- *Materialized*
- *Permanent*

**Answer :**

\*Secure (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 43

**A task tree can have task spanning multiple schemas ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Virtual warehouses cannot be resized while they are running. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which one of these actions is an automated process**

Options are :

- *Scaling OUT (Correct)*
- *Scaling UP*

**Answer :**

\*Scaling OUT (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 28

**What factors impact credit consumption of the compute layer ? Select all that apply**

Options are :

- *Amount of data processed*
- *Size of warehouse (Correct)*
- *Number of clusters (Correct)*
- *Number of user*

**Answer :**

\*Size of warehouse (Correct)

\*Number of clusters (Correct)

**How long does Snowpipe History Load last ?**

Options are :

- *7 days*
- *14 days (Correct)*
- *21 days*
- *28 days*

**Answer :**

\*14 days (Correct)

**A company operating in the health sector wants to move to Snowflake. The company asks you about the edition of Snowflake to choose to comply with HIPAA regulation. Select the minimum edition of Snowflake to recommend ?**

Options are :

- *Enterprise Edition*

- *Business Critical Edition (Correct)*
- *Standard Edition*
- *Virtual Private Snowflake*

**Answer :**

\*Business Critical Edition (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 26**

**When you share data with a consumer account, who pays for compute time when data is access by the consumer?**

Options are :

- *The data provider does*
- *The data consumer does. (Correct)*
- *Both provider and consumer does*
- *None of the above*

**Answer :**

\*The data consumer does. (Correct)

**Snowflake Web Interface data loading wizard is designed to load small amounts of data where each data file is smaller than 50MB in size. TRUE or FALSE load**

Options are :

- *TRUE (Correct)*
- *FLASE*

**Answer :**

\*TRUE (Correct)

**Which of the following terms are NOT associated with the Cloud Services Layer ?**

Options are :

- *Query Planning*
- *Query Optimization*
- *Query Processing (Correct)*

- *Query Design*
- *Query Compilation*

**Answer :**

\*Query Processing (Correct)

### Snowflake Certification Preparation Set 23

**Is Snowflake HIPAA compliant ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**When loading JSON file inside Snowflake table, which file format option remove the outer array structure and load the records into separate table rows ? (**

Options are :

- *STRIP\_NULL\_VALUES = TRUE*
- *STRIP\_OUTER\_ARRAY = TRUE (Correct)*
- *TRIM\_SPACE = TRUE*
- *None of the above*

**Answer :**

\*STRIP\_OUTER\_ARRAY = TRUE (Correct)

**Which compute resource can be used by Snowflake for data loading jobs ? Select all that apply**

Options are :

- *User-managed virtual warehouse (Correct)*
- *Hardware provisionned by user directly from cloud provider*
- *Snowflake-managed serverless compute (Correct)*
- *Cloud Provider Runtime Environment*

**Answer :**

- \*User-managed virtual warehouse (Correct)
- \*Snowflake-managed serverless compute (Correct)

**Which common tasks for traditional on-premises database and IT staff, are not required with Snowflake ?**

Options are :

- *Maintaining Metadata (Correct)*
- *Maintaining statistics (Correct)*
- *Maintaining the physical security of a server room (key cards, door locks, etc.) (Correct)*
- *Maintaining database objects*

**Answer :**

- \*Maintaining Metadata (Correct)
- \*Maintaining statistics (Correct)
- \*Maintaining the physical security of a server room (key cards, door locks, etc.) (Correct)

**A Snowflake session can have more than one virtual warehouse. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**Your organization wants to use Column-level security to apply masking policies to columns in tables or views. Which Snowflake's edition offers this features ? Select all that apply.**

Options are :

- *Enterprise Edition (Correct)*
- *Business Critical Edition (Correct)*
- *Virtual Private Snowflake (Correct)*
- *Standard*

- *None of the above*

**Answer :**

- \*Enterprise Edition (Correct)
- \*Business Critical Edition (Correct)
- \*Virtual Private Snowflake (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 16

**What ROLE should create users and roles in Snowflake ? Select all that apply**

Options are :

- *ACCOUNTADMIN*
- *USERADMIN (Correct)*
- *SYSADMIN*
- *SECURITYADMIN (Correct)*
- *PUBLIC*

**Answer :**

- \**USERADMIN (Correct)*
- \**SECURITYADMIN (Correct)*

**What are the actions cloud services layer is responsible for ? (Select all that apply)**

Options are :

- *User authentication (Correct)*
- *Infrastructure management (Correct)*
- *Metadata management (Correct)*
- *Query parsing and optimization (Correct)*
- *Access control (network policies) (Correct)*
- *Query processing*

**Answer :**

- \*User authentication (Correct)
- \*Infrastructure management (Correct)
- \*Metadata management (Correct)
- \*Query parsing and optimization (Correct)
- \*Access control (network policies) (Correct)

**A company operating in the health sector wants to move to Snowflake. The company asks you about the edition of Snowflake to choose to comply with HIPAA regulation. Select the minimum edition of Snowflake to recommend ?**

Options are :

- *Enterprise Edition*
- *Business Critical Edition (Correct)*
- *Standard Edition*
- *Virtual Private Snowflake*

**Answer :**

\*Business Critical Edition (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 8

**When running a SELECT COUNT(\*) on a table, which of the following statements is TRUE ?**

Options are :

- *Only XS warehouse is needed because Count statistics are stored in Results Cache*
- *No warehouse is needed because Count statistics are stored in Metadata Cache (Correct)*
- *No warehouse will be needed because all users share count statistics*
- *Only XS warehouse will be needed because all roles share count statistics*

**Answer :**

\*No warehouse is needed because Count statistics are stored in Metadata Cache (Correct)

**What privileges does SYSADMIN have by default?**

Options are :

- *CREATE DATABASE (Correct)*
- *CREATE SHARE*
- *CREATE WAREHOUSE (Correct)*
- *CREATE USER*

- *CREATE ROLE*

**Answer :**

- \**CREATE DATABASE* (Correct)
- \**CREATE WAREHOUSE* (Correct)

**If you find a data-related tool that is not listed as part of the Snowflake ecosystem, what industry standard options could you check for as a way to easily connect to Snowflake?**

Options are :

- *Check see if the tool can connect to other solution via JDBC (Correct)*
- *Check see if the tool can connect to other solution via ODBC (Correct)*
- *Check see if there is a petition in the community to create a driver*
- *Check see if you can develop a driver and put it on github*

**Answer :**

- \*Check see if the tool can connect to other solution via JDBC (Correct)
- \*Check see if the tool can connect to other solution via ODBC (Correct)

### Snowflake Snowpro Certification Practice Test Set 3

**For how long does a query remain in the query history tab ?**

Options are :

- *1 day*
- *7 days*
- *14 days (Correct)*
- *48 days*
- *None of the above*

**Answer :**

- \**14 days (Correct)*

**What is not a Snowflake system-defined role ?**

Options are :

- *PUBLIC*

- *ANALYST (Correct)*

- *USERADMIN*

- *SYSADMIN*

- *ACCOUNTADMIN*

- *SECURITYADMIN*

**Answer :**

\*ANALYST (Correct)

**Select all Snowflake stages.**

Options are :

- *Internal Named Stage (Correct)*

- *External Named Stage (Correct)*

- *User Stage (Correct)*

- *Table Stage (Correct)*

- *Account Stage (Correct)*

**Answer :**

\*Internal Named Stage (Correct)

\*External Named Stage (Correct)

\*User Stage (Correct)

\*Table Stage (Correct)

\*Account Stage (Correct)

### [Snowflake Certification Preparation Set 18](#)

**What are the two mechanisms that can be used with Snowpipe for detecting that staged files are available ?**

Options are :

- *Calling Snowpipe REST endpoints (Correct)*

- *Calling Snowpipe SOAP endpoints*

- *Automating Snowpipe using CloudWatch*

- *Automating Snowpipe using cloud messaging (Correct)*

**Answer :**

- \*Calling Snowpipe REST endpoints (Correct)
- \*Automating Snowpipe using cloud messaging (Correct)

**What action caused a Data cache to be purged ? Select the best option.**

Options are :

- Waiting 24 hours
- Resuming the warehouse
- Suspending the warehouse (Correct)
- None of the above

**Answer :**

- \*Suspending the warehouse (Correct)

**Snowflake Virtual Warehouses cannot be resized while running. TRUE or FALSE ?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

- \*FALSE (Correct)

### [Snowflake Certification Preparation Set 23](#)

**Which of the following best describes Snowflake's processing engine ?**

Options are :

- Native SQL (Correct)
- Leverages Apache Spark
- Based on map reduce
- A derivative of presto

**Answer :**

- \*Native SQL (Correct)

**Who can use the data cache ? Select the best option.**

Options are :

- *Anyone who uses the same warehouse (Correct)*
- *Only the user who execute the query*
- *Only user with the same role*
- *Only account administrator*
- *None of the above*

**Answer :**

\*Anyone who uses the same warehouse (Correct)

**Which cache runs on a 24 hour clock?**

Options are :

- *Result cache (Correct)*
- *Metadata Cache*
- *Warehouse cache*

**Answer :**

\*Result cache (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 25**

**If a length is not specified when defined a VARCHAR column the default length is ?**

Options are :

- *256 characters*
- *1 character*
- *0 character*
- *Maximum length possible (Correct)*
- *None of the above*

**Answer :**

\*Maximum length possible (Correct)

**In snowflake container hierarchy, what container is represented as a URL ( e.g. <https://jz96798.us-east-1.azure.snowflakecomputing.com> ) ?**

Options are :

- *Database*
- *Schema*
- *Role*
- *Account (Correct)*
- *Region*

**Answer :**

\*Account (Correct)

**What attributes make Snowflake a true SaaS solution ?**

Options are :

- *No hardware to purchase or configure (Correct)*
- *No creation of user accounts or roles is required*
- *No maintenance upgrades or patches to install (Correct)*
- *No data storage costs*
- *No query processing costs*
- *Transparent releases don't required user intervention (Correct)*

**Answer :**

\*No hardware to purchase or configure (Correct)

\*No maintenance upgrades or patches to install (Correct)

\*Transparent releases don't required user intervention (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 42](#)

**In the Query History tab, how can you identify if a query pulled data from long-term centralized storage. Select the best option**

Options are :

- *Look in the Bytes scanned column for a green bar (Correct)*

- Look in the Bytes scanned column for a green bar
- Look in the Total Duration column for a orange bar
- Look in the Total Duration column for a green bar

**Answer :**

\*Look in the Bytes scanned column for a green bar (Correct)

**What kind of metadata does Snowflake stores in micro-partition for a table**

Options are :

- MIN (Correct)
- MAX (Correct)
- Table's privileges
- Table's owner
- Number of distinct values (Correct)
- None of the above

**Answer :**

\*MIN (Correct)

\*MAX (Correct)

\*Number of distinct values (Correct)

**It is possible to install Snowflake on-premise. TRUE or FALSE**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Snowflake SnowPro Core Certification Practice Tests Set 11**

**Streams support repeatable read isolation. TRUE or FALSE ?**

Options are :

- TRUE (Correct)

- *FALSE*

**Answer :**

\*TRUE (Correct)

**Snowflake has scalar and tabular user-defined functions (UDFs), with support for both SQL, JavaScript, Java. TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Snowflake can be run on private cloud infrastructures (on-premises or hosted). TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### **Snowflake SnowPro Core Certification Practice Test Set 3**

**Which installment option version of Snowflake are available ?**

Options are :

- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)*
- *Microsoft Cloud Native Accounts*
- *Hybrid On-Premise + Cloud Installation*
- *Enterprise In-House VPC Installation*
- *On-Premise Custom Installation*

**Answer :**

- \*Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)
- \*Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)

**When a warehouse is resized, which queries make use of the new size ?**

Options are :

- *Only currently running queries*
- *Only subsequent queries (Correct)*
- *Both current and subsequent queries*

**Answer :**

- \*Only subsequent queries (Correct)

**In which Snowflake editions is it possible to extended Time Travel (up to 90 days) ? (Select two or more)**

Options are :

- *Standard Edition*
- *Enterprise Edition (Correct)*
- *Business Critical Edition (Correct)*
- *Virtual Private Snowflake (Correct)*

**Answer :**

- \*Enterprise Edition (Correct)
- \*Business Critical Edition (Correct)
- \*Virtual Private Snowflake (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 37**

**Snowflake account hosted on AWS cannot load data from a file that is staged on GCP or Azure.  
TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \*FALSE (Correct)

**When data is staged to an internal stage using PUT command, the data is encrypted on client machine. TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Who can use the metadata cache ?**

Options are :

- *Only the user who execute the query*
- *Only user with the same role*
- *Only account administrator*
- *Everyone (Correct)*
- *None of the above*

**Answer :**

\**Everyone (Correct)*

### **Snowflake Practice Questions - SnowPro Core Certified Set 3**

**Which of the following have drivers/connectors (or information about where to find them) available via Help->Downloads in the Snowflake WebUI?**

Options are :

- *Go (Correct)*
- *R*
- *Node.Js (Correct)*
- *JDBC (Correct)*
- *Hive*
- *Spark*

**Answer :**

- \*Go (Correct)
- \*Node.Js (Correct)
- \*JDBC (Correct)

**Which type of stages are automatically available in Snowflake and do not need to be created or configured ?**

Options are :

- *User (Correct)*
- *Table (Correct)*
- *Named Internal*
- *Named External*
- *None of the above*

**Answer :**

- \*User (Correct)
- \*Table (Correct)

**What is the maximum size for a VARIANT data type column ?**

Options are :

- *64MB uncompressed*
- *16MB uncompressed (Correct)*
- *32MB uncompressed*
- *128MB uncompressed*

**Answer :**

- \*16MB uncompressed (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 11

**What is the account identifier inside this Snowflake URL ' <https://jz96798.us-east-1.azure.snowflakecomputing.com> ' ?**

Options are :

- *jz96798.us-east-1.azure*

- *us-east-1.azure*
- *jz96798 (Correct)*
- *snowflakecomputing.com*
- *Azure*

**Answer :**

\*jz96798 (Correct)

**Is it possible to disable Time Travel for a Snowflake Account ?**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**Which file formats is not supported by Snowflake to load data ?**

Options are :

- *Delimited files*
- *JSON*
- *Avro*
- *ORC*
- *Parquet*
- *PDF (Correct)*

**Answer :**

\*PDF (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 2**

**Which of the following is true regarding a tree of tasks? Select all that apply.**

Options are :

- *Root tasks must have a schedule, or the tree will never execute (Correct)*

- Child tasks can have schedules
- Child tasks cannot have schedules (Correct)
- Child tasks can have an AFTER clause that triggers them to execute once a previous task has been completed (Correct)

**Answer :**

- \*Root tasks must have a schedule, or the tree will never execute (Correct)
- \*Child tasks cannot have schedules (Correct)
- \*Child tasks can have an AFTER clause that triggers them to execute once a previous task has been completed (Correct)

**To create a SHARE what is the minimum required role ?**

Options are :

- SYSADMIN
- SECURITYADMIN
- ACCOUNTADMIN (Correct)
- USERADMIN

**Answer :**

- \*ACCOUNTADMIN (Correct)

**What services are managed by Cloud Services layer ? (Select all that apply)**

Options are :

- Authentication (Correct)
- Infrastructure Management (Correct)
- Metadata Management (Correct)
- Query parsing and optimization
- Data storage
- Query processing

**Answer :**

- \*Authentication (Correct)
- \*Infrastructure Management (Correct)
- \*Metadata Management (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 25

**Which of the following data types are treated as VARCHAR by Snowflake ?**

Options are :

- *TEXT (Correct)*
- *STRING (Correct)*
- *CHAR (Correct)*
- *VARCHAR (Correct)*
- *NVARCHAR*

**Answer :**

\*TEXT (Correct)  
\*STRING (Correct)  
\*CHAR (Correct)  
\*VARCHAR (Correct)

**What actions can be triggered by a resource monitor ? Select all that apply**

Options are :

- *Notify and Suspend (Correct)*
- *Auto-resume*
- *Notify (Correct)*
- *Auto-scale*
- *Notify and Suspend Immediately (Correct)*
- *None of the above*

**Answer :**

\*Notify and Suspend (Correct)  
\*Notify (Correct)  
\*Notify and Suspend Immediately (Correct)

**How long does Bulk data load 'Load history' last ?**

Options are :

- *14*

- 28
- 32
- 64 (*Correct*)

**Answer :**

\*64 (Correct)

**Which table do not have Fail-safe period ? Select all that apply**

Options are :

- Permanent
- Temporary (*Correct*)
- Transient (*Correct*)
- External
- None of the above

**Answer :**

\*Temporary (Correct)  
\*Transient (Correct)

**It is possible to the latest ODBC connector for use with Snowflake through Snowflake Web UI console. TRUE or FALSE ?**

Options are :

- TRUE (*Correct*)
- FALSE

**Answer :**

\*TRUE (Correct)

**In which situation would you 'scale up' your warehouse ? Select the best option**

Options are :

- To make some credit economy when the warehouse is underutilized
- To provide more resources for a complex queries, to improve performance (*Correct*)
- To support multi-concurrent users or queries

- *None of the above*

**Answer :**

\*To provide more resources for a complex queries, to improve performance (Correct)

## Snowflake Certification Preparation Set 9

**Which of the following statements about Snowflake are true?**

Options are :

- *Storage can increase or decrease without any effect on virtual warehouse sizes.*  
*(Correct)*
- *Cloud services can be increased to multiple regions without any user intervention*
- *Two Virtual Warehouses can access the same data at the same time without causing contention issues*  
*(Correct)*
- *Virtual Warehouses have no effect on compute costs*
- *Compute can be scaled up, down, out, or in and there is no effect on storage used*  
*(Correct)*

**Answer :**

\*Storage can increase or decrease without any effect on virtual warehouse sizes.  
(Correct)

\*Two Virtual Warehouses can access the same data at the same time without causing contention issues (Correct)

\*Compute can be scaled up, down, out, or in and there is no effect on storage used  
(Correct)

**What does "Query Load" measure ? Select the best option**

Options are :

- *average number of queries that were running within a specific interval*
- *average number of queries that were queued within a specific interval*
- *average number of queries that were running or queued within a specific interval*  
*(Correct)*
- *None of the above*

**Answer :**

\*average number of queries that were running or queued within a specific interval  
(Correct)

**Which installment option version of Snowflake are NOT available ? Select all that apply.**

Options are :

- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure)*
- *Microsoft Cloud Native Accounts (Correct)*
- *Hybrid On-Premise + Cloud Installation (Correct)*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure)*
- *Enterprise In-House VPC Installation (Correct)*
- *On-Premise Custom Installation (Correct)*

**Answer :**

- \*Microsoft Cloud Native Accounts (Correct)
- \*Hybrid On-Premise + Cloud Installation (Correct)
- \*Enterprise In-House VPC Installation (Correct)
- \*On-Premise Custom Installation (Correct)

### [Snowflake Certification Preparation Set 5](#)

**A Snowflake session can NOT have more than one virtual warehouse. TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**How long does Bulk data load 'Load history' last ?**

Options are :

- *14*
- *28*
- *32*
- *64 (Correct)*

**Answer :**

\*64 (Correct)

**Which scaling policy, starts additional cluster only if the system estimates there is enough query load to keep the cluster busy for at least 6 minutes.**

Options are :

- *Economy (Correct)*
- *Standard*

**Answer :**

\*Economy (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 23

**What are the elements of a worksheet context ?**

Options are :

- *Role, Database, Schema, Warehouse (Correct)*
- *User, Role, Database, Schema*
- *Account, User, Role, Warehouse*
- *Acount, Database, Schema, Warehouse*

**Answer :**

\*Role, Database, Schema, Warehouse (Correct)

**A database specialist is planning to use Snowflake. He wants to know which of the following Snowflake objects can be cloned? (Choose four.)**

Options are :

- *Tables (Correct)*
- *Named File Formats (Correct)*
- *Schemas (Correct)*
- *Shares*
- *Databases (Correct)*
- *Users*

**Answer :**

- \*Tables (Correct)
- \*Named File Formats (Correct)
- \*Schemas (Correct)
- \*Databases (Correct)

**Snowflake allows its customers to directly access the micro-partition files that make up its tables.  
Is this statement True or False?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

### Tackling Snowflake Certification - Practice Questions Set 4

**When a Snwoflake Pipe object is recreated using the CREATE OR REPLACE PIPE command:**

Options are :

- *The Pipe load history is reset to empty*
- *The REFRESH parameter is set to TRUE (Correct)*
- *Previously loaded files will be ignored*
- *All of the above*

**Answer :**

- \**The REFRESH parameter is set to TRUE (Correct)*

**Which of the following are best practices for Snowflake users with the  
SYSADMIN/ACCOUNTADMIN roles? Choose 3 answers**

Options are :

- *Their default role should be set to SYSADMIN (the lower of the two) (Correct)*
- *They should not set up multi\_Factor Authentication (MFA)?as administrator they may need to change the MFA settings and those enrolled in MFA are unable to do so*
- *They should only access and step into the ACCOUNTADMIN role temporarily, as needed to complete a specific operation (Correct)*

- They should ensure all database objects in the account are owned by the ACCOUNTADMIN role
- They should use the SYSADMIN role to perform administrative work on database objects (Correct)

**Answer :**

- \*Their default role should be set to SYSADMIN (the lower of the two) (Correct)
- \*They should only access and step into the ACCOUNTADMIN role temporarily, as needed to complete a specific operation (Correct)
- \*They should use the SYSADMIN role to perform administrative work on database objects (Correct)

**If a Snowflake Small Warehouse is made up of 2 servers/cluster, How many servers/cluster make up a Medium Warehouse in Snowflake? Warehouse?**

Options are :

- 4 (Correct)
- 16
- 32
- 128

**Answer :**

- \*4 (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**For a Snowflake multi-cluster Warehouse, the number of credits billed is calculated based on?**

Options are :

- The number of queries that ran using the Warehouse.
- The size of the Warehouse and the number of clusters that ran within a given time period. (Correct)
- The size of the Warehouse and the maximum number of clusters configured for the Warehouse.
- The number of users who accessed the Warehouse.

**Answer :**

- \*The size of the Warehouse and the number of clusters that ran within a given time period. (Correct)

**If auto-suspend is enable for a Snowflake Virtual Warehouse, the Warehouse is automatically suspended when:**

Options are :

- All Snowflake sessions using the warehouse are terminated.
- The last query using the warehouse completes.
- There are no users loaned into Snowflake.
- The Warehouse is inactive for a specified period of time. (Correct)

**Answer :**

\*The Warehouse is inactive for a specified period of time. (Correct)

**All Snowflake table types include fail-safe storage. Is this statement True or False?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 7

**Which statements are true about Snowflake micro-partitions? Choose 2 answers**

Options are :

- They are approximately 16MB in size (Correct)
- They are stored compressed only if COMPRESS=TRUE on Table
- They are Immutable (Correct)
- They are only encrypted in the Enterprise edition and above

**Answer :**

\*They are approximately 16MB in size (Correct)

\*They are Immutable (Correct)

**Which of the following are true about Snowflake multi-cluster Warehouses? Select all that apply**

Options are :

- *multi-cluster Warehouse can add clusters automatically based on query activity (Correct)*
- *A multi-cluster Warehouse can automatically turn itself off after a period of inactivity (Correct)*
- *A multi-cluster Warehouse can scale down when query activity slows (Correct)*
- *A multi-cluster Warehouse can automatically turn itself on when a query is executed against it (Correct)*

**Answer :**

\*multi-cluster Warehouse can add clusters automatically based on query activity (Correct)  
\*A multi-cluster Warehouse can automatically turn itself off after a period of inactivity (Correct)  
\*A multi-cluster Warehouse can scale down when query activity slows (Correct)  
\*A multi-cluster Warehouse can automatically turn itself on when a query is executed against it (Correct)

**Which of the following are Snowflake operations that require a Virtual Warehouse to complete, assuming no queries have been executed previously? Choose 3 answers**

Options are :

- *MIN(< < column value>>)*
- *COPY (Correct)*
- *SUM(<< column value >>) (Correct)*
- *UPDATE (Correct)*

**Answer :**

\*COPY (Correct)  
\*SUM(<< column value >>) (Correct)  
\*UPDATE (Correct)

### [Snowflake SnowPro Core Certification Practice Test Set 3](#)

**It is possible to query data from a Snowflake named External stage without loading the data into Snowflake. Is this statement True or False?**

Options are :

- *TRUE (Correct)*

- FALSE

**Answer :**

\*TRUE (Correct)

**Which of following statements are correct about Snowflake PUT command to load data? (Choose two.)**

Options are :

- Automatically creates a File Format object
- Automatically uses the last Stage created
- Automatically compresses files using Gzip (Correct)
- Automatically encrypts files (Correct)

**Answer :**

\*Automatically compresses files using Gzip (Correct)

\*Automatically encrypts files (Correct)

**When can a Snowflake Virtual Warehouse start running queries?**

Options are :

- 12am-5am
- Only during administrator defined time slots
- When its provisioning is complete (Correct)
- After replication

**Answer :**

\*When its provisioning is complete (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 4

**Increasing the maximum number of clusters in a Snowflake Multi-Cluster Warehouse is an example of:**

Options are :

- Scaling rhythmically
- Scaling max

- *Scaling out (Correct)*
- *Scaling Up*

**Answer :**

\*Scaling out (Correct)

**When scaling up Snowflake Virtual Warehouse by increasing Virtual Warehouse size, you are primarily scaling for improved: Select one.**

Options are :

- *Concurrency*
- *Performance (Correct)*

**Answer :**

\*Performance (Correct)

**When data share is established between a Snowflake Data Provider and a data Consumer, the Data Consumer can extend that data share to other Data Consumers. Is this statement True or False?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 1

**Snowflake enforces unique, primary key, and foreign key constraints during DML operations. Is this statement True or false?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following are common use cases for Snowflake zero-copy cloning? Choose 3 answers**

Options are :

- *Quick provisioning of Dev and Test/QA environments (Correct)*
- *Data backups (Correct)*
- *Point in time snapshots (Correct)*
- *Performance optimization*

**Answer :**

\*Quick provisioning of Dev and Test/QA environments (Correct)  
\*Data backups (Correct)  
\*Point in time snapshots (Correct)

**A Snowflake Virtual Warehouse's auto-suspend and auto-resume settings apply to:**

Options are :

- *The primary duster in the virtual warehouse*
- *The entire Virtual Warehouse (Correct)*
- *The database the Virtual Warehouse resides in*
- *The queries currently being run by the Virtual Warehouse*

**Answer :**

\*The entire Virtual Warehouse (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 1

**The FLATTEN function is used to query which type of data in Snowflake?**

Options are :

- *Structured data*
- *Semi-structured data (Correct)*
- *Both of the above*
- *None of the above*

**Answer :**

\*Semi-structured data (Correct)

**Which of the following statements are true of Snowflake data loading? Choose 3 answers**

Options are :

- *VARIANT "null" values are not the same as SQL Null values (Correct)*
- *It is recommended to do frequent, single row DMLS*
- *It is recommended to validate the data before loading into the Snowflake target table (Correct)*
- *It is recommended to use staging tables to manage MERGE statements (Correct)*

**Answer :**

- \*VARIANT "null" values are not the same as SQL Null values (Correct)
- \*It is recommended to validate the data before loading into the Snowflake target table (Correct)
- \*It is recommended to use staging tables to manage MERGE statements (Correct)

**True or False:** When Snowflake is configured to use Single Sign-on (sso), Snowflake receive the usernames and credentials from the sso service and loads them into the customer's Snowflake account.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

### Snowflake SnowPro Core Certification Practice Tests Set 15

**Which of the following are valid approaches for loading data into a snowflake table? select all the below that apply.**

Options are :

- *Bulk copy from an External Stage (Correct)*
- *Continuous load using Snowpipe REST API (Correct)*
- *The Snowflake Web Interface (UT) data loading wizard (Correct)*
- *Bulk copy from an Internal Stage (Correct)*

**Answer :**

- \**Bulk copy from an External Stage (Correct)*
- \**Continuous load using Snowpipe REST API (Correct)*

\*The Snowflake Web Interface (UT) data loading wizard (Correct)

\*Bulk copy from an Internal Stage (Correct)

**Which of the following accurately represents how a table fits into Snowflake's logical container hierarchy?**

Options are :

- *Account -> Schema -> Database -> Table*
- *Account -> Database -> Schema -> Table (Correct)*
- *Database -> Schema -> Table -> Account*
- *Database -> Table -> Schema -> Account*

**Answer :**

\*Account -> Database -> Schema -> Table (Correct)

**What is the maximum compressed row size in Snowflake?**

Options are :

- *8KB*
- *16MB (Correct)*
- *50MB*
- *4000GB*

**Answer :**

\*16MB (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 35**

**Which of the following items does Snowflake Cloud services Layer manage? Choose 4 answers**

Options are :

- *user authentication (Correct)*
- *Metadata (Correct)*
- *Query compilation and optimization (Correct)*
- *external blob storage*

- *Data security (Correct)*

**Answer :**

- \*user authentication (Correct)
- \*Metadata (Correct)
- \*Query compilation and optimization (Correct)
- \*Data security (Correct)

**It is possible to unload structured data to semi-structured formats such as JSON and parquet in Snowflake environment. Is this statement True or False?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**Which type of table corresponds to a single Snowflake session?**

Options are :

- *Temporary (Correct)*
- *Translent*
- *Provisional*
- *Permanent*

**Answer :**

- \*Temporary (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 31**

**Which of the following statements is true of Snowflake zero-copy cloning?**

Options are :

- *Zero-copy clones objects inherit*
- *All zero-copy clone objects inherit the privileges of their original objects*
- *Zero-copy cloning is licensed as an additional Snowflake feature*

- At the instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)

**Answer :**

\*At the instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)

**Which of the following connectors are available in the Downloads section of the Snowflake web Interface (UI) Choose 2 answers**

Options are :

- SnowSQL (Correct)
- ODBC (Correct)
- R
- HIVE

**Answer :**

\*SnowSQL (Correct)

\*ODBC (Correct)

**True or False: Each worksheet in the Snowflake Web Interface (UI) can be associated with different roles, databases, schemas, and Virtual Warehouses.**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

### Snowflake Certification Preparation Set 8

**What happens when a Snowflake Data Provider revokes privileges to a Share on an object in their source database?**

Options are :

- The object immediately becomes unavailable for all Data Consumers (Correct)
- Any additional data arriving after this point in time will not be visible to Data Consumers

- The Data Consumers stop seeing data updates and become responsible for storage charges for the object
- A static copy of the object at the time the privilege was revoked is created In the Data Consumers accounts

**Answer :**

\*The object immediately becomes unavailable for all Data Consumers (Correct)

**Which of the following languages can be used to implement Snowflake User Defined Functions (UDFs)? Choose 2 answers**

Options are :

- Java (Correct)
- JavaScript (Correct)
- SQL (Correct)
- Python
- C#

**Answer :**

\*Java (Correct)

\*JavaScript (Correct)

\*SQL (Correct)

**Each incremental increase in Snowflake Virtual Warehouse size (e.g. Medium to Large) generally results in what?**

Options are :

- More micro-partitions
- Better query scheduling
- Double the numbers of servers In the compute duster (Correct)
- Higher storage costs

**Answer :**

\*Double the numbers of servers In the compute duster (Correct)

**How does Snowflake billed for compute ?**

Options are :

- *Per hour, with one-minute minimum*
- *Per second, with one-minute minimum (Correct)*
- *Per day*
- *Per month*

**Answer :**

\*Per second, with one-minute minimum (Correct)

**What Snowflake objects can be shared ? (Select all that apply)**

Options are :

- *Tables (Correct)*
- *Views*
- *Roles*
- *Secure UDFs (Correct)*
- *Materialized views*
- *None of the above*

**Answer :**

\*Tables (Correct)

\*Secure UDFs (Correct)

**Which cloud platforms can a Snowflake account be hosted on ?**

Options are :

- *Google Cloud Platform (GCP) (Correct)*
- *OVHcloud*
- *Microsoft Azure (Azure) (Correct)*
- *Alibaba Cloud*
- *Amazon Web Services (AWS) (Correct)*
- *None of the above*

**Answer :**

\*Google Cloud Platform (GCP) (Correct)

\*Microsoft Azure (Azure) (Correct)  
\*Amazon Web Services (AWS) (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 2

**What is called the action of adding more servers to the same cluster of a virtual warehouse ?**

Options are :

- *Scale UP (Correct)*
- *Scale DOWN*
- *Scale IN*
- *Scale OUT*
- *None of the above*

**Answer :**

\*Scale UP (Correct)

**Snowflake can automatically scale your warehouse UP when needed (make your warehouse larger). TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**A client has an important number of users that access a BI report every working day at 9AM. Every time a user refreshes the report it requests Snowflake for the data. The report performances are not optimal.\n What suggestion could you make to improve the performance ? Select the best option**

Options are :

- *Create a single warehouse for each user*
- *Regroup the users by age and associate a warehouse to each group*
- *Use a multi-cluster warehouse (Correct)*
- *None of the above*

**Answer :**

\*Use a multi-cluster warehouse (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 2

**Select the two terms that refer to the same Snowflake layer ?**

Options are :

- *Metadata Layer*
- *Virtual Storage Layer*
- *Catalog Layer*
- *Data Layer (Correct)*
- *Storage Layer (Correct)*

**Answer :**

\*Data Layer (Correct)

\*Storage Layer (Correct)

**Metadata is created when you upload data into a Snowflake internal. TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**How can you disable < auto-suspend > option for a virtual warehouse**

Options are :

- *Choose 'Never' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI (Correct)*
- *Choose 'Disable' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI*
- *ALTER WAREHOUSE set AUTO\_SUSPEND = 0 (Correct)*
- *ALTER WAREHOUSE set AUTO\_SUSPEND = NULL (Correct)*

- *ALTER WAREHOUSE set AUTO\_SUSPEND = -1*
- *None of the above*

**Answer :**

\*Choose 'Never' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI (Correct)  
\**ALTER WAREHOUSE set AUTO\_SUSPEND = 0* (Correct)  
\**ALTER WAREHOUSE set AUTO\_SUSPEND = NULL* (Correct)

### Snowflake Certification Preparation Set 4

**Select the three terms referring to the same layer.**

Options are :

- *Compute Layer (Correct)*
- *Virtual Warehouse Layer (Correct)*
- *Catalog Layer*
- *Metadata Layer*
- *Query Processing Layer (Correct)*

**Answer :**

\*Compute Layer (Correct)  
\*Virtual Warehouse Layer (Correct)  
\*Query Processing Layer (Correct)

**Which transformations can NOT be applied on data while loading it into a table using COPY command ? (Select all that apply)**

Options are :

- *Column reordering*
- *Aggregation (Correct)*
- *Casts*
- *Filtering (Correct)*
- *Column omission*
- *Truncating text strings that exceed the target column length*

**Answer :**

- \*Aggregation (Correct)
- \*Filtering (Correct)

**What type of view is most like a table**

Options are :

- *Non-materialized views*
- *Materialized views (Correct)*

**Answer :**

- \*Materialized views (Correct)

### Snowflake Certification Preparation Set 19

**What does this cron expression (USING CRON 0 0-23 \* \* \*) express ?**

Options are :

- *Every minute of every hour of every day*
- *At 00:00 am on the 1st of every month*
- *The first minute of every hour of every day (Correct)*
- *At 00:00 am on the 1st of every month*

**Answer :**

- \*The first minute of every hour of every day (Correct)

**In which situation would you 'scale up' your warehouse ? Select the best option**

Options are :

- *To make some credit economy when the warehouse is underutilized*
- *To provide more resources for a complex queries, to improve performance (Correct)*
- *To support multi-concurrent users or queries*
- *None of the above*

**Answer :**

- \*To provide more resources for a complex queries, to improve performance (Correct)

**Which scaling policy, starts additional cluster immediately when either a query is queued or the system detects that there is one more query than the currently-running clusters can execute.**

Options are :

- *Economy*
- *Standard (Correct)*

**Answer :**

\*Standard (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 26

**Snowflake internal named stages can be cloned. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Select the three Snowpipe REST API endpoints.**

Options are :

- *insertFiles (Correct)*
- *insertData*
- *insertReport (Correct)*
- *loadHistoryScan (Correct)*
- *loadHistory*

**Answer :**

\*insertFiles (Correct)  
\*insertReport (Correct)  
\*loadHistoryScan (Correct)

**Which of the following Snowflake objects allows to query data without loading it first inside Snowflake ?**

Options are :

- *Snowpipe*
- *COPY*
- *External table (Correct)*
- *Virtual table*

**Answer :**

\*External table (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 42](#)

**What is the region identifier inside this Snowflake URL << https://jz96798.us-east-1.azure.snowflakecomputing.com >> ?**

Options are :

- *jz96798.us-east-1.azure*
- *us-east-1.azure (Correct)*
- *jz96798*
- *snowflakecomputing.com*

**Answer :**

\*us-east-1.azure (Correct)

**Scaling OUT is an automated process (e.g. Min clusters to max clusters). TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**A client ask your support to choose a geographic deployment region for a Snowflake account. What factors should you consider ? Select the two best options.**

Options are :

- Proximity to the point of service (Correct)
- Additional fees charged for regions with geo-political unrest
- Number of availability zones within a region (Correct)
- End-user perceptions of glamorous or trendy geographic locations

**Answer :**

\*Proximity to the point of service (Correct)  
\*Number of availability zones within a region (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 35**

**With enterprise edition, what is the longest time you can specify for Time Travel?**

Options are :

- 1 day
- 30 days
- 90 days (Correct)
- 365 days
- None of the above

**Answer :**

\*90 days (Correct)

**What kind of stream can be created on external table ?**

Options are :

- Standard
- Append-only
- Insert-only (Correct)
- None of the above

**Answer :**

\*Insert-only (Correct)

**What attributes make Snowflake a true SaaS solution?**

Options are :

- *No hardware to install or configure (Correct)*
- *No maintenance upgrades or patches to install (Correct)*
- *Transparent releases don't require user intervention (Correct)*
- *No creation of user accounts or roles is required*
- *No query processing costs*
- *No data storage costs*

**Answer :**

- \*No hardware to install or configure (Correct)
- \*No maintenance upgrades or patches to install (Correct)
- \*Transparent releases don't require user intervention (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 9

**Multiple warehouses can access the same data simultaneously without resource contention issues.  
TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**To improve the query performance, which of the below techniques can be used in snowflake?**

Options are :

- *Query hints*
- *Distribution keys*
- *Indexes*
- *Cluster keys/Reclustering (Correct)*

**Answer :**

- \*Cluster keys/Reclustering (Correct)

**Snowpipe is recommended to load data in which of the below scenario?**

Options are :

- Huge volume of data generated as part of batch schedule
- When data is in small volume with high frequency (Correct)

**Answer :**

\*When data is in small volume with high frequency (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 33**

**To download files from the stage location loaded through COPY INTO <LOCATION> command, you will use**

Options are :

- COPY INTO
- UNLOAD
- GET (Correct)

**Answer :**

\*GET (Correct)

**Select the options that are TRUE about 'Snowflake Partner Connect' partners.**

Options are :

- Can be connected from the WEB UI (Correct)
- Includes a partner trial account signup (Correct)
- Includes automated role, user and staging database setup (Correct)
- None of the above
- Connect with snowflake through a wizard (Correct)

**Answer :**

\*Can be connected from the WEB UI (Correct)

\*Includes a partner trial account signup (Correct)

\*Includes automated role, user and staging database setup (Correct)

\*Connect with snowflake through a wizard (Correct)

**When cloning a table, if the COPY GRANTS keywords are not included in the CREATE <object> statement, then the new object does not inherit any explicit access privileges granted on the original table but does inherit any future grants defined for the object type in the schema**

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

**Which of the below are valid SNOWPIPE REST APIs. Select three.**

Options are :

- *loadHistoryScan (Correct)*
- *loadHistoryProgress*
- *insertFiles (Correct)*
- *insertReport (Correct)*

**Answer :**

\**loadHistoryScan (Correct)*

\**insertFiles (Correct)*

\**insertReport (Correct)*

**Is standard Time Travel (up to 1 day) for accessing/restoring modified and deleted data available in all editions of snowflake?**

Options are :

- Yes (Correct)
- No

**Answer :**

\*Yes (Correct)

**In your snowflake environment, you have a medium warehouse. The medium warehouse is used by business to run adhoc queries. The warehouse has auto suspend set at 15 minutes. You have noticed that all of the queries that run on this warehouse finishes within a minute. What will you do to optimize cost of compute in this case?**

Options are :

- You will tell the users that they do not need snowflake to run their queries, they should load their data in another on premise database and run query from there
- Since all the queries are completing in a minute, you should delete this warehouse and tell the users to use another existing warehouse
- You will recommend to reduce the auto suspend time to 1 minute (Correct)

**Answer :**

\*You will recommend to reduce the auto suspend time to 1 minute (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

A successful response from the SNOWPIPE insertFiles API means that the files are ingested.

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

Looker is a snowflake business intelligence partner

Options are :

- FALSE
- TRUE (Correct)

**Answer :**

\*TRUE (Correct)

What is the recommended approach for making a variant column accessible in a BI tool?

Options are :

- pre-defined mapping
- BI tool cannot access json
- creating a view (Correct)

- Leveraging a json parser

**Answer :**

\*creating a view (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 21

**Which constraints can be enforced in snowflake?**

Options are :

- NOT NULL constraint (Correct)
- UNIQUE Constraint
- Referential integrity constraints

**Answer :**

\*NOT NULL constraint (Correct)

**You ran a query in snowflake and went to query history tab. The query history shows you the below columns 1. QueryID 2. SQL TEXT 3. WAREHOUSE NAME 4. WAREHOUSE SIZE 5. SESSION ID 6. START TIME 7. END TIME Which of the above column will indicate if an compute cost was incurred to run the query?**

Options are :

- CREDIT
- WAREHOUSE SIZE (Correct)
- WAREHOUSE NAME
- SESSION ID

**Answer :**

\*WAREHOUSE SIZE (Correct)

**In which editions of snowflake users get extended Time Travel option (up to 90 days)?**

Options are :

- Standard
- VPS (Correct)
- Enterprise (Correct)

- Business Critical (Correct)

**Answer :**

- \*VPS (Correct)
- \*Enterprise (Correct)
- \*Business Critical (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 13

**Which of the below are considered as best practices while loading data into snowflake?**

Options are :

- Compress the source files (Correct)
- Split large files into smaller files (Correct)
- If format is in CSV, convert them into ORC
- Isolate data loading workload into its own virtual warehouse (Correct)

**Answer :**

- \*Compress the source files (Correct)
- \*Split large files into smaller files (Correct)
- \*Isolate data loading workload into its own virtual warehouse (Correct)

**Is it possible to re-size virtual warehouse once created?**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

- \*TRUE (Correct)

**Which of the below activities can be done by leveraging time-travel feature of snowflake?**

Options are :

- To backup data for compliance and regulatory requirements
- Restoring data-related objects (tables, schemas, and databases) that may have been accidentally or intentionally deleted (Correct)
- Analyzing data usage/manipulation over specified periods of time (Correct)

- Duplicating and backing up data from key points in the past (Correct)

**Answer :**

\*Restoring data-related objects (tables, schemas, and databases) that may have been accidentally or intentionally deleted (Correct)

\*Analyzing data usage/manipulation over specified periods of time (Correct)

\*Duplicating and backing up data from key points in the past (Correct)

### **Snowflake SnowPro Core Certification Exam Questions Set 5**

**Select the options that are true for snowflake kafka connector**

Options are :

- Kafka connector guarantees reprocessing of messages
- Kafka connector guarantees that rows are inserted in the order
- Kafka connector guarantees exactly-once delivery (Correct)

**Answer :**

\*Kafka connector guarantees exactly-once delivery (Correct)

**Choose which architecture is followed by Snowflake:**

Options are :

- Multi-Cluster Shared disk architecture
- Shared disk architecture
- Shared nothing architecture
- Hybrid of traditional shared-disk and shared-nothing database architectures (Correct)

**Answer :**

\*Hybrid of traditional shared-disk and shared-nothing database architectures (Correct)

**Snowflake adds the system defined roles to each new account. Check the system defined roles below.**

Options are :

- SUPERUSER
- SECURITYADMIN (Correct)

- *ACCOUNTADMIN* (*Correct*)
- *PUBLIC* (*Correct*)
- *USERADMIN* (*Correct*)
- *SYSADMIN* (*Correct*)

**Answer :**

\**SECURITYADMIN* (*Correct*)  
 \**ACCOUNTADMIN* (*Correct*)  
 \**PUBLIC* (*Correct*)  
 \**USERADMIN* (*Correct*)  
 \**SYSADMIN* (*Correct*)

**Which of the following commands cannot be executed from worksheets?**

Options are :

- *PUT* (*Correct*)
- *SHOW*
- *LIST*
- *GET* (*Correct*)

**Answer :**

\**PUT* (*Correct*)  
 \**GET* (*Correct*)

**How do you set a return value in a task? Hint: Tasks are used to automate sql statements either by running at a schedule or triggering it from an event. Return value would help to understand if task was completed or has raised error.**

Options are :

- *create task set\_return\_value warehouse = return\_task\_wh schedule = '1 minute' as call set\_return\_value('This task has executed successfully');*
- *create task set\_return\_value warehouse = return\_task\_wh schedule = '1 minute' as call system\$set\_return\_code('This task has executed successfully');*
- *create task set\_return\_value warehouse = return\_task\_wh schedule = '1 minute' as call system\$set\_return\_value('This task has executed successfully');* (*Correct*)

**Answer :**

\**create task set\_return\_value warehouse = return\_task\_wh schedule = '1 minute' as call system\$set\_return\_value('This task has executed successfully');* (*Correct*)

## How SQL queries can be monitored on Snowflake?

Options are :

- By clicking on History tab on Snowflake and access all queries which are executed in last 14 days (Correct)
- Through snowflake python connector
- By connecting with a visualization tool

### Answer :

\*By clicking on History tab on Snowflake and access all queries which are executed in last 14 days (Correct)

## Tackling Snowflake Certification - Practice Questions Set 3

### A user can be defaulted to a role which user does not have access to

Options are :

- TRUE
- FALSE (Correct)

### Answer :

\*FALSE (Correct)

You can get customer dedicated virtual warehouse, which is the lowest snowflake edition that you should opt for

Options are :

- Business Critical
- Standard (Correct)
- Enterprise

### Answer :

\*Standard (Correct)

What technique does snowflake use to limit the number of micro-partitions scanned by each query?

Options are :

- *Pruning (Correct)*
- *B-tree*
- *Indexing*
- *Map Reduce*

**Answer :**

\*Pruning (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 39

**Which Snowflake role can see billing on the account, administer users, create reader accounts and manage all database objects?**

Options are :

- *SECURITYADMIN*
- *SUPERROLE*
- *SYSADMIN*
- *ACCOUNTADMIN (Correct)*

**Answer :**

\*ACCOUNTADMIN (Correct)

**When unloading data into multiple files, which copy option is used to specify the maximum size of each file created**

Options are :

- *MAX\_FILE\_BYTES*
- *MAX\_SIZE*
- *MAX\_FILE\_SIZE (Correct)*

**Answer :**

\*MAX\_FILE\_SIZE (Correct)

**Which of the below security features are supported by Snowflake?**

Options are :

- *TRI-SECRET SECURE ENCRYPTION (Correct)*

- *NONE OF THE ABOVE*
- *MULTI-FACTOR AUTHENTICATION (Correct)*
- *ROLE-BASED ACCESS CONTROL (RBAC) (Correct)*

**Answer :**

- \**TRI-SECRET SECURE ENCRYPTION (Correct)*
- \**MULTI-FACTOR AUTHENTICATION (Correct)*
- \**ROLE-BASED ACCESS CONTROL (RBAC) (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 16

**How to resume a created TASK on snowflake?**

Options are :

- *ALTER TASK sample\_task RESUME; (Correct)*
- *No need to resume, the task is automatically enabled at the time of creation*
- *ALTER TASK sample\_task START;*

**Answer :**

- \**ALTER TASK sample\_task RESUME; (Correct)*

**If you attempt to PUT a file but cannot because a file with the same name already exists in the stage, what can you do to load the file?**

Options are :

- *Set FORCE COPY = TRUE*
- *Set OVERWRITE = TRUE (Correct)*
- *Wait until the existing file is loaded, then retry PUT (Correct)*
- *Rename the file that you want to PUT (Correct)*

**Answer :**

- \**Set OVERWRITE = TRUE (Correct)*
- \**Wait until the existing file is loaded, then retry PUT (Correct)*
- \**Rename the file that you want to PUT (Correct)*

**Scenario:** A transaction is executed in a snowflake session. Due to some reason the session disconnects and the transaction now is in a detached state. The transaction cannot be committed or rolled back. The object on which the transaction was applied is also now locked. If you do not do anything and let snowflake eventually abort the transaction, how long will you need to wait?

Options are :

- 4 hours (Correct)
- 15 minutes
- 60 minutes

**Answer :**

\*4 hours (Correct)

### **Snowflake SnowPro Core Certification Exam Questions Set 6**

**Which of the below are true with respect to snowflake web ui?**

Options are :

- You must use the same role and warehouse for all worksheets
- Each worksheet is a separate session (Correct)
- You will be able to run a 'PUT' command from worksheet
- Each worksheet can have different role, warehouse, schema and database (Correct)

**Answer :**

\*Each worksheet is a separate session (Correct)

\*Each worksheet can have different role, warehouse, schema and database (Correct)

**What does snowflake use for monitoring network traffic and use activity?**

Options are :

- Lacework (Correct)
- Threat Stack
- Sumo logic

**Answer :**

\*Lacework (Correct)

**Which file format option would you use when you want to load a JSON file after removing the outer array structure and load the records into separate table rows?**

Options are :

- *UNSET\_OUTER\_ARRAY*
- *STRIP\_OUTER\_ARRAY (Correct)*
- *DROP\_OUTER\_ARRAY*
- *TRUNCATE\_OUTER\_ARRAY*

**Answer :**

\**STRIP\_OUTER\_ARRAY (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 14

**Auto clustering can be switched off at an account level**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\**FALSE (Correct)*

**Scenario:** You are running your query in the worksheet using a large warehouse named as **LOAD\_WH**. You want to switch the warehouse to **COMPUTE\_WH**. How will you do that? Choose 2 answers

Options are :

- *SET WAREHOUSE CONTEXT = COMPUTE\_WH;*
- *USE WAREHOUSE COMPUTE\_WH; (Correct)*
- *Update the Warehouse field in the worksheet Context Menu located above the worksheet (Correct)*
- *Go to the Worksheet properties page and set the Warehouse field to COMPUTE\_WH*

**Answer :**

\**USE WAREHOUSE COMPUTE\_WH; (Correct)*

\**Update the Warehouse field in the worksheet Context Menu located above the worksheet (Correct)*

**Each time a persisted result for a query is reused, Snowflake resets the 24-hour retention period for the result, up to a maximum of 31 days from the date and time that the query was first executed. After 31 days, the result is purged and the next time the query is submitted, a new result is generated and persisted.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 13

**Snowflake supports multi-factor authentication (MFA) to provide increased login security for users connecting to Snowflake. Which statements are true about MFA security?**

Options are :

- *MFA can be used to connect to SNOWFLAKE using Snowflake JDBC Driver (Correct)*
- *MFA is an integrated feature powered by DUO Security Service (Correct)*
- *MFA Login is designed to connect to snowflake only through the web interface*
- *MFA is not enabled by default. it is available to all, but user need to activate it (Correct)*

**Answer :**

\**MFA can be used to connect to SNOWFLAKE using Snowflake JDBC Driver (Correct)*

\**MFA is an integrated feature powered by DUO Security Service (Correct)*

\**MFA is not enabled by default. it is available to all, but user need to activate it (Correct)*

**Snowflake supports specifying a SELECT statement instead of a table in the COPY INTO <location> command**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which of the given statements are true with respect to SMT(simple message transformation) when neither key.converter or value.converter is set when working with Kafka?**

Options are :

- *regex.router is not supported (Correct)*
- *All SMTs are supported*
- *Most SMTs are supported (Correct)*

**Answer :**

\**regex.router is not supported (Correct)*  
\**Most SMTs are supported (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 42](#)

**When data is staged to a Snowflake internal staging area using the PUT command, the data is automatically encrypted**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Snowflake waits till all the servers are provisioned for a new virtual warehouse, before it executes query on that warehouse**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Snowpipe does not guarantee loading of files in the order that they are staged**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

## SnowFlake SnowPro Advanced Architect Practice Exam Set 7

**Loading data using snowpipe REST API is supported for external stage only**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**What is the expiration time (in days) for Load history of Snowpipe?**

Options are :

- *64*
- *365*
- *30*
- *14 (Correct)*

**Answer :**

\**14 (Correct)*

**How do you remove files from a stage?**

Options are :

- *PURGE*
- *REMOVE (Correct)*
- *DELETE*
- *TRUNCATE*

**Answer :**

\**REMOVE (Correct)*

## Snowflake SnowPro Core Certification Practice Tests Set 16

**How query processing is done in Snowflake?**

Options are :

- *It is done using Storage layer*
- *It is done using Snowflake's Virtual warehouses (Correct)*
- *It is done through an external compute such as Azure virtual machine or AWS EC2*
- *It is done using Cloud services layer*

**Answer :**

\*It is done using Snowflake's Virtual warehouses (Correct)

**Resource monitors can be used to control credit usage for the Snowflake-provided warehouses, including the snowpipe warehouse.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**In a multi-cluster warehouse, you will need to manually scale the warehouses(minimum cluster to maximum cluster)**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

## Snowflake Certification Preparation Set 2

**Which of the two statements are true about the variant data type in SnowFlake?**

Options are :

- *Stored in a separate file format from structured data*
- *Optimized storage based on repeated elements (Correct)*

- Can be queried using json path notation (Correct)
- Requires a custom mapping for each record type

**Answer :**

- \*Optimized storage based on repeated elements (Correct)
- \*Can be queried using json path notation (Correct)

**Select two choices that are correct about roles on snowflake**

Options are :

- Snowflake users has a limit on the number of roles that they can assume
- Snowflake user can have one or more roles (Correct)
- For a particular session, only one role can be active at a given time (Correct)

**Answer :**

- \*Snowflake user can have one or more roles (Correct)
- \*For a particular session, only one role can be active at a given time (Correct)

**If you clone a permanent table(employees) into a transient table(employees\_new) using the below command create transient table employees\_new clone employees copy grants; What will happen to the partitions?**

Options are :

- Only the old partitions will be affected
- All the partitions will be affected
- Old partitions will not be affected, however new partitions added to the clone will follow the transient lifecycle (Correct)

**Answer :**

- \*Old partitions will not be affected, however new partitions added to the clone will follow the transient lifecycle (Correct)

### [\*\*Snowflake SnowPro Core Certification Practice Test Set 5\*\*](#)

**When you use this parameter with SHOW STREAMS, The output also includes an additional dropped\_on column, which displays 1. Date and timestamp (for dropped streams). 2.NULL (for active streams).**

Options are :

- Stream last used value

- *Terse*
- *DROPPED value as 'TRUE'*
- *HISTORY (Correct)*

**Answer :**

\*HISTORY (Correct)

**Which OPTION will you set on the copy if you would like to ignore the load metadata?**

Options are :

- *LOAD\_UNCERTAIN\_FILES*
- *IGNORE\_METADATA*
- *FORCE (Correct)*

**Answer :**

\*FORCE (Correct)

**How is the data storage cost computed for snowflake?**

Options are :

- *Data size stored in compressed format (Correct)*
- *Based on daily average of data stored (Correct)*
- *Data size stored in un-compressed format*
- *Amount of data stored on first day of the month*
- *Amount of data stored on last day of the month*

**Answer :**

\*Data size stored in compressed format (Correct)

\*Based on daily average of data stored (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 27**

**The COPY command is more performant than the INSERT statement?**

Options are :

- *TRUE (Correct)*

- FALSE

**Answer :**

\*TRUE (Correct)

**Can you create internal stage without linking it to the cloud providers?**

Options are :

- No
- Yes (Correct)

**Answer :**

\*Yes (Correct)

**Select which of the given options will incur data storage cost on snowflake?**

Options are :

- Semi-structured data stored on cloud storages
- Data retained as history (Time travel and Failsafe) to enable data recovery (Correct)
- Data stored in tables (Correct)

**Answer :**

\*Data retained as history (Time travel and Failsafe) to enable data recovery (Correct)

\*Data stored in tables (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 30

**Select the options that are true about data sharing through data shares.**

Options are :

- Providers can establish revocable, fine-grained access grants to shares (Correct)
- Shared data is always live, real-time and immediately available to consumers (Correct)
- Data Sharing is supported between accounts in the same or different Snowflake Provider and Region
- There is only one copy of data, which lives in the data provider's account (Correct)

**Answer :**

\*Providers can establish revocable, fine-grained access grants to shares (Correct)

\*Shared data is always live, real-time and immediately available to consumers (Correct)

\*There is only one copy of data, which lives in the data provider's account (Correct)

#### **What is the best way to get the latest ODBC connector to connect with Snowflake?**

Options are :

- Download it from snowflake web UI (Correct)
- Download from external sites available on Google
- There is no ODBC connector available for Snowflake

#### **Answer :**

\*Download it from snowflake web UI (Correct)

**Suppose two data warehouses are running on your snowflake account. Data has been updated in the storage layer using one of the virtual warehouses. After how much time, updated data will be available to the other warehouse?**

Options are :

- 1 Hour
- 15 minutes
- Immediately (Correct)
- 30 minutes

#### **Answer :**

\*Immediately (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 19**

#### **Which one of these is used by snowflake for computing?**

Options are :

- Azure Virtual Machine
- AWS EC2 Instance
- Virtual Warehouses (Correct)

#### **Answer :**

\*Virtual Warehouses (Correct)

**Scenario: Your organization need Column-level security by applying masking policies to columns in tables or views. Which editions will be suitable for this requirement?**

Options are :

- *VPS (Correct)*
- *Enterprise (Correct)*
- *Standard*
- *Business Critical (Correct)*

**Answer :**

\*VPS (Correct)  
\*Enterprise (Correct)  
\*Business Critical (Correct)

**You need snowflake instances in different regions. You will need to have separate account for each region**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What is the column datatype used for ingesting semi-structured data?**

Options are :

- *VARCHAR*
- *BINARY*
- *VARIANT (Correct)*

**Answer :**

\*VARIANT (Correct)

**In which layer of snowflake, the security information is stored?**

Options are :

- *Storage*
- *Cloud Services (Correct)*
- *Compute*

**Answer :**

\*Cloud Services (Correct)

**PUT command in snowflake can upload files from a local folder/directory of a client machines to the below locations:**

Options are :

- *External Stage*
- *Named internal stage (Correct)*
- *Internal stage for a specific table (Correct)*
- *Internal stage for the current user (Correct)*

**Answer :**

\*Named internal stage (Correct)

\*Internal stage for a specific table (Correct)

\*Internal stage for the current user (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 12

**On which key service/s, snowflake pricing model is based upon?**

Options are :

- *Virtual Warehouse (Compute) only*
- *Storage only*
- *Storage, Virtual Warehouse (Compute) and Data Transfer across regions & across clouds (Correct)*
- *Storage and Virtual Warehouse (Compute)*

**Answer :**

\*Storage, Virtual Warehouse (Compute) and Data Transfer across regions & across clouds (Correct)

**Which feature is only available in Virtual Private Snowflake (VPS) edition of snowflake?**

Options are :

- *SOC 2 Type II certification*
- *Automatic encryption of all data*
- *Dedicated metadata store and pool of virtual servers (used in virtual warehouses). (Correct)*

**Answer :**

\*Dedicated metadata store and pool of virtual servers (used in virtual warehouses).  
(Correct)

#### **Metadata columns for staged files can only be queried by name**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

#### **SnowFlake SnowPro Advanced Architect Practice Exam Set 1**

##### **Which of the below statement will you use to recreate a specified object?**

Options are :

- *SHOW SCHEMA*
- *GET\_DDL (Correct)*
- *FETCH\_DDL*
- *GET\_SCHEMA*

**Answer :**

\*GET\_DDL (Correct)

##### **Snowflake recommends to compress your data files when you are loading large data sets.**

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**When a network policy includes values in both the allowed and blocked IP address lists, Snowflake applies the blocked IP address list first.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 5

**In case multiple instances of the kafka connector is started on the same topic or partitions, duplicate records may flow into snowflake table.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Snowflake is a Data Migration/ETL Tool**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**How do you truncate a date (from a timestamp) down to the year, month, and day**

Options are :

- *Split the date into year, month and day*

- Use concatenation of the date part or substring
- Use `date_trunc()` function and execute a query as below: `select to_date('2021-05-30T05:05:23.134-09:00') as "DATE1", date_trunc('YEAR','DATE1') as "TRUNCATED TO YEAR", date_trunc('MONTH','DATE1') as "TRUNCATED TO MONTH", date_trunc('DAY','DATE1') as "TRUNCATED TO DAY"; (Correct)`

**Answer :**

\*Use `date_trunc()` function and execute a query as below: `select to_date('2021-05-30T05:05:23.134-09:00') as "DATE1", date_trunc('YEAR','DATE1') as "TRUNCATED TO YEAR", date_trunc('MONTH','DATE1') as "TRUNCATED TO MONTH", date_trunc('DAY','DATE1') as "TRUNCATED TO DAY"; (Correct)`

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 37](#)

**Client has ODBC or JDBC available in their system but they do not have snowflake drivers. Client will still be able to connect to snowflake?**

Options are :

- FALSE (Correct)
- TRUE

**Answer :**

\*FALSE (Correct)

**Scenario:** In your organization, you use a 2XL warehouse to load data into Snowflake as part of your ETL workload. Your BI team reached out to you and wanted access to the data stored in Snowflake. You gave them access to the tables required for them to query the data. What will be your recommendation with respect to the warehouse that they will need to use?

Options are :

- It is recommended to use a multi-warehouse approach for workload isolation in snowflake. Since you would not like the load workload to interfere with the read workload of the BI team, you will create a separate warehouse for them (Correct)
- Since you already have a 2XL warehouse for loading data, you will let them use that warehouse
- Warehouse is not required for querying data in snowflake

**Answer :**

\*It is recommended to use a multi-warehouse approach for workload isolation in snowflake. Since you would not like the load workload to interfere with the read workload of the BI team, you will create a separate warehouse for them (Correct)

**Which are the key concepts that will need to be considered while loading data into snowflake**

Options are :

- *FILE SIZE*
- *TRANSFORMATION AND ERROR VALIDATION (Correct)*
- *FILE FORMAT (Correct)*
- *STAGE OBJECT (Correct)*

**Answer :**

\*TRANSFORMATION AND ERROR VALIDATION (Correct)  
\*FILE FORMAT (Correct)  
\*STAGE OBJECT (Correct)

#### Snowflake Certification Preparation Set 4

**Which are the two metadata columns that can be used for staged files**

Options are :

- *METADATA\$FILENAME (Correct)*
- *METADATA\$FILEFORMAT*
- *METADATA\$FILE\_ROW\_NUMBER (Correct)*

**Answer :**

\*METADATA\$FILENAME (Correct)  
\*METADATA\$FILE\_ROW\_NUMBER (Correct)

**In snowflake, what are the two major cost categories**

Options are :

- *Snowflake Maintenance*
- *Cloud Services*
- *Data Security*
- *Compute (Correct)*
- *Storage (Correct)*

**Answer :**

\*Compute (Correct)  
\*Storage (Correct)

**In load metadata, Snowflake maintains detailed metadata for each table into which data is loaded. What else is stored in this metadata?**

Options are :

- *Name of each file from which data was loaded (Correct)*
- *File size (Correct)*
- *ETag for the file (Correct)*
- *Timestamp of the last load for the file (Correct)*
- *Number of columns in the file*
- *Number of rows parsed in the file (Correct)*

**Answer :**

\*Name of each file from which data was loaded (Correct)  
\*File size (Correct)  
\*ETag for the file (Correct)  
\*Timestamp of the last load for the file (Correct)  
\*Number of rows parsed in the file (Correct)

### Snowflake Certification Preparation Set 13

**The query profiler view is only available for completed queries.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**Scenario: You had a permanent table named EMPLOYEE but you have dropped the table. Now, you have created another table with the same name EMPLOYEE. What will happen If you execute the UNDROP command?**

Options are :

- *The new table will be renamed with a suffix of v1*
- *UNDROP command will fail (Correct)*
- *It will overwrite the newly created table*

**Answer :**

\*UNDROP command will fail (Correct)

**You have two types of named stages, one is external stage and the other one is internal stage. External stage will always require a cloud storage provider.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 7

**Using COPY INTO <location> command, you can unload data from a table (or query) into the below locations.**

Options are :

- *Named external stage (Correct)*
- *Named internal stage (or table/user stage) (Correct)*
- *External location (Amazon S3, Google Cloud Storage, or Microsoft Azure) (Correct)*
- *Local drive/disk*

**Answer :**

\*Named external stage (Correct)

\*Named internal stage (or table/user stage) (Correct)

\*External location (Amazon S3, Google Cloud Storage, or Microsoft Azure) (Correct)

**Which are the two variant columns available in a snowflake table loaded by kafka connector?  
NOTE: Using Kafka connector, data is loaded in form of JSON in both VARIANT columns**

Options are :

- *RECORD\_METADATA (Correct)*
- *RECORD\_TIMESTAMP*
- *RECORD\_CONTENT (Correct)*
- *RECORD\_KEY*

**Answer :**

- \*RECORD\_METADATA (Correct)
- \*RECORD\_CONTENT (Correct)

**What are the two mechanisms to detect if new stage file is there in a snowpipe?**

Options are :

- *Calling the custom APIs exposed through AWS EKS*
- *Calling Snowpipe REST endpoints (Correct)*
- *Automating Snowpipe using cloud messaging (Correct)*

**Answer :**

- \*Calling Snowpipe REST endpoints (Correct)
- \*Automating Snowpipe using cloud messaging (Correct)

## Tackling Snowflake Certification - Practice Questions Set 2

**What can you expect if the filters specified in an INFORMATION\_SCHEMA query are not sufficiently selective?**

Options are :

- *WILL SWITCH TO A LARGE WAREHOUSE*
- *WILL SHOW RESULTS WITH AN WARNING*
- *AN ERROR (Correct)*

**Answer :**

- \*AN ERROR (Correct)

**Scenario: When you are trying to set a variable by using the set variable function SET TMP\_VAR ='XXXX', you got an error as below: Assignment to 'TMP\_VAR' not done because value exceeds size limit of variables. Its size is 312; the limit is 256. What is the reason for this error?**

Options are :

- *Your warehouse is not big enough to accomodate this*
- *This is a temporary error due to insufficient memory in VM instance*
- *The size of string or binary variables is limited to 256 bytes (Correct)*

**Answer :**

- \*The size of string or binary variables is limited to 256 bytes (Correct)

**The VALIDATION\_MODE parameter does not support COPY statements that transform data during a load.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### [Snowflake SnowPro Core Certification Practice Tests Set 13](#)

**Only named stages (internal or external) and user stages are supported for COPY transformations**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**In which cloud providers, OVERWRITE option for PUT is supported?**

Options are :

- *Amazon Web Services (Correct)*
- *Google cloud platform*
- *Microsoft Azure (Correct)*

**Answer :**

\**Amazon Web Services (Correct)*

\**Microsoft Azure (Correct)*

**The FLATTEN command in snowflake have two versions. One version uses a join and the other version uses an object keyword. Please select two words that represent the options used with the command?**

Options are :

- *TABLE (Correct)*

- *TRY\_CAST*
- *LATERAL (Correct)*
- *OBJECT\_CONSTRUCT*

**Answer :**

\*TABLE (Correct)  
\*LATERAL (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 8

**Snowflake has scalar and tabular user-defined functions (UDFs), with support for both SQL and JavaScript. Is the statement True or False?**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**Data can be loaded using PUT command through worksheets in snowflake web UI**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**To help control the usage of credits in Auto-scale mode, Snowflake provides a property, \_\_\_\_\_, that determines the scaling policy to use when automatically starting or shutting down additional clusters. Select the property.**

Options are :

- *Max no of cluster*
- *Min no of cluster*
- *Scaling Policy (Correct)*
- *Auto Scale*

**Answer :**

\*Scaling Policy (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

**Is snowflake available on premise and can be downloaded from snowflake's website?**

Options are :

- Yes
- No (Correct)

**Answer :**

\*No (Correct)

**Snowpipe is a serverless service**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**You are hired as snowflake architect by an organization and you want to design their virtual warehouse strategy. Which one of the strategy will you recommend?**

Options are :

- Recommendation would be to use multi-warehouse strategy. Separate warehouses will be used to perform each type of activity such as daily warehouse load, adhoc queries, reporting, etc (Correct)
- You do not need to use warehouse for loading data. It is only required for query processing
- Multiple warehouses would incur high cost. You would keep just one virtual warehouse for all types of workloads

**Answer :**

\*Recommendation would be to use multi-warehouse strategy. Separate warehouses will be used to perform each type of activity such as daily warehouse load, adhoc queries, reporting, etc (Correct)

## Snowflake Certification Preparation Set 7

**Which of the below terms is associated with compute layer of snowflake architecture?**

Options are :

- *Query planning*
- *Query optimization*
- *Query processing (Correct)*

**Answer :**

\*Query processing (Correct)

**By executing the 'SHOW TABLES' command, we can list all the tables in all the schemas even if we do not have access to all the tables**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Column level security in Snowflake is applied using masking policies to one or more column of a table or view. Which two features are related to column level security?**

Options are :

- *External Tokenization (Correct)*
- *Data Loss prevention (DLP)*
- *Dynamic Data masking (Correct)*

**Answer :**

\*External Tokenization (Correct)

\*Dynamic Data masking (Correct)

## Snowflake Practice Questions - SnowPro Core Certified Set 2

**Does snowflake allow to create multiple virtual warehouses?**

Options are :

- *No*
- *Yes (Correct)*

**Answer :**

\*Yes (Correct)

**When a table is cloned, which of the below queries will work on the cloned table?**

Options are :

- *ALL OF THE ABOVE (Correct)*
- *DROP TABLE*
- *SELECT \*FROM*
- *SHOW TABLES LIKE "*

**Answer :**

\*ALL OF THE ABOVE (Correct)

**How often does snowflake release new features?**

Options are :

- *Yearly*
- *Weekly (Correct)*
- *Monthly*
- *Daily*

**Answer :**

\*Weekly (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 2

**A COUNT(\*) and COUNT(<COLUMN\_NAME>) on a table will always give you the same result**

Options are :

- *FALSE (Correct)*

- *TRUE*

**Answer :**

\**FALSE (Correct)*

**A stream is an object used for CDC which itself contains data of base table**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**There is no query performance difference between a column with a maximum length declaration (e.g. VARCHAR(16777216)), and a smaller precision. Still it is recommended to define an appropriate column length because of the below reasons**

Options are :

- *Data loading operations are more likely to detect issues such as columns loaded out of order, e.g. a 100-character string loaded erroneously into a VARCHAR(50) column. Such issues will produce errors. (Correct)*
- *When the column length is unspecified, some third-party tools may anticipate consuming the maximum size value, which can translate into increased client-side memory usage or unusual behavior (Correct)*
- *Data unloading will be performant if appropriate column lengths are defined*

**Answer :**

\**Data loading operations are more likely to detect issues such as columns loaded out of order, e.g. a 100-character string loaded erroneously into a VARCHAR(50) column. Such issues will produce errors. (Correct)*

\**When the column length is unspecified, some third-party tools may anticipate consuming the maximum size value, which can translate into increased client-side memory usage or unusual behavior (Correct)*

**Which are the system defined roles in snowflake?**

Options are :

- *USERADMIN (Correct)*
- *SUPERUSER*

- *ACCOUNTADMIN* (*Correct*)
- *SECURITYADMIN* (*Correct*)
- *POWERUSER*
- *SYSADMIN* (*Correct*)

**Answer :**

\**USERADMIN* (*Correct*)  
\**ACCOUNTADMIN* (*Correct*)  
\**SECURITYADMIN* (*Correct*)  
\**SYSADMIN* (*Correct*)

**Following transformations are not supported in a COPY command. Select three.**

Options are :

- *NVL*
- *FLATTEN* (*Correct*)
- *JOIN* (*Correct*)
- *GROUP BY* (*Correct*)

**Answer :**

\**FLATTEN* (*Correct*)  
\**JOIN* (*Correct*)  
\**GROUP BY* (*Correct*)

**When working with a cloned table, you can use the below SQL statements**

Options are :

- *ALL OF THE ABOVE* (*Correct*)
- *SHOW*
- *SELECT*
- *DROP*

**Answer :**

\**ALL OF THE ABOVE* (*Correct*)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 38**

**How are virtual warehouse credits charged?**

Options are :

- *per hour*
- *per second*
- *per minute*
- *per-second, with a 60-second (i.e. 1-minute) minimum: (Correct)*

**Answer :**

\*per-second, with a 60-second (i.e. 1-minute) minimum: (Correct)

**Which of the statements are TRUE about streams?**

Options are :

- *Stream itself does not contain any table data (Correct)*
- *The hidden columns used by a stream do not consume any storage*
- *A stream only stores the offset for the source table (Correct)*

**Answer :**

\*Stream itself does not contain any table data (Correct)

\*A stream only stores the offset for the source table (Correct)

**When working on a small table which has only 10,000 rows, will specifying a clustering key further improve the queries that runs on that table?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Snowflake Practice Questions - SnowPro Core Certified Set 2**

**Mark all the true statements**

Options are :

- *Snowflake Stages can be defined as either FTP or SFTP*

- External Stages require a cloud storage provider (Correct)
- FTP Stages require three cloud storage providers
- Cloud storage providers must be linked to internal stage objects
- Snowflake Stages can be defined as either External or Internal (Correct)

**Answer :**

- \*External Stages require a cloud storage provider (Correct)
- \*Snowflake Stages can be defined as either External or Internal (Correct)

**Select 3 statements which are true for an external table**

Options are :

- Data can be updated in external tables
- External tables can be used for query and join operations (Correct)
- External tables are read-only (Correct)
- Views can be created against external tables (Correct)

**Answer :**

- \*External tables can be used for query and join operations (Correct)
- \*External tables are read-only (Correct)
- \*Views can be created against external tables (Correct)

**Which approach would result in improved performance through linear scaling of data ingestion workload?**

Options are :

- Consider practice of splitting input file batch within the recommended size of 10 MB to 100 MB
- Consider practice of organizing data by granular path
- Resize virtual warehouse (Correct)
- All of the above

**Answer :**

- \*Resize virtual warehouse (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 31**

**You want to get the DDL statement of a snowflake table. What is the command that you will use?**

Options are :

- *show table 'empldb.emplschema.employee';*
- *select fetch\_ddl('table', 'empldb.emplschema.employee');*
- *show DDL like 'empldb.emplschema.employee';*
- *select get\_ddl('table', 'empldb.emplschema.employee'); (Correct)*

**Answer :**

\**select get\_ddl('table', 'empldb.emplschema.employee');* (Correct)

**Query load is calculated by dividing the execution time (in seconds) of all queries in an interval by the total time (in seconds) for the interval.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Which is the model in which each object has an owner, who can in turn grant access to that object?**

Options are :

- *Object ownership model*
- *Discretionary Access Control (DAC) (Correct)*
- *Role-based Access Control (RBAC)*

**Answer :**

\**Discretionary Access Control (DAC) (Correct)*

### [Snowflake SnowPro Core Certification Practice Tests Set 14](#)

**After how many days does the load metadata expire**

Options are :

- *365*
- *30*

- 64 (Correct)

- 60

**Answer :**

\*64 (Correct)

**You can use the query profiler view only for completed queries**

Options are :

- FALSE (Correct)

- TRUE

**Answer :**

\*FALSE (Correct)

**Zero-copy cloning in snowflake is powered by which service?**

Options are :

- Cache of virtual warehouse
- Metadata store of cloud services layer (Correct)
- Query result cache

**Answer :**

\*Metadata store of cloud services layer (Correct)

## Snowflake Certification Preparation Set 12

**What are the options that are available to load data on snowflake?**

Options are :

- SnowSql (Correct)
- 3rd party ETL tools (Correct)
- Snowflake web UI (Correct)
- Snowflake data uploader
- Snowpipe (Correct)

**Answer :**

- \*SnowSql (Correct)
- \*3rd party ETL tools (Correct)
- \*Snowflake web UI (Correct)
- \*Snowpipe (Correct)

**Which of the below features are only available to Business Critical edition and higher?**

Options are :

- *Support for secure, direct proxy to your other virtual networks or on-premises data centers using AWS PrivateLink or Azure Private Link. (Correct)*
- *Support for PCI DSS (Correct)*
- *Extended Time Travel (up to 90 days)*
- *Customer-managed encryption keys through Tri-Secret Secure. (Correct)*
- *Support for PHI data (in accordance with HIPAA regulations). (Correct)*

**Answer :**

- \*Support for secure, direct proxy to your other virtual networks or on-premises data centers using AWS PrivateLink or Azure Private Link. (Correct)
- \*Support for PCI DSS (Correct)
- \*Customer-managed encryption keys through Tri-Secret Secure. (Correct)
- \*Support for PHI data (in accordance with HIPAA regulations). (Correct)

**What can you do to restore the table if one of your users have accidentally dropped a table EMPLOYEE in production environment?**

Options are :

- *Call Snowflake team and request them to recover the data from cloud storage*
- *Execute statement "UNDROP TABLE EMPLOYEE;" (Correct)*
- *Nothing can be done, whatever is lost cannot be recovered*
- *Execute statement "UNDELETE TABLE EMPLOYEE;"*

**Answer :**

- \*Execute statement "UNDROP TABLE EMPLOYEE;" (Correct)

**Snowflake Certification Preparation Set 1**

**Before any PHI data can be stored in Snowflake, a signed business associate agreement (BAA) must be in place between your agency/organization and Snowflake Inc**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**What is the maximum size of a micro-partition?**

Options are :

- *500MB (Correct)*
- *100MB*
- *1024MB*

**Answer :**

\**500MB (Correct)*

**Can PUT command be run from Snowflake worksheets?**

Options are :

- *Yes*
- *No (Correct)*

**Answer :**

\**No (Correct)*

**Snowflake SnowPro Core Certification Exam Questions Set 12**

**Using which datatype, semi-structured data (XML/JSON/Parquet) can be stored in a column?**

Options are :

- *Column of JSON datatype*
- *Column of Varchar datatype*
- *Column of Parquet datatype*
- *Column of Variant datatype (Correct)*

**Answer :**

\*Column of Variant datatype (Correct)

**create or replace table result\_scan\_table\_1 like result\_scan\_table; Will the above query cost compute credits?**

Options are :

- *No, since it is a metadata operation only (Correct)*
- *It is hard to say without looking at query profile*
- *Yes, it will need compute as the table structure is also getting created*

**Answer :**

\*No, since it is a metadata operation only (Correct)

**Which snowflake edition supports private communication between Snowflake and your other VPCs through AWS PrivateLink**

Options are :

- *Enterprise*
- *Advanced*
- *Business Critical (Correct)*
- *Standard*

**Answer :**

\*Business Critical (Correct)

**Snowflake Practice Questions - SnowPro Core Certified Set 2**

**Scenario: A virtual warehouse ran for 62 seconds and then was suspended. After few minutes, it resumed and ran for 20 seconds. For how many seconds will you be billed?**

Options are :

- *82*
- *20*
- *122 (Correct)*
- *62*

**Answer :**

\*122 (Correct)

**An entity to which access can be granted. Unless allowed by a grant, access will be denied. in snowflake terms, what is this called?**

Options are :

- *Privilege*
- *Role*
- *Securable object (Correct)*
- *User*

**Answer :**

\*Securable object (Correct)

**To share data with other users, what should be the minimum snowflake edition to enable data sharing?**

Options are :

- *Business Critical*
- *Enterprise*
- *Standard (Correct)*

**Answer :**

\*Standard (Correct)

### Snowflake Snowpro Certification Practice Test Set 3

**Scenario: Your organization requires a snowflake instance that is a completely separate Snowflake environment, isolated from all other Snowflake accounts. Which edition will you suggest?**

Options are :

- *Business Critical*
- *Enterprise*
- *Standard*
- *Virtual Private Snowflake (VPS) (Correct)*

**Answer :**

\*Virtual Private Snowflake (VPS) (Correct)

**What is the name of the Snowflake's Command Line Interface tool?**

Options are :

- *SnowSQL (Correct)*
- *SnowPowerShell*
- *SnowCMD*
- *SnowCLI*

**Answer :**

\*SnowSQL (Correct)

**Which of the following performance optimization techniques are supported by Snowflake?**

Options are :

- *QUERY RESULT CACHING (Correct)*
- *SSD CACHING/DATA CACHING (Correct)*
- *B-TREE INDEXES*

**Answer :**

\*QUERY RESULT CACHING (Correct)

\*SSD CACHING/DATA CACHING (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 16](#)

**Scaling down a virtual warehouse(e.g from a large warehouse to a small one) is an automated process.**

Options are :

- *FALSE (Correct)*
- *TRUE*

**Answer :**

\*FALSE (Correct)

**Business Critical Edition formerly known as \_\_\_\_.**

Options are :

- *Business Edition*
- *Enterprise Edition*
- *Standard Edition*
- *Enterprise for Sensitive Data (ESD) (Correct)*

**Answer :**

\*Enterprise for Sensitive Data (ESD) (Correct)

**Snowflake recommends relying more heavily on insertReport than loadHistoryScan to help avoid exceeding the rate limit (error code 429)**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 24](#)

**You need to contact Snowflake Support team if retrieval of data is required from fail safe**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What are the installation options available for snowflake?**

Options are :

- *Snowflake hosted account on AWS (Correct)*
- *Snowflake hosted account on GCP (Correct)*
- *Hybrid cloud and on-premise*

- *Snowflake hosted account on Azure (Correct)*
- *On-Premise*

**Answer :**

- \*Snowflake hosted account on AWS (Correct)
- \*Snowflake hosted account on GCP (Correct)
- \*Snowflake hosted account on Azure (Correct)

**A network policy includes values in both allowed and blocked IP address lists, snowflake applies the blocked IP address list first.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

- \*TRUE (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 8](#)

**What is the best method to convert an existing permanent table to a transient table (or vice versa) while preserving data and other characteristics such as column defaults and granted privileges?**

Options are :

- *Unload the data from the existing table into a CSV file. Create the new table and then load the data.*
- *Run an ALTER TABLE command to convert the tables*
- *Create a new table and use the COPY GRANTS clause (Correct)*

**Answer :**

- \*Create a new table and use the COPY GRANTS clause (Correct)

**Do COPY and INSERT operations block other operations on snowflake table?**

Options are :

- *No (Correct)*
- *Yes*

**Answer :**

- \*No (Correct)

**With respect to Snowflake UI, which of the following are true?**

Options are :

- *A single session can be shared between multiple worksheets*
- *Every worksheet has its own session (Correct)*
- *Every worksheet can have a different role, warehouse and a database (Correct)*
- *Worksheets cannot have different role, warehouse and database*

**Answer :**

\*Every worksheet has its own session (Correct)

\*Every worksheet can have a different role, warehouse and a database (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 15

**Which command can be used to list streams for the current/specify database or schema, or across your entire account.**

Options are :

- *SHOW STREAMS (Correct)*
- *DISPLAY STREAMS*
- *LIST STREAMS*

**Answer :**

\*SHOW STREAMS (Correct)

**Select all transformations that are available when using 'COPY INTO' command to load data files from a stage into Snowflake tables**

Options are :

- *Aggregation*
- *Column data type conversion (Correct)*
- *Filters*
- *Column concatenation (Correct)*

**Answer :**

\*Column data type conversion (Correct)

\*Column concatenation (Correct)

**By default, result re-use is enabled in snowflake. What can be done by the users to disable result re-use?**

Options are :

- Execute 'ALTER SESSION SET QUERY\_CACHE=FALSE'
- Execute 'ALTER SESSION SET USE\_CACHED\_RESULT=FALSE' (Correct)
- Execute 'ALTER SESSION SET QUERY\_RESULT\_CACHE=FALSE'

**Answer :**

\*Execute 'ALTER SESSION SET USE\_CACHED\_RESULT=FALSE' (Correct)

#### Snowflake Certification Preparation Set 4

**Which objects do we explicitly refer when using the COPY INTO command to load data from external stages? HINT: 3 options are correct**

Options are :

- Stage (Correct)
- File Format (Correct)
- View
- Table (Correct)

**Answer :**

\*Stage (Correct)

\*File Format (Correct)

\*Table (Correct)

**How many numbers of clusters a Snowflake warehouse can have at maximum?**

Options are :

- 7
- 10 (Correct)
- 4
- 2

**Answer :**

\*10 (Correct)

**Files that are already copied from the stage to the source table can be loaded again into a table cloned from the source table**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**SNOWPIPE REST APIs supports both internal and external stages**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\**TRUE (Correct)*

**To ingest semi structured data in snowflake tables, which data type will you use?**

Options are :

- *NUMBER*
- *VARCHAR*
- *VARIANT (Correct)*

**Answer :**

\**VARIANT (Correct)*

**Which factors influence the unit cost of snowflake credits and data storage?**

Options are :

- *Whether account is an On Demand or Capacity account (Correct)*
- *Snowflake Edition (Correct)*
- *Number of users of snowflake*
- *The region where the Snowflake account is provisioned (Correct)*

**Answer :**

- \*Whether account is an On Demand or Capacity account (Correct)
- \*Snowflake Edition (Correct)
- \*The region where the Snowflake account is provisioned (Correct)

**SnowFlake SnowPro Advanced Architect Practice Exam Set 7**

**A warehouse can be assigned to a single resource monitor only**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \**TRUE (Correct)*

**When you load data using Snowpipe, are loads always performed in a single transaction?**

Options are :

- *No (Correct)*
- *Yes*

**Answer :**

- \**No (Correct)*

**From a snowflake perspective what is common about Fivetran, Informatica, Matillion, Segment, Stitch and Talend**

Options are :

- *All are snowflake data integration partners (Correct)*
- *All are owned by Snowflake*
- *All are snowflake competitors*

**Answer :**

- \**All are snowflake data integration partners (Correct)*

**Snowflake Certification Preparation Set 4**

**Snowflake is compliant with which certifications?**

Options are :

- *HIPAA only*
- *HIPAA, PCI DSS and FedRAMP (Correct)*
- *HIPAA, PCI DSS*

**Answer :**

\*HIPAA, PCI DSS and FedRAMP (Correct)

**Data providers must sign an agreement before their data is made public on the data marketplace of snowflake.**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**When a warehouse does not have enough resources available to process the queries, what happens to the incoming queries?**

Options are :

- *The queries are queued and then executed when the resources are available again (Correct)*
- *Snowflake automatically resizes the warehouse*
- *Queries are aborted immediately*

**Answer :**

\*The queries are queued and then executed when the resources are available again (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 36](#)

**Which snowflake edition provides a Tri-Secret Secure option?**

Options are :

- *All editions*

- Business critical or higher (Correct)
- Enterprise Edition

**Answer :**

\*Business critical or higher (Correct)

**How the data is loaded on Snowflake?**

Options are :

- Data is loaded in structured format in the underlying storage
- Snowflake loads data in its optimized storage in columnar format which is saved as compressed data (Correct)
- Snowflake always loads the data from an external source
- Data is loaded in Semi-structured format in the underlying storage

**Answer :**

\*Snowflake loads data in its optimized storage in columnar format which is saved as compressed data (Correct)

**Which are the supported data types for a JSON name/value pair?**

Options are :

- A string (in double quotes) (Correct)
- A Boolean (true or false) (Correct)
- Null (Correct)
- An object (in curly braces) (Correct)
- A number (integer or floating point) (Correct)
- An array (in square brackets) (Correct)

**Answer :**

\*A string (in double quotes) (Correct)  
\*A Boolean (true or false) (Correct)  
\*Null (Correct)  
\*An object (in curly braces) (Correct)  
\*A number (integer or floating point) (Correct)  
\*An array (in square brackets) (Correct)

[\*\*Snowflake Snowpro Certification Practice Test Set 4\*\*](#)

**How to identify the potential performance bottlenecks and improvement opportunities of a query?**

Options are :

- *Use Query Profile (Correct)*
- *View table statistics*
- *Call snowflake support*
- *Use Explain Profile*

**Answer :**

\*Use Query Profile (Correct)

**A transient table cannot be cloned to a permanent table**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What information can be seen in the RECORD\_METADATA column once it is created and loaded with data using Kafka?**

Options are :

- *Topic (Correct)*
- *CreateTime / LogAppendTime (Correct)*
- *Partition (Correct)*
- *Value*
- *Key (Correct)*

**Answer :**

\*Topic (Correct)

\*CreateTime / LogAppendTime (Correct)

\*Partition (Correct)

\*Key (Correct)

## Tackling Snowflake Certification - Practice Questions Set 4

All of the below are valid execution states of a snowpipe except:

Options are :

- STOPPED (Correct)
- PAUSED
- RUNNING
- STALLED\_EXECUTION\_ERROR
- STOPPED\_FEATURE\_DISABLED

**Answer :**

\*STOPPED (Correct)

When using Snowpipe REST APIs, to help avoid exceeding the rate limit (error code 429), snowflake recommends relying more heavily on insertReport than loadHistoryScan.

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

The following factors affect data load rates

Options are :

- Virtual warehouse RAM
- Gzip compression efficiency (Correct)
- Thread size
- Physical location of the stage (Correct)

**Answer :**

\*Gzip compression efficiency (Correct)

\*Physical location of the stage (Correct)

## Tackling Snowflake Certification - Practice Questions Set 1

Which scenarios can be considered to use materialized views?

Options are :

- *None of the above*
- *The query results contain a small number of rows and/or columns relative to the base table (Correct)*
- *View's base table does not change frequently (Correct)*
- *Query is on an external table (Correct)*
- *Query results contain results that require significant processing (Correct)*

**Answer :**

\*The query results contain a small number of rows and/or columns relative to the base table (Correct)

\*View's base table does not change frequently (Correct)

\*Query is on an external table (Correct)

\*Query results contain results that require significant processing (Correct)

**Snowflake has three types of stages USER, TABLE, NAMED. Named stage is again divided into External and Internal. Which of the below two stages are automatically created and does not need explicit configuration by the user?**

Options are :

- *EXTERNAL*
- *TABLE (Correct)*
- *USER (Correct)*
- *INTERNAL*

**Answer :**

\**TABLE (Correct)*

\**USER (Correct)*

**You have a Snowflake table which is defined as below CREATE OR REPLACE TABLE EMPLOYEE(EMPLOYEE\_NUMBER NUMBER, DEPARTMENT VARCHAR, SALARY NUMBER); If you would like to convert the SALARY column to be a decimal with a certain precision and scale, which command will you run?**

Options are :

- *SELECT SALARY::DECIMAL(10,5) FROM EMPLOYEE; (Correct)*
- *SELECT SALARY AS DECIMAL(10,5) FROM EMPLOYEE;*
- *SELECT FRUIT\_NUMBER.DECIMAL(10,5) FROM FRUITS;*
- *SELECT SALARY(DECIMAL(10,5)) FROM EMPLOYEE;*

**Answer :**

\**SELECT SALARY::DECIMAL(10,5) FROM EMPLOYEE; (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 24

**What are the snowflake editions available as of today?**

Options are :

- *Standard Edition,Enterprise Edition,Business Critical Edition,Virtual Private Snowflake (Correct)*
- *Standard Edition, Enterprise Edition*
- *Enterprise Edition,Business Critical Edition*

**Answer :**

\**Standard Edition,Enterprise Edition,Business Critical Edition,Virtual Private Snowflake (Correct)*

**On which factors, compute cost is calculated on Snowflake?**

Options are :

- *Query execution time + Query wait time + Warehouse start time*
- *Query execution time*
- *Query execution time and Virtual Warehouse size (Correct)*

**Answer :**

\**Query execution time and Virtual Warehouse size (Correct)*

**The table functions in INFORMATION\_SCHEMA can be used to return account-level usage and historical information for storage, warehouses, user logins, and queries**

Options are :

- *FALSE*

- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

### Tackling Snowflake Certification - Practice Questions Set 1

**Snowflake does not support nested transactions**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Snowflake offers inbuilt tools to extract data from external source systems**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which are the three layers of Snowflake architecture**

Options are :

- *Extract, Load, Transform*
- *Storage Layer, Virtual Warehouse/Query Processing Layer, Cloud Services Layer (Correct)*
- *Storage Layer, Transformation Layer, Data Load layer*

**Answer :**

\*Storage Layer, Virtual Warehouse/Query Processing Layer, Cloud Services Layer (Correct)

### Snowflake Certification Preparation Set 5

**The insertReport SNOWPIPE API can retrieve file ingestion events and report it. The events are retained for infinite time until deleted manually.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**When you clone a table in snowflake, does the storage requirements gets doubled as it needs to now store two copies of data?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**With respect to snowflake security, Tri-secret secure option requires customers to manage their keys.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 12

**Which snowflake edition allows only a maximum of 1 day of time travel?**

Options are :

- *Standard (Correct)*
- *Enterprise*
- *Virtual Private Snowflake (VPS)*

- *Business Critical*

**Answer :**

\*Standard (Correct)

**GET does not support downloading files from external stages**

Options are :

- *FALSE*
- *TRUE (Correct)*

**Answer :**

\*TRUE (Correct)

**After how many days do the load metadata of a table expire?**

Options are :

- 7
- 14
- 30
- 64 (*Correct*)
- 365

**Answer :**

\*64 (Correct)

### **Snowflake Certification Preparation Set 21**

**Does Snowpipe guarantees that files are loaded in the same order they are staged?**

Options are :

- *No (Correct)*
- *Yes*

**Answer :**

\*No (Correct)

**Select two options which are true about variant data type in snowflake?**

Options are :

- Optimizes storage based on repeated elements (Correct)
- A separate file format is used to store variant data
- Supports querying using JSON path notation (Correct)
- Requires custom mapping for each type of record

**Answer :**

\*Optimizes storage based on repeated elements (Correct)  
\*Supports querying using JSON path notation (Correct)

#### **Which snowflake features are available for enabling continuous data pipelines**

Options are :

- Continuous data loading / Snowpipe (Correct)
- Table Pipes
- Recurring tasks (Correct)
- Change data tracking / Streams (Correct)
- Kafka - using connector for Snowflake (Correct)

**Answer :**

\*Continuous data loading / Snowpipe (Correct)  
\*Recurring tasks (Correct)  
\*Change data tracking / Streams (Correct)  
\*Kafka - using connector for Snowflake (Correct)

#### **Snowflake Certification Preparation Set 1**

**Only structured data is allowed to be loaded on Snowflake.**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**You are an account administrator and you want to use an external tool with snowflake. How you can use Snowflake with a third-party (external) tool?**

Options are :

- Visit external tool website and register
- Visit partner connect on snowflake account and activate or enable the tool (Correct)
- Contact snowflake support for handling of such requests

**Answer :**

\*Visit partner connect on snowflake account and activate or enable the tool (Correct)

**Which of the below options are layers of Snowflake Architecture?**

Options are :

- STORAGE (Correct)
- SECURITY
- COMPUTE/VIRTUAL WAREHOUSE (Correct)
- DATA CATALOG
- CLOUD SERVICES (Correct)

**Answer :**

\*STORAGE (Correct)

\*COMPUTE/VIRTUAL WAREHOUSE (Correct)

\*CLOUD SERVICES (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 5

**Which snowpipe rest APIs Fetches a report about ingested files whose contents have been added to table?**

Options are :

- insertfiles
- insertPipeReport
- loadHistoryScan and insertReport (Correct)

**Answer :**

\*loadHistoryScan and insertReport (Correct)

**When you sign up for a snowflake trial account, what are three snowflake editions that are offered**

Options are :

- *Enterprise Advanced*
- *Standard (Correct)*
- *Enterprise (Correct)*
- *Business Critical (Correct)*
- *Free-Trial*

**Answer :**

\*Standard (Correct)  
\*Enterprise (Correct)  
\*Business Critical (Correct)

**Using default settings, how long a query can run on snowflake?**

Options are :

- *Snowflake will cancel the query if it runs more than 48 hours (Correct)*
- *Snowflake will cancel the query if the warehouse runs out of memory and hard disk storage*
- *Snowflake will cancel the query if it runs more than 24 hours*
- *Snowflake will cancel the query if the warehouse runs out of memory*

**Answer :**

\*Snowflake will cancel the query if it runs more than 48 hours (Correct)

### [Snowflake Certification Preparation Set 12](#)

**If file format options are specified in multiple locations, the load operation applies the options in the following order of precedence.**

Options are :

- 1. Stage definition. 2. COPY INTO TABLE statement. 3. Table definition.
- 1. COPY INTO TABLE statement. 2. Table definition. 3. Stage definition.
- 1. COPY INTO TABLE statement. 2. Stage definition. 3. Table definition. (Correct)

**Answer :**

\*1. COPY INTO TABLE statement. 2. Stage definition. 3. Table definition. (Correct)

**You have created a secure view sample\_secured\_view. Now you want to convert it back to a regular view. Which of the below can be used for the conversion?**

Options are :

- *alter view sample\_secured\_view set secure;*
- *None of the above*
- *alter view sample\_secured\_view set type = unsecure;*
- *after view sample\_secured\_view unset secure; (Correct)*

**Answer :**

\**after view sample\_secured\_view unset secure; (Correct)*

**What is the technique called which snowflake uses to limit the number of micro-partitions retrieved as part of a query?**

Options are :

- *Deleting*
- *Filtering*
- *Clustering*
- *Indexing*
- *Pruning (Correct)*

**Answer :**

\**Pruning (Correct)*

**Tools      Topics**

## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 35

SnowPro Core Certification 3 Full Practice Exams 2021 Set 35

**Materialized view can be created on top of external tables to improve query performance without any additional cost. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which factors affect your snowflake data loading rate ?**

Options are :

- *Physical location of the stage (Correct)*
- *RAM on the virtual warehouse*
- *Thread size*
- *GZIP compression efficiency (Correct)*

**Answer :**

\**Physical location of the stage (Correct)*

\**GZIP compression efficiency (Correct)*

**To which of theses Snowflake locations you CANNOT upload files from a local machine with PUT command ? Select all that apply.**

Options are :

- *Named internal stage*
- *Internal stage for specified table*
- *Named External stage (Correct)*
- *Internal stage for the current user*

**Answer :**

\**Named External stage (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 19**

**Snowflake Web Interface data loading wizard is designed to load small amounts of data where each data file is smaller than 50MB in size. TRUE or FALSE load**

Options are :

- TRUE (Correct)

- FALSE

**Answer :**

\*TRUE (Correct)

**What privileges does SYSADMIN have by default?**

Options are :

- CREATE DATABASE (Correct)
- CREATE SHARE
- CREATE WAREHOUSE (Correct)
- CREATE USER
- CREATE ROLE

**Answer :**

\*CREATE DATABASE (Correct)

\*CREATE WAREHOUSE (Correct)

**Snowflake can automatically scale your warehouse UP when needed (make your warehouse larger). TRUE or FALSE ?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Snowflake SnowPro Core Certification Practice Test Set 5**

**What type of view is most like a table**

Options are :

- Non-materialized views
- Materialized views (Correct)

**Answer :**

\*Materialized views (Correct)

**What stages can you use when setting up your Snowflake account with AWS as your cloud platform provider ? Select all that apply**

Options are :

- *Named External - Using S3 Buckets (Correct)*
- *Named External - Hadoop Cloudera File System*
- *Named External - Using Azure Blob storage (Correct)*
- *Named External - Using GCS/GCP Buckets (Correct)*
- *Named External - AWS Redshift tables*
- *None of the above*

**Answer :**

\*Named External - Using S3 Buckets (Correct)  
\*Named External - Using Azure Blob storage (Correct)  
\*Named External - Using GCS/GCP Buckets (Correct)

**What services are managed by Cloud Services layer ? (Select all that apply)**

Options are :

- *Authentication (Correct)*
- *Infrastructure Management (Correct)*
- *Metadata Management (Correct)*
- *Query parsing and optimization*
- *Data storage*
- *Query processing*

**Answer :**

\*Authentication (Correct)  
\*Infrastructure Management (Correct)  
\*Metadata Management (Correct)

**Snowflake SnowPro Core Certification Practice Test Set 2**

**Select the term that is associated with compute layer?**

Options are :

- *Query planning*
- *Query optimization*
- *Query processing (Correct)*
- *Query Design*
- *None of the above*

**Answer :**

\*Query processing (Correct)

**Which transformations can be applied on data while loading it into a table using COPY command ?  
Select all that apply**

Options are :

- *column reordering (Correct)*
- *Aggregation*
- *Casts (Correct)*
- *Filtering*
- *Column omission (Correct)*
- *Truncating text strings that exceed the target column length (Correct)*

**Answer :**

\*column reordering (Correct)

\*Casts (Correct)

\*Column omission (Correct)

\*Truncating text strings that exceed the target column length (Correct)

**Select all Snowflake stages.**

Options are :

- *Internal Named Stage (Correct)*
- *External Named Stage (Correct)*
- *User Stage (Correct)*
- *Table Stage (Correct)*

- *Account Stage (Correct)*

**Answer :**

- \*Internal Named Stage (Correct)
- \*External Named Stage (Correct)
- \*User Stage (Correct)
- \*Table Stage (Correct)
- \*Account Stage (Correct)

### **Snowflake Snowpro Certification Practice Test Set 5**

**With Economy scaling, how long might queries be queued before another cluster is added to the warehouse ?**

Options are :

- *3 minutes*
- *6 minutes (Correct)*
- *9 minutes*
- *12 minutes*
- *None of the above*

**Answer :**

- \*6 minutes (Correct)

**What are the two mechanisms that can be used with Snowpipe for detecting that staged files are available ?**

Options are :

- *Calling Snowpipe REST endpoints (Correct)*
- *Calling Snowpipe SOAP endpoints*
- *Automating Snowpipe using CloudWatch*
- *Automating Snowpipe using cloud messaging (Correct)*

**Answer :**

- \*Calling Snowpipe REST endpoints (Correct)
- \*Automating Snowpipe using cloud messaging (Correct)

**What kind of solution is Snowflake**

Options are :

- *Infrastructure as a Service (IaaS)*
- *Platform as a Service (PaaS)*
- *Software as a Service (SaaS) (Correct)*

**Answer :**

\*Software as a Service (SaaS) (Correct)

### **Snowflake SnowPro Advanced Architect Practice Exam Set 7**

**In snowflake container hierarchy, what container is represented as a URL ( e.g. <https://jz96798.us-east-1.azure.snowflakecomputing.com> ) ?**

Options are :

- *Database*
- *Schema*
- *Role*
- *Account (Correct)*
- *Region*

**Answer :**

\*Account (Correct)

**How can you disable < auto-suspend > option for a virtual warehouse**

Options are :

- *Choose 'Never' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI (Correct)*
- *Choose 'Disable' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI*
- *ALTER WAREHOUSE set AUTO\_SUSPEND = 0 (Correct)*
- *ALTER WAREHOUSE set AUTO\_SUSPEND = NULL (Correct)*
- *ALTER WAREHOUSE set AUTO\_SUSPEND = -1*
- *None of the above*

**Answer :**

- \*Choose 'Never' for Auto Suspend in the Warehouses configuration tab in Snowflake Web UI (Correct)
- \*ALTER WAREHOUSE set AUTO\_SUSPEND = 0 (Correct)
- \*ALTER WAREHOUSE set AUTO\_SUSPEND = NULL (Correct)

**External stages do not require a cloud provider. TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

- \**FALSE (Correct)*

**Snowflake Certification Preparation Set 3**

**In which mode Snowflake multi-cluster warehouses can run ? Select all that apply.**

Options are :

- *Maximized (Correct)*
- *Auto-resume*
- *Auto-suspended*
- *Auto-scale (Correct)*
- *None of the above*

**Answer :**

- \**Maximized (Correct)*
- \**Auto-scale (Correct)*

**Which transformations can NOT be applied on data while loading it into a table using COPY command ? (Select all that apply)**

Options are :

- *Column reordering*
- *Aggregation (Correct)*
- *Casts*

- *Filtering (Correct)*
- *Column omission*
- *Truncating text strings that exceed the target column length*

**Answer :**

\*Aggregation (Correct)  
\*Filtering (Correct)

**How many clusters are in standard Medium warehouse ?**

Options are :

- *One (Correct)*
- *Two*
- *Three*
- *Four*

**Answer :**

\*One (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 12**

**A task does not need a warehouse and Snowflake provides compute for task ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which command returns Snowflake current version ?**

Options are :

- *SELECT version(); (Correct)*
- *SHOW version();*
- *DESC version();*

- `SELECT current_version();`

**Answer :**

\*`SELECT version();` (Correct)

**Which of the following have drivers/connectors (or information about where to find them) available via Help->Downloads in the Snowflake WebUI?**

Options are :

- `Go` (Correct)
- `R`
- `Node.Js` (Correct)
- `JDBC` (Correct)
- `Hive`
- `Spark`

**Answer :**

\*`Go` (Correct)

\*`Node.Js` (Correct)

\*`JDBC` (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 9

**Which syntax allows you to concat the column named 'EMPLOYEE' between two % signs ? Select all that apply**

Options are :

- `%' || EMPLOYEE || '%'` (Correct)
- `% + EMPLOYEE + %`
- `CONCAT('%', EMPLOYEE, '%')`
- `concat('%', concat(EMPLOYEE, '%'))` (Correct)
- All of the above

**Answer :**

\*`%' || EMPLOYEE || '%'` (Correct)

\*`concat('%', concat(EMPLOYEE, '%'))` (Correct)

**Snowflake can be run on private cloud infrastructures (on-premises or hosted). TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**What Snowflake objects can be shared ? (Select all that apply)**

Options are :

- *Tables (Correct)*
- *Views*
- *Roles*
- *Secure UDFs (Correct)*
- *Materialized views*
- *None of the above*

**Answer :**

\**Tables (Correct)*

\**Secure UDFs (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 2**

**For how long does a query remain in the query history tab ?**

Options are :

- *1 day*
- *7 days*
- *14 days (Correct)*
- *48 days*
- *None of the above*

**Answer :**

\**14 days (Correct)*

**Which of the following data types are treated as VARCHAR by Snowflake ?**

Options are :

- *TEXT (Correct)*
- *STRING (Correct)*
- *CHAR (Correct)*
- *VARCHAR (Correct)*
- *NVARCHAR*

**Answer :**

\*TEXT (Correct)  
\*STRING (Correct)  
\*CHAR (Correct)  
\*VARCHAR (Correct)

**Select the three terms refering to the same layer.**

Options are :

- *Compute Layer (Correct)*
- *Virtual Warehouse Layer (Correct)*
- *Catalog Layer*
- *Metadata Layer*
- *Query Processing Layer (Correct)*

**Answer :**

\*Compute Layer (Correct)  
\*Virtual Warehouse Layer (Correct)  
\*Query Processing Layer (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 3

**A company operating in the health sector wants to move to Snowflake. The company asks you about the edition of Snowflake to choose to comply with HIPAA regulation. Select the minimum edition of Snowflake to recommend ?**

Options are :

- *Enterprise Edition*

- Business Critical Edition (Correct)
- Standard Edition
- Virtual Private Snowflake

**Answer :**

\*Business Critical Edition (Correct)

**Time travel is available for which table types ?**

Options are :

- Permanent (Correct)
- Temporary (Correct)
- Transient (Correct)
- External

**Answer :**

\*Permanent (Correct)

\*Temporary (Correct)

\*Transient (Correct)

**When a warehouse is resized, which queries make use of the new size ?**

Options are :

- Only currently running queries
- Only subsequent queries (Correct)
- Both current and subsequent queries

**Answer :**

\*Only subsequent queries (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 36**

**Snowflake Virtual Warehouses cannot be resized while running. TRUE or FALSE ?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**What is the account identifier inside this Snowflake URL ' <https://jz96798.us-east-1.azure.snowflakecomputing.com> ' ?**

Options are :

- *jz96798.us-east-1.azure*
- *us-east-1.azure*
- *jz96798 (Correct)*
- *snowflakecomputing.com*
- *Azure*

**Answer :**

\**jz96798* (Correct)

**To create a SHARE what is the minimum required role ?**

Options are :

- *SYSADMIN*
- *SECURITYADMIN*
- *ACCOUNTADMIN (Correct)*
- *USERADMIN*

**Answer :**

\**ACCOUNTADMIN* (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 8**

**Which of the following terms are associated with the Compute/Warehouse Layer ?**

Options are :

- *Query Processing (Correct)*
- *Query Planning*
- *Query Optimization*

- *Query Design*
- *Query Compilation*

**Answer :**

\*Query Processing (Correct)

**Which cache runs on a 24 hour clock?**

Options are :

- *Result cache (Correct)*
- *Metadata Cache*
- *Warehouse cache*

**Answer :**

\*Result cache (Correct)

**Snowflake data storage costs include which type of data ?**

Options are :

- *Metadata*
- *Persistent data stored in permanent tables (Correct)*
- *Data retained to enable recovery (time travel, fail-safe) (Correct)*
- *Cached results*
- *Semi-structured additional fees*

**Answer :**

\*Persistent data stored in permanent tables (Correct)

\*Data retained to enable recovery (time travel, fail-safe) (Correct)

## [Snowflake Certification Preparation Set 2](#)

**When data is staged to an internal stage using PUT command, the data is encrypted on client machine. TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What is called the action of adding more servers to the same cluster of a virtual warehouse ?**

Options are :

- *Scale UP (Correct)*
- *Scale DOWN*
- *Scale IN*
- *Scale OUT*
- *None of the above*

**Answer :**

\*Scale UP (Correct)

**Multiple warehouses can access the same data simultaneously without resource contention issues.  
TRUE or FALSE ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 20**

**A stream advances the offset only when it is used in a DML transaction**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What influences Snowflake pricing ?**

Options are :

- *Storage and Compute (Correct)*
- *Amount of data queried from Snowflake*
- *Amount of data scanned during querying Snowflake*
- *None of the above*

**Answer :**

\*Storage and Compute (Correct)

**A client ask your support to choose a geographic deployment region for a Snowflake account. What factors should you consider ? Select the two best options.**

Options are :

- *Proximity to the point of service (Correct)*
- *Additional fees charged for regions with geo-political unrest*
- *Number of availability zones within a region (Correct)*
- *End-user perceptions of glamorous or trendy geographic locations*

**Answer :**

\*Proximity to the point of service (Correct)

\*Number of availability zones within a region (Correct)

**Which metadata is created when you upload data into Snowflake internal stage ? Select all that apply**

Options are :

- *File name (Correct)*
- *File Owner*
- *File size (Correct)*
- *DATE\_OF\_CREATION*
- *LAST\_MODIFIED date (Correct)*

**Answer :**

\*File name (Correct)

- \*File size (Correct)
- \*LAST\_MODIFIED date (Correct)

#### **What additional columns can we access when using a stream ?**

Options are :

- METADAT\$ACTION (Correct)
- METADATA\$DELETED
- METADATA\$ISUPDATE (Correct)
- METADATA\$ROW\_ID (Correct)
- METADATA\$COLUMN\_ID (Correct)
- None of the above

**Answer :**

- \*METADAT\$ACTION (Correct)
- \*METADATA\$ISUPDATE (Correct)
- \*METADATA\$ROW\_ID (Correct)
- \*METADATA\$COLUMN\_ID (Correct)

#### **Which Snowflake option enable to query data stored in external cloud storage without loading it into Snowflake**

Options are :

- External views
- Materialized views
- External Tables (Correct)

**Answer :**

- \*External Tables (Correct)

#### **Snowflake Snowpro Certification Practice Test Set 1**

#### **What is the size of data store inside Snowflake micro-partition ?**

Options are :

- 25KB - 50KB uncompressed
- 50MB - 500MB uncompressed (Correct)

- 10MB - 100MB uncompressed
- 1GB - 10GB uncompressed

**Answer :**

\*50MB - 500MB uncompressed (Correct)

**Scaling policy for a multi-cluster warehouse ONLY apply if it is running in Auto-scale mode. TRUE or FALSE**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Select all statements that are TRUE about FILE FORMAT.**

Options are :

- Stores data from a loaded file in a specific, named format
- Is securable (Correct)
- Stores information needed to parse files during load/unload (Correct)
- Can only be of type CSV

**Answer :**

\*Is securable (Correct)

\*Stores information needed to parse files during load/unload (Correct)

### Snowflake SnowPro Advanced Architect Practice Exam Set 3

**How many servers are in standard Large warehouse ?**

Options are :

- 2
- 4
- 8 (Correct)
- 16

**Answer :**

\*8 (Correct)

**Scaling OUT is an automated process (e.g. Min clusters to max clusters). TRUE or FALSE**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What does this cron expression (USING CRON 0 0-23 \* \* \*) express ?**

Options are :

- *Every minute of every hour of every day*
- *At 00:00 am on the 1st of every month*
- *The first minute of every hour of every day (Correct)*
- *At 00:00 am on the 1st of every month*

**Answer :**

\*The first minute of every hour of every day (Correct)

### **Snowflake Certification Preparation Set 3**

**Which of the following are TRUE about Snowpipe ?**

Options are :

- *You must scale a virtual warehouse to manage the compute available to Snowpipe*
- *The resource for Snowpipe are automatically scaled up and down by Snowflake (Correct)*
- *Snowpipe makes use of the active virtual warehouse for compute resources*
- *Snowpipe uses serverless compute resources managed by Snowflake (Correct)*

**Answer :**

\*The resource for Snowpipe are automatically scaled up and down by Snowflake (Correct)

\*Snowpipe uses serverless compute resources managed by Snowflake (Correct)

**What is the maximum number of child a single task can have in Snowflake ?**

Options are :

- 1
- 100 (*Correct*)
- 1000
- 10000

**Answer :**

\*100 (*Correct*)

**How does Snowflake billed for storage ?**

Options are :

- *Per Gigabyte, per month*
- *Per Terabyte, per day*
- *Per Gigabyte, per day (Correct)*
- *Per Terabyte, per month*

**Answer :**

\**Per Gigabyte, per day (Correct)*

### Snowflake SnowPro Core Certification Exam Questions Set 11

**Where do Snowflake supports loading data into ? Select all that apply.**

Options are :

- *Internal stage on cloud storage platform (Correct)*
- *External stage on cloud storage platform (Correct)*
- *Internal stage on local system*
- *Direct loading to the table*
- *None of the above*

**Answer :**

- \*Internal stage on cloud storage platform (Correct)
- \*External stage on cloud storage platform (Correct)

**Is Snowflake HIPAA compliant ?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \**TRUE (Correct)*

**Snowpipe is best suited to load what kind of data ? Select the best option**

Options are :

- *Small volume of frequent data (Correct)*
- *Large volume of data on a batch schedule*
- *Any kind of volume small or batch or continuous*
- *None of the above*

**Answer :**

- \**Small volume of frequent data (Correct)*

**Snowflake SnowPro Core Certification Practice Tests Set 10**

**What factors impact credit consumption of the compute layer ? Select all that apply**

Options are :

- *Amount of data processed*
- *Size of warehouse (Correct)*
- *Number of clusters (Correct)*
- *Number of user*

**Answer :**

- \**Size of warehouse (Correct)*
- \**Number of clusters (Correct)*

### **How does Snowflake billed for compute ?**

Options are :

- *Per hour, with one-minute minimum*
- *Per second, with one-minute minimum (Correct)*
- *Per day*
- *Per month*

### **Answer :**

\*Per second, with one-minute minimum (Correct)

### **A Snowflake session can have more than one virtual warehouse. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

### **Answer :**

\*FALSE (Correct)

## **Snowflake Certification Preparation Set 11**

### **Who can use the data cache ? Select the best option.**

Options are :

- *Anyone who uses the same warehouse (Correct)*
- *Only the user who execute the query*
- *Only user with the same role*
- *Only account administrator*
- *None of the above*

### **Answer :**

\*Anyone who uses the same warehouse (Correct)

### **Which is Snowflake Architecture layer ? (Select all that apply)**

Options are :

- Database storage Layer (Correct)
- Account Layer
- Cloud services Layer (Correct)
- Query processing (Correct)
- Administration Layer

**Answer :**

\*Database storage Layer (Correct)  
\*Cloud services Layer (Correct)  
\*Query processing (Correct)

**Select all statements that are TRUE about fail safe.**

Options are :

- Data can be recovered from fail-safe storage only by snowflake employee (Correct)
- Tables with fail-safe turned on incurs additional storage costs compared to tables where fail-safe is not turned on (Correct)
- Fail-safe can be used to create DEV/TEST environments
- Fail safe period is reconfigurable (Correct)
- None of the above

**Answer :**

\*Data can be recovered from fail-safe storage only by snowflake employee (Correct)  
\*Tables with fail-safe turned on incurs additional storage costs compared to tables where fail-safe is not turned on (Correct)  
\*Fail safe period is reconfigurable (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 21](#)

**How many credits consume a Large-size multi-cluster warehouse with 3 warehouses which runs in Maximized mode for 2 hours**

Options are :

- 24
- 48 (Correct)

- 16

- 56

**Answer :**

\*48 (Correct)

**Scaling policy for a multi-cluster warehouse can be apply if it is running in Maximized mode. TRUE or FALSE**

Options are :

- *TRUE*

- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Who can use the metadata cache ?**

Options are :

- *Only the user who execute the query*

- *Only user with the same role*

- *Only account administrator*

- *Everyone (Correct)*

- *None of the above*

**Answer :**

\*Everyone (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 14**

**Tri-Secret Secure option is available in which snowflake edition ?**

Options are :

- *Business critical (or higher) (Correct)*

- *Enterprise edition*

- *All editions*

**Answer :**

\*Business critical (or higher) (Correct)

**What is the maximum number of clusters that can be added in multi cluster warehouse ?**

Options are :

- 5
- 10 (Correct)
- 20
- 30

**Answer :**

\*10 (Correct)

**What ROLE should create users and roles in Snowflake ? Select all that apply**

Options are :

- ACCOUNTADMIN
- USERADMIN (Correct)
- SYSADMIN
- SECURITYADMIN (Correct)
- PUBLIC

**Answer :**

\*USERADMIN (Correct)

\*SECURITYADMIN (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**In which Snowflake editions is it possible to extended Time Travel (up to 90 days) ? (Select two or more)**

Options are :

- Standard Edition
- Enterprise Edition (Correct)
- Business Critical Edition (Correct)

- *Virtual Private Snowflake (Correct)*

**Answer :**

- \*Enterprise Edition (Correct)
- \*Business Critical Edition (Correct)
- \*Virtual Private Snowflake (Correct)

**Which of the following Snowflake Editions automatically store data in an encrypted state?**

Options are :

- *Standard Edition (Correct)*
- *Business Critical Edition (Correct)*
- *Virtual Private Snowflake (Correct)*
- *Enterprise Edition (Correct)*
- *None of the above*

**Answer :**

- \*Standard Edition (Correct)
- \*Business Critical Edition (Correct)
- \*Virtual Private Snowflake (Correct)
- \*Enterprise Edition (Correct)

**Which of the following best describes Snowflake's processing engine ?**

Options are :

- *Native SQL (Correct)*
- *Leverages Apache Spark*
- *Based on map reduce*
- *A derivative of presto*

**Answer :**

- \*Native SQL (Correct)

### Snowflake Certification Preparation Set 6

**Snowflake account hosted on AWS cannot load data from a file that is staged on GCP or Azure.  
TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which file formats is not supported by Snowflake to load data ?**

Options are :

- *Delimited files*
- *JSON*
- *Avro*
- *ORC*
- *Parquet*
- *PDF (Correct)*

**Answer :**

\**PDF (Correct)*

**Virtual warehouses cannot be resized while they are running. TRUE or FALSE ?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

#### Snowflake Certification Preparation Set 4

**In which situation would you 'scale up' your warehouse ? Select the best option**

Options are :

- *To make some credit economy when the warehouse is underutilized*
- *To provide more resources for a complex queries, to improve performance (Correct)*
- *To support multi-concurrent users or queries*

- *None of the above*

**Answer :**

\*To provide more resources for a complex queries, to improve performance (Correct)

**Select all TRUE statements about Snowpipe.**

Options are :

- *Snowpipe can load a file with the same name if it has been modified after it was last loaded (Correct)*
- *Snowpipe REST APIs support both internal and external stage (Correct)*
- *insertFiles is a Snowpipe REST API endpoint (Correct)*
- *Load history of Snowpipe expires after 14 days (Correct)*

**Answer :**

\*Snowpipe REST APIs support both internal and external stage (Correct)

\*insertFiles is a Snowpipe REST API endpoint (Correct)

\*Load history of Snowpipe expires after 14 days (Correct)

**It is possible to install Snowflake on-premise. TRUE or FALSE**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 32**

**Which cloud platforms can a Snowflake account be hosted on ?**

Options are :

- *Google Cloud Platform (GCP) (Correct)*
- *OVHcloud*
- *Microsoft Azure (Azure) (Correct)*
- *Alibaba Cloud*
- *Amazon Web Services (AWS) (Correct)*

- *None of the above*

**Answer :**

- \*Google Cloud Platform (GCP) (Correct)
- \*Microsoft Azure (Azure) (Correct)
- \*Amazon Web Services (AWS) (Correct)

**Which data type should you define to load semi-structured data inside Snowflake**

Options are :

- *JSON*
- *PARQUET*
- *VARIANT (Correct)*
- *VARCHAR*
- *None of the above*

**Answer :**

- \**VARIANT (Correct)*

**A client has an important number of users that access a BI report every working day at 9AM. Every time a user refreshes the report it requests Snowflake for the data. The report performances are not optimal.\n What suggestion could you make to improve the performance ? Select the best option**

Options are :

- *Create a single warehouse for each user*
- *Regroup the users by age and associate a warehouse to each group*
- *Use a multi-cluster warehouse (Correct)*
- *None of the above*

**Answer :**

- \**Use a multi-cluster warehouse (Correct)*

#### Snowflake SnowPro Core Certification Practice Test Set 4

**When loading JSON file inside Snowflake table, which file format option removes the outer array structure and loads the records into separate table rows ? (**

Options are :

- *STRIP\_NULL\_VALUES = TRUE*
- *STRIP\_OUTER\_ARRAY = TRUE (Correct)*
- *TRIM\_SPACE = TRUE*
- *None of the above*

**Answer :**

\**STRIP\_OUTER\_ARRAY = TRUE (Correct)*

**Each Snowflake account is limited to a maximum of 10 shares**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**How many tasks can we have in a simple tree ?**

Options are :

- *10*
- *100*
- *1000 (Correct)*
- *10000*

**Answer :**

\**1000 (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 22**

**What is the maximum size for a VARIANT data type column ?**

Options are :

- *64MB uncompresssed*
- *16MB uncompressed (Correct)*
- *32MB uncompressed*

- 128MB uncompressed

**Answer :**

\*16MB uncompressed (Correct)

**What are the actions cloud services layer is responsible for ? (Select all that apply)**

Options are :

- User authentication (Correct)
- Infrastructure management (Correct)
- Metadata management (Correct)
- Query parsing and optimization (Correct)
- Access control (network policies) (Correct)
- Query processing

**Answer :**

\*User authentication (Correct)  
\*Infrastructure management (Correct)  
\*Metadata management (Correct)  
\*Query parsing and optimization (Correct)  
\*Access control (network policies) (Correct)

**How long does the query result cache last ?**

Options are :

- 72 hours
- 48 hours
- 24 hours (Correct)
- 1 hour

**Answer :**

\*24 hours (Correct)

# 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 37

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 37

**What kind of stream can be created on external table ?**

Options are :

- *Standard*
- *Append-only*
- *Insert-only (Correct)*
- *None of the above*

**Answer :**

\**Insert-only (Correct)*

**What does "Query Load" measure ? Select the best option**

Options are :

- *average number of queries that were running within a specific interval*
- *average number of queries that were queued within a specific interval*
- *average number of queries that were running or queued within a specific interval (Correct)*
- *None of the above*

**Answer :**

\**average number of queries that were running or queued within a specific interval (Correct)*

**Which installment option version of Snowflake are available ?**

Options are :

- *Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)*
- *Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)*
- *Microsoft Cloud Native Accounts*
- *Hybrid On-Premise + Cloud Installation*

- *Enterprise In-House VPC Installation*
- *On-Premise Custom Installation*

**Answer :**

\*Snowflake-Hosted Accounts (on Amazon cloud infrastructure) (Correct)

\*Snowflake-Hosted Accounts (on Azure cloud infrastructure) (Correct)

### [Snowflake Practice Questions - SnowPro Core Certified Set 6](#)

**Which of the following terms describes Snowflake's Architecture ?**

Options are :

- *Shared Disk*
- *Shared Nothing*
- *Shared Data (Correct)*
- *Shared Memory*

**Answer :**

\*Shared Data (Correct)

**Which compute resource can be used by Snowflake for data loading jobs ? Select all that apply**

Options are :

- *User-managed virtual warehouse (Correct)*
- *Hardware provisionned by user directly from cloud provider*
- *Snowflake-managed serverless compute (Correct)*
- *Cloud Provider Runtime Environment*

**Answer :**

\*User-managed virtual warehouse (Correct)

\*Snowflake-managed serverless compute (Correct)

**If you find a data-related tool that is not listed as part of the Snowflake ecosystem, what industry standard options could you check for as a way to easily connect to Snowflake?**

Options are :

- *Check see if the tool can connect to other solution via JDBC (Correct)*

- Check see if the tool can connect to other solution via ODBC (Correct)
- Check see if there is a petition in the community to create a driver
- Check see if you can develop a driver and put it on github

**Answer :**

\*Check see if the tool can connect to other solution via JDBC (Correct)  
\*Check see if the tool can connect to other solution via ODBC (Correct)

### **Snowflake SnowPro Core Certification Exam Questions Set 9**

**What actions can be triggered by a resource monitor ? Select all that apply**

Options are :

- Notify and Suspend (Correct)
- Auto-resume
- Notify (Correct)
- Auto-scale
- Notify and Suspend Immediately (Correct)
- None of the above

**Answer :**

\*Notify and Suspend (Correct)  
\*Notify (Correct)  
\*Notify and Suspend Immediately (Correct)

**Is it possible to disable Time Travel for a Snowflake Account ?**

Options are :

- True
- False (Correct)

**Answer :**

\*False (Correct)

**Who is responsible of materialized views maintenance ?**

Options are :

- *Account administrator*
- *System administrator*
- *Snowflake (Correct)*
- *Everyone*
- *None of the above*

**Answer :**

\*Snowflake (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 20](#)

**What is the name of the Snowflake-produced Command Line Interface tool?**

Options are :

- *SnowCLI*
- *SnowCommand*
- *SnowSQL (Correct)*
- *SnowSpark*
- *None of the above*

**Answer :**

\*SnowSQL (Correct)

**If your organizations cloud services are hosted on Amazon Web Services, then you must also host your Snowflake account on Amazon Web Services.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**When data is loaded into Snowflake, Snowflake does the following to optimize data storage and query performance. (select all that apply)**

Options are :

- *The data is ordered based on size.*
- *The data is shared between multiple regions.*
- *The data is compressed. (Correct)*
- *The data is converted into a structured file format. (Correct)*
- *The data is converted into a columnar format.*
- *Data is written into micro-partitions based on ingestion order. (Correct)*

**Answer :**

- \*The data is compressed. (Correct)
- \*The data is converted into a structured file format. (Correct)
- \*Data is written into micro-partitions based on ingestion order. (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 2

**Which of the following are data loading best practices? (choose 3)**

Options are :

- *Compress the files before loading (Correct)*
- *Convert to CSV before loading*
- *Use the same warehouse that you will use to query the data to load the data*
- *Aggregate smaller files (Correct)*
- *Combine small files into one large file before loading*
- *Split large files into smaller files (Correct)*

**Answer :**

- \*Compress the files before loading (Correct)
- \*Aggregate smaller files (Correct)
- \*Split large files into smaller files (Correct)

**Which stage types can be used by one or more users to load data into one or more tables? (select all that apply)**

Options are :

- *Multi-stage*
- *Named external (Correct)*

- *Named internal (Correct)*
- *Table*
- *Snowpipe stage*
- *User*

**Answer :**

\*Named external (Correct)  
\*Named internal (Correct)

**What is the minimum edition of Snowflake you would choose if you needed extended time travel and support for multi-factor authentication?**

Options are :

- *Standard*
- *Virtual Private Snowflake*
- *Enterprise (Correct)*
- *Professional*
- *Government*
- *Business Critical*

**Answer :**

\*Enterprise (Correct)

**SnowFlake SnowPro Advanced Architect Practice Exam Set 7**

**What is the largest warehouse size available? Assuming the account is hosted on AWS and preview features enabled.**

Options are :

- *6X-Large (Correct)*
- *4X-Macroscale*
- *7X-Large*
- *4X-Large*
- *5X-Large*

- 6X-Microscale

**Answer :**

\*6X-Large (Correct)

**What are the three main layers that make up the Snowflake architecture?**

Options are :

- Compute (Correct)
- Authentication
- Query optimizer
- Metadata
- Cloud services (Correct)
- Storage (Correct)

**Answer :**

\*Compute (Correct)

\*Cloud services (Correct)

\*Storage (Correct)

**What is a stage? (select all that apply)**

Options are :

- A type of task used to load data into Snowflake
- A database object used to store credentials needed to connect to a cloud provider
- A first class Snowflake object used to configure virtual warehouse usage
- A location to store files before loading them into a table (Correct)
- A reference to the location of data files in cloud storage (Correct)
- A step in the data ingestion pipeline

**Answer :**

\*A location to store files before loading them into a table (Correct)

\*A reference to the location of data files in cloud storage (Correct)

**Snowflake SnowPro Core Certification Practice Test Set 4**

**Which of the following warehouse sizes would you recommend for running queries with multiple joins on large tables?**

Options are :

- *Small*
- *X-Small*
- *X-Large*
- *Medium*
- *2X-Large (Correct)*

**Answer :**

\**2X-Large (Correct)*

**When a table is cloned, the micro-partitions of the original table are duplicated and used by the cloned table. Causing increased storage costs.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\**False (Correct)*

**How many databases and schemas are you allowed to create in Snowflake?**

Options are :

- *1000*
- *10,000*
- *Unlimited (Correct)*
- *4000*
- *100*

**Answer :**

\**Unlimited (Correct)*

**Snowflake SnowPro Core Certification Exam Questions Set 10**

**What encryption method does Snowflake use to encrypt ingested data?**

Options are :

- *AES-128*
- *RSA*
- *Data masking*
- *3DES*
- *AES-256 strong (Correct)*
- *Twofish*

**Answer :**

\*AES-256 strong (Correct)

**What solutions does Snowflake provide for handling semi-structured data? (select all that apply)**

Options are :

- *Transforming and loading the data into separate columns (Correct)*
- *Loading the data into a single column of type VARIANT (Correct)*
- *Automatically creating a table with the correct column definitions and loading the data into this table*
- *Automatically unnesting the data and loading into the appropriate columns*
- *Referencing the data directly using external tables (Correct)*
- *Automatically detecting the schema and column definitions (Correct)*

**Answer :**

\*Transforming and loading the data into separate columns (Correct)  
\*Loading the data into a single column of type VARIANT (Correct)  
\*Automatically creating a table with the correct column definitions and loading the data into this table  
\*Referencing the data directly using external tables (Correct)  
\*Automatically detecting the schema and column definitions (Correct)

**Suppose you run a query on a warehouse that is currently suspended and has auto-resume on. Your query takes 41s to complete. How many seconds were you billed for running this query?**

Options are :

- *101s (Correct)*
- *71s*
- *41s*
- *It depends on whether the query result had been previously cached or not*
- *It depends on whether AUTO\_SUSPEND was set to TRUE or FALSE*

**Answer :**

\*101s (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 11](#)

**On what cadence does Snowflake release new features?**

Options are :

- *Monthly*
- *Daily*
- *Weekly (Correct)*
- *Quarterly*
- *Bi-weekly*
- *There is no defined cadence*

**Answer :**

\*Weekly (Correct)

**Which SQL functions can be used together to detect the schema of semi-structured data and create a table with the column definitions? (choose 2)**

Options are :

- *DETECT\_SCHEMA*
- *GET\_SCHEMA*
- *GENERATE\_COLUMN\_DESCRIPTION (Correct)*
- *GET\_COLUMN\_METADATA*
- *GENERATE\_COLUMN\_DEFINITIONS*

- *INFER\_SCHEMA* (Correct)

**Answer :**

- \**GENERATE\_COLUMN\_DESCRIPTION* (Correct)
- \**INFER\_SCHEMA* (Correct)

**What is the minimum edition of Snowflake needed to use the Failover/Failback feature?**

Options are :

- *Enterprise*
- *Secure*
- *Virtual Private Snowflake*
- *Premier*
- *Standard*
- *Business Critical* (Correct)

**Answer :**

- \**Business Critical* (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 8](#)

**When staging uncompressed files in a Snowflake stage, the files are automatically compressed unless compression is disabled.**

Options are :

- *True* (Correct)
- *False*

**Answer :**

- \**True* (Correct)

**What factors affect the amount of credits charged for a warehouse? (choose 3)**

Options are :

- *The amount of data scanned in GB*
- *The number of warehouses (if multi-cluster)* (Correct)
- *The warehouse size* (Correct)

- *The length of time each warehouse is running (Correct)*
- *The number of users using the warehouse*
- *The number of queries ran per hour*

**Answer :**

- \*The number of warehouses (if multi-cluster) (Correct)
- \*The warehouse size (Correct)
- \*The length of time each warehouse is running (Correct)

**What is the minimum Snowflake edition you would choose if you need 24-hour early access to releases and SOC 2 Type II certification?**

Options are :

- *Virtual Private Snowflake*
- *Snowflake secure*
- *Business Critical*
- *Enterprise (Correct)*
- *Premier*
- *Standard*

**Answer :**

- \*Enterprise (Correct)

### [SnowFlake SnowPro Advanced Architect Practice Exam Set 1](#)

**Which of the following statements about external stages are true? (choose 2)**

Options are :

- *Can be used to upload files from your local file system*
- *External stages use more credits than internal stages*
- *A named external stage is a database object created in a schema (Correct)*
- *External stages are only available to Enterprise accounts or higher*
- *When loading data using an external stage you will incur additional storage costs*
- *An external stage stores the URL to files in cloud storage (Correct)*

**Answer :**

- \*A named external stage is a database object created in a schema (Correct)
- \*An external stage stores the URL to files in cloud storage (Correct)

**What are the advantages to using a storage integration? (choose 2)**

Options are :

- *They make it simpler for users to access data in an S3 bucket (Correct)*
- *They reduce storage costs*
- *They can modify micropartitions to make querying faster*
- *The can be used to integrate data on your local disk*
- *They avoid the need for users to use login credentials when accessing data (Correct)*

**Answer :**

- \*They make it simpler for users to access data in an S3 bucket (Correct)
- \*They avoid the need for users to use login credentials when accessing data (Correct)

**Load metadata expires after 32 days.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

- \*False (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 14****What are the benefits of Snowflake's architecture? (choose 3)**

Options are :

- *It is easy for the user to install updates.*
- *Performance can be easily scaled up or out. (Correct)*
- *Data management is simple. (Correct)*
- *Virtual warehouses can be created for specific use cases, allowing teams to run multiple workloads without queries entering a queue. (Correct)*

- Permissions can be assigned to specific users based on the access they need.
- You are charged compute credits based on the amount (in GBs) of data scanned by your queries.

**Answer :**

- \*Performance can be easily scaled up or out. (Correct)
- \*Data management is simple. (Correct)
- \*Virtual warehouses can be created for specific use cases, allowing teams to run multiple workloads without queries entering a queue. (Correct)

**How long is standard Time Travel?**

Options are :

- 7 days
- 2 days
- 1 month
- 1 week
- 1 day (Correct)

**Answer :**

- \*1 day (Correct)

**Which of the following are options supported by the COPY command for loading files from a stage? (choose 3)**

Options are :

- A list of files (Correct)
- By size
- By date staged
- By number of rows
- Using pattern matching (Correct)
- By path (Correct)

**Answer :**

- \*A list of files (Correct)
- \*Using pattern matching (Correct)
- \*By path (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 11

**What function is recommended to load in JSON data into separate rows of the target table?**

Options are :

- *UNNEST\_JSON*
- *REMOVE\_OUTER\_ARRAY*
- *JSON\_FLATTEN*
- *STRIP\_OUTER\_ARRAY (Correct)*
- *ARRAY\_STRIP*
- *JSON\_PARSE\_TO\_ROW*

**Answer :**

\**STRIP\_OUTER\_ARRAY* (Correct)

**Suppose you are running a small (2 credits/hr) multi-cluster auto-scaling warehouse with 3 warehouses. In the first hour only warehouse 1 runs. In the second hour warehouse 1 and warehouse 2 both run. In the third hour, all three warehouses run. How many total credits will be consumed in these 3 hours?**

Options are :

- *It depends on what the auto-scaling policy is*
- *18*
- *6*
- *12 (Correct)*
- *4*
- *9*

**Answer :**

\**12* (Correct)

**You can host a Snowflake account in multiple regions.**

Options are :

- *False (Correct)*

- *True*

**Answer :**

\*False (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 2

**What is the minimum edition of Snowflake that supports government regions for cloud providers?**

Options are :

- *Virtual Private Snowflake*
- *Enterprise*
- *Government*
- *Business Critical (Correct)*
- *Premium*
- *Standard*

**Answer :**

\*Business Critical (Correct)

**How many credits per hour does a X-large warehouse consume?**

Options are :

- 8
- 4
- 16 (*Correct*)
- 28
- 32

**Answer :**

\*16 (Correct)

**Suppose you run a query (query A) that scans 10Gb of data on a medium warehouse that takes 2 mins. Then you change to a large warehouse and run a query (query B) that scans 20GB of data and takes 1 min. Which query would have consumed more credits?**

Options are :

- *Query B, since it scanned more data*
- *They both consumed the same amount of credits, since the query on the large warehouse took half the time than the query on the medium warehouse (Correct)*
- *Query A, since it took longer to run*

**Answer :**

\*They both consumed the same amount of credits, since the query on the large warehouse took half the time than the query on the medium warehouse (Correct)

### Tackling Snowflake Certification - Practice Questions Set 1

**What are the Snowflake editions? (choose 4)**

Options are :

- *Government*
- *Standard (Correct)*
- *Enterprise (Correct)*
- *Virtual Private Snowflake (Correct)*
- *Business Critical (Correct)*
- *Premium*

**Answer :**

\*Standard (Correct)  
 \*Enterprise (Correct)  
 \*Virtual Private Snowflake (Correct)  
 \*Business Critical (Correct)

**The following url format is valid for an account in the US East region. ira67891.us-east-2.aws.snowflakecomputing.com**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**You will incur additional storage costs when loading data using an external stage.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)



## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 38

SnowPro Core Certification 3 Full Practice Exams 2021 Set 38

What should you consider when choosing the region for your Snowflake account? (choose 4)

Options are :

- *Cost implications (Correct)*
- *Feature availability (Correct)*
- *User access*
- *Data storage costs (Correct)*
- *Latency concerns (Correct)*

**Answer :**

\*Cost implications (Correct)  
\*Feature availability (Correct)  
\*Data storage costs (Correct)  
\*Latency concerns (Correct)

Each successive edition of Snowflake builds on the previous edition.

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**How long does it take Snowpipe to load new data after receiving a notification?**

Options are :

- *3 minutes*
- *10 milliseconds*
- *10 seconds*
- *1 second*
- *1 minute (Correct)*

**Answer :**

\*1 minute (Correct)

### Snowflake Certification Preparation Set 20

**The PUT command can be used to upload local files to tables in Snowflake.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Which of the following is not a viable way to connect to Snowflake?**

Options are :

- *Third-party connectors that can be used to connect to other applications.*
- *Web-based user interface*
- *Native connectors that can be used to develop applications*
- *On private cloud infrastructures (Correct)*
- *Command line clients*
- *ODBC and JDBC drivers that can be used by other applications*

**Answer :**

\*On private cloud infrastructures (Correct)

**Data clustering is the process of collecting metadata on the micro-partitions and sorting it in a way that allows Snowflake to only scan certain micro-partitions when a query is ran.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 1

**Chose the statement that is true regarding how tables are represented in Snowflake?**

Options are :

- *Each schema belongs to a virtual warehouse and each virtual warehouse belongs to a table*
- *Each table belongs to a schema which in turn belongs to a database (Correct)*
- *Each table belongs to a data share and each data share belongs to a schema*
- *Each table belongs to a stage which in turn belongs to a schema*
- *Each table belongs to a database which in turn belongs to a schema*

**Answer :**

\*Each table belongs to a schema which in turn belongs to a database (Correct)

**Snowflake's architecture is a combination of which two database architectures? (choose 2)**

Options are :

- *Shared-nothing (Correct)*
- *Cloud computing*
- *Multi-cluster*
- *Distributed*
- *Shared-disk (Correct)*

**Answer :**

- \*Shared-nothing (Correct)
- \*Shared-disk (Correct)

**Snowflake supports loading data from files staged in which of the following locations? Regardless of the cloud platform of your Snowflake account. (select all that apply)**

Options are :

- *Warehouse stages*
- *Through SnowSQL*
- *Internal stages (Correct)*
- *Microsoft Azure blob storage (Correct)*
- *Google Cloud Storage (Correct)*
- *Amazon S3 (Correct)*

**Answer :**

- \*Internal stages (Correct)
- \*Microsoft Azure blob storage (Correct)
- \*Google Cloud Storage (Correct)
- \*Amazon S3 (Correct)

## [Snowflake SnowPro Core Certification Practice Tests Set 10](#)

**What is a virtual warehouse?**

Options are :

- *A space to store external stages*
- *A cluster of S3 buckets used to store persisted query results*
- *A cluster of compute resources in Snowflake (Correct)*
- *A cluster of databases in Snowflake*
- *A space to store internal stages*

**Answer :**

- \*A cluster of compute resources in Snowflake (Correct)

**How is warehouse load calculated in Snowflake?**

Options are :

- *The average amount of data scanned within a specific time interval*
- *The average amount of time spent executing queries within a specific time interval (Correct)*
- *The average number of credits used within a specific time interval*
- *The average number of queries that were running or queued within a specific time interval*
- *The average number of tables queried within a specific time interval*

**Answer :**

\*The average amount of time spent executing queries within a specific time interval  
(Correct)

**For each table into which data is loaded, what metadata does Snowflake maintain? (select all that apply)**

Options are :

- *Number of columns parsed in the file*
- *File size (Correct)*
- *File size (Correct)*
- *Timestamp of the last load for the file (Correct)*
- *Number of rows parsed in the file (Correct)*
- *ETag for the file (Correct)*

**Answer :**

\*File size (Correct)  
\*File size (Correct)  
\*Timestamp of the last load for the file (Correct)  
\*Number of rows parsed in the file (Correct)  
\*ETag for the file (Correct)

[\*\*SnowPro Core Certification 3 Full Practice Exams 2021 Set 34\*\*](#)

**What are the security layers in Snowflake? (select all that apply)**

Options are :

- *Access (Correct)*
- *Data security (Correct)*
- *Dynamic data masking*
- *Authentication (Correct)*
- *Object security (Correct)*
- *Cloud services*

**Answer :**

- \*Access (Correct)
- \*Data security (Correct)
- \*Authentication (Correct)
- \*Object security (Correct)

**Which of the following contribute to costs in Snowflake? (select all that apply)**

Options are :

- *Cloud services (Correct)*
- *Data storage (Correct)*
- *Average amount of data scanned (in bytes)*
- *Virtual warehouse usage (Correct)*
- *Number of queries ran each day*
- *Number of users in the account*

**Answer :**

- \*Cloud services (Correct)
- \*Data storage (Correct)
- \*Virtual warehouse usage (Correct)

**What are the different stage types in Snowflake? (choose 4)**

Options are :

- *Table stage (Correct)*
- *Named internal stage (Correct)*
- *Named external stage (Correct)*
- *Role stage*

- *Snowpipe stage*
- *User stage (Correct)*

**Answer :**

- \*Table stage (Correct)
- \*Named internal stage (Correct)
- \*Named external stage (Correct)
- \*User stage (Correct)

### [Snowflake Certification Preparation Set 18](#)

**Snowflake runs completely on cloud infrastructure.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

- \*True (Correct)

**A Snowflake account can be hosted on which of the following cloud platforms? (choose 3)**

Options are :

- *Google Cloud Platform (Correct)*
- *VMWare*
- *Microsoft Azure (Correct)*
- *SAP*
- *IBM Watson*
- *Amazon Web Services (Correct)*

**Answer :**

- \*Google Cloud Platform (Correct)
- \*Microsoft Azure (Correct)
- \*Amazon Web Services (Correct)

**Fail-safe is meant to restore data in which of the following events? (choose 3)**

Options are :

- You want compare the changes to a table between two points in time
- A schema is dropped intentionally
- Hardware failure (Correct)
- Natural disaster (Correct)
- System failure (Correct)
- A table is dropped on accident

**Answer :**

- \*Hardware failure (Correct)
- \*Natural disaster (Correct)
- \*System failure (Correct)

### Tackling Snowflake Certification - Practice Questions Set 3

**Which cloud providers currently have limitations? (choose 2)**

Options are :

- SAP
- Microsoft Azure
- Amazon Web Services (Correct)
- Google Cloud Platform (Correct)
- Databricks

**Answer :**

- \*Amazon Web Services (Correct)
- \*Google Cloud Platform (Correct)

**The larger the warehouse, the larger it's cache.**

Options are :

- True (Correct)
- False

**Answer :**

- \*True (Correct)

**Security & Governance is a category of 3rd party data tools that can connect to Snowflake.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

### Snowflake Certification Preparation Set 20

**The data objects stored by Snowflake are directly accessible and visible to customers.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)

**The Snowflake Python connector is a native, pure Python package.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**What services are managed by the Cloud Services layer? (select all that apply)**

Options are :

- *Authentication (Correct)*
- *Query parsing and optimizing (Correct)*
- *Data loading*
- *Metadata management (Correct)*
- *Infrastructure management (Correct)*

- Access control (Correct)

**Answer :**

- \* Authentication (Correct)
- \* Query parsing and optimizing (Correct)
- \* Metadata management (Correct)
- \* Infrastructure management (Correct)
- \* Access control (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 43

**Insert-only streams are only supported for which table type?**

Options are :

- Transient
- Permanent
- Temporary
- Snowflake
- External (Correct)

**Answer :**

- \* External (Correct)

**Which factors impact unit costs for credits? (select all that apply)**

Options are :

- The average number of queries ran each day
- How many users will be accessing the account
- The region where your account is hosted (Correct)
- The Snowflake edition you are using (Correct)
- Whether it is an On Demand or Capacity account (Correct)
- Amount of data scanned by queries (in TBs)

**Answer :**

- \* The region where your account is hosted (Correct)
- \* The Snowflake edition you are using (Correct)
- \* Whether it is an On Demand or Capacity account (Correct)

**Which of the following would you use to create an object with certain rows of a base table?**

Options are :

- *Tri-secret secure*
- *Discretionary access control*
- *Secure view (Correct)*
- *Dynamic data masking*
- *Role based access control*

**Answer :**

\*Secure view (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 41

**To stop being charged storage costs, data needs to leave which state?**

Options are :

- *Time Travel*
- *Materialized*
- *Current*
- *Fail-safe (Correct)*
- *Active*

**Answer :**

\*Fail-safe (Correct)

**What is the minimum edition of Snowflake that supports HITRUST CSF Certification?**

Options are :

- *Business Critical (Correct)*
- *Premium*
- *Certified*
- *Standard*

- *Enterprise*
- *Virtual Private Snowflake*

**Answer :**

\*Business Critical (Correct)

**As long as you are within the defined time period, you can use Time Travel to query data that has been updated or deleted.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 41

**When loading data into Snowflake, Snowflake recommends to use files of roughly which size?**

Options are :

- *100-250 MB uncompressed*
- *150 MB compressed*
- *100-250 KB compressed*
- *100-250 KB uncompressed*
- *150 MB uncompressed*
- *100-250 MB compressed (Correct)*

**Answer :**

\*100-250 MB compressed (Correct)

**Which stage types cannot be altered or dropped? (choose 2)**

Options are :

- *Table (Correct)*
- *Named external*

- *Warehouse*
- *User (Correct)*
- *Named internal*
- *Role*

**Answer :**

\*Table (Correct)  
\*User (Correct)

**What is Snowpipe used for? (choose 2)**

Options are :

- *Loading large volumes of data*
- *Loading small volumes of data frequently (Correct)*
- *Storage optimization*
- *Converting from CSV to JSON*
- *Continuous data loading (Correct)*
- *Unloading data*

**Answer :**

\*Loading small volumes of data frequently (Correct)  
\*Continuous data loading (Correct)

**Snowflake Certification Preparation Set 4**

**What existing data platform is Snowflake built on?**

Options are :

- *Node.js*
- *Apache Arrow*
- *None of the above (Correct)*
- *Hadoop*
- *Oracle DB*

**Answer :**

\*None of the above (Correct)

**What feature of Snowflake's architecture allows data to be accessed easily?**

Options are :

- Administrators can enable SSO to make it easier for users to login and access the data they need.
- Snowflake's virtual warehouses are available to all users.
- Data loaded into Snowflake is not encrypted.
- Snowflake can be accessed via a Web UI.
- Snowflake uses a central repository for data that is accessible by all compute nodes. (Correct)

**Answer :**

\*Snowflake uses a central repository for data that is accessible by all compute nodes. (Correct)

**Which copy options could you use to load files whose load metadata has expired? (choose 2)**

Options are :

- FORCE\_LOAD
- LOAD\_OLD\_FILES
- FORCE (Correct)
- RELOAD\_FILES
- LOAD\_EXPIRED\_FILES
- LOAD\_UNCERTAIN\_FILES (Correct)

**Answer :**

\*FORCE (Correct)

\*LOAD\_UNCERTAIN\_FILES (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 14**

**What Snowflake edition has a dedicated metadata store and pool of compute resources?**

Options are :

- *Virtual Secure Snowflake*
- *Virtual Private Snowflake (Correct)*
- *Professional*
- *Standard*
- *On-demand*
- *Enterprise*

**Answer :**

\*Virtual Private Snowflake (Correct)

**What parameter would you use to load files that match a certain regular expression?**

Options are :

- *PATTERN (Correct)*
- *PATTERN\_REGEX*
- *REGEX\_PATTERN*
- *REGEX*
- *MATCH\_BY*
- *PATTERN\_MATCH*

**Answer :**

\*PATTERN (Correct)

**How often does Snowflake deploy a behavior change release?**

Options are :

- *Daily*
- *Quarterly*
- *Weekly*
- *Bi-weekly*
- *Monthly (Correct)*

**Answer :**

\*Monthly (Correct)

## SnowFlake SnowPro Advanced Architect Practice Exam Set 5

**Which data types are currently supported for schema inference? (choose 3)**

Options are :

- Parquet (Correct)
- Avro (Correct)
- ORC
- TSV
- CSV
- JSON

**Answer :**

\*Parquet (Correct)  
\*Avro (Correct)

**What role is needed to view account level storage use?**

Options are :

- PUBLIC
- SECURITYADMIN
- SYSADMIN
- STORAGEADMIN
- TASKADMIN
- ACCOUNTADMIN (Correct)

**Answer :**

\*ACCOUNTADMIN (Correct)

**Which of the following are examples of semi-structured file formats? (choose 4)**

Options are :

- ORC (Correct)
- JSON (Correct)

- *TSV*
- *CSV*
- *XML (Correct)*
- *Parquet (Correct)*

**Answer :**

- \**ORC (Correct)*
- \**JSON (Correct)*
- \**XML (Correct)*
- \**Parquet (Correct)*

### Tackling Snowflake Certification - Practice Questions Set 2

**Why is Snowflake considered a true SaaS (Software-as-a-Service) offering? (select all that apply)**

Options are :

- *There is no hardware to install. (Correct)*
- *There is virtually no software to install or manage. (Correct)*
- *Snowflake will respond within 24 hours to questions marked as business critical*
- *Snowflake supports role based access control.*
- *The customer doesn't need to worry about ongoing maintenance and tuning. (Correct)*
- *Snowflake gathers metadata on files loaded into tables.*

**Answer :**

- \**There is no hardware to install. (Correct)*
- \**There is virtually no software to install or manage. (Correct)*
- \**The customer doesn't need to worry about ongoing maintenance and tuning. (Correct)*

**When creating a table, what parameter can be used to specify the data retention period?**

Options are :

- *DATA\_RETENTION\_TIME*
- *DATA\_RETENTION\_TIME\_IN\_HOURS*
- *DATA\_RETENTION\_TIME\_IN\_DAYS (Correct)*
- *DATA\_RETENTION\_PERIOD*

- *RETENTION\_TIME*

**Answer :**

\*DATA\_RETENTION\_TIME\_IN\_DAYS (Correct)

**As you increase warehouse size to the next larger size, credit usage is doubled.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)



## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 39

SnowPro Core Certification 3 Full Practice Exams 2021 Set 39

**Virtual warehouses only come in one size.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**To get the latest Snowflake features, users need to make sure to download and install the latest updates.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)

**What determines how many data files can be processed in parallel? (select all that apply)**

Options are :

- *How efficient the micro partitions are pruned*
- *The size of the warehouse you are using (Correct)*
- *The amount of micro partitions in the table*
- *The number of queries currently queued*
- *The amount of compute resources in a warehouse (Correct)*
- *The number of queries being ran currently*

**Answer :**

\*The size of the warehouse you are using (Correct)

\*The amount of compute resources in a warehouse (Correct)

### Snowflake Certification Preparation Set 23

**It is best to use Snowpipe to load a large number of small files rather than a small number of large files.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**You run a query that applies a filter on a specific column and the only micro-partitions scanned are those that contain data that matches the filter. This is an example of what?**

Options are :

- *Query clustering*
- *Micro-partition trimming*
- *Micro-partition pruning*
- *Query pruning (Correct)*
- *Query trimming*
- *Materialized views*

**Answer :**

\*Query pruning (Correct)

**When analyzing warehouse load, it is best to execute queries with similar size and complexity on the same warehouse.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 27

**Files are automatically encrypted when staged in a Snowflake internal location.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**How long is data stored in Fail-safe?**

Options are :

- *1 month*
- *14 days*
- *48 hours*
- *1 day*
- *7 days (Correct)*
- *7 weeks*

**Answer :**

\*7 days (Correct)

**Select the following statements that are true regarding Snowflake Partner Connect. (select all that apply)**

Options are :

- *Muse be set up using the ACCOUNTADMIN role (Correct)*
- *Lets you connect a Snowflake session to another users Snowflake session*
- *Only available to Enterprise accounts or higher*
- *You are required to stay connected to the 3rd-party tool for at least 90 days.*
- *Lets you easily create trial accounts with Snowflake business partners (Correct)*

**Answer :**

\*Muse be set up using the ACCOUNTADMIN role (Correct)

\*Lets you easily create trial accounts with Snowflake business partners (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 38**

**Which of the following describes an On Demand account?**

Options are :

- *Dedicated compute clusters that run on demand*
- *Pricing based on amount of data scanned (in TBs)*
- *Usage-based pricing with no long term commitment (Correct)*
- *Pricing based on daily average of queries ran*
- *Discounted pricing based on an up-front commitment*

**Answer :**

\*Usage-based pricing with no long term commitment (Correct)

**You can create and manage multiple Snowflake accounts, each hosted on a different cloud platform.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Early access is recommended for all organizations.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

### **Snowflake SnowPro Core Certification Practice Test Set 5**

**Which of the following are scaling policies in Snowflake? (select all that apply)**

Options are :

- *Economy (Correct)*
- *Performance*
- *Standard (Correct)*
- *Auto-scale*
- *Maximized*
- *Legacy*

**Answer :**

\*Economy (Correct)

\*Standard (Correct)

**Which cloud platforms have cloud regions that are compliant with US federal privacy and security standards? (choose 2)**

Options are :

- *FedRAMP*
- *Amazon Web Services (Correct)*
- *Oracle Cloud*
- *Microsoft Azure (Correct)*

- Google Cloud Platform

**Answer :**

- \*Amazon Web Services (Correct)
- \*Microsoft Azure (Correct)

**Which transformations are available when loading data using the COPY command? (choose 4)**

Options are :

- Column reordering (Correct)
- Column concatenation
- Column omission (Correct)
- Array flattening
- Truncating text that exceeds the target column length (Correct)
- Casts (Correct)

**Answer :**

- \*Column reordering (Correct)
- \*Column omission (Correct)
- \*Truncating text that exceeds the target column length (Correct)
- \*Casts (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 15](#)

**Suppose you are running queries on a warehouse named compute\_wh\_1 and someone else is loading data using another warehouse named compute\_wh\_2. If they start loading a large file (Multiple Terabytes) then you may start to encounter poor query performance due to the warehouses competing for compute resources.**

Options are :

- True
- False (Correct)

**Answer :**

- \*False (Correct)

**What is the minimum edition of Snowflake needed to use AWS Private link?**

Options are :

- *Business Critical (Correct)*
- *Premier*
- *Standard*
- *Enterprise*
- *Secure*
- *Virtual Private Snowflake*

**Answer :**

\*Business Critical (Correct)

**Which Snowflake editions support data sharing? (select all that apply)**

Options are :

- *Standard (Correct)*
- *Premium*
- *Enterprise (Correct)*
- *Government*
- *Business Critical (Correct)*
- *Virtual Private Snowflake*

**Answer :**

\*Standard (Correct)

\*Enterprise (Correct)

\*Business Critical (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 16**

**When loading a CSV into Snowflake, it is recommended to make sure the number of columns in each row is consistent.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**Snowflake credits are used to pay for using virtual warehouses.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Which types of tables will experience the most benefit from clustering?**

Options are :

- *Tables in the multi-gigabyte range*
- *Tables with more columns than rows*
- *Tables with a small range of values*
- *Tables with more rows than columns*
- *Tables in the multi-terabyte range (Correct)*
- *Tables in the multi-megabyte range*

**Answer :**

\*Tables in the multi-terabyte range (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 2**

**How long does Snowflake retain metadata on COPY INTO commands?**

Options are :

- *14 days (Correct)*
- *64 days*
- *1 year*
- *1 month*
- *Snowflake does not retain metadata for COPY INTO commands*
- *7 days*

**Answer :**

\*14 days (Correct)

**Snowpipe requires a user-specified warehouse.**

Options are :

• *False (Correct)*

• *True*

**Answer :**

\*False (Correct)

**Snowpipe loads files in the same order they are staged.**

Options are :

• *False (Correct)*

• *True*

**Answer :**

\*False (Correct)

### [Snowflake Snowpro Certification Practice Test Set 3](#)

**Resource monitors can be used to restrict the number of credits used by virtual warehouses in your account.**

Options are :

• *True (Correct)*

• *False*

**Answer :**

\*True (Correct)

**The REMOVE command can be used to remove files from a stage after the load completes.**

Options are :

• *True (Correct)*

• *False*

**Answer :**

\*True (Correct)

**User stages have which of the following limitations? (select all that apply)**

Options are :

- They do not support setting file format options (Correct)
- The files can only be accessed by a single user (Correct)
- They cannot be used to upload semi-structured data
- They cannot be altered or dropped (Correct)
- The files can only be loaded into certain tables
- It takes longer to load data into tables using a user stage

**Answer :**

\*They do not support setting file format options (Correct)

\*The files can only be accessed by a single user (Correct)

\*They cannot be altered or dropped (Correct)

### [Snowflake SnowPro Core Certification Practice Tests Set 15](#)

**Which of the following objects is used to track the delta of changes to a table?**

Options are :

- Table\_changes
- Task
- CDC view
- Stream (Correct)
- View
- Storage integration

**Answer :**

\*Stream (Correct)

**It is possible to use network policies to allow certain IP addresses while blocking others.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Multi-cluster warehouses improve query concurrency by allocating additional warehouses to make more compute resources available than you would have with a normal warehouse.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 29

**What data types are supported by Snowpipe?**

Options are :

- *Parquet*
- *TSV*
- *All the above (Correct)*
- *Avro*
- *CSV*
- *JSON*

**Answer :**

\*All the above (Correct)

**What can you do with data exchange in Snowflake? (choose 2)**

Options are :

- *Collaborate with customers and vendors (Correct)*
- *Easily convert data from semi-structured to structured*

- Exchange data between internal business units (Correct)
- Host your Snowflake account across multiple regions
- Modify micro-partitions for use by other business units

**Answer :**

- \*Collaborate with customers and vendors (Correct)
- \*Exchange data between internal business units (Correct)

**Which of the follow can you not do with worksheets as part of the web UI? (choose 2)**

Options are :

- Run ad hoc queries and other DML commands
- Specify a different role for each worksheet
- Resize the current warehouse to increase query completion speed
- Use the PUT command to load data from your local file system (Correct)
- View a chart of credits used each day (Correct)
- Share a saved worksheet with another user (Correct)

**Answer :**

- \*Use the PUT command to load data from your local file system (Correct)
- \*View a chart of credits used each day (Correct)
- \*Share a saved worksheet with another user (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 26](#)

**Which of the following views will have the worst query performance?**

Options are :

- Secure stage view
- Stage view
- Secure materialized view
- Secure non-materialized view (Correct)
- Materialized view
- Non-materialized view

**Answer :**

\*Secure non-materialized view (Correct)

**The process of recording changes (and metadata about the changes) made to a table is referred to as change data capture.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**Which of the following can be used to determine pruning efficiency?**

Options are :

- *Analyzing the clustering depth*
- *Analyzing number of rows returned vs total rows*
- *Analyzing the partitions scanned vs total partitions (Correct)*
- *By querying the PRUNING\_EFFICIENCY table function*
- *Analyzing number of rows scanned vs total number of rows*

**Answer :**

\*Analyzing the partitions scanned vs total partitions (Correct)

### **Snowflake Certification Preparation Set 21**

**A stream allows you to query changes made to a table between two points in time.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Clustering depth measures the average number of overlapping micro partitions for specific columns in a table.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**What table type would you use if you need the table to exist beyond each session but does not need Fail-safe enabled?**

Options are :

- *Provisional*
- *Temporary*
- *External*
- *Permanent*
- *Snowflake*
- *Transient (Correct)*

**Answer :**

\*Transient (Correct)

### **[SnowPro Core Certification 3 Full Practice Exams 2021 Set 9](#)**

**Snowflake automatically saves worksheets as you type.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**What command can be used to see the files that have been uploaded to a stage?**

Options are :

- *DESCRIBE\_STAGE*

- *SHOW\_FILES*
- *SHOW*
- *LIST (Correct)*
- *SHOW\_STAGES*
- *LIST\_FILES*

**Answer :**

\**LIST (Correct)*

**Loading data requires a warehouse.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\**True (Correct)*

### Snowflake Certification Preparation Set 9

**How would you set up a multi-cluster warehouse in auto-scale mode?**

Options are :

- *By setting your session to auto-scale = TRUE*
- *By specifying different values for max and min number of warehouses (Correct)*
- *By setting SCALING\_POLICY = AUTO-SCALE*
- *By specifying the same value for max and min number of warehouses*
- *By changing the warehouse options to auto-scale = TRUE*

**Answer :**

\**By specifying different values for max and min number of warehouses (Correct)*

**There is no event that can trigger a root task, root tasks can only run on a schedule.**

Options are :

- *True (Correct)*

- *False*

**Answer :**

\*True (Correct)

If you wanted to view metadata about the objects in a specific database, where would you look?

Options are :

- *The INFORMATION\_SCHEMA for the database (Correct)*
- *The Web UI*
- *SnowSQL*
- *Run the SHOW OBJECTS command*
- *Run the DESCRIBE OBJECTS command*

**Answer :**

\*The INFORMATION\_SCHEMA for the database (Correct)

**Tools**      **Topics**

Advertisement

## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

Multi\_Factor Authentication (MFA) in Snowflake is only supported in conjunction with single Sign-on (sso). Is this statement True or False?

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

On which of the following cloud platform can a Snowflake account be hosted? Choose 2 answers

Options are :

- *Amazon Web Services (Correct)*

- *Private Virtual Cloud*
- *Oracle Cloud*
- *Microsoft Azure Cloud (Correct)*

**Answer :**

- \*Amazon Web Services (Correct)
- \*Microsoft Azure Cloud (Correct)

**Increasing the size of a Snowflake Virtual Warehouse from an X-Small to an X-Large is an example of:**

Options are :

- *Scaling rhythmically*
- *Scaling max*
- *Scaling out*
- *Scaling up (Correct)*

**Answer :**

- \*Scaling up (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 11](#)

**Fail-safe is unavailable on which Snowflake table types?**

Options are :

- *Temporary (Correct)*
- *Translent (Correct)*
- *Provisional*
- *Permanent*

**Answer :**

- \*Temporary (Correct)
- \*Translent (Correct)

**Which of the following statements is true of data loading in Snowflake?**

Options are :

- Resizing the virtual warehouse from x-Small to Small will process a single file twice as fast
- The "ideal file size for loading is 16MB to match micro-partition size
- Many files in the 10-100 MB range tend to land in the ♦sweet spot" for load parallelism Once loaded, there is no option to force a reload of an already loaded file (Correct)

**Answer :**

\*Many files in the 10-100 MB range tend to land in the ♦sweet spot" for load parallelism Once loaded, there is no option to force a reload of an already loaded file (Correct)

**As a best practice, clustering keys should only be defined on tables of which minimum size?**

Options are :

- Multi-Kilobyte (KB) Range
- Multi-Megabyte (MB) Range
- Multi-Gigabyte (GB) Range
- Multi-Terabyte (TB) Range (Correct)

**Answer :**

\*Multi-Terabyte (TB) Range (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 20

**Data storage for individual tables can be monitored using which Snowflake commands and/or object(s)? Choose 2 answers**

Options are :

- SHOW TABLES; (Correct)
- SHOW STORAGE BY TABLE;
- Information Schema -> TABLE\_STORAGE\_METRICS (Correct)
- Information Schema -> TABLE\_HISTORY

**Answer :**

\*SHOW TABLES; (Correct)

\*Information Schema -> TABLE\_STORAGE\_METRICS (Correct)

**True or False: The user has to specify which cluster a query will run on in Snowflake multi-clustering Warehouse.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**What is the recommended method for loading data into Snowflake?**

Options are :

- *Load row by row*
- *Load data in batch (Correct)*
- *Load data by writing it In the Snowflake Web Interface (UI)*
- *Load data via frequent, angle row DML ↘s*

**Answer :**

\**Load data in batch (Correct)*

### Snowflake Practice Questions - SnowPro Core Certified Set 8

**In which layer of its architecture does Snowflake store its metadata statistics?**

Options are :

- *Storage Layer*
- *Compute Layer*
- *Database Layer*
- *Cloud Service Layer (Correct)*

**Answer :**

\**Cloud Service Layer (Correct)*

**When creating a Snowflake user it is advisable to: Choose 2 answers**

Options are :

- Set the user to be initially disabled
- Force an immediate password change (Correct)
- Set a default role for the user (Correct)
- Set the number of minutes to unlock to 15 minutes
- Set the users access to expire within a specified timeframe

**Answer :**

\*Force an immediate password change (Correct)  
 \*Set a default role for the user (Correct)

**Which Snowflake role has access to the Account and Notifications section in Web UI?**

Options are :

- Public
- Accountadmin (Correct)
- Sysadmin
- Securityadmin

**Answer :**

\*Accountadmin (Correct)

**Snowflake Practice Questions - SnowPro Core Certified Set 6**

**A Snowflake engineer wants to change the warehouse(e.g HR\_Warehouse) that will be used to run a SQL command within a specific worksheet What are the two ways in which engineer can achieve this?**

Options are :

- Run USE WAREHOUSE HR\_Warehouse in Snowflake worksheet or Snowsql (Correct)
- Run SET WAREHOUSE CONTEXT=HR\_Warehouse in Snowflake worksheet or Snowsql
- In Snowflake Worksheet Context Menu select SMALL\_WH warehouse in drop down. (Correct)
- Set the default warehouse at snowflake account level

**Answer :**

\*Run USE WAREHOUSE HR\_Warehouse in Snowflake worksheet or Snowsql (Correct)

\*In Snowflake Worksheet Context Menu select SMALL\_WH warehouse in drop down.  
(Correct)

**What is the max limit on Snowflake child task?**

Options are :

- 50
- 100 (Correct)
- 500
- 1000

**Answer :**

\*100 (Correct)

**Clustering keys for a table can contain only a single column**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 30

**Which of the following operation is NOT performed by a virtual warehouse?**

Options are :

- INSERT
- DELETE
- SNOWPIPE (Correct)
- SELECT

**Answer :**

\*SNOWPIPE (Correct)

**A snowflake user can see the query results of a query executed by another user.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**What are the types of Snowflake queries that a virtual warehouse can execute. Select all that apply.**

Options are :

- *SNOWPIPE*
- *SELECT (Correct)*
- *INSERT (Correct)*
- *COPY INTO (Correct)*

**Answer :**

\**SELECT (Correct)*

\**INSERT (Correct)*

\**COPY INTO (Correct)*

### Snowflake SnowPro Core Certification Exam Questions Set 7

**The snowflake credits used by a virtual warehouses increase proportionately as the size of the virtual warehouse is increased**

Options are :

- *NO, The virtual warehouse size has no implicationon the number of credits being used*
- *YES, The credits usage is tied with the warehouse size (Correct)*

**Answer :**

\**YES, The credits usage is tied with the warehouse size (Correct)*

**What is a benefit of a multi cluster virtual warehouse?**

Options are :

- A multi cluster virtual warehouse will execute more queries for the same amount of Snowflake credit usage
- A multi cluster virtual warehouse can accommodate increased number of users without performance degradation (Correct)
- It is cheaper to run your queries on a multi cluster virtual warehouse
- It removes the need to manually start & stop virtual warehouse to accommodate fluctuating workloads (Correct)

**Answer :**

\*A multi cluster virtual warehouse can accommodate increased number of users without performance degradation (Correct)

\*It removes the need to manually start & stop virtual warehouse to accommodate fluctuating workloads (Correct)

**The Snowflake reader account belongs to?**

Options are :

- The producer (Correct)
- shared ownership
- The consumer
- Neither

**Answer :**

\*The producer (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 21**

**What best describes Snowflake architecture?**

Options are :

- Single Cluster Shared Data
- Multiple Clusters Distributed Data
- Multiple Clusters Shared Data (Correct)
- None of the above

**Answer :**

\*Multiple Clusters Shared Data (Correct)

**When loading data into Snowflake Environment, the Snowflake COPY command supports? Choose 2 answers**

Options are :

- *Joins*
- *Fitters*
- *Data type conversions (Correct)*
- *Column reordering (Correct)*
- *Aggregates*

**Answer :**

\*Data type conversions (Correct)  
\*Column reordering (Correct)

**To run a Multi-Cluster Warehouse in auto-scale mode, a Snowflake user would:**

Options are :

- *Configure the Maximum Clusters setting to ?Auto-Scale?*
- *Set the Warehouse type to ?Auto ?*
- *Set the Minimum Clusters and Maximum Clusters settings to the same value*
- *Set the Minimum Clusters and Maximum Clusters settings to the different values (Correct)*

**Answer :**

\*Set the Minimum Clusters and Maximum Clusters settings to the different values (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 6

**True or false: Snowflake enforces primary key, foreign key and unique constraints during DML operations.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**When a user creates a role, they are initially assigned ownership of the role and they maintain ownership until it is transferred to another user. Is this statement True or False?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which formats are supported for unloading data from Snowflake Environment? Choose 2 answers**

Options are :

- *Delimited (CSV, TSV, etc.) (Correct)*
- *Avro*
- *JSON (Correct)*
- *ORC*

**Answer :**

\*Delimited (CSV, TSV, etc.) (Correct)

\*JSON (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 24**

**Snowflake charges a premium for storing semi-structured data. Is this statement True or False?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following are options when creating a Snowflake Virtual Warehouse?**

Options are :

- *Auto-suspend (Correct)*
- *Auto-resume (Correct)*
- *Local SSD size*
- *User count*

**Answer :**

- \*Auto-suspend (Correct)
- \*Auto-resume (Correct)

**Which of the following are main sections of the top navigation of the Snowflake web Interface (UI)?**

Options are :

- *Database (Correct)*
- *Tables*
- *Warehouses (Correct)*
- *Worksheets (Correct)*

**Answer :**

- \*Database (Correct)
- \*Warehouses (Correct)
- \*Worksheets (Correct)

### [SnowFlake SnowPro Advanced Architect Practice Exam Set 1](#)

**Which are true of Snowflake roles?**

Options are :

- *All grants to objects are given to roles, and never to users (Correct)*
- *In order to do DML/DOL, a user needs to have selected a single role that has that specific access to the object and operation (Correct)*
- *The public role controls all other roles*
- *Roles are a subset of users and users own objects In Snowflake*

**Answer :**

- \*All grants to objects are given to roles, and never to users (Correct)
- \*In order to do DML/DOL, a user needs to have selected a single role that has that specific access to the object and operation (Correct)

**In Snowflake environment Account-level storage usage can be monitored using?**

Options are :

- *The snowflake wet Interface (UI) in the Databases section.*
- *The Snowflake web interface (UI) in the Account -> Billing a usage section (Correct)*
- *The Information Schema -> ACCOUNT\_USAGE\_HISTORY View*
- *The Account usage Schema ↗ > ACCOUNT\_USAGE\_METRICS View*

**Answer :**

\*The Snowflake web interface (UI) in the Account -> Billing a usage section (Correct)

**A Company is planning to use Snowflake. Its engineers wants to use Snowflake Cloing features. Which of the following are common use cases for zero-copy cloning? Choose 3 answers**

Options are :

- *Quick provisioning of Dev and Test/QA environments (Correct)*
- *Data backups (Correct)*
- *Point in time snapshots (Correct)*
- *Performance optimization*

**Answer :**

\*Quick provisioning of Dev and Test/QA environments (Correct)

\*Data backups (Correct)

\*Point in time snapshots (Correct)

#### Snowflake Snowpro Certification Practice Test Set 4

**A solution engineer in a company is trying to minimized Snowflake Cost. Which object will help enginner in limiting the number of credits consumed within a Snowflake account?**

Options are :

- *Account usage Tracking*
- *Resource Monitor (Correct)*
- *Warehouse Limit Parameter*
- *Credit Consumption Tracker*

**Answer :**

\*Resource Monitor (Correct)

**Which of the following statements describes a benefit of Snowflake's separation of compute and storage? (Choose all that apply.)**

Options are :

- *Growth of storage and compute are tightly coupled together (Correct)*
- *Storage expands without the requirement to add more compute*
- *Compute can be scaled up or down without the requirement to add more storage*
- *Multiple compute clusters can access stored data without contention*

**Answer :**

\*Growth of storage and compute are tightly coupled together (Correct)

**What is true of Snowflake datawarehouse network policies? A Snowflake network policy: (Choose two.)**

Options are :

- *Is available to all Snowflake Editions (Correct)*
- *Is only available to customers with Business Critical Edition*
- *Restricts or enables access to specific IP addresses (Correct)*
- *Is activated using an `ALTER DATABASE` command*

**Answer :**

\*Is available to all Snowflake Editions (Correct)

\*Restricts or enables access to specific IP addresses (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 19

**Reader Accounts incur no additional Compute costs to the Data Provider since they are simply reading the shared data without making changes. Is this True or False?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**An engineer is planning to use Snowflake stage to load temporary data. What command is used to load files into an Internal Stage within Snowflake?**

Options are :

- *PUT (Correct)*
- *COPY INTO*
- *TRANSFER*
- *INSERT*

**Answer :**

\*PUT (Correct)

**The number of queries that a Snowflake warehouse can concurrently process is determined by:  
Choose 2 answers**

Options are :

- *The complexity of each query (Correct)*
- *The CONCURRENT\_QUERY\_UMIT parameter set on the Snowflake account*
- *The size of the data required for each query (Correct)*
- *The tool that is executing the query*

**Answer :**

\*The complexity of each query (Correct)

\*The size of the data required for each query (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 37**

**A user has to specify which Snowflake cluster a query will run on in multi-clustering Warehouse.  
Is this statement True or False?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**A company is migrating a Data Warehouse from Oracle to Snowflake. Which item in the Data Warehouse migration process does not apply in Snowflake?**

Options are :

- *Migrate Users*
- *Migrate Schemas*
- *Migrate Indexes (Correct)*
- *Build the Data pipeline*

**Answer :**

\*Migrate Indexes (Correct)

**Snowflake COPY command must specify a File Format directly or indirectly(using stage or during Table definition) in order to execute successfully. Is this statement True or False?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### [Snowflake Certification Preparation Set 9](#)

**A third-party tool that supports standard JDBC or ODBC but has no Snowflake-specific driver will be unable to connect to Snowflake. Is this statement True or False?**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Snowflake Micro-partition metadata enables some operations to be completed without requiring Compute. Is this statement True or False?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**Snowflake database is based on the traditional shared disk architecture used by RDBMS like MySQL, Postgres.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*



## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 40

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 40**

**MFA (Multi-Factor Authentication) is only available in the Enterprise edition or higher.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\**False (Correct)*

**Which view can be used to see the history of data loaded into tables with the COPY INTO command?**

Options are :

- *COPY\\_INTO\\_HISTORY*

- *SUCCESSFUL\_LOAD\_HISTORY*
- *DATA\_LOAD\_HISTORY*
- *TABLE\_STORAGE\_METRICS*
- *LOAD\_HISTORY (Correct)*
- *TABLE\_HISTORY*

**Answer :**

\**LOAD\_HISTORY (Correct)*

**What metadata does Snowflake store about rows stored in a micro-partition? (choose 3)**

Options are :

- *The date the file was loaded*
- *The number of distinct values (Correct)*
- *Min and max values (Correct)*
- *The date modified*
- *The type of compression*
- *The range of values for each of the columns in the micro-partition (Correct)*

**Answer :**

\**The number of distinct values (Correct)*

\**Min and max values (Correct)*

\**The range of values for each of the columns in the micro-partition (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 20](#)

**Which column in an external table is of type VARIANT?**

Options are :

- *Variant*
- *METADATA\$FILENAME*
- *Val*
- *Value (Correct)*
- *Src*

- *Source*

**Answer :**

\*Value (Correct)

**Warehouse resizing is primarily intended for what purpose?**

Options are :

- *Bulk data loading*
- *Improving query performance (Correct)*
- *Handling multiple users*
- *Scaling query concurrency*
- *Scanning different amounts of data*

**Answer :**

\*Improving query performance (Correct)

**An individual task can have more than one parent task.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 34**

**Which of the following is the correct syntax to traverse semi-structured data when querying?  
Assuming you are querying a VARIANT column.**

Options are :

- *: (Correct)*
- *::*
- *..*
- *.value.*

•.

**Answer :**

\*: (Correct)

**What is the name of the command line client for connecting to Snowflake?**

Options are :

- *SnowCLI*
- *Snowpark*
- *Snowpipe*
- *Snowflake connect*
- *Snowflake CLI*
- *SnowSQL (Correct)*

**Answer :**

\*SnowSQL (Correct)

**If you include a named file format in a stage's definition, you don't need to specify file format options when loading data of the appropriate file format.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

### [Snowflake SnowPro Core Certification Practice Test Set 6](#)

**Materialized views are meant to improve performance for processes that use repeated complex queries.**

Options are :

- *Flase*
- *True (Correct)*

**Answer :**

\*True (Correct)

**You are required to create a named file format to copy data into Snowflake.**

Options are :

• *False (Correct)*

• *True*

**Answer :**

\*False (Correct)

**Increasing the size of your warehouse will always improve data loading performance.**

Options are :

• *False (Correct)*

• *True*

**Answer :**

\*False (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 7

**Snowpipe follows a serverless compute model.**

Options are :

• *False*

• *True (Correct)*

**Answer :**

\*True (Correct)

**What happens when a task is still running during the next scheduled execution time?**

Options are :

• *The original task is suspended and the scheduled task is skipped*

• *Both instances of the task are paused until resumed by the user*

- *The original task continues while the scheduled task also starts*
- *The original task continues to run and the next scheduled task is skipped (Correct)*
- *The original task is suspended and the scheduled task starts*

**Answer :**

\*The original task continues to run and the next scheduled task is skipped (Correct)

**Which of the follow table types has the lowest storage costs?**

Options are :

- *Snowflake*
- *Permanent*
- *Transient*
- *Materialized*
- *Temporary*
- *External (Correct)*

**Answer :**

\*External (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 22**

**What is the maximum time-travel retention period for a permanent table in an enterprise edition account?**

Options are :

- *7 weeks*
- *1 month*
- *14 days*
- *90 days (Correct)*
- *7 days*

**Answer :**

\*90 days (Correct)

**What copy option would you use if you wanted to specify what happens when an error is found during data loading?**

Options are :

- *WITH\_ERROR*
- *ERROR*
- *HANDLE\_ERROR*
- *ON\_ERROR (Correct)*
- *ERROR\_OPTIONS*

**Answer :**

\**ON\_ERROR (Correct)*

**Which of the following are advantages of using views? (choose 2)**

Options are :

- *They can improve the clustering of the underlying table*
- *They allow you to show different parts of a table to users with a certain role (Correct)*
- *They can make query writing more simple (Correct)*
- *The definition of the view can be changed easily*
- *They reduce storage costs of the underlying table*

**Answer :**

\**They allow you to show different parts of a table to users with a certain role (Correct)*

\**They can make query writing more simple (Correct)*

### Snowflake SnowPro Core Certification Exam Questions Set 12

**What file formats does Snowflake support for query export? (select all that apply)**

Options are :

- *JSON*
- *TSV (Correct)*
- *Orc*

- CSV (Correct)
- Avro
- Parquet

**Answer :**

- \*TSV (Correct)
- \*CSV (Correct)

Credits are billed on a per-second basis while the warehouse is running, with a 1 min minimum each time the warehouse is resumed.

Options are :

- True (Correct)
- False

**Answer :**

- \*True (Correct)

A stream returns CDC data by storing table data from the source table.

Options are :

- False (Correct)
- True

**Answer :**

- \*False (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 6

Which of the following are semi-structured data formats? (choose 4)

Options are :

- CSV
- JSON (Correct)
- DAG
- Parquet (Correct)
- ORC (Correct)

- Avro (Correct)

**Answer :**

- \*JSON (Correct)
- \*Parquet (Correct)
- \*ORC (Correct)
- \*Avro (Correct)

**Which table function allows you to validate the files loaded in a past execution of the COPY INTO command?**

Options are :

- *VALIDATION\_MODE*
- *VALIDATE (Correct)*
- *SHOW\_VALIDATIONS*
- *VALIDATION*
- *VIEW\_DATA\_LOADS*

**Answer :**

- \**VALIDATE (Correct)*

**Suppose you need to create a materialized view over an external table, what is the minimum Snowflake edition needed?**

Options are :

- *Enterprise (Correct)*
- *Premium*
- *Virtual Private Snowflake*
- *Standard*
- *Business Critical*

**Answer :**

- \**Enterprise (Correct)*

## Snowflake Certification Preparation Set 8

**The query profile shows execution details and statistics for a selected query.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Reclustering in Snowflake is a manual process that needs to be performed by the ACCOUNTADMIN.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**You can write queries more quickly with Snowsight.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 26**

**Which character combination identifies a user stage?**

Options are :

- @\$
- @-
- @
- @% (*Correct*)
- @~

**Answer :**

\*@% (Correct)

**Suppose you need to query a small amount of data within your current session and you don't need other users to be able to access the data. What table type would you use?**

Options are :

- *Provisional*
- *External*
- *Temporary (Correct)*
- *Permanent*
- *Transient*
- *Snowflake*

**Answer :**

\*Temporary (Correct)

**Suppose you load new data into Snowflake, how long will you need to wait until the data is accessible by every virtual warehouse in your account?**

Options are :

- *It depends on the edition of Snowflake you are using*
- *30 mins*
- *After you assign each warehouse to the new table*
- *1 min*
- *The data will be available immediately (Correct)*

**Answer :**

\*The data will be available immediately (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 20**

**How is a materialized view different from a standard view? (choose 3)**

Options are :

- Queries on materialized views are generally faster than queries on non-materialized views (Correct)
- Materialized views are only available for business critical accounts and higher
- Materialized views are only available to users with the SYSADMIN role
- It costs more to maintain a materialized view (Correct)
- A materialized view's results are stored (Correct)

**Answer :**

- \*Queries on materialized views are generally faster than queries on non-materialized views (Correct)
- \*It costs more to maintain a materialized view (Correct)
- \*A materialized view's results are stored (Correct)

**Which of the following are true about masking policies? (choose 2)**

Options are :

- They are currently only available for accounts hosted on AWS
- They provide column level security (Correct)
- They are Schema-level objects (Correct)
- Only available in Business Critical accounts and higher
- They are automatically applied to all temporary tables

**Answer :**

- \*They provide column level security (Correct)
- \*They are Schema-level objects (Correct)

**Which of the following are benefits of micro-partitioning in Snowflake? (select all that apply)**

Options are :

- They don't need to be explicitly defined up-front or maintained (Correct)
- The small size enables fine-grained pruning (Correct)
- The columnar format enables efficient scanning (Correct)
- Users can view micro partitions to improve performance
- They reduce storage costs

**Answer :**

- \*They don't need to be explicitly defined up-front or maintained (Correct)
- \*The small size enables fine-grained pruning (Correct)
- \*The columnar format enables efficient scanning (Correct)

**Snowflake Certification Preparation Set 7**

**Which of the following are advantages of using semi-structured data? (select all that apply)**

Options are :

- *The raw data is easy to view*
- *It is easier to load semi-structured data into Snowflake*
- *A fixed schema is required*
- *Hierarchical data can be easily stored (Correct)*
- *A fixed schema is not required (Correct)*

**Answer :**

- \*Hierarchical data can be easily stored (Correct)
- \*A fixed schema is not required (Correct)

**It is always necessary to specify the file format in the COPY command when loading data.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

- \*False (Correct)

**Which of the following is recorded by a stream object? (choose 3)**

Options are :

- *Table inserts (Correct)*
- *Queries ran on the source table*
- *Table reclustering*
- *Row deletion (Correct)*

- *Table value updates (Correct)*

**Answer :**

- \*Table inserts (Correct)
- \*Row deletion (Correct)
- \*Table value updates (Correct)

### Snowflake Certification Preparation Set 3

**The current role determines which operations can be performed and which objects you can access, based on the privileges granted to the role.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

- \*True (Correct)

**What is the name of the SQL worksheet replacement, designed for data analysis, that was introduced in June 2020?**

Options are :

- *Snowpipe*
- *Snowsight (Correct)*
- *SnowSQL*
- *Snowpark*
- *Snowflake data analysis*
- *Snowchart*

**Answer :**

- \*Snowsight (Correct)

**If you were to load semi-structured data into a Snowflake table with a single column, what data type could that column be?**

Options are :

- *VARCHAR*

- *Semi-structured*
- *VARIANT (Correct)*
- *JSON*
- *STRING*

**Answer :**

\**VARIANT (Correct)*

### [Snowflake SnowPro Core Certification Exam Questions Set 6](#)

**How long does the History page in the classic web UI store details of queries executed?**

Options are :

- *1 year*
- *14 days (Correct)*
- *7 weeks*
- *14 weeks*
- *10 days*
- *7 days*

**Answer :**

\**14 days (Correct)*

**Which of the following does not have a connector or driver for Snowflake?**

Options are :

- *R (Correct)*
- *JDBC*
- *Python*
- *Go*
- *Spark*
- *ODBC*

**Answer :**

\*R (Correct)

The Snowflake web interface wizard is ideal for loading a large number of files.

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 20](#)

It is possible to configure the Fail-safe period using the SET\_FAILSAFE\_TIME\_IN\_DAYS parameter.

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)

What are some advantages to using Snowsight over the classic web interface? (choose 3)

Options are :

- *Smart autocomplete when writing queries (Correct)*
- *Queries complete faster in Snowsight*
- *No upper size limit when downloading results (Correct)*
- *Credits cost slightly less when using Snowsight*
- *Snowsight allows micropartitions to be modified*
- *Query results are cached for up to 1 year (Correct)*

**Answer :**

\*Smart autocomplete when writing queries (Correct)

\*No upper size limit when downloading results (Correct)

\*Query results are cached for up to 1 year (Correct)

**It is not recommended to compress your data files before a large volume of data.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)



## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 41

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 41

**What happens if a virtual warehouse does not have enough resources to process a query?**

Options are :

- *A new virtual warehouse is automatically created to process the query*
- *If the user has multiple Snowflake accounts, the query load is split between warehouses in both accounts*
- *The user is charged additional compute credits due to overloading the warehouse*
- *The query is queued, pending available resources (Correct)*
- *The query is aborted*

**Answer :**

\*The query is queued, pending available resources (Correct)

**What stage type has the greatest degree of flexibility?**

Options are :

- *Named external*
- *Table*
- *Warehouse*
- *Named internal (Correct)*

- *User*
- *Snowstage*

**Answer :**

\*Named internal (Correct)

**What is the recommended way to configure access to a private S3 bucket?**

Options are :

- *Upload the access keys to a table in Snowflake*
- *You cannot configure Snowflake to access data in a private S3 bucket*
- *Using a user stage*
- *Configure an AWS IAM role with the required policies*
- *Configure a storage integration object (Correct)*

**Answer :**

\*Configure a storage integration object (Correct)

**Snowflake SnowPro Core Certification Practice Tests Set 17**

**What is the minimum edition of Snowflake needed to use Tri-secret secure?**

Options are :

- *Secure*
- *Standard*
- *Business Critical (Correct)*
- *Virtual Private Snowflake*
- *Enterprise*

**Answer :**

\*Business Critical (Correct)

**Which of the following is proper syntax to cast a column as a string?**

Options are :

- *column:string*

- *column::string* (*Correct*)
- *column['string']*
- *cast(column) as string*
- *column.string*

**Answer :**

\**column::string* (*Correct*)

**For better query performance, it is recommended to flatten semi-structured data rather than storing it in a VARIANT column.**

Options are :

- *False*
- *True* (*Correct*)

**Answer :**

\**True* (*Correct*)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 11

**Which feature in Snowsight allows you to set keywords for a subquery or list of values?**

Options are :

- *Custom filters* (*Correct*)
- *Keywords*
- *Snowfilter*
- *Query objects*
- *Configurator*

**Answer :**

\**Custom filters* (*Correct*)

**Which command can be used to change the default warehouse for a session?**

Options are :

- *SWITCH WAREHOUSE*

- *SET DEFAULT WAREHOUSE*
- *USE WAREHOUSE (Correct)*
- *CHANGE WAREHOUSE*
- *UPDATE WAREHOUSE*

**Answer :**

\**USE WAREHOUSE (Correct)*

**Which of the following describes a view? (select all that apply)**

Options are :

- *A type of join between 3 or more tables*
- *A common table expression*
- *The query history of a certain user*
- *A named query definition (Correct)*
- *A dashboard used to view the results of a query as a chart*
- *An object that allows the results of a query to be used as if they were in a table (Correct)*

**Answer :**

\**A named query definition (Correct)*

\**An object that allows the results of a query to be used as if they were in a table (Correct)*

### [Snowflake SnowPro Core Certification Practice Tests Set 14](#)

**When does a stream advance it's offset?**

Options are :

- *When the stream is queried*
- *When the SET\_OFFSET parameter is updated*
- *When a file is added the source table's stage*
- *When it is used in a DML transaction (Correct)*
- *When a new table version is created*

**Answer :**

\*When it is used in a DML transaction (Correct)

**Which of the following is not a type of Tech Partner?**

Options are :

- *Data integration*
- *Security and governance*
- *Data labeling (Correct)*
- *Machine learning/data science*
- *Business intelligence*
- *SQL development and management*

**Answer :**

\*Data labeling (Correct)

**A snapshot (taken by a stream) of a table at a specific point in time, is referred to as an offset.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

### Snowflake Certification Preparation Set 17

**Data stored in Fail-safe is only accessible by Snowflake.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**Snowsight supports viewing query results as a chart.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Which of the follow are true regarding micro-partitions? (select all that apply)**

Options are :

- *They are duplicated when a table is cloned*
- *The are organized in a columnar fashion (Correct)*
- *They contain between 50MB and 500MB of uncompressed data (Correct)*
- *They are modified as rows are inserted or removed*
- *The are contiguous units on storage (Correct)*
- *All data in Snowflake tables is stored as micro-partitions (Correct)*

**Answer :**

\*The are organized in a columnar fashion (Correct)

\*They contain between 50MB and 500MB of uncompressed data (Correct)

\*The are contiguous units on storage (Correct)

\*All data in Snowflake tables is stored as micro-partitions (Correct)

## SnowFlake SnowPro Advanced Architect Practice Exam Set 2

**To use the search optimization service, you need to have at least an enterprise edition account.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**The search optimization service is best for tables that have which type of queries ran on them?**

Options are :

- *Queries that select a small number of columns*
- *Queries that don't have a WHERE clause*
- *Queries that return a large number of rows*
- *Queries that perform joins*
- *Queries that return a small number of distinct rows (Correct)*

**Answer :**

\**Queries that return a small number of distinct rows (Correct)*

**When querying a stream, which column would you use to identify the DML operations recorded on the source table?**

Options are :

- *METADATA\$TYPE*
- *METADATA\$ACTION (Correct)*
- *METADATA\$ROW\_ID*
- *METADATA\$DML\_OPS*
- *METADATA\$ISUPDATE*
- *METADATA\$HISTORY*

**Answer :**

\**METADATA\$ACTION (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 33

**Which command would you use to modify the properties of an existing pipe object?**

Options are :

- *ALTER\_PIPE (Correct)*
- *SET\_PIPE*
- *UPDATE\_PIPE*
- *PIPE\_OPTIONS*

- *MODIFY\_PIPE*

**Answer :**

\**ALTER\_PIPE* (Correct)

**MFA (Multi-Factor Authentication) is enabled by default.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\**False* (Correct)

**Which of the following defines a schedule for executing a SQL statement or calling a stored procedure?**

Options are :

- *Task (Correct)*
- *DAG*
- *Cron*
- *Data pipeline*
- *Routine*
- *Stream*

**Answer :**

\**Task* (Correct)

### Snowflake Certification Preparation Set 16

**Which copy option would you use to delete files from the stage after they have been loaded successfully?**

Options are :

- *TRUNCATE*
- *REMOVE*
- *DROP\_FILES*

- *DELETE*
- *PURGE (Correct)*

**Answer :**

\**PURGE (Correct)*

**The @% character combination identifies what kind of stage?**

Options are :

- *Named internal*
- *User*
- *Role*
- *Table (Correct)*
- *Named external*
- *Database*

**Answer :**

\**Table (Correct)*

**Which of the follow are a type of stream? (select all that apply)**

Options are :

- *Append-only (Correct)*
- *Truncate-only*
- *Insert-only (Correct)*
- *Standard (Correct)*
- *Update-only*

**Answer :**

\**Append-only (Correct)*  
\**Insert-only (Correct)*  
\**Standard (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 3**

**What is the minimum Snowflake edition needed to use multi-cluster warehouses?**

Options are :

- *Government*
- *Business Critical*
- *Virtual Private Snowflake*
- *Premium*
- *Standard*
- *Enterprise (Correct)*

**Answer :**

\*Enterprise (Correct)

**You can export query results from the History page in the classic web UI.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**What parameter would you use to check that your files will load correctly?**

Options are :

- *TRY*
- *INSPECT\_FILES*
- *FILE\_CHECK*
- *VALIDATION*
- *CHECK\_FILES*
- *VALIDATION\_MODE (Correct)*

**Answer :**

\*VALIDATION\_MODE (Correct)

**Snowflake SnowPro Core Certification Exam Questions Set 7**

**When a role is dropped, what happens to ownership of objects that were owned by that role?**

Options are :

- Ownership is transferred to the ACCOUNTADMIN role
- The object is unassigned and will be dropped
- Ownership is transferred to the user that executes the DROP ROLE command
- Ownership is unassigned until the ACCOUNTADMIN assigns ownership to a role
- Ownership is transferred to the role that executes the DROP ROLE command  
*(Correct)*

**Answer :**

\*Ownership is transferred to the role that executes the DROP ROLE command (Correct)

**It is always necessary to load data into Snowflake before executing queries.**

Options are :

- False (Correct)
- True

**Answer :**

\*False (Correct)

**How long are query results cached (persisted)?**

Options are :

- 14 hours
- 4 hours
- 7 hours
- 7 days
- 14 days
- 24 hours (Correct)

**Answer :**

\*24 hours (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 26

**What is the minimum auto-suspend time you can set for a warehouse?**

Options are :

- *5 mins*
- *0 or NULL (meaning the warehouse will never suspend) (Correct)*
- *1 min*
- *15 mins*
- *30 seconds*
- *10 mins*

**Answer :**

\*0 or NULL (meaning the warehouse will never suspend) (Correct)

**What copy option would you use to unload a table as a single file?**

Options are :

- *MAX\_FILE\_BYTES*
- *FILE\_SIZE*
- *MIN\_FILE\_SIZE*
- *OVERRIDE\_SIZE*
- *SINGLE (Correct)*
- *MULTIPLE*

**Answer :**

\*SINGLE (Correct)

**A virtual warehouse must be provided by the user in which of the following scenarios? (select all that apply)**

Options are :

- *Automatic clustering*
- *Updating rows in a table (Correct)*

- Counting the number of distinct values in a column
- Running SHOW TABLES
- Loading data using Snowpipe
- Loading data from an external stage (Correct)

**Answer :**

- \*Updating rows in a table (Correct)
- \*Loading data from an external stage (Correct)

### Snowflake Certification Preparation Set 7

**All database objects shared between accounts are read-only.**

Options are :

- True (Correct)
- False

**Answer :**

- \*True (Correct)

**If you want your query to be processed by a certain warehouse, which of the following commands would you run first?**

Options are :

- PROVISION
- SWITCH TO
- USE (Correct)
- SET
- SET WAREHOUSE

**Answer :**

- \*USE (Correct)

**Micro-partitions are immutable.**

Options are :

- False

- *True (Correct)*

**Answer :**

\*True (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 10

**Meeting all the conditions for persisted query results to be used doesn't guarantee that they will be used.**

Options are :

- *True (Correct)*

- *False*

**Answer :**

\*True (Correct)

**Which of the following are true about clustering? (select all that apply)**

Options are :

- *Clustering keys can only be defined by one column or expression (Correct)*
- *Automatic clustering consumes credits*
- *The larger the average clustering depth, the better clustered a table is*
- *Automatic clustering requires a virtual warehouse*
- *Automatic clustering is available in all editions of Snowflake (Correct)*

**Answer :**

\*Clustering keys can only be defined by one column or expression (Correct)

\*Automatic clustering is available in all editions of Snowflake (Correct)

**Which of the following are data types used by Snowflake to represent semi-structured data? (select all that apply)**

Options are :

- *VARIANT (Correct)*
- *Object (Correct)*
- *Key*

- *Dictionary*
- *Array (Correct)*
- *Blob*

**Answer :**

- \*VARIANT (Correct)
- \*Object (Correct)
- \*Array (Correct)

## Snowflake Certification Preparation Set 2

**What command would you use to remove all the rows from a table while preserving the table?**

Options are :

- *TRUNCATE TABLE (Correct)*
- *DELETE*
- *DROP ROWS*
- *CLOSE*
- *CLEAR TABLE*
- *MERGE TABLE*

**Answer :**

- \*TRUNCATE TABLE (Correct)

**You can unload data to an internal stage or external stage.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

- \*True (Correct)

**Snowflake task runs are executed by the task owner.**

Options are :

- *False (Correct)*

- *True*

**Answer :**

\*False (Correct)

### Snowflake Certification Preparation Set 22

**In a simple tree of tasks, all tasks don't need to have the same owner.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**Which of the following commands can be used to view the execution plan for a query?**

Options are :

- *EXPLAIN (Correct)*
- *VIEW\_PLAN*
- *TEST*
- *QUERY\_PLAN*
- *SHOW\_EXECUTION*

**Answer :**

\*EXPLAIN (Correct)

**Which of the following roles is automatically granted to every Snowflake user?**

Options are :

- *DEFAULT*
- *PUBLIC (Correct)*
- *STANDARD*
- *ACCOUNTADMIN*
- *SYSUSER*

- SECURITYADMIN

**Answer :**

\*PUBLIC (Correct)



## 1. SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

Suppose there is a subset of rows in a base table that require significant processing and are used often. Which of the following would have the best performance when accessing this data?

Options are :

- Secure materialized view
- Secure view
- External table
- Materialized view (Correct)
- Non-materialized view

**Answer :**

\*Materialized view (Correct)

Which of the following are serverless features that don't require a virtual warehouse? (choose 3)

Options are :

- Automatic clustering (Correct)
- Search optimization service (Correct)
- Bulk data loading
- Query processing
- Data unloading
- Snowpipe (Correct)

**Answer :**

\*Automatic clustering (Correct)

\*Search optimization service (Correct)

\*Snowpipe (Correct)

**Creating a view of an external table is one way to access data that doesn't incur any Snowflake storage costs.**

Options are :

- *True (Correct)*

- *False*

**Answer :**

\*True (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 1

**In what type of stage is the file path set by concatenating the URL in the stage definition with the file name?**

Options are :

- *Snowpipe*

- *Table*

- *External (Correct)*

- *Internal*

- *User*

**Answer :**

\*External (Correct)

**What is the minimum edition of Snowflake needed to use the search optimization service?**

Options are :

- *Virtual Private Snowflake*

- *Business Critical*

- *Advanced*

- *Standard*

- *Enterprise (Correct)*

- Premier

**Answer :**

\*Enterprise (Correct)

**Which of the following objects contains all the information needed to share a database?**

Options are :

- Export
- Share (Correct)
- External config
- Stage
- Sharing integration
- External stage

**Answer :**

\*Share (Correct)

### Snowflake Certification Preparation Set 9

**Which of the following best describes Snowflakes architecture?**

Options are :

- Multi-cluster shared data (Correct)
- Multi-layered
- Shared data
- SaaS
- Shared disk
- Data warehouse

**Answer :**

\*Multi-cluster shared data (Correct)

**Suppose you needed to access historical query data from 7 months ago, where would you look?**

Options are :

- Contact Snowflake to access the query history
- Use the history tab on the web UI
- Query the QUERY\_HISTORY table function in the information schema
- Query the QUERY\_STATS view using SnowSQL
- Query the ACCOUNT\_USAGE schema inside the SNOWFLAKE system-defined shared database (Correct)

**Answer :**

\*Query the ACCOUNT\_USAGE schema inside the SNOWFLAKE system-defined shared database (Correct)

**A reader account can be used to share data with a consumer who does not have a Snowflake account.**

Options are :

- False
- True (Correct)

**Answer :**

\*True (Correct)

### Snowflake Certification Preparation Set 3

**Defining a clustering key is not recommended for small tables.**

Options are :

- True (Correct)
- False

**Answer :**

\*True (Correct)

**Which of the following are ways to change your current role? (choose 2)**

Options are :

- Using the Snowflake web UI (Correct)
- GRANT ROLE

- *USE ROLE (Correct)*
- *SWITCH ROLE*
- *Ask the ACCOUNTADMIN to assign the role to you*

**Answer :**

- \*Using the Snowflake web UI (Correct)
- \**USE ROLE (Correct)*

**What is the default time limit for a single run of a task?**

Options are :

- *7 days*
- *24 hours*
- *60 mins (Correct)*
- *2 hours*
- *30 mins*

**Answer :**

- \**60 mins (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 30**

**Which function can be used to return the current offset for a stream?**

Options are :

- *SYSTEM\$CURRENT\_OFFSET*
- *SYSTEM\$OFFSET*
- *SYSTEM\$STREAM\_TIMESTAMP*
- *SYSTEM\$STREAM\_GET\_TABLE\_TIMESTAMP (Correct)*
- *GET\_OFFSET*

**Answer :**

- \**SYSTEM\$STREAM\_GET\_TABLE\_TIMESTAMP (Correct)*

**It is possible to load Parquet data into separate columns.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**When does a stream become stale? (choose 2)**

Options are :

- *When the source table is updated*
- *When the micro-partitions of the stream need to be reclustered*
- *When the source table is renamed*
- *When its offset is outside the data retention period for the source table (Correct)*
- *When the historical data of the source table isn't accessible (Correct)*
- *When the offset is hasn't been advanced for more than 2 days*

**Answer :**

\*When its offset is outside the data retention period for the source table (Correct)

\*When the historical data of the source table isn't accessible (Correct)

## Snowflake Certification Preparation Set 21

**What metadata columns can be queried or copied into tables? (choose 2)**

Options are :

- *METADATA\$FILENAME (Correct)*
- *METADATA\$FILE\_ROW\_NUMBER (Correct)*
- *METADATA\$DATE*
- *METADATA\$LOCATION*
- *METADATA\$FILE\_TYPE*
- *METADATA\$FILE\_SIZE*

**Answer :**

- \*METADATA\$FILENAME (Correct)
- \*METADATA\$FILE\_ROW\_NUMBER (Correct)

**For how long does Snowflake retain task activity with the TASK\_HISTORY table function?**

Options are :

- 1 day
- 14 days
- 7 days (Correct)
- 1 week
- 1 month

**Answer :**

- \*7 days (Correct)

**What unloading data from Snowflake, which file formats can you specify? (choose 4)**

Options are :

- TSV (Correct)
- Parquet (Correct)
- Avro
- Orc
- CSV (Correct)
- JSON (Correct)

**Answer :**

- \*TSV (Correct)
- \*Parquet (Correct)
- \*CSV (Correct)
- \*JSON (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 21**

**The data marketplace utilizes Snowflake secure data sharing.**

Options are :

- *True (Correct)*

- *False*

**Answer :**

\*True (Correct)

**A window function allows you to apply a function to a group of related rows.**

Options are :

- *True (Correct)*

- *False*

**Answer :**

\*True (Correct)

**Which of the following are types of data loading in Snowflake? (choose 2)**

Options are :

- *Bulk loading (Correct)*

- *Metadata loading*

- *Rhythmic loading*

- *Piecewise loading*

- *Continuous loading (Correct)*

**Answer :**

\*Bulk loading (Correct)

\*Continuous loading (Correct)

### Snowflake Snowpro Certification Practice Test Set 1

**When using secure data sharing, the consumer account will be charged for storage of the copied data.**

Options are :

- *False (Correct)*

- *True*

**Answer :**

\*False (Correct)

**Which of the following are actions that can be performed by a resource monitor when the threshold has been reached? (choose 3)**

Options are :

- Suspend assigned warehouses after all current queries complete (Correct)
- Suspend assigned warehouses immediately (Correct)
- Send a notification (Correct)
- Change the scaling policy for the warehouse from standard to economy
- Create a new warehouse for load balancing

**Answer :**

\*Suspend assigned warehouses after all current queries complete (Correct)

\*Suspend assigned warehouses immediately (Correct)

\*Send a notification (Correct)

**The default behavior for Snowpipe states that the load will be aborted upon detection of the first error.**

Options are :

- False (Correct)
- True

**Answer :**

\*False (Correct)

### [Snowflake Certification Preparation Set 23](#)

**Which of the following are DML commands? (select all that apply)**

Options are :

- SHOW
- ALTER
- INSERT (Correct)
- UPDATE (Correct)

- *DROP*
- *MERGE (Correct)*

**Answer :**

- \**INSERT (Correct)*
- \**UPDATE (Correct)*
- \**MERGE (Correct)*

**Which of the following are true about file formats in Snowflake? (choose 3)**

Options are :

- *An active warehouse is required to create a file format*
- *They instruct Snowflake how to parse files during a data load (Correct)*
- *They are named objects (Correct)*
- *Only users with the SYSADMIN role can create file formats*
- *It is not recommended to create a file format for bulk loading data*
- *They exist inside of a schema (Correct)*

**Answer :**

- \**They instruct Snowflake how to parse files during a data load (Correct)*
- \**They are named objects (Correct)*
- \**They exist inside of a schema (Correct)*

**What languages can you use to create UDFs in Snowflake? (select all that apply)**

Options are :

- *JavaScript (Correct)*
- *SQL (Correct)*
- *Scala*
- *Python*
- *SnowSQL*
- *Java (Correct)*

**Answer :**

- \**JavaScript (Correct)*

\*SQL (Correct)

\*Java (Correct)

## Snowflake Snowpro Certification Practice Test Set 2

If the conditions are met, the results cache will always be used.

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

Which of the following is not a default function in Snowflake?

Options are :

- *ARRAY\_TO\_STRING*
- *LATERAL FLATTEN*
- *ARRAY\_SIZE*
- *ARRAY\_AGG*
- *ARRAY\_SORT (Correct)*

**Answer :**

\**ARRAY\_SORT (Correct)*

The COPY INTO command allows using a query to further transform data during a load.

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

## SnowFlake SnowPro Advanced Architect Practice Exam Set 7

Time travel is necessary for which of the following functions? (choose 3)

Options are :

- *UNDROP (Correct)*
- *SELECT AT|BEFORE (Correct)*
- *UNCLONE*
- *CLONE AT|BEFORE (Correct)*
- *REFRESH*

**Answer :**

\*UNDROP (Correct)  
\*SELECT AT|BEFORE (Correct)  
\*CLONE AT|BEFORE (Correct)

**Database failover involves replicating your databases to another Snowflake account in a different region due to an outage in your current region.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**What is the size limit for individual rows of the VARIANT data type?**

Options are :

- *128 MB compressed*
- *16 MB uncompressed*
- *256 MB uncompressed*
- *32 MB uncompressed*
- *32 MB compressed*
- *16 MB compressed (Correct)*

**Answer :**

\*16 MB compressed (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 4

**What happens if the predecessor for a child task is removed?**

Options are :

- You cannot remove a predecessor task without removing the child task first
- The former child task is paused
- The former child task becomes either a standalone task or a root task (Correct)
- The former child task is removed
- The former child task is dropped

**Answer :**

\*The former child task becomes either a standalone task or a root task (Correct)

**Which of the following are securable objects? (select all that apply)**

Options are :

- USERS
- FILE FORMAT (Correct)
- TABLE (Correct)
- ROLES
- FUNCTION (Correct)

**Answer :**

\*FILE FORMAT (Correct)

\*TABLE (Correct)

\*FUNCTION (Correct)

**By default, the COPY INTO command will stop loading data when the first error is encountered.**

Options are :

- False
- True (Correct)

**Answer :**

\*True (Correct)

## Tackling Snowflake Certification - Practice Questions Set 2

**Stream can be used to track changes made to materialized views.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

**Which of the following functions is used to extract the string that matches a regular expression pattern?**

Options are :

- *STRING\_EXTRACT*
- *SUBSTRING\_EXTRACT*
- *REGEX\_SUBSTR (Correct)*
- *PATTERN*
- *REGEX\_EXTRACT*

**Answer :**

\**REGEX\_SUBSTR (Correct)*

**What are the layers of caching in Snowflake? (choose 3)**

Options are :

- *External Cache*
- *Remote Disk (Correct)*
- *Local Disk Cache (Correct)*
- *Cloud Cache*
- *Micro-partition Cache*
- *Result Cache (Correct)*

**Answer :**

- \*Remote Disk (Correct)
- \*Local Disk Cache (Correct)
- \*Result Cache (Correct)

### Snowflake Snowpro Certification Practice Test Set 5

**Snowflake users a hybrid model of DAC (discretionary access control) and RBAC (role-based access control) to control access to all objects.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

- \*True (Correct)

**Which feature below dramatically improves the performance of point lookup queries?**

Options are :

- *Snowpipe*
- *Micro-partition pruning*
- *Search optimization service (Correct)*
- *Creating a view of the table*
- *Automatic clustering*

**Answer :**

- \*Search optimization service (Correct)

**DDL commands are used to create, manipulate and modify objects in Snowflake.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

- \*True (Correct)

## Snowflake SnowPro Core Certification Practice Tests Set 10

**Which command is used to unload data files from Snowflake?**

Options are :

- *DOWNLOAD RESULTS*
- *UNLOAD INTO*
- *EXPORT INTO*
- *COPY INTO (Correct)*
- *EXPORT TO*
- *UNLOAD TO*

**Answer :**

\**COPY INTO* (Correct)

**Which construct/operator is required when writing a query?**

Options are :

- *GROUP BY*
- *HAVING*
- *SELECT (Correct)*
- *WHERE*
- *ORDER BY*
- *LIMIT*

**Answer :**

\**SELECT* (Correct)

**If you want to query a table that is not in the schema of your session, you need to fully qualify the table name.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**How long is load metadata stored for each table into which data is loaded?**

Options are :

- *7 days*
- *14 days*
- *1 year*
- *24 hours*
- *64 days (Correct)*

**Answer :**

\*64 days (Correct)

**Once a temporary table has been created, it can only be converted to a transient table.**

Options are :

- *False (Correct)*
- *True*

**Answer :**

\*False (Correct)

**Which of the following is most similar to data integration?**

Options are :

- *Data cleaning*
- *Data migration*
- *Data sanitization*
- *ETL (Correct)*
- *Data ingestion*

**Answer :**

\*ETL (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 4

**Copy options to perform column reordering, column omission and casts using a select statement are only supported for which data type?**

Options are :

- *Parquet*
- *Avro*
- *JSON*
- *Orc*
- *CSV (Correct)*

**Answer :**

\**CSV (Correct)*

**Tasks do not support account or user parameters.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\**True (Correct)*

**The history tab on the web UI only allows you to view successfully completed queries.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\**True (Correct)*

## Tackling Snowflake Certification - Practice Questions Set 3

**Which of the following are the best candidates to be chosen as a clustering key? (choose 2)**

Options are :

- *Columns of the VARIANT data type*
- *Columns that contain dates*
- *Columns used in selective filters (Correct)*
- *Columns frequently used in joins (Correct)*
- *Columns that have a small range of values*
- *Rows that have a small range of values*

**Answer :**

\*Columns used in selective filters (Correct)  
\*Columns frequently used in joins (Correct)

### How is data stored in Snowflake?

Options are :

- *Snowflake stores data in the Parquet file format*
- *Snowflake stores data as compressed JSON*
- *Snowflake stores data in many small CSV files*
- *Snowflake stores the data in an internal, compressed, columnar format (Correct)*
- *Snowflake stores the data in the cloud services layer*

**Answer :**

\*Snowflake stores the data in an internal, compressed, columnar format (Correct)

### What determines where compute resources are provisioned?

Options are :

- *The region the account is hosted (Correct)*
- *The size of the virtual warehouse you are using*
- *The types of queries being ran*
- *The COMPUTE\_RESOURCE\_LOCATION warehouse parameter*
- *The Snowflake edition*

**Answer :**

\*The region the account is hosted (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 10

**Which of the following are valid parameters for the ON\_ERROR copy option? (choose 3)**

Options are :

- *ABORT\_STATEMENT (Correct)*
- *SKIP\_FILE (Correct)*
- *CONTINUE (Correct)*
- *PURGE*
- *FORCE*
- *SKIP*

**Answer :**

\**ABORT\_STATEMENT (Correct)*

\**SKIP\_FILE (Correct)*

\**CONTINUE (Correct)*

**What happens when a virtual warehouse is suspended? (select all that apply)**

Options are :

- *The user can run queries that use the results cache (Correct)*
- *The warehouse can be resumed by running the command USE WAREHOUSE*
- *Queries that were already running are aborted*
- *The cache for the warehouse is cleared (Correct)*
- *You are charged compute credits to process suspending the warehouse*

**Answer :**

\**The user can run queries that use the results cache (Correct)*

\**The cache for the warehouse is cleared (Correct)*

**What happens when query results can't be kept in memory?**

Options are :

- They will spill to the cloud services layer
- They will spill to local storage (Correct)
- The query results are truncated
- They will spill to the remote disk
- The query will be suspended until the warehouse is resized

**Answer :**

\*They will spill to local storage (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 16

**Clustering speeds up query execution time by allowing the query processor to avoid scanning certain micro-partitions.**

Options are :

- True (Correct)
- False

**Answer :**

\*True (Correct)

**Which of the following database objects can be shared with other Snowflake accounts? (select all that apply)**

Options are :

- Stages
- Storage integrations
- Secure views (Correct)
- Materialized views
- External tables (Correct)
- Tables (Correct)

**Answer :**

\*Secure views (Correct)  
 \*External tables (Correct)  
 \*Tables (Correct)

**It is possible to map semi-structured data into separate columns in a target table.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 22**

**What is the top level container in Snowflake?**

Options are :

- *Account (Correct)*
- *Database*
- *Schema*
- *Table*
- *Region*

**Answer :**

\*Account (Correct)

**Where is Snowpipe load history stored?**

Options are :

- *In the metadata of the target table*
- *In the stage used by the pipe*
- *In the information schema*
- *In the metadata of the pipe (Correct)*
- *Load history for Snowpipe is not recorded*

**Answer :**

\*In the metadata of the pipe (Correct)

**What is the maximum number of warehouses in a multi-cluster warehouse?**

Options are :

- 5
- 10 (*Correct*)
- 20
- 15
- 9

**Answer :**

\*10 (Correct)

#### SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

**Dynamic data masking policies are applied before query run time.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**When sharing data between accounts, no actual data is copied between accounts.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Change tracking metadata can be recorded for a table, without creating a stream, by setting CHANGE\_TRACKING = TRUE.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

## Snowflake Certification Preparation Set 2

**Which of the following Snowflake editions has the highest price per credit?**

Options are :

- *Premier*
- *Standard*
- *Professional*
- *Business Critical (Correct)*
- *Enterprise*

**Answer :**

\*Business Critical (Correct)

**What metadata is stored at the table level for permanent tables? (choose 2)**

Options are :

- *Number of distinct rows*
- *Micro-partition count*
- *Table size (in bytes) (Correct)*
- *Row count (Correct)*
- *Column count*

**Answer :**

\*Table size (in bytes) (Correct)

\*Row count (Correct)

**Including the root task, what is the max number of tasks that can make up a simple tree of tasks?**

Options are :

- *500*
- *100*

- 1000 (Correct)
- 999
- 90

**Answer :**

\*1000 (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 29](#)

**Resource monitors can only be created by account administrators and security admins.**

Options are :

- False (Correct)
- True

**Answer :**

\*False (Correct)

**To list files in a table stage, you must use a role that has what privilege on the table?**

Options are :

- USAGE
- OWNERSHIP (Correct)
- MODIFY
- OPERATE
- READ

**Answer :**

\*OWNERSHIP (Correct)

**By default, the COPY INTO command does not purge loaded files from the stage.**

Options are :

- False
- True (Correct)

**Answer :**

\*True (Correct)

## Snowflake Certification Preparation Set 8

**Which function can be used to explode nested values into separate columns?**

Options are :

- *STRIP\_OUTER\_ARRAY*
- *UNNEST*
- *JSON\_FLATTEN*
- *FLATTEN (Correct)*
- *CROSS JOIN*

**Answer :**

\*FLATTEN (Correct)

**It is possible to unload data directly into a S3 bucket.**

Options are :

- *False*
- *True (Correct)*

**Answer :**

\*True (Correct)

**Which of the following is the correct way to use the qualified name of an object?**

Options are :

- ..
- ..
- .. (*Correct*)
- ..
- ..

**Answer :**

\*.. (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 32

**When bulk unloading data using a SELECT statement instead of a tabe, you can use full query syntax including JOINs.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\*True (Correct)

**Which of the following are privileges that can be granted on a warehouse? (choose 4)**

Options are :

- *MODIFY (Correct)*
- *OPERATE (Correct)*
- *SUSPEND*
- *MONITOR (Correct)*
- *RESIZE*
- *USAGE (Correct)*

**Answer :**

\*MODIFY (Correct)  
\*OPERATE (Correct)  
\*MONITOR (Correct)  
\*USAGE (Correct)

**You cannot create a clone of a clone of a table.**

Options are :

- *True*
- *False (Correct)*

**Answer :**

\*False (Correct)

## Snowflake Snowpro Certification Practice Test Set 2

**Which table function would you use to retrieve the 100 most recent task executions in the account?**

Options are :

- *VALIDATE\_TASK*
- *SHOW\_TASKS*
- *TASK\_HISTORY (Correct)*
- *TASK\_EXECUTION\_HISTORY*
- *TASK\_DEPENDENCIES*
- *LIST\_TASK*

**Answer :**

\**TASK\_HISTORY (Correct)*

**Each database in Snowflake contains a default schema named PUBLIC.**

Options are :

- *True (Correct)*
- *False*

**Answer :**

\**True (Correct)*

**Which of the following are not supported functions for COPY transformations? (select all that apply)**

Options are :

- *GROUP BY (Correct)*
- *PARSE\_JSON*
- *TO\_ARRAY*
- *JOIN (Correct)*

- *FLATTEN* (*Correct*)

- *CONCAT*

**Answer :**

\*GROUP BY (*Correct*)

\*JOIN (*Correct*)

\*FLATTEN (*Correct*)

#### Snowflake SnowPro Core Certification Exam Questions Set 4

**You have to use the same role across all your worksheets in Snowflake.**

Options are :

- *True*

- *False* (*Correct*)

**Answer :**

\*False (*Correct*)

**Which command would you use to download files from an internal stage to your local file system?**

Options are :

- *EXPORT*

- *IMPORT*

- *GET* (*Correct*)

- *UNLOAD*

- *DOWNLOAD*

**Answer :**

\**GET* (*Correct*)

**What must happen before a child task in a tree of tasks can be modified?**

Options are :

- *All tasks in the tree of tasks must be suspended*

- *The warehouse running the task must be suspended*

- *The child task must be suspended*

- *The root task must be modified*
- *The root task must be suspended (Correct)*

**Answer :**

\*The root task must be suspended (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 26

**How are storage costs calculated?**

Options are :

- *Average amount of storage used per day (in TBs) before compression*
- *Total amount of storage used (in TBs) per month after compression*
- *Average amount of storage used (in TBs) per month after compression (Correct)*
- *Average amount of micro-partitions used per month*
- *Average amount of storage used (in TBs) per month before compression*

**Answer :**

\*Average amount of storage used (in TBs) per month after compression (Correct)

**A virtual warehouse must be running in order to process SQL queries. For which type of queries results may be produced without requiring a running virtual warehouse?**

Options are :

- *Queries that have previously run and their results are stored in the result cache (Correct)*
- *Queries that make use of clustering keys*
- *Queries that process less than 10 MB of data*
- *Queries for which results can be fulfilled from the metadata cache (Correct)*

**Answer :**

\*Queries that have previously run and their results are stored in the result cache (Correct)

\*Queries for which results can be fulfilled from the metadata cache (Correct)

**The COPY command can load data from which of the following?**

Options are :

- *Detachable hard disks*
- *External stage i.e cloud storage (Correct)*
- *Tape drivers*
- *Internal stage (Correct)*

**Answer :**

\*External stage i.e cloud storage (Correct)  
 \*Internal stage (Correct)

**Queries that have the following characteristics will benefit from clustering. Select all that apply**

Options are :

- *Join on the columns which are part of the cluster key (Correct)*
- *Group on the columns which are part of the cluster key (Correct)*
- *Filter on the columns which are part of the cluster key (Correct)*
- *Sort on the columns which are part of the cluster key (Correct)*
- *Select all rows in the table*
- *Select all columns in the table*

**Answer :**

\*Join on the columns which are part of the cluster key (Correct)  
 \*Group on the columns which are part of the cluster key (Correct)  
 \*Filter on the columns which are part of the cluster key (Correct)  
 \*Sort on the columns which are part of the cluster key (Correct)

**Snowflake SnowPro Core Certification Practice Tests Set 10**

**You have shared a table with another Snowflake account. A user in the consumer account has executed a query on the shared table. Who will be charged for the query cost?**

Options are :

- *Consumer (Correct)*
- *Producer*

**Answer :**

\*Consumer (Correct)

**Snowflake COPY command provides the capability to validate the data inside a file, without actually loading it.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**The size of your Snowflake based data warehouse is virtually unlimited because it uses cloud storage as the underlying storage mechanism.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### [\*\*SnowPro Core Certification 3 Full Practice Exams 2021 Set 4\*\*](#)

**A virtual warehouse can be resized any time, regardless if it is in suspended state or active and executing queries.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**What is the minimum Snowflake edition in which Federated Authentication & MFA are available?**

Options are :

- *Standard (Correct)*
- *Enterprise*
- *Business critical*

- VPS

**Answer :**

\*Standard (Correct)

**A cloned object doesn't contribute to the over all storage unless..**

Options are :

- *The configuration CREATE\_COPY is set to true while cloning data.*
- *Operations that modify data are performed on the cloned table e.g. update, insert or delete data. (Correct)*
- *New populated tables are created in one of the cloned schemas. (Correct)*
- *The cloning is performed by a role which does't have privillages on the source objects.*

**Answer :**

\*Operations that modify data are performed on the cloned table e.g. update, insert or delete data. (Correct)

\*New populated tables are created in one of the cloned schemas. (Correct)

### Snowflake Certification Preparation Set 23

**What are the languages supported by Snowflake user defined functions? Pick 3 options**

Options are :

- SQL (Correct)
- PYTHON
- JAVA (Correct)
- JAVASCRIPT (Correct)

**Answer :**

\*SQL (Correct)

\*JAVA (Correct)

\*JAVASCRIPT (Correct)

**It is mandatory to define a clustering key for permanent tables in Snowflake?**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\**TRUE (Correct)*

**When loading a file through the COPY command, there is no way to partially load a file if errors are encountered. Either the whole file is loaded or nothing is loaded if errors are encountered.**

Options are :

- *TRUE*

- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

### Snowflake Certification Preparation Set 23

**Select all that is true about Snowflake Failsafe.**

Options are :

- *Failsafe is the same as time travel.*

- *Failsafe ensures historical data is protected in the event of a catastrophic failure. (Correct)*

- *Failsafe provides up to 90 days of historical data protection for permanent tables.*

- *Failsafe is distinct & separate from the time travel feature. (Correct)*

- *Failsafe provides up to 7 days of historical data protection for permanent tables. (Correct)*

**Answer :**

\*Failsafe ensures historical data is protected in the event of a catastrophic failure. (Correct)

\*Failsafe is distinct & separate from the time travel feature. (Correct)

\*Failsafe provides up to 7 days of historical data protection for permanent tables. (Correct)

**Micro Partitions are small in size and (before compression) are generally of size?**

Options are :

- 1-10Mb
- 10-50Mb
- 500Mb-1Gb
- 50-500Mb (Correct)

**Answer :**

\*50-500Mb (Correct)

**A Snowflake share must have at least one consumer added to it**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### Snowflake Certification Preparation Set 11

**Please select the correct options that can be used to bring semi structured data into Snowflake.**

Options are :

- Transform the data while using the COPY command. (Correct)
- Load the data as you would load a CSV file, snowflake will make sense of the unstructured data itself
- Load semi structured data into a VARIANT column. (Correct)

**Answer :**

\*Transform the data while using the COPY command. (Correct)

\*Load semi structured data into a VARIANT column. (Correct)

**Micro Partitioning can be disabled for tables if needed.**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Snowflake stores data for each table in a proprietary columnar format.**

Options are :

• *TRUE (Correct)*

• *FALSE*

**Answer :**

\*TRUE (Correct)

### **SnowFlake SnowPro Advanced Architect Practice Exam Set 3**

**When loading data through the COPY command, it is a requirement that your table and the file from which the data is being loaded should have same number of columns.**

Options are :

• *TRUE*

• *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Snowflake stores the following metadata about rows in a micro-partition. Select all that apply**

Options are :

• *The number of distinct value*

• *The range of values for each of the column in the micro partition.*

• *Additional properties for optimization and efficient processing.*

• *All the above (Correct)*

**Answer :**

\*All the above (Correct)

**Which of the following Snowflake commands sets the Virtual Warehouse for a session?**

Options are :

- *COPY WAREHOUSE FROM <>;*
- *SET WAREHOUSE = <>;*
- *USE WAREHOUSE <>; (Correct)*
- *USE VIRTUAL\_WAREHOUSE <>;*

**Answer :**

\**USE WAREHOUSE <>; (Correct)*

## Snowflake Certification Preparation Set 2

**Which object allows you to limit the number of credits consumed within a Snowflake account?**

Options are :

- *Account Usage Tracking*
- *Resource Monitor (Correct)*
- *Warehouse Limit Parameter*
- *Credit Consumption Tracker*

**Answer :**

\**Resource Monitor (Correct)*

**Snowflake is designed for which type of workloads? (Choose two.)**

Options are :

- *OLAP (Analytics) workloads (Correct)*
- *OLTP (Transactional) workloads*
- *Concurrent workloads (Correct)*
- *On-premise workloads*

**Answer :**

\**OLAP (Analytics) workloads (Correct)*

\**Concurrent workloads (Correct)*

**Why would a customer size a Snowflake Virtual Warehouse from an X-Small to a Medium?**

Options are :

- *To accommodate more queries (Correct)*
- *To accommodate more users*
- *To accommodate fluctuations in workload*
- *To accommodate a more complex workload*

**Answer :**

\*To accommodate more queries (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 7

**True or False:** Reader Accounts incur no additional Compute costs to the Data Provider since they are simply reading the shared data without making changes.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which of the following statements describes a benefit of Snowflake's separation of compute and storage? (Choose all that apply.)**

Options are :

- *Growth of storage and compute are tightly coupled together*
- *Storage expands without the requirement to add more compute (Correct)*
- *Compute can be scaled up or down without the requirement to add more storage (Correct)*
- *Multiple compute clusters can access stored data without contention (Correct)*

**Answer :**

\**Storage expands without the requirement to add more compute (Correct)*

\**Compute can be scaled up or down without the requirement to add more storage (Correct)*

\**Multiple compute clusters can access stored data without contention (Correct)*

**True or False:** It is possible to unload structured data to semi-structured formats such as JSON and Parquet.

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 42

**In which layer of its architecture does Snowflake store its metadata statistics?**

Options are :

- *Storage Layer*
- *Compute Layer*
- *Database Layer*
- *Cloud Services Layer (Correct)*

**Answer :**

\**Cloud Services Layer (Correct)*

**True or False: Snowflake's data warehouse was built from the ground up for the cloud in lieu of using an existing database or a platform, like Hadoop, as a base.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**The PUT command: (Choose two.)**

Options are :

- *Automatically creates a File Format object*
- *Automatically uses the last Stage created*
- *Automatically compresses files using Gzip (Correct)*

- Automatically encrypts files (Correct)

**Answer :**

- \*Automatically compresses files using Gzip (Correct)
- \*Automatically encrypts files (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 11**

**Which type of table corresponds to a single Snowflake session?**

Options are :

- Temporary (Correct)
- Transient
- Provisional
- Permanent

**Answer :**

- \*Temporary (Correct)

**Which Snowflake interfaces can be used to create and/or manage Virtual Warehouses?**

Options are :

- The Snowflake Web Interface (UI)
- SQL commands
- Data integration tools
- All of the above (Correct)

**Answer :**

- \*All of the above (Correct)

**When a Pipe is recreated using the CREATE OR REPLACE PIPE command:**

Options are :

- The Pipe load history is reset to empty
- The REFRESH parameter is set to TRUE (Correct)
- Previously loaded files will be ignored

- *All of the above*

**Answer :**

\*The REFRESH parameter is set to TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 25

**What is the minimum Snowflake edition that customers planning on storing protected information in Snowflake should consider for regulatory compliance?**

Options are :

- *Standard*
- *Premier*
- *Enterprise*
- *Business Critical Edition (Correct)*

**Answer :**

\*Business Critical Edition (Correct)

**True or False: Snowpipe via REST API can only reference External Stages as source.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**True or False: It is possible to load data into Snowflake without creating a named File Format object.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### Snowflake Certification Preparation Set 2

**True or False: A table in Snowflake can only be queried using the Virtual Warehouse that was used to load the data.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which of the following statements are true of Snowflake data loading? (Choose three.)**

Options are :

- *VARIANT ◊null◊ values are not the same as SQL NULL values (Correct)*
- *It is recommended to do frequent, single row DMLs*
- *It is recommended to validate the data before loading into the Snowflake target table (Correct)*
- *It is recommended to use staging tables to manage MERGE statements (Correct)*

**Answer :**

\**VARIANT ◊null◊ values are not the same as SQL NULL values (Correct)*

\**It is recommended to validate the data before loading into the Snowflake target table (Correct)*

\**It is recommended to use staging tables to manage MERGE statements (Correct)*

**Which statements are true of micro-partitions? (Choose two.)**

Options are :

- *They are approximately 16MB in size (Correct)*
- *They are stored compressed only if COMPRESS=TRUE on Table*
- *They are immutable (Correct)*
- *They are only encrypted in the Enterprise edition and above*

**Answer :**

\**They are approximately 16MB in size (Correct)*

\**They are immutable (Correct)*

[\*\*Snowflake SnowPro Core Certification Practice Test Set 5\*\*](#)

**A deterministic query is run at 8am, takes 5 minutes, and the results are cached. Which of the following statements are true? (Choose two.)**

Options are :

- *The exact query will ALWAYS return the precomputed result set for the RESULT\_CACHE\_ACTIVE = time period (Correct)*
- *The same exact query will return the precomputed results if the underlying data hasn't changed and the results were last accessed within previous 24 hour period*
- *The same exact query will return the precomputed results even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period (Correct)*
- *The 24 hour timer on the precomputed results gets renewed every time the exact query is executed*

**Answer :**

\*The exact query will ALWAYS return the precomputed result set for the RESULT\_CACHE\_ACTIVE = time period (Correct)

\*The same exact query will return the precomputed results even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period (Correct)

**Increasing the maximum number of clusters in a Multi-Cluster Warehouse is an example of:**

Options are :

- *Scaling rhythmically*
- *Scaling max*
- *Scaling out (Correct)*
- *Scaling up*

**Answer :**

\*Scaling out (Correct)

**Which statement best describes Snowflake tables?**

Options are :

- *Snowflake tables are logical representations of underlying physical data (Correct)*
- *Snowflake tables are the physical instantiation of data loaded into Snowflake*

- Snowflake tables require that clustering keys be defined to perform optimally
- Snowflake tables are owned by a user

**Answer :**

\*Snowflake tables are logical representations of underlying physical data (Correct)

### Snowflake Certification Preparation Set 21

**Which item in the Data Warehouse migration process does not apply in Snowflake?**

Options are :

- Migrate Users
- Migrate Schemas
- Migrate Indexes (Correct)
- Build the Data Pipeline

**Answer :**

\*Migrate Indexes (Correct)

**What is the maximum compressed row size in Snowflake?**

Options are :

- 8KB
- 16MB (Correct)
- 50MB
- 4000GB

**Answer :**

\*16MB (Correct)

**Which of the following are main sections of the top navigation of the Snowflake Web Interface (UI)? (Choose three.)**

Options are :

- Databases (Correct)
- Tables

- Warehouses (Correct)
- Worksheets (Correct)

**Answer :**

- \*Databases (Correct)
- \*Warehouses (Correct)
- \*Worksheets (Correct)

**What is the recommended Snowflake data type to store semi-structured data like JSON?**

Options are :

- VARCHAR
- RAW
- LOB
- VARIANT (Correct)

**Answer :**

- \*VARIANT (Correct)

**If a Small Warehouse is made up of 2 servers/cluster, how many servers/cluster make up a Medium Warehouse?**

Options are :

- 4 (Correct)
- 16
- 32
- 128

**Answer :**

- \*4 (Correct)

**True or False: Snowflake charges additional credits base on usage to Data Providers for the SHARE they create for a Consumer who are not Snowflake Customer?**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 8

**Snowflake provides a mechanism for its customers to override its natural clustering algorithms. This method is:**

Options are :

- *Micro-partitions*
- *Clustering keys (Correct)*
- *Key partitions*
- *Clustered partitions*

**Answer :**

\*Clustering keys (Correct)

**Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)**

Options are :

- *Custom*
- *Economy (Correct)*
- *Optimized*
- *Standard (Correct)*

**Answer :**

\*Economy (Correct)

\*Standard (Correct)

**True or False: A single database can exist in more than one Snowflake account.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 12

**Which of the following roles is recommended to be used to create and manage users and roles?**

Options are :

- *SYSADMIN*
- *SECURITYADMIN (Correct)*
- *PUBLIC*
- *ACCOUNTADMIN*

**Answer :**

\**SECURITYADMIN (Correct)*

**Which statement best describes ◊clustering◊?**

Options are :

- *Clustering represents the way data is grouped together and stored within Snowflake’s micro-partitions (Correct)*
- *The database administrator must define the clustering methodology for each Snowflake table*
- *The clustering key must be included on the COPY command when loading data into Snowflake*
- *Clustering can be disabled within a Snowflake account*

**Answer :**

\**Clustering represents the way data is grouped together and stored within Snowflake’s micro-partitions (Correct)*

**True or False: The COPY command must specify a File Format in order to execute.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

## Snowflake Certification Preparation Set 16

**Query results are stored in the Result Cache for how long after they are last accessed, assuming no data changes have occurred?**

Options are :

- 1 Hour
- 3 Hours
- 12 hours
- 24 hours (Correct)

**Answer :**

\*24 hours (Correct)

**Which of the following connectors are available in the Downloads section of the Snowflake Web Interface (UI)? (Choose two.)**

Options are :

- SnowSQL (Correct)
- ODBC (Correct)
- R
- HIVE

**Answer :**

\*SnowSQL (Correct)

\*ODBC (Correct)

**Which of the following DML commands isn't supported by Snowflake?**

Options are :

- UPSERT (Correct)
- MERGE
- UPDATE
- TRUNCATE TABLE

**Answer :**

\*UPSERT (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 3

**Which of the following statements is true of zero-copy cloning?**

Options are :

- *Zero-copy clones increase storage costs as cloning the table requires storing its data twice*
- *All zero-copy clone objects inherit the privileges of their original objects*
- *Zero-copy cloning is licensed as an additional Snowflake feature*
- *At the instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)*

**Answer :**

\*At the instance a clone is created, all micro-partitions in the original table and the clone are fully shared (Correct)

**A Virtual Warehouse's auto-suspend and auto-resume settings apply to:**

Options are :

- *The primary cluster in the Virtual Warehouse*
- *The entire Virtual Warehouse (Correct)*
- *The database the Virtual Warehouse resides in*
- *The queries currently being run by the Virtual Warehouse*

**Answer :**

\*The entire Virtual Warehouse (Correct)

**Which of the following is true of Snowpipe REST API Option? (Choose two.)**

Options are :

- *You can only use it on Internal Stages*
- *All COPY INTO options are available during pipe creation (Correct)*
- *Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands (Correct)*

- Snowpipe keeps track of which files it has loaded (Correct)

**Answer :**

- \*All COPY INTO options are available during pipe creation (Correct)
- \*Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands (Correct)
- \*Snowpipe keeps track of which files it has loaded (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 12

**Snowflake recommends, as a minimum, that all users with the following role(s) should be enrolled in Multi-Factor Authentication (MFA):**

Options are :

- SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN
- SECURITYADMIN, ACCOUNTADMIN, SYSADMIN
- SECURITYADMIN, ACCOUNTADMIN
- ACCOUNTADMIN (Correct)

**Answer :**

- \*ACCOUNTADMIN (Correct)

**When can a Virtual Warehouse start running queries?**

Options are :

- 12am-5am
- Only during administrator defined time slots
- When its provisioning is complete (Correct)
- After replication

**Answer :**

- \*When its provisioning is complete (Correct)

**True or False: Users are able to see the result sets of queries executed by other users that share their same role.**

Options are :

- TRUE

- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 19

**True or False: Can a Snowflake use pick a cluster to run his query in a multi-cluster Warehouse?**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**True or False: Snowflake Snow Pipes can be Paused and resumed?**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**When should you consider disabling auto-suspend for a Virtual Warehouse? (Choose two.)**

Options are :

- When users will be using compute at different times throughout a 24/7 period
- When managing a steady workload (Correct)
- When the compute must be available with no delay or lag time (Correct)
- When you do not want to have to manually turn on the Warehouse each time a user needs it

**Answer :**

\*When managing a steady workload (Correct)

\*When the compute must be available with no delay or lag time (Correct)

### Snowflake SnowPro Core Certification Practice Test Set 3

**Which of the following simple transformations can be used while loading data through the COPY command?**

Options are :

- *Cast (Correct)*
- *Transpose*
- *Pivot*
- *Truncate (Correct)*
- *Re-order columns (Correct)*
- *Omit columns (Correct)*

**Answer :**

\*Cast (Correct)  
\*Truncate (Correct)  
\*Re-order columns (Correct)  
\*Omit columns (Correct)

**Which of the following scaling type would result in Snowflake preserving credits over performance?**

Options are :

- *Economy (Correct)*
- *Performance*
- *Cost*
- *Standard*

**Answer :**

\*Economy (Correct)

**True or False: Snowflake maintains the partitions using the order of the data in which it is inserted**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 32

**True or False:** If you have multiple virtual warehouses in your Snowflake system, they will each be accessing the same shared data.

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**True or False:** A new custom role will automatically be assigned to all existing users.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following are valid Casting function available in Snowflake?**

Options are :

- *TO\_VARIANT (Correct)*
- *TO\_BLOB*
- *TO\_TIMESTAMP (Correct)*
- *TO\_NUMERIC (Correct)*
- *TO\_NUMBER (Correct)*

**Answer :**

\*TO\_VARIANT (Correct)  
\*TO\_TIMESTAMP (Correct)  
\*TO\_NUMERIC (Correct)  
\*TO\_NUMBER (Correct)

## SnowPro Core Certification 3 Full Practice Exams 2021 Set 23

**What is the maximum time travel period allowed for transient & temporary tables?**

Options are :

- 7
- 90
- 0
- 1 (*Correct*)

**Answer :**

\*1 (Correct)

**Which role is most powerful role in a Snowflake system?**

Options are :

- PUBLIC
- SYSADMIN
- SECURITYADMIN
- ACCOUNTADMIN (*Correct*)
- USERADMIN

**Answer :**

\*ACCOUNTADMIN (Correct)

**True or False: REST API calls can be used to trigger a Snowpipe regardless of whether it is defined over an Internal Stage or an External Stage.**

Options are :

- TRUE (*Correct*)
- FALSE

**Answer :**

\*TRUE (Correct)

## Snowflake Snowpro Certification Practice Test Set 1

**When a database or a schema is cloned, which of the following statements are true for the stages in that database?**

Options are :

- *Table Internal stages are cloned (Correct)*
- *External named stages are cloned (Correct)*
- *Named internal stages are NOT closed (Correct)*
- *Named internal stages are cloned*
- *External named stages are Not cloned*

**Answer :**

\*Table Internal stages are cloned (Correct)  
\*External named stages are cloned (Correct)  
\*Named internal stages are NOT closed (Correct)

**Which of the following objects can be shared? Select all that apply**

Options are :

- *Secure UDFs (Correct)*
- *Database (Correct)*
- *Tables (Correct)*
- *Schema (Correct)*
- *Secure Views (Correct)*
- *Views*

**Answer :**

\*Secure UDFs (Correct)  
\*Database (Correct)  
\*Tables (Correct)  
\*Schema (Correct)  
\*Secure Views (Correct)

**An Engineer increased the size of a virtual warehouse to improve performance of a query? This is a case of?**

Options are :

- *Scale Out*

- *Scale UP (Correct)*
- *Scale High*
- *Scale Max*

**Answer :**

\*Scale UP (Correct)

### [Snowflake Snowpro Certification Practice Test Set 5](#)

**To create a new Task a user should be using a role which**

Options are :

- *Has the CREATE TASK privilege on the schema where the Task is being created (Correct)*
- *Has the SYSADMIN privilege*
- *Has the ACCOUNTADMIN privilege*
- *Has the SECURITYADMIN privilege*

**Answer :**

\*Has the CREATE TASK privilege on the schema where the Task is being created (Correct)

**Snowflake automatically determines the most efficient algorithm to compress columns in a micro partitions**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Which of the following is required in order for a role to be able to successfully execute a Task?**

Options are :

- *The role should have the permissions required to run the SQL statement executed in the Task (Correct)*
- *The role should have ACCOUNTADMIN privillages*

- *The role should have SYSADMIN privillages*
- *The role should have the global EXECUTE TASK privillage on the Account.*

**Answer :**

\*The role should have the permissions required to run the SQL statement executed in the Task (Correct)

**Snowflake Snowpro Certification Practice Test Set 5**

**Snowflake SQL UDFs can return either Scalar or Tabular results?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**Assume a virtual warehouse of size X-Large (128 servers) running for an hour. What is the amount of Snowflake credits that will be billed?**

Options are :

- *64*
- *60*
- *256*
- *128 (Correct)*

**Answer :**

\*128 (Correct)

**Compared to a single cluster virtual warehouse, a multi-cluster virtual warehouse supports which of the following? Select all that apply**

Options are :

- *Auto Suspend (Correct)*
- *Auto Resume (Correct)*
- *Suspend (Correct)*

- *Resize (Correct)*

**Answer :**

- \*Auto Suspend (Correct)
- \*Auto Resume (Correct)
- \*Suspend (Correct)
- \*Resize (Correct)

### SnowFlake SnowPro Advanced Architect Practice Exam Set 7

**A cloned database doesn't inherit the source permissions, but schemas, tables & views contained inside the cloned database will inherit the source permissions.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**Which of the following is not a layer in Snowflake Architecture?**

Options are :

- *Cloud services*
- *Virtual warehouse*
- *Database Storage*
- *Virtual Machines (Correct)*

**Answer :**

- \*Virtual Machines (Correct)

**Snowflake Clustering key helps query performance by?**

Options are :

- *Distributing the data over multiple cluster*
- *Pre-calculating query results*
- *Improving column compression*
- *Pruning unnecessary partitions (Correct)*

**Answer :**

\*Pruning unnecessary partitions (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 4

**True or False:** Except VPS, any snowflake account regardless of the edition, can produce and consume shared data.

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**True or False:** If you create more than one virtual warehouse they will share the memory and CPU resources.

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Snowflake javaScript UDFs can return either scalar or Tabular results?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**A reader account can consume data from sources other than the producer that created the reader account.**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

\*FALSE (Correct)

**Which of the following are true for Snowpipe?**

Options are :

- You must scale a virtual warehouse yourself in order to manage the compute available to Snowpipe
- The resource for Snowpipe are automatically scaled up and down by Snowflake (Correct)
- Snowpipe makes use of the active virtual warehouse for compute resources
- Snowpipe uses server-less compute resources managed by Snowflake (Correct)

**Answer :**

\*The resource for Snowpipe are automatically scaled up and down by Snowflake (Correct)

\*Snowpipe uses server-less compute resources managed by Snowflake (Correct)

**Which of the following are options when creating a new multi-cluster virtual warehouse in Snowflake. Select all that apply**

Options are :

- SUSPEND\_AND\_RESUME
- MAX\_CLUSTER\_COUNT (Correct)
- SCALING\_POLICY (Correct)
- MIN\_CLUSTER\_COUNT (Correct)

**Answer :**

\*MAX\_CLUSTER\_COUNT (Correct)

\*SCALING\_POLICY (Correct)

\*MIN\_CLUSTER\_COUNT (Correct)

**Snowflake SnowPro Core Certification Practice Test Set 4**

**In Snowflake, Cloning a schema will clone which of the following**

Options are :

- *All roles within the schema*
- *All other clone-able objects in the schema (Correct)*
- *All tables in the schema (Correct)*
- *The schema itself (Correct)*

**Answer :**

\*All other clone-able objects in the schema (Correct)  
\*All tables in the schema (Correct)  
\*The schema itself (Correct)

**Snowflake customers can control what format Snowflake stores the data for a table.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**When setting up replication for cross cloud or cross region data sharing, the data provider must replicate data once for each data consumer.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### **Snowflake SnowPro Core Certification Exam Questions Set 7**

**Which function is used to find the query id of the second MOST RECENT query executed in Snwoflake current session?**

Options are :

- *SELECT last\_query\_id(2) (Correct)*

- *SELECT min(query\_id)*
- *SELECT max(query\_id)*
- *SELECT last\_query\_id()*

**Answer :**

\**SELECT last\_query\_id(2)* (Correct)

**To see existing virtual warehouse which names start with bi\_reporting\_vwh, what is the correct command?**

Options are :

- *Show warehouse where name = bi\_reporting\_vwh%*
- *Show warehouse where name like bi\_reporting\_vwh%*
- *Show warehouse like bi\_reporting\_vwh%* (Correct)
- *none of the above*

**Answer :**

\**Show warehouse like bi\_reporting\_vwh%* (Correct)

**The Snowflake credit usage for multi cluster virtual warehouses is based on? Select 2 correct options**

Options are :

- *The number of hour for the warehouse is running* (Correct)
- *The number of cluster running during each hour for the warehouse is running* (Correct)
- *The value that is configured as the maximum in the multi cluster virtual warehouse configuration.*
- *None of the above*

**Answer :**

\**The number of hour for the warehouse is running* (Correct)

\**The number of cluster running during each hour for the warehouse is running* (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 10

**What happens when a table is shared by a snowflake account to another snowflake account? Select all that applies**

Options are :

- Data sharing is managed through Snowflake metadata services layer (Correct)
- The target Snowflake account is charged for the shared data storage
- No actual data is copied or transferred between accounts (Correct)
- Data is copied to the target Snowflake account

**Answer :**

\*Data sharing is managed through Snowflake metadata services layer (Correct)  
\*No actual data is copied or transferred between accounts (Correct)

**To clone a table your current role must have which privilege on the source table?**

Options are :

- Select (Correct)
- Usage
- Write
- Execute

**Answer :**

\*Select (Correct)

**An engineer is creating a virtual warehouse he wants warehouse to run all the time. What is the correct syntax for creating such a virtual warehouse?**

Options are :

- Create or replace warehouse finance\_wh warehouse\_size=large auto\_suspend=0;
- Create or replace warehouse finance\_wh warehouse\_size=large auto\_suspend=1;
- Create or replace warehouse finance\_wh warehouse\_size=large auto\_suspend=null; (Correct)
- Create or replace warehouse finance\_wh warehouse\_size=large never\_suspend=true;

**Answer :**

\*Create or replace warehouse finance\_wh warehouse\_size=large auto\_suspend=null; (Correct)

[\*\*Snowflake SnowPro Core Certification Practice Test Set 5\*\*](#)

**For a snowflake session, can a user specify more than one virtual warehouse at a time for executing a query.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which minimum Snowflake Edition provides up to 90 days of time travel?**

Options are :

- *Enterprise (Correct)*
- *Business Critical*
- *Standard*
- *All Edition Support this feature*

**Answer :**

\**Enterprise (Correct)*

**The privileges provided by the SYSADMIN & SECURITYADMIN role are automatically contained in the ACCOUNTADMIN role since the ACCOUNTADMIN role sits on the top of the role hierarchy.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 14**

**After a clustering key has been defined on a table, no additional administration is required on the table to maintain the clustring.**

Options are :

- *TRUE (Correct)*

- *FALSE*

**Answer :**

\*TRUE (Correct)

If Snowflake data is accidentally deleted and retrieval of data from fail safe is required, you must

Options are :

- *Use SQL to retrieve the data yourself*
- *Use Account icon from WebUI to retrieve the data*
- *Contact snowflake support team to facilitate the retrieval (Correct)*
- *None of the above*

**Answer :**

\*Contact snowflake support team to facilitate the retrieval (Correct)

Snowflake SHOW command requires a running virtual warehouse. Is this statement True or False?

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 40**

On Snowflake Web UI which button on the top bar should you select to see the query history?

Options are :

- *Databases*
- *History (Correct)*
- *Account*
- *Query Profile*
- *Virtual warehouse*
- *Worksheets*

**Answer :**

\*History (Correct)

**Which of the following are options when creating a new Snowflake Virtual Warehouse. Select all that apply**

Options are :

- *resume\_on\_demand*
- *auto\_suspend (Correct)*
- *Suspend\_and\_resume*
- *scaling\_policy (Correct)*
- *resource\_monitor (Correct)*
- *max\_cluster\_count (Correct)*

**Answer :**

\**auto\_suspend (Correct)*  
\**scaling\_policy (Correct)*  
\**resource\_monitor (Correct)*  
\**max\_cluster\_count (Correct)*

**A Snowflake Virtual warehouse can be resized at any time whether they are serving a query or idle.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 20](#)

**A database is created from a share by the consuming account, the access to this database is configurable using the role-based access the Snowflake provides?**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

**What method does Snowflake use to limit the number of micro-partitions scanned during a query execution?**

Options are :

- *In memory operations*
- *Indexing*
- *Table scan*
- *Pruning (Correct)*

**Answer :**

\*Pruning (Correct)

**Snowflake Multi cluster virtual warehouses are designed to handle ... Select all that apply**

Options are :

- *Large Complex Queries*
- *Queuing issues (Correct)*
- *Large number of concurrent queries (Correct)*
- *Large number of concurrent users (Correct)*

**Answer :**

\*Queuing issues (Correct)

\*Large number of concurrent queries (Correct)

\*Large number of concurrent users (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 10**

**Each snowflake share can consist of ...**

Options are :

- *Privileges that grant access to the objects i.e tables, secure views etc. which will be shared. (Correct)*
- *The virtual warehouse that will be used to execute queries on the shared objects.*

- The consumer account with which the database & the objects are to be shared. (Correct)
- Privileges that grant access to a database & a schema from which objects will be shared. (Correct)

**Answer :**

\*Privileges that grant access to the objects i.e tables, secure views etc. which will be shared. (Correct)

\*The consumer account with which the database & the objects are to be shared. (Correct)

\*Privileges that grant access to a database & a schema from which objects will be shared. (Correct)

**What are some general indicators that a clustering key is required on a table? Select all that apply.**

Options are :

- The query performance has slowed down over time. (Correct)
- The size of the table is in multi terabytes (Correct)
- The table has large number of columns
- Snowflake automatically does, no clustering is ever needed.

**Answer :**

\*The query performance has slowed down over time. (Correct)

\*The size of the table is in multi terabytes (Correct)

**Which one of the following can NOT be cloned in Snowflake?**

Options are :

- Internal named stage (Correct)
- external named stage
- table stage

**Answer :**

\*Internal named stage (Correct)

**Snowflake Certification Preparation Set 12**

**For a multi cluster virtual warehouse the administrator set the minimum cluster count and the maximum cluster count to diffrent values. This can be best described as?**

Options are :

- *Manual Scaling*
- *Fixed Scaling*
- *Resizing Scaling*
- *Auto Scaling (Correct)*

**Answer :**

\*Auto Scaling (Correct)

**Which minimum Snowflake edition allows Multicloud Virtual Warehouses capability?**

Options are :

- *Standard*
- *Business critical*
- *Enterprise (Correct)*
- *VPS*

**Answer :**

\*Enterprise (Correct)

**Snowflake uses a unique architecture in which data and compute has been decoupled and both can be scaled independent of each other.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\*TRUE (Correct)

### [Snowflake Certification Preparation Set 4](#)

**Following are the key reasons that make Snowflake a ♦♦True SaaS♦♦ product.**

Options are :

- Snowflake completely runs on cloud infrastructure and is accessible over the internet.
- There is no hardware or software to install, configure, maintain or upgrade at the user's end.
- Snowflake supports key properties of as-a-service offering such as on-demand and instance elasticity, unlimited scalability of compute and storage and pay-per-use.
- All of the above (Correct)

**Answer :**

\*All of the above (Correct)

**Select the correct example of a typical Snowflake URL from the options given below. (select all that apply)**

Options are :

- <https://xu42345.snowflakecomputing.com/> (Correct)
- <https://er35324.us-east-1.snowflakecomputing.com/> (Correct)
- <https://pv53924.ap-south-1.aws.snowflakecomputing.com/> (Correct)
- <https://er35324.us-east-1.azure.snowflakecomputing.com/> (Correct)
- <https://er35324.europe-west2.gcp.snowflakecomputing.com/> (Correct)

**Answer :**

\*<https://xu42345.snowflakecomputing.com/> (Correct)  
 \*<https://er35324.us-east-1.snowflakecomputing.com/> (Correct)  
 \*<https://pv53924.ap-south-1.aws.snowflakecomputing.com/> (Correct)  
 \*<https://er35324.us-east-1.azure.snowflakecomputing.com/> (Correct)  
 \*<https://er35324.europe-west2.gcp.snowflakecomputing.com/> (Correct)

**What are the key considerations for choosing a specific region for your Snowflake account? (Select all that apply)**

Options are :

- Cost (Correct)
- Proximity to the end-users (Correct)
- Data sovereignty (Correct)
- Regulatory requirements (Correct)

**Answer :**

- \*Cost (Correct)
- \*Proximity to the end-users (Correct)
- \*Data sovereignty (Correct)
- \*Regulatory requirements (Correct)

## Snowflake Certification Preparation Set 20

**Which of the following constraints are NOT enforced by Snowflake? (select all that apply)**

Options are :

- Primary Key (Correct)
- Foreign Key (Correct)
- NOT NULL
- Unique (Correct)
- Check constraint (Correct)

**Answer :**

- \*Primary Key (Correct)
- \*Foreign Key (Correct)
- \*Unique (Correct)
- \*Check constraint (Correct)

**Following programming languages/drivers, support is available out-of-the-box (natively) in Snowflake (select all that apply).**

Options are :

- JDBC/ODBC driver (Correct)
- Google Flutter
- Go, Node.js driver (Correct)
- Scala connector
- Python connector (Correct)
- Spark connector (Correct)

**Answer :**

- \*JDBC/ODBC driver (Correct)
- \*Go, Node.js driver (Correct)

\*Python connector (Correct)

\*Spark connector (Correct)

**In the context of a Snowflake virtual warehouse, an example of scaling out is ..... and an example of scaling up is .....**

Options are :

- An example of scaling out is changing the size of compute from M to L; an example of scaling up is adding a new virtual warehouse (VW).
- An example of scaling out is adding a new virtual warehouse (VW); an example of scaling up is changing the size of computing from M to L. (Correct)
- An example of scaling out is adding a new data storage layer; an example of scaling up is changing the size of the storage layer from M to L.
- Both a and c
- Both b and c

**Answer :**

\*An example of scaling out is adding a new virtual warehouse (VW); an example of scaling up is changing the size of computing from M to L. (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 37

**Snowflake follows a 3 stage release process for...**

Options are :

- All releases
- Patch releases
- New feature releases (Correct)
- None of the above

**Answer :**

\*New feature releases (Correct)

**◆ History◆ tab in the Snowflake Web Interface supports viewing of user queries submitted in the past..... day(s)**

Options are :

- 24 hours or 1 day

- 7 days
- 14 days (Correct)
- 90 days

**Answer :**

\*14 days (Correct)

**Which layer of the Snowflake architecture is referred to as the brain of the snowflake data warehouse?**

Options are :

- Compute Layer
- Data Storage Layer
- Cloud Services Layer (Correct)
- Query processing layer

**Answer :**

\*Cloud Services Layer (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 12

**Identify the key features of the Snowflake Micro-partition system. (select all that apply)**

Options are :

- Each micro-partition contains 50-500MB of uncompressed data. (Correct)
- Data stored in micro-partition are compressed and optimized. (Correct)
- Groups of rows in tables are organized in the columnar format in micro-partitions. (Correct)
- Table rows are organized in micro-partitions using the order of insertion. (Correct)
- End users can directly access micro-partitions for their queries.

**Answer :**

\*Each micro-partition contains 50-500MB of uncompressed data. (Correct)

\*Data stored in micro-partition are compressed and optimized. (Correct)

\*Groups of rows in tables are organized in the columnar format in micro-partitions. (Correct)

\*Table rows are organized in micro-partitions using the order of insertion. (Correct)

**Informatica, Matillion, Azure data factory are examples of ..... partners of the Snowflake partner ecosystem.**

Options are :

- *Data Management*
- *Business Intelligence*
- *Data Integration (Correct)*
- *Data Science*

**Answer :**

\*Data Integration (Correct)

**To execute a query in Snowflake, you need to provide context. The context consists of the following. (select four)**

Options are :

- *Virtual Warehouse (VW) (Correct)*
- *Database (Correct)*
- *Schema (Correct)*
- *Role (Correct)*
- *User*
- *Session*

**Answer :**

\*Virtual Warehouse (VW) (Correct)

\*Database (Correct)

\*Schema (Correct)

\*Role (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 8](#)

**What is Snowflake architecture known as?**

Options are :

- *Shared disk architecture*
- *Shared nothing architecture*

- Multi-cluster shared data architecture (Correct)
- None of the above

**Answer :**

\*Multi-cluster shared data architecture (Correct)

**Identify system-defined roles in Snowflake from the roles given below. (select four)**

Options are :

- AccountAdmin (Correct)
- SysAdmin (Correct)
- Auditor
- RoleAdmin
- SecurityAdmin (Correct)
- Public (Correct)

**Answer :**

\*AccountAdmin (Correct)  
\*SysAdmin (Correct)  
\*SecurityAdmin (Correct)  
\*Public (Correct)

**Select the correct object container hierarchy for the objects given below. (down from the highest level object in the hierarchy) 1. Account 2. Schema 3. Database 4. Table**

Options are :

- B >> A >> C >> B
- A >> C >> B >> D (Correct)
- D >> A >> B >> C
- A >> B >> C >> D

**Answer :**

\*A >> C >> B >> D (Correct)

**Select the correct statements from the given set of statements that indicates differences between a view and a materialized view in Snowflake. (select three)**

Options are :

- A view is a named query; a materialized view is unnamed.
- A view does not store query results; a materialized view stores query results for later use. (Correct)
- A view has slower performance; a materialized view has faster performance. (Correct)
- A view does not consume credit; a materialized view consumes credit. (Correct)

**Answer :**

\*A view does not store query results; a materialized view stores query results for later use. (Correct)

\*A view has slower performance; a materialized view has faster performance. (Correct)

\*A view does not consume credit; a materialized view consumes credit. (Correct)

**As a Snowflake data engineer, you have a task to ingest data generated by sensors located on electric poles. As more sensors are getting installed in the city, the ingestion workload is increasing continuously. You have been tasked to support the growth in data loading volume while optimizing the cost of ingestion. Which best practices would you employ to increase the performance of your data ingestion workload? (select all options that apply)**

Options are :

- Increase the size of the virtual warehouse (S >> M >> L) (Correct)
- Split input files into small batches (Correct)
- Using SnowPipe as SnowPipe method is better in ingesting continuous micro-batch workloads (Correct)
- Use SnowSQL as the SnowSQL method is better in ingesting continuous micro-batch workloads

**Answer :**

\*Increase the size of the virtual warehouse (S >> M >> L) (Correct)

\*Split input files into small batches (Correct)

\*Using SnowPipe as SnowPipe method is better in ingesting continuous micro-batch workloads (Correct)

**Which of the following statements are true about Snowflake Query History View? (Snowflake Web UI --> History Tab)**

Options are :

- Snowflake query history view stores queries and profile information for 14 days. (Correct)
- You can download query results from the query history view any time within 24 hours only after the query was executed. (Correct)
- To download query results from query history view, you will need a running Snowflake virtual warehouse.
- You can view other people's queries in the query history view but cannot download results. (Correct)

**Answer :**

\*Snowflake query history view stores queries and profile information for 14 days.

(Correct)

\*You can download query results from the query history view any time within 24 hours only after the query was executed. (Correct)

\*You can view other people's queries in the query history view but cannot download results. (Correct)

### Snowflake SnowPro Core Certification Practice Tests Set 12

**Which of the following statements about Cloning features in Snowflake are true? (select all that apply)**

Options are :

- The cloned data is automatically updated to synchronize with the original data.
- Zero-copy cloning in Snowflake allows users to have multiple copies of data without incurring additional storage costs. (Correct)
- All Snowflake objects can be cloned.
- Zero-copy cloning feature is available in Enterprise edition and above.

**Answer :**

\*Zero-copy cloning in Snowflake allows users to have multiple copies of data without incurring additional storage costs. (Correct)

**Which of the following statements are true about the Snowflake Cache feature?** 1. Snowflake Metadata cache contains object definitions and statistics. 2. Snowflake results cache stores exact results from exact queries for 24 hours. 3. Snowflake will return the resultset from the results cache if the exact same query is executed by different users, even when they have different roles. 4. Snowflake warehouse cache is available only when the Snowflake warehouse is running.

Options are :

- All of the above

- Only 1 and 2
- Only 1, 2 and 3
- Only 1, 2 and 4 (Correct)
- Only 1 and 4

**Answer :**

\*Only 1, 2 and 4 (Correct)

**Which of the following are valid context functions in Snowflake? (select all that apply)**

Options are :

- Select current\_client() (Correct)
- Select current\_account() (Correct)
- Select current\_cloud()
- select current\_user() (Correct)
- select current\_role() (Correct)

**Answer :**

\*Select current\_client() (Correct)  
\*Select current\_account() (Correct)  
\*select current\_user() (Correct)  
\*select current\_role() (Correct)

### Snowflake Practice Questions - SnowPro Core Certified Set 8

**Choose Snowflake interfaces that support multi-factor authentication-based login.**

Options are :

- Snowflake Web UI
- SnowSQL
- Snowflake JDBC and ODBC drivers
- All of the above (Correct)

**Answer :**

\*All of the above (Correct)

**Namespace in Snowflake is comprised of:**

Options are :

- *Database and schema (Correct)*
- *Account, database, and schema*
- *Database, schema, and table*
- *Warehouse, database, and schema*

**Answer :**

\*Database and schema (Correct)

**With Snowflake's time travel feature, you can**

Options are :

- *retrieve and query your data that has been deleted or updated.*
- *create clones (or duplicates) of your tables, schemas, or entire databases as they were at a specific time in the past.*
- *restore tables, schemas, or databases that have been deleted (dropped).*
- *All of the above (Correct)*

**Answer :**

\*All of the above (Correct)

### Snowflake Certification Preparation Set 5

**Default data retention period for time travel in Snowflake is**

Options are :

- *0*
- *1 Day (24 hours) (Correct)*
- *90 Days*
- *It depends on the Snowflake edition*

**Answer :**

\*1 Day (24 hours) (Correct)

**Select the features supported by the Enterprise edition of Snowflake. (select all that apply)**

Options are :

- *Materialized Views (Correct)*
- *Time Travel up to 90 days (Correct)*
- *HIPAA compliance*
- *Customer Managed Keys*

**Answer :**

\*Materialized Views (Correct)  
\*Time Travel up to 90 days (Correct)

**Which of the following security features are supported by the Business-Critical Edition of Snowflake?**

Options are :

- *PCI-DSS, HIPAA, and HITRUST compliance (Correct)*
- *Customer Managed keys through Tri-Secret Secure feature (Correct)*
- *The exclusive Cloud Services layer*
- *Support for Private Link between Snowflake and your VPC (Correct)*

**Answer :**

\*PCI-DSS, HIPAA, and HITRUST compliance (Correct)  
\*Customer Managed keys through Tri-Secret Secure feature (Correct)  
\*Support for Private Link between Snowflake and your VPC (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 34**

**Which Snowflake edition supports 24/7 coverage and a 1-hour response time SLA for Severity 1 issues?**

Options are :

- *Standard Edition*
- *Enterprise Edition*
- *Business Critical Edition*
- *All of the above (Correct)*
- *None of the above*

**Answer :**

\*All of the above (Correct)

**Once the time-travel period has expired, it is possible to request Snowflake support to retrieve historical data for a period of**

Options are :

- *1 Day*
- *7 Days (Correct)*
- *90 Days*
- *It depends on the Snowflake edition*
- *It is user-configurable*

**Answer :**

\*7 Days (Correct)

**In which of the following real-life situations will you use time travel? (select all that apply)**

Options are :

- *To protect against accidental updates of data*
- *To restore past data*
- *To analyze the changes made to data over a period of time*
- *All of the above (Correct)*

**Answer :**

\*All of the above (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 42**

**A BOOLEAN data type in Snowflake can represent (select all that apply)**

Options are :

- *TRUE (Correct)*
- *FALSE (Correct)*
- *NULL (Correct)*

- *BOOLEAN datatype is not supported*

**Answer :**

- \*TRUE (Correct)
- \*FALSE (Correct)
- \*NULL (Correct)

**What would be the output of the following query? select to\_date ('2015-02-28') + INTERVAL '1 year, 1 day';**

Options are :

- *Syntax error: #from # keyword missing*
- *2016-03-01*
- *2016-02-29 (Correct)*
- *None of the above*

**Answer :**

- \*2016-02-29 (Correct)

**When a virtual warehouse is started or resized or resumed, the minimum billing charge is 1 minute.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

### [Snowflake SnowPro Core Certification Exam Questions Set 12](#)

**Your users are executing a large and complex query that takes 60 seconds to complete. They are unhappy with query performance. You are tasked to find out solutions to speed up the query. You look at the virtual warehouse that#s executing the query; it#s an M-size warehouse. You propose to change it to XL- size. What sort of approximate performance benefit do you expect to get by making this change?**

Options are :

- *When the query is run next time, it will complete in approximately 30 seconds.*

- When the query is run next time, it will complete in approximately 15 seconds.  
(Correct)
- It cannot be determined
- None of the above

**Answer :**

\*When the query is run next time, it will complete in approximately 15 seconds.  
(Correct)

**It is a good practice to create multiple virtual warehouses of various sizes in Snowflake based on the workload.**

Options are :

- TRUE (Correct)
- FALSE

**Answer :**

\*TRUE (Correct)

**Following are key features of a multi-cluster virtual warehouse in Snowflake. (select all that apply)**

Options are :

- Each individual cluster in a multi-cluster warehouse can be resized individually and independently of other clusters in the warehouse.
- A multi-cluster warehouse cannot have the same value for the minimum and maximum number of clusters.
- A user can choose when to spin up and spin down an individual cluster in a multi-cluster warehouse.
- A multi-cluster warehouse supports all properties of a single cluster warehouse.  
(Correct)

**Answer :**

\*A multi-cluster warehouse supports all properties of a single cluster warehouse.  
(Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 38

**Which of the following are correct configuration options that you can use while creating a multi-cluster virtual warehouse? (select all that apply)**

Options are :

- *Size of compute (Correct)*
- *Min and Max number of clusters (Correct)*
- *Auto-suspend and Auto-resume (Correct)*
- *Scaling\_Policy (Correct)*
- *Storage Size*

**Answer :**

- \*Size of compute (Correct)
- \*Min and Max number of clusters (Correct)
- \*Auto-suspend and Auto-resume (Correct)
- \*Scaling\_Policy (Correct)

**Multi-cluster warehouses are very effective for improving query concurrency but not as effective in improving individual queries♦ performance.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

- \*TRUE (Correct)

**Following are the two modes in which you can run a multi-cluster virtual warehouse in Snowflake.**

Options are :

- *Dynamic*
- *Maximized (Correct)*
- *Auto-Scale. (Correct)*
- *Multi-threaded*

**Answer :**

- \*Maximized (Correct)
- \*Auto-Scale. (Correct)

**Snowflake SnowPro Core Certification Practice Tests Set 16**

**You will choose a ♦standard♦ scaling policy in Snowflake multi-cluster environment if you expect a high query load and the queries need to be processed with little or no lag.**

Options are :

- *TRUE (Correct)*
- *FALSE*

**Answer :**

\**TRUE (Correct)*

**♦Economy♦ scaling policy in a multi-cluster warehouse... (select all that apply)**

Options are :

- *is applicable when a multi-cluster warehouse is running in auto-scale mode. (Correct)*
- *starts a new cluster as soon as a query is queued.*
- *is applicable when a multi-cluster warehouse is running in maximized mode.*
- *can be specified when a warehouse is created. It cannot be changed afterward.*

**Answer :**

\**is applicable when a multi-cluster warehouse is running in auto-scale mode. (Correct)*

**Query load chart in Snowflake Web interface allows you to**

Options are :

- *find out details about concurrent queries being processed and queries in queued state at any given time in the last two weeks.*
- *identify time periods where compute resource contention occurs (high # of queries in the queued state).*
- *identify sub-optimally utilized warehouses (total query load < 1).*
- *All of the above (Correct)*

**Answer :**

\**All of the above (Correct)*

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 42**

**A multi-cluster virtual warehouse can be configured to suspend automatically after a period of inactivity. However, in order to resume it, SYSADMIN must issue the `ALTER WAREHOUSE .... RESUME` command.**

Options are :

- *TRUE*
- *FALSE (Correct)*

**Answer :**

\**FALSE (Correct)*

**Which of the following programming languages are supported in Snowflake to write user-defined functions (UDFs)?**

Options are :

- *SQL (Correct)*
- *Java*
- *JavaScript (Correct)*
- *Python*

**Answer :**

\**SQL (Correct)*

\**JavaScript (Correct)*

**Select the following statements that are NOT true about transaction management in Snowflake. (select all that apply)**

Options are :

- *Snowflake supports nested transactions. (Correct)*
- *A transaction in Snowflake can span multiple user sessions. (Correct)*
- *All Snowflake transactions must begin with BEGIN TRANSACTION keyword. (Correct)*
- *By default AUTOCOMMIT in Snowflake is ON.*

**Answer :**

\**Snowflake supports nested transactions. (Correct)*

\**A transaction in Snowflake can span multiple user sessions. (Correct)*

\**All Snowflake transactions must begin with BEGIN TRANSACTION keyword. (Correct)*

## Snowflake SnowPro Core Certification Exam Questions Set 4

**Select which of the following statements are NOT true with regards to multi-factor authentication in Snowflake. (select all that apply)**

Options are :

- MFA is available for the Enterprise edition of Snowflake and above.
- MFA is a fully managed service in Snowflake. (Correct)
- ACCOUNTADMIN role can enable MFA feature for your account and enforce at an account level.
- All users need to self-enroll for MFA individually. (Correct)

**Answer :**

\*MFA is a fully managed service in Snowflake. (Correct)

\*All users need to self-enroll for MFA individually. (Correct)

**Select which of the following statements are true with regards to multi-factor authentication in Snowflake. (select all that apply)**

Options are :

- Duo Security Service enables MFA in Snowflake. (Correct)
- Two Snowflake users can have the same cell phone number registered for MFA. (Correct)
- One user can register more than one cellphone number for MFA. So in case, if one cellphone is discharged, the other can be used to log on to Snowflake.
- If you need to change your registered cellphone number, you need to contact Snowflake support.

**Answer :**

\*Duo Security Service enables MFA in Snowflake. (Correct)

\*Two Snowflake users can have the same cell phone number registered for MFA. (Correct)

**Which of the following table types in Snowflake are the most appropriate for storing transitory data (i.e., data that need not be stored for an extended period)? (select all that apply)**

Options are :

- Transient table (Correct)

- *Temporary Table (Correct)*
- *Clone Table*
- *System Table*

**Answer :**

- \*Transient table (Correct)
- \*Temporary Table (Correct)

### [Snowflake Practice Questions - SnowPro Core Certified Set 6](#)

**A transient table in Snowflake (select all that apply)**

Options are :

- *offers the same level of persistence as a permanent table but incurs less cost for large tables. (Correct)*
- *has slower performance compared to a permanent table.*
- *is not accessible to users except the owner of the table.*
- *All of the above*

**Answer :**

- \**offers the same level of persistence as a permanent table but incurs less cost for large tables. (Correct)*

**The following object in Snowflake is also known as the ♦data dictionary♦ of a database.**

Options are :

- *PUBLIC schema*
- *Data Catalog*
- *INFORMATION\_SCHEMA schema (Correct)*
- *None of the above*

**Answer :**

- \**INFORMATION\_SCHEMA schema (Correct)*

**A secure view in Snowflake.... (select all that apply)**

Options are :

- *does not expose view definition. (Correct)*
- *has slower performance than a regular view. (Correct)*
- *cannot be converted to a regular view and vice-versa.*
- *can only be created by the SECURITYADMIN role.*

**Answer :**

\*does not expose view definition. (Correct)  
\*has slower performance than a regular view. (Correct)

### Snowflake Certification Preparation Set 1

**Which of the following statmenets are true regarding Snowflake Scale up and Scale out features?  
Select 2 options**

Options are :

- *Scaling down is a automated process.*
- *Scaling up is a manual process. (Correct)*
- *Scaling out is an automated process. (Correct)*
- *Snapping back is a manual process.*

**Answer :**

\*Scaling up is a manual process. (Correct)  
\*Scaling out is an automated process. (Correct)

**Which of the following can be used to optimize the performance of Snowflake Queries?**

Options are :

- *Table Statistics*
- *Query Profile (Correct)*
- *Explain Plan*
- *Indexes*

**Answer :**

\*Query Profile (Correct)

**Which Snowflake or third party tools can be used to load data to Snowflake?**

Options are :

- *Datapump*
- *Snowsql (Correct)*
- *Snowpipe (Correct)*
- *Sqlloader*
- *Informatica (Correct)*
- *Matillion (Correct)*

**Answer :**

- \**Snowsql (Correct)*
- \**Snowpipe (Correct)*
- \**Informatica (Correct)*
- \**Matillion (Correct)*

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 43](#)

**Which of the following terms are associated with Snowflake Cloud Services layer**

Options are :

- *Query Optimization (Correct)*
- *Query Planning (Correct)*
- *Query Processing*
- *Query Design*
- *Query Compilation (Correct)*

**Answer :**

- \**Query Optimization (Correct)*
- \**Query Planning (Correct)*
- \**Query Compilation (Correct)*

**Which of the following attributes of Snowflake makes it a true SaaS cloud solution**

Options are :

- *No hardware to purchase (Correct)*
- *No user or roles management*

- Patches and Upgrades are automatic (Correct)
- No Storage Costs
- No query processing costs
- Transparent releases (Correct)

**Answer :**

- \*No hardware to purchase (Correct)
- \*Patches and Upgrades are automatic (Correct)
- \*Transparent releases (Correct)

**What kind of Stream can be created on Snowflake External Table?**

Options are :

- Standard
- Append Only
- Insert Only (Correct)
- None of them

**Answer :**

- \*Insert Only (Correct)

### Snowflake Certification Preparation Set 12

**A task does not need a warehouse and Snowflake provides compute for task. Is this statement True or False**

Options are :

- TRUE
- FALSE (Correct)

**Answer :**

- \*FALSE (Correct)

**Which of following describes Snowflake Architecture**

Options are :

- Shared Nothing

- *Shared Data (Correct)*
- *Shared Disk*
- *Shared Memory*
- *Shared Compute*

**Answer :**

\*Shared Data (Correct)

**Which of the following are true about Snowflake Objects? Select all that apply**

Options are :

- *Tables are stored in Schemas (Correct)*
- *Databases are stored in Warehouse*
- *Schemas are stored within warehouse*
- *Warehouse are stored with database*
- *Schemas are stored within database (Correct)*

**Answer :**

\*Tables are stored in Schemas (Correct)

\*Schemas are stored within database (Correct)

**Which of these statement are true about Snowflake Resource Monitor? Select all that apply**

Options are :

*Resource monitor send a notificaiton when a warehouse reaches a credit consumption threshold (Correct)*

*Resource monitor can suspend a warehouse at a specific data and time (Correct)*

*Monitor Snowflake provided serverless warehouse used by snowpipe and Clustering*

*Resource Monitor is used to monitor the credit usage for entire Account as well as individual warehouse (Correct)*

**Answer :**

\*Resource monitor send a notificaiton when a warehouse reaches a credit consumption threshold (Correct)

\*Resource monitor can suspend a warehouse at a specific data and time (Correct)

\*Resource Monitor is used to monitor the credit usage for entire Account as well as individual warehouse (Correct)

**Which of the following are valid tab/menu options in Snowflake Web UI Database section? Select 4 correct option**

Options are :

*Pipes (Correct)*

*Tasks*

*Indexes*

*Tables (Correct)*

*Stages (Correct)*

*File Formats (Correct)*

**Answer :**

\*Pipes (Correct)

\*Tables (Correct)

\*Stages (Correct)

\*File Formats (Correct)

**What is the maximum no of Clusters that can be part of Snowflake Multi-Cluster warehouse?**

Options are :

5

8

20

*10 (Correct)*

16

**Answer :**

\*10 (Correct)

### Snowflake Certification Preparation Set 23

**What are the best practices for loading huge data sets in Snowflake? Select all that apply**

Options are :

*Split large files into files roughly 100-250 MB compressed. (Correct)*

*Aggregate very smaller files to minimize the processing overhead for each file.  
(Correct)*

*Use fewer large files of 10-250 GB to load Terabytes data*

*Dedicate separate warehouses for loading (Correct)*

*Use Multiple Warehouse for the load*

**Answer :**

\*Split large files into files roughly 100-250 MB compressed. (Correct)

\*Aggregate very smaller files to minimize the processing overhead for each file.  
(Correct)

\*Dedicate separate warehouses for loading (Correct)

**Which of the following editions of snowflake supports extended Time Travel of 90 days?**

Options are :

*Developer*

*Enterprise (Correct)*

*Standard*

*Business Critical (Correct)*

*Evaluation*

**Answer :**

\*Enterprise (Correct)

\*Business Critical (Correct)

**Which of the following are a Snowflake feature?**

Options are :

*Zero Copy Cloning & Time Travel (Correct)*

*Data Migration Service*

*Real Time Data Ingestion*

*Row and Column level Security (Correct)*

*Materialized Views (Correct)*

*Data and Result Cache (Correct)*

**Answer :**

- \*Zero Copy Cloning & Time Travel (Correct)
- \*Row and Column level Security (Correct)
- \*Materialized Views (Correct)
- \*Data and Result Cache (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 21**

**Which of the following are valid options to improve performance of Snowflake queries when working on large Tables of Terabyte size? Select 3 Correct options**

Options are :

*Creating indexes on large tables*

*Clustering large tables (Correct)*

*Increasing Size of Warehouse (Correct)*

*Partitioning Large Tables*

*Creating Materialized Views (Correct)*

**Answer :**

- \*Clustering large tables (Correct)
- \*Increasing Size of Warehouse (Correct)
- \*Creating Materialized Views (Correct)

**What is the correct way to find out Snowflake Virtual Warehouse credit usage information?**

Options are :

*Go to Account -> Billing & Usage in the WebUI (Correct)*

*Query WAREHOUSE\_METERING\_HISTORY view in the account\_usage Schema (Correct)*

*Query WAREHOUSE\_METERING\_HISTORY Function in the information Schema (Correct)*

*Query WAREHOUSE\_CREDITS\_USED in the information Schema*

*Query WAREHOUSE\_CREDITS\_USED in the Account Usage Schema*

**Answer :**

- \*Go to Account -> Billing & Usage in the WebUI (Correct)
- \*Query WAREHOUSE\_METERING\_HISTORY view in the account\_usage Schema (Correct)
- \*Query WAREHOUSE\_METERING\_HISTORY Function in the information Schema (Correct)

**What are the resources that a Snowflake Virtual Warehouse provides? Select all that apply.**

Options are :

*CPU (Correct)*

*Temporary Storage (Correct)*

*Memory (Correct)*

*User Management*

**Answer :**

\*CPU (Correct)

\*Temporary Storage (Correct)

\*Memory (Correct)

#### Snowflake Certification Preparation Set 4

**It is possible to share data with any other Snowflake customer regardless of what edition of Snowflake they are running?**

Options are :

*TRUE*

*FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following are valid parameters when create a new virtual warehouse?**

Options are :

*WAREHOUSE\_SIZE (Correct)*

*CPU\_COUNT*

*AUTO\_RESUME (Correct)*

*RAM-SIZE*

*AUTO\_SUSPEND (Correct)*

**Answer :**

\*WAREHOUSE\_SIZE (Correct)

\*AUTO\_RESUME (Correct)  
\*AUTO\_SUSPEND (Correct)

**Snowflake can load from data staged in which of the following location?**

Options are :

- Oracle Cloud Storage*
- Azure Blob Storage (Correct)*
- Internal Stage (Correct)*
- AWS S3 Bucket (Correct)*
- Google Cloud Storage (Correct)*
- VMWare Storage*

**Answer :**

\*Azure Blob Storage (Correct)  
\*Internal Stage (Correct)  
\*AWS S3 Bucket (Correct)  
\*Google Cloud Storage (Correct)

### **Snowflake Practice Questions - SnowPro Core Certified Set 2**

**Snowflake does not start executing any queries on a new virtual warehouse untill all of the servers are provisioned.**

Options are :

- TRUE (Correct)*
- FALSE*

**Answer :**

\*TRUE (Correct)

**The "loadmetadata" (which maintains which files have already been loaded) for a table expires after how many days?**

Options are :

- 30*
- 1*

*64 (Correct)*

*365*

**Answer :**

\*64 (Correct)

**Which of the following are options when created a new multi-cluster virtual warehouse. Select all that apply.**

Options are :

*MAX\_CLUSTER\_COUNT (Correct)*

*SUSPEND\_AND\_RESUME*

*MIN\_CLUSTER\_COUNT (Correct)*

*SCALING\_POLICY (Correct)*

**Answer :**

\*MAX\_CLUSTER\_COUNT (Correct)

\*MIN\_CLUSTER\_COUNT (Correct)

\*SCALING\_POLICY (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 17](#)

**What is correct about multi cluster virtual warehouse?**

Options are :

*Only one multi cluster virtual warehouse can be created per Snowflake customer.*

*Multi cluster virtual warehouse can not be suspended or resumed.*

*They will remove the additional clusters as query demand decreases. (Correct)*

*They will add additional clusters automatically when under high query workload which is beyond the current capacity. (Correct)*

**Answer :**

\*They will remove the additional clusters as query demand decreases. (Correct)

\*They will add additional clusters automatically when under high query workload which is beyond the current capacity. (Correct)

**When a file has been loaded into a table, Snowflake marks that file as loaded in metadata so that the file doesn't get processed again.**

Options are :

*TRUE (Correct)*

*FALSE*

**Answer :**

\*TRUE (Correct)

**A consumer of a shared database can add new tables or views in the shared database?**

Options are :

*TRUE*

*FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 3

**Select the cloud services provider supported by Snowflake architecture. (three choices)**

Options are :

*Amazon Web Services (AWS) (Correct)*

*Google (GCP) (Correct)*

*HP Virtual Private Cloud*

*Cloud Foundry*

*Microsoft Azure (Correct)*

**Answer :**

\*Amazon Web Services (AWS) (Correct)

\*Google (GCP) (Correct)

\*Microsoft Azure (Correct)

**Which of the following statements are true with respect to the Snowflake release process? 1. Snowflake deploys new feature releases and/or patch releases every week. 2. Snowflake deploys patch releases every week, but new feature releases happen once a month. 3. There's minimal downtime associated with Snowflake depending upon the content of the release. 4. It is possible for you as a user to request 24-hour early access to the upcoming releases so that you can do additional release testing before the release is rolled out.**

Options are :

1, 3 and 4

2, 3 and 4

1 and 3

1 and 4 (Correct)

**Answer :**

\*1 and 4 (Correct)

**When opening a snowflake account, what do you need to decide (select three)?**

Options are :

*Edition (Correct)*

*Deployment Environment (Public Cloud or Private Cloud)*

*Cloud Service Provider (Correct)*

*Region (Correct)*

*Availability Zone*

**Answer :**

\*Edition (Correct)

\*Cloud Service Provider (Correct)

\*Region (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 10**

**A resource monitor in Snowflake....(select all that apply)**

Options are :

*is created to monitor and control credits consumed by virtual warehouse and storage space consumed by database storage.*

*can monitor credits consumed by user-managed virtual warehouses and Snowflake-provisioned virtual warehouses. (Correct)*

*can be created by individual users to monitor the warehouses created by them.*

*can be created using Snowflake Web UI or a SQL command. (Correct)*

**Answer :**

- \*can monitor credits consumed by user-managed virtual warehouses and Snowflake-provisioned virtual warehouses. (Correct)
- \*can be created using Snowflake Web UI or a SQL command. (Correct)

**Each warehouse can be assigned more than one resource monitor.**

Options are :

*TRUE*

*FALSE (Correct)*

**Answer :**

- \*FALSE (Correct)

**Select appropriate configuration parameters to create a resource monitor in Snowflake. (select all that apply)**

Options are :

*Credit quota and cycle (daily, weekly, monthly, yearly, never) (Correct)*

*Credit quota distribution %*

*Monitoring Level (account or warehouse) (Correct)*

*Monitoring threshold and corresponding notification and action when the threshold is reached (Correct)*

*Notification email address (e.g. Email, SMS or both)*

**Answer :**

- \*Credit quota and cycle (daily, weekly, monthly, yearly, never) (Correct)
- \*Monitoring Level (account or warehouse) (Correct)
- \*Monitoring threshold and corresponding notification and action when the threshold is reached (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 16**

**A ♦notify and suspend immediately♦ action will cancel all statements being executed by the warehouses at that time and notify account administrators.**

Options are :

*TRUE (Correct)*

*FALSE*

**Answer :**

\*TRUE (Correct)

**After a warehouse is suspended by a credit monitor action, the warehouse will not resume until the following event(s) occur. (select all that apply)**

Options are :

*Snowflake account administrator manually resumes the warehouse.*

*Credit quota is increased in the resource monitor configuration. (Correct)*

*Credit threshold is increased in the resource monitor configuration. (Correct)*

*Next monitoring cycle (monitoring interval) starts. (Correct)*

**Answer :**

\*Credit quota is increased in the resource monitor configuration. (Correct)

\*Credit threshold is increased in the resource monitor configuration. (Correct)

\*Next monitoring cycle (monitoring interval) starts. (Correct)

**Which of the following options are supported in Snowflake for bulk/continuous loading of data? (select all that apply)**

Options are :

*Snowflake Web UI*

*SnowSQL (Correct)*

*Snowpipe (Correct)*

*ETL (Correct)*

**Answer :**

\*SnowSQL (Correct)

\*Snowpipe (Correct)

\*ETL (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 39**

**Which of the following file types are supported in Snowflake for bulk loading of data? (select all that apply)**

Options are :

*CSV (Correct)*

*JSON (Correct)*

*JPG*

*Avro (Correct)*

**Answer :**

- \*CSV (Correct)
- \*JSON (Correct)
- \*Avro (Correct)

**To execute a data loading task in Snowflake, you will need to have ...(select all that apply)**

Options are :

*A stage (Correct)*

*A data file (Correct)*

*Adequate storage space*

*A target table (Correct)*

*A virtual warehouse (Correct)*

*A SYSADMIN role*

**Answer :**

- \*A stage (Correct)
- \*A data file (Correct)
- \*A target table (Correct)
- \*A virtual warehouse (Correct)

**Snowflake provides the following types of internal stages to load data. (select all that apply)**

Options are :

*Table Stage (Correct)*

*Data Stage*

*User Stage (Correct)*

*Named Stage (Correct)*

**Answer :**

- \*Table Stage (Correct)
- \*User Stage (Correct)
- \*Named Stage (Correct)

## Tackling Snowflake Certification - Practice Questions Set 2

**Following is the correct sequence of commands to bulk load data from a Snowflake internal stage into a table. (select all that apply)**

Options are :

*CREATE STAGE ... PUT... COPY INTO.... (Correct)*

*PUT... COPY INTO ... (Correct)*

*CREATE STAGE.... COPY INTO... PUT....*

*Any of these sequences of commands will work.*

**Answer :**

\**CREATE STAGE ... PUT... COPY INTO.... (Correct)*

\**PUT... COPY INTO ... (Correct)*

**Assume that there is a table in a Snowflake database called STUDENT containing student data, and it has a column named S\_NAME that stores student names. You need to retrieve the first 10 student names after sorting them alphabetically. Which one of the following will be a correct SQL statement for this task? (select all that apply)**

Options are :

*Select student\_name from student limit 10 order by student\_name;*

*Select student\_name from student order by student\_name limit 10; (Correct)*

*Select student\_name from student order by student\_name fetch 10; (Correct)*

*Select top 10 student\_name from student order by student\_name; (Correct)*

**Answer :**

\**Select student\_name from student order by student\_name limit 10; (Correct)*

\**Select student\_name from student order by student\_name fetch 10; (Correct)*

\**Select top 10 student\_name from student order by student\_name; (Correct)*

**An external stage in Snowflake... (select all that apply):**

Options are :

*is useful when data is already available in cloud storage. (Correct)*

*does not require a PUT command. (Correct)*

*supports loading data stored in Amazon AWS, Microsoft Azure and Google Cloud Platform (GCP) only. (Correct)*

*is not visible in the Stages tab of Snowflake Web UI.*

**Answer :**

\*is useful when data is already available in cloud storage. (Correct)

\*does not require a PUT command. (Correct)

\*supports loading data stored in Amazon AWS, Microsoft Azure and Google Cloud Platform (GCP) only. (Correct)

### Snowflake Snowpro Certification Practice Test Set 1

**While data loading in a Snowflake provided stage (i.e., internal stage), which of the following statement(s) is true (with regard to data encryption)? (select all that apply)**

Options are :

*If the data is unencrypted, the user will need to encrypt the data before loading.*

*If the data is unencrypted, client software (e.g., SnowSQL) will be used to automatically encrypt the data on the client machine before being transmitted to the Snowflake stage. (Correct)*

*In Snowflake, data at rest is always encrypted. (Correct)*

*When data is unloaded in the customer environment, data is automatically decrypted by Snowflake.*

**Answer :**

\*If the data is unencrypted, client software (e.g., SnowSQL) will be used to automatically encrypt the data on the client machine before being transmitted to the Snowflake stage. (Correct)

\*In Snowflake, data at rest is always encrypted. (Correct)

**While loading data into Snowflake, which of the following transformations are supported? (select all that apply)**

Options are :

*Column reordering (Correct)*

*Column omission (Correct)*

*Derived columns*

*Cast operation (Correct)*

*Truncation (Correct)*

*Transpose*

**Answer :**

- \*Column reordering (Correct)
- \*Column omission (Correct)
- \*Cast operation (Correct)
- \*Truncation (Correct)

**Which of the following statements are true with respect to how Snowflake handles data encryption while loading data into Snowflake internal stage? (select all that apply)**

Options are :

*Automatic encryption mode in Snowflake is 128-bit encryption. (Correct)*

*Automatic encryption mode in Snowflake is 256-bit encryption.*

*It is possible to disable automatic encryption while loading data into Snowflake (for example, to improve the performance of data load).*

*If the file is already encrypted, the encryption key must be supplied to Snowflake. (Correct)*

**Answer :**

- \*Automatic encryption mode in Snowflake is 128-bit encryption. (Correct)
- \*If the file is already encrypted, the encryption key must be supplied to Snowflake. (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 15**

**If the file being loaded into Snowflake internal stage is uncompressed, Snowflake automatically compresses files using gzip.**

Options are :

*TRUE (Correct)*

*FALSE*

**Answer :**

- \*TRUE (Correct)

**SnowSQL (COPY statement) is a more appropriate approach when you need to load data in bulk. On the other hand, using Snowpipe is more appropriate when you need to load data from a continuous data stream or micro-batches.**

Options are :

*TRUE (Correct)*

*FALSE*

**Answer :**

\*TRUE (Correct)

**Which of the following are true with respect to the COPY statement and Snowpipe for data loading? (select all that apply)**

Options are :

*Data loading using COPY statement will require a user-provisioned warehouse, whereas Snowpipe uses Snowflake provided compute (serverless compute model) (Correct)*

*Snowpipe does not support transformations.*

*COPY statement supports loading data from external and internal stages; snowpipe supports data loading from internal stages only.*

*None of the above*

**Answer :**

\*Data loading using COPY statement will require a user-provisioned warehouse, whereas Snowpipe uses Snowflake provided compute (serverless compute model) (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 11**

**It is NOT possible to query data stored in the cloud in Snowflake without first loading data into a Snowflake table.**

Options are :

*TRUE*

*FALSE (Correct)*

**Answer :**

\*FALSE (Correct)

**Which of the following will contribute to your account level storage usage? Select all that apply.**

Options are :

*User Defined Functions*

*Snowflake Internal stage (Correct)*

*Views*

*Snowflake External Stages*

*Database (and the tables in it) (Correct)*

**Answer :**

\*Snowflake Internal stage (Correct)

\*Database (and the tables in it) (Correct)

**As a consumer you can create only one database per share?**

Options are :

*TRUE (Correct)*

*FALSE*

**Answer :**

\*TRUE (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 28**

**Which role is most powerful role in a Snowflake system?**

Options are :

*PUBLIC*

*SYSADMIN*

*SECURITYADMIN*

*ACCOUNTADMIN (Correct)*

*USERADMIN*

**Answer :**

\*ACCOUNTADMIN (Correct)

**What is the lowest snowflake edition that allows customer dedicated virtual warehouses?**

Options are :

*Business Critical*

*Enterprise*

*Standard (Correct)*

*VPS*

**Answer :**

\*Standard (Correct)

**You have a virtual warehouse named my\_vwh which is sized as Medium. You need the size of the virtual warehouse to be Large for future queries. What is the best way to make this change.**

Options are :

*After the existing virtual warehouse and change the size to Large (ALTER WAREHOUSE my\_vwh WAREHOUSE\_SIZE=large) (Correct)*

*Drop the existing virtual warehouse and create a new one with the required parameters*

*It is not permitted to re-size a virtual warehouse once it has been created.*

**Answer :**

\*After the existing virtual warehouse and change the size to Large (ALTER WAREHOUSE my\_vwh WAREHOUSE\_SIZE=large) (Correct)

**Customer working Working Snowflake needs to perform hardware maintenance, upgrades or regular patches**

Options are :

*True*

*False (Correct)*

**Answer :**

\*False (Correct)

**Which Snowflake editions provides data replication and failover capabilities**

Options are :

*Standard*

*Enterprise*

*Business Critical (Correct)*

*VPS - Virtual Private Snowflake (Correct)*

**Answer :**

- \*Business Critical (Correct)
- \*VPS - Virtual Private Snowflake (Correct)

**Which of the below is not a role available when Snowflake account is created initially**

Options are :

*Useradmin*

*Accountadmin*

*Sysadmin*

*Developer (Correct)*

**Answer :**

- \*Developer (Correct)

### [SnowPro Core Certification 3 Full Practice Exams 2021 Set 2](#)

**Which of the following statement describes a benefit of snowflake's separation of computer and storageC**

Options are :

*Storage can expand without requirement to add more compute (Correct)*

*Compute Cluster can access stored data without contention (Correct)*

*Compute Cluster can be scaled up and down without requirement of adding any new storage (Correct)*

*Compute and Storage are tightly coupled together*

**Answer :**

- \*Storage can expand without requirement to add more compute (Correct)
- \*Compute Cluster can access stored data without contention (Correct)
- \*Compute Cluster can be scaled up and down without requirement of adding any new storage (Correct)

**Which snowflake edition supports Query Statement Encryption?**

Options are :

*VPS (Virtual Private Snowflake )*

*Standard*

*Enterprise*

*None (Correct)*

**Answer :**

\*None (Correct)

**In the Snowflake container hierarchy, what container is represented as a URL . For example <https://xxxpqr12345.ap-southeast-2.snowflakecomputing.com>**

Options are :

*Account (Correct)*

*Database*

*Virtual Warehouse*

*Schema*

**Answer :**

\*Account (Correct)

### **Snowflake SnowPro Core Certification Exam Questions Set 5**

**VPS - Virtual Private Snowflake Edition have customer-dedicated metadata store**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

**Snowflake offers tools to extract data from source systems?**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

**Snowflake is HIPAA compliant?**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

### Snowflake SnowPro Core Certification Exam Questions Set 6

**How often does Snowflake release new features?**

Options are :

*Weekly (Correct)*

*Monthly*

*Bi-Annually*

*Yearly*

**Answer :**

\*Weekly (Correct)

**In how many regions Snowflake is available on Google Cloud**

Options are :

1

4

2

5

*3 (Correct)*

**Answer :**

\*3 (Correct)

**Which layer contains the data in compressed, columnar format?**

Options are :

*Query Processing Layer*

*None*

*Cloud Service Layer*

*Storage Layer (Correct)*

**Answer :**

\*Storage Layer (Correct)

### **SnowPro Core Certification 3 Full Practice Exams 2021 Set 5**

#### **How does snowflake achieves High Availability**

Options are :

*High availability is limited to Storage Layer*

*High availability is limited to compute layer only due to multi-cluster architecture*

*It provide built-in, resilient capabilities in all 3 layers (Correct)*

*It does not support fault tolerant*

**Answer :**

\*It provide built-in, resilient capabilities in all 3 layers (Correct)

#### **Which of the following best describes Snowflake's Processing Engine**

Options are :

*Presto*

*Native SQL (Correct)*

*Apache Spark*

*MapReduce Framework*

**Answer :**

\*Native SQL (Correct)

**Snowflake's architecture is a hybrid of traditional shared-disk and shared-nothing database architectures.**

#### **Select all applicable service layers in snowflake architecture**

Options are :

*Query Processing Layer (Correct)*

*Cloud Service Layer (Correct)*

*Data Management Layer*

*Data Storage Layer (Correct)*

**Answer :**

- \*Query Processing Layer (Correct)
- \*Cloud Service Layer (Correct)
- \*Data Storage Layer (Correct)

### Snowflake Certification Preparation Set 6

**Cloud service layer is responsible for ? Select all that apply?**

Options are :

*User authentication (Correct)*

*Data security (Correct)*

*Availability zone management*

*Metadata storage (Correct)*

*Data Partitioning*

**Answer :**

- \*User authentication (Correct)
- \*Data security (Correct)
- \*Metadata storage (Correct)

**MFA can be used for connecting to Snowflake via the Snowflake JDBC driver**

Options are :

*False*

*True (Correct)*

**Answer :**

- \*True (Correct)

**Snowflake processes queries based on which cluster processing architecture**

Options are :

*Leader Follower Architecture*

*Master Slave Architecture*

*MPP (massively parallel processing) (Correct)*

*Publisher-Subscriber Architecture*

**Answer :**

\*MPP (massively parallel processing) (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 15

**While setting up a snowflake account, which out of below you do not specify**

Options are :

*Availability Zone (AZ) (Correct)*

*Cloud Infra Provider*

*Snowflake Edition*

*Geographic Deployment Region*

**Answer :**

\*Availability Zone (AZ) (Correct)

**Which one of the following is not a layer in Snowflake Architecture?**

Options are :

*Data Catalog (Correct)*

*Virtual Warehouse*

*Cloud Services*

*Storage*

**Answer :**

\*Data Catalog (Correct)

**Snowflake cannot be run on private cloud infrastructures (on-premises or hosted)**

Options are :

*True (Correct)*

*False*

**Answer :**

\*True (Correct)

### SnowPro Core Certification 3 Full Practice Exams 2021 Set 3

**Snowflake technology has been build from scratch, specifically designed for execution on cloud platforms**

Options are :

*True (Correct)*

*False*

**Answer :**

\*True (Correct)

**In snowflake, which layer is responsible for storing security information**

Options are :

*Query Processing Layer*

*Storage Layer*

*None of the Above*

*Cloud Service Layer (Correct)*

**Answer :**

\*Cloud Service Layer (Correct)

**Credits unit costs and data storage determined by the region where snowflake services needs to be hosted/invoked?**

Options are :

*True (Correct)*

*False*

**Answer :**

\*True (Correct)

## Snowflake Practice Questions - SnowPro Core Certified Set 2

**An agreement is required and must be signed by Data Providers before a listing will be made public on Data Marketplace**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

**In which layer of its architecture does snowflake stores its metadata statistics**

Options are :

*Data Storage Layer*

*Query Processing Layer or Computer Layer*

*Data Catalog Layer*

*Cloud Service Layer (Correct)*

**Answer :**

\*Cloud Service Layer (Correct)

**Micro-partitions are the physical (not logical) data files that comprise Snowflake's logical tables**

Options are :

*True (Correct)*

*False*

**Answer :**

\*True (Correct)

## Snowflake Snowpro Certification Practice Test Set 5

**Snowflake is SaaS Based Data warehouse for data Analytics.**

Options are :

*True*

*False (Correct)*

**Answer :**

\*False (Correct)

**Elastic Data Warehousing is not available in all Snowflake Editions. Check all Snowflake editions that have Elastic Data Warehousing enabled**

Options are :

*Premium Editions*

*VPS (Virtual Private Snowflake) (Correct)*

*Enterprise Editions (Correct)*

*Standard Editions*

*Business Critical Editions (Correct)*

**Answer :**

\*VPS (Virtual Private Snowflake) (Correct)

\*Enterprise Editions (Correct)

\*Business Critical Editions (Correct)

**What all attribute makes up Snowflake a true SaaS Solution**

Options are :

*No compute cost*

*Transparent release which does not need user intervention (Correct)*

*No hardware to purchase or configure (Correct)*

*No maintenance in terms of patching or upgrades (Correct)*

*No storage Cost*

**Answer :**

\*Transparent release which does not need user intervention (Correct)

\*No hardware to purchase or configure (Correct)

\*No maintenance in terms of patching or upgrades (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 6

**MFA Is applicable for which of the following connectors. Select all which apply**

Options are :

*Python (Correct)*

*ODBC (Correct)*

*JAVA*

*JDBC (Correct)*

*Web Interface (Correct)*

**Answer :**

\*Python (Correct)

\*ODBC (Correct)

\*JDBC (Correct)

\*Web Interface (Correct)

**Characteristics of Snowflake's Multi-Cluster environment. Select all that apply**

Options are :

*All of the Above*

*Option available to user to choose the warehouse to process the query*

*Multi-cluster warehouse supports all the same properties and actions similar to single cluster warehouse (Correct)*

*Individual warehouses automatically scale up and down based on the query load (Correct)*

**Answer :**

\*Multi-cluster warehouse supports all the same properties and actions similar to single cluster warehouse (Correct)

\*Individual warehouses automatically scale up and down based on the query load (Correct)

**Snowflake billing mechanism works per**

Options are :

*every 5 minutes*

*Hour*

*Second (Correct)*

*Minutes*

**Answer :**

\*Second (Correct)

### **Snowflake Certification Preparation Set 23**

**Data storage and compute is completely independent in snowflake**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

**Snowflake has which type of Architecture**

Options are :

*None of the above*

*Hybrid of Shared Disk and Shared Nothing (Correct)*

*Shared Nothing*

*Shared Disk*

**Answer :**

\*Hybrid of Shared Disk and Shared Nothing (Correct)

**Which layer contains virtual warehouse for data access**

Options are :

*None*

*Cloud services*

*Database storage*

*Query processing (Correct)*

**Answer :**

\*Query processing (Correct)

### Snowflake Snowpro Certification Practice Test Set 3

**How many Availability Zone does Snowflake replicate to for disaster recovery (DR) and high availability (HA)**

Options are :

*Two*

*As specified by the user at the time of configuration*

*3 (Correct)*

*One*

**Answer :**

\*3 (Correct)

**Different editions of snowflake Instances require separate accounts**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

**What are the various cloud platform supported for Snowflake**

Options are :

*GCP (Correct)*

*AWS (Correct)*

*IBM Cloud*

*Oracle Cloud*

*Azure (Correct)*

**Answer :**

\*GCP (Correct)

\*AWS (Correct)

\*Azure (Correct)

## Snowflake SnowPro Core Certification Exam Questions Set 1

**Which of the following terms are associated with Cloud Service Layer**

Options are :

*Query Pipelining*

*Query Standardization*

*Query Compilation (Correct)*

*Query Planning (Correct)*

*Query Optimization (Correct)*

**Answer :**

\*Query Compilation (Correct)

\*Query Planning (Correct)

\*Query Optimization (Correct)

**Which type of cache gets purged regularly?**

Options are :

*Warehouse Cache*

*Result Cache (Correct)*

*Metadata Cache*

**Answer :**

\*Result Cache (Correct)

**What are 2 modes of multi-cluster warehouse?**

Options are :

*Dynamically*

*Maximized (Correct)*

*Auto-Scale (Correct)*

*Statistically*

**Answer :**

- \*Maximized (Correct)
- \*Auto-Scale (Correct)

#### [Snowflake SnowPro Core Certification Practice Test Set 4](#)

**You can start or stop Virtual Warehouses at any time**

Options are :

*False*

*True (Correct)*

**Answer :**

- \*True (Correct)

**You can resize a Virtual Warehouse while queries are running**

Options are :

*True (Correct)*

*False*

**Answer :**

- \*True (Correct)

**Scaling in is a manual process (e.g. Max Clusters to min clusters)**

Options are :

*False (Correct)*

*True*

**Answer :**

- \*False (Correct)

#### [Snowflake SnowPro Core Certification Practice Test Set 7](#)

**The number of queries that a warehouse can currently process is determined by the size and complexity of the query.**

Options are :

*False*

*True (Correct)*

**Answer :**

\*True (Correct)

**When configuring a Warehouse using a Snowflake edition that has Elastic Data Warehousing enabled, what facets or components will you need to configure that are not needed in accounts where Elastic Data Warehousing is not enabled**

Options are :

*Auto Resume*

*Minimum and Maximum Clusters (Correct)*

*Auto Suspend*

*Minimum and Maximum Servers*

*Scaling Policy (Correct)*

**Answer :**

\*Minimum and Maximum Clusters (Correct)

\*Scaling Policy (Correct)

**There is no cost saving if you stop the Virtual warehouse before first 60 seconds period**

Options are :

*True (Correct)*

*False*

**Answer :**

\*True (Correct)

**SnowPro Core Certification 3 Full Practice Exams 2021 Set 6**

**All Virtual Warehouses will have access to all data?**

Options are :

*True (Correct)*

*False*

**Answer :**

\*True (Correct)

**Snowpipe uses compute resources provided by user provided VW (Virtual Warehouses)**

Options are :

*True*

*False (Correct)*

**Answer :**

\*False (Correct)

**A warehouse was started and it was up and running for 75 seconds. For how many seconds will the warehouse be billed for?**

Options are :

*120 Seconds*

*75 Seconds (Correct)*

*70 Seconds*

*60 Seconds*

**Answer :**

\*75 Seconds (Correct)

## [Snowflake Certification Preparation Set 6](#)

**Which of the following refers the same layer**

Options are :

*Metadata Layer*

*Compute Layer (Correct)*

*Catalog Layer*

*Query Processing Layer (Correct)*

*Virtual Warehouse Layer (Correct)*

**Answer :**

\*Compute Layer (Correct)

\*Query Processing Layer (Correct)

\*Virtual Warehouse Layer (Correct)

Question 1: **Incorrect**

Which primary tool loads data to Snowflake from a local file system?

- ETL tools  
**(Incorrect)**
- Snowflake UI
- SnowSQL  
**(Correct)**
- External Stage

#### Explanation

**SnowSQL is the primary tool to load data to Snowflake from a local file system.** You can run it in either interactive shell or batch mode.

**Note:** Don't get confused between SnowSQL and SnowCD. SnowCD (i.e. Snowflake Connectivity Diagnostic Tool) helps users to diagnose and troubleshoot their network connection to Snowflake.

Question 2:

**Skipped**

Select the correct statements for Table Clustering. (Select 3)

- Automatic clustering can not be suspended or resumed
- Snowflake recommends a maximum of three or four columns (or expressions) per key  
**(Correct)**
-

Tables in multi-terabytes range are good candidate for clustering keys

(Correct)

- 

Clustering keys are not for every table

(Correct)

- 

Automatic Clustering doesn't consume credit

- 

Snowflake doesn't charge for Reclustering

#### Explanation

Clustering keys are not for every table. Tables in the multi-terabyte range are good candidates for clustering keys. Both automatic clustering and reclustering consume credit. A single clustering key can contain one or more columns or expressions. Snowflake recommends a maximum of three or four columns (or expressions) per key for most tables. Adding more than 3-4 columns tends to increase costs more than benefits.

Question 3: **Correct**

How much uncompressed data does a micro-partition contain in Snowflake?

- 

Between 1 GB to 10 GB

- 

Between 1 MB to 100 MB

- 

Between 5 MB to 50 MB

- 

Between 50 MB to 500 MB

(Correct)

#### Explanation

Each micro-partition contains between 50 MB and 500 MB of uncompressed data (Note that the actual size in Snowflake is smaller because data is always stored compressed.). Groups of rows in tables are mapped into individual micro-partitions, organized in a columnar fashion. This size between 50 MB and 500 MB, and the structure allows for extremely granular pruning of very large tables, which can be comprised of millions, or even hundreds of millions, of micro-partitions. It enables extremely efficient DML and fine-grained pruning for faster queries.

Question 4:

**Skipped**

Which command can be used to suspend Automatic Clustering for a table?

- 

**ALTER TABLE**

**(Correct)**

- 

**STOP TABLE**

- 

**SUSPEND RECLUSTER**

- 

**DROP CLUSTERING**

**Explanation**

Example - ALTER TABLE EMPLOYEE SUSPEND RECLUSTER; please note, SUSPEND RECLUSTER is a clause here, not a command.

Question 5:

**Skipped**

Which object parameter can users with the ACCOUNTADMIN role use to set the default retention period for their account?

- 

**DATA\_RETENTION\_TIME\_MAX**

- 

**DATA\_RETENTION\_TIME\_IN\_DAYS**

**(Correct)**

-

## **DATA\_RETENTION\_TIME\_IN\_HOURS**

- 

## **DATA\_RETENTION\_IN\_TIME\_TRAVEL**

### **Explanation**

Users can use the **DATA\_RETENTION\_TIME\_IN\_DAYS** object parameter with the ACCOUNTADMIN role to set the default retention period for their account.

Question 6:

**Skipped**

**Snowflake data providers can share data from one database per share. Data from multiple databases can not be shared with a share. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

**Snowflake data providers can share data that resides in different databases by using secure views.** A secure view can reference objects such as schemas, tables, and other views from one or more databases, as long as these databases belong to the same account.

Question 7:

**Skipped**

**Which of the Snowflake editions provides HIPPA Support feature? (Select 2)**

- 

**Enterprise**

- 

**Virtual Private Snowflake (VPS)**

**(Correct)**

- 

**All of the Snowflake Editions**

-

**Business Critical**

**(Correct)**

- 

**Standard**

**Explanation**

**Business Critical and Virtual Private Snowflake (VPS) editions provide HIPPA support.**

Question 8:

**Skipped**

**Which is not the DML (Data Manipulation Language) command?**

- 

**MERGE**

- 

**INSERT**

- 

**UNDROP**

**(Correct)**

- 

**TRUNCATE**

- 

**UPDATE**

- 

**DELETE**

**Explanation**

**UNDROP is Snowflake's DDL (Data Definition Language) command.**

Question 9:

**Skipped**

**Which of these are types of Snowflake releases? (Select 3)**

-

## Bug Fix Release

- 

## Behavior Change Release

(Correct)

- 

## Patch Release

(Correct)

- 

## Full Release

(Correct)

- 

## Part Release

### Explanation

There are three types of releases:

**Full Release:** A full release may include any of the following:

- New features
- Feature enhancements or updates
- Fixes

**Patch Release:** A patch release includes fixes only.

**Behavior Release:** Every month, Snowflake deploys one behavior change release. Behavior change releases contain changes to existing behaviors that may impact customers.

Question 10:

**Skipped**

**Monica has successfully created a task with the 5 minutes schedule. It has been 30 minutes, but the task did not run. What could be the reason?**

- 

**Task schedule should not be less than 60 minutes**

- 

Monica doesn't have the authority to run the task

- 

Monica should run the ALTER TASK command to RESUME the task

(Correct)

- 

Monica should run the ALTER TASK command to SUSPEND the task, and then again run the ALTER TASK command to RESUME the task

#### Explanation

The first time we create the TASK, we must run the ALTER TASK command to RESUME the task.

Question 11:

Skipped

Which of these are applicable for Snowflake Connector for Kafka? (Select all that apply)

- 

Reads data from one or more Kafka topics and loads the data into a Snowflake table

(Correct)

- 

If the topics are not mapped, then the Kafka connector creates a new table for each topic using the topic name

(Correct)

- 

Kafka topics can be mapped to existing Snowflake tables in the Kafka configuration

(Correct)

- 

Kafka connector required a pre-configured Snowflake table to map the topics with that Snowflake table

- 

**The Kafka connector subscribes to one or more Kafka topics**

**(Correct)**

#### **Explanation**

Kafka topics can be mapped to existing Snowflake tables in the Kafka configuration. **If the topics are not mapped, then the Kafka connector creates a new table for each topic using the topic name.** The Kafka connector subscribes to one or more Kafka topics based on the configuration information provided via the Kafka configuration file or command line (Or the Confluent Control Center; Confluent only). Question 12:

#### **Skipped**

**In which of the cloud platforms a Snowflake account can be hosted? (Select 3)**

- 

**IBM Cloud**

- 

**Oracle Cloud**

- 

**AZURE**

**(Correct)**

- 

**AWS**

**(Correct)**

- 

**GCP**

**(Correct)**

#### **Explanation**

A Snowflake account can be hosted on any of the following cloud platforms:

- Amazon Web Services (AWS)
- Google Cloud Platform (GCP)
- Microsoft Azure (Azure).

On each platform, Snowflake provides one or more regions where the account is provisioned.

Question 13:

**Skipped**

**A stored procedure can simultaneously run the caller's and the owner's rights.**

(True / False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

#### **Explanation**

**A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. A caller's rights stored procedure runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query. An**

**owner's rights stored procedure** runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table.

At the time that the stored procedure is created, the creator specifies whether the procedure runs with the owner's rights or the caller's rights. The default is owner's rights.

Question 14:

**Skipped**

**Which of these are Snowgrid's capabilities? (Select all that apply)**

- 

**ETL dependent**

- 

**Live, ready to query data**

**(Correct)**

- 

**Zero-copy cloning**

**(Correct)**

- 

**Share internally with private data exchange or externally with public data exchange**

**(Correct)**

- 

**Secure, governed data sharing**

**(Correct)**

### **Explanation**

Snowgrid allows you to use Secure Data Sharing features to provide access to live data, **without any ETL** or movement of files across environments.

Question 15:

**Skipped**

**What are the key benefits of The Data Cloud? (Select 3)**

- 

**Action**

**(Correct)**

- 

**Backup**

- 

**Governance**

**(Correct)**

- 

**Maintenance**

-

## Access

(Correct)

### Explanation

The benefits of The Data Cloud are Access, Governance, and Action. Access means that organizations can easily discover data and share it internally or with third parties without regard to geographical location. Governance is about setting policies and rules and protecting the data in a way that can unlock new value and collaboration while maintaining the highest levels of security and compliance. Action means you can empower every part of your business with data to build better products, make faster decisions, create new revenue streams and realize the value of your greatest untapped asset, your data.

Question 16:

Skipped

If the micro-partitions are constant, how much is the Clustering Overlap Depth?

20

1

(Correct)

2

10

0

### Explanation

When there is no overlap in the range of values across all micro-partitions, the micro-partitions are considered to be in a constant state (i.e. they cannot be improved by clustering).

Question 17:

Skipped

If `DATA_RETENTION_TIME_IN_DAYS` is set to a value of 0, and `MIN_DATA_RETENTION_TIME_IN_DAYS` is set higher at the account level and is greater than 0, which value (0 or higher) setting takes precedence?

- 

**Higher value (set in MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS)**

**(Correct)**

- 

**0 (set in DATA\_RETENTION\_TIME\_IN\_DAYS)**

### **Explanation**

If DATA\_RETENTION\_TIME\_IN\_DAYS is set to a value of 0, and MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS is set at the account level and is greater than 0, the higher value setting takes precedence. The data retention period for an object is determined by **MAX(DATA\_RETENTION\_TIME\_IN\_DAYS, MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS)**.

Question 18:

**Skipped**

**Cloning a table replicates the source table's structure, data, load history, and certain other properties (e.g., STAGE FILE FORMAT). (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

Cloning a table replicates the source table's structure, data, and certain other properties (e.g., STAGE FILE FORMAT). **A cloned table does not include the load history of the source table.** One consequence is that data files loaded into a source table can be loaded again into its clones.

Question 19:

**Skipped**

**UDF runs with either the caller's or the owner's rights. (True / False)**

- 

**FALSE**

**(Correct)**

-

**TRUE**

**Explanation**

**UDF only runs as the function owner.** A stored procedure runs with either the caller's rights or the owner's rights. It cannot run with both at the same time. **A caller's rights stored procedure** runs with the privileges of the caller. The primary advantage of a caller's rights stored procedure is that it can access information about that caller or about the caller's current session. For example, a caller's rights stored procedure can read the caller's session variables and use them in a query. **An**

**owner's rights stored procedure** runs mostly with the privileges of the stored procedure's owner. The primary advantage of an owner's rights stored procedure is that the owner can delegate specific administrative tasks, such as cleaning up old data, to another role without granting that role more general privileges, such as privileges to delete all data from a specific table.

At the time that the stored procedure is created, the creator specifies whether the procedure runs with the owner's rights or the caller's rights. The default is owner's rights.

Question 20:

**Skipped**

**Which objects are not available for replication in the Standard Edition of Snowflake? (Select 3)**

- 

**Shares**

- 

**Integrations**

**(Correct)**

- 

**Database**

- 

**Roles**

**(Correct)**

- 

**Users**

(Correct)

### Explanation

Database and share replication are available in all editions, including the Standard edition. Replication of all other objects is only available for Business Critical Edition (or higher).

Question 21:

Skipped

If a warehouse runs for 61 seconds, shuts down, and then restarts and runs for less than 60 seconds, for how much duration will the billing be charged?

- 61 seconds
- 
- 60 seconds
- 
- 180 seconds
- 
- 120 seconds
- 
- 121 seconds

(Correct)

### Explanation

It will be billed for 121 seconds ( $60 + 1 + 60$ ).

Question 22:

Skipped

Micro-partitioning is the on-demand feature of Snowflake. It is required to be enabled explicitly by ACCOUNTADMIN. (True / False)

- 
- FALSE

(Correct)

-

**TRUE**

**Explanation**

Micro-partitioning is automatically performed on all Snowflake tables. Tables are transparently partitioned using the Ordering of the data as inserted or loaded.

Question 23:

**Skipped**

**Which services are managed by Snowflake's cloud services layer? (Select all that apply)**

- 

**Query Parsing and Optimization**

**(Correct)**

- 

**Authentication**

**(Correct)**

- 

**Access Control**

**(Correct)**

- 

**Infrastructure Management**

**(Correct)**

- 

**Only Infrastructure Management**

- 

**Metadata Management**

**(Correct)**

**Explanation**

The cloud services layer is a collection of services that coordinate activities across Snowflake. These services tie together all of the different components of Snowflake in order to process user requests, from login to query dispatch. The cloud service

layer manages Authentication, Infrastructure Management, Metadata Management, Query parsing and optimization, and Access control services.

Question 24:

**Skipped**

**Snowflake automatically and transparently maintains materialized views.**

(True/False)

- 

**TRUE**

**(Correct)**

- 

**FALSE**

#### **Explanation**

**Snowflake automatically and transparently maintains materialized views.** A background service updates the materialized view after changes to the base table. This is more efficient and less error-prone than manually maintaining the equivalent of a materialized view at the application level.

Question 25:

**Skipped**

You have a table with a 30-day retention period. If you increase the retention period to 40 days, how would it affect the data that would have been removed after 30 days?

- 

**The data will now retain an additional 10 days before moving into Fail-safe**

**(Correct)**

- 

**The data will still be moved to Fail-safe at the end of the 30-day retention period**

#### **Explanation**

Increasing Retention causes the data currently in Time Travel to be retained for a more extended time. For example, suppose you have a table with a 30-day retention period and increase the period to 40 days. In that case, **data that would have been removed after 30 days is now retained for an additional 10 days before moving into Fail-safe.**

Note that this does not apply to any data that is older than 30 days and has already moved into Fail-safe.

Question 26:

Skipped

Monica wants to delete all the data from table t1. She wants to keep the table structure, so she does not need to create the table again. Which command will be appropriate for her need?

- UNDROP

- REMOVE

- TRUNCATE

(Correct)

- DROP

- DELETE

Explanation

TRUNCATE will delete all of the data from a single table. So, once Monica truncates table t1, table t1's structure remains, but the data will be deleted. DELETE is usually used for deleting single rows of data.

Question 27:

Skipped

How can you view the data storage across your entire Snowflake account? (Select 2)

- 

Using Classic Web Interface: Click on Account > Billing & Usage > Average Storage Used

(Correct)

- 

Using Classic Web Interface: Click on Account > Resource Monitors > Average Storage Used

- 

**Using Snowsight: Select Admin > Usage > Storage**

**(Correct)**

- 

**Using Snowsight: Select Data > Usage > Storage**

### **Explanation**

Suppose you have been assigned the ACCOUNTADMIN role (i.e., you serve as the top-level administrator for your Snowflake account). In that case, you can use Snowsight or the classic web interface to view data storage across your entire account:

- Using Snowsight: Select Admin > Usage > Storage,
- Using Classic Web Interface: Click on Account > Billing & Usage > Average Storage Used

Question 28:

**Skipped**

**Which database objects can be shared using the Snowflake Secure Data Sharing feature? (Select all that apply)**

- 

**External Tables**

**(Correct)**

- 

**Roles**

- 

**Secure UDFs**

**(Correct)**

- 

**Secure Materialized View**

**(Correct)**

- 

**Tables**

**(Correct)**

- 

### Secure Views

**(Correct)**

#### Explanation

Secure Data Sharing enables sharing selected objects in a database in your account with other Snowflake accounts. The following Snowflake database objects can be shared:

- Tables
- External tables
- Secure views
- Secure materialized views
- Secure UDFs

Snowflake enables the sharing of databases through shares created by data providers and “imported” by data consumers.

Question 29:

**Skipped**

Which of these Snowflake Editions automatically stores data in an encrypted state?

- 

**Enterprise**

- 

**Virtual Private Snowflake(VPS)**

- 

**Business Critical**

- 

**Standard**

- 

**All of the Snowflake Editions**

**(Correct)**

#### Explanation

**All of the Snowflake Editions (Standard, Enterprise, Business Critical, Virtual Private Snowflake) automatically store data in an encrypted state.**

Question 30:

**Skipped**

The Kafka connector creates Snowflake Objects for each topic.

- 

**One pipe to ingest the data files for each topic partition**

- 

**One table for each topic. If the table specified for each topic does not exist**

- 

**All of these**

**(Correct)**

- 

**One internal stage to temporarily store data files for each topic**

### Explanation

The connector creates the following objects for each topic:

- One internal stage to temporarily store data files for each topic.
- One pipe to ingest the data files for each topic partition.
- One table for each topic. If the table specified for each topic does not exist, the connector creates it; otherwise, the connector creates the RECORD\_CONTENT and RECORD\_METADATA columns in the existing table and verifies that the other columns are nullable (and produces an error if they are not).

Question 31:

**Skipped**

**Direct data sharing can only be done with accounts in the same region and the same cloud provider. (TRUE/FALSE)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### Explanation

**Direct data sharing can only be done with accounts in the same region and the same cloud provider.** Suppose you want to share with someone outside of your region. In that case, you replicate that database into the region you want to share with and share from there.

Question 32:

**Skipped**

The search optimization service speeds only equality searches. (True/False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### Explanation

The search optimization service speeds Equality and IN predicates searches.

Question 33:

**Skipped**

You have a table t1 with a column j that gets populated by a sequence s1. s1 is defined to start from 1 and with an increment of 1. create or replace sequence s1 start = 1 increment = 1 ; create or replace table t1 ( i int, j int default s1.nextval ); You inserted 3 records in table t1: insert into t1 values (1,s1.nextval), (2,s1.nextval), (3,s1.nextval); After that insert statement, you altered the sequence s1 to set the increment to -4: alter sequence s1 set increment = -4; You again inserted 2 records in table t1: insert into t1 values (4,s1.nextval), (5,s1.nextval); What would be the result of the following query? select j from t1 where i = 4;

- 

**5**

- 

**4**

**(Correct)**

- 

**-1**

-

3

- 

0

### Explanation

**ALTER SEQUENCE command takes effect after the second use of the sequence after executing the ALTER SEQUENCE command.** So, if you fetch row where i = 5, you will find j = 0 [row 4 value of j i.e., 4 + (-4) = 0]

Question 34:

**Skipped**

**Which of these stages can not be dropped or altered? (Select 2)**

- 

**Table Stage**

**(Correct)**

- 

**User Stage**

**(Correct)**

- 

**Named Stage**

### Explanation

**User Stage:** User stages cannot be altered or dropped. A user stage is allocated to each user for storing files. This stage type is designed to store staged and managed files by a single user but can be loaded into multiple tables.

**Table Stage:** Table stages cannot be altered or dropped. A table stage is available for each table created in Snowflake. This stage type is designed to store staged and managed files by one or more users but only loaded into a single table. Note that a table stage is not a separate database object but an implicit stage tied to the table itself. A table stage has no grantable privileges of its own.

**Named Stage:** A named internal stage is a database object created in a schema. This stage type can store files staged and managed by one or more users and loaded into one or more tables. Because named stages are database objects, the ability to create, modify, use, or drop them can be controlled using security access control privileges.

Question 35:

**Skipped**

Tasks require compute resources to execute code. Either Snowflake-managed or User-managed compute models can be chosen for individual tasks. (True / False)

- 

**TRUE**

**(Correct)**

- 

**FALSE**

**Explanation**

Tasks require compute resources to execute SQL code. Either of the following compute models can be chosen for individual tasks:

- Snowflake-managed (i.e. serverless compute model)
- User-managed (i.e. virtual warehouse)

Question 36:

**Skipped**

The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake. (True/False)

- 

**TRUE**

**(Correct)**

- 

**FALSE**

**Explanation**

Snowflake manages all aspects of how this data is stored – the organization, file size, structure, compression, metadata, statistics, and other aspects of data storage are handled by Snowflake. The data objects stored by Snowflake are not directly visible nor accessible by customers; they are only accessible through SQL query operations run using Snowflake.

Question 37:

**Skipped**

Which of the following is the correct hierarchy for the Snowflake objects?

- 

**ORGANIZATION > ACCOUNT > DATABASE > SCHEMA > TABLE**

(Correct)

- ACCOUNT > ORGANIZATION > ROLE > USER > DATABASE > SCHEMA > TABLE
- ORGANIZATION > ACCOUNT > ROLE > USER > DATABASE > SCHEMA > STAGE > TABLE
- ORGANIZATION > ACCOUNT > DATABASE > SCHEMA > TABLE > STAGE

#### Explanation

The topmost container is the customer organization. Securable objects such as tables, views, functions, and stages are contained in a schema object, which are in turn contained in a database. All databases for your Snowflake account are contained in the account object. USER, ROLE, DATABASE, WAREHOUSE are at same level and contained in a Snowflake Account Object.

Question 38:

Skipped

What is the maximum data retention period for permanent databases, schemas, and tables for Snowflake Enterprise Edition (and higher)?

- 1 day
- 90 days

(Correct)

- 0 days
- 30 days

#### Explanation

For Snowflake Enterprise Edition (and higher):

- For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also true for temporary tables.
- **For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days.**

Question 39: **Correct**

**Fail-safe helps access historical data after the Time Travel retention period has ended.(True/False)**

- 
- **TRUE**
- 
- **FALSE**

**(Correct)**

#### **Explanation**

**Fail-safe is not provided as a means for accessing historical data after the Time Travel retention period has ended.** It is for use only by Snowflake to recover data that may have been lost or damaged due to extreme operational failures. Data recovery through Fail-safe may take from several hours to several days to complete.

Question 40: **Correct**

**Snowpark is a new developer framework for Snowflake. It allows data engineers, data scientists, and data developers to code in their familiar way with their language of choice and execute the pipeline, ML workflow, and data apps faster and more securely in a single platform. Which of these following languages does Snowpark support? (Select 3)**

- **Python**
- **Scala**
- **C#**

**(Correct)**

**(Correct)**

**(Correct)**

- 

C++

- 

Java

(Correct)

### Explanation

Snowpark is a new developer framework for Snowflake. It allows data engineers, data scientists, and data developers to code in their familiar way with their language of choice and execute the pipeline, ML workflow, and data apps faster and more securely in a single platform. It brings deeply integrated, DataFrame-style programming to the languages developers like to use and functions to help you efficiently expand more data use cases. Now all these can be executed inside Snowflake using the elastic performance engine. **Snowpark support starts with Scala API, Java UDFs, and External Functions and expands to Java & Python.**

Question 41:

Skipped

SQL clause that helps define the clustering key:

- 

CLUSTERING BY

- 

CLUSTER ON

- 

CLUSERTING ON

- 

CLUSTER BY

(Correct)

### Explanation

Example - create or replace table t1 (c1 date, c2 string, c3 number) **cluster by** (c1, c2);

Question 42:

Skipped

Which of these objects do not clone? (Select 2)

- **Internal (Snowflake) stages**  
**(Correct)**
- **External Table**  
**(Correct)**
- **Schemas**
- **Databases**

#### **Explanation**

Databases and Schemas can be cloned. **External Table and Internal (Snowflake) stages do not get cloned.**

Question 43:

**Skipped**

**Snowflake supports multiple ways of connecting to the service. (Select 3)**

- **Command line clients (e.g. SnowSQL)**  
**(Correct)**
- **Only ODBC**
- **Only JDBC**
- **A web-based user interface**  
**(Correct)**
-

## ODBC and JDBC drivers

(Correct)

### Explanation

Snowflake supports following multiple ways of connecting to its service: A web-based user interface from which all aspects of managing and using Snowflake can be accessed. Command line clients (e.g. SnowSQL) which can also access all aspects of managing and using Snowflake. ODBC and JDBC drivers that can be used by other applications (e.g. Tableau) to connect to Snowflake. Native connectors (e.g. Python, Spark) that can be used to develop applications for connecting to Snowflake. Third-party connectors that can be used to connect applications such as ETL tools (e.g. Informatica) and BI tools (e.g. ThoughtSpot) to Snowflake.

Question 44:

Skipped

What are the three layers in Snowflake's unique architecture? (Select 3)

- 

Database Storage

(Correct)

- 

Computation Services

- 

Cloud Services

(Correct)

- 

Query Processing

(Correct)

### Explanation

Snowflake's unique architecture consists of three key layers:

- Database Storage
- Query Processing
- Cloud Services

Question 45:

Skipped

How many maximum columns (or expressions) are recommended for a cluster key?

- Higher the number of columns (or expressions) in the key, better will be the performance
  - 7 to 8
  - 12 to 16
  - 3 to 4
- (Correct)

#### Explanation

A single clustering key can contain one or more columns or expressions. **Snowflake recommends a maximum of 3 or 4 columns (or expressions) per key for most tables.** Adding more than 3-4 columns tends to increase costs more than benefits.

Question 46:

Skipped

Snowflake supports SQL UDFs that return a set of rows. Which keyword in CREATE FUNCTION statement does need to be specified to enable UDF (i.e., UDTF) to return a set of rows?

- TABLE
- (Correct)
- ROWS
- SCALAR
- MULTIPLE

#### Explanation

TABLE keyword after RETURNS needs to be specified to create a UDTF (user-defined table function). Example :

```
create function t()  
returns table(msg varchar)  
as  
$$  
select 'Hello'  
union  
select 'World'  
$$;
```

**Note:** UDF returns a singular scalar value or if defined as a TABLE function, a set of rows. If you see UDTF in the exam, that simply means UDF that returns a set of rows.

Question 47:

**Skipped**

**What sized tables will experience the most benefit from clustering?**

- Tables with sizes between the range of 1 GB to 10 GB compressed**
- Tables with sizes between the range of 100 MB to 1 GB compressed**
- Tables in the multi-terabyte (TB) range**  
**(Correct)**
- All sizes of tables**

#### **Explanation**

Generally, **tables in the multi-terabyte (TB) range will experience the most benefit from clustering**, mainly if DML is performed regularly/continually on these tables.

Question 48:

**Skipped**

**Which Snowsight interface does help in setting up Multi-factor authentication (MFA)?**

- 

**Admin Interface**

- 

**Account Selector Interface**

- 

**User Menu Interface**

**(Correct)**

- 

**You can not setup Multi-factor authentication (MFA) using Snowsight interface**

- 

**Left Nav interface**

**Explanation**

**There are three interfaces in Snowsight. Left Nav, User Menu, and Account Selector.**

- Left Navigation consists of Worksheets, Dashboards, Data, Marketplace, Activity, Admin, Help & Support.
- User Menu lets you Switch Role, Profile including multi-factor authentication (MFA) , Partner Connect, Documentation, Support and Sign Out.
- The account selector, located at the bottom of the left nav, lets you sign in to other Snowflake accounts.

Question 49:

**Skipped**

**If we make any changes to the original table, then**

- 

**The changes do not reflect in the cloned table**

**(Correct)**

-

**The changes get immediately reflected in the cloned table**

- 

**The cloned table data get refreshed with the entire new data of the source table**

### Explanation

Zero-copy cloning allows us to make a snapshot of any table, schema, or database without actually copying data. **A clone is writable and is independent of its source (i.e., changes made to the source or clone are not reflected in the other object).** A new clone of a table points to the original table's micro partitions, using no data storage. If we make any changes in the cloned table, then only its changed micro partitions are written to storage.

Question 50:

**Skipped**

**Which table function in the Snowflake Information Schema can be used to query the replication history for a specified database within a specified date range?**

- 

**DATABASE\_REFRESH\_HISTORY**

- 

**REPLICATION\_USAGE\_HISTORY**

**(Correct)**

- 

**REPLICATION\_GROUP\_REFRESH\_HISTORY**

- 

**DATA\_TRANSFER\_HISTORY**

### Explanation

The table function **REPLICATION\_USAGE\_HISTORY** in Snowflake Information Schema can be used to query the replication history for a specified database within a specified date range. The information returned by the function includes the database name, credits consumed and bytes transferred for replication.

Question 51:

**Skipped**

**Which of these are not supported by the Search Optimization Service? (Select all that apply)**

-

## Casts on table columns

(Correct)

- 

## Columns defined with COLLATE clause

(Correct)

- 

## External Tables

(Correct)

- 

## Column Concatenation

(Correct)

- 

## Materialized Views

(Correct)

- 

## Analytical Expressions

(Correct)

## Explanation

**None of these are currently supported by the Search Optimization**

**Service.** Additionally, Tables and views protected by row access policies cannot be used with the Search Optimization Search.

The search optimization service can improve the performance of queries that use:

- **Equality predicates** (for example, <column\_name> = <constant>).
- Predicates that use IN

Question 52:

**Skipped**

**Which privileges are provided with a share by the provider? (Select 2)**

- 

**Grant access(SELECT) to the specific tables in the database**

**(Correct)**

- 

**Grant access(OPERATE) to the database and the schema containing the tables to share**

- 

**Grant access(USAGE) to the database and the schema containing the tables to share**

**(Correct)**

- 

**Grant access(USAGE) to the specific tables in the database**

- 

**Grant access(MODIFY) to the specific tables in the database**

### **Explanation**

Shares are named Snowflake objects that encapsulate all of the information required to share a database. Each share consists of:

- The privileges that grant access to the database(s) and the schema containing the objects to share.
- The privileges that grant access to the specific objects in the database.
- The consumer accounts with which the database and its objects are shared.

Example:

```
CREATE SHARE "SHARED_DATA" COMMENT="";
```

```
GRANT USAGE ON DATABASE "DEMO_DB" TO SHARE "SHARED_DATA";
```

```
GRANT USAGE ON SCHEMA "DEMO_DB"."TWITTER_DATA" TO SHARE "SHARED_DATA";
```

```
GRANT SELECT ON VIEW "DEMO_DB"."TWITTER_DATA"."FOLLOWERS" TO SHARE "SHARED_DATA";
```

Question 53:

**Skipped**

The LIST command returns a list of files that have been staged. Which of these stages supports the LIST command?

- 
- All of these**  
**(Correct)**
- 
- Stage for the current user.**
- 
- Stage for a specified table.**
- 
- Named internal stage.**
- 
- Named external stage.**

#### Explanation

The LIST command returns a list of files staged from these specified snowflake stages.

Question 54:

**Skipped**

User-managed Tasks is recommended when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources. (True /False)

- 
- TRUE**  
**(Correct)**
- 
- FALSE**

#### Explanation

User-managed Tasks is recommended when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources.

Serverless Tasks is recommended when you cannot fully utilize a warehouse because too few tasks run concurrently or they run to completion quickly (in less than 1 minute).

Question 55:

**Skipped**

**User-managed tasks are recommended when you cannot fully utilize a warehouse because only a few tasks run concurrently or they run to completion quickly (in less than 1 minute). (True / False)**

- **TRUE**
- **FALSE**

**(Correct)**

**Explanation**

**Serverless Tasks is recommended when you cannot fully utilize a warehouse because too few tasks run concurrently or they run to completion quickly (in less than 1 minute).**

User-managed Tasks is recommended when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources.

Question 56:

**Skipped**

**Which of these Snowflake tasks can be performed by Time Travel? (Select 3)**

- **Create clones of entire tables, schemas, and databases at or before specific points in the past**  
**(Correct)**
- **Query data in the past that has since been updated or deleted**  
**(Correct)**
-

Share the restored data objects over a specified period of time

- 

Restore tables, schemas, and databases that have been dropped.

(Correct)

### Explanation

Using Time Travel, you can perform the following actions within a defined period:

- Query data in the past that has since been updated or deleted.
- Create clones of entire tables, schemas, and databases at or before specific points in the past.
- Restore tables, schemas, and databases that have been dropped.

Question 57:

Skipped

Suppose we have a table t1. We drop the table t1 and then create a new table t1 again. What will happen if we execute the UNDROP command to restore dropped t1 table now?

- 

The dropped table t1 will be restored with a new arbitrary name set by Snowflake

- 

The dropped table t1 will be restored with name t1

- 

UNDROP command will fail

(Correct)

### Explanation

If an object with the same name already exists, UNDROP fails. We must rename the existing object, which then enables us to restore the previous version of the object.

Question 58:

Skipped

What types of accounts are involved in data sharing? (Select 3)

- 

Reader Accounts

(Correct)

- 
- Data Publishers**
- 
- Data Consumers**
- (Correct)**
- 
- Shared Accounts**
- 
- Data Providers**
- (Correct)**

### Explanation

There are three types of accounts involved in data sharing.

**Data Providers:** Share data with others

**Data Consumers:** Accesses shared data with their own Snowflake account.

**Reader Accounts:** Query data using compute from the data provider's account.  
Reader Accounts are what you can use to share data with somebody who does not already have a Snowflake account.

Question 59:

**Skipped**

**Which of these Snowflake Connectors are available? (Select all that apply)**

- 
- Snowflake Connector for Spark**
- (Correct)**
- 
- Snowflake Connector for ODBC**
- 
- Snowflake Connector for Kafka**

(Correct)

- 

**Snowflake Connector of Python**

(Correct)

- 

**Snowflake Connector for JDBC**

#### Explanation

ODBC and JDBC are drivers. **Connectors available for Snowflake are Python, Kafka, and Spark.** Snowflake also provides several drivers like ODBC, JDBC, Node.js, Go,.Net, and PHP PDO. The Snowflake SQL API is a REST API that you can use to access and update data in a Snowflake database.

Question 60: **Correct**

**Which is the default timestamp in Snowflake?**

- 

**TIMESTAMP\_TZ**

- 

**TIMESTAMP\_LTZ**

- 

**TIMESTAMP\_NTZ**

(Correct)

- 

**None of these**

#### Explanation

**TIMESTAMP\_NTZ** is the default timestamp type if you just define a column as a timestamp. Hint to remember : NTZ represents NO TIME ZONES.

Question 61:

**Skipped**

**Which of the following data storage does incur the cost? (Select 3)**

- 

**Only Active and Fail-Safe storage**

- All storage except Fail-Safe storage

- Fail-Safe Storage

(Correct)

- Time Travel Storage

(Correct)

- Active data Storage

(Correct)

- Only Active and Time Travel Storage

#### Explanation

Storage is calculated and charged for data regardless of whether it is in the **Active, Time Travel, or Fail-safe state.**

Question 62:

Skipped

How long does Snowflake keep batch load history (from Stage) using COPY statement?

- 30 days

- 14 days

- 64 days

(Correct)

**31 days**

- 

**1 day**

#### **Explanation**

**Snowflake keeps the batch load history for 64 days.**

Question 63:

**Skipped**

**Which of the following Data Types are supported by Snowflake? (Select all that apply)**

- 

**BOOL**

- 

**INTEGER**

**(Correct)**

- 

**FLOAT**

**(Correct)**

- 

**NUMERIC**

**(Correct)**

- 

**VARCHAR**

**(Correct)**

- 

**CHAR**

**(Correct)**

#### **Explanation**

All of these data types are supported by Snowflake except BOOL. BOOLEAN is the correct data type.

Question 64:

Skipped

Which of these are types of the stream? (Select 3)

- 

Standard

(Correct)

- 

Append-only

(Correct)

- 

External

- 

Update-only

- 

Insert-only

(Correct)

### Explanation

The following stream types are available based on the metadata recorded by each:

**Standard** - Supported for streams on tables, directory tables, or views. A standard (i.e. delta) stream tracks all DML changes to the source object, including inserts, updates, and deletes (including table truncates).

**Append-only** - Supported for streams on standard tables, directory tables, or views. An append-only stream tracks row inserts only. Update and delete operations (including table truncates) are not recorded.

**Insert-only** - Supported for streams on external tables only. An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops).

Question 65:

### Skipped

Snowflake stores metadata about all rows stored in a micro-partition, including  
(Select 3)

- 

**The number of distinct values**

**(Correct)**

- 

**The range of values for the first column in the micro-partition**

- 

**Additional properties used for both optimization and efficient query processing**

**(Correct)**

- 

**The number of similar values**

- 

**The range of values for each of the columns in the micro-partition**

**(Correct)**

### Explanation

Micro-partitioning is automatically performed on all Snowflake tables. Tables are transparently partitioned using the Ordering of the data as inserted/loaded.

Snowflake stores metadata about all rows stored in a micro-partition, including:

- The range of values for each of the columns in the micro-partition.
- The number of distinct values.
- Additional properties used for both optimization and efficient query processing.

Question 66:

### Skipped

Which data types are not supported by the Search Optimization Service? (Select 4)

- 

**Floating-point data types**

(Correct)

- 

Fixed-point numbers (e.g. INTEGER, NUMERIC)

- 

VARCHAR

- 

BINARY

- 

Semi-structured data types

(Correct)

- 

DATE, TIME, and TIMESTAMP

#### Explanation

The search optimization service currently supports **equality predicate and IN list predicate** searches for the following data types: **Fixed-point numbers (e.g. INTEGER, NUMERIC). DATE, TIME, and TIMESTAMP. VARCHAR. BINARY.** Currently, the search optimization service does not support floating point data types, semi-structured data types, or other data types not listed above.

Question 67:

Skipped

Search optimization is a Database-level property applied to all the tables within the database with supported data types. (True/False)

- 

FALSE

(Correct)

- 

TRUE

#### Explanation

Search optimization is a table-level property and applies to all columns with supported data types. The search optimization service aims to significantly improve

the performance of selective point lookup queries on tables. A point lookup query returns only one or a small number of distinct rows. A user can register one or more tables to the search optimization service.

Question 68:

**Skipped**

**What actions can a consumer perform on a share? (Select 2)**

- 

**Query the shared data and join it with an existing table in their own account**

**(Correct)**

- 

**Import the same share to more than one database**

- 

**Execute Time Travel on a share**

- 

**Copy shared data into another table in their own account with CREATE TABLE AS**

**(Correct)**

- 

**Re-share the share**

- 

**Clone a share**

### **Explanation**

Shared databases are read-only. A consumer cannot UPDATE a share. However, the consumer can do a CREATE TABLE AS to make a point-in-time copy of the data that's been shared. The consumer cannot clone and re-share a share or forward it. And also, time travel data on a share is not available to the consumer. A share can be imported into one database.

**Note: Very important for the exam. You can expect 2-3 questions on what a consumer can or can not do with a share.**

Question 69:

**Skipped**

Which of these types of VIEW does Snowflake support? (Select 3)

- 

#### MATERIALIZED VIEW

(Correct)

- 

#### EXTERNAL VIEW

- 

#### SECURE VIEW

(Correct)

- 

#### STANDARD VIEW

(Correct)

- 

#### TEMPORARY VIEW

- 

#### PERMANENT VIEW

#### Explanation

Snowflake supports three types of views.

Standard View, Secure View, and Materialized View.

- **Standard View:** It is a default view type. Its underlying DDL is available to any role with access to the view. When you create a standard view, Snowflake saves a definition of the view. Snowflake does not run the query. When someone accesses the view, that is when the query is run. The standard view will always execute as the owning role.
- **Secure View:** The secure view is exactly like a standard view, except users cannot see how that view was defined. Sometimes a secure view will run a

little slower than a standard view to protect the information in a secure view. Snowflake may bypass some of the optimizations.

- **Materialized View:** A materialized view is more like a table. Unlike a standard or secure view, Snowflake runs the query right away when you create a materialized view. It takes the results set and stores that result set as a table in Snowflake. Because Snowflake is storing that materialized view as a table, creating micro partitions. Snowflake is creating metadata about those micro partitions. So when you query a materialized view, if you put a filter on the view, you get the same benefit of micro partition pruning that you would get from a table. With Snowflake, the materialized view is automatically refreshed every time there is a transaction against the base table. So it is always going to be in sync. If you want, you can also create a secure materialized view, which again will hide the logic from the user. A note about materialized views, because Snowflake is auto-refreshing them in the background, they use some credits, so there is a little bit of a cost there. Moreover, there is some storage, and Snowflake stores the result set as a table in Snowflake. So materialized views use more storage and compute than standard or secure views.

Question 70:

**Skipped**

**In what situations should you consider User-Managed Tasks over Serverless Tasks? (Select 2)**

- 

**Consider when adherence to the schedule interval is highly important.**

- 

**Consider when you cannot fully utilize a warehouse because too few tasks run concurrently or they run to completion quickly (in less than 1 minute).**

- 

**Consider when adherence to the schedule interval is less important.**

**(Correct)**

- 

**Consider when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources.**

**(Correct)**

**Explanation**

**User-managed Tasks** is recommended when you can fully utilize a single warehouse by scheduling multiple concurrent tasks to take advantage of available compute resources. Also, recommended when adherence to the schedule interval is less critical. Serverless Tasks is recommended when you cannot fully utilize a warehouse because too few tasks run concurrently or they run to completion quickly (in less than 1 minute). Also, recommended when adherence to the schedule interval is critical.

Question 71: **Incorrect**

If you make any changes (e.g., insert, update) in a cloned table, then \_\_

- Only the changed micro partitions are written to the data storage  
**(Correct)**
- The entire table is written to data storage
- Cloned tables are read-only, you can not make any changes  
**(Incorrect)**
- The source table also gets updated with the new changes in the cloned table

### Explanation

Zero-copy cloning allows us to make a snapshot of any table, schema, or database without actually copying data. A clone is writable and is independent of its source (i.e., changes made to the source or clone are not reflected in the other object). A new clone of a table points to the original table's micro partitions, using no data storage. **If we make any changes in the cloned table, then only its changed micro partitions are written to storage.**

Question 72:

**Skipped**

How long does Snowflake keep Snowpipe's load history?

- 1 day
- 30 days

- 14 days  
**(Correct)**
- 31 days
- 64 days

#### Explanation

Snowflake keeps the Snowpipe's load history for 14 days.

[Note / Important for exam]: If you recreate the PIPE then the load history will reset to empty.

Question 73:

**Skipped**

Snowflake stores data into its

- internal optimized, uncompressed, columnar format
- internal optimized, compressed, columnar format  
**(Correct)**
- internal optimized, uncompressed, row format
- internal optimized, compressed, row format

#### Explanation

When data is loaded into Snowflake, Snowflake reorganizes that data into its internal optimized, compressed columnar format. Snowflake stores this optimized data in cloud storage.

Question 74:

**Skipped**

Which of these Snowflake features does enable accessing historical data (i.e., data that has been changed or deleted) at any point within a defined period?

- 

**Zero Copy Cloning**

- 

**Time Travel**

**(Correct)**

- 

**Search Optimization Service**

- 

**Data Sharing**

**Explanation**

**Snowflake Time Travel** enables accessing historical data (i.e. data that has been changed or deleted) at any point within a defined period. It serves as a powerful tool for performing the following tasks:

- Restoring data-related objects (tables, schemas, and databases) that might have been accidentally or intentionally deleted.
- Duplicating and backing up data from key points in the past.
- Analyzing data usage/manipulation over specified periods of time.

Question 75:

**Skipped**

**Snowflake architecture is**

- 

**None of these**

- 

**Shared-nothing architecture**

- 

**Hybrid of Shared-disk and Shared-nothing database architectures**

**(Correct)**

- 

### Shared-disk architecture

#### Explanation

Snowflake's architecture is a hybrid of traditional shared-disk and shared-nothing database architectures. Like shared-disk architectures, Snowflake uses a central data repository for persisted data accessible from all compute nodes in the platform. But similar to shared-nothing architectures, Snowflake processes queries using MPP (massively parallel processing) compute clusters where each node in the cluster stores a portion of the entire data set locally. This approach offers the data management simplicity of a shared-disk architecture but with the performance and scale-out benefits of a shared-nothing architecture. It is also termed as Multi-Cluster Shared Data Architecture.

Question 76:

**Skipped**

**Snowflake is available in four editions. Which are those? (Select 4)**

- 

**Business Critical**

**(Correct)**

- 

**Standard**

**(Correct)**

- 

**Professional Plus**

- 

**Virtual Private Snowflake (VPS)**

**(Correct)**

- 

**Professional**

- 

**Enterprise**

**(Correct)**

**Explanation**

**Snowflake is available in four editions: Standard, Enterprise, Business Critical, and Virtual Private Snowflake (VPS).**

Standard comes with most of the available features.

Enterprise adds on to Standard with things like extra days of time travel, materialized view support, and data masking.

Business Critical brings to the table: HIPAA support, Tri-secret Secure, and more.

Virtual Private Snowflake is everything that Business Critical has, but with the ability to have customer-dedicated metadata stores and customer-dedicated virtual service.

Question 77:

**Skipped**

**Which database objects are currently not supported for replication? (Select 3)**

- 

**Stages**

**(Correct)**

- 

**Views**

- 

**Temporary tables**

**(Correct)**

- 

**Transient tables**

- Streams
- (Correct)

#### Explanation

Temporary tables, stages, tasks, streams, pipes, and external tables are not currently supported for replication.

Question 78:

Skipped

Which is not the DDL (Data Definition Language) command?

- UNDROP
  - 
  - DROP
  - 
  - SHOW SHARES
  - 
  - CREATE
  - 
  - ALTER
  - 
  - TRUNCATE
- (Correct)

#### Explanation

TRUNCATE is DML (Data Manipulation Language) command.

Question 79:

Skipped

How can an ACCOUNTADMIN view the billing for Automatic Clustering? (Select all that apply)

-

**Classic Web Interface: Click on Account > Billing & Usage under warehouse named 'AUTOMATIC\_CLUSTERING'**

(Correct)

- 

**Snowsight: Select Admin > Usage**

(Correct)

- 

**Query - AUTOMATIC\_CLUSTERING\_HISTORY View (in Account Usage)**

(Correct)

- 

**There is no way to check the Automatic Clustering billing without contacting Snowflake Support Team**

- 

**Classic Web Interface: Click on Account > Billing & Usage under storage named 'AUTOMATIC\_CLUSTERING'**

- 

**Query - AUTOMATIC\_CLUSTERING\_HISTORY table function (in the Snowflake Information Schema)**

(Correct)

### Explanation

Users with the ACCOUNTADMIN role can view the billing for Automatic Clustering using Snowsight, the classic web interface, or SQL:

**Snowsight:** Select Admin » Usage.

**Classic Web Interface:** Click on Account tab » Billing & Usage

The billing for Automatic Clustering shows up as a separate Snowflake-provided warehouse named AUTOMATIC\_CLUSTERING.

**SQL:** Query either of the following: AUTOMATIC\_CLUSTERING\_HISTORY table function (in the Snowflake Information Schema).

AUTOMATIC\_CLUSTERING\_HISTORY View (in Account Usage).

Question 80:

**Skipped**

Which command can be used to resume Automatic Clustering for a table?

- 

**START TABLE**

- 

**RESUME RECLUSTER**

- 

**TRIGGER CLUSTERING**

- 

**ALTER TABLE**

**(Correct)**

### Explanation

Example - ALTER TABLE EMPLOYEE RESUME RECLUSTER; please note that RESUME RECLUSTER is a clause, not a command.

Question 81:

**Skipped**

You have a table with a 30-day retention period. If you decrease the retention period to 20 days, how would it affect the data that would have been removed after 30 days?

- 

**The data will still retain for 30-day before moving to Fail-safe**

- 

**The data will now retain for a shorter period of 20 days**

**(Correct)**

### Explanation

Decreasing Retention reduces the amount of time data is retained in Time Travel:

- For active data modified after the retention period is reduced, the new shorter period applies.
- For data that is currently in Time Travel:
  - If the data is still within the new shorter period, it remains in Time Travel.

- If the data is outside the new period, it moves into Fail-safe.

For example, if you have a table with a 30-day retention period and you decrease the period to 20-day, data from days 21 to 30 will be moved into Fail-safe, leaving only the data from day 1 to 20 accessible through Time Travel.

However, the process of moving the data from Time Travel into Fail-safe is performed by a background process, so the change is not immediately visible. Snowflake guarantees that the data will be moved, but does not specify when the process will complete; until the background process completes, the data is still accessible through Time Travel.

Question 82:

**Skipped**

**Snowflake supports extended DMLs. Which of the following commands are extended DMLs? (Select all that apply)**

- 

**UNDROP**

- 

**GET**

**(Correct)**

- 

**LIST**

**(Correct)**

- 

**COPY INTO <table>**

**(Correct)**

- 

**PUT**

**(Correct)**

- 

**REMOVE**

(Correct)

### Explanation

#### Extended DMLs:

**File Staging Commands** – PUT (to a stage), GET (from a stage), LIST and REMOVE. These commands are specific for working with stages.

**Data Loading / Unloading DML** – COPY INTO <table>, COPY INTO <location>

Question 83:

Skipped

Which of these columns gets appended on creating a stream on a table? (Select 3)

- **METADATA\$ISDELETE**
- **METADATA\$ISUPDATE**

(Correct)

- **METADATA\$ISINSERT**
- **METADATA\$ACTION**

(Correct)

- **METADATA\$ROW\_ID**

(Correct)

### Explanation

Adding a stream to a table appends three metadata columns: **METADATA\$ACTION**, **METADATA\$ISUPDATE**, **METADATA\$ROW\_ID**. These columns track the CDC records and their type: appends, deletes, or both (updates = inserts + deletes).

**METADATA\$ACTION** - Indicates the DML operation (INSERT, DELETE) recorded.

**METADATA\$ISUPDATE** - Indicates whether the operation was part of an UPDATE statement.

**METADATA\$ROW\_ID** - Specifies the unique and immutable ID for the row, which can be used to track changes to specific rows over time.

Question 84:

**Skipped**

**Which systems function can help find the overlap depth of a table's micro-partitions?**

- 

**SYSTEM\$CLUSTERING\_DEPTH**

**(Correct)**

- 

**SYSTEM\$CLUSTERING\_INFORMATION**

**(Correct)**

- 

**SYSTEM\$CLUSTERING\_ALL**

- 

**SYSTEM\$CLUSTERING\_INFO**

- 

**SYSTEM\$CLUSTERING\_WEIGHT**

### **Explanation**

For example, if you have an EMPLOYEE table - you can run any of these queries to find the depth - `SELECT SYSTEM$CLUSTERING_INFORMATION('EMPLOYEE');` `SELECT SYSTEM$CLUSTERING_DEPTH('EMPLOYEE');`

Question 85:

**Skipped**

**Which command will list the pipes for which you have access privileges?**

- 

**LIST PIPES;**

- 

**SHOW PIPES();**

- LIST PIPES();
- DESCRIBE PIPES;
- SHOW PIPES;

(Correct)

#### Explanation

**SHOW PIPES Command lists the pipes for which you have access privileges.** This command can list the pipes for a specified database or schema (or the current database/schema for the session) or your entire account.

Question 86:

Skipped

**Which roles can use SQL to view the task history within a specified date range?**  
(Select all that apply)

- 

**Role that has the global MONITOR EXECUTION privilege**

(Correct)

- 

**Task Owner having OWNERSHIP privilege on a task**

(Correct)

- 

**Account Administrator (ACCOUNTADMIN)**

(Correct)

#### Explanation

All of these roles can use SQL to view the task history within a specified date range. To view the run history for a single task: Query the TASK\_HISTORY table function (in the Snowflake Information Schema). To view details on a DAG run that is currently scheduled or is executing: Query the CURRENT\_TASK\_GRAPHS table function (in the Snowflake Information Schema). To view the history for DAG runs that executed successfully, failed, or were canceled in the past 60 minutes: Query the

COMPLETE\_TASK\_GRAPHS table function (in the Snowflake Information Schema).  
Query the COMPLETE\_TASK\_GRAPHS View (in Account Usage).

Question 87:

**Skipped**

Which view in the Account Usage Schema can be used to query the replication history for a specified database?

- 
- DATABASE\_REFRESH\_HISTORY**
- 
- DATA\_TRANSFER\_HISTORY**
- 
- REPLICATION\_GROUP\_REFRESH\_HISTORY**
- 
- REPLICATION\_USAGE\_HISTORY**

**(Correct)**

#### Explanation

This **REPLICATION\_USAGE\_HISTORY** view in the Account Usage Schema can be used to query the replication history for a specified database. The returned results include the database name, credits consumed, and bytes transferred for replication. Usage data is retained for 365 days (1 year).

Question 88:

**Skipped**

Time Travel can be disabled for an account by ACCOUNTADMIN. (True/False)

- 
- TRUE**
- 
- FALSE**

**(Correct)**

#### Explanation

Time Travel cannot be disabled for an account. A user with the ACCOUNTADMIN role can set DATA\_RETENTION\_TIME\_IN\_DAYS to 0 at the account level, which means that all databases (and subsequently all schemas and tables) created in the

account have no retention period by default; however, this default can be overridden at any time for any database, schema, or table.

Question 89:

**Skipped**

**The data retention period for a database, schema, or table can not be changed once ACCOUNTADMIN sets it at the account level. (True/False)**

- 
- TRUE
- 
- FALSE

**(Correct)**

**Explanation**

**The data retention period for a database, schema, or table can be changed at any time.** DATA\_RETENTION\_TIME\_IN\_DAYS parameter can be used to explicitly override the default when creating a database, schema, and individual table. For example: CREATE TABLE t1 (c1 int) DATA\_RETENTION\_IN\_DAYS=90;

Question 90: **Incorrect**

**A task can execute any one of the following types of SQL code: (Select 3)**

- Call to a stored procedure  
**(Correct)**
- Single SQL Statement  
**(Correct)**
- Procedural logic using Snowflake Scripting  
**(Correct)**
- Multiple SQL statements

**Explanation**

A task can execute any one of the following types of SQL code:

- Single SQL statement
- Call to a stored procedure
- Procedural logic using Snowflake Scripting.

Question 91:

**Skipped**

**What is the maximum data retention period for transient databases, schemas, and tables for Snowflake Enterprise Edition (and higher)?**

- **0 days**
- 
- **30 days**
- 
- **1 day**
- 
- **90 days**

### **Explanation**

For Snowflake Enterprise Edition (and higher):

- **For transient databases, schemas, and tables, the retention period can be set to 0 (or unset back to the default of 1 day). The same is also true for temporary tables.**
- For permanent databases, schemas, and tables, the retention period can be set to any value from 0 up to 90 days.

Question 92:

**Skipped**

**In which layer does Snowflake perform query execution?**

- 
- **Query Processing**
-

## Cloud Services

- 

**None of these**

- 

**Database Storage**

## Explanation

Query execution is performed in the processing layer. Snowflake processes queries using “virtual warehouses.”

Question 93:

**Skipped**

**Which stream type is supported for streams on the external table only?**

- 

**Insert-only**

**(Correct)**

- 

**External**

- 

**Append-only**

- 

**Update-only**

- 

**Standard**

## Explanation

**Insert-only is supported for streams on external tables only.** An insert-only stream tracks row inserts only; they do not record delete operations that remove rows from an inserted set (i.e. no-ops).

Question 94:

**Skipped**

**What is the default standard data retention period automatically enabled for all Snowflake accounts?**

- 1 day  
**(Correct)**
- 0 days
- 30 days
- 90 days

#### Explanation

The standard retention period is 1 day (24 hours) and is automatically enabled for all Snowflake accounts.

Question 95:

**Skipped**

The snowflake data warehouse is not built on an existing database or “big data” software platform like Hadoop.(True/False)

- FALSE
- TRUE

**(Correct)**

#### Explanation

Snowflake is a 100% cloud-native data platform.

Question 96:

**Skipped**

John has a SECURITYADMIN role. He created a custom DBA\_ROLE. He granted the SYSADMIN role to DBA\_ROLE. He created a user, 'Monica'. John then granted DBA\_ROLE to Monica. Monica creates a Database Monica\_DB. Monica then created a Table T1 in Monica\_DB under the PUBLIC schema. What should John do to access Table T1, created by Monica?

-

```
GRANT ROLE DBA_ROLE TO John; USE ROLE DBA_ROLE; USE DATABASE monica_db; Select * from t1;
```

(Correct)

- 

```
GRANT ROLE DBA_ROLE TO John; USE DATABASE monica_db; Select * from t1;
```

- 

```
USE ROLE dba_role; USE DATABASE monica_db; Select * from t1;
```

- 

```
USE ROLE SECURITYADMIN; USE DATABASE monica_db; Select * from t1;
```

#### Explanation

It does not matter if John has created the DBA\_ROLE. If John wants to access the object created by DBA\_ROLE, he needs to grant DBA\_ROLE to himself.

Question 97:

**Skipped**

Which capabilities are available in Snowsight (the new Snowflake web interface)?  
(Select all that apply)

- 

Snowflake Marketplace is not available with Snowsight currently

- 

Sharing data with other Snowflake accounts

(Correct)

- 

You can display visual statistics on columns (SUM, MIN, MAX, etc.) without re-running the query

(Correct)

- 

Creating and managing users and other account-level objects

(Correct)

- 

**The smart autocompletes feature suggests SQL or object syntax to insert**

**(Correct)**

### **Explanation**

Snowsight is the new Snowflake Web Interface. It can be used to perform the following operations:

- Building and running queries.
- Loading data into tables.
- Monitoring query performance and copy history.
- Creating and managing users and other account-level objects.
- Creating and using virtual warehouses.
- Creating and modifying databases and all database objects.
- Sharing data with other Snowflake accounts.
- Exploring and using the Snowflake Marketplace.
- One of the cool features is the smart autocomplete, which suggests SQL or object syntax to insert.

Question 98:

**Skipped**

**Which table types does Snowflake support? (Select all that apply)**

- 

**MATERIALIZED TABLE**

- 

**TEMPORARY TABLE**

**(Correct)**

- 

**TRANSIENT TABLE**

**(Correct)**

- 

**EXTERNAL TABLE**

**(Correct)**

- 

**SECURED TABLE**

- 

## PERMANENT TABLE

(Correct)

### Explanation

Snowflake supports four different table types: Permanent Table, Temporary Table, Transient Table, and External Table.

- **Permanent Table:** It persists until dropped. It is designed for data requiring the highest data protection and recovery level and is the default table type. Permanent Tables can be protected by up to 90 days of time travel with Enterprise Edition or above. Moreover, the failsafe is covered on all the Permanent Tables.
- **Temporary Table:** A Temporary table is tied to a specific session, which means it is tied to a single user. Temporary tables are used for things like materializing subquery. You can only cover temporary tables by up to one day of time travel, and they are not covered by a failsafe.
- **Transient Table:** A Transient table is essentially a temporary table that more than one user can share because multiple users share a transient table. You have to drop it when you are finished with it, and it also is only covered by up to one day of time travel and is not covered by a failsafe. NOTE - WE CAN ALSO HAVE TRANSIENT DATABASES AND SCHEMAS.
- **External Table:** An External Table is used to access data in a data lake. It is always read-only because it is based on files that live outside of Snowflake and are not managed by Snowflake, and Time Travel and Failsafe do not cover it.

Question 99:

Skipped

In the case of cloning massive databases or schemas, the original databases and schemas get locked while the cloning operation is running. While cloning is in progress, no DML operation can be done on the original databases and schemas. (True/False)

- 

TRUE

- 

FALSE

(Correct)

### Explanation

Cloning is not instantaneous, particularly for large objects (databases, schemas, tables), and **does not lock the object being cloned**. A clone does not reflect any DML

statements applied to table data, if applicable, while the cloning operation is still running.

Question 100:

**Skipped**

Which of the following is not a type of Snowflake's Internal stage?

- 

**User Stage**

- 

**Schema Stage**

**(Correct)**

- 

**Table Stage**

- 

**Name Stage**

### Explanation

An internal stage is a cloud repository that resides within a Snowflake account and is managed by Snowflake. An external stage is a pointer to a cloud file repository outside a Snowflake account, which the customer manages independently. **There are three types of stages, and they are table stage, user stage, and named stage.** **Table Stage:** When you create a table, the system will create a table stage with the same name but with the prefix @%. **User Stage:** A user stage is created whenever you create a new user in Snowflake. The user stage uses the @~. **Named Stage:** Named stages are created manually. They can be internal or external and are prefixed with an @ and then the stage's name.

Question 101:

**Skipped**

Which object parameter can users with the ACCOUNTADMIN role use to set the minimum retention period for their account?

- 

**MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS**

**(Correct)**

- 

**DATA\_RETENTION\_TIME\_IN\_MIN\_DAYS**

- MIN\_DATA\_RETENTION\_TIME\_IN\_HOURS
- DATA\_RETENTION\_TIME\_IN\_DAYS

#### Explanation

The **MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS** account parameter can be set by users with the **ACCOUNTADMIN** role to set a minimum retention period for the account. This parameter does not alter or replace the **DATA\_RETENTION\_TIME\_IN\_DAYS** parameter value. However, it may change the effective data retention time. When this parameter is set at the account level, the effective minimum data retention period for an object is determined by **MAX(DATA\_RETENTION\_TIME\_IN\_DAYS, MIN\_DATA\_RETENTION\_TIME\_IN\_DAYS)**.

Question 102:

**Skipped**

While choosing the clustering key, what should we consider? (Select 3)

- Ordering the columns from highest cardinality to lowest cardinality

- Columns which are more less used in join conditions

- Columns which are more often used in join conditions

**(Correct)**

- Ordering the columns from lowest cardinality to highest cardinality

**(Correct)**

- Columns which are less often used in where clause

- Columns which are more often used in where clause

**(Correct)**

### **Explanation**

Best Practices for choosing clustering key: 1. Columns which are more often used in where clause 2. Columns that are more often used in join conditions 3. Order you specify the clustering key is important. As a general rule, Snowflake recommends ordering the lowest to highest cardinality columns.

Question 103:

**Skipped**

**Which of these are Snowflake Cloud Partner Categories? (Select 3)**

- 

**Machine Learning & Data Science**

**(Correct)**

- 

**Application Integration**

- 

**Data Integration**

**(Correct)**

- 

**Native Programmatic Interfaces**

**(Correct)**

### **Explanation**

Snowflake has the following Cloud Partner Categories:

- Data Integration
- Business Intelligence (BI)
- Machine Learning & Data Science
- Security Governance & Observability
- SQL Development & Management
- Native Programmatic Interfaces.

Question 104:

**Skipped**

**Which products does Snowflake offer for secure data sharing? (Select 3)**

- 

**Data Exchange**

(Correct)

- 

### Data Replication

- 

### Indirect share

- 

### Data Marketplace

(Correct)

- 

### Direct share

(Correct)

## Explanation

Snowflake provides three product offerings for data sharing that utilize Snowflake Secure Data Sharing to connect providers of data with consumers. **Direct Share:** It is the simplest form of data sharing that enables account-to-account sharing of data utilizing Snowflake's Secure Data Sharing. As a data provider, you can easily share data with another company so that your data shows up in their Snowflake account without having to copy it over or move it. **Data Exchange:** With a Snowflake data exchange, you actually set up a private exchange between partners that you want to have in this exchange, and any member of that exchange can share data in this private exchange. And any member of the exchange can also consume data from that exchange. So instead of one-to-one or one-to-many, it's many-to-many. But it's a very exclusive club. Only people who are invited into this exchange can access any of that data. **Data Marketplace:** The Snowflake Data Marketplace is where companies can publish their data to be consumed by anybody who has a Snowflake account and wants to connect to the marketplace and download that data.

Question 105:

Skipped

What happens to the data when the retention period ends for an object?

- 

Data is permanently lost

- 

Data can be restored by increasing the retention period

- ○  
**SYSADMIN can restore the data from Fail-safe**
- ○  
**Data is moved to Snowflake Fail-safe**  
**(Correct)**

#### **Explanation**

**When the retention period ends for an object, the historical data is moved into Snowflake Fail-safe.** Snowflake support needs to be contacted to get the data restored from Fail-safe.

- 1 .MFA support :snowsql,jdbc,odbc,python.
- 2.MFA will only go with SSO federation.
- 3.MFA need which software: duo
- 4.Object access by which view(scenario based)
5. Login history by which view(scenario based)prior 90 days
6. Role1 is owner of table , table is drop by useradmin who will be the owner?  
Opt: Useradmin, account admin, table will be owener less
- 7.Table is droped 30 min ago , what will happen when create clone clone at offset= -60  
Error , table will be clone , table will Clonewith empty row
8. Which obj will be undrop:(que repeated with diff opt) role,user ,schema,db,table,wh
9. Function name for masking col with token.
10. Select 2 for materialized view
  - Can query two table
  - Join can be joined with view
  - Join can defined with view
  - Etc
11. Question on stored procedure. (Que Repeated with diff opt)  
Answer: can run wil caller or owner rights
- 12.Architecture of snowflake.
13. Table name mytable, stage name mystage stage declaration: answer @ mystage
14. Wh cache contain: row access by result/ query , col access by result/ query
15. Query takes 2 min to fetch first then it take sec, from which layer it takes.  
Answer: opt contain both query cache and cloud service layer (que is asking about layer)
16. (Scenario based.)Magner want to info about schema. grant on schema./ Role/ user
17. (Scenario based.)want to info about all user and role in company. grant on schema./ Role/user and grant of role
18. Maximized mode= same on min and Max
- 19.lagre wh to 4 x wh expect to run query twice fast true or false
20. Change in wh size effect new or old query/ both/ only new
21. Query speed is slow?scale up to large, scale out to large
22. Want to change primary to secondary db which min edition support
23. Min edition for materialized view
24. Min edition for data sharing
25. When is data encrypted on client machine.
- 26.Reader account charge no stroge Cost on provider account. true or false
27. Sf function used for restoring.
28. Why Zero clone is better than  
creat tb as select
  - Create instantly
  - No storage cost
- 29.stage contain 10 file then it copy into tb
  - Again 10 file is added with update in previous file then
  - Error
  - Add data of new 10 file.

All old file will load again with new one

Only add file with updated value

30.non block command: update insert merge copy

31. Connectivity diagnostic tool : snowcd

32. One user created worksheet in snowsql want to share with teammates.

Create clone

Share with market place

Create privileged and assign to user

Question 1:

**Skipped**

Which of these configurations will set up a warehouse in maximized mode?

- 

**Minimum Clusters = 6 and Maximum Clusters = 6**

**(Correct)**

- 

**Minimum Clusters = 1 and Maximum Clusters = 1**

- 

**Minimum Clusters = 9 and Maximum Clusters = 10**

- 

**Minimum Clusters = 1 and Maximum Clusters = 10**

### Explanation

Maximized mode is enabled by specifying the same value for both maximum and minimum number of clusters (note that the specified value must be larger than 1). In this mode, when the warehouse is started, Snowflake starts all the clusters so that maximum resources are available while the warehouse is running. This mode is effective for statically controlling the available compute resources, particularly if you have large numbers of concurrent user sessions and/or queries and the numbers do not fluctuate significantly.

Question 2:

**Skipped**

ACCOUNTADMIN role cannot view the results for a query run by another user. (True/False)

- 

**TRUE**

**(Correct)**

- 

**FALSE**

### Explanation

A user cannot view the result set from a query that another user executed. This behavior is intentional. For security reasons, only the user who executed a query can access the query results. This behavior is not connected to the Snowflake access control model for objects. **Even a user with the ACCOUNTADMIN role cannot view the results for a query run by another user.**

Question 3:

Skipped

Which of these are kind of Cache in Snowflake?

- 
- Query Result Cache**
- 
- Data/Local Disk Cache**
- 
- All of these**
- (Correct)**
- 
- Metadata Cache**

### Explanation

Snowflake has three types of cache.

- The metadata cache that lives in the cloud services layer.
- The data cache/local disk cache that lives on the SSD drives in the virtual warehouses
- The query result cache. If a result is small, it will be stored in the cloud services layer, but larger results are going to be stored in the storage layer.

Question 4:

Skipped

Which privilege is required to enable altering any properties of the resource monitor, such as changing the monthly credit quote?

- 
- MODIFY**
- (Correct)**
-

## USAGE

- 

## MONITOR

### Explanation

**MODIFY** - Enables altering any properties of a resource monitor, such as changing the monthly credit quota.

**MONITOR** - Enables viewing a resource monitor.

Question 5:

### Skipped

Monica has run a query **SELECT \* FROM t1**; After a couple of hours, John ran the same query. John has the same role as Monica and has the **SELECT** permissions on table **t1**. John got the result sooner than Monica. What could be the reason for the faster result?

- 

**John's query resulted from the Query Result cache.**

**(Correct)**

- 

**John's query resulted from the Metadata cache.**

- 

**John's query resulted from the Local Disk cache.**

- 

**John's query resulted from Remote disk.**

### Explanation

It is a typical use case of Query Result Cache. It is stored and managed by the Cloud Services Layer. It is used if the identical query is run and base tables (**t1** in this case) have not changed. Query Result Cache doesn't require Virtual Warehouse and is available for other users in the same role with **SELECT** permissions on all tables in the query.

Question 6:

### Skipped

Which features of Snowflake provide Column-level security? (Select 2)

-

## Column Masking

- 

### Dynamic Data Masking

(Correct)

- 

### External Tokenization

(Correct)

- 

### Internal Tokenization

#### Explanation

**Dynamic Data Masking** is a Column-level Security feature that uses masking policies to selectively mask plain-text data in table and view columns at query time.

**External Tokenization** enables accounts to tokenize data before loading it into Snowflake and detokenize the data at query runtime. Tokenization is the process of removing sensitive data by replacing it with an undecipherable token. External Tokenization makes use of masking policies with external functions.

Question 7:

Skipped

Permissions on database objects such as databases or tables are granted to:

- 

**Virtual Warehouses**

- 

**Users**

- 

**Schemas**

- 

**Roles**

**(Correct)**

### Explanation

**Snowflake supports Role-Based Access control.** Permissions on database objects such as databases or tables are granted to Roles.

Question 8:

**Skipped**

**A user can be assigned multiple roles. (True / False)**

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### Explanation

Roles are the entities to which privileges on securable objects can be granted and revoked. Roles are assigned to users to allow them to perform actions required for business functions in their organization. **A user can be assigned multiple roles.** It allows users to switch roles (i.e., choose which role is active in the current Snowflake session) to perform different actions using separate sets of privileges.

Question 9:

**Skipped**

**The suspended warehouse cannot be resized until they resume. (True / False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### Explanation

The suspended warehouse can be easily resized. Resizing a suspended warehouse does not provision any new compute resources for the warehouse. It simply instructs Snowflake to provision the additional compute resources when the warehouse is next resumed, at which time all the usage and credit rules associated with starting a warehouse apply.

Question 10:

**Skipped**

Readers accounts enable providers to share data with consumers who are not already Snowflake customers without requiring the consumers to become Snowflake Customers. Which role can create the Reader account?

- 

**ACCOUNTADMIN**

**(Correct)**

- 

**USERADMIN**

- 

**SECURITYADMIN**

- 

**SYSADMIN**

**Explanation**

**ACCOUNTADMIN** role (or a role granted the CREATE ACCOUNT global privilege) only can create the Reader account.

Question 11:

**Skipped**

Both non-materialized and materialized views can be defined as secure. (True / False)

- 

**FALSE**

- 

**TRUE**

**(Correct)**

**Explanation**

**Yes, both non-materialized and materialized views can be defined as secure.**

Question 12:

**Skipped**

Federated Authentication is supported by the following:

(Select all that apply)

- 

**Standard Edition**

**(Correct)**

- 

**Business Critical**

**(Correct)**

- 

**Enterprise Edition**

**(Correct)**

- 

**VPS**

**(Correct)**

#### **Explanation**

**Federated authentication is supported by all of the Snowflake editions.**

Question 13:

**Skipped**

A DBA\_ROLE created a database. Later the DBA\_ROLE was dropped. Who will own the database now, which was created by the DBA\_ROLE?

- 

**The role that dropped the DBA\_ROLE will own the database.**

**(Correct)**

- 

**The database will get dropped too.**

- 

**The DBA\_ROLE can't get dropped as it is the database owner.**

- 

**No one will be able to access the database.**

#### **Explanation**

**The role that dropped the DBA\_ROLE will own the database.**

Question 14:

#### **Skipped**

Materialized views are particularly useful when:

(Select 3)

- 

**Query results contain results that require significant processing.**

**(Correct)**

- 

**Query results contain results that require simple processing.**

- 

**The view's base table changes frequently.**

- 

**The view's base table does not change frequently.**

**(Correct)**

- 

**Query results contain a small number of rows and/or columns relative to the base table (the table on which the view is defined).**

**(Correct)**

#### **Explanation**

Materialized views are particularly useful when:  
- Query results contain a small number of rows and/or columns relative to the base table (the table on which the view is defined).  
- Query results contain results that require significant processing, including: Analysis of semi-structured data. Aggregates that take a long time to calculate.  
- The query is on an external table (i.e., data sets stored in files in an external stage), which might have a slower performance compared to querying native database tables.  
- The view's base table does not change frequently.

Question 15:

**Skipped**

**What is the minimum billing charge for provisioning compute resources?**

- **120 seconds**
- **60 seconds**
- **(Correct)**  **1 second**
- **30 seconds**

**Explanation**

The minimum billing charge for provisioning compute resources is 1 minute (i.e. 60 seconds). There is no benefit to stopping a warehouse before the first 60-second period is over because the credits have already been billed for that period.

Question 16:

**Skipped**

Snowflake network policies currently support both Internet Protocol versions 4 and 6 (i.e., IPv4 as well as IPv6). (True/False)

- **FALSE**
- **(Correct)**  **TRUE**

**Explanation**

**Network policies currently support only Internet Protocol version 4 (i.e. IPv4) addresses.**

Question 17:

### Skipped

Suppose you have an auto-scaling mode setup with a Economy policy. In what situation does Snowflake spin up an additional cluster?

- 

The first cluster starts immediately when either a query is queued or the system detects that there's one more query than the currently-running clusters can execute.

- 

Only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes.

(Correct)

### Explanation

In Economy Scaling policy, Snowflake spins up an additional cluster only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes

Question 18:

### Skipped

What is a key benefit of scaling out a warehouse?

- 

Scaling out improves performance.

- 

Scaling out improves concurrency.

(Correct)

### Explanation

Scaling out is explicitly designed for handling queuing and performance issues related to large numbers of concurrent users and/or queries.

Question 19:

### Skipped

Snowflake prunes micro-partitions based on a predicate with a subquery, even if the subquery result is constant. (TRUE/FALSE)

- 

TRUE

- 

**FALSE**

(Correct)

#### Explanation

Please note, not all predicate expressions can be used to prune. **Snowflake does not prune micro-partitions based on a predicate with a subquery, even if the subquery results in a constant.**

Question 20:

Skipped

Which of these security features are supported in Snowflake? (Select all that apply)

- 

**Role-Based Access Control**

(Correct)

- 

**Network Policy**

(Correct)

- 

**Multi-Factor Authentication**

(Correct)

- 

**Tri-Secret Secure Encryption**

(Correct)

#### Explanation

**Snowflake is a highly secured platform and provides multi-level security like Multi-Factor Authentication (MFA), provision to set up Network policy to block access by unwanted IPs, Single Sign On (SSO), Role-Based Access Control, and Tri Secret Secure, and so on.** Tri-Secret Secure is the combination of a Snowflake-maintained key and a customer-managed key in the cloud provider platform that hosts your Snowflake account to create a composite master key to protect your Snowflake data.

Question 21:

### Skipped

Which AWS service is used to create private VPC endpoints that allow direct, secure connectivity between your AWS VPCs and the Snowflake VPC without traversing the public internet?

- 

**Snowflake PrivateLink**

- 

**AWS Direct Connect**

- 

**AWS PrivateVPC**

- 

**AWS PrivateLink**

**(Correct)**

### Explanation

**AWS PrivateLink** is an AWS service for creating private VPC endpoints that allow direct, secure connectivity between your AWS VPCs and the Snowflake VPC without traversing the public internet. The connectivity is for AWS VPCs in the same AWS region.

For External Functions, you can also use AWS PrivateLink with private endpoints.

In addition, if you have an on-premises environment (e.g. a non-hosted data center), you can choose to use AWS Direct Connect, in conjunction with AWS PrivateLink, to connect all your virtual and physical environments in a single, private network.

### Question 22:

### Skipped

If you have privileges to view queries executed by another user, the Query Detail page displays the details for the query and the actual query result. (TRUE/FALSE)

- 

**FALSE**

**(Correct)**

-

**TRUE**

**Explanation**

If you have privileges to view queries executed by another user, the Query Detail page displays the details for the query, but, for data privacy reasons, the page does not display the actual query result.

Question 23:

**Skipped**

**Is it possible to create a user without a password?**

- 

**No**

- 

**Yes**

**(Correct)**

**Explanation**

**Yes, it is possible to create a user in Snowflake without a password.** We cannot use the Snowflake web interface to create users with no passwords or remove passwords from existing users, and **we must use CREATE USER or ALTER USER**. Without a password in Snowflake, a user cannot log in using Snowflake authentication and must use federated authentication instead.

Question 24:

**Skipped**

Which of these system-defined roles can manage operations at the organizational level?

- 

**ORGADMIN**

**(Correct)**

- 

**SYSADMIN**

- 

**USERADMIN**

-

## ACCOUNTADMIN

- 

## SECURITYADMIN

### Explanation

**ORGADMIN** role manages operations at the organizational level.

More specifically, this role:

- Can create accounts in the organization.
- Can view all accounts in the organization (using SHOW ORGANIZATION ACCOUNTS) and all regions enabled for the organization (using SHOW REGIONS).
- Can view usage information across the organization.

Question 25:

### Skipped

Which privilege is required to enable changing the state of a warehouse (stop, start, suspend, resume)?

- 

## OPERATE

(Correct)

- 

## MODIFY

- 

## MONITOR

- 

## USAGE

### Explanation

MODIFY - Enables altering any properties of a warehouse, including changing its size.

Required to assign a warehouse to a resource monitor. Note that only the ACCOUNTADMIN role can assign warehouses to resource monitors. MONITOR - Enables viewing current and past queries executed on a warehouse as well as usage statistics on that warehouse.

OPERATE - Enables changing the state of a warehouse (stop, start, suspend, resume). In addition, enables viewing current and past queries executed on a warehouse and aborting any executing queries. USAGE - Enables using a virtual warehouse and, as a result, executing

queries on the warehouse. If the warehouse is configured to auto-resume when a SQL statement (e.g. query) is submitted to it, the warehouse resumes automatically and executes the statement. OWNERSHIP - Grants full control over a warehouse. Only a single role can hold this privilege on a specific object at a time. ALL [ PRIVILEGES ] - Grants all privileges, except OWNERSHIP, on the warehouse.

Question 26:

**Skipped**

The closer the ratio of scanned micro-partitions and columnar data is to the ratio of actual data selected, the more efficient the pruning performed on the table.  
(TRUE/FALSE)

- 

**FALSE**

- 

**TRUE**

**(Correct)**

### **Explanation**

Snowflake uses columnar scanning of partitions so that an entire partition is not scanned if a query only filters by one column. The closer the ratio of scanned micro-partitions and columnar data is to the ratio of actual data selected, the more efficient is the pruning performed on the table

Question 27:

**Skipped**

**Which role has privileges to create warehouses and databases (and other objects) in an account?**

- 

**ORGADMIN**

- 

**ACCOUNTADMIN**

- 

**USERADMIN**

-

## SYSADMIN

(Correct)

- 

## SECURITYADMIN

### Explanation

**SYSADMIN role has privileges to create warehouses and databases (and other objects) in an account.** This role also has the ability to grant privileges on warehouses, databases, and other objects to other roles.

Question 28:

Skipped

Which schema can be used to find out about storage, compute, and objects in a Snowflake account?

- 

## INFORMATION\_SCHEMA

(Correct)

- 

## USAGE\_SCHEMA

- 

## RESOURCE\_SCHEMA

- 

## SNOWFLAKE\_SCHEMA

### Explanation

INFORMATION\_SCHEMA can be used to find out about storage, compute, and objects in a Snowflake account. Every database that you create on Snowflake has a schema called INFORMATION\_SCHEMA that's automatically created, and inside that schema, you can find views and table functions that provide metadata information about objects in your account.

Question 29:

Skipped

Which of the Snowflake editions provides a federated authorization feature?

-

## Enterprise

- 
- **Virtual Private Snowflake(VPS)**
- 
- **Standard**
- 
- **Business Critical**
- 
- **All of the Snowflake Editions**

**(Correct)**

## Explanation

**All Snowflake Editions (Standard, Enterprise, Business Critical, Virtual Private Snowflake) provide Federated Authentication.**

Question 30:

**Skipped**

If you create a Network Policy by providing both 'Allowed IP Addresses' and 'Blocked IP Addresses', which is applied first by Snowflake while validating the access?

- 
- **Blocked IP Addresses**
- 
- **Allowed IP Addresses**

## Explanation

If you provide both Allowed IP Addresses and Blocked IP Addresses, Snowflake applies the Blocked List first.

Question 31:

**Skipped**

The user access history can be found by querying the

-

## Information Schema ACCESS\_REPORT view

- 

## Account Usage ACCESS\_REPORT view

- 

## Account Usage ACCESS\_HISTORY view

(Correct)

- 

## Information Schema ACCESS\_HISTORY view

### Explanation

Access History in Snowflake refers to when the user query reads column data and when the SQL statement performs a data write operation, such as INSERT, UPDATE, and DELETE, along with variations of the COPY command, from the source data object to the target data object. **The user access history can be found by querying the Account Usage ACCESS\_HISTORY view.**

Question 32:

Skipped

Which of these SQL Queries can be answered completely by Metadata? (Select 3)

- 

MIN, MAX (integers and dates)

(Correct)

- 

SHOW Commands

(Correct)

- 

None of these

- 

AVG

-

COUNT

(Correct)

**Explanation**

SHOW Commands, MIN, MAX (integers and dates), and COUNT SQL queries take advantage of the Metadata cache and do not require the virtual warehouse, but you still have some cloud service charges.

Question 33:

Skipped

What is a key benefit of scaling up a warehouse?

- 

**Scaling up improves performance.**

(Correct)

- 

**Scaling up improves concurrency.**

**Explanation**

Resizing a warehouse generally improves query performance, particularly for larger, more complex queries. It can also help reduce the queuing that occurs if a warehouse does not have enough compute resources to process all the queries that are submitted concurrently. Note that warehouse resizing is not intended for handling concurrency issues; instead, use additional warehouses to handle the workload or use a multi-cluster warehouse.

Question 34:

Skipped

Suppose we resize a warehouse to a smaller size while it is executing SQL statements. There will be no impact on already executing SQL statements. (True / False)

- 

**FALSE**

- 

**TRUE**

(Correct)

**Explanation**

Resizing a warehouse does not impact the statements that the warehouse is executing. When resizing to a smaller size, compute resources are removed from the warehouse only when they are no longer used to execute any current statements.

Question 35:

**Skipped**

If an account has federated authentication enabled. Can Snowflake admins still maintain user IDs and passwords in Snowflake?

- 

**No**

- 

**Yes**

**(Correct)**

### **Explanation**

With federated authentication enabled on an account, Snowflake still allows maintaining and using Snowflake user credentials (login name and password). In other words:

- Account and security administrators can still create users with passwords maintained in Snowflake.
- Users can still log into Snowflake using their Snowflake credentials.

However, if federated authentication is enabled for an account, Snowflake does not recommend maintaining user passwords in Snowflake. Instead, user passwords should be maintained solely in your IdP.

Question 36:

**Skipped**

Which objects are the securable objects in Snowflake? (Select all that apply)

- 

**Warehouse**

**(Correct)**

- 

**Table**

**(Correct)**

- 

### File Format

(Correct)

- 

### Database

(Correct)

### Explanation

All of these are securable objects in Snowflake. A securable Object is an entity to which access can be granted. Unless allowed by a grant, access will be denied.

Question 37:

Skipped

How many servers are available in a large-sized cluster warehouse?

- 

8

(Correct)

- 

32

- 

64

- 

16

- 

128

- 

1

### Explanation

There are eight servers available in a large-sized cluster warehouse.

Question 38:

Skipped

A role inherits all the privileges of those higher in the hierarchy. (True / False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### Explanation

A role inherits all the privileges of its underlying roles (those "lower" in the hierarchy). ACCOUNTADMIN inherits privileges from SECURITYADMIN, USERADMIN, SYSADMIN, and PUBLIC. SECURITYADMIN inherits privileges from USERADMIN and PUBLIC. USERADMIN and SYSADMIN inherit privileges from PUBLIC. PUBLIC inherits nothing.

Question 39:

Skipped

Which of these is not a valid authentication method supported by Snowflake?

- 

**Username and one-time generated pin**

**(Correct)**

- 

**Oauth, Key Pair**

- 

**Muti-factor authentication**

- 

**Username and password**

- 

**SCIM (System for Cross-domain Identity Management specification)**

-

## Federated Authentication (SAM 2.0)

### Explanation

Snowflake supports the following authentication methods:

- Username and password
- MFA
- SAML
- other authentication methods such as OAuth, Key-pair, and SCIM.

Question 40:

### Skipped

Which IdP vendors provide native Snowflake support For federated authentication and SSO? (Select 2)

- 

**Microsoft ADFS**

**(Correct)**

- 

**Microsoft Azure Active Directory**

- 

**Google G Suite**

- 

**Okta**

**(Correct)**

- 

**Onelogin**

### Explanation

**Okta and Microsoft ADFS** provide native Snowflake support for federated authentication and SSO.

Question 41:

### Skipped

Which privilege is required to execute queries using a virtual warehouse?

-

## OPERATE

- 

## MONITOR

- 

## USAGE

(Correct)

- 

## MODIFY

### Explanation

Virtual Warehouse Privileges: **USAGE:** Enables using a virtual warehouse and, as a result, executing queries on the warehouse. If the warehouse is configured to auto-resume when a SQL statement (e.g. query) is submitted to it, the warehouse resumes automatically and executes the statement.

**MODIFY:** Enables altering any properties of a warehouse, including changing its size. Required assigning a warehouse to a resource monitor. Note that only the ACCOUNTADMIN role can assign warehouses to resource monitors.

**MONITOR:** Enables viewing of current and past queries executed on a warehouse as well as usage statistics on that warehouse.

**OPERATE:** Enables changing the state of a warehouse (stop, start, suspend, resume). In addition, enables viewing current and past queries executed on a warehouse and aborting any executing queries.

**OWNERSHIP:** Grants full control over a warehouse. Only a single role can hold this privilege on a specific object at a time.

Question 42:

Skipped

Which SQL command determines whether a network policy is set on the account or for a specific user?

- 

## SHOW POLICY

-

## SHOW POLICIES

- 

## SHOW PARAMETERS

(Correct)

- 

## SHOW PARAMETER

- 

## SHOW NETWORK\_POLICIES

### Explanation

The SHOW PARAMETERS command determines whether a network policy is set on the account or for a specific user.

**For Account level:** SHOW PARAMETERS LIKE 'network\_policy' IN ACCOUNT;

**For User level:** SHOW PARAMETERS LIKE 'network\_policy' IN USER <username>;

**Example -** SHOW PARAMETERS LIKE 'network\_policy' IN USER john;

Question 43:

Skipped

Choose the true statements about Secure views. (Select 2)

- 

Secure views provide improved data privacy and data sharing

(Correct)

- 

Secure views allow faster access than Standard views

- 

Only non-materialized views can be defined as secure

- 

Both non-materialized and materialized views can be defined as secure

**(Correct)**

- 

**Only materialized views can be defined as secure**

#### **Explanation**

Both non-materialized and materialized views can be defined as secure. Secure views have advantages over standard views, including improved data privacy and data sharing; however, they also have some performance impacts to take into consideration.

Question 44:

**Skipped**

In a federated Snowflake environment, can a user still log into Snowflake using their Snowflake credentials?

- 

**Yes**

**(Correct)**

- 

**No**

#### **Explanation**

Users can still log into Snowflake using their Snowflake credentials.

Question 45:

**Skipped**

Which of these roles is granted the MANAGE GRANTS security privilege to be able to modify any grant globally, including revoking it?

- 

**ORGADMIN**

- 

**SYSADMIN**

- 

**ACCOUNTADMIN**

-

## USERADMIN

- 

## SECURITYADMIN

(Correct)

### Explanation

SECURITYADMIN role can manage any object grant globally, as well as create, monitor, and manage users and roles. More specifically, this role:

- Is granted the MANAGE GRANTS security privilege to be able to modify any grant, including revoking it.
- Inherits the privileges of the USERADMIN role via the system role hierarchy (i.e. USERADMIN role is granted to SECURITYADMIN).

Question 46:

Skipped

Suppose we resize a warehouse to a larger size while it is executing SQL statements. In that case, the already running SQL statements will finish faster. (True / False)

- 

FALSE

(Correct)

- 

TRUE

### Explanation

Resizing a warehouse doesn't have any impact on statements that are currently being executed by the warehouse. When resizing to a larger size, the new compute resources, once fully provisioned, are used only to execute statements that are already in the warehouse queue, as well as all future statements submitted to the warehouse.

Question 47:

Skipped

How many network policies can be activated for a user at a time?

- 

31

- 

1

(Correct)

- 100
- 99
- 16

#### Explanation

**Only a single network policy can be activated for each user at a time;** however, different network policies can be activated for different users for granular control. Associating a network policy with a user automatically removes the currently-associated network policy (if any)

Question 48:

Skipped

Which command will help you to view the current permissions on a Schema?

- SHOW ALL PREIVILEGE ON SCHEMA <database\_name>. <schema\_name>;
- SHOW GRANTS OF SCHEMA <database\_name>. <schema\_name>;
- SHOW ROLES ON SCHEMA <database\_name>. <schema\_name>;
- SHOW GRANTS ON SCHEMA <database\_name>. <schema\_name>;

(Correct)

#### Explanation

To view the current set of privileges granted on an object, you can execute the **SHOW GRANTS** command. To view the current permissions on a schema, execute the following command: **SHOW GRANTS ON SCHEMA <database\_name>. <schema\_name>;**

Question 49:

## Skipped

Network policies allow restricting access to your account based on\_\_\_\_ (Select all that apply)

- 

User Operating System Type (example - MAC, Windows)

- 

CIDR Notation based IP ranges

(Correct)

- 

IP address

(Correct)

## Explanation

Network policies allow restricting access to your account based on user IP address. Effectively, a network policy enables you to create an IP allowed list, as well as an IP blocked list, if desired.

Question 50:

## Skipped

Materialized views can improve the performance of queries that use the same subquery results repeatedly. (True/False)

- 

TRUE

(Correct)

- 

FALSE

## Explanation

Materialized views are designed to improve query performance for workloads composed of common, repeated query patterns. However, materializing intermediate results incur additional costs. As such, before creating any materialized views, you should consider whether the costs are offset by the savings from re-using these results frequently enough.

Question 51:

## Skipped

Which of these roles can configure a network policy? (Select 2)

- **SYSADMIN**
- **SECURITYADMIN**  
**(Correct)**
- **USERADMIN**
- **PUBLIC**
- **ACCOUNTADMIN**

**(Correct)**

#### Explanation

Only **security administrators** (i.e., users with the **SECURITYADMIN** role) or higher or a role with the global CREATE NETWORK POLICY privilege can create network policies.

Question 52:

**Skipped**

What is the best practice after creating a custom role in a Snowflake account?

- **Grant PUBLIC to the role so all database objects owned by PUBLIC will be available to the new role.**
- **Grant the custom role to the SYSADMIN role so administrators can manage all objects in the account.**  
**(Correct)**
-

**Grant ownership of the role to itself so a member of the role can control access to the role.**

- 

**Grant the role to the USERADMIN role so users can be added to the role.**

### **Explanation**

The custom role gets mainly created for specific access to specific objects. **As per best practice, grant ownership to SYSADMIN or a child role under SYSADMIN so that SYSADMIN can control access to the objects.**

Question 53:

### **Skipped**

There are two modes to set up a multi-cluster warehouse. Select those from the given choices.

- 

**Minimized mode**

- 

**Maximized mode**

**(Correct)**

- 

**Maximum mode**

- 

**Auto-scaling mode**

**(Correct)**

- 

**Minimum mode**

### **Explanation**

There are two ways to set up a multi-cluster warehouse: in maximized mode, or auto-scaling mode. With maximized mode, you simply set your minimum equal to your maximum, and those values are something greater than one.

Question 54:

### **Skipped**

All files stored in internal stages for data loading and unloading are automatically encrypted using AES-256 strong encryption. (True/False)

- 

**FALSE**

- 

**TRUE**

**(Correct)**

#### **Explanation**

**All files stored in internal stages for data loading and unloading are automatically encrypted using AES-256 strong encryption.**

Question 55:

**Skipped**

**Snowflake Query history page allows you to view the details of all the queries executed in the last 31 days. (True/False)**

- 

**TRUE**

- 

**FALSE**

**(Correct)**

#### **Explanation**

Snowflake Query history page allows you to view the details of all the queries executed in the last 14 days. You can query the Query\_History view in Snowflake's Account Usage schema for older queries.

Question 56:

**Skipped**

**Which of these configurations will set up a warehouse in auto-scale mode?**

- 

**Minimum Clusters = 6 and Maximum Clusters = 6**

- 

**Minimum Clusters = 2 and Maximum Clusters = 6**

**(Correct)**

### **Explanation**

Auto-scale mode is enabled by specifying different values for maximum and minimum number of clusters.

Question 57:

**Skipped**

**Multi-cluster warehouses are beneficial in improving the performance of slow-running queries or data loading. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

Multi-cluster warehouses are best utilized for scaling resources to improve concurrency for users/queries. They are not as beneficial for improving the performance of slow-running queries or data loading. For these types of operations, resizing the warehouse provides more benefits.

Question 58:

**Skipped**

**If you want to create a warehouse that remains in a suspended state initially, which property do you need to set for that warehouse?**

- 

**INITIALLY\_SUSPENDED = TRUE**

**(Correct)**

- 

**AUTO\_SUSPEND = 0**

- 

**AUTO\_RESUME = TRUE**

-

**AUTO\_SUSPEND = TRUE**

- 
- 

**AUTO\_RESUME = FALSE**

**Explanation**

INITIALLY\_SUSPENDED = TRUE | FALSE Specifies whether the warehouse is created initially in the 'Suspended' state. The valid values are TRUE and FALSE. TRUE: The warehouse is created, but suspended. FALSE: The warehouse starts running after it is created. Default is FALSE

Question 59:

**Skipped**

**What will happen if a policy is assigned to a user who is already signed in?**

- 
- 

**There will be no interruption until the user logoffs and signs in again.**

- 
- 

**The user can continue running the SQL queries in the currently opened session.**

- 
- 

**The user can't do anything else until signed in and signed back in again.**

**(Correct)**

**Explanation**

If a policy is assigned to a user who already signed in, they **can't do anything else until they sign and signed back in again** to make use of the new policy

Question 60:

**Skipped**

**ACCOUNTADMIN role should not be used to create objects in Snowflake? (True/False)**

- 
- 

**FALSE**

- 
- 

**TRUE**

**(Correct)**

### Explanation

The ACCOUNTADMIN role is intended for performing initial setup tasks in the system and managing account-level objects and tasks on a day-to-day basis. It should not be used to create objects in your account unless you absolutely need these objects to have the highest level of secure access.

Question 61:

### Skipped

A user cannot view the result set from a query that another user executed except for the ACCOUNTADMIN role. (True / False)

- **TRUE**
- **FALSE**

**(Correct)**

### Explanation

A user cannot view the result set from a query that another user executed. This behavior is intentional. For security reasons, only the user who executed a query can access the query results. This behavior is not connected to the Snowflake access control model for objects. Even a user with the ACCOUNTADMIN role cannot view the results for a query run by another user.

Question 62:

### Skipped

The major benefits of defining Clustering keys: (Select 2)

- **To help in organizing small tables (<1 GB)**
- **To help in faster data sharing**
- **To help improve query performance**

**(Correct)**

To help optimize table maintenance

(Correct)

#### Explanation

Defining clustering keys for very large tables (in the multi-terabyte range) helps optimize table maintenance and query performance. Small tables are not a good candidate for clustering.

Question 63:

Skipped

You can create a user-level network policy using \_\_\_\_\_

- - **Classic Web Interface**
  - 
  - **Snowsight**
  - 
  - **SQL**
- (Correct)
- 
- Only Snowflake Support can create the Account level Network Policy**

#### Explanation

You can set a policy at the user level, but this can only be done through the SQL commands. To activate a network policy for an individual user, set the NETWORK\_POLICY parameter for the user using ALTER USER.

Question 64:

Skipped

What is the best way to analyze the optimum warehouse size?

- 
- Execute relatively homogeneous queries (size, complexity, data sets, etc.) on the same warehouse**
- (Correct)
-

**Execute queries of widely-varying size and/or complexity on the same warehouse**

### Explanation

To achieve the best results, try to execute relatively homogeneous queries (size, complexity, data sets, etc.) on the same warehouse; executing queries of widely-varying size and/or complexity on the same warehouse makes it more difficult to analyze warehouse load, which can make it more difficult to select the best size to match the size, composition, and number of queries in your workload.

Question 65:

**Skipped**

**What key insights can we get from the Explain plan in Snowflake? (Select 3)**

- 

**Partition Pruning**

**(Correct)**

- 

**Estimated Query Time**

- 

**Join Ordering**

**(Correct)**

- 

**Exact Query Time**

- 

**Join Types**

**(Correct)**

### Explanation

The key insights that the explain plan gives us in its results output are information on partition pruning, join ordering, and join types. The explain plan is a useful tool for determining the efficiency of your query. It's a command that compiles your query to figure out all the steps Snowflake would have to work through if it were actually to run the query.

Question 66:

**Skipped**

**When the Virtual Warehouse data cache gets filled up, in which fashion does the data get flushed out from the data cache?**

- 

**LEAST-RECENTLY USED (LRU)**

**(Correct)**

- 

**First In First Out (FIFO)**

- 

**MOST-RECENTLY USED (MRU)**

- 

**Last In Last Out (LILO)**

### **Explanation**

If the cache does fill up, it is flushed out in a least-recently used fashion.

Question 67:

**Skipped**

A user's default role is

- 

**the role a user gets set to each time the user logs in to Snowflake.**

**(Correct)**

- 

**always the default PUBLIC role.**

- 

**changed each time the user logs in to Snowflake.**

- 

**the name used to log in to the Snowflake WebUI.**

### **Explanation**

**A user's default role is the role a user gets set to each time the user logs in to Snowflake.** Snowflake uses roles to control the objects (virtual warehouses, databases, tables, etc.) that users can access:

Snowflake provides a set of predefined roles, as well as a framework for defining a hierarchy of custom roles. All Snowflake users are automatically assigned the predefined PUBLIC role, which enables login to Snowflake and basic object access.

In addition to the PUBLIC role, each user can be assigned additional roles, with one of these roles designated as their default role. A user's default role determines the role used in the Snowflake sessions initiated by the user; however, this is only a default. Users can change roles within a session at any time.

Roles can be assigned at user creation or afterward.

Question 68:

**Skipped**

**Snowflake blocks certain IPs by default to ensure that customer is getting the highest level of Network security. (TRUE / FALSE)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

**Explanation**

**By default, Snowflake allows users to connect to the service from any computer or device IP address.** A security administrator (or higher) can create a network policy to allow or deny access to a single IP address or a list of addresses.

Question 69:

**Skipped**

**Snowflake supports multi-factor authentication (i.e., MFA) to provide increased login security for users connecting to Snowflake. Which role is strongly recommended for using MFA?**

- 

**ACCOUNTADMIN**

**(Correct)**

- **USERADMIN**
- **SYSADMIN**
- **USERADMIN**

#### Explanation

Snowflake strongly recommends that all users with the **ACCOUNTADMIN** role be required to use MFA.

Question 70:

**Skipped**

**Which of these system-defined roles encapsulates the SYSADMIN and SECURITYADMIN roles?**

- **ORGADMIN**
- **SYSADMIN**
- **SECURITYADMIN**
- **USERADMIN**
- **ACCOUNTADMIN**

**(Correct)**

#### Explanation

**ACCOUNTADMIN** role **encapsulates the SYSADMIN and SECURITYADMIN system-defined roles**. It is the top-level role in the system and should be granted only to a limited/controlled number of users in your account.

Question 71:

### Skipped

When deciding whether to suspend a warehouse or leave it running, what should you consider?

- Consider suspending the warehouse if the warehouse is large and there are no active queries.
- Consider the trade-off between saving credits by suspending the warehouse versus the operational cost of resuming the warehouse when needed.
- Consider the trade-off between saving credits by suspending the warehouse versus maintaining the cache of data from the previous queries to help with performance.

(Correct)

### Explanation

Consider the trade-off between saving credits by suspending a warehouse versus maintaining the cache of data from previous queries to help with performance.

Question 72:

### Skipped

Which of the Snowflake shared view can be used to query the Snowflake Query History? (Select 1)

- QUERY\_HISTORY view in INFORMATION\_SCHEMA
- QUERY\_HISTORY\_VIEW view in ACCOUNT\_USAGE
- QUERY\_HISTORY\_VIEW view in INFORMATION\_USAGE
- QUERY\_HISTORY view in ACCOUNT\_USAGE

(Correct)

### Explanation

**QUERY\_HISTORY view in ACCOUNT\_USAGE view** can be used to query Snowflake query history by various dimensions (time range, session, user, warehouse, etc.) within the last 365 days (1 year).

Question 73:

Skipped

At what frequency does Snowflake rotate the object keys?

- 
- 60 Days**
- 
- 1 Year**
- 
- 16 Days**
- 
- 30 Days**

(Correct)

### Explanation

All Snowflake-managed keys are automatically rotated by Snowflake when they are more than 30 days old. Active keys are retired, and new keys are created. When Snowflake determines the retired key is no longer needed, the key is automatically destroyed. When active, a key is used to encrypt data and is available for usage by the customer. When retired, the key is used solely to decrypt data and is only available for accessing the data.

Question 74:

Skipped

You can create an account level network policy using \_\_\_\_ (Select all that apply)

- 
- Only Snowflake Support can create the Account level Network Policy**
- 
- Snowsight**

**(Correct)**

- 

**SQL**

**(Correct)**

- 

**Classic Web Interface**

**(Correct)**

### **Explanation**

Only security administrators (i.e., users with the SECURITYADMIN role) or higher or a role with the global CREATE NETWORK POLICY privilege can create network policies **using Snowsight, Classic Web Interface, and SQL**.

Question 75:

**Skipped**

A row access policy contains an expression that can specify Snowflake database objects (e.g., table or view) and use functions to determine which rows should be visible in a given context. Which of these functions are used in determining row access policy? (Select 2)

- 

**Conversion Functions**

- 

**Metadata Functions**

- 

**Conditional Expression Functions**

**(Correct)**

- 

**Context Functions**

**(Correct)**

### **Explanation**

A row access policy uses **Conditional Expression Functions and Context Functions** to determine which rows should be visible in a given context. Context

Functions such as CURRENT\_USER(), CURRENT\_ROLE(), and CURRENT\_ACCOUNT(), which act as dynamic filters and are commonly used with secure views to limit row access in a table.

Question 76:

Skipped

If you create a user with MUST\_CHANGE\_PASSWORD = TRUE in a Snowflake federated environment, will that user be forced to change the password while logging through IdP the first time?

- 

Yes

- 

No

(Correct)

Explanation

The MUST\_CHANGE\_PASSWORD user property does not apply for federated authentication and should not be used. In particular, if you choose not to maintain passwords in Snowflake for users, ensure this property is set to FALSE for these users.

Question 77:

Skipped

Choose the false statements. (Select 2)

- 

Results are stored in SSD in Virtual Warehouse in case of Local Disk Cache.

- 

Group and Execute similar queries on the different virtual warehouses to maximize local disk cache reuse, for performance and cost optimization.

(Correct)

- 

If Virtual Warehouse is suspended, then results in Local Disk Cache will be lost

-

**Group and Execute similar queries on the same virtual warehouse to maximize local disk cache reuse, for performance and cost optimization.**

- 

**Results are stored in the Cloud Storage layer in the case of Local Disk Cache.**

**(Correct)**

### **Explanation**

As a best practice, Group and Execute similar queries on the same virtual warehouse to maximize local disk cache reuse for performance and cost optimization. The results get stored in the SSD of Virtual Warehouse. So, if the Virtual Warehouse gets suspended, then results get lost.

Question 78:

**Skipped**

**If you drop or disable a user in Snowflake in an Okta IdP federated environment, the user can still access Snowflake login through Okta. (True/False)**

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### **Explanation**

**Users who are dropped or disabled in Snowflake are still able to log into their Okta accounts, but they will receive an error message when they attempt to connect to Snowflake. You must recreate or enable the user before they can log in.**

Question 79:

**Skipped**

**Will these queries be considered as same to get the benefit of the Query Result cache?**

**Query 1 : SELECT \* FROM t1; Query 2 : select \* FROM t1;**

- 

**NO**

**(Correct)**

-

**YES**

**Explanation**

These queries will not be considered as same to get the benefit of the Query Result cache. First Query has 'SELECT' is upper case, and the Second query has 'select' in lower case.

Question 80:

**Skipped**

**John has a table EMPLOYEE\_DATA, and he wants to create another table EMPLOYEE\_DATA\_OTHER, which should be the same as EMPLOYEE\_DATA table with the same data. What is the best option for John?**

- 

**CREATE SHARE EMPLOYEE\_DATA;**

- 

**Clone the table with same data with SQL command as follows - CREATE TABLE EMPLOYEE\_DATA\_OTHER CLONE EMPLOYEE\_DATA;**

**(Correct)**

- 

**Create the table with same data with SQL command as follows - CREATE TABLE EMPLOYEE\_DATA\_OTHER AS SELECT \* FROM EMPLOYEE\_DATA;**

- 

**Create the table with LIKE SQL command as follows - CREATE TABLE EMPLOYEE\_DATA\_OTHER LIKE EMPLOYEE\_DATA;**

**Explanation**

The best option is to Clone the table as EMPLOYEE\_DATA and EMPLOYEE\_DATA\_OTHER have the same structure and data. It will help save the storage cost. LIKE command only creates the empty table. CREATE TABLE ... AS SELECT (also referred to as CTAS) Creates a new table populated with the data returned by a query but consumes additional storage.

Question 81:

**Skipped**

**Which are the required parameters for creating a Network Policy? (Select 2)**

- 

**Policy Name**

**(Correct)**

- 

### **Allowed IP Addresses**

**(Correct)**

- 

### **Comment**

- 

### **Blocked IP Addresses**

#### **Explanation**

**Policy Name and Allowed IP Addresses are mandatory for creating a Network Policy.** Blocked IP Addresses and Comment Parameters are optional.

**Please Note** - If you provide both Allowed IP Addresses and Blocked IP Addresses, Snowflake applies the Blocked List first.

Question 82:

**Skipped**

**Which of these objects are not replicated? (Select 2)**

- 

### **Transient Tables**

- 

### **Temporary Tables**

**(Correct)**

- 

### **External Tables**

**(Correct)**

- 

### **Permanent Tables**

- 

## Views

### Explanation

**Temporary Tables, External Tables, Stages, Temporary Stages, Streams, and Tasks do not get replicated.**

Question 83:

### Skipped

In a Snowflake federated environment, **Snowflake serves as the Identity provider (IdP).** (True/False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

### Explanation

In a Snowflake federated environment, **Snowflake serves as the Service Provider (SP).** The external, independent entity like Okta serves as the Identify Provider (IdP)

Question 84:

### Skipped

An account-level resource monitor overrides the resource monitor assignment for individual warehouses. (True/False)

- 

**TRUE**

- 

**FALSE**

**(Correct)**

### Explanation

An account-level resource monitor does not override resource monitor assignments for individual warehouses. If either the account resource monitor or the warehouse resource monitor reaches its defined threshold and a suspend action has been defined, the warehouse is suspended.

Question 85:

## Skipped

Suppose you have an auto-scaling mode setup with a Standard policy. In what situation does Snowflake spin up an additional cluster?

- 

**Only if the system estimates there's enough query load to keep the cluster busy for at least 6 minutes.**

- 

**The first cluster starts immediately when either a query is queued or the system detects that there's one more query than the currently-running clusters can execute.**

**(Correct)**

## Explanation

In Standard Scaling policy, the first cluster starts immediately when either a query is queued, or the system detects that there's one more query than the currently-running clusters can execute. Each successive cluster waits to start 20 seconds after the prior one has started. For example, if your warehouse is configured with ten max clusters, it can take 200+ seconds to start all 10 clusters.

Question 86:

## Skipped

During Warehouse provisioning, if any of the compute resources fail to provision, then Snowflake kills the entire warehouse provisioning and tries to provision a new warehouse of the same requested size. (True/False)

- 

**FALSE**

**(Correct)**

- 

**TRUE**

## Explanation

Snowflake does not begin executing SQL statements submitted to a warehouse until all of the compute resources for the warehouse are successfully provisioned, unless any of the resources fail to provision: If any of the compute resources for the warehouse fail to provision during start-up, Snowflake attempts to repair the failed resources. During the repair process, the warehouse starts processing SQL statements once 50% or more of the requested compute resources are successfully provisioned.

Question 87:

**Skipped**

What are the security layers that Snowflake takes care of? (Select 4)

- 

**Data Protection**

**(Correct)**

- 

**Access**

**(Correct)**

- 

**Infrastructure**

- 

**Authorization**

**(Correct)**

- 

**Authentication**

**(Correct)**

**Explanation**

**Infrastructure Security is managed by the cloud provider.**

Question 88:

**Skipped**

Snowflake's approach to access control combines aspects from which of the following models? (Select 2)

- 

**Rule-based access control (RBAC)**

- 

**Role-based Access Control (RBAC)**

**(Correct)**

- 

### **Discretionary Access Control (DAC)**

**(Correct)**

- 

### **Mandatory Access Control (MAC)**

#### **Explanation**

Snowflake's approach to access control combines aspects from both of the following models:

**Discretionary Access Control (DAC):** Each object has an owner, who can in turn grant access to that object.

**Role-based Access Control (RBAC):** Access privileges are assigned to roles, which are in turn assigned to users.

Question 89:

**Skipped**

**How can we turn off the query result cache?**

- 

**Setting the parameter USE\_QUERY\_CACHED to FALSE**

- 

**Query result cache can be turned off.**

- 

**Setting the parameter USE\_CACHED\_RESULT to FALSE**

**(Correct)**

- 

**Setting the parameter USE\_CACHED\_INFO to FALSE**

#### **Explanation**

We can turn off the query result cache by setting the parameter USE\_CACHED\_RESULT to FALSE. Though the only reason we would really want to do this is if we are doing performance testing.

Question 90:

Skipped

Which one is not the Snowflake System-Defined role?



**DATABASEADMIN**

(Correct)



**SYSADMIN**



**USERADMIN**



**ORGADMIN**



**SECURITYADMIN**



**ACCOUNTADMIN**

### Explanation

The following are the Snowflake System-Defined roles:

ORGADMIN

ACCOUNTADMIN

SECURITYADMIN

USERADMIN

SYSADMIN

PUBLIC.

System-defined roles cannot be dropped. In addition, the privileges granted to these roles by Snowflake cannot be revoked.

Question 91:

Skipped

Which type of object key is only used for decryption?



**Retired Key**

(Correct)



**Active key**



**None of these**



**Destroyed key**

**Explanation**

Active Key is used for both encryption and decryption. **The retired Key is used for decryption only.** The destroyed Key is no longer used.

Question 92:

Skipped

You have a dashboard that connects to Snowflake via JDBC. The dashboard is refreshed hundreds of times per day. The data is very stable, only changing once or twice per day. The query run by the dashboard user never changes. How will Snowflake manage changing and non-changing data? Mark all true statements.



**Snowflake will re-use data from the Results Cache as long as it is still the most up-to-date data available.**

(Correct)



**Snowflake will spin up a warehouse only if the underlying data has changed.**

(Correct)

- 

Snowflake will show the most up-to-date data each time the dashboard is refreshed.

(Correct)

- 

Snowflake will compile result cache data from all user results, so no warehouse is needed.

#### Explanation

Until data has not changed and the query is the same - Snowflake reuses the data from the cache. Please note, Each time the persisted result for a query is reused, Snowflake resets the 24-hour retention period for the result up to a maximum of 31 days from the date and time that the query was first executed. After 31 days, the result is purged, and the next time the query is submitted, a new result is generated and persisted.

Question 93:

Skipped

The BI group is complaining about their queries taking too long to run. Checking the virtual warehouse information shows that the queued time is pretty high. What is the best way to fix this issue?

- 

Set the STATEMENT\_QUEUED\_TIMEOUT\_IN\_SECONDS parameter to a low value to cancel those queries if they get in the queue

- 

Increase the size of the virtual warehouse

- 

Increase the virtual warehouse MAX\_CLUSTER\_COUNT property

(Correct)

- 

Determine which users have the high priority queries and set the other users

-

## Provide a virtual warehouse for every user in the group

### Explanation

Queuing can be solved by SCALE-OUT (provision new clusters), i.e., increasing MAX\_CLUSTER\_COUNT helps in additional cluster provisioning to handle the concurrent workloads.

Question 94:

### Skipped

David ran a query that took around 30 mins to complete. He referred to the Query profiler and found the 'Bytes spilled to local storage' has a big number. What could be the issue?

- 

**Warehouse size has no impact on Bytes spilling.**

- 

**David is using a comparatively smaller warehouse.**

**(Correct)**

- 

**David should contact Snowflake Personnel.**

- 

**David is using a comparatively larger warehouse.**

### Explanation

If a node has insufficient memory to complete its portion of a query, it will "spill" to local SSD storage. This can negatively impact performance but is sometimes acceptable. If a node has insufficient local SSD storage to complete its portion of a query, it will "spill" to remote cloud storage. This is almost always very bad for performance. The solution, in either case, is to simplify the SQL query or increase the warehouse size.

Question 95:

### Skipped

Which of these roles is dedicated to user and role management only?

- 

**ORGADMIN**

-

## USERADMIN

(Correct)

- 

SYSADMIN

- 

ACCOUNTADMIN

- 

SECURITYADMIN

### Explanation

**USERADMIN** role is dedicated to user and role management only. More specifically, this role:

- Is granted the CREATE USER and CREATE ROLE security privileges.
- Can create users and roles in the account.

This role can also manage users and roles that it owns. Only the role with the OWNERSHIP privilege on an object (i.e. user or role), or a higher role, can modify the object properties.

Question 96:

Skipped

How can you create a "Super-User" or "Super-Role" in Snowflake who can bypass all the authorization checks?

- 

ACCOUNTADMIN role is same as Super-Role

- 

Contact Snowflake personnel to create a Super-Role or Super-User for your account

- 

CREATE ROLE SUPER\_ROLE;

- 

There is no concept of SUPER-ROLE or SUPER-USER in Snowflake

**(Correct)**

### Explanation

There is no concept of a “super-user” or “super-role” in Snowflake that can bypass authorization checks. All-access requires appropriate access privileges.

Question 97:

**Skipped**

The warehouse performance can be evaluated by querying the

- Information Schema QUERY\_HISTORY view
  - Account Usage QUERY\_HISTORY view
- (Correct)**
- Information Schema LOAD\_HISTORY view
  - Account Usage LOAD\_HISTORY view

### Explanation

The warehouse performance can be evaluated by querying the Account Usage QUERY\_HISTORY view.

Question 98:

**Skipped**

What would happen if we suspend the warehouse while it is executing the SQL statement?

- All compute resources of the warehouse will be up until the statement is complete.
- All the compute resources of the warehouse will be shut down immediately, and the running statement will be canceled.
-

**Only idle compute resources of the warehouse will be shut down, allowing any compute resources executing statements to continue until the statement is complete.**

**(Correct)**

- 

**When trying to suspend the warehouse, we will get an error while the same warehouse is executing SQL statements.**

### **Explanation**

When we suspend a warehouse, Snowflake immediately shuts down all idle compute resources for the warehouse. However, it allows any compute resources executing statements to continue until the statements are complete. At this time, the resources are shut down, and the warehouse status changes to "Suspended". Compute resources waiting to shut down are considered to be in "quiesce" mode.

Question 99:

**Skipped**

Dynamic Data Masking is supported by (Select all that apply)

- 

**Business Critical**

**(Correct)**

- 

**Standard Edition**

- 

**Enterprise Edition**

**(Correct)**

- 

**VPS**

**(Correct)**

### **Explanation**

**Dynamic Data Masking features require Enterprise Edition (or higher).**

Question 100:

## Skipped

Monica ran a **SELECT** query on a large table t1. The query took longer than expected. She looked into the query profile and found that 'Bytes spilled to local storage' and 'Bytes spilled to remote storage' are very high. What advice will you give to her to improve the query performance? (Select 3)

- 

**Using a larger warehouse (effectively increasing the available memory/local disk space for the operation)**

**(Correct)**

- 

**Processing data in larger batches**

- 

**Trying to split the processing into several steps**

**(Correct)**

- 

**Increasing the number of parallel queries running in the warehouse**

- 

**Processing data in smaller batches**

**(Correct)**

## Explanation

When Snowflake warehouse cannot fit an operation in memory, it starts spilling (storing) data first to the local disk of a warehouse node and then to remote storage. In such a case, Snowflake first tries to temporarily store the data on the warehouse's local disk. As this means extra IO operations, any query that requires spilling will take longer than a similar query running on similar data that is capable to fit the operations in memory. Also, if the local disk is insufficient to fit the spilled data, Snowflake further tries to write to the remote cloud storage, which will be shown in the query profile as "Bytes spilled to remote storage".

**The spilling can't always be avoided, especially for large batches of data, but it can be decreased by:**

- Reducing the amount of data processed. For example, by trying to improve partition pruning or projecting only the columns that are needed in the output.

- Decreasing the number of parallel queries running in the warehouse.
- Trying to split the processing into several steps (for example, by replacing the CTEs with temporary tables).
- Using a larger warehouse - effectively means more memory and more local disk space.

Question 101:

Skipped

If a user is logged in to Snowflake in a federated environment and IdP times out, what happens to the user's snowflake session?

- 

**It does not affect the user's Snowflake sessions. However, to initiate any new Snowflake sessions, the user must log into the IdP again.**

**(Correct)**

- 

**The Snowflake web interface is disabled, and the prompt for IdP authentication is displayed.**

### Explanation

After a specified period of time (defined by the IdP), a user's session in the IdP automatically times out, but this does not affect their Snowflake sessions. Any Snowflake sessions that are active at the time remain open and do not require re-authentication. However, to initiate any new Snowflake sessions, the user must log into the IdP again.

Question 102:

Skipped

Which role is inherited to every other role in the account?

- 

**PUBLIC**

**(Correct)**

- 

**ACCOUNTADMIN**

- 

**SECURITYADMIN**

- 

SYSADMIN

- 

USERADMIN

#### Explanation

The PUBLIC role is Pseudo-role, which is automatically granted to every user and every role in your account. This role is typically used in cases where explicit access control is not needed, and all users are viewed as equal with regard to their access rights.

Question 103:

Skipped

How long do results remain in the Query results cache?

- 

16 hours

- 

24 hours

(Correct)

- 

31 hours

- 

1 hours

- 

12 hours

#### Explanation

Results are retained for 24 hours in Query Result Cache. Snowflake resets the 24-hour retention period for the result, up to a maximum of 31 days from the date and time that the query was first executed. After 31 days, the result is purged and the next time the query is submitted, a new result is generated and persisted.

Question 104:

Skipped

Which of these are types of Scaling policies? (Select 2)

- 

**Minimized**

- 

**Economy**

**(Correct)**

- 

**Business**

- 

**Maximized**

- 

**Standard**

**(Correct)**

### Explanation

There are two different scaling policies, one is the **Standard** policy, and one is the **Economy** policy.

Question 105: **Incorrect**

What will happen if you add 0.0.0.0/0 to BLOCKED\_IP\_LIST and your IP address to ALLOWED\_IP\_LIST of a Network policy?

- 

**You will not be able to access the Snowflake account from your IP address.**

**(Correct)**

- 

**You will be able to access the Snowflake account from your IP address**

**(Incorrect)**

### Explanation

If you provide both Allowed IP Addresses and Blocked IP Addresses, **Snowflake** applies the Blocked List first. This would block your own access.

Additionally, in order to block all IP addresses except a select list, you only need to add IP addresses to ALLOWED\_IP\_LIST. Snowflake automatically blocks all IP addresses not included in the allowed list.

All below questions in this file are select one , unless mentioned select multiple

1.Need to check Login History upto last 90 days .

1 option of Information Schema table UDF

3 options from Usage Schema

2.Need to check login history prior to 90 days.

2 options from Information schema.

2 option from Usage schema.

3.Where to check all the objects access by query in last 1 year.

4 options from Usage schema . Need exact view name.

Objects\_History , Access\_history etc.

4.Materialized view limitation

All aggregate operations not supported

Only 2 tables can be referenced

Cannot be joined with other tables

Cannot be defined with a Join

5.Shared secure view needs which access on referenced database

Usage

Reference Usage

Select

Create

6.Virtual warehouse cache stores what information

All rows of result

Columns from result query

columns accessed from query  
all columns in a micro partition

7.Query history for Web UI

14 days

8.Metadata history of load for Snowpipe

7

9.Metadata history of load for Snowpipe

64

10.When does a query does not fit into memory

Spiling to storage  
Partitions scanned = total partition  
Aggregate Node present  
Join Node present  
50% of time taken by join operation

11.In standard policy , new cluster is provisioned only when

Has 6 minutes of workload  
10 minutes  
2 min  
1 day

12.Primary to secondary database URL switch supported by which minimum edition

Standard  
Ent  
VPS

\*\*refer fail over strategy / replication

13.Materialized view available with which min edition

14.Column masking available with which min edition

15.Best recommended practice while loading from Cloud store

use external stages caching

use logical paths

load all files in one large file etc.

16.If a role1 grants ownership to role 2 on an object , what all privileges on object with role1

Role 1 and rol2 both are Owners

Role 1 has limited access

17.Network policies can be applied at (select 2)

Accout level

Database level

User level

Role level

Schema level

Table level

18. Information Schema can be used to view sharing details

True

False

19. User1 has been granted role 1 and he/she creates a schema schema1 in database db1. If we want to see all the grants associtaed with schema1 , what shall we use

Show Grants on Schema Schema1

Show Grants for Role1

Show Grants of Role1

Show Grants on Role1

20.Snowflake ML and Data Science Partner

Data Robot

Informatica

Adobe

21.Feature to select and restore historic data

Time Travel

Fail Safe

Cloning

Clustering

22.Only way to control organization of data in micropartiton

Clustering

Cloning

Search Optimization

23.Unique feature of Stored Procedure of Snowflake

Option to execuet stored Prod with caller's or creator's rights

some other options

24.MFA requires which thrid party tool/application

OKTA

Duo

other options

25.MFA needs to implemented with SSO

true

false

26. MFA can be implemented with (multi select)

Snowflake Web UI

jdbc

odbc

python

snowsql

27. Which tool is provided by snowflake to deal with connectivity issues

SnowCD

Snowpipe

SnowCLI

SNOWSQL

28.Which Consumer accounts incur costs to Provide

Consumer account

Reader Account

Full Account

29.Provider account needs to incur additional storage costs for objects shared with consumer

True

False

30.Which features allow snowflake to replace sensitive data with Random data before the data is loaded to snowflake

- Column Masking
- External Tokenization
- External Function
- Row access policy

31. External Functions

- Can return multiple rows
- Are a type of UDF
- Have separate SQL code
- Store code inside Snowflake

32. Results cache are retrieved from which layer

- Metadata
- Cloud Services
- Storage
- Compute

33. Choose below correct statement

- Show commands needs running Warehouse but Information schema doesnt need
- Show commands doesnt need running Warehouse but Information schema needs

34. Which datatype can used to query SemiStructured data (select 2)

- Variant
- XML
- Array
- Varchar

BLOB

35. Which option can be given while creating Snowpipe along with Copy Into

VALIDATION\_MODE

FORCE

Auto\_Ingest

PURGE

*Q 1. You are the owner of a table T1 which is in schema S1. The schema is in database D1. In order to grant read-only permissions of this table to a newly created role R1, you will need to... (select all that apply)*

- A. Grant ‘USAGE’ on database D1
- B. Grant ‘SELECT’ on database D1
- C. Grant ‘USAGE’ on schema S1
- D. Grant ‘SELECT’ on schema S1
- E. Grant ‘USAGE’ on table T1
- F. Grant ‘SELECT’ on table T1

**Correct Answers: A, C, and F**

To allow a role to use your database objects, you, as the owner must grant appropriate privileges on the database, schema, and objects. First up, you will need to grant USAGE privilege on the parent database and schema. USAGE privilege enables the grantee to refer to the database and schema in the query. Then, you will need to grant SELECT privilege on the table. The SQL statements are given below:

Grant usage on the database:

```
grant usage on database D1 to role R1;
```

Grant usage on the schema:

```
grant usage on schema D1.S1 to role R1;
```

Grant the ability to query an existing table:

```
grant select on table D1.S1.T1 to role R1;
```

*Q 2. Identify system-defined roles in Snowflake from the roles given below. (select four)*

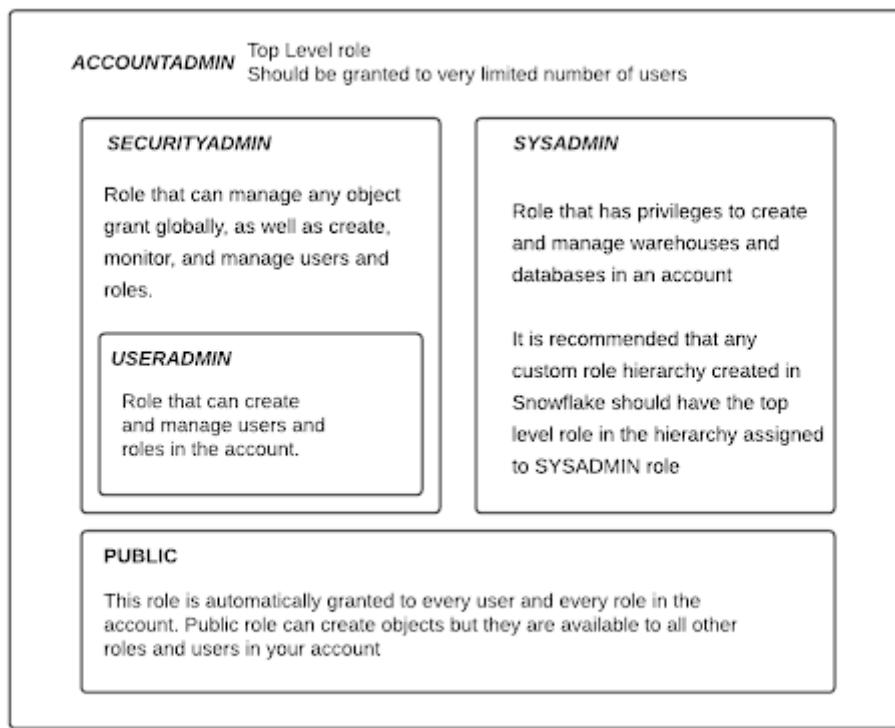
- A. AccountAdmin
- B. SysAdmin
- C. Auditor
- D. RoleAdmin

E. SecurityAdmin

F. Public

### Correct Answers: A, B, E, and F

A role in snowflake is essentially a container of privileges on objects. Roles are assigned to users to allow them to perform actions on the objects. A role can be directly assigned to the user, or a role can be assigned to a different role leading to the creation of role hierarchies. The role at a higher level in the hierarchy inherits the privileges of the role at a lower level in the hierarchy. Snowflake defined 5 system-defined roles. These roles, the key features, and the hierarchy of system-defined roles are shown in the diagram below.



account admin snowflake

## Snowflake Virtual Warehouses

*Q 3. What are the minimum and the maximum number of clusters in a multi-cluster warehouse?*

A. Minimum: 1, Maximum: 99

- B. Minimum: 1, Maximum: 100
- C. Minimum: 1, Maximum: 10
- D. Minimum: 1, Maximum: unlimited

**Correct Answer: C**

- **Answer choice-C is correct.** In a multi-cluster warehouse, the minimum number of clusters can be 1 and the maximum number of clusters can be 10.

*Q 4. When a virtual warehouse is started or resized or resumed, the minimum billing charge is 1 minute.*

- A. True
- B. False

**Correct Answer: A**

The above statement is True. When a new warehouse is started afresh, or when you resize an existing running warehouse or when a warehouse is resumed (automatically or manually), the warehouse is billed for a minimum of 1 minute's worth of usage. After the 1st minute has elapsed, all subsequent billing is charged per second.

- **Practical Info** – There is no benefit in stopping a warehouse before the first 1st minute is over because you are anyway going to pay for the first 60-seconds, as the discussion given above shows.
- **Exam Tip**– It is NOT necessary to stop a virtual warehouse to resize it. You can resize a virtual warehouse even when it is executing a query.

*Q 5. When a multi-cluster warehouse is suspended, which of the following Snowflake cache will be purged?*

- A. Metadata cache
- B. Remote disk cache (Resultset cache)
- C. Local disk cache
- D. All of the above

## **Correct Answer: C**

- When a multi-cluster warehouse (or a single cluster warehouse) is suspended, only the local disk cache is dropped. This cache is also known as warehouse cache. This is because this cache is essentially the fast SSD of the warehouse. So when the warehouse is gone, this cache is also gone with it. **Answer choice-C is correct.**
- Metadata cache contains Snowflake account object information and statistics. It is always on and it is never dropped. **Answer choice-A is incorrect.**
- Remote disk cache (resultset cache) caches exact query results including aggregations etc. and it persists the exact resultset for 24-hours after the query is executed irrespective of the state of the warehouse. **Answer choice-C is incorrect.**

## **Snowflake Overview and Architecture**

*Q 6. From the given options, select the cloud services provider supported by Snowflake architecture. (three choices)*

- A. Amazon Web Services (AWS)
- B. Google (GCP)
- C. HP Virtual Private Cloud
- D. Cloud Foundry
- E. Microsoft Azure

## **Correct Answers: A, B, and E**

Snowflake is a SaaS product. It supports three public cloud platforms given below.

1. Amazon Web Services (AWS)
2. Microsoft Azure (Azure)
3. Google Cloud Platform (GCP)

Snowflake does NOT support any other cloud platform, any private cloud deployment, or on-prem deployment. Therefore, option choices – HP Virtual Private Cloud and Cloud Foundry – are not correct.

- **Practical Info:** It is not necessary that you have to host Snowflake on the same cloud platform where your organization's other IT applications are hosted. For example, it is perfectly okay to have Snowflake hosted on Azure and your other IT applications in your organization hosted on AWS/GCP.
- **Exam Tip:** It must be noted that each Snowflake account is hosted in one cloud platform and one region, and you must choose cloud platform and region while opening an account. You cannot change it thereafter. If you want to change region and/or cloud platform at a later point, you will need to open a new Snowflake account in that region and/or cloud platform.  
Your organization may open multiple Snowflake accounts (such as Dev, Test, Prod), and they can all reside on the same cloud platform and the same region, OR you may choose to open them on different cloud platforms and/or regions.

*Q 7. What are the key considerations for choosing a specific region for your Snowflake account? (Select all that apply)*

- A. Cost
- B. Proximity to the end-users
- C. Data sovereignty
- D. Regulatory requirements

**Correct Answers: A, B, C, and D**

- **Option 1 is correct.** Snowflake offers different pricing for computing and storage depending upon the region you choose. For example, a Snowflake instance running in the Mumbai region costs \$3.30/credit, whereas the same instance running in the Frankfurt region costs \$4.00/credit. So if you are looking to save costs, you would choose a low-cost region.
- **Option 2 is correct.** The closer you are to your end-users, data transfer latency would be lower so you can transfer and access data

faster. Hence, it is always a good practice to set up your Snowflake account in the region closer to your end-users.

- **Option 3 and 4 are correct.** As is often the case in most countries, there are data sovereignty and regulatory restrictions on where the data of users is stored. For example, in the EU, certain sensitive user data is now allowed to cross the EU boundary. India also has similar restrictions enforced by the law. Hence, if you build software that will store such sensitive user data, you will need to check the regulatory requirements and choose the Snowflake region accordingly.

*Q 8. Informatica, Matillion, Azure data factory are examples of ..... partners of the Snowflake partner ecosystem.*

- A. Data Management
- B. Business Intelligence
- C. Data Integration
- D. Data Science

### **Correct Answer: C**

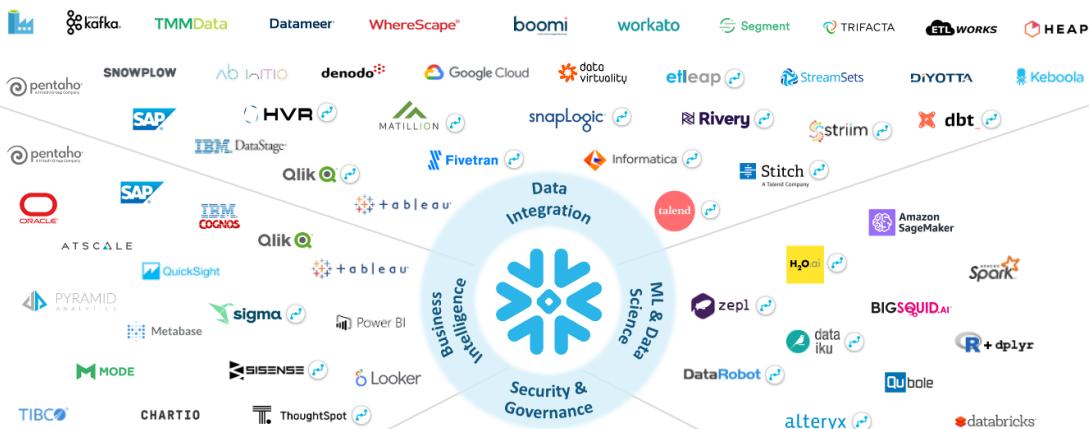
Informatica, Matillion, and Talend are **data integration** partners that provide the following functionalities:

Extract – Exporting data from the data source.

Transform – Modifying data according to business rules.

Load – Loading data into Snowflake.

The below diagram shows Snowflake partner ecosystem players divided based on the functional categories.





## ecosystem-overview

Snowflake partner eco-system (image source: Snowflake official documentation)

**Q 9. Which of the following programming languages are supported in Snowflake to write user-defined functions (UDFs)?**

- A. SQL
- B. Java
- C. JavaScript
- D. Python

### Correct Answers: A and C

User-defined functions (UDFs) can be used in Snowflake to write custom business logic which is hard to encapsulate within individual SQL statements. UDFs (together with Stored Procedures) enable database-level programming in Snowflake.

At present, Snowflake supports two programming languages to write UDFs: Javascript UDFs and SQL UDFs.

**► Practical Info-** Stored procedures in Snowflake are written in Javascript. The javascript code can issue SQL statements to process data in the Snowflake database. An example of a Snowflake stored procedure that clears a log table using SQL is given below: (this is only for your understanding, we are not expecting coding related questions in the exam)

```
create or replace procedure clear_log()
```

```
returns int not null
language javascript
as
$$
var delete_log_sql_command = "delete from user_log";
var return_value = 0;
var statement1 = snowflake.createStatement( {sqlText:
delete_log_sql_command} );
var result_set1 = statement1.execute();
return_value = 1;
return return_value;
$$
```

*Q 10. Which Snowflake edition supports transmitting data in encrypted form over the network between VPCs (virtual private cloud)?*

- A. All editions
- B. Enterprise edition and above
- C. Business Critical edition and Above
- D. All except for the Virtual Private Snowflake (VPS) edition

**Correct Answer: C**

A Snowflake account on AWS (or Azure) is implemented as a VPC. There are two ways to establish communication between your Snowflake VPC and other VPCs (e.g. your organization's VPC). One is to transmit the traffic over the public internet. Other (and safer) option is to establish an exclusive, highly secure network between your Snowflake account and your other AWS VPCs (in the same AWS region), fully protected from unauthorized access. To implement this secure channel of communication between VPCs, AWS supports a feature called AWS PrivateLink (Azure also supports a similar feature called Azure PrivateLink). Snowflake offers support for AWS PrivateLink (and Azure PrivateLink) based communication in Business Critical Edition and above.

**Q 11. More than one clustering key can co-exist in a Snowflake table**

- A. True
- B. False

**Correct Answer: B**

This statement is false. You can define at most one clustering key in a Snowflake table to organize micro-partitions. When you define a clustering key, Snowflake will reorganize the naturally clustered micro-partitions and will relocate related rows to the same micro-partition and group them according to the clustering key. This process is called **Reclustering**.

► **Practical Info** – Reclustering happens automatically once a clustering key is defined for a table. The process consumes credits. So be cognizant of the cost when you go for reclustering.

**Q 12. Which of the following statements will you use to change the warehouse for workload processing to a warehouse named 'COMPUTE\_WH\_XL'?**

- A. SET CURRENT\_WAREHOUSE = COMPUTE\_WH\_XL
- B. USE WAREHOUSE COMPUTE\_WH\_XL;
- C. USE CURRENT\_WAREHOUSE('COMPUTE\_WH\_XL');
- D. SET CURRENT\_WAREHOUSE = COMPUTE\_WH, SIZE = XL;

**Correct Answer: B**

A session context in Snowflake consists of 4 objects:

1. Role
2. Warehouse
3. Database
4. Schema

You can set appropriate session context using a set of SQL statements given

below. These statements specify the role, warehouse, database, or schema to use for the current session:

You can set appropriate session context using a set of SQL statements given below. These statements specify the role, warehouse, database, or schema to use for the current session:

Use warehouse <>warehouse-name>>

Use role <>role-name>>

Use database <>database-name>>

Use schema <>schema-name>>

So the correct answer of this question is:

Use warehouse COMPUTE\_WH\_XL;

*Q 13. In the case of a Snowflake account created on AWS, ..... is responsible for the management of Availability Zones?*

- A. Customer
- B. Snowflake
- C. Cloud Provider
- D. It is a shared responsibility

**Correct Answer: C**

An Availability Zone (AZ) is essentially a group of one or more physically separated data centers with redundant power, networking, and connectivity and located within a single cloud region. Each cloud region has multiple AZs (most have 3 or more). The AZs are connected with high-bandwidth, low-latency network infrastructure and support synchronous replication. Due to the automatic, synchronous replication and physical separation plus isolation, AZs enable your applications and databases to be highly available. When one AZ within a region fails, another one should remain active and the switch between the two zones is silent and transparent to customers. The cloud providers are responsible for the maintenance of AZ infrastructure and replication of your databases and switching to alternate AZ in case of any

failure.

## Snowflake Storage and Protection

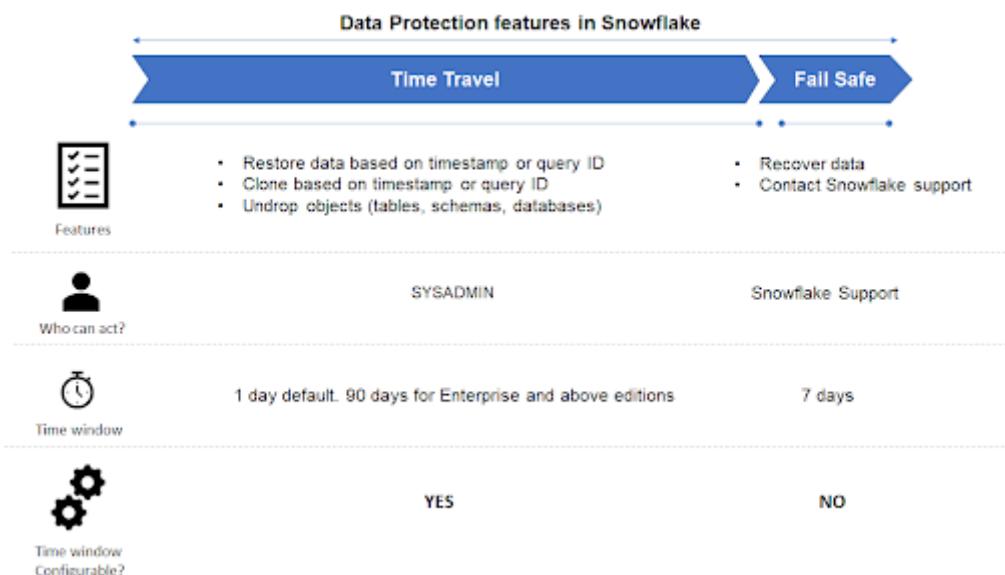
*Q 14. Once the time-travel period has expired, it is possible to request Snowflake support to retrieve historical data for a period of*

- A. Day
- B. Days
- C. Days
- D. It depends on the Snowflake edition
- E. It is user-configurable

### Correct Answer: B

After the time travel data retention period is over, you can use Snowflake's fail-safe feature to recover your data. The duration of the fail-safe period is 7 days (Answer choice 2 is correct). This is a fixed duration and cannot be changed. Only Snowflake support personnel can help recover data during the fail-safe period. The fail-safe feature is available to all customers irrespective of the Snowflake edition.

The below diagram succinctly summarizes key differences between Snowflake's two important data protection features – time-travel and fail-safe.



**Q 15. Which of the following statements are TRUE concerning a stream object in Snowflake? (select all that apply)**

- A. A stream object provides a record of DML changes (inserts, updates, deletes) made to a table at row level.
- B. A stream object can keep track of DML changes for the entire life of a table.
- C. Streams on materialized views are not supported.
- D. Streams on external tables are not supported.

**Correct Answers: A and C**

**Answer choice-A is correct.** A stream object provides change tracking over a source table. It records DML changes made to tables, (inserts, updates, and deletes) as well as metadata about each change. This is referred to as Change Data Capture (CDC), and this feature is extensively used in data warehousing scenarios to create data pipelines. Please note that the stream object itself does not store this data. It relies on the version history of source data maintained in the metadata layer.

**Answer choice-B is incorrect.** Stream object keeps track of DML changes of a source table up until the data retention period of the source table. After that, the DML changes are no longer accessible.

**Answer choice-C is correct.** Currently, Snowflake does not support creating stream objects on materialized views.

**Answer choice-D is incorrect.** Snowflake supports creating insert-only stream objects on external tables.

**Q 16. Only one stream object can be created on a source table**

- A. True
- B. False

### **Correct Answer: B**

- **The above statement is false.** You can create any number of streams on a source table. These streams can have the same or different offset positions. One example of creating multiple streams is when you want to report month-on-month changes, week-on-week changes, and day-on-day changes happening in a product inventory table of your POS database. In this case, you may create three streams on the table to record monthly, weekly and daily changes. All three streams exist independently of each other with their respective offset positions.

*Q 17. Snowflake replicates the following layer(s) across availability zones (select two)*

- A. Cloud Services Layer
- B. Warehouse Layer
- C. Storage Layer
- D. Cloud Agnostic Layer
- E. Data Sharing Layer

### **Correct Answers: A and C**

Snowflake replicates the cloud services layer (**Answer choice-A**) and the storage layer (**Answer choice-C**) across the availability zones of your cloud and region.

## **Data loading/unloading**

*Q 18. Which Snowflake edition supports Search Optimization Service to improve performance of point lookup queries?*

- A. All editions
- B. Enterprise edition and above
- C. Business Critical edition and Above
- D. All except Virtual Private Snowflake (VPS) edition

**Correct Answer: B**

Search Optimization Service in Snowflake is a background service that is designed to improve the performance of point lookup queries. A point lookup query is essentially a SELECT statement that returns only a small number of distinct rows from a large dataset. The service runs transparently from an end-user's standpoint. Only the Enterprise edition and above provide this feature.

*Q 19. Consider the XML given below. The XML file is loaded in a User Stage:*

```
<bookshelf>
<book>
<name>Famous Five</name>
<author> Roald dahl</author>
</book>
<book>
<name>Secret Seven</name>
<author> Enid Blayton</author>
</book>
<book>
<name>Mahashweta</name>
<author> Sudha Murthy</author>
</book>
</bookshelf>
```

*To remove the top-level element <bookshelf/> and load <book/> elements into separate rows in a Snowflake table, which of the following file format options will be used in the COPY INTO command?*

- A. STRIP\_TOP\_ELEMENT
- B. STRIP\_OUTER\_ELEMENT
- C. STRIP\_OUTER\_ARRAY
- D. STRIP\_FIRST\_ELEMENT

**Correct Answer: B**

Enabling the STRIP\_OUTER\_ELEMENT file format option for the COPY INTO <table> command to remove the outer element and load the <book/> records into separate table rows:

```
copy into <table> from @~/<file>.xml  
file_format = (type = 'XML' strip_outer_element = true);
```

► **Exam Tip:** If the semi-structured file format is JSON (instead of XML), you will need to enable STRIP\_OUTER\_ARRAY file format option for the COPY INTO command.

## Performance and Tuning

*Q 20. History' tab in the Snowflake Web Interface supports viewing of user queries submitted in the past..... day(s)*

- A. 24 hours or 1 day
- B. 7 days
- C. 14 days
- D. 90 days

### Correct Answer: C

'Query History' view in Snowflake Web UI stores queries executed and query profile information for **all queries, executed by all users, using any interface (e.g., Web UI, SnowSQL, Business Application... or any other interface) for a period of 14 days**

► **Exam Tip:** In addition to storing queries and profiler, it also allows users to download the resultset of queries within 24 hours. However, the most important point here is: You cannot view other people's query results if your role and the role of the user who executed the query are different.

*Q 21. What is the optimal file sizing recommendation for loading data into Snowflake?*

- A. 10 MB to 100 MB, uncompressed

- B. 100 MB to 250 MB, compressed
- C. 100 MB to 250 MB, uncompressed
- D. 10 GB to 100 GB, compressed

**Correct Answer: B**

Snowflake recommends that to optimize the number of parallel operations for a load, data files should **roughly be 100-250 MB in size, compressed**. In case input data files are smaller than this, it is better to aggregate multiple input files. On the other hand, if the input data files are larger than this, it is better to split them into smaller files to match this guidance.

Snowflake also recommends against loading very large files (e.g., 100 GB+ sizes) without splitting them as splitting of files will take advantage of parallelism.

► **Practical Info** – The recommendation applies to both modes of data loading – bulk load using SnowSQL (COPY INTO) and continuous/micro-batch using Snowpipe.

*Q 22. Selecting the right clustering key can dramatically improve query performance. Following are the good yardsticks to choose an appropriate clustering key:*

- A. Use columns that are most frequently found in WHERE clause
- B. A date column is usually not a good candidate to be in the clustering key
- C. A join column can be considered as clustering key
- D. Using expressions in the clustering key may help reduce cardinality.

**Correct Answers: A, C, and D**

**Answer choice-A is correct.** Choosing the clustering keys from dimensions (i.e. columns in the WHERE clause) used in often repeating queries will significantly improve the performance of these queries.

**Answer choice-B is incorrect.** If you have queries that filter the data based on date criteria, you can use the date column in the clustering key. However, you will need to check the cardinality of the date column before using it in the clustering key.

**Answer choice-C is correct.** Columns used in the join clauses with other tables are worthy candidates to be in the clustering key.

**Answer choice-D is correct.** In case you would like to use a particular column in a clustering key but the column has very high cardinality, using expressions may help reduce cardinality. For example, you can use SUBSTR expression to get a relevant portion of a string column for use in the clustering key.

*Q 23. Selecting the right clustering key can dramatically improve query performance. Following are the good yardsticks to choose an appropriate clustering key:*

- A. Use columns that are most frequently found in WHERE clause
- B. A date column is usually not a good candidate to be in the clustering key
- C. A join column can be considered as clustering key
- D. Using expressions in the clustering key may help reduce cardinality.

**Correct Answers: A, C, and D**

**Answer choice-A is correct.** Choosing the clustering keys from dimensions (i.e. columns in the WHERE clause) used in often repeating queries will significantly improve the performance of these queries.

**Answer choice-B is incorrect.** If you have queries that filter the data based on date criteria, you can use the date column in the clustering key. However, you will need to check the cardinality of the date column before using it in the clustering key.

**Answer choice-C is correct.** Columns used in the join clauses with other tables are worthy candidates to be in the clustering key.

**Answer choice-D is correct.** In case you would like to use a particular column in a clustering key but the column has very high cardinality, using expressions may help reduce cardinality. For example, you can use SUBSTR expression to get a relevant portion of a string column for use in the clustering key.

*Q 24. Which cache does the query optimizer refer to first?*

## *Second? And Third?*

- A. First: warehouse cache. Second: results cache. Third: metadata cache
- B. First: warehouse Cache. Second: metadata Cache. Third: results cache
- C. First: metadata Cache. Second: results Cache. Third: warehouse cache
- D. First: metadata Cache. Second: warehouse Cache. Third: results cache

### **Correct Answer: C**

The query optimizer in the Snowflake cloud services layer examines the metadata cache first. Thereafter it refers to the results cache next and finally the warehouse cache.

## **Semi-Structured data**

*Q 25. Which of the following data types are supported in Snowflake for semi-structured data? (select all that apply)*

- A. VARIANT
- B. OBJECT
- C. TEXT
- D. CLOB
- E. ARRAY

### **Correct Answers: A, B, and E**

Snowflake supports three data types for processing semi-structured data.

**VARIANT** – The basic (and most used) data type for storing semi-structured data.

**OBJECT** – Used to store key-value pairs, where the key is a non-empty string, and value is semi-structured data in VARIANT data type.

**ARRAY** – Used to store arrays of semi-structured data where array index is non-negative integer and value is semi-structured data in VARIANT data type.

When are you planning to take SnowFlake SnowPro Core Exam? Do share your feedback in comments and help other learners with real exam experience.

Whizlabs [practice tests](#) are based on real exam pattern; we have received good feedback from exam takers.

### **Reference links:**

- <https://www.youtube.com/channel/UCs10x-muRrTQMJ4Ya-fmIlw>
- <https://docs.snowflake.com/en/>

## Domain : Snowflake Overview and Architecture

*Q26 : Following are the key reasons that make Snowflake a “True SaaS” product.*

- A.** Snowflake completely runs on cloud infrastructure and is accessible over the internet
- B.** There's no hardware or software to install, configure, maintain or upgrade at the user's end
- C.** Snowflake supports key properties of as-a-service offering such as on-demand and instance elasticity, unlimited scalability of compute and storage and pay-per-use
- D.** All of the above

**Correct Answer: D**

### **Explanation**

Snowflake is fully managed, cloud-native, on-demand software offered as a service (SaaS – software as a service) over the internet. It is built from the ground up for the cloud. It is fully hosted on public cloud infrastructure and uses cloud resources for compute and storage. At the user's end, there's no hardware or software installation required. There are no patches or upgrades to be done by the user. New releases are rolled out transparently and don't require user intervention. The software is available on-demand over the internet. You can log on to [snowflake.com](http://snowflake.com), open an account, and you are ready

to use Snowflake in your organization.

Snowflake offers unlimited and instant scalability of compute and storage resources. You can scale up and down on-demand based on your need. You need to only pay for the Snowflake resources that you use. Owing to these unique design elements, Snowflake is currently the most popular cloud data warehouse solution for enterprises worldwide.

**Further Reading:** [Key Concepts & Architecture – Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q27: Following programming languages/drivers, support is available out-of-the-box (natively) in Snowflake (select all that apply)*

- A.** JDBC/ODBC driver
- B.** Google Flutter
- C.** Go, Node.js driver
- D.** Scala connector
- E.** Python connector
- F.** Spark connector

**Correct Answers:** **A, C, E and F**

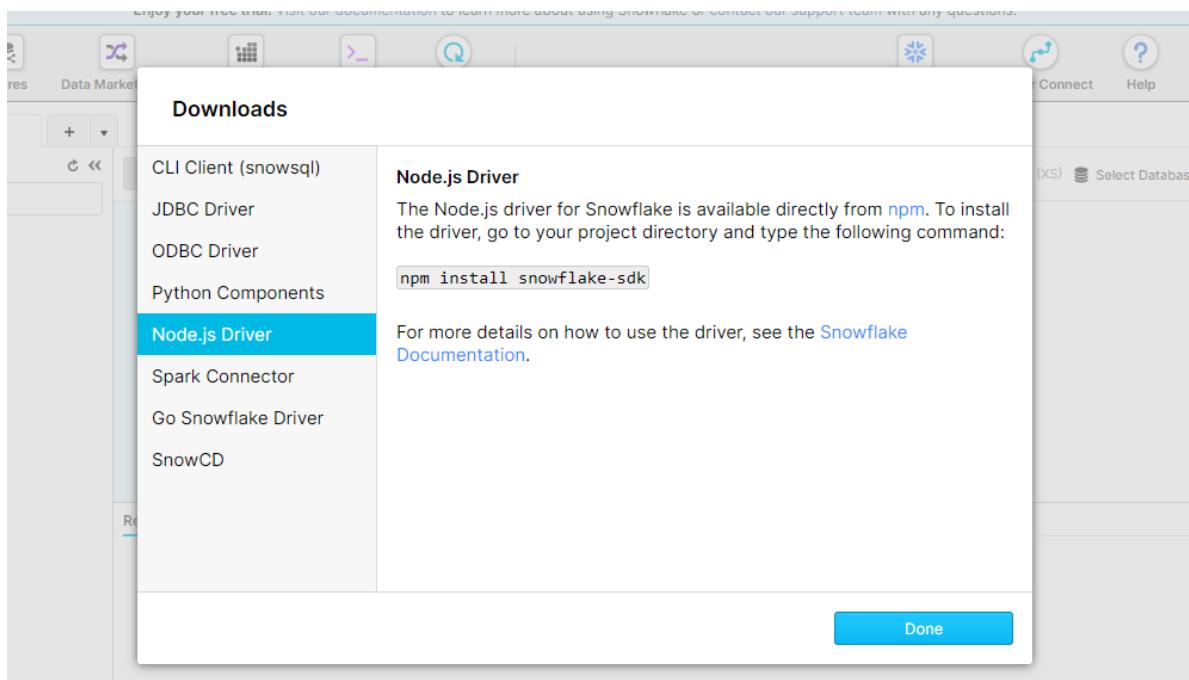
### Explanation

Snowflake provides two types of native clients to develop business applications.

1. Connectors – For Python and Spark
2. Drivers – For JDBC, ODBC, Node.js, Go

These components are available for download in Snowflake Web UI. Go to Help

→ Downloads in the Web UI.



**Further Reading:** [Native Programmatic Interfaces – Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q28 : To execute a query in Snowflake, you need to provide context. The context consists of the following. (select four)*

- A. Virtual Warehouse (VW)
- B. Database
- C. Schema
- D. Role
- E. User
- F. Session

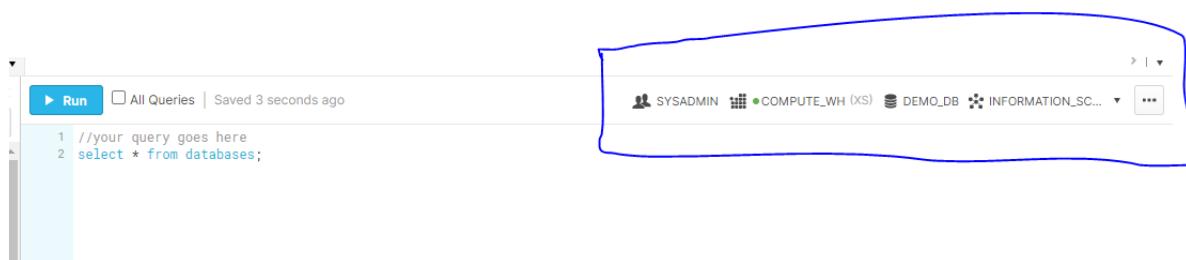
**Correct Answers:** A, B, C and D

### Explanation

A query context in Snowflake consists of 4 objects:

1. Role – Every query you execute should have a role that it uses to access the objects referred to in the query.
2. Warehouse – You need to assign a warehouse to use to a query that will do the necessary computation required for the query and return a resultset
3. Database – Every query must be assigned a database.
4. Schema – Every query must be assigned a specific schema within the database. Objects referred to in the query are looked up in the schema.

In Snowflake WebUI, you can set the context at the worksheet level by selecting appropriate values in the context drop-down as shown below.



Alternatively, you can set the context using a set of SQL statements given below:

Use warehouse <><>

Use role <><>

Use database <><>

Use schema <><>

## Domain : Snowflake Virtual Warehouses

*Q29: As a Snowflake data engineer, you have a task to ingest data generated by sensors located on electric poles. As more*

*sensors are getting installed in the city, the ingestion workload is increasing continuously. You have been tasked to support the growth in data loading volume while optimizing the cost of ingestion. Which best practices would you employ to increase the performance of your data ingestion workload? (select all options that apply)*

- A.** Increase the size of the virtual warehouse (S >> M >> L)
- B.** Split input files into small batches
- C.** Using SnowPipe as SnowPipe method is better in ingesting continuous micro-batch workloads
- D.** Use SnowSQL as the SnowSQL method is better in ingesting continuous micro-batch workloads

**Correct Answers: A, B and C**

### **Explanation**

**Option choice 1 is correct** because Scaling up a virtual warehouse (i.e., resizing the virtual warehouse to increase its computing power) helps in faster processing of ingestion workload. This is because as you increase the size of the warehouse, the performance increases **by a factor of two**. That means your data ingestion will run twice as fast.

**Option choice 2 is correct** because when you split the file into batches and submit them for processing, Snowflake will take advantage of multiple cores of the CPUs and process the files in parallel. Snowflake recommends that instead of submitting 1 large file for processing, you should submit multiple small files in 10-100 MB. Refer to the video lesson in the further reading section for an experiment that shows this phenomenon.

**Option choice 3 is correct, and Option choice 4 is not correct** because the SnowPipe feature in Snowflake is designed to support the ingestion of a continuous stream of micro-batches. Because it has per-second billing, it is more cost-effective than SnowSQL.

**Further Reading:** [Concurrency in Snowflake – YouTube](#), [Getting Started – Introduction to Virtual Warehouses – YouTube](#)

# Domain : Snowflake Overview and Architecture

*Q30 : Namespace in Snowflake is comprised of:*

- A.** Database and schema
- B.** Account, database, and schema
- C.** Database, schema, and table
- D.** Warehouse, database, and schema

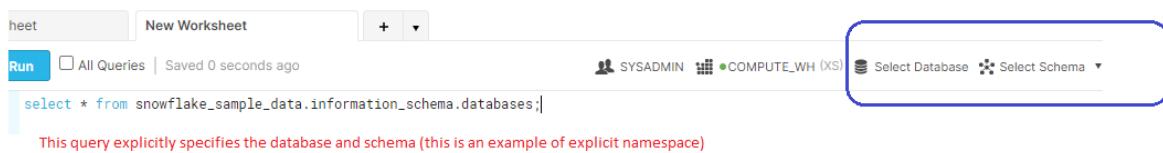
**Correct Answer: A**

## Explanation

When you want to perform any operation (e.g., DDL or DML) on Snowflake's database objects, you need to specify the namespace. Namespace in Snowflake consists of the name of the database and the schema's name within that database. **Therefore answer choice-A is correct.** There are two ways you can specify the namespace. 1) By setting the database and schema in the session context. In this case, Snowflake will infer the namespace from the context. This is referred to as an 'inferred namespace' 2) By explicitly specifying the namespace in the query. This is called an 'explicit namespace.'

Inferred namespace:

Explicit namespace :



The screenshot shows a Snowflake worksheet interface. At the top, there are tabs for 'Sheet' and 'New Worksheet', and a button to add a new worksheet. Below the tabs, there are buttons for 'Run' (highlighted in blue), 'All Queries' (unchecked), and 'Saved 0 seconds ago'. On the right side, there are user status indicators for 'SYSADMIN' and 'COMPUTE\_WH (XS)', and dropdown menus for 'Select Database' and 'Select Schema'. A blue box highlights the 'Select Database' and 'Select Schema' dropdowns. In the main query editor area, the following SQL code is written:

```
select * from snowflake_sample_data.information_schema.databases;
```

Below the code, a red message is displayed: "This query explicitly specifies the database and schema (this is an example of explicit namespace)".

► **Practical Info** – In Snowflake, The concept of namespace makes it super easy to write a query that processes data across databases within the same account. You just need to specify the right namespace of the database objects

you are interested in querying. Many traditional databases may require you to write expensive and effort-intensive ETL routines to do such operations.

See below example where an explicit namespace is used to join tables across two separate databases – Student\_db and course\_db.

```
select * from student_db.public.students_master student,  
course_db.public.course_master course, course_db.public.student_course  
mapping_table where
```

```
student.rollno = mapping_table.student_id and mapping_table.course_id =  
course.course_id;
```

**Further Reading:** [Database, Schema, & Share DDL – Snowflake Documentation](#)

The YouTube video given below explains how Snowflake simplifies querying across databases:

- [Query Multiple Databases Within Snowflake – YouTube](#)

## Domain : Snowflake Overview and Architecture

*Q31: A BOOLEAN data type in Snowflake can represent (select all that apply)*

- A. TRUE
- B. FALSE
- C. NULL
- D. BOOLEAN datatype is not supported

**Correct Answers: A, B and C**

**Explanation**

Snowflake supports BOOLEAN data type. You can define a column in a table that can store boolean values. See the example given below.

```
Create table german_language_students_table (rollno int, name string,  
knows_basic_german boolean);
```

BOOLEAN data type here can have one of the three values: TRUE, FALSE, or NULL. NULL represents ‘unknown.’

- the student knows basic german

```
insert into german_language_students_table values (1, 'amit khanna', TRUE);
```

- the student does not know basic german

```
insert into german_language_students_table values (2, 'shardul mandalo',  
FALSE);
```

- it is not known whether the student knows german

```
insert into german_language_students_table values (2, 'siddharth shah',  
NULL);
```

**► Exam Tip:**

STRING and NUMBER datatype can be implicitly cast into BOOLEAN datatype as follows:

A string ‘true’ is implicitly cast into boolean TRUE, ‘false’ is implicitly cast into boolean FALSE. The numeral 0 is implicitly cast into boolean FALSE. Any non-zero numeral is implicitly cast into boolean TRUE.

**Further Reading:** [Logical Data Types – Snowflake Documentation](#)

# Domain : Snowflake Virtual Warehouses

*Q32 : It is a good practice to create multiple virtual warehouses of various sizes in Snowflake based on the workload.*

- A. True
- B. False

**Correct Answer: A**

## Explanation

A snowflake virtual warehouse is one or more compute clusters that enable users to load data, execute queries and perform other operations on the data.

Snowflake has a unique pay-per-use model to charge the cost of compute. The charge is based on the size of virtual warehouses and the time they are run. When you create a new virtual warehouse, Snowflake only stores the meta-data of the virtual warehouse. Only when this warehouse is in a running state, you are charged. Hence, creating multiple virtual warehouses doesn't cost money.

Secondly, the workload processing of each warehouse is completely isolated from the workload processing of other warehouses. Therefore, if an L size warehouse is loading data and an M size warehouse is processing user queries, each workload's performance is completely unaffected by the presence of the other.

Owing to these reasons, it is a good practice to tailor each warehouse size to the expected workload. Therefore, the answer to this question is A (True).

**Further reading:** A nice 6-min video lesson by Snowflake on key features of virtual warehouses. [Getting Started – Introduction to Virtual Warehouses – YouTube](#)

# Domain : Snowflake Virtual Warehouses

*Q33 : Following are the two modes in which you can run a multi-cluster virtual warehouse in Snowflake.*

- A.** Dynamic
- B.** Maximized
- C.** Auto-Scale
- D.** Multi-threaded

**Correct Answers: B and C**

## Explanation

A multi-cluster virtual warehouse consists of a group of compute clusters provisioned either based on the user/query load. There are two modes in which you can create a multi-cluster warehouse. These modes govern the way compute Snowflake provisions clusters.

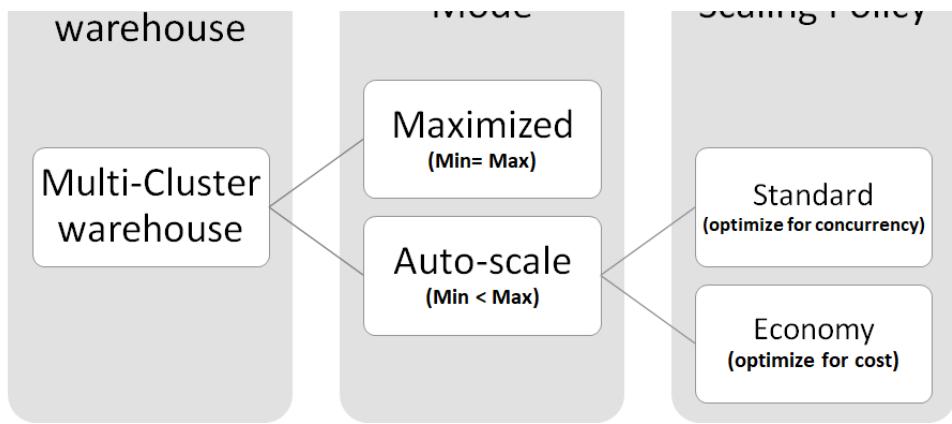
- **Maximized** – In this mode, Snowflake starts all compute clusters within the warehouse when the warehouse is started. This mode is enabled by setting the **same value for Min. and Max. clusters**
- **Auto-Scale** – In this mode, Snowflake starts the warehouse with Min. number of clusters and spins up new clusters incrementally and automatically as and when the user/query load increases. The maximum number of clusters Snowflake can add is governed by the value given in Max. clusters. This mode is enabled by setting different values for Min. and Max. clusters.

► **Exam Tip:** To provide further control over the way the ‘Auto-scale’ mode operates, Snowflake provides a property called ‘SCALING\_POLICY’. This property can have two values: 1) Standard 2) Economy.

Virtual

Mode

Scaling Policy



**Further Reading:** [Multi-cluster Warehouses – Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q34 : The following object in Snowflake is also known as the 'data dictionary' of a database.*

- A. PUBLIC schema
- B. Data Catalog
- C. INFORMATION\_SCHEMA schema
- D. None of the above

**Correct Answer: C**

### Explanation

When you create a database in Snowflake, two schemas get automatically created within the database. 1. INFORMATION\_SCHEMA schema and 2. PUBLIC schema.

INFORMATION\_SCHEMA is a read-only schema. It contains only views. The views have all meta-data information about the database objects e.g., all tables of that database, all columns, all stages, all sequences etc. These views can be queried to retrieve meta-data information about the database. Hence, the INFORMATION\_SCHEMA schema is also known as the data dictionary of a database.

The screenshot shows the Snowflake UI with the database 'MYDATABASE' selected. Under it, the 'INFORMATION\_SCHEMA' is expanded, showing a list of views:

- APPLICABLE\_ROLES
- COLUMNS
- DATABASES
- ENABLED\_ROLES
- EXTERNAL\_TABLES
- FILE\_FORMATS
- FUNCTIONS
- INFORMATION\_SCHEMA\_CATALOG...
- LOAD\_HISTORY
- OBJECT\_PRIVILEGES
- PIPES
- PROCEDURES
- REFERENTIAL\_CONSTRAINTS
- REPLICATION\_DATABASES

► **Exam Tip** – Another potential exam question: Name the two schemas that every database has in Snowflake.

**Further Reading:** [Information Schema – Snowflake Documentation](#)

## Domain : Snowflake Virtual Warehouses

*Q35: After a warehouse is suspended by a credit monitor action, the warehouse will not resume until the following event(s) occur. (select all that apply)*

- A. Snowflake account administrator manually resumes the warehouse
- B. Credit quota is increased in the resource monitor configuration
- C. Credit threshold is increased in the resource monitor configuration
- D. Next monitoring cycle (monitoring interval) starts

**Correct Answers:** B, C and D

**Explanation**

If a monitor has a ‘Suspend’ or ‘Suspend Immediately’ action defined and its used credits reach the threshold for the action, all warehouses assigned to the monitor are suspended and cannot be resumed until (any) one of the following event(s) take place.

You alter the resource monitor and increase the credit quota value (**Answer choice-B**) or

You can increase the credit threshold for the ‘suspend’ action to a higher value. This will allow warehouses to function again (**Answer choice-C**). Alternatively, the warehouses will become available again for query processing when the new monitoring cycle (monitoring interval) starts. The new cycle will reset the credit quota. Therefore the warehouses can now be run again. (**Answer choice-D**)

**Answer choice-A is invalid.** A warehouse suspended due to resource monitor action will not resume until the resource monitor configuration is changed.

► **Exam Tip** – In addition to the above, you can also take any of the below actions.

- You alter the resource monitor configuration and remove the warehouse monitoring. This will stop the monitoring of the credit consumption, or
- You drop the resource monitor altogether. This will release the credit quota threshold. Therefore, the warehouse can run again.

**Further Reading:** [Working with Resource Monitors – Snowflake Documentation](#)

## Domain : Data loading/unloading

*Q36 : Snowflake provides the following types of internal stages to load data. (select all that apply)*

A. Table Stage

- B.** Data Stage
- C.** User Stage
- D.** Named Stage

**Correct Answers: A, C and D**

**Explanation**

A stage is essentially a location in the cloud where you will place your data to load in Snowflake. A stage can be internal (i.e., within Snowflake) or external (outside Snowflake in supported cloud environments i.e., AWS, Azure or GCP). Snowflake provides 3 types of internal stages to load data: table stage, user stage, named stage. A comparison of these stages and corresponding features are given below. You may get a variety of questions in the exam about the features of each stage.

	<b>Table stage</b>	<b>User Stage</b>	<b>Named stage</b>
Purpose	This is a table-specific stage. Automatically available for every Snowflake table. Can copy data into that specific table only. Can be used by any user.	This is a user-specific stage. Automatically available for every user. Can copy data into any table to which the user has access. One user cannot access another user's stage.	Most flexible of the three. Can copy data into any table and can be used by any user having the privileges to access that stage.
Name	Has the same name as the table name.	Has the same name as the user name	first-class database objects and can have any user-defined

			name
Referenced as	@%<stage-name>	@~<stage-name>	@<stage-name>
Data transformations	Not supported	Not supported	Supported
Web UI	Cannot be seen in the Stages tab of Web UI	Cannot be seen in the Stages tab of Web UI	Can be seen in the Stages tab of Web UI

**Further Reading:** [Choosing a Stage for Local Files – Snowflake Documentation](#), [Copying Data from an Internal Stage – Snowflake Documentation](#)

## Domain : Data loading/unloading

*Q37: While loading data into Snowflake, which of the following transformations are supported? (select all that apply)*

- A. Column reordering
- B. Column omission
- C. Derived columns
- D. Cast operation
- E. Truncation
- F. Transpose

**Correct Answers: A, B, D and E**

### Explanation

Snowflake supports 4 types of simple transformations while doing data loading.

The supported transformations are – reordering of columns, omission of columns, data type conversion (also called casting or typecasting) and truncating text strings that exceed the length in the target column. Derived columns and transpose or any other complex data transformation operations are not supported while data loading.

## Domain : Data loading/unloading

*Q38 : It is NOT possible to query data stored in the cloud in Snowflake without first loading data into a Snowflake table.*

- A. True
- B. False

**Correct Answer: B**

### Explanation

**The above statement is false.** Snowflake supports direct querying of data stored in a cloud location.

Imagine a large organization that has invested time and energy to build a data lake on the Azure cloud (or Amazon cloud or GCP). The data lake contains petabytes of data accumulated over many years across multiple channels. Now, the organization wants to query a portion of this data to derive useful insights without reloading data in Snowflake. For these kind of scenarios, Snowflake provides a feature called ‘External Tables’.

An external table in Snowflake references to your data files in external cloud storage (e.g., data lake). External tables store file-level metadata about the data files in the data lake, such as the filename, a version identifier and related properties. Now when you query an external table, Snowflake executes that query against the data stored in external cloud storage, as if it were inside a database.

► **Practical Info** – When you query an external table, it may appear as if data is being fetched from the database. But in reality, the data is fetched directly from external cloud storage every time you run a query on an external table. Therefore the source of truth for the data is always the external cloud store.

**Further Reading:** [Overview of Data Loading – Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q39 : From the statements given below, select the statements which are TRUE for reader accounts. (select all that apply)*

- A. A reader account facilitates data sharing with consumers that do not have a Snowflake account
- B. One reader account can be used to access multiple shares provided by different provider accounts
- C. A reader account consumes credit of the provider account
- D. A reader account has the same access to Snowflake support as the provider account

**Correct Answers:** A and C

### Explanation

**Answer choice-A is correct.** If consumers of your share do not have a Snowflake account (e.g., a consumer using other database products such as Oracle or PostgreSQL), Snowflake provides a functionality called ‘reader account’ which can be used to access the share. A reader is a special Snowflake account created by the provider of the share for the sole purpose of providing access to the share to a consumer. The provider fully manages the reader accounts.

**Answer choice-B is incorrect.** Each reader account belongs to the provider account that created it. This means that if you need to access multiple shares from different providers, you will need to have multiple reader accounts – one

reader account per provider.

**Answer choice-C is correct.** A reader account will require a warehouse to execute queries on the data shared with it. As the reader account is not a licensed user of Snowflake, the warehouse needs to be provisioned by the provider account. Consequently, the provider will be charged for the credits consumed by the queries executed by the reader accounts it has created.

**Answer choice-D is incorrect.** The reader account does not have access to Snowflake support. The provider that creates and manages the reader account will be able to route questions to Snowflake support on behalf of the reader account.

► **Practical Info:** Technically, a reader account can consume an unlimited number of credits, which will be charged to the provider account. Therefore, to monitor and control the credit consumption by the reader accounts, the provider account should create resource monitors.

**Further Reading:** [Managing reader accounts](#)

## Domain : Snowflake Storage and Protection

*Q40 : Match the appropriate data protection feature given on the right with the scenario given on the left. (Assume Business-critical edition of Snowflake)*

	<i>Scenario</i>	<i>Data Protection Feature</i>
S1	<i>The latest release rolled out today morning has corrupted data in the production environment.</i>	1. Fail-Safe
S2	<i>A data issue in the production database was discovered 3 days after time travel has expired.</i>	2. Triple Redundancy
S3	<i>Snowflake instance has crashed, but the cloud</i>	3. Time-Travel
S4		4. Database replication
		5. Automatic AZ

	<i>region is available.</i>	
S5	<i>One availability zone in the cloud region where Snowflake was hosted has failed.</i>	<i>failover</i>
	<i>The entire cloud region where Snowflake was hosted has failed (for example, due to some natural calamity or network outage)</i>	

- A. S1 – 2, S2 – 3, S3 – 5, S4 – 1, S5 – 4
- B. S1 – 3, S2 – 2, S3 – 1, S4 – 5, S5 – 4
- C. S1 – 4, S2 – 1, S3 – 5, S4 – 4, S5 – 3
- D. S1 – 3, S2 – 1, S3 – 2, S4 – 5, S5 – 4

**Correct Answer:** D

### Explanation

**S1** – Latest release rolled out today morning has corrupted data in the production environment. **Answer: Time-Travel**

As data got corrupted in the morning, you have time to recover this data yourself by using the Snowflake time-travel feature. Time-travel is available for every Snowflake edition. For the Standard Edition, you have 1 day (24 hours) of default time travel. For Enterprise edition and above, you get up to 90 days of time travel.

**S2** – A data issue in the production database was discovered 3 days after time travel has expired

As it is given in this scenario that the time-travel window has expired, you can rely on the fail-safe feature of Snowflake and retrieve the older version of your data. The fail-safe period is 7 days after the expiry of time travel. **Answer: Fail-safe**

**S3** – Snowflake instance has crashed, but the cloud region is available. This scenario indicates that the Snowflake instance has crashed, but the underlying cloud provider/region is up. Therefore this is a Snowflake-specific

problem. Snowflake provides a triple redundancy feature that will get activated in such scenarios to mitigate the impact on client workloads. This is completely transparent to the end-users. **Answer: Triple Redundancy**

**S4** – One availability zone in the cloud region where Snowflake was hosted has failed.

Snowflake provides failover protection across three availability zones of your cloud provider region. Your data is synchronously replicated across the three availability zones. Therefore when your primary availability zone fails, the workload is automatically and transparently moved to one of the other availability zones. **Answer: Automatic AZ Fail-over**

**S5** – The entire cloud region where Snowflake was hosted has failed.

To ensure resiliency in this scenario, You will use the 'Database replication and failover/failback' feature available in the Business Critical edition. This feature allows you to create a secondary database in another cloud provider/region. The secondary database is automatically synced with the primary database. In case of a cloud region failure, you can promote your secondary database to serve as the primary database and continue to process data. **Answer:**

**Database replication**

**Further Reading:** [How Snowflake protects your data and services](#)

## Domain : Performance and Tuning

*Q41: Select the statements that apply to clustering depth in a micro-partition. (select all that apply)*

- A.** Clustering depth is the average depth of overlapping micro-partitions for a specified set of columns
- B.** The lower the clustering depth, the better clustered the table is
- C.** The higher the clustering depth, the better clustered the table is
- D.** Is an indication of whether the table needs an explicit clustering key

**Correct Answers: A, B and D**

## Explanation

**Answer choice-A is correct.** Clustering depth is the average depth of overlapping micro-partitions. Clustering depth indicates how well the data is dispersed across micro-partitions.

**Answer choice-B is correct and Answer choice-C is incorrect.** The ideal clustering depth is 1 and it indicates that no micro-partitions overlap for a given set of columns. Therefore, when a search query is executed, Snowflake can directly find the right micro-partition. Any number greater than 1 indicates that when a search query is executed, Snowflake will find several micro-partitions and it will need to process these micro-partitions to prepare the resultset. For an empty table, the clustering depth is 0.

**Answer choice-D is correct.** In general, Snowflake organizes your data in micro-partitions in such a way that clustering depth is low. However, over time, as insert/updates are performed on the data, the data may no longer stay clustered optimally and clustering depth becomes large. This is specifically true for very large tables. When this happens, the search query performance becomes slower over time.

To improve the clustering of the data in micro-partitions, Snowflake allows you to designate one or more table columns/expressions as an explicit clustering key for the table. A clustering key co-locates rows in the same micro-partitions thereby improving search efficiency. Hence, high clustering depth indicates that the queries on that table may benefit by defining an explicit clustering key.

**Further Reading:** [Clustering depth illustration – Snowflake documentation](#)

## Domain : Snowflake Virtual Warehouses

*Q42 : Select statement(s) which are TRUE from the given set of statements concerning Snowflake virtual warehouses.  
(select all that apply)*

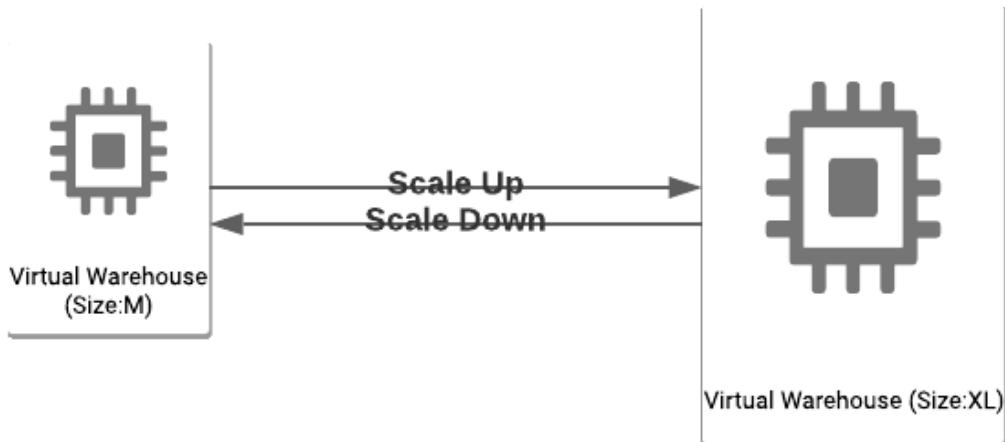
- A. Scaling up a virtual warehouse is an automatic process

- B.** Scaling up a virtual warehouse is a manual process
- C.** Scaling down a virtual warehouse is an automatic process
- D.** Scaling down a virtual warehouse is a manual process
- E.** Scaling out a virtual warehouse is an automatic process
- F.** Scaling out a virtual warehouse is a manual process
- G.** Scaling back a virtual warehouse is an automatic process
- H.** Scaling back a virtual warehouse is a manual process

**Correct Answers: B, D, E and G**

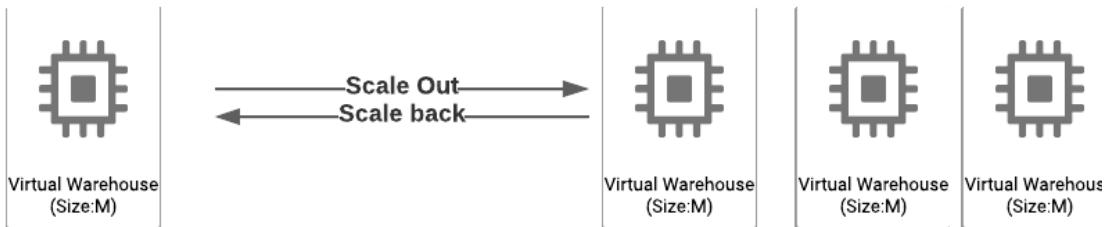
**Explanation**

Scaling up or scaling down a virtual warehouse is also known as warehouse re-sizing. When you re-size a warehouse, you change the T-shirt size of the warehouse (e.g. from S to L or from XL to M). This is a manual activity. You may want to increase the T-shirt size of your warehouse (i.e. scale up) if your query performance or data loading performance is slow. You will want to decrease the T-shirt size of your warehouse (i.e. scale down) if your query workload is low.



Scaling out or scaling back a virtual warehouse is adding or removing new warehouse clusters (without regard to the size of the warehouse). This is an automatic process managed transparently by Snowflake. Snowflake will automatically increase the number of warehouse clusters as the number of

queries increases to prevent queries from queuing. When Snowflake determines that additional clusters are no longer needed, it will shut them down. In summary, Snowflake ensures that a multi-cluster warehouse dynamically adapts to increase or decrease in the number of queries without any user intervention.



Summary:

Scaling model	How?	For?
Scale-Up/Down	Manual	Query/ Data Loading Performance
Scale-Out/In	Automatic	Query concurrency/

**Further Reading:** [Working with virtual warehouses in Snowflake – YouTube video](#)

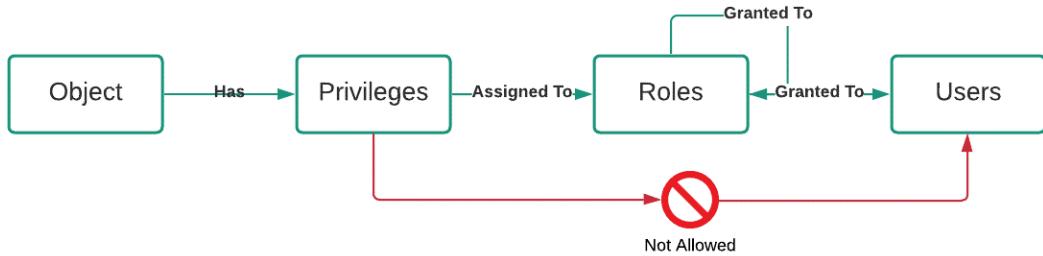
## Domain : Account and Security

*Q43: In Snowflake, For each securable object, there is a set of privileges. The privileges provide fine-grained access control on the object. Each securable object has an owner that can assign these privileges directly to a user or a group of users.*

- A. True
- B. False

**Correct Answer: B**

The above statement is false. This is because the owner of a securable object CANNOT grant privileges over the object directly to a user. The privileges must be assigned to roles. The roles can be granted to other roles creating a role hierarchy OR the role can be granted to a user or a group of users.



**Further Reading:** [Access Control Framework – Snowflake Documentation](#)

## Domain : Snowflake Overview and Architecture

*Q44 : Snowflake being a SaaS software, which of the following activities of an on-prem/hosted data warehouse are not required in Snowflake from the user's perspective? (select all that apply)*

- A. Hardware sizing, purchase, or configuration
- B. Hardware scaling
- C. User and access management
- D. Software installation, maintenance, or upgrades
- E. Database Tuning
- F. Site-level disaster recovery (due to loss of data center)

**Correct Answers: A, B, D, E and F**

### Explanation

Snowflake is a true SaaS, cloud-native, data warehouse product. As it completely runs on public cloud infrastructure, it does not require end-users to purchase, configure or manage any on-prem hardware (**answer choice A**).

Practically unlimited scalability available on tap means end users will not need to carry out costly and effort-intensive hardware scaling by themselves (**answer choice B**). All software maintenance, bug fixes, feature updates, etc. are rolled out completely by Snowflake without any end-user intervention (**answer choice D**). Database tuning is no longer a task for the DBA in Snowflake as Snowflake does not support the creation of indexes, database partitions, (**Answer choice-E**). Your data is automatically replicated across three availability zones in the cloud region without any user intervention (**Answer choice-F**). Out of the given list of answer choices, only 'User and access management is a user responsibility.

**Further Reading:** [Snowflake data warehouse pros and cons – Snowflake community blog](#)

## Domain : Snowflake Storage and Protection

*Q45 : Zero-copy cloning operation is supported for the following objects in Snowflake. (select all that apply)*

- A.** A temporary Table
- B.** A transient Table
- C.** An external table
- D.** A permanent Table
- E.** A database created from a Share
- F.** Internal (named) Stage

**Correct Answers: B and D**

### Explanation

- Zero-copy cloning is available only for the permanent and transient table types. It is not available for temporary and external table types.
- You cannot clone a database that is created from a Share.
- Also, zero-copy cloning is not available for stage objects.

## Domain : Snowflake Virtual Warehouses

*Q46 : A resource monitor can be created by...*

- A.** Owner of the virtual warehouse
- B.** ACCOUNTADMIN role
- C.** SYSADMIN role
- D.** USERADMIN role
- E.** Any of the above

**Correct Answer:** **B**

### Explanation

Resource monitors can only be created by account administrators (i.e. users with the ACCOUNTADMIN role). However, account administrators can choose to grant MONITOR and MODIFY privileges on resource monitors to the users with other roles as needed.

**Further Reading:** [Access Control Privileges for Resource Monitors – Snowflake Documentation](#)

## Domain : Data loading/unloading

*Q47 : You have a business-critical edition of Snowflake on AWS cloud. The data lake implementation of your organization is in Azure and utilizes Azure Data Lake Gen2 service. As your Snowflake account and your data lake are on different clouds, it is not possible to do bulk loading from Azure data lake into Snowflake using the COPY command.*

- A.** True
- B.** False

**Correct Answer: B**

### **Explanation**

This statement is not correct. Bulk loading using COPY INTO supports data loading into Snowflake from the data files on your local file system or in cloud storage external to Snowflake (Amazon S3, Microsoft Azure, or Google Cloud Storage) irrespective of where your Snowflake account is hosted.

**Further Reading:** [Introduction to Data Loading – Youtube Video](#)

## Domain : Semi Structured data

*Q48 : Consider a table vehicle\_inventory that stores vehicle information of all vehicles in your dealership. The table has only one VARIANT column called vehicle\_data which stores information in JSON format. The data is given below:*

```
{  
  "date_of_arrival": "2021-04-28",  
  "supplier_name": "Hillside Honda",  
  "contact_person": {  
    "name": "Derek Larssen",  
    "phone": "8423459854"  
  },  
  "vehicle": [  
    {  
      "make": "Honda",  
      "model": "Civic",  
      "variant": "GLX",  
      "year": "2020"  
    }  
  ]  
}
```

*Which of the following are valid SQL queries that retrieve supplier\_name? (select all that apply)*

- A.** select vehicle\_data.supplier\_name::string from vehicle\_inventory
- B.** select vehicle\_data:supplier\_name::string from vehicle\_inventory
- C.** select vehicle\_data(supplier\_name(string)) from vehicle\_inventory
- D.** select vehicle\_data(supplier\_name)::string from vehicle\_inventory

**Correct Answer:** B

### Explanation

In this example, supplier\_name is the first level element in the JSON. To access this in SQL you will need to use the below syntax in the SQL if you are using .(dot) notation.

<column\_name>:<key\_name>::<cast\_datatype>

Therefore, the correct SQL statements would be –

select vehicle\_data:supplier\_name::string from vehicle\_inventory

**Further Reading:** [Querying Semi-Structured Data – Snowflake Documentation](#)

## Domain : Snowflake Virtual Warehouses

*Q49: Which of the following statements are true about multi-cluster Warehouses in Snowflake? (Select all that apply).*

- A.** Multi-cluster warehouses support all properties of a single-cluster warehouse
- B.** Increasing the number of compute resources in a multi-cluster warehouse is an example of scaling up
- C.** A single cluster warehouse can be reconfigured to turn into a multi-cluster warehouse but vice-versa is not true
- D.** None of the above

## **Correct Answer: A**

### **Explanation**

**Answer choice-A is correct.** Multi-cluster warehouses support all the same properties and actions as single warehouses, including:

- Specifying warehouse size ( e.g. XS, S, M, L...)
- Resizing a warehouse at any time.
- Auto-suspending a running warehouse due to inactivity;
- Auto-resuming a suspended warehouse when new queries are submitted.

**Answer choice-B is incorrect.** Increasing the number of compute resources in a multi-cluster warehouse is an example of scaling out. Increasing the size of compute clusters is an example of scaling up.

**Answer choice-C is incorrect.** A single cluster warehouse can be reconfigured to turn into a multi-cluster warehouse. To do this, you will have to set the max\_cluster\_count to a value greater than 1 as shown in the example below:

```
ALTER WAREHOUSE "COMPUTE_WH" SET MAX_CLUSTER_COUNT = 3;
```

Similarly, a multi-cluster warehouse can be reconfigured to turn into a single cluster warehouse. To do this, you will have to set the max\_cluster\_count to a value equal to 1 as shown in the example below:

```
ALTER WAREHOUSE "COMPUTE_WH" SET MAX_CLUSTER_COUNT = 1;
```

## **Domain : Snowflake Overview and Architecture**

*Q50 : Using materialized views can be beneficial in the following scenarios. (select all that apply)*

- A.** when you have a complex query that is frequently used
- B.** When the underlying data changes frequently
- C.** when the query results are smaller than the base table
- D.** When the query includes an external table

**Correct Answers: A, C and D**

**Explanation**

**Answer choice-A is correct.** Encapsulating a frequently used, complex query that normally takes a long time to execute into a materialized view can improve performance. This is because the query is executed automatically in the background and the computed result set is stored in the materialized view. Querying the view, then, becomes significantly faster compared to re-running the query.

**Answer choice-B is incorrect.** A background service in Snowflake refreshes the materialized view after changes are made to the base table. This service consumes credits. If the underlying data changes frequently, the background service has to be executed repeatedly which leads to higher credit consumption. Hence, it is not recommended to create materialized views on a base table that changes frequently.

**Answer choice-C is correct.** When the query results are smaller than the base table, the background service that refreshes the materialized view needs to compute results for a small set of records. This results in lower consumption of credits to keep the view up to date.

**Answer choice-D is correct.** Querying an external table is likely to be slower than querying native database tables as the data in an external table is stored in an external stage in the cloud (outside of Snowflake). Creating a materialized view brings the resultset locally within your Snowflake account thereby improving query performance.